

The screenshot shows the Ignition software interface. At the top, there's a navigation bar with the Ignition logo, a Help button, and a prominent blue "Get Designer" button with a red border. On the left, a vertical sidebar has three tabs: "Home" (selected), "Status", and "Config". The main content area has a dark blue header with the text "WELCOME" and "Meet Ignition". Below this is a large orange icon with a white exclamation mark. The text "Build It" is centered below the icon. A callout box on the right contains the heading "Get the Designer", a description about the Designer tool, and a "Download Designer Launcher" button. The "Download Designer Launcher" button is highlighted with a red border.

Starting and Stopping the Gateway

After installation, the Gateway starts automatically. The Gateway runs as a service, so you can use your operating system's normal mechanisms to start or stop the service. You can also start or stop the Gateway from command line.

Windows

Ignition's installation directory contains `start-ignition.bat` and `stop-ignition.bat`, which can start or stop the service. Example:

```
C:\Program Files\Inductive Automation\Ignition> start-ignition.bat
```

However, you can also use Windows native service commands to control the running state of the Gateway:

```
net start ignition
```

```
net stop ignition
```

Linux

You can control the service using the `ignition.sh` script. It can be called with the `start` and `stop` parameters to perform the relevant operations.

For example:

```
/usr/local/bin/ignition/ignition.sh start
```

Additionally, you can use native terminal commands to start or stop the service:

```
service Ignition-Gateway start
```

```
service Ignition-Gateway stop
```

Mac OS X

You can access the service from the install directory using the "ignition.sh" script. On a typical Mac install using the dmg installer, the full command (without a custom location specified) is the following:

```
/usr/local/ignition/ignition.sh start
```

Gateway Command-line Utility

The [Gateway Command-line Utility - gwcmd](#) (GCU) is a lightweight standalone application that provides information about the Gateway. It performs high-level tasks that aren't available inside the Gateway webpage such as stopping and restarting the Gateway server, and setting ports used between the Gateway and clients. It shows the status of the Tomcat web server and the Ignition Gateway application. You can reset the Gateway password, and even launch the web browser to the Gateway webpage.

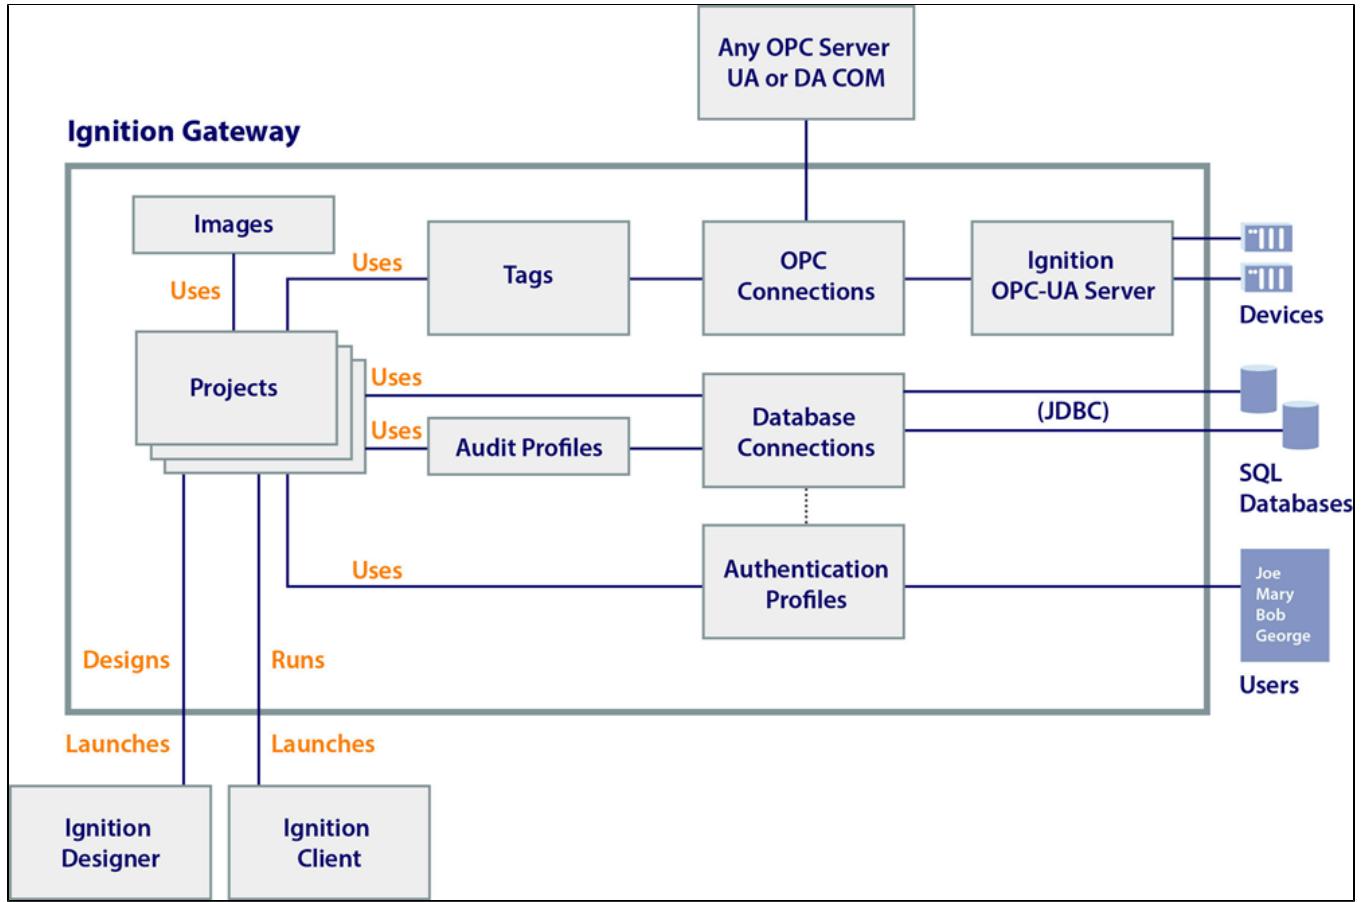
For more information, refer to the page on the [Gateway Command-line Utility - gwcmd](#).

Gateway Architecture

The Gateway's architecture contains many parts, each one letting you perform a specific task. The Gateway parts are as follows:

- System Management
- Project and Module Management
- Security and Auditing Settings
- Database Connectivity
- Alarming and Modification Settings
- Tags Provider Settings
- OPC Connectivity
- Enterprise Administration Settings
- Projects

The following image shows the different parts of the Gateway and how projects work within the Gateway. Note the parts that are outside the Ignition Gateway box, non-Ignition OPC servers, databases, devices, etc., are all separate from Ignition but can be connected.



Home

The Home tab provides you with all of the tools you need to get started designing projects and launching Perspective Sessions and Vision Clients. There are four launchers on the Homepage that will quickly get you started:

- The [Designer Launcher](#) allows you to create or modify a project. Download the launcher and create a shortcut so it's always on your desktop when you need it.
- The [Vision Client Launcher](#) browses all Gateways for Vision projects that are available on your local network. Once your projects are added to the Vision Client Launcher, they will be displayed and available to open in a Vision Client. Simply download the launcher and create a shortcut on your desktop so it's always at your fingertips when you want to launch a Vision Client.
- The [Perspective Session Launcher](#) allows you to easily launch a session directly in your browser. Simply click the link and Ignition will display all your Perspective projects.

The following feature is new in Ignition version **8.1.0**

[Click here](#) to check out the other new features

- The [Perspective Workstation](#) is a single application that acts as both a launcher, and desktop "wrapper" for Perspective Sessions.

While the top navigation bar is present no matter what tab you are in, the Get Designer button in the upper right corner does not require that a user login to the Gateway. This allows a user to launch the Designer right away, only logging in once the Designer has fully launched. See the [Designer](#) section for more information on launching and using the Designer.

On the Homepage, we also provide quick links to several resources: [Inductive University](#), Product Documentation, and the Appendix which contains a complete reference for [components](#), [expressions](#), and [scripting functions](#) in Ignition.

On this page ...



Customizing the Gateway Homepage

[Watch the Video](#)

[Home > Get Started](#)

WELCOME

Meet Ignition



Build It

Get the Designer

The Designer brings all your data, systems, and developers together into one beautifully simple, integrated development environment specifically designed to help you build industrial applications more quickly.

[Download Designer Launcher](#)

Run It

Vision Client Launcher

Download the Native Client Launcher to open Vision clients from any Ignition Gateway.

[Download Vision Client Launcher](#)**Perspective Workstation**

Download Perspective Workstation to open sessions from any Ignition Gateway.

[Download Perspective Workstation](#)**Perspective Session Launcher**

Launch a Perspective session directly in your browser or download the native application.

[View Projects](#)

Related Topics ...

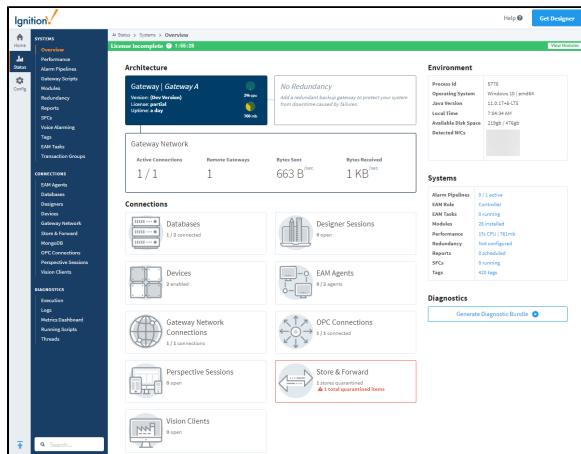
- [Status](#)
- [Config](#)
- [Inductive University](#)

Status

Your Gateway at a Glance

The Status tab provides both an 'at a glance' overview of all of the systems in your Gateway, while also giving you the ability to drill down into specific systems and get a more in depth view of what is currently happening in your system. The Systems and Connections sections that are displayed on the left side of the Status page are based on what modules are currently installed. Some 3rd party modules could potentially add sections that are not discussed in this manual.

When first opening the Status tab, you will be taken to the Overview page. This page provides an overview of all Ignition systems, while also acting as a launchpad to all other sections in the Status tab. The pages in the Status tab are built around the idea of quickly drawing attention to problem areas by highlighting them. As you can see, this Gateway has a faulted database. The pages in the Status tab also allow you to 'drill down' into sections to see more information. Most objects can be clicked on, like the faulted database. Clicking either of these will take you to the appropriate section, allowing you to quickly find out what is wrong with a particular system.



The Status section of the Gateway Webpage provides detailed information relating to the following parts:

Systems	
Overview	Provides a top-down view of many of the components of your Gateway. This view is also useful for determining what step might be next when setting up your Ignition Gateway for the first time. You can view the status of your database connections, device connections, OPC connections, the number of open Clients and the number of open Designers. The Status > Overview page displays the number of Perspective Sessions currently running.
Performance	Displays the performance status for the Ignition system such as CPU, Memory and Threads.
Alarm Pipelines	Shows the status details about the alarm notification pipelines . You can see the status of an alarm and where the alarm is in the pipeline.
Gateway Scripts	Shows status details about all the Gateway event scripts running in Gateway. You can see information such as their execution status, whether the scripts are running or not, and so on.
Modules	A list of installed modules , their status, as well information about their version and current license state.
Redundancy	Lists information about the current state of Redundancy in Ignition. This information is only helpful when connected to another redundant Ignition server.
Reports	Displays information about the current and scheduled reports on the Gateway if you are using the Reporting module .
SFCs	Displays information about Sequential Function Chart (SFC) instances.
Tags	Lists information and statistics about all configured Tag Providers as well as a view into the Tag subscription model, scan classes, and what tags it is currently subscribed to.
Transaction Groups	Displays information about the current Transaction Groups .

On this page ...

- [Your Gateway at a Glance](#)
 - [Systems](#)
 - [Connections](#)
 - [Diagnostics](#)
- [Search](#)
- [Diagnostic Bundle Export](#)



INDUCTIVE
UNIVERSITY

Gateway Status Page Troubleshooting

[Watch the Video](#)

Connections	
EAM Agents	Displays a list of information from configured EAM agents.
Databases	Displays a list of configured databases, and if they have a valid connection or not. Shows active queries, long running queries, the number of queries a second that are running, as well as a trend showing the percentage of queries that completed in that time.
Designers	Displays information on currently running Designer sessions along with some information about each session.
Devices	Displays a list of currently configured OPC UA devices and which are connected and which have a faulty connection. It also shows how many Tags Ignition is requesting from the device along with how often it is requesting them. This information is used to determine if the device is overloaded with too many requests too quickly, or if the device is being under utilized.
Gateway Network	<p>Shows an overview of the status of all Gateways within the Gateway Network. It also provides some metrics for each Gateway, giving an idea of the rate of data transfer between two Gateways, as well as a list of connection events.</p> <div style="border: 1px solid orange; padding: 5px; margin-top: 10px;"> <p>The following feature is new in Ignition version 8.1.25 Click here to check out the other new features</p> </div> <p>The "Live Diagram" tab will take you to the Gateway Network Diagram. This tool is a visual representation of your Gateway Network and contains relevant information from the local Gateway's perspective.</p>
Store & Forward	Displays a list of Store and Forward engines, including status, as well as the number of records currently in each Store and Forward system.
MongoDB	Displays metrics from a list of created MongoDB connections.
OPC Connections	Displays a list of all current OPC connections and their status.
Perspective Sessions	Shows a list of current Perspective sessions and details about each session.
Vision Clients	Shows a list of current Vision Clients and details about each client.
Diagnostics	
Execution	Displays a status of all tasks that your Gateway runs on a schedule, such as duration and execution time of an alarm journal update or the average time it takes a Gateway to execute a Tag Group.
Logs	Displays errors caused by Gateway events like database or device connections, authentication profiles, alarm journals, and pipelines. Logs include a wealth of information about the running state of the Gateway.
Metrics Dashboard	<div style="border: 1px solid orange; padding: 5px; margin-bottom: 10px;"> <p>The following feature is new in Ignition version 8.1.13 Click here to check out the other new features</p> </div> <p>Displays user-configurable dashboards featuring gauges, counters, meters, and other metrics visualizations. Available metrics include information about performance, tags, alarming, projects, and databases.</p>
Running Scripts	Shows all actively running Gateway scripts, as well as providing a way to terminate any running script. In addition, Vision client and Designer consoles have a Running Scripts tab, which also lists running scripts and provides a way to terminate them.
Threads	Shows what each thread is doing in the Gateway, including their state and CPU usages.

Search

The Status tab also displays a search bar at the bottom left of each page. This search bar allows you to type in a word or phrase, and it will list all appropriate pages in both the Status and Config tabs to easily find all pages related to a specific system.

The screenshot shows the Ignition Controller interface. The left sidebar has tabs for Home, Status (which is selected and highlighted in blue), and Config. Under the Status tab, there's a sidebar menu with options like Overview, Performance, Alarm Pipelines, Gateway Scripts, Modules, Redundancy, Reports, and SCC. A search bar at the top of this sidebar says "Search results for 'database'...". Below it is a list of links: Databases in > Connections, Store & Forward in > Connections, Transaction Groups in > Systems, Connections in > Databases, Drivers in > Databases, and Store and Forward in > Databases. A red box highlights the search bar and the list of links. On the right, there are two cards: "Valid Connections" (1 / 2) and "Total Throughput" (2.9 queries/sec). Below these is a table with columns: Driver, Status, Connections, Throughput, and Actions. It shows two rows: MySQL ConnectorJ (Valid, 0 / 8, 2.9 queries/sec, Details) and Oracle JDBC Driver (ERROR, 0 / 8, 0.0 queries/sec, Details). A red box highlights the search bar in the sidebar.

Diagnostic Bundle Export

The following feature is new in Ignition version **8.1.21**
[Click here](#) to check out the other new features

The Status tab's Overview page contains an option to generate and download different sets of diagnostic information at once. This information includes:

- Gateway Information
- Thread dump
- Wrapper log file(s)
- System log IDB
- Metrics IDB
- Memory dump

Note: The memory dump will only be generated if the `-Dignition.diagnostic-bundle.enable-memory-dump` parameter is set to true, and available disk space is **greater than or equal to** twice the max heap size. For more information on adding system properties, see the [Gateway Configuration File Reference](#) page.

- The following feature is new in Ignition version **8.1.27**
[Click here](#) to check out the other new features

Jetty server dump



Exporting the Diagnostics Bundle with an increased number of wrapper logs and/or larger-sized wrapper logs may impact Gateway performance or cause the Gateway to restart.

Environment

Process Id	28
Operating System	Linux amd64
Java Version	11.0.15+10-LTS
Local Time	2:53:58 PM
Available Disk Space	239gb / 288gb
Detected NICs	172.17.0.5

Systems

Alarm Pipelines	0 active
EAM Role	Unknown
Modules	24 installed
Performance	0% CPU 182mb
Redundancy	Not configured
Reports	0 scheduled
SFCs	0 running
Tags	196 tags

Diagnostics

[Generate Diagnostic Bundle](#)

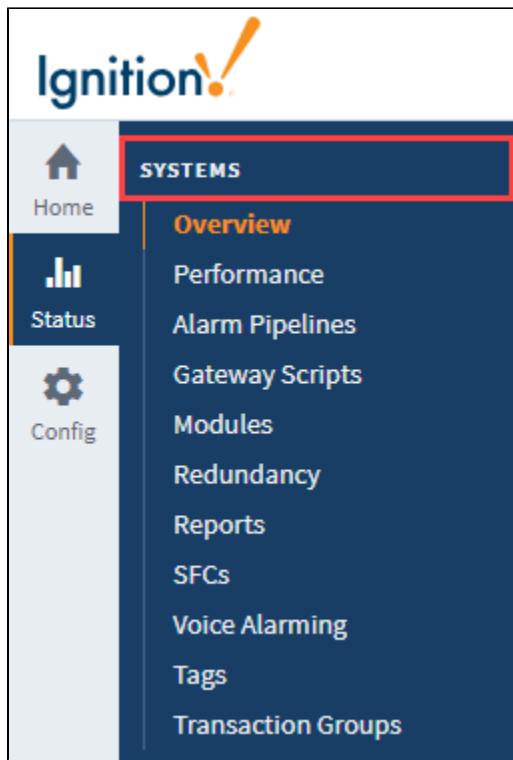
Related Topics ...

- [Config](#)
- [Home](#)

In This Section ...

Systems

The Systems section of the Status tab displays information regarding the internal Gateway systems. This includes systems such as the currently running Transaction Groups as well as active Alarm Pipelines. These sections are not necessarily concerned with any connections to external systems.



On this page ...

- Overview Section
 - Architecture
 - Environment
 - Systems
 - Connections
- Performance
- Alarm Pipelines
- Gateway Scripts
- Modules
- Redundancy
 - Force Re-Sync
 - Request Failover
- Reports
- SFCs
- Voice Alarming
- Tags
- EAM Tasks
- Transaction Groups



Gateway Status Page Troubleshooting

[Watch the Video](#)

Overview Section

The Overview page is the first page that you see when navigating to the Status tab. It provides a visual synopsis of everything in the Gateway, as well as some basic information about the server Ignition is installed on. There are four sections on the page: **Architecture**, **Environment**, **Systems**, and **Connections**.

Ignition!

Status > Systems > Overview

Architecture

Gateway | Ignition-ignition8-ubuntu-64bit
Version: 8.0.0-beta0 (b2019032009)
License: Licensed
Uptime: a day

No Redundancy

Add a redundant backup gateway to protect your system from downtime caused by failures.

Gateway Network

Active Connections	Remote Gateways	Bytes Sent	Bytes Received
1 / 2	1	2 KB/sec	1 KB/sec

Connections



Designer Sessions
2 open



Databases
1 / 1 connected



Gateway Network Connections
2 / 2 connections



OPC Connections
1 / 1 connected



Store & Forward
0 stores quarantined



Devices
5 enabled



Vision Clients
0 open

Environment

Process Id	37117
Operating System	Linux amd64
Java Version	11.0.2+7-LTS
Local Time	10:04:17 AM
Available Disk Space	3gb / 16gb
Detected NICs	10.10.115.3

Systems

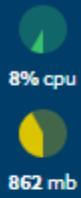
Performance	9% CPU 512mb
Redundancy	Not configured
Modules	18 installed
Tags	1,103 tags
Alarm Pipelines	0 / 1 active
SFCs	0 running
EAM Role	Unknown
Reports	4 scheduled

Architecture

The Architecture segment provides a top down view of your Gateway. You can see your Ignition version and resource usage, as well as similar information with a redundant Gateway, if one is connected. You also get an overview of the Gateway Network including how many active connections there are, and what is the amount of data being transferred between them.

Architecture

Gateway | Ignition-
ignition8-ubuntu-64bit



Version: 8.0.0-beta0
(b2019032009)

License: licensed
Uptime: a day

No Redundancy

Add a redundant backup gateway to protect your system from downtime caused by failures.

Gateway Network

Active Connections

2 / 2

Remote Gateways

2

Bytes Sent

2 KB/sec

Bytes Received

1 KB/sec

Environment

The Environment segment provides some basic information about your server such as the Operating System (OS), Java version, available disk space, and IP address of detected network interface cards (NIC)..

Environment

Process Id	37117
Operating System	Linux amd64
Java Version	11.0.2+7-LTS
Local Time	11:18:31 AM
Available Disk Space	3gb / 16gb
Detected NICs	10.10.115.3

Systems

The Systems segment summarizes each of Ignition's systems. Here, you can see how many modules you have installed as well as how many Tags are configured in your Gateway, along with a handful of other information. Click on any of the links in blue to find out more information about a

particular system.

Systems

Performance	6% CPU 504mb
Redundancy	Not configured
Modules	18 installed
Tags	1,103 tags
Alarm Pipelines	0 / 1 active
SFCs	0 running
EAM Role	Unknown
Reports	4 scheduled

Connections

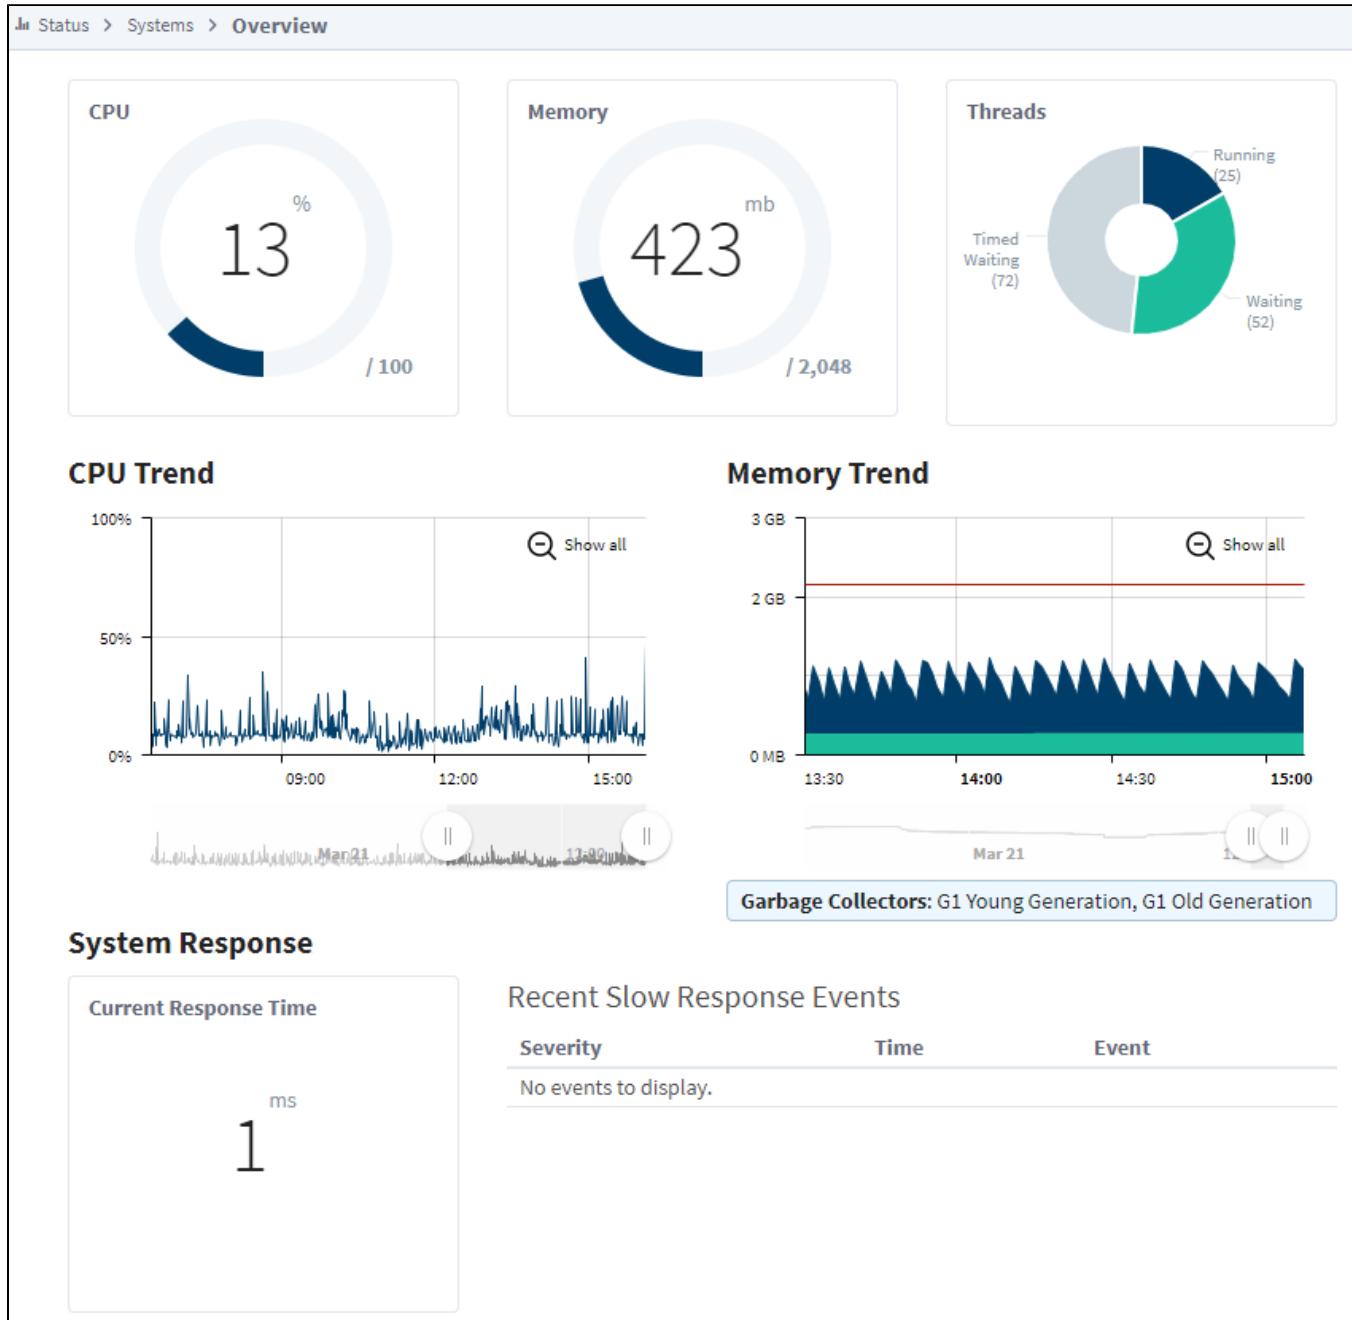
The Connections segment shows all of the systems that Ignition is connected to. If any of these have issues, they will be outlined in red. Clicked on the connection to be taken to the relevant page within the Status tab.

Connections

 Designer Sessions 3 open	 Databases 1 / 1 connected
 Gateway Network Connections 2 / 2 connections	 OPC Connections 1 / 1 connected
 Store & Forward 0 stores quarantined	 Devices 5 enabled
 Vision Clients 0 open	

Performance

The Performance page displays information on the resource usage of the Gateway. This page enables you to quickly assess your Gateway's overall health and determine if the Gateway is being taxed too heavily for the server it's running on. You can see the current CPU usage, the current memory usage (in megabytes) out of the total allocated for the Gateway, and how CPU threads are performing. In the CPU Trend and Memory Trend sections there are sliders you can move to set the time range anywhere within in the previous 24 hours. The lower portion of the page shows the current system response time as well as a log of any recent slow response events.



Alarm Pipelines

The Alarm Pipelines page lists the currently configured alarm pipelines in the Gateway. Each pipeline displays the number of alarm items currently running through it. Click on the Details button to the right to open up a new page that will show the alarm pipeline status and logs. There will also be a set of tabs at the top of the page. One for the Pipeline Status that is already open, the other for Pipeline Logs. This [log viewer](#) will act as a miniature log viewer that will only show logs that pertain to this particular pipeline.

The screenshot shows a web-based interface for managing alarm pipelines. At the top, there is a navigation bar with links for Status, Systems, and Alarm Pipelines. The 'Alarm Pipelines' link is highlighted with a red box. In the center, there is a large white box labeled 'Active' with the text '0/2'. To the right, there is a 'Configuration' button. Below this, there is a search/filter bar with a 'Filter' dropdown and a 'View' dropdown set to 20 items. A table follows, displaying two rows of pipeline information. The columns are 'Name', 'Active?', 'Items', and 'Actions'. The first row is for 'project:alarm-pipelines:/pipeline:Loop Pipeline' with 'Active?' set to false and 'Items' at 0. The second row is for 'project:NewProject_SJP:/pipeline:Test Pipeline' with 'Active?' set to false and 'Items' at 0. Each row has a blue 'Details' button in the 'Actions' column.

Name	Active?	Items	Actions
project:alarm-pipelines:/pipeline:Loop Pipeline	false	0	Details
project:NewProject_SJP:/pipeline:Test Pipeline	false	0	Details

Gateway Scripts

The Gateway Scripts page displays a list of all currently configured Gateway scripts. These are scripts that are configured in the [Gateway Event Scripts](#) section of the project. The tabs at the top of the page will swap between the different types of Gateway Event Scripts: Timer, Tag Change, Message Handler, Startup, and Shutdown. They provide useful information such as the name of the project that it is running in as well as the last time it ran, the duration and status (if it was successful or not).

With this page you can quickly verify that your scripts are running properly. If a script has an error, click on the error for more details about what went wrong with the script. All of the logs for that section are also contained in the log viewer at the bottom of the page, to make it easy to see a list of all of the logged errors for that set of Gateway Event Scripts.

Status > Systems > **Gateway Scripts**

Timer Tag Change Message Handler Startup Shutdown

« < 1 of 1 > »

Filter type to filter View 20 ▾

Project ▾	Name	Rate	Last Execution	Duration	Status
No scripts found.					

« < 1 of 1 > »

Log Activity

Min level **INFO** ▾ i Live Values **on** Merge to Logs ↗

Logger	Time	Message
No log entries found		

Modules

The Modules page shows a list of the currently installed modules, their current version, and if they are running properly with a license. The top of the page shows a quick count of the running and licensed modules, so you can easily tell if anything is not working. It is then easy to scroll down through the list to see which modules are having issues.

Running Modules

18 / 18

Licensed Modules

18 / 18

License Details

License Key	LRN-IGN
Version	8
edition	standard

Inductive Automation, LLC

Name	Version	License	Status
Alarm Notification	5.0.0-beta0 (b2019032010)	Activated 🔍	✓ RUNNING
Allen-Bradley Driver	5.0.0-beta0 (b2019032010)	Activated 🔍	✓ RUNNING
DNP3 Driver	3.0.0-beta0 (b2019032010)	Activated 🔍	✓ RUNNING
Enterprise Administration	3.0.0-beta0 (b2019032010)	Activated 🔍	✓ RUNNING
Logix Driver	4.0.0-beta0 (b2019032010)	Activated 🔍	✓ RUNNING
Modbus Driver	6.0.0-beta0 (b2019032010)	Activated 🔍	✓ RUNNING

Redundancy

The Redundancy section displays information regarding the redundant system, if one is configured. It easily shows the Role of the Gateway you are viewing, the status of the connected Gateway, and their IP addresses. The trends on this page give a snapshot of the last few minutes of communication between the two Gateways. The first trend shows the data that is being sent and received between the two Gateways, the second shows the state updates that have recently occurred. The bottom of the page contains a log of the last system events, to easily track major events between the Gateways.

Force Re-Sync

The **Force Re-Sync** button forces a full synchronization of the redundant configuration state. The backup node will be forced to restart.

Request Failover

The **Request Failover** button switches the active node in a redundant pair.

The following feature is new in Ignition version **8.1.17**
[Click here](#) to check out the other new features

Failover to the other redundant node is now allowed if the nodes have different platform versions, which will allow attached clients to remain connected to at least one node during a redundant pair upgrade.

Status > Systems > **Redundancy**

Configuration

Role	Peer Connected	Redundancy Properties
Backup	Yes	Activity Level Cold Synchronization Status Out of Date Local Address 10.10.115.7 Peer Address Ignition-ignition8-ubuntu-64bit-Master
		Force Re-Sync Request Failover

System Events

Severity	Time	Event
Info	22Mar2019 16:43:04	Initiating a forced configuration re-sync
Info	22Mar2019 16:43:04	Forced configuration sync initiated.
Info	22Mar2019 16:42:28	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=OutOfDate, History level=Full

Reports

The Reports page shows information on any Reports that have at least one scheduled action set up. The top of the page gives a quick count of reports that are executing an action, have executed an action, or are going to execute an action. Below you can also find a list of Reports in those states as well, to quickly see what project they are located in, or what their last execution time was. As with most other pages, there is also a log of all events related to reports at the bottom of the page.

Status > Systems >

Report Stats

Currently Executing

0

Completed Scheduled Executions

1

Upcoming Scheduled Executions

1

Currently Executing Reports

Project	Report Path	Status	Execution Start	Elapsed Time
No items to display				

Completed Scheduled Executions

Project	Report Path	Status	Execution Start	Elapsed Time
TestProject	Report	Finished	30Sep2016 00:00:00	283 milliseconds

Upcoming Scheduled Executions

Project	Report Path	Scheduled Start	Actions
TestProject	Report	01Oct2016 00:00:00	Save File

Log Activity

Min levelINFO(i)Live Values Merge to Logs

LoggerTimeMessage

No log entries found

SFCs

The [Sequential Function Charts](#) (SFCs) section displays a list of all of the currently configured SFCs along with a count of the currently running charts at the top of the page. Click the **Details** button next to any SFC for details about that particular Chart.

» Status > Systems > Sequential Function Charts

Chart Stats

Running Charts
0 / 1

Defined Charts

Filter type to filter View 20 ▾

Chart ▾	Project	Run Mode	Running Instances	Actions
New Chart	test	✓ Callable	0	Details

« ‹ 1 of 1 › »

On the Details for an individual SFC, there are two tabs: Chart Instances and Chart Logs. It contains a list of all currently running Chart instances, and allows you to swap tabs to look at any logged events for that particular chart.

» Status > Systems > Sequential Function Charts

Chart Instances Chart Logs

Chart Instances

Status	Execution Start ▾	Current Step	Parent	Started By	Actions
No items to display.					

Voice Alarming

The Voice Alarming page provides details about the currently configured [Voice Notification Profiles](#) on the Gateway. This status can be used to quickly see which Notification Profiles are not working. Clicking the details button enables you to see the current calls with that Voice Notification profile and the current queue of calls waiting to be completed.

Status > Systems > **Voice Alarming**

[Configuration](#)

Voice Alarming

Available Voices

Name	Locale
No voice modules installed.	

Profiles

Name	Status	Pending Calls	Action
Skype	⚠️ unknown [-1]	0	Details
voiceProfile1	✓ Registered with VOIP Host	0	Details

Tags

The Tags page contains a lot of information about the Gateway's configured Tags and Tag Providers. The first page will show a list of all realtime and historical Tag Providers that make it easy to see if any of them are having issues. Clicking on an errored provider will show an error message that may help fix the problem. Clicking on the **Details** button to the left will open up a page with more information about the Tags in that particular provider.

Status > Systems > **Tags**

[Realtime Configuration](#) [Historical Configuration](#)

Realtime Providers

Name	State	Tag Count	Actions
default	● Running	299	Browse
ExportTags	● Running	11	Browse
System	● Running	104	Browse
ubuntu-64bit_default-tags	● Unknown	Not Available	Browse

Historical Providers

Name	State	Actions
MySQL	✓ Good	Details
vm_db	✓ Good	Details



**INDUCTIVE
UNIVERSITY**

Tag Diagnostics

[Watch the Video](#)

The details of the Tag Provider will show a list of all currently configured Tags in that Tag Provider, and show some basic information about them. You can browse through the Tag structure to get more information about each Tag. When clicking the Details button for an individual Tag, you will be brought to a Tag Diagnostics screen that will provide more detailed information about that particular Tag, such as if

it has a script written in its Tag event scripts and what its last value was, among other information.

The screenshot shows the 'Tags' tab for the system 'AU 1'. There are two entries listed:

Name	Type	Value	Quality	Actions
Fan1 HOA	AtomicTag	0	Good	<button>Details</button>
Fan2 HOA	AtomicTag	0	Good	<button>Details</button>

The screenshot shows the 'Fan 1 HOA' tag diagnostics. The following parameters are displayed:

Path	Value
Last Value	n/a
Last Quality	Good
Last Execution	n/a
Value Timestamp	09Jul2019 15:41:26
Last Error	n/a
Last Error Cause	n/a
Alarmed?	false
Bound?	false
Historical?	false
Secure?	false
Scaled?	false
Scripted?	false

The Groups Tab allows you to see the Tag Groups configured. It also shows the number of executions for each Tag Group, the last time the Group was executed and the average duration in milliseconds. Provider Logs display any events that relate to that particular provider.

The screenshot shows the 'Groups' tab. The table displays the following groups:

Name	Mode	Rate	Leased Rate	Is Leased	Executions	Last Duration	Avg Duration
Direct	Direct	1000	1000	false	73787	0	0
Direct 5 Seconds	Direct	5000	1000	false	14769	0	0
Driven	Driven	10000	1000	false	81172	0	0
Driven Machine State	Driven	10000	1000	false	81172	0	0
Driven Manual	Driven	4000	500	false	165852	0	0
Driven One Shot	Driven	120000	1000	true	1	0	0
Leased	Leased	10000	1000	false	81172	0	0
Time Driven	Driven	60000	1000	false	75018	0	0
default	Direct	1000	1000	false	73787	0	0

EAM Tasks

The EAM Tasks section will only show up once the [EAM module](#) is configured. This page shows information regarding the currently running and scheduled tasks, as well as tasks that were recently executed and whether they were successful or not. From here you can pause a scheduled task, or see the error that caused a previous execution to fail.

Job Status > Systems > **EAM Tasks**

Configuration

Executing Agent Tasks	Scheduled Agent Tasks
0	0

Executing Agent Tasks

Task Name	Task Type	Task State	Execution Start	Progress	Message
No items to display.					

Scheduled Agent Tasks

Task Name	Task Type	Repeats	Next Execution Start	Status
No items to display				

Recently Completed Tasks

Task Name	Task Type	Task Start	Agent	Task Result
Collect Backup (forced)	backup	25Mar2019 12:39:32	Ignition-2	✓ Success

Transaction Groups

The Transaction Groups page makes it easy to get a quick count of the currently running [Transaction Groups](#) as well as the currently errored groups. Below the counts is a list of Projects that contain Transaction Groups and some basic information like how many are running and how many database queries it generating. Clicking the details button to the right for one of the Projects brings up more information about that Project's Groups. There are some counts on the number of Groups in various states, as well as a list of the groups. Any Groups that have problems can be clicked on to see what the error is to get a better idea of why the Group is unable to run.

Gateway Stats

Running Groups

0 / 5

Errored Groups

0 / 5

« ‹ 1 of 1 › »

Filter type to filter

View 20 ▾

Project ▾	Total	Running	Errored	Tag Writes	DB Queries	Actions
Tank_Control	1	0	0	0 / sec	0 / sec	<button>Details</button>
NewProject_SJP	4	0	0	0 / sec	0 / sec	<button>Details</button>

« ‹ 1 of 1 › »

Related Topics ...

- [Connections](#)
- [Diagnostics - Logs](#)

Connections

Status of Gateway Connections

The Connections section of the Status tab on the Gateway Webpage contains information regarding the status of Gateway connections to external systems. The list of systems displayed under the Connections section is based on what modules are installed. It can include EAM Agents, Databases, Devices, Store and Forward, the Gateway Network including other Gateways within the Gateway Network. The Connections section allows you to drill down and open up more specific information to easily find problems with anything connected to the Gateway, and even find crucial information about a faulty connection.

On the left side of the Status section of the Gateway Webpage, you'll see a list of all your connections. Click on any system to open a detailed page to see all the available options, and to get more detailed information. Some third party modules could potentially add sections that are not discussed in this manual.



On this page ...

- [Status of Gateway Connections](#)
- [Log Activity](#)

Log Activity

Most of the systems in the Connection section have a Log Activity area somewhere within the details of the connection. This Log Activity area functions much like the [Logs in the Diagnostics](#) section of the Status page in that you see a list of all log messages depending on the logging level you are looking at as well as the filters you have set up. However, what makes the Log Activity areas in each Connections page unique is that they are filtered to only show loggers for that particular system. This helps you narrow down any potential problems by showing you only the information that pertains to the section you are looking at.

A screenshot of the 'Log Activity' interface. At the top, there are filter controls: 'Min level' set to 'INFO', a dropdown menu, and a help icon. To the right are buttons for 'Live Values' (on) and 'Merge to Logs'. Below the filters is a table with columns 'Logger', 'Time', and 'Message'. A message at the bottom says 'No log entries found'. There are also navigation icons at the top right.

Related Topics ...

- [Diagnostics - Logs](#)

In This Section ...

Connections - EAM Agents

EAM Agents

The EAM Agents page shows a list of all the currently configured Agents, including information on the connection status. Agents can be organized into groups allowing you to group agents by location or agent function. If no groups exist, all Agents will be in the Default Group.



On this page ...

- [EAM Agents](#)
- [EAM Agents Page](#)

EAM Agents Page

The EAM Agents page contains some useful information about all of your Agent connections.

Attribute	Description
EAM Agent Connections	
EAM Agents	Number of connected agents out of the number of configured agents.
Default Group and all groups	
Gateway	Name of the Agent.
Edition	Identifies a Standard vs Edge Gateway. Edge edition shows "edge," and if it's a Standard Gateway, this field is left blank.
Status	Current status of the Agent connection.
Last Comm	Date and time recorded for the last communication with the Agent.
Last Event	Name of the event last recorded with the Agent.
Event Date	Date and time recorded for the last occurring event with the Agent.
Log Activity	Shows any EAM Agent activity along with the Time the log was created and a brief message. You can find more of these same type of messages in the Gateway Logger which is found in the Status section under Diagnostics > Logs , and the Wrapper Logger file under Program Files > Inductive Automation> Ignition > logs > wrapper.log .

[Configuration](#)

Connected Agents

1 / 1

Default Group

Gateway ▾	Edition	Status	Last Comm	Last Event	Event Date
Ignition-2		✓ Connected	25Mar2019 14:26:30	cpu [NORMAL]	25Mar2019 12:39:20

Log Activity



Min level	INFO	▼	?	Live Values	<input checked="" type="checkbox"/> ON	Merge to Logs
Logger						
E AgentLicenseDetailPage	25Mar2019 12:41:53			license.ipl cannot be loaded from the archive!		
I AgentEventRecorder	25Mar2019 12:37:43			EAM event record table successfully verified.		
I Agent	25Mar2019 12:37:43			[Ignition-2] Agent state has transitioned to Running		
I AgentModel	25Mar2019 12:36:59			New agent 'Ignition-2' has connected to this controller. It needs to be approved.		

Related Topics ...

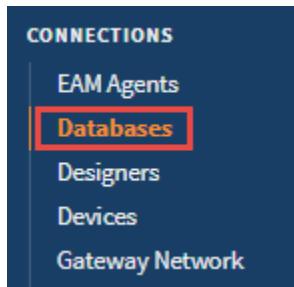
- [Enterprise Administration](#)
- [Creating a Controller](#)
- [Adding an Agent](#)

Connections - Databases

The Databases page shows a list of configured databases, and if they have a valid connection or not. Clicking on the **Details** button to the right of a connection will either show the full error if the connection is faulted, or it will bring you to a Details page for that database connection. On the Details page, you can easily see any active queries, long running queries, the number of queries a second that are running, as well as a trend showing the percentage of queries that completed in that time.

On this page ...

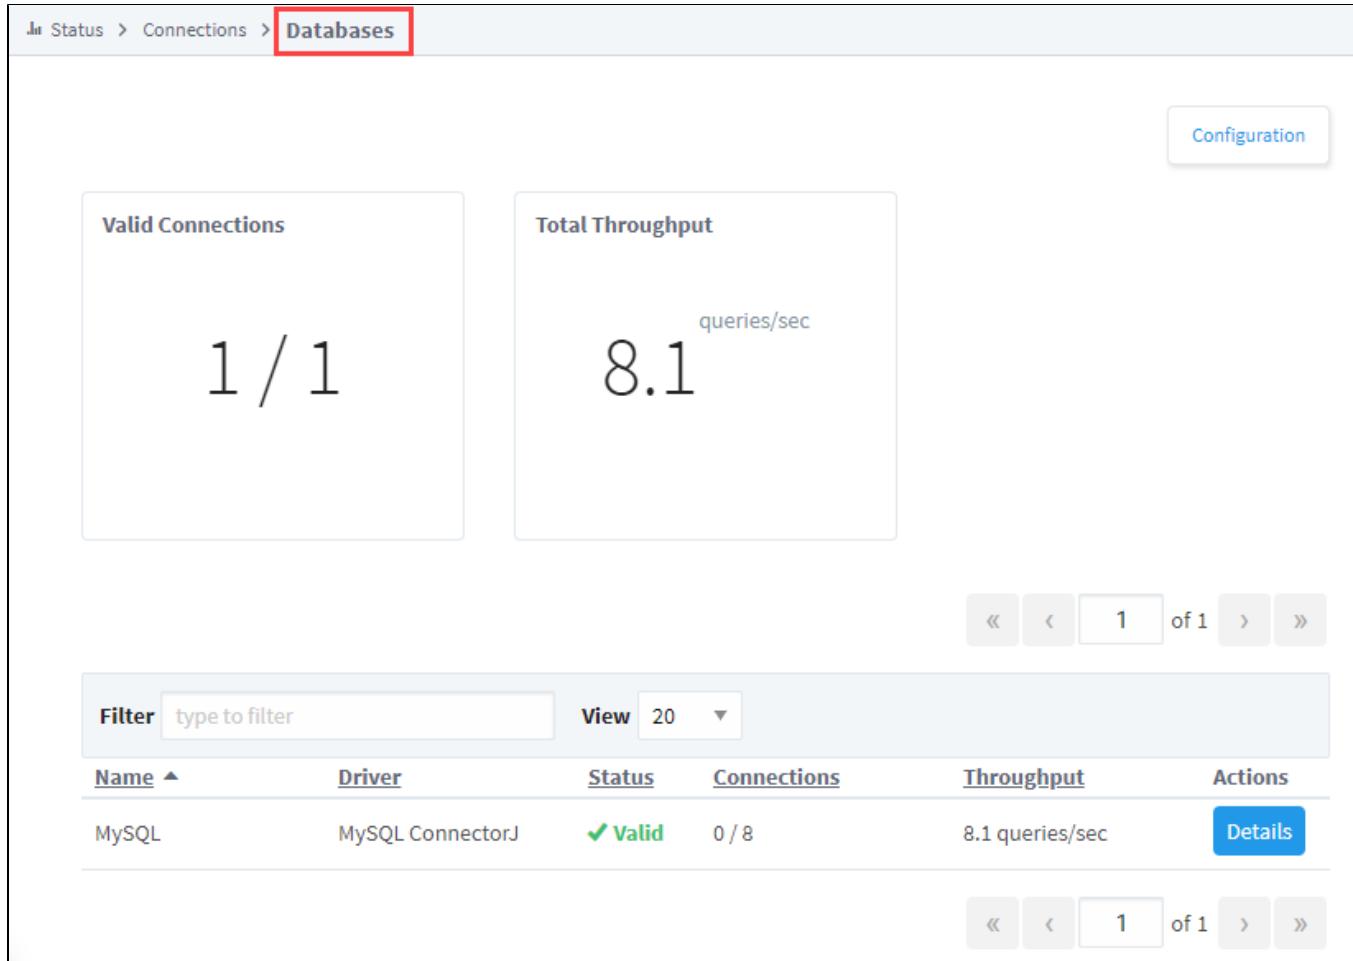
- [Databases Page](#)
- [Database Connection Details](#)



Databases Page

The main database page contains some useful information about all of your database connections.

Attributes	Description
Valid Connections	Number of valid connections.
Total Throughput	Number of queries and their status.
Name	Name of the query.
Driver	Name of the driver.
Status	Current status of the database connection.



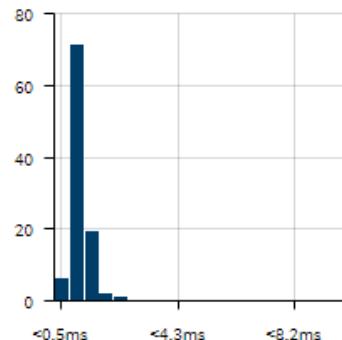
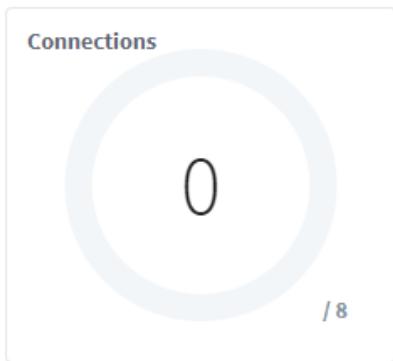
Database Connection Details

The **Details** page for an individual connection provides more in depth information for that particular database connection.

Attribute	Description
Database Stats - Stats about the database connection	
Connections	Number of database connections out of configured databases.
Queries / Sec	Number of queries running per second along with a trend showing the percentage of queries that completed in that time.
Active Queries - A list of currently active queries	
Query	Currently running queries.
Started	When the query was started.
Actions	The ability to cancel a query.
Longest Recent Queries - A list of the longest running queries.	
Query	Displays the actual query.
Duration	Amount of time the query ran.
Started	When the query started running.

[Configuration](#)

Database Stats



Active Queries

Query ▲	Started	Actions
No items to display.		

Longest Recent Queries

Query ▼	Duration	Started
SELECT a.`id`,a.`tagpath`,a.`datatype`,a.`scid`,a.`querymode`,a.`created`, b.`drvid` FROM sqlth_te a left join sqlth_scinfo b on a.`scid`=b.`id` WHERE b.`drvid` in (?, ?, ?) and a.`retired` is null or a.`retired`=0	30 ms	7 minutes ago
SELECT `id` FROM sqlth_scinfo WHERE `scname`=? AND `drvid`=?	8 ms	7 minutes ago

Related Topics ...

- [Database Connections](#)

Connections - Designers

Designers

The Designers page displays information on currently running designer sessions. All the open designers are displayed on the page along with some basic information about each session, such as what user is logged into each designer session and the project they are currently working on. Clicking on the **Details** button to the right of a designer session will display more information about that particular designer. On the Details page, you can see session information, as well as what designer locks the session currently has set. Locks are when a designer is working on a particular page or set of pages, the system places a lock on those resources (i.e., window, pipelines, etc.) to prevent other designers from working on the same resource. There is also a log at the bottom of the page displaying any errors pertaining to that designer session.



On this page ...

- [Designers](#)
 - [Designers Page](#)
 - [Designer Details](#)

Designers Page

The main Designers page has a list of all currently running designer sessions, and some basic information about all of them.

Attribute	Description
Designer Stats	
Active Designers	Number of active Designer connections.
Requests / Sec	Number of requests running per second along with a trend showing the percentage of requests that completed in that time.
Designer Sessions	
Filter	Search criteria to filter for specific designer sessions.
View	The number of designer sessions to preview.
Id	Designer session id number.
User	Name of the user logged into the Designer.
Project	Name of project the user is currently working on.
Uptime	Amount of time the user is logged into the designer session.
Status	If the Designer is actively connected or not.
Address	The IP Address and name of the computer running the designer.
Memory	Current memory usage for the session.

Status > Connections > Designers

Designer Stats

Active Designers

1

Requests / sec

5.1

Latency Bin	Requests / sec
<0.1ms	~42
<0.8ms	~20
<1.6ms	~1

Designer Sessions

Filter type to filter
View 20 ▾

Id	User	Project	Uptime	Status	Address	Memory	Actions
697B7559	admin	alarm-pipelines	11 minutes	✓ Active	TR-89MC8R2-WS (10.10.110.54)	128mb	<button style="border: 1px solid blue; padding: 2px;">Details</button>

«
‹
1
of 1
›
»

Designer Details

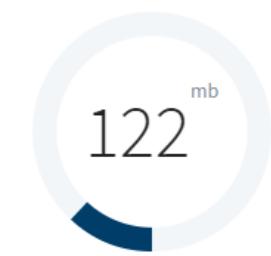
The Designer Details page shows more in depth information about that particular designer session, with the most important being the list of Designer Locks. The locks are project resources that the designer session is using. Because that designer session is using those project resources, no other designers will be able to access those resources, to prevent two designers from working on the same thing.

Attribute	Description
Session Details	
User	Name of the user logged into the Designer.
Project	Name of project the user is currently working on.
Address	The IP Address and name of the computer running the designer.
Uptime	Amount of time the user is logged into the designer session.
Memory	Current memory usage for the session.
Timezone	Local time of the user.
Log Activity	
Min Level	Dropdown menu with options Info, Debug, and Trace.
Live Values	Toggle switch to turn live values on or off.
Logger	Name of the logger that describes the context of the message.
Time	Time of log.
Message	Message for the log.

Session Details

User	admin
Project	alarm-pipelines
Address	TR-89MC8R2-WS (10.10.110.54)
Uptime	22 minutes
Timezone	America/Los_Angeles [GMT-8:00]

Memory



Log Activity



Min level

INFO



Live Values



Merge to Logs



Logger

Time

Message

T Gateway

26Mar2019 14:42:27

Writing push notifications

Related Topics ...

- Designer

Connections - Devices

Devices

The Devices page lists the currently configured OPC UA devices, and lets you know which are connected and which have a faulty connection. From here, you can drill into a device connection to see how many Tags Ignition is requesting from the device, along with how often it is requesting them. This information can be used to determine if we are overloading the device with too many requests too quickly, or if we can request more from our device.



On this page ...

- [Devices](#)
- [Devices Page](#)
- [Device Details](#)
 - [Device Statistics](#)

Devices Page

The main Devices page lists out all OPC UA devices, as well as how many currently have a valid connection. Note: this will only show the devices connected through an Ignition device connection. For information about devices connected through other OPC Servers, see those programs.

Attribute	Description
Connected Devices	Number of devices connected out of configured devices.
Name	Name of the device.
Driver	Name of the device driver.
Status	Current device status.

Connected Devices

1 / 1

« ‹ 1 of 1 › »

Filter type to filter **View** 20 ▾

Name	Driver	Status	Actions
Micro_1	MicroLogix	Connected: Protocol: EIP	Details

« ‹ 1 of 1 › »

Device Details

Clicking the **Details** button to the right will display a diagnostics page for that device which provides metrics for the device configuration. The page contains an Aggregate Statistics table and lists additional statistics to help determine if the device is overloaded with requests. There are values for each subscription (such as those created by Tag Groups or Transaction Groups using OPC items) from the specified device, as well as aggregate statistics which pull from all subscriptions to get an average for the device.

Device Statistics

The table below summarizes the statistics available on the Device Details page.

Attribute	Description
Request Count	Tracks the number of requests that are coming in from the device. A request is a group of tags/items that the driver has grouped together to be read at the same time. Each driver forms these groups based on the protocol being used and occasionally configuration settings in the driver.
Throughput (Mean)	Represents the average amount of requests that come through per second since the device was last started. If the device connection is edited and saved, this will cause the device connection to reinitialize and this value will be reset.
Throughput (1 min)	Represents the average amount of requests that come through per second for the last minute.
Mean Response Time	The average time it takes for Ignition to get a response from the device. This number is an average based of the graph on the right of the page.
Monitored Item Count	Represents the total number of items that are subscribed to.

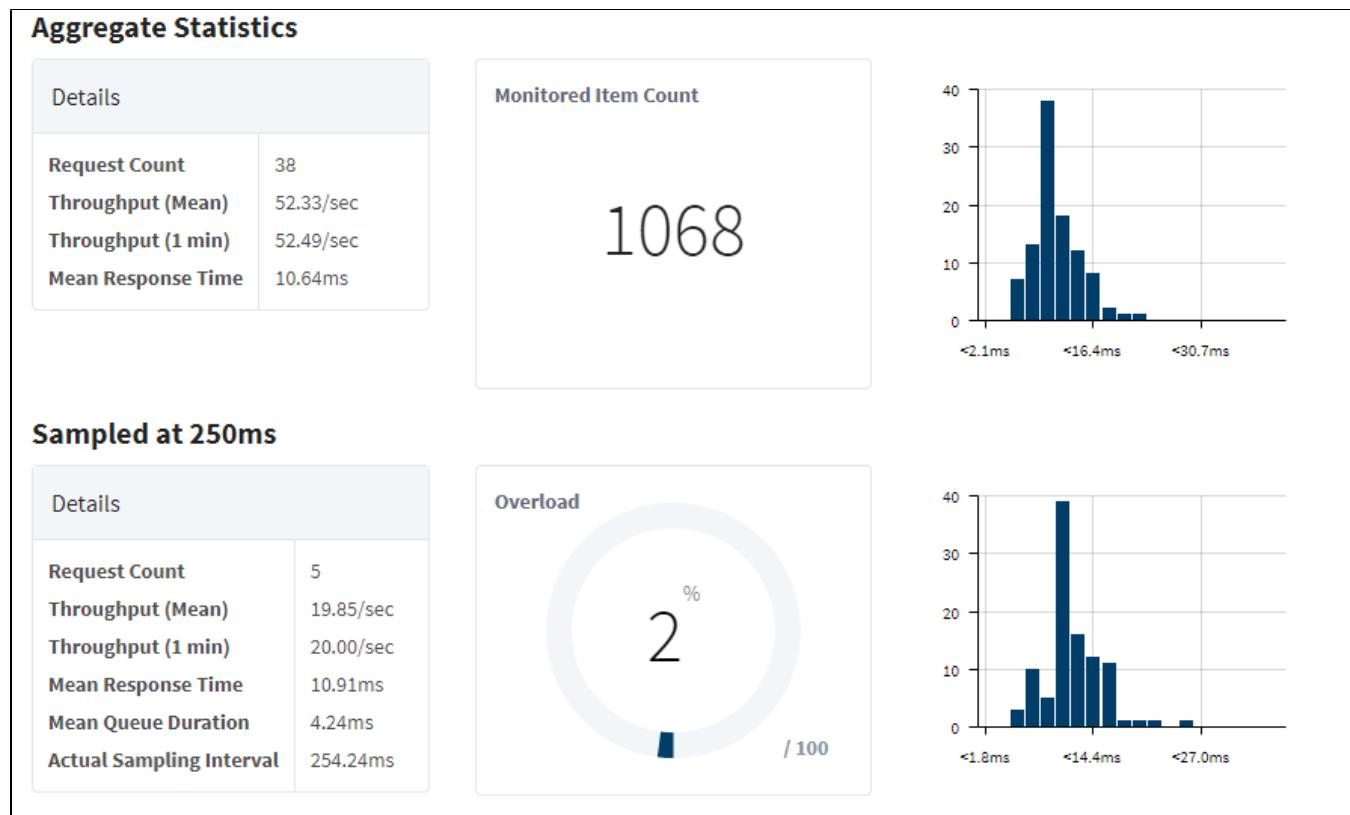
	Note that this count includes device diagnostic tags. However, other statistics on the device status page do not account for diagnostic tags.
Mean Queue Duration	<p>The following feature is new in Ignition version 8.1.6 Click here to check out the other new features</p> <p>Represents the average amount of time a request has spent in the request queue. Not available under aggregate statistics.</p>
Actual Sampling Interval	<p>The following feature is new in Ignition version 8.1.6 Click here to check out the other new features</p> <p>Represents the actual rate the driver is sampling the device at. Used to determine Overload. Not available under aggregate statistics.</p>



Note that the metrics above have corresponding [diagnostics tags](#) available, allowing other areas of Ignition to access these values

In the image below, we see there are a total of 1068 items being monitored, which are optimized into 38 different requests.

Of those requests, five are being sampled at a 250ms rate. At this same rate, requests sit in a request queue for an average of 4.24ms before being processed. Thus the actual sampling interval for this sampling rate is 254.24ms.



Overload

The following feature is new in Ignition version **8.1.6**
[Click here](#) to check out the other new features



The overload metric was introduced in 8.1.6, replacing the prior Load Factor metric. See the [8.0 Connections - Devices page](#) for more information.

Overload represents how well the device is able to keep up with the requests at the sample rate. It is calculated using the formula:

$$100 * (\text{Mean Queue Duration} / \text{Sampling Interval}) = \text{Overload}$$

If overload exceeds 100% then the requests are sampling at a slower than ideal rate.

Based on the image above, overload for our 250ms sample group is determined by:

$$100 * (4.24 / 250) = 1.696$$

Our value of 1.696 can then be rounded up to 2%. In this case, it doesn't seem as if the device is having much trouble keeping up with our requests.

Connections - Gateway Network

The Gateway Network Status page is designed to give a quick overview of the status of all Gateways within the Gateway Network. If a Gateway connection is faulted, the status message in red can be selected to see the error that pertains to why it is faulted. Any Gateway connections with a status of 'Connected' can be drilled into by clicking the **Details** button to the right. On the Details page, metrics for the selected Gateway connection are displayed, giving an idea of the rate of data transfer between the two Gateways, as well as a list of recent connection events.

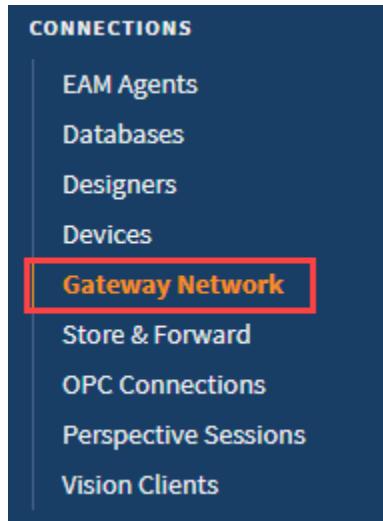
This page gives a general overview of the Connections - Gateway Network webpage. See the [Gateway Network Connection Details](#) page for more in-depth information.

On this page ...

- [Gateway Network Page](#)

The following feature is new in Ignition version **8.1.25**
[Click here](#) to check out the other new features

The Gateway Network Status page also has a tab which visualizes data about your Gateway Network. See the [Gateway Network Diagram](#) page for more information.



Gateway Network Page

The main Gateway Network page displays a list of all current Gateway Network connections both incoming and outgoing, as well a list of all Remote Gateways that the Local Gateway can see both from its Gateway Network connections and through proxy connections. Each list has some basic information along with the ability to see more details on a specific Gateway Network connection or a Remote Gateway.

Attribute	Description
Gateway Network Connections	
Remote Gateway	Name of the Remote Gateway connection.
Direction	The direction of the Gateway Network connection. Can either be Incoming or Outgoing.
Redundancy Role	The redundancy role of the Gateway. Can either be Independent, Backup, or Master.
Last Comm	The time of the last communication with the Gateway.
Ping Time	Reaction time of Gateway connection. How fast you get a response after you've sent out a request.
Status	Current state of the Gateway Network Connection.
Fault Count	Number of times the connection has faulted since the Gateway has been started.
Remote Gateways	
Gateway Name	Name of the Remote Gateway.
Outgoing Msg/Sec	The number of outgoing messages per second.
Incoming Msg/Sec	The number of incoming messages per second.

Pending	Number of messages pending in a queue that are waiting to be dispatched to the Gateway Network connection.
Connected Through	How the Gateway is connected to the Remote Gateway.
Status	Current state of the Remote Gateway Connection.

Home Status > Connections > **Gateway Network**

Configuration

Network Live Diagram

Connections

« < 1 of 1 > »

Filter	type to filter	View	20				
Remote Gateway	Direction	Redundancy Role	Last Comm	Ping Time	Status	Fault Count	Actions
gateway+b	Outgoing	Independent	a few seconds ago	25 ms	Running	0	<button>Details</button>

« < 1 of 1 > »

Remote Gateways

« < 1 of 1 > »

Filter	type to filter	View	20				
Gateway Name	Outgoing Msg/Sec	Incoming Msg/Sec	Pending	Active	Routed Through	Status	Actions
Gateway B	0.2	0.4	0	0	gateway+b	Connected	<button>Details</button>
	0.0	0.0	0	0		Defined_NeverConnected	<button>Details</button>
	0.0	0.0	0	0		Defined_NeverConnected	<button>Details</button>

« < 1 of 1 > »

Related Topics ...

- [Gateway Status](#)
- [Diagnostics - Logs](#)

In This Section ...

Gateway Network Connection Details

Gateway Network Connection Details

Ignition's Gateway Network system shares information across Gateways using threads to send and receive information. For example, if you have a Remote Tag Provider configured between Gateway A and Gateway B, messages containing live tag information will be sent between these two Gateways using threads to send and receive live tag data. Similarly, a Remote Historical Tag Provider will send/receive messages with historical tag data between Gateway A and Gateway B using this same set of threads. Each Ignition sub system that uses the Gateway Network will utilize these threads in some way.

Ignition's Gateway Network also has a queue associated with each Ignition subsystem. These queues allow for Ignition to have a way to prioritize which sub system should have access to a send or receive thread. Prioritization here is especially important because if all send threads are in use, messages cannot be sent between two Gateways. The Gateway Network page includes additional information that will help better monitor Gateway interactions.

On this page ...

- [Gateway Network Connection Details](#)
- [Active Outgoing and Incoming Tasks List](#)
- [Gateway Network Statistics](#)
- [Outgoing Queues](#)
- [Temporary Queue Actions](#)
- [Outgoing and Incoming Tasks Statistics](#)

The screenshot shows the Ignition interface for managing gateway network connections. At the top, there's a navigation bar: Status > Connections > Gateway Network. On the right, there's a 'Configuration' button. Below the navigation, there are two tabs: 'Network' (selected) and 'Live Diagram'. Under 'Connections', there's a table with columns: Remote Gateway, Direction, Redundancy Role, Last Comm, Ping Time, Status, Fault Count, and Actions. One row is shown: gateway+b, Outgoing, Independent, a few seconds ago, 25 ms, Running, 0, and a 'Details' button (which is highlighted with a red box). There are also navigation arrows and a page number '1 of 1'. Below this, under 'Remote Gateways', is another table with columns: Gateway Name, Outgoing Msg/Sec, Incoming Msg/Sec, Pending, Active, Routed Through, Status, and Actions. It lists three entries: Gateway B (Connected), and two entries with status 'Defined_NeverConnected'. Each entry has a 'Details' button. Navigation arrows and a page number '1 of 1' are also present here.

Remote Gateway	Direction	Redundancy Role	Last Comm	Ping Time	Status	Fault Count	Actions
gateway+b	Outgoing	Independent	a few seconds ago	25 ms	Running	0	Details

Gateway Name	Outgoing Msg/Sec	Incoming Msg/Sec	Pending	Active	Routed Through	Status	Actions
Gateway B	0.2	0.4	0	0	gateway+b	Connected	Details
	0.0	0.0	0	0		Defined_NeverConnected	Details
	0.0	0.0	0	0		Defined_NeverConnected	Details

The image above shows the basic Gateway Network Connection Status page where we can see there is an outgoing Gateway Network connection from Gateway A to Gateway B. Pressing the Details button will bring up detailed information about this connection as below:

Status

Running

Average Outgoing Bytes

1 KB /sec

Average Incoming Bytes

2 KB /sec

Ping Time

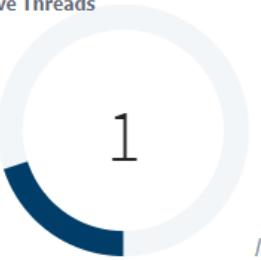
1 ms

Send Threads



0 / 5

Receive Threads



1 / 5

Active Outgoing Tasks

<u>Task Name</u> ▾	<u>Source Queue</u>	<u>Duration Secs</u>
CallResult:Services/Get Service State	Call Results Queue	30354
CallResult:Services/Get Service State	Call Results Queue	401254
CallResult:Services/Get Service State	Call Results Queue	554102

Active Incoming Tasks

<u>Task Name</u> ▾	<u>Duration Secs</u>
No items to display.	

Attribute	Description
Status	Current state of the Gateway Network Connection.
Average Outgoing Bytes	Average bytes of data going from the Local Gateway to the Remote Gateway per second.
Average Incoming Bytes	Average bytes of data going from the Remote Gateway to the Local Gateway per second.
Ping Time	Reaction time of Gateway connection. How fast you get a response after you've sent out a request.
Active	Number of messages being actively processed by the Gateway network connection.
Send Threads	A thread that is used by the Gateway Network to upload messages from one Gateway to another.
Receive Threads	A thread that is used by the Gateway Network to download messages from one Gateway to another.
Local Id	Id of the local Gateway.
Remote Id	Id of the remote Gateway.
Remote Gateway	Name of the Gateway on the Gateway network.
Network Address	Physical address used to communicate with all devices on the Gateway network.
Redundant Role	The redundancy role of the Gateway. Can either be Independent, Backup, or Master.
Direction	The direction of the Gateway Network connection. Can either be Incoming or Outgoing.
Session Id	Gateway connection session Id number. Connected Gateways use the same session Id on both Gateways.
Last Comm	The time of the last communication with the Gateway.

Fault Count	Number of times the connection has faulted since the Gateway has been started.
Connection Events	Displays a list of recent connection events.

Active Outgoing and Incoming Tasks List

Attribute	Description
Task Name	Name of the task that is using a thread.
Source Queue	The Ignition sub system queue that dispatched this task.
Duration Secs	This is how long in seconds it takes for a task to be performed.

Gateway Network Statistics

In addition to showing live thread and task information, users can also keep track of Gateway Network Statistics. By clicking on the Details button below, the Gateway Network Statistics Page appears for Gateway B:

The screenshot shows the Ignition software interface with the following details:

- Header:** Status > Connections > Gateway Network
- Configuration:** A button in the top right corner.
- Network Tab:** Active tab (highlighted in blue).
- Live Diagram:** A small button next to the Network tab.
- Connections Section:**
 - Filter:** Type to filter.
 - View:** 20 items.
 - Table Headers:** Remote Gateway, Direction, Redundancy Role, Last Comm, Ping Time, Status, Fault Count, Actions.
 - Data:** One row for "gateway+b" with "Outgoing" direction, "Independent" redundancy role, "a few seconds ago" last comm, 25 ms ping time, "Running" status, 0 faults, and a "Details" button.
 - Pagination:** 1 of 1.
- Remote Gateways Section:**
 - Filter:** Type to filter.
 - View:** 20 items.
 - Table Headers:** Gateway Name, Outgoing Msg/Sec, Incoming Msg/Sec, Pending, Active, Routed Through, Status, Actions.
 - Data:** Three rows for "Gateway B". The first row has 0.2 outgoing, 0.4 incoming, 0 pending, 0 active, and "Connected" status. It has a red-bordered "Details" button. The other two rows have 0.0 outgoing, 0.0 incoming, 0 pending, 0 active, and "Defined_NeverConnected" status. They have standard blue "Details" buttons.
 - Pagination:** 1 of 1.

Outgoing Queues

The Gateway Network Statistics page has three sections associated with it. First, the Outgoing Queues section. The Outgoing Queues section shows tasks that are both in a pending and active state. A pending task is a task that has not yet been dispatched to the Gateway Network thread pool. An active task is a task that is being processed by the Gateway Network thread pool and should show up under the Gateway Network Connection Status page as either an outgoing or incoming task.

Outgoing Queues							
Name ▲	Priority	Inserts/Sec	Pending	Active	Avg Pending Secs	Total	Actions
Call Results Queue	Highest	0.5	0	0	0.0	333,299	<button>More ▾</button> <button>Clear</button>
Default Queue	Normal	0.0	0	0	0.0	2	<button>More ▾</button> <button>Clear</button>
Long Wait Queue	Low	0.0	0	0	0.0	0	<button>More ▾</button> <button>Clear</button>
Proxy Queue	AboveNormal	0.0	0	0	0.0	0	<button>More ▾</button> <button>Clear</button>
Tag Value Publishing	Normal	0.0	0	0	0.0	0	<button>More ▾</button> <button>Clear</button>

Attribute	Description
Name	Name of the queue.
Priority	Level of priority for a queue.
Inserts/Sec	Rate of task inserts per second for a queue
Pending	Number of pending tasks in a queue that have not yet been dispatched to a Send/Receive thread.
Active	Number of messages being actively processed by the Gateway Network connection.
Avg Pending Secs	Average number of seconds that a task has been pending in a queue.
Total	Total number of tasks executed from this queue.
Actions	Set of actions associated with a queue. Users can both Pause and Clear a queue.

Temporary Queue Actions

The Gateway Network Statistics page has controls for pausing and clearing a queue. To **pause** a queue means no new tasks will be allowed to be inserted into the paused queue. To **clear** a queue means that all pending tasks will be purged. These actions are designed to help the user deal with a possibly overloaded Gateway Network connection due to a specific sub system flooding the queue with more tasks than the connection can handle. Note that neither of these actions will have any effect on active tasks, as they have already been dispatched to the Gateway Network connection and cannot be cancelled.

Outgoing Queues							
Name ▲	Priority	Inserts/Sec	Pending	Active	Avg Pending Secs	Total	Actions
Call Results Queue	Highest	0.5	0	0	0.0	333,313	<button>More ▾</button> <button>Clear</button>
Default Queue	Normal	0.0	0	0	0.0	2	<button>More ▾</button> <button>Clear</button>
Long Wait Queue	Low	0.0	0	0	0.0	0	<button>More ▾</button> <button>Clear</button>

Outgoing and Incoming Tasks Statistics

The Outgoing and Incoming Task Statistics section shows individual tasks that have been processed by the Gateway Network Connection.

Outgoing Task Statistics

Name ▲	Description	Queue	Invocations/Sec	Avg Duration Secs	Total
Remote Tags (v7)/getProviders	Returns a list of tag providers	Remote Tags (v7)	0.0	0.3	1
Remote Tags/cancelSubscription	Cancel a remote tag subscription	Remote Tags	0.0	0.4	1
Remote Tags/getProperties	Return a list of tag provider properties for a specified provider	Remote Tags	0.0	0.4	2
Remote Tags/modifySubscription	Modify a remote tag subscription	Remote Tags	0.0	0.4	1
Remote Tags/validateSubscription	Validates that tag subscriptions between gateways is synchronized	Remote Tags	0.1	0.6	39
Services/Enumerate Services	Returns a list of available services on a remote machine	Default Queue	0.0	0.7	12
Tag History Storage Service/canAcceptData	Indicates whether the local historian can store new data	Tag History Storage Service	0.0	0.4	1
Tag History/getHistoricalProviders	Returns a list of historical providers	Tag History	0.0	0.3	1

Incoming Task Statistics

Name ▲	Description	Invocations/Sec	Avg Duration Secs	Total
CallResult:Remote Tags (v7)/getProviders	Returns a list of tag providers	0.0	0.0	1
CallResult:Remote Tags/cancelSubscription	Cancel a remote tag subscription	0.0	0.0	1
CallResult:Remote Tags/getProperties	Return a list of tag provider properties for a specified provider	0.0	0.0	2
CallResult:Remote Tags/modifySubscription	Modify a remote tag subscription	0.0	0.0	1
CallResult:Remote Tags/validateSubscription	Validates that tag subscriptions between gateways is synchronized	0.1	0.0	39
CallResult:Services/Enumerate Services	Returns a list of available services on a remote machine	0.0	0.0	10
CallResult:Tag History Storage Service/canAcceptData	Indicates whether the local historian can store new data	0.0	0.0	1
CallResult:Tag History/getHistoricalProviders	Returns a list of historical providers	0.0	0.0	1

Outgoing and Incoming Tasks Attributes	Description
Name	Name of a task.
Description	Description for a task.
Queue	Queue invoking a task (For Outgoing Tasks only)
Invocation/Sec	Rate at which a task is invoked per second.
Avg Duration Secs	Average duration in seconds of the time it takes for this task to execute.
Total	Number of times the task has been executed in total.

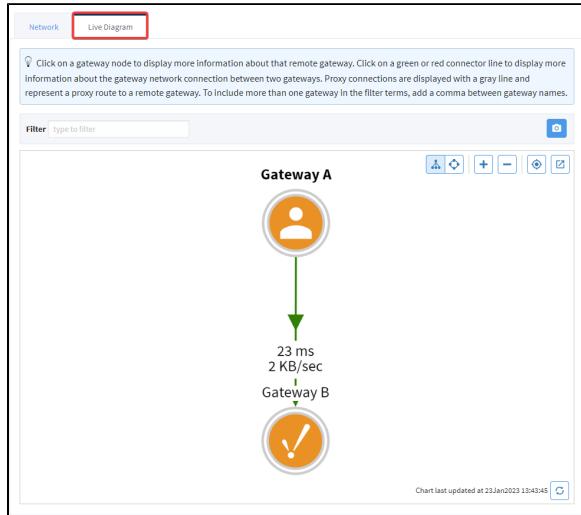
Note: If you are looking for connection details for Ignition version 8.0.14 and earlier, refer to [Connections - Gateway Network](#) page in version 8.0 of the Ignition User Manual.

Gateway Network Diagram

The following feature is new in Ignition version **8.1.25**
[Click here](#) to check out the other new features

Gateway Network Diagram Overview

The **Gateway Network Diagram** is a visual representation of your Gateway Network. You can access the **Live Diagram** tab by going to your Gateway webpage's Status tab > Connections > Gateway Network > Live Diagram.

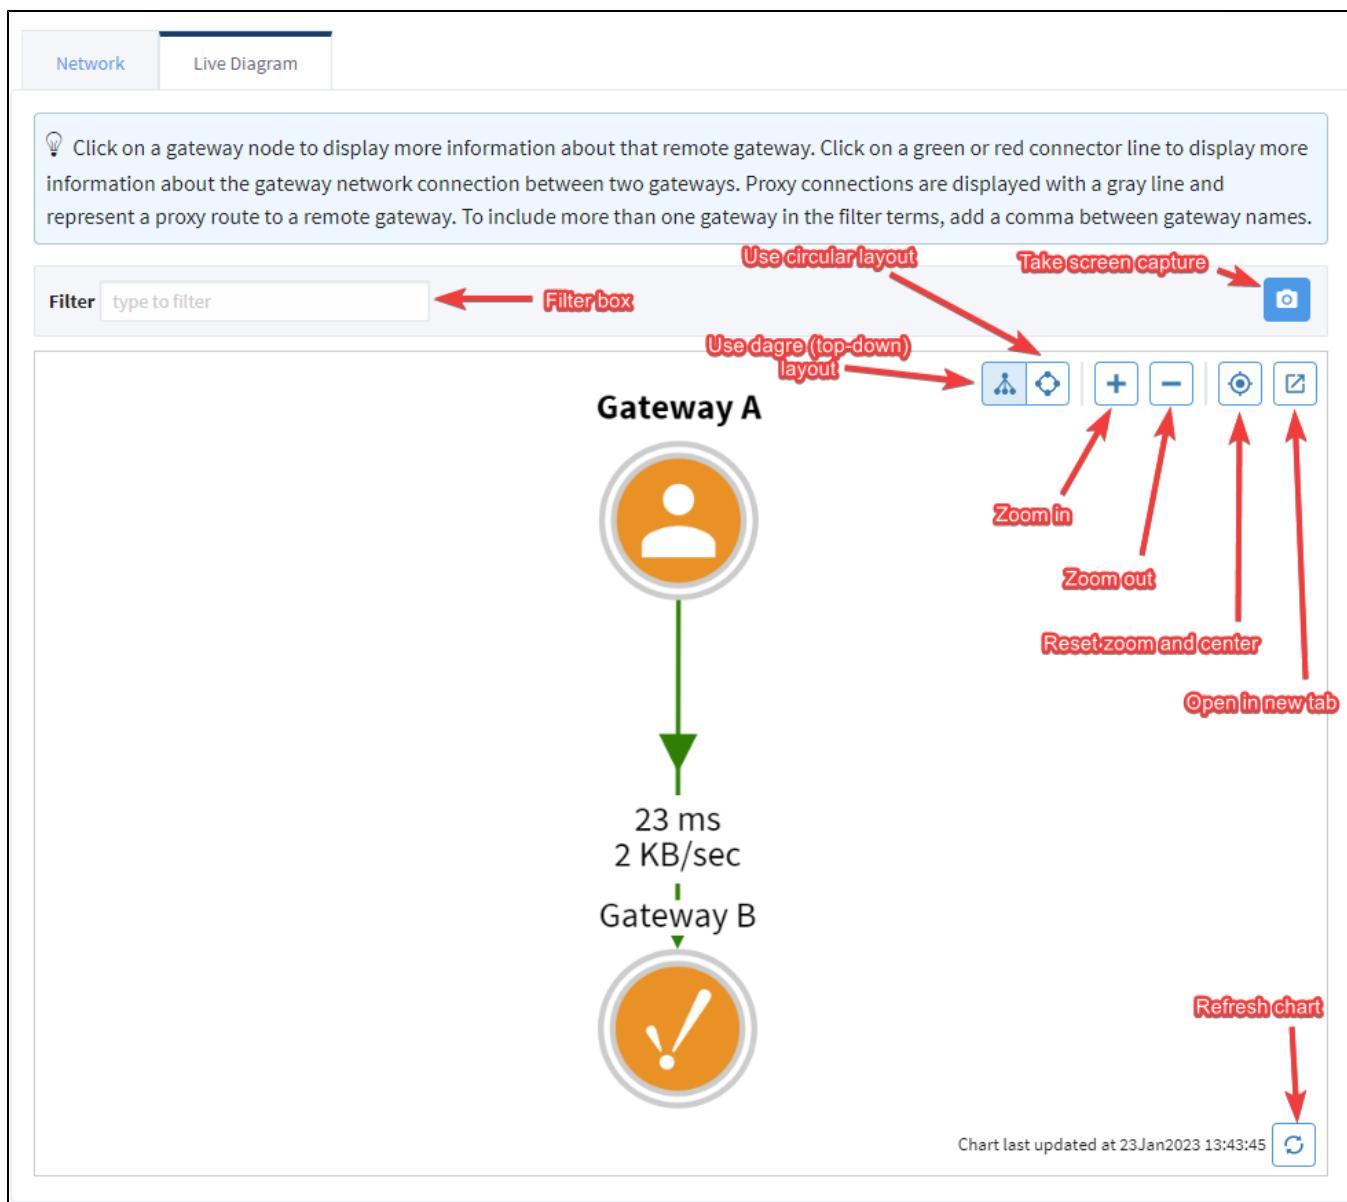


On this page ...

- [Gateway Network Diagram Overview](#)
 - [Gateway Network Diagram User Interface](#)
- [Gateway Network Diagram Details](#)
 - [Live Diagram Icons](#)
 - [Connection Details](#)
 - [Gateway Nodes](#)
 - [Redundancy Representation](#)
 - [Proxy Connections](#)

Gateway Network Diagram User Interface

To navigate the Diagram, you can use your mouse's left-click to pan around and your mouse's scroll wheel to zoom in or out. Additionally, you can drag and rearrange different nodes as you like. There are also different buttons on the Diagram interface that may be helpful:



Icon	Element	Description
Filter type to filter	Filter box	Filter by Gateway name. You can use comma-delimited Gateway names so that multiple Gateways show through the filter.
	Take screen capture	Takes a screenshot of the currently displayed Gateway Network Diagram.
	Use dagre layout	Display the Gateway Network Diagram in a dagre, or top-down layout from the Perspective of the local Gateway.
	Use circular layout	Display the Gateway Network Diagram in a circular layout, with the center being the local Gateway.
	Zoom in	Zoom into the Gateway Network Diagram.
	Zoom out	Zoom out of the Gateway Network Diagram.

	Reset zoom and center	Resets the camera panning and zoom and centers the Gateway Network Diagram.	
	Open in new tab	Opens the Gateway Network Diagram in a new, separate tab.	
	Refresh chart	Refreshes the Gateway Network Diagram so that the most recent Gateway Network configuration is shown.	

Gateway Network Diagram Details

Live Diagram Icons

The following table shows each Gateway Network Diagram icon and its description.

Icon	Description
	The local node from which you are viewing the Gateway Network Diagram. This icon's orange color also indicates that the node is running Standard Edition.
	The Gateway node that the local Gateway is connected to. This icon's orange color also indicates that the connected node is running Standard Edition.
	The local node from which you are viewing the Gateway Network Diagram. This icon's green color also indicates that the node is running Edge Edition.
	The Gateway node that the local Gateway is connected to. This icon's green color also indicates that the connected node is running Edge Edition.
	The local node from which you are viewing the Gateway Network Diagram. The icon's multicolor also indicates that the node is running Maker Edition.
	The Gateway node that the local Gateway is connected to. The icon's multicolor also indicates that the connected node is running Maker Edition.
	The local node from which you are viewing the Gateway Network Diagram. The icon's blue color also indicates that the node is running Cloud Edition
	The Gateway node that the local Gateway is connected to. The icon's blue color also indicates that the connected node is running Cloud Edition.

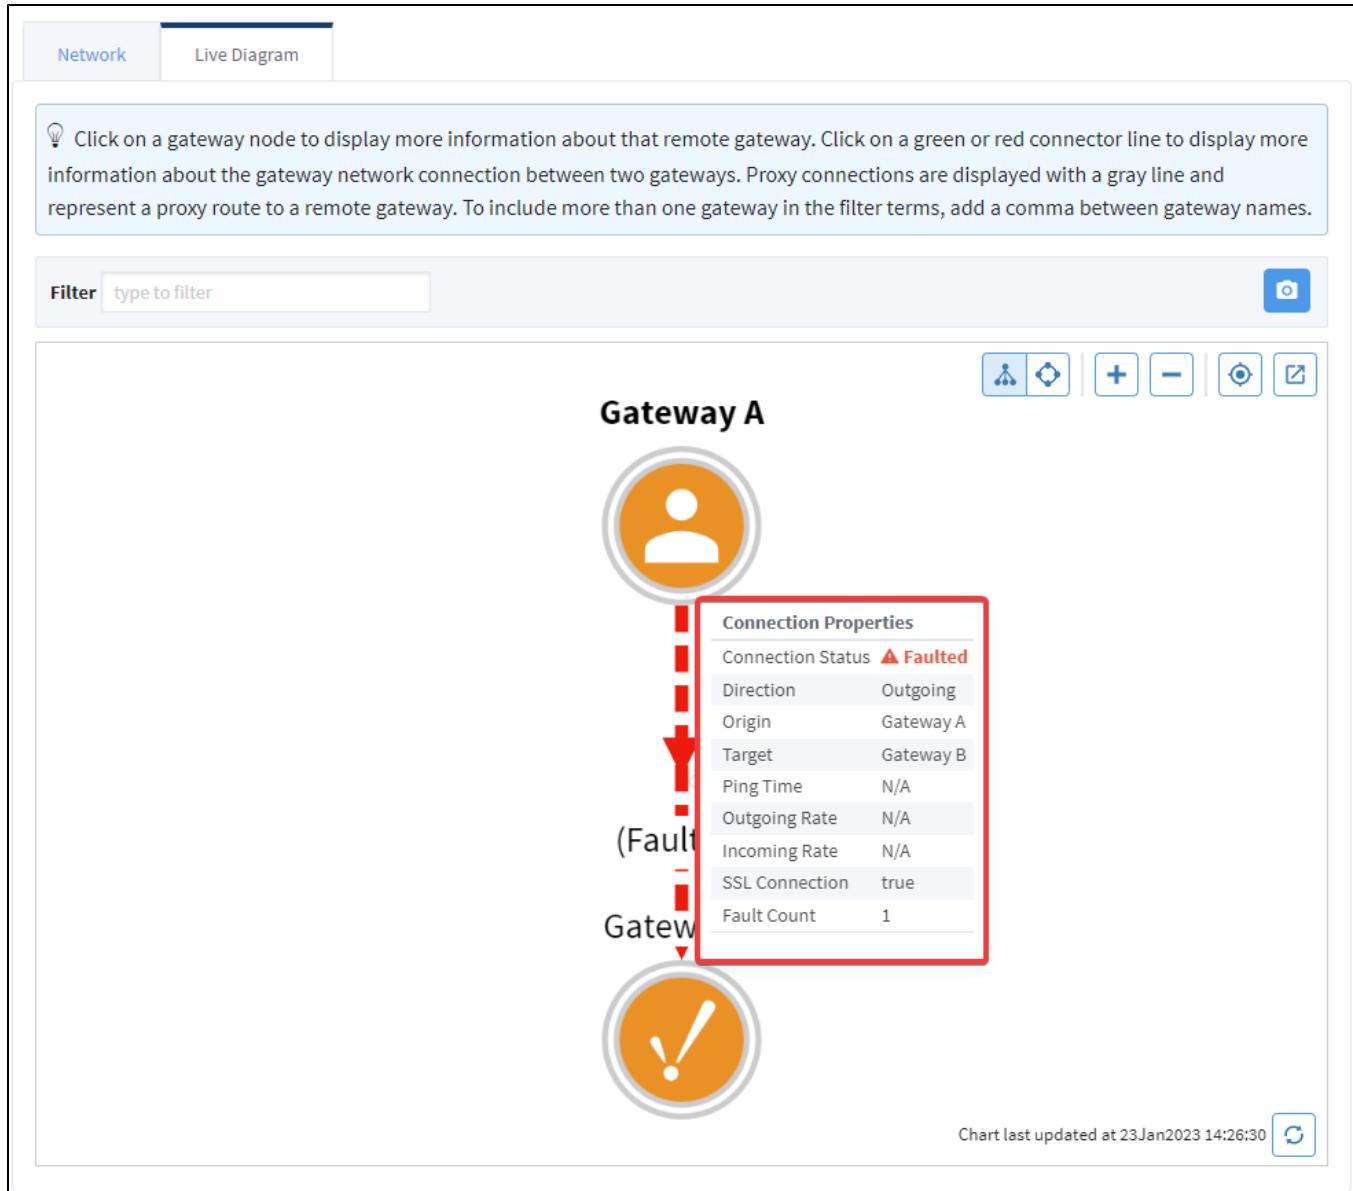
Connection Details

You can find out more information about each node and connection by clicking on them in the diagram. For example, clicking on the connection itself will display information such as:

- Connection Status
- Direction (Outgoing versus Incoming)
- Origin Gateway and Target Gateway
- Ping time in milliseconds
- Rates (Outgoing and Incoming bytes per second)
- Whether the connection is using SSL
- How many times the connection has faulted



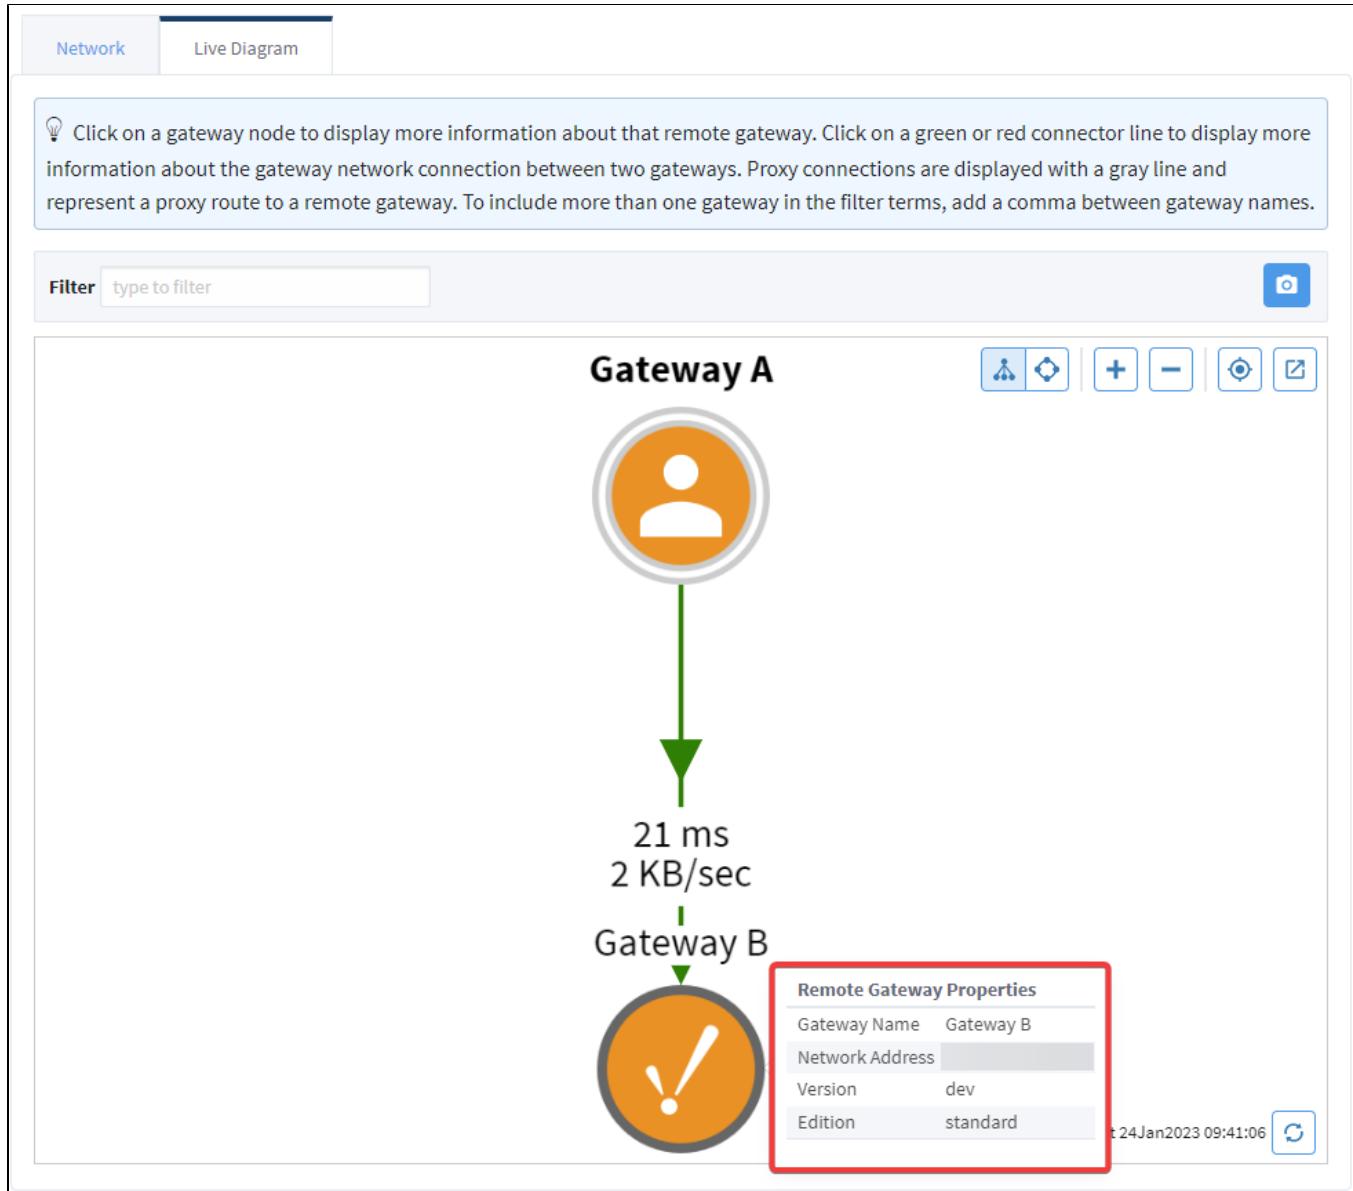
If the connection is in another state besides **Running**, the Diagram will update accordingly:



Gateway Nodes

On the other hand, clicking on the Gateway node itself will display information such as:

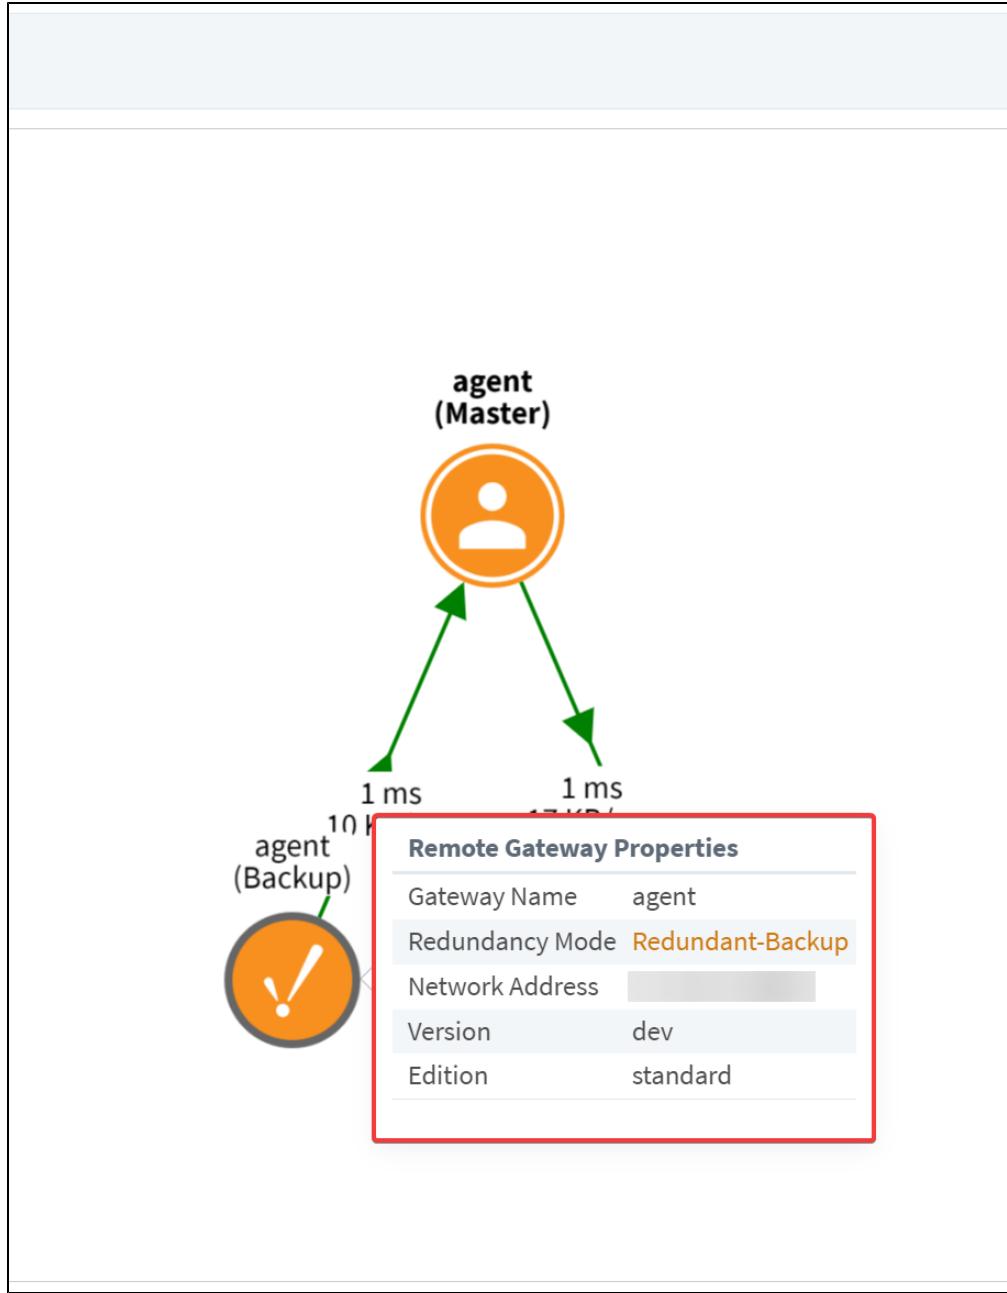
- Gateway Name
- Network Address
- Version
- Edition



Note: If a Gateway in your Gateway Network is running Ignition version 8.1.24 or below, the Version and Edition fields will be N/A.

Redundancy Representation

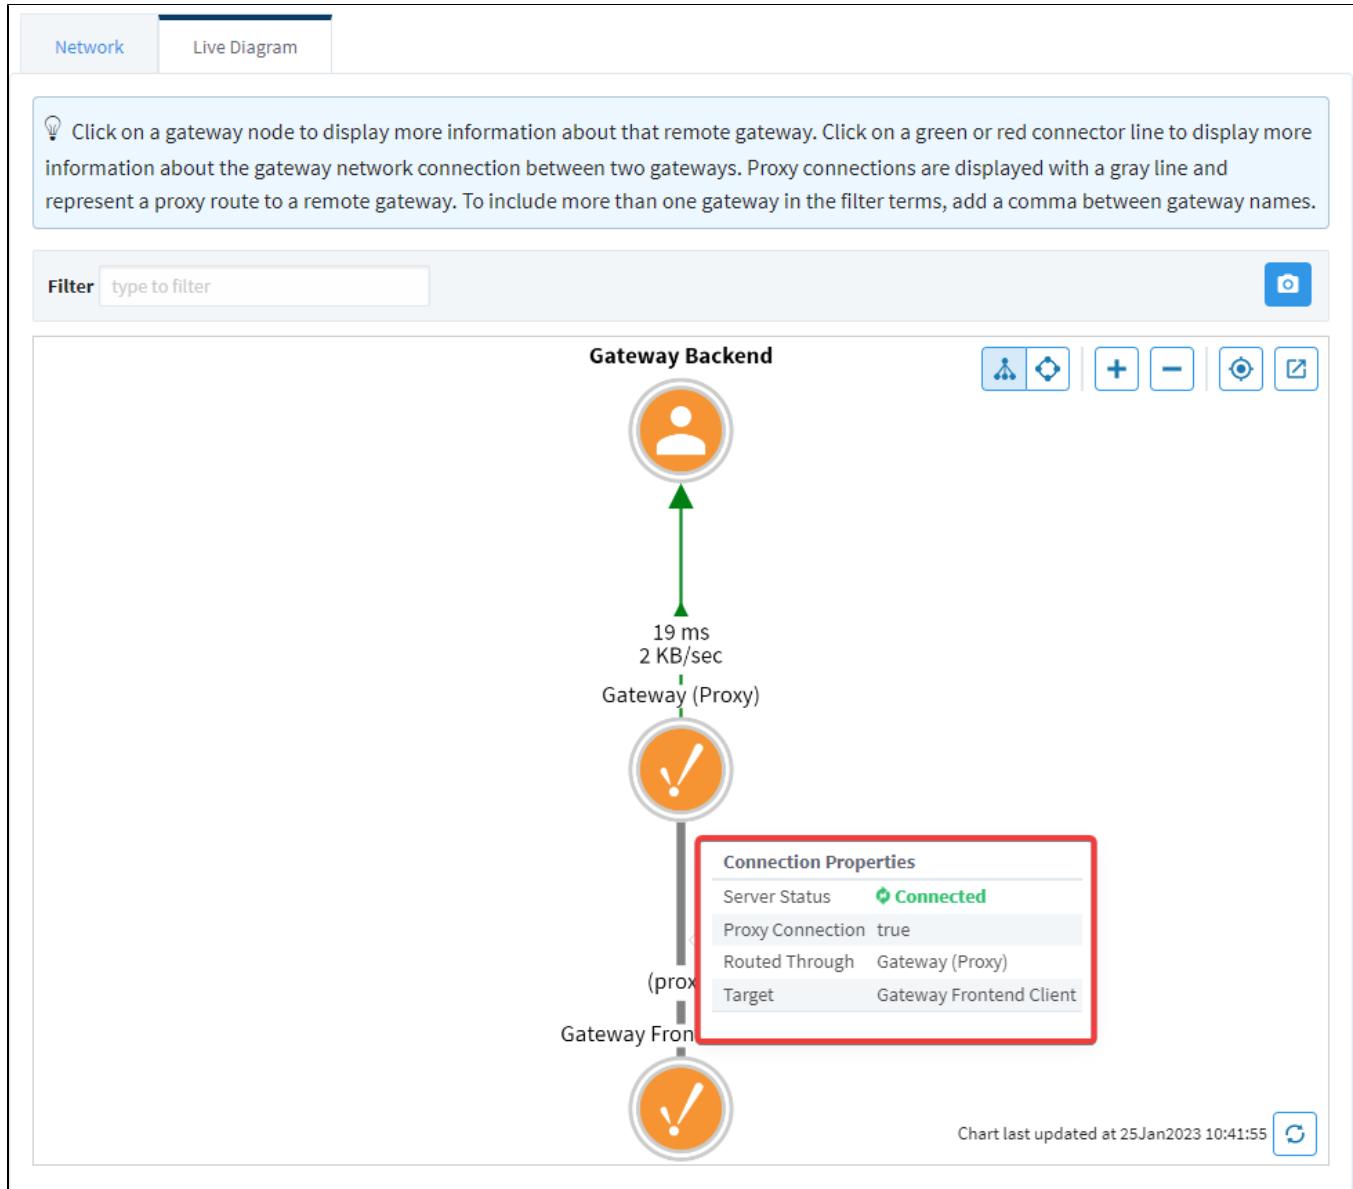
If two Gateways are part of a redundant pair, the diagram will show each Gateway node with an orange circle around it. In addition to the Gateway Name, Network Address, Version, and Edition, the Gateway Node will also display the **Gateway's role** in the redundant pair:



Proxy Connections

The Gateway Network Diagram will also display proxy connections differently. Instead of a green arrow indicating status and direction, the connection will be grey with no direction indicator. Clicking on a proxy connection will give the following information:

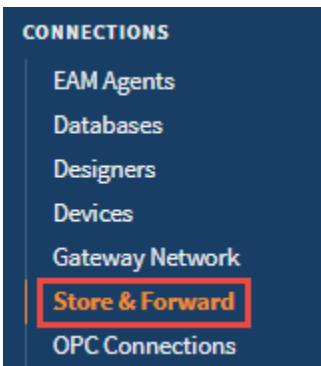
- The status of the server
- Whether or not the two Gateways are using a Proxy connection
- Which Gateway the proxy is routed through
- The proxy connection's target Gateway



Connections - Store & Forward

Store & Forward

The [Store and Forward](#) page displays a list of the Store and Forward engines, including their status, as well as the number of records currently in each Store and Forward system. If the database connection becomes faulted, the database records wait in the Store and Forward system until the database connection is restored.



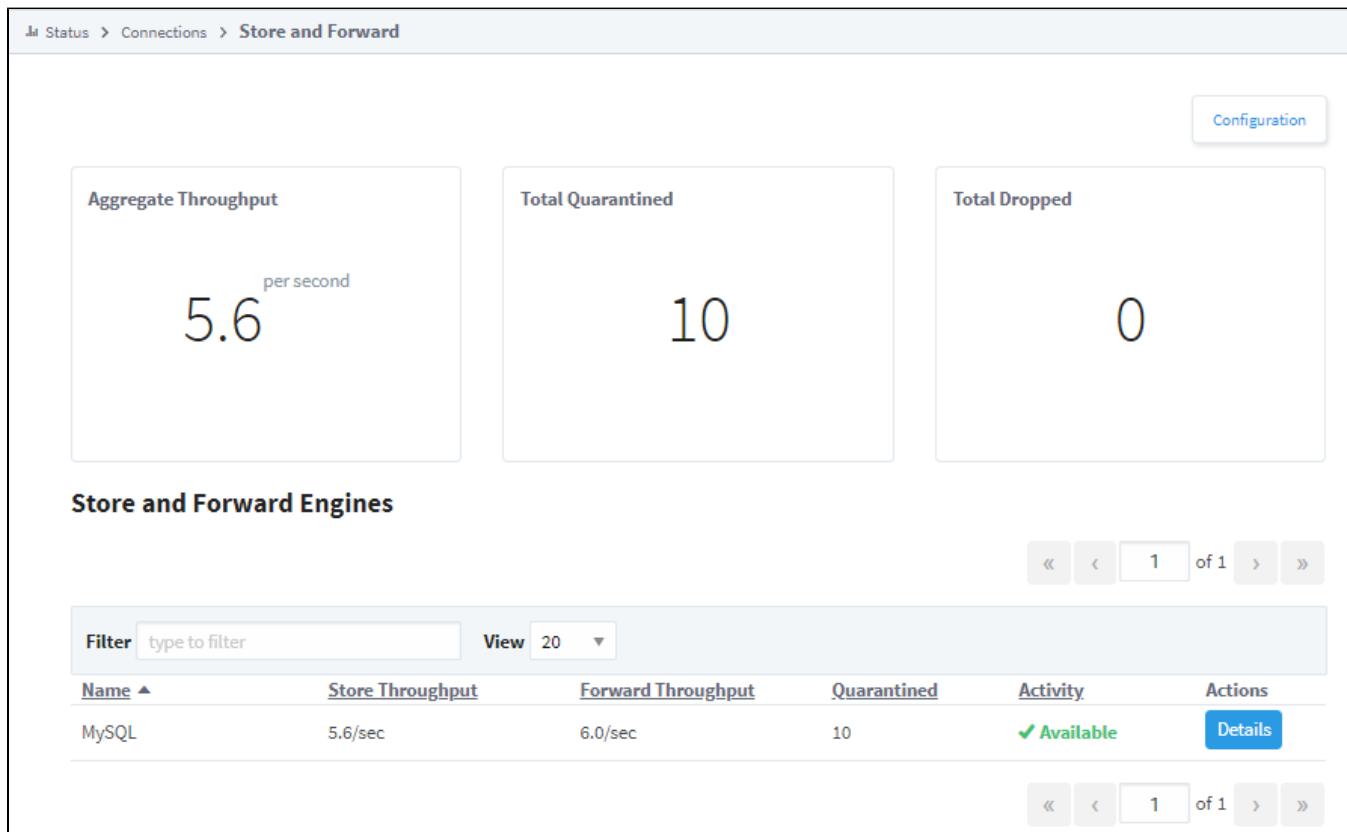
On this page ...

- [Store & Forward](#)
 - [Store & Forward Page](#)
 - [Store & Forward Details](#)

Store & Forward Page

The main Store and Forward page lists out all store and forward engines. Typically, each database connection gets its own store and forward engine, so there should be one engine for each database connection. In addition to displaying some basic stats for each engine, you can also find some totals for all store and forward systems, to get an idea of how much data is being pushed through the system to databases and if any records are being dropped.

Attributes	Description
Store and Forward Connections	
Aggregate Throughput	Aggregate number of records inserted into a database from any Store and Forward engine per second.
Total Quarantine	Number of quarantined items for all Store and Forward engines.
Total Dropped	Number of records dropped from the Store and Forward engines. A record is considered dropped if it can not be added to one of the buffers, (i.e., when a buffer is full and the Store and Forward engine can no longer accept new records).
Store and Forward Engines	
Name	Name of the Store and Forward engine.
Store Throughput	Number of records that go through the Store and Forward Engine per second.
Forward Throughput	Number of records to be forwarded on to the database per second.
Quarantined	Is data that has errored-out multiple times during attempts to forward it, or data that could not be stored because of some configuration issues.
Activity	Current state of the Store and Forward engine.
Actions	By clicking Details, shows additional information about Store and Forward engines.



Store & Forward Details

Clicking the **Details** button brings up a new window that will show even more details about the records in the selected Store and Forward Engine. Here, we can see a count of the number of records in the memory buffer and local cache, as well as the number of quarantined records. The quarantined items at the bottom of the Details page will have some buttons that allow you to [control the data that is in the quarantine](#). The quarantined item can be retried, where it will be thrown back through the Store and Forward system to see if it will go through properly, assuming the original reason why it was quarantined has been fixed. It can also be deleted so that it is no longer taking up space in the Store and Forward system, or exported to your local machine where you can save it to try again later. You can then import the file back from the same page when you resolved the issue that caused the data to be quarantined in the first place.

Attribute	Description
Memory Buffer	Number of records entering the Memory Buffer per second. The progress bar shows the percent of the buffer being utilized, along with the current and max number of records.
Local Cache	Number of records entering the Local Cache per second. An "Idle" state means the engine is able to successfully store all records into the database before the Write Size or Write Time values have been reached. The progress bar shows the percent of the buffer being utilized, along with the current and max number of records. <p>Note: On Edge Gateways, the meter for the local cache represents the rows stored within the Edge Historian database. This value will update when pruning is triggered by Edge's one week storage policy or when storing 10 million data points.</p>
Database Storage	Displays the number of records pushed from either buffer to the database per second.
Quarantine d Items	Shows a list of quarantine items and allows you to choose the quarantined file and import it.
ID	Identification number of the quarantine item.
Count	Number of occurrences for the quarantined item.
Description	Description of where the quarantine item originated from.
Reason	Explanation why the record was placed into quarantine.

Actions | Provides and opportunity to retry, delete, or export the items from quarantine.

All Status > Connections > **Store and Forward**

Configuration

Store Details

Memory Buffer	Local Cache	Database Storage
6 /sec	1 /sec	6 /sec
2% (5/250)	0% (0/25000)	

Quarantined Items

Import quarantine file

Retry All

Delete All

« ‹ 1 of 1 › »

ID	Count	Description	Reason	Actions
1	10	SQLTag History Data	This data sink does not accept data of the given type.	Retry Delete Export

« ‹ 1 of 1 › »

Related Topics ...

- [Using Store and Forward](#)
- [Controlling Quarantine Data](#)

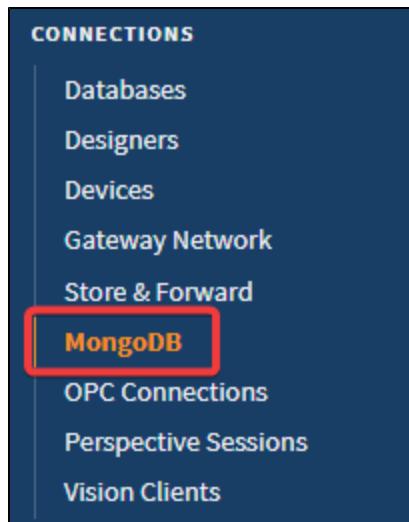
Connections - MongoDB

MongoDB

The MongoDB page shows a list of configured MongoDB connections. To view this page, go to your **Gateway webpage > Status > Connections > MongoDB**. You can filter by name to search for a specific MongoDB connection, and change how many connections are displayed on this page.

Note:

The MongoDB option will only show up if the MongoDB Cloud Connector Module is installed.



On this page ...

- [MongoDB](#)
- [MongoDB Page](#)

MongoDB Page

The main MongoDB page contains useful information about your configured MongoDB connections.

Attribute	Description
Name	Name of the MongoDB connection,
Status	Current status of the MongoDB connection.
Connections	Current number of connections between Ignition and the specified MongoDB database.
Throughput	Number of requests per second.

» Status > Connections > **MongoDB**

License Incomplete ? 1:45:36 [View Modules](#)

« < 1 of 1 > »

Filter type to filter **View** 20 ▾

Name	Status	Connections	Throughput
MongoDBTEST	✓ Valid	0 / 100	0.0 request/sec

« < 1 of 1 > »

Connections - OPC Connections

OPC Connections

The OPC Connections page displays all currently configured OPC (both UA and DA) connections.



If you are using an OPC UA connection, you should be using the [OPC UA Module](#). If you are using the DA connection, you should be using the [OPC COM Module](#).



On this page ...

- [OPC Connections](#)
 - [OPC Connections Page](#)
 - [Nodes](#)

OPC Connections Page

Here on the main OPC Connections page, we can see a list of all current OPC connections, as well as their status. If any are faulted, you can click the red faulted status to get an error message popup with a full description of the error.

Attribute	Description
Connected Servers	Displays the list of OPC servers out of configured servers, and their status.
Name	OPC server name.
Filter	Search criteria to filter for specific server names.
Type	OPC server type - UA or DA
Uptime	Total time OPC server is connected.
Status	Current status of OPC server.
Diagnostics	Displays diagnostic information for any connected OPC UA server. <ul style="list-style-type: none">• Server - Shows server diagnostics.• Client - Shows client connection subscription diagnostics.

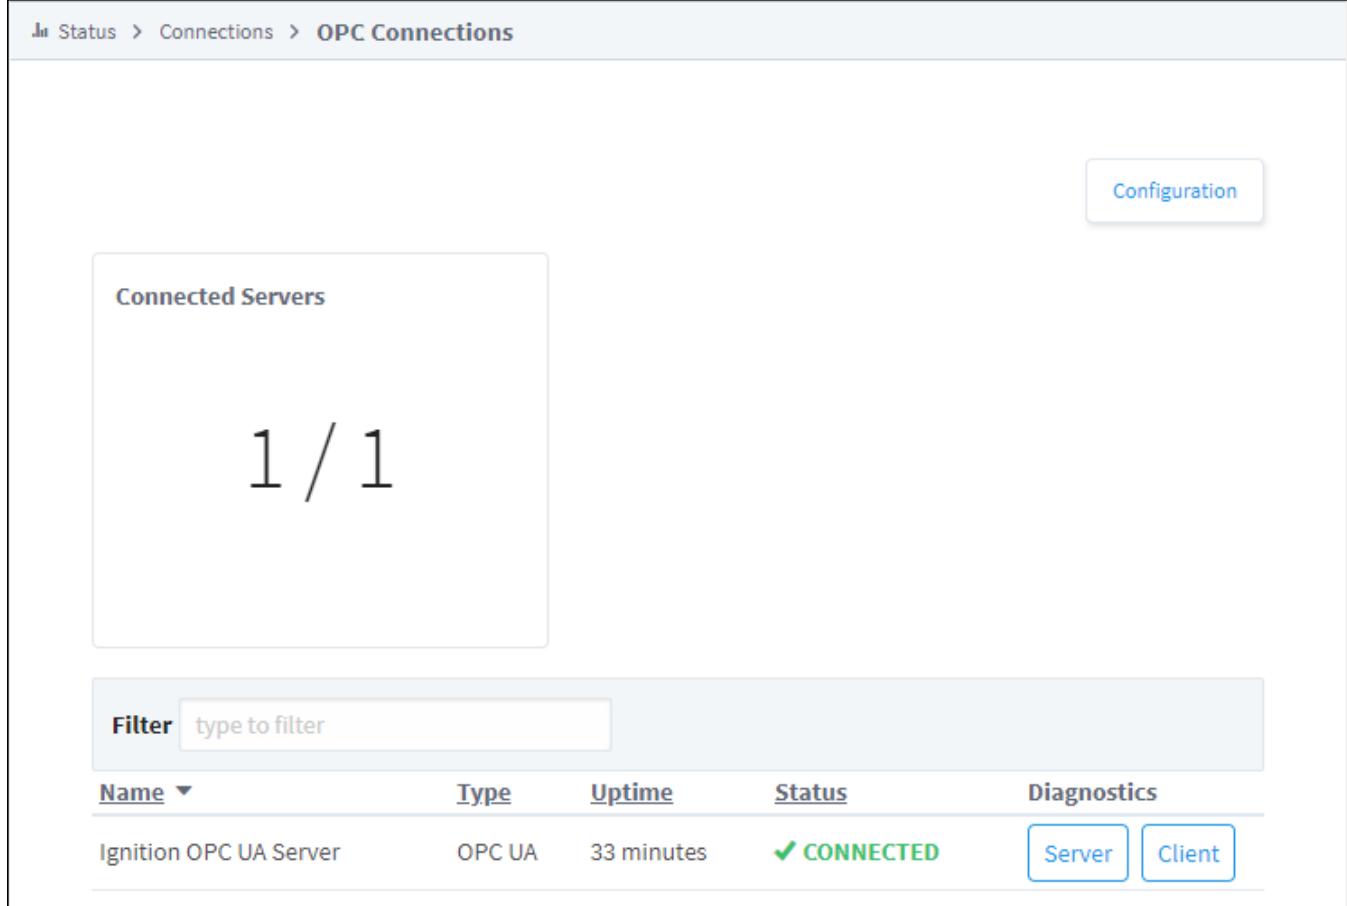
» Status > Connections > OPC Connections

Configuration

Connected Servers
1 / 1

Filter type to filter

Name ▾	Type	Uptime	Status	Diagnostics
Ignition OPC UA Server	OPC UA	33 minutes	✓ CONNECTED	Server Client



Server

Click the **Server** button to get information on the server. By default, diagnostics are set to off since they can generate a considerable amount of network overhead and impact performance. To turn on diagnostics, click **Enable Diagnostics**. Not all OPC UA servers support diagnostics.

[Configuration](#)

Ignition OPC UA Server

[← Return to Servers](#)[Server Information](#)[Session Diagnostics](#)

Server Status

Parameter ▾	Value
currentTime	2020-07-10T10:37:36.900-07:00
startTime	2020-07-09T09:35:29.474-07:00
state	CONNECTED

Build Info

Parameter ▾	Value
buildDate	2020-07-09T09:35:28.127-07:00
buildNumber	dev
manufacturerName	Inductive Automation
productUri	urn:inductiveautomation:ignition:opcua:server
softwareVersion	dev

Server Diagnostics

[Enable Diagnostics](#)

This server does not have diagnostics enabled.

Clients

By clicking on the **Client** button, it brings up the subscription information for that particular server. It will list out all of the subscriptions to that server with the publishing rates, as well as the number of items within that subscription.

Attribute	Description
Filter	Search criteria to filter for specific subscriptions.
Refresh	Refreshes subscription data.
Name	Displays the subscription name.
Rate	The rate defined in the Tag Group, used as the requested sampling interval for monitored Items belonging to the corresponding subscription.
Request Publishing Interval	The rate a subscription will report accumulated change notifications at. The interval is derived from Rate, unless explicitly defined.
Revised Publishing Interval	The server's revised rate for accumulated change notifications.
Tag Count	The number of Tags currently subscribed to that Tag Group.

! Status > Connections > OPC Connections > Ignition OPC UA Server - Subscriptions

Configuration

Ignition OPC UA Server

Name	Rate	Request Publishing Interval	Revised Publishing Interval	Tag Count	
tag-group-exporttags_default	10000	1000	1000	3	<button>Nodes</button>
tag-group-default_time driven	60000	1000	1000	2	<button>Nodes</button>
tag-group-default_driven machine state	10000	1000	1000	2	<button>Nodes</button>
tag-group-default_direct 5 seconds	5000	1000	1000	4	<button>Nodes</button>
tag-group-default_direct	1000	500	500	2	<button>Nodes</button>
tag-group-default_default	10000	1000	1000	188	<button>Nodes</button>

Nodes

Clicking the **Nodes** button for one of the subscriptions will bring up the list of subscribed OPC items.

Attribute	Description
Filter	Search criteria to filter for a specific node.
Refresh	Refreshes node data.
Node ID	The OPC item path.
Requested Sampling Interval	The rate the underlying tag/node will be polled at.
Revised Sampling Interval	The revised rate at which the underlying tag/node will be polled at.
Requested Queue Size	Determines how many data points can be stored and transferred to the client once the sampling rate elapses when the publishing rate is slower than the publishing rate.
Status Code	Status of the node.

Configuration

Ignition OPC UA Server > tag-group-default_driven machine state-unleased

Node ID ▾	Requested Sampling Interval	Revised Sampling Interval	Requested Queue Size	Revised Queue Size	Status Code
ns=1;s=[Generic]_Meta:Sine/Sine9	10000	10000	1	1	Good
ns=1;s=[Generic]_Meta:Sine/Sine8	10000	10000	1	1	Good
ns=1;s=[Generic]_Meta:Sine/Sine0	10000	10000	1	1	Good
ns=1;s=[Generic]_Meta:Ramp/Ramp9	10000	10000	1	1	Good
ns=1;s=[Generic]_Meta:Ramp/Ramp8	10000	10000	1	1	Good
ns=1;s=[Generic]_Meta:Ramp/Ramp7	10000	10000	1	1	Good

Related Topics ...

- [OPC UA](#)
- [OPC UA Client Connection Settings](#)

Connections - SECS/GEM Equipment

SECS/GEM Equipment

The SECSGEM Equipment page displays a list of all equipment connections, their status, as well as the number of sent requests and received messages. If any piece of equipment becomes faulted, it will show a status of "Not Connected." If you click the **Details** button, it opens a new page with some Connection Stats. Note: the SECSGEM Module is not standard and will be missing for most installs of Ignition. The SECSGEM Equipment Connection will only be displayed in the Status section of the Gateway webpage under Connections when the module is installed.



On this page ...

- SECS/GEM Equipment
 - SECS/GEM Equipment Page
 - SECS/GEM Equipment Details

SECS/GEM Equipment Page

The main SECS/GEM Equipment page lists out all equipment connections and displays the number of messages sent between them and Ignition.

Attribute	Description
Connections	The total number of active SECS/GEM Equipment connections.
Aggregate Throughput	The messages per second the system is sending and receiving from all devices at that time.
Name	The name of the equipment connection in the list.
Sent Messages	The number of messages sent to the equipment.
Received Messages	The number of messages received from the equipment.
Status	The status of the equipment.

A screenshot of the SECSGEM Equipment page. At the top, there is a breadcrumb navigation: Status > Connections > SECSGEM Equipment. On the right side, there is a 'Configuration' button. Below the navigation, there is a section titled 'Equipment Stats' with two cards: 'Connections' (value: 1) and 'Aggregate Throughput' (value: 0.2 per second). At the bottom, there is a table with columns: Name, Sent Messages, Received Messages, and Status. One row in the table shows 'SimEquipOne' with '154' in the Sent Messages column, '77' in the Received Messages column, and a green checkmark in the Status column indicating 'Communicating'. There is also a 'Details' button next to the status column.

SECS/GEM Equipment Details

Clicking the **Details** button to the right of a piece of equipment will take you to a page that shows more detailed information for that particular piece of equipment.

Attribute	Description
Sent Messages	
Throughput	The messages per second the system is currently sending.
Average	The average messages sent per second from when the equipment was first enabled.
Total	The total messages sent.
Received Messages	
Throughput	The messages per second the system is currently receiving.
Average	The average messages received per second from when the equipment was first enabled.
Total	The total messages received.

Status > Connections > SECS/GEM Equipment

Configuration

Connection Stats

Sent

Last to: S1F14

Throughput	Average	Total
0.1 per second	0.0 per second	166

Received

Last from: S1F13

Throughput	Average	Total
0.1 per second	0.0 per second	83

This screenshot shows the 'Connection Stats' section of the SECS/GEM Equipment Details page. It displays two sets of metrics: 'Sent' and 'Received'. Under 'Sent', it shows 'Last to: S1F14' with three cards: 'Throughput' (0.1 per second), 'Average' (0.0 per second), and 'Total' (166). Under 'Received', it shows 'Last from: S1F13' with three cards: 'Throughput' (0.1 per second), 'Average' (0.0 per second), and 'Total' (83). The page also includes a 'Configuration' button in the top right corner.

Related Topics ...

- SECS/GEM

Connections - Perspective Sessions

Perspective Sessions

Much like the Designers page, the Perspective Sessions page shows a ton of information regarding currently open Sessions. All the open Perspective Sessions are displayed on the page along with some basic information about each session, such as what user is logged into each session and the project they are currently working on, and their IP address. Each session has a Details button that allows you navigate all the elements of a session: pages, view instances, and components.



On this page ...

- Perspective Sessions
 - Perspective Sessions Page
 - Perspective Session Details
 - View Instances
 - Components

Perspective Sessions Page

The main Perspective Sessions page displays a list of all currently running Designer and Client Sessions to show how many are open, and the name of the users that are currently using them. By clicking on the **Details** button to the right of a session will display even more information about that particular session, including any Log Activity pertaining to errors recorded while the session was active.

Attributes	Description
Sessions Details	
Total Sessions	The number of currently active Designer and Perspective Sessions.
Device	Type of device running the Session.
User	The user logged into the Session.
Project	The name of the project open in the Session.
Uptime	The total uptime of the Session.
Last Comm.	Date and time recorded for the last communication with the Session.
Address	The IP Address and computer name where the Session is launched.
Session Scope	Indicates what browser opened the Session, and the Perspective icon indicates the Designer is open.
Actions	<p>The following feature is new in Ignition version 8.1.23 Click here to check out the other new features</p> <p>More > Terminate: Allows users to terminate the session. Note the Terminate option is not available for any Designer Perspective sessions and is only visible for other session scopes if the user has Config-level permissions.</p> <p>Details: Accesses a new Perspective Sessions page that displays performance stats, pages information, session details, and log activity.</p>

The screenshot shows the Ignition software interface. The left sidebar has a dark blue background with white text and icons. It includes sections for Home, Status (which is highlighted in orange), and Config. Under SYSTEMS, there are links for Overview, Performance, Alarm Pipelines, Gateway Scripts, Modules, Redundancy, Reports, SFCs, Voice Alarming, Tags, and Transaction Groups. Under CONNECTIONS, there are links for Databases, Designers, Devices, Gateway Network, Store & Forward, OPC Connections, Perspective Sessions (which is also highlighted in orange), and Vision Clients.

The main content area has a light gray header with the breadcrumb navigation: Status > Connections > Perspective Sessions. Below this is a section titled "Session Stats" with a box containing the number "2".

Under "Session Details", there is a table with two rows of data:

Device	User	Authorized?	Project	Uptime	Last Comm.	Address	Session Scope	Actions
Windows 10	Unauthenticated	Yes	samplequickstart	a minute	a few seconds		Browser, Chrome	<button>More</button> <button>Details</button>
Unknown	admin	Yes	samplequickstart	a minute	a few seconds		Designer	<button>Terminate</button> <button>Details</button>

Perspective Session Details

The Performance Page within the session displays more in depth information such as the Page Id and how many views are on each Page. Hit the **Details** button to get more detailed information about each page in the session.

Attributes	Description
Performance	
Total Views	The number of currently active pages.
Total Bytes Sent	Displays total number of bytes sent out.
Pages	
Id	Page Id
Views	Displays total number of views on a page.
Actions	Displays the details of the view instances in the session.
Log Activity	
Min level	Dropdown menu with options Info, Debug, and Trace.
Live Values	Toggle switch to turn live values on or off.
Merge to Logs	Merge the settings of the current view with the main Diagnostics Log Viewer.
Logger	Name of the logger that describes the context of the message.
Time	Time of log.
Message	Message for the log.
(Log Properties)	Displays log properties of the logged event.

Performance

Total Views

2

Total Bytes Sent

60 MB

Pages

<u>Id</u> ▾	<u>Views</u>	<u>Actions</u>
2c3e945	2	Details

Details

Session Details

User	Unauthenticated
Project	Perspective
Device	Windows NT 10.0; Win64; x64
Address	0:0:0:0:0:0:1
Uptime	20 hours
Last Comm.	a few seconds
Session Scope	

Log Activity



Min level	INFO	▼	i	Live Values <input checked="" type="checkbox"/> ON	Merge to Logs
Logger	Time		Message		
I DesignSession	29Mar2019 10:37:26		Socket connected to session. pageld=nav toolbar		
I DesignSession	29Mar2019 10:37:21		Socket connected to session. pageld>New View		
I ClientSession	29Mar2019 09:45:04		WebSocket disconnected from session.		

View Instances

The View Instances page shows the number of View Instances on a Page and the number of Components in each view instance. Press the **Details** button next to each View Instance to get even more information about the type of components used in each view.

Attributes	Description
Sessions Details	

InstanceID	Displays View Instances for each view.
Components	Total number of components used in the specified view.
Actions	By clicking on the Details button shows more information for each view instance.

.. Status > Connections > **Perspective Sessions**

View Instances

Instance ID ▾	Components	Actions
nav topbar@T[0]	7	<button>Details</button>
New View@C	7	<button>Details</button>

Components

The Components page lists all the components used in a View Instance of Page along with the following information about each component.

Attributes	Description
Name	Name of the component.
AddressPath	Address of the component.
Bindings	Displays the number of bindings for each component.
Children	Displays the number of children for each component.
Properties	Displays the number of properties for each component.
PropertyChangeScripts	Displays the number of property change scripts for each component
Actions	Displays the number of actions for each component.

.. Status > Connections > **Perspective Sessions**

Components

Name ▾	Address Path	Bindings	Children	Properties	Property Change Scripts	Actions
LedDisplay_2	0:15	1	0	4	0	0
LedDisplay_1	0:9	1	0	4	0	0
LedDisplay_0	0:8	1	0	4	0	0
LedDisplay	0:7	1	0	4	0	0
Label_3	0:14	0	0	4	0	0
Label_2	0:10	0	0	4	0	0
Label_1_0	0:6	0	0	4	0	0
Label_1	0:5	0	0	4	0	0
Label_0	0:4	0	0	4	0	0
Label	0:3	0	0	4	0	0
CylindricalTank_3	0:2	1	0	4	0	0
CylindricalTank_0	0:1	1	0	4	0	0
CylindricalTank	0:0	1	0	4	0	0
Button_0_0	0:13	0	0	4	0	0
Button_0	0:12	0	0	4	0	1
Button	0:11	0	0	4	0	1

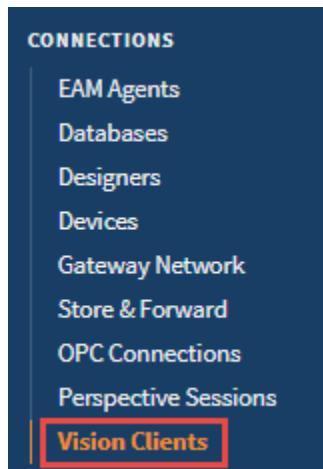
Related Topics ...

- [Launching a Perspective Session](#)

Connections - Vision Clients

Vision Clients

Much like the Designers page, the Vision Clients page shows information regarding currently opened Clients. The Clients are listed and show some basic information such as the address of the Client. From here, the Client session can be terminated by selecting the **More** button and hitting **Terminate**, or select the **Details** option to see more details about the Client session such as the number of Tags that the session is currently subscribed to, as well as a log of errors that may have happened with that Client.



On this page ...

- Vision Clients
 - Vision Clients Page
 - Client Session Details

Vision Clients Page

The main Vision Clients page displays a list of all currently running Clients which show how many clients are open and what users are currently using them.

Attributes	Description
Clients Stats and Details	
Active Clients	The number of currently active Clients.
Requests /sec	The number of requests per second coming from all Clients.
Id	The Client Id.
User	The user logged in on the Client session.
Project	The name of the project open in the Client session.
Uptime	The total uptime of the Client session.
Activity	Denotes whether the Client session is currently connected and Active, or Dormant. A Dormant status means the Gateway hasn't received any activity from the Vision Client, but the Client hasn't been logged out yet. Dormant statuses can also be caused by a user opening multiple Clients without closing or logging out of previously opened ones. Note: Dormant Clients should automatically clear after a period of time, but they can also be terminated manually or resolved by logging out of the Vision Client, either by the user or through a Client Event Timer script. You can utilize system.util.getInactivitySeconds() and system.security.logout() for the Client Event Timer script.
Address	The IP Address and computer name where the Client is launched.
Memory	The Client's current memory usage.
Actions	The option to terminate the Client session, or see more details about the session.

Status > Connections > Vision Clients

Client Stats

Active Clients

1

Requests / sec

5.1

1 items << < 1 of 1 > >>

Client Details

Filter type to filter						View 10	
<u>Id</u>	<u>User</u>	<u>Project</u>	<u>Uptime</u>	<u>Activity</u>	<u>Address</u>	<u>Memory</u>	<u>Actions</u>
38638402	admin	NewProject	18 minutes	✓ Active	ws3 (127.0.0.1)	68mb	More Details

1 items << < 1 of 1 > >>

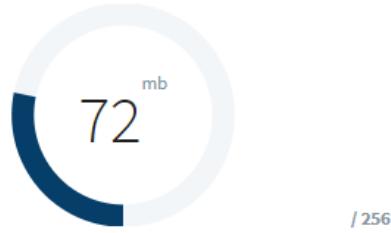
Client Session Details

Clicking on the **Details** button for a Client session will take you to a page that displays more in depth information for that particular Client session. It also gives a logger at the bottom where errors coming from that particular Client can be seen.

Attributes	Description
Client Performance and Details	
Memory	The Client's current memory use.
Subscriptions	The number of Tags the Client is currently subscribed to.
User	The user logged in on the Client session
Project	The name of the project open in the Client session.
Address	The IP Address and computer name where the Client is launched.
Uptime	The total uptime of the Client session.
Last Comm	The last time the Gateway communicated with the Client.
Client JVM Version	The Java version that the Client is currently running on.

Performance

Memory



Subscriptions

8 tags

Details

Session Details

User	admin
Project	NewProject
Address	ws3 (127.0.0.1)
Uptime	34 minutes
Last Comm	485ms
Client JVM Version	11.0.2

Related Topics ...

- [Vision Client Launcher](#)

Diagnostics - Execution

The Execution page reports the status of all tasks that your [Gateway](#) runs on a schedule. Here you can find helpful information such as the duration and execution time of an [alarm](#) journal update or the average time it takes your [Gateway](#) to execute a [Tag Group](#).

Attributes	Description
Throughput	Number of executions that come through per second.
Total Executions	Total number of times the task executed.
Delay	Amount of time waiting to execute the task.
Avg. Duration	Average time to run the task.
Last Duration	Amount of time to run the last task.

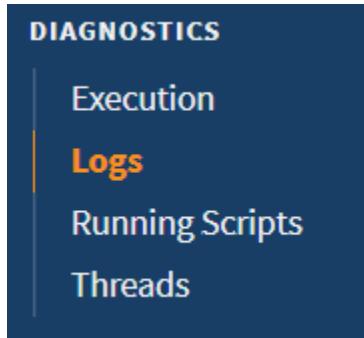
Shared Execution Engine						
Owner / Tasks	Throughput	Total Executions	Delay	Avg. Duration	Last Duration	
databaseconnectionmanager						
txtimoutdaemon	0.500 exec/sec	3,639	2 sec	0 millis	0 millis	
eam						
agentupdate	0.999 exec/sec	7,258	1 sec	0 millis	0 millis	
emailackmanager						
orphanedacknowledgementcleanup	idle	122	1 min	0 millis	0 millis	
gateway						
license-checker	idle	114	1 min	4 sec	4 sec	
logfile cleanup	idle	1	2 hours	0 millis	8 millis	
performance monitor	0.200 exec/sec	1,451	5 sec	4.7 millis	8 millis	
gatewayareanetworkconnectionmanager						
connection monitor	0.200 exec/sec	1,453	5 sec	0.1 millis	0 millis	
gatewaysystemtags						
basictags	0.999 exec/sec	7,274	1 sec	0 millis	0 millis	

Related Topics ...

- [Diagnostics - Logs](#)

Diagnostics - Logs

One of the most important troubleshooting tools of the [Gateway Webpage](#) is the Logs page. This console shows errors caused by [Gateway](#) events including things like Database or Device connections, Authentication profiles, Alarm Journals and Pipelines, and anything else that is [Gateway](#) scoped. You can find logs in the Status tab of the [Gateway Webpage](#) under **Diagnostics > Logs**.



On this page ...

- [Changing Logging Levels](#)
- [Downloading the Logs](#)
- [Printing to the Logs](#)
- [Mapped Diagnostic Context Keys](#)

The [Gateway](#) Logs also provide a wealth of information about the running state of the [Gateway](#). To learn more, refer to the section on [Troubleshooting the Gateway](#). The Logs page is where the [Gateway](#) Console is located that allows you to see a live flow of log events in the system.

Attributes	Description
Logger	Name of the logger that describes the context of the message.
Time	Time of log.
Message	Message for the log.
Filter	Search criteria to filter for specific tasks or events. Click the Calendar icon to set a filter date and time range. Options are Last Month, Last Week, Last 24 Hours, or Custom.
Min Level	Dropdown menu with options: All, Trace, Debug, Info, Warn and Error.
Live Values	Toggle switch to turn live values on or off.

Status > Diagnostics > Logs			
37978 items < < 1 of 380 > >			
Logger	Time	Message	
I ModuleManager	26Mar2019 10:46:23	Uninstalling module "secsgem"	
I ModuleManager	26Mar2019 10:30:52	Starting up module 'secsgem' (v2.9.10 (b2018112813))...	
W ModuleManager	26Mar2019 10:30:51	Module "SECS/GEM" requires Ignition 7.8.0 (b0) and is not compatible with Ignition 8.0.0-beta0 (b2019032202)	
I ModuleManager	26Mar2019 10:30:51	Installing module: "secsgem"	
I DesignSession	26Mar2019 09:15:58	Socket connected to session. pageld=Test/FirstView	
I DesignSession	26Mar2019 09:15:20	Socket connected to session. pageld=session	
I InternalDatabase	26Mar2019 09:08:07	Created auto-backup of internal database "config.idb" in 1 seconds	
I InternalDatabase	26Mar2019 09:08:06	Creating auto-backup of internal database "config.idb"...	
W LicenseManager	26Mar2019 09:04:20	Could not read license from HASP key. HASP login() failed. Last error = 14. Error executing global timer script: Perspective/Print to Diagnostic Logs @1,000ms . Repeat errors of this type will be logged as 'debug' messages.	
E TimerScriptTask	26Mar2019 09:03:37		
W SubscriptionManager	26Mar2019 09:03:37	[Ignition-ignition8-ubuntu-64bit] Subscription path model set detected to be out of sync. Will re-synchronize.	
I IgnitionGateway	26Mar2019 09:03:36	Ignition[state=STARTING] ContextState = RUNNING	
I IgnitionGateway	26Mar2019 09:03:36	Gateway started in 31 seconds.	

Changing Logging Levels

The logs can be filtered by using a search term or by date.

1. To change logging levels, go to the Gateway's **Status > Diagnostics > Logs** page and click on the icon.
2. On the Log Configuration popup that will open, search for the logger name in the Filter box, and select the desired logging level from the dropdown to the right of the logger name

Downloading the Logs

Gateway logs can also be exported using the **Download Logs** icon the right.

Printing to the Logs

You can print to Gateway logs by using the `system.util.getLogger()` function. Below is a simple example, you can look in the [appendix](#) for more options.

```
logger = system.util.getLogger("My Logger Name")
logger.warn("My Warning Message")
```

Mapped Diagnostic Context Keys

Mapped Diagnostic Context Keys (MDC Keys) allow you to specify a specific context, such as a particular project, and then set a logging level for it. This will set all loggers that pertain to the specified project to the logging level. This is useful to help diagnose an issue with a specific system within

the Gateway. To use MDC Keys, click the icon on the Logs page and navigate to the Context tab. Here, you can choose specific Key-Value pairs that match a particular system. Selecting the Key text field will bring up a list of possible keys in your system. Once a Key has been selected, selecting the Value text field will bring up a list of possible values that relate to the selected Key.

Log Configuration

Loggers **Context** X

Key	Value	Set Level	INFO	Add Level
store-forward-name				
connection-name				
project-name				
alarm-notification-profile	MySQL		DEBUG	
device-name				
gw-name				
module-name	Symbol Factory		INFO	
project				
route-path				
session-project				

Filter by Mapped Diagnostic Context Key

The Logs can also be filtered to show only logs that pertain to an MDC Key. This can be useful when altering the logging level of one or more MDC Keys. To set an MDC filter on the logs, simply click the **Add Mapped Diagnostic Context Filter** icon on the Logs page to open a window where MDC Key filters can be set. The filter can be specified for all values in a Key by not specifying a Value.

MDC Filter Configuration

X

Key	Value	Add Filter
enter key	enter values	

Key **Values** **Action**

No Log Filters Configured

Merging to Logs

Many of the other pages in the Status tab also contain a logger that will just show logs that are relevant for that particular section. These log views all contain a button on the right side called **Merge to Logs**. This button allows the current filter to be applied to the main logs page, allowing you to still see only the logs for that section but use all of the tools available.

◀ ▶ ⏪ ⏩

Live Values ON	Merge to Logs [x]
---	--

[Related Topics ...](#)

- [Diagnostics - Execution](#)
- [Diagnostics - Running Scripts](#)
- [Diagnostics - Threads](#)

[In This Section ...](#)

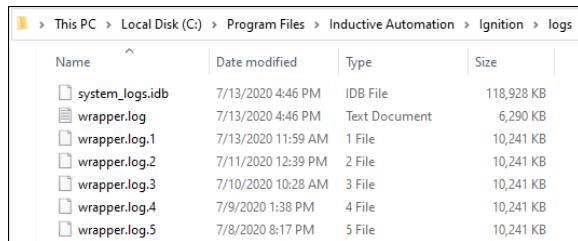
Wrapper Logs

Wrapper logs are a set of plain text log files located inside the installation directory on the computer where Ignition Gateway is installed. The exact path may vary depending on the installation directory, but the typical path to the location of the wrapper logs is C:\Program Files\Inductive Automation\Ignition\logs\ on a Windows operating system and /usr/local/bin/ignition on Linux installations.

Wrapper Log Location

There are typically six wrapper log files. The file containing the most recent information is the one called wrapper.log. Older information is contained in the wrapper.log.1, wrapper.log.2, wrapper.log.3, etc., the higher the number, the older the information. A very recent installation of Ignition may have fewer than six wrapper log files.

Some operating system settings may cause the file extensions to be hidden. In that case, check your operating system's documentation on how to show file extensions, or simply look for the files with the name starting with "wrapper".



A screenshot of a Windows File Explorer window. The path shown is This PC > Local Disk (C:) > Program Files > Inductive Automation > Ignition > logs. Inside the logs folder, there are several files listed in a table:

Name	Date modified	Type	Size
system_logs.idb	7/13/2020 4:46 PM	IDB File	118,928 KB
wrapper.log	7/13/2020 4:46 PM	Text Document	6,290 KB
wrapper.log.1	7/13/2020 11:59 AM	1 File	10,241 KB
wrapper.log.2	7/11/2020 12:39 PM	2 File	10,241 KB
wrapper.log.3	7/10/2020 10:28 AM	3 File	10,241 KB
wrapper.log.4	7/9/2020 1:38 PM	4 File	10,241 KB
wrapper.log.5	7/8/2020 8:17 PM	5 File	10,241 KB

On this page ...

- [Wrapper Log Location](#)
- [Wrapper Log Content](#)

Wrapper Log Content

Wrapper logs are similar to Gateway logs in that they contain information about the operation of the Ignition Gateway and its sub-systems, as well as the errors generated by Gateway scoped resources. However, there are some important differences between the Gateway log and the wrapper logs:

- The information in the wrapper logs persists across Gateway restarts, while Gateway logs are cleared by Gateway restart
- The wrapper logs contain information about issues during the Gateway startup, including any issues that may have prevented the Gateway from starting. The most recent wrapper log file is the first place to look for the cause if the Gateway has failed to start, or if the Gateway stopped or restarted unexpectedly.
- The wrapper logs (but not the Gateway log) contain the output of any print statements in Gateway-scoped scripts.
- Unlike the Gateway log, the wrapper log files can only be obtained from their directory on the computer where Ignition Gateway is installed. There is no way to view them via the Gateway's Web interface.

Diagnostics - Running Scripts

The Running Scripts page shows all actively running Gateway scripts, as well as providing a way to terminate any running script. In addition, the Vision client and Designer consoles have a Running Scripts tab, which also lists running scripts and provides a way to terminate them.

Attributes	Description
Thread Id	Thread Id number.
Description	Name and description of the script.
Execution Start	Time script started running.
Elapsed Time	Amount of time the script has been running.
Actions	Lists actions that can be performed against the running script, such as canceling it.

The screenshot shows a table titled 'Currently Running Scripts' with one row of data. The columns are Thread Id, Description, Execution Start, Elapsed Time, and Actions. The 'Actions' column contains a single button labeled 'Cancel'. Navigation buttons for the table are visible at the top and bottom right.

Thread Id	Description	Execution Start	Elapsed Time	Actions
3188	TimerScript - project:Production GatewayLogger @10,000ms	26Mar2020 10:56:07	24 seconds	<button>Cancel</button>

The following feature is new in Ignition version **8.1.29**
[Click here](#) to check out the other new features

Note: Users who have Status webpage access but not Config webpage access can view Running Scripts without the ability to terminate them.

Related Topics ...

- [Diagnostics - Execution](#)
- [Diagnostics - Logs](#)
- [Diagnostics - Running Threads](#)

Diagnostics - Threads

The Threads page displays a snapshot of information about Gateway's execution threads. This information may be important when troubleshooting certain types of issues. If you are working with Support, you may be asked for a thread dump.

Their state and CPU usages are displayed to easily find problem threads, as well as a chart of what systems are using the threads and the ability to filter the threads based on a keyword. Each thread can be expanded to give details on what it is currently doing, and that individual thread process can be copied to the clipboard.



On this page ...

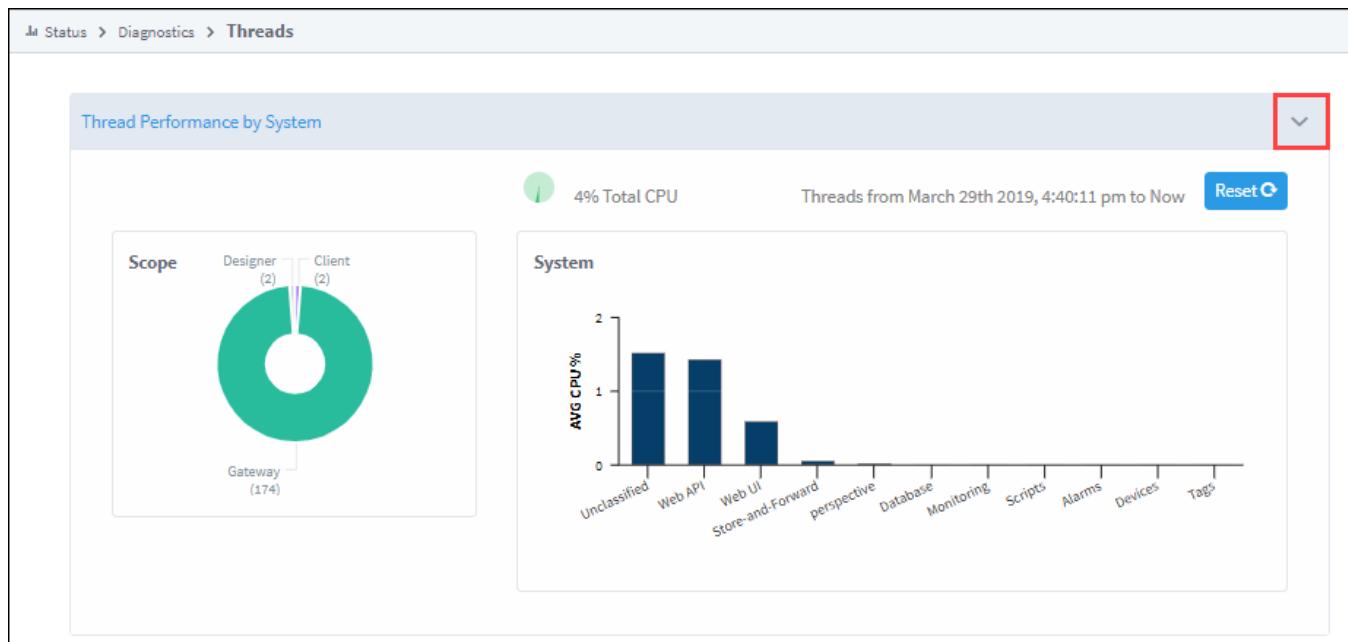
- [Thread Performance by System](#)
- [Individual Threads](#)

Thread Performance by System

The upper section of the Diagnostics - Threads page depicts graphical representation of the current threads including percentage of total CPU used, overall scope, and average CPU usage per system.

Click the **Expand** icon to see the display.

Click the **Reset** icon to reset the information.



Individual Threads

The lower portion of the Diagnostics - Threads page displays statistics for individual threads.

To download a thread dump, click the **Download**  icon on the right of the page. In order to be useful, a thread dump must be taken at the time the issue is observed. Unlike logs, which maintain a record of past messages, a thread dump is a time-specific snapshot and will not contain a record of past states. To see more information on thread dumps, including other methods of obtaining them such as [system functions](#) or the command line, click [here](#).

Attributes	Description
Thread Name	Name of the thread.
Filter	Search criteria to filter for specific threads.
CPU(%)	Percentage of CPU usage.
State	Current state / status of the system.
System	Search criteria to filter for specific tasks or events.
Actions	View thread details.
Live Values	Toggle switch to turn live values on or off.

Status > Diagnostics > **Threads**

Thread Performance by System

Thread Name	CPU (%)	State	System	Actions
platform-executor-266, id=706212	4.14	TIMED_WAITING	Unclassified	Show 
ws-ping-executor, id=142	0.94	TIMED_WAITING	Unclassified	Show 
WebSocketClient@123122752-706903, id=706903	0.08	TIMED_WAITING	Unclassified	Show 
webserver-706021, id=706021	0.08	RUNNABLE	Web API	Show 
gateway-shared-exec-engine-4, id=46	0.08	WAITING	Unclassified	Show 
gateway-logging-sqlite-appender, id=20	0.08	TIMED_WAITING	Unclassified	Show 
ws-ping-executor, id=141	0.00	TIMED_WAITING	Unclassified	Show 
ws-ping-executor, id=139	0.00	TIMED_WAITING	Unclassified	Show 
WrapperSimpleAppMain, id=19	0.00	WAITING	Unclassified	Show 
Wrapper-Control-Event-Monitor, id=14	0.00	TIMED_WAITING	Unclassified	Show 

Click the **Show**  icon to display additional details for a single thread. You can click the **Clipboard**  icon to copy that thread to the clipboard.

Thread Performance by System				
<input type="text"/> Filter type to filter				
Thread Name	CPU (%)	State	System	Actions
platform-executor-267, id=706857	5.15	TIMED_WAITING	Unclassified	Hide Copy
Thread [platform-executor-267] id=706857, (TIMED_WAITING for java.util.concurrent.SynchronousQueue\$TransferStack@7a8acf5) java.base@11.0.6/jdk.internal.misc.Unsafe.park(Native Method) java.base@11.0.6/java.util.concurrent.locks.LockSupport.parkNanos(Unknown Source) java.base@11.0.6/java.util.concurrent.SynchronousQueue\$TransferStack.awaitFill(Unknown Source) java.base@11.0.6/java.util.concurrent.SynchronousQueue\$TransferStack.transfer(Unknown Source) java.base@11.0.6/java.util.concurrent.SynchronousQueue.poll(Unknown Source) java.base@11.0.6/java.util.concurrent.ThreadPoolExecutor.getTask(Unknown Source) java.base@11.0.6/java.util.concurrent.ThreadPoolExecutor.runWorker(Unknown Source) java.base@11.0.6/java.util.concurrent.ThreadPoolExecutor\$Worker.run(Unknown Source) java.base@11.0.6/java.lang.Thread.run(Unknown Source)				
ws-ping-executor, id=142	0.94	TIMED_WAITING	Unclassified	Show Copy

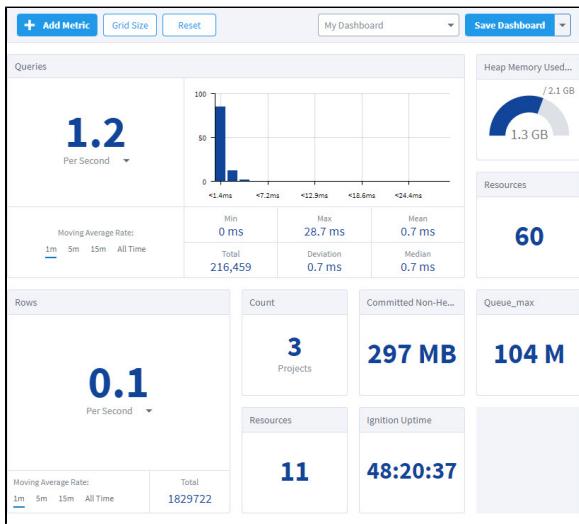
Related Topics ...

- [Diagnostics - Execution](#)
- [Diagnostics - Logs](#)
- [Diagnostics - Running Scripts](#)

Diagnostics - Metrics Dashboard

The following feature is new in Ignition version 8.1.13
[Click here](#) to check out the other new features

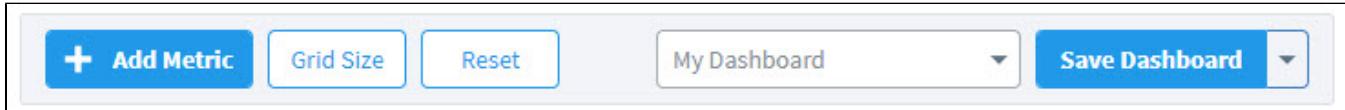
The Metrics Dashboard allows users to visualize a variety of metrics on their own customizable dashboards. Users may browse and search for any registered metric and add it to a dashboard, expanding the diagnostic capabilities available on the Status page. With the addition of this dashboard, developers can contribute diagnostics to the Status page without any additional UI design.



On this page ...

- User Interface
- Types of Metrics
 - Gauge
 - X/Y
 - Meter
 - Histogram
 - Timer
- Available Metrics
- Examples
 - Adding Metrics to a Dashboard

User Interface



Attributes	Description
Add Metric	Select from a list of available metrics to drag and drop onto the dashboard grid.
Grid Size	Control the number of rows and columns on the grid.
Reset	Clear all metrics tiles from the current dashboard.
Dashboard Dropdown	Switch between saved dashboards.
Save Dashboard Dropdown	Save changes, rename a saved dashboard, or delete the current dashboard.

Types of Metrics

Gauge

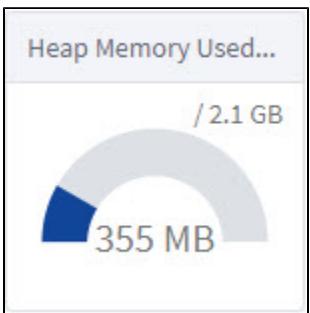
Gauges can be any kind of instant-read measurement, such as the number of megabytes of memory consumed, the count of faulted database connections, or the length of a pending task queue. Gauges are used to display a dynamic number.

Resources	Count
60	3 Projects

Property	Description
value	A numeric value.

X/Y

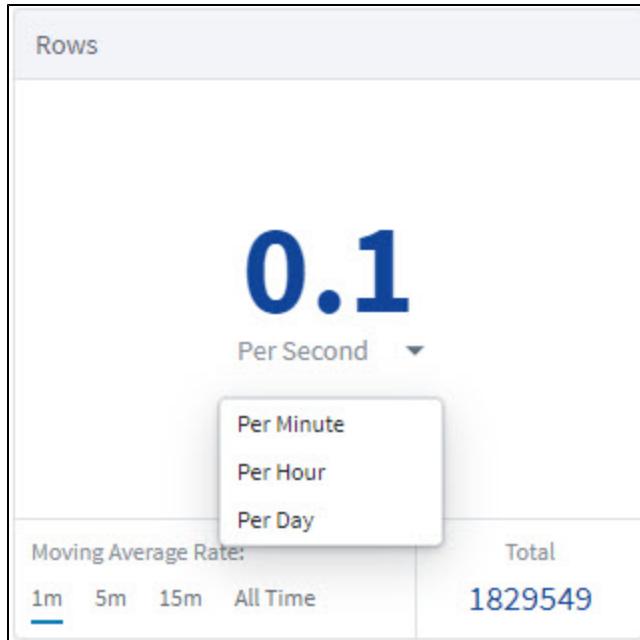
An X/Y metric is a pair of gauges that can be interpreted as a ratio. This is an instant-read gauge on a value that has a maximum, such as heap size or a fixed-size thread pool.



Property	Description
value	A numeric value.
max	The maximum value.

Meter

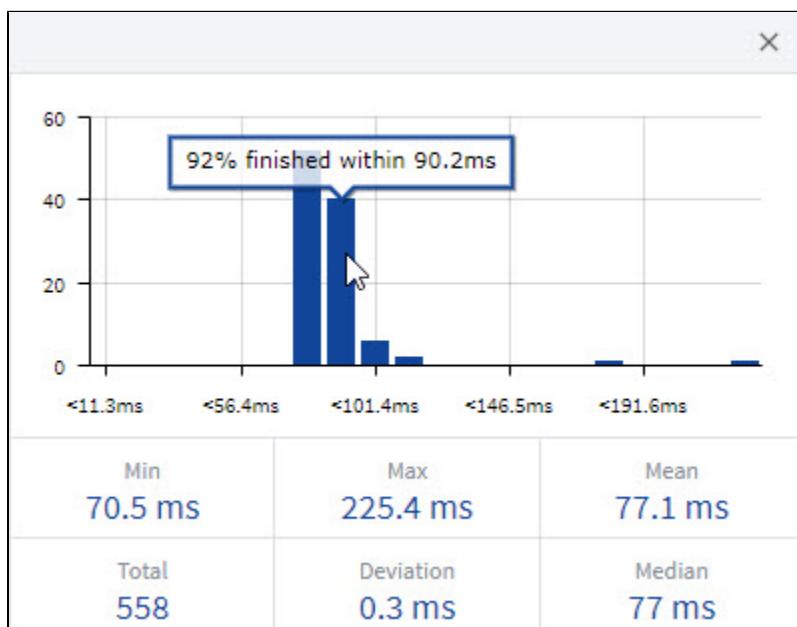
A meter measures the rate at which something occurs, such as the number of messages sent per second. Users can select the units as well as which of the moving average rates to display.



Property	Description
count	The number of events seen.
meanRate	The mean rate, overall.
oneMinuteRate	The one-minute moving average rate.
fiveMinuteRate	The five-minute moving average rate.
fifteenMinuteRate	The fifteen-minute moving average rate.

Histogram

A histogram measures the statistical distribution of values in a stream of data. Hovering over a quantile reveals the value at that quantile.

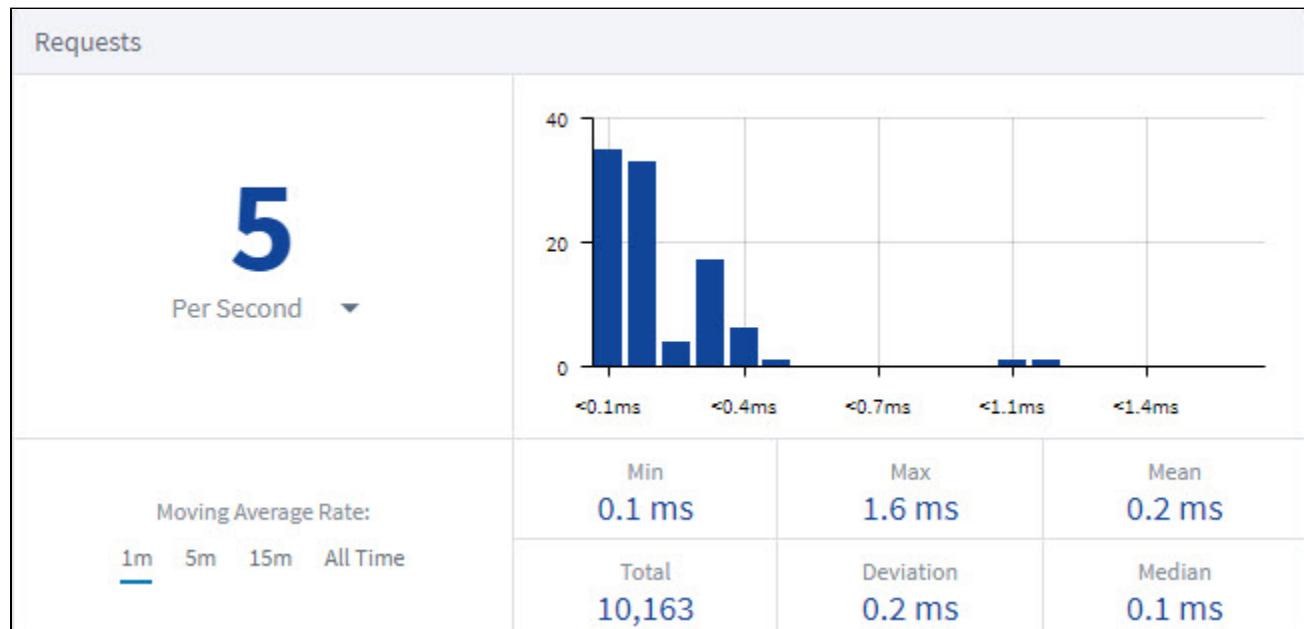


Property	Description
count	The number of values or events seen.

min	The minimum value (in milliseconds).
max	The maximum value (in milliseconds).
mean	The mean value (in milliseconds).
stdDev	The standard deviation (in milliseconds).
median	The median value (in milliseconds).

Timer

A timer is used to understand both the frequency of and duration of events. It is a combination of a meter (rate of occurrence) and a histogram (distribution of duration). A timer combines the functions and properties of both a meter and a histogram.



Available Metrics

The table below includes all metrics currently available to view on the Metrics Dashboard. New metrics may be made available with the development of new systems, modules, and features.

Metric	Description				Size	Type																																					
authTokenRecords	<p>The following feature is new in Ignition version 8.1.37 Click here to check out the other new features</p>																																										
There are authTokenRecords for both the Client and Designer.																																											
Metric	Description				Size	Type																																					
changeEventDistributionOperationAccumulator	<table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchProcessDuration</td> <td>The amount of time that the ChangeEventDistributionOperationAccumulator takes to process a batch.</td> <td>4x2</td> <td>Timer</td> </tr> <tr> <td>changeEventDistributions</td> <td> <table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchSize</td> <td>Number of ChangeEvent Distributions in the ChangeEventDistributionOperationAccumulator batch.</td> <td>1x1</td> <td>Gauge</td> </tr> <tr> <td>dispatchedThroughput</td> <td>The throughput of ChangeEvent Distributions dispatched to the Execution task for batch processing.</td> <td>2x2</td> <td>Meter</td> </tr> <tr> <td>queuedThroughput</td> <td>The throughput of ChangeEvent Distributions added to the batch.</td> <td>2x2</td> <td>Meter</td> </tr> </tbody> </table> </td> <td></td> <td></td> </tr> <tr> <td>changeEventListenerMutations</td> <td> <table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchSize</td> <td>Number of ChangeEvent Listener Mutations in the</td> <td>1x1</td> <td>Gauge</td> </tr> </tbody> </table> </td><td></td><td></td> </tr> </tbody> </table>	Metric	Description	Size	Type	batchProcessDuration	The amount of time that the ChangeEventDistributionOperationAccumulator takes to process a batch.	4x2	Timer	changeEventDistributions	<table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchSize</td> <td>Number of ChangeEvent Distributions in the ChangeEventDistributionOperationAccumulator batch.</td> <td>1x1</td> <td>Gauge</td> </tr> <tr> <td>dispatchedThroughput</td> <td>The throughput of ChangeEvent Distributions dispatched to the Execution task for batch processing.</td> <td>2x2</td> <td>Meter</td> </tr> <tr> <td>queuedThroughput</td> <td>The throughput of ChangeEvent Distributions added to the batch.</td> <td>2x2</td> <td>Meter</td> </tr> </tbody> </table>	Metric	Description	Size	Type	batchSize	Number of ChangeEvent Distributions in the ChangeEventDistributionOperationAccumulator batch.	1x1	Gauge	dispatchedThroughput	The throughput of ChangeEvent Distributions dispatched to the Execution task for batch processing.	2x2	Meter	queuedThroughput	The throughput of ChangeEvent Distributions added to the batch.	2x2	Meter			changeEventListenerMutations	<table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchSize</td> <td>Number of ChangeEvent Listener Mutations in the</td> <td>1x1</td> <td>Gauge</td> </tr> </tbody> </table>	Metric	Description	Size	Type	batchSize	Number of ChangeEvent Listener Mutations in the	1x1	Gauge				
Metric	Description	Size	Type																																								
batchProcessDuration	The amount of time that the ChangeEventDistributionOperationAccumulator takes to process a batch.	4x2	Timer																																								
changeEventDistributions	<table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchSize</td> <td>Number of ChangeEvent Distributions in the ChangeEventDistributionOperationAccumulator batch.</td> <td>1x1</td> <td>Gauge</td> </tr> <tr> <td>dispatchedThroughput</td> <td>The throughput of ChangeEvent Distributions dispatched to the Execution task for batch processing.</td> <td>2x2</td> <td>Meter</td> </tr> <tr> <td>queuedThroughput</td> <td>The throughput of ChangeEvent Distributions added to the batch.</td> <td>2x2</td> <td>Meter</td> </tr> </tbody> </table>	Metric	Description	Size	Type	batchSize	Number of ChangeEvent Distributions in the ChangeEventDistributionOperationAccumulator batch.	1x1	Gauge	dispatchedThroughput	The throughput of ChangeEvent Distributions dispatched to the Execution task for batch processing.	2x2	Meter	queuedThroughput	The throughput of ChangeEvent Distributions added to the batch.	2x2	Meter																										
Metric	Description	Size	Type																																								
batchSize	Number of ChangeEvent Distributions in the ChangeEventDistributionOperationAccumulator batch.	1x1	Gauge																																								
dispatchedThroughput	The throughput of ChangeEvent Distributions dispatched to the Execution task for batch processing.	2x2	Meter																																								
queuedThroughput	The throughput of ChangeEvent Distributions added to the batch.	2x2	Meter																																								
changeEventListenerMutations	<table border="1"> <thead> <tr> <th>Metric</th> <th>Description</th> <th>Size</th> <th>Type</th> </tr> </thead> <tbody> <tr> <td>batchSize</td> <td>Number of ChangeEvent Listener Mutations in the</td> <td>1x1</td> <td>Gauge</td> </tr> </tbody> </table>	Metric	Description	Size	Type	batchSize	Number of ChangeEvent Listener Mutations in the	1x1	Gauge																																		
Metric	Description	Size	Type																																								
batchSize	Number of ChangeEvent Listener Mutations in the	1x1	Gauge																																								

			batch.						
	batch.	batch.	batch.						
	batch.	batch.	batch.						
	batch.	batch.	batch.						
	changeEventListeners. count	The number of ChangeEvent Listeners registered with the system.							
	changeEventPersistence OperationAccumulator	Metric	Description			Size			
		batchProcessDuration	The amount of time that the ChangeEventPersistenceOperationAccumulator takes to process a batch.			4x2			
		batchSize	The number of ChangeEvent Contexts in the ChangeEventPersistenceOperationAccumulator batch.			1x1			
		dispatchedThroughput	The throughput of ChangeEvent Contexts dispatched to the Execution task for batch processing.			2x2			
		queuedThroughput	The throughput of ChangeEvent Contexts added to the batch.			2x2			
		synchronizerDuration	The amount of time that the File Synchronizer takes to synchronize files with memory.			4x2			
	maintenance	Metric	Description			Size			
		cleaner.duration	The amount of time that the Cleaner takes.			2x4			
		pruner.duration	The amount of time that the Pruner takes.			2x4			
		reaper.duration	The amount of time that the Reaper takes.			2x4			
	records.count	The number of AuthTokenRecords in the system.							
	touchEventAccumulator	Using the metrics below, you can identify that a system is processing the incoming volume of touch events efficiently if the dispatchedThroughput keeps up with the queuedThroughput minus the collisionThroughput.							
		Metric	Description			Size			
		batchProcessDuration	The amount of time that the TouchEventAccumulator takes to process a batch.			4x2			
		batchSize	The number of touch events in the TouchEventAccumulator batch.			1x1			
		collisionThroughput	The throughput of touch events added to the batch resulting in a collision.			2x2			
		dispatchedThroughput	The throughput of touch events dispatched to the Execution task for batch processing.			2x2			
		queuedThroughput	The throughput of touch events added to the batch.			2x2			
ClockDriftDetector	Metric	Description			Size	Type			
	driftGauge	The most recent clock drift, in milliseconds.			1x1	Gauge			
	driftTimer	A timer that tracks the distribution of the duration of clock drifts.			4x2	Timer			
databases	Metric	Description							
	connection- {dbName}	Each connected database will have the following associated metrics:							
		Metric	Description			Size			
		activeConnections	Count of active connections.			1x1			
		queries	Measures the rate and distribution of query execution.			4x2			
		rows	Measures how many rows the queries being executed against the given database connection are returning			2x2			
		throughput	Throughput of queries per second.			1x1			
gateway-network	The following feature is new in Ignition version 8.1.37 Click here to check out the other new features								
	Metric	Description							
	activeMsgCountAll	Total count of all active messages currently being sent over the Gateway Network.							
	IncomingBytesSecAll	Accumulation of the number of bytes received per second across all Gateways.							
	IncomingMsgsPer	Accumulation of messages received per second across all Gateways.							

	SecAll																																						
	Outgoing BytesSec All	Accumulation of the number of bytes sent per second to all Gateways.		1x1	Gauge																																		
	Outgoing MsgsPer SecAll	Accumulation of messages sent per second to all Gateways.		1x1	Gauge																																		
	pendingAckExpired	Indicates how many pending acknowledgements were purged during the last cleanup execution. This metric is relevant for outgoing connections.		1x1	Gauge																																		
		<p>Note: Some messages create a pending acknowledgement object. When the message is uploaded to another Gateway and that Gateway acknowledges receiving the message, the pending acknowledgement object is cleared. If no acknowledgement is received, these objects can become orphaned. Orphaned pending message acknowledgements are regularly purged.</p>																																					
	pendingMessagesExpired	Records how many messages were purged during the last cleanup execution, giving an estimate of potentially lost messages. This metric is relevant for incoming connections. Messages that were dispatched, but not downloaded by the other Gateway are regularly purged.		1x1	Gauge																																		
	Pending MsgCountAll	Total count of all messages waiting to be sent over the Gateway Network		1x1	Gauge																																		
	proxyMsgsPerSec	The number of messages per second that are forwarded through this Gateway. Only applies when this Gateway is configured as a proxy Gateway.		2x2	Meter																																		
	services	incomingServiceCalls																																					
		<table border="1"> <thead> <tr> <th>Metric</th><th>Description</th><th>Size</th><th>Type</th></tr> </thead> <tbody> <tr> <td>Get Service State</td><td>Number of calls per second. This metric is queried by other Gateways to see if a specific service is available.</td><td>2x2</td><td>Meter</td></tr> <tr> <td>Enumerate Services</td><td>Measures the rate and distribution of enumerate service calls.</td><td>2x2</td><td>Meter</td></tr> </tbody> </table>	Metric	Description	Size	Type	Get Service State	Number of calls per second. This metric is queried by other Gateways to see if a specific service is available.	2x2	Meter	Enumerate Services	Measures the rate and distribution of enumerate service calls.	2x2	Meter																									
Metric	Description	Size	Type																																				
Get Service State	Number of calls per second. This metric is queried by other Gateways to see if a specific service is available.	2x2	Meter																																				
Enumerate Services	Measures the rate and distribution of enumerate service calls.	2x2	Meter																																				
ignition	Metric	Description		Size	Type																																		
	designer.requests	Measures the rate and distribution of designer requests.		4x2	Timer																																		
	performance	<table border="1"> <thead> <tr> <th>Metric</th><th>Description</th><th>Size</th><th>Type</th></tr> </thead> <tbody> <tr> <td>committed-non-heap</td><td>Committed non-heap memory.</td><td>1x1</td><td>Gauge</td></tr> <tr> <td>cpu</td><td>CPU usage as percent utilized.</td><td>1x1</td><td>Gauge</td></tr> <tr> <td>disk</td><td>Disk used out of total space available.</td><td>1x1</td><td>X/Y</td></tr> <tr> <td>heap</td><td>Heap memory used out of total memory available.</td><td>1x1</td><td>X/Y</td></tr> <tr> <td>heap-max</td><td>Maximum heap memory.</td><td>1x1</td><td>Gauge</td></tr> <tr> <td>heap-used</td><td>Current heap memory used.</td><td>1x1</td><td>Gauge</td></tr> <tr> <td>non-heap</td><td>Current non-heap memory used.</td><td>1x1</td><td>Gauge</td></tr> <tr> <td>non-heap-pct</td><td>Non-heap memory used out of total non-heap memory available.</td><td>1x1</td><td>X/Y</td></tr> </tbody> </table>	Metric	Description	Size	Type	committed-non-heap	Committed non-heap memory.	1x1	Gauge	cpu	CPU usage as percent utilized.	1x1	Gauge	disk	Disk used out of total space available.	1x1	X/Y	heap	Heap memory used out of total memory available.	1x1	X/Y	heap-max	Maximum heap memory.	1x1	Gauge	heap-used	Current heap memory used.	1x1	Gauge	non-heap	Current non-heap memory used.	1x1	Gauge	non-heap-pct	Non-heap memory used out of total non-heap memory available.	1x1	X/Y	
Metric	Description	Size	Type																																				
committed-non-heap	Committed non-heap memory.	1x1	Gauge																																				
cpu	CPU usage as percent utilized.	1x1	Gauge																																				
disk	Disk used out of total space available.	1x1	X/Y																																				
heap	Heap memory used out of total memory available.	1x1	X/Y																																				
heap-max	Maximum heap memory.	1x1	Gauge																																				
heap-used	Current heap memory used.	1x1	Gauge																																				
non-heap	Current non-heap memory used.	1x1	Gauge																																				
non-heap-pct	Non-heap memory used out of total non-heap memory available.	1x1	X/Y																																				
uptime	Total uptime.	1x1	Gauge																																				
Metric	Description		Size	Type																																			
bindings	Perspective Binding count.		1x1	Gauge																																			
components	Perspective Components count.		1x1	Gauge																																			
expressions	Measures the rate and distribution of expressions.		4x2	Timer																																			
fetches	Measures the rate and distribution of fetches.		4x2	Timer																																			
pages	Perspective pages count.		1x1	Gauge																																			
property-changes	Measures the rate of property changes.		2x2	Meter																																			
perspective	scripts	Measures the rate and distribution of script executions.		4x2	Timer																																		
	session-(sessionName)	Each connected Perspective session will have the following associated metrics:																																					
		<table border="1"> <thead> <tr> <th>Metric</th><th>Description</th><th>Size</th><th>Type</th></tr> </thead> <tbody> <tr> <td>expressions</td><td>Measures the rate and distribution of expressions.</td><td>4x2</td><td>Timer</td></tr> <tr> <td>fetches</td><td>Measures the rate and distribution of fetches.</td><td>4x2</td><td>Timer</td></tr> <tr> <td>messages-received</td><td>Count of messages received.</td><td>2x2</td><td>Counter</td></tr> <tr> <td>messages-sent</td><td>Count of messages sent.</td><td>2x2</td><td>Counter</td></tr> <tr> <td>property-changes</td><td>Measures the rate of property changes.</td><td>2x2</td><td>Meter</td></tr> <tr> <td>queue-length</td><td>Current length of queue.</td><td>1x1</td><td>Gauge</td></tr> </tbody> </table>	Metric	Description	Size	Type	expressions	Measures the rate and distribution of expressions.	4x2	Timer	fetches	Measures the rate and distribution of fetches.	4x2	Timer	messages-received	Count of messages received.	2x2	Counter	messages-sent	Count of messages sent.	2x2	Counter	property-changes	Measures the rate of property changes.	2x2	Meter	queue-length	Current length of queue.	1x1	Gauge									
Metric	Description	Size	Type																																				
expressions	Measures the rate and distribution of expressions.	4x2	Timer																																				
fetches	Measures the rate and distribution of fetches.	4x2	Timer																																				
messages-received	Count of messages received.	2x2	Counter																																				
messages-sent	Count of messages sent.	2x2	Counter																																				
property-changes	Measures the rate of property changes.	2x2	Meter																																				
queue-length	Current length of queue.	1x1	Gauge																																				

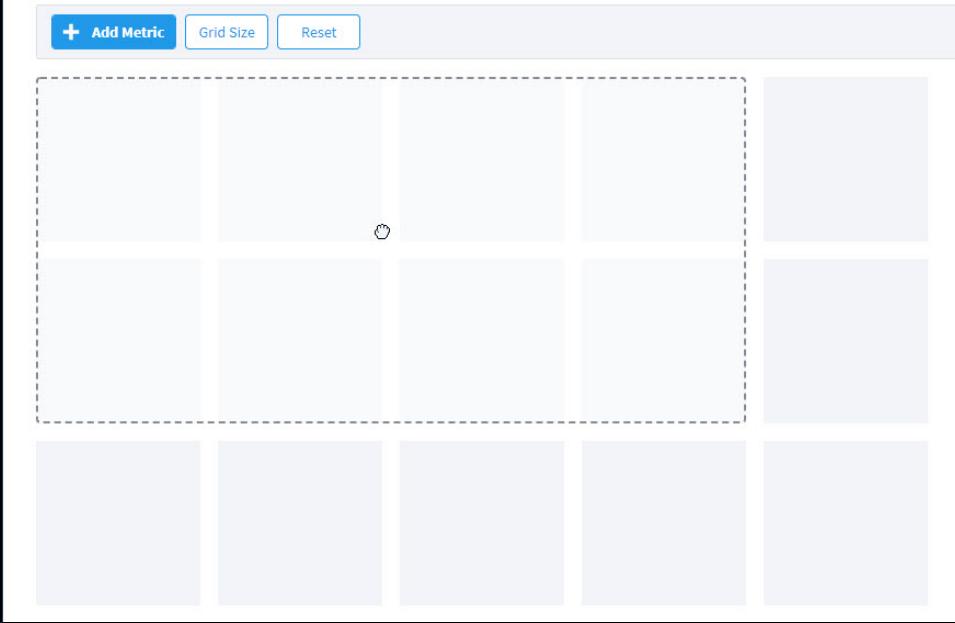
		queue-tasks	Measures the rate and distribution of queued tasks.	4x2	Timer	
		reconnects	Count of reconnects,	2x2	Counter	
		scripts	Measures the rate and distribution of script executions.	4x2	Timer	
		uptime	Total uptime for the Perspective session.	1x1	Gauge	
	sessions	Count of active sessions.			1x1	Gauge
	views	Count of views.			1x1	Gauge
projects	Metric	Description			Size	Type
	(projectName).resources	Count of resources per project.			1x1	Gauge
	count	Count of projects.			1x1	Gauge
	disk-sync	Measures the frequency and duration of projects directory synchronization.			4x2	Timer
	saves	Measures the rate and distribution of project saves.			4x2	Timer
redundancy	Metric	Description			Size	Type
	config.updateThroughput	Rate of updates.			2x2	Meter
	alarmpipelinemanager	Metric	Description	Size	Type	
		taskSendThroughput	Rate of tasks sent.	2x2	Meter	
		taskRecvThroughput	Rate of tasks received.	2x2	Meter	
	sync.queue_max	queueSize	Current queue size.	1x1	Gauge	
		Maximum queue size.			1x1	Gauge
scripts	Metric	Description			Size	Type
	gateway	Metric	Description	Size	Type	
		compile	Measures the rate and distribution of script compiles.	4x2	Timer	
		execute	Measures the rate and distribution of script executions.	4x2	Timer	
	project-(projectName)	Metric	Description	Size	Type	
		compile	Measures the rate and distribution of script compiles.	4x2	Timer	
		execute	Measures the rate and distribution of script executions.	4x2	Timer	
vision.requests	Measures the rate and distribution of Vision requests.					

Examples

Adding Metrics to a Dashboard

1. Click **Add Metric**.
2. Select a metric from the list and click **Add**.

3. Hover over the grid to position the metric.



4. Click to snap the metric to the dashboard. To reposition your metric, you can drag and drop it anywhere else on the grid.



Config

Gateway Configuration

The Config tab provides access to configuration options for [Gateway settings](#). This is where most of the settings that affect the whole Gateway are set up. We can add database and device connections, users and roles, adjust alarm settings, set up security, and create a schedule for a Gateway backup to be taken automatically at specific times.

The list of Config options on the left menu change based on what modules are installed on your Gateway. Third-party modules have settings that are not discussed on this page.

Once you have the [Gateway](#) up and running, you start by configuring some or all of the general services in Ignition. You make the configuration changes from the [Config](#) section of the Gateway Webpage. The different broad categories of what you can configure are as follows:

- [System](#) (Overview, Backup/Restore, Licensing, Modules, Projects, Redundancy and Gateway Settings)
- [Networking](#) (Web Server, Gateway Network and Email Settings)
- [Security](#) (Auditing, Users, Roles, Service Security, Identity Providers, OAuth2 Clients, Security Levels, and Security Zones)
- [Databases](#) (Connections, Drivers, and Store and Forward)
- [Alarming](#) (General, Journal, Notification, On-Call Rosters, and Schedules)
- [Tags](#) (History and Realtime)
- [OPC Client](#) (OPC Connections and OPC Quick Client)
- [OPC UA](#) (Device Connections, Security, and Server Settings)
- [Enterprise Administration](#) (Event Thresholds, Controller Settings, Agent Management, License Management, and Agent Tasks)
- [Sequential Function Charts](#) (Settings)

- The following feature is new in Ignition version **8.1.20**
[Click here](#) to check out the other new features

[Perspective](#) (Branding Customization)

- The following feature is new in Ignition version **8.1.28**
[Click here](#) to check out the other new features

[Connectors](#) (MongoDB)

Note: Depending on what modules you have installed, some categories may be missing.

On this page ...

- [Gateway Configuration](#)
- [System](#)
- [Networking](#)
- [Security](#)
- [Databases](#)
- [Alarming](#)
- [Tags](#)
- [OPC Client](#)
- [OPC UA](#)
- [Enterprise Administration](#)
- [Sequential Function Charts](#)
- [Perspective](#)
- [Connectors](#)

The screenshot shows the Ignition Configuration interface. The left sidebar has a dark blue header with the Ignition logo and navigation tabs: Home, Status, Config (which is selected). Below the tabs are sections for SYSTEM, NETWORKING, SECURITY, DATABASES, and ALARMING, each with a list of sub-options. A search bar at the bottom of the sidebar contains the placeholder "Search...". The main content area has a header "Config" > "System" > "Overview". The title "Configuration" is centered above a section of text: "From the Configure section you can set up all connections, projects, and settings. Here are some common actions to get you started." Below this are six cards arranged in a grid:

- PLATFORM**: Update System Name, Configure Redundancy, Install or Upgrade a Module, Create New Project, Activate a License, Download Gateway Backup.
- NETWORKING**: Change Web Server Settings, Enable SSL for the Gateway Network, Create an Email Profile, Manage incoming/outgoing Gateway Network connections.
- SECURITY**: Change General Gateway Security Settings, Create a new user, Assign a user a new role, View the logs of an audit profile, Define a Security Zone, Set access levels on a Security Policy.
- CONNECTIONS**: Create a new database connection, Connect to a 3rd party OPC server, Create a new device connection.
- SYSTEMS**: Create an alarm journal profile, Manage schedules and holidays, Create a new alarm notification profile, Test an alarm notification pipeline, Add users to an on-call roster.
- DATA ACQUISITION**: Define a new realtime tag provider, Manage tag historians, Quickly read or write tags in a device.

System

The System section is a sort of catch all section that can do a lot of different things. The first window is the Overview page, which is what will first show up when navigating to the Configure tab. The Overview page only has links to other sections of the Config tab, but it is useful if you aren't exactly sure where the setting you are looking for is located, because the Overview page lists the most common configuration changes.

Page	Description
Backup /Restore	The Backup/Restore page is where you can manually take a backup, and restore a previous backup.
Ignition Exchange	On the Ignition Exchange page you can access the Ignition Exchange to browse for resources. You can also import an Ignition Exchange Package that you've downloaded.
Licensing	The Licensing page allows you to control any and all licenses currently activated on the Gateway. It will show all the modules that the license is currently good for, and have some options for activating a new or additional license.
Modules	The Modules page displays all currently installed modules and if they are active or not. This is where you can install a new module or upgrade an existing module.
Projects	The Projects page can easily manage all of the projects currently configured in the Gateway, with settings that change the Name and Title, or the default database and authentication profile of the project.
Redundancy	The Redundancy page is where all of Ignition's redundancy settings are configured. This is where the master and backup nodes are configured as well as the network settings to make sure the two nodes can properly communicate.
Gateway Settings	The Gateway Settings page is where settings are located for the system name, homepage redirect URL, launch settings, scheduled backups, error reporting and other miscellaneous.

Networking

The Networking section deals with setup and management of the [Gateway Network](#).

Page	Description
Web Server	<p>The Web Server page is for configuring the http and https ports, setting up the SSL / TLS certificate, redirecting traffic through a known address, and whether or not all http traffic should be forcefully redirecting to https. If you are allowing users to access your Gateway from outside your network (through the internet), you will need to configure the Public HTTP Address settings.</p> <p>You can find out more about SSL in Secure Communication (SSL / TLS).</p>
Gateway Network	<p>The Gateway Network allows you to connect multiple Gateways together over a wide area network. The Gateway's connection settings can also be changed to only allow certain connections. The Gateway Settings set the basic rules for the system.</p>
Email Settings	<p>The Email Settings section allows you to create an SMTP server connection in the Gateway that can be used by several different resources such as Alarm Notification and Report Schedules.</p>

Security

Ignition provides several [Security](#) options to safeguard data and applications in Ignition which fall into the following categories, and are set up in the Config section.

Page	Description
General	<p>The following feature is new in Ignition version 8.1.0 Click here to check out the other new features</p> <p>The Gateway General Security Settings page is new for 8.1.0. It determines security permissions for the Gateway and Designer. For more information, see Security.</p>
Auditing	Here you can set up an Audit Profile to record details about specific events that occurred.
Users and Roles	This is where you set up users and roles . Security policies are defined in terms of the roles set up in the system.
Service Security	A security policy can be defined for each Security Zone and is set up in Service Security . The Security Policy has four sections: Alarm Notification, Alarm Status, History Provider Access, and Tag Access. They work together to define how the local Gateway gives access to incoming Gateway connections.
Identity Providers	Identity Providers (IdP) provide a way for users to log in to Ignition using credentials stored outside of Ignition. An IdP creates, maintains, and manages identity (login) information while providing authentication services to Ignition. This provides a secure login that allows Ignition to use SSL and two-factor authentication (2FA).
OAuth 2 Clients	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>Set up OAuth 2.0 Clients here.</p>
Security Levels	Security Levels define a hierarchy for access inside a Perspective Session or Vision Client using IdP authentication. This authorization system provides a way to map roles from an Identity Provider (IdP) to Ignition roles.
Security Zones	A Security Zone is a list of Gateways, Computers, or IP addresses that are defined and grouped together. This group now becomes a zone on the Gateway Network , which can have additional policies and restrictions placed on it. Security Zones provide this functionality to the Gateway Network, limiting locations instead of people to be read-only for specific actions. This allows for greater control over the type of information that is passing over the network, improving security and helping to keep different areas of the business separate, while still allowing them to interconnect.

Databases

The Databases section is where [database connections](#) are set up. Databases are used in historical data logging, reporting, storing alarm logs, and Tag storage.

Page	Description
Connections	Many of the advanced features of Ignition, such as the Transaction Groups and Tags Historian require a connection to an external database, and most databases require special permissions for each computer that wants to connect. Ignition takes care of all of this. You create a connection to your database once, and every system in Ignition will use that central connection. From here, you can create new database connections and edit existing connections.
Drivers	JDBC drivers used in database connections are imported and configured in Drivers of the Database section.
Store and Forward	The Store-and-Forward system provides a reliable way for Ignition to store data to the database. The Store-and-Forward system settings offer a good deal of flexibility in tuning. Different types of situations and goals will likely require different configurations.

Alarming

The Alarming section provides general [alarm configuration](#) settings to provide up-to-date status of alarms, store alarm history, build the logic for how, why, and when alarm notifications are delivered, manage alarm notifications for user groups, and send Email, SMS, or Voice notifications. With all these features and functions in Alarming, you can easily create alarms, and design and manage your alarm notifications any way you choose.

Page	Description
General	This General setting provides of some basic alarm configuration settings.
Journal	Alarm Journals are configured in the Alarming section to store basic historical information in a database about alarms that occurred, such as their source and timestamp, associated data on the alarm, and the values of the alarm's properties at the time the event occurred.
Notification	Notification Profiles are configured in the Alarming section to allow for Email , SMS , or Voice notifications to be sent out when an alarm event occurs.
On-Call Rosters	The On-Call Roster is where you create user groups to be notified when an alarm occurs. When an alarm is triggered, it is sent to a designated On-Call Roster where it evaluates the users schedules, and only notifies those users that have an active schedule. Users that are off-schedule will not be notified.
Schedules	Defines the times of users on-call availability and unavailability by configuring Schedules .

Tags

The Tags section is where both [Realtime Tag Providers](#) as well as [Historical Tag Providers](#) are configured. Note that this is not where individual Tags are set up.

Page	Description
History	Configure the settings for each of the configured Historical Tag Providers, or create new remote or split providers.
Realtime	Configure the settings for each of the configured Realtime Tag Providers, or create new standard or remote providers.

OPC Client

The OPC Client section is where connections from Ignition's internal OPC UA server to other OPC servers are located.

Page	Description
OPC Connections	Configure OPC Connections to Ignition's built in OPC Server.
OPC Quick Client	The OPC Quick Client allows for quick and simple testing of any OPC Connections connected to the OPC server.

OPC UA

The OPC UA Server is where Ignition's internal OPC UA server is configured.

Page	Description

Device Connections	This is where all device connections to our internal OPC UA server are configured.
Security	Upload and trust client and server OPC UA certificates .
Server Settings	Configure Ignition's internal OPC UA Server 's settings.

Enterprise Administration

The [Enterprise Administration](#) section controls the majority of [EAM functions](#). Setting the Gateway to be a Controller or Agent as well as creating Agent Tasks and managing the various Agents from the Controller Gateway can all be done in this section.

Sequential Function Charts

A [Sequential Function Chart \(SFC\)](#) is a series of scripts that are defined in a single location and then called in sequential order.

Page	Description
Settings	This is where the SFC Settings are configured.

Perspective

The following feature is new in Ignition version **8.1.20**

[Click here](#) to check out the other new features

Perspective co-branding customization allows you to add your own flair on your Perspective apps by incorporating custom colors and logos.

Page	Description
Branding Customization	This is where Co-Branding settings are configured.

Connectors

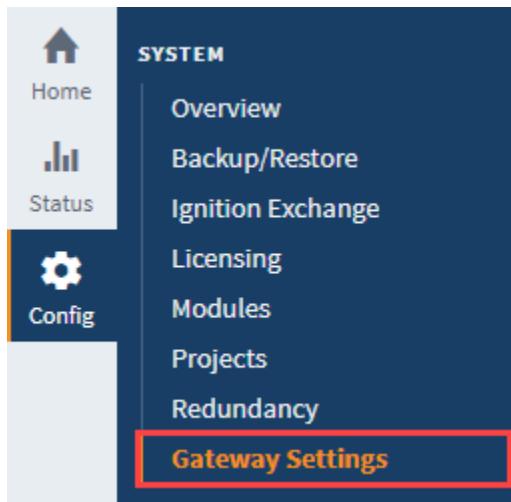
Cloud Connectors are geared towards Ignition Cloud Edition and cloud infrastructure, although the connectors can also be used on a standard Ignition Gateway.

Page	Description
MongoDB	This is where MongoDB Connections are configured.

Gateway Settings

Gateway Settings Property Reference

After you [launch the Gateway](#), you can define the high-level settings that apply to the entire Gateway by going to the **Config > Gateway Settings** on the Gateway Webpage. From this page, you can use the default values or define a new setting. The Gateway is at the heart of the Ignition software. It runs as a web server and you can access it through a web browser. Once running, you can get various status information about the Gateway and access important functions.



On this page ...

- [Gateway Settings Property Reference](#)
 - [Local Client Fallback Settings](#)
 - [Scheduled Backups](#)
 - [Automatic Thread Dump Settings](#)
 - [Error Reporting](#)
 - [Multicast Settings](#)

The following feature is new in Ignition version **8.1.0**
[Click here](#) to check out the other new features

The Gateway Settings Properties were updated significantly for release 8.1.0. Many security related settings have been moved to [Gateway General Security Settings](#).

The following tables describe all the properties on the Gateway Settings.

Gateway Settings	
System Name	<p>Is a unique name for this Ignition installation. It is used to distinguish this server from others on the network when working with multiple Ignition installations.</p> <p>Caution: It is not recommended to change the Gateway's System Name after initial Gateway setup. Renaming your Gateway may break critical dependencies, affect data collection, and halt production. To see a list of features dependent on the Gateway name, click here.</p>
Homepage Redirect URL	The URL this gateway will redirect to when http://ip:port/ is visited. Can either be a relative path (e.g., /web/home), or fully qualified (e.g., https://inductiveautomation.com).
Gateway Scripting Project	The Gateway Scripting Project is a Project in which Gateway-scoped scripts with no project affiliation can access user script libraries.
Enable Tag Reference Tracker Store	<p>The following feature is new in Ignition version 8.1.34 Click here to check out the other new features</p> <p>Enables the storing of Tag Reference entries to a database on the local Gateway for analysis in a Designer for third party modules. Default is true.</p>

Launch Settings	
Designer Memory	The maximum amount of memory that the Designer has access to.
Disable Direct3D	Disables the Direct3D rendering pipeline launched clients. Direct3D can cause performance problems with XOR painting. (affects clients on Windows only)
Disable DirectDraw	Disables the DirectDraw system for launched clients. Can be useful for some video cards that don't support DirectDraw well. (affects clients on Windows only)
Use Condensed Font	<p>The following feature is new in Ignition version 8.1.4 Click here to check out the other new features</p> <p>Designers and Vision Clients will use a condensed version of the "Dialog" font for better compatibility with pre-8.0 Vision project design. (affects clients on Windows and Linux only). Defaults to false.</p>

Local Client Fallback Settings

Ignition provides a [Local Vision Client](#) Fallback mechanism that lets you use a Gateway running on the local machine. If the Gateway is lost, the Client can automatically retarget to a project that you specify in the local in the local Gateway.

Local Vision Client Fallback	
Enable Local Fallback	Enables a client to fall back to a project in a local Gateway if communication is lost to the central Gateway. Note that port 6501 must be open on the local machine.
Seconds Before Failover	The number of seconds to wait before switching to the local Gateway project after communication loss.
Fallback Project	The local project to use during fallback.

Scheduled Backups

The Scheduled Backup Settings control the Gateway's scheduled backup system. This system is capable of automatically making a [Gateway backup](#) and storing it to the specified location, which can be a network path. When you enable this system, you must specify a destination folder. This can be a local folder, for example `C:\backups` or `/var/backups`, or a network path such as `\fileserver\backups`.

The scheduled backup system works on a schedule that is specified using UNIX Crontab syntax. See the [Crontab Formatting Reference](#) page for more details.

Scheduled Backup Settings							
Enable Scheduled Backups	Enables the scheduled backup system which automatically makes backups at a scheduled time.						
Backup Folder	A path to a folder in which to put the scheduled backups.						
Backup Schedule	A UNIX 'crontab' format scheduling string representing when to make the backups.						
Retention Count	The number of backups to keep in the backup folder.						
Filename Pattern	<p>The following feature is new in Ignition version 8.1.18 Click here to check out the other new features</p> <p>The filename pattern used for creating scheduled and manually downloaded Gateway backups. Default pattern is <code> \${gatewayName}_Ignition-backup-\${edition}\${timestamp}.gwbk</code>. There are a few possible placeholder pattern values:</p> <table border="1"> <thead> <tr> <th>Filename Pattern</th><th>Description</th></tr> </thead> <tbody> <tr> <td><code> \${os}</code></td><td>The operating system the Gateway is installed on.</td></tr> <tr> <td><code> \${version}</code></td><td>The Ignition Gateway version.</td></tr> </tbody> </table>	Filename Pattern	Description	<code> \${os}</code>	The operating system the Gateway is installed on.	<code> \${version}</code>	The Ignition Gateway version.
Filename Pattern	Description						
<code> \${os}</code>	The operating system the Gateway is installed on.						
<code> \${version}</code>	The Ignition Gateway version.						

\${gatewayName}	The System Name of your Gateway.
\${edition}	The Ignition version you are using. Standard Ignition versions will be blank.
\${timestamp}	The current date and time at which the backup initializes.

Automatic Thread Dump Settings

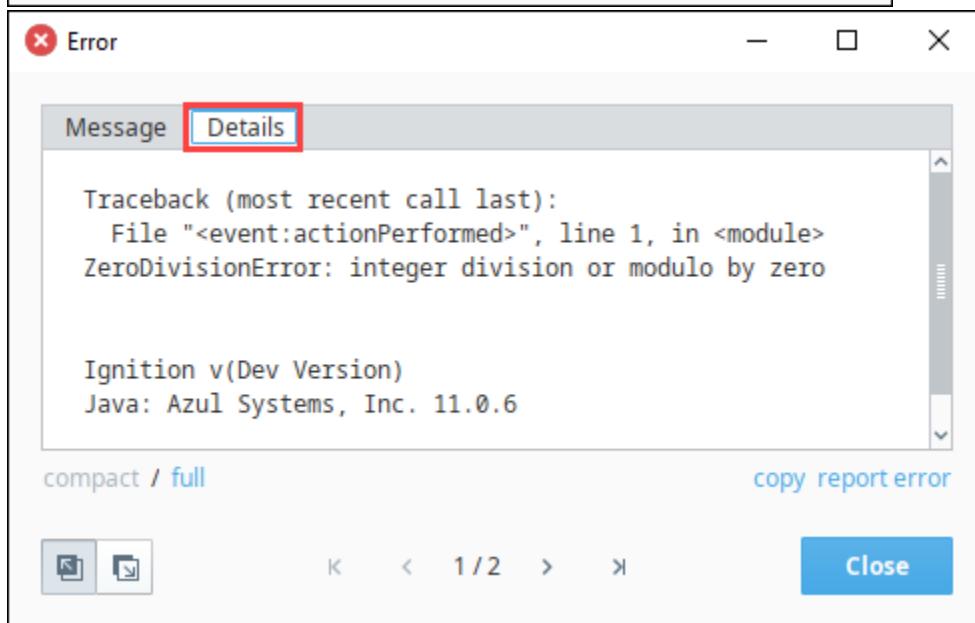
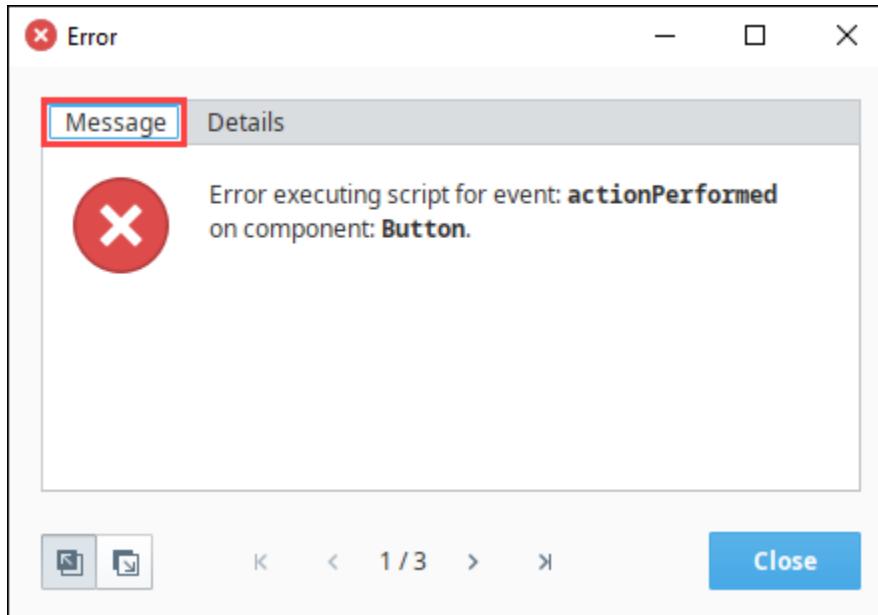
The following feature is new in Ignition version **8.1.13**
[Click here](#) to check out the other new features

When enabled, Ignition will conduct automatic thread dumps at the specified interval when above the target CPU Usage Threshold. If the threshold is met or exceeded, a timestamped thread dump will be saved to the logs folder. When the number of thread dumps in the logs folder exceeds the Retention Count, old files will be automatically removed. The default sample rate (60 seconds) can be overridden by adding an [additional parameter](#) to the Ignition configuration file.

Automatic Thread Dump Settings	
Enable Automatic Thread Dumps	Enables automated thread dump generation when CPU utilization exceeds the threshold below.
Thread Dump Capture Interval	Number of seconds between thread dumps when utilization threshold+duration is met.
CPU Usage Threshold	CPU Usage Threshold (%) that when exceeded for a given duration (below) will trigger a thread dump capture.
CPU Usage Exceedance Duration	Duration (seconds) during CPU Usage Threshold exceedance to delay a thread dump capture.
Retention Count	The number of thread dumps to keep in the logs folder.

Error Reporting

When an error occurs in the Client or Designer, the users can click a link on the Details tab to report the error via email.



These settings define how the errors are reported.

Error Reporting	
SMTP Server	When not blank, user-reported errors are emailed using this SMTP server.
To Address	The email address(es) that will receive the error notification. Separate multiple email addresses with a semicolon (;)
From Address	The email address that the error notification is from.
SMTP Username	A username for the SMTP server, if required.
Change Password?	Check this box to change the existing password.
Password	A password for the SMTP server, if required.
Password	Re-type password for verification.

Multicast Settings

These properties allow the Gateway to broadcast information about itself via multicast UDP packets. This allows the Gateway to be discoverable by any components that are also listening to the same multicast address. For example, native client launchers listen on a multicast address to provide a list of available Gateways on the network. Verify that the send ports and receive ports are open on the Gateway machine in order to be able to broadcast multicast message.

Multicast Settings	
Enable Multicast	Allows this Gateway to be discoverable on your local network.
Multicast IP Address	Gateway messages are broadcast on this address.
Send Port	This port must be open on this machine to send multicast messages.
Receive Port	This port must be open on any machine that will receive multicast messages.
Message Interval	The interval in milliseconds at which multicast messages will be sent.

Related Topics ...

- [Config](#)
- [Status](#)
- [Gateway Security](#)
- [Gateway Command-line Utility - gwcmd](#)

Email Settings

Email Profiles

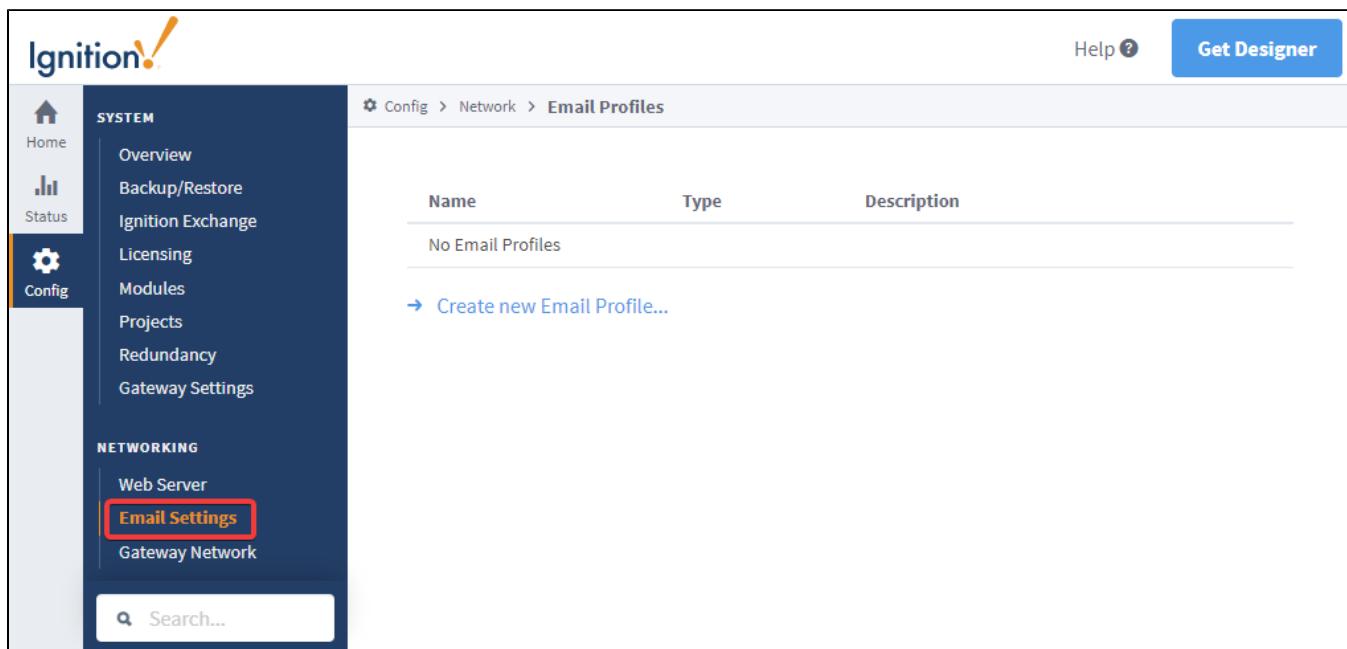
In Ignition, there are several places that you might want to send an email from. Instead of setting up a new email server connection at each one, you can add Email Profiles in the Gateway Config section and reference them in other places. For example, you can use a pre-configured connection in any of these places:

- **Alarm Notification** - While the alarm notification system allows you to set up an Email Notification Profile that is separate from the Email Profile, you can instead choose to use the settings configured in the Email Profile.
 - **Report Schedules** - When Scheduling an Email Action in a report, the Action requires a configured Email Profile before the report can be emailed out.
 - **system.net.sendEmail()** - Instead of manually entering in values for the Email Profile inside each of your scripts, this function can instead use the settings from an existing Email Profile.

Once your Email Profile is set up, you only need to reference the name of the connection and Ignition will take care of the rest.

Setting Up an Email Profile

The Email Settings page allows you to select either basic authentication or OAuth 2.0 and configure an Email Profile from there. These Email Profiles can be used by several different resources in the Gateway. This means that instead of setting up an Email Profile within each resource (i.e., Alarm Notification Profile, Reports, etc.), you can configure them once here, and have those resources all use the same SMTP server. You can find the Email Settings in the Config section of the [Gateway](#) under the Networking heading. Here, you can create a new Email Profile, or manage your existing profiles.



Classic SMTP Profile Settings

Below is a list of properties available on a classic SMTP server connection.

Property Name	Property Description
Main	
Name	The name of the SMTP profile.
Description	A description of the SMTP profile.
SMTP Server Settings	

On this page ...

- Email Profiles
 - Setting Up an Email Profile
 - Classic SMTP Profile Settings
 - OAuth2 SMTP
 - OAuth2 SMTP Profile Settings
 - Manage Token Example
 - Test OAuth2 Email Profile Example

Hostname	Hostname of the SMTP server to send email through.
Port	Port SMTP service is running on. Default is 25.
Use SSL/TLS	Connect using dedicated SSL/TLS. Default is false.
Use STARTTLS	Enable use of the STARTTLS command, allowing the connection to be upgraded to an SSL or TLS connection if supported by the server. This is not necessary for connections that are already SSL/TLS. Default is false.
Username	The username the Gateway will use when authenticating against the mailserver. This is only required if the SMTP server expects authentication.
Password	The password the Gateway will use when authenticating against the mailserver. This is only required if the SMTP server expects authentication.
Password	Retype password for verification.

Advanced Properties

SMTP Timeout	Timeout (in milliseconds) to use when connecting to, reading from, and writing to the SMTP server. Default is 10,000.
Debug Mode Enabled	Enable email session debugging. Information is printed to standard output (wrapper.log). Default is false.
SSL Protocols	A comma separated list of protocols that will be allowed if connecting via SSL/TLS. Default is TLSv1.2.

OAuth2 SMTP

The following feature is new in Ignition version **8.1.24**

[Click here](#) to check out the other new features

Besides setting up classic SMTP Email Profiles, you can also set up OAuth2 SMTP Profiles. To set up an OAuth2 SMTP Profile, go to the Email Settings properties of the Gateway's **Config** section.

OAuth2 SMTP Profile Settings

Below is a list of properties available on an OAuth2 SMTP server connection.

Property Name	Property Description
Main	
Name	The name of the SMTP profile.
Description	A description of the SMTP profile.
SMTP Server Settings	
Hostname	Hostname of the SMTP server to send email through.
Port	Port SMTP service is running on. Default is 25.
Use SSL/TLS	Connect using dedicated SSL/TLS. Default is false.
Use STARTTLS	Enable use of the STARTTLS command, allowing the connection to be upgraded to an SSL or TLS connection if supported by the server. This is not necessary for connections that are already SSL/TLS. Default is false.
OAuth2 Settings	
Username	The username that the Gateway will use along with the access token which together make up the credential in the SMTP protocol's XOAUTH2 SASL authentication mechanism.
OAuth2 Client	The name of the OAuth2 Client which will be used by this Email Profile to obtain access tokens.
Interactive	If enabled, user interaction is required to obtain access and refresh tokens (3-legged OAuth 2.0 flow). If disabled, user interaction is not required to obtain access tokens - the Gateway will obtain access tokens using its own client credentials (2-legged OAuth 2.0 flow). Default is false.
Scope	The OAuth 2.0 scope values representing the access token's set of permissions requested by the Gateway.

Advanced	
SMTP Timeout	Timeout (in milliseconds) to use when connecting to, reading from, and writing to the SMTP server. Default is 10,000.
Debug Mode Enabled	Enable email session debugging. Information is printed to standard output (wrapper.log). Default is false.
SSL/TLS Protocols	A comma separated list of protocols that will be allowed if connecting via SSL/TLS. Default is TLSv1.2.
Advanced OAuth2 Settings	
Access Token Timeout	Amount of time in seconds to wait for a valid access token before giving up on sending an email. Values less than or equal to zero will cause the profile to expect a valid access token to be available immediately and will fail fast if one is not available at the time an email needs to be sent. Default is 120.
Default Access Token Expiration	The amount of time in seconds before an access token expires by default. This value only applies for access tokens which do not have an explicit expiration from the authorization server. Values less than or equal to zero means access tokens do not expire by default. Default is 0.
Refresh Before Expiration	The amount of time in seconds before an access token expires when a new token should be requested. Values less than or equal to zero results in new access token requests exactly when the last access token expires. Default is 300.
Additional Authorization Request Parameters	Add any additional parameters which should be encoded on the OAuth 2.0 authorization request (interactive only).
Additional Authorization Code Access Token Request Parameters	Add any additional parameters which should be encoded on the OAuth 2.0 access token request granted using an authorization code (interactive only).
Additional Client Credentials Access Token Request Parameters	Add any additional parameters which should be encoded on the OAuth 2.0 access token request granted using client credentials (non-interactive only).
Additional Refresh Token Access Token Request Parameters	Add an additional parameters which should be encoded on the OAuth 2.0 access token request granted using a refresh token (interactive only).

Note: When using OAuth2 Email Profiles, you must set up predefined OAuth2 Clients for authorization. An OAuth2 Client may be used for multiple OAuth2 Email Profiles. See the [OAuth 2.0 Client page](#) for more information.

Once you set up OAuth2 SMTP Email Profile, you will have options to manage your token, delete the SMTP Profile, or test it by composing and sending a test email.

The screenshot shows the 'Email Profiles' configuration screen. At the top, there's a breadcrumb navigation: Config > Network > Email Profiles. Below this, a table lists an existing profile: 'Name' is 'OAuth2SMTP', 'Type' is 'OAuth2 SMTP'. To the right of the table are two buttons: 'More ▾' and 'edit'. A red box highlights a dropdown menu that appears when hovering over the 'More' button. The dropdown menu contains three items: 'manage token', 'delete', and 'test'.

Name	Type	Description
OAuth2SMTP	OAuth2 SMTP	More ▾ edit

→ Create new Email Profile...

[manage token](#)
[delete](#)
[test](#)

Manage Token Example

This page will allow you to view the current state of the access token managed by the OAuth2 SMTP Email Profile. If the Email Profile has any issues obtaining or refreshing tokens, details about the issue(s) will be shown here. There will be actions you may take to correct the issue(s), such as granting a new access token, or manually refreshing an access token. In some cases, you may need to adjust your Email Profile configuration settings to fix a problem, depending on the error shown.

Note: The "Refresh Token" button will only be available if the OAuth2 SMTP Email Profile is not interactive OR if it is interactive and a refresh token is available. Additionally, the "Grant New Token" option will only be available when the Email Profile is interactive.

This option will allow you to grant the Client a new access token, or refresh the current token and its expiration time.

1. Set up your OAuth2 SMTP Email Profile if you haven't already. In this example, we will be using the OAuth2 Email Profile in the screenshot above.
2. Click on "More", and then on "manage token".

The screenshot shows the Ignition configuration interface. The path is 'Config > Network > Email Profiles'. A table lists an email profile named 'OAuth2SMTP' of type 'OAuth2 SMTP'. To the right of the profile are 'More' and 'edit' buttons. A context menu is open, with the 'manage token' option highlighted by a red box. Other options in the menu include 'delete' and 'test'.

3. If you are using an interactive-based token, your access token will show as invalid the first time you click on "manage token". To fix this, click on "Grant New Token".

The screenshot shows the 'Manage Token' page for the 'OAuth2SMTP' profile. The page title is 'Config > Network > Email Profiles > OAuth2SMTP > Manage Token'. A message box displays 'Access Token is Invalid' with the sub-instruction 'Access token is not yet set. Click the "Grant New Token" button to grant Ignition a new access token.' Below the message is a large blue 'Grant New Token' button.

4. Since we are using a Google Email Profile, we are redirected to a Google permission prompt. Click "Allow" to continue.

 Sign in with Google

[REDACTED] wants to access
your Google Account

D [REDACTED]

This will allow [REDACTED] to:



Read, compose, send, and permanently delete all [\(i\)](#)
your email from Gmail

Make sure you trust [REDACTED]

You may be sharing sensitive info with this site or app. You
can always see or remove access in your [Google Account](#).

Learn how Google helps you [share data safely](#).

See [REDACTED] Privacy Policy and Terms of Service.

[Cancel](#)

[Allow](#)

5. After the authorization server grants the initial access token, you will be redirected back to the Gateway. After a few seconds, the Gateway will refresh, showing a valid access token.

Config > Network > Email Profiles > OAuth2SMTP > Manage Token

 **Access Token is Valid**

Access token fetched 29m53s ago. Expires in 30m5s. Next refresh attempt in 25m5s.

[Grant New Token](#)

Test OAuth2 Email Profile Example

This option will allow you to send a test email using the OAuth2 Email Profile and OAuth2 Client you set up.

1. Navigate to your OAuth2 Email Profile. Click on "More", and then "test".

The screenshot shows a list of email profiles. One profile is selected, and a context menu is open next to it. The menu items are: 'More', 'edit', 'manage token', 'delete', and 'test'. The 'test' button is highlighted with a red rectangle.

Name	Type	Description
OAuth2SMTP	OAuth2 SMTP	

[Create new Email Profile...](#)

2. This will bring you to a page where you can compose your test email.

The screenshot shows a form for composing a test email. The fields are: From (empty), To (empty), Subject (Test Email), Body (empty), and Content Type (text/plain). At the bottom is a large blue 'Send Test Email' button and a link to 'Back to Email Profiles'.

From	
To	
Subject	Test Email
Body	
Content Type	text/plain

Send Test Email

[< Back to Email Profiles](#)

3. Fill out the form. The "From" field should be populated with the email address that will be sending the email, and the "To" field should be the test email's recipient. Click on "Send Test Email" when you are done.

Compose Test Email

From	
To	
Subject	Test Email
Body	Test OAuth2
Content Type	text/plain ▾

Send Test Email

< Back to Email Profiles

4. If the email sent successfully, you should receive a notification below the "Send Test Email" button.

Send Test Email

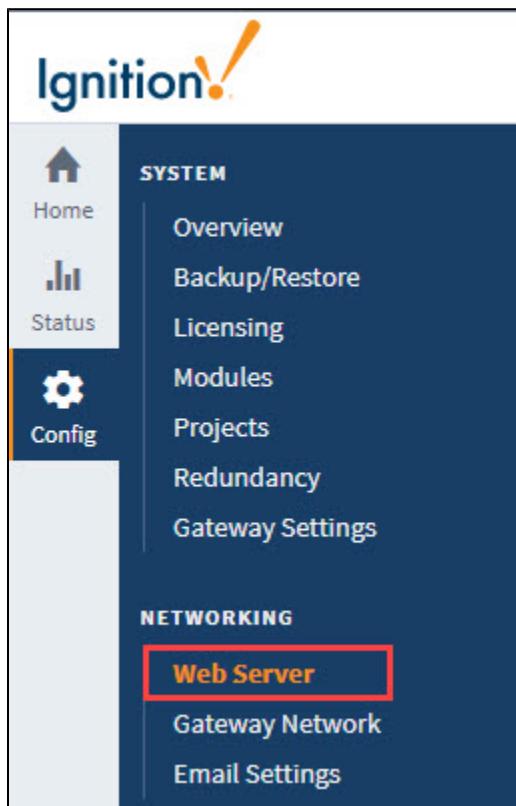
Email sent in 2 seconds

< Back to Email Profiles

Web Server Settings

The Web Server page is for configuring the HTTP and HTTPS ports, enabling SSL/TLS, redirecting traffic through a known address, and specifying whether or not all HTTP traffic should be forcefully redirecting to HTTPS.

If you are allowing users to access your Gateway from outside your network (through the Internet), you will need to configure the Public HTTP Address settings.



On this page ...

- [SSL/TLS Settings](#)
 - [HTTP and HTTPS Settings](#)
 - [Public HTTP Address settings](#)
- [Cipher Support](#)

SSL/TLS Settings

On the Web Server screen you can view details of an SSL certificate details, export keys, remove the installed SSL certificate, and transition to a CA signed certificate.

From the Gateway Webpage, click on **Config > Networking > Web Server**. From the Web Server page, click on the **View Details** button. See [Adding a Signed Security Certificate](#) process for more information on enabling SSL/TLS and installing security certificates.



SSL / TLS ENABLED

CA-Signed SSL Certificate installed.

SSL / TLS

SSL / TLS
Certificate

[View Details](#)

View the certificate details, export keys, remove the installed SSL certificate, and transition to a CA-signed Certificate

HTTP Settings

HTTP Port

8088

The port to which Ignition will listen for incoming HTTP traffic. *Example: 8088*

HTTPS Settings

On the Certificate Details page, you can also generate a Certificate Signing Request (CSR) by clicking the Generate CSR button in the upper right.

 **SSL / TLS ENABLED**
CA-Signed SSL Certificate installed.

Active SSL Certificate	
Subject	US California Folsom Inductive Automation inductiveautomation.com
Issuer	US California Folsom Inductive Automation Ignition Test Intermediate Certificate Authority
Version	3
Signature Algorithm	SHA256withRSA
Not Valid After	Wed Oct 23 2019 08:47:17 GMT-0700 ⚠ Expires Soon
Not Valid Before	Mon Jul 15 2019 08:47:17 GMT-0700

→ [Upload Trusted CA-signed SSL Certificate...](#)

→ [Return to Web Server...](#)

For more information, see [Secure Communication \(SSL / TLS\)](#).

HTTP and HTTPS Settings

HTTP Settings	
HTTP Port	The port Ignition will listen for incoming HTTP traffic, for example: 8088.
Use Proxy Forwarded Headers	<p>The following feature is new in Ignition version 8.1.10 Click here to check out the other new features</p> <p>When enabled, the Gateway inspects each incoming HTTP request in search for headers that indicate it has been forwarded by one or more proxies. If these headers are present, then the request is updated so that the proxy is not seen as the other end point of the connection from which the request originated.</p>

	<p>Caution: Enabling this setting when users can directly connect to the Gateway is a security risk. This setting is intended to be used in scenarios where untrusted users will not be able to bypass a trusted proxy that is responsible for setting the appropriate headers.</p> <p>For more information, see Use Proxy Forwarded Headers Explained.</p>
Resolve Client Hostnames	<p>The following feature is new in Ignition version 8.1.10 Click here to check out the other new features</p> <p>When enabled, Ignition's web server will attempt to resolve the remote HTTP client's hostname by performing a reverse DNS lookup using the remote HTTP client's IP address where appropriate. Enabling this setting could have a performance impact as the Gateway may attempt an expensive hostname lookup when handling requests. When disabled, Ignition's web server will not attempt to resolve hostnames and any queries for the remote HTTP client's hostname will result in their IP address instead.</p> <p>When enabling this setting, it is highly recommended that reverse DNS is configured to prevent host lookup failures. This includes configuring valid mappings from user IP addresses to their hostnames, and from user hostnames back to their IP addresses. If reverse DNS isn't configured, then DNS queries could block certain requests until the queries time out (default of 10 seconds).</p>
HTTPS Settings	
HTTP S Port	The port Ignition will listen for incoming HTTPS traffic, for example: 8043.
Force Secure Redirect	When enabled, and if SSL / TLS is enabled, all http traffic will be redirected to its https counterpart. (Default: disabled)
Included Cipher Suites	Whitelist of included cipher suites for clients connecting to Ignition using SSL/TLS.
Excluded Cipher Suites	Blacklist of excluded cipher suites for clients connecting to Ignition using SSL/TLS. Takes precedence over allowed cipher suites.
<p>Note: The Excluded and Included Cipher Suites settings only apply to the port specified in the HTTPS Port field. If you need to Whitelist or Blacklist cipher suites for port 8060, refer to the Gateway Network > HTTPS Settings information on the Gateway Configuration File Reference page.</p>	

HTTP and HTTPS Connectors Restart

Certain actions will cause the HTTP port and/or the HTTPS port to restart. Refer to the following table for details.

Configuration Change	HTTP Port Restarted?	HTTPS Port Restarted?
HTTPS Port	Yes	Yes
HTTP Port	Yes	Yes
Force Secure Redirect	No	Yes
User Included Cipher Suites	No	Yes
User Excluded Cipher Suites	No	Yes
SSL/TLS Setup	No	Yes

Use Proxy Forwarded Headers Explained

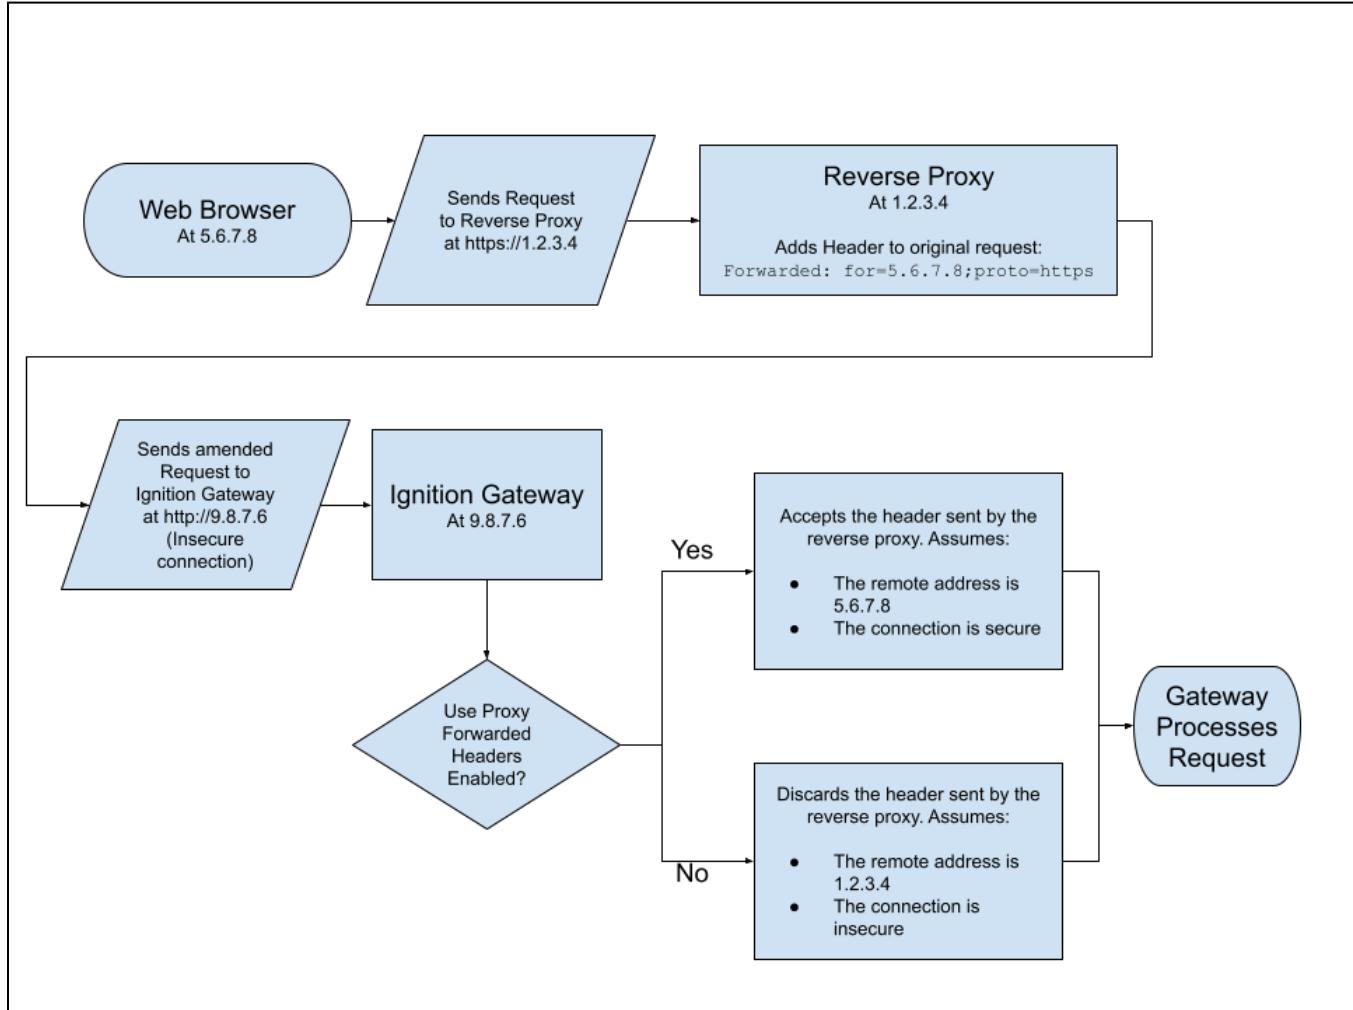
While enabled, the Gateway's web server will look for request headers mentioned on this page: [ForwardRequestCustomizer](#). Depending on what headers come in, the web server will alter its view of the remote client's connection on the incoming http request. The following is a list of which parts of the request that can be altered, although it's not exhaustive:

- The remote HTTP client's IP address

- The remote HTTP client's port
- The scheme used by the remote HTTP client when connected to the Gateway through one or more proxies (i.e., http/https)
- Whether or not the connection is considered secure
- The host/IP and port that the remote HTTP client used to connect to the Gateway through one or more proxies.

While this setting is enabled, if the Gateway does not see any of the mentioned headers, then the request will not be altered, effectively acting as if the setting is disabled.

The diagram below represents a request originating from a browser, and demonstrates how this setting can impact the request.



Public HTTP Address settings

If you are allowing users to access your Gateway from outside your network (through the Internet), you will need to configure the Public HTTP Address settings.

Public HTTP Address	
Auto Detect HTTP Address	To specify an explicit HTTP address that Vision Clients and Perspective Sessions will use, turn this off. Most users will leave autodetect on. (Default: enabled)
Public Address	The public facing address that Vision Clients and Perspective Sessions must use to connect. If Force Secure Redirect is enabled, redirected connections will use this address, for example: yourcompany.com .
Public HTTP Port	The public facing HTTP port that Vision Clients and Perspective Sessions must use to connect, for example: 80
Public HTTPS Port	The public facing HTTPS port that Vision Clients and Perspective Sessions must use to connect. If Force Secure Redirect is enabled, redirected connections will use this port, for example: 443

Cipher Support

Listed below are the supported cipher suites for both 8043 and 8060 TLS ports.

Cipher suites enabled by default:

```
TLS_AES_128_GCM_SHA256
TLS_AES_256_GCM_SHA384
TLS_CHACHA20_POLY1305_SHA256
TLS_DHE_DSS_WITH_AES_128_CBC_SHA256
TLS_DHE_DSS_WITH_AES_128_GCM_SHA256
TLS_DHE_DSS_WITH_AES_256_CBC_SHA256
TLS_DHE_DSS_WITH_AES_256_GCM_SHA384
TLS_DHE_RSA_WITH_AES_128_CBC_SHA256
TLS_DHE_RSA_WITH_AES_128_GCM_SHA256
TLS_DHE_RSA_WITH_AES_256_CBC_SHA256
TLS_DHE_RSA_WITH_AES_256_GCM_SHA384
TLS_DHE_RSA_WITH_CHACHA20_POLY1305_SHA256
TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA256
TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256
TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA384
TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384
TLS_ECDHE_ECDSA_WITH_CHACHA20_POLY1305_SHA256
TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA256
TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256
TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA384
TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384
TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305_SHA256
TLS_ECDH_ECDSA_WITH_AES_128_CBC_SHA256
TLS_ECDH_ECDSA_WITH_AES_128_GCM_SHA256
TLS_ECDH_ECDSA_WITH_AES_256_CBC_SHA384
TLS_ECDH_ECDSA_WITH_AES_256_GCM_SHA384
TLS_ECDH_RSA_WITH_AES_128_CBC_SHA256
TLS_ECDH_RSA_WITH_AES_128_GCM_SHA256
TLS_ECDH_RSA_WITH_AES_256_CBC_SHA384
TLS_ECDH_RSA_WITH_AES_256_GCM_SHA384
TLS_EMPTY_RENEGOTIATION_INFO_SCCSV
```

Cipher suites disabled by default:

```
TLS_DHE_DSS_WITH_AES_128_CBC_SHA
TLS_DHE_DSS_WITH_AES_256_CBC_SHA
TLS_DHE_RSA_WITH_AES_128_CBC_SHA
TLS_DHE_RSA_WITH_AES_256_CBC_SHA
TLS_ECDHE_ECDSA_WITH_AES_128_CBC_SHA
TLS_ECDHE_ECDSA_WITH_AES_256_CBC_SHA
TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA
TLS_ECDHE_RSA_WITH_AES_256_CBC_SHA
TLS_ECDH_ECDSA_WITH_AES_128_CBC_SHA
TLS_ECDH_ECDSA_WITH_AES_256_CBC_SHA
TLS_ECDH_RSA_WITH_AES_128_CBC_SHA
TLS_ECDH_RSA_WITH_AES_256_CBC_SHA
TLS_RSA_WITH_AES_128_CBC_SHA
TLS_RSA_WITH_AES_128_GCM_SHA256
TLS_RSA_WITH_AES_256_CBC_SHA
TLS_RSA_WITH_AES_256_GCM_SHA256
TLS_RSA_WITH_AES_256_GCM_SHA384
```

In This Section ...

Secure Communication (SSL / TLS)

Inductive Automation recommends enforcing secure communication in Ignition using digital certificates. The Ignition Gateway web server can provide modern end-to-end security using Transport Level Security (TLS) technologies. This protects externally-originated connections such as: Perspective sessions, Vision clients, Designers, and Ignition web configuration. Users should be familiar with the browser padlock icon (secure session) from online banking, shopping, or medical portals. TLS assures users of the distant end identity and offers protection from attackers and eavesdroppers through strong encryption. This configuration is different from Gateway-originated outbound communication such as database and device (OPC UA) connections, alarming, and web services (REST) calls, which are secured separately and have configuration that depends on distant nodes.

Ignition versions 8.0.4 and later default to TLS versions 1.2 and 1.3 with a valid certificate. Older versions of Ignition should be upgraded to offer protection against known vulnerabilities. If you are looking for additional security settings, check out the [Ignition Security Hardening Guide](#) and the [Let's Encrypt Guide](#) for recommendations.

Enabling Force Secure Redirect

Normally, Clients, Sessions, Designers, and Web browsers that communicate with the Gateway will do so over an HTTP. However, you can force these communications to be redirected to the more secure HTTPS.

1. Go to the **Config** section of the Gateway Webpage.
2. Choose **Networking > Web Server** from the menu on the left.
3. Select the checkbox for **Force Secure Redirect**, and click the **Save** button at the bottom of the page.

HTTPS Settings	
HTTPS Port	8043 The port to which Ignition will listen for incoming HTTPS traffic. Example: 8043
Force Secure Redirect	<input checked="" type="checkbox"/> When enabled, and if SSL / TLS is enabled, all http traffic will be redirected to its https counterpart. (Default: disabled)
Included Cipher Suites	<input type="button" value="Add"/> Whitelist of included cipher suites for clients connecting to Ignition using SSL/TLS
Excluded Cipher Suites	<input type="button" value="Add"/> Blacklist of excluded cipher suites for clients connecting to Ignition using SSL/TLS. Takes precedence over allowed cipher suites.

When enabled, all requests between the host Gateway and any Clients, Sessions, Designers, or web browsers will be redirected to the HTTPS port (by default, port 8043), and thus encrypted. However, you will likely want to install a security certificate signed by a certificate authority.

Note that although the Gateway Network may share the same port as HTTP traffic, Force Secure Redirect settings will not apply.

Adding a CA Signed Security Certificate

We are not able to ship a real certificate with Ignition because security certificates have to be obtained individually from a Certificate Authority (CA). Ignition supports certificates from both your organization's internal CA, as well as commercial CAs (Verisign, GoDaddy, Comodo, etc.). In either case, the procedure for how to install a certificate is listed below.

Note: After you have added a certificate, the KeyStore will automatically refresh every 15 minutes. You can disable this in the `ignition.conf` file by altering the `ignition.ssl.refresh` entry (Set to 0 to not refresh).

Get a Certificate Signing Request

Since SSL/TLS requires the installation of a security certificate, filling out the form below will generate a certificate signing request (CSR) to provide to a certificate authority.

1. Go to the **Config** tab of the Gateway Webpage and choose **Networking > Web Server**.
2. You'll see a warning message indicating that SSL/TLS is not enabled.
3. Click on **Setup SSL/TLS**.

On this page ...

- [Enabling Force Secure Redirect](#)
- [Adding a CA Signed Security Certificate](#)
 - [Get a Certificate Signing Request](#)
 - [Demonstrate Proof of Ownership over your Domain](#)
 - [Install Security Certificates](#)
- [What is a Self-Signed Certificate](#)
- [How to Install a Self-Signed Certificate](#)



INDUCTIVE
UNIVERSITY

Requiring SSL

[Watch the Video](#)

**SSL / TLS NOT ENABLED**

SSL / TLS allows for secure communication with the gateway. In order to be enabled, it requires an SSL certificate to be installed.
[Click here](#) to be guided through the process of installing an SSL certificate in order to enable SSL / TLS.

SSL / TLS	
Setup SSL / TLS	Setup SSL / TLS
Guided process for installing an SSL certificate in order to enable SSL / TLS.	

4. Click on the **I don't have all the items above** button. The Create Certificate screen is displayed.
5. Fill in the required fields on the screen, then click the **Generate Certificate Signing Request** button. This can be brought to a Certificate Authority.

Basic Details	
Field	Definition
Common Name	Full DNS name (required). This is typically what you type in your browser URL bar in order to navigate to this Gateway, for example: yourdomain.com
Organization Name	Name of company (required). For example: Inductive Automation.
Organization Department	Department or section (required). For example: Engineering
Email	Email address. For example: your@email.com .
Country	Typically an ISO 3166 2 character code (required). For example: US
State / Province	State, province or region, for example: California
Locality (City)	Name of city. For example: Folsom
Street	Street number and street name. For example: 90 Parkshore Dr
Postal Code	Postal Code Example: 95630
Key Type	The algorithm of the key pair which will be generated for the self signed certificate. Options are RSA or EC. Recommended: RSA
Key Size	The strength of the generated Key. Recommended: 2048 bits
Expires in	The number of days the generated Certificate will be valid. This number will only apply to a self-signed certificate. The certificate authority will determine a final expiration date as needed for the CA signed certificate
Subject Alternative Names	
Field	Definition
IP Addresses	The IP addresses of all the servers you plan on installing the certificate. Click the Add button for each additional IP address. <div style="border: 1px solid #ccc; padding: 5px; background-color: #f9f9f9;"><p>This feature was changed in Ignition version 8.1.17: In versions 8.1.17 and newer, click the Add IP Address button for each additional IP address.</p></div>
DNS Names	DNS names which map to the list of IP addresses above. Click the Add button for each additional DNS name. <div style="border: 1px solid #ccc; padding: 5px; background-color: #f9f9f9;"><p>This feature was changed in Ignition version 8.1.17: In versions 8.1.17 and newer, click the Add DNS Name button for each additional DNS name.</p></div>

Demonstrate Proof of Ownership over your Domain

After providing a CSR to your CA and before your CA issues a new [domain-validated \(DV\) certificate](#), your CA will require you to demonstrate proof of control over the domain(s) used to access the Ignition Gateway Web Server. These domains should be set as the subject common name and/or subject alternative names in your CSR. Your CA will require you to perform additional actions and to provide additional information if you are seeking

an organization validation (OV) class certificate or an [extended-validation \(EV\) certificate](#). See the Wikipedia article on [Public Key Certificates](#) and consult your CA for more details.

Methods for determining control over a domain varies based on the CA. Typical methods include:

1. Responding to email sent to the email contact in the domain's whois details
2. Responding to email sent to a well-known administrative contact in the domain, e.g. (admin@, postmaster@, etc.)
3. Publishing a DNS TXT record
4. Publishing a nonce provided by an automated certificate issuing system

The first three methods are performed out-of-band as far as Ignition is concerned. Let's drill a little bit deeper into the fourth method. Typically, the CA will provide a randomly generated cryptographic string of data which may only be used once. [This value is called a "nonce"](#). The CA expects you to configure your web server to serve the nonce as a special HTTP resource referenced using a well-known path. Once you configure your web server to serve this resource, you give your CA permission to make an HTTP request with a URL containing the domain you are trying to validate as the hostname component of the URL (which should resolve to an IP address of the host running your web server), an agreed-upon port (usually standard HTTP port 80 or HTTPS port 443, though sometimes this is configurable), and the well-known path as the path component of the URL.

There exists software purpose-built to stand-up a temporary web server for the purpose of exposing the special nonce value to the CA. The idea is that this web server is only listening for HTTP requests until you confirm that the CA has issued an HTTP request to the server and has validated the nonce value as correct, then the web server is shut down. In some cases, using such software is not possible if the Ignition Gateway is already listening on whatever port is required by the validation method. In these cases, shutting down the Ignition Gateway for the purpose of starting up one of these other tools would cause potentially unacceptable downtime.

To work around the downtime issue mentioned in the previous paragraph, Ignition supports two open standard specifications which accomplish the fourth validation method mentioned above. The first is the ACME Protocol's "HTTP Challenge", [detailed in section 8.3 of RFC 8555](#) (ACME is the protocol used by [Let's Encrypt](#)). Ignition's web server will map HTTP requests with a path pattern of `/.well-known/acme-challenge/<challenge-nonce>` to files located at `$GATEWAY_HOME/.well-known/acme-challenge/<challenge-nonce>` (where `$GATEWAY_HOME` is the path to the home directory of your Ignition Gateway installation). No Gateway restart is required to start serving new files recently added to the above directory to the web. The `challenge-nonce` placeholder is the nonce value provided by the ACME server which should be base64-url-encoded.

The second method is defined as the "Agreed-Upon Change to Website v2" process and procedure of validation of domain authorization or control, [detained in section 3.2.2.4.18 of version 1.81 of the CA / Browser Forum's Baseline Requirements for the Issuance and Management of Publicly-Trusted Certificates](#). Ignition's web server will map HTTP requests with a path pattern of `/.well-known/pki-validation/<challenge-file>` to files located at `$GATEWAY_HOME/.well-known/pki-validation/<challenge-file>`. Again, no Gateway restart is required to start serving new files recently added to the above directory to the web. The `challenge-file` placeholder in this case is arbitrary—check with your CA for more detailed instructions as to what the value should be.

For both methods above, it is best practice to stop exposing the challenge value to the web by deleting the file on the Gateway filesystem as soon as the CA is done verifying your control over the domain.

These standard challenge mechanisms can be used to automate certificate renewal—in fact, that's what the ACME Protocol and Let's Encrypt is all about. To learn more about how to use Ignition with ACME and Let's Encrypt, see the [Let's Encrypt Guide for Ignition](#).

Install Security Certificates

Once you have an SSL certificate, it needs to be added to Ignition.

1. Go to the **Config** tab of the Gateway Webpage and choose **Networking > Web Server**.
2. You'll see a warning message indicating that SSL/TLS is not enabled. Click on the [Click here](#) link.
3. The Setup SSL/TLS screen is displayed. Review the following list:
 - Private Key
 - Certificate Signed By A Certificate Authority (CA)
 - Any Intermediate CA Certificates (Provided by your CA)
 - Root CA Certificate (Provided by your CA)
4. If you have the items, click on the [I have all the items above](#) button. If you don't have all the items, click on the [I don't have all the items above button](#), and follow the previous procedure, [Get a Certificate Signing Request](#).

Setup SSL / TLS

Securing web communications requires the installation of a *SSL Certificate* and requires the following items:

- Private Key
- Certificate Signed By A Certificate Authority (CA)
- Any Intermediate CA Certificates (Provided by your CA)
- Root CA Certificate (Provided by your CA)

Do you currently have all items listed above? If you don't, we can create a certificate signing request (CSR) to provide to a CA for signing and a private key.

I have all the items above

I don't have all the items above

Return

5. The Certificate Wizard is displayed. The first step is to import your private key in one of the following three ways.

- Drag and Drop your certificate from your computer onto the screen.
- Click anywhere on the grey box to browse for the private key.
- Click **Manually enter data** button to type in the private key information

The screenshot shows the 'Certificate Wizard' step 1: Private Key. The top navigation bar includes 'Config > Network > Web Server > Certificate Wizard'. A note box states: 'Note: If you upload your own private key, it will be done over an unsecured network. It is important to understand the risks and mitigations before proceeding. The certificates provided by a CA build a chain of trust. Follow the steps below in order to build this chain of trust.' Below the note, a red box highlights the '1) Private Key*' label. To its right is a large dashed box containing a central download icon with the text 'Drag and drop or click to browse' and 'Manually enter data' below it. To the left of the dashed box is a text area: 'The DER or PEM encoded private key in SSLeay or PKCS8 format. Provide a password if the private key is encrypted.' At the bottom of the page are 'Cancel' and 'Continue' buttons.

6. If the private key is encrypted, click the **checkbox** to enable a password for this certificate and enter the password in the field. Click **Continue**.
7. The next step is to import the server certificate. This is the The DER or PEM encoded X.509 SSL Certificate that Ignition will use for SSL / TLS. Drag and drop the certificate file, browse for it, or manually enter the data.

Note: If you upload your own private key, it will be done over an unsecured network. It is important to understand the risks and mitigations before proceeding. The certificates provided by a CA build a chain of trust. Follow the steps below in order to build this chain of trust.

1) Private Key*
The DER or PEM encoded private key in SSLeay or PKCS#8 format. Provide a password if the private key is encrypted.

2) Server Certificate*
The DER or PEM encoded X.509 SSL Certificate that Ignition will use for SSL / TLS.

Private Key Successfully Uploaded

Drag and drop or click to browse
Manually enter data

Cancel **Continue**

8. The next step is to import the certificate chain. This gives you the Intermediate CA Certificate. Drag and drop the certificate file or bundle, browse for it, or manually enter the data. You'll see a message that the Intermediate CA Certificate was successfully uploaded.
9. Finally, import the root CA certificate: Drag and drop the certificate file, browse for it, or manually enter the data. You'll see a message that the Root CA Certificate was successfully uploaded.
10. Click the **Continue** button.
11. You'll see a confirmation message that the certificate is installed and SSL/TLS is enabled.

SSL / TLS ENABLED
CA-Signed SSL Certificate installed.

SSL / TLS

SSL / TLS Certificate	View Details
View the certificate details, export keys, remove the installed SSL certificate, and transition to a CA-signed Certificate	

HTTP Settings

HTTP Port	8088
The port to which Ignition will listen for incoming HTTP traffic. <i>Example: 8088</i>	

12. If you have a redundant installation, you'll need to repeat this procedure on your backup server.

What is a Self-Signed Certificate

SSL/TLS can still be setup without a CA signed certificate by installing a self-signed certificate. This gives you the encryption benefits of SSL, but not the identity validation, and it isn't a "real" certificate. When a self-signed certificate is used to enable SSL/TLS, a web browser will display warnings to users that they shouldn't trust the website. Users will then have to choose to proceed to the "unsafe" connection through Advanced options.



Your connection is not private

Attackers might be trying to steal your information from (for example, passwords, messages, or credit cards). [Learn more](#)

NET::ERR_CERT_AUTHORITY_INVALID



To get Chrome's highest level of security, [turn on enhanced protection](#)

[Hide advanced](#)

[Back to safety](#)

This server could not prove that it is its security certificate is not trusted by your computer's operating system. This may be caused by a misconfiguration or an attacker intercepting your connection.

[Proceed to \[REDACTED\] \(unsafe\)](#)



How to Install a Self-Signed Certificate

1. Go to Config section of the Gateway Webpage.
2. Choose **Networking >> Web Server** from the menu on the left.
3. Under SSL/TLS click on **Setup SSL/TLS**.
4. Select "**I don't have all the items above**".
5. Under the Basic Details section, fill in all of the required fields (*). This includes the Common Name, Organization Name, Organization Department, and Country.



The Common Name is typically what you input into your browser URL bar in order to navigate to this gateway. If you do not have a dedicated domain for your gateway, use the IP address.

6. Under Subject Alternative Names , enter in the IP Address along with any DNS associated with your gateway server.
7. Select **Show Advanced Properties** and click on **Install Self-Signed Certificate**.

Show advanced properties

Advanced

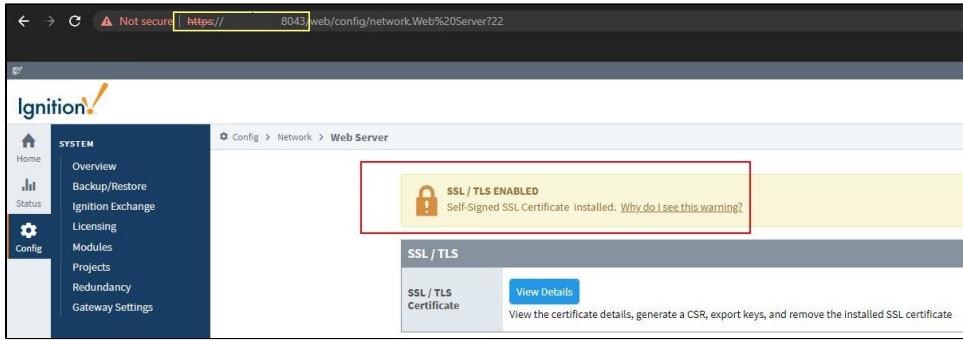
[Install Self-Signed Certificate](#)

[Install Self-Signed Certificate](#)

Installing a self-signed certificate will allow for encrypted communications with the Gateway. However, browsers will not trust the self-signed certificate since it will not be signed by a trusted Certificate Authority (CA). It is recommended to generate a certificate signing request (CSR) to use with a trusted Certificate Authority (CA) of your choosing in order to install a trusted CA-signed certificate.

8. This will take you to a page that displays the privacy warning that all connections will now see.

9. SSL/TLS is now enabled on your Gateway. The URL for your gateway now uses HTTPS and port 8043, which is the default port for SSL.



Related Topics ...

- [Security](#)
- [Security in Perspective](#)
- [Security in Vision](#)

Gateway Backup and Restore

Gateway Backups vs Project Exports

There are two main types of backups available in Ignition: Gateway backups and Project exports. Gateway backups are all inclusive while Project exports are simply a backup of an individual project. This page discusses making Gateway backups and restoring from a Gateway backup. To learn more about project backups and restoring projects, refer to the [Project Export and Import](#) page.

Gateway Backup and Restore

Creating Gateway backups and restoring from a Gateway backup are super easy in Ignition. Gateway backups are all inclusive, and typically takes less than a minute to run. It includes everything you find in the Ignition Gateway Webpage. Everything gets backed up - all your projects, Gateway settings, authentication profiles, Tags, database connections, OPC and device connections, alarm pipelines, scripts, sequential function charts, reports, and Image Management Library (i.e., png, jpg, and jpeg files). The only data that is not included in a Gateway backup is data stored in other programs such as SQL databases, PLC programs, other files you manually added to the install directory, and any additional files you may be using. This information needs to be backed up separately.

Perspective custom assets such as fonts, icons, and themes are included in a Gateway backup. In addition, any custom assets in a backup file will be restored when performing a Gateway restore.

Note: The [Tag Reference Store](#) made available in 8.1.34 is not included in Gateway backups. To save the recorded tag reference data, access the data\diagnostics\tags folder within the Data directory of your Ignition installation.

You have the option of creating a Gateway backup and restoring a backup from the Gateway Webpage or [Gateway Command-line Utility - gvcmd](#). We recommend using Gateway Webpage since it's much easier.

Scheduled Gateway Backups

In addition to manually taking a gateway backup, an Ignition Gateway can be configured to manually collect gateway backups on a schedule. See the [Gateway Settings page](#) for more details.

Contents of a Gateway Backup

A Gateway backup represents a number of resources and configurations to apply during restoration. Restoring from a Gateway backup replaces the following resources on the target Gateway:

- Database connections
- JDBC drivers
- Gateway settings
- Web server settings
- Gateway network settings
- Email settings
- Auditing settings
- User Sources
- IdP configurations
- Security Levels
- Security Zones
- Alarming settings
- Alarm Journal profiles
- Alarm notification settings
- On-call rosters
- Alarm schedules
- SECS/GEM equipment configurations and module settings
- Tag history providers
- All Tag providers, and the Tags, UDT definitions, and Tag groups within
- Device connections
- OPC UA client connections
- OPC UA Settings
- BACnet local device configurations
- EAM settings
- SFC settings
- All projects, and their resources within

Notably, when restoring a Gateway backup on an existing Gateway, the following items are **not** changed or replaced by the backup:

On this page ...

- [Gateway Backups vs Project Exports](#)
- [Gateway Backup and Restore](#)
 - [Scheduled Gateway Backups](#)
 - [Contents of a Gateway Backup](#)
 - [Manually Collecting a Gateway Backup](#)
 - [Gateway Restore](#)



INDUCTIVE
UNIVERSITY

Project Backup vs. Ignition Gateway Backup

[Watch the Video](#)



INDUCTIVE
UNIVERSITY

Making Ignition Gateway Backups

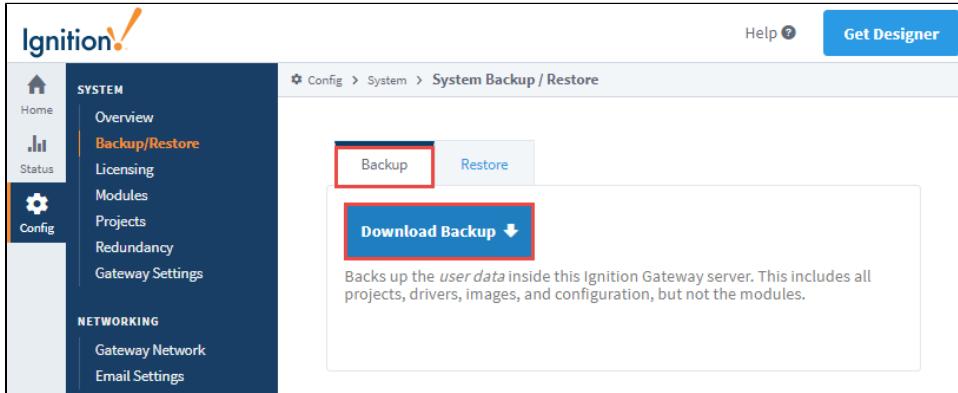
[Watch the Video](#)

- Modules
- License information/grants
- Redundancy settings

Manually Collecting a Gateway Backup

The easiest way to create a backup of the Gateway is using the Gateway Webpage.

1. Go to the **Config** tab of the Gateway Webpage, and click on **System > Backup/Restore**.
2. The System Backup/Restore page will be displayed. Make sure the **Backup** tab is selected, then click **Download Backup**.



3. By default, this downloads a **.gwbk** file extension to your local file system in your **Downloads** folder.

Note: The Gateway Backup default filename will look like **GatewayName_Ignition-backup-YYYYMMDD-HHMM.gwbk** where **YYYYMMDD-HHMM** is the timestamp of when it was created.

Command-line Utility

In Windows and Linux, you can use the command-line utility to create a Gateway backup. To run the Command-line Utility, open a shell and enter the command below.

Note: If you want to add a timestamp, you need to enter the date and time in the filename. Refer to the [Command-line Utility](#) page for a complete list of 'gwcmd' options.

```
gwcmd -b C:\Backups\Ignition\IgnitionBackup.gwbk
```

Gateway Restore

Restoring a Gateway backup is just as easy as backing it up and can also be done from the Gateway Webpage.

Caution: When you perform a Gateway Restore, **ALL** of the server's current configuration will be permanently lost. Restoring a Gateway backup overwrites all of the existing settings including your projects. There is no merge option for a Gateway backup. We recommend you always make a backup of the existing server immediately before performing a Gateway Restore.



Restoring Ignition Gateway Backups

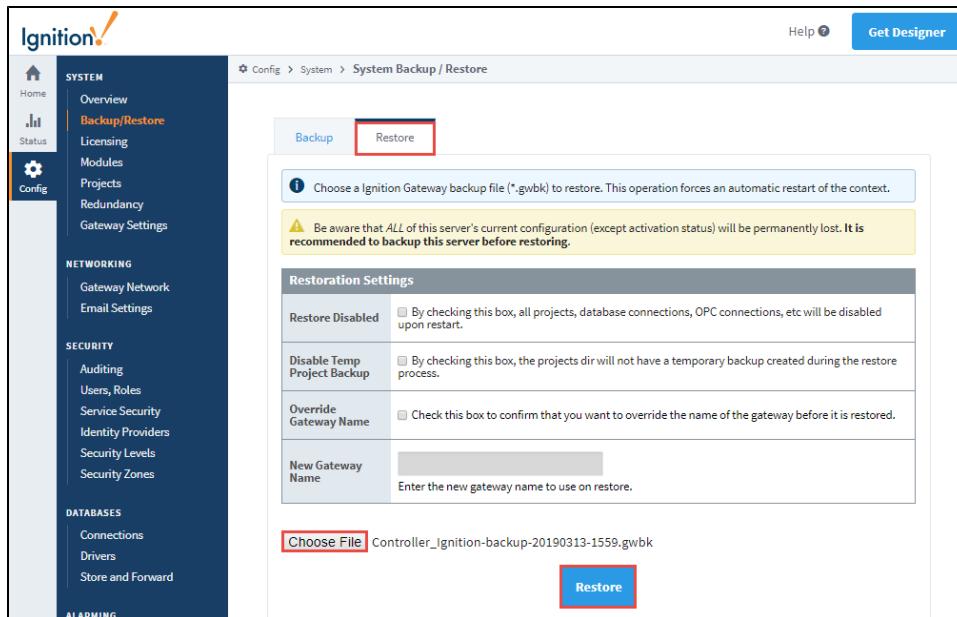
[Watch the Video](#)

1. Go to the **Config** tab and click on **System > Backup/Restore**.
2. The System Backup/Restore screen will be displayed. Click on the **Restore** tab. Read carefully through the restoration settings, and check the ones you wish to enable.

Restore Disabled Checkbox

If you check the **Restore Disabled** box, the imported resources will be disabled upon restoration.

3. Click **Choose File**, then navigate to your Gateway backup file (*.gwbk). By default, all your Gateway backup files are saved in your Downloads folder unless you select another folder location to choose an existing Gateway Backup file (*.gwbk) to restore.
4. Choose your Ignition backup file (*.gwbk), and click **Open**.
5. Click **Restore** at the bottom of the System Backup/Restore screen. The Gateway stops while restoring the backup file. When restoring is complete, the Gateway restarts itself to apply the restored settings. The Gateway Webpage is refreshed and your projects are loaded and your whole Gateway is restored in less than a minute.



The following table describes different Gateway restoration settings.

Gateway Restore	
Restore Disabled	Checkbox for toggling if the Gateway Backup file's contents such as projects, database connections, OPC connections, Tag Providers, etc will be disabled upon restore.
Disable Temp Project Backup	Checkbox for toggling if the projects directory will have a temporary backup created during the restore process.
Override Gateway Name	Checkbox for toggling if the name of the Gateway should be overridden before the Gateway is restored. Caution: It is not recommended to change the Gateway's System Name after initial Gateway setup. Renaming your Gateway may break critical dependencies, affect data collection, and halt production. To see a list of features dependent on the Gateway name, click here .
New Gateway Name	The name of the new Gateway. Requires "Override Gateway Name" to be enabled.

Command-line Utility

In Windows and Linux, you can use the command-line utility to restore a Gateway backup. To run the Command-line Utility, open a shell and enter the command below. Refer to the [Command-line Utility](#) page for a complete list of '**gwcmd**' options.

```
gwcmd -s C:\Backups\Ignition\IgnitionBackup.gwbk
```

Related Topics ...

- [Exporting and Importing Tags](#)

Ignition Exchange

In the Ignition Exchange you can access resources, templates, and tools that you can use in your own Ignition projects. The Ignition Exchange offers a wide variety of ways to store and access these resources. Ignition assets designed by others in various industries can be shared through the Exchange. This collection encompasses anything that can be built inside of Ignition including screens, graphics, templates, views, reports, alarm pipelines, scripting functions, database backups, projects, full systems, and Ignition demos.

As an individual, you can upload, browse, or download your own tools. If you'd like to keep these tools private, you can store them for your own use, or you can make the resource public, so anyone can access and use it.

Organizations or teams can use the Exchange to store and access private resources they want to make available to a select user group. This gives them the opportunity to collaborate and access company templates, and use them across corporate projects.

On this page ...

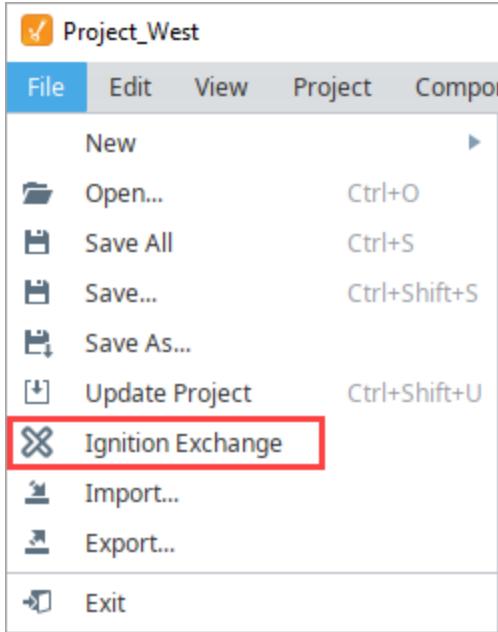
- [Access the Exchange](#)
- [Import an Ignition Exchange Package from the Gateway](#)
- [Import an Ignition Exchange Project Package from the Designer](#)
- [Upload a Resource to Ignition Exchange](#)
- [Restore a Gateway Backup from Ignition Exchange](#)

The screenshot shows the Ignition Exchange website. At the top left is the Ignition Exchange logo. Below it is a large, colorful graphic with overlapping blue, red, and purple arrows pointing to the right. To the left of the graphic, there is a callout box containing text about discovering, sharing, and downloading community-made Ignition resources. Below the graphic are three buttons: '+ New Resource' (blue), 'Sort by: New | Popular | Recently Updated' (grey), and 'Search' with a magnifying glass icon (grey). On the left side, there is a sidebar with filters: 'All Resources' (selected, highlighted in grey), 'My Resources', 'My Organization', 'Watched', and 'Downloaded'. In the center, there are three search results cards: 'SCRIPTS Advanced' (ICC Machine), 'PERSPECTIVE VIEW Intermediate' (Perspective Historical), and 'VISION TEMPLATE Beginner' (Vision Ad Hoc Trend). Each card has a small thumbnail icon and a brief description.

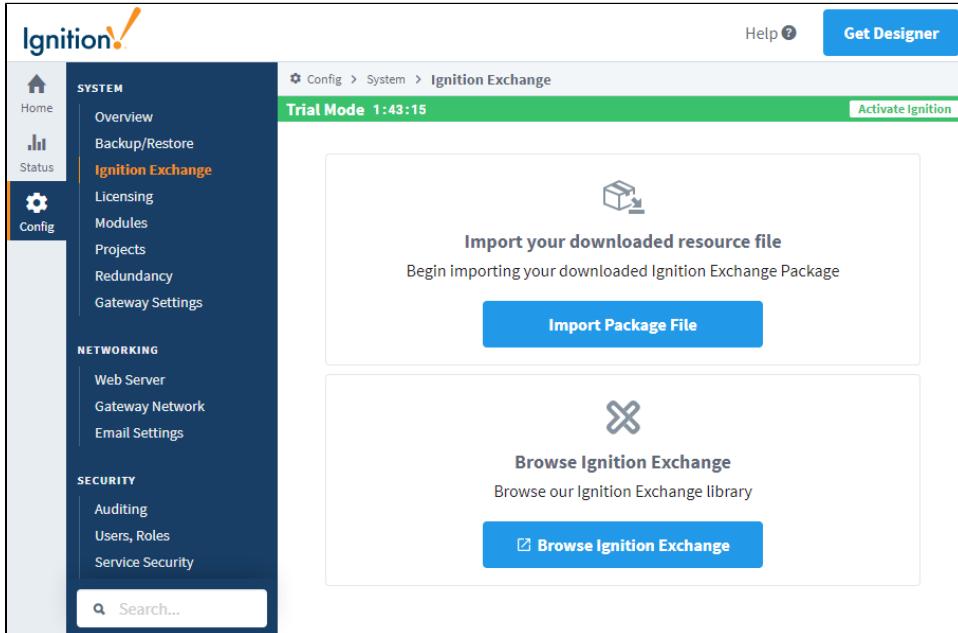
Access the Exchange

You can access the Ignition Exchange, in several ways:

- Go to <https://inductiveautomation.com/exchange/>.
- Within the Designer, select **File > Ignition Exchange**.



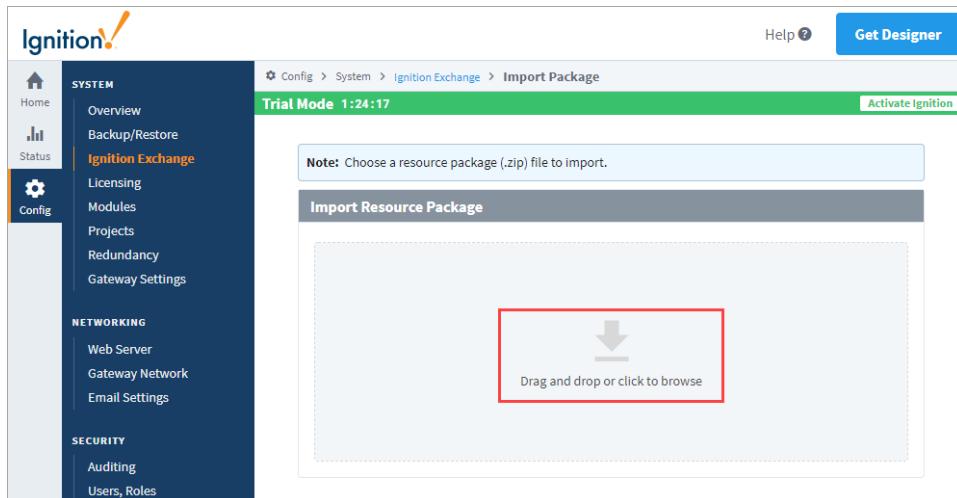
- Within the Gateway webpage, click on **Config > Ignition Exchange**. Then click **Browse Ignition Exchange**.



Import an Ignition Exchange Package from the Gateway

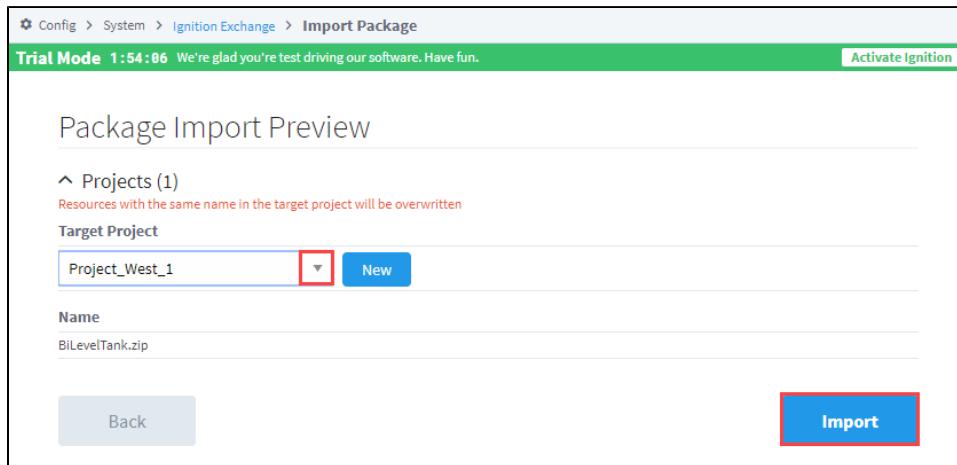
When you import a project from an exported file in the Gateway, it will be merged into your existing Gateway. To import a downloaded resource file:

1. Go to the Gateway, and click **Config > Ignition Exchange**.
2. Click on **Import Package File**.
3. Drag and drop your resource package file, or click to browse for the file you want to import and click **Open**.

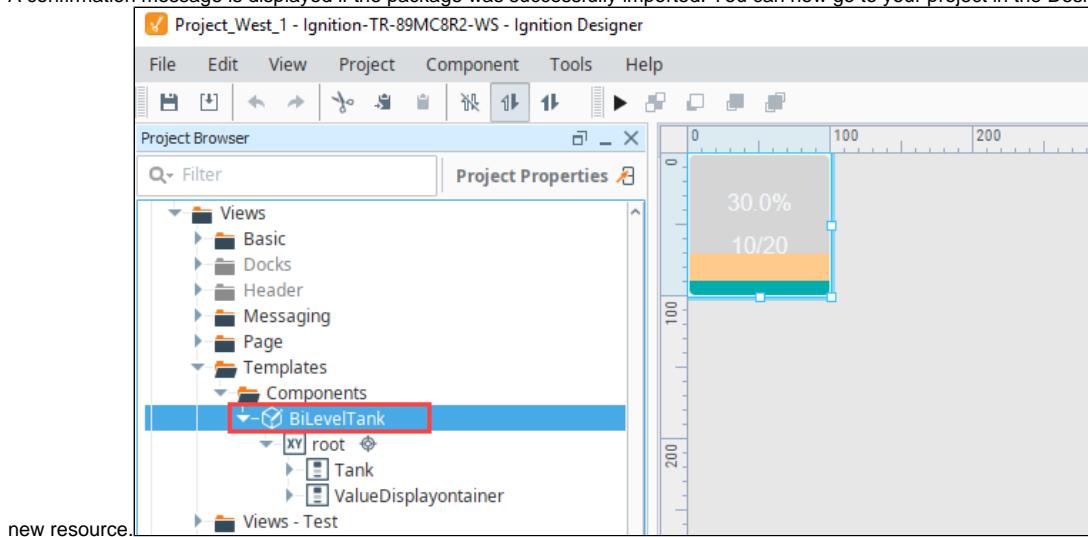


4. The Gateway will display a confirmation message if the package is valid. It will then display a README file with instructions. Click **Next**.
5. Use the down arrow in the Target Project box to choose a project. Then click **Import**.

Caution: Resources with the same name in the target project will be overwritten.



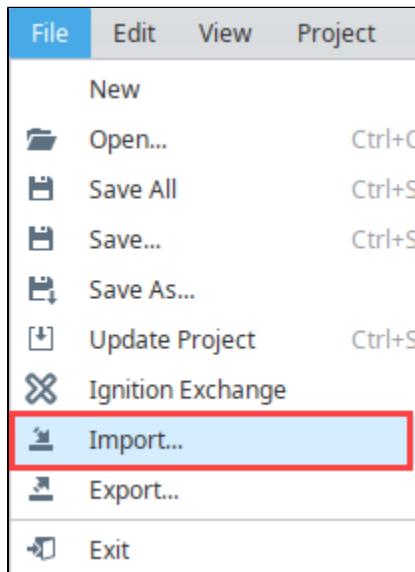
A confirmation message is displayed if the package was successfully imported. You can now go to your project in the Designer and see the



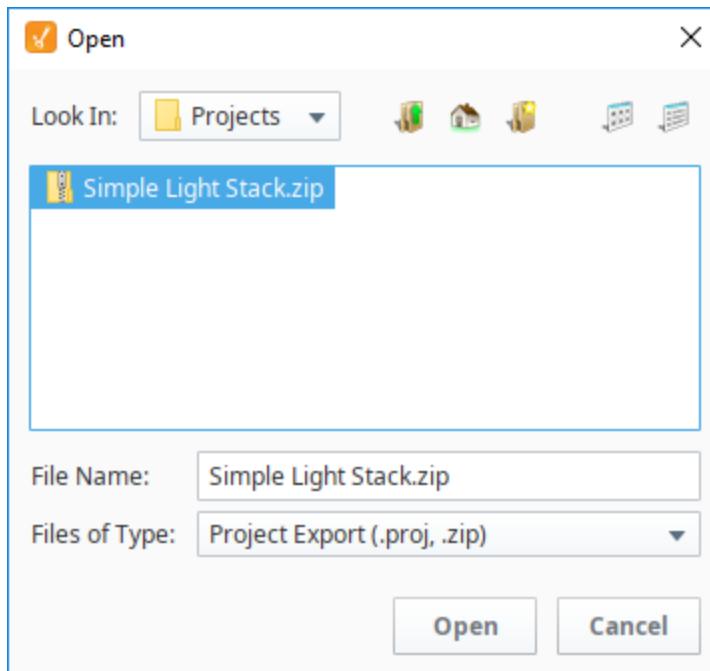
Import an Ignition Exchange Project Package from the Designer

In this example, we're going to import the Simple Light Stack project that we already downloaded from Ignition Exchange. Note that we also extracted the initial .zip file that was downloaded. This project contains a Perspective view with a light stack graphic. To import this view into a project that is open in the Designer, do the following.

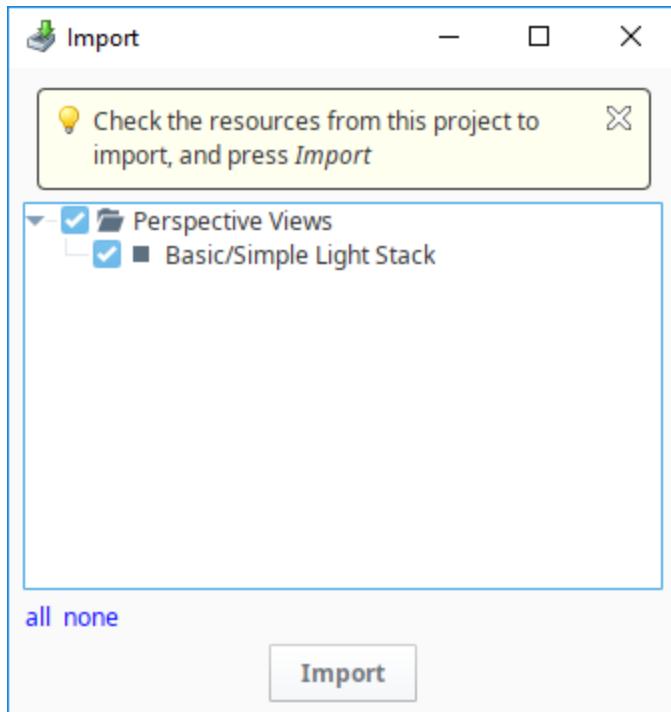
1. In the Designer, click **File > Import**.



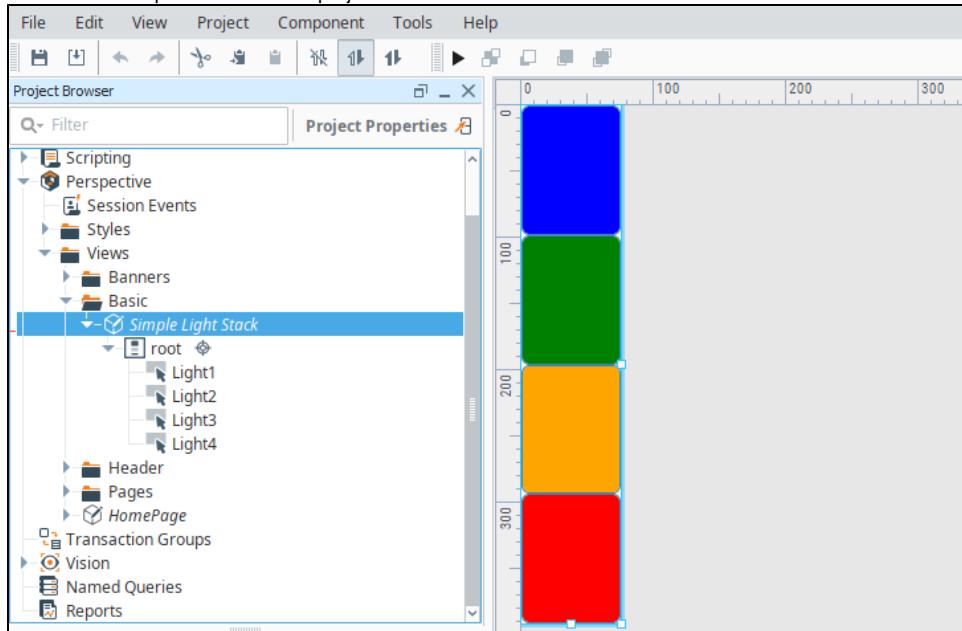
2. Navigate to the project file (.proj or .zip) you want to import, then click **Open**.



3. In the Import popup, select the resources that you want to import from the project. Click **Import**.



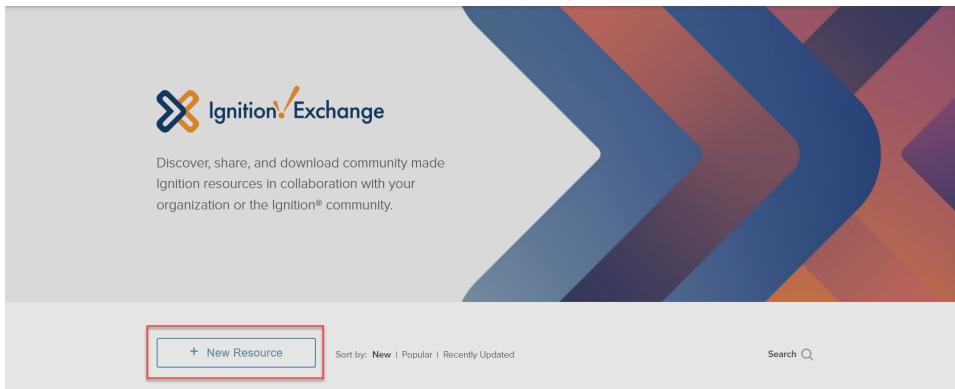
The view is now part of the current project.



Upload a Resource to Ignition Exchange

1. Use your credentials to log into your account. If you do not have an account, you can sign up at <https://account.inductiveautomation.com/create-account>.
2. Click on **+ New Resource** to open a popup that will start the resource upload draft.

Note: The **+ New Resource** option will only appear once signed in.



3. Fill out the **Overview** tab based on your resource requirements. The following table shows the description for each section of the Overview:

Section	Description
Visibility	Choose how your resource should appear on the Exchange.
Title	Create an easily identifiable title for your resource. Organization resources must have a unique title.
Tagline	Create a short but descriptive tagline to help identify your resource.
Description	Create a detailed description of your resource. Include information like how you envision it being used, what types of people might benefit from this package, any relevant industries, etc.
Resource Type	Choose a primary resource type that best describes your resource.
Skill Level	Choose a skill level that best describes the complexity of installation and use.
Background Image (optional)	Customize the background of your resource with a colorful design.
Category	Resources are displayed by category. You can select up to three.
Contact the Developer (optional)	This feature allows contributors to contact you via email directly from the resource page.
Tags (optional)	Tags make your resource more searchable. You can create up to 10 tags.
Images & Screenshots (optional)	Choose any additional images that showcase your resource in use. Be sure to follow the Exchange Guidelines .

4. Click **Save & Continue** when finished. This will take you to the Package tab.

Version	Release Tagline	Last Updated	Ignition Platform	Status
1	0	0	8.1	.30

Version will auto-increment by patch unless otherwise specified.

Choose the minimum version of Ignition required.

Release Tagline: Test release tagline

5. Fill out the **Package** tab based on your resource requirements. The following table shows the description for each section of the Package:

Section	Description
Version	Version will auto-increment by patch unless otherwise specified.
Ignition Platform	Choose the minimum version of Ignition required.
Release Tagline	Create a short but descriptive tagline to help identify what changes were made in this version.
Release Notes	List detailed release notes explaining what changed in this version. Think about what improvements were made and how they will impact the resource.
Required Modules (optional)	Choose any modules that are required when using this resource.
Maker Edition (optional)	Let people know this resource is Maker Edition compatible. Note: This option is unavailable for resources requiring unsupported Maker Edition modules. Learn more about Ignition Maker Edition™
Other Requirements (optional)	List any other external or custom requirements as separate line items.
Package Files	Choose a package file (up to 10mb) to upload. Note: Uploaded files will be packaged into a .zip format. A README file with your documentation will be created and added to the package.

6. Review your resource information and then click **Publish** to complete the resource upload.

Note: Use the **Update** option if you are not ready to publish yet or if you need to revise your published resource.

Your resource will be submitted to Ignition Exchange for a review process before becoming available to the community. The review process may take up to a few weeks and only resources available to the public will be reviewed. Make sure to adhere to the guidelines and provide accurate details to ensure a successful upload.

Restore a Gateway Backup from Ignition Exchange

Ignition Exchange can also have Gateway backup files as resources. After you've downloaded the file from Ignition Exchange, follow the steps for a Gateway Restore in [Gateway Backup and Restore](#).

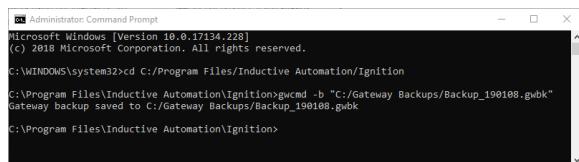
Caution: When you perform a Gateway Restore, **ALL** of the server's current configuration will be permanently lost. Restoring a Gateway backup overwrites all of the existing settings including your projects. There is no merge option for a Gateway backup. We recommend you always make a backup of the existing server immediately before performing a Gateway Restore.

Gateway Command-line Utility - gwcmd

The Gateway Command-line Utility provides a list of commands you can use to perform specific functions in the Gateway. The Gateway Command-line Utility or **gwcmd** provides basic commands, such as resetting the main password, changing the Gateway's port, or restarting the Gateway.

Invoking **gwcmd** can only be done from command line, so you'll need to utilize a command line interface of some sort (Power Shell, Terminal, etc). Because gwcmd is a file sitting in the Gateway's installation directory, these commands can only ever be invoked from where the Gateway is installed. Furthermore, interacting with gwcmd requires administrative privilege.

The gwcmd file sits at the root of the Gateway's installation directory. See the [Installing and Upgrading Ignition](#) page for more details on default installation directories.



```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.17134.228]
(c) 2018 Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>cd C:/Program Files/Inductive Automation/Ignition
C:\Program Files\Inductive Automation\Ignition>gwcmd -b "C:/Gateway Backups/Backup_190108.gwbk"
gateway backup saved to C:/Gateway Backups/Backup_190108.gwbk

C:\Program Files\Inductive Automation\Ignition>
```

On this page ...

- [Command-line Utility Options](#)
 - [Use the Command-line Utility to Start or Stop the Gateway](#)
- [Command-line Utility 'gwcmd' Options](#)
- [Gateway Password Reset](#)



A note to our legacy users...

Older versions of Ignition featured a visual Gateway Control Utility or GCU that could start and stop the Ignition service. This visual element of the GCU, as well as the ability to start and stop the service have since been removed in Ignition 8.0. For more information on starting or stopping the service, please see [Use the Command-line Utility to Start or Stop the Gateway](#) below.

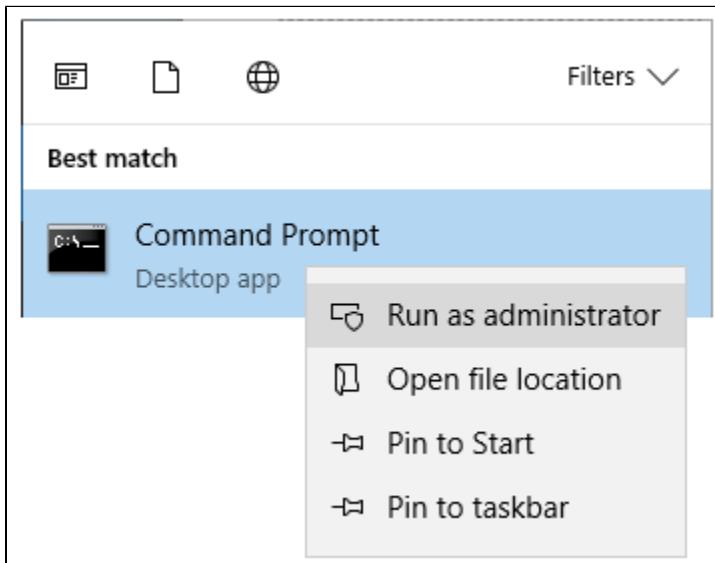
More information on the older version of the GCU can be found in [Deprecated Features](#) of the user manual.

Command-line Utility Options

The Gateway command-line utility supports Windows, Linux, and Mac OS platforms. The '**gwcmd**' only runs on the same machine as the Ignition Gateway and requires administrative privileges.

To run the Gateway Command-line Utility, open a command shell and type: **gwcmd <option>**. This example uses the Windows platform:

1. Open the **Command Prompt** with admin privileges. In the search bar, enter **cmd** then right click next to the **Command Prompt** to select **Run as administrator**.



2. Any time you run any of the Gateway command utility options, you need to run them from the directory that Ignition is installed in. The default install directory is: **C:/Program Files/Inductive Automation/Ignition**.
3. From here, you can enter any of the command options listed in the table below. This example is of a Gateway backup using the following command: **gwcmd -b "C:/Gateway Backups/Backup_190108.gwbk"**.



Using spaces in a file path

When using spaces in a file path name, use quotation marks around the full path name, as shown in the example above. Forward or backward slashes can be used to separate folders. The file path includes the disk name, folder path, and file name.

4. The above command created a Gateway Backup on the C drive, in the Gateway Backups folder with a file name of Backup_190108.gwbk. The image below shows all the commands used in Steps 2 and 3.

A screenshot of a Windows Command Prompt window titled "Administrator: Command Prompt". The window shows the following text:
Microsoft Windows [Version 10.0.17134.228]
(c) 2018 Microsoft Corporation. All rights reserved.
C:\WINDOWS\system32>cd C:/Program Files/Inductive Automation/Ignition
C:\Program Files\Inductive Automation\Ignition>gwcmd -b "C:/Gateway Backups/Backup_190108.gwbk"
Gateway backup saved to C:/Gateway Backups/Backup_190108.gwbk
C:\Program Files\Inductive Automation\Ignition>

Use the Command-line Utility to Start or Stop the Gateway

One of the common uses for the Command-Line Utility is to start or stop the Gateway.

Windows

Ignition's installation directory contains `start-ignition.bat` and `stop-ignition.bat`, which can start or stop the service. Example:

```
C:\Program Files\Inductive Automation\Ignition> start-ignition.bat
```

However, you can also use Windows native service commands to control the running state of the Gateway:

```
net start ignition
```

```
net stop ignition
```

Linux

You can control the service using the `ignition.sh` script. It can be called with the `start` and `stop` parameters to perform the relevant operations.

For example:

```
/usr/local/bin/ignition/ignition.sh start
```

Additionally, you can use native terminal commands to start or stop the service:

```
service Ignition-Gateway start
```

```
service Ignition-Gateway stop
```

Mac OS X

You can access the service from the install directory using the "ignition.sh" script. On a typical Mac install using the dmg installer, the full command (without a custom location specified) is the following:

```
/usr/local/ignition/ignition.sh start
```

Command-line Utility 'gwcmd' Options

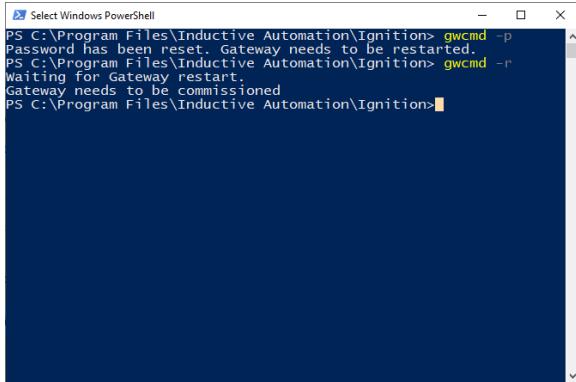
The following table lists all available **'gwcmd'** options.

Options	Description
<code>-a,--activate <license-key></code>	Creates an <code>activation_request.txt</code> file that can be used to request a <code>license.ipl</code> file from the Inductive Automation website. You must specify the license key to use for activation. The <code>activation_request.txt</code> file is saved in the current directory.
<code>-b,--backup <new filepath></code>	Downloads a Gateway backup <code>.gwbk</code> file and saves the file to the specified path. The path can be either an absolute path or a relative path. If another <code>.gwbk</code> file with the same name already exists, you will be prompted whether it is OK to overwrite the file. You can override with the <code>-y</code> option to force the file to always be overwritten.
<code>-c,--clearks</code>	Clears the gateway's SSL / TLS setup. The gateway's SSL / TLS connector will be immediately shut down.
<code>-d, --disabled</code>	Use with the <code>--restore</code> flag to disable all items after gateway restoration and restart.
<code>-e, --exportks <new filepath></code>	Exports the gateway's SSL KeyStore in PKCS #12 format and saves to the specified path.
<code>-f, --exportpk <new filepath></code>	Exports the private key from the gateway's SSL KeyStore in PEM format and saves to the specified path.
<code>-g,--reloadks</code>	Reloads the Gateway's SSL KeyStore from disk. Any update to the KeyStore will be automatically applied to any new connections.
<code>-h,--help</code>	Shows the usage for this command.
<code>-i,--info</code>	Retrieves server status and port information from the Gateway if it is running.
<code>-k,--port <new port></code>	Changes the Gateway http port.
<code>-l,--sslport <new port></code>	Changes the Gateway https port.
<code>-m, --skip-gateway-contact</code>	<p>The following feature is new in Ignition version 8.1.6 Click here to check out the other new features</p> <p>Use with the <code>--restore</code> flag to skip contacting a running gateway and stage the restore file directly.</p>
<code>-n, --nocrypt</code>	Add to the export private key command to not encrypt the private key.
<code>-o,--name <new gateway name></code>	Specifies a Gateway name while restoring a backup. Additionally, the <code>-y</code> command now skips prompts asking for a Gateway name override.
<code>-p,--passwd</code>	Enables a password reset command, which will allow you to create a temporary user that can access the gateway again. Requires a gateway restart to take effect. See Gateway Password Reset below.
<code>-r,--restart</code>	Restarts the Gateway.
<code>-s,--restore <backup file path></code>	Restores from a Gateway backup, using the file specified at the path.
<code>-t,--tdump</code>	Performs a thread dump in the Gateway and prints the dump to the command-line.
<code>-u,--unactivate</code>	Creates an <code>unactivation_message.txt</code> file that you can use to unactivate a license via the Inductive Automation website. The <code>unactivation_message.txt</code> file is saved in the current directory.
<code>-w,--uselicense <license.ipl path></code>	Applies a <code>license.ipl</code> file that was downloaded from the Inductive Automation website. You must supply the location of the <code>license.ipl</code> file. If it is in the current directory, use <code>license.ipl</code> for the location.
<code>-y,--promptyes</code>	Automatically answers <code>yes</code> to any prompt that may appear in the above commands, such as permission to overwrite an existing file.
<code>-z, --timeout <seconds></code>	<p>The following feature is new in Ignition version 8.1.17 Click here to check out the other new features</p>

Number of seconds to wait for a backup to be generated. Starting in 8.1.17, the default timeout value is 60 seconds. In older versions, the default timeout value is 30 seconds.

Gateway Password Reset

If you can no longer access the Gateway (due to, say, a forgotten password), you can use the `-p` command to cause a password reset. During a password reset, instead of just changing the initial user's password, a partial commissioning process will trigger upon the next Gateway restart (which can be accomplished with the `-r` command), allowing you to create a new user that can access the Gateway. From there you'll be able to address any issues that prevented you from using your normal credentials.



```
PS C:\Program Files\Inductive Automation\Ignition> gwcmd -p
Password has been reset. Gateway needs to be restarted.
PS C:\Program Files\Inductive Automation\Ignition> gwcmd -r
Waiting for Gateway restart...
Gateway needs to be commissioned
PS C:\Program Files\Inductive Automation\Ignition>
```



Password Reset with GWCMD

[Watch the Video](#)

The following feature is new in Ignition version 8.1.0
[Click here](#) to check out the other new features

However when performing this process, several things will happen to the Gateway:

- During commissioning, you'll be asked to provide a user name and password for a new user.
- A "temp" **user source** is created.
- The user you provided credentials for will be added to the "temp" user source.
- The new user will be assigned the role "Administrator".
- A "temp" Ignition Identity Provider will be created. The "temp" user source will be assigned as the provider backing the Identity Provider.
- On the **General Gateway Security Settings**, the following properties will be changed:
 - **System User Source** will be set to the "temp" user source.
 - **System Identity Provider** will be set to the "temp" identity provider.
 - **Gateway Config Permissions** will be set to the "Administrator" Security Level.

Thus, if you trigger a password reset and are able to use your normal credentials again, you'll want to make sure you change the values on the modified Gateway Security Settings to their property value. Also, you'll likely want to remove the "temp" user source and Identity Providers.

Licensing and Activation

How Licensing Works

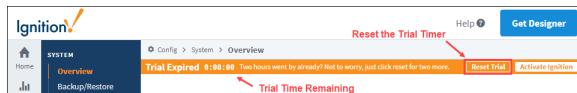
Ignition's licensing is unique and easy to use because Ignition is licensed by the server, not the client. You only need one license for your server; any clients are automatically included. In addition to that, an Ignition license is unlimited and sold based on which modules you want. There are unlimited clients, Tags, and projects. Buy only the modules you need, and don't worry about running into limits. If you want to test other modules, you don't need to do anything extra because of our built-in Trial mode. All unlicensed modules can be reset in 2 hour trial mode.

Trial Period

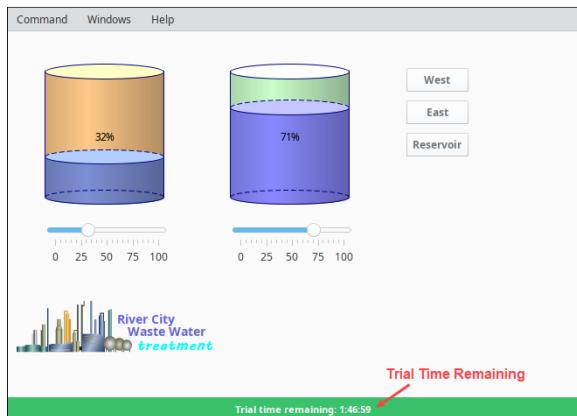
Our goal at Inductive Automation is to provide an easy way to access and learn Ignition. We want everyone to have the opportunity to try Ignition and start working with it with no restrictions. You can download Ignition from our webpage, install it, and start using it for two hours at a time. At the end of the two hour time period, all modules will stop running, but don't panic, you can reset the timer to run for another two hours. You can reset the timer as many times as you want, so go for it! The Gateway Webpage and the Designer are not affected by this trial, so you can develop for as long as you want without interruption. If you do have a license, any unlicensed modules will run in this Trial mode, but licensed modules will never timeout.

You can re-start the Trial period by logging into the Gateway, and clicking **Reset Trial** to enable another two hours of execution. The Trial Timer may be restarted any number of times. Depending on the module, you may need to take some additional actions. For example, the Vision Clients requires you to log out and back in again in order to continue the Trial.

In the Gateway, the Trial Time banner is displayed near the top of the screen. To reset the trial timer, click the Reset Trial button:



In a Vision Client (for Edge or Ignition Standard), the banner appears near the bottom of the screen.



On this page ...

- [How Licensing Works](#)
- [Trial Period](#)
- [Activating a Standard License \(Six-Character Key\)](#)
 - [Online Activation Example](#)
 - [Offline Activation Example](#)
- [Leased Licenses \(Eight-character keys\)](#)
 - [Leased License Activation](#)
- [Adding Multiple Licenses to a Single Gateway](#)
 - [Effective vs Applied Licenses](#)
- [Module Status](#)
 - [License Incomplete](#)
- [Updating a License](#)
 - [Online License Update](#)
 - [Offline License Update](#)
- [Unactivating a License](#)
 - [Online Unactivation](#)
 - [Offline Unactivation](#)
- [Emergency Activation](#)
 - [How to Activate in Emergency Mode](#)



INDUCTIVE
UNIVERSITY

About Licensing

[Watch the Video](#)



INDUCTIVE
UNIVERSITY

About the Trial Period

[Watch the Video](#)

Activating a Standard License (Six-Character Key)

When you purchase a license for Ignition Standard edition or Edge, you receive a license key, a six-digit code that identifies your purchase. Use this license key to activate the software online through the Ignition Gateway. If you later want to add any additional modules, your account is updated and you



INDUCTIVE
UNIVERSITY

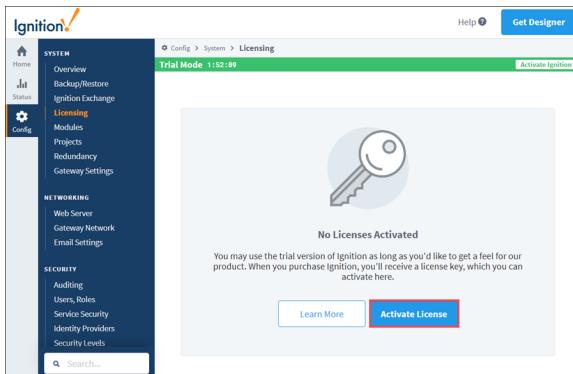
can re-use your existing license key to activate the new features. You can also deactivate your license key, and reuse it to activate Ignition on a different machine as many times as you want, allowing you to transfer a license from one Gateway to another.

You can activate your license in two ways:

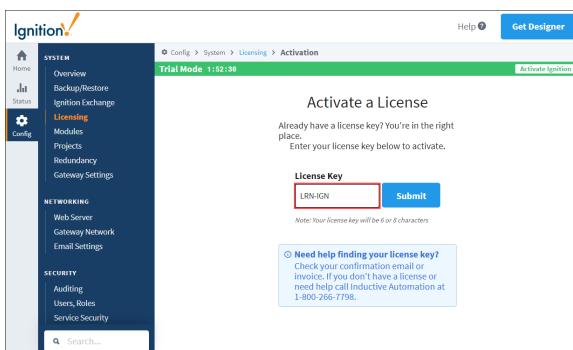
- **Online Activation** - from the **System > Licensing** section on the Gateway Webpage, your request to activate your Gateway is activated over the internet.
- **Offline Activation** - if you don't have an Internet connection, you can follow the manual activation process in the Offline Activation section.

Online Activation Example

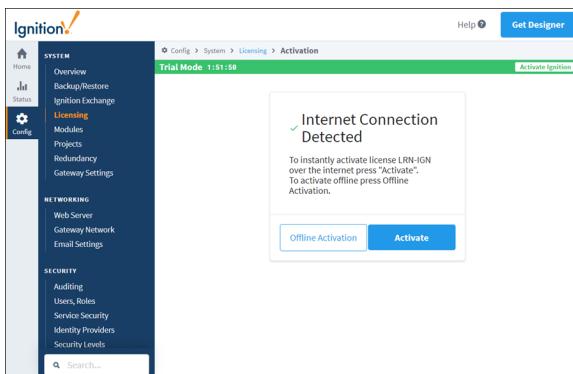
1. Go to the Gateway Webpage and select the **Config** section.
2. From the menu on the left, select **System > Licensing**. The Licensing page will appear. If you already have a license key, click the **Activate License** button.



3. The **Licensing / Activate Online** page will be displayed. Enter your **License Key** and click **Submit**.



- a. If you are connected to the Internet, click the **Activate** button.
- b. If you want to activate the license offline, click the **Offline Activation** button.



4. The Licensing page will refresh, and your Current License will be successfully activated.

Online Activation

[Watch the Video](#)



INDUCTIVE
UNIVERSITY

Offline Activations

[Watch the Video](#)

Offline Activation Example

When you do not have an internet connection, you can do the following steps to activate your license manually.

Note: If you're planning on performing an offline activation to update a license, make sure to [unactivate the license](#) first. Otherwise your system may enter [emergency mode](#) after attempting another activation.

1. Go to the Config tab on the Gateway Webpage.
2. Scroll down to **System > Licensing**. The Licensing Activation page appears.
3. Click on **Activate License**.
4. Enter your **License Key** and click **Submit**.

5. Click the **Offline Activation** button.

6. Click the **Download Activation Request**. An activation request file, called `activation_message.txt` is generated and downloaded.

The screenshot shows the Ignition software interface. On the left, there is a dark sidebar with various configuration options like Home, Status, Config (selected), SYSTEM, NETWORKING, SECURITY, DATABASES, ALARMING, and TWILIO. The main content area is titled "Offline Activation". It contains three steps: Step 1 (Download Activation Request), Step 2 (Transfer activation request file to a computer with internet access), and Step 3 (Upload license.ipl file). A red box highlights the "Download Activation Request" button in Step 1.

7. Take the `activation_message.txt` file to a machine with Internet access and go to <https://links.inductiveautomation.com/activation>.
 8. Click **Choose File** and select the `activation_message.txt` file.

The screenshot shows the Inductive Automation support website. At the top, there is a navigation bar with links for Call Us, Schedule a Demo, Product, Pricing, Resources, Partners, Support, About, and Download Ignition. Below the navigation bar, there is a search bar and a breadcrumb trail: Home / Support Home. The main content area is titled "Activate or Unactivate Ignition". It has a section for "Activate Ignition" with instructions to download the activation request file and upload it. A red box highlights the "Choose File" button, which is currently set to "activation_message (1).txt". Below it is an orange "Upload Activation File" button.

Note: Should any trouble occur with the upload process, please contact our Support department.

9. Once you upload the `activation_message.txt` file, a license file called `license.ipl` is generated.
 10. Bring the `license.ipl` file back to the computer on which you're licensing Ignition. Back on the Offline Activation page, upload the `license.ipl` file. Then click **Activate** to finish the process.

11. The Licensing page will refresh showing your license was successfully activated.
12. Click **View Modules** to see a list of all your activated modules.

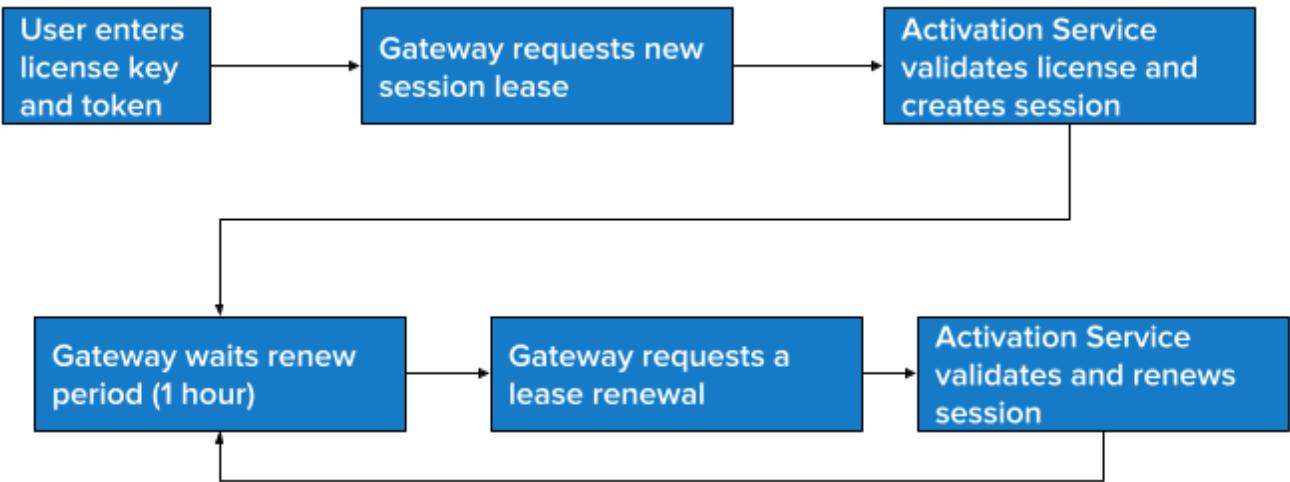
Leased Licenses (Eight-character keys)

A license key with eight characters signifies a **Leased License**. These types of licenses also require an **activation token**, which is a very long, unique, string of characters that must also be provided before a Gateway is licensed. Leased licenses are ideal for containerized environments (Docker, Kubernetes), as well as cloud-based deployments.

This type of license requires that the Gateway frequently check in with an activation service to validate the license.

Validation Feature	Description
Validation Timer	The Gateway will attempt to renew the lease every hour. A renewed lease will also reset the timeout period. This creates a buffer, allowing a Gateway to remain activated in cases where the Gateway is unable to reach the activation service in a timely manner.
Timeout Period	Leased sessions have a default timeout period of four hours. If a longer timeout period is needed for specific special applications, please discuss your requirements with your sales representative.
After Timeout	After the timeout period has expired, the Gateway reverts to trial mode.

In order to reach the activation service, the Gateway must have internet access at all times.



Leased License Activation

Activating a Gateway with a leased license follows a similar flow to the [standard license's online activation](#).

In addition, when selecting Maker Edition during installation, a leased license and activation token are required.



Activate License

Enter your Maker Edition license key here to continue this installation. If you don't already have a license, one can be generated online from your Inductive Account profile.

License Key

Your key will be 8 digits, AAAA-BBBB.

Activation Token

ⓘ Need help finding your key?

To generate a free Maker Edition key, sign in to your [Inductive Automation Account](#) online or create a profile. Within your account, navigate to the 'Licenses' page from the main menu.

Step 2 of 4

Next →

Adding Multiple Licenses to a Single Gateway

In order to better support our community of third-party module authors, we also allow for multiple license keys to be installed on a single Gateway. A third-party module author can issue a license key for their module directly to a customer, whereby, they can immediately do an install of the module. To learn more, check out the [Third-Party Module Showcase](#) to find and purchase modules that extend Ignition's functionality.

It is important to note that there may only be one license on it with a platform version per Gateway. The platform will look similar to the picture below. If you try and activate a second license with a platform onto a Gateway that already has a license with a platform, the new license will overwrite the previous license.

In the image below, two licenses have been applied to this Gateway, but only one is active.

The screenshot shows the Ignition Gateway Webpage. The left sidebar has a dark blue background with white text. It includes links for Home, Status, and Config (which is highlighted with an orange border). Under SYSTEM, there are links for Overview, Backup/Restore, Ignition Exchange, and Licensing (which is also highlighted). Under NETWORKING, there are links for Web Server, Gateway Network, and Email Settings. Under SECURITY, there is a link for Gateway Settings. The main content area has a light gray background. At the top, there's a green banner with the text "License Incomplete 0:47:49" and a "View Modules" button. Below the banner, the title "Effective License" is displayed, followed by a list item "+ Ignition". Then, the title "Applied Licenses" is shown, followed by two items: "+ LRN-IGN" and "+ 3PSW-F2P7". There are small blue icons next to each license entry. At the bottom of the content area, there's a link "→ Activate new license...".

Effective vs Applied Licenses

This page has two sections, which detail different information.

Applied Licenses shows all licenses that have been applied to the gateway. Entries here are typically purchased or given, and have a license key associated with them.

Effective Licenses provides a summary of modules active on the gateway. Modules are active if they're provided by an applied license. This area is the result of potentially multiple licenses across various activation mechanisms. For example, if **license A** provides **module A**, and **license B** provides **module B**, then Effective Licenses would show that **both** modules A and B are active on the gateway. The caveat here is that the modules will only appear if they're installed on the Gateway. In this same scenario with licenses A and B, if module B was not installed, then it will not appear under Effective Licenses.

In addition, the Effective Licenses area may have "synthetic" items that are translated to modules. For example, in the case of Edge gateways, their Applied Licenses will show products. These products will be translated to modules under the Effective License section.

Module Status

It's important to know which version of a module you have installed and which version of the module you are licensed for. You could have a license for another version of a module so it's not going to work correctly until you have the correct version of the license. To verify the versions on your installed modules, go to the **Modules** page in the **Status** section of the Gateway Webpage or by clicking **View Modules** in the green banner.

Here you can see all the modules that are currently running with their version numbers and license status. The licenses are either in Trial, Activated, or Free mode. On this page, you can add or remove modules from your Ignition. If you are licensed for a module you are running, it will run in Activated mode. If you are not licensed for a module you are running, it will run in Trial mode until the Trial time is expired.

For more information on installing new modules, see the [Installing or Upgrading a Module](#) page.

License Incomplete

Don't be alarmed by a 'License Incomplete' message on the green banner. Click on the View Modules link within the banner. It opens an informational message box showing modules that are installed. Some of the modules are in Trial Mode.

The screenshot shows the Ignition web interface with the following details:

- Banner:** License Incomplete ? 1:59:03
- Running Modules:** 19 / 22
- Licensed Modules:** 19 / 22
- Licensing Configuration:** A button labeled "View Modules" is visible.
- Table:** Inductive Automation, LLC

Name	Version	License	Status
Alarm Notification	5.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Allen-Bradley Driver	5.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
DNP3 Driver	3.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Enterprise Administration	3.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Logix Driver	4.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Modbus Driver	6.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
OPC-UA	8.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Omron Driver	3.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
OpcCom	5.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Perspective	1.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Reporting	5.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Vision	10.0.0-beta0 (b2019032202)	Activated	✓ RUNNING
Voice Notification	4.9.10 (b2018112821)	Trial	✓ RUNNING

Updating a License

Online License Update

If you added one or more new module to an existing license, then you'll need to update (Reactivate) the license. You update a license from the **Config > Licensing** section on the Gateway webpage. Simply click the Reactivate button next to the license that was changed, as shown below.

The screenshot shows the Ignition Config > Licensing page with the following details:

- Effective License:** Ignition
- Applied Licenses:**
 - LRN-IGN (with a red box around the "R" button)
 - Activate new license...



Reloading a License

[Watch the Video](#)

Pressing the Reactivate button will cause the gateway to attempt to reach out to our licensing server, and will update shortly after.

Offline License Update

If the gateway does not have internet access, it will be unable to update automatically. In this case you'll need to do the following:

1. Unactivate the license following the [Offline Unactivation](#) steps mentioned in the next section.
2. Activate the license via an [Offline Activation](#).

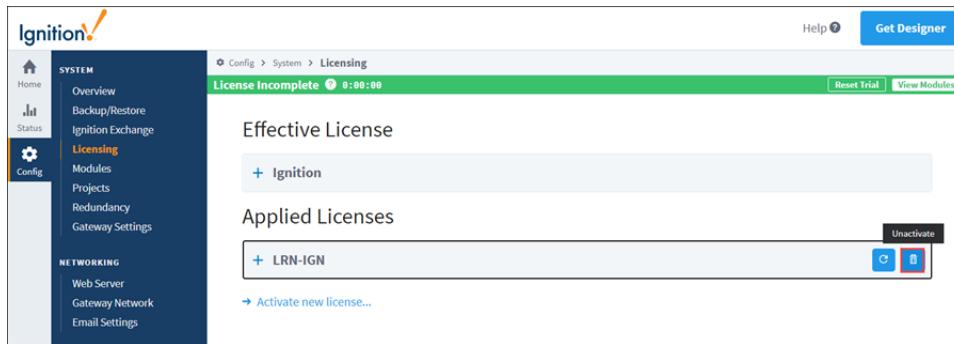
Unactivating a License

For a given license key, a limited amount of simultaneous activations are allowed at a given time. If you want to activate [Ignition](#) on a different [server](#), you must first unactivate it on the [current server](#). You can unactivate the license on one Gateway, and then activate it on a different Gateway if needed. Unactivation occurs immediately over the Internet, and makes this license available for activation on another machine.

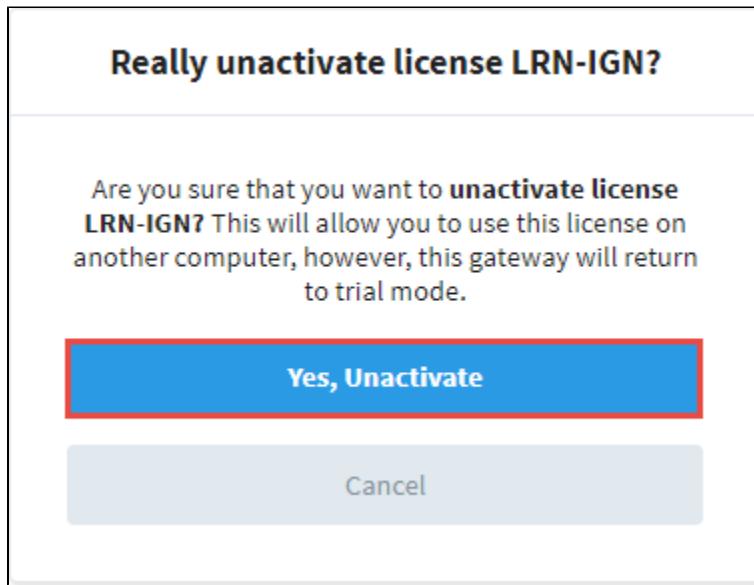
Online Unactivation

To unactivate the [Gateway](#), do the following steps:

1. From the **Config** section of [Gateway](#), go to **System > Licensing**. The **Licensing** page is displayed and you can see the currently installed license key.
2. Click on **Unactivate License** icon.



3. A Licensing / Confirm Unactive window will appear asking you to confirm the unactivation. Click the **Yes, Unactivate** button. It may take a minute or so for the request to finish.



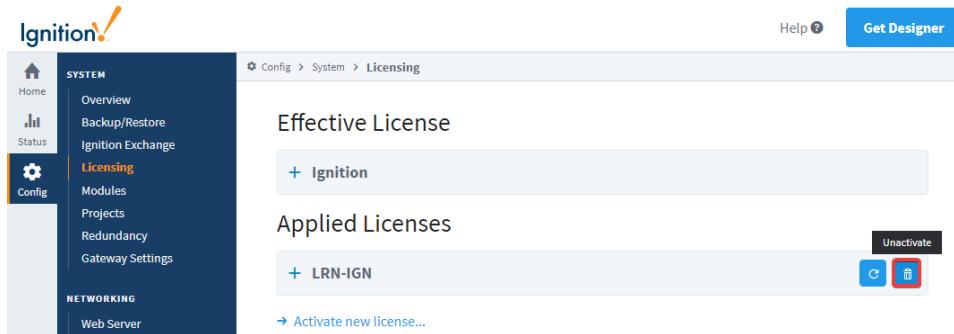
4. The unactivation request will be sent to [Inductive Automation's](#) licensing servers, and the [license](#) will again be available for activation on another [Gateway](#).

Offline Unactivation

In the event your Gateway is unable to reach our licensing server, you will need to perform an offline unactivation.

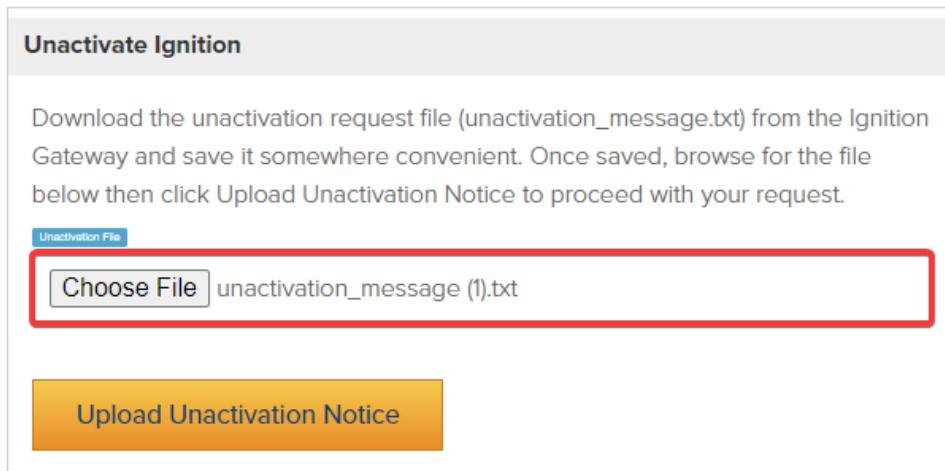
To perform an offline unactivation, do the following steps:

1. Go to your Gateway's **Config** section > **System** > **Licensing**. You will be able to see any licenses that are currently applied.
2. Click the **Unactivate** button (trash can icon) under **Applied Licenses**.



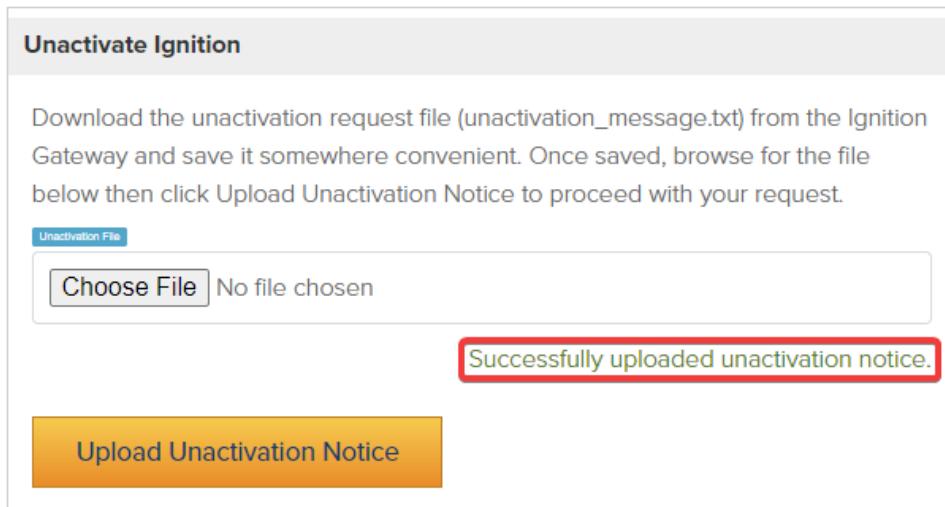
The screenshot shows the Ignition Configuration interface. On the left, there's a sidebar with 'Config' selected. The main area is titled 'Licensing'. It shows an 'Effective License' section with '+ Ignition' and an 'Applied Licenses' section with '+ LRN-IGN'. Below the applied license, there are two buttons: a blue 'Activate new license...' button and a black 'Unactivate' button with a trash can icon. The trash can icon is highlighted with a red box.

3. A confirmation window will appear, asking you to confirm the unactivation. Click the **Yes, Unactivate** button. It may take a minute or so for the request to finish.
4. An unactivation request file, called `unactivation_message.txt` will be generated and downloaded.
5. Take the `activation_message.txt` file to a machine with Internet access and go to <https://links.inductiveautomation.com/activation>.
6. Click **Choose File** and select the `unactivation_message.txt` file.



The screenshot shows a web-based form titled 'Unactivate Ignition'. It instructs the user to download the unactivation request file and save it somewhere convenient, then click 'Upload Unactivation Notice' to proceed. There is a 'Unactivation File' input field containing 'Choose File unactivation_message (1).txt', which is highlighted with a red box. Below it is a large orange 'Upload Unactivation Notice' button.

7. You will see a visual confirmation that your `unactivation_message.txt` file has been successfully uploaded, and the license key can now be activated on a different Gateway.



The screenshot shows the same web-based form after the file has been uploaded. The 'Unactivation File' input field now says 'Choose File No file chosen'. Below it is a green success message 'Successfully uploaded unactivation notice.' and a large orange 'Upload Unactivation Notice' button.

Emergency Activation

In cases where you may have a hardware or OS failure and you cannot unactivate a license, Ignition provides an [Emergency Activation mode](#). In this mode, you can temporarily activate your license for 7 days giving you enough time to contact [Inductive Automation Support](#).

How to Activate in Emergency Mode

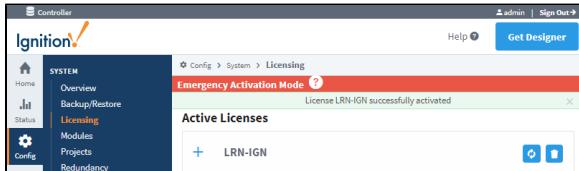
Activating your license in emergency mode is exactly the same as activating with your normal license, you don't have to do anything different because Ignition handles it all for you. See [License Activation](#) above for details on how to activate. The Gateway will know to run in emergency activation mode and it will display a timer stating how many days, hours, and minutes you have remaining in the Emergency Activation Mode banner. Any time before it expires, you can contact [Inductive Automation Support](#) to get your license fixed.

[In This Section ...](#)

Emergency Activation

What Is Emergency Activation?

After activating your License Key in Ignition, you may see a banner that says Ignition is in **Emergency Activation Mode**.



What Does this Mean?

This means the License Key that you used has already been activated. Either it was used previously on another computer, or something went wrong during activation and it registered twice. Fortunately, you have plenty of time to fix this problem. Emergency Activation Mode means **you are fully licensed for 7 days**, so even at night or over the weekend, you can still run without interruption until it is fixed.

On this page ...

- [What Is Emergency Activation?](#)
 - [What Does this Mean?](#)
 - [How Do I Fix it?](#)
 - [Why Does this Exist?](#)
- [Quick Disaster Recovery Plan](#)
 - [Emergency Restore Steps](#)

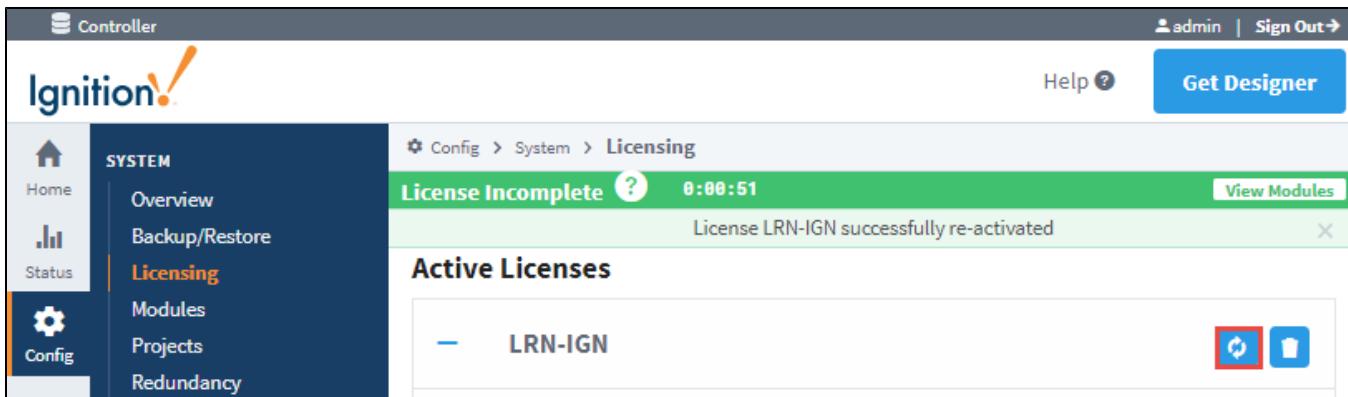


Emergency Activation

[Watch the Video](#)

How Do I Fix it?

In order to fix your License Key, you must [contact Inductive Automation](#). You can reach us by email or by phone and we will walk you through fixing your license. It is usually a very quick procedure that involves you re-activating a license after we fixed your License Key from our end. Once your license key is fixed, press the **Refresh** button and your license will be re-activated.



Why Does this Exist?

Basically, computers sometimes fail and you need to get running again, fast.

When you initially activate your copy of Ignition, it will run as long as your computer does. As we all know, sometimes there are hardware failures or your computer just stops working. There are many things you can do to minimize this, but once it happens you need to be able to get Ignition up and running again quickly. This Emergency Activation Mode allows you to do this without having to involve Inductive Automation until after your facility is back on track.

Quick Disaster Recovery Plan

If you already have Ignition installed and running in your facility, it's easy to get back to running quickly. Just make sure you have the following stored on a computer or shared drive **that is not your Ignition Gateway**:

- A [Gateway Backup](#). If you haven't already, you should set up [Scheduled Backups](#).

- Your Ignition License Key. Make sure you store the 6 digit code somewhere that you can easily retrieve it if your computer fails.
- The [Ignition software Installer](#). In an emergency, it's best to have the installer for the version you are currently using. You can always download an archived version from our website, but having one on hand is preferred. *Pro tip:* Edit the installer filename to include your License Key so it's easily available.

Emergency Restore Steps

1. Find another computer or create a new Virtual Machine. You need to get something up and running to put Ignition on first. We recommend having something on standby that is the same as the original computer. Further considerations:
 - a. Set your new computer to the same IP Address of the original Ignition Gateway. This way your clients won't need to do anything special to start running again.
 - b. Make sure your firewall is set correctly. At minimum, you must have the Ignition port open (8088 by default). In an emergency situation, you may want to just disable the firewall temporarily.
2. [Install Ignition](#) on the new computer.
3. [Load a Gateway Backup](#). Everything in Ignition is in the Gateway Backup, once it's loaded, you are running again!
4. [Activate your License](#). Your existing License Key will only work normally the first time it is used. After that, if you try to use it again, it will instead go into an **Emergency 7-Day Trial**. This gives you plenty of time to get back up and running before you have to deal with the License Key, even if your failure is on the weekend or outside Inductive Automation's business hours.
5. [Contact Inductive Automation](#). You can reach us by email or by phone (**1-800-266-7798**), and we will walk you through fixing your license.

Transfer a License Key Between Two Gateways

License Key Transfer Process

Situations sometimes arise where your current Gateway or server is no longer able to fulfill your needs. In these cases, you may be required to set up a new server and move your license key over from your old Gateway.

Transferring your license key combines two parts: the unactivation steps and activation steps, which are covered on the [Licensing and Activation](#) User Manual page. License keys can be transferred from an **online (connected to the Internet) Gateway** to an **offline (not connected to the Internet) Gateway**, and vice versa.

On this page ...

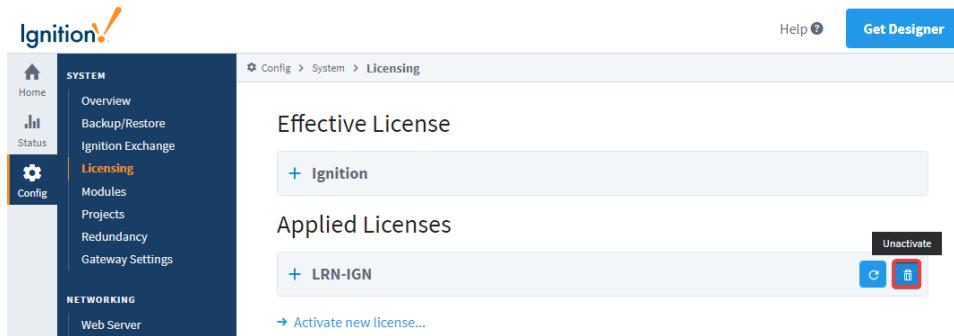
- License Key Transfer Process
- Part 1: Removing the Active License
 - Part 1a: Online Unactivation
 - Part 1b: Offline Unactivation
- Part 2: Moving the License
 - Part 2a: Online Activation
 - Part 2b: Offline Activation

Part 1: Removing the Active License

To move your license key to a different Gateway, do the following steps, depending on if your Gateway is online or offline:

Part 1a: Online Unactivation

1. Unactivate your license key from your old Gateway by going to the **Config** section > **System** > **Licensing**. You will be able to see any licenses that are currently applied.
2. Click the **Unactivate** button (trash can icon) under **Applied Licenses**.



3. A confirmation window will appear, asking you to confirm the unactivation. Click the **Yes, Unactivate** button. It may take a minute or so for the request to finish.
4. The unactivation request will be sent to [Inductive Automation's](#) licensing servers, and the license will again be available for activation on another Gateway.

Part 1b: Offline Unactivation

In the event your Gateway is unable to reach our licensing server, you will need to perform an offline unactivation. Be sure not to skip any of the following steps, as skipping a step may result in your license key going into [Emergency Activation Mode](#) when activating it on the new Gateway.

1. Go to your Gateway's **Config** section > **System** > **Licensing**. You will be able to see any licenses that are currently applied.
2. Click the **Unactivate** button (trash can icon) under **Applied Licenses**.

The screenshot shows the Ignition software's 'Licensing' page. The left sidebar has 'Config' selected. The main area shows the 'Effective License' section with '+ Ignition' and the 'Applied Licenses' section with '+ LRN-IGN'. A red box highlights the 'Unactivate' button next to the LRN-IGN license entry.

3. A confirmation window will appear, asking you to confirm the unactivation. Click the **Yes, Unactivate** button. It may take a minute or so for the request to finish.
4. An unactivation request file, called `unactivation_message.txt` will be generated and downloaded.
5. Take the `unactivation_message.txt` file to a machine with Internet access and go to <https://links.inductiveautomation.com/activation>.
6. Click **Choose File** and select the `unactivation_message.txt` file.

Unactivate Ignition

Download the unactivation request file (`unactivation_message.txt`) from the Ignition Gateway and save it somewhere convenient. Once saved, browse for the file below then click Upload Unactivation Notice to proceed with your request.

Unactivation File

Choose File unactivation_message (1).txt

Upload Unactivation Notice

7. You will see a visual confirmation that your `unactivation_message.txt` file has been successfully uploaded, and the license key can now be activated on a different Gateway.

Unactivate Ignition

Download the unactivation request file (`unactivation_message.txt`) from the Ignition Gateway and save it somewhere convenient. Once saved, browse for the file below then click Upload Unactivation Notice to proceed with your request.

Unactivation File

Choose File No file chosen

Successfully uploaded unactivation notice.

Upload Unactivation Notice

Part 2: Moving the License

Now that your license has been unactivated from the old Gateway, you are ready to activate the license on a separate Gateway.

Part 2a: Online Activation

If the new Gateway that will house your license key is online, do the following steps:

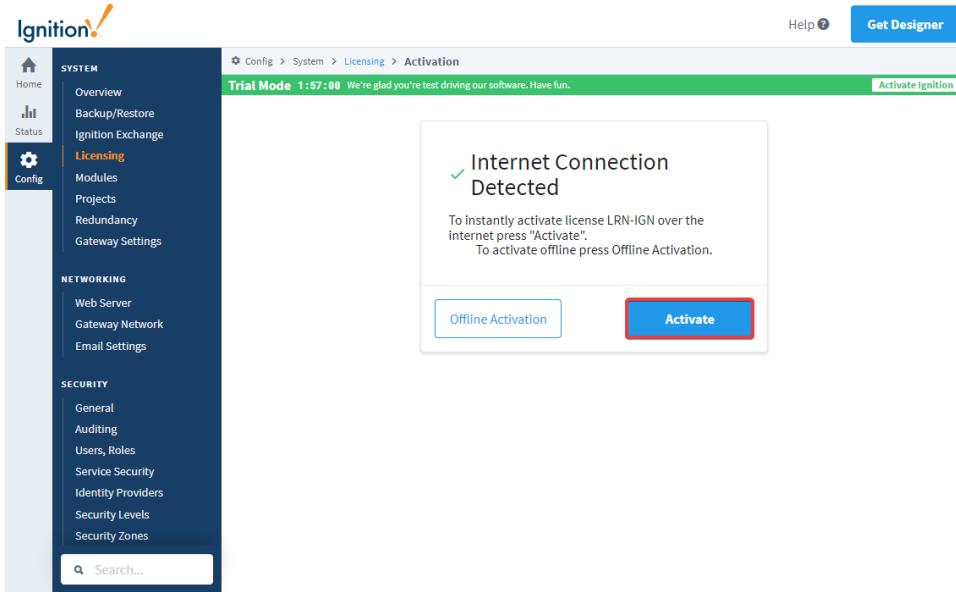
1. Go to the Gateway webpage and select the **Config** section.
2. From the menu on the left, select **System > Licensing**. The Licensing page will appear, from which you will need to click **Activate License**.

The screenshot shows the Ignition software interface. The left sidebar has sections for Home, Status, Config (which is selected), Networking, and Security. The main content area is titled 'Licensing' under 'SYSTEM'. It displays a message: 'Trial Mode 1:59:15 We're glad you're test driving our software. Have fun.' A large icon of a key is centered, with the text 'No Licenses Activated' below it. A note says: 'You may use the trial version of Ignition as long as you'd like to get a feel for our product. When you purchase Ignition, you'll receive a license key, which you can activate here.' There are 'Learn More' and 'Activate License' buttons at the bottom. A green bar at the top right says 'Activate Ignition'.

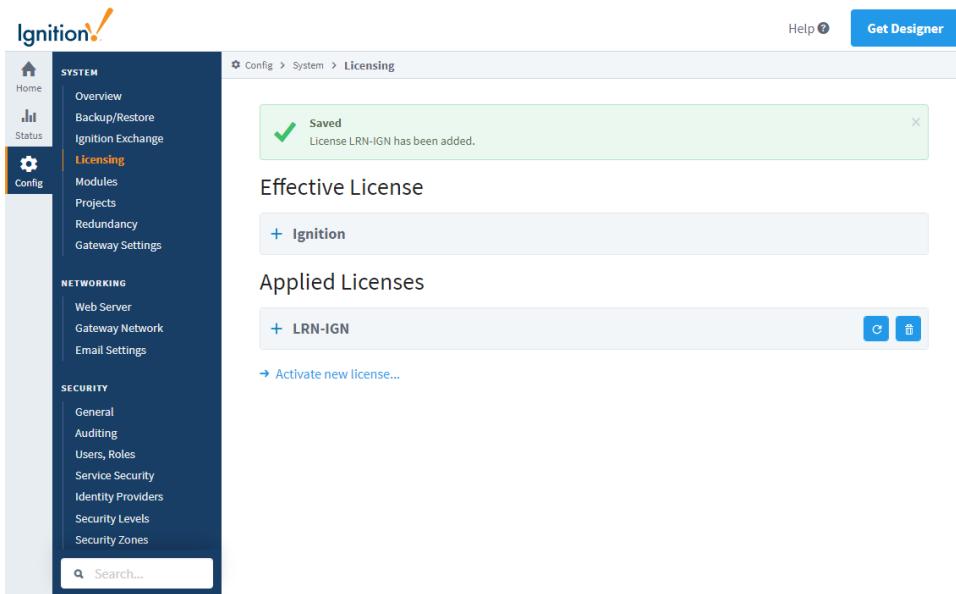
3. The **Activate a License** page will be shown. Enter the license key from the old Gateway and click **Submit**.

The screenshot shows the Ignition software interface. The left sidebar has sections for Home, Status, Config (selected), Networking, and Security. The main content area is titled 'Activation' under 'Licensing'. It displays a message: 'Trial Mode 1:57:57 We're glad you're test driving our software. Have fun.' Below it says 'Activate a License' and 'Already have a license key? You're in the right place. Enter your license key below to activate.' There is a 'License Key' input field containing 'LRN-IGN' and a 'Submit' button. A note below the input field says: 'Note: Your license key will be 6 or 8 characters'. A callout box says: 'Need help finding your license key? Check your confirmation email or invoice. If you don't have a license or need help call Inductive Automation at 1-800-266-7798.' A green bar at the top right says 'Activate Ignition'.

4. If you are connected to the Internet, click the **Activate** button.



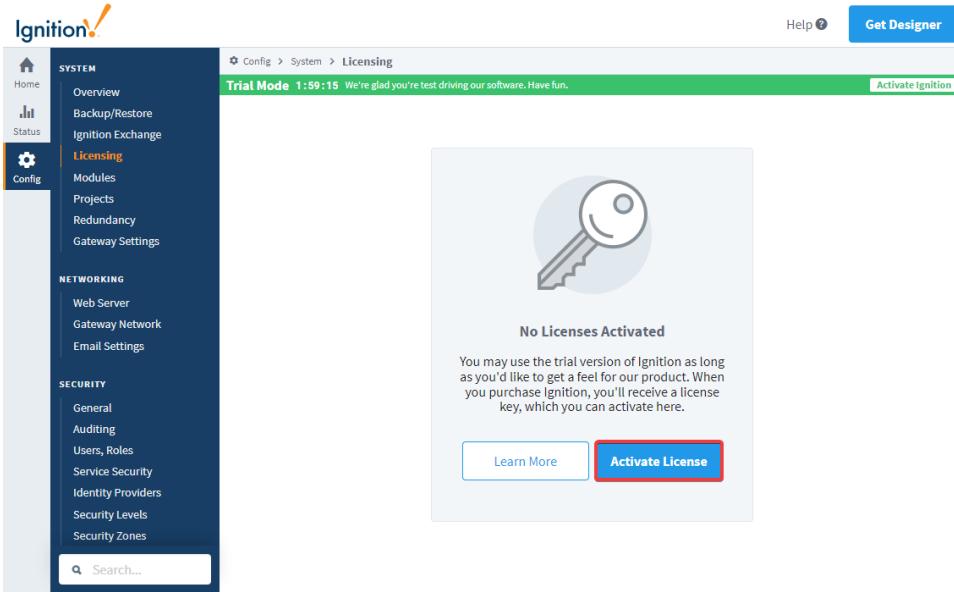
5. The **Licensing** page will refresh, and your license key will be successfully activated.



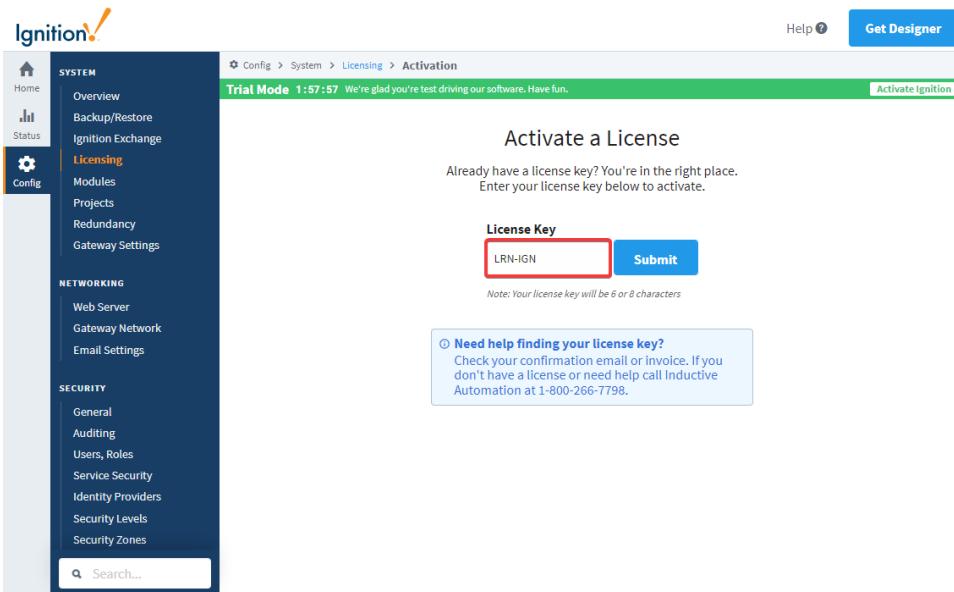
Part 2b: Offline Activation

If the new Gateway that will house your license key is offline, or if you want to perform an offline activation, do the following steps:

1. Go to the Gateway webpage and select the **Config** section.
2. From the menu on the left, select **System > Licensing**. The Licensing page will appear, from which you will need to click **Activate License**.



3. The **Activate a License** page will be shown. Enter the license key from the old Gateway and click **Submit**.



4. If you are not connected to the Internet, or if you want to activate your Gateway manually, click the **Offline Activation** button.

5. Click the **Download Activation Request** button. An activation request file, called `activation_message.txt` will be generated and downloaded.

6. Take the `activation_message.txt` file to a machine with Internet access and go to <https://links.inductiveautomation.com/activation>.
 7. Click **Choose File** and select the `activation_message.txt` file in the **Activate Ignition** section.

Activate or Unactivate Ignition

Activate Ignition

Download the activation request file (activation_message.txt) from the Ignition Gateway and save it somewhere convenient. Once saved, browse for the file below then click Upload Activation File to proceed with your request.

Activation File

Choose File activation_message.txt

Upload Activation File

Note: Should any trouble occur with the upload process, please [contact our Support department](#).

8. Once you upload the activation_message.txt file, a license file called license.ipl is generated.
9. Bring the license.ipl file back to the computer on which you're licensing Ignition. Back on the Offline Activation page, upload the license.ipl file. Then click **Activate** to finish the process.

The screenshot shows the Ignition Config interface with the 'Config' tab selected. The left sidebar has 'Licensing' highlighted. The main content area is titled 'Offline Activation'. It includes a 'Trial Mode' banner, a 'Step 1' section with a 'Download Activation Request' button, a 'Step 2' section with a URL link, and a 'Step 3' section with a file upload input field and an 'Activate' button. A red box highlights the 'Choose File' input field in the Step 3 section.

10. The **Licensing** page will refresh, showing your license was successfully activated.
11. Click **Modules** to see a list of all your activated modules.

Note: If something goes wrong during the unactivation or activation step of the license key transfer, your license key may end up in [Emergency Activation Mode](#). If this happens, you will need to contact [Inductive Automation Support](#) to get your license fixed.

Projects

What Is a Project?

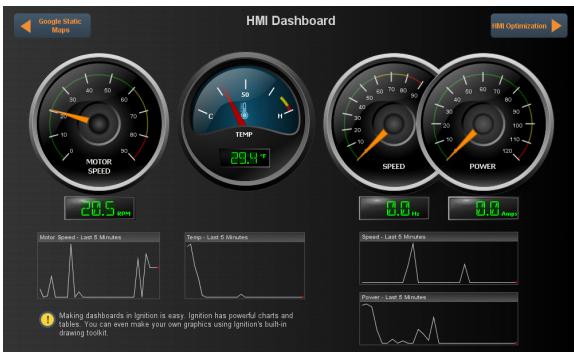
Ignition has two main parts, the Gateway and your projects. The Gateway holds all the shared information like database connections, device connections, and Tags. Projects hold all the designed elements that do the real work. Your projects can hold both interactive elements (like controls, charts, reports, and entry forms) and persistent elements (like historical loggers, automated reports, etc.).

Projects are predominantly used to create the screens that your users can interact with, the visualization part of Ignition. Here you can create any system you want, from copying existing HMI/SCADA applications to whole new systems with anything you could want to do. The windows in your project can be used for history charts, reports, database forms, alarms, drawing components, scripting, and templates.

In Ignition, a project is a unit of configuration that contains:

- **Windows, views, and components:** The HMI and SCADA controls to interact with Tags and databases
- **Transaction Groups:** A bi-directional link between databases and PLCs
- **Templates:** A collection of components that can be re-used and quickly updated
- **Reports:** PDF reports for displaying and recording data
- **Scripts:** Timer and event based scripts used throughout the system
- **General settings and properties:** The settings that control access, resource connections, layout, and timing

You use the Designer to configure and create projects. The projects are then viewed in the runtime ([Vision Clients](#) or [Perspective Sessions](#)). You can create as many projects as you want, and users can easily jump between projects on the fly or open multiple projects at the same time.



On this page ...

- [What Is a Project?](#)
- [Visualization Systems - Vision and Perspective](#)
 - [When to Use Vision for a Project](#)
 - [When to Use Perspective for a Project](#)
- [Designing a Project](#)
 - [What Is in a Project?](#)
- [Switching Between Projects](#)
- [Project Workflow](#)
- [Project Export and Import](#)
- [Project Versioning and History](#)
- [Project Examples](#)
- [What Is not in a Project?](#)

Visualization Systems - Vision and Perspective

When starting a new project, there are many things to consider, such as who the users will be, how much data you need, what kind of time and resources you have to work with, visualization needs, and so forth. Deciding whether to use Vision or Perspective really comes down to which module best fits your project at the visualization level.

When to Use Vision for a Project

Vision is the best choice for traditional industrial plant-floor and desktop screens, standalone HMIs, and the like. If you need a full, production-ready application right now, then it is best to choose Vision. If you're using multi-monitor or multi-desktop workstations, dedicated control panels and applications, desktop-dedicated or dedicated-access applications, terminals, or parallel screens, then Vision is the best way to go.

When to Use Perspective for a Project

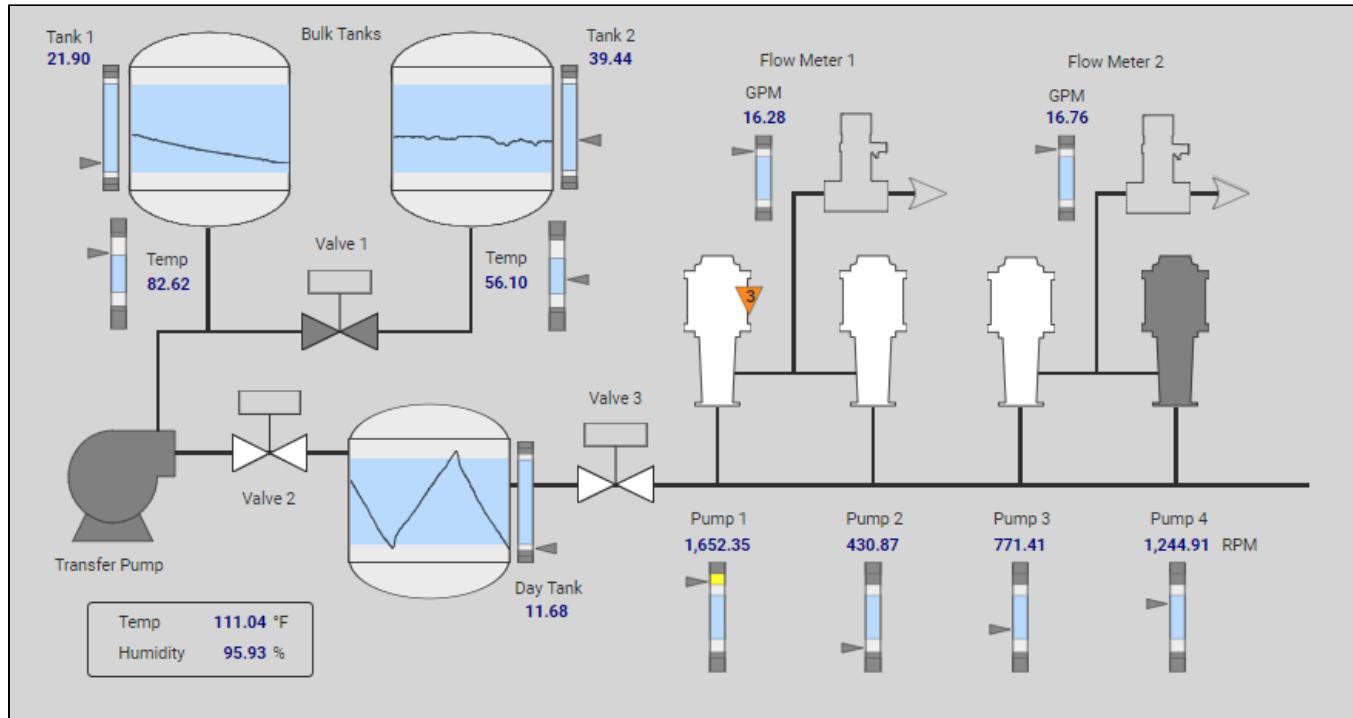
If you need to build mobile-responsive applications, then Perspective is the recommended way to go. If your application needs to run on a mobile OS, namely iOS or Android, then choose Perspective. If you need Two-Factor Authentication or federated identities for the application you're building, choose Perspective.

With Perspective, you can automatically adapt to fit any screen size using mobile-optimized container types. It provides the ability to use your device's sensors and intuitive touch commands, as well as message handling, flexible property bindings, and CSS3 styles.

Designing a Project

When you launch the Designer, you're prompted to select or create a project. The Designer then launches the project and you can set and modify the different types of project settings and resources. In the Designer, you can create any number of projects using either Vision or Perspective. Projects that contain viewable elements, such as Vision Windows and Perspective Pages, will have a launch link on the Gateway homepage. Non-viewable elements such as [Transaction Groups](#), [Named Queries](#), [SFCs](#) and [Reports](#) exist in a project and execute in the Gateway. These resources do not have a runtime, and run independently of any [Vision Clients](#) or [Perspective Sessions](#) being open.

There are no limits to the number of projects that can be created on a Gateway, but each runtime Client, Session, or Designer can operate only on one project at a time.



What Is in a Project?

You use the different tools in Ignition such as components, shapes, images, Symbol Factory graphics, and Scalable Vector Graphics (SVGs) to create the components. Configuring components is the bulk of the designer's work when designing a project. The basic workflow is to take a component from the palette and drop it into a container on a Vision Window or Perspective View. You can use the **Property Editor** panel to alter the component's properties which changes the component's appearance and behavior.

To make the component do something useful, like display dynamic information or control a device register, you configure property bindings for the component. To make the component react to user interaction, you configure event handlers for it. It is primarily through property bindings that you bring windows to life, and have them do useful things. A property binding simply links one component's property to another on the same window.



Switching Between Projects

When you launch an Ignition Client or Perspective Session, it opens a single project to display. If you want to open multiple projects you can have multiple clients open on the same computer, or you can use Ignition's [Retargeting](#) system to make seamless transitions between projects. The Retargeting feature allows you to jump from one project to another without closing the client. This allows your users to jump from area to area while still allowing you to keep your designs compartmentalized in multiple projects.

Your projects may all use different authentication sources or role sets for [security](#). This means as your users jump from project to project, they may be able to use some or all of the controls on one project, but only see what is happening on another project.

Project Workflow

Getting up and running quickly with your project is simple:

1. Open the Designer
2. Choose either Perspective or Vision.
3. Start designing!

The only challenge is figuring out exactly what you want to make with your system. For additional information, see [Creating a Project](#).

Project Export and Import

In Ignition, a project backup and restore is referred to as Project Export and Import. Projects are exported individually and only include **project-specific** resources such as Perspective Views, Perspective properties, Vision Windows, Vision Templates, client event scripts, alarm pipelines, named queries to name a few. They **do not** include any Gateway resources like database connections and Tag Providers. A project is exported to a **.zip** file, and you can import it to any other Gateway that you have permission to access.

When you import a project from the **.zip** file, it will be merged into the existing Gateway. If a project already exists on the Gateway you are attempting to import it on, you have the option of renaming the project or overwriting the existing project.

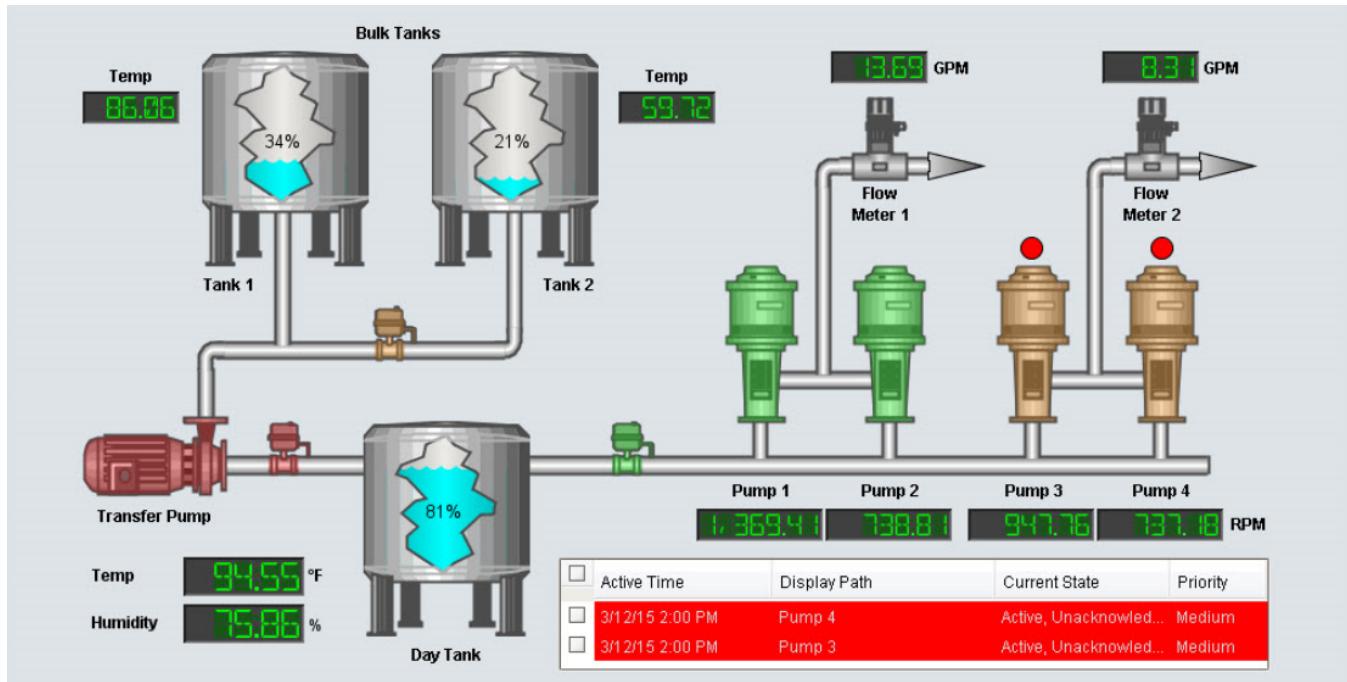
Project Export and Import are covered in detail on the [Project Export and Import](#) page.

Project Versioning and History

Project versioning is handled outside of Ignition. The file system that stores data in Ignition stores everything as a series of files. You can use any tool you'd like to save or export these files outside of Ignition, including using versioning software to keep track of your changes.

Project Examples

Ignition is such a diverse application that it can work in any field you can imagine. Just a few examples: Oil and gas, food and beverage, automotive, medical, air traffic control, water and waste water. You can get anything you want set up in Ignition. To see a few live example screens, check out our online demo project examples for [Vision](#) and [Perspective](#). They are packed with samples of the features and functions you might want to use from reporting, to history, to HMI optimization.



What Is not in a Project?

The Designer allows you to create and modify several types of resources that are shared by ALL projects. Depending on the resource, this means that either they run independently of the projects, or they are available for use by any project. Here are a few of them:

- [Alarm Pipelines](#) - Control alarm notification (runs independently).
- [Sequential Function Charts \(SFC\)](#) - Logic to step through a process (runs independently).
- [Transaction Groups](#) - Perform various actions such as storing data historically, synchronizing database values to OPC, or loading recipe values.
- [Project Library](#) - Blocks of scripting code (available to all projects).
- [Tags](#) - Basic or UDT Tags provide realtime data (available to all projects).
- [Alarming](#) - Alarms exist on Tags, and so are not in a project (available to all projects).

In addition to these resources, the Gateway connections and settings are available to all projects and are set up in the [Gateway](#). These include resources such as [database connections](#), [OPC server](#) and [device connections](#).

Related Topics ...

- [Designer](#)
- [Perspective Sessions](#)
- [Property Bindings in Perspective](#)
- [Scripting in Perspective](#)
- [Vision Client Launcher](#)
- [Property Bindings in Vision](#)
- [Scripting in Vision](#)
- [Alarming](#)
- [SQL Bridge \(Transaction Groups\)](#)

In This Section ...

Project Inheritance

Project Inheritance allows one project to inherit resources from another project. If you have project resources that you want other projects to use such as views, windows, scripts, templates, or pipelines, you can create an inheritable project allowing other projects to inherit those resources. The project inheriting the resources can also overwrite the resources and let you re-define them specifically for that project.

Project inheritance is extremely flexible in that it allows inherited projects to also be configured as 'inheritable' and become a parent project itself. This allows for complex hierarchies of re-usable resources to be designed. Within each project, inherited resources may be used by other, "local" resources. For example, an inherited Vision template could be embedded in a window, or an inherited script could be executed by a button.

This page demonstrates how project inheritance allows you to share resources across multiple projects as you deem fit by pointing one project to another project, and overriding resources to re-define resources specifically for the project you're working on.

Configuring Project Inheritance

In order to make a project inheritable, you need to enable project inheritance on the project containing the resources that you want to share. In a similar way, these parent projects need to be assigned to the child projects that will be inheriting the resources to populate those resources. The following example demonstrates how to configure project inheritance for existing projects, but if you know a project needs to be inheritable or inherit a parent project's resources before creating, you can configure inheritance settings upon creation. This is available in the Gateway after selecting Create new project on the Config > System > Projects page and choosing the settings called out below, or in the [Designer](#).

In this example, we have two projects: one project called 'global' which contains some project resources, and another project that is not yet inheriting any resources called 'Project_X.' Since we will be sharing the resources from the 'global' project with Project_X, let's first configure the 'global' project to be inheritable.

Configuring Parent Project

1. Go to the **Config** tab on the Gateway Webpage, and select **Systems > Projects**. This brings up a list of all your projects. You'll notice that for each project listed, you'll see at a glance if a project is Inheritable, and if so, the name of the Parent Project will be displayed.

Name	Description	Enabled	Inheritable	Parent project	Actions
global		true	true		Edit
Project_X		true	false		Edit
Controller1		true	false	global	Edit

2. Find the project you want to make inheritable, and click the **Edit** next to the project name (i.e., global) to open the Project Settings window.
3. Enable the **Inheritable** property, and click **Save**.

Allow Overrides

When you make a project inheritable, the Allow Overrides function is set by default on all project resources in the inheritable project. This allows all project resources to be propagated to all inherited projects

On this page ...

- [Configuring Project Inheritance](#)
 - [Configuring Parent Project](#)
 - [Configuring Child Project](#)
 - [Using Inherited Resources](#)
- [Inheritance and "Runnable" Resources](#)
 - ["Runnable" Resources](#)
- [Overriding Inherited Resources](#)
 - [Discarding Inherited Resource Overrides](#)
 - [Renaming an Inherited Resource](#)
 - [Opening Read-only Named Queries](#)
- [Inheritable Project Examples](#)
- [Project Export](#)



Project Inheritance

[Watch the Video](#)

Project Settings

Name *	global
Choose a name to identify this project.	
Description	
Title	
The title for the project. This can contain more characters than the name (space, etc), and will be used to represent the project to users. If empty, the name will be used.	
Enabled	<input checked="" type="checkbox"/> A disabled project will not be active on the Gateway, but will remain editable in the Designer.
Inheritable	<input checked="" type="checkbox"/> Inheritable projects are not runnable as a stand-alone project, but are intended to provide shared resources to one or more child projects.
Parent Project	-None-

Cannot Launch Inheritable Projects

When you have a project that's flagged as inheritable, you can not launch it as a stand-alone project (i.e., Perspective Session or Vision Client). You will get a 'Project Not Runnable' error message. If you have an inheritable project that you want to launch, you must have another project to inherit from it.



Configuring Child Project

Now that we have an inheritable project, let's setup an existing project (i.e., Project_X) so it inherits resources from the 'global' project.

1. Under the **Config** tab select **System > Projects** page
2. Find your project and click the **Edit** button.
3. Select the inheritable project from the **Parent Project** dropdown list (i.e. global), and click **Save**.

Project Settings

Name *	Project_X
Choose a name to identify this project.	
Description	
Title	The title for the project. This can contain more characters than the name (space, etc), and will be used to represent the project to users. If empty, the name will be used.
Enabled	<input checked="" type="checkbox"/> A disabled project will not be active on the Gateway, but will remain editable in the Designer.
Inheritable	<input type="checkbox"/> Inheritable projects are not runnable as a stand-alone project, but are intended to provide shared resources to one or more child projects.
Parent Project	global

Using Inherited Resources

You can treat an inheritable project as a library of resources for use in other projects. When project resources are changed in the original project, these changes will get passed down to the inherited projects. The same thing is true if new resources are added to the inheritable project. By using inheritable projects, you can create a resource library that will help designers build their projects more quickly, ensure consistency and reusability across all their projects.

To view your inherited resources, go to the Designer and open your child project (i.e., Project_X). Expand the folders that contain resources, and you'll notice that the inherited project resources are grayed out, including the Perspective Views and Vision Windows folders. Grayed out resources mean that those resources are inherited, and can only be edited from their parent project.

The following feature is new in Ignition version **8.1.26**
[Click here](#) to check out the other new features

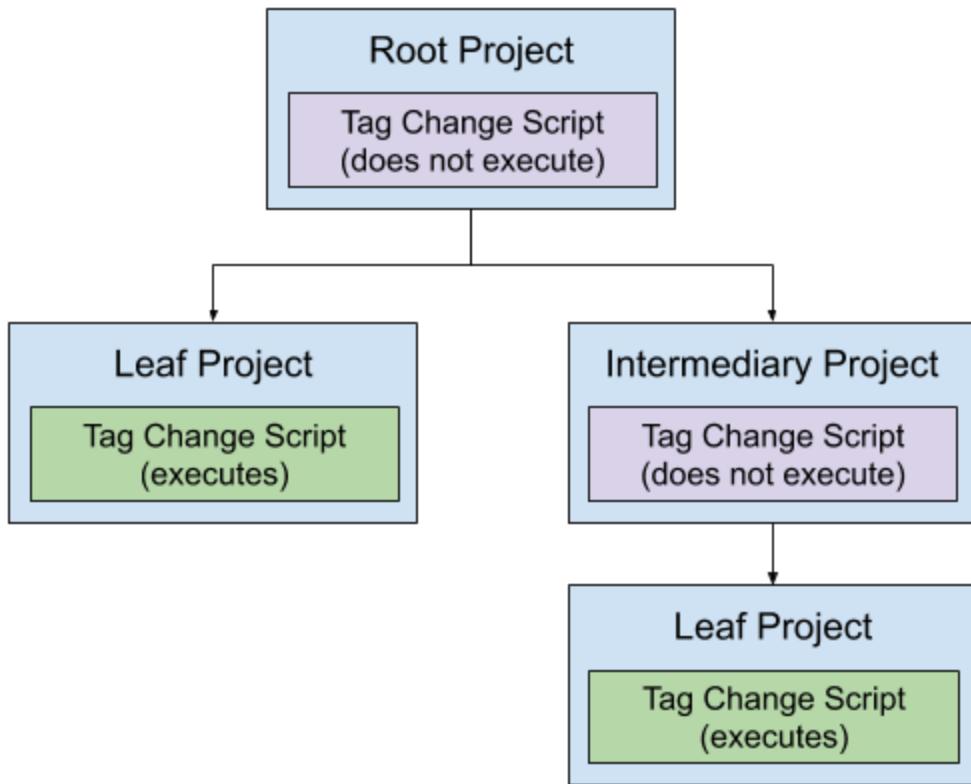
Inherited resources will have one of two icons following the resource name. The Inherited Resource  icon is displayed when the resource is grayed out and indicates the resource is inherited and unchanged from the parent project resource. The Inherited Resource Overridden  icon indicates the resource has been overridden and is now editable.

Anytime a new resource is added or an existing resource is changed in the parent project, it will propagate down to the inherited project and will appear grayed out with an Inherited Resource  icon.

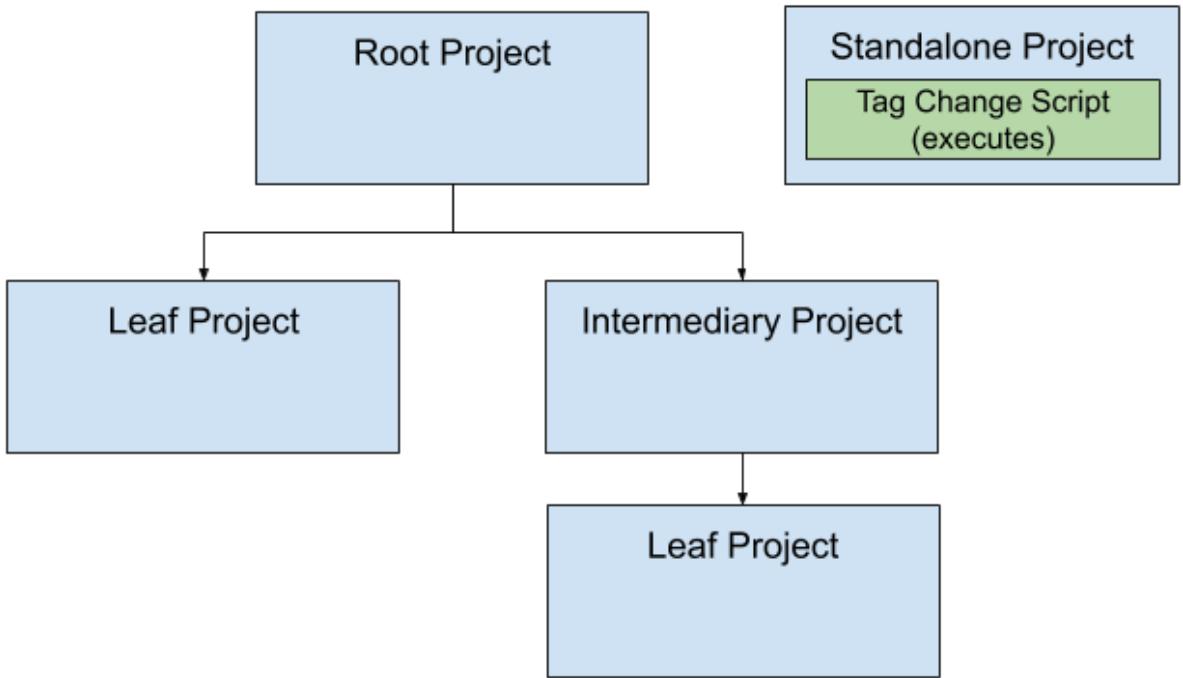
Inheritance and "Runnable" Resources

Some resources in a project "run" or execute in every leaf project (that is, a project that is **not** inherited by any other project) in the inheritance chain. Thus, if multiple projects inherit from the same parent project, then each leaf project will contain runnable resources, potentially resulting in duplicate executions.

The diagram below represents an inheritance hierarchy on a single gateway. A "Root" project contains a Gateway Tag Change script. Two projects inherit from the Root, so they'll inherit the script. In both cases, the inheritance chain leads to leaf projects. In this single gateway, two running instances of the same tag change script exist, meaning there will be duplicate script executions.



In the case of resources that "run", it is highly advised that they exist in a leaf project, or a standalone project (a project that does not participate in inheritance at all). In regards to the diagram above, we could prevent duplicate execution of our Tag Change Script easily by moving the Tag Change Script to a separate, standalone project:



"Runnable" Resources

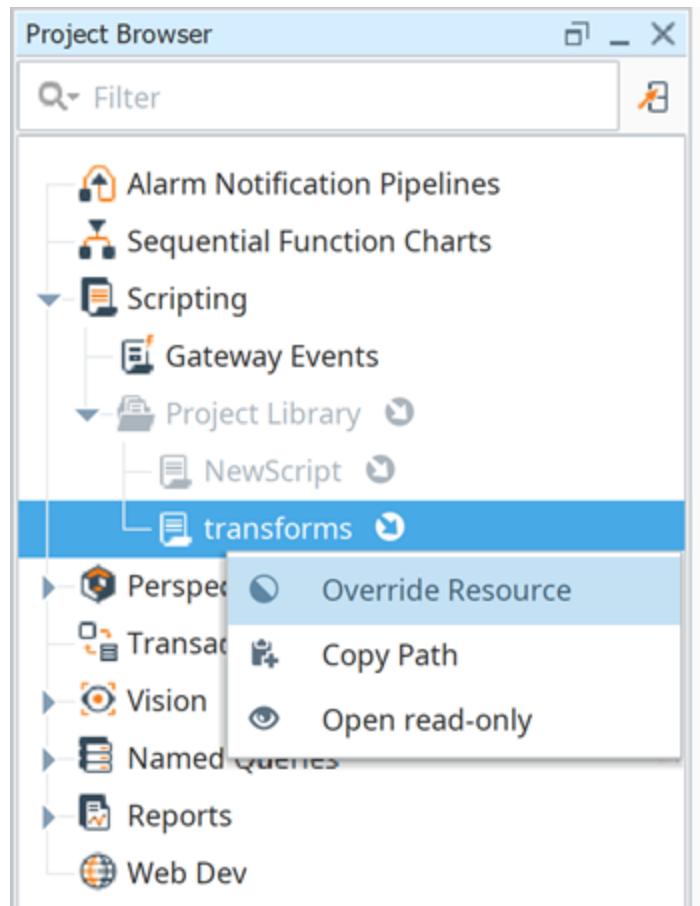
The following resources are considered "runnable":

- [Gateway Event Scripts](#)
- [Alarm Notification Pipelines](#)
- [Sequential Function Charts](#)
- [Transaction Groups](#)

Overriding Inherited Resources

To edit a resource in an inherited project, you need to override the resource by right-clicking the resource and selecting **Override Resource**. By overriding the inherited resource, the resource is recreated in the inherited project, and any future changes made to the original project resource will *not* propagate down to the inherited resource.

If new resources are added in the parent project, they will automatically propagate down to inherited projects. New resources that are added in the inherited project will initially display grayed out. You will need to override each new resource to edit it.

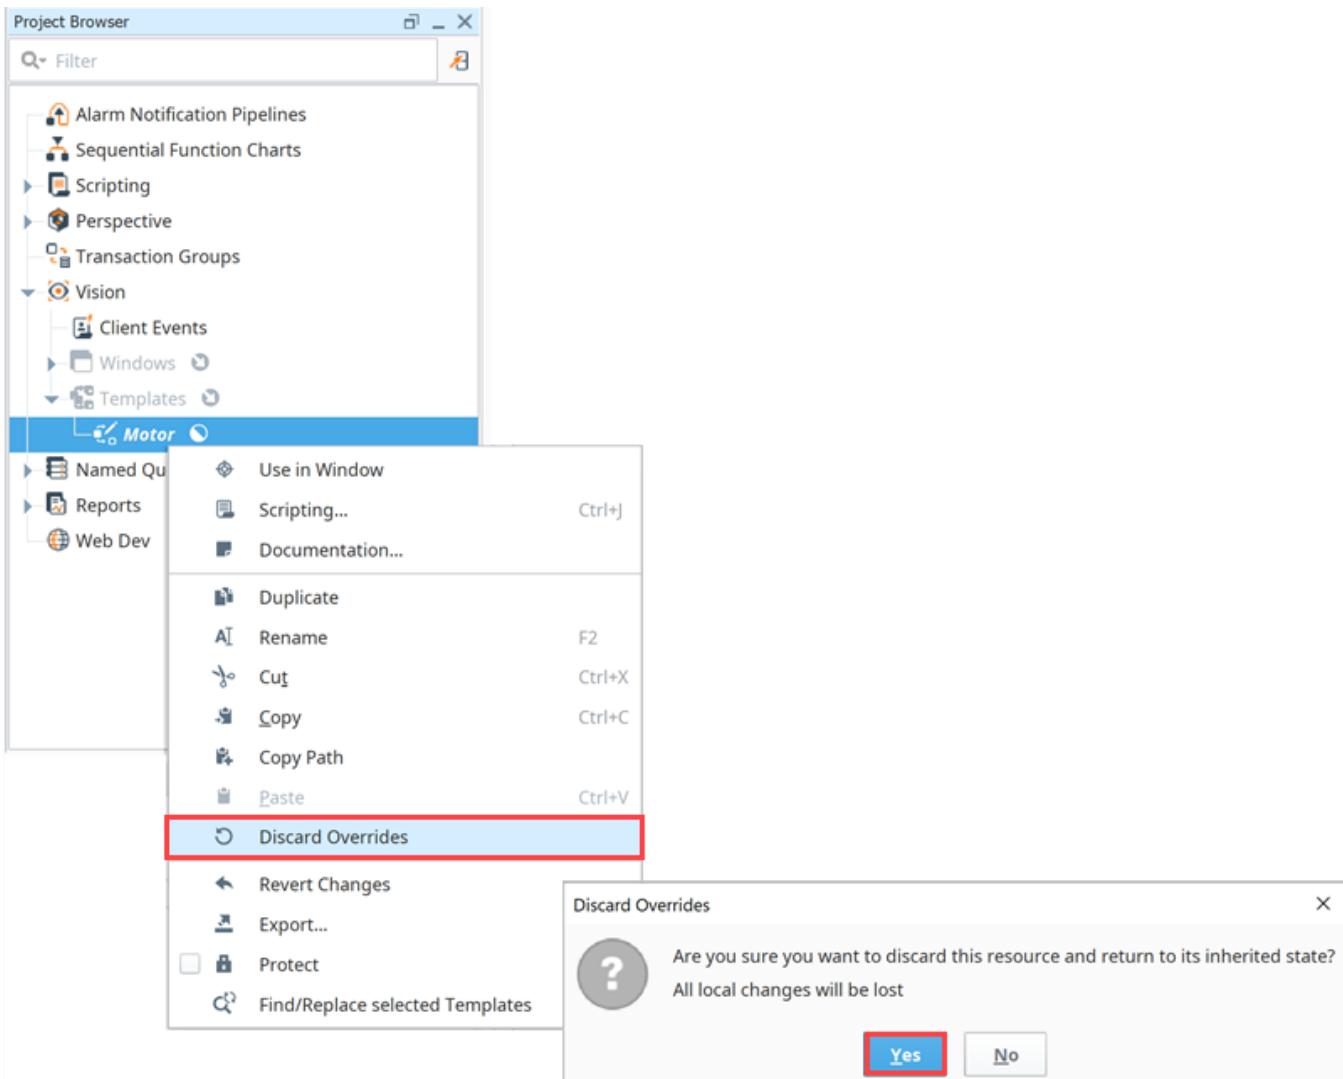


Discarding Inherited Resource Overrides

This feature was changed in Ignition version **8.1.26**:

The Discarding Overrides function was previously represented by selecting a Delete option in the right-click menu. Since this selection results in the resource being grayed out to prevent editing, the function name was updated to Discarding Overrides for clarity.

Changes made to inherited resources that have been overridden can be discarded. Selecting **Discard Overrides** removes the override placed on the inherited resource. The inherited project resource will return to its original state inherited from the parent project (without your edits), and the resource becomes grayed out in the Project Browser with the Inherited Resource icon. To discard the override on an inherited resource, right-click on the resource and select **Discard Overrides**. A dialog box will pop up confirming you want to discard the resource.



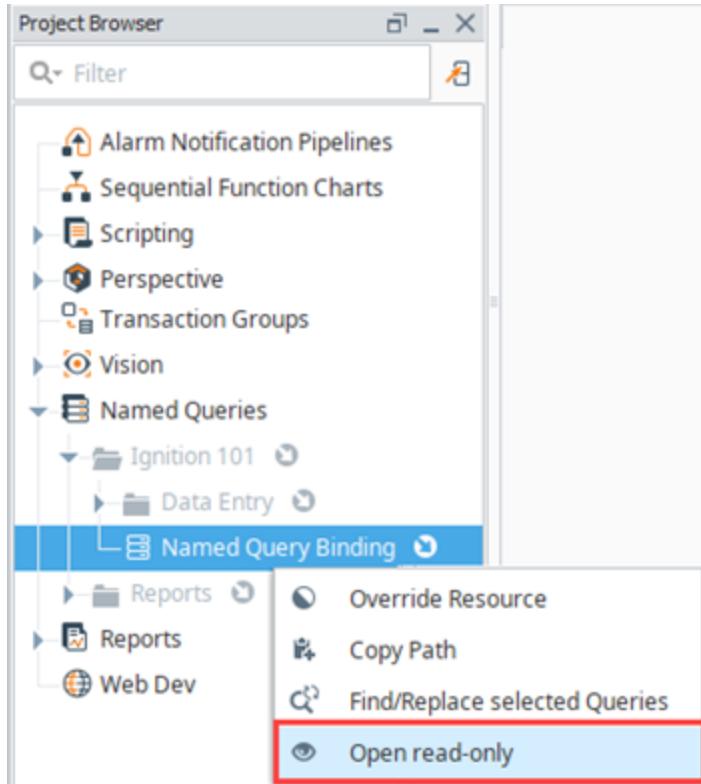
Renaming an Inherited Resource

You can easily rename an inherited resource, but beware that when you rename a resource, the inheritable project will propagate the original resource to the inherited project. For example, say you have an inheritable project resource named 'Map_Transform' and you renamed it to 'CA_Map_Transform.' Ignition knows that the original project resource is no longer there, and because the project is flagged as inheritable, it will propagate that original resource to the inherited project. Now you have both the renamed 'CA_Map_Transform' and the original inheritable resource 'Map_Transform.'

Opening Read-only Named Queries

The following feature is new in Ignition version 8.1.26
[Click here](#) to check out the other new features

Overriding an inherited Named Query used to be the only way to view the Named Query, even if you did not want to make any changes. This option is still available (and will allow you to edit the Named Query), but you can now select the **Open read-only** option from the Named Query right-click menu to view the Named Query Settings, Authoring, and Testing tabs. Since the read-only option does not allow changes, you'll notice the Inherited Resource  icon remains next to the selected Named Query.



Inheritable Project Examples

There are many ways how you might want to configure and organize your inheritable projects. It's whatever works best for your organization and design projects. Here are a couple of common ways to organize your shareable resources.

- You can create one inheritable project that contains many project different resources: scripts, pipelines, views, templates, windows, SFCs, etc. This is one inheritable project containing all your inheritable project resources.
- Another option is to create several inheritable projects. You can have one inheritable project dedicated for each type of resource: one for scripts, one for views, one for templates, one for pipelines, etc.

Project Export

When a project is [exported](#) from the Designer you can choose to export strictly local resources or all resources (including those inherited). Exporting inherited resources is a great way to create a "flattened" project export, containing all resources used by the project regardless of which project the original resources are located in.

Note that [exporting a project from the Gateway's UI](#) will only export a project's local resources, and not contain inherited resources.

Related Topics ...

- [Project Inheritance - Upgraded Features in Ignition 8](#)
- [Project Export and Import](#)

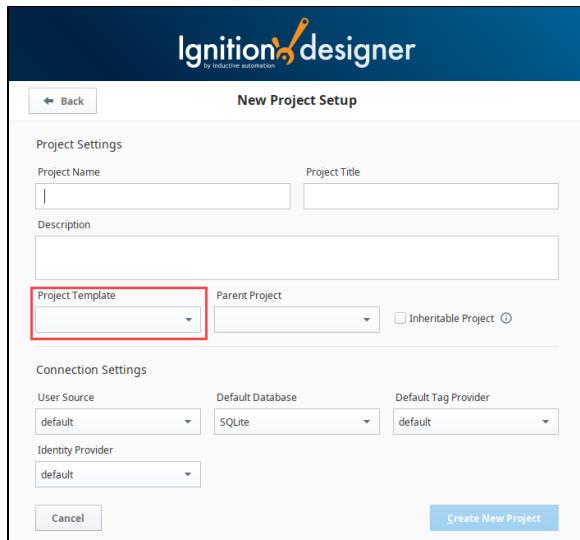
Project Templates

When you [launch the Designer](#), you will be presented with the Open/Create Project window. This window lets you hit the ground running almost immediately by populating a few fields, the most important ones being your project Name and Title. You can use the defaults in the remaining fields and change them in the **Config > Project** page later, but don't overlook the **Project Template** field. Ignition provides several project templates for you to choose from to help you quickly get started developing your project in either Perspective or Vision. The templates get you started with a basic navigation structure and let you add more items as your project matures. Additional project templates can be found on the [Ignition Exchange website](#).

Note: You must be connected to the Internet in order to see the templates available in the Project Template list.

On this page ...

- [Perspective Project Templates](#)
 - [Web Nav](#)
 - [Perspective Menu Nav](#)
 - [Editing Properties and Page Configuration Settings](#)
- [Vision Project Templates](#)
 - [Tab Nav](#)
 - [2-Tier Tab Nav](#)
 - [Tree Nav](#)
 - [Adding Windows to the Navigation](#)



Perspective Project Templates

Here are two Perspective Project Templates you can choose from to get you off to a quick start: **Menu Nav** and **Web Nav**.

The Perspective Menu Nav and Web Nav templates contain several pre-defined views: Home, Charts, Alarms, and Settings. These views can easily be edited by selecting the view in the Project Browser and editing its Props to anything you want them to be. You can also go to the [Page Configuration](#) and edit the page configuration settings and [docked view](#) properties.

The images below show what the Menu Nav and Web Nav look like in a Perspective Session.

Web Nav

The Web Nav is good for small size project structures where you only have a few main Views. It allows you to navigate the areas of your project using the tabs at the top of the screen. It has a docked view that contains tabs that are always open to do navigation, and the main view which fills the rest of the space. It is a flat structure similar to what you see on many webpages.

The screenshot shows the Ignition Perspective application interface. At the top is a header bar with the Ignition logo on the left and navigation links for HOME, CHARTS, ALARMS, and SETTINGS on the right. Below the header is a dark blue navigation bar with the word "Home" in white. The main content area has a light gray background. At the top of this area, the text "Welcome to Perspective!" is displayed in bold. Below it is a paragraph of text: "The skeleton project is intended to help you get started with Perspective. The project is designed to be mobile-responsive and fully utilizes styles. Feel free to modify as desired." In the bottom right corner of the main content area, there is a small dark button with a white upward-pointing arrow.

Perspective Menu Nav

The Menu Nav template is a hierarchical view of groups of information that can be configured to expand submenu branches and menu items and is docked on the left side of the screen. The Menu Nav is good for medium and large project structures because you can expose important information at a glance and allow users to navigate and expand submenus if any exist. The menu automatically hides itself on smaller screens devices.

This screenshot shows the same Ignition Perspective interface as the first one, but with a vertical sidebar on the left labeled "Menu Nav". The sidebar contains four items: "HOME" (with a house icon), "CHARTS" (with a chart icon), "ALARMS" (with an alarm clock icon), and "SETTINGS" (with a gear icon). The main content area is identical to the first screenshot, featuring the "Welcome to Perspective!" message and the paragraph about the skeleton project. The small button in the bottom right corner is also present.

Editing Properties and Page Configuration Settings

To edit the Menu Nav and Web Nav menus, first you have to have pages configured in your project. Once you have pages configured, you can go to the Navigation view in the Project Browser, select the component and change the corresponding component props in the Perspective Property Editor to add, remove, or update the menu items.

Page Configuration Settings

Page Configuration

Shared settings

- / → Page/Home
- /alarms → Page/Alarms
- /charts → Page/Charts
- /settings → Page/Settings

Header BP Large

Corner Priority

left-right **top-bottom**

Menu

Configure Docked View

View: Docks/Menu

Display: auto

Resizable?: false

Content: push

Modal?: false

Size: 260 x 768

Auto Breakpoint

Dock ID: menu

Handle: hide

Handle Icon

View Parameters: + Add Object Member...

Project Browser and Project Editor

File Edit View Project Component Tools Help

Project Browser

Perspective Property Editor

Perspective Components

Home

Welcome to Perspective!

The skeleton project is intended to help you get started with Perspective. The project is designed to be mobile-responsive and fully utilizes styles. Feel free to modify as desired.

Views

- Banners
- Header
- Pages
- Alarms
- Charts
- Home
- root
- Container
- Banner
- Spacer Banner
- Content
- Spacer Footer
- Settings

Vision Project Templates

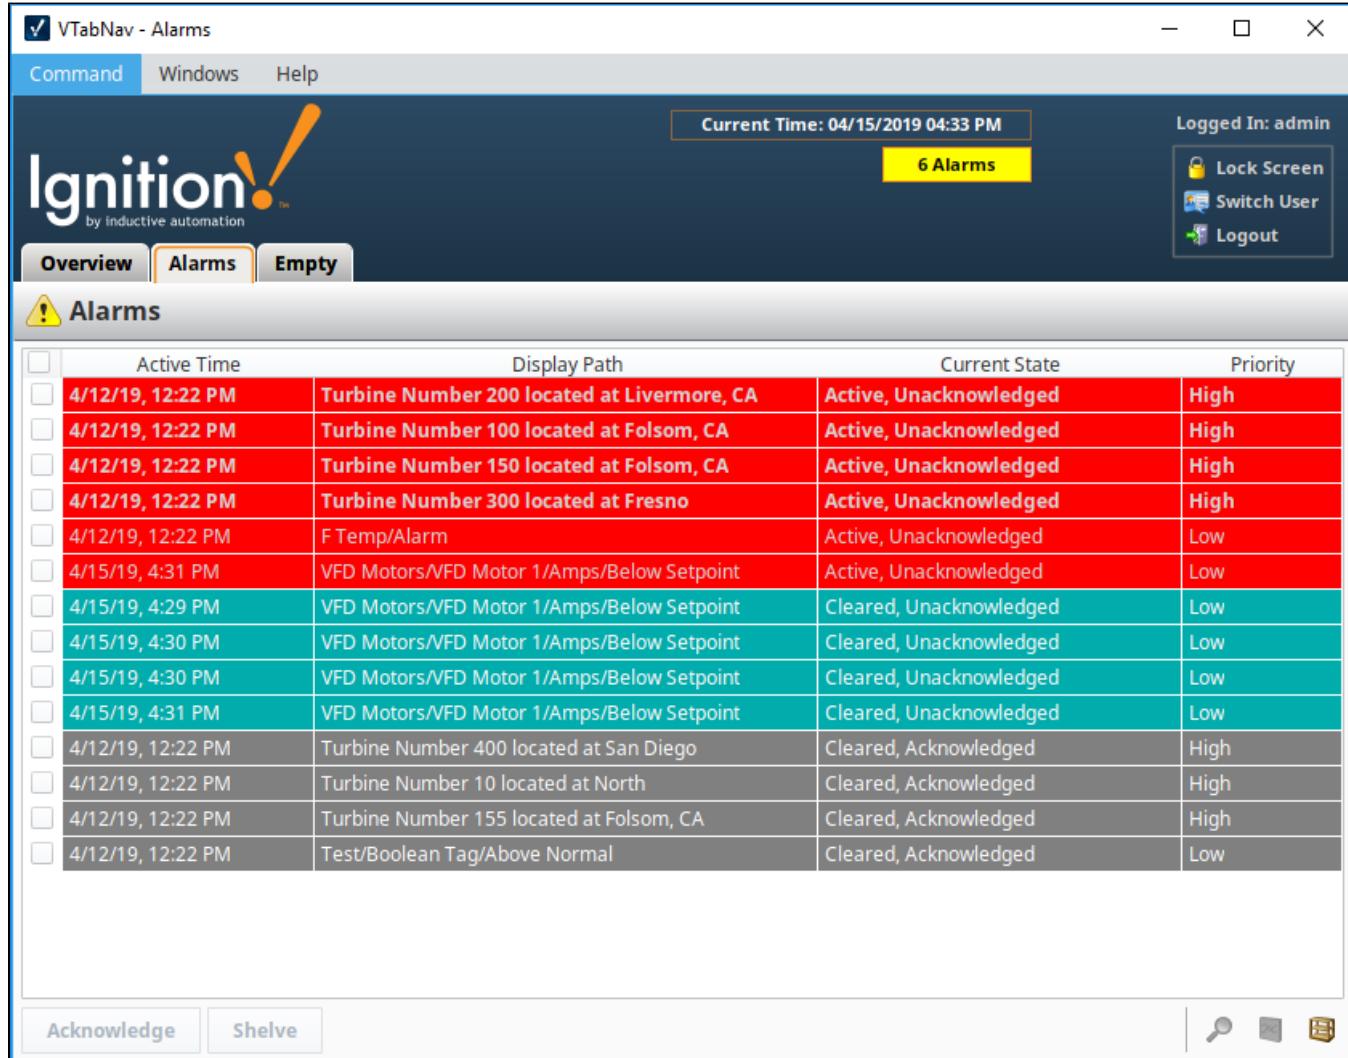
There are three Vision Project Templates to choose from to help kick start your project: **2-Tier Tab Nav**, **Tab Nav** and **Tree Nav**.

The Vision templates, just like the Perspective templates contain several pre-defined menu objects. These objects can easily be edited by selecting the object in the Project Browser by editing its properties and changing them to anything you want them to be.

The images below show what each of the Vision project templates look like in a Vision Client.

Tab Nav

The Tab Nav is good for small size project structures where you only have a few main windows. It allows you to navigate the areas of your project using the tabs. The Tab Nav project template has a docked window that contains tabs that are always open to do navigation, and the main window which fills the rest of the space.



2-Tier Tab Nav

2-Tier Tab is good for small and regular size project structures where windows are grouped. The Tab Nav project template is similar to the Tab Nav template only it has a second tier of tabs added. It contains a second level of tabs allowing you to navigate around various areas of your project. The 2-Tier Tab Nav project template has a docked window that contains tabs that show and hide based on selection that are always open to do navigation, and the main window which fills the rest of the space.

The screenshot shows the Ignition V2TierTabNav - Overview application window. The top menu bar includes 'Command', 'Windows', and 'Help'. The header features the Ignition logo, the current time '04/15/2019 03:48 PM', and a 'Logged In: admin' status with options to 'Lock Screen', 'Switch User', and 'Logout'. Below the header, there are two tabs: 'HMI Screens' (selected) and 'Administration'. A sub-menu bar below shows 'Overview' (selected), 'Alarms', and 'Empty'. The main content area is titled 'Overview' and contains a button labeled 'Open Popup'. The text in the content area reads:

Welcome to the Skeleton project (Dual Tier Tabbed)

This project is meant to be a template for a large variety of applications. We wanted to give new Ignition project designers a jumping off point.

This project uses a dual tier tabbed navigation structure. It follows the typical navigation strategy outlined in the user manual (http://www.inductiveautomation.com/support/usermanuals/ignition/typical_navigation_strategy.htm).

You will notice a *docked* north window called "Navigation" that contains tabs to switch between different sections. Once a section is selected a different set of tabs appear below to switch between different windows. The *Tab Strip* component provides automatic window navigation using the built-in "Swap to Window" mode. To configure the tabs, right click on the tab strip component and select *Customizers* -> *Tab Strip Customizer* in the Ignition designer. Each tab points to a different main window. Remember: for this skeleton project there are two sets of tabs that are visible when the appropriate section is selected. Those tabs have a visible property that is bound to an expression based on the first set of tabs. In the Ignition designer use the *Preview Mode* [F5] to switch between tabs to make the second set of tabs visible.

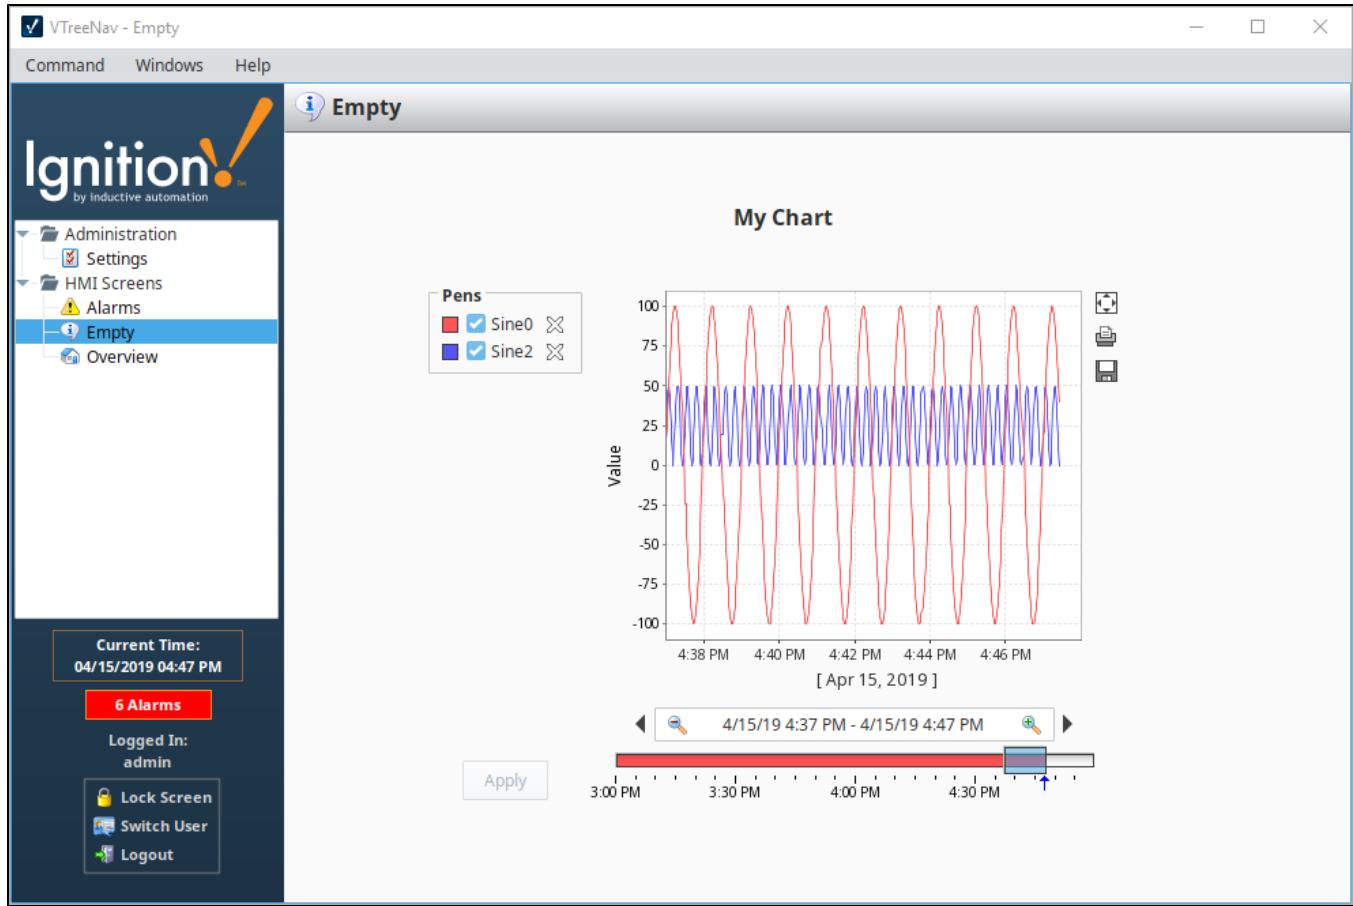
This window, "Overview", has a button to open a popup window. Popup windows are great for displaying more information about particular sections. To configure which window the button opens right click on the button and select *Event Handlers...*. Under the *Navigation* tab you can select which window to open. Popup windows can be parameterized so you can re-use it for multiple purposes, depending on the values that you pass in. For more information on parameterizing windows see the user manual (http://www.inductiveautomation.com/support/usermanuals/ignition/parameterized_windows.htm).

Lastly, this project is set to automatically lock the screen after 5 minutes of inactivity (ie no mouse or keyboard activity). To disable or change the amount of time select *Project* -> *Event Scripts (Client)* from the file menu in the Ignition designer. On the left select *Timer* and choose the *Check Inactivity* script. Feel free to delete it or change the 300 to any number of seconds.

Enjoy the skeleton project!

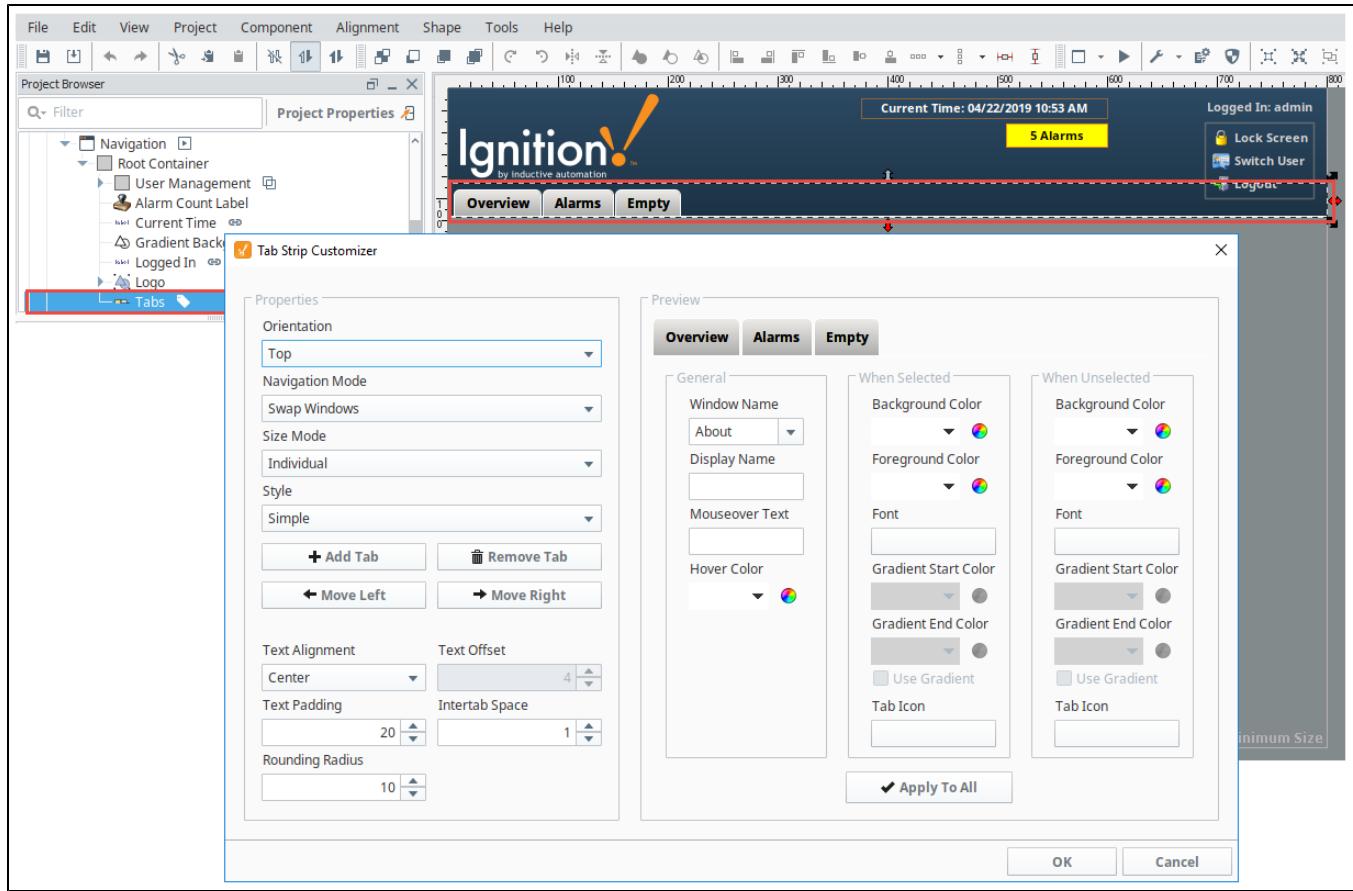
Tree Nav

The Tree Nav project template is a hierarchical view of groups of information that can be configured to expand submenu branches and menu items and is docked on the left side of the screen. It is great for medium and large project structures because you can view the entire project structure at a glance allowing you to navigate to any structure within the tree view. The tree structure uses folders to group main windows, and can be as many levels deep as you need.



Adding Windows to the Navigation

To edit the Tab Nav and Tree Nav menus, you just need to edit the properties of the navigation components. You can go to the Navigation window in the Project Browser, select the Tab Strip, and open the Tab Strip Customizer to change the corresponding component props to add, remove, or update the menu items.



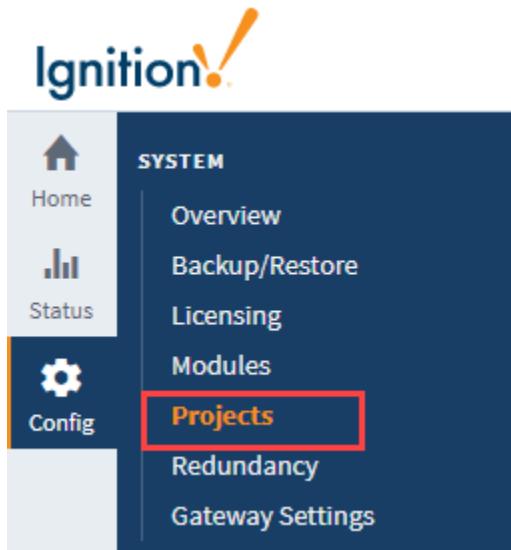
Related Topics ...

- [Quick Start Guide](#)
- [Designer](#)
- [Pages in Perspective](#)
- [Navigation Strategies in Vision](#)

Project Settings

Project Settings

When you create a new project, the Open/Create Project window captures most of the project settings. See [Creating a Project in the Designer](#). If you want to view or edit your project settings once your project is saved, go to the Config section on the Gateway Webpage and select **System > Projects**.



On this page ...

- [Project Settings](#)
 - [Project Settings Table](#)
- [Managing Projects](#)
 - [Viewing the Details of a Project](#)
 - [Editing a Project](#)
 - [Deleting a Project](#)
 - [Copying a Project](#)
 - [Renaming a Project](#)
 - [Exporting a Project](#)

Click on the **Edit** to the right of the project name.

Config > System > Projects					
« < 1 of 1 > »					
Filter <input type="text" value="type to filter"/>		View 20 ▾			
Name	Description	Enabled	Inheritable	Parent project	Actions
P_Class_Test_SJP		true	false		More ▾ Edit
Parent_Project		true	true		More ▾ Edit
Project_East_A	East plant test project	true	false		More ▾ Edit

This opens up the Project Settings page.

Project Settings

Name *	Project_East_A Choose a name to identify this project.
Description	East plant test project
Title	The title for the project. This can contain more characters than the name (space, etc), and will be used to represent the project to users. If empty, the name will be used.
Enabled	<input checked="" type="checkbox"/> A disabled project will not be active on the Gateway, but will remain editable in the Designer.
Inheritable	<input type="checkbox"/> Inheritable projects are not runnable as a stand-alone project, but are intended to provide shared resources to one or more child projects.
Parent Project	Project_West_Templates

Connections

User Source	default
Default Database	DB
Default Tag Provider	default

Project Settings Table

Project Settings	
Property	Description
Name	Name of the project (read only).
Description	Brief description of the project (optional). This description can be viewed on the Open/Create Project screen when you hover over the Information icon.
Title	Title for the project (optional). This is the name that will be displayed on the launch page of the Gateway and in the runtime Client or Session. There are no restrictions to special characters or spaces. If no title is specified, the project name will be used instead.
Enabled	A disabled project will not be active on the Gateway, but will remain editable in the Designer.
Inheritable	Inheritable projects are not runnable as a stand-alone project, but are intended to provide shared resources to one or more child projects.
Parent	Each project can have a parent project, and will inherit all of the resources of that parent project.

Project	
Connections	
User Source	A group of users with their associated roles. Security policies are defined in terms of these roles.
Default Database	Database to be used for historical data logging, reporting, storing alarm logs, and Tags storage. You can also query existing data and update data in the database.
Default Tag Provider	Identifies a Tag database (a collection of Tags) and a name.

Managing Projects

Once you have a project or two (or twenty) set up, you might want to change some of the base settings. You can manage the projects from the Config section of Gateway by going to **System > Projects**. The Projects window displays all your projects. To the right of the project name, click the **More** button, and here you can Edit, Delete, view Details, Copy, Rename or Export your project.

Note: You can manage many of these settings in the Designer. See the [Project Properties](#) page for more information.

Name	Description	Enabled	Inheritable	Parent project	Actions
Compressor		true	false		More Edit
Controller		true			More Delete Edit
Controller1		true			More Edit
ExportTags		true			More Edit
NewProject		true			More Edit
NewProject1		true			More Edit
NewProject2		true	false	global	More Edit

Viewing the Details of a Project

To view the project details, go to the Config section of the Gateway webpage, and select **System > Projects**. Find your project, press the **More** button to the right of the project name, and select **Details**. The Details link on the Project window takes you to the Project details for '**<Project Name>**' page where you can choose to view the [Project Settings](#).

Editing a Project

To edit some of the project settings, go to the Config section of Gateway, and then to **System > Projects**. You will see a list of all your projects. Click **Edit** to the right of project name in the list. The Project - Edit page is displayed. You can now change some of the project settings such as project Description, Title, Enabled, etc. You can also change connections to the User Source, Default Database, and Default Tag Provider.

Caution: Recommend Changing the Title, not the Project Name

It is not advisable to change the Project Name after it's been created, instead, change the Title property if you want to change how the project appears. Shortcuts that refer to the project will no longer work if the project name is changed.

Many of these settings can be modified in the Designer too, but it is important to note that you cannot change any of these project settings from the Gateway if the project is currently open in a Designer.

Deleting a Project

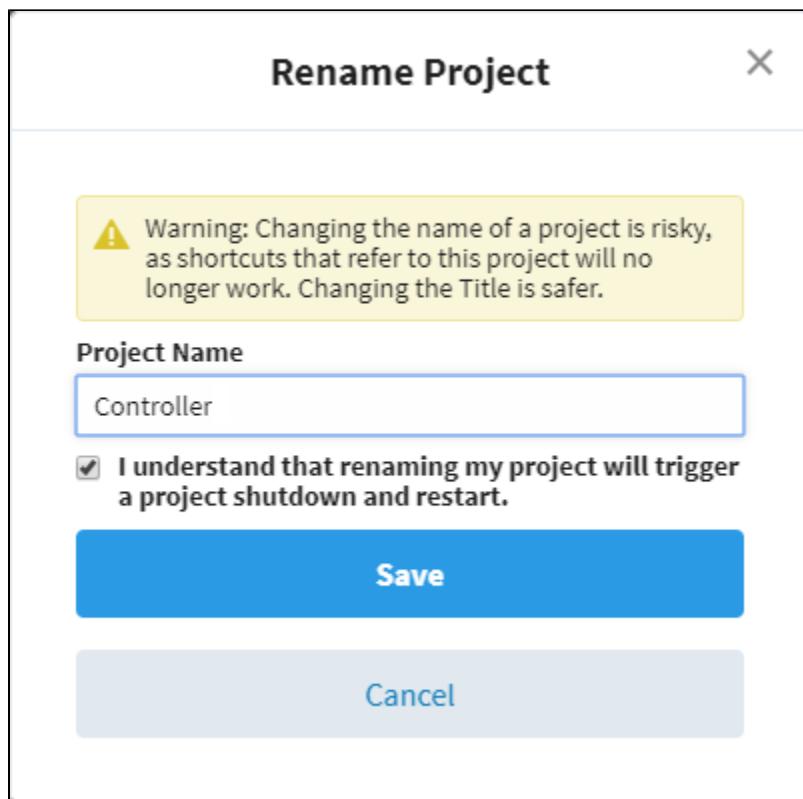
Be aware that once you delete a project, it cannot be undone, a deleted project is gone forever (unless it can be recovered from a [project export](#), or a [Gateway backup](#)). Always make a [project export](#) or [Gateway Backup](#) before deleting a project. To delete a project, go to the Config section of Gateway, and then to **System > Projects**. You will now see a list of all your projects. To the right of the project name in the list, look for **Delete** and click on it. This deletes your project.

Copying a Project

Copying your project is useful when you need a *snapshot* of a project before starting major changes, or for creating a starting point for a new project based on an old one. To copy a project, go to the Config section of Gateway, and then to **System > Projects**. You will now see a list of all your projects. To the right of the project name in the list, click **Copy**. This creates a copy of your project.

Renaming a Project

Changing the name of a project is risky. Shortcuts that refer to the project will no longer work. It is advisable to change the Title which is much safer. If you decide to rename a project, go to the Config section of Gateway and then to **System > Projects**. To the right of the project name in the list, click **Rename**. This will open a Rename Project window. Read it carefully, and if you want to proceed, enter a new project name, mark the checkbox acknowledging that the project will shutdown and restart, then click **Save**.



Note: Renaming a project triggers a project shutdown and restart.

Exporting a Project

Project Export is a project backup. It takes less time than a Gateway backup and it's smaller. This exports your project as a **.zip** file. Once you have an exported file of your project, you can take it to any other Gateway and merge it in with other projects. Refer to the [Project Export and Import](#) page to learn exactly what is included in a project export.

Related Topics ...

- [Project Properties](#)

- Project Export and Import

Project Export and Import

Project backup and restoring from a project backup is referred to as Project Export and Import. Projects are exported individually, and only include project-specific elements shown in the list below. They **do not** include Gateway level configurations, like database connections, Tag Providers, and Tags. The exported file (.zip) is used to restore / import a project.

The resources listed below are included in a project export.

- Alarm Pipelines
- Named Queries
- Perspective Properties
- Perspective Views
- Project Properties
- Reports
- Sequential Function Charts
- Transaction Groups
- Vision Client Tags
- Vision Windows
- Vision Templates
- Client Event Scripts
- Gateway Event Scripts

There are two primary ways to export and import a project:

- **Gateway Webpage** - exports and imports the entire project.
- **Designer** - exports and imports only those resources that are selected.

This page describes how to create a project export and import a project from an exported file.

On this page ...

- Project Export
 - Export a Project from the Gateway
 - Export a Project from the Designer
- Project Import
 - Import a Project from the Gateway Webpage
 - Import a Project from the Designer



INDUCTIVE
UNIVERSITY

Making Project Backups

[Watch the Video](#)

Project Export

A Project Export is a little smaller and takes even less time to run than a Gateway backup! Once you have an exported file of your project, you can take it to any other Gateway and merge it in with the other projects. This makes it simple to keep a development server and push your projects after you complete them. The [Enterprise Administration Module](#) can even do this for you. You can perform a project export from two locations, the Gateway Webpage and the Designer. This section describes how to create a project export from the Gateway Webpage and the Designer.

Export a Project from the Gateway

Making an export from the Gateway Webpage, exports the entire project to a (.zip file). The exported file only includes resources from the project. Notably missing from the project export are any Gateway resources, even if they are presented in the Designer such as Ignition Tags. They need to be exported separately. Refer to [Exporting and Importing Tags](#).

Caution: Remember that Tags and Gateway level configurations (such as device connections, database connections, tags, etc) are not included in a Project Backup. Those resources are only exported in a [gateway backup](#).

1. Go to the **Config** tab of the Gateway Webpage, and click on **System > Projects**.
2. The **Projects** screen will be displayed and you can see your existing projects.

To the right of project name, click the **More** button and select **Export**. This exports your project as a .zip file. The exported file will have the project name, date, and a 4 digit unique number followed by the file extension (i.e., Compressor_2019-03-21_1123.zip).

The screenshot shows the Ignition Config interface with the 'Config' tab selected. In the left sidebar, under the 'SYSTEM' section, the 'Projects' option is highlighted. The main area displays a table of projects with columns: Name, Description, Enabled, Inheritable, Parent project, and Actions. The 'Actions' column contains buttons for Delete, Details, Copy, Rename, More, and Edit. A red box highlights the 'Export' button for the project named 'NewProject'.

3. By default, the .zip file is saved in your web browser's default Downloads folder. This export of the project can be imported to any other instance of Ignition that is running the same version or later.

Export a Project from the Designer

When making a project export from the Designer, you get to choose which project resources are added to the export file. The export does not include any Gateway resources. When restoring a project export in the Designer, you also get to choose which project resources are restored into the currently open project from the exported .zip file.

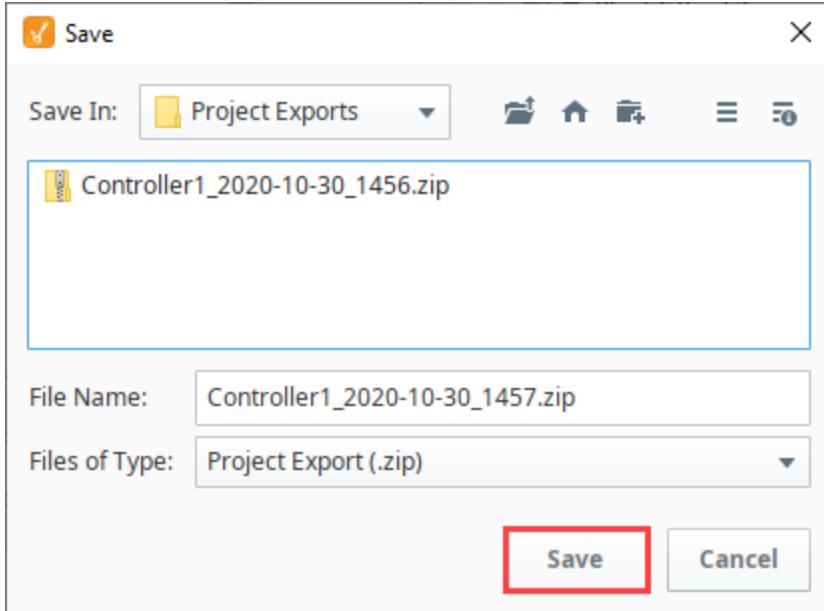
1. From the Designer, go to the top menubar, click on the **File** menu and select **Export**.
2. The Export screen will be displayed and you'll see a list of project resources to export. By default, local project resources are selected. Unselect any resources you don't want to include in the export.

When exporting a project, there is a **Send to Project** button that allows you to directly send the project (or resources in a project) to another project on the same gateway. There is a '**local**' project export option, meaning only local (non-inherited) project resources are selected and will be exported.

Click **Export**.

The screenshot shows the Ignition Designer application. On the left, the 'File' menu is open, with the 'Export...' option highlighted. To the right, a modal dialog titled 'Export Project Resources' is displayed. The 'Select Resources' section contains a tree view of project resources, with several items checked. A red box highlights the 'Perspective Views' item in the tree. At the bottom of the dialog are buttons for 'Select All', 'Local', and 'None'. Below these are 'Export' and 'Send to Project' buttons.

3. This opens a Save dialog window. Select a folder and click **Save**.



4. This export of the project can be imported to any other instance of Ignition that is running the same version or later.

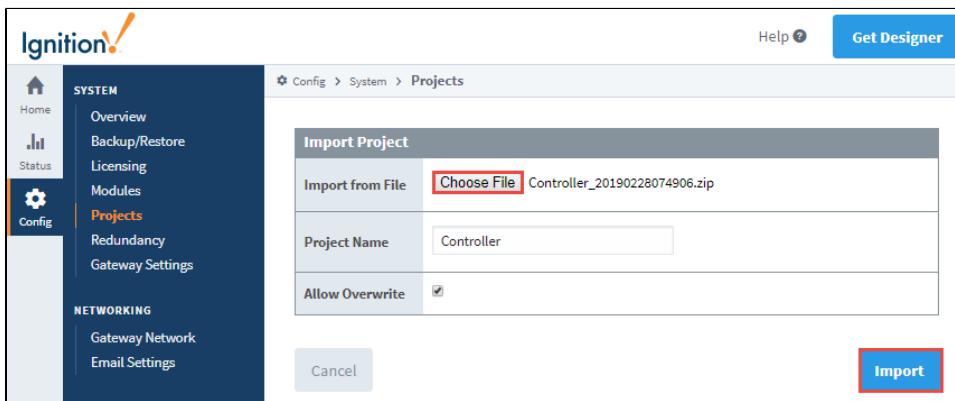
Project Import

When you restore / import a project from an exported file in the Gateway Webpage, it will be merged into your existing Gateway. If there is a naming collision, you have the option of renaming the project or overwriting the project. Project exports can also be restored / imported in the Designer. This will even allow you to select which parts of the project import you want to include and will merge them into the currently open project.

Import a Project from the Gateway Webpage

When restoring a project from the Gateway Webpage, Ignition imports the entire project from an exported file.

1. Go to the **Config** tab of the Gateway Webpage, and click on **System > Projects**.
2. The **Projects** screen will be displayed and you can see your existing projects. At the bottom of the screen, click the **Import project...** link.
3. A second Projects screen will open. Click on **Choose File**, find your exported .zip file from your browser, and then click **Open**.
4. Enter the **Project Name**. If there is already a project with the same name, the **Import button** will be grayed out preventing you from importing the file. The system gives you the option to **Rename** the project you want to import, or **Overwrite** the existing project.
5. Click **Import**.



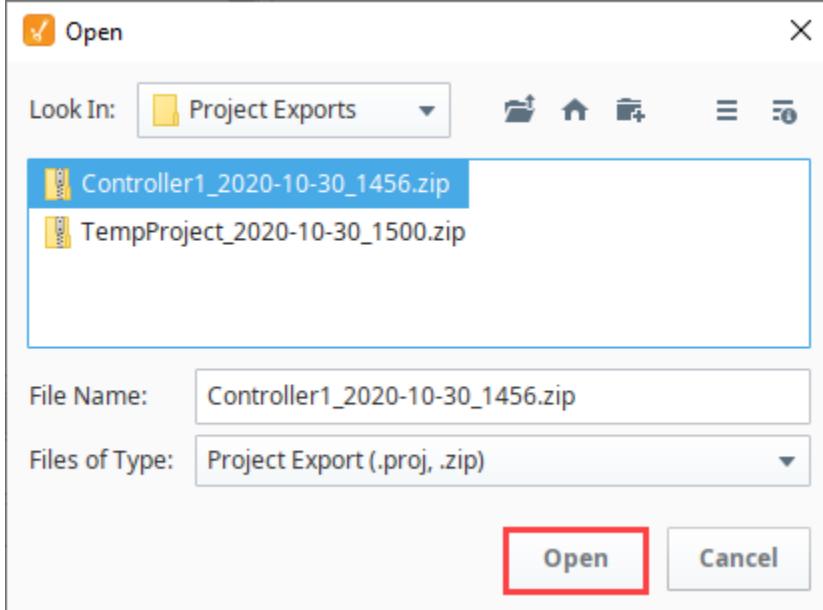
6. If you are importing a new project, you will see your new project added to the list along with the other projects.

Import a Project from the Designer

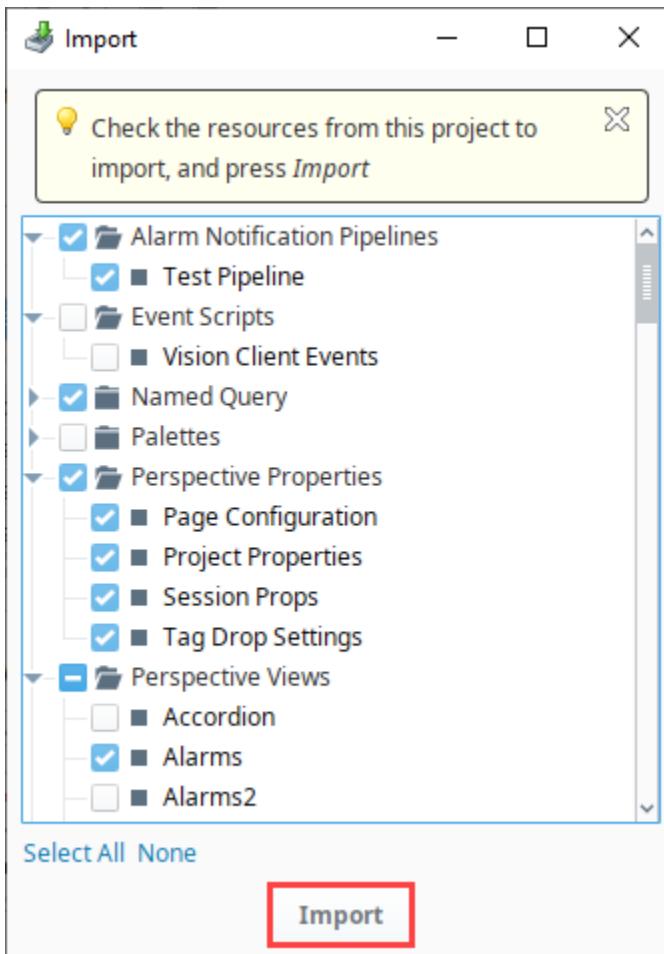
When restoring a project from the Designer, you can select the entire project or parts of the project, and merge them into the one you are currently working on. Ignition will let you choose which resources to import.

1. In Designer, select **File > Import**.

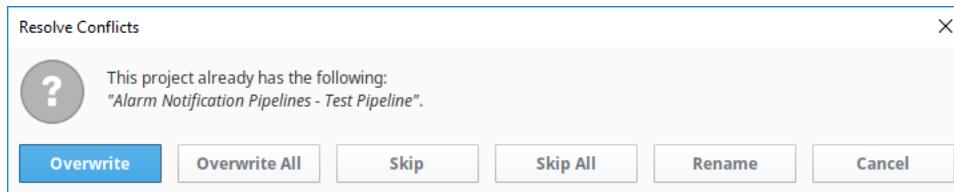
2. The dialog screen will open, select the project export file (.zip) from the default folder, and click **Open**.



3. By default, all resources are selected. Unselect the resources from the project that you do not want to import. Click the **Import** button.



4. If there are any conflicts when importing your project resources, a dialog box will appear and ask you to either **Overwrite**, **Overwrite All**, **Skip**, **Skip All**, **Rename** the affected items, or **Cancel**. Once all conflicts are resolved, **Save** your project.

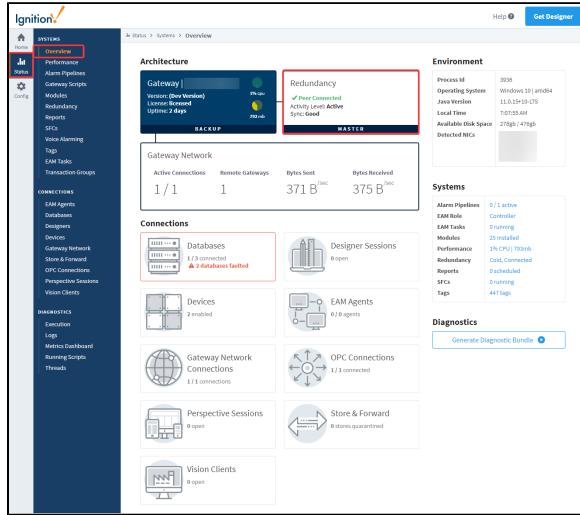


Related Topics ...

- [Gateway Backups and Restore](#)
- [Exporting and Importing Tags](#)
- [Project Inheritance](#)

Ignition Redundancy

Ignition redundancy supports a 2-node system, meaning there are two copies of the Gateway running. One node is the Master Gateway and the other is the Backup Gateway or backup node. All projects, Gateway settings, etc., are shared between nodes. The master node manages the configuration then replicates it to the backup node.



On this page ...

- [Redundancy Status Page Overview](#)
- [Node Communication](#)
- [Configuration Synchronization](#)
- [Runtime State Synchronization](#)
- [Status Monitoring](#)
- [Historical Logging](#)
- [Client Failover](#)
 - [Vision Clients](#)
 - [Perspective Sessions](#)



INDUCTIVE
UNIVERSITY

How Redundancy Works

[Watch the Video](#)

Updating or Patching a Redundant Ignition Pair

Learn about updating redundant servers and how to make the process a success.

[Link to Knowledge Base Article](#)

When you have redundant systems in place, you can get detailed status information by going to Gateway webpage and selecting **Status > Redundancy** to view the system's status and events.

Ignition!

Help ? Get Designer

SYSTEMS

- Home
- Status (Redacted)
- Config
- Overview
- Performance
- Alarm Pipelines
- Gateway Scripts
- Modules
- Redundancy**
- Reports
- SFCs
- Voice Alarming
- Tags
- EAM Tasks
- Transaction Groups

CONNECTIONS

- EAM Agents
- Databases
- Designers
- Devices
- Gateway Network
- Store & Forward
- OPC Connections
- Perspective Sessions
- Vision Clients

DIAGNOSTICS

- Execution
- Logs
- Metrics Dashboard
- Running Scripts
- Threads

Status > Systems > Redundancy

Configuration

Role: Backup

Peer Connected: Yes

Active Uptime: N/A

Redundancy Properties

Activity Level	Cold
Synchronization Status	Good
Local Address	[Redacted]
Peer Address	[Redacted]

Force Re-Sync | **Request Failover**

Redundancy Providers

Provider Name	Last Pull	Last Pull Duration	Last Apply Duration	Full Sync Needed?	System restart required
Alarm Events	2 days	12 ms	3 ms	false	N/A
Alarm Notification	2 days	10 ms	0 ms	false	N/A
Alarm Shelf	2 days	10 ms	6 ms	false	N/A
Internal DB	2 days	25 ms	13 ms	false	N/A
Project 'child'	never	0 ms	0 ms	false	N/A
Project 'global'	never	0 ms	0 ms	false	N/A
Project 'samplequickstart'	never	0 ms	0 ms	false	N/A
Project 'test'	never	0 ms	0 ms	false	N/A
Project Manager	2 days	12 ms	109 ms	false	N/A
SFC	2 days	10 ms	6 ms	false	N/A
Tag provider 'Sample_Tags'	never	0 ms	0 ms	false	N/A
Tag provider 'System'	2 days	71 ms	43 ms	false	N/A
Tag provider 'default'	never	0 ms	0 ms	false	N/A

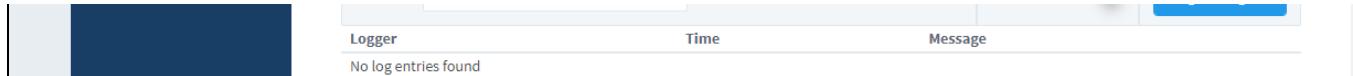
System Events

Severity	Time	Event
Info	16Sep2022 06:45:54	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Good, History level=Full
Info	16Sep2022 06:45:54	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
Info	16Sep2022 06:42:35	Redundancy state changed: Role=Backup, Activity level=Active, Project state=OutOfDate, History level=Full
Info	15Sep2022 20:45:33	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
Info	15Sep2022 07:16:53	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Good, History level=Full
Info	15Sep2022 07:16:53	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
Info	15Sep2022 07:15:39	Redundancy state changed: Role=Backup, Activity level=Active, Project state=OutOfDate, History level=Full
Info	14Sep2022 15:57:09	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
Info	14Sep2022 10:34:06	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Good, History level=Full
Info	14Sep2022 10:34:05	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Unknown, History level=Full
Info	14Sep2022 10:34:05	[alarm_notification_] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
Info	14Sep2022 10:34:05	[com.inductiveautomation.sfc.ChartManager] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
Info	14Sep2022 10:34:05	[alarm_system_] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
Info	14Sep2022 10:34:05	[alarmshelf_] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
Info	14Sep2022 10:34:05	[tagprovider-System] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).

Log Activity

Min level: INFO

Live Values: Merge to Logs



Redundancy Status Page Overview

The following feature is new in Ignition version **8.1.21**
[Click here](#) to check out the other new features

The Redundancy Status page has been overhauled starting in 8.1.21. This section breaks down the Redundancy Status Page, describing each feature.

Metrics and Information

The screenshot shows the Redundancy Status Page. At the top right is a 'Configuration' button. Below it are three cards: 'Role' (Backup), 'Peer Connected' (Yes), and 'Active Uptime' (N/A). A large box labeled 'Redundancy Properties' contains a table with four rows: Activity Level (Cold), Synchronization Status (Good), Local Address (redacted), and Peer Address (redacted). At the bottom of this box are two buttons: 'Force Re-Sync' and 'Request Failover'.

Redundancy Properties	
Activity Level	Cold
Synchronization Status	Good
Local Address	[REDACTED]
Peer Address	[REDACTED]

The top section of the Redundancy Status page gives information about the Gateway and your redundancy setup, including:

- Redundancy configuration settings
- The current Gateway's role in the Redundant pair
- If the Gateway has a peer connected
- The current node's uptime after failover
- Current redundancy settings
- Force re-sync and failover options

Redundancy Providers Statistics

Redundancy Providers

<u>Provider Name</u>	<u>Last Pull</u>	<u>Last Pull Duration</u>	<u>Last Apply Duration</u>	<u>Full Sync Needed?</u>	<u>System restart required</u>
Alarm Events	2 days	12 ms	3 ms	false	N/A
Alarm Notification	2 days	10 ms	0 ms	false	N/A
Alarm Shelf	2 days	10 ms	6 ms	false	N/A
Internal DB	2 days	25 ms	13 ms	false	N/A
Project 'child'	never	0 ms	0 ms	false	N/A
Project 'global'	never	0 ms	0 ms	false	N/A
Project 'samplequickstart'	never	0 ms	0 ms	false	N/A
Project 'test'	never	0 ms	0 ms	false	N/A
Project Manager	2 days	12 ms	109 ms	false	N/A
SFC	2 days	10 ms	6 ms	false	N/A
Tag provider 'Sample_Tags'	never	0 ms	0 ms	false	N/A
Tag provider 'System'	2 days	71 ms	43 ms	false	N/A
Tag provider 'default'	never	0 ms	0 ms	false	N/A

The next section of the Redundancy Status page details metrics about applicable redundancy providers.

Data that is presented include:

- The name of the provider
- When the provider was last pulled
- The latest time duration (or how long it takes) for the Gateway to get data from the provider
- The latest time duration (or how long it takes) for the Gateway to apply the data received from the provider
- Whether a full sync is needed
- Whether a system restart is needed

System Event Information

System Events

Severity	Time	Event
1	16Sep2022 06:45:54	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Good, History level=Full
1	16Sep2022 06:45:54	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
1	16Sep2022 06:42:35	Redundancy state changed: Role=Backup, Activity level=Active, Project state=OutOfDate, History level=Full
1	15Sep2022 20:45:33	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
1	15Sep2022 07:16:53	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Good, History level=Full
1	15Sep2022 07:16:53	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
1	15Sep2022 07:15:39	Redundancy state changed: Role=Backup, Activity level=Active, Project state=OutOfDate, History level=Full
1	14Sep2022 15:57:09	Redundancy state changed: Role=Backup, Activity level=Active, Project state=Good, History level=Full
1	14Sep2022 10:34:06	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Good, History level=Full
1	14Sep2022 10:34:05	Redundancy state changed: Role=Backup, Activity level=Cold, Project state=Unknown, History level=Full
1	14Sep2022 10:34:05	[_alarm_notification_] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
1	14Sep2022 10:34:05	[com.inductiveautomation.sfc.ChartManager] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
1	14Sep2022 10:34:05	[_alarm_system_] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
1	14Sep2022 10:34:05	[_alarmshelf_] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).
1	14Sep2022 10:34:05	[tagprovider-System] Will perform full pull. Reason: major change detected between nodes (sync provider UUIDs did not match).

The third section of the Redundancy Status page displays a table that will log system events whenever a full sync is required, helping to establish a timeline of when full sync events were requested.

Information displayed in this table includes:

- How severe the system event was/is
- When the system event occurred
- The reason for the system event

Logging Activity

Logger	Time	Message
No log entries found		

The final section of the Redundancy Status page shows logger activity and allows users to enable DEBUG and TRACE logs for a specific redundancy provider.

Features of the log activity table include:

- Minimum logging level. Options are:
 - INFO
 - DEBUG
 - TRACE
- An option to merge logs to the main diagnostic log viewer
- The specified logger
- The log's timestamp
- The issue being logged

Node Communication

The master and backup nodes communicate over TCP/IP. Therefore, they must be able to see each other over the network, through any firewalls that might be in place. All communication goes from the backup to the master node over the gateway network (default **port 8088** without SSL, **port 8060** with SSL). Therefore, that port must allow TCP listening on the master machine.

Configuration Synchronization

The master node maintains the official version of the system configuration. You must make all changes to the system on the master Gateway, the backup Gateway does not allow you to edit properties. Similarly, the Designer only connects to the master node.

When changes are made on the master, they are queued up to be sent to the backup node. When the backup connects, it retrieves these updates, or downloads a full system backup if it is too far out of date.

If the master node has modules that aren't present on the backup, they are sent across. Both types of backup transfers, **data only** and **full**, will trigger the Gateway to perform a soft reboot.

Runtime State Synchronization

Information that is only relevant to the running state, such as current alarm states, is shared between nodes on a differential basis so that the backup can take over with the same state that the master had.

On first connection or if the backup node falls too far out of sync, a full state transfer is performed. This information is light-weight and does not trigger a Gateway restart.

The following feature is new in Ignition version **8.1.19**
[Click here](#) to check out the other new features

After the Master Gateway and Backup Gateway reestablishes a redundancy connection, the Backup Gateway will check if it has any conflicting data compared to the Master Gateway's data. If the Backup Gateway has instances of conflicting data, the Backup Gateway will drop those instances in favor of the Master Gateway's data.

This feature was changed in Ignition version **8.1.31**:

The above behavior changes depending on if the Use Active Uptime to Resolve Conflicts property is toggled. If this property is toggled, the redundancy system will compare the Master and Backup nodes' active uptimes to each other, and use the data from the node with higher active uptime. This behavior can result in the Master overwriting its own data with data from an active Backup node.

Status Monitoring

Once connected, the nodes begin monitoring each other for liveliness and configuration changes. While the master is up, the backup runs according to the **stand by activity level** in the settings.

When the master cannot be contacted by the backup for the specified amount of time, it is determined to be down and the backup assumes responsibility. When the master becomes available again, responsibility is dictated by the recovery mode and the master either takes over immediately or waits for user interaction.

Historical Logging

Historical data presents a unique challenge when working with redundancy because it is never possible for the backup node to know whether the master is truly down or simply unreachable. If the master was running, but unreachable due to a network failure, the backup node becomes active and begins to log history at the same time as the master, who is still active.

In some cases this is OK because the immediate availability of the data is more important than the fact that duplicate entries are logged. But in other cases, it's desirable to avoid duplicates, even at the cost of not having the data available until information about the master state is available.

Ignition redundancy provides for both of these cases, with the **backup history level**, which can be either **Partial** or **Full**.

- In **Full** mode, the backup node logs data directly to the database.
- In **Partial** mode, however, all historical data is cached until a connection is reestablished with the master. At that time, the backup and master communicate about the uptime of the master, and only the data that was collected while the master was truly down is forwarded to the database.

Client Failover

This feature was changed in Ignition version 8.1.17:

Failover to the other redundant node is now allowed if the nodes have different platform versions, which will allow attached clients to remain connected to at least one node during a redundant pair upgrade.

Vision Clients

All Vision clients connect to the active node. When this system fails and is no longer available, they automatically re-target to the other node. The reconnection and session establishment procedures are handled automatically, but the user is notified that they have been transferred to a different node so that they can notify the system administrator that the system may need attention.

Perspective Sessions

Like Vision clients, Perspective sessions connect to the active node. When connection to the active node is lost, or the activity level of the Gateway changes from **active**, the session will simultaneously attempt to:

1. Re-establish the connection to the Gateway it was connected to, and check to make sure its activity level is **active**.
2. Monitor the backup Gateway. If the backup Gateway becomes reachable and **active** before the connection to the active Gateway can be re-established, the Perspective session navigates in the browser to the same project and page on the backup Gateway.

In This Section ...

Setting Up Redundancy

In redundancy, both nodes will share the exact same configuration state. When a Backup node connects to a Master node, the Backup will attempt to synchronize itself with the Master. Therefore, before you set up for redundancy the following should be considered:

1. Start with a fresh install for the Backup node.

Because the current configuration of the Backup node will be overwritten, make sure that it does not contain anything valuable. It is a good idea to export any projects that are unique to the Backup before enabling redundancy.

2. All system configurations relative to the Master node must also resolve on the Backup node.

For example, OPC UA connections and database connections must use addresses that resolve from both nodes, or any OPC-COM servers must be installed and configured identically on both nodes. This means using "localhost" in any of the database connections won't work. You should use the IP address of the computer instead.

3. Configure firewalls between the redundancy nodes.

Redundant systems need TCP connectivity between each other on the default Gateway network ports. Turning off software firewalls or adding special exception rules for each others' addresses is required. The default Gateway Network port is **port 8088** (without SSL), and **port 8060** (with SSL), and the Backup node must be able to send outgoing data on that port. The port can be changed from Gateway Network settings.

Note: Two [Edge Gateways](#) can be set up with redundancy. An Edge Gateway can only failover to another Edge Gateway (not a standard Ignition Gateway). Also, an Edge Gateway cannot be used as backup to a Standard Ignition Gateway.

Note:

While the OS platform (i.e., Windows, OS X, Linux) for the Master and Backup can differ, it is recommended to have similar OS platforms. If the OS platforms do differ, the Windows machine should be the master system or else the Force Failover option will not work.

However, different versions of the same operating system such as Windows 10 and Windows 8 or OSX 10 and OSX 11 have full functionality.

On this page ...

- [On the Master Gateway](#)
- [On the Backup Gateway](#)
- [On the Master Gateway](#)
- [Redundancy Settings](#)
- [Troubleshooting](#)
 - [Redundancy Connectivity](#)
 - [Advanced Troubleshooting](#)



INDUCTIVE
UNIVERSITY

Setting Up Redundancy

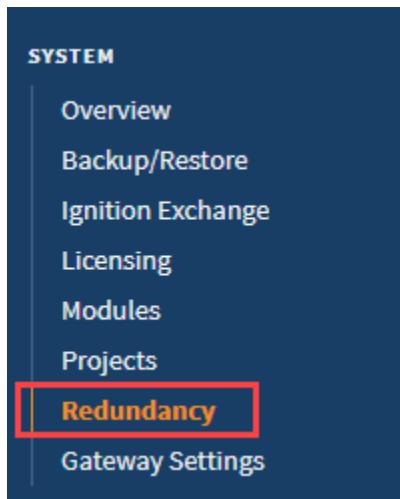
[Watch the Video](#)

On the Master Gateway

1. Go to the [Config](#) section of the Master Gateway Webpage.

2. Select [System > Redundancy](#).

The Redundancy and Network Configuration page is displayed showing different sections and settings. See the table below for a description of all settings.



3. Change the following settings:
- Under Redundancy Settings, set Mode to **Master**.
 - Optionally, configure any desired settings under Master Node Settings.

The screenshot shows the 'Redundancy Settings' section of the configuration interface. The 'Mode' dropdown is open, displaying four options: 'Independent', 'Independent', 'Backup', and 'Master'. The 'Master' option is highlighted with a red box. Below the dropdown, a tooltip provides information about the mode selection: 'This node's role. There should be one master and one backup. "Independent" turns off redundancy.' Other settings visible include 'Standby Activity Level' (set to 'Cold'), 'Startup Connection Allowance' (set to 30000), and 'Sync Timeout' (set to 60).

4. Click **Save Changes**. The Confirm change to Redundancy Settings page is displayed.

The screenshot shows a confirmation dialog titled 'Confirm Change to Redundancy Settings'. The message inside the dialog reads: 'Altering redundancy settings can potentially have a large impact on the system, and may require a restart. Are you sure you want to continue?'. At the bottom of the dialog are two buttons: a blue 'Confirm' button and a grey 'Cancel' button.

5. Click **Confirm** to apply your settings.
 6. Go to the Config tab and select **System > Redundancy** to ensure the redundancy mode and state is properly set.

On the Backup Gateway

Do the exact same steps 1-6 above on the Backup Gateway Webpage, except replace step 3 with the following:

- Under Redundancy Settings, set **Mode to Backup**.
- Under Backup Node Settings, configure the Master Node Address and Port to point to the Master Gateway. The Master Node address should be a hostname or IP address. The Port setting (assuming default configurations) should be **8060** if using SSL,

otherwise **8088**.

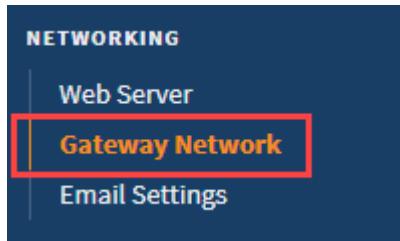
Config > System > Redundancy and Network Configuration

Redundancy Settings

Mode	Independent Backup Master	Defines this node's role. There should be one master and one backup. Independent turns off redundancy.
Standby Activity Level	Cold	How the node should run when it is not currently the active node. If cold, the node will perform minimal operations until it becomes active. If warm, the node will run at a high level, reducing failover times.
Startup Connection Allowance	30000	The time in milliseconds that the system will wait at startup for a connection before making a decision on the node's responsibility level. (default: 30000)
Sync Timeout	60	The maximum time (in seconds) allowed for a redundancy sync operation. Sync operations that exceed this value will time out. (default: 60)

On the Master Gateway

1. Return to the Config section of the Master Gateway Webpage.
2. Select **Networking > Gateway Network**.



3. Navigate to the **Incoming Connections** tab. You should see a new incoming connection from the Backup Gateway. Find the connection, select **More**.



If you are not using SSL and your connection isn't displayed, make sure the **Require SSL** box is unchecked under the **General Settings** tab. Clearing the checkmark and returning to the **Incoming Connections** tab will now populate your new incoming connection.

General Settings Outgoing Connections Incoming Connections Diagnostics Queue Management

Main

Enabled	<input checked="" type="checkbox"/> (default: true)
Require SSL	<input type="checkbox"/> If true, only connections that use SSL to encrypt traffic will be allowed. This setting only applies to incoming connections. (default: true)

4. Select **approve**.

5. To verify the redundancy setup, that is, to ensure the Master and the Backup Gateways are connected, go to the Status tab of the Gateway Webpage and click on **System > Redundancy**. The Redundancy page will show the connected nodes and their current states.

After approving the connection, the Backup connects to the Master and downloads a system backup, then restarts. Once the restart is complete, the Backup node is synchronized and in communication with the Master.

Redundancy Settings

All redundancy settings are configured in the Gateway Webpage under the Config tab, **Systems > Redundancy**. Most settings are used by both the Master and Backup nodes, with their individual settings broken out into separate categories.

It is important to know that while the full system configuration is shared between nodes, redundancy settings are not shared between nodes. Therefore, it is perfectly acceptable to have different values for the same settings on the two nodes. For example, it is possible to have a different Standby Activity Level on both nodes, and, of course, the network settings will often be different.

Note: The Master node shares all configuration with the Backup node, and this means that changes cannot be made to your project from the Backup. In fact, the Designer can never be opened from a Backup node, even if the Master is currently offline.

Redundancy Settings	
Mode	<p>Enable or disable redundancy, and specify this node's role. There should be one master and one backup node per redundant pair. Independent turns off redundancy.</p> <p>Independent - Redundancy is not enabled and this Ignition system runs as an independent node.</p> <p>Master - This is the Master node, who listens for a connection from the Backup node, and is in charge of managing system synchronization.</p> <p>Backup - This is the Backup node, who will connect to the Master and receive system updates.</p>
Standby Activity Level	<p>How the node should run when it is not currently the Active node.</p> <p>Cold - The system connects to all OPC servers but does not subscribe to Tag values. The Ignition OPC UA server does not communicate with any device, but third party OPC UA servers may still have device connections. This allows the system to standby without putting additional load on the devices and network. Failover takes slightly longer, as Tags must be subscribed and initialized.</p> <p>Warm - The system runs as if it were active, with the exception of logging data or writing to devices, allowing for faster fail-over.</p>
Failover Timeout	<p>This feature was removed from Ignition in version 8.1.37</p> <p>The time of inactivity, in milliseconds, before the backup assumes responsibility. Default is 10000 milliseconds.</p>
Startup Connecti	

on Allowance	<p>The time in milliseconds that the system will wait at startup for a connection before making a decision on the node's responsibility level. This is used to prevent unnecessary switch over caused by a node starting as active, only to connect and find that the other node is active, resulting in one of the nodes being deactivated. Default is 30000 milliseconds.</p> <div style="border: 1px solid #ccc; padding: 10px;"> <p>Note:</p> <p>It is important to notice that this setting can interfere with the Master Recovery mode:</p> <ul style="list-style-type: none"> • If the Master is active, it will always request the Backup to de-activate. • If this setting is low, or 0, the Master will always become active before connecting to the Backup, and thus "manual recovery" will not be possible. </div>
Sync Timeout	<p>The following feature is new in Ignition version 8.1.22 Click here to check out the other new features</p> <p>The maximum time in seconds allowed for a redundancy sync operation. Sync operations that exceed this value will time out. Default is 60 seconds</p>

Network Settings	
Auto Detect Network Interface	If true, the system will automatically select which network interface to use. Most commonly disabled on systems with multiple network cards, in order to explicitly specify which interface to use. If false, the system will bind itself to the interface of the specified address.
Network Bind Interface	The IP address of the network interface to use for redundancy. Only used if "Auto Detect" is false.

Master Node Settings	
Recovery Mode	<p>How the Master node resumes responsibility after starting again.</p> <p>Automatic - The Master automatically takes back responsibility, and becomes active. The Backup node goes to standby. Manual - The Backup node is allowed to stay active. The Master will become active if the Backup node fails, or if the user requests a switchover from the Gateway configuration page.</p>
Runtime Update Buffer Size	<p>This feature was removed from Ignition in version</p> <p>8.1.21</p> <p>How many "runtime state" updates can be queued in memory before the system stops tracking and a full transfer is performed. These updates represent information that the other node should have in order to have the same running state as the Master when it's forced to take over.</p> <p>This is most often the values of static Tags and the current alarm state. Given that the update buffer is only used once the nodes are connected, the default value is usually fine, and only needs to be increased on systems that may have many alarms that change together, or many static Tag writes.</p>
Config Update Queue Size	The maximum size (in megabytes) of config updates allowed before a full transfer is performed.

Backup Node Settings	
Master Node Address	The address of the Master Ignition system.
Port	The Gateway Network port used by the Master to listen on. For the Backup, the port to connect to on the Master.
Use SSL	Use SSL to connect to the remote machine.

Ping Rate	How often, in milliseconds, to send a message from the Backup to the Master.
Ping Timeout	The maximum time, in milliseconds, allowed for a ping response. Pings that time out are counted as missed pings.
Ping Max Missed	The amount of missed pings that will force the connection to the master to be considered faulted.
Websocket Timeout	The maximum time, in milliseconds, allowed for a new web socket to connect to the Master.
HTTP Connect Timeout	The maximum time, in milliseconds, allowed to establish an HTTP connection to the Master.
HTTP Read Timeout	The maximum time, in milliseconds, allowed to read or send HTTP data to the Master.
History Mode	How history is treated by the Backup system. If Full , history will be stored normally, as it would be on the Master system. If Partially , history will be cached until the Master is available again and the Backup node is able to determine the exact time that the Master was down.
Use Active Uptime to Resolve Conflicts	<p>The following feature is new in Ignition version 8.1.31 Click here to check out the other new features</p> <p>When enabled, the system will resolve data conflicts by examining if the Master node or Backup node have been active for longer. The redundancy data from the longer running active node will be selected, and will overwrite the data on the other node.</p>

Troubleshooting

Redundancy Connectivity

When the two redundant nodes are connected, you will be able to see their state details in the [Status](#) section of the Gateway Webpage. There are also various other places where the redundancy state is shown as **connected**.

If the two nodes cannot connect, check the following:

- Verify that the Master address is correct in the Backup. Try to ping the Master machine from the Backup machine, and verify that you're using the correct address for the network card that the Master is connected through.
- If using system names (or domain names), verify that the name is resolving to the correct address by performing a ping.
- Verify that the firewall on the Master is set to allow TCP traffic to the designated port.
- Verify that the Backup is not connecting and then immediately disconnected for some reason.
- Viewing the error log in the Gateway console section should show this. If errors are occurring at regular intervals, look at the message for an indication of what is happening. An example of a potential problem is when the failover time is set too low for the given network, which results in many socket read timeout exceptions, which in turn leads to many disconnect/reconnect attempts.
- If errors are occurring, but the cause isn't clear, contact [Inductive Automation Support](#).

Advanced Troubleshooting

A variety of loggers can be found under the Gateway console section by going to "Levels" and searching for "Redundancy". By setting these loggers to a finer level, more information will be logged to the console. This is generally only useful under the guidance of Inductive Automation support personnel, though more advanced users may find the additional logged information helpful.

Related Topics ...

- [Database Considerations](#)

Database Considerations

Ignition Database Requirements

Given that many parts of the Ignition system interact with the database, it's important to give some thought as to how it will be used when redundancy is turned on, and the different database architectures that are possible.

When evaluating database architectures for use with Ignition, it's important to look carefully at how the system will use the database. Which pieces are critical? Which pieces are "optional" so that the system continues to function while the database is down? Which pieces can operate in "read-only" mode if necessary?

Ignition uses the database for many purposes. Here are some common areas where they are used, and how availability can impact the system:

Tags

Tags rely on the database for Tags that execute queries. These Tags will error out if the database is unavailable, but the status and control functionality of the system will function on the whole.

History - Tags and Other

All history in Ignition goes through the [Store-and-Forward](#) system, meaning that it will be cached until the database is available. However, while the data is cached, it will be unavailable to view or analyze on the clients. Therefore, when looking at how the database should be set up, it is necessary to determine how crucial rapid-availability of the data is.

Alarming

The alarm status system does not reside in the database, so it will continue to function if the connection is down. [Alarm Journal](#) information will go through the [Store-and-Forward](#) system as history data.

Project Screens

Almost all projects use database access for providing information on screens. These queries will error out as long as the database is unavailable. Screens that only use Tags (in an internal provider) will continue to function, so it would be beneficial to make a distinction between status screens and history screens, if a failover database is not used.

Database Architectures

Single Shared Server

A single database server is used. Any Ignition Gateways will use it, so it is expected to be available even when one of the nodes is not. For that reason, it almost always resides externally, on a separate server machine. This arrangement is the easiest to use with Ignition. A single database connection configured on the master will be replicated to the backup, and both nodes will use the connection as necessary.

Clustered/Replicated Database Servers

There is a wide variety of capabilities supported by the different brands of database servers. To obtain fault-tolerance on the database front, it is usually necessary to have some sort of cluster/replication system in place. However, it can be very important to examine how Ignition is using the databases, and what capabilities the clustering solution provides.

For example, in many replication scenarios, the master database copies data to the backup. The backup can be used for read purposes, but new data inserted will not be replicated back to the master. Therefore, it is possible to have a failover connection to the backup database, so that clients will continue to receive data, but it would be necessary to run in partial history mode, so that the historical data was cached and inserted only to the master database. The failover connection would be set to standard mode, so the primary connection would be used when possible.

In a more complete cluster environment, where writes to either node would be replicated, a sticky failover connection could be used with full history mode.

Pertinent Settings

When working with various database architectures, there are a few settings in various parts of the system that are important.

Database Connection Settings - Failover Datasource

On this page ...

- [Ignition Database Requirements](#)
 - [Tags](#)
 - [History - Tags and Other](#)
 - [Alarming](#)
 - [Project Screens](#)
- [Database Architectures](#)
 - [Single Shared Server](#)
 - [Clustered/Replicated Database Servers](#)
- [Pertinent Settings](#)
 - [Database Connection Settings - Failover Datasource](#)
 - [Clustering Settings - History Mode](#)

Any database connection can have a failover datasource. If the main connection is unavailable, any queries executed on it will pass through to the secondary connection. In this way, a secondary database can be used when the first is not available, and the system will continue to function. It is important to note that everything passed through to the failover will function normally- no special considerations will be made. For example, the system won't cache data for the primary connection, it will forward it to the secondary. In cases where you want to allow reading from the secondary database, but not writing, you can set up another connection directly to the first database, with no failover, and set all of your write operations to use that.

Clustering Settings - History Mode

The history mode dictates how history will be treated when the node is not active. If partial, the data will be cached, and only forwarded when the master node is available. This mode can be used to prevent data from being inserted into a backup database in some cases. This setting can be found on the [Redundancy](#) page under the [Config](#) section of the Gateway.

Related Topics ...

- [Redundant Licensing](#)
- [Setting Up Redundancy](#)

Redundant Licensing

Types of Redundant Licensing

When working with Redundancy, both nodes will require a license. However, there are two approaches that are detailed below. In both cases, the license for the Backup node generally contains most, if not all, the same modules as the Master node. For example, if the Master is storing history with the SQL Bridge module, then the Backup would also require the module in the event a failover occurs, otherwise, data would be lost.

However, non-critical modules do not need to be added to the Backup License. The Master license could make use of the Symbol Factory module, and since the Designer only connects to the Master node, there is no reason to add this module to the Backup node.

On this page ...

- Types of Redundant Licensing
 - Two Standard Licenses - Classic Redundancy
 - Backup-Only Licenses

Two Standard Licenses - Classic Redundancy

Traditionally, Redundancy involves two standard licenses: one license will be applied to the Master Gateway, and the other license will be applied to the Backup Gateway. Since two standard licenses are being used, this approach allows users to disable Redundancy, and set both Gateways to independent modes.

Backup-Only Licenses

Backup-Only licenses are available for purchase. This license forces the Gateway into a Backup mode, and the mode can not be changed while the license is applied. The benefit of this type of license is that they come at a discounted price. For more details, contact your Account Representative.



Redundant Licensing

[Watch the Video](#)

If using a Backup only license, the Mode property on the Backup Gateway's Redundancy page will look like the image below.

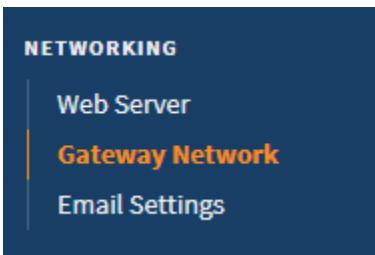
Redundancy Settings	
Mode	<input type="button" value="Backup"/> Enable or disable redundancy, and specify this node's role. There should be one master and one backup node per redundant pair. <input type="button" value="Independent"/> turns off redundancy.
Standby Activity Level	<input type="button" value="Cold"/> How the node should run when it is not currently the active node. If <input type="button" value="cold"/> , the node will perform minimal operations until it becomes active. If <input type="button" value="warm"/> , the node will run at a high level, reducing failover times.
Failover Timeout	<input type="text" value="10000"/> The time of inactivity, in milliseconds, before the backup assumes responsibility. (default: 10000)
Startup Connection Allowance	<input type="text" value="30000"/> The time in milliseconds that the system will wait at startup for a connection before making a decision on the node's responsibility level. (default: 30000)

Instead of the dropdown list that normally appears on this property, the license forced this Gateway into a Backup mode.

If you would like to change the mode of this Gateway, the backup-only license must first be unactivated. Once unactivated, a new license will need to be applied, otherwise, the Gateway will operate in 2 hour trial mode.

Gateway Network

The Gateway Network allows you to connect multiple Gateways together over a wide area network, and opens up many distributed features between Gateways.



The Gateway Network provides the following features:

- A dedicated HTTP data channel that can handle multiple streams of message data.
- The ability to set up a node to act as a proxy for another node.
- Security settings that restrict incoming connections based on a white list or on manual approval of the connection. Incoming connections can also be disabled entirely.
- An available Secure Socket Layer (SSL) mode. When enabled, connections must send SSL certificates to prove their identity. A connection will not be accepted until its SSL certificate is approved.

Gateway Network Features

The Gateway Network opens up certain services for use that make managing multiple Gateways and having them effectively communicate with each other a snap. It also has special security that can restrict certain services from happening in certain zones of the Gateway Network.

Enterprise Administration

The [Enterprise Administration Module](#) (EAM) uses the Gateway Network for message and file transfer, and can monitor network connections for availability. The EAM reports whenever communications are lost via alarm events and system Tags.

Distributed Services

Distributed services included the following:

- Remote Providers: Remote [Realtime](#) and [Historical](#) Tag providers make remotely controlling and storing Tag data even easier.
- Remote Alarming: [Remote Alarming](#) makes notifying all Gateways in the network possible, to quickly and effortlessly track down issues.

Security Zones and Service Security

[Security Zones](#) can be set up to lock down or prevent access to certain parts of Gateways within the Gateway Network.

On this page ...

- [Gateway Network Features](#)
 - [Enterprise Administration](#)
 - [Distributed Services](#)
 - [Security Zones and Service Security](#)
- [Outgoing vs. Incoming Connections](#)
 - [Connections and Servers](#)
 - [Which Server Should I Configure the Outgoing Connection On?](#)
- [General Settings](#)
 - [Main](#)
 - [Security](#)
- [Setting Up a Gateway Network Connection](#)
 - [Main](#)
 - [Ping](#)
 - [Timeouts](#)
 - [Gateway Network Connection Example](#)
- [Deleting Connections](#)
- [Gateway Network Diagnostics](#)
- [Gateway Network Queue Management](#)
 - [Queue Settings](#)
 - [Adjust Queue Message Capacity](#)
- [Gateway Network Proxy Rules](#)
 - [Proxy Rule Settings](#)
 - [Setting up a Proxy Rule](#)



INDUCTIVE
UNIVERSITY

Gateway Network Overview

[Watch the Video](#)

Outgoing vs. Incoming Connections

When using the Gateway Network, you will be working with two type of connections.

- **Outgoing Connections:** To establish communications, create an outgoing connection on the local machine. The outgoing connection always begins the connection process to a remote machine. After the outgoing connection is created, the local machine will attempt to use the connection to establish communications with the remote machine.
- **Incoming Connections:** On the remote machine, an incoming connection will automatically be created when the new connection attempt is detected. For connections where security settings require manual approval, you will need to approve the incoming connection before it can be used. If no security controls have been set, the incoming connection will automatically accept the connection from the local machine and begin sharing data.

Connections and Servers

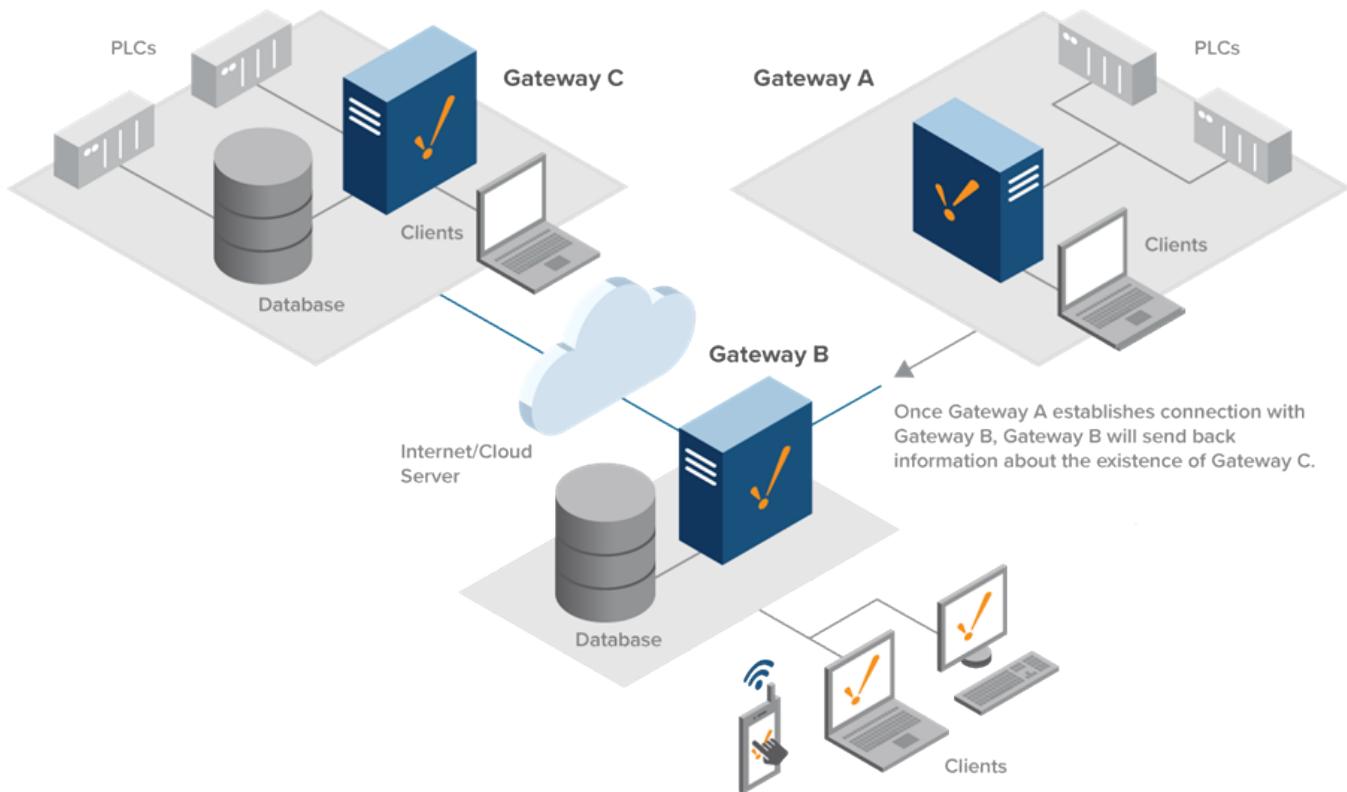
Every machine on the Gateway Network is known as a Server. When you establish a connection to a remote machine, the remote Server sends data about itself and also sends data about any other Servers known to that machine. For example, assume your local machine is GatewayA. The remote machine is known as GatewayB. GatewayB also knows about another remote machine named GatewayC. As soon as your local GatewayA establishes a connection with GatewayB, GatewayB also sends information about the existence of GatewayC.

Modules such as the Enterprise Administration Module (EAM) are aware of this relationship and allow communication between GatewayA and GatewayC, even though there is no direct connection from the local machine to GatewayC.

If you are cloning Gateways to then be connected via Gateway Network, it is important to notice that there is a Gateway unique identifier in %IgnitionInstallationDirectory%/data/.uuid. No two Gateways connected via Gateway Network should share a .uuid. Generally, Gateways are cloned by restoring the same Gateway backup on multiple servers. Since Gateway backups carry their .uuid with them, restoring the same Gateway on multiple servers will result in multiple Gateways having the same .uuid. To get around this, you must stop your Ignition service, delete %IgnitionInstallationDirectory%/data/.uuid, then start your Ignition service so that a new, unique .uuid is generated. Doing this before connecting two cloned Gateways will prevent any .uuid collisions.

Which Server Should I Configure the Outgoing Connection On?

In regards to connecting multiple Gateways over the Gateway Network, there is little difference between an Outgoing and Incoming connection: these terms simply indicate which server the connection was configured on, and are mostly ignored by the rest of Ignition. Thus, assuming GatewayA and GatewayB, configuring an outgoing connection from **A** to **B** is equivalent to configuring an outgoing connection from **B** to **A**. When connecting two Gateways, only a single connection is required between them.



General Settings

Main

The Gateway Network General Settings set the basic rules for the system. By default, these settings are lenient to allow for easy setup, but can be set for security.

Setting	Description

Enabled	Uncheck this checkbox to disable using the Gateway Network.
Require SSL	If true, only connections that use SSL to encrypt traffic will be allowed. This setting only applies to incoming connections. Default is true.
Require Two Way Auth	<p>Enforces two-way SSL authentication. If true, you will need to install the remote machine's certificate on this machine, in addition to manual approval of this machine's certificate on the remote machine.</p> <p>To provide the remote machine's certificate, manually export a certificate from the remote machine's metro KeyStore, located in <installdir>/webserver/metro-keystore. Default KeyStore password is metro, and the alias is metro-key. Place the certificate on the local machine in data/certificates/gateway_network</p> <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> <u>This feature was changed in Ignition version 8.1.14:</u> </div> <p>On Ignition versions 8.1.14 and greater, the certificate is stored at data/gateway-network/client/security/pki/trusted/certs</p>
Send Threads	<p>The maximum number of threads that will be used to upload messages. Applies to outgoing connections. Default is 5.</p> <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> <u>This feature was changed in Ignition version 8.1.37:</u> This property was moved to the Gateway Network Settings > Outgoing Connections tab to be set when creating or editing an outgoing connection. See Connection Send Threads under Setting up a Gateway Network Connection for more details. </div>
Receive Threads	<p>The maximum number of threads that will be used to download messages. Applies to outgoing connections. Default is 5.</p> <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> <u>This feature was changed in Ignition version 8.1.37:</u> This property was moved to the Gateway Network Settings > Outgoing Connections tab to be set when creating or editing an outgoing connection. See Connection Receive Threads under Setting up a Gateway Network Connection for more details. </div>
Processing Queue Limit	Number of received messages that can be held until they are processed by the local system. When this capacity is exceeded, new messages are rejected and errors are reported to the remote Gateway. Applies to incoming connections.
Websocket Idle Timeout	<div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;"> The following feature is new in Ignition version 8.1.3 Click here to check out the other new features </div> <p>The maximum number of milliseconds that a websocket is allowed to remain idle before it is closed. This value should always be set higher than outgoing connection ping rates to avoid premature connection termination.</p>
Proxy Service Call Intercept	<div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;"> The following feature is new in Ignition version 8.1.34 Click here to check out the other new features </div> <p>If enabled on a proxy Gateway, this setting will intercept recurring service enumeration calls that the proxy Gateway normally passes along. If the service enumeration information is available locally on a proxy Gateway, the information is returned directly instead of forwarding the call. Useful when too many service enumeration calls are occurring. Default is false.</p>

Security

You have several options to control security from the Gateway Network settings.

Note: These settings are independent of SSL mode, which is detailed below. To change security settings, go to the Gateway Webpage and navigate to **Config > Networking > Gateway Network** and select the **General Settings** tab.

Setting	Description
Allow Incoming Connections	If false, only outward connections defined on this gateway will be allowed. Uncheck this checkbox to disable all remote machines from being able to establish an incoming connection. To establish any connections with remote machines, you will need to create outgoing connections from this machine. Default is true.
Connection Policy	Dictates what connections are allowed. Options as follows: <ul style="list-style-type: none"> ApprovedOnly - Default mode, where incoming connections are created, but cannot be used to send or receive data until you approve the connection under Gateway Network > Incoming Connections tab. To approve an incoming connection, click the Approve link on the right side of the connection. You can also deny a previously approved connection by clicking the deny link. The approve and deny links will appear next to a connection only if you have enabled the ApproveOnly setting. Unrestricted - Allows all incoming connections unless the Allow Incoming Connections checkbox is unchecked. SpecifiedList - An incoming connection will only be allowed if its server name is on this list. Separate server names with a comma.
Specified List	Connections with a Gateway Name in this list are automatically allowed if the Connection Policy is set to SpecifiedList . Separate Gateway names with a comma.
Allow Proxying	If enabled, this Gateway will be allowed to act as a proxy, and forward requests between Gateways that do not have direct connections. Default is false. <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> <p><u>This feature was changed in Ignition version 8.1.14:</u></p> </div> <p>This setting was replaced in 8.1.14 with Allowed Proxy Hops.</p>
Allowed Proxy Hops	<div style="background-color: #ffccbc; padding: 10px; border: 1px solid #ff9800; margin-bottom: 10px;"> <p>The following feature is new in Ignition version 8.1.14 Click here to check out the other new features</p> </div> <p>The maximum number of proxy hops which could be used to reach the destination Gateway. Any number less than or equal to zero is equivalent to no proxy hops allowed.</p>

Setting Up a Gateway Network Connection

When you create a new outgoing Gateway Network connection, you need to specify the address for the remote server. There are also settings for ping rates and timeouts. The defaults can be used for these fields. The following are all the available settings for setting up an outgoing Gateway Network.

Main

Main	
Host	The address of the remote server, not including the port.
Port	8060 The port of the remote server. (default: 8,060)
Enabled	<input checked="" type="checkbox"/> (default: true)
Use SSL	<input checked="" type="checkbox"/> Use SSL to connect to the remote machine. (default: true)
Description	Documentation about how the connection is used. (default:)
Connection Send Threads	5 The maximum number of threads that will be used to upload messages. Applies to outgoing connections. (default: 5)
Connection Receive Threads	5 The maximum number of threads that will be used to download messages. Applies to outgoing connections. (default: 5)



**INDUCTIVE
UNIVERSITY**

**Setting up a
Gateway Network
Connection**

[Watch the Video](#)

Setting	Description

Host	The address of the remote server, not including the port. Example: 1.2.3.4
Port	The port of the remote server. Default is 8060.
Enabled	Whether this connection is enabled. Default is true.
Use SSL	Use SSL to connect to the remote machine. Default is true.
Description	<p>The following feature is new in Ignition version 8.1.26 Click here to check out the other new features</p> <p>Documentation about how the connection is used. The description will not be sent to other Gateways.</p>
Connection Send Threads	<p>The following feature is new in Ignition version 8.1.37 Click here to check out the other new features</p> <p>The maximum number of threads that will be used to upload messages for this outgoing connection. Default is 5.</p>
Connection Receive Threads	<p>The following feature is new in Ignition version 8.1.37 Click here to check out the other new features</p> <p>The maximum numbers of threads that will be used to download messages for this outgoing connection. Default is 5.</p>

Ping

Ping	
Ping Rate	1000 How often, in milliseconds, to send a ping to a remote machine. (default: 1,000)
Ping Timeout	300 The maximum time, in milliseconds, allowed for a ping response. Pings that time out are counted as missed pings. (default: 300)
Missed Pings	30 The amount of missed pings that will force the connection to be considered faulted. (default: 30)

Setting	Description
Ping Rate	How often, in milliseconds, to send a ping to a remote machine. Default is 1,000.
Ping Timeout	The maximum time, in milliseconds, allowed for a ping response. Pings that time out are counted as missed pings. Default is 300.
Missed Pings	The amount of missed pings that will force the connection to be considered faulted. Default is 30.

Timeouts

Timeouts	
Websocket Connect Timeout	10000 The maximum time, in milliseconds, allowed for a new web socket to connect to a remote machine. (default: 10,000)
Http Connect Timeout	10000 The maximum time, in milliseconds, allowed to establish an HTTP connection to a remote machine. (default: 10,000)
HTTP Read Timeout	60000 The maximum time, in milliseconds, allowed to read or send HTTP data to a remote machine. (default: 60,000)

Setting	Description
Websocket Connect Timeout	The maximum time, in milliseconds, allowed for a new web socket to connect to a remote machine. Default is 10,000.

HTTP Connect Timeout	The maximum time, in milliseconds, allowed to establish an HTTP connection to a remote machine. Default is 10,000.
HTTP Read Timeout	The maximum time, in milliseconds, allowed to read or send HTTP data to a remote machine. Default is 60,000.

Gateway Network Connection Example

To establish a basic communication link between two Gateways, first log into the Gateway where you want to establish the outgoing connection. For this example, we use an SSL connection.

1. On the Gateway Webpage, navigate to **Config>Networking>Gateway Network**.
2. Click on the **Outgoing Connections** tab. Click the **Create new Outgoing Gateway Connection** link.

3. In the **Host** field, enter the network address of the remote server.
4. In the **Port** field, enter the SSL port used by the remote server. By default, this is set to **8060** (which is defined /data/gateway.xml).

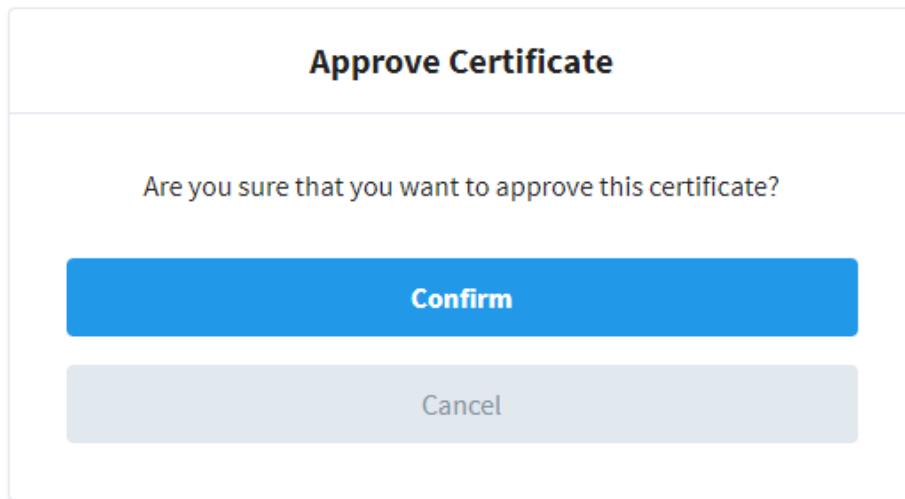
Note: This port is different from the default SSL port an Ignition Gateway would use when communicating to a client (default port 8043).

5. Check the **Use SSL** checkbox.
6. Enter any desired information in the **Description** field to easily identify how the connection is used from the Outgoing Connections list.

7. Use the default settings in the **Ping** section and **Timeouts** section of the page.
8. Click the **Create New Outgoing Gateway Connection** button at the bottom of the page.
9. You'll see a confirmation message that the connection was created.
10. At this point, your Gateway transmitted its certificate to the connected Gateway, but the incoming connection is not yet allowed. The Gateway's connection will not show up under the Incoming Connections section until after the certificate has been approved.
11. Log into the other Gateway. Navigate to **Config>Networking>Gateway Network**.
12. Click on the **Incoming Connections** tab. The first Gateway's certificate should be present. The certificate Common Name field holds the network address of the machine that transmitted the certificate. The Serial field holds a numeric string that is automatically generated when the certificate is created, and is unique to every certificate.

The screenshot shows the Ignition configuration interface. The left sidebar has sections for SYSTEM (Overview, Backup/Restore, Licensing, Modules, Projects, Redundancy, Gateway Settings) and NETWORKING (Gateway Network, Email Settings). The SECURITY section is also visible. The main panel is titled 'Config > Network > Gateway Network Settings'. The 'Incoming Connections' tab is selected. A table lists a remote certificate: 'Common Name' is 'backup:8060', 'Serial' is '8830218834790436753', 'Issuer' is 'Self', and 'Status' is 'Not Approved'. A red box highlights the 'approve' button.

13. Click the **approve** button to accept the certificate. You'll see a confirmation message. Click the **Confirm** button.



14. Now, the Gateway's connection will appear under the Incoming Connections section with a **PendingApproval** Security status. Click **More>Approve** to establish the connection.

Deleting Connections

Outgoing and incoming connections can be deleted for cases when the connection no longer exists on the other side.

1. To delete a connection, navigate to **Config>Networking>Gateway Network**.
2. Click on either the **Outgoing Connections** tab or the **Incoming Connections** tab.
3. Click **More**, and then select **Delete** next to the connection.

Note: For incoming connections, if a remote machine is still connected to the local machine with an outgoing connection, a new incoming connection will be created after deletion. For these cases, you must log into the remote Gateway and delete the outgoing connection. Then you can delete the local incoming connection.

The screenshot shows the Ignition configuration interface. The left sidebar has sections for SYSTEM (Overview, Backup/Restore, Licensing, Modules, Projects, Redundancy, Gateway Settings), NETWORKING (Gateway Network, Email Settings), and SECURITY. The main panel is titled 'Config > Network > Gateway Network Settings'. The 'Outgoing Connections' tab is selected. A table lists a connection: 'Host' is '10.10.115.3', 'Port' is '8060', 'Enabled' is 'true', 'Gateway Name' is 'Ignition-ignition8-ubuntu-64bit', 'Redundancy Role' is 'Independent', and 'Status' is 'Running'. A red box highlights the 'delete' button in the 'More' dropdown menu.

Gateway Network Diagnostics

The Diagnostics tab on the Gateway Network Settings page gives you insight to the Gateway and remote server response times.

1. To test the response time of a remote server, select the server name from the **Server** dropdown list.
2. Click the **Submit** button.

General Settings Outgoing Connections Incoming Connections Diagnostics Queue Management

Gateway Network Diagnostics

Gateway Network Port
This server is listening on port **8060** for incoming SSL certificates/connections.

Test Remote Server Response

Server	Serial	Issuer
<input type="button" value="Choose One"/> Controller Ignition-dartmouth-backup Ignition-dartmouth-backup-B Ignition-mrob-lt4	ignition8-ubuntu-64bit:8060	5585666939040193040 Self

Submit

3. The results will be displayed indicating if the call to the remote server was successful, what the response time was, and if there were any errors.

General Settings Outgoing Connections Incoming Connections Diagnostics Queue Management

Gateway Network Diagnostics

Gateway Network Port
This server is listening on port **8060** for incoming SSL certificates/connections.

Test Remote Server Response

Server

Controller

Results

Call Result	SUCCESS
Response Time	24 ms
Error	None

Submit

Local Certificate Information

Common Name	Serial	Issuer
ignition8-ubuntu-64bit:8060	5585666939040193040	Self

Gateway Network Queue Management

The Gateway Network system shares information across Gateways using a configurable number of send and receive threads. Gateway Network Queues are named to reflect their purpose and enable Ignition to prioritize which subsystem should have access to send or receive a thread at any given time. The Queue Management tab allows users to further direct how a queue should behave.

Queue Name	Description	Synchronous Delivery	Priority	Max Active	Settings Overridden?	
Call Results Queue	Handles results for remote service calls over the Gateway Network	false	Highest	Unlimited	false	<button>modify</button>
Default Queue	Generic Gateway Network queue	false	Normal	Unlimited	false	<button>modify</button>
Long Wait Queue	Handles messages that take up to an hour to deliver	false	Low	Unlimited	false	<button>modify</button>
Proxy Queue	Fowards requests through a proxy Gateway	false	AboveNormal	Unlimited	false	<button>modify</button>
Tag Value Publishing	Handles tag value change events	true	Normal	1	false	<button>modify</button>

Queue Settings

Clicking **modify** for a queue type accesses the Queue Settings page where the following information is available to view or edit.

Settings	Description
Queue Name	Name of the queue you are modifying (read only).
Description	Description for the queue you are modifying (read only).
Synchronous Delivery	This setting is configured by the queue and is unchangeable. If true, the queue will not dispatch another task until the current active task has completed. When a queue uses synchronous delivery, the maximum number of allowed active tasks is fixed at 1 and cannot be changed. Default is false. Note: Some queues are hard-coded as "Synchronous Delivery" queues, for example the Tag Value Update queue. For these queues, the Max Active setting is fixed at 1 and cannot be changed by the user. The user can only change the priority of the queue.
Max Active	The maximum number of active tasks allowed at a time. A task is considered active when it has been dispatched to the Gateway Network connection. You can set a limit to ensure that the Gateway Network connection will not become overloaded. Set this value to -1 to not enforce a limit on active tasks. Default is -1.
Priority	Determines the queue's priority in relation to other queues. A lower priority may result in messages in this queue taking longer to send, but can help prevent a Gateway Network connection from being overloaded.

Adjust Queue Message Capacity

Overloading a Gateway's queue active messages can starve the send or receive threads for Gateway Network connections. Starving the send or receive threads for the connection can potentially prevent other subsystem messages from getting through. To solve this issue, edit the **Max Active** field to set a limit that will allow the Gateway Network connection to continue functioning correctly.

Since this is a known problem with Tag History Queues, we will use an example with a GatewayA and GatewayB, where GatewayA requests tag history data from GatewayB to display on a chart. A user has noticed other items related to the Gateway Network connection, such as remote tags, are no longer displaying properly. They use the Gateway Network Diagnostics tab to confirm the network is not the issue, which indicates a slow database may be causing a cascading effect where other messages are delayed by all the Tag History Query calls.

This cascading effect happens because the Gateway Network asynchronously dispatches some functions. One Gateway thread handles the initial call, and a different Gateway thread dispatches the call's results. In our example, GatewayA is unaware that GatewayB is taking a long time to run the queries and will continue to dispatch more tag history requests. When the limit of send and receive threads available to a Gateway Network connection is reached, other messages can no longer get through. This can be verified as the issue by checking GatewayB's current Incoming /Outgoing Tasks and Results Response Time.

Gateway Network Diagnostics

Gateway Network Port

This server is listening on port **8060** for incoming SSL certificates/connections.

Test Remote Server Response

Server

controller	▼	Submit
------------	---	--------

Results

Call Result	SUCCESS
Response Time	10098 ms
Error	None

Adjusting queue message capacity controls the flow of requests to prevent thread starvation across the network. Select **Modify** for the Tag History Queue to access Queue Settings and adjust capacity. Make sure the Gateway responsible for the thread requests is where the queue capacity is updated. In this example, GatewayA Tag History Queue Max Active count was set to 2 and saved by clicking **Create New Queue Override Settings**.

Queue Settings	
Queue Name	Tag History
Description	Retrieves tag historical data over the Gateway Network
Synchronous Delivery	<input type="checkbox"/> This setting is configured by the queue and is unchangeable. If true, the queue will not dispatch another task until the current active task has completed. Note that when a queue uses synchronous delivery, the maximum number of allowed active tasks is fixed at 1 and cannot be changed. (default: false)
Max Active	<div style="border: 1px solid red; padding: 2px; width: 15px; height: 15px; display: inline-block; vertical-align: middle;">2</div> <p>The maximum number of active tasks allowed at a time. A task is considered active when it has been dispatched to the Gateway Network connection. You can set a limit to ensure that the Gateway Network connection will not become overloaded. Set this value to -1 to not enforce a limit on active tasks. (default: -1)</p>
Priority	<div style="border: 1px solid #ccc; padding: 2px; width: 150px; display: inline-block; vertical-align: middle;"> Normal <div style="float: right; font-size: small;">▼</div> </div> <p>Determines the queue's priority in relation to other queues. A lower priority may result in messages in this queue taking longer to send, but can help prevent a Gateway Network connection from being overloaded. (default: Normal)</p>

[Create New Queue Override Settings](#)

The Gateway Network Status page Remote Gateway Details now shows pending items since the Tag History Active upper limit has been met. Note that pending tasks are ones that have not yet been picked up for dispatch by the Gateway Network connection. These will be discarded if not dispatched within 1 minute. Be aware these requests being deleted may show more errors on components like the GatewayA chart, but all other subsystem messages can now get through without issue.

Outgoing Queues

Name	Priority	Inserts/Sec	Pending	Active	Avg Pending Secs	Total	Actions
Alarm Query Provider Service	Normal	0.0	0	0	0.0	0	Pause Clear
Call Results Queue	Highest	0.0	0	0	0.0	0	Pause Clear
Default Queue	Normal	0.0	0	0	2.3	44	Pause Clear
Long Wait Queue	Low	0.0	0	0	0.0	0	Pause Clear
Proxy Queue	AboveNormal	0.0	0	0	0.0	0	Pause Clear
Remote Tags	Normal	1.8	0	0	0.2	2,612	Pause Clear
Remote Tags (v7)	Normal	0.0	0	0	0.0	0	Pause Clear
Tag History	Normal	3.8	2	2	0.6	1,949	Pause Clear

Gateway Network Proxy Rules

The following feature is new in Ignition version **8.1.34**
[Click here](#) to check out the other new features

Gateway Network Proxy Rules allow you to control the amount of service enumeration calls over your Gateway Network. Service enumeration calls run every 60 seconds against every other Gateway the local Gateway is aware of, including Gateways on the opposite side of a proxy connection. Given enough Gateways, this may result in a large amount of network traffic, negatively impacting your Gateway Network.

Proxy Rule Settings

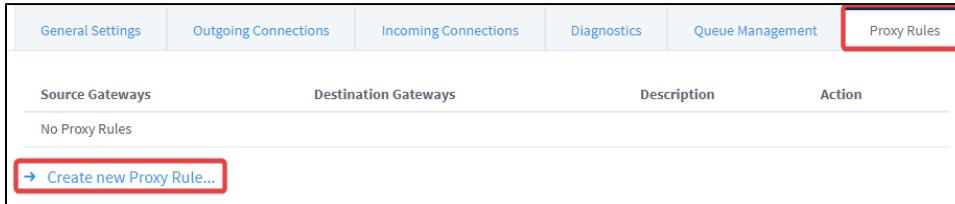
The following table represents the various configurable settings within a proxy rule. In addition to specifying individual Gateways, you can use wildcards to list out Gateways you want.

Main	
Name	Description
Source Gateways	The list of Gateways that are asking for available routes. If the asking Gateway is on this list either by name or wildcard, the proxy Gateway will check the list of Destination Gateways to determine what to do next.
Destination Gateways	The list of Gateways the proxy Gateway will use to see if there is a name or wildcard match. The Action property will then either Allow or Deny the route.
Description	Description of the proxy rule. Useful for users to determine what the rule is used for.
Action	The action to take when a proxy route between a source Gateway and a destination Gateway matches this rule. Options include Allow or Deny . If set to Deny, the proxy route will not be reported to remote Gateways. Default is Allow.

If a route does not match a proxy rule, that route will be allowed and be visible to the asking Gateway by default.

Setting up a Proxy Rule

1. Navigate to your Gateway Network configuration settings, located at **Config > Networking > Gateway Network**.
2. Go to the **Proxy Rules** tab and click on **Create new Proxy Rule...**



General Settings	Outgoing Connections	Incoming Connections	Diagnostics	Queue Management	Proxy Rules
Source Gateways	Destination Gateways	Description	Action		
No Proxy Rules					
Create new Proxy Rule...					

3. Set up your list of Source Gateways and Destination Gateways. See [Proxy Rule Settings](#) for a description of each setting.



Source Gateways	A comma separated list of gateway names. May use the wildcard '*' , such as 'Manufacturing_*'.
Destination Gateways	A comma separated list of gateway names. May use the wildcard '*' , such as 'Manufacturing_*'.

4. Give your rule a description. This will make it easier to identify the purpose of your rule in the future.

Description	
-------------	--

5. Set the Action setting to either allow the proxy route or deny the proxy route, then click **Create New Proxy Rule**.

Action	<p>Allow ▾ The action to take when a proxy route between a source gateway and a destination gateway matches this rule. If set to Deny, the proxy route will not be reported to remote gateways. (default: Allow)</p>
--------	--

Create New Proxy Rule

A few things to note about proxy rules:

- Proxy rules are followed in a top-down order, with the rule at the top of the list having the highest priority.
- Proxy rules are bidirectional. If a controller Gateway can access an agent Gateway with proxy rules, the agent Gateway can also access the controller Gateway.
- Proxy rules view redundant Gateway pairs as a single entity.

In This Section ...

Gateway Network Certificates and SSL

When a remote machine establishes an incoming connection, its server name is transmitted and appears in the Server Name field under **Gateway Network -> Incoming Connections**. However, no identity authentication is performed when the connection is created. The local system accepts the remote system id without question. To perform identity authentication on a connection, you must use Secure Socket Layer (SSL) and certificates. By default, SSL is enabled.

Note: When using the Gateway Network and Redundancy, SSL Certificates are automatically pushed from the redundant Master to the Backup.

Client Certificates

The following feature is new in Ignition version **8.1.14**

[Click here](#) to check out the other new features

Client certificates are the certificates of peer Gateways that the current Gateway trusts when it is making outgoing connections. Client certificates live under:

```
$GATEWAY_HOME/data/gateway-network/client/security/pki/
```

When the Gateway makes an outgoing connection to a peer Gateway whose certificate is not yet trusted, the peer certificate (or its certificate chain if one is configured) is copied into:

```
$GATEWAY_HOME/data/gateway-network/client/security/pki/rejected/
```

This model allows users to configure the Gateway Network client to trust the peer Gateways on outgoing connections by moving the certificate on the file system from `$GATEWAY_HOME/data/gateway-network/client/security/pki/rejected/` to `$GATEWAY_HOME/data/gateway-network/client/security/pki/trusted/certs/`. This file system change will be picked up immediately by the Gateway and the connection will be trusted when it attempts to reconnect again.

Server Certificates

The following feature is new in Ignition version **8.1.14**

[Click here](#) to check out the other new features

Server certificates are the certificates of peer Gateways that the current Gateway trusts when it is handling incoming connections. Server certificates live under:

```
$GATEWAY_HOME/data/gateway-network/server/security/pki/
```

The Gateway Network config UI's **Incoming Connections** tab was made compatible with this new model so that incoming connection certificates may continue to be approved, denied, or deleted there.

Note:

If you are using your own CA to sign Gateway Area Network certificates, add the CA public key to:

```
$GATEWAY_HOME/data/gateway-network/server/security/pki/trusted/certs/
```

Denying a Certificate

To deny a certificate, navigate to **Config -> Networking -> Gateway Network -> Incoming Connections**. The certificate **More** dropdown displays **deny** and **delete** options. If deny is selected, the connection that has been using that certificate will no longer be allowed to connect. Select delete for certificates that are no longer in use. Keep in mind that if you delete a certificate, and a remote machine is still using that certificate, it will reappear on the Certificates page. In this case, you must navigate to the remote Gateway and delete its outgoing connection. Then you can permanently delete the certificate from the Certificates page.

Regenerating Gateway Network Certificates

On this page ...

- [Client Certificates](#)
- [Server Certificates](#)
- [Denying a Certificate](#)
- [Regenerating Gateway Network Certificates](#)

Ignition generates a self-signed certificate for the Gateway Network on start up if no existing certificate is found. These self-signed certificates have a lifespan of 10 years. Unlike trusted certificates, self-signed certificates cannot simply be reuploaded and replaced. Regenerating the certificates creates a new certificate with an expiration date set for ten years from the date the certificate is regenerated. If you need to regenerate a self-signed certificate, remove the `$INSTALL_LOCATION/webserver/metro-keystore` file and restart the Gateway. The certificate will need to be trusted again by all other gateways that trusted the expired certificate.

Database Connections

How Are Databases Used in Ignition?

While connecting to a database is not required for basic status and control functionality, it can dramatically increase the possibilities that the system offers. There are a few places where databases are used in Ignition, such as historical data logging, reporting, and storing alarm logs.

Historical Data Logging

Logging data for historical analysis, either through [Tags Historian](#) or with the [SQL Bridge module](#), require s a database connection. Databases are great at handling historical data, and by using a standard relational database your data is stored in an open format that can be used in many ways.

Reports, Graphs, and Charts

The Vision module makes it easy to present data stored in databases in a variety of ways. You can quickly create charts that show performance over time, locate anomalies and detect trends.. Furthermore, it's important to remember that it is possible to pull data from any database that Ignition is connected to, even if the data wasn't placed there by Ignition. This means you can tie in data from other sources or areas of your company, such as pulling in inventory and staff information, as well.

Storing Alarm Logs

Store alarm information historically and examine it later for patterns or trouble spots.

Getting Started with Databases

The first step in using a database with Ignition is to identify a database server. Many companies already have database servers maintained by their IT departments. If you do not, or wish to set up your own database server for Ignition, the [Supported Databases](#) section below offers some advice on choosing a database vendor.

Once you've identified a server, all you need to do is [create a connection](#) to that server to get up and running.

On this page ...

- [How Are Databases Used in Ignition?](#)
 - [Historical Data Logging](#)
 - [Reports, Graphs, and Charts](#)
 - [Storing Alarm Logs](#)
- [Getting Started with Databases](#)
 - [Supported Databases in Ignition](#)
 - [Database](#)
 - [Version](#)
 - [Installing and Connecting to a Database](#)
- [Database Drivers and Translators](#)
 - [What Is JDBC?](#)
 - [JDBC in Ignition](#)
- [Monitoring Connection Status](#)

Supported Databases in Ignition

Ignition has been tested with the following databases, and can connect to them directly after installation. You can connect to other databases by installing additional JDBC drivers (the Java database connection specification), which are often provided by database vendors.

Database	Version
Full Support	
MySQL	5.0+ for full support. Ignition can connect to 4.x, but many features such as Tags are not tested.
Microsoft SQL Server	2005, 2008, 2012, 2014, 2016, 2017, 2019, 2022 (full and express editions). Ignition can connect to 2000, but has not been fully tested.
Oracle	10g, 11g, 12c (full and express). The letters stand for "grid" and "cloud"
PostgreSQL	8.0+
Firebird	All versions.
IBM DB2	9.5+
SQLite	A driver for the popular embedded database system. This can be used to connect to an existing SQLite database, or create a new database: setting the connect URL property to a file that doesn't exist will result in the driver attempting to create the database.
Limited support	
Other JDBC drivers	Due to variances in databases, some features may not work fully through other non-tested JDBC drivers. However, it is usually possible to get full functionality though the careful use of the database translator feature. For example, the JDBC driver for MariaDB could be downloaded and added to Ignition.

Installing and Connecting to a Database

Once you've identified a server, all you need to do is create a connection to that server to get up and running. See the [Installing Databases](#) and [Connecting to Databases](#) sections for details about how to install and connect to different databases through Ignition.

If we don't already have a connector for your database type, you can [simply add it in yourself](#).

Database Drivers and Translators

What Is JDBC?

JDBC stands for the Java DataBase Connectivity API. It is a standardized way for Java-based applications to interact with a wide range of databases and data sources. A JDBC Driver enables Ignition to connect to, and use data from, a particular database system.

JDBC in Ignition

Ignition, being a Java-based application, leverages JDBC in order to connect to a variety of data sources. This enables Ignition to offer a standardized set of functionality on a wide range of different systems and databases. This includes not only commonly-used databases such as MySQL, Microsoft SQL Server, and Oracle, but additionally other lesser-known systems as well, provided the manufacturer offers a JDBC driver for the system.

Name	Driver Type	Default Translator	Status	Actions
MariaDB	MySQL	MYSQL	Installed	delete edit
Microsoft SQLServer	Microsoft SQL Server	MSSQL	Installed	delete edit
MySQL	MySQL	MYSQL	Installed	delete edit
Oracle Database	Oracle	ORACLE	Error - either required files are missing or classname is incorrect	delete edit
PostgreSQL	PostgreSQL	POSTGRES	Installed	delete edit

Monitoring Connection Status

The state or status of a database can be monitored from the [Status](#) section of the Gateway Webpage, under [Connections > Databases](#). The status panels show the current state and a fault message, if applicable, or throughput statistics if the connection is active.

When a connection is not available, it is re-tested every 10 seconds, and the status is updated.

Config > Database > Database Connections

Name	Description	JDBC Driver	Translator	Status		
DB		MariaDB	MYSQL	Faulted	delete	edit
MSSQL		MySQL	MYSQL	Valid	delete	edit
SQLServer		Microsoft SQLServer	MSSQL	Reconnecting	delete	edit

→ Create new Database Connection...

Note: For details about a connection's status, see the [Database Connection Status page](#).

Related Topics ...

- [SQL in Ignition](#)
- [JDBC Drivers and Translators](#)

In This Section ...

Installing Databases

Why Install a Database?

A lot of additional functionality becomes simple or is only accessible when Ignition is connected to a database. Storing Historical data, storing notes or files, and creating dynamic lists to name a few. It is important to note that Ignition does not install any databases for you. There are many types that you can connect to, but you need to choose the database that is best for you. Installing your own database means you have complete control over it, anything Ignition adds to it can be accessed by another program easily.

You can install as many database systems as you like, and each of them allow you to create as many schemas (or groups of data) as you want. You get to decide where your database is installed, or where you want to install all of them.

On this page ...

- [Why Install a Database?](#)
- [Which Database Should You Use?](#)
- [Where to Install a Database](#)

Which Database Should You Use?

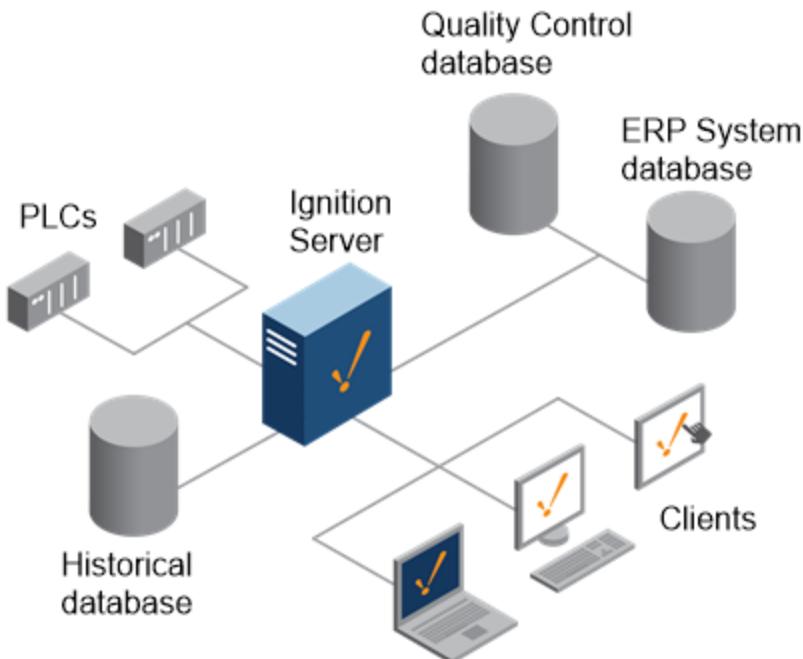
There is a lot of debate on this subject but the short answer to this question is 'whichever database your IT department already uses.' Modern relational database all have the same basic functionality, but slightly different ways of doing things. If your IT department already supports Microsoft SQL Server (MSSQL), then they already have the tools and knowledge to manage those databases. Because they are so similar (and Ignition takes care of so much for you), it is almost always easier to learn to use an existing database than to add IT support for a second type. Not to mention that adding a second type might mean hiring new personnel in the IT department.

If your company does not already have a database preference, then it's up to you to decide which is best based on your needs. They all have a free version, but different limitations. For example: MSSQL has a cap on how much data can be stored in their free version. MySQL does not have a cap like this, but also does not allow phone support for their free version.

Where to Install a Database

You can install your database anywhere that Ignition has access to through the network. There are two main options for installing your database: one is on the computer that Ignition is installed on, and the other is installed on a different server in your network. Technically, there is a third option to connect Ignition to a database that is in a remote location using a VPN or some other way to access it. This third option will work, but because of latency and the data being physically very far away, it is not recommended for storing data that will be accessed often like Tag History.

For production systems, we recommend that your database is on its own server, not installed on the computer with Ignition. This is helpful for many reasons, but mostly because databases can potentially take up a lot of resources on a computer. If the database is on its own computer, you don't have to worry about other programs starving for memory or CPU. If you do this and install your database on another computer, just make sure to adjust your firewalls and pay attention to the database connection security. Most databases don't allow the default username to connect remotely.



[In This Section ...](#)

Installing IBM DB2

You need to install **IBM DB2 Community Edition** before you connect to the database.

To Download IBM DB2 Community Edition

1. Go to the [Db2 Download](#) page and sign in with your IBM credentials.
2. Download the IBM® Db2 11.5 Edition for your Operating System.
3. Log on to your system using a local administrator account.
4. Locate the `v11.5_ntx64_dec.zip` file that you downloaded and extract it to your desired location.
5. In the extracted folder, right-click on the `setup.exe` file and select **Run as administrator**.
6. Click **Install a Product**
7. Scroll down to the end of the *Db2 Version 11.5.0.0 Server Editions* section and click **Install New**, then click **Next** to begin the installation process.
8. Review and accept the software license agreement and click **Next**.
9. Select the **Typical** Installation option and then click **Next**.
10. Select **Install Db2 Server Edition on this computer and save my settings in a response file** and click **Next**.
11. Set your installation path and click **Next**.
12. Set your **DB2 copy name** and click **Next**.
13. Select **Local user or Domain user account** and enter your user credentials.
14. Select **DB2** as the default instance and click **Next**.
15. Uncheck the **Set up your DB2 server to send notifications checkbox** and click **Next**.
16. Select the **Enable operating system security** checkbox and click **Next**.
17. Click **Finish** to finalize your installation.



For detailed instructions, refer to the links below:

[Installing the Db2 server by using the Db2 Setup wizard \(Windows\)](#)

[Installing the Db2 server by using the Db2 Setup wizard \(Linux and UNIX\)](#)

[Downloading and Installing Docker Editions](#)

Note: The links above are for the IBM DB2 Database version 11.5.

You can change the instruction version by using the drop-down list located on the top left side

After you have successfully installed IBM DB2, you will need to create a database.

Creating a database in Windows

1. Open Command Prompt with Administrator privileges.
2. Type `db2cmd` (This will open the DB2 Command Line Tool)
3. Type `db2` (This will open the Command Line Processor for DB2 Client)
4. To create a database:
`CREATE DATABASE database_name`



IBM Documentation on creating databases:

On this page ...

- [To Download IBM DB2 Community Edition](#)
- [Creating a database in Windows](#)

Creating Databases

Installing MySQL

Install MySQL Server and MySQL Workbench

The goal of this page is to demonstrate how to install MySQL Server, and a helpful tool called MySQL workbench. This guide is not an exhaustive listing of all of the various installation steps or scenarios for MySQL. For more information, take a look at MySQL's documentation: [MySQL Documentation](#).

1. Go to the MySQL website at <https://dev.mysql.com/downloads/mysql/>
2. Scroll-down to **Windows (x86, 32-bit), MSI Installer**. You will notice multiple download options. Both allow you to install MySQL

Note: MySQL Installer is 32-bit, but will allow you to install the 64-bit version of MySQL.

The screenshot shows the MySQL Community Downloads page for MySQL Installer 8.0.20. The 'General Availability (GA) Releases' tab is selected. Under 'Select Operating System:', 'Microsoft Windows' is chosen. Two download links are visible:

- Windows (x86, 32-bit), MSI Installer** (mysql-installer-web-community-8.0.20.0.msi) - 8.0.20, 24.4M, Download button. This link is highlighted with a red box.
- Windows (x86, 32-bit), MSI Installer** (mysql-installer-community-8.0.20.0.msi) - 8.0.20, 420.6M, Download button.

A note below the links suggests using MD5 checksums and GnuPG signatures for integrity verification.

On this page ...

- [Install MySQL Server and MySQL Workbench](#)
- [Running the Installer](#)

The Inductive University logo features a blue hexagon with 'IU' and a green laurel wreath. Below it, the text 'INDUCTIVE UNIVERSITY' is displayed. To the right, there are two sections: 'Installing MySQL' and 'Watch the Video'.

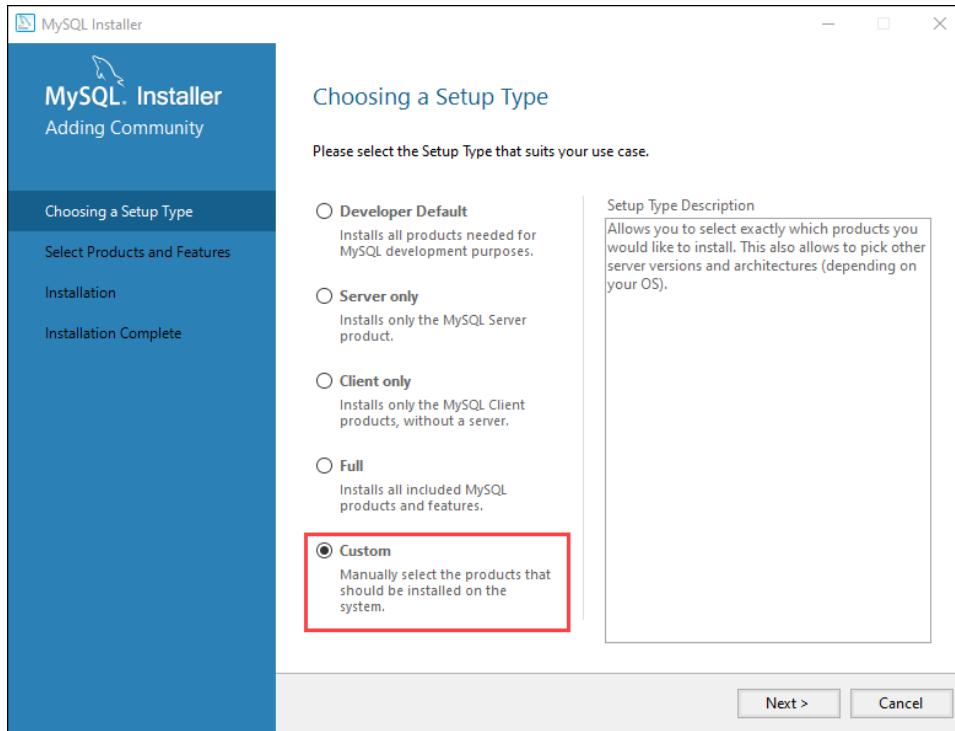
3. Click on the **Download** button.
4. On the next page, you can login or create an account if you'd like. Otherwise click **No thanks, just start my download**.

The screenshot shows the MySQL Community Downloads page after clicking the 'Download' button for the 32-bit MSI installer. It now displays a 'No thanks, just start my download.' button at the bottom, which is highlighted with a red box.

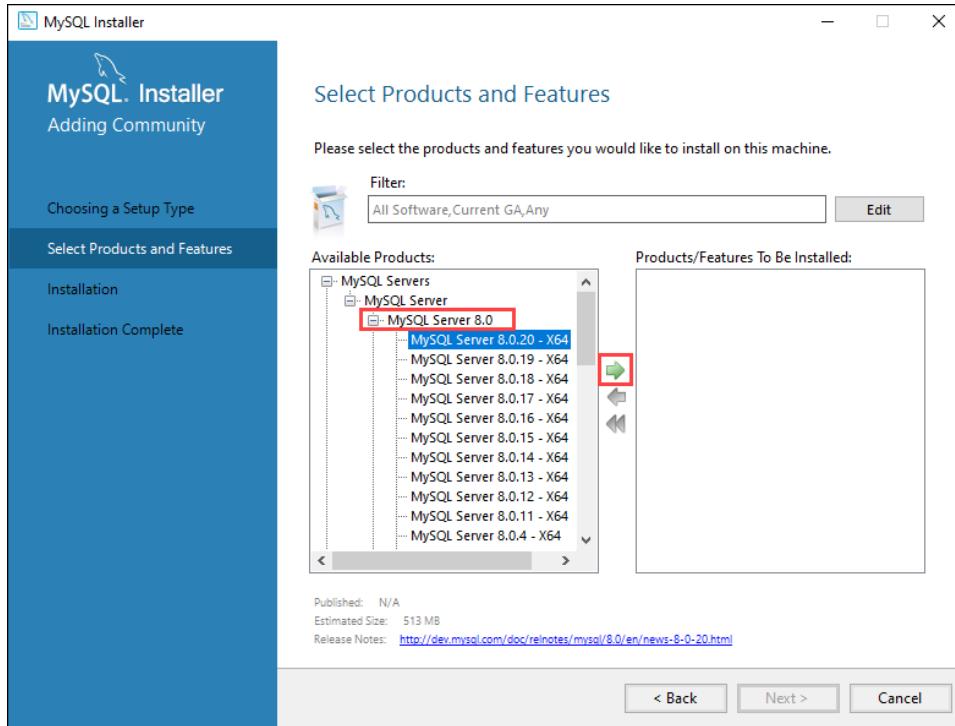
Running the Installer

1. Once the **.msi** file is downloaded, run the file to begin the installation process.
2. The **Welcome** window is displayed. Select the **Install MySQL Products** action.
3. On the Choosing a Setup Type page, select **Custom** and click **Next**. While you can select one of the other options, at minimum you'll want to install both the server (the actual database) and MySQL Workbench (an application that allows you to quickly and easily interact with the database, without using a command-line client). Any other items beyond these two are generally unnecessary in most environments.

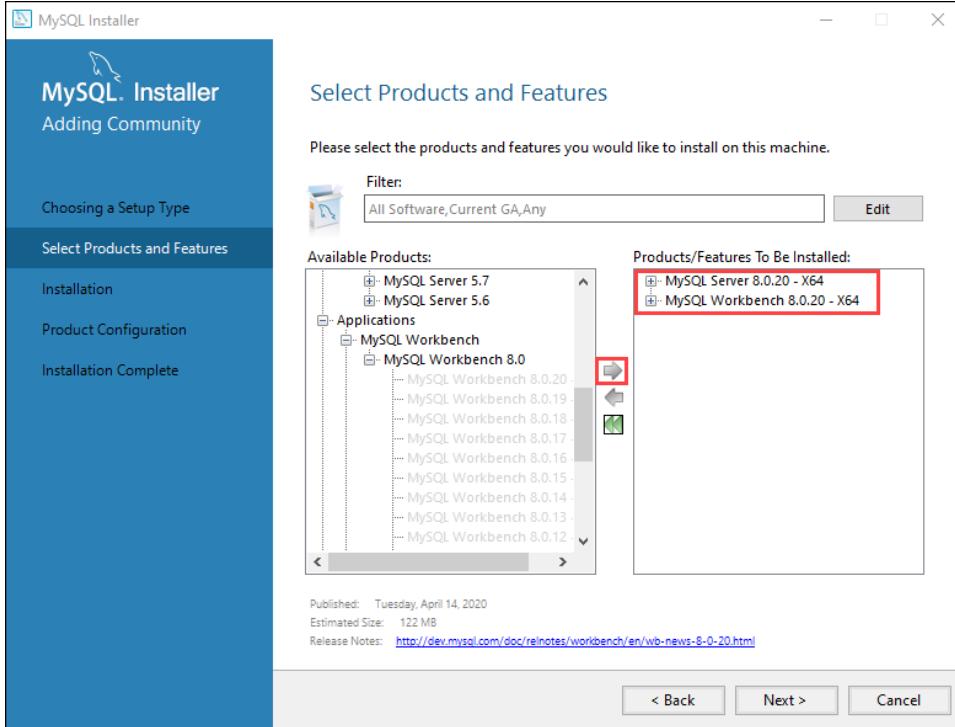
If you want to learn more, MySQL does have some additional information on their documentation if you're curious about the other options: [MySQL Docs](#). This example will continue with a Custom installation.



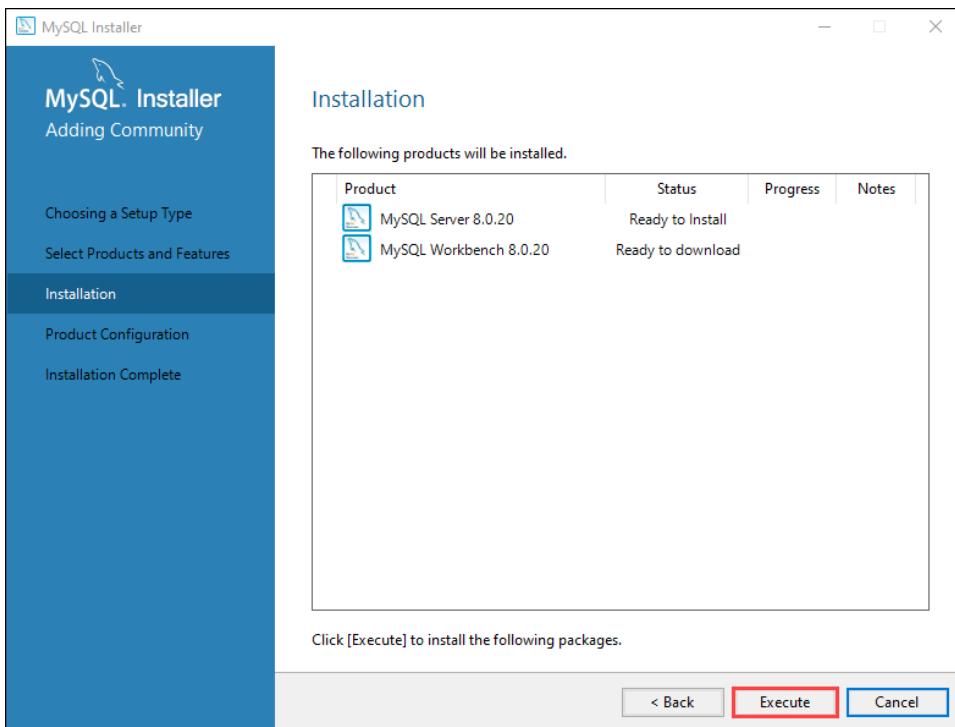
4. On the Select Products and Features page, scroll down to choose **MySQL Server 8.0**. Click the right arrow to move it to the "Products /Features To Be Installed" column.



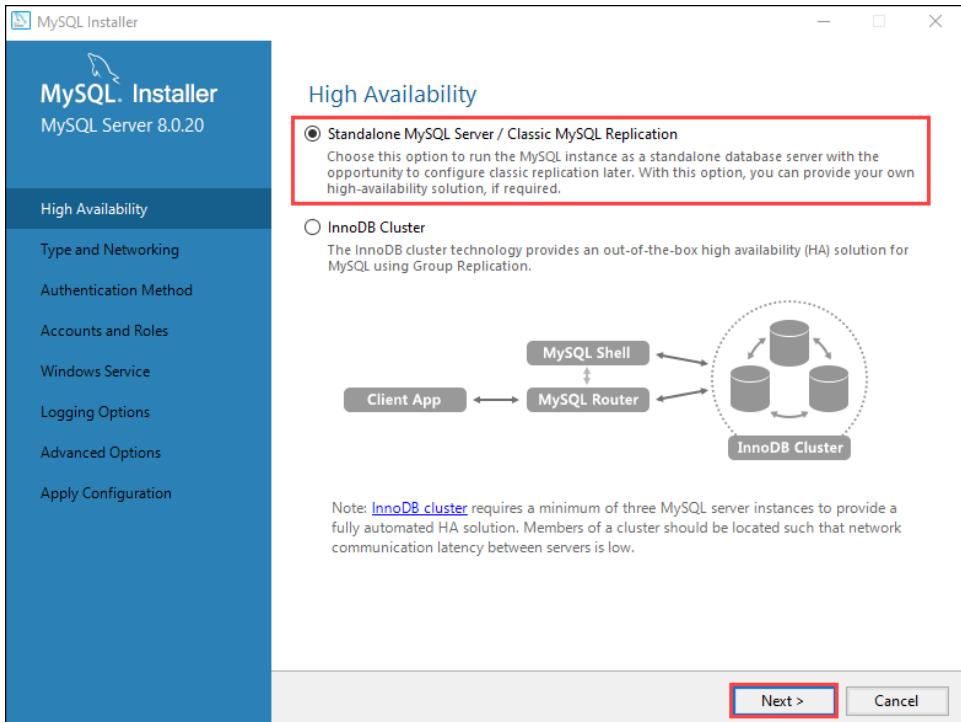
5. Scroll down to MySQL Workbench, select a version, and click the right arrow to move it to the "Products/Features To Be Installed" column.
6. Click **Next**.



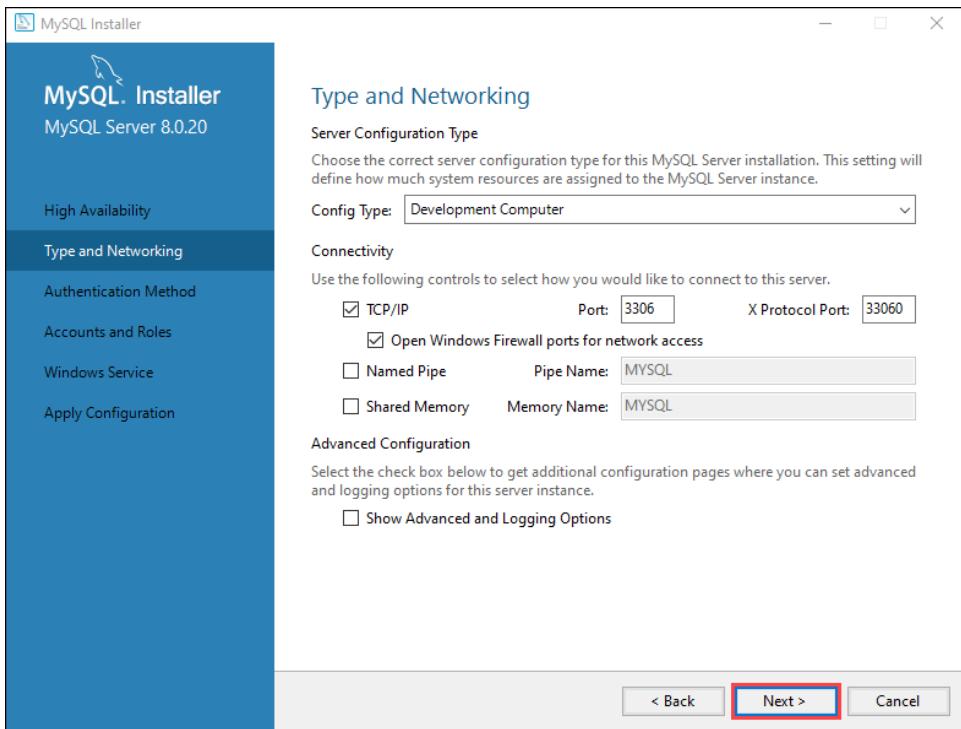
7. On the Installation screen, click **Execute**.



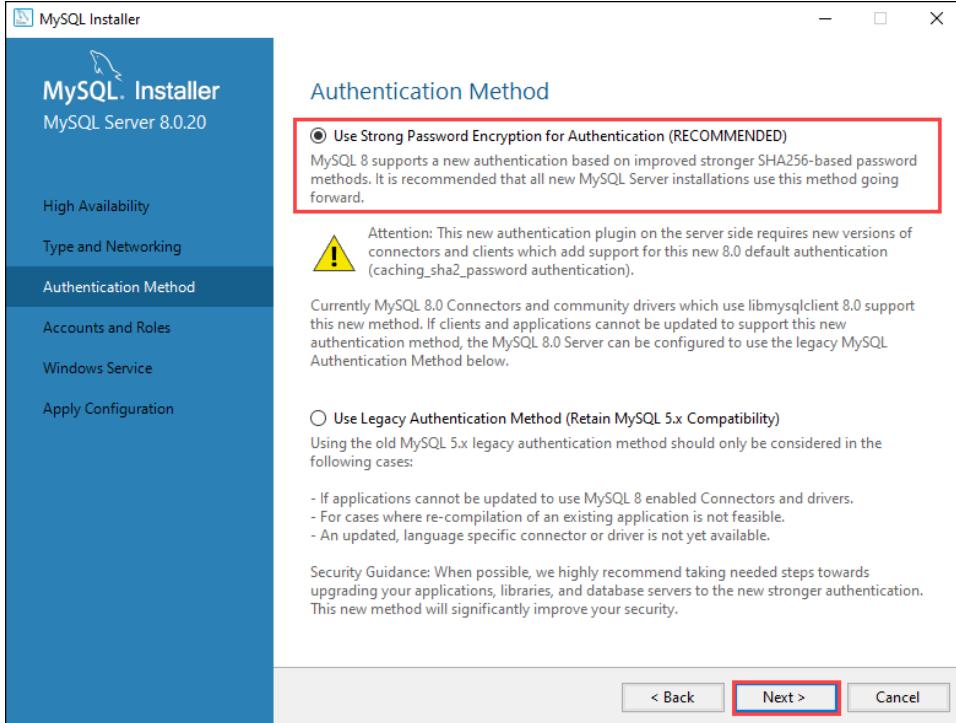
8. After these are downloaded and installed, you see the Product Configuration page. Click **Next**.
9. In this guide, we're going to use a standalone server. On the High Availability page, select **Standalone MySQL Server / Classic MySQL Replication**, and click **Next**.



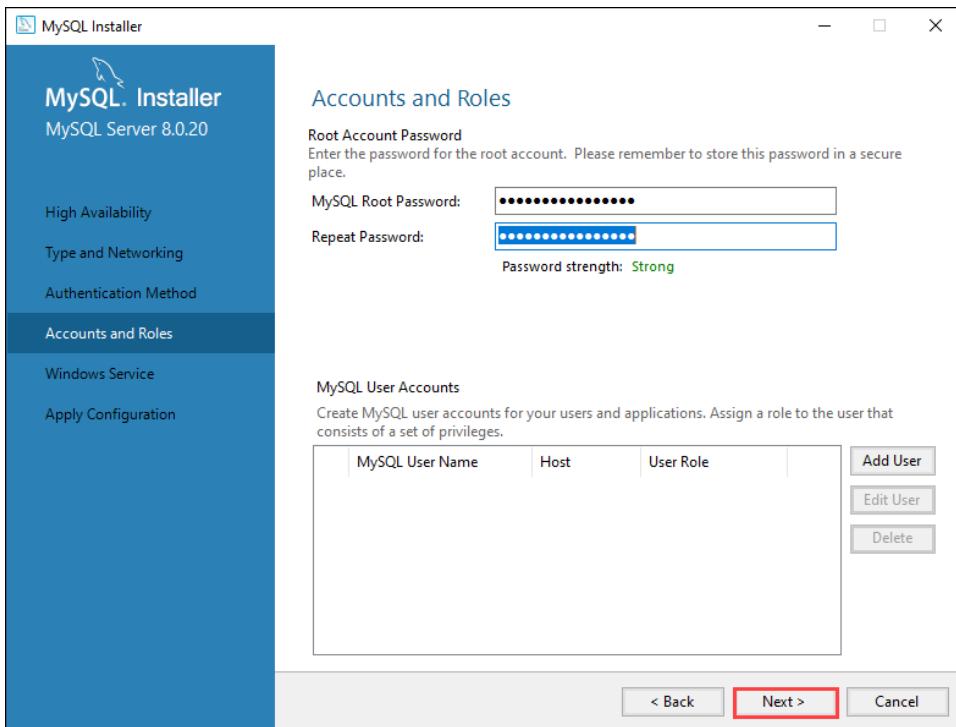
10. Leave the default settings on the Type and Networking page. Click **Next**.



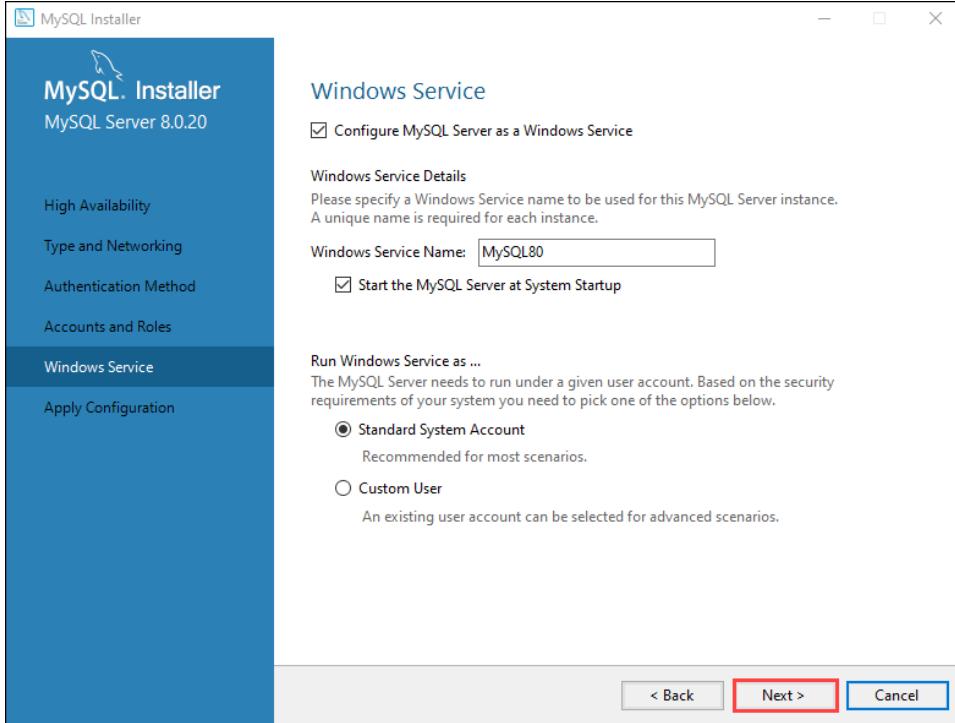
11. On the Authentication Method page, choose **Use Strong Password Encryption for Authentication**. Click **Next**.



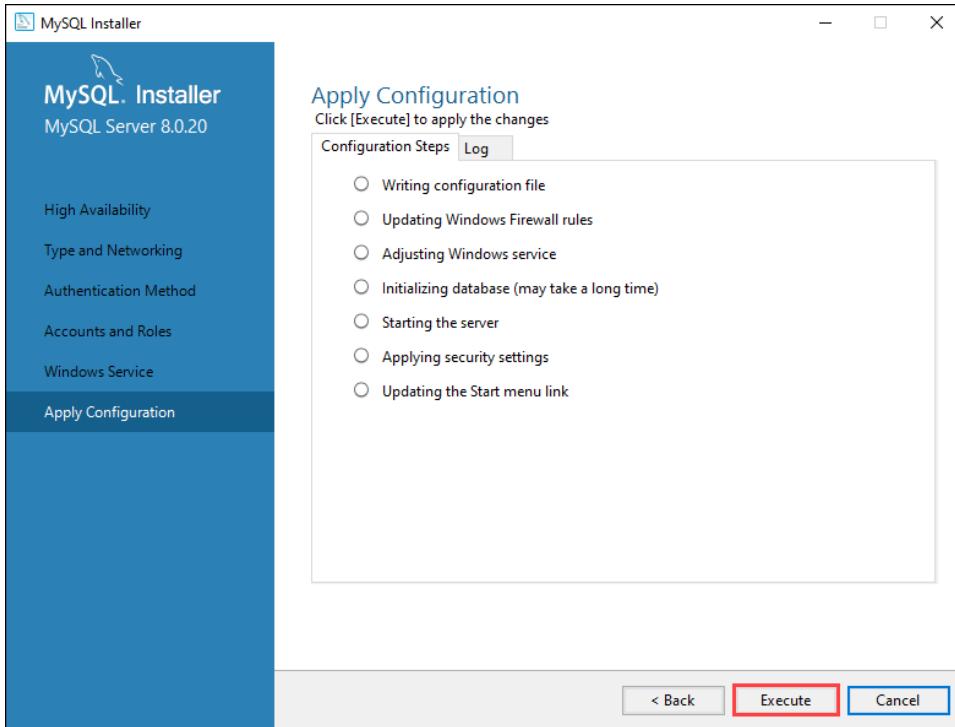
12. Create a strong password and click **Next**.



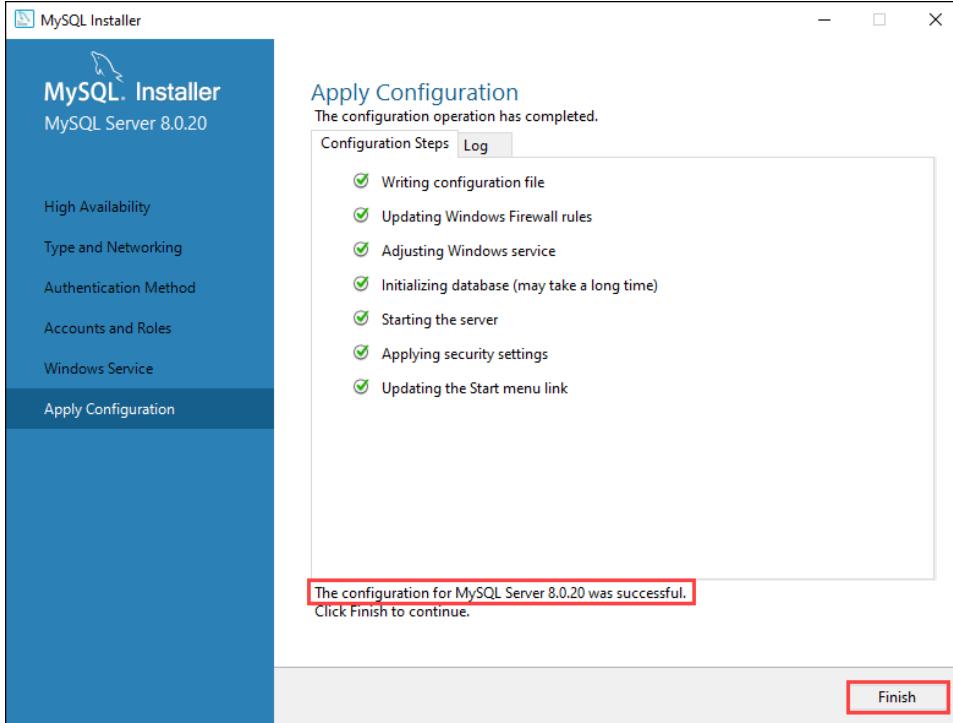
13. On the Windows Service page, leave the default settings and click **Next**.



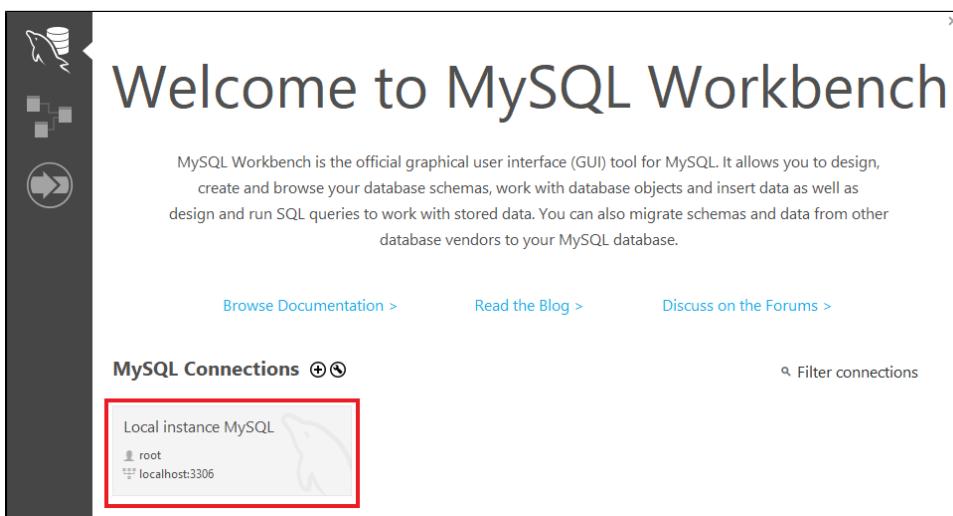
14. On the Apply Configuration page, click Execute.



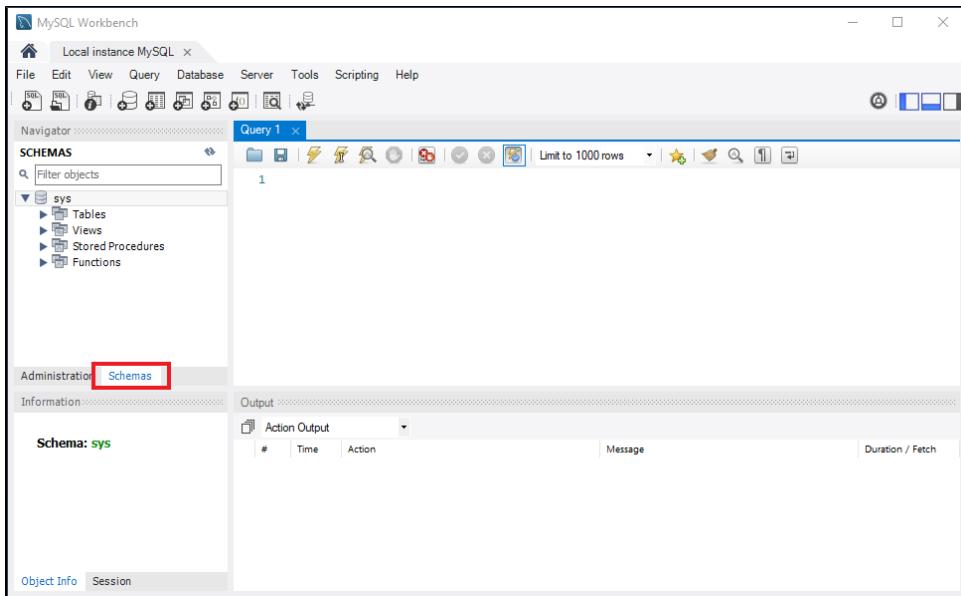
15. Once the configuration is applied and the database is initialized, you'll see a confirmation message. Click **Finish** to complete the install process.



16. The **MySQL Workbench** window is displayed. MySQL Workbench lets you to administrate the MySQL server.
17. Click on **Local instance MySQL** to connect to the newly installed MySQL server.

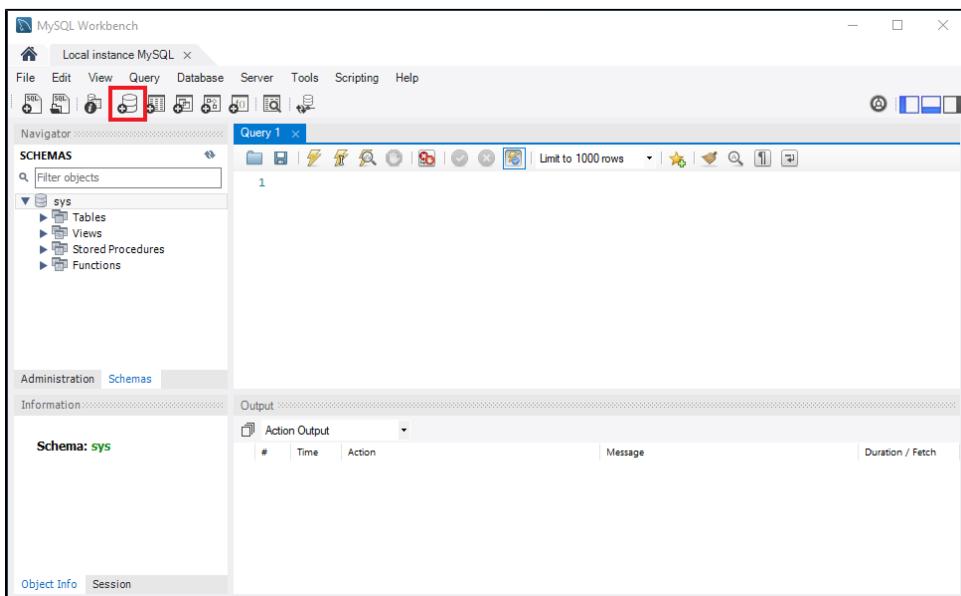


18. Enter the root **password** you earlier entered into the installer and click **OK**.
19. You are now connected and can see, in the Navigator, the default **sys** schema. Click on the **Schemas** tab to see a listing of schemas

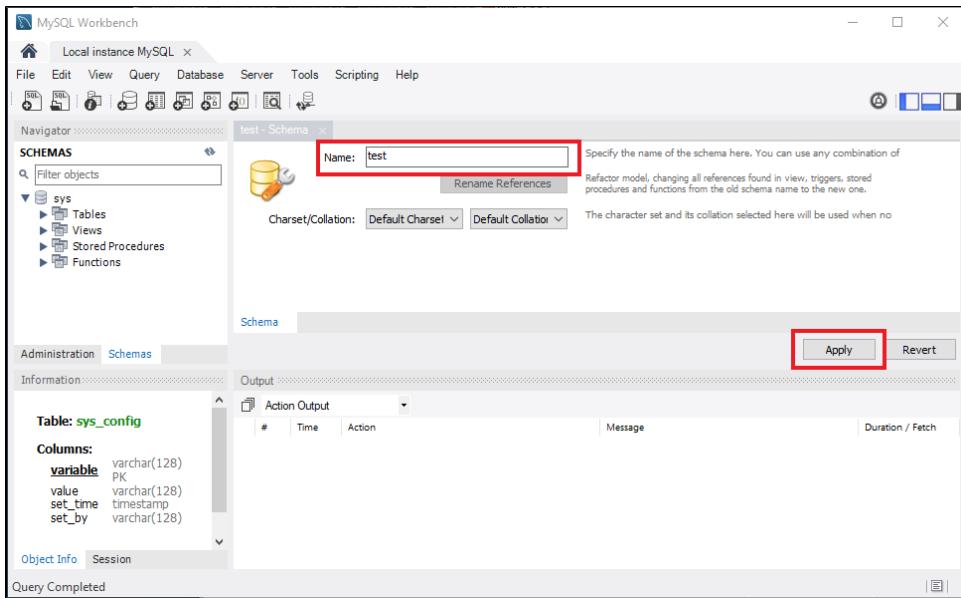


20. A "schema" is a collection of tables and other database objects. Ignition needs a schema to connect to. Instead of using the **sys** schema, we can create a schema dedicated to Ignition.

To create the schema, click the **Create New Schema** icon.

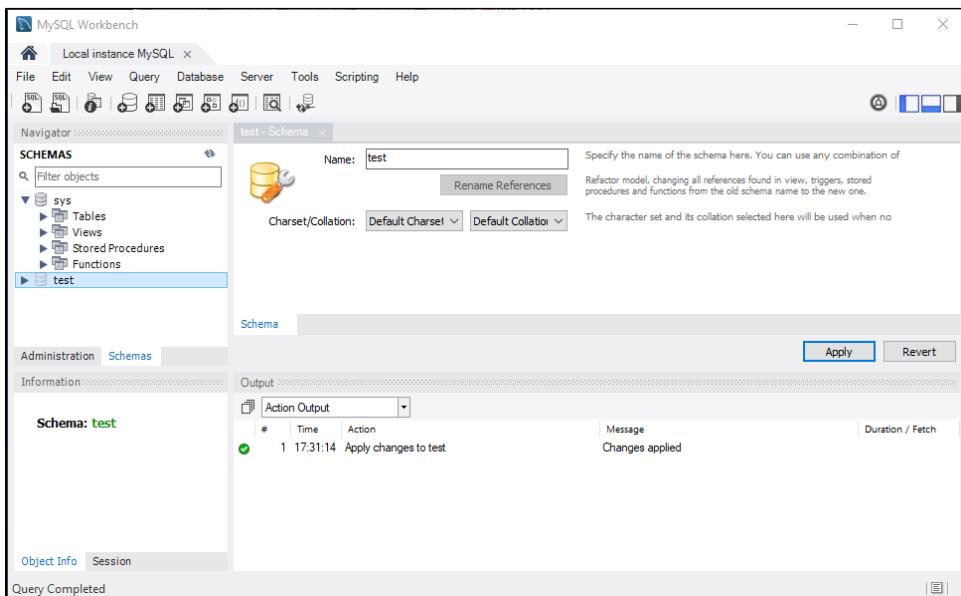


21. You'll see the Schema creation tab. This allows you to make a new **Schema**. Type in **test** as the name, and click **Apply**. Technically you can call the schema anything the database allows, but the default MySQL connection in Ignition assumes a schema named "test". If you name the schema something else here, you'll need to remember the name when creating the database connection in Ignition later.



22. You'll see the **Apply SQL Script to Database** window. Click **Apply**.

23. If there was an issue, the following window will state the issue. Go back and address the issue. Otherwise, click the **Finish** button. You should now see the **test** schema in the Navigator.



Now that the database is installed, you can connect Ignition to it. Learn more here: [Connecting to MySQL](#).

Installing Microsoft SQL Server Express

You need to download and install both the **SQL Server Express** and the **SQL Management Studio** before you connect to the database.

To Download SQL Server Express

This section walks through the process of installing a new instance of SQL Server Express.

1. Go to <https://www.microsoft.com/en-us/sql-server/sql-server-downloads>
2. Look for the link/button to download the **Express Edition**.
3. Run the installer
4. When given a choice between installation type, select **Custom**.
5. Select an installation directory. If you don't have a preference, simply use the default.
6. The installer will unpack and download required files.
7. Once the installer is ready, the **SQL Server Installation Center** window will appear. Select the **I nstallation** heading on the side bar, and click on the **New SQL Server stand-alone installation or add features to an existing installation** link.
8. On the **License Terms** window, choose **I accept the license terms** and click **Next**.
9. On the **Product Updates** window, click **Next** to start installing the SQL Server.
10. On the **Feature Selection** window, select the features you require and click **Next**.

Note:

Be mindful that you may find some of the default packages extraneous. You can uncheck many of the options, such as the optional Machine Learning Services and Language Extensions. The Server and SQL Server Management Studio are the main tools you'll need when interacting with an Ignition installation. You can always remove additional components from the SQL Server installer later if you choose so.

Check [Microsoft's documentation](#) for more details on what each package does.

11. The **Instance Configuration** window shows **Named instance: SQL Express**, you can keep it or change it if you like. Click **Next**.
12. When asked about a JRE to use, you can use the provided JRE. Click **Next**.
13. On the **Server Configuration** window, choose **Automatic** from the dropdown under **Startup Type** for the **SQL Server Browser** service, and click **Next**.
14. On the **Database Engine Configuration** window, for **Authentication Mode** choose **Mixed Mode**, enter a **password** for the SA account. Note that you can use a Windows Authentication Mode with Ignition, but it does require some [additional configuration](#) when connecting later on. In either case, click **Next**.
15. On the next few windows, continue clicking **Next** until it shows installation is **Complete**. **SQL Server Express** is now installed.

To Download SQL Management Studio

1. Go to <http://www.microsoft.com/en-us/download/details.aspx?id=8961>
2. Click on the **Download** button.
3. Click on the **SQLManagementStudio_x64_ENU.exe** file to run the executable.
4. On the **SQL Server Installation Center** window, click on the **New SQL Server stand-alone installation or add features to an existing installation** link.

On this page ...

- [To Download SQL Server Express](#)
- [To Download SQL Management Studio](#)



INDUCTIVE
UNIVERSITY

Installing Microsoft SQL Server Express

[Watch the Video](#)

5. On the **Installation Type**, keep the defaults and click **Next**.
6. On the **License Terms** window, choose **I accept the license terms** and click **Next**.
7. On the **Product Updates** window, click **Next** to start installing the SQL Server.
8. On the **Feature Selection** window, stay with the default selection of **Management Tools** and click **Next**.
9. On the **Database Engine Configuration** window, for **Authentication Mode** choose **Mixed Mode**, enter a **password**, and click **Next**.
10. On the next few windows, continue clicking **Next** until it shows installation is **Complete**.
The **Management Tools** is now installed.
11. To run the program, go to **Start > Programs > Microsoft SQL Server > SQL Server Management Studio**.
12. Click on **Connect** to connect to the Microsoft SQL Server.
13. In **Object Explorer**, you can now see some databases under **Databases > System Databases**.
14. Right-click on **Databases** and select **New Databases....**
The **New Databases** window is displayed.
15. In **Database name**, enter **test**, click **Add**, and then **OK**.
Now you can see the **test** database in the **Databases** folder and can connect Ignition to it, see [C onnecting to Microsoft SQL Server](#).

Installing PostgreSQL

To Install the PostgreSQL Database

1. Go to the **PostgreSQL** website at <http://www.postgresql.org>
2. Click on **Downloads**, look for and click on the **Windows** link.
3. On the **Windows installers** page, find the **Download** link and click on it, and on the next page select the installer you are interested in.
For example, you can select the **Wins x86-32**.
4. From your Download folder on your computer, click on the Postgre .exe file to install the database.
5. Go through all the windows of the **Install Wizard** until installation is complete.
6. From the Windows **Start** menu, open the **PostgreSQL** pgAdmin database.
7. In the **Object browser** of the pgAdmin window, right-click on **PostgreSQL**, select **Connect**, enter your **password** to connect to the Server, and click **OK**.
You will now see the **Databases** folder in the Object browser.

On this page ...

- [To Install the PostgreSQL Database](#)



INDUCTIVE
UNIVERSITY

Installing PostgreSQL

[Watch the Video](#)

Connecting to Databases

Connect Once

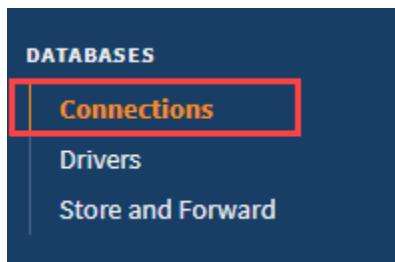
Many of the advanced features of Ignition, such as the Transaction Groups and Tags Historian require a connection to an external database and most databases require special permissions for each computer that wants to connect. Fortunately, Ignition takes care of all of this for us. You can create a connection to your database once and every system in Ignition will use that central connection. There's no need to worry about updating your database settings to add another client.

This central database connection also makes it easy to swap between databases or schemas. You can tell every query to use the default connection, then just change the default to update everything. Alternatively, you can force specific queries or systems to use a particular connection. Create as many database connections as you want and start designing using all of them.

Add a Database Connection

Now that we've installed your database, let's connect to it. You can find detailed descriptions for many database connections in this User Manual, however, they all include the same steps:

1. On the Gateway Webpage, go to the Gateway **Config** tab. Scroll down to the **Databases > Connections** section.



2. On at the Database Connections page, click on the **Create new Database Connection...** link at the bottom of the table.

A screenshot of the Ignition Gateway Config interface under the 'Connections' tab. The left sidebar shows 'Config' selected. The main area displays a table of database connections with columns: Name, Description, JDBC Driver, Translator, and Status. Three rows are listed: DB (MariaDB, MySQL, Faulted), MSSQL (MySQL, MySQL, Valid), and SQLServer (Microsoft SQLServer, MSSQL, Faulted). At the bottom of the table, there is a blue button labeled 'Create new Database Connection...' with a red box around it. A note below the table says: 'Note: For details about a connection's status, see the Database Connection Status page.' A green banner at the top says 'Trial Mode 1:24:26 We're glad you're test driving our software. Have fun.' and 'Activate Ignition'.

3. The next step is to choose a JDBC Driver.

Ignition connects to databases using JDBC drivers that are unique to each database. Drivers for the most popular databases are included so there is usually no need to install the JDBC driver manually.

Ignition ships with drivers for Microsoft SQL Server, MySQL, Oracle, and PostgreSQL. Pick the JDBC driver for your database, and click on the **Next** button.

If a suitable driver is not available in the list, you need to add a new JDBC driver for other databases, like IBM DB2, which is not very difficult to do, see [Adding a JDBC Driver](#).

4. **Configure the Connection**

After selecting the driver, you'll configure the settings for the connection. Some settings, such as the Connect URL are specific to the driver that you're using.

On this page ...

- [Connect Once](#)
- [Add a Database Connection](#)

Main Database Connection Properties

Name	Each database connection needs a unique name, which consists of letters, numbers and underscores.
Description	A brief description of the database.
JDBC Driver	The JDBC driver dictates the type of database that this connection can connect to. It cannot be changed once created.
Connect URL	A string that instructs the driver how to connect to the database. This string is the server address, and may include the port, instance name, database name, and so on. The format and parameters depend on the driver being used.
Username	The username to use when connecting. Some databases support other authentication methods, such as Windows authentication, in which case this field is not used.
Change Password?	Check the box to change the existing password.
Password	Enter password.
Password	Re-type password for verification.
Extra Connection Properties	Depending on which database you are connecting to, there will be different default values placed in this box. MS SQL Server requires you to place your database name here, but for other databases you can usually leave this at its default values. Each database has its own set of available extra connection properties so you must refer to your Database documentation to determine what is valid here.
Enabled	Lets you to enable or disable a database connection.
Failover Datasource	The connection that is automatically used when this connection is not available.
Failover Mode	Lets you select how to handle the database connection failing and recovering. Database connections support <i>failover</i> . This means that the objects which use a database connection will use a different connection if the one they are using becomes unavailable. The Failover Datasource property determines which connection is used, and the Failover Mode determines when, if ever, the connection is switch back to the primary connection. There are two failover modes: <ul style="list-style-type: none">• STANDARD mode means that this datasource will fail over when a connection cannot be retrieved, but when connectivity is restored, connections will again come from this datasource.• STICKY mode means that once this datasource fails over, connections will continue coming from the failover datasource until the failover datasource itself fails or the Gateway is restarted.
Slow Query Log Threshold	Queries that take longer than this amount of time, in milliseconds, will be logged. This helps to find queries that are not performing well. (default: 60,000)
Validation Timeout	The time in milliseconds between database validation checks. (default: 10,000)
Advanced Settings	
There are many advanced settings that you don't need to change under normal circumstances. See the description for each property on the settings page.	

In This Section ...

Connecting to IBM DB2

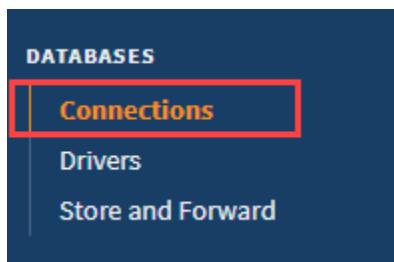
IBM DB2 Connection Requirements

In order to connect Ignition to IBM DB2 you must have a Translator, a Driver, and a Database Connection.

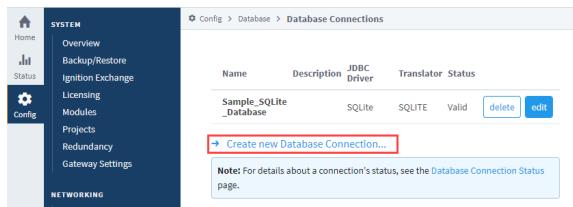
The Translator and Driver only need to be configured once. You can make as many connections as you want to any compatible IBM DB2 database.

Connect to IBM DB2

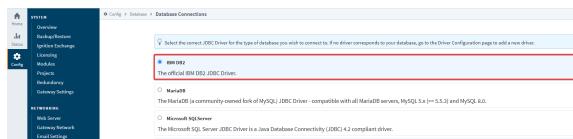
1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. The Database Connections page is displayed. Click on **Create new Database Connection....**



4. Select **IBM DB2 JDBC Driver**, and click **Next**.



i For more information, see [JDBC Drivers and Translators](#) for instructions on obtaining the requisite JAR file, and then follow the steps for [Upgrading a JDBC Driver](#) before continuing with these instructions.

5. On the New Database Connection page, enter the following information:

Name: **IBM_DB2**

Connect URL: **jdbc:db2://localhost:25000/SAMPLE**

- i**
- If on Db2 Version 11.5.5 and older, the default port number is 50000.
 - If on Db2 Version 11.5.6 and newer, the default port number is 25000.

Username: **your_username**

Password: **your_password**

On this page ...

- [IBM DB2 Connection Requirements](#)
- [Connect to IBM DB2](#)
- [Troubleshooting](#)
 - [Troubleshooting Tools](#)
 - [Installation Issues](#)
 - [Other Issues](#)
- [JDBC Drivers and Translators](#)

Main Properties

Name	IBM_DB2 Choose a name for this database connection.
Description	
JDBC Driver	IBM DB2 The JDBC driver dictates the type of database that this connection can connect to. It cannot be changed once created.
Connect URL	jdbc:db2://localhost:50000/SAMPLE The Connect URL is JDBC-driver specific. It usually contains the address of the machine that the database is running on. The format of the DB2 connect URL is: db2://host:port/database With the three parameters in bold: host: The host name or IP address of the database server. port: The port that the database server is running on. DB2 default port is 50000. database: The name of the logical database that you are connecting to on the DB2 server.
Username	your_username
Password	*****
Re-type Password	*****

6. At the bottom of the form, click on **Create New Database Connection**.

Your connection is now created. The Database Connections page is displayed and will show the status of Reconnecting, then Valid.

The screenshot shows a success message: "Successfully created new Database Connection \"IBM_DB2\"". Below it is a table with columns Name, Description, JDBC Driver, Translator, and Status. One row is shown: IBM_DB2, IBM DB2, IBM DB2, Valid. A link to "Create new Database Connection..." is also present.

7. To display the details about the status of your database connection, see the **Note** on the above window and click on the **Database Connection Status** link.

The screenshot shows a summary section with "Valid Connections: 1 / 1" and "Total Throughput: 0.1 queries/sec". Below is a table with columns Name, Driver, Status, Connections, and Throughput. One row is shown: IBM_DB2, IBM DB2, Valid, 0 / 0, 0.1 queries/sec. A "Details" button is next to the row.

Troubleshooting

Troubleshooting Tools

[IBM DOC - Troubleshooting Tools](#)

Installation Issues

DB2 installation fails to create instance and DAS: [IBM DOC - DB2 installation fails to create instance and DAS](#)

Other Issues

Error Code	Error Description	Resources
SQL30081N	Various TCP/IP communication errors. Each error has its own definition and action plan.	SQL30081N TCP/IP communication errors
SQL1092N	"USERID does not have the authority to perform the requested command or operation"	SQL1092N Error

JDBC Drivers and Translators

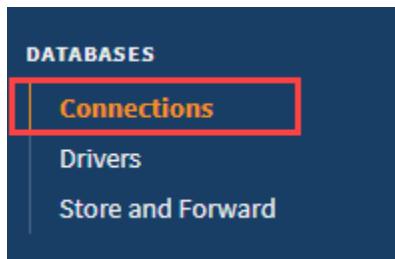
In some cases, you may need to add your own JDBC Driver, or configure a Translator. More information on configuring these can be found on the [JDBC Drivers and Translators](#) page. However, you may need to check the JDBC driver's documentation for information on how to configure them.

Related Topics ...

- [Connecting to MariaDB](#)
- [Store and Forward](#)
- [OPC UA](#)
- [Designer](#)

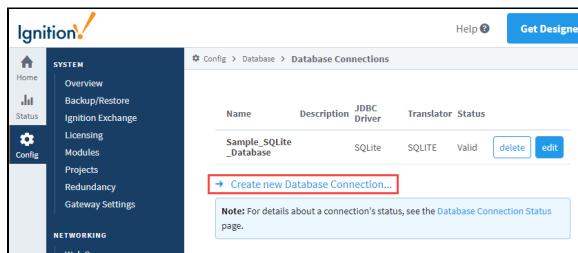
Connecting to MariaDB

1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



The Inductive University logo is displayed, featuring a green hexagon with 'IU' and a laurel wreath. To the right, the text 'INDUCTIVE UNIVERSITY' is written in a serif font. Below this, the title 'Connecting to MariaDB' is shown in bold, followed by a blue link 'Watch the Video'.

3. Click on **Create new Database Connection....**



4. Select the **MariaDB** driver and click **Next**.
5. On the New Database Connection page, enter the following information:

Name: **MariaDB** (or use a meaningful name for this connection)

Connect URL: **jdbc:mariadb://localhost:3306/myDB**

This is a screenshot of the 'New Database Connection - Main Properties' form. The 'Name' field is set to 'MariaDB'. The 'JDBC Driver' dropdown is set to 'MariaDB'. The 'Connect URL' field contains the value 'jdbc:mariadb://localhost:3306/myDB'. A detailed description of the connect URL is provided below the input field, explaining the format: 'jdbc:mariadb://host:port/database'. It specifies that 'host' is the host name or IP address, 'port' is the port number (3306 by default), and 'database' is the logical database name.

Where **localhost** is the IP address or hostname of the machine with MariaDB installed, for example: localhost, 192.168.1.1, db-server, etc.

myDB specifies the name of the database in MariaDB you want Ignition to connect to. If you're uncertain which database to connect to, speak with your local database administrator.

6. To configure the connection, Ignition needs credentials to connect to MariaDB. The **Username** and **Password** fields are where you provide credentials for a user that Ignition will use to authenticate against the database. The user should be able to do the following:
 - a. Create and drop tables within the schema
 - b. Alter tables within the schema
 - c. Insert, update, select, and delete rows from tables in the schema
 - d. Create, alter, and execute stored procedures within the schema
7. After entering a username and password, click on **Create New Database Connection** at the bottom of the form. Your connection is now created and the Database Connections page is displayed showing the **Status** of your connection as **Valid**.

Connecting to Microsoft Azure SQL

Azure SQL Server Requirements

Microsoft Azure SQL Server uses the same **Java Database Connectivity** (JDBC) driver as **Microsoft SQL Server** (MSSQL). As long as you are able to connect to an MSSQL database in Ignition, you can use the same JDBC driver to connect to Azure SQL Server. The steps below will show you how to set up Ignition to connect to Azure SQL Server.

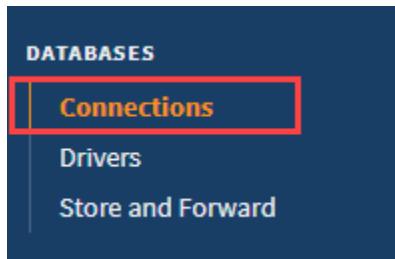
On this page ...

- [Azure SQL Server Requirements](#)
- [Connect Ignition to Azure SQL Server](#)
- [Troubleshooting and Tips](#)
 - [TCP/IP Communication](#)
 - [Windows Firewall](#)
 - [JDBC Drivers](#)

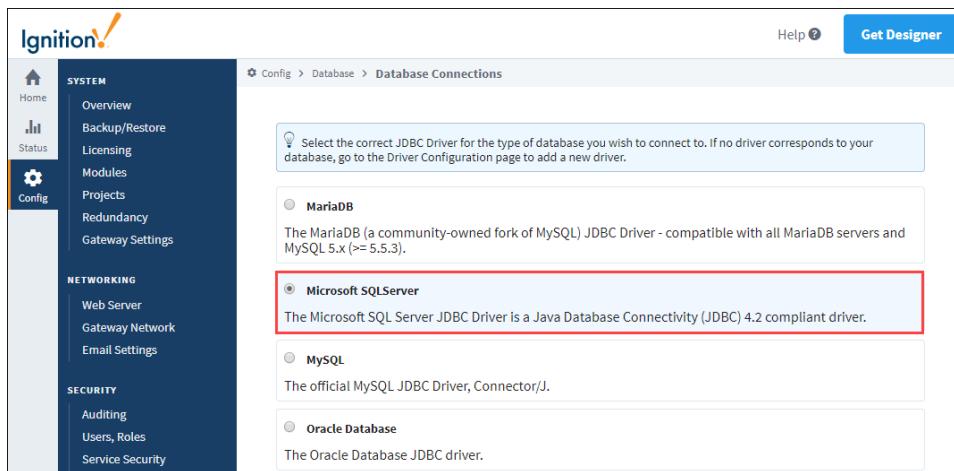
Connect Ignition to Azure SQL Server

In this example, we are using SQL Authentication. For information on using Azure Active Directory authentication with Java-based applications, see [Microsoft's official documentation](#).

1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. The Database Connections page is displayed. Click on **Create new Database Connection...**
4. Select **Microsoft SQLServer**, and click **Next**.



5. On the Database Connections page, enter the following information:
Name: The name of the connection. In this example, we are naming our connection **AzureSQLServer**.
Connect URL: The address that your Azure SQL Server instance is hosted on. In this example, we are using **jdbc:sqlserver://sudev.database.windows.net:1433**.
Username: The username of your database login credentials. In this example, we are using **dev_admin@sudev**.
Password: The password of your database login credentials.
Extra Connection Properties (Optional): Any other properties you may need to configure to connect to your Azure SQL Server instance will be listed here, such as the database name and SSL.

Note:

In this example, we are using port 1433 to connect to our Azure instance. This is the default port Ignition uses and is called **Proxy** since it uses Microsoft's proxy servers to connect. Connecting to an Azure SQL Server instance while using Microsoft's proxy servers is easier to configure as you only need to have one IP/port open. However, this method can increase latency and reduce throughput.

If you want to connect directly to your Azure SQL Server instance, you can instead use port 3306. This policy is called **Redirect** and requires many more open ports and IP addresses. However, since the connection is direct, performance will be much better with decreases in latency and increases in throughput.

You can learn more about Microsoft's proxy settings in their [Azure SQL documentation](#).

Main Properties	
Name	AzureSQLServer <input type="button" value="..."/>
Description	<input type="text"/>
JDBC Driver	Microsoft SQLServer <input type="button" value="..."/> The JDBC driver dictates the type of database that this connection can connect to. It cannot be changed once created.
Connect URL	<pre>jdbc:sqlserver://sudev.database.windows.net:1433</pre> <p>The Connect URL is JDBC-driver specific. It usually contains the address of the machine that the database is running on. The format of the SQL Server connect URL is: jdbc:sqlserver://host\instanceName[:port]</p> <p>With the three parameters (in bold) host: The host name or IP address of the database server. instanceName: (optional) the instance to connect to on the host. If not specified, a connection to the default instance is made. port: (optional) the port to connect to. The default is 1433. If you are using the default, you can omit the port and the preceding ':'. For SQL Server, you specify the <i>database name</i> to connect to using the <code>databaseName</code> property in the <i>Extra Connection Properties</i>.</p>
Username	dev_admin@sudev
Password	<input type="password"/> <input type="button" value="..."/>
Re-type Password	<input type="password"/> <input type="button" value="..."/>
Extra Connection Properties	databaseName=dev Use <code>databaseName=YOUR_DATABASE</code> to specify the database to connect to.

6. At the bottom of the form, click in **Create New Database Connection**.

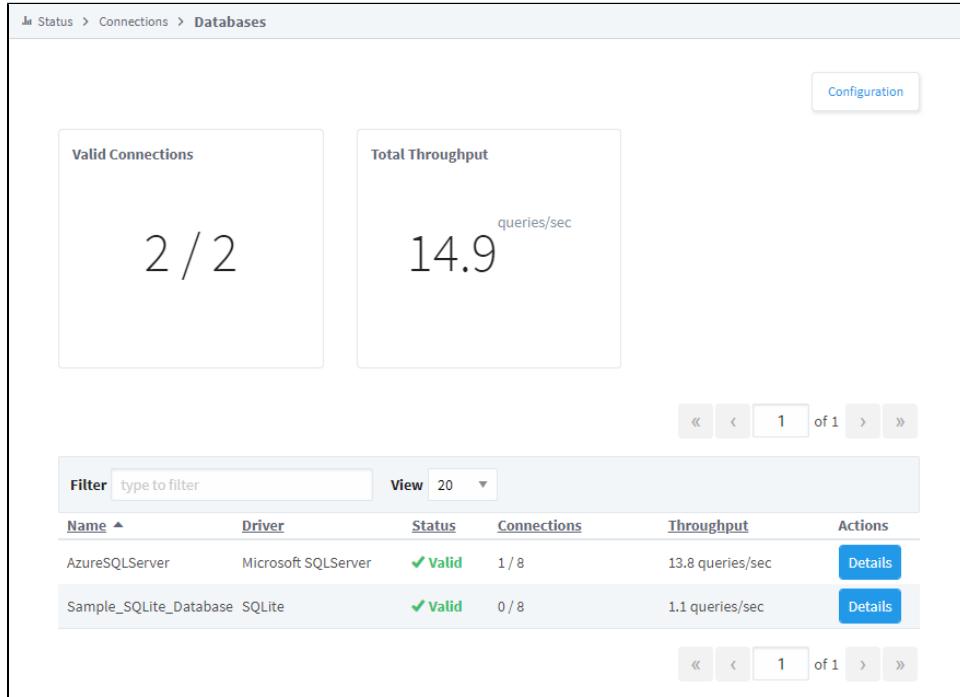
Your connection is now created. The Database Connections page is displayed and will show the status of **Reconnecting**, then **Valid**.

Name	Description	JDBC Driver	Translator	Status
AzureSQLServer		Microsoft SQLServer	MSSQL	Valid <input type="button" value="delete"/> <input type="button" value="edit"/>
Sample_SQLite_Database		SQLite	SQLITE	Valid <input type="button" value="delete"/> <input type="button" value="edit"/>

→ Create new Database Connection...

Note: For details about a connection's status, see the [Database Connection Status page](#).

7. To display the details about the status of your database connection, see the **Note** on the above window and click on the **Database Connection Status** link.



Troubleshooting and Tips

TCP/IP Communication

If your database connection is not successfully connecting, you may need to check your TCP/IP settings. See Microsoft's documentation on [configuring IP addresses for an Azure network interface](#).

Windows Firewall

In some situations, you may need to configure your firewall to set up any required ports. Typically, ports 1433 and 1434 need to be open for TCP traffic, but other ports may also need to be configured, such as when connecting to your Azure SQL Server instance directly using port 3306.

JDBC Drivers

When using a JDBC driver to connect Ignition to Azure SQL Server, there are two things you may need to check:

- Make sure your JDBC driver is up to date.
- Make sure the JDBC driver version you are using is the same as the Java version Ignition is using. Using an incompatible JDBC driver may result in Gateway errors and crashes.

Connecting to Microsoft SQL Server

SQL Server Connection Requirements

In order to get connected to SQL Server, you must have a Translator, a Driver, and a Connection. The Translator and Driver only needs to be installed once, and after that you can make as many connections as you want to any compatible SQL Server databases.

Note: When you Upgrade Ignition, any existing drivers are carried over. This means only a fresh install of Ignition will not have a SQL Server Connector.

On this page...

- [SQL Server Connection Requirements](#)
- [Connect to Microsoft SQL Server](#)
- [Microsoft SQL Server Connection Guide](#)
- [Different Ways of Connecting to SQL Server](#)
- [Troubleshooting](#)
- [JDBC Drivers and Translators](#)

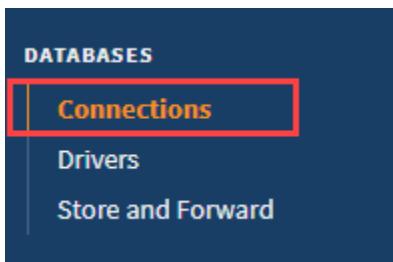


Connecting to Microsoft SQL Server Express

[Watch the Video](#)

Connect to Microsoft SQL Server

1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. The Database Connections page is displayed. Click on **Create new Database Connection....**
4. Select **Microsoft SQLServer JDBC Driver**, and click **Next**.

The screenshot shows the Ignition Config interface. On the left, there's a sidebar with 'Config' selected. The main area is titled 'Database Connections'. It has a note: 'Select the correct JDBC Driver for the type of database you wish to connect to. If no driver corresponds to your database, go to the Driver Configuration page to add a new driver.' Below this, there are four options: 'MariaDB', 'Microsoft SQLServer' (which is selected and highlighted with a red border), 'MySQL', and 'Oracle Database'.

5. On the **Database Connections** page, enter the following information:

Name: **SQLServer**

Connect URL: **jdbc:sqlserver://localhost\SQLEXPRESS**

Note: We are connecting to the express edition of SQL Server using the default instance name. If you have the full SQL Server with default settings, replace SQLEXPRESS with the instance name of your SQL Server installation. Non-express versions of SQL Server tend to use MSSQLSERVER as a default instance name.

username: **sa**

password: **sqlserver** (password is what you entered during the SQL Server installation. For this example, password is **sqlserver**)

The screenshot shows the 'Main Properties' section of the Database Connections configuration. The 'Name' field is set to 'SQLServer'. The 'JDBC Driver' is set to 'Microsoft SQLServer'. The 'Connect URL' field contains 'jdbc:sqlserver://localhost\SQLEXPRESS'. The 'Username' field is set to 'sa'. The 'Password' field is filled with '.....' and the 'Re-type password for verification' field also contains '.....'.

6. At the bottom of the form, click on **Create New Database Connection**.

Your connection is now created. The Database Connections page is displayed and will show the status of Reconnecting, then Valid.

Config > Database > Database Connections

Successfully created new Database Connection "SQLServer"

Name	Description	JDBC Driver	Translator	Status		
DB		MariaDB	MySQL	Valid	delete	edit
SQLServer		Microsoft SQLServer	MSSQL	Reconnecting	delete	edit

[Create new Database Connection...](#)

Note: For details about a connection's status, see the [Database Connection Status](#) page.

- To display the details about the status of your database connection, see the **Note** on the above window and click on the **Database Connection Status** link.

Status > Connections > Databases

Help ? [Get Designer](#)

Valid Connections: 1 / 2

Total Throughput: 12.2 queries/sec

Name	Driver	Status	Connections	Throughput	Actions
SQLServer	Microsoft SQLServer JDBC Driver	✓ Valid	0 / 8	12.0 queries/sec	Details

Microsoft SQL Server Connection Guide

This guide helps you with any difficulties you may have in getting the correct settings and parameters when connecting Ignition to Microsoft SQL Server, a popular and robust relational database.

Multiple Instances of Database

Microsoft SQL Server supports multiple instances of the database running concurrently on the same computer. Each instance has its own name and set of system and user databases that are not shared between instances. Applications, such as Ignition, can connect to each instance on a computer in much the same way they connect to databases running on different computers. By default, each instance gets assigned a dynamic TCP/IP port on startup that listens for any incoming requests. Since the port is dynamic and the application does not know what the new port is, it must connect using the instance name.

So if the communication is over TCP/IP and the application knows the instance name, how does the application find which port to communicate to?

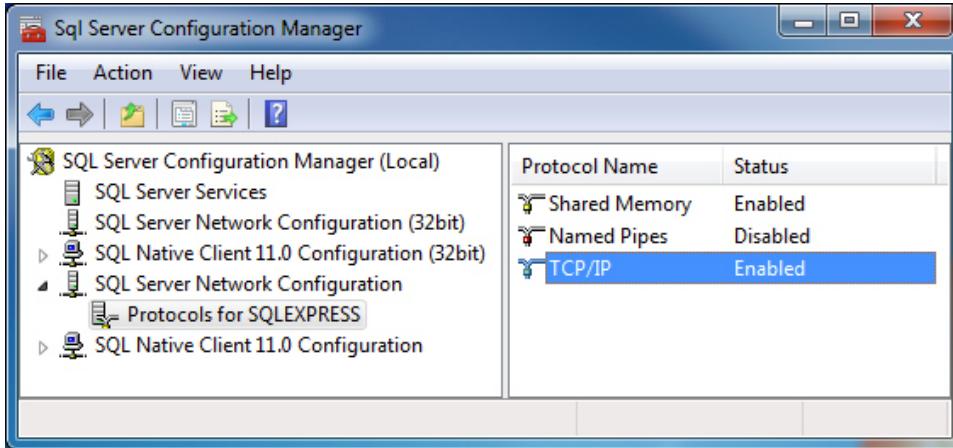
The answer is the **Microsoft SQL Server Browser** service. The Microsoft SQL Server Browser program runs as a Windows service and listens for all incoming requests for resources and provides information, such as the TCP/IP port, about each instance installed on the computer. Microsoft SQL Server Browser also contributes to these two actions: browsing a list of available servers and connecting to the correct server instance.

If the Microsoft SQL Server Browser service is not running, you can still connect to SQL Server if you provide the correct port number. For example, you can connect to the default instance of SQL Server with TCP/IP if it is running on port 1433.

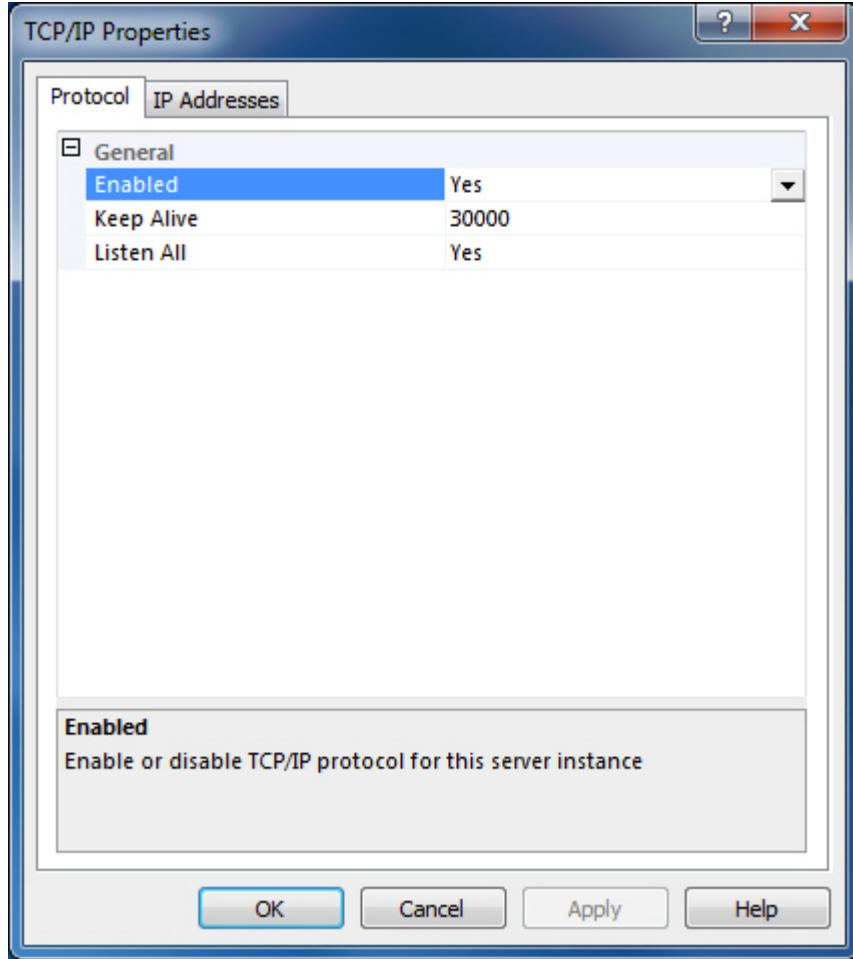
Check 1: Make Sure the Database has TCP/IP Enabled

Ignition connects using TCP/IP, therefore make sure your database has TCP/IP enabled.

1. Open the SQL Server Configuration Manager from **Start > All Programs > Microsoft SQL Server Version # > Configuration Tools > SQL Server Configuration Manager**.
The Sql Server Configuration Manager window is displayed.
2. To see all the instances set up on that machine, expand **SQL Server Network Configuration**.
3. Find the database (or instance) you plan on using. To the right, all of the protocols the database supports are shown. Find the **TCP/IP** protocol and select it.



4. Make sure the **Status** next to TCP/IP is **Enabled**. If not, double-click **TCP/IP** and choose **Yes** from the drop-down next to Enabled and click **OK**

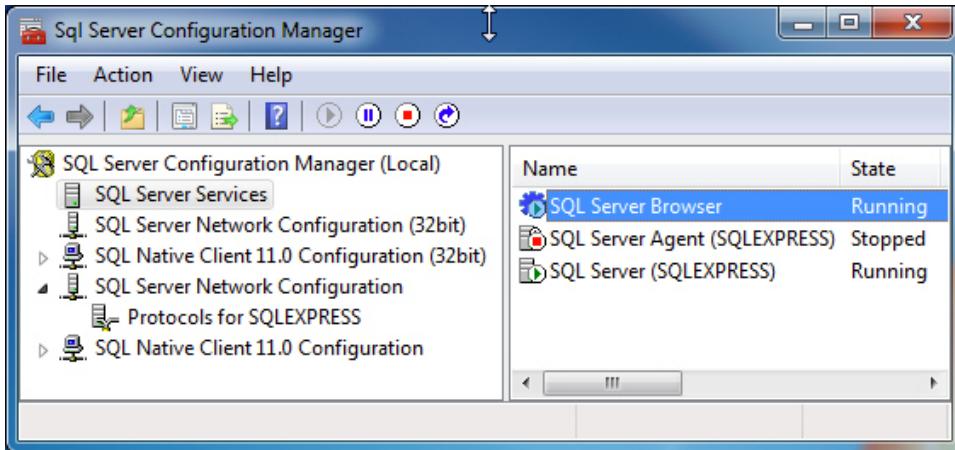


Check 2: Make Sure Microsoft SQL Server Browser is Running

If you ARE connecting to your database using a NAMED INSTANCE, you must make sure that the Microsoft SQL Server Browser is running. As mentioned earlier, the Microsoft SQL Server Browser translates the instance name to a TCP/IP port in order for Ignition to connect to it.

1. Open the **SQL Server Configuration Manager** from Start > All Programs > Microsoft SQL Server Version # > Configuration Tools > **SQL Server Configuration Manager**.
2. Select the **SQL Server Services** section.
3. On the right, see all of the services installed. One of the services is **SQL Server Browser**. Make sure this service is in fact running. If the service is not running, right-click and select **Start**.

Note: The service could be disabled, so you may need to double-click it to enable the service before starting it up.



Check 3: Make Sure There is a Database Created

If you are connecting to an existing SQL Server installation, you only need to know the name of the database and you can skip this step. In newer versions of SQL Server, a fresh installation does not include a database so you must first create a new database.

1. Open the **SQL Server Management Studio** program. This program was an option when you installed SQL Server.
2. Log into the Management Studio using either SQL Server Authentication or Windows Authentication.
3. In the **Object Explorer** (on the left), expand the instance folder to find the **Databases** folder. The instances folder is usually expanded by default.
4. Right Click on the **Databases** folder and select the **New Database...** option.
5. Type in a Database Name.
6. Click the OK button in the lower right. You can then expand the **Databases** folder in the Object Explorer to see the new database. If it doesn't show up right away you can right-click to refresh.
- Security** - Make your user a db_owner of the new database.
7. In the Object Explorer, expand to the **Security > Logins** folder, right click on your username, and select **Properties**.
8. In the **User Mapping** page click the checkbox next to your new database, then click the **db_owner** checkbox below.
9. Click OK. Now your user has access to a database.

Different Ways of Connecting to SQL Server

Now that you have ensured that TCP/IP is enabled and the Microsoft SQL Server Browser is running, you can connect to Microsoft SQL Server in four different ways (all using TCP/IP communication) as follows:

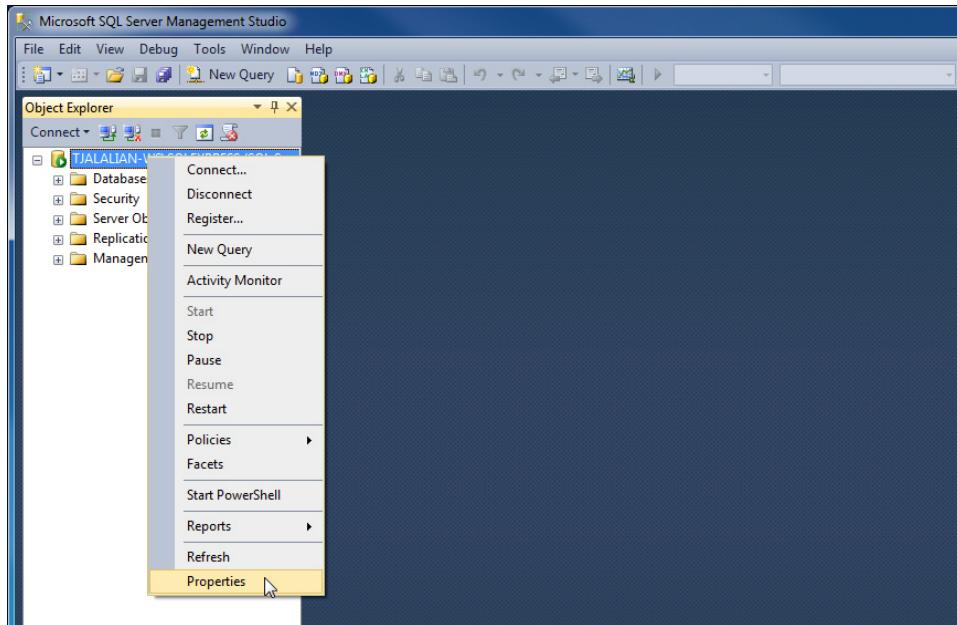
1. Connect using an Instance Name and SQL Authentication.
2. Connect using an Instance Name and Windows Authentication (this is the most common method).
3. Connect using a Port and SQL Authentication.
4. Connect using a Port and Windows Authentication.

Scenario 1: Connect By Using an Instance Name and SQL Authentication

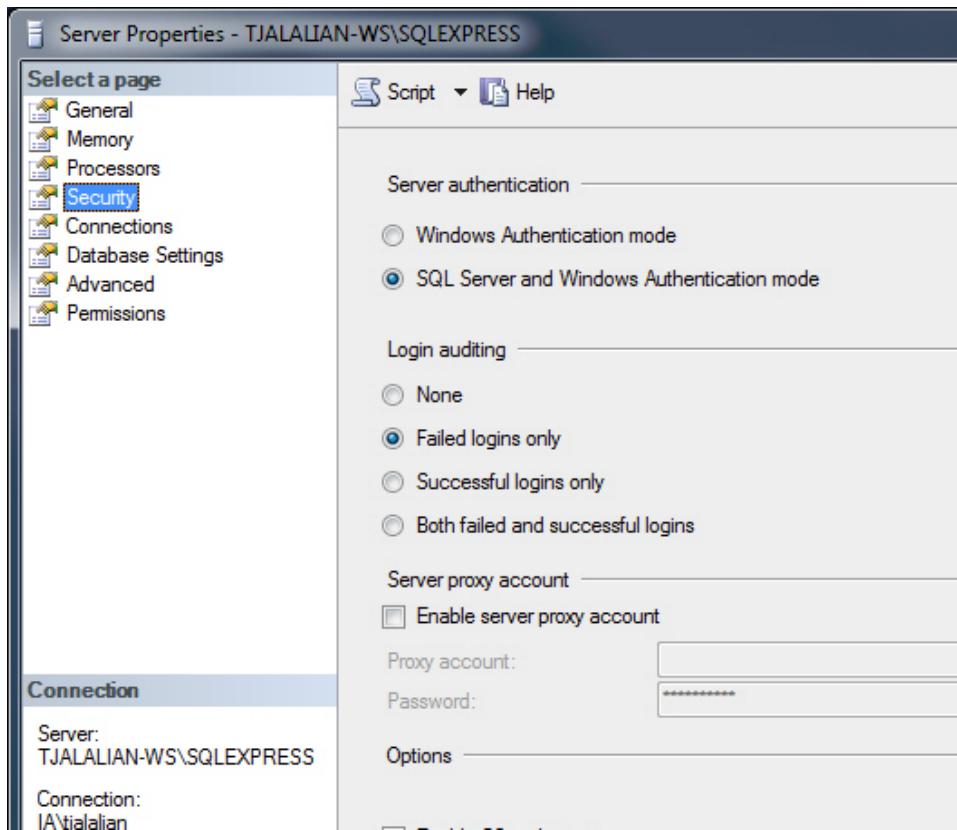
By default, Microsoft SQL Server only allows Windows authentication since it is more secure. But because we are using SQL authentication, we must enable Microsoft SQL to allow this type of authentication.

Enable SQL Authentication

1. Open the **Microsoft SQL Server Management Studio** window from **Start > All Programs > Microsoft SQL Server Version # > SQL Server Management Studio**.
The window is displayed showing connections to your database.
2. Right-click the top-level database in the **Object Explorer** and select **Properties**.



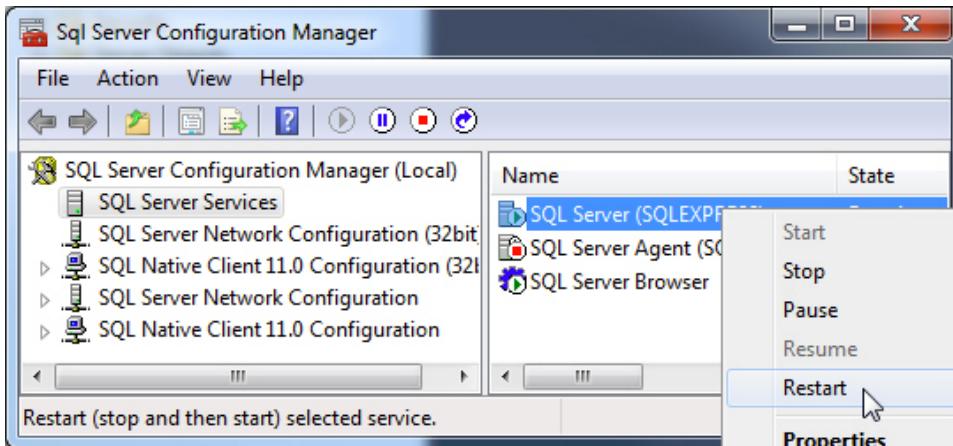
3. From the Server Properties window, on the left side, select **Security**.



4. Verify that **SQL Server and Windows Authentication mode** is selected.
If not, select it and click **OK**.

Now you need to restart the **SQL Server Windows** service so that this setting takes effect.

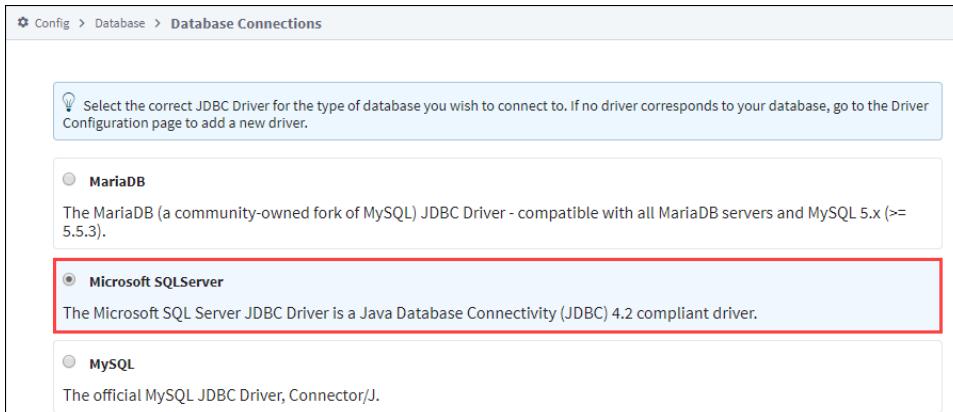
5. Open the **SQL Server Configuration Manager** at
Start > All Programs > Microsoft SQL Server Version # > Configuration Tools > SQL Server Configuration Manager.
6. Select the **SQL Server Services** section and restart the **SQL Server (Instance Name)** item.



Now that Microsoft SQL Server accepts SQL authentication, we can configure Ignition.

Configure the Database Connection in Ignition

1. Go to and login to the Ignition Gateway Config page from your webbrowser at <http://hostname:8088/main/web/config/>
2. Select **Databases > Connections** from the menu.
3. Click on **Create new Database Connection**.
4. Select **Microsoft SQL Server JDBC Driver** and click **Next**.



5. In the **New Database Connection** window, enter the following information:

Name: **SQLServer_SQLAuth** (no spaces)

Connect URL: **jdbc:sqlserver://Hostname\InstanceName**

where **Hostname** is your databases IP address or hostname and **InstanceName** is your databases instance name, for example:
jdbc:sqlserver://localhost\SQLEXPRESS
jdbc:sqlserver://10.10.1.5\MSSQLSERVER

6. Set the username and password to a valid SQL authentication user. For example, **sa** is the default administrator account you can use.

The screenshot shows the 'Main Properties' section of a database connection configuration. The connection is named 'SQLServer_SQLAuth'. The 'JDBC Driver' is set to 'Microsoft SQLServer'. The 'Connect URL' is 'jdbc:sqlserver://localhost\SQLEXPRESS'. The 'Username' is 'sa' and the 'Extra Connection Properties' include 'databaseName=test'.

Main Properties	
Name	SQLServer_SQLAuth <input type="button" value="..."/>
Description	<input type="text"/>
JDBC Driver	Microsoft SQLServer <input type="button" value="..."/>
Connect URL	<p>jdbc:sqlserver://localhost\SQLEXPRESS</p> <p>The Connect URL is JDBC-driver specific. It usually contains the address of the machine that the database is running on.</p> <p>The format of the SQL Server connect URL is: jdbc:sqlserver://host\instanceName[:port]</p> <p>With the three parameters (in bold)</p> <ul style="list-style-type: none"> host: The host name or IP address of the database server. instanceName: (optional) the instance to connect to on the host. If not specified, a connection to the default instance is made. port: (optional) the port to connect to. The default is 1433. If you are using the default, you can omit the port and the preceding ':'. For SQL Server, you specify the <i>database name</i> to connect to using the <code>databaseName</code> property in the <i>Extra Connection Properties</i>.
Username	sa <input type="text"/>
Password	<input type="password"/> <input type="button" value="..."/>
Re-type Password	<input type="password"/> <input type="button" value="..."/>
Extra Connection Properties	databaseName=test Use <code>databaseName=YOUR_DATABASE</code> to specify the database to connect to.

7. To add your own user account, open the SQL Server Management Studio and expand the **Security > Logins** folder. You will see all the current logins including **sa** and you can add a new login.

- To add a new login, right-click on the **Logins** folder and click **New Login...**. The **Login** window is displayed.
- Choose the **SQL Server authentication** mode and type in a Login name and password.

Note: You will also have to add permissions to your database by mapping db_datareader and db_datawriter to the new user in the User Mapping section of the Login window. If you want Ignition to be able to create tables (ie: for Tag History), you also need to give table creation access such as db_owner.

8. Go back to the **New Database Connection** page in the Gateway, enter the name of your database, for example, in the **Extra Connection Properties** enter: **databaseName=test** (replace **test** with your database name, not the instance name).
9. Click **Create New Database Connection**. The Database Connection page is displayed showing the Status as **Valid** after a couple of seconds. If the connection is **Faulted**, click on the Database Connection Status link to find out why. Typically, the username/password is incorrect or the user doesn't have the right permissions.

Scenario 2: Connect By Using Instance Name and Windows Authentication

In Windows authentication mode, the username and password used to connect comes from the Ignition Windows Service logon. By default, the Ignition Windows Service is set to local system account which usually doesn't have privileges to connect.

Set Up the Service to Use Windows Authentication

- Download a copy of the [SQL Server JDBC driver](#). Specifically, download a ZIP or tar.gz file (NOT an installer), as you will need to extract a specific file and relocate it to the Gateway's installation directory. The exact version required depends on the version of Java your Gateway is using. Ignition 8.1.33+ uses Java 17, but previous 8.0/8.1 versions use Java 11.



Although, it's recommended you use the most recent driver available for your system, make sure the version of the JDBC Driver Ignition is using matches the downloaded DLL file. Your current driver version can be seen in the following directory:

```
{installDirectory}/user-lib/jdbc
```

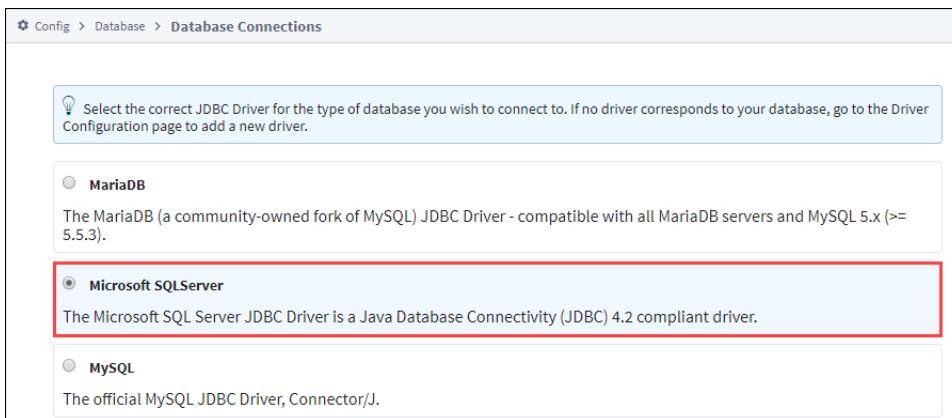
2. Locate the DLL file from the correct architecture folder ("x64" for 64-bit JDBC) inside of the `enu/auth` folders in the zip file.
3. Copy the DLL file to the `lib` folder in your install directory. If you have the default install directory, it's in the following location:
`C:\Program Files\Inductive Automation\Ignition\lib`

Note: Older versions of the JDBC driver (such as version 7.2.1) need to be renamed to `sqljdbc_auth.dll` before Ignition can utilize the file. However, in more modern versions of the driver, this is no longer the case.

4. The account used to connect will be the account that Ignition is running under in the services menu. To set up Ignition to logon using the right Windows account, open the **Services Control Panel** from Start > Control Panel > Administrative Tools > Services
5. Right-click the **Ignition** service (or whatever service name your Ignition installation is using) and choose **Properties**.
6. Select the **Log On** tab.
7. Choose the **This account** radio button and enter in your Windows username and password.
8. Click **OK** to save.
9. Now restart the Ignition service to make this change take effect. Click the **Action > Restart** button in the menubar to restart the Ignition service (or you can stop and start from the right-click menu).

Configure the Database Connection in Ignition

1. Go to and login to the Ignition Gateway Config page from your webbrowser at <http://hostname:8088/main/web/config/>
2. Select **Databases > Connections** from the menu.
3. Click on **Create new Database Connection**.
4. Select **Microsoft SQL Server JDBC Driver** and click **Next**.



5. On the **New Database Connection** page, enter the following information:
Name: **SQLServer_WinAuth** (no spaces)

Connect URL: **jdbc:sqlserver://Hostname\InstanceName**

where **Hostname** is your databases IP address or hostname and **InstanceName** is your databases instance name, for example:
`jdbc:sqlserver://localhost\SQLEXPRESS`
`jdbc:sqlserver://10.10.1.5\MSSQLSERVER`

username: leave blank
password: leave blank

Extra Connection Properties:
databaseName=test; integratedSecurity=true; (replace test with your database name)

Config > Database > Database Connections

Main Properties	
Name	<input type="text" value="SQLServer_WinAuth"/> [i] Choose a name for this database connection.
Description	<input type="text"/>
JDBC Driver	Microsoft SQLServer [▼] The JDBC driver dictates the type of database that this connection can connect to. It cannot be changed once created.
Connect URL	<input type="text" value="jdbc:sqlserver://localhost:MSSQLSERVER"/> The Connect URL is JDBC-driver specific. It usually contains the address of the machine that the database is running on. The format of the SQL Server connect URL is: <code>jdbc:sqlserver://host\instanceName[:port]</code> With the three parameters (in bold) host : The host name or IP address of the database server. instanceName : (optional) the instance to connect to on the host. If not specified, a connection to the default instance is made. port : (optional) the port to connect to. The default is 1433 . If you are using the default, you can omit the port and the preceding '!'. For SQL Server, you specify the <i>database name</i> to connect to using the <code>databaseName</code> property in the <i>Extra Connection Properties</i> .
Username	<input type="text"/>
Password	<input type="password"/> [i]
Password	<input type="password"/> [i] Re-type password for verification.
Extra Connection Properties	<input type="text" value="databaseName=test;integratedSecurity=true"/> Use <code>databaseName=YOUR_DATABASE</code> to specify the database to connect to.
Enabled	<input checked="" type="checkbox"/> Disabling a connection will prevent communication to the target database. (default: true)
Failover Datasource	<input type="text" value="- none -"/> [▼] Another datasource that will be used to handle queries if this datasource faults.

6. Click on **Create New Database Connection**.

The Status should be Valid after a couple of seconds. Again, if the connection is Faulted, click the Database Connection Status link to find out why.

Scenario 3: Connect By Using Port and SQL Authentication

1. Connecting by using a port and SQL authentication is just like scenario 1 above except you specify a port instead of the instance name in the **New Database Connection** page.
2. Enter the following:

Connect URL: `jdbc:sqlserver://Hostname:Port`

where **Hostname** is your databases IP address or hostname and **Port** is your databases TCP/IP port (SQLSERVER default port is 1433), for example:

`jdbc:sqlserver://localhost:1433`
`jdbc:sqlserver://10.10.1.5:1433`

Scenario 4: Connect By Using Port and Windows Authentication

Connecting by using a port and Windows authentication is just like scenario 2 above except you specify a port instead of the instance name in the **New Database Connection** page. Don't forget to download the `sqljdbc_auth.dll` file if you need it.

Enter the following:

Connect URL: `jdbc:sqlserver://Hostname:Port`

where **Hostname** is your databases IP address or hostname and **Port** is your databases TCP/IP port (SQLSERVER default port is 1433), for example:
`jdbc:sqlserver://localhost:1433`
`jdbc:sqlserver://10.10.1.5:1433`

Troubleshooting

TCP/IP Communication Not Enabled

SQL Server requires that you explicitly turn on TCP connectivity. To do this, use the SQL Server Configuration Manager, located in the **Start** menu under **Microsoft SQL Server > Configuration Tools**. Under **SQL Server Network Configuration**, select your instance, and then enable TCP/IP in the panel to the right. You need to restart the server for the change to take affect.

Window Firewall

When connecting remotely, make sure that Windows Firewall is disabled, or set up to allow the necessary ports. Normally ports 1434 and 1433 must be open for TCP traffic, but other ports may be required based on configuration.

SQL Server Browser Process Not Running

To connect to a named instance, the **SQL Server Browser** service must be running. It is occasionally disabled by default, so you need to verify that the service is not only running, but set to start automatically on bootup. The service can be found in the Windows Service Manager (**Control Panel > Administrative Tools > Services**).

Mixed Mode Authentication Not Enabled

Unless selected during setup, **mixed mode** or **SQL authentication** is not enabled by default. This mode of authentication is the **username/password** scheme that most users are used to. When not enabled, SQL Server only allows connections using Windows Authentication. Due to the ease of using SQL Authentication over Windows Authentication, we recommend enabling this option and defining a user account for Ignition.

1. To enable this, open the SQL Server Management Studio.
2. Connect to the server.
3. Right click on the instance and select **Properties**.
4. Under **Security**, select **SQL Server and Windows Authentication mode**.

TLS/Security

An update starting in Java 8 disables TLS 1.0 and 1.1, which may affect your connection to an MSSQL database. The [Support](#) department has written a [guide on how to resolve this issue](#).

JDBC Drivers and Translators

In some cases, you may need to add your own JDBC Driver, or configure a Translator. More information on configuring these can be found on the [JDBC Drivers and Translators](#) page. However, you may need to check the JDBC driver's documentation for information on how to configure them.

If you are using a more recent JDBC driver such as 10.2 or 11.2, you will need to add the following argument to the extra connection properties: **trustServerCertificate=true**. This will allow the transport layer to use SSL to encrypt the channel and bypass going through the certificate chain to validate trust.

Related Topics ...

- [Connecting to Oracle Express](#)
- [Store and Forward](#)
- [OPC UA](#)
- [Designer](#)

Connecting to MySQL

On this page we'll demonstrate how to connect Ignition to MySQL.

MariaDB Connections to MySQL Databases

Ignition can use the built-in MariaDB driver to connect to MySQL 5.7 and prior databases. This circumvents the need to manually provide a JAR file to the MySQL JDBC Driver configuration on the Gateway.

The following feature is new in Ignition version **8.1.2**.
[Click here](#) to check out the other new features

Ignition version 8.1.2 includes a **MariaDB** driver that can connect to **MySQL 8** databases. Note that upgrading Ignition does not replace existing JDBC drivers. See the [JDBC Drivers and Translators](#) page for more information.

MySQL Connector/J Connections to MySQL Databases

Should you choose to connect to a MySQL database using the official MySQL JDBC Driver, you will need to install the driver. New Ignition installations do not have the driver (a JAR file), so you will have to acquire the file yourself. See the [JDBC Drivers and Translators](#) page for more details on obtaining the required file.

Once acquired, you can follow the steps for [Upgrading a JDBC Driver](#). Once the JAR file has been provided, you can follow the steps listed on this page to configure a connection between Ignition and MySQL.

On this page ...

- [MariaDB Connections to MySQL Databases](#)
- [MySQL Connector/J Connections to MySQL Databases](#)
- [Connect Ignition to MySQL Database](#)

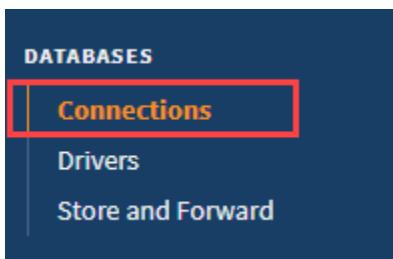


Connecting to MySQL

[Watch the Video](#)

Connect Ignition to MySQL Database

1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. Click on **Create new Database Connection....**

A screenshot of the 'Database Connections' table in the Ignition Config. The table has columns: Name, Description, JDBC Driver, Translator, and Status. One row is shown: 'Sample_SQLite_Database', 'SQLite', 'SQLITE', 'Valid', with 'delete' and 'edit' buttons. Below the table is a button labeled 'Create new Database Connection...'. A note at the bottom states: 'Note: For details about a connection's status, see the [Database Connection Status](#) page.'

4. Select the **MySQL** driver and click **Next**. You may need to install a JDBC driver for MySQL.

The screenshot shows the Ignition configuration interface for database connections. A red warning box at the top states: "⚠ The MySQL driver is missing required files. Please see [this help page](#) for information about how to install the necessary driver files for your database." Below this, there is a note: "💡 Select the correct JDBC Driver for the type of database you wish to connect to. If no driver corresponds to your database, go to the Driver Configuration page to add a new driver." Three options are listed: "MariaDB" (radio button), "Microsoft SQLServer" (radio button), and "MySQL" (radio button, which is selected). Below each option is a brief description. At the bottom of the page, there is a "Next Step" button.

For more information, see [JDBC Drivers and Translators](#) for instructions on obtaining the requisite JAR file, and then follow the steps for [Upgrading a JDBC Driver](#) before continuing with these instructions.

Note: Alternatively, you may select the **MariaDB** driver, and use that to connect to MySQL. See the [MariaDB Connections to MySQL Databases](#) section above.

5. On the New Database Connection page, enter the following information:

Name: **MySQL** (use a meaningful name such as **MySQL**)

Connect URL: **jdbc:mysql://localhost:3306/test** (By default, MySQL creates an empty database called test)

The screenshot shows the 'Main Properties' tab of the database connection configuration. It includes fields for Name (MySQL), Description (empty), JDBC Driver (MariaDB), and Connect URL (jdbc:mariadb://localhost:3306/test). The 'Connect URL' field has a detailed description explaining the format: "The Connect URL is JDBC-driver specific. It usually contains the address of the machine that the database is running on. The format of the MariaDB connect URL is: [jdbc:mariadb://]host[:port]/database]. With the three parameters (in bold): host: The host name or IP address of the database server. port: The port that the database server is running on. MariaDB default port is 3306. database: The name of the logical database that you are connecting to on the MariaDB server."

As you see in the example above, MySQL uses the following **Connect URL** format:
jdbc:mysql://hostaddress:3306/database

Where **hostaddress** is the address of the machine with MySQL installed, for example: localhost, 192.168.1.1, db-server, etc. **Database** specifies the database schema the connection will target. The connection will target one schema (a collection of tables and other objects) in the database.

6. To configure the connection, Ignition needs credentials to connect to MySQL. The **Username** and **Password** fields are where you provide credentials for a user that Ignition will use to authenticate against the database. The user should be able to do the following:
- Create and drop tables within the schema
 - Alter tables within the schema
 - Insert, update, select, and delete rows from tables in the schema
 - Create, alter, and execute stored procedures within the schema
7. After entering a username and password, click on **Create New Database Connection** at the bottom of the form. Your connection is now created and the Database Connections page is displayed showing the **Status** of your connection as **Valid**.

Config > Database > Database Connections

Successfully created new Database Connection "MySQL"

Name	Description	JDBC Driver	Translator	Status		
MySQL		MySQL	MYSQL	Valid	delete	edit
Sample_SQLite_Database		SQLite	SQLITE	Valid	delete	edit

→ Create new Database Connection...

8. To display the details about the status of your database connection, see the **Note** on the above window and click on the **Database Connection Status** link. This will display any errors if your status is Faulted, in this example it shows the status as being Valid.

Connecting to Oracle Express

This page documents how to configure a database connection to an Oracle Express instance.

Oracle User Grants

When using an Oracle Express database connection, it is required to provide user credentials that have grants for "CREATE TRIGGER" and "CREATE SEQUENCE". Some of Ignition's subsystems, such as the [Tag Historian](#), will fail to work properly if the user defined in the Database Connection does not have these grants. In addition, any manual queries (for example, those called by `system.db.runNamedQuery`) that need to insert records or create sequences may fail if the gateway does not have the grants.

On this page ...

- [Oracle User Grants](#)
- [Connect Ignition to the Oracle Express Database](#)
- [JDBC Drivers and Translators](#)

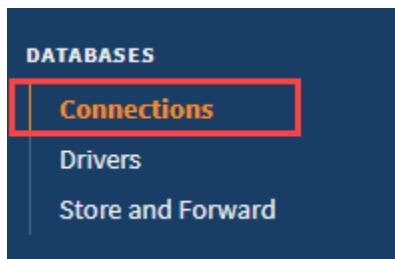


Connecting to Oracle Express

[Watch the Video](#)

Connect Ignition to the Oracle Express Database

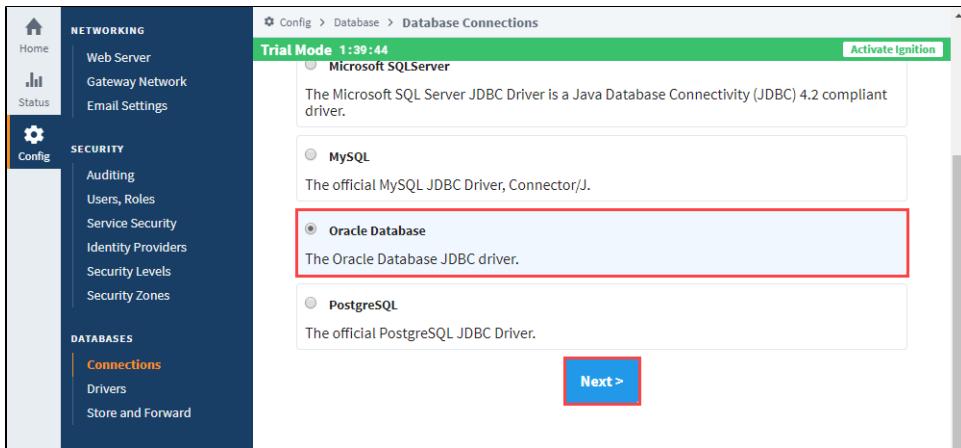
1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. The Database Connections page is displayed. Click on **Create new Database Connection...**

A screenshot of the Ignition Config Database Connections page. A red box highlights the 'Create new Database Connection...' button. Below it, a note states: 'Note: For details about a connection's status, see the Database Connection Status page.' A table lists one existing connection: Name: DB, Description: MariaDB, JDBC Driver: MySQL, Translator: None, Status: Faulted. There are 'delete' and 'edit' buttons for this row.

4. Select the **Oracle JDBC Driver** and click **Next**.



5. Enter the following information:

Property Name	How should it be configured
Name	The name of the database connection. Other features in Ignition will reference this connection by the name specified here.
Connect URL	<p>A URL that describes where on the network the database is, as well as information about which schema to connect to. The initial value should look something like the following</p> <pre>jdbc:oracle:thin:@localhost:1521:XE</pre> <p>However, you may need to make some changes. If we examine the key pieces of the URL, it would look like this:</p> <pre>jdbc:oracle:thin://1:2:3</pre> <p>Where:</p> <ul style="list-style-type: none"> 1 = The IP Address or hostname of the computer/server that the database is installed on. 2 = The port that the database is running on. The default is 1521, but this could have been changed during installation, or sometime afterwards. When in doubt, ask the person that installed the database. 3 = The System ID for the database. More information on the System ID can be found in Oracle's documentation.
Username	The username that the Gateway will use to connect to the database.
Password	The password for the user specified under the Username property.

For our example, we entered information relevant to our installation. The values seen below may differ from each installation.

The screenshot shows the Ignition Configuration interface. On the left, there's a sidebar with various system configurations like Overview, Backup/Restore, and Database. Under the Database section, 'Connections' is selected. The main area is titled 'Database Connections' and shows a table with two rows: MySQL and Oracle. The Oracle row has a red box around its 'Status' column, which is labeled 'Valid'. At the bottom of the page, there's a note: 'Successfully created new Database Connection "Oracle"'.

Name	Description	JDBC Driver	Translator	Status
MySQL		MySQL	MYSQL	Faulted
Oracle		Oracle Database	ORACLE	Valid

6. Click **Create New Database Connection** at the bottom of the form.
7. The Database Connections page is displayed showing the **Status** of your connection as **Valid**.

The screenshot shows the Ignition Configuration interface. The sidebar and main content area are identical to the previous screenshot, but the status for the Oracle connection has changed from 'Faulted' to 'Valid', indicated by a green checkmark icon and the word 'Valid' in the status column.

Name	Description	JDBC Driver	Translator	Status
MySQL		MySQL	MYSQL	Faulted
Oracle		Oracle Database	ORACLE	Valid

8. To display the details about the status of your database connection, see the **Note** on the above window and click on the **Database Connection Status** link.



JDBC Drivers and Translators

In some cases, you may need to add your own JDBC Driver, or configure a Translator. More information on configuring these can be found on the [JDBC Drivers and Translators](#) page.

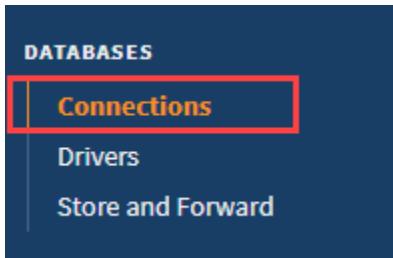
Related Topics ...

- [Store and Forward](#)
- [OPC UA](#)
- [Designer](#)

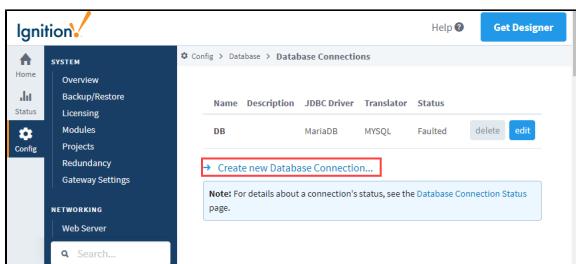
Connecting to PostgreSQL

Connect Ignition to the PostgreSQL Database

1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. The Database Connections page is displayed. Click on **Create new Database Connection....**



4. Select the **PostgreSQL JDBC Driver** and click **Next**.



5. Next, you'll need to enter information that will allow the Gateway to connect to a Postgres instance. In most cases, the following properties should be updated:

Property Name	How should it be configured
Name	The name of the database connection. Other features in Ignition will reference this connection by the name specified here.
Connect URL	A URL that describes where on the network the database is, as well as information about which database to connect to. The initial value should look something like the following: <code>jdbc:postgresql://localhost:5432/database</code> However, you may need to make some changes. If we examine the key pieces of the URL, it would look like this: <code>jdbc:postgresql://1:2/3</code> Where: 1 = The IP Address or hostname of the computer/server that Postgres is installed on

On this page ...

- [Connect Ignition to the PostgreSQL Database](#)
- [JDBC Drivers and Translators](#)



INDUCTIVE
UNIVERSITY

Connecting to PostgreSQL

[Watch the Video](#)

	2 = The port that Postgres is running on. The default is 5432, but this could have been changed during installation, or sometime afterwards. When in doubt, ask the person that installed the database.
	3 = The database name that this connection will provide access to.
Username	The username that the Gateway will use to connect to the database.
Password	The password that the Gateway will use to connect to the database.

6. Click **Create New Database Connection** at the bottom of the form.
7. The Database Connections page is displayed showing the **Status** of your connection as **Valid**.
8. To display the details about the status of your database connection, see the **Note** on the above window and click on the **Database Connection Status** link.

JDBC Drivers and Translators

In some cases, you may need to add your own JDBC Driver, or configure a Translator. More information on configuring these can be found on the [JDBC Drivers and Translators](#) page.

Related Topics ...

- [Store and Forward](#)
- [OPC UA](#)
- [Designer](#)

Connecting to SQLite

An Ignition Gateway can create a SQLite database, allowing for data collection without installing a separate SQL database. This is ideal for small scale applications, as well as testing and demonstrations.

Note: SQLite connections offer convenience, but Relational Database Management Systems (RDBMS) generally offer better performance, especially in cases where queries are frequently executed such as systems utilizing a historian system. As a result, the SQLite connection is **not** recommended for production systems as a historian database. Learn more about appropriate uses of SQLite in their [official documentation](#).

Unlike other database connections, SQLite databases only support a single connection at a time. Configuring this type of Database Connection means the Gateway will constantly be connected to the SQLite database, preventing other systems from connecting.

SQLite Connect URL

Ignition Gateways can configure a SQLite database connection without any additional installation. This is because the driver will create a file (which is the SQLite database) at the directory specified in the Connect URL property. Pointing the Connect URL to an existing IDB file will cause the connection to access the existing tables within the file. All Connect URLs must start with the following pattern:

```
jdbc:sqlite:
```

Followed by the destination to the database file.

```
# On Windows  
jdbc:sqlite:C:/Path/To/File.db  
  
# On Mac  
jdbc:sqlite:/path/on/mac/File.bd  
  
# On Linux  
jdbc:sqlite:/path/on/linux/File.db
```

Note: For Windows machines, either forward / or backslashes\ can be used.

Storing to Memory

Instead of creating a file, the Connect URL can specify that data should be stored in memory. As expected, this is not an ideal solution for systems that require any sort of long term storage, but can be useful in cases where data should be ephemeral.

```
jdbc:sqlite::memory
```

Relative Paths

The following feature is new in Ignition version **8.1.10**
[Click here](#) to check out the other new features

SQLite database connections feature several keywords that represent relative locations within Ignition's installation directory.

Keyword	Description
<code>\$(data)</code>	Represents the Gateway's "data" directory at <code>/usr/local/ignition/data/</code> .
<code>\$(local)</code>	Represents the Gateway's "local" directory at <code>/usr/local/</code> .

On this page ...

- [SQLite Connect URL](#)
 - [Storing to Memory](#)
 - [Relative Paths](#)
- [Connect Ignition to SQLite Database](#)

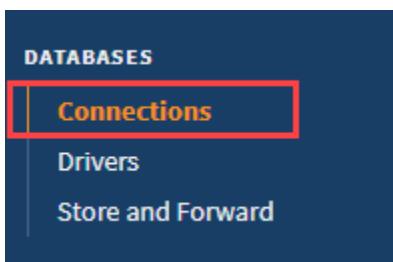
Note: SQLite files created in this directory will not be included in gateway backup files. If you wish for a GWBK to include the SQLite file, then set the path to another location within the data directory, but outside of the local file.

In both cases, additional folders can be added after the keyword. The example below would create "Folder" in the gateway's data directory, and place the DB file in the Folder.

```
jdbc:sqlite:${data}/Folder/File.db
```

Connect Ignition to SQLite Database

1. On the Gateway Webpage, go to the **Config** section.
2. Scroll down to **Databases > Connections**.



3. The Database Connections page is displayed. Click on **Create new Database Connection....**

A screenshot of the Ignition Database Connections page. The left sidebar is visible with the 'Config' tab selected. The main area shows a table of database connections with two entries: 'DB' (MariaDB, MySQL, Faulted) and 'MSSQL' (MySQL, MySQL, Valid). Below the table is a red box around the 'Create new Database Connection...' button. A note at the bottom right says: 'Note: For details about a connection's status, see the [Database Connection Status page](#)'.

4. Select the **SQLite** driver and click **Next**.

The screenshot shows the 'Database Connections' configuration page. It lists several database options: MySQL, Oracle Database, Oracle JDBC, PostgreSQL, and SQLite. The 'SQLite' option is selected, indicated by a red border around its row. Below the list is a blue 'Next >' button.

5. On the **New Database Connection** page, enter a name for your connection. We used **SQLite_Connect**.
6. Next, enter the connect URL for the database. For SQLite this setting must lead to a location on a filesystem. In our example here we used **jdbc:sqlite:C:Program Files/SQLite/File.db**. This path ultimately must lead to a local directory, or a locally mapped drive.

The screenshot shows the 'Main Properties' section of the 'New Database Connection' form. The 'Name' field contains 'SQLite_Connect'. The 'Description' field contains 'My SQLite connection'. The 'JDBC Driver' dropdown is set to 'SQLite'. The 'Connect URL' field contains 'jdbc:sqlite:C:/Program Files/SQLite/File.db', with additional examples provided below it: 'jdbc:sqlite:/path/on/linux/File.db' and 'jdbc:sqlite::memory'.

7. Click on **Create New Database Connection** at the bottom of the form.

Your connection is now created and the Database Connections page is displayed showing the **Status** of your connection as **Valid**.

Config > Database > Database Connections

Successfully created new Database Connection "SQLite_Connect"

Name	Description	JDBC Driver	Translator	Status		
DB		MariaDB	MYSQL	Faulted	delete	edit
MSSQL		MySQL	MYSQL	Valid	delete	edit
SQLServer		Microsoft SQLServer	MSSQL	Reconnecting	delete	edit
SQLite_Connect		SQLite	SQlite	Valid	delete	edit

[Create new Database Connection...](#)

Note: For details about a connection's status, see the [Database Connection Status](#) page.

JDBC Drivers and Translators

In most cases, the default JDBC drivers and Translator settings in Ignition will not need to be modified. However, there are cases where drivers or translators may need modification.

JDBC Drivers and Ignition Upgrades

When upgrading Ignition, JDBC drivers are **not** modified during the upgrade process. This is to prevent database connection issues on upgrade: you should only need to upgrade your JDBC drivers when the database is updated, not when Ignition is upgraded. Only new Ignition installations will always use the JDBC drivers that are included with the installer.



The following feature is new in Ignition version **8.1.33**
[Click here](#) to check out the other new features

The bundled version of Java in Ignition has been upgraded from Java 11 to Java 17.0.8. If upgrading, make sure to check if any JDBC drivers need to be modified for Java 17 compatibility.

On this page ...

- [JDBC Drivers and Ignition Upgrades](#)
 - [Common Third Party Drivers](#)
 - [Bundled JDBC Drivers and Ignition Upgrades](#)
- [Database Translator and JDBC Driver Settings](#)
- [Add a New JDBC Driver](#)
- [Upgrade a JDBC Driver](#)
- [Database Translators](#)

Running the installer to upgrade a preexisting Ignition installation will not modify the installed JDBC drivers, even if newer drivers are included in the installer. In addition, restoring a Gateway backup from an older version will replace any new versions of the drivers with the versions from the backup. In this case, you will need to manually update the JDBC drivers.

Common Third Party Drivers

Commonly, you will have to download the official JAR file from the creator's website. We have a few links here to make it easy to find.

Database	Link	Notes
IBM DB2	https://www.ibm.com/support/pages/db2-jdbc-driver-versions-and-downloads	<ol style="list-style-type: none">1. Select the GA version of your desired driver.2. Select your desired package.3. Log in to your IBM account and complete your package download.
MySQL	https://dev.mysql.com/downloads/connector/j/	<ol style="list-style-type: none">1. Select Your Operating system (or Platform Independent if you are on Windows).2. After the file has been downloaded, unzip the archive. On Windows you can right-click and select the Extract All option.3. The location of the JAR we need should be in the extracted folder under mysql-connector-java-X.X.XX where the Xs are the version number. You are looking for a file that is named like mysql-connector-java-X.X.XX.jar. <p>Compatibility with different JRE versions can be found in the MySQL connector docs.</p>
MSSQL	https://github.com/microsoft/mssql-jdbc/releases	You only need the .jar file, so you can select just the mssql-jdbc-X.X.XX.jreX.jar file. The Xs will be replaced with the version numbers of both the jdbc driver and jre used. Check the SQL Server docs for more details which jdbc drivers contain jre versions compatible with your Ignition server's embedded Java version.
Oracle	https://www.oracle.com/technetwork/database/application-development/jdbc/downloads/index.html	<ol style="list-style-type: none">1. You will need to create an Oracle account to download the JAR files.2. You can select the Unzipped version of the newest JAR.3. You are looking for a file that is named like odbcX.jar. Where X is the version number.

Bundled JDBC Drivers and Ignition Upgrades

Ignition installers come with the latest version of some JDBC drivers. During the installation process, the installer will use these drivers. However, during upgrade, the installer will not replace existing drivers with those in the installer. This is to preserve your existing connections, since newer drivers may not work with older database installations.

The following feature is new in Ignition version **8.1.8**
[Click here](#) to check out the other new features

As of 8.1.8, upgrading Ignition will create a directory at `installDirectory/user-lib/jdbc-bundled` and place updated JDBC drivers in the directory, allowing you to manually update your system's JDBC drivers at a later time without having obtain more recent drivers yourself.

Database Translator and JDBC Driver Settings

In some cases, you may need to add your own JDBC Driver, or configure a Translator. However, you may need to check the JDBC driver's documentation for information on how to configure them. The [JDBC Driver and Database Translator Settings page](#) contains some recommended settings, but the vendor's documentation should always supersede any suggestions here.

Add a New JDBC Driver

To add a new JDBC driver to Ignition, do the following steps:

1. On the Gateway Webpage **Config** section, click on **Databases > Drivers**.



The Database Drivers & Settings page is displayed.

2. Click on the **Create new JDBC Driver...** link at the bottom of the page.

Note: Ignition comes preconfigured with some popular JDBC drivers already. In some cases you may only need to provide a JAR file to an existing driver configuration, instead of creating a new driver configuration. See the [Common Third Party Drivers table](#) for a list of preconfigured drivers that need a JAR file.

A screenshot of the 'Database Drivers & Settings' page. The left sidebar shows the 'Config' section selected. The main area displays a table of installed JDBC drivers. The table columns are Name, Driver Type, Default Translator, and Status. The drivers listed are MariaDB, Microsoft SQLServer, MySQL, Oracle JDBC, PostgreSQL, and SQLite. Each row has 'delete' and 'edit' buttons. Below the table is a red-bordered button labeled 'Create new JDBC Driver...'. A note at the bottom states: 'Note: Please see [this help page](#) for information about installing driver files that are not able to be bundled with Ignition.'

3. In the **Name** field, type the full name of the JDBC driver, see the manufacturer's documentation to get the name.
4. In the **JAR File(s)** field, specify the JAR file that contains the driver, as well as any other required JARs. If you do not have the JAR file needed, see above for download links.
5. Use the default settings for the following properties:

Driver Defaults and Instructions	
Driver Type	Is the brand of database. This is used for optimizations in the Gateway, if in doubt, select GENERIC.
URL Format	Is a default value for the connect URL. This provides a hint to the format of the connect URL that this driver requires while adding a datasource connection. For example, the hint for the format can be, <code>jdbc:dbtype://host:port</code>

	/database.
URL Instructions	Free form instructions that are shown to help the user to create a connection.
Default Connection Properties	Any additional properties to add by default to the connection string.
Connection Properties Instructions	Tips about which connection properties might be useful.
Default Validation Query	The default query that is used to verify that the connection is available.
SQL Language Compatibility	
Default Translator	The database translator that is used by default for connections from this driver.

6. Click the **Create New JDBC Driver** button, located at the very bottom of the page, to create the new driver.

Upgrade a JDBC Driver

In some cases you may need to upgrade a driver. The steps below detail where this would take place

1. You will need to obtain the new driver. These are typically provided by the same organization that made the database. The driver will be a JAR file.
2. Once you have the new driver, head to your Ignition Gateway's **Config** section, click on **Databases > Drivers**.
3. The Database Drivers & Settings page will be displayed. These are the currently configured JDBC drivers on the Gateway, and can be modified from the **Edit** button. Click the Edit button for the driver you need to upgrade.
4. You will need to pass in the new driver to the **JAR File(s)** property.
 - a. Click the **Choose File** button
 - b. Navigate to the driver, and click **Open**.
5. **[Optional]** Update any other properties. In most cases, you may skip this step. However you may need to update some other properties when a new driver is in place. This step depends on the driver, and what it changes. Refer to the driver's documentation to determine if any connection properties need to be changed. For example, users upgrading to MySQL 8.0 from legacy versions will need to change the Default Connection Properties value from:

```
zeroDateTimeBehavior=convertToNull;
```

to:

```
zeroDateTimeBehavior=CONVERT_TO_NULL;useSSL=false;allowPublicKeyRetrieval=true;
```

Again, this step depends on the driver, and in some cases you may be able to skip it.

6. Click the **Save Changes** button at the bottom of the page.

Database Translators

Despite the presence of a SQL standard, many database system vary in how they implement or accomplish various tasks. The JDBC driver system tries to hide these differences as much as possible, but unfortunately some differences persist.

The database translator system in Ignition navigates these differences as they apply to the system. It provides a way to define certain key operations that are commonly different between database vendors, such as creating auto-incrementing index columns, and the keywords used for different data types.

Translator Management

Database translators are managed in the Gateway from the **Databases > Drivers > Translators** tab. Ignition comes pre-configured with translators for the major supported databases, but you can edit and remove them, as well as create new translators. It is necessary to create a new translator only when adding a new JDBC driver for a database that does not share syntax with any of the existing translators.

Creating a New Translator

To add a new database translator to Ignition, do the following steps:

1. In the Gateway Config section, click on **Databases > Drivers**.
The Database Drivers & Settings page is displayed.
2. Go to the Translators tab, find the blue arrow, and click on the **Create new Database Translator...** link.
The New Database Translator page is displayed showing a list of all the translator properties.
3. Define the tokens used with the translator properties on the New Database Translator page.

For most of the properties, you need to define special token markers to indicate places where other values are placed. For example, the default **Create Table Syntax** entry looks as follows:

```
CREATE TABLE {tablename} ({creationdef}{primarykeydef})
```

Where:

`tablename`, `creationdef`, and `primarykeydef` are all tokens that are expanded. `tablename` is replaced directly with the table, `creationdef` is a list of columns, and `primarykeydef` is the phrase created by the **Primary Key Syntax** entry in the translator.

The possible tokens are as follows:

Token	Description
<code>tablename</code>	The name of the table being created.
<code>indexname</code>	The name of the index to create, when adding a column index to the table.
<code>primarykeydef</code>	A clause that defines a primary key for a new table.
<code>creationdef</code>	The list of columns to create in the table.
<code>alterdef</code>	A list of columns to add/remove/modify in the table.
<code>columnname</code>	The name of a column.
<code>type</code>	The data type of a column.
<code>limit</code>	The value of the limit clause.
Other Properties	
<code>Limit Position</code>	Defines where the limit clause should be placed. Back , the limit is placed at the end of the query. Front , places it directly after the SELECT keyword.
<code>Column Quote Character</code>	All columns are created and accessed with the defined quote, which tells the database to use a specific casing, as well as avoiding collisions between the column name and database keywords.
<code>Supports Returning Auto-generated Keys? / Fetch Key Query</code>	Indicates whether the JDBC driver supports the return of generated keys. If the driver does not support this feature, the Fetch Key Query is used to retrieve the last key.
Date Type Mapping	
<code>All data types</code>	The keywords that are used when creating columns of the given types.

4. Click the **Create New Database Translator** button, located at the very bottom of the page, to create the translator.

In This Section ...

JDBC Driver and Database Translator Settings

This page details various JDBC driver and translator settings. See the [JDBC Drivers and Translators page](#) for information on adding and configuring JDBC drivers and translators.

On this page ...

- IBM
 - IBM Default JDBC Driver Settings
 - IBM Default Translator Settings
- MySQL
 - MySQL Default JDBC Driver Settings
 - MySQL Default Translator Settings
- MSSQL
 - MSSQL Default JDBC Driver Settings
 - MSSQL Default Translator Settings
- Oracle Express
 - Oracle Express Default JDBC Driver Settings
 - Oracle Express Default Translator Settings
- PostgreSQL
 - PostgreSQL Default JDBC Driver Settings
 - PostgreSQL Default Translator Settings
- Snowflake
 - Snowflake Default JDBC Driver Settings
 - Snowflake Default Translator Settings

IBM

IBM Default JDBC Driver Settings

Main Properties	
Name	IBM DB2
Description	The official IBM DB2 JDBC Driver.
Classname	com.ibm.db2.jcc.DB2Driver
JAR File(s)	<Click on the Choose File button to select and upload the JAR(s). This is the part where you upload the JDBC driver.>
Driver Defaults & Instructions	
Driver Type	DB2
URL Format	jdbc: db2 ://localhost:50000/SAMPLE OR jdbc: db2 ://localhost:25000/SAMPLE
URL Instructions	 The format of the DB2 connect URL is: <code>db2://host:port/database</code> With the three parameters (in bold) <ul style="list-style-type:none; margin-left:10px;">host: The host name or IP address of the database server.port: The port that the database server is running on. DB2 default port is 50000.database: The name of the logical database that you are connecting to on the DB2 server.
Default Conn	<Leave Blank>

Action Properties	
Connection Property Instructions	<Leave Blank>
Default Validation Query	select 1 from sysibm.sysdummy1
SQL Language Compatibility	
Default Translator	<The Translator this driver should use> Default: IBM DB2

IBM Default Translator Settings

Main Properties	
Name	<Name of the Translator. The JDBC driver will reference the settings below by the name specified here>
Create Table Syntax	CREATE TABLE {tablename} ({creationdef}{primarykeydef})
Create Sequence Syntax	<leave empty>
Create Trigger Syntax	<leave empty>
Create Index Syntax	CREATE INDEX {indexname} ON {tablename}({columnname})
Auto Increment Field Definition	{type} GENERATED ALWAYS AS IDENTITY PRIMARY KEY
Alter Table Syntax	ALTER TABLE {tablename} {alterdef}
Add Column Syntax	ADD COLUMN {columnname} {type}
Primary Key Syntax	<leave empty>
Limit Syntax	FETCH FIRST {limit} ROWS ONLY
Limit Position	Back
Current Timestamp Query	values current timestamp
Column Quote Character	"
Supports Returning Auto-generated Keys?	False
Fetch Key Query	SELECT max({columnname}) FROM {tablename}
Table List Filter	<leave empty>
Data Type Mapping	
Byte (I1)	int
Short (I2)	int
Integer (I4)	int
Long (I8)	bigint
Boolean	int
Datetime	timestamp
Float (R4)	float
Double (R8)	double
String	varchar(255)

Binary	blob
Long Text	clob(65536)

MySQL

MySQL Default JDBC Driver Settings

Main Properties	
Name	MySQL
Description	The official MySQL JDBC Driver, Connector/J.
Classname	com.mysql.cj.jdbc.Driver
JAR File(s)	<Click on the Choose File button to select and upload the JAR(s). This is the part where you upload the JDBC driver.>
Driver Defaults & Instructions	
Driver Type	MySQL
URL Format	jdbc:mysql://localhost:3306/test
URL Instructions	<p>
The format of the MySQL connect URL is:
<code>jdbc:mysql://host:port/database</code>
With the three parameters (in bold) <ul style="list-style-type:none; margin-left:10px;">host: The host name or IP address of the database server.port: The port that the database server is running on. MySQL default port is 3306.database: The name of the logical database that you are connecting to on the MySQL server.</p>
Default Connection Properties	zeroDateTimeBehavior=CONVERT_TO_NULL;connectTimeout=120000;socketTimeout=120000;useSSL=false;allowPublicKeyRetrieval=true;
Connection Properties Instructions	<p>There is an extensive list of extra connection properties available for MySQL Connector/J. See the documentation for a table describing all connection properties.
A default <tt>serverTimezone</tt> value (taken from the Gateway) will be appended to the connection string if one is not specified.</p>
Default Validation Query	SELECT 1
SQL Language Compatibility	
Default Translator	MySQL

MySQL Default Translator Settings

Main Properties	
Name	MySQL
Create Table Syntax	CREATE TABLE {tablename} ({creationdef}{primarykeydef})
Create Sequence Syntax	<leave empty>
Create Trigger Syntax	<leave empty>
Create Index Syntax	CREATE INDEX {indexname} ON {tablename}({columnname})
Auto Increment Field Definition	{type} NOT NULL AUTO_INCREMENT

Alter Table Syntax	ALTER TABLE {tablename} {alterdef}
Add Column Syntax	ADD COLUMN {columnname} {type}
Primary Key Syntax	PRIMARY KEY ({columnname})
Limit Syntax	LIMIT {limit}
Limit Position	Back
Current Timestamp Query	SELECT CURRENT_TIMESTAMP
Column Quote Character	'
Supports Returning Auto-generated Keys?	True
Fetch Key Query	<leave empty>
Table List Filter	<leave empty>

Data Type Mapping

Byte (I1)	int
Short (I2)	int
Integer (I4)	int
Long (I8)	bigint
Boolean	int
Datetime	datetime
Float (R4)	float(10)
Double (R8)	double
String	varchar(255)
Binary	varbinary
Long Text	text

MSSQL

MSSQL Default JDBC Driver Settings

Main Properties	
Name	<Name of the driver, as you would like it to appear on the Gateway>
Description	<Enter a useful description you would like to see next to the driver>
Classname	com.microsoft.sqlserver.jdbc.SQLServerDriver
JAR File(s)	<Click on the Choose File button to select and upload the JAR(s). This is the part where you upload the JDBC driver.>
Driver Defaults & Instructions	
Driver Type	Microsoft SQL Server
URL Format	jdbc:sqlserver://localhost\SQLEXPRESS

UR L Inst ructi ons	 The format of the SQL Server connect URL is: <code>jdbc:sqlserver://host\instanceName[:port]</code> With the three parameters (in bold) <ul style="list-style-type:none; margin-left:10px;">host: The host name or IP address of the database server.instanceName: (optional) the instance to connect to on the host. If not specified, a connection to the default instance is made.port: (optional) the port to connect to. The default is 1433. If you are using the default, you can omit the port and the preceding ':'. For SQL Server, you specify the <i>database name</i> to connect to using the <code>databaseName</code> property in the <i>Extra Connection Properties</i>.
Def ault Con nect ion Pro pert ies	databaseName=test
Con nect ion Pro pert ies Inst ructi ons	Use <i>databaseName=YOUR_DATABASE</i> to specify the database to connect to.
Def ault Vali dati on Que ry	SELECT 1
SQL Language Compatibility	
Def ault Tra nsla tor	The Translator this driver should use. If you're adding a new Driver, then you may

MSSQL Default Translator Settings

Main Properties	
Name	<Name of the Translator. The JDBC driver will reference the settings below by the name specified here>
Create Table Syntax	CREATE TABLE {tablename} ({creationdef}{primarykeydef})
Create Sequence Syntax	<Blank>
Create Trigger Syntax	<Blank>
Create Index Syntax	CREATE INDEX {indexname} ON {tablename}({columnname})
Auto Increment Field Definition	{type} IDENTITY(1,1)
Alter Table Syntax	ALTER TABLE {tablename} ADD {alterdef}
Add Column Syntax	{columnname} {type}
Primary Key Syntax	PRIMARY KEY CLUSTERED ({columnname})
Limit Syntax	TOP {limit}
Limit Position	Front
Current Timestamp Query	SELECT CURRENT_TIMESTAMP
Column Quote Character	"
Supports Returning Auto-generated Keys?	True
Fetch Key Query	<Blank>

Table List Filter	<Blank>
Data Type Mapping	
Byte (I1)	int
Short (I2)	int
Integer (I4)	int
Long (I8)	bigint
Boolean	int
Datetime	datetime
Float (R4)	float(10)
Double (R8)	double precision
String	varchar(255)
Binary	varbinary
Long Text	nvarchar(max)

Oracle Express

Oracle Express Default JDBC Driver Settings

Main Properties	
Name	<Name of the driver, as you would like it to appear on the Gateway>
Description	<Enter a useful description you would like to see next to the driver>
Classname	oracle.jdbc.driver.OracleDriver
JAR File(s)	<Click on the Choose File button to select and upload the JAR(s). This is the part where you upload the JDBC driver.>
Driver Defaults & Instructions	
Driver Type	Oracle
URL Format	jdbc:oracle:thin:@localhost:1521:test
URL Instructions	 The format of the Oracle connect URL is: <code>jdbc:oracle:thin:@host:port:SID</code> With the three parameters (in bold) <ul style="list-style-type:none; margin-left:10px;">host: The host name or IP address of the database server.port: The port that the database server is running on. Oracle's default port is 1521.SID: the system ID that identifies the database to connect to.
Default Connection Properties	<Blank>
Connection Properties Instructions	<Blank>
Default Validation Query	SELECT 1 FROM DUAL
SQL Language Compatibility	

Default Translator	The Translator this driver should use. If you're adding a new Driver, then you may
--------------------	--

Oracle Express Default Translator Settings

Main Properties	
Name	<Name of the Translator. The JDBC driver will reference the settings below by the name specified here>
Create Table Syntax	CREATE TABLE {tablename} ({creationdef}{primarykeydef})
Create Sequence Syntax	CREATE SEQUENCE {tablename}seq START WITH 1 INCREMENT BY 1
Create Trigger Syntax	CREATE TRIGGER {tablename}trig BEFORE INSERT ON {tablename} REFERENCING NEW AS NEW FOR EACH ROW BEGIN select {tablename}seq.nextval INTO :NEW.{columnname} FROM dual; END;
Create Index Syntax	CREATE INDEX {indexname} ON {tablename}({columnname})
Auto Increment Field Definition	{type} NOT NULL
Alter Table Syntax	ALTER TABLE {tablename} ADD ({alterdef})
Add Column Syntax	{columnname} {type}
Primary Key Syntax	PRIMARY KEY ({columnname})
Limit Syntax	rownum<={limit}
Limit Position	Where
Current Timestamp Query	SELECT CURRENT_TIMESTAMP FROM DUAL
Column Quote Character	"
Supports Returning Auto-generated Keys?	False
Fetch Key Query	SELECT {tablename}SEQ.CURRVAL FROM DUAL
Table List Filter	<leave empty>
Data Type Mapping	
Byte (I1)	int
Short (I2)	int
Integer (I4)	int
Long (I8)	int
Boolean	int
Datetime	timestamp
Float (R4)	float
Double (R8)	double precision
String	varchar2(255)
Binary	varbinary
Long Text	nclob

PostgreSQL

PostgreSQL Default JDBC Driver Settings

Main Properties

Name	<Name of the driver, as you would like it to appear on the Gateway>
Description	<Enter a useful description you would like to see next to the driver>
Classname	org.postgresql.Driver
JAR File(s)	<Click on the Choose File button to select and upload the JAR(s). This is the part where you upload the JDBC driver.>
Driver Defaults & Instructions	
Driver Type	PostgreSQL
URL Format	jdbc:postgresql://localhost:5432/test
URL Instructions	 The format of the PostgreSQL connect URL is: <code>jdbc:postgresql://host:port/database</code> With the three parameters (in bold) <ul style="list-style-type:none; margin-left:10px;">host: The host name or IP address of the database server.port: The port that the database server is running on. PostgreSQL default port is 5432.database: The name of the logical database that you are connecting to on the PostgreSQL server.
Default Connection Properties	<Blank>
Connection Properties Instructions	No extra connection parameters are recommended for PostgreSQL. For possible parameter values, see the documentation at http://jdbc.postgresql.org >the PostgreSQL JDBC driver website.
Default Validation Query	SELECT 1
SQL Language Compatibility	
Default Translator	The Translator this driver should use. If you're adding a new Driver, then you may

PostgreSQL Default Translator Settings

Main Properties	
Name	<Name of the Translator. The JDBC driver will reference the settings below by the name specified here>
Create Table Syntax	CREATE TABLE {tablename} ({creationdef}{primarykeydef})
Create Sequence Syntax	<Blank>
Create Trigger Syntax	<Blank>
Create Index Syntax	CREATE INDEX {indexname} ON {tablename}({columnname})
Auto Increment Field Definition	SERIAL NOT NULL
Alter Table Syntax	ALTER TABLE {tablename} {alterdef}
Add Column Syntax	ADD COLUMN {columnname} {type}
Primary Key Syntax	PRIMARY KEY ({columnname})
Limit Syntax	LIMIT {limit}
Limit Position	Back
Current Timestamp Query	SELECT CURRENT_TIMESTAMP

Column Quote Character	"
Supports Returning Auto-generated Keys?	True
Fetch Key Query	<Blank>
Table List Filter	<Blank>
Data Type Mapping	
Byte (I1)	int
Short (I2)	int
Integer (I4)	int
Long (I8)	bigint
Boolean	int
Datetime	timestamp
Float (R4)	float
Double (R8)	double precision
String	varchar(255)
Binary	bytea
Long Text	text

Snowflake

Snowflake Default JDBC Driver Settings

Note: Many of the fields for the Snowflake JDBC driver can be left blank if your custom Snowflake translator is set up properly and referenced in the "Default Translator" field.

Main Properties	
Name	<Name of the driver, as you would like it to appear on the Gateway>
Description	<Enter a useful description you would like to see next to the driver>
Classname	net.snowflake.client.jdbc.SnowflakeDriver
JAR File(s)	<Click on the Choose File button to select and upload the JAR(s). This is the part where you upload the JDBC driver.>
Driver Defaults & Instructions	
Driver Type	Generic
URL Format	<Blank>
URL Instructions	<Blank>
Default Connection Properties	<Blank>
Connection Properties Instructions	<Blank>
Default Validation Query	SELECT 1
SQL Language Compatibility	
Default Translator	<Your custom Snowflake Translator>

Snowflake Default Translator Settings

Note: If you have a database table called "NO_OP_TABLE", you should use a different table name for the "Create Index Syntax" field.

Main Properties	
Name	<Name of the translator, as you would like it to appear on the Gateway>
Create Table Syntax	CREATE TABLE IF NOT EXISTS {tablename} ({creationdef}{primarykeydef})
Create Sequence Syntax	CREATE SEQUENCE IF NOT EXISTS {tablename}seq
Create Trigger Syntax	<Blank>
Create Index Syntax	ALTER TABLE IF EXISTS NO_OP_TABLE SET COMMENT='no-op table due to forced value for TRANSLATORS > ALTER TABLE SYNTAX';
Auto Increment Field Definition	{type} NOT NULL AUTOINCREMENT
Alter Table Syntax	ALTER TABLE {tablename} {alterdef}
Add Column Syntax	ADD COLUMN {columnname} {type}
Primary Key Syntax	PRIMARY KEY ({columnname})
Limit Syntax	LIMIT {limit}
Limit Position	Back
Current Timestamp Query	SELECT CURRENT_TIMESTAMP
Column Quote Character	<Leave Field Blank>
Supports Returning Auto-generated Keys?	False
Fetch Key Query	SELECT NEXT_VALUE FROM INFORMATION_SCHEMA.SEQUENCES WHERE SEQUENCE_NAME='{tablename}'
Table List Filter	<Blank>
Data Type Mapping	
Byte (I1)	int
Short (I2)	int
Integer (I4)	int
Long (I8)	bigint
Boolean	int
Datetime	datetime
Float (R4)	float
Double (R8)	double
String	string
Binary	binary
Long Text	text

Note: In addition to a custom JDBC driver and custom translator, you may need to also configure the Connection Initialization Commands property in your Snowflake [database connection](#) to properly store Tag History. Add the following values to the Connection Initialization Commands property:

```
USE DATABASE ignition_db
USE SCHEMA ignition_db.time_series
```

Store and Forward

The store-and-forward system provides a reliable way for Ignition to store data to the database. In Ignition, systems such as [Tag Historian](#) and [SQL Bridge \(Transaction Groups\)](#) use store-and-forward to ensure that data reaches its destination in the database, and is stored in an efficient manner. The store-and-forward system can be configured in a number of ways, offering both memory buffering for performance and local disk caching for safe storage.

Note: Store-and-forward engines are automatically created for each Database Connection.

Primary Features and Benefits

The store-and-forward system offers a number of benefits over other systems that log directly to the database, such as:

- **Data loss prevention**
Data is removed from the system only when the write to the database has executed successfully.
- **Guaranteed ordering**
Data is forwarded in the same order that it arrived, even if a database connection is not currently available.
- **Enhanced performance**
By first buffering the data in memory, the store-and-forward system can optimize writes, and prevent the originating systems from blocking. This means that the system is less likely to lose data samples in the event of system slow downs.

On this page ...

- [Primary Features and Benefits](#)
- [Store and Forward Data Flow](#)
 - [Understanding the Forward Triggers](#)
 - [Single Connection Policy](#)
 - [Store and Forward for Reliability](#)



INDUCTIVE
UNIVERSITY

Using Store and Forward

[Watch the Video](#)

Store and Forward Data Flow

Although the system offers settings that can affect the pipeline, by default the data flow occurs as follows:

1. Data is generated in some system.
2. Data is placed in a memory buffer.
3. If not removed from memory buffer in some time (the **Write Time**), or if a certain amount of data accumulates (**Write Size**), it is placed in the local cache.
4. The data sink, based on a database connection, pulls data in first from the local store, and then the memory buffer, based on the **Write Time** and **Write Size** settings under **Forward Settings**.
5. If the data fails to forward, either due to an error in the connection or in the data itself, it is returned to the buffer or cache.
6. If the data errors out too many times, it becomes quarantined.
7. Quarantined data can be managed through the Gateway, and can be deleted or un-quarantined, once the error is resolved.

Understanding the Forward Triggers

Data is forwarded from one stage to the next based on the **Write Time** and **Write Size** triggers. These settings work as an **either/or** manner, meaning that if either of them is surpassed, the data is forwarded. One important point to note is that the **Write Size** setting influences the transaction size of similar data to be forwarded, and therefore can have a big impact on performance. As a result, the **Write Time** should normally be used as the controlling factor, with the **Write Size** set to something that will provide reasonable transactions, like 100.

Single Connection Policy

While database connections have a pool of multiple connection (8 by default), the Store and Forward engine only uses one of those connections. The system heavily optimizes queries by grouping multiple queries into a single transaction before sending the data off.

Store and Forward for Reliability

The store-and-forward system settings, while seemingly limited, offer a good deal of flexibility in tuning. Different types of situations and goals will likely require different configurations.

When the safety of the data is a concern, the goal is to get the data stored to disk as quickly as possible in order to minimize risk of loss due to a power outage or system failure. The local cache plays a crucial role in this, allowing the system to store data locally for any amount of time until the remote database can accept it. This protects against network failures and database failures, as well.

By setting the **write size** and **write time** of both the local cache and forwarder to low values, the data spends less time in the memory buffer. While the memory buffer can be set to 0 to bypass it completely, this is not usually recommended, as the buffer is used to create a loose coupling between the history system and other parts of Ignition that report history. This disconnect improves performance and protects against temporary system slowdowns. In fact, it is recommended that for reliable logging, this value be set to a high value, to allow the maximum possible amount of data to enter the system in the case of a storage slowdown.

[In This Section ...](#)

Using Store and Forward

The current status of the Store and Forward system can be viewed from the Status section of the [Gateway Web Interface](#). These pages provide detailed analysis on each Store and Forward engine.

To Monitor the Store and Forward Engine

1. Go to the **Status** section of the Gateway.
2. Click on **Store & Forward** from the left menus.
The **Store & Forward Connections** page is displayed showing each store-and-forward engine along with the current throughput and capacity of its **Memory Buffer** and **Disk Cache**.

On this page, there are several notable items:

Name	Description
Aggregate Throughput	The aggregated number of records inserted into a database from any engine, per second.
Total Quarantined	The current count of quarantined items across all engines.
Total Dropped	The number of records that have been dropped from all store and forward engines. A record is considered dropped if it can not be added to one of the buffers, such as when a buffer is full, and the engine can no longer accept new records.

On this page ...

- [To Monitor the Store and Forward Engine](#)



INDUCTIVE
UNIVERSITY

Using Store and Forward

[Watch the Video](#)

My-Gateway

Ignition!

admin | Sign Out

Help ? Get Designer

SYSTEMS

- Home
- Status
- Config
- Overview
- Performance
- Alarm Pipelines
- Gateway Scripts
- Modules
- Redundancy
- Reports
- SFCs
- Voice Alarming
- Tags
- EAM Tasks
- Transaction Groups

CONNECTIONS

- EAM Agents
- Databases
- Designers
- Devices
- Gateway Network
- Store & Forward**
- OPC Connections
- Perspective Sessions
- Vision Clients

DIAGNOSTICS

- Execution
- Logs
- Running Scripts
- Threads

Status > Connections > **Store and Forward**

Configuration

Aggregate Throughput per second: 0.5

Total Quarantined: 0

Total Dropped: 0

Store and Forward Engines

Filter type to filter View 20

Name	Store Throughput	Forward Throughput	Quarantined	Activity	Actions
DB	0.5/sec	0.5/sec	0	✓ Available	Details
Local Edge Historian	N/A	N/A	0	✓ Available	Details
SQLite	N/A	N/A	0	✓ Available	Details

You can click on **Details** under the **Store and Forward Engines** section to refresh and update the displayed values. This page provides in-depth information on the current status of the engine.

Name	Description
Memory Buffer	The number of records entering the Memory Buffer per second. The progress bar shows the percent of the buffer being utilized, along with the current and max number of records.
Disk Buffer	The number of records entering the Disk Buffer per second. Note that a state of "idle" means the engine is able to successfully store all records into the database before the Write Size or Write Time values have been reached. The progress bar shows the percent of the buffer being utilized, along with the current and max number of records.
Database	Shows the number of records pushed from either buffer to the database per second.
Quarantined Items	Lists all quarantined items in the engine. Includes the number of occurrences, a description of the where the items originated from, and the reason why the record was placed into the quarantine. Provides an opportunity to retry, export, or delete the items.

My-Gateway

admin | Sign Out

Ignition!

Home

Status

Config

SYSTEMS

- Overview
- Performance
- Alarm Pipelines
- Gateway Scripts
- Modules
- Redundancy
- Reports
- SFCs
- Voice Alarming
- Tags
- EAM Tasks
- Transaction Groups

CONNECTIONS

- EAM Agents
- Databases
- Designers
- Devices
- Gateway Network
- Store & Forward
- OPC Connections
- Perspective Sessions
- Vision Clients

Help ?

Get Designer

Configuration

Store Details

Memory Buffer /sec 0 0% (0/250)

Local Cache Idle 0% (0/25000)

Database Storage /sec 0

Quarantined Items

Import quarantine file < Retry All

Delete All

The screenshot shows the Ignition My-Gateway web interface. The left sidebar has a dark blue background with white text. It includes sections for SYSTEMS (Overview, Performance, Alarm Pipelines, etc.) and CONNECTIONS (EAM Agents, Databases, Designers, Devices, Gateway Network, Store & Forward, OPC Connections, Perspective Sessions, Vision Clients). The 'Store & Forward' option is highlighted with a red box. The main panel has a light gray background. At the top, it shows the current status path: Status > Connections > Store and Forward. On the right, there are three boxes labeled 'Memory Buffer', 'Local Cache', and 'Database Storage'. Each box contains a large number '0' followed by '/sec' and a progress bar indicating '0% (0/250)' for Memory Buffer and '0% (0/25000)' for Local Cache. Below these is a section titled 'Quarantined Items' with buttons for 'Import quarantine file', 'Retry All', and 'Delete All'.

Related Topics ...

- Configuring Store and Forward

Configuring Store and Forward

Store and Forward for High-speed Buffering

When configuring the store-and-forward system for high-speed buffering, you are expecting the case that data will come in quick bursts. By buffering the data, the system can accommodate more information than would be possible going directly against the database.

The key points in configuring a buffering system is to avoid expensive operations like storing and reading from the local cache, and to set the memory buffer large enough to accommodate the expected burst sizes.

Each database connection has its own store and forward settings. Store-and-forward engines are directly correlated to database connections, and are automatically managed so that each connection has an engine defined.



You can create multiple database connections pointing to the same database if you wish to configure multiple store-and-forward engines for different purposes.

Configuring Store and Forward

To configure the Store and Forward engine for your database, do the following steps:

1. Go to the Gateway **Config** section and select **Databases > Store and Forward**.
The **Store and Forward** page is displayed and you will see a store-and-forward setting for each of your database connections and you can edit these settings.
2. On the **Store and Forward** page, look for **edit** at the far right of the table and click on it to see all the store-and-forward settings.
The settings of a store-and-forward engine define how and when data is moved through the system. You must understand these settings so that you can carefully set them according to your goals.

On this page ...

- [Store and Forward for High-speed Buffering](#)
- [Configuring Store and Forward](#)
- [Store and Forward Settings](#)



INDUCTIVE
UNIVERSITY

Configuring Store and Forward

[Watch the Video](#)

The screenshot shows the Ignition software interface with the following details:

- Header:** Ignition-47220cc42c3d, admin, Log Out, Get Designer.
- Left Sidebar:** Home, Status, Config (selected), SYSTEM: Overview, Backup/Restore, Ignition Exchange, Licensing, Modules, Projects, Redundancy, Gateway Settings; NETWORKING: Web Server, Gateway Network, Email Settings.
- Top Bar:** Help, Activate Ignition.
- Breadcrumbs:** Config > Database > Store and Forward > Overview.
- Text:** Trial Mode 0:09:55.
- Table:** Displays a single row of data:

Name	Memory Buffer Size	Disk Cache Size	Disk Cache Enabled	Actions
Sample_SQLite_Database	250	25000	true	More Edit
- Note:** Store and Forward engines are automatically created for configured database connections and remote history providers. For details about a store's status, see the [Store and Forward Status](#) page.

Store and Forward Settings

Buffer Settings

Memory Buffer Size	The number of records that can be stored in the memory buffer, the first stage of the store-and-forward chain. Other settings define when the data will move from the memory buffer forward, this setting only determines the maximum size. If the max size is reached, additional data will error out and be discarded. The memory buffer cannot quarantine data, so if there are errors and the disk cache is not enabled, the data will be lost. If set to 0, the memory buffer will always be considered full, dropping records.
--------------------	---

Store Settings

These settings apply to the local disk storage cache.

Disk Cache Enabled	Turns on the hard-disk cache. Data is stored here if it cannot be forwarded in a timely manner. The cache also stores quarantined data (that is, data with errors).
Max Records	The maximum size of the cache. After the max is reached, data is backed up into the memory buffer, and once that is full, it is dropped. A 'record' is an insert or update statement. These statements may be batches, thus it is possible for a single 'record' to impact multiple rows. The default maximum is 25,000 records. Note: The disk cache is intended to be a short term storage system. It is not recommended to increase the Max Records beyond 50,000.
Write Size	The number of records that should be accumulated in the memory store before written to the cache. Writing data in blocks can increase performance, but too large of a size increases the risk of data being lost in the event of a power outage or system failure.
Write Time	The max age of records in the memory buffer before they are stored to the cache. This setting is used in combination with the write size in order to give the forwarder the opportunity to retrieve data directly from the memory store and avoid the write to disk entirely.

Forward Settings

These settings govern when data is forwarded to the database. The data is pulled first from the local cache, and then from the memory store. When no data is present in the cache, it is pulled directly from the memory store.

Write Size	Same as disk cache setting above.
Write Time	Same as disk cache setting above.
Enable Schedule	If enable schedule is selected, the forward engine will only be enabled during the times specified by the pattern. The pattern can specify specific times and ranges using a simple syntax.
Schedule Pattern	The schedule is specified as a comma separated list of times or time ranges. You can use the following formats: 24-hour times, that is 8:00-15:00 (for 8am through 3pm) or 21:00-24:00 (9pm through midnight). 12-hour with am/pm (if not specified, 12 is considered noon): 8am-3pm or 9pm-12am Note: When the time period is over, any queued data will remain cached until the next execution period. That is, the forward engine does not run until all data is forwarded.

Once you made the changes you want, click **Save Changes** at the bottom of the page. This will take you back to the Store and Forward page.

Related Topics ...

- [Controlling Quarantine Data](#)

Controlling Quarantine Data

Quarantined data is data that has erred-out multiple times during attempts to forward it or data that could not be stored because of some configuration issues. It is removed from the forward queue to allow other data to pass. The most common reason for data quarantining is an invalid schema in the database for the data that is being stored. Quarantined data is held indefinitely until the issue is resolved, then you can either delete it or re-insert it into the queue.

Handle the Quarantined Data

1. From **Status** section of the Gateway, go to **Connections > Store and Forward**.
2. Click the **Details** tab next to a Store and Forward engine that you would like to see the quarantine data for. Here you will see any quarantined data, including the number of occurrences, a description of the where the items originated from, and the reason why the record was placed into the quarantine. Each set of data has the option to retry it, delete it, or export it for later use. If there are a lot of quarantined records, it may be a good idea to export and delete them so that the store and forward engine won't fill up and drop records.
3. Fix the problem/error you found that caused the quarantine data.
4. Click on **retry**, or if the data had been exported, import the data using the import tool and then **retry**. This way, you can ensure no data gets lost.

On this page ...

- [Handle the Quarantined Data](#)
- [Disk Cache Management](#)
 - [Archive Disk Cache](#)
 - [Load Disk Cache](#)



INDUCTIVE
UNIVERSIT

Controlling Quarantine Data

[Watch the Video](#)

The screenshot shows the Ignition software interface under the 'Status > Connections > Store and Forward' path. In the top right corner, there is a 'Configuration' button.

Store Details

Memory Buffer	Local Cache	Database Storage
6 /sec	1 /sec	6 /sec
1% (2/250)	0% (0/25000)	

Quarantined Items

Import quarantine file:

- Choose File: No file chosen
- Submit

Actions: Retry All, Delete All

Filter: type to filter, View: 20

ID	Count	Description	Reason	Actions
1	10	SQLTag History Data	Duplicate entry '252-1548372275630' for key 'PRIMARY'	Retry, Delete, Export

Disk Cache Management

The following feature is new in Ignition version **8.1.21**
[Click here](#) to check out the other new features

In the Gateway's Config > Databases > Store and Forward section, there are two options regarding a database connection's disk cache. Note that the following options will only appear if the Disk Cache Enabled property for the specified database connection is set to true:

- Archive Disk Cache
- Load Disk Cache

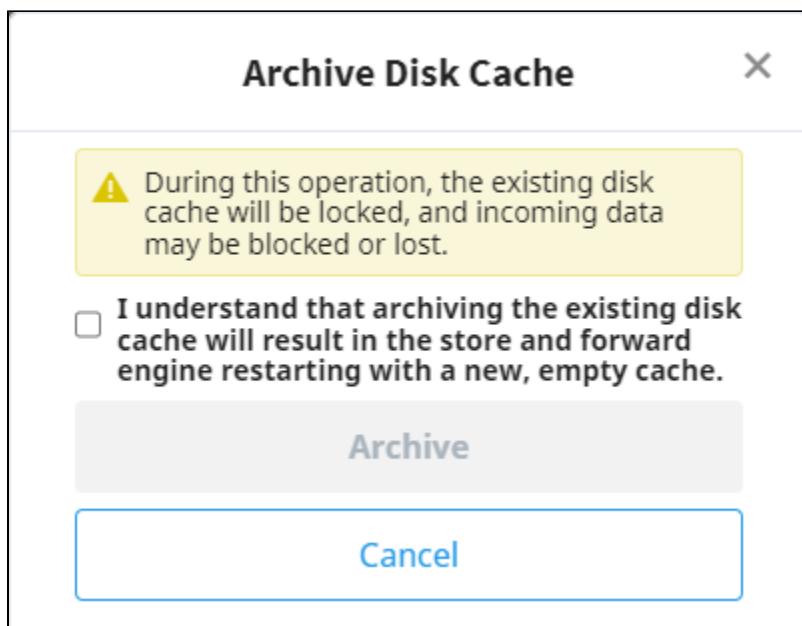
Archive Disk Cache

The screenshot shows the Ignition web interface with the URL 'Ignition-47220cc42c3d'. The left sidebar has 'Config' selected. The main content area is titled 'Store and Forward > Overview' under 'Config > Database'. It shows a table with one row:

Name	Memory Buffer Size	Disk Cache Size	Disk Cache Enabled	Actions
Sample_SQLite_Database	25000			Archive Disk Cache (highlighted with a red box)

A note at the bottom states: 'Note: Store and Forward engines are automatically created for configured database connections and remote history providers. For details about a store's status, see the [Store and Forward Status](#) page.'

Attempting to archive a disk cache will result in Ignition throwing a warning, saying that the existing cache will be locked and incoming data may be blocked or lost.



During the archiving process:

1. The disk cache is shut down.
2. Archived data is moved to a folder located at `%IgnitionInstallationDirectory%/data/datacache/archives`, making it available to load to replace an existing disk cache. The cache folder will use the following naming convention: `{database name}_{timestamp}`.
3. The disk cache is started again, and a new cache folder is created.

Load Disk Cache

The screenshot shows the Ignition web interface with the URL 'Ignition-47220cc42c3d'. The top navigation bar includes 'admin | Log Out', 'Help ?' (with a question mark icon), and a blue button 'Get Designer'. The left sidebar has sections for 'SYSTEM' (Overview, Backup/Restore, Ignition Exchange, Licensing, Modules, Projects, Redundancy, Gateway Settings) and 'NETWORKING' (Web Server, Gateway Network, Email Settings). A search bar at the bottom of the sidebar contains the placeholder 'Search...'. The main content area is titled 'Config > Database > Store and Forward > Overview'. It displays 'Trial Mode 1:59:38' and a green button 'Activate Ignition'. Below this is a table with the following data:

Name	Memory Buffer Size	Disk Cache Size	Disk Cache Enabled	Actions
Sample_SQLite_Database	25000			Archive Disk Cache More Edit

A note at the bottom states: 'Note: Store and Forward engines are automatically created for configured database connections and remote history providers. For details about a store's status, see the [Store and Forward Status](#) page.'

Attempting to load a disk cache will result in Ignition throwing a warning, saying that the existing disk cache will be overwritten.

The modal dialog has a title 'Load Disk Cache' and a close button 'X'. Inside, there is a yellow warning box containing the text: '⚠ During this operation, the existing disk cache will be overwritten with the selected, archived cache.' Below this is a dropdown menu labeled 'Select an Option'. Underneath is a checkbox with the text: 'I understand that loading the selected disk cache will result in the store and forward engine restarting with the archived cache.' At the bottom are two buttons: a grey 'Load' button and a blue 'Cancel' button.

During the loading process:

1. The local disk cache is stopped and deleted.
2. The chosen archived disk cache is renamed.
3. The disk cache is started again.

Related Topics ...

- [Store and Forward](#)
- [Connections - Store & Forward](#)

Security

Security options in Ignition provide many ways to safeguard access to your data and applications. You control not only who accesses your systems, but when and where they can access them. Ignition offers two authentications strategies: [Classic Authentication Strategy](#) or [Identity Provider Authentication Strategy](#).

Gateway Security

Security in Ignition falls into a few categories, tying into the various scopes (Designer, Gateway, Vision Clients and Perspective Sessions). In the Gateway scope, the bulk of security setup happens under the **Config** section of the Gateway Webpage, under the Security header, you'll find pages for authentication, role mappings, and zones.



On this page ...

- [Gateway Security](#)
- [Authentication Strategies](#)
 - [Classic Authentication Strategy \(Designer and Vision Only\)](#)
 - [Identity Provider Authentication Strategy](#)

The primary purpose of Gateway security is to protect access to the two most critical areas of Ignition: the Designer and the Gateway. Many important resources are configured in these areas, so access to each Gateway section (Status and Config), as well as the Designer, can be limited by Security Level.

Authentication Strategies

In regard to authentication and permissions, there are two approaches.

Classic Authentication Strategy (Designer and Vision Only)

[Classic Authentication Strategy](#) involves a concept known as a User Source, which is a configuration that contains multiple roles and users. Users are assigned roles, and security restrictions within a project can be used to check if a user has one or more roles. User Sources can be "internal", meaning all users and roles are contained within an Ignition Gateway, or externally stored in an SQL database. Furthermore, User Sources offer integration with [Active Directory](#).

Identity Provider Authentication Strategy

Ignition can also integrate with Federated [Identity Providers](#) (IdP), allowing users to authenticate against a trusted third party. The Identity Provider Authentication Strategy works by assigning [Security Level](#) restrictions to various features within Ignition, and utilizing [User Attribute Mapping](#) and [Security Level Rules](#) to assign Security Levels to users.

Ignition can integrate with both OpenID Connect and Security Assertion Markup Language (SAML) providers. In addition, Ignition can act as an Identity Provider for isolated systems.

Gateway General Security Settings

The following feature is new in Ignition version **8.1**
[Click here](#) to check out the other new features

The Gateway General Security Settings page is new for release 8.1. This page determines security permissions for the Gateway and Designer.

Config > Security > General

General Gateway Security Settings * Required Field

System Identity Provider *	default	This Identity Provider controls access to the Gateway's web configuration interface and the Designer when the Designer Authentication Strategy is set to Identity Provider. <input checked="" type="checkbox"/> Always ask the IdP to re-authenticate users by default When enabled, Ignition will always ask the IdP to re-authenticate the user by default. This effectively disables Single Sign-On.
Designer Authentication Strategy *	Classic	Controls how Designer users are authenticated. The Classic strategy requires the user to enter their username and password in an embedded login form in the Designer. Classic authentication is performed against the System User Source. The Identity Provider strategy redirects the user to their IdP in their web browser in order to authenticate. The System Identity Provider setting controls which Identity Provider the user is redirected to.
System User Source *	default	This user source controls access to the Designer when the Designer Authentication Strategy is set to Classic.
Designer Role(s) *	Administrator	Users must belong to one of these roles in order to log into the Designer. Multiple roles can be specified by separating them with commas. Example: Administrator, Operator
Create Project Role(s)		Users must belong to one of these roles in order to create a new Designer project. Multiple roles

On this page ...

- [Gateway Security Settings Table](#)



INDUCTIVE
UNIVERSITY

Restricting Gateway Access

[Watch the Video](#)

Gateway Security Settings Table

Setting	Description
System Identity Provider	<p>Dropdown list to select the Identity Provider that controls access to the Gateway's web configuration interface and the Designer (only when the Designer Authentication Strategy is set to Identity Provider).</p> <p>Additional option to always ask the IdP to re-authenticate users by default. When enabled, Ignition will always ask the IdP to re-authenticate the user by default. This effectively disables Single Sign-On.</p> <p>This field is required.</p>
Designer Authentication Strategy	<p>Controls how the Designer authenticates users. Options are Classic or Identity Provider.</p> <ul style="list-style-type: none">Classic: The Classic strategy requires the user to enter their username and password in an embedded login form in the Designer. Classic authentication is performed against the System User Source.Identity Provider: The Identity Provider strategy redirects the user to their IdP in their web browser in order to authenticate. The System Identity Provider setting controls which Identity Provider the user is redirected to. Required. <p>This field is required. Additional options on this screen will change depending on the Designer Authentication Strategy that is selected here.</p>
Designer Auth Token Inactivity Timeout	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>(<i>Identity Provider strategy only</i>) The number of minutes which must elapse before expiring a designer user's auth token due to inactivity caused by a disconnected session. Must be greater than zero. Default value is 10.</p>
Designer Auth	

Token Time-To-Live	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>(Identity Provider strategy only) The maximum number of minutes a designer user's auth token may exist before it expires. If set to any number less than or equal to zero, auth tokens may live forever, as long as the auth token has not expired due to inactivity.</p>
Designer Permissions	<p>(Identity Provider strategy only) Select one of the following options:</p> <ul style="list-style-type: none"> • Users must belong to all of these security levels in order to login to the Designer. • Users must belong to at least one of these security levels in order to login to the Designer. <p>Caution: Empty value in this field means "Public" security level: Access will be unrestricted.</p>
Create Project Permissions	<p>(Identity Provider strategy only) Enter the security levels required to create a new project, for example, Authenticated/Roles /Administrator, SecurityZones/localhost.. Then select one of the following options:</p> <ul style="list-style-type: none"> • Users must belong to all of these security levels in order to create a new Designer project. • Users must belong to at least one of these security levels in order to create a new Designer project. <p>Caution: Empty value in this field means "Public" security level: Access will be unrestricted.</p>
System User Source	<p>(Classic authentication strategy only) This user source controls access to the Designer. This field is required.</p>
Designer Role(s)	<p>(Classic authentication strategy only) Enter the roles required for access to the Designer. Users must belong to at least one of these roles in order to log into the Designer. Multiple roles can be specified by separating them with commas, for example: Administrator, Operator.</p>
Create Project Role(s)	<p>(Classic authentication strategy only) Enter the roles required for create a new Designer project. Users must belong to at least one of these roles in order to create a new Designer project. Multiple roles can be specified by separating them with commas, for example: Administrator, Operator.</p>
Gateway Config Permissions	<p>Enter the security levels required for access to the Gateway Config section. Then select one of the following options:</p> <ul style="list-style-type: none"> • Users must belong to all of these security levels in order to login to the configuration section. • Users must belong to at least one of these security levels in order to login to the configuration section. <p>Multiple security level paths can be specified by separating them with commas. For example, Authenticated/Roles/Administrator, SecurityZones/localhost</p> <p>Caution: Empty value in this field means "Public" security level: Access will be unrestricted.</p>
Status Page Permissions	<p>Enter the security levels required for access to the Gateway Status section. Then select one of the following options:</p> <ul style="list-style-type: none"> • Users must belong to all of these security levels in order to login to the configuration section. • Users must belong to at least one of these security levels in order to login to the configuration section. Multiple security level paths can be specified by separating them with commas, for example: Authenticated/Roles/Administrator, SecurityZones /localhost. <p>Caution: Empty value in this field means "Public" security level: Access will be unrestricted.</p>
Home Page Permissions	<p>Sets the security levels required to access the Gateway Home section. Then select one of the following options:</p> <ul style="list-style-type: none"> • Users must belong to all of these security levels in order to login to the home section. • Users must belong to at least one of these security levels in order to login to the home section. Multiple security level paths can be specified by separating them with commas, for example: Authenticated/Roles/Administrator, SecurityZones/localhost.

	<p>Caution: Empty value in this field means "Public" security level: Access will be unrestricted.</p>
User Inactivity Timeout	<p>The following feature is new in Ignition version 8.1.1 Click here to check out the other new features</p> <p>The number of minutes which must elapse before expiring a user's gateway web interface session to inactivity. Sessions will not timeout if set to any number less than or equal to zero.</p>
Allow User Admin	Allows the administration of the gateway's system user source from the Designer and client. Unless this is enabled, the Vision module's 'User Management Component' will be prevented from altering the gateway's system user source and scripts will not be able to alter users or roles. (Default is false.)
Allow Designer SSO	Allows single-sign-on authentication for logging into the Designer if the System User Source supports it. The Designer SSO capability is only available when the Designer Authentication Strategy is set to Classic. (Default is false.)
Gateway Audit Profile	Dropdown list to select the The name of the audit profile that Gateway-scoped actions will log to.

Related Topics ...

- [Users, Roles](#)
- [Service Security](#)

Classic Authentication Strategy

The Classic Authentication strategy authenticates users against a User Source. Both the Designer and Vision Client can authenticate users with this strategy.

Note: If you have Ignition 8.1 with the Perspective module, authentication is handled instead by the [Identity Provider Authentication Strategy](#).

User Sources

User sources are a collection of users, roles, and other user data, such as contact information or schedule. When a [new user or role is created](#), it is applied and stored in the user source. Projects and the Gateway are assigned a User Source to authenticate against. This determines which users have access to which project(s).

There are several types of user sources: single-storage types with varying storage mediums, "hybrids" that combine features of the other types, and a cache type used in Local Client Fallback systems.

Single-Storage

Users and roles are stored in a single location. The single-storage users sources are:

- [Internal Authentication](#) - Users and roles are stored internally to Ignition.
- [Database Authentication](#) - Users and roles are stored in a SQL database. Managing users is done via direct interaction with the database.
- [Active Directory Authentication](#) - Users and roles are managed by Active Directory. Users are authenticated through the LDAP protocol.

Hybrid

Users in hybrid user sources authenticate against Active Directory, meaning that user names and passwords are checked against those stored in Active Directory. However, roles are stored either internally in Ignition or in a SQL database, so it is possible to make a role change without having to contact your Active Directory administrator. This way, Active Directory can be consulted to see if a user is valid, but the management of roles does not require coordination with the IT department, who typically control the Active Directory system. This "best of both worlds" approach is popular for many users of Active Directory.

- [Active Directory-Internal Hybrid](#) - Users managed by Active Directory and roles stored to Ignition internally.
- [Active Directory-Database Hybrid](#) - Users managed by Active Directory and roles stored in an SQL database.

Fallback Cache

This User Source was developed specifically for a system that is using Local Client Fallback, and allows you to cache the login credentials from a remote user source. This means your users can still log in with their normal username/password on a Local Client Fallback project, even when the network connection is unavailable.

More information can be found on the [Fallback Cache Authentication](#) page.

Creating a User Source

All of the described User Source types can be created by navigating to the **Config > Security > Users, Roles** section of the Gateway. To begin the process, select **Create new User Source...** and choose the User Source type you need. Once you've finished entering the User Source properties, click **Create New User Source**.

On this page ...

- [User Sources](#)
 - [Single-Storage](#)
 - [Hybrid](#)
 - [Fallback Cache](#)
 - [Creating a User Source](#)
- [Shared Functionality](#)
- [Main Properties](#)
- [The Default User Source](#)
 - [Editing a User](#)

The screenshot shows the Ignition software interface with the following details:

- Left Sidebar (Config Tab):**
 - SYSTEM:** Overview, Backup/Restore, Ignition Exchange, Licensing, Modules, Projects, Redundancy, Gateway Settings.
 - NETWORKING:** Web Server, Email Settings, Gateway Network.
 - SECURITY:** General, Auditing, **Users, Roles** (highlighted in orange), Service Security.
- Top Navigation:** Config > Security > User Sources
- User Sources Table:**

Name	Type	Description
default	Internal	This user source profile was automatically created during a clean startup
- Buttons:** More ▾, edit, Create new User Source..., Verify a User Source...

Note: The User Source page also includes a link to **Verify a User Source...** See the [Verify a User on a User Source](#) page for information on how to verify an existing user's profile.

Shared Functionality

Regardless of type, all User Sources have the following functionality:

- Failover Source:** If the User Source is unavailable for authentication, then a backup User Source can be specified. The type of the fail-over User Source can differ from the primary, so configurations where an internal-type fails over to a database-type are possible.
- Schedule Restrictions:** The User Source can prevent users from logging in when they are off schedule, meaning that the schedule assigned to the user determines when the user may login.

Main Properties

All User Sources have a section of **Main** properties that are listed before their unique type properties. Below are the descriptions of the Main properties. See the individual User Source pages for the remaining property descriptions.

Name	Description
Name	The name of the User Source. This is how other systems in Ignition reference the user source. Note that every User Source must have a unique name.
Description	An optional description of the user source. Useful for noting which database connection or AD server the User Source may be referencing.
Enabled	<p>The following feature is new in Ignition version 8.1.27 Click here to check out the other new features</p> <p>Disabling a User Source profile will prevent it from being used completely, including authentication and user/role management. User Sources are enabled by default.</p>
Schedule Restricted	Forces schedule restrictions on users. If true, users are only able to log in when their assigned schedule is active. This means a login will fail for users attempting to log into a client while they are off schedule.
Failover Source	Allows authentication attempts against this User Source to failover to another User Source in the event of a network outage, or some other connection issue. Useful with database or Active Directory user sources, as connection failures to the database/AD server will prevent users from logging in.
	This property is initially set to None , meaning a failover User Source is not configured.
Failover Mode	When a Failover Source is configured, this property determines when the failover User Source should be consulted. The following options are available:

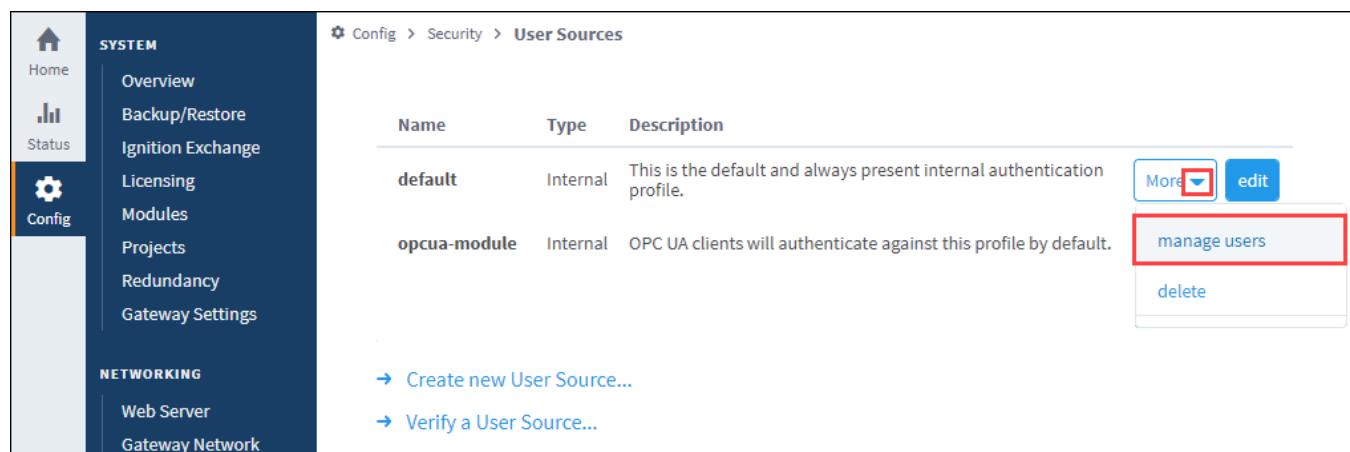
	<p>Hard: The Failover User Source is only consulted when this User Source is unreachable.</p> <p>Soft: The Failover User Source will be consulted if the user's credentials fail authentication, meaning that the user typed in credentials that are unrecognized or incorrect.</p>
Cache Validation Timeout	The amount of time between cache updates of the User Source. If you set this value to -1, the cache validation timeout is turned off.
Lockout Enabled	Lock out a user's account after more than the maximum allowed number of failed authentication attempts occur within the lockout window. Default is true. Note that access can be restored to all locked out users by editing the user source, and clicking the Save Changes button.
Lockout Attempts	Maximum number of failed authentication attempts allowed within the lockout window before locking the user out. Default is 5. If this value is set to something less than zero (for example, -1), then the lockout functionality will be entirely disabled, regardless of what the Lockout Enabled property is set to.
Lockout Window	The duration of the lockout window in minutes. Default is 15. Setting this property to a value of less than zero (for example, -1) will disable the lockout functionality entirely, regardless of what the Lockout Enabled property is set to.

Details on the Password Policy Properties can be found on the [Internal Authentication](#) page.

The Default User Source

When Ignition is installed for the first time, an internal User Source named 'default' is created. You can manage the default User Source by navigating to the **Config > Security > Users, Roles** section of the Gateway.

The **manage users** option under the 'default' user source **More** dropdown allows you to add new users, modify roles and passwords for existing users, remove users, and add/remove roles from the user source.



Name	Type	Description
default	Internal	This is the default and always present internal authentication profile.
opcua-module	Internal	OPC UA clients will authenticate against this profile by default.

More edit
manage users
delete

→ Create new User Source...
→ Verify a User Source...

When you open the 'default' user source for the first time, you will see the first user that was created at installation. This is the administrator account that has full privileges. If this user source has been modified before, a list of existing users is displayed.

Editing a User

Clicking **Edit** on the **User Source > More > manage users** page will access the corresponding User Properties page, which allows you to make any necessary changes to that user.

Config > Security > User Sources

Users Roles

✓ Changes to user "admin" saved.

Username	Name	Roles	Contact Info	Schedule	
admin	Administrator		email: admin@mycompany.com, phone: 555-555-5555	Always	Edit Delete
Bob	Driver		email: bob@mycompany.com	Always	Edit Delete

→ Add User

Fill out the fields for that user then click **Save Changes**.

User Properties

Name	Description
Username	The name of the user.
Change Password?	Check this box to change the existing password.
Password	New password. The following feature is new in Ignition version 8.1.26 Click here to check out the other new features The password policies set for the internal User Source this user belongs to will be listed under the password field for better requirement visibility when creating user passwords. If any password requirement changes are desired, they must be made to the User Source.
Password	Re-type password for verification.
First Name	First name of the user.
Last Name	Last name of the user.
Roles	Role(s) assigned to this user. Check the box next to each role you want this user to have.
Schedule	Schedule for the user. Choose from a dropdown list of schedules that are already defined.
Language	Language to be used for the user. Choose from a dropdown list of languages that are already defined.
Notes	Any notes for this user.
Badge	A string that represents the value set for the user's badge.

Contact Info	
Type	Choose email or SMS.
Value	The email value or SMS number.

Config > Security > User Sources

Users Roles

User Properties

Username	<input type="text"/>
Password	<input type="password"/> <ul style="list-style-type: none"> • 8 characters minimum. • Include at least 2 of the following: Lowercase, uppercase, digits, punctuation.
Password	<input type="password"/> Re-type password for verification.
First Name	<input type="text"/>
Last Name	<input type="text"/>

In This Section ...

Managing Users and Roles

Users and Roles

Security is based on the roles that are assigned to specific users. Roles do not have any structure or hierarchy by default, but it can be created. You can create a hierarchy based on users with a greater role being assigned all matching lesser roles.

There isn't a built-in restriction to the number of roles a user can have, so each user can be assigned many roles or none at all.

It's important to think about the different roles in your project and how they affect the security of your project. For instance, what level of access a particular area of a project needs may determine the functional type roles that you create, and the different users assigned to each role.



When using role-based security in a project, the project stores the name of the role as a string. This means that if you were to modify the name of the role in the Gateway, the role-based security in your project will not update to reflect the new name, and instead will try searching for a role with the original name. Be very careful when modifying the names of roles.

You can manage users and roles using either the Gateway interface, or using the [User Management component](#) inside the Designer or Client. This section shows how to manage users and roles using the Gateway interface under **Config > Security > Users, Roles**. Click on the **manage users** link for the **User Source** you want to administer.

Manage Users

User Sources support managing the users and roles from within Ignition to varying degrees. Some User Sources are fully manageable, meaning that you can administer the users, roles, contact info, and so on from within the Ignition Gateway, as well as inside a Vision Client. Other User Sources do not support this at all or only partially support it. Make sure you understand how and where the administration takes place before you choose a User Source type.

Often, it is desirable to let some management or administrative users of a Vision project manage other users without having to log into the Gateway's Configure section. To do this, you can simply use the built-in [User Management Component](#) that comes with the Vision Module.

Create a Role

When a project is first created, the **Administrator** role is the only role available. Additional roles can be added from the User Sources page on the Gateway.

1. On the Gateway Webpage, go to the **Config** section.
2. Navigate to **Security > Users, Roles**.
3. Click on the **manage users** link for the **User Source** you want to manage.

Name	Type	Description	Actions
default	Internal	This is the default and always present internal authentication profile.	manage users edit
opcua-module	Internal	OPC UA clients will authenticate using this profile by default.	manage users More edit

→ [Create new User Source...](#)
→ [Verify an User Source...](#)

4. Click the **Roles** tab. Look for the blue arrow at the bottom, and click the **Add Role** link.
5. Name the role by entering it in the **Role Name** field, and click on the **Add Role** button. The role is now available to be associated with specific users.

On this page ...

- [Users and Roles](#)
- [Manage Users](#)
- [Create a Role](#)
 - [Assigning Roles](#)
 - [Role Hierarchy](#)
- [User Management Component](#)

Role Properties

Role Name	Maintenance
-----------	-------------

< Cancel Add Role

Assigning Roles

Although it is not required for a user to have a role, be aware users without role assignments might not have access to an area of the project that requires a role.

Existing Users

1. On the Gateway Webpage, go to the **Config** tab, and navigate to **Security > Users, Roles** to access the User Sources page.
2. Click **more > manage users** for the **User Source** you want to manage.
3. Click the **Edit** link for the User you want to edit. The User Settings page is displayed.

Username	Name	Roles	Contact Info	Schedule
admin	Administrator	Administrator, Operator		Always Edit Delete
guest				Always Edit Delete

4. Select the role(s) that you want this user to have from the supplied Roles list.

Roles	<input checked="" type="checkbox"/> Administrator <input type="checkbox"/> guests	<input type="checkbox"/> Maintenance <input checked="" type="checkbox"/> Operator	<input type="checkbox"/> Prod_Sups <input type="checkbox"/> Test_Group
-------	--	--	---

5. Click **Save Changes**. The user now has the privileges associated with the selected role(s).

New Users

1. On the Gateway Webpage, go to the **Config** tab, and navigate to **Security > Users, Roles** to access the User Sources page.
2. Click on **more > manage users** for the **User Source** you want to manage.
3. Click **Add User** to add a new user. The User Settings page is displayed.

Username	Name	Roles	Contact Info	Schedule
admin	Administrator, Operator		Always	Edit Delete
guest			Always	Edit Delete

Add User

- Enter the [user's properties](#), including the roles you want this user to have.

Roles	Administrator	Maintenance	Prod_Sups
guests	<input type="checkbox"/>	<input checked="" type="checkbox"/> Operator	<input type="checkbox"/>
Test_Group	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>

- Click **Add User**. The user now has the privileges associated with the selected role(s).

Role Hierarchy

Often you might want to have one role that includes all the permissions for another role, i.e., Supervisor can do everything that Administration and Maintenance roles can do. In the Designer, access to components can be restricted to specific security roles. You can give any Supervisor both of the Administration and Maintenance.

User Management Component

Ignition has a special User Management component in the Vision Module that allows you to add, modify, and delete users and roles (and more) inside the Designer and the Client. By default, changes to a User Source may not be made from this component to prevent users from locking themselves out of the Gateway, or give themselves access to the Gateway. You can override this behavior by enabling the **Allow User Admin** property located under the [Gateway's general security settings](#). Once the Allow User Admin property is enabled, the User Management component is simple to set up and use.

- In Designer, go to the Project Browser and then to Vision.
- Create a new Window or open an existing one.
- Drag a [User Management](#) component to your window. This component will automatically point to the default user source being used by your project. You can change the User Source property if needed.
- If you already have some users and roles set up using the Gateway Webpage, you will see them in the User Management component. If you don't have any users or roles set up, you can create them here. Use the icons on the right side to add, edit, or delete a user or role.
- To add a new user, put the Designer in **Preview Mode**. Click the the plus  icon next to the user section.

The screenshot shows a 'Users' management interface. At the top, there's a table titled 'Users' with columns: Username, Name, Roles, Contact I..., and Schedule. The data includes:

Username	Name	Roles	Contact I...	Schedule
Jane_D	Jane Doe	Administration, Main...		Always
Maria	Maria Trejo	Administration, Main...		Always
Min_C	Min Chan	Supervisor		Always
opcuauser		ReadWrite		Always

To the right of the table is a vertical toolbar with a green plus icon and a button labeled 'Add New User'. Below the table is a section titled 'Roles' with a table showing the count of members for each role:

Role name	# of Members
Administration	2
Maintenance	2
Maintenance - East	0
ReadWrite	1

6. The Add User window will open. At a minimum, enter the **Username** and **Password**. All other properties are optional. When finished, click **Save**.

The screenshot shows the 'Add User' form. It has sections for 'User Properties' and 'Contact Info'.

User Properties:

- Username: Jerry_A
- Password: (two password fields, both marked with green checkmarks)
- First Name: Jerry
- Last Name: Anderson
- Schedule: Always
- Language: English
- Notes: (empty text area)

Roles: A list of roles with checkboxes next to them. The 'Administration' checkbox is checked.

Contact Info: A table with columns 'Type' and 'Value'. There is one empty row in the table.

7. To add a new role, make sure the Designer is in **Preview Mode**. Click the the plus **+** icon next to the role. The Add Role window will open.
8. Enter the name of the new role. Click **Save**.

Users > Add Role

< back 

Role name



9. Now you can see the user and role that were just added in the User Management window.

Users

Username	Name	Roles	Con...	Sch...
Jane_D	Jane Doe	Administration, Maintenance, ...		Alw...
Jerry_A	Jerry Anders...	Maintenance - East		Alw...
Maria	Maria Trejo	Administration, Maintenance, ...		Alw...
Min_C	Min Chan	Supervisor		Alw...
opcuauser		ReadWrite		Alw...

Roles

Role name	# of Members
Administration	2
Maintenance	2
Maintenance - East	1
ReadWrite	1

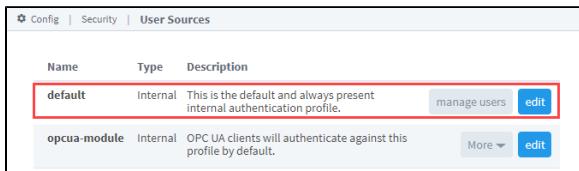
Internal Authentication

Internal User Sources

An Internal type User Source stores user information internally in the Gateway's database. This means that Internal User Sources are included in Gateway Backup files, and don't require an external SQL database, or other external user management system.

When Ignition is first installed, the [default User Source](#) that initially grants access to the Gateway and Design is an Internal type User Source. You can, of course, continue to use this default internal User Source for your project(s), or you may choose to use other User Sources instead.

The Internal User Source is fully [manageable](#) from within Ignition. You can access User Sources from the Gateway Webpage under the [Config](#) section, [Security > User, Roles](#), and click the [edit](#) button.



Name	Type	Description	Actions
default	Internal	This is the default and always present internal authentication profile.	manage users edit
opcua-module	Internal	OPC UA clients will authenticate against this profile by default.	More edit

On this page ...

- [Internal User Sources](#)
- [Property Reference](#)
 - [Main Properties](#)
 - [Password Policies Properties](#)



Internal Authentication

[Watch the Video](#)

Property Reference

This section details Internal User Source properties, organized by category.

Main Properties

Details on the Main Properties can be found on the [User Sources](#) page.

Password Policies Properties

The Internal User Source has password policies that are configurable from within the Gateway to provide an extra layer of security by ensuring that good password practices are used.

1. From the [Config](#) tab in the [Gateway Webpage](#), select [Users, Roles](#).
2. Click the Edit button for the User Source you want to update.
3. Scroll down to the Password Policy section. You can change the default password policies by entering the appropriate password values to support your password policies.

Password Policy	
Password Max Age	90 The maximum age (in days) that a password is valid for. A value of zero disables password expiration. (default: 0)
Password Min Length	10 Passwords must be at least this many characters long. (default: 1)
Password Complexity	2 The number of character types (lowercase letters, uppercase letters, digits, punctuation) each password must contain. For example, a value of 3 means passwords must have 3 of the 4 character types to be considered valid. (default: 1)
Password History	3 The number of previous passwords to store. Passwords in the history list may not be re-used. A value of zero disables this feature. (default: 0)

Below is a description of the Password Policy properties.

Name	Description
Password Maximum Age	<p>The maximum age in days that the password will still be valid. After the number of days has past, when the user tries to login, it will prompt them to change their password. A value of 0 will disable this feature.</p> <div style="border: 1px solid orange; padding: 5px; margin-top: 10px;"> <p>The following feature is new in Ignition version 8.1.23 Click here to check out the other new features</p> </div> <p>This property will also affect password age and expiration when you are using Ignition's Internal Identity Provider (IdP). If you try to log in after the expiration period, you will be redirected and prompted to update your password.</p>
Password Minimum Length	The minimum amount of characters that a password must contain to be considered valid. If the user tries to make a shorter password, it will not allow it, and let them know that it does not meet the minimum length requirements.
Password Complexity	This determines how complex a password must be. There are four character types: lowercase letter, uppercase letter, digits, and special characters. The value here determines how many of those character types must be present at least once in the password for it to be considered valid.
Password History	Determines the number of previously used passwords to store. When users make a new password, old passwords can not be re-used. A value of zero disables this functionality. History is only stored while this setting is turned on, so any passwords used while this is off can be re-used when history is turned back on.

Related Topics ...

- [Database Authentication](#)
- [Active Directory Authentication](#)
- [AD Internal Hybrid Authentication](#)
- [AD Database Hybrid Authentication](#)
- [Identity Provider Authentication Strategy](#)
- [User Management](#)
- [Managing Users and Roles](#)
- [Security](#)

Database Authentication

Database User Source

The Database Authentication type uses an external database instead of storing data inside Ignition. Managing users is done via direct interaction with the database. This section addresses how to set up a database user source. The Database Authentication type requires you have a connection to an existing database, like SQL Server, Oracle, or MySQL. It stores all users, roles, schedules, and more in the database, and uses queries to check login credentials. When you create a database user source, you have the option of setting it up in Automatic or Manual mode.

Automatic Mode

In Automatic mode, Ignition will create and manage the database tables for you. You can specify a prefix for the tables that are created automatically for you, but their names after the prefix are chosen by the user source. In this mode, the user source will be fully [manageable](#) in Ignition.

Manual Mode

In Manual mode, you must provide SQL queries for various functions of the user source. In this mode, the user source will not be manageable from the Gateway or the Clients. You'll have to manage the users directly through the database. Examples for each of the queries are given on the user source setup page. Read each query description carefully to make sure you design your queries to return all the columns that are defined in the query's description as shown below.

Property Reference

Database User Sources have the following properties, organized by category

Main Properties

Details on the Main Properties can be found on the [Classic Authentication Strategy](#) page. The Database User Source also has the following properties:

Name	Description
Database	The database connection this User Source will retrieve user information from.
Mode	How the Gateway should manage the database tables. Has two settings: Automatic: The gateway will automatically create the database tables necessary, and all interactions with the table will use the built-in queries. When this option is set, the Tablename Prefix property is utilized. Manual: The Gateway will not automatically create any database tables, nor will it automatically modify users or roles. When set to manual, it is assumed that you want to manually write the queries that update the tables, or are utilizing another system. When Mode is set to this option, the Manual Mode properties are used to determine how the Gateway should query user data.

Automatic Mode Properties

Name	Description
Tablename Prefix	When set to Automatic mode, this property determines the prefix that will be used on all automatically created tables. Useful when multiple database User Sources are connected to the same database scheme.

Manual Mode

Name	Description
Authentication Query	A query that must return a row if the given username and password combination provided is valid. The query will run as a prepared statement, so use the question mark character (?) to represent username first and then password. The returned row may contain the user's basic properties under the column names: [firstname, lastname, schedule, language, notes] Note that the Gateway will pass both the username and password the user typed in, so this query MUST utilize exactly two question marks, otherwise an exception will occur.

On this page ...

- [Database User Source](#)
- [Property Reference](#)
 - [Main Properties](#)
 - [Automatic Mode Properties](#)
 - [Manual Mode](#)
- [To Create a Database User Source](#)



INDUCTIVE
UNIVERSITY

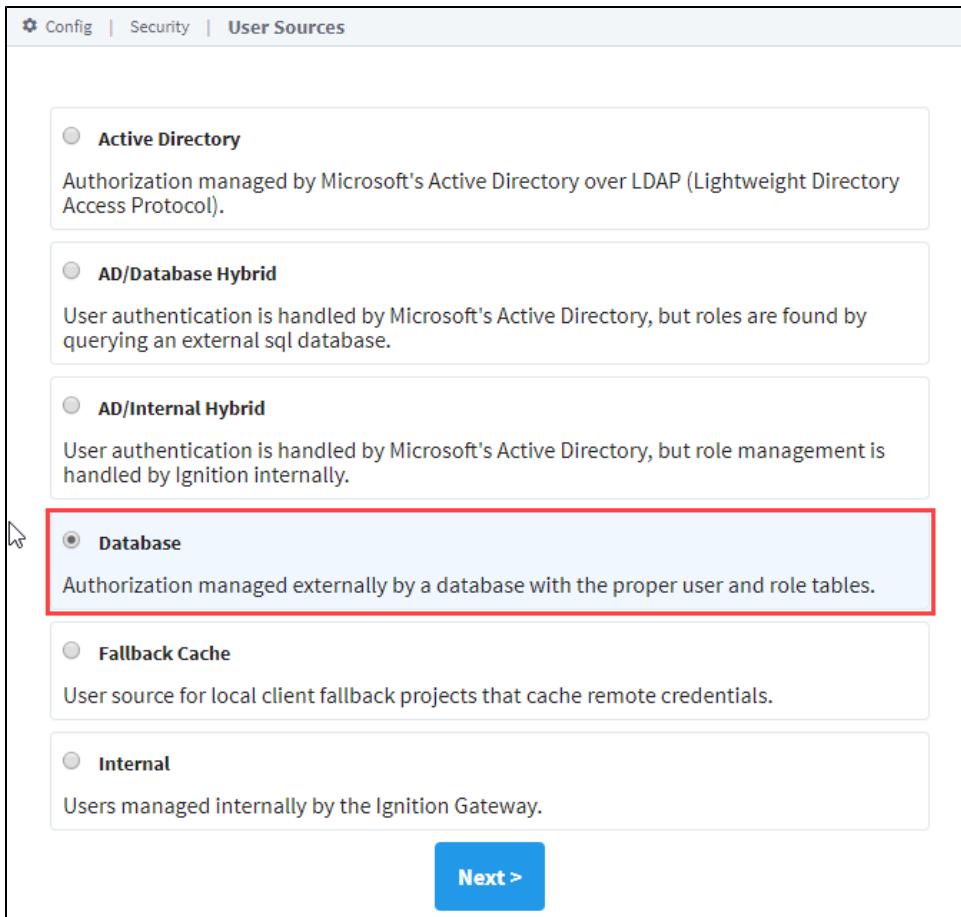
Database Authentication

[Watch the Video](#)

Badge Authentication Query	A query that must return a row if the given badge provided is valid. The query will be run as a prepared statement, so use question mark (?) to represent the badge. The returned row must contain the username. Example: SELECT username FROM USERS WHERE badge = ?
List Roles Query	A query that returns all possible roles that any user could be a member of. The role names must be returned in the first column of the query's results.
User's Roles Query	A query that returns all of the roles that the provided user belongs to. The roles must be strings and must be in the first column of the query's results. The query will be run as a prepared statement with one parameter: the username.
Find User Query	<p>The following feature is new in Ignition version 8.1.6 Click here to check out the other new features</p> <p>A query that must return a row if a user with the given username exists. The query will be run as a prepared statement, so use question mark (?) to represent username. There must be at least one column: the username. Other columns are optional, supported columns are: [username, firstname, lastname, schedule, language, notes, badge].</p> <p>Example: SELECT username, firstname, lastname, schedule, FROM USERS WHERE username = ?</p> <p>When the User Source is paired with an Ignition Identity Provider, enabling this option can provide a modest performance improvement during session login.</p>
List Users Query	A query that returns a row containing each username. There must be at least one column: the username. Other columns are optional, supported columns are: [username, firstname, lastname, schedule, language, notes].
Contact Info Query	A query that returns all of the contact info for the user. The first column must be the contact type, the second column the contact value. Optional, may be blank.
Schedule Adjustment Query	<p>A query that returns the upcoming schedule adjustments for the user. This property is optional, and may be left blank.</p> <p>The results set expects the following columns:</p> <ul style="list-style-type: none"> • Start(date) • End(date) • Available(boolean) • Note(string)
Extra Properties Query	A query that returns name, value pairs of extra properties for the user. Will be run with one parameter: the username. Optional, may be blank.

To Create a Database User Source

1. On the Gateway Webpage under the **Config** tab, go **Security > Users, Roles**.
The User Sources page will be displayed. Click the blue arrow, **Create new User Source**.
2. Choose the **Database** authentication type, and click **Next**.



3. The New User Source window will open. Some properties are optional, but if you're using Automatic mode, enter the following properties as appropriate.

- **Name:** DBAuth - name of the user source.
- **Failover Source:** default - failover user source ('default' is the internal user source).
- **Failover Mode:** Hard - if the source is unreachable, then use the failover source. (Can choose the Hard or Soft option).
- **Database:** MySQL - external database.
- **Mode:** Automatic - tables in the external database will be automatically created when needed.
- **Tablename Prefix:** 'auth_' is the prefix for all the tables that get created. (You can leave this field blank, but if you use a prefix when the tables get created, they will contain the specified prefix in their name).
- When finished, click **Create New User Source**.

Config | Security | User Sources

Main

Name	DBauth <input type="button" value=""/>
Description	Database authentication
Schedule Restricted	<input type="checkbox"/> Users are only able to log in when their assigned schedule is active. (default: false)
Failover Source	default <input type="button" value=""/> If this source is unreachable for authentication, this failover source will be used instead.
Failover Mode	Hard <input type="button" value=""/> The failover mode to use if a failover source is set. Hard: Failover only if this source is un-reachable. Soft: Try the failover source when a user fails to authenticate with this source. (default: HARD)
Cache Validation Timeout	60000 The amount of time between cache updates of the user source. (default: 60,000)

Main Properties

Database	MySQL <input type="button" value=""/> Choose the database connection this user source will use.
Mode	Automatic <input type="button" value=""/> In Automatic mode, tables will be created for you and all interaction with the database is handled automatically. In Manual mode, you write queries by hand against tables that you've created. Management of users (adding, removing etc) is not supported in manual mode. (default: Automatic)

Automatic Mode

Tablename Prefix	auth_ <input type="button" value=""/> When in automatic mode, tables will be created to store the users and roles. They will have this prefix. (default: scada_)
------------------	---



Table Creation

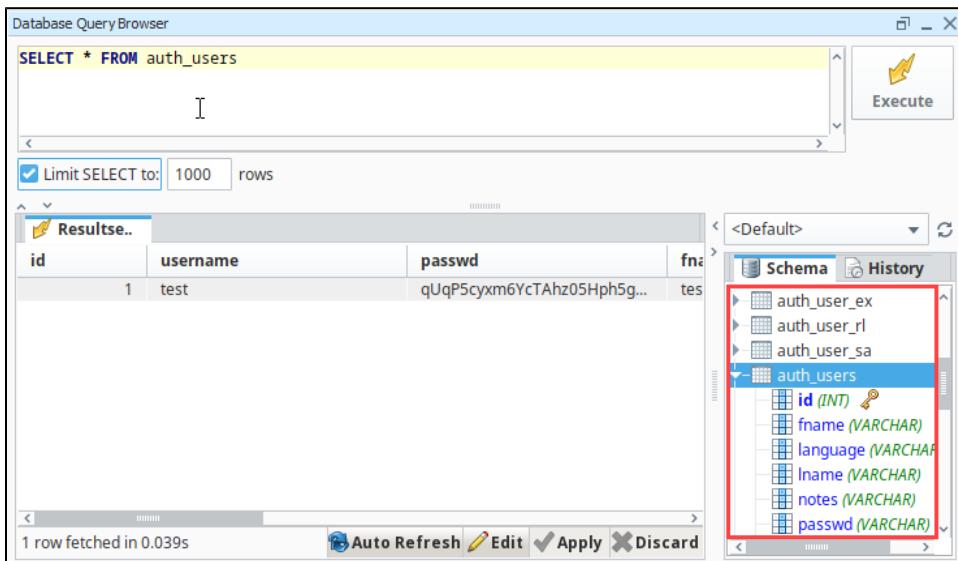
The tables in the database will not be created in the database until they are needed. For example, as soon as a user or role is added, the associated tables will automatically get created.

- Now that your Authentication profile is created, add a user. On the right, click on the **More > Manage Users** link. Click on the **Add User** link and fill in the required fields.

5. Now that your tables are created we can verify them. To view the tables, go into **Designer** and from the menu bar, select **Tools > Database Query Browser**.

You will see all the tables that were created beginning with 'auth_' when the user and role get created.

6. Double click on any of the tables beginning with 'auth_', and click **Execute**. In this example, you will see the tables associated with 'roles' and 'users' displayed in the Schema area.



Active Directory Authentication

Active Directory User Source

The Active Directory Authentication profile uses Microsoft's Active Directory over **LDAP** (Lightweight Directory Access Protocol) to store all the users, roles, and more that make up an Authentication profile. Active Directory **Groups** are used for Ignition's **roles** and user-role mappings.

While using an Active Directory User Source, administration of users and roles is through Active Directory itself, and not manageable within Ignition. Thus adding new users to an Active Directory User Source, or modifying pre existing users, requires the modifications be made from Active Directory, usually through an AD Administrator.

Active Directory User Sources supports SASL (Simple Authentication and Security Layer). SASL is a framework for authentication and data security in Internet protocols such as LDAP.

Property Reference

Active Directory User Sources have the following properties shown in the table below, organized by category.

On this page ...

- [Active Directory User Source](#)
- [Property Reference](#)
 - [Main Properties](#)
 - [Active Directory Properties](#)
 - [LDAP Search Properties](#)
 - [SASL Properties](#)
- [To Create an Active Directory User Source](#)
- [Connect AD over SSL](#)

Note: Certain properties in the Active Directory User Source allow you to filter users, such as the **User List Filter**. These filters only determine which users will be displayed on screen. They are not authentication filters, so even if a user does not show in the list they can still authenticate and may have access to unintended areas. Be sure to configure [project security](#) appropriately to prevent this from happening!

Main Properties

Details on the Main Properties can be found on the [Classic Authentication Strategy](#) page.

Active Directory Properties

Name	Description
Domain	The Windows Domain your active Active Directory server is running on. If you aren't sure of your domain, ask your network administrator. Leave blank to set advanced properties manually.
Gateway Username	The login name for the Gateway to use when querying Active Directory. Used for retrieving the list of users and roles via LDAP.
Password	The password for the above username.
Secondary Domain Controller Host	The IP address or hostname of your primary domain controller. Example: "192.168.1.4" or "MainServer"
Secondary Domain Controller Port	The port number for the primary domain controller's LDAP interface.
SSO Domain	The domain that Windows users must match in order to use SSO. If blank, the main "Domain" property will be used. Not case-sensitive.

This feature was changed in Ignition version **8.1.17**:

This **SSO Enabled** setting was disabled and deprecated in 8.1.17 to protect against a potential security vulnerability. While the property is still visible, it cannot be enabled without setting a special system property. This is not recommended. See the [Active Directory Deprecated Properties](#) page for more information.

LDAP Search Properties

Name	Description
Username Prefix	This prefix will be prepended to the username before an Active Directory bind is attempted for authentication.
Username Suffix	This suffix will be appended to the username before an Active Directory bind is attempted for authentication.
Automatic Suffix	If this option is checked, and the suffix is left blank, then the suffix will automatically be assigned a value of "@<domain>".
Use prefix and suffix for Gateway username	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>If this option is checked, the username prefix and suffix will be applied to the Gateway username before a bind is attempted. This option is checked by default.</p>
User Search Base	<p>The base folder to search for users under, such as:</p> <p>DC=MyCompany,DC=com</p> <p>The entire subtree under this folder will be searched using the User Search Filter.</p> <p>Multiple subtrees can be specified by putting them in parenthesis, like so:</p> <p>(OU=Administrators,DC=MyCompany,DC=com)(OU=Operators,DC=MyCompany,DC=com)</p>
User Search Filter	The LDAP search filter that will be used to find a specific user. Use the placeholder {0} as a standin for the login name.
User List Filter	The LDAP search filter used when querying for the list of all users. Should restrict the type to user.
User Name Attribute	The attribute on the User object to define the username.
User Role Attribute	Attributes of this name on the User object will define the user's roles.
Role Name Attribute	The attribute of this name on the Role object will define the role's name. Leave blank to use the raw value of the attribute defined by the User Role Attribute property.
Full Name Attribute	The attribute on the User object to define the full name of the user.
Phone Attribute	The attribute name on the user object that represents the user's phone number.
Email Attribute	The attribute name on the user object that represents the user's email address.
SMS Attribute	The attribute name on the user object that represents the phone number that this user receives text messages on.
Badge Attribute	<p>The following feature is new in Ignition version 8.1.25 Click here to check out the other new features</p> <p>The attribute on the User object to define the badge. This setting is required to enable badge-based authentication.</p>
Read Timeout	The read timeout in milliseconds for LDAP operations.
Results Page Size	The number of entries returned per page of results in a query.
Role Search Base	<p>The base folder to search for roles under, such as:</p> <p>OU=Roles,DC=MyCompany,DC=com</p> <p>The entire subtree under this folder will be searched using the Role Search Filter. If you specify the root of your tree structure, the search may take a very long time.</p> <p>Multiple subtrees can be specified by putting them in parenthesis, like so:</p> <p>(OU=Builtin,DC=MyCompany,DC=com)(OU=Users,DC=MyCompany,DC=com)</p>

	If you leave this blank the whole subtree of the domain controller will be searched.
Role Search Filter	The LDAP search filter that will be used to locate roles.
Badge Search Filter	<p>The following feature is new in Ignition version 8.1.25 Click here to check out the other new features</p> <p>The LDAP search filter to use to find a specific user given a badge. Use the placeholder {0} as a stand-in for the user's badge. Example: (&(objectClass=user)(badge={0}))</p>
Allow Anonymous	<p>If enabled, authentication attempts with blank passwords will be passed through to LDAP, which may choose to accept them.</p> <p>Caution: It is highly recommended to disable this setting unless you know it is required. AD servers may allow logging in as any user with a blank password when Security Authentication is set to "None" or "Simple" (even if a provided username does not exist in AD), which is a major security risk.</p>
Security Protocol	<p>Specifies the security protocol between the Gateway and AD server. The following options are available:</p> <p>AUTO: No security protocol is explicitly used or requested by the Gateway.</p> <p>SSL: SSL should be used for the connection.</p>
Security Authentication	<p>This property specifies how usernames and passwords are used to bind to LDAP. The following options are available:</p> <p>AUTO: Unspecified from the Gateway side, meaning the LDAP implementation will choose.</p> <p>NONE: Anonymous access. (Not recommended due to security risks)</p> <p>SIMPLE: Plaintext username and passwords will be used. (Not recommended due to security risks)</p> <p>STRONG: Usernames and passwords will be encrypted.</p> <p>SASL: Simple Authentication and Security Layer. See the SASL Properties table below for additional SASL authentication configuration settings.</p>
Referral	<p>The following feature is new in Ignition version 8.1.1 Click here to check out the other new features</p> <p>Specifies how referrals are to be processed. Possible options are:</p> <p>Follow: Always automatically follow referrals. This is the default option.</p> <p>Ignore: Ignores referrals.</p> <p>Throw: Throws a ReferralException whenever a referral is encountered.</p>

SASL Properties

These settings are utilized when Security Authentication is set to **SASL**.

Name	Description
Mechanism	An ordered list of space-separated mechanism names. The LDAP provider will use the first mechanism for which it finds an implementation. A blank value will leave this setting unspecified. (Default is DIGEST-MD5 CRAM-MD5.)
Realm	A realm defines the namespace from which the user is selected. A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. Default is blank.
Quality of Protection	A comma-separated list of Quality-of-Protection (QoP) values, the order of which specifies the preference order. There are three well-known values: "auth" (authentication only), "auth-int" (authentication with integrity protection), and "auth-conf" (authentication with integrity and privacy protection). A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. (Default is auth-conf,auth-int,auth.)
Protect	A comma-separated list of privacy protection strength values, the order of which specifies the preference order. The three possible

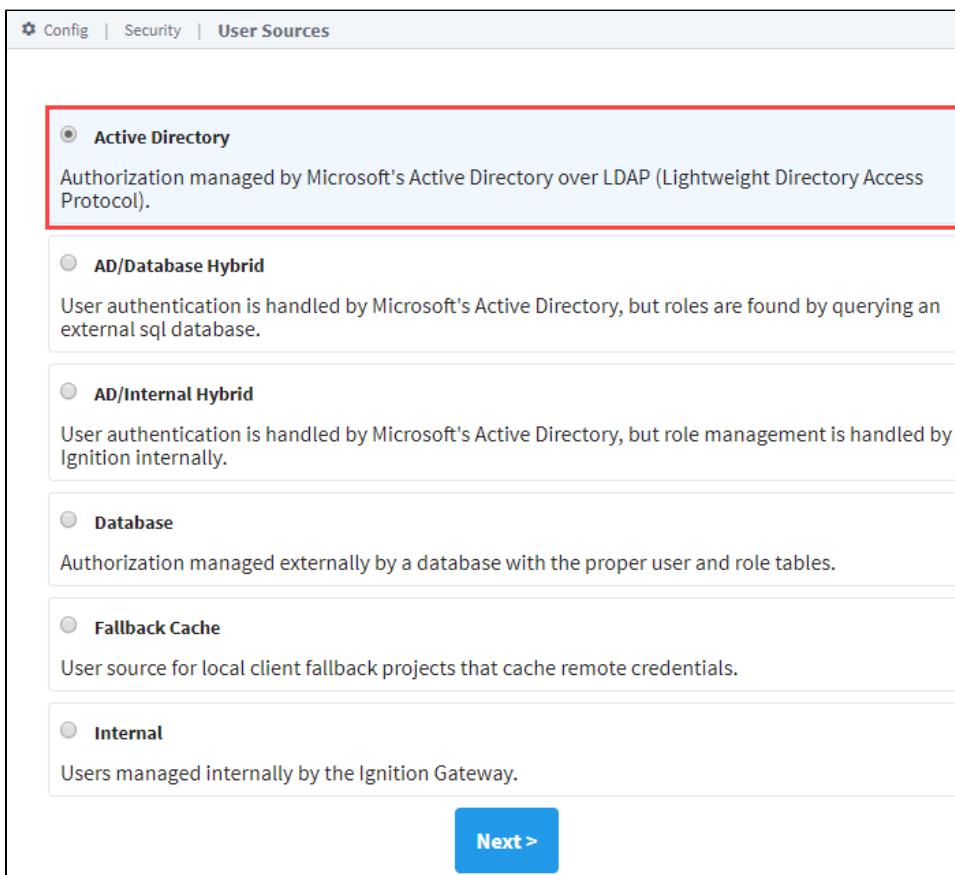
on Strength	strength values are "low", "medium", and "high". A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. Default is high,medium,low.
Mutual Authentication	Enable or disable mutual authentication. This setting will only be used by mechanisms which support it. Default is disabled.

To Create an Active Directory User Source

To configure an Active Directory User Source, you must specify the host that is acting as your primary domain controller. You can also use a secondary domain controller in case the primary is unavailable. You'll also need to specify the name of the domain and credentials for the Gateway itself to use: the Gateway needs a user account to interact with the AD server, even when it's simply querying for a list of roles.

Note: When using Active Directory User Source, you may need to consult with your internal IT Department to get the required information to complete your user source setup. These settings are common to AD (not specific to Ignition), and your IT department will know what values to supply to each property.

1. On the [Gateway Webpage](#), under the Config tab, go to **Security > Users, Roles**.
The User Sources page will be displayed. Click the blue arrow, **Create new User Source**.
2. Choose the **Active Directory** authentication type, and click **Next**.



3. The New User Source window will open. Some properties are optional. In the very least, you must specify the following: **Domain, Gateway Username, Password, Primary Domain Controller Host**.
4. Click the **Create New User Source** button to create the User Source.

Connect AD over SSL

For additional security, you can adjust the Active Directory settings to enable SSL since LDAP is not encrypted by default. The default port for LDAP is port 389, but LDAPS uses port 636 and establishes SSL upon connecting.

1. On the Gateway Webpage, under the Config tab, go to **Security > Users, Roles**.

- Click the **Create new User Source** blue arrow on the User Sources page.
- Choose the **Active Directory** authentication type, and click **Next**.
- Change the **Primary Domain Controller Port** to 636.

Primary Domain Controller Host	The IP address or hostname of your primary domain controller. Example: "192.168.1.4" or "MainServer"
Primary Domain Controller Port	636 The port number for the primary domain controller's LDAP interface. (default: 389)
Secondary Domain Controller Host	The IP address or hostname of your secondary domain controller (optional). Example: "192.168.1.4" or "MainServer"
Secondary Domain Controller Port	389 The port number for the secondary domain controller's LDAP interface. (default: 389)
Use SSL	<input checked="" type="checkbox"/> Disable to use "ldap://" protocol, enabled to use "ldaps://" (default: false)

- Check **Use SSL** to enable "ldaps://".
- Check **Show advanced properties** to expand.
- Change the **Security Protocol** to SSL.

Security Protocol	SSL (default: AUTO)
-------------------	----------------------------

If you try to query or authenticate against the AD server at this point, you will receive the following error:

```
Caused by: javax.net.ssl.SSLHandshakeException: sun.security.validator.ValidatorException: PKIX path building failed: sun.security.provider.certpath.SunCertPathBuilderException: unable to find valid certification path to requested target
```

This error indicates Ignition was unable to find a valid certificate generated from the AD server, and therefore cannot validate the AD server's identity. Work with IT to obtain a certificate from the AD server. This certificate must be added to the data/certificates/supplemental directory and then imported into the Java `cacerts` keystore. Once the certificate is added, restart the Gateway.

Related Topics ...

- [Internal Authentication](#)
- [Database Authentication](#)
- [AD Internal Hybrid Authentication](#)
- [AD Database Hybrid Authentication](#)
- [Identity Provider Authentication Strategy](#)
- [Classic Authentication Strategy](#)
- [Project Security in the Designer](#)

AD Internal Hybrid

AD/Internal User Source

The Active Directory/Internal Hybrid authentication profile type combines the Internal User Source type with the Active Directory User Source type. Active Directory is used to find all of the users, and to check their credentials when they attempt to log in. However, it allows assigning of roles, contact info, and other meta-information about a user through Ignition, then stores all this information as if it were an Internal User Source. This way, Active Directory can be consulted to see if a username/password is valid, but the management of roles does not require coordination with your IT Department, who typically controls the Active Directory system. This "best of both worlds" approach is popular for many users of Active Directory.

The AD/Internal Hybrid User Source is partially manageable in Ignition. Users cannot be added or removed, and their usernames and passwords cannot be changed. This is because this information resides in Active Directory, not within Ignition. Other information, such as user roles, contact info, schedules, are [manageable](#) in Ignition.

Gateway Settings

Before you can use the User Management component to manage roles, contact info, etc., you first have to go into [Gateway Settings](#), and mark the checkbox to '**Allow User Admin.**' This allows for the administration of the Gateway's system user source from the Designer and the Client. Unless this is enabled, the Vision Module's User Management component is prevented from modifying the Gateway system's user source.

Property Reference

This [User Source](#) shares many properties with the AD User Source. Please see the [Active Directory Authentication](#) page for a list of properties.

Creating an AD/Internal Hybrid User Source

To set up an AD/Internal Hybrid User Source, you must specify the host that is acting as your primary domain controller. You can also use a secondary domain controller in case the primary is unavailable. You'll also need to specify the name of the domain and credentials for the Gateway itself to use for authentication for when it queries the list of roles.

May need to contact your internal IT Department for...

When using AD/Internal Hybrid User Source, you may need to consult with your internal IT Department to get the required information to complete your user source setup.

1. On the [Gateway Webpage](#), under the Config tab, go to **Security > Users, Roles**. The User Sources page will be displayed. Click the blue arrow, **Create new User Source**.
2. Choose the **AD/Internal Hybrid** authentication type, and click **Next**.

On this page ...

- [AD/Internal User Source](#)
- [Property Reference](#)
- [Creating an AD/Internal Hybrid User Source](#)
 - [Active Directory Properties](#)
 - [Advanced Properties](#)
 - [SASL Properties](#)



INDUCTIVE
UNIVERSITY

AD Internal Hybrid

[Watch the Video](#)

The screenshot shows a configuration interface for user sources. At the top, there are tabs for 'Config' and 'Security', with 'User Sources' selected. Below the tabs, there are six options listed in boxes:

- Active Directory**: Authorization managed by Microsoft's Active Directory over LDAP (Lightweight Directory Access Protocol).
- AD/Database Hybrid**: User authentication is handled by Microsoft's Active Directory, but roles are found by querying an external sql database.
- AD/Internal Hybrid**: User authentication is handled by Microsoft's Active Directory, but role management is handled by Ignition internally. This option is highlighted with a red border.
- Database**: Authorization managed externally by a database with the proper user and role tables.
- Fallback Cache**: User source for local client fallback projects that cache remote credentials.
- Internal**: Users managed internally by the Ignition Gateway.

At the bottom right of the list area is a blue 'Next >' button.

3. The New User Source window will open. Some properties are optional depending on how you set up your profile. Details on the Main Properties can be found on the [User Sources](#) page. Active Directory properties are listed in the table below.
4. Click **Create New User Source** to save the new user source.

Active Directory Properties

Name	Description
Domain	The Windows Domain your active <u>Active Directory server</u> is running on. If you aren't sure of your domain, ask your <u>network administrator</u> . Leave blank to set advanced properties manually.
Primary Domain Controller Host	The IP address or hostname of your primary domain controller. Example: "192.168.1.4" or "MainServer"
Primary Domain Controller Port	The port number for the primary domain controller's <u>LDAP</u> interface.
List Users from Active Directory	If true, Active Directory will be queried for the list of all users. If false, users must be added manually. Default is true.
Populate Users On-Demand	<p>The following feature is new in Ignition version 8.1.6 Click here to check out the other new features</p> <p>If true, a user record will be created in the Internal Database from the AD entry of any user who successfully authenticates into Active Directory if the record does not already exist in the Internal Database at the time of authentication.</p> <p>When the User Source is paired with an <u>Ignition Identity Provider</u>, enabling this option can provide a modest performance improvement during session login.</p>
Gateway Username	The login name for the Gateway to use when querying Active Directory. Used for retrieving the list of users and roles via LDAP.
Password	The password for the above username.
Re-type Password	Re-type password for verification.

SSO Enabled	Whether or not to use Single-Sign-On (SSO) to authenticate AD users. Note that projects must also have this option enabled for SSO to work. Default is false.
	<p>This feature was changed in Ignition version 8.1.17: This setting was disabled and deprecated in 8.1.17 to protect against a potential security vulnerability. While the property is still visible, it cannot be enabled without setting a special system property. This is not recommended. Check out the Active Directory SSO Disabled for 8.1.17 & 7.9.20 article for more information.</p>
SSO Domain	The domain that Windows users must match in order to use SSO . If blank, the main "Domain" property will be used. Not case-sensitive.

Advanced Properties

Name	Description
Secondary Domain Controller Host	The IP address or hostname of your secondary domain controller (optional). Example: "192.168.1.5" or "BackupServer"
Secondary Domain Controller Port	The port number for the secondary domain controller's LDAP interface. Default: 389.
Read Timeout	The read timeout in milliseconds for LDAP operations. Default is 60,000.
Results Page Size	The number of entries returned per page of results in a query. Default is 1,000.
User Listing Base	The base folder to search for users under, such as "DC=MyCompany,DC=com". The entire subtree under this folder will be searched using the User List Filter. Multiple subtrees can be specified by putting them in parenthesis, like so: "(OU=Administrators, DC=MyCompany,DC=com)(OU=Operators,DC=MyCompany,DC=com)" If you leave this blank the whole subtree of the domain controller will be searched.
User List Filter	The LDAP search filter used when querying for the list of all users. Should restrict the type to user. Default is (& (objectClass=user) (! (objectClass=computer))).
User Search Filter	The LDAP search filter to use to find a specific user. Use the placeholder {0} as a standin for the login name. Default is (& (objectClass=user) (sAMAccountName={0})).
Username Attribute	The attribute on the User object to define the username. Default is sAMAccountName.
Username Prefix	This prefix will be prepended to the username before an Active Directory bind is attempted for authentication.
Username Suffix	This suffix will be appended to the username before an Active Directory bind is attempted for authentication.
Automatic Suffix	If this option is checked, and the suffix is left blank, then the suffix will automatically be assigned a value of "@<domain>". Default is true.
Use prefix and suffix for Gateway username	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>If this option is checked, the username prefix and suffix will be applied to the Gateway username before a bind is attempted. This option is checked by default.</p>
Allow Anonymous	If enabled, authentication attempts with blank passwords will be passed through to LDAP, which may choose to accept them.
	<p>Caution: It is highly recommended to disable this setting unless you know it is required. AD servers may allow logging in as any user with a blank password when Security Authentication is set to "None" or "Simple" (even if a provided username does not exist in AD), which is a major security risk.</p>
Security	Auto or SSL. Default is Auto.

Protocol	
Security Authentication	<p>This property specifies how usernames and passwords are used to bind to LDAP. The following options are available:</p> <p>AUTO: Unspecified from the Gateway side, meaning the LDAP implementation will choose.</p> <p>NONE: Anonymous access.</p> <p>SIMPLE: Plaintext username and passwords will be used.</p> <p>STRONG: Usernames and passwords will be encrypted.</p> <p>SASL: Simple Authentication and Security Layer. See the SASL Properties table below for additional SASL authentication configuration settings.</p>
Referral	<p>The following feature is new in Ignition version 8.1.1 Click here to check out the other new features</p> <p>Specifies how referrals are to be processed. Possible options are:</p> <p>Follow: Always automatically follow referrals. This is the default option.</p> <p>Ignore: Ignores referrals.</p> <p>Throw: Throws a ReferralException whenever a referral is encountered.</p>

SASL Properties

These settings are utilized when Security Authentication is set to **SASL**.

Name	Description
Mechanism	An ordered list of space-separated mechanism names. The LDAP provider will use the first mechanism for which it finds an implementation. A blank value will leave this setting unspecified. Default is DIGEST-MD5 CRAM-MD5.
Realm	A realm defines the namespace from which the user is selected. A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. Default is blank.
Quality of Protection	A comma-separated list of Quality-of-Protection (QoP) values, the order of which specifies the preference order. There are three well-known values: "auth" (authentication only), "auth-int" (authentication with integrity protection), and "auth-conf" (authentication with integrity and privacy protection). A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. Default is auth-conf,auth-int,auth.
Protection Strength	A comma-separated list of privacy protection strength values, the order of which specifies the preference order. The three possible strength values are "low", "medium", and "high". A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. Default is high,medium,low.
Mutual Authentication	Enable or disable mutual authentication. This setting will only be used by mechanisms which support it. Default is disabled.

AD Database Hybrid

AD/Database User Source

This AD/Database Hybrid User Source is not [manageable](#) from within Ignition. Users/passwords must be administered through [Active Directory](#), and roles, contact info, and so on, must be administered directly through the database. The way AD/Database Hybrid works, is it has all the same information requirements as the other authentication profiles, but it also has a number of Database properties. You need to specify a database (i.e., MySQL) to store information, and set up queries that you want to use. You must also specify the host that is acting as your primary domain controller, and a secondary domain controller in case the primary is unavailable. You'll also need to specify the name of the domain and credentials for the Gateway itself to use for authentication when it queries the list of roles.

AD/Database User Sources now support SASL (Simple Authentication and Security Layer). SASL is a framework for authentication and data security in Internet protocols such as LDAP.

Property Reference

This [User Source](#) shares many properties with both the AD User Source and Database User Source.

- See the [Active Directory Authentication](#) page for a list of Active Directory User Source related properties.
- See the [Database Authentication](#) page for a list of Database User Source related properties.

On this page ...

- [AD/Database User Source](#)
- [Property Reference](#)
- [Creating an AD/Database Hybrid User Source](#)
 - [Active Directory Properties](#)
 - [Database Properties](#)
 - [Advanced Properties](#)
 - [SASL Properties](#)



INDUCTIVE
UNIVERSITY

AD Database Hybrid

[Watch the Video](#)

Creating an AD/Database Hybrid User Source

1. On the [Gateway Webpage](#), under the Config tab, go to **Security > Users, Roles**.
The [Classic Authentication Strategy](#) page will be displayed. Click the blue arrow, **Create new User Source**.
2. Choose the **AD/Database Hybrid** authentication type, and click **Next**.

Config | Security | User Sources

- Active Directory
Authorization managed by Microsoft's Active Directory over LDAP (Lightweight Directory Access Protocol).
- AD/Database Hybrid
User authentication is handled by Microsoft's Active Directory, but roles are found by querying an external sql database.
- AD/Internal Hybrid
User authentication is handled by Microsoft's Active Directory, but role management is handled by Ignition internally.
- Database
Authorization managed externally by a database with the proper user and role tables.
- Fallback Cache
User source for local client fallback projects that cache remote credentials.
- Internal
Users managed internally by the Ignition Gateway.

Next >

Note: When using AD/Database Hybrid User Source, you may need to consult with your internal IT Department to get the required information to complete your user source setup.

3. The New User Source window will open. Some properties are optional depending on how you set up your profile. Details on the Main Properties can be found on the [User Sources](#) page. The Active Directory Properties and Database properties are listed in the tables below.
4. Click **Create New User Source** to save the new user source.

Active Directory Properties

Name	Description
Domain	The Windows Domain your active Active Directory server is running on. If you aren't sure of your domain, ask your network administrator . Leave blank to set advanced properties manually.
Primary Domain Controller Host	The IP address or hostname of your primary domain controller. Example: "192.168.1.4" or "MainServer"
Primary Domain Controller Port	The port number for the primary domain controller's LDAP interface.
List Users from Active Directory	If true, Active Directory will be queried for the list of all users. If false, users must be added manually. (Default is true.)
Gateway Username	The login name for the Gateway to use when querying Active Directory. Used for retrieving the list of users and roles via LDAP.
Password	The password for the above username.
Password	Re-type password for verification.
SSO Enabled	Whether or not to use Single-Sign-On (SSO) to authenticate AD users. Note that projects must also have this option enabled for SSO to work. (Default is false.)

This feature was changed in Ignition version [8.1.17](#):

	This setting was disabled and deprecated in 8.1.17 to protect against a potential security vulnerability. While the property is still visible, it cannot be enabled without setting a special system property. This is not recommended. Check out the Active Directory SSO Disabled for 8.1.17 & 7.9.20 article for more information.
SSO Domain	The domain that Windows users must match in order to use SSO . If blank, the main "Domain" property will be used. Not case-sensitive.

Database Properties

Name	Description
Database	Dropdown list. Choose the database connection this authentication profile will use.
User Properties Query	A query that returns the basic properties for a single user. Supported return columns are [username, firstname, lastname, schedule, language, notes].
Role List Query	A query that returns all possible roles that any user could have. The role names must be returned in the first column of the query's results
User's Roles Query	A query that returns all of the roles that the provided user belongs to. The roles must be strings (i.e., the role names), and must be in the first column of the query's results. The username will be inserted into this query as a parameter.
Contact Info Query	A query that returns all of the contact info for the user. The first column must be the contact type, the second column the contact value, and the third column the name of a schedule. Optional, may be blank.
Schedule Adjustment Query	A query that returns the upcoming schedule adjustments for the user. Columns must be Start(date), End(date), Available(boolean), Note(string). Optional, may be blank.
Extra Properties Query	A query that returns name, value pairs of extra properties for the user. Will be run with one parameter: the username. Optional, may be blank.
Find User Query	<p>The following feature is new in Ignition version 8.1.6 Click here to check out the other new features</p> <p>A query that must return a row if a user with the given username exists. The query will be run as a prepared statement, so use question mark (?) to represent username. There must be at least one column: the username. Other columns are optional, supported columns are: [username, firstname, lastname, schedule, language, notes, badge].</p> <p>Example: <code>SELECT username, firstname, lastname, schedule, FROM USERS WHERE username = ?</code></p> <p>When the User Source is paired with an Ignition Identity Provider, enabling this option can provide a modest performance improvement during session login.</p>
List Users Query	A query that returns a row containing each username. Only used if "List Users from Active Directory" is false. There must be at least one column: the username. Other columns are optional, supported columns are: [username, firstname, lastname, schedule, language, notes].
Badge Authentication Query	A query that must return a row if the given badge provided is valid. The query will be run as a prepared statement, so use question mark (?) to represent the badge. The returned row must contain the username. Example: <code>SELECT username FROM USERS WHERE badge = ?</code>

Advanced Properties

Name	Description
Secondary Domain Controller Host	The IP address or hostname of your secondary domain controller (optional). Example: "192.168.1.5" or "BackupServer"
Secondary Domain Controller Port	The port number for the secondary domain controller's LDAP interface. Default: 389.

Read Timeout	The read timeout in milliseconds for LDAP operations. Default is 60,000.
Results Page Size	The number of entries returned per page of results in a query. Default is 1,000.
User Listing Base	The base folder to search for users under, such as "DC=MyCompany,DC=com". The entire subtree under this folder will be searched using the User List Filter. Multiple subtrees can be specified by putting them in parenthesis, like so: "(OU=Administrators, DC=MyCompany,DC=com)(OU=Operators,DC=MyCompany,DC=com)" If you leave this blank the whole subtree of the domain controller will be searched.
User List Filter	The LDAP search filter used when querying for the list of all users. Should restrict the type to user. Default is (& (objectClass=user) (! (objectClass=computer))).
User Search Filter	The LDAP search filter to use to find a specific user. Use the placeholder {0} as a standin for the login name. Default is (& (objectClass=user) (sAMAccountName={0})).
Username Attribute	The attribute on the User object to define the username. Default is sAMAccountName.
Username Prefix	This prefix will be prepended to the username before an Active Directory bind is attempted for authentication.
Username Suffix	This suffix will be appended to the username before an Active Directory bind is attempted for authentication.
Automatic Suffix	If this option is checked, and the suffix is left blank, then the suffix will automatically be assigned a value of "@<domain>". Default is true.
Use prefix and suffix for Gateway username	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>If this option is checked, the username prefix and suffix will be applied to the Gateway username before a bind is attempted. This option is checked by default.</p>
Allow Anonymous	<p>If enabled, authentication attempts with blank passwords will be passed through to LDAP, which may choose to accept them.</p> <p>Caution: It is highly recommended to disable this setting unless you know it is required. AD servers may allow logging in as any user with a blank password when Security Authentication is set to "None" or "Simple" (even if a provided username does not exist in AD), which is a major security risk.</p>
Security Protocol	Auto or SSL. Default is Auto.
Security Authentication	<p>This property specifies how usernames and passwords are used to bind to LDAP. The following options are available:</p> <p>AUTO: Unspecified from the Gateway side, meaning the LDAP implementation will choose.</p> <p>NONE: Anonymous access.</p> <p>SIMPLE: Plaintext username and passwords will be used.</p> <p>STRONG: Usernames and passwords will be encrypted.</p> <p>SASL: Simple Authentication and Security Layer. See the SASL Properties table below for additional SASL authentication configuration settings.</p>
Referral	<p>The following feature is new in Ignition version 8.1.1 Click here to check out the other new features</p> <p>Specifies how referrals are to be processed. Possible options are:</p> <p>Follow: Always automatically follow referrals. This is the default option.</p> <p>Ignore: Ignores referrals.</p> <p>Throw: Throws a ReferralException whenever a referral is encountered.</p>

SASL Properties

These settings are utilized when Security Authentication is set to **SASL**.

Name	Description
Mechanism	An ordered list of space-separated mechanism names. The LDAP provider will use the first mechanism for which it finds an implementation. A blank value will leave this setting unspecified. (Default is DIGEST-MD5 CRAM-MD5.)
Realm	A realm defines the namespace from which the user is selected. A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. (Default is blank.)
Quality of Protection	A comma-separated list of Quality-of-Protection (QoP) values, the order of which specifies the preference order. There are three well-known values: "auth" (authentication only), "auth-int" (authentication with integrity protection), and "auth-conf" (authentication with integrity and privacy protection). A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. (Default is auth-conf,auth-int,auth.)
Protection Strength	A comma-separated list of privacy protection strength values, the order of which specifies the preference order. The three possible strength values are "low", "medium", and "high". A blank value will leave this setting unspecified. This setting will only be used by mechanisms which support it. (Default is high,medium,low.)
Mutual Authentication	Enable or disable mutual authentication. This setting will only be used by mechanisms which support it. (Default is disabled.)

Fallback Cache Authentication

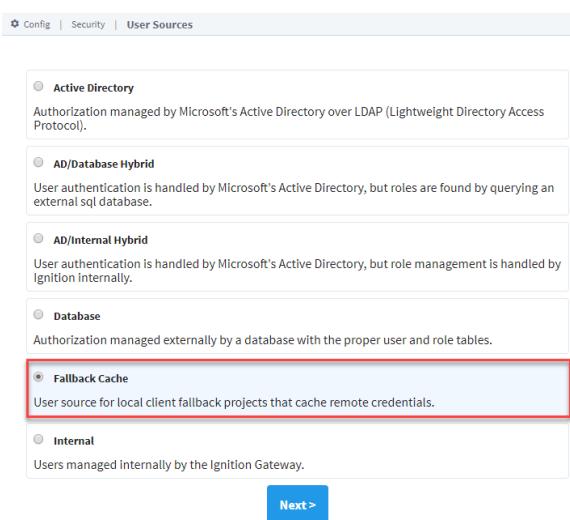
This User Source was developed specifically for a system that is using [Local Vision Client Fallback](#), and allows you to cache the login credentials from a remote user source. This means your users can still log in with their normal username/password on a Local Vision Client Fallback project, even when the network connection is unavailable.

Note: Fallback Cache Authentication does not work with [Identity Provider Authentication Strategy](#). This User Source will only function with Vision clients and user sources using [Classic Authentication Strategy](#)

Creating the Fallback Cache User Source

The Fallback Cache User Source is created in a similar fashion to any other User Source:

1. On the [Gateway Webpage](#), go to the **Config** tab. Select **Security > User, Roles**.
2. Click on the [Create new User Source...](#) link.
3. Select the Fallback Cache option and click the **Next** button.



4. Type in a name for the new User Source and click [Create New User Source](#).
5. Details on the Main Properties can be found on the [User Sources](#) page. You can also set Cache Retention as follows.

Fallback Cache Properties	
Cache Retention	Number of days that the cache will retain recently used credentials. This property determines the number of days credentials will be stored in the cache. Once exceeded, the credentials will be removed from the cache. (Default: 15)

6. Click the [Create New User Source](#) button.

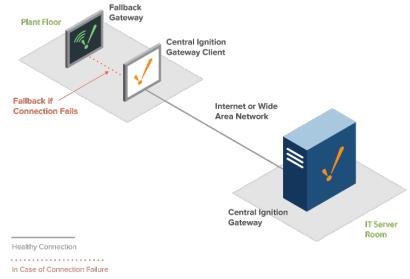
Populating the Cache

Users and Roles can not be manually added to the Fallback Cache. Instead, they are automatically copied from remote Gateways. This type of User Source is normally configured on an [Edge Gateway](#), but can be utilized on an Ignition Gateway.

- Before the Fallback Cache will populate, both the **Central Ignition Gateway** and **Fallback Gateway** must be connected over the [Gateway Network](#). The credentials are passed from a User Source on the **Central Ignition Gateway** to the Fallback Cache over the Gateway Network.
- A Fallback Cache User Source must exist on the **Fallback Gateway**.
- A client must be launched on the **Fallback Gateway** from the **Central Ignition Gateway**. If the user successfully authenticates against the **Central Ignition Gateway's** User Source, then the credentials are cached into the Fallback Cache.

On this page ...

- [Creating the Fallback Cache User Source](#)
- [Populating the Cache](#)



Related Topics ...

- [User Source](#)

Verify a User on a User Source

You can verify that a user exists in a given [User Source](#), if the password is correct, what roles a user has, and any other information added about the user.

Verify a User

1. On the Gateway Webpage, go to the **Config** tab.
2. Choose **Security > Users, Roles** from the menu on the left.
The **User Sources** page is displayed.
3. Select the **Verify an Authentication Profile** link.

The screenshot shows a table of user sources. The first row is 'default' (Internal) with a description: 'This is the default and always present internal authentication profile.' It has a 'manage users' button and an 'edit' button. The second row is 'opcua-module' (Internal) with a description: 'OPC-UA clients will authenticate against this profile by default.' It has a 'More' button and an 'edit' button. Below the table are two links: 'Create new User Source...' and 'Verify an Authentication Profile...'. The 'Verify an Authentication Profile...' link is highlighted with a red border.

4. The Verify Authentication Profile window will appear. Choose a **Profile** from the dropdown list. Enter a **Username** and **Password**. Click **Test Login**.
Ignition will test the credentials then display the results of the validity test.

The screenshot shows the 'Verify Authentication Profile' window. A message at the top says: '💡 Use this page to test your Authentication Profile's settings.' Below is a 'Test Login' form with the following fields:

- Profile: default
- Username: admin
- Password: (redacted)

A blue 'Test Login' button is at the bottom right. The entire window has a light gray background.

If the Login is incorrect, an error message will appear stating that the Login failed for a specified user.

If the Login is correct, a successful message will appear for a specified user along with their information.

The screenshot shows the 'Login Successful' window. It displays a green success message: '✓ Login succeeded for user "admin"' and a list of user details:

- User's roles:**
 - Username: admin
 - First Name: Admin
 - Last Name:
 - Schedule: Always
 - Operator
 - Administrator
 - Contact Info
 - email / admin@inductiveautomation.com
 - phone / 800-266-7798
 - Extended Properties
- [Test again...](#)

On this page ...

- [Verify a User](#)



Verifying an Authentication Profile

[Watch the Video](#)

Identity Provider Authentication Strategy

An Identity Provider (IdP) offers a way for users to log in to Ignition using credentials stored outside of Ignition. An IdP creates, maintains, and manages identity (login) information while providing authentication services to Ignition. This provides a secure login that allows Ignition to use SSL and two-factor authentication (2FA). Note that launching a project from an IdP-initiated SSO response is not currently supported.

Identity Providers (IdPs) offer user authentication as a service. An IdP creates, maintains, and manages identity information for principals while providing authentication services to relying party applications within a federation or distributed network. Authentication of the user is handled by the IdP. Ignition can connect to these three different types of IdPs:

- Ignition's internal IdP
- OpenID Connect 1.0
- Security Assertion Markup Language (SAML)

Your organization's IT may have some sort of existing integration with an Identity Provider. Some popular Identity Providers are listed below.

- Ping Identity
- Okta
- Active Directory Federation Services
- Duo

IdPs are set up at the Gateway level. [Security Levels](#) are also set through the Gateway. The Security Levels enable you to define a hierarchy of access inside a Perspective Session.

The following feature is new in Ignition version **8.1.0**
[Click here](#) to check out the other new features

As of release 8.1, Identity Providers can also be used in the Vision module, the Designer, and on the Gateway. The Identity Provider strategy redirects the user to their IdP in their web browser in order to authenticate. The System Identity Provider setting controls which Identity Provider the user is redirected to.

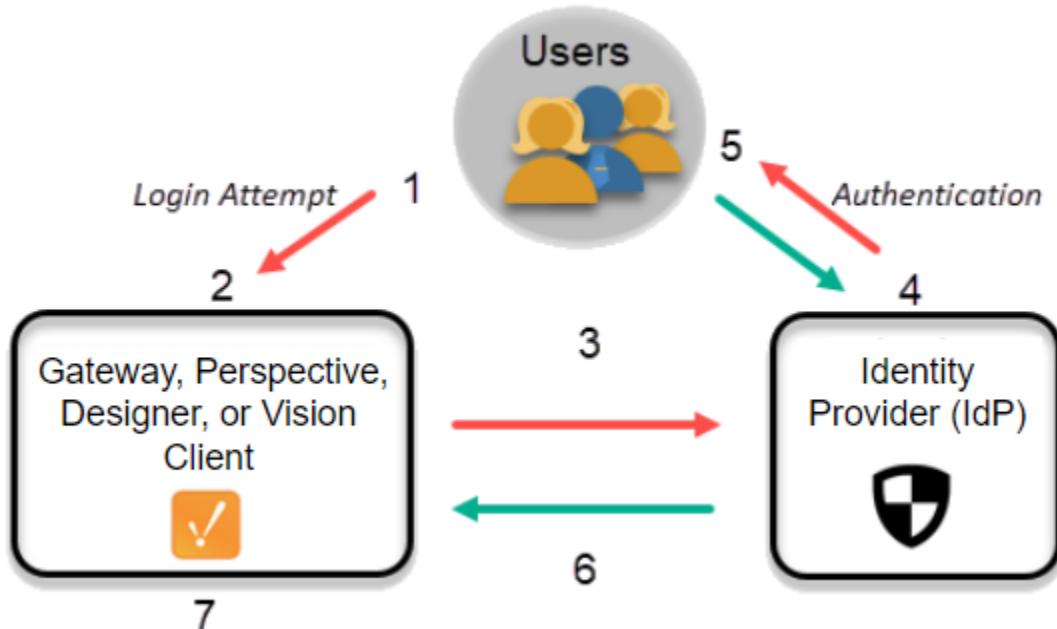
Note: If your browser is not supported, you will get an error message.

On this page ...

- [Identity Provider Authentication Workflow](#)
- [Using Identity Providers](#)
 - [Auth Token Connection Recovery](#)

Identity Provider Authentication Workflow

The following diagram illustrates how IdP authentication works.



1. User make a login attempt to the Gateway, Perspective, the Designer, or a Vision Client.
2. Ignition sees that IdP authentication is required.
3. Ignition redirects the User to Identity Provider for authentication of their credentials:
4. IdP Authenticates the User: The IdP prompts the user with a security challenge, such as requesting a username and password. The extent of the challenge depends entirely on the provider, but many providers may offer support for multi-factor authentication (MFA).
5. User Responds: The user correctly responds to the security challenge.
6. Redirect back to Ignition: If the IdP successfully validates the user, it will redirect the user back to the Perspective Session. Some IdPs may have an additional workflow they will guide the user through, such as re-verifying an email address or replacing an expired password. The IdP will also return information about that user to the Session. This provides some context about the user that the Session can use to assign Security Levels.
7. Update the User's Security Level: Once back at the session, the user will be mapped to the specified Security Level, giving the user access to the restricted action.

Using Identity Providers

The first step in using Identity Providers is to configure them. For the steps for configuring Internal Ignition IdP, OpenID Connect 1.0, or Security Assertion Markup Language (SAML), go to [Configuring Identity Providers](#).

Once an Identity Provider has been configured, there are a few things that can be done to test and adjust how it works. You can [map the attributes](#) that are returned in the IdP response document to more familiar user properties that are available to use within the project. You can add [rules to custom security levels](#) that determine when a user falls into the level. Overrides can be given to users in the form of [User Grants](#), so that they are granted certain security levels regardless of the rules. Finally, you can [test out the IdP](#) by logging in with a user to confirm what is returned in the response document.

The screenshot shows the Ignition Config interface with the following details:

- Left Sidebar (Config Tab):**
 - Home
 - Status
 - Config** (selected)
 - Overview
 - Backup/Restore
 - Ignition Exchange
 - Licensing
 - Modules
 - Projects
 - Redundancy
 - Gateway Settings
 - Networking**
 - Web Server
 - Gateway Network
 - Email Settings
 - Security**
 - General
 - Auditing
 - Users, Roles
 - Service Security
 - Identity Providers** (selected)
 - Security Levels
- Top Header:** Config > Security > Identity Providers
- Table Headers:** Filter, type to filter, View 20, Name, Type, Description, Action
- Data Rows:**

Name	Type	Description	Action
default	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "default".	More ▾ Settings
Okta_2018	OpenID Connect 1.0		More ▾ Settings
- Buttons:** Create new Identity Provider..., Import Identity Provider...
- Pagination:** 1 of 1

Auth Token Connection Recovery

The following feature is new in Ignition version **8.1.24**
[Click here](#) to check out the other new features

After logging into the IdP, a special auth token is generated with the session on the Gateway and is saved in the Designer and Vision Client instance memory after authenticating with an IdP. If a connection is lost and later recovered, Designers and Vision Client instances may securely resume their sessions without having to completely restart by passing the Gateway a valid auth token. Note that auth tokens are not included in Gateway Backups. Any existing auth tokens are cleared when a Gateway Backup is restored.

You can further configure auth tokens by adjusting settings that control the auth token lifecycle. To see these settings, make sure **Identity Provider** is selected as the Authentication Strategy as these settings do not apply to the Classic Authentication Strategy.

- User Inactivity Timeout: The number of minutes which must elapse before expiring a user's auth token due to inactivity caused by a disconnected session. Must be greater than zero. Default: 10 minutes.
- Time-To-Live (TTL): The maximum number of minutes a user's auth token may exist before it expires. If set to any number less than or equal to zero, auth tokens will not expire, as long as the auth token has not expired due to inactivity. Default: 0 minutes (does not expire).

For Designer Auth Tokens, these settings can be found on the [Gateway General Security Settings](#) page by navigating to **Gateway > Config > Security > General**.

General Gateway Security Settings

System Identity Provider *: default

Designer Authentication Strategy *: Identity Provider

Designer Auth Token Inactivity Timeout *: 10

Designer Auth Token Time-To-Live *: 0

For Vision Client auth tokens, these settings can be found in the Designer by opening the **Project Menu**, selecting **Project Properties** and navigating to **Vision > Login**.

Project Properties

Vision / Login

Authentication Strategy: Identity Provider

Login Screen

Welcome Message

Welcome Image

Login Button Text: Login

Show Locale Selector: Automatic

Login Prompt: Log In to Continue

Login Message: Your security is our top priority. Click below to log into your account in the browser.

Auth Token

Inactivity Timeout: 10

Time-To-Live (TTL): 0

Buttons: OK, Apply, Cancel

When redundancy is enabled, Vision Client auth tokens are synchronized from the Master to the Backup so that IdP-authenticated Vision Client sessions may be resumed seamlessly during failover by using an auth token.

[In This Section ...](#)

Configuring Identity Providers

Registering the Ignition Gateway

Before configuring an Identity Provider on the Ignition Gateway, it must first be registered as an Identity Provider Client. Your Identity Provider will have a workflow to register, and it will most likely request something called a **return URL** or **redirect URI**. The paths provided utilize your Gateway's address /hostname, and they change depending on the type of provider.

Note: The same redirect URI is used for login and logout.

OpenID Connect Providers (OP):

OpenID Connect Providers

`http://yourGatewayAddress:Port/data/federate/callback/oidc`

SAML Providers:

SAML Providers

`http://yourGatewayAddress:Port/data/federate/callback/saml`

Secure Integration with IdPs

You should always use the secure versions of those redirect URIs (`https`) in production environments. To do this you must [enable SSL/TLS](#) in Ignition and install a valid certificate. This is the best practice for maintaining a secure integration with third party Identity Providers.

On this page ...

- [Registering the Ignition Gateway](#)
- [Secure Integration with IdPs](#)
- [Configure an Identity Provider](#)
- [Common Properties](#)
- [Ignition Identity Provider](#)
 - [Badge \(RFID\) Settings](#)
 - [Built-In Attributes](#)
 - [Remember Me Example](#)
- [OpenID Connect Providers](#)
 - [Importing Metadata from the Provider](#)
 - [Configuring the Provider](#)
 - [JSON Web Key Configuration](#)
- [Security Assertion Markup Language \(SAML\) Providers](#)
 - [Importing Metadata from the Provider](#)
 - [Configuring the Provider](#)
 - [SAML Signature Verifying Key Configuration](#)
- [Import Provider Metadata for Redundant Backup](#)
- [IdP Examples and Troubleshooting](#)



INDUCTIVE
UNIVERSITY

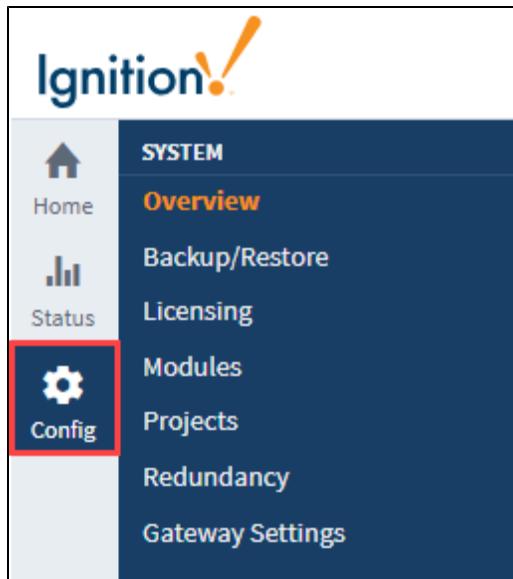
Configuring Identity Providers

[Watch the Video](#)

Configure an Identity Provider

Although there are several types, the general workflow for creating an Identity Provider is the same.

1. On the Gateway Webpage, click on the **Config** tab.



2. Under the Security section, click on **Identity Providers**. The Identity Providers screen is displayed. This screen will list all IdPs that have been configured. You can change filter by name or adjust the number of IdPs displayed in the view.

This screenshot shows the 'Identity Providers' page within the Ignition configuration. The left sidebar includes 'Config' under the SECURITY section, which is highlighted with a red box. The main area displays a table of configured identity providers with columns for Name, Type, Description, and Action. Two entries are listed: 'Default_User_Source' (Ignition) and 'Okta_2018' (OpenID Connect 1.0). At the bottom of the table, there are two buttons: 'Create new Identity Provider...' and 'Import Identity Provider...', both of which are highlighted with red boxes.

3. Click on **Create New Identity Provider...**
4. Choose the type of provider. The current options are **Ignition**, **OpenID Connect 1.0** or **Security Assertion Markup Language 2.0 (SAML)**.

This screenshot shows the 'Create New' dialog for an identity provider. The left sidebar shows 'Config' under SECURITY. The main area has three radio button options: 'Ignition' (selected), 'OpenID Connect 1.0', and 'Security Assertion Markup Language 2.0'. Each option has a brief description below it. At the bottom of the dialog are 'Cancel' and 'Next' buttons, with 'Step 1 of 2' centered between them.

5. Click the **Next** button.

- Configure the adapter. This step varies based on the type of provider. Please see the reference tables below for a description of properties.
- Once you've filled in the properties, click **Save**.

Common Properties

All Identity Provider types share the following properties:

Property Name	Description
Provider Name	<p>The name of the adapter. Adapter names must be unique, so no two adapters on the same Gateway may have the same name.</p> <p>The naming conventions for IdPs are as follows:</p> <ul style="list-style-type: none"> IdP names must begin with an underscore or alpha character. The remaining characters in the name must be either underscores or alphanumeric. IdP names are not case sensitive.
Provider Description	A description of the provider.
Provider Type	The type of Identity Provider. The value for this field comes from the previous screen. It cannot be changed here.

Ignition Identity Provider

The Ignition Identity Provider has the following properties:

Property Name	Description									
User Source	The User Source for this IdP. In order to properly authenticate users, the Ignition Identity Provider must be able to query the list of users from the underlying user source profile.									
Session Inactivity Timeout	<p>The following feature is new in Ignition version 8.1.0 Click here to check out the other new features</p> <p>The number of minutes which must elapse before expiring a session due to user inactivity. Sessions will not timeout if set to any number less than or equal to zero.</p>									
Session Expiration	<p>The following feature is new in Ignition version 8.1.0 Click here to check out the other new features</p> <p>The maximum number of minutes a session may exist before it is expired. Sessions will not have a max lifetime if set to any number less than or equal to zero.</p>									
Remember Me Expiration	<p>The following feature is new in Ignition version 8.1.0 Click here to check out the other new features</p> <p>The maximum number of hours a user will be remembered if they elect to be remembered. Remember Me is disabled when this value is set to any number less than or equal to zero. For more information on this option, see the Remember Me section below.</p>									
Authentication Methods	You can opt into Badge based authentication for the IdP by enabling the "Badge" Authentication Method. The "Default" radio button determines which option users first see when attempting to authenticate against the IdP.									
	<table border="1"> <thead> <tr> <th>Name</th> <th>Enabled</th> <th>Default</th> </tr> </thead> <tbody> <tr> <td>Username and Password</td> <td><input checked="" type="checkbox"/></td> <td><input checked="" type="radio"/></td> </tr> <tr> <td>Badge</td> <td><input type="checkbox"/></td> <td><input type="radio"/></td> </tr> </tbody> </table>	Name	Enabled	Default	Username and Password	<input checked="" type="checkbox"/>	<input checked="" type="radio"/>	Badge	<input type="checkbox"/>	<input type="radio"/>
Name	Enabled	Default								
Username and Password	<input checked="" type="checkbox"/>	<input checked="" type="radio"/>								
Badge	<input type="checkbox"/>	<input type="radio"/>								
Badge Secret	Choose whether or not the user is required to enter a secret (password) along with their badge scan. Additional option can be checked to require the user to enter their password in addition to scanning their badge.									

Badge (RFID) Settings

Property Name	Description	Required?
Badge Secret	Choose whether or not the user is required to enter a secret (password) along with their badge scan. Check to require the user to enter their password in addition to scanning their badge. Default is false (not checked).	Yes

Built-In Attributes

The following attributes are available in the Ignition IdP.

Attribute	Type	Description	Example
auth_time	Date	Represents the time the user last authenticated.	<pre>// Check if it has been within 15 minutes since the last // authentication attempt dateDiff ({idp-attributes: auth_time}, now(), "minutes") <= 15</pre>
challenged	Boolean	<p>Signifies if the user provided credentials at the last login.</p> <p>If true, then the user was asked to re-validate their credentials the last time they attempted to login.</p> <p>If false, then they were not challenged to re-validate their credentials during the last login attempt. This can happen when a login request was made after a user was already authenticated. For example, if a user was already authenticated in a Perspective Session, and a separate call to <code>system.perspective.login</code> function was made with the <code>forceAuth</code> parameter set to false, meaning the user did not provide credentials during the last authentication challenge.</p>	<pre>// Returns True or False, depending on whether or not the user // provided credentials at the last login. {idp-attributes: challenged}</pre>

Remember Me Example

The following feature is new in Ignition version **8.1.0**
[Click here](#) to check out the other new features

The Remember Me option allows your login to be remembered for a set amount of time, even if you close your browser or restart your Gateway. When set, you will be remembered on this device for the specified number of hours without needed to log in again.

Note:

The following items will prevent this feature from working successfully:

- Updating the Ignition Gateway to a new version will cause the device to "forget" the remembered user for some update versions, namely 8.1.2.
- Checking the **Always ask the IdP to re-authenticate users by default** option for the **General Gateway Security Settings > System Identity Provider** setting will still require re-authentication regardless of Remember Me time allotments.

Caution: This option is not recommended if you are using a public or shared device.

To set up Remember me, do the steps that follow:

1. On the Gateway Webpage, click on the **Config** tab. Scroll down to **Security > Identity Providers**.
2. For the Ignition Identity Provider you'd like to configure, click on the **More** option and choose **Settings**.

The screenshot shows the Ignition Configuration interface under the Security > Identity Providers section. A table lists one provider: 'default' (Ignition type). The 'More' button for this provider is highlighted with a red box. Below the table, a context menu is open for the 'default' provider, with the 'Settings' option highlighted with a red box. Other options in the menu include 'User Attribute Mapping' and 'User Grants'.

3. On the Settings page, scroll down to the **Provider Configuration** section.
4. For the Remember Me Expiration option, enter a value greater than zero. For this example, we set the option to **two** hours.

The screenshot shows the 'Provider Configuration' settings for the 'default' provider. The 'Remember Me Expiration' field is highlighted with a red box and contains the value '2'. The 'Authentication Methods' section shows 'Username and Password' as enabled and the default method.

Name	Enabled	Default
Username and Password	<input checked="" type="checkbox"/>	<input checked="" type="radio"/>
Badge	<input type="checkbox"/>	<input type="radio"/>

5. Click Save to save your changes.

To enable Remember Me for your login, do a test login:

1. On the Gateway Webpage, click on the **Config** tab. Scroll down to **Security > Identity Providers**.
2. For the Ignition Identity Provider you'd like to configure, click on the **More** option and choose **Test Login**.
3. Enter your password and select the Remember Me option.

4. Click the **Continue** button.
5. Your login will now be remembered for the amount of hours that were specified in the Gateway setting (in this example, it is 2 hours).

OpenID Connect Providers

OpenID Connect Providers (OP) properties are listed in the following tables. The values on many of these properties may require you to refer to information from your third-party IdP.

Importing Metadata from the Provider

This method is preferred because of its ease-of-use and accuracy. After importing, you will only need to add your client ID and secret manually. (However you can revise the imported data if needed as well.)

Property Name	Description
Import from URL	URL to the OpenID Provider Configuration document. Typically, if the issuer is "https://example.org/foo" then the metadata URL would be "https://example.org/foo/.well-known/openid-configuration"
Import From File	File must be a JSON document with the properties described in section 3 (OpenID Provider Metadata) of the OpenID Connect Discovery 1.0 specification.

Configuring the Provider

Most OpenID Providers will require registering Ignition as a client. After the registration process is complete, the provider will generate a client ID and secret for Ignition, which is required below. This gives Ignition the ability to communicate securely with the provider. Most providers will also require a set of redirect URIs. An example redirect URI would look like: `http://hostname:port/data/federate/callback/oidc`

Property Name	Description
Client ID	The client identifier registered within the identity provider. This value is provided by the Identity Provider.
Client Secret	The client secret registered within the identity provider. This value is provided by the Identity Provider.
Authorization URL	URL of the OP's OAuth 2.0 Authorization Endpoint.
Token URL	URL of the OP's OAuth 2.0 Token Endpoint.

Logout URL	Optional URL at the OP to which an RP can perform a redirect to request that the end user be logged out at the OP.
JSON Web Keys URL	URL of the OP's JSON Web Key Set document.
Use Json Web Keys URI	If checked, then identity provider public keys will be automatically downloaded from given JSON Web Keys URL. New keys will be automatically fetched when the identity provider generates new keys. If unchecked, then the static set of JSON Web Keys (configured below) are used, so when the identity provider rotates keys, they must be manually added to this configuration.
User Info URL	Optional URL to retrieve UserInfo claims from the provider. Resulting claims are typically determined by the scopes listed under the Scope setting.
User Info HTTP Request Method	<p>The following feature is new in Ignition version 8.1.27 Click here to check out the other new features</p> <p>The HTTP method used for sending User Info API requests to the User Info URL. Available options are POST and GET.</p>
Issuer	Entity that issues a set of claims.
Supported ID Token Signing Algorithm Values	A list of the JSON Web Signature (JWS) signing algorithms supported by the OP for the ID Token to encode the claims in a JWT.
Scope	A list of scopes which will be sent for each auth request to the OP. Commonly used scopes would be <code>email</code> and <code>profile</code> but check your Identity Provider's documentation for more information.
JSON Web Key Config	A list of signing key(s) the RP uses to validate signatures from the OP.

JSON Web Key Configuration

Property Name	Description
Key Type	The cryptographic algorithm family used with the key. Options are EC, RSA or oct.
Public Key Use	The intended use of the public key. Options are sig or eng.
Key Operations	The operation(s) for which the key is intended to be used.
Algorithm	The algorithm intended for use with the key.
Key ID	Used to match a specific key.
X.509 URL	A URI that refers to a resource for an X.509 public key certificate or certificate chain. The identified resource MUST provide a representation of the certificate or certificate chain that conforms to RFC 5280 in PEM-encoded form, with each certificate delimited as specified in Section 6.1 of RFC 4945.
X.509 Certificate Chain	The "x5c" (X.509 certificate chain) parameter contains a chain of one or more PKIX certificates. Each entry must be a base64-encoded (Section 4 of RFC4648 -- not base64url-encoded) DER PKIX certificate value.
X.509 Certificate SHA-1 Thumbprint	A base64url-encoded SHA-1 thumbprint (a.k.a. digest) of the DER encoding of an X.509 certificate.
X.509 Certificate SHA-256 Thumbprint	A base64url-encoded SHA-256 thumbprint (a.k.a. digest) of the DER encoding of an X.509 certificate.

There are some additional properties, that depend on which Key Type is selected.

Key Type: EC

Property Name	Description
crv (Curve)	The cryptographic curve used with the key.
x (X Coordinate)	The x coordinate for the Elliptic Curve point represented as the base64url encoding of the octet string representation of the

	coordinate.
y (Y Coordinate)	The y coordinate for the Elliptic Curve point represented as the base64url encoding of the octet string representation of the coordinate.
d (ECC Private Key)	The Elliptic Curve private key value represented as the base64url encoding of the octet string representation of the private key value.

Key Type: RSA

Property Name	Description
n (Modulus)	The modulus value for the RSA public key represented as a Base64urlUInt-encoded value.
e (Exponent)	The exponent value for the RSA public key represented as a Base64urlUInt-encoded value.
d (Private Exponent)	The private exponent value for the RSA public key represented as a Base64urlUInt-encoded value.
p (First Prime Factor)	The first prime factor represented as a Base64urlUInt-encoded value.
q (Second Prime Factor)	The second prime factor represented as a Base64urlUInt-encoded value.
dp (First Factor CRT Exponent)	The Chinese Remainder Theorem (CRT) exponent of the first factor represented as a Base64urlUInt-encoded value.
dq (Second Factor CRT Exponent)	The CRT exponent of the second factor represented as a Base64urlUInt-encoded value.
qi (First CRT Coefficient)	The CRT coefficient of the second factor represented as a Base64urlUInt-encoded value.
oth (Other Primes Info)	Information about any third and subsequent primes, should the exist. Each new Prime added will provide users with new Prime Factor, Factor CRT Exponent, and Factor CRT Coefficient properties, all of which are required.

Key Type: oct

Property Name	Description
k (Key Value)	The value of the symmetric (or other single-values) key represented as the base64url encoding of the octet sequence containing the key value.

Import Provider Metadata for Redundant Backup

The following feature is new in Ignition version **8.1.11**
[Click here](#) to check out the other new features

See the [Import Provider MetaData for Redundant Backup](#) section below.

Security Assertion Markup Language (SAML) Providers

The properties for Security Assertion Markup Language (SAML) are listed in the following tables. The values on many of these properties may require you to refer to information from your third-party IdP.

Importing Metadata from the Provider

This method is preferred because of its ease-of-use and accuracy. After importing, you will have the opportunity to revise the imported data if needed before saving.

Property Name	Description
Import from URL	URL to the SAML Identity Provider Metadata document.
Import From File	File must be an XML document which conforms to the SAML 2.0 metadata schema described in saml-metadata-2.0-os.

The SAML Service Provider (SP) metadata for an Ignition Gateway may be accessed at the following URL: <http://<ipaddress>:<port>/data/saml/metadata/sp>.

The Assertion Consumer Service (ACS) URL for this Ignition Gateway is: <http://<ipaddress>:<port>/data/federate/callback/saml>

Both of these addresses assume you know the IP Address and port of your Ignition install. For example, if you are on the computer Ignition is installed on, you could use: <http://localhost:8088/data/saml/metadata/sp> for the SP metadata.

The following feature is new in Ignition version 8.1.1

[Click here](#) to check out the other new features

SAML IdPs may send the Base64-encoded SAML Response in a line-wrapped form (with new line characters such as \r and \n separating each line). As of release 8.1.1, Ignition's 2.0 SP implementation can handle both line-wrapped and non-line-wrappedBase64-encoded SAML responses.

Configuring the Provider

Property Name	Description
IdP Entity ID	The Identity Provider's Entity ID.
SP Entity ID	The Service Provider's Entity ID. In this case, the Service Provider is the current Ignition Gateway. By default, the Identity Provider will automatically generate the SP Entity ID based on the hostname that the client uses to connect to this Gateway. You may opt out of this setting and input a different SP Entity ID for the Gateway.
Assertion Consumer Service (ACS) Binding	The expected binding used by the Identity Provider when interacting with Ignition's Assertion Consumer Service.
Name ID Format	The expected name ID format for subjects of assertions resulting from Authn Requests. Options are UNSPECIFIED, EMAIL_ADDRESS, X509_SUBJECT_NAME, WINDOWS_DOMAIN_QUALIFIED_NAME, KERBEROS_PRINCIPAL_NAME, ENTITY_IDENTIFIER, PERSISTENT_IDENTIFIER, TRANSIENT_IDENTIFIER.
Single Sign-On (SSO) Service URL	The Identity Provider's Single Sign-On (SSO) Service URL.
Single Sign-On (SSO) Service Binding	The binding Ignition will use for sending Authn Requests to the Identity Provider's Single Sign-On (SSO) Service.
Force Authn	Check this box to force complying Identity Providers to authenticate the user each time instead of relying on a previous security context. See section 3.4.1 of saml-core-2.0-os for more details.
Validate Response Signatures	Check this box to validate the signature of the response from the Identity Provider.
Validate Assertion Signatures	Check this box if it is expected that assertions will be signed. Ignition will validate the signatures of each assertion.
IdP Metadata URL	<p>The following feature is new in Ignition version 8.1.16 Click here to check out the other new features</p> <p>URL to the SAML Identity Provider Metadata document.</p> <p>If Use IdP Metadata URL is checked, the identity provider's signature verifying keys and certificates will be automatically downloaded from the given metadata URL. New keys and certificates will be automatically fetched when the identity provider generates them. If unchecked, then the static set of keys and certificates (configured below) are used, so when the identity provider rotates them, they must be manually added to this configuration.</p>

Signature Verifying Keys	A list of signing key(s) that Ignition uses to validate signatures from the IdP.
Signature Verifying Certificates	A base64-encoded DER PKIX certificate value.

SAML Signature Verifying Key Configuration

Property Name	Description
Key Algorithm	The algorithm identifier for this signature verifying key. Options are DSA, RSA, or EC.
Key Value	A base64-encoded DER key value.

Import Provider Metadata for Redundant Backup

The following feature is new in Ignition version **8.1.11**

[Click here](#) to check out the other new features

See the [Import Provider Metadata for Redundant Backup section](#) below.

Import Provider Metadata for Redundant Backup

The following feature is new in Ignition version **8.1.11**

[Click here](#) to check out the other new features

Both OIDC and SAML provider configurations contain a section that allows you to define metadata which will be used by the backup gateway in a redundant Ignition configuration. When a backup gateway becomes the active node, it will interact with the identity provider via the backup settings. By default, the backup meta data will use the same configuration settings that the master node uses, unless "Provider Metadata for Redundant Backup" is set to false, in which case the IdP configuration will allow you to define settings for the backup gateway.

IdP Examples and Troubleshooting

The [OpenID Connect 1.0 Example](#) page will show you how to configure an external IdP that used OpenID Connect 1.0 with your Ignition system. Go to [Troubleshooting Identity Providers](#) for helpful examples to help you diagnose and troubleshoot issues with configuring IdPs.

Refer to [SAML Example](#) page for how to configure an Identity Provider that is using the SAML protocol.

User Attribute Mapping

The User Attribute Mapping page allows you to map information in the Identity Providers (IdP) response document to easily understandable properties. These properties are then made available as [Session Properties](#) in the Perspective Session. To work, this requires that the Gateway already has a valid IdP configuration which returns a response document when attempting to login.

On this page ...

- [Configuring a User Attribute Mapping](#)
- [User Attribute Mapping Property Reference Table](#)
- [Direct Mapping](#)
 - [Attribute Sources](#)
- [Expression Mapping](#)
- [Built-in Expression Objects](#)
 - [Single-Valued Object Example](#)
 - [Multi-Valued Object Example](#)
 - [Tag Path Example](#)

Configuring a User Attribute Mapping

1. From the Gateway Webpage **Config** tab, click on **Security > Identity Providers**.
2. Select the **Identity Provider** and click on **User Attribute Mapping** under the **More** button.

The screenshot shows the Ignition Gateway's Configuration interface. In the top navigation bar, 'Config' is selected. Under 'Security', 'Identity Providers' is chosen. The main content area displays a table of identity providers:

Name	Type	Description	Action
default	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "default".	More Settings
Okta_2018	OpenID Connect 1.0		More Settings

Below the table, there are two links: [Create new Identity Provider...](#) and [Import Identity Provider...](#). To the right of the table is a vertical 'More' dropdown menu. The 'User Attribute Mapping' option is highlighted with a red box.

3. The Name, Description, and Provider Type are not editable here, but are listed on the page to make clear which IdP the User Attribute Mapping is being configured on.

Main		* Required Field
Provider Name *	Okta_2018	Nickname for the provider.
Provider Description	Provider Description	Description for the provider.
Provider Type *	OpenID Connect 1.0	The type of the provider.

4. Under the **User Attributes** section of properties, you'll find settings for each of the mappable user attributes.

Config > Security > Identity Providers > Okta_2018 - User Attribute Mapping

User Attributes

* Required Field

ID *	Type direct direct expression ID Token Claims The name of the attribute source.
	Path Path Path to the attribute to map.
Username *	Type expression The type of mapping to apply for usernames. Expression Expression The source code for the expression used to map the targeted user attribute.
First Name	Type -None- The type of mapping to apply for first names.
Last Name	Type -None- The type of mapping to apply for last names.
Email	Type -None- The type of mapping to apply for email addresses.
Roles	Type -None- The type of mapping to apply for user roles.

Cancel **Save**

5. If you make changes to any of the mappings, be sure to click **Save** when finished.

User Attribute Mapping Property Reference Table

Mappings can be configured on several user attributes.

Attribute	Description	Data Type
ID	The Type of mapping to apply for user IDs.	String
Username	The Type of mapping to apply for usernames.	String
First Name	The type of mapping to apply for first names.	String

Last Name	The type of mapping to apply for last names.	String
Email	The type of mapping to apply for email addresses.	String
Roles	The type of mapping to apply for user roles. Note: While using an Expression type attribute mapping, you can use a runScript call to create a list of roles through scripting. However the script must return an ArrayList, as opposed to Jython collection types: <pre>import java.util.ArrayList as ArrayList roles = ArrayList() roles.add("Operator")</pre>	Java ArrayList of Strings

Each attribute mapping has similar properties, which are listed below

Property Name	Description
Type	The Type of mapping to apply for the attribute. Options are direct or expression .
Source	The name of the attribute source. In cases where the identity provider has several Attribute Sources , this property allows you to specify which source the mapping should use. (for direct type only)
Path	Path to the attribute map based on the selected source. Each node in the path is delimited by a slash character ("/"). (for direct type only)
Expression	The source code for the expression used to map the targeted user attribute. See the Security Level Rules . (for expression type only)

Direct Mapping

Direct mappings require that you enter in the path to an attribute in the response document for the given property. Values in the Identity Provider's response document are dereferenced using the configured paths and are copied to the respective user properties in Ignition. The Source dropdown allows a particular mapping to target a specific object in the response document. The Path property then determines what object inside of source should be.

In the image below, the ID of the user, as represented on the Ignition Gateway, will be determined by the **sub** attribute, located in the **ID Token Claims** source in the response document.

User Attributes

ID *

Type	direct
The type of mapping to apply for user IDs.	
Source	ID Token Claims
The name of the attribute source.	
Path	sub
Path to the attribute to map.	

Attribute Sources

Each Identity Provider has a different list of possible attribute sources. Below are several tables that list each possible source, as well as the expression path name for the source (used by {attribute-source:X:Y} and {multi-attribute-source:X:Y} expression paths. See [Security Level Rules](#).)

IdP Type	Attribute Source(s)
Ignition	ID Token Claims (default)
	Token Endpoint Response
OIDC	ID Token Claims (default)
	Token Endpoint Response
	User Info Claims
SAML	Authentication Response

For Ignition and OIDC IdPs, the correct source to use will vary with the referenced attribute, and is included in the response document on the [Test Login](#) page.

Note: The user info claims attribute source is present only when the user info URL is configured and Ignition receives a valid user info response from the IdP.

Expression Mapping

The Expression type allows you to use the [expression language](#) to derive the mapping from contextual data, such as the IdP response document or Tags. To assist with parsing the response document in an expression, there are several built-in objects.

Username *	Type expression
	The type of mapping to apply for usernames. Expression Expression
	The source code for the expression used to map the targeted user attribute.

Built-in Expression Objects

In IdP contexts, it is possible to reference IdP response document elements using a three-part format:

```
{<Attribute Type>:<Attribute Source>:<Attribute Path>}
```

In this format, "<Attribute Type>" is the attribute type we're attempting to reference, "<Attribute Source>" is the name of the top-level attribute source, and "<Attribute Path>" is a path to the property inside the attribute source. Here, we'll discuss the possible values for each part of this expression, and when they should be used.

Attribute Type

Type	Expression Value	Description
Single-Value Object	attribute-source	Provides the ability to reference an attribute (such as the username) returned from an IdP attribute source where the attribute's value is expected to be a single atomic value, usually a string.
Multi-Value Object	multi-attribute-source	Provides the ability to reference an attribute (such as the user's roles) returned from an IdP attribute source where the attribute's value is expected to be a collection of single atomic values, usually a collection of strings.

Attribute Source

Possible values for the attribute source will vary by IdP type:

IdP Type	Attribute Source(s)	Expression Path Name
Ignition	ID Token Claims (default)	idTokenClaims
	Token Endpoint Response	tokenEndpointResponse
OIDC	ID Token Claims (default)	idTokenClaims
	Token Endpoint Response	tokenEndpointResponse
	User Info Claims	userInfo
SAML	Authentication Response	authnResponse

For Ignition and OIDC IdPs, the correct source to use will vary with the referenced attribute, and is included in the response document on the [Test Login](#) page.

The following feature is new in Ignition version **8.1.5**
[Click here](#) to check out the other new features

Prior to Ignition 8.1.5, the expression path for SAML Authentication Response did not have an explicit name, so the attribute-source expression path would simply default to Authentication Response in all cases. As of 8.1.5, authnResponse must be explicitly stated when attempting to access the Authentication Response source with expression paths. See [Expression Mapping](#).

Attribute Path

The final element in an attribute reference is a path to the attribute within the attribute source. Here, the [Test Login](#) page is helpful for accessing the attribute source's structure in the response document.

If the attribute source is a JSON document (as is the case with Ignition and OIDC IdPs), then a JSON Path may be used to reference the target attribute.

The following feature is new in Ignition version **8.1.5**
[Click here](#) to check out the other new features

If the attribute source is an XML document (as is the case with SAML IdPs), then an [XPath 1.0 expression](#) may be used to reference the target attribute. See the [Multi-Valued Object Example](#) further down.

Single-Valued Object Example

This example makes use of the Single-Valued Object on an JSON response document from an OpenID Connect IdP. The document used by this example is listed in the code block below.

Sample OIDC Response Document

```
{
  "idTokenClaims": {
    "sub": "00u4z166kjWSxpKyH357",
    "name": "Paul Person",
    "email": "person@company.com",
    "ver": 1,
    "iss": "https://573779.okta.com",
    "aud": "0oa2uqi4odPcNHRpS357",
    "iat": 1617831866,
    "exp": 1617835466,
    "jti": "ID.itbv0TrKsyxqJpEsqwkVHaVa07xBD_haRQMnGUQ5tz4",
    "amr": [
      "pwd"
    ],
    "idp": "00ob8o86nj2kN48V4356",
    "nonce": "a2Zq9QA1xL7t_apS58YIPaegbGD9w6FN60qSgP0i7fg",
    "preferred_username": "person@company.com",
    "auth_time": 1617831865,
    "at_hash": "TEbuSuF8aMhWGYDuHsC0jQ"
  },
  "tokenEndpointResponse": {
    "token_type": "Bearer",
    "expires_in": 3600,
    "access_token": ""
  }
}
"eyJraWQiOjJCvlowd3JweUdiRnB5cG5EQ05HbUpnMnhtbw5sc0h3T1BrDvp3MjQxsJdrIiwiYWxnIjoiu1MyNTYifQ.
eyJ2ZXIiOjEsImp0aSI6IkFULmhmaEVRLUFizzIyVU1Bb3NWNm9wWld1NHFGWmIzzUdlR25VVUNhX1lWWCiLCJpc3MiOjJodHRwcovL2Rldi01NzM3Nzkub2t0YS5jb20iLCJhdWQiOjJodHRwcovL2Rldi01NzM3Nzkub2t0YS5jb20iLCJzdWIoiJwc2NvdHRAaW5kdWN0aXZ1YXV0b21hdGlvbi5jb20iLCJpYXQiOjE2MTc4MzE4NjYsImV4cci6MTYxNzgzNTQ2NiwiY2lkIjoimG9hMnVxaTRvZFBjTkhsCFMzNTciLCJ1aWQioiIwmHU0emw2NmtqV1N4cEt5SDM1NyIsInNjcCI6WyJvcGVuaWqiLCJ1bWFpbCISinByb2ZpbGUIXX0.
GR3JotmbNtD7FDV59cuiE3wxsNJAARN_zz5pmtnhLr19BGgiRTKJOnYXBo7E_KBj82RAEKsGRKFLzSDIzaB2QUCPyzJSSNg2LKvht5yoQoguoY6HqgeU73Ukf6gbEhE3c_WfW1reeewN8eaziOnvYJQhmVzYrhlf04XY0cB8o_1raNku3yUw76MDyAlWpaXuYQ2b9UYir3CL1b7AwpYTHdRHak7zdfv713m4aR5IgVt0Pch8pFarX4k3o6Sg4ZghLW-
4eioSzs0CrVLk3j09vQIEKn0jdlKqY0i9_9HBiEGUrVOXuiJ587ocRTjvxeNCpr_B7DBczBbtTkOFNzw3xRhw",
  "scope": "openid email profile",
  "id_token": "eyJraWQiOjJUNEhmZDQ0LXVVQ09sMGJpT19fWnprQkJYVwdSSx1QYXBIZ1lQeFRFTHYwIiwiYWxnIjoiu1MyNTYifQ.
eyJzdWliOiwMHU0emw2NmtqV1N4cEt5SDM1NyIsIm5hbWUiOijQYXVsIFNjb3R0IiwiZW1haWwiOijwc2NvdHRAaW5kdWN0aXZ1YXV0b21hdGlvbi5jb20iLCJ2ZXIiOjEsIm1zcI6Imh0dBzOi8vZGV2LTU3Mzc3OS5va3RhLmNvbSISimF1ZCI6IjBvYTJ1cWk0b2RQY05IUnBTMzU3IiwiWF0IjoxNjE3ODMxODY2LCJleHAIoje2MTc4MzU0NjYsImp0aSI6Ik1ELml0YnYwVHJLc314cUpwRXNxd2tWSGFwyTA3eEJEX2hhU1FnbkdVUTV0ejQjLCJhbXiiolsicHdkIl0sIm1lcCI6IjAwb2I4bzg2bkoya0400FY0MzU2Iiwbm9uY2UiOijhMpxOVFBMxhMN3rfYXBTNThZSVbhZWdir0Q5dzZGTjYwcVNnUDBpN2ZnIiwichJlZmVycmVx3VzzXJuYw11IjoichNjb3R0QGluzHVjdG12ZWFl1dG9tYXRpb24uY29tIiwiYXV0aF90aW11IjoxNjE3ODMxODY1LCJhdF9oYXNoIjoiVEVidvNVZjhhtWhxR11EdUhzQzBqUSJ9.
```

```

b5wi6CH0p1OigL6aaamZu3aVbOHrgiKdHFKqEgU0PLfbTILGtabhLktic843XaBTiQm8s3oXz153P3rdsUP4VPj1jXF8C221iEjP-
55PyBX2kQ-ggyqXVrWlywmOWFNn4kSjViuzMIT2gXeW8yW-
OBtxBTZr7QPDF6XKwZk6V183K0hT3ZbrzHcqZ0juIEkxrrfDU8N8Fo8b094A2rGPP0mfSPttkrxxZBJ3bO6oY5e7NtcBnt4IsbSzfvI3_1KUo
Qt0SylljsescEZGxRJCFnDE3dJyOaQZuBKLN8ouxibD9Jdmn9_raeLGKOsxNHKOpNMRbptJ9UD4VemyezQvza"
},
"userInfo": {
  "sub": "00u4zl66kjWSxpKyH357",
  "name": "Paul Person",
  "locale": "en-US",
  "email": "person@company.com",
  "preferred_username": "person@company.com",
  "given_name": "Paul",
  "family_name": "Person",
  "zoneinfo": "America/Los_Angeles",
  "updated_at": 1596495564,
  "email_verified": true
}
}

```

If, for example, we wanted to configure a mapping on the First Name, which would allow us to retrieve the name of the user provided by the response document. The **given_name** under **userInfo** appears to have the value we're looking for.

We can use the attribute-source object to directly reference the node.

```
{attribute-source:userInfo:given_name}
```

This is equivalent to using a direct mapping type with Source set to "ID Token Claims" and Path set to "name". However, since we're using the expression language, we have access to more tools, which is useful in cases where the attribute value we're looking for isn't as readily accessible.

For example, say the response document did not contain a node that represented just the users first name. However, it does contain a **name** under **idTokenClaims** (line 4) to that has the user's full name. We could access the full name, and then use the expression language's **split** function to split the first and last names.

```
split({attribute-source:idTokenClaims:name}, ' ')[0,0]
```

Multi-Valued Object Example

This example makes use of the Multi-Valued Object on an XML response document from a SAML IdP. Here we can use an [XPath expression](#) to examine the document and reference a specific node. We'll configure an expression type mapping that will retrieve attribute values that we want mapped to a user's Roles. Below we see the results of a [test login](#). Our user does not appear to have any roles mapped.

Test Results - Security Assertion Markup Language 2.0

IdP Response Data

Mapped User Attributes

Security Level Grants

Attribute Name

Attribute Value

ID	person@company.com
Username	person@company.com
First Name	Paul
Last Name	Scott
Email	
Roles	

[Test Login](#)

[Test Logout](#)

If we click on the **IdP Response Data** tab, we see the results listed in the code block below.

Sample SAML Response Document

```
<?xml version="1.0" encoding="UTF-8" standalone="no"?>
<saml2p:Response xmlns:saml2p="urn:oasis:names:tc:SAML:2.0:protocol" xmlns:xs="http://www.w3.org/2001/XMLSchema" Destination="http://localhost:8088/data/federate/callback/saml" ID="id508417152055433477622124" InResponseTo="A-tiacipjeAZM4e4xHRRgFkV9Kr0x5Mv8tHrdowuL7zi0" IssueInstant="2021-04-06T21:42:33.776Z" Version="2.0">
  <saml2:Issuer xmlns:saml2="urn:oasis:names:tc:SAML:2.0:assertion" Format="urn:oasis:names:tc:SAML:2.0:nameid-format:entity">http://www.okta.com/exkj0b0vtttYvPzH10h7</saml2:Issuer>
    <ds:Signature xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
      <ds:SignedInfo>
        <ds:CanonicalizationMethod Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
        <ds:SignatureMethod Algorithm="http://www.w3.org/2001/04/xmldsig-more#rsa-sha256" />
        <ds:Reference URI="#id508417152055433477622124">
          <ds:Transforms>
            <ds:Transform Algorithm="http://www.w3.org/2000/09/xmldsig#enveloped-signature" />
            <ds:Transform Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
            <ec:InclusiveNamespaces xmlns:ec="http://www.w3.org/2001/10/xml-exc-c14n#" PrefixList="xs" />
          </ds:Transforms>
        </ds:Reference>
      </ds:SignedInfo>
    </ds:Signature>
    <saml2p:Status>
      <saml2p:StatusCode Value="urn:oasis:names:tc:SAML:2.0:status:Success" />
    </saml2p:Status>
    <saml2:Assertion xmlns:saml2="urn:oasis:names:tc:SAML:2.0:assertion" ID="id508417152165556256586917" IssueInstant="2021-04-06T21:42:33.776Z" Version="2.0">
      <saml2:Issuer Format="urn:oasis:names:tc:SAML:2.0:nameid-format:entity">http://www.okta.com/exkj0b0vtttYvPzH10h7</saml2:Issuer>
        <ds:Signature xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
          <ds:SignedInfo>
            <ds:CanonicalizationMethod Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
            <ds:SignatureMethod Algorithm="http://www.w3.org/2001/04/xmldsig-more#rsa-sha256" />
            <ds:Reference URI="#id508417152165556256586917">
```

```

<ds:Transforms>
    <ds:Transform Algorithm="http://www.w3.org/2000/09/xmldsig#enveloped-signature"/>
    <ds:Transform Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#">
        <ec:InclusiveNamespaces xmlns:ec="http://www.w3.org/2001/10/xml-exc-c14n#">
PrefixList="xs" />
        </ds:Transform>
    </ds:Transforms>
    <ds:DigestMethod Algorithm="http://www.w3.org/2001/04/xmlenc#sha256"/>
    <ds:DigestValue>xNlFDtmRTPXYqSLpruOZoPSq2Lz6q1rPrGiROR6MXdM=</ds:DigestValue>
</ds:Reference>
</ds:SignedInfo>
</ds:Signature>
<saml2:Subject>
    <saml2:NameID Format="urn:oasis:names:tc:SAML:1.1:nameid-format:unspecified">person@company.com</saml2:NameID>
        <saml2:SubjectConfirmation Method="urn:oasis:names:tc:SAML:2.0:cm:bearer">
            <saml2:SubjectConfirmationData InResponseTo="A-tiacipjeAZM4e4xHRRgFkV9KrOx5Mv8tHrdowuL7zi0" NotOnOrAfter="2021-04-06T21:47:33.776Z" Recipient="http://localhost:8088/data/federate/callback/saml"/>
        </saml2:SubjectConfirmation>
    </saml2:Subject>
    <saml2:Conditions NotBefore="2021-04-06T21:37:33.776Z" NotOnOrAfter="2021-04-06T21:47:33.776Z">
        <saml2:AudienceRestriction>
            <saml2:Audience>http://localhost:8088</saml2:Audience>
        </saml2:AudienceRestriction>
    </saml2:Conditions>
    <saml2:AuthnStatement AuthnInstant="2021-04-06T21:26:07.499Z" SessionIndex="A-tiacipjeAZM4e4xHRRgFkV9KrOx5Mv8tHrdowuL7zi0">
        <saml2:AuthnContext>
            <saml2:AuthnContextClassRef>urn:oasis:names:tc:SAML:2.0:ac:classes:PasswordProtectedTransport</saml2:AuthnContextClassRef>
                </saml2:AuthnContext>
            </saml2:AuthnStatement>
            <saml2:AttributeStatement>
                <saml2:Attribute Name="FirstName" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
                    <saml2:AttributeValue xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:type="xs:string">Paul</saml2:AttributeValue>
                </saml2:Attribute>
                <saml2:Attribute Name="LastName" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
                    <saml2:AttributeValue xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:type="xs:string">Scott</saml2:AttributeValue>
                </saml2:Attribute>
                <saml2:Attribute Name="Email" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
                    <saml2:AttributeValue xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:type="xs:string">person@company.com</saml2:AttributeValue>
                </saml2:Attribute>
                <saml2:Attribute Name="Role" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
                    <saml2:AttributeValue xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:type="xs:string"/>Operator</saml2:AttributeValue>
                    <saml2:AttributeValue xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:type="xs:string"/>Supervisor</saml2:AttributeValue>
                </saml2:Attribute>
                <saml2:Attribute Name="Groups" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
                    <saml2:AttributeValue xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xsi:type="xs:string"/>
                </saml2:Attribute>
            </saml2:AttributeStatement>
        </saml2:Assertion>
    </saml2p:Response>

```

In the example document, lines 68 and 69 contain `saml2:AttributeValue` elements, which contain the roles for our user. Starting from the roles, we can work our way back towards the top of the document, but making the following observations.

- We see both elements are under a `saml2:Attribute` element (line 67).

- The saml2:Attribute element is under a saml2:AttributeStatement element (line 57).
- The saml2:AttributeStatement element is under a saml2:Assertion element (line 23).
- Finally, the saml2:Assertion element is under the saml2p:Response element, which seems to be the root for the entire path (line 2).

Thus, we're looking for all text values of all saml2:AttributeValue elements that are under a specific saml2:Attribute element. We see that the saml2SAttribute element we're looking for also has name attribute set to "Role", which we can use to help focus our path. Based on our observations, the path looks like the following:

Note that we start at the root of the saml2p:Reponse element (on the left), and build the rest of the path in reverse order of our observations.

We specified that the saml2:Attribute must contain an attribute by the name of "Name", with a value of "Role", with which filters out the other saml2:Attribute elements in the document (such as the elements with FirstName and LastName).

In addition, we ended the path with text(), which access the text values of the nodes.

```
/saml2p:Response/saml2:Assertion/saml2:AttributeStatement/saml2:Attribute[@Name="Role"]/saml2:AttributeValue  
/text()
```

Writing the Path

Now that we know what our path looks like, we can write an expression. We'll type the following into the "Expression" field for the Roles setting:

```
{multi-attribute-source:authnResponse:/saml2p:Response/saml2:Assertion/saml2:AttributeStatement/saml2:  
Attribute[@Name="Role"]/saml2:AttributeValue/text()}
```

The screenshot shows a configuration interface for mapping user attributes. On the left, there's a sidebar with a 'Roles' section. The main area has two sections: 'Type' and 'Expression'. The 'Type' section shows 'expression' selected. The 'Expression' section contains the XPath '/saml2p:Response/saml2:Assertion/saml2:AttributeStatement/saml2:Attribute[@Name="Role"]/saml2:AttributeValue/text()'. Below the expression, a tooltip-like text says 'The source code for the expression used to map the targeted user attribute.'

If we then **Save**, we can attempt a [test login](#), and check the mapped user attributes, where we see our roles from the response document now appear mapped to our user.

Test Results - Security Assertion Markup Language 2.0

IdP Response Data

Mapped User Attributes

Security Level Grants

Attribute Name

Attribute Value

ID	person@company.com
Username	person@company.com
First Name	Paul
Last Name	Scott
Email	
Roles	[Operator,Supervisor]

[Test Login](#)

[Test Logout](#)

Tag Path Example

In addition, Tag values can be referenced by mappings. Note that the values are only checked when the user logs in. So if the value of any tags used in a mapping change after the user logs in, the result of the mapping won't retroactively change.

```
{[default]A_Folder/A_Tag.value}
```

In the example above, the expression would attempt to return the value of the Tag at path "A_Folder/A_Tag", in the Tag Provider named "default".

OpenID Connect 1.0 Example

This section provides an example of how to connect an Identity Provider that is using the OpenID Connect 1.0 protocol. This example uses the Okta IdP service. Your IdP vendor may differ and the specific links will differ.

Prerequisites

Before you begin configuring Ignition there are some preliminary requirements that need to be done outside of Ignition:

- A configured remote IdP (Okta in this example)
- The metadata file specific to your IdP
- The scope data specific to your IdP
- Login credentials to use as a test

On this page ...

- Prerequisites
 - Configured IdP
 - Metadata File
 - Scope Data
 - Test Login Credentials
 - Configure Ignition Gateway

Configured IdP

An IT department is usually the one to set up and configure a remote IdP. You need a configured remote IdP that is compatible with OpenID Connect 1.0. protocol. At minimum there needs to be an account set up with the IdP, users added to the IdP account, and applications added to the IdP.

Metadata File

You will need the metadata file specific to your IdP. This document defines how to communicate with the IdP. It is usually a web page that allows the metadata file to be exported to a JSON file. Often it is a URL that ends with ".well-known/openid-configuration".

You will need the URL link to this page or a JSON export of this page. For example, if your IdP user login URL is something like this:

<https://dev-123456.oktapreview.com/login/login.htm>

Then the metadata import URL may look like something like this:

<https://dev-123456.oktapreview.com/.well-known/openid-configuration>

Here is an example of part of a metadata file for Okta. Notice that the URL link has "/.well-known/openid-configuration" at the end and is very similar to the login URL. It is recommended to use the URL specific to your IdP. Your IT department may choose to export this JSON file of this page and provide it to you. Either option will work.

The screenshot shows a browser window displaying the JSON configuration for an Okta Identity Provider (IdP). The URL in the address bar is <https://dev-123456.oktapreview.com/.well-known/openid-configuration>. The JSON structure includes fields such as issuer, authorization_endpoint, token_endpoint, userinfo_endpoint, registration_endpoint, jwks_uri, response_types_supported, response_modes_supported, grant_types_supported, subject_types_supported, id_token_signing_alg_values_supported, and scopes_supported.

```
issuer: "https://dev-123456.oktapreview.com"
authorization_endpoint: "https://dev-123456.oktapreview.com/oauth2/v1/authorize"
token_endpoint: "https://dev-123456.oktapreview.com/oauth2/v1/token"
userinfo_endpoint: "https://dev-123456.oktapreview.com/oauth2/v1/userinfo"
registration_endpoint: "https://dev-123456.oktapreview.com/oauth2/v1/clients"
jwks_uri: "https://dev-123456.oktapreview.com/oauth2/v1/keys"
response_types_supported:
  0: "code"
  1: "id_token"
  2: "code id_token"
  3: "code token"
  4: "id_token token"
  5: "code id_token token"
response_modes_supported: [...]
grant_types_supported: [...]
subject_types_supported: [...]
id_token_signing_alg_values_supported:
  0: "RS256"
scopes_supported:
  0: "openid"
  1: "email"
  2: "profile"
  3: "address"
  4: "phone"
  5: "offline_access"
  6: "groups"
```

Note: It is recommended to use the URL specific to your IdP. Your IT dept. may choose to export this JSON file of this page and provide it to you. Either option will work.

Scope Data

When a user is verified by the IdP a lot of the user specific data is not returned in the response file by default (i.e., username, email, firstname, lastname, etc). This user specific data is called the scope, and it can be returned if the **Scope** section of the Ignition Gateway is configured. The list of available scope definitions may be available from your IT department or available from the developer documentation of the IdP you are using.

For this Okta example, the scope data is in the developer notes at <https://developer.okta.com/docs/api/resources/oidc#scope-dependent-claims-not-alwaysReturned>

Test Login Credentials

You need an account specific to the IdP for testing purposes (Okta in this example). To test and verify the IdP account, login to your IdP. For our example, the Okta login page is shown here:

Configure Ignition Gateway

1. On the Gateway Webpage, click on the **Config** tab. You will need to log in if you aren't already.
2. Under the Security section, click on **Identity Providers**. The Identity Providers screen is displayed. This screen will list all IdPs that have been configured. You can filter by name or adjust the number of IdPs displayed per page in the view.



3. Click on the **Create a New Identity Provider...** link.

The screenshot shows the 'Identity Providers' list page. It includes a header with 'Config > Security > Identity Providers', a navigation bar with a 'Filter' input and a 'View' dropdown set to '20', and a table with columns 'Name', 'Type', 'Description', and 'Action'. There are two entries:

Name	Type	Description	Action
default	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "default".	More Settings
Okta_2018	OpenID Connect 1.0		More Settings

At the bottom left, there are two buttons: 'Create new Identity Provider...' (highlighted with a red box) and 'Import Identity Provider...'. Below these buttons is another navigation bar with a 'View' dropdown set to '20'.

4. Select the **OpenID Connect 1.0** option and click **Next**.

The screenshot shows the 'Create New Identity Provider' step 1 screen. It has a header 'Config > Security > Identity Providers > Create New'. The main content area contains three options:

- Ignition**: Describes the instance of Ignition hosting the configuration as an identity provider.
- OpenID Connect 1.0**: Describes OpenID Connect 1.0 (OIDC) as an authentication layer on top of OAuth 2.0.
- Security Assertion Markup Language 2.0**: Describes SAML 2.0 as an open standard for exchanging authentication and authorization data.

At the bottom, there are 'Cancel', 'Step 1 of 2', and 'Next' buttons.

5. On the Basic Details screen, provide an **Provider Name**. You can also add an **Provider Description** if desired. The Provider Type field will fill in automatically from the previous screen.

Config > Security > Identity Providers > Create New

Basic Details		* Required Field
Provider Name *	Okta_Example	
Give the provider a name.		
Provider Description	Test setup for Okta IdP.	
A description for the provider.		
Provider Type *	OpenID Connect 1.0	
The type of the provider.		

6. The next section is **Import Provider Metadata**. In the **Import from URL** section, enter in the URL from earlier that shows the ".well-known/openid-configuration" link specific to your IdP. You can also import a file below if it was provided by your IT department.
7. Click on the **Import** button.

Import Provider Metadata

Import from URL	<input type="text" value="https://dev.123456.oktapreview.com/"/> Import
Import from File	<input type="button" value="Choose File"/> No file chosen File must be a JSON document with the properties described in section 3 (OpenID Provider Metadata) of the OpenID Connect Discovery 1.0 specification. Import

8. Ignition will now generate a URI redirect address for your Ignition server. It is listed just below the "Import Provider Metadata" area of the configuration page.
In our example it is <http://10.10.115.3:8088/data/federate/callback/oidc>. You need to provide this URI to your IdP (usually this means giving it to your IT department).

Note: The URI should be a web address that is accessible from the end user's web browser.

💡 Most OpenID Providers will require registering Ignition as a client. After the registration process is complete, the provider will generate a client ID and secret for Ignition, which is required below. This gives Ignition the ability to communicate securely with the provider. Most providers will also require a set of redirect URIs. This Ignition Gateway's Redundancy Role is set to **Independent**. The redirect URI for this Ignition Gateway is: <http://localhost:8088/data/federate/callback/oidc> 🚀

9. Once you've given your IT department the redirect address, they can add your Ignition server as an application to the IdP. Once they have done this, they can provide you with a "Client ID" and "Client Secret". This is needed for the Ignition Gateway to properly communicate with the IdP.

Note: The IdP can use the same redirect address for the Login, Logout, and Initiate Login.

10. The next section is Provider Configuration. Most of the fields below should now be filled in when you imported the IdP Metadata. Fill in the **Client ID** and **Client Secret** fields with the information obtained from your IdP (or IT department). If you don't know them yet, you can put in bogus values for now and edit them later once the correct values are provided to you.

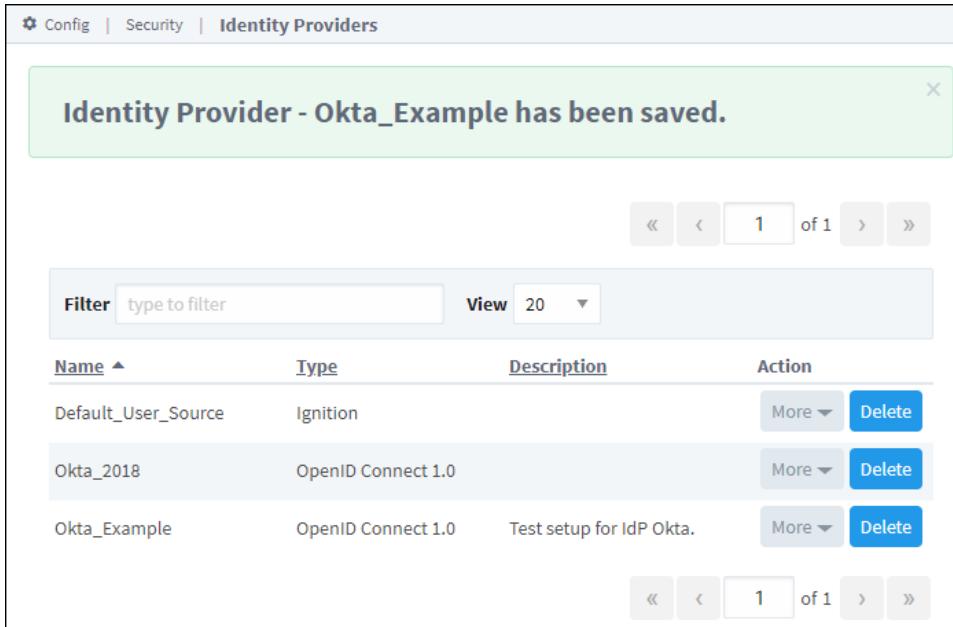
Provider Configuration		* Required Field				
Client ID *	0oaifp0dlnt1234560h7 The client identifier registered within the identity provider.					
Client Secret *	Client Secret The client secret registered within the identity provider. <input type="checkbox"/> Check this box to change the client secret					
Authorization URL *	https://dev-123456.oktapreview.com/ URL of the OP's OAuth 2.0 Authorization Endpoint.					
Token URL *	https://dev-123456.oktapreview.com/ URL of the OP's OAuth 2.0 Token Endpoint.					
Logout URL	https://dev-123456.oktapreview.com/ Optional URL at the OP to which an RP can perform a redirect to request that the End-User be logged out at the OP.					
Json Web Keys URL	https://dev-123456.oktapreview.com/ URL of the OP's JSON Web Key Set document. <input checked="" type="checkbox"/> Use Json Web Keys URI If checked, then identity provider public keys will be automatically downloaded from given Json Web Keys URL. New keys will be automatically fetched when the identity provider generates new keys. If unchecked, then the static set of Json Web Keys (configured below) are used, so when the identity provider rotates keys, they must be manually added to this configuration.					
Issuer *	https://dev-123456.oktapreview.com/ Entity that issues a set of Claims.					
Supported ID Token Signing Algorithm Values	<p>-None-</p> <p>A list of the JWS signing algorithms supported by the OP for the ID Token to encode the Claims in a JWT.</p> <table border="1"> <thead> <tr> <th>Name</th> <th>Action</th> </tr> </thead> <tbody> <tr> <td>RS256</td> <td>Remove</td> </tr> </tbody> </table>		Name	Action	RS256	Remove
Name	Action					
RS256	Remove					
Scope	Scope	Add A list of default scopes which will be sent for each auth request to the OP.				
JSON Web Key Config	<p>A list of signing key(s) the RP uses to validate signatures from the OP.</p> <table border="1"> <thead> <tr> <th>Name</th> <th>Action</th> </tr> </thead> <tbody> <tr> <td>W4IYgV7xggajMkl5T1234561234567Of-SsKViOM_Jal</td> <td>Edit Remove</td> </tr> </tbody> </table> <p>→ Create new JSON Web Key</p>		Name	Action	W4IYgV7xggajMkl5T1234561234567Of-SsKViOM_Jal	Edit Remove
Name	Action					
W4IYgV7xggajMkl5T1234561234567Of-SsKViOM_Jal	Edit Remove					

11. Providing scope is optional. These fields are specific to your IdP, and you may need to find the developer documentation specific to your IdP.

Enter "email" in the **Scope** field and press the **Add** button. Repeat for each scope you want returned. For our example, the list of available scopes is in the Okta developer documentation: <https://developer.okta.com/docs/api/resources/oidc#scope-dependent-claims-not-always-returned>.

Scope	<input type="text" value="email"/> Add A list of default scopes which will be sent for each auth request to the OP.				
	<table border="1" style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th style="background-color: #333; color: white;">Name</th> <th style="background-color: #333; color: white;">Action</th> </tr> </thead> <tbody> <tr> <td>profile</td> <td style="text-align: right;">Remove</td> </tr> </tbody> </table>	Name	Action	profile	Remove
Name	Action				
profile	Remove				

12. Press the **Save** button at the bottom right of the page. You'll see a confirmation message.

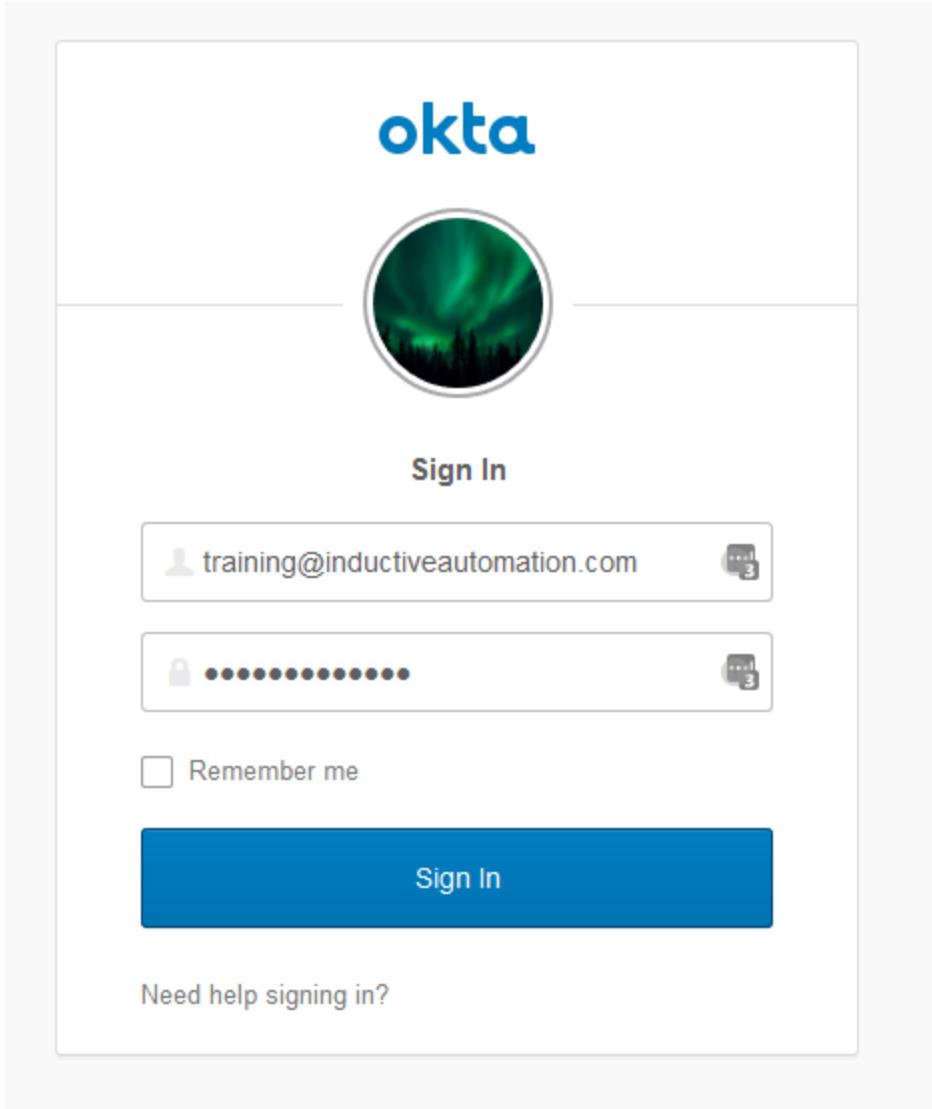


The screenshot shows the Ignition configuration interface under the 'Identity Providers' tab. At the top, there's a navigation bar with 'Config' and 'Security' tabs, and the 'Identity Providers' tab is active. Below the navigation is a green success message box with the text 'Identity Provider - Okta_Example has been saved.' In the center, there's a table listing three identity providers:

Name	Type	Description	Action
Default_User_Source	Ignition		More Delete
Okta_2018	OpenID Connect 1.0		More Delete
Okta_Example	OpenID Connect 1.0	Test setup for IdP Okta.	More Delete

13. The next step is to perform a test login. From the Identity Providers screen, select **More** and then **Test Login**.

14. You will be re-directed to the Okta login. Enter in your test login credentials and click the **Sign In** button.



15. If the login is successful, you will be returned to the Identity Provider Test Login screen. The returned results will be displayed in the Results section.

Note: In this example, we did not add the username or email to our scope. Thus they have not been returned.

Test Login	
Adapter Name	Okta_Example Nickname for the adapter.
Adapter Description	Adapter Description Description for the adapter.
Provider Type	OpenID Connect 1.0 The type of the adapter.
Results	<pre>{ "sub": "00uifxlfn2yjSXQzC0h7", "ver": 1, "iss": "https://dev-997763.oktapreview.com/oauth2/default", "aud": "0oaifp0dlntvhYxaX0h7", "iat": 1546555091, "exp": 1546558691, "jti": "ID.tTIUVnMSyWGzHyCsBf8vf95XdZN9HCzShEps08b2L5U", "amr": ["pwd"], "idp": "00oiezmbtjcSqDHJM0h7", "nonce": "D11EhKChylzTog2i36FNsqU_PEMCod946X1rfyo52mE", "auth_time": 1546555084, "at_hash": "yOe0W~-OT0xdWeYC2s8M8w" }</pre>

Editor notes are only visible to logged in users

Admin Username/password for Training's OKTA instance are available through LastPass. See Bobby.

OKTA USER LOGIN INFO:

Okta Login Page: <https://dev-997763.oktapreview.com/login/login.htm>

Okta ia User: training@ia.io

Okta ia Pass: P4ssw0rd

Okta Metadata Definition: <https://dev-997763.oktapreview.com/.well-known/openid-configuration>

Okta Scope Definition: <https://developer.okta.com/docs/api/resources/oidc#scope-dependent-claims-not-alwaysReturned>

Okta Client ID: 0oaifp0dlntvhYxaX0h7

Okta Client Secret: ZvMy_gt-XqiILB4qQH6AacpsEqA_M2Y8gwkvoj3Q

SAML Example

This section provides an example of how to connect an Identity Provider that is using the SAML protocol. This example uses the Okta IdP service. Your IdP vendor may differ and the specific links will differ.

Prerequisites

Before you begin configuring Ignition there are some preliminary requirements that need to be done outside of Ignition:

- A configured remote IdP (Okta in this example)
- The metadata file specific to your IdP
- The scope data specific to your IdP
- Login credentials to use as a test

On this page ...

- [Prerequisites](#)
 - [Configured IdP](#)
 - [Metadata File](#)
 - [Test Login Credentials](#)
 - [Configure Ignition Gateway](#)

Configured IdP

An IT department is usually the one to set up and configure a remote IdP. You need a configured remote IdP that is compatible with SAML protocol.

At minimum there needs to be an account set up with the IdP, users added to the IdP account, and applications added to the IdP.

Metadata File

You will need the metadata file specific to your IdP. This document defines how to communicate with the IdP. It is usually a web page that allows the metadata file to be exported to an XML file.

You will need the URL link to this page or an XML export of this page. For example, the metadata import URL may look like something like this:

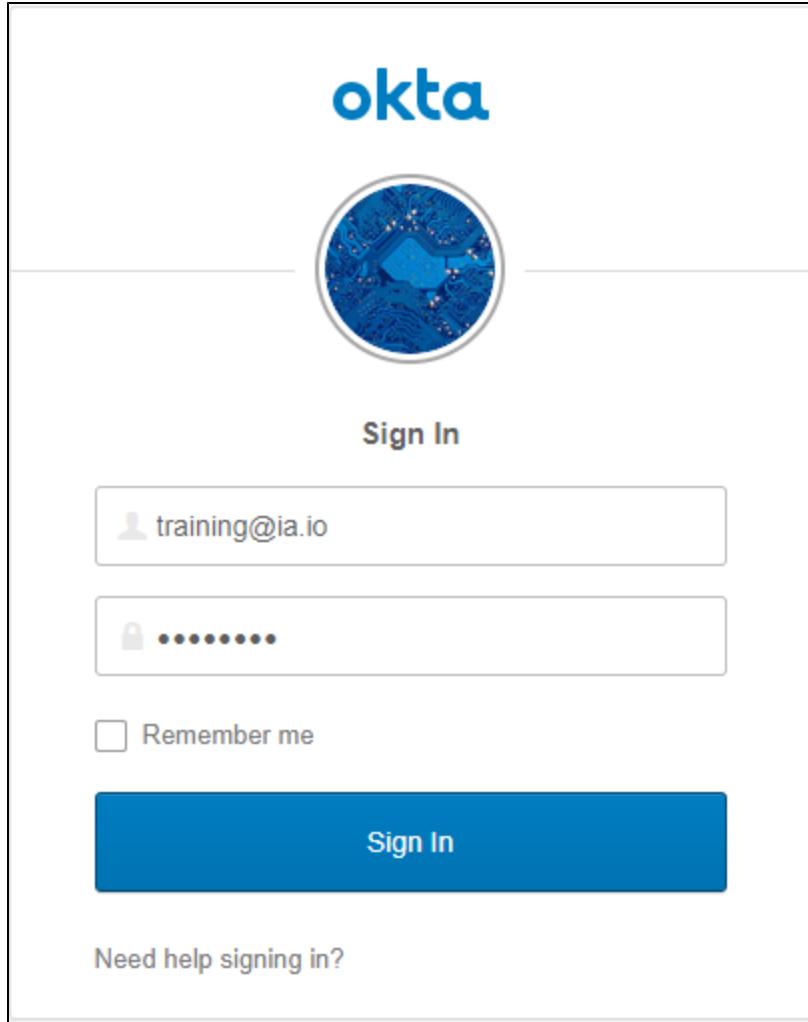
```
https://dev-123456.okta.com/app/esdfdsf7886sd6723hjkdf/sso/saml/metadata
```

Here is an example of part of a metadata file for Okta. Notice that file is XML format. You can use the file or the URL to automatically import the configuration into Ignition. Otherwise it will need to be manually typed in.

```
<?xml version="1.0" encoding="UTF-8"?>
- <md:EntityDescriptor entityID="http://www.okta.com/exkinzujwtqNw0fzI356" xmlns:md="urn:oasis:names:tc:SAML:2.0:metadata">
  - <md:IDPSSODescriptor protocolSupportEnumeration="urn:oasis:names:tc:SAML:2.0:protocol" WantAuthnRequestsSigned="false">
    - <md:KeyDescriptor use="signing">
      - <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
        - <ds:X509Data>
          <ds:X509Certificate>MIIDpDCCAoygAwIBAgIGAWppx7LZMA0GCSqGSIb3DQEBCwUAMIGSMQswCQYDVQQGEwJVUzETMBEG
A1UECAwKQ2FsaWZvcm5pYTEwMBQGA1UEBwwNU2FuIEZyYW5jaXNjbzENMAsGA1UECgwET2t0YTEU
MBIGA1UECwvLU1NPNUHjvdmlkZXIxEzARBgNVBAMMCmRldi1MzAzMDIxHDAABgkqhkiG9w0BCQEW
DWluZm9Ab2t0YSjzb20wHhcNMTkwNDI5MTU0NjM2WhcNMjkwNDI5MTU0NzM2WjCBkjELMAkGA1UE
BhMCVVMeEzARBgNVBAgMCkNhbgImb3JuaWExFjAUUBgNVBACMDVNhbIBGcmFuY2lzcY28xDTALBgNV
BBoMBE9rdGExFDASBgNVBAAsMC1NTT1Byb3ZpZGVyMRMwEQYDVQDDApkZXxtNTMwMzAyMRowwGgYJ
KoZlhvcNAQkBFg1pbmZvQG9rdEuY29tMIIBjANBgkqhkiG9w0BAQEFAAOCAQ8AMIIIBCgKCAQE
26clu1q2eNnUp4VbhnsJ5mucEbpoSReDzaG8/CLNj73mp5dCicLZ/zrw/6SO+VbIV2drShOJDQzTx
sBfPOCm7cn2QV8CARW8Ci/oPPnqrXp/8yetElp4tea1g7WQbkZxO0uStjRlOrW6ISKgKrWEo1LII
6iU7kdbu7GIUUhY6KpSuMpT3iQq0MiQV7v+VFVx03VpGiAVFGIPY4hbDc0PwURN3MBw4cnhCK
kTsTAqwsSwH0N+9wsjwYWhEEMDYq/Eo0pbA2sGkChkz/vYJSUXjWMKnD432T3tYNYxEm95HzAIJn3
ixdn8IPTbiERH60LyYWQtn55TAf5oCCq4UxGoWIDAQABMA0GCSqGSIb3DQEBCwUA4IBAQ4PWYQ
QBdfPnkWx2Tqacsm3BnliPvQaslzejk/9+jL3PKvR/GQIQMFaFis2Fc+ucB426d0u3EbWSB8e9v4
deSVWytwpMEHQijKHV8zjaX2AcWWRgiv/1VR4Klp2cahX3h2tU+n87A3a3su201VuJ+inSIxquTu
```

Test Login Credentials

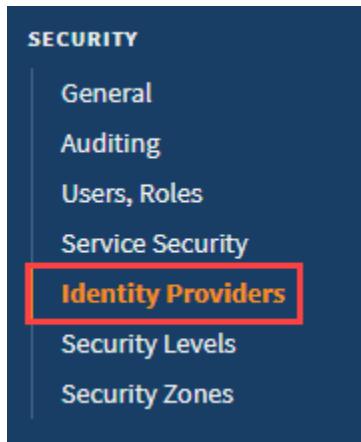
You need an account specific to the IdP for testing purposes (Okta in this example). To test and verify the IdP account, login to your IdP. For our example, the Okta login page is shown here:



You should now have a IdP credentials to test with, a metadata URL or metadata XML file. The next step is to configure Ignition to communicate with your IdP.

Configure Ignition Gateway

1. On the Gateway Webpage, click on the **Config** tab. You will need to log in if you aren't already.
2. Under the Security section, click on **Identity Providers**. The Identity Providers screen is displayed. This screen will list all IdPs that have been configured. You can filter by name or adjust the number of IdPs displayed per page in the view.



3. Click on the **Create a New Identity Provider...** link.

The screenshot shows the Ignition Config interface under the Security section, specifically the Identity Providers page. At the top, there is a navigation bar with links for Config, Security, and Identity Providers. Below the navigation is a search bar labeled "Filter type to filter" and a "View" dropdown set to 20 items. A table lists two identity providers:

Name	Type	Description	Action
default	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "default".	More Settings
Okta_2018	OpenID Connect 1.0		More Settings

Below the table are two buttons: "Create new Identity Provider..." (highlighted with a red box) and "Import Identity Provider...". At the bottom of the page is a footer with navigation icons.

4. Select the **Security Assertion Markup Language 2.0** option and click **Next**.

The screenshot shows the "Create New" configuration screen for an identity provider. The title bar indicates the path: Config > Security > Identity Providers > Create New. The screen displays three options:

- Ignition**: Describes the instance of Ignition hosting the configuration as an identity provider, useful when no external provider is available.
- OpenID Connect 1.0**: Describes OpenID Connect 1.0 (OIDC) as an authentication layer on top of OAuth 2.0.
- Security Assertion Markup Language 2.0**: Describes SAML as an open standard for exchanging authentication and authorization data between parties.

The third option is highlighted with a red box. At the bottom of the screen are "Cancel", "Step 1 of 2", and "Next" buttons.

5. On the Basic Details screen, provide an **Provider Name**. You can also add an **Provider Description** if desired. The Provider Type field will fill in automatically from the previous screen.

Config > Security > Identity Providers > Create New

Basic Details		* Required Field
Provider Name *	Okta_SAML_Example Give the provider a name.	
Provider Description	Test setup for SAML IdP. A description for the provider.	
Provider Type *	Security Assertion Markup Language 2 The type of the provider.	

6. The next section is **Import Provider Metadata**. In the **Import from URL** section, enter in the URL from earlier specific to your IdP. You can also import a file below if it was provided by your IT department.
7. Click on the **Import** button.

Import Provider Metadata

Import from URL	<input type="text" value="https://dev-123456.okta.com/app/esd"/> URL to the SAML Identity Provider Metadata document. Import
Import from File	<input type="button" value="Choose File"/> No file chosen File must be an XML document which conforms to the SAML 2.0 metadata schema described in saml-metadata-2.0-os. Import

8. Ignition will now generate a URI redirect address for your Ignition server. It is listed just below the “Import Provider Metadata” area of the configuration page.
In our example it is <http://10.10.110.86:8088/data/federate/callback/saml>. You need to provide this URI to your IdP (usually this means giving it to your IT department).

Note: The URI should be a web address that is accessible from the IdP server.

💡 The SP Entity ID for this Ignition Gateway is: <http://localhost:8088>
 The Assertion Consumer Service (ACS) URL for this Ignition Gateway
 is: <http://localhost:8088/data/federate/callback/saml>
 The SP Metadata for this Ignition Gateway may be accessed at the following
 URL: <http://localhost:8088/data/saml/metadata/sp>

9. Once you have given your IT department the redirect address, they can add your Ignition server as an application to the IdP.

Note: The IdP can use the same redirect address for the Login, Logout, and Initiate Login.

10. The next section is Provider Configuration. Most of the fields below should now be filled in when you imported the IdP Metadata.

Provider Configuration		* Required Field
IdP Entity ID *	<input type="text" value="http://www.okta.com/exkinzujwtqNwt"/> The Identity Provider's Entity ID.	
SP Entity ID *	<input checked="" type="checkbox"/> Automatically generate the SP Entity ID based on the hostname that the client uses to connect to this Gateway. <input type="text" value="SP Entity ID"/> The Service Provider's Entity ID. In this case, the Service Provider is the current Ignition Gateway	
Assertion Consumer Service (ACS) Binding	<input type="text" value="HTTP_POST"/> The expected binding used by the Identity Provider when interacting with Ignition's Assertion Consumer Service.	
Name ID Format	<input type="text" value="UNSPECIFIED"/> The expected name ID format for subjects of assertions resulting from Authn Requests.	
Single Sign-On (SSO) Service URL *	<input type="text" value="https://dev-530302.okta.com/app/indi"/> The Identity Provider's Single Sign-On (SSO) Service URL.	
Single Sign-On (SSO) Service Binding	<input type="text" value="HTTP_POST"/> The binding Ignition will use for sending Authn Requests to the Identity Provider's Single Sign-On (SSO) Service.	
Force Authn	<input type="checkbox"/> Check this box to force Ignition to authenticate the user each time instead of relying on a previous security context. See section 3.4.1 of saml-core-2.0-os for more details.	
Validate Response Signatures	<input checked="" type="checkbox"/> Check this box to validate the signature of the response from the Identity Provider.	
Validate Assertion Signatures	<input type="checkbox"/> Check this box if it is expected that assertions will be signed. Ignition will validate the signatures of each assertion.	

11. Press the Save button at the bottom right of the page. You'll see a confirmation message.

The screenshot shows the Ignition configuration interface under the 'Security' section, specifically the 'Identity Providers' page. A green success message at the top states 'Identity Provider - Okta_SAML_Example has been saved.' Below this, there is a table listing three identity providers:

Name	Type	Description	Action
default	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "default".	More Delete
Okta_SAML_Example	Security Assertion Markup Language 2.0	Test setup for SAML IdP	More Delete
opcua-module	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "opcua-module".	More Delete

12. The next step is to perform a test login. From the Identity Providers screen, select **More** and then **Test Login**.
13. You will be re-directed to the Okta login. Enter in your test login credentials and click the **Sign In** button.

The screenshot shows the Okta sign-in page. At the top, the Okta logo is displayed above a circular profile picture. Below the profile picture, the word 'Sign In' is centered. The page contains two input fields: one for email (containing 'training@inductiveautomation.com') and one for password (containing a masked string). Below the password field is a 'Remember me' checkbox. At the bottom of the form is a large blue 'Sign In' button. Below the button, a link 'Need help signing in?' is visible.

14. If the login is successful, you will be returned to the Identity Provider Test Login screen. The returned results will be displayed in the Results section.

Test Results - Security Assertion Markup Language 2.0

[IdP Response Data](#)

[Mapped User Attributes](#)

[Security Level Grants](#)

```
<?xml version="1.0" encoding="UTF-8" standalone="no"?>
<saml2p:Response xmlns:saml2p="urn:oasis:names:tc:SAML:2.0:protocol" Destination="http://10.10.110.86:8088/d
  <saml2:Issuer xmlns:saml2="urn:oasis:names:tc:SAML:2.0:assertion" Format="urn:oasis:names:tc:SAML:2.0:na
  <ds:Signature xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
    <ds:SignedInfo>
      <ds:CanonicalizationMethod Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
      <ds:SignatureMethod Algorithm="http://www.w3.org/2001/04/xmldsig-more#rsa-sha256" />
      <ds:Reference URI="#id33568854137229631961225447">
        <ds:Transforms>
          <ds:Transform Algorithm="http://www.w3.org/2000/09/xmldsig#enveloped-signature" />
          <ds:Transform Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
        </ds:Transforms>
        <ds:DigestMethod Algorithm="http://www.w3.org/2001/04/xmlenc#sha256" />
        <ds:DigestValue>voI6f9hmc30bTzHTm+onwR6QdmPaPF4zYFLSs+7Qous=</ds:DigestValue>
      </ds:Reference>
    </ds:SignedInfo>
    <ds:SignatureValue>ijLjE5ZiGe1YVXf3BWPF2naBrYEmaHoaHooESPiEcsVlRvMLnuoijmJ9ZUMCn+iqcAsCW+UbYdkamtPro
    <ds:KeyInfo>
```

User Grants

A User Grant is a way to directly assign a user to a [Security Level](#), even if they do not meet the requirements of the [Security Level Rules](#). User Grants essentially act as an override to the original rules of the Security Level.

User Grants are accessed from the Gateway Webpage **Config** section in **Security > Identity Providers**. Users can be added and edited using the buttons in the Users table so that Security Levels can then be granted to them.

Note: When adding, editing, and deleting users in User Grants, you are only modifying the User Grant (whether the user is granted permission that overrides the Security Rules). The user is not changed in the Identity Provider.

Users are identified by either their username or their ID from the provider response document. Once you have identified a user, you can assign them any number of grants to Security Levels. Selecting a level will automatically select all security levels above it. The User Grants can only be applied to a user after they authenticate with the Identity Provider, though the grants do not have to be for levels within the Authenticated branch.

Note: The system can't validate any user created here against actual users in the Identity Provider (IdP). Instead, the username or ID needs to be entered exactly, and when a user logs in, the system will check to see if they match any of the configured usernames/IDs to give User Grants to.

On this page ...

- [Configuring a User Grant](#)



INDUCTIVE
UNIVERSITY

User Grants

[Watch the Video](#)

Configuring a User Grant

There are two parts to configuring a User Grant: Adding a user then applying User Grants.

1. From the Gateway Webpage **Config** tab, click on **Security > Identity Providers**. The screen will refresh and you will see a list of all your IdPs.
2. Choose the IdP and click the **More** button to see the actions in the dropdown list.
3. Select **User Grants**.

The screenshot shows the 'Identity Providers' configuration screen. At the top, there are tabs for 'Config', 'Security', and 'Identity Providers'. Below the tabs is a search bar labeled 'Filter' with a placeholder 'type to filter' and a 'View' dropdown set to '20'. A table lists identity providers with columns for 'Name', 'Type', 'Description', and 'Action'. The 'Action' column contains a 'More' dropdown and a 'Delete' button. The 'More' dropdown menu for 'Okta_2018' includes 'Settings', 'User Attribute Mapping', and 'User Grants'. The 'User Grants' option is highlighted with a red box. The 'More' dropdown menu for 'Okta_TroubleShooting' also includes 'User Grants', which is also highlighted with a red box. At the bottom of the table, there are links for 'Create new Identity Provider...' and 'Import Identity Provider...'. To the right of the table, there is a sidebar with links for 'Test Login' and 'Export'.

4. To add a new user, click the **Add** icon.
5. Choose how you will identify the user; either with a username or an ID. Click **Confirm** to save the changes.

Add New User

ID
jsmith

Type
 ID Username

Confirm

Cancel

6. With the user created and highlighted in the Users table, select Security Levels to grant them when they Authenticate with this Identity Provider.

Users	
ID / Username	Type
jsmith	ID

- Public
- Authenticated
 - Custom Roles
 - PlantA
 - Floor1
 - Intern
 - Manager
 - Operator
 - NewLevel
 - Weekend Access
 - Roles
 - Administrator
 - Operator
 - Weekend Security

7. Click **Save**.

Now you can test this user through the [Test Login and Logout](#) screen to verify the new roles have been assigned.

Test Login and Logout

On the Identity Providers screen you can test a username and password combination against an Identity Provider (IdP).

When you select the Test Login option for your IdP, it will confirm the IdP name and Type that you are testing against. It gives you a way to test your attribute mapping configuration and your security level rules / direct user grants configuration

Clicking the Test Login button will redirect you to the IdP where you can login. Upon successful authentication with the IdP, the page navigates back to Ignition, and Ignition displays the response document as the results. These results can vary between IdPs, so it can be useful to test out a login to see what your IdP returns in its response document.

You can use Test Logout option to log out of the ID you were testing.

Ignition's IdP returns an 'amr' attribute that indicates how the user was authenticated.

- If the user was authenticated with a username and password challenge, the amr returns ["uname", "pwd"]
- If the user was authenticated with a badge challenge, the amr returns: ["badge"]
- If the user was authenticated with a badge and password challenge, the amr returns: [:badge", "pwd"]

With this feature, you can enable different security levels based on how the user authenticated. In a security level rule, you could enter:

```
containsAll ({idp-attributes:amr}, 'uname', 'pwd')
```

On this page ...

- [Test a Login](#)
- [Test a Logout](#)

Test a Login

1. From the Gateway Webpage **Config** tab, go to **Security > Identity Providers**. The window will refresh and your list of Identity Providers will be displayed.
2. Choose the Identity Provider and click the **More** button to see the actions in the dropdown list, and select **Test Login**.

The screenshot shows the Ignition Gateway Webpage with the 'Identity Providers' tab selected. The table lists several providers:

Name	Type	Description	Action
Default_User_Source	Ignition		More ▾ Delete
Okta_2018	OpenID Connect 1.0		More ▾ Test Login Delete
Okta_Example	OpenID Connect 1.0	Test s	More ▾ User Attribute Mapping Delete
Okta_TroubleShooting	OpenID Connect 1.0		More ▾ User Grants Delete

At the bottom of the list, there are links to 'Create new Identity Provider...' and 'Import Identity Provider...'. A red box highlights the 'Test Login' option in the dropdown menu for the 'Okta_2018' provider.

3. Log in at your IdP's login screen.

4. If the login is successful, you will be returned to the Identity Provider Test Login screen. The returned results will be displayed under the IdP Response Data tab.

The screenshot shows the 'Test Results - Ignition' page with the 'IdP Response Data' tab selected. The content area displays a JSON object representing the access token response:

```
{  
  "tokenEndpointResponse": {  
    "access_token": "pE7Z8qg-06EJIrLyi59ceheFHljEZFnM4FEx5dBoieY",  
    "id_token": "eyJraWQiOiJrMSIsImFsZyI6IlJTMyU2In0.eyJpc3MiOiJkZWZhWx0Iiw1YXVkJoi  
    "token_type": "Bearer",  
    "expires_in": 3600,  
    "scope": "openid"  
  },  
  "idTokenClaims": {  
    "iss": "default",  
    "aud": "ignition",  
    "exp": 1602202511,  
    "jti": "pB7xNnbS93vARr_ijgygTQ",  
    "iat": 1602201911,  
    "nbf": 1602201701  
  }  
}
```

5. Click on the Mapped User Attributes tab to view the user attributes for the currently logged in user.

The screenshot shows the 'Test Results - Ignition' page with the 'Mapped User Attributes' tab selected. A red box highlights this tab. The content area displays a table of user attributes:

Attribute Name	Attribute Value
ID	1
Username	admin
First Name	
Last Name	
Email	admin@mycompany.com
Roles	Administrator, Operator

Below the table are two buttons: 'Test Login' and 'Test Logout'.

6. Click on the Security Level Grants tab to view the Security Levels for the roles of the currently logged in user.

Config > Security > Identity Providers > default - Test Login

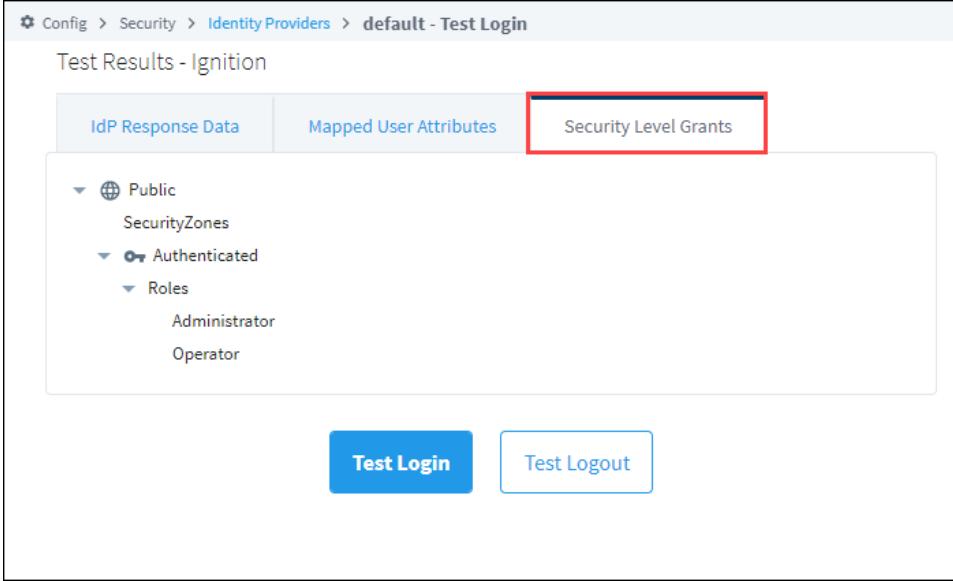
Test Results - Ignition

IdP Response Data	Mapped User Attributes	Security Level Grants
-------------------	------------------------	-----------------------

▼ Public
 SecurityZones

▼ Authenticated
 ▼ Roles
 Administrator
 Operator

Test Login **Test Logout**



Test a Logout

After testing a User ID, you do not want to stay logged in as the user. You can use the Test Logout function to log out. For the Ignition IdP, this function also logs you out of the IdP. For an OpenID Connect IdP, this function will also log you out of the IdP if you have a Logout URL.

1. To log out of the ID you were testing, click the **Test Logout** button on the Test Login page.

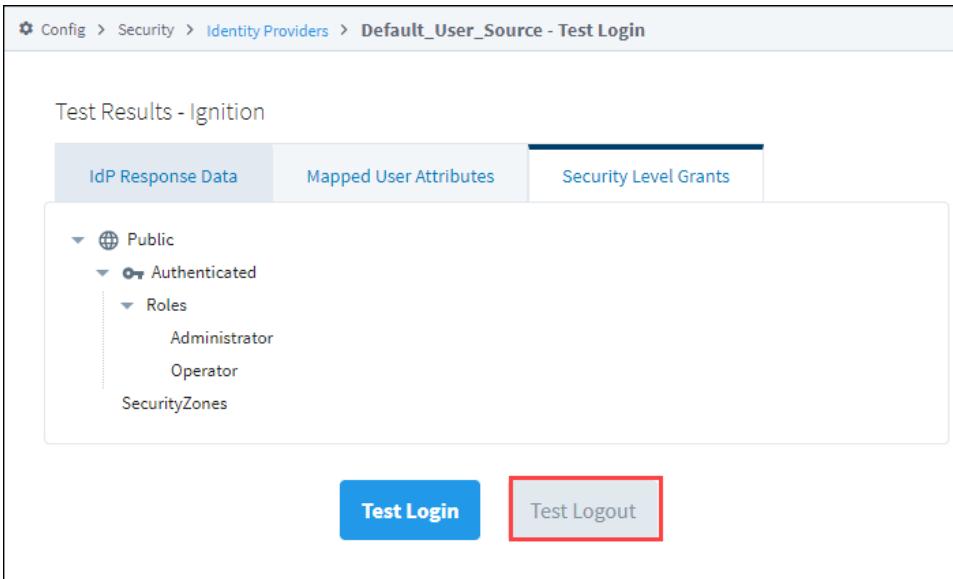
Config > Security > Identity Providers > Default_User_Source - Test Login

Test Results - Ignition

IdP Response Data	Mapped User Attributes	Security Level Grants
-------------------	------------------------	-----------------------

▼ Public
 ▼ Authenticated
 ▼ Roles
 Administrator
 Operator
 SecurityZones

Test Login **Test Logout**



2. You will get a confirmation message of a successful logout.

The screenshot shows the Ignition configuration interface under the 'Security' section, specifically the 'Identity Providers' page. At the top, there is a navigation bar with 'Config > Security > Identity Providers'. Below the navigation, a green success message box is displayed with the text 'Successfully logged out of Identity Provider default.' A red rectangle highlights this message box. Below the message, there is a search/filter bar with 'Filter type to filter' and a 'View 20' dropdown. A pagination control shows '1 of 1'. The main content area contains a table with the following columns: 'Name' (sorted), 'Type', 'Description', and 'Action'. There are two rows in the table:

Name	Type	Description	Action
default	Ignition	Automatically generated Ignition Identity Provider which uses the User Source Profile named "default".	More Settings
Okta_2018	OpenID Connect 1.0		More Settings

At the bottom left, there are two blue links: 'Create new Identity Provider...' and 'Import Identity Provider...'. The bottom right shows another pagination control with '1 of 1'.

Security Levels

With Security Levels, you define a hierarchy for access inside a Perspective Session or Vision Client. This authorization system provides a way for you to map roles from an Identity Provider (IdP) to Ignition roles. Any IdP can be used to provide roles, and security levels are independent of the type of IdP being used. Any role from the IdP is automatically granted to the user as a role, but only roles in your Security Levels are available to the security screens in the Designer. You can also use the [User Grants](#) option to grant additional access for each user.

Security Levels are defined at the Gateway and they are arranged in a tree structure. Each child (nested) level of the tree inherits the security of its parent levels. There are four reserved Security Levels in the platform:

- Public
- Authenticated
- Authenticated/Roles
- Security Zones.

To access Security Levels, go to the Gateway Webpage under the **Config** tab, and choose **Security > Security Levels**.

The screenshot shows the Ignition Config interface with the 'Security Levels' section selected. On the left, a tree view shows the hierarchy: Public, Authenticated, Roles (Administrator, Manager, Operator), and SecurityZones. On the right, a 'Security Level Details' panel is open for the 'Authenticated' node. It contains fields for 'Name' (set to 'Roles'), 'Description' (set to 'Represents the roles that a user has'), and 'Path' (set to 'Authenticated/Roles'). A note below the path says 'Editing not allowed for this security level'. At the bottom left of the panel is a '+ Add Security Level' button.

The following feature is new in Ignition version **8.1.16**
[Click here](#) to check out the other new features

You may now copy the selected Security Level Path to your clipboard by clicking **Copy Path** in the Security Level Details panel:

The screenshot shows a close-up of the 'Security Level Details' panel. The 'Path' field is highlighted with a red box and contains the text 'Authenticated/Roles'. To the right of the path is a blue 'Copy Path' button, which is also highlighted with a red box.

Reserved Security Levels

The reserved security levels are mostly created for you, and have special rules that determine when a user is granted that level. They can't be renamed or deleted.

Public

All users are always granted the Public security level, even if they are not authenticated (logged in). Public security level indicates open access and the least amount of security. A session that only has the Public security level is not authenticated. This is similar to being a guest or anonymous. Unless another security level is required, guest access will be allowed. The Public security level is the ancestor of all other security levels in the hierarchy.

Authenticated

On this page ...

- Reserved Security Levels
 - Public
 - Authenticated
 - Authenticated/Roles
 - SecurityZones
- Custom Security Levels
- Add a New Security Level
- Edit a Security Level
- Delete a Security Level
- Import a Security Levels Configuration
- Export a Security Levels Configuration



Security Levels

[Watch the Video](#)

The Authenticated Security Level is a child of the Public Security Level. If a session has authenticated against the configured IdP successfully, the Authenticated Security Level is granted. Users are required to be logged in in order to have access to this level.

Authenticated/Roles

The Roles level is a special level which itself has no special rules, but it acts as a parent placeholder for potential roles returned from the IdP. This particular level is not configurable; however there can be levels added underneath the Roles level as children. These levels should correspond to the names of roles that would be expected from the IdP. If the IdP provides role information, these roles are automatically mapped to the child security levels underneath Authenticated/Roles. The names of the roles must match exactly for them to be correctly mapped to. For example, if you authenticate against the Ignition IdP configured to delegate to the Internal user source, and your user was granted the roles "A", and "B", you would have (at a minimum) the following security levels granted to you:

- Public
 - Authenticated
 - Roles
 - A
 - B

Note: You can only add one level of children to the Roles Security Level. Custom Roles can be nested as deeply as you want.

SecurityZones

The SecurityZones level is another special placeholder level that itself has no rules but is a parent for all of the Security Zones on the Gateway. Security Zones are automatically pulled in from the Gateway. A [Security Zone](#) is a list of Gateways, Computers, or IP addresses that are defined and grouped together. This group is a zone on the [Gateway Network](#), which can have additional policies and restrictions placed on it. Security Zones provide a way to bridge the IdP method of permissions with location-based permission modeling. You cannot add, edit, or remove the SecurityZones node or any node in the SecurityZones sub-tree.

Custom Security Levels

Custom Security Levels can be added to almost anywhere within the tree. When these levels are granted to a user is determined by the [Security Level Rules](#), which can pull information from the IdP, Security Zones, and even Tags. The placement of custom Security Levels can affect when they may be potentially granted to a user. Any custom levels set under the Public level, but not within Authenticated, do not need to have a user authenticate against the IdP to be granted to a user. However, custom levels within Authenticated do need to have the user authenticate to be granted to the user, even if the rule for that level does not use any of the IdP attributes.

Add a New Security Level

1. From the Gateway Webpage **Config** tab, click on **Security > Security Levels**.
2. In the Security Level tree, select the level that will be a parent for the new level.
3. Click the **Add Security Level** button.

 Add Security Level

4. In the **Security Level Details** screen area, enter the **Name** for the level.

Note: Security Level names within the same parent must be unique.

5. The path for the parent is filled in automatically. Use the **Parent** dropdown list if you want to change the parent for this new level.

Security Level Details

Name	NewLevel
Parent	Public (root)
Description	
	<input type="button" value="Delete"/>
	Parent
	Public (root) Public (root) Authenticated Authenticated/Custom Roles Authenticated/Custom Roles/PlantA Authenticated/Custom Roles/PlantA/Floor1 Authenticated/Custom Roles/PlantA/Floor1/Intern Authenticated/Custom Roles/PlantA/Floor1/Manager Weekend Access Weekend Access/Roles Weekend Access/Roles/Administrator Weekend Access/Roles/Operator Weekend Access/Roles/Weekend Security

6. Add a **Description** for the new level (optional).
7. If you make changes to Security Levels, but decide not to save them, you can use the Reset button  to return the tree view to its currently saved configuration.
8. To save the changes, press **Save**.

Edit a Security Level

1. From the Gateway Webpage **Config** section, click on choose **Security > Security Levels**.
2. In the Security Level Tree, select the level that you want to edit.

Security Levels

   	Security Level Details <table border="1"> <tr> <td>Name</td> <td>Manager - West</td> </tr> <tr> <td>Parent</td> <td>Authenticated/Custom Roles/PlantA/Floor1/Manager</td> </tr> <tr> <td>Description</td> <td></td> </tr> <tr> <td></td> <td style="text-align: right;"><input type="button" value="Delete"/></td> </tr> </table>	Name	Manager - West	Parent	Authenticated/Custom Roles/PlantA/Floor1/Manager	Description			<input type="button" value="Delete"/>
Name	Manager - West								
Parent	Authenticated/Custom Roles/PlantA/Floor1/Manager								
Description									
	<input type="button" value="Delete"/>								
 <ul style="list-style-type: none"> Public <ul style="list-style-type: none"> Authenticated <ul style="list-style-type: none"> Roles <ul style="list-style-type: none"> Administrator Operator Custom Roles <ul style="list-style-type: none"> PlantA <ul style="list-style-type: none"> Floor1 <ul style="list-style-type: none"> Manager - West Operator Supervisor SecurityZones <ul style="list-style-type: none"> Default zone 1 zone 2 Custom <ul style="list-style-type: none"> Roles <ul style="list-style-type: none"> Administrator Operator 	<input type="button" value="+ Add Security Level"/>								
	<input type="button" value="Save"/>								

Editor notes are only visible to logged in users

The screenshot shows the 'Security Levels' screen. On the left, there is a tree view of security levels. The 'Authenticated' node is expanded, showing 'Roles' (which contains 'Administrator', 'Manager', and 'Operator'), and 'SecurityZones' (which contains 'Default'). The 'Administrator' node is selected and highlighted in blue. On the right, there is a 'Security Level Details' panel. It contains fields for 'Name' (set to 'Administrator'), 'Parent' (set to 'Authenticated/Roles'), and 'Description'. Below these fields is a 'Path' field containing 'Authenticated/Roles/Administrat'. At the bottom of the panel are buttons for 'Delete' (with a trash icon) and 'Save' (with a checkmark icon). Above the 'Save' button is a 'Reset' button (with a circular arrow icon).

3. Make the desired changes in the **Security Levels Details** screen area.
4. If you make changes to Security Levels, but decide not to save them, you can use the Reset button to return the tree view to its currently saved configuration.
5. To save the changes, click the **Save** button.

Delete a Security Level

Caution: When you delete a Security Level, all children under that level will also be deleted.

1. From the Gateway Webpage **Config** tab, click on **Security > Security Levels**.
2. In the Security Level tree, select the level that you want to delete.
3. Click the **Delete** button on the **Security Level Details** screen.

Security Level Details

Name
Child Item

Parent
~Public/Authenticated...

Description
Area to provide a description about a specific Security Level.

 Delete

4. In the confirmation box, click **Delete** to confirm the delete.

The following feature is new in Ignition version **8.1.25**
[Click here](#) to check out the other new features

Note: Deleted security levels will need to be unchecked in the Designer permissions configurations with new settings saved to avoid permission errors. [Warning indicators](#) for selected security levels that no longer exist will be visible on the Project Properties, Event Configuration, Edit Permissions, and Tag Editor windows when applicable.

Import a Security Levels Configuration

1. From the Gateway webpage **Config** tab, click on **Security > Security Levels**.
2. Click the Import icon 
3. Choose **Import** on the confirmation screen.

Import Security Levels Config

Are you sure you want to import a new Security Levels configuration?

Import

Cancel

4. Choose a security levels configuration file to import.
5. Click **Open**.

Export a Security Levels Configuration

1. From the Gateway Webpage **Config** tab, choose **Security > Security Levels**.
2. Click the Export icon .
3. The security levels configuration will be saved as a .json file with a unique number, for example:
 [security-levels-config-1544547744394.json](#)

Security Level Rules

When a user accesses a project, they can fall into one or more of many [Security Levels](#) that you can set up. The Security Level Rules determine if the user has that Security Level or not. The rules come in the form of expressions that can access any of the Expression Language's functions, Tag values, Security Zone information, or Identity Provider attributes. Security Level Rules are accessible from the Gateway Webpage [Config](#) tab in **Security > Identity Providers**.

Predefined Rulesets

One thing you may notice right away is that on the Security Level Rules page, many of the built in Security Levels are missing. Some Security Levels don't allow you to create an expression that defines their rules. These particular Security Levels already have a set of rules that govern how a user gets them. The Security Levels corresponding to the different Security Zones are automatically given to users depending on which zones they fall into. Similarly, the Security Levels that correspond to the Roles a user gets while authenticated are automatically given to users depending on what roles we receive from the Identity Provider. These Security Levels are removed from the tree here because rules can't be defined on them. The "Public" and "Authenticated" Security Levels also can't have Security Level Rules defined on them, but are present in the tree because they can potentially be a parent to custom nodes which can have Security Level Rules. The Public level is granted to every user when they open the project, and the Authenticated level is granted when the user authenticates against an Identity Provider, regardless of what roles they may have.

On this page ...

- [Predefined Rulesets](#)
- [Defining Security Level Rules](#)
- [Special Object Reference](#)
 - [Special Function Reference](#)
 - [Evaluating Tag Values](#)
 - [Special Considerations for Rules](#)
- [Configuring Security Level Rules](#)



INDUCTIVE
UNIVERSITY

Security Level Rules

[Watch the Video](#)

Security Levels

The screenshot shows the 'Public' security level configuration. On the left, a tree view lists security components under 'Public': Authenticated (Roles: Administrator, Operator; Custom Roles: PlantA (Floor1: Intern, Manager)), SecurityZones (Default, zone 1, zone 2, zone 3, zone 4), and Weekend Access (Roles: Administrator, Operator). A button '+ Add Security Level' is visible. On the right, the 'Security Level Details' panel shows the name 'Public' and a note 'Editing not allowed for this security level'. A large grayed-out area covers the description field. At the bottom is a blue 'Save' button.

Defining Security Level Rules

For the Security Levels that can have rules defined, the rules are defined in the form of an expression which should return either True or False, the results of which determine whether a user falls into that level or not. The rules can take advantage of everything the [expression language](#) has to offer, including the built in expression functions and any Tag values. The expressions here also have the unique ability to access attributes from the Identity Providers response document for the authenticated user, as well as what Security Zone the user falls into.

Special Object Reference

These special objects can be used to reference information gathered from the IdP response document, mapped user attributes, or the Security Zone that the user falls under.

Object Type	Reference	Description
Response Attributes	Varies, see Expression Mappings for syntax info.	References Response Document Attributes directly.

User Attributes	{user:id} {user:userName} {user:firstName} {user:lastName} {user:email} {user:roles}	The following feature is new in Ignition version 8.1.5 Click here to check out the other new features References mapped User Attributes. Relyes on mappings configured on the User Attribute Mappings page. The <code>roles</code> attribute can be handled using the <code>containsAll</code> or <code>containsAny</code> .
Security Zones	{security-zones}	References the collection of security zones that the user currently has. The collection can be handled using one of the unique functions: <code>containsAll</code> or <code>containsAny</code> .
Security Level Name	{security-level-name}	The following feature is new in Ignition version 8.1.5 Click here to check out the other new features References the name of the security level being configured. For example, if the Security Level is "Authenticated/PlantA/Administrator", then the security level name is "Administrator"
Security Level Path	{security-level-path}	The following feature is new in Ignition version 8.1.5 Click here to check out the other new features References the full path to the security level being configured, beginning with the first node below Public . For example, "Authenticated/PlantA/Administrator".

Note: Response and User Attributes are only available to Security Levels that fall within the Authenticated Security Level. See below under Special Considerations for Rules.

Special Function Reference

When writing an expression to determine Security Level Rules, there are functions available that are not a part of the normal set available to Expression Bindings. These additional functions are:

Function Name	Description	Example
containsAll(collection, element 0, ..., element N)	Checks to see if all of the listed elements are present in the collection object. The function requires at least two arguments, a collection and an element.	containsAll({security-zones}, 'PlantA', 'Floor1', 'Press Room')
containsAny(collection, element 0, ..., element N)	Checks to see if any of the listed elements are present in the collection object. The function requires at least two arguments, a collection and an element.	containsAny({attribute-source:idTokenClaims:roles}, 'Manager', 'Operator')

Evaluating Tag Values

Tag values can be accessed in the Security Level Rules expression area by encasing the Tag Path (including the Tag Provider) in braces ("{}")

```
{[tagProvider]path/to/tag}
```

Note: Security Levels are determined on initial login for each session, so if a Security Level is using an expression that references a Tag value, changing the value while the session is running won't change the Security Levels applied to the users already logged in.

Special Considerations for Rules

When defining rules for a Security Level, it is important to notice where in the Security Level tree you are. If you want to access information out of the Identity Provider such as the username, you will need to ensure that the Security Level is located in the Authenticated branch. User information is only captured once a user logs in, so that information will only fall under Security Levels that come from the user being Authenticated. If a Security Level lies outside of the Authenticated branch, then the level will only have access to information such as Tag values and Security Zones.

Configuring Security Level Rules

1. From the Gateway Webpage under the **Config** tab, go to **Security > Identity Providers**.
2. A list of the Identity Providers will be displayed. Click the **More** button for the Identity Provider you want to edit, and select **Security Level Rules**.

Name	Type	Description	Action
Default_User_Source	Ignition		More Delete
Okta_2018	OpenID Connect 1.0		More Delete
Okta_Example	OpenID Connect 1.0	Test s	More Delete
Okta_TroubleShooting	OpenID Connect 1.0		More Delete

→ Create new Identity Provider...
→ Import Identity Provider...

3. Select the **Security Level Name** and, if a rule is defined, it will appear in the Rule field. If not, you can create one. We copied the expression "containsAny ({security-zones}, 'PlantA', 'Floor1', 'Press Room')" from the example above for the following example.

The screenshot shows a left-hand navigation pane and a right-hand configuration pane. The navigation pane displays a hierarchical structure of security levels and roles:

- Public
 - Authenticated
 - Custom Roles
 - PlantA
 - Floor1
 - Intern (highlighted in blue)
 - Manager
 - Operator
 - NewLevel
 - Weekend Access
 - Roles
 - Administrator
 - Operator
 - Weekend Security

4. After you enter your rule, click **Save**.

Troubleshooting Identity Providers

Editor notes are only visible to logged in users

QA items:

Logout URL seems to have no effect

JSON Web Key seems to have no effect

This Troubleshooting section has a compilation of examples to help you diagnose and troubleshoot issues with configuring IdPs.

The Save Button Is not Selectable

All required fields must be entered on the Identity Providers screen before Save can be selected. Required fields have an asterisk (*) next to their name.

Note:

If you are waiting for values for the ClientID and Client Secret fields, you can enter fake values and return when you have the correct value.

After Importing Metadata from a File, Values Did not Auto Populate

1. Verify that import file is JSON format.
2. Verify that the import URL is valid.
3. Re-import the file.

Login Testing: IdP Login Page Does not Appear

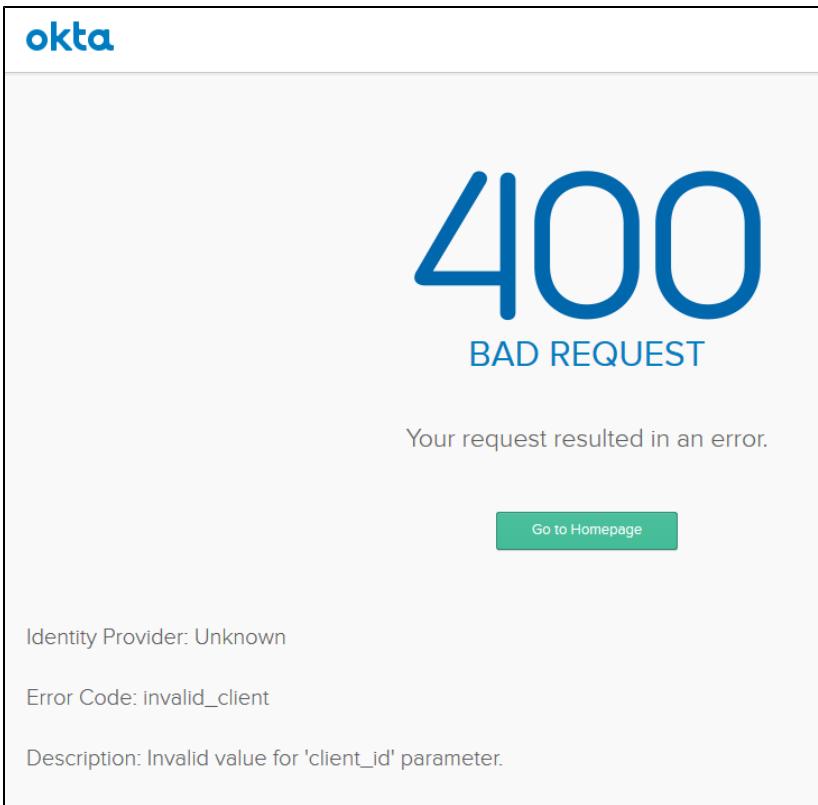
1. Confirm that the values in **Client ID** and **Client Secret** fields are correct. These come from the IdP when your Ignition Gateway is added as an application.
2. Check all other configuration settings.

Login Testing: The IdP login Is Displayed but the Login Attempt Fails

This issue is outside of Ignition. Check with your IT department and verify the login credentials (username and password) for your IdP.

On this page ...

- [The Save Button Is not Selectable](#)
- [After Importing Metadata from a File, Values Did not Auto Populate](#)
- [Login Testing: IdP Login Page Does not Appear](#)
- [Login Testing: The IdP login Is Displayed but the Login Attempt Fails](#)
- [Login Testing: The IdP Login Accepts the User but the IDP Redirect Fails \(HTTP ERROR 500\)](#)
- [Login Testing: The Test Is Successful, but Results Do not Show Useful data \(i.e., user name, eMail\)](#)
- [Login Testing: Revised User Attributes Are not Shown in the Results of a Successful Test](#)
- [You Are not Re-directed Back to Ignition after a Successful IdP Login](#)



Login Testing: The IdP Login Accepts the User but the IDP Redirect Fails (HTTP ERROR 500)

1. Go to Settings for your IdP.

A screenshot of the Okta 'Identity Providers' configuration screen. The left sidebar shows 'CONFIG' is selected under 'SYSTEM'. The main table lists four identity providers: 'Default_User_Source' (Ignition), 'Okta_2018' (OpenID Connect 1.0), 'Okta_Example' (OpenID Connect 1.0), and 'Okta_TroubleShooting' (OpenID Connect 1.0). The 'Okta_Example' row has a red box around its 'Settings' button. The table has columns for Name, Type, Description, and Action (More and Delete buttons).

2. Verify the setting for **Supported ID Token Signing Algorithm Values**.
3. If the URL for the IdP's metadata is available, try re-importing it.
4. Verify and re-enter the Client Id and Client Secret.
5. Verify the Token URL. Then re-import the IdP's metadata.
6. Verify the JSON web keys URL. (Default is to leave the check box checked.) Then re-import the IdP's metadata.
7. Verify the Issuer URL, then re-import the IdP's metadata.

HTTP ERROR 500

Problem accessing /data/federate/callback/oidc. Reason:

com.inductiveautomation.ignition.gateway.auth.idp.IdpAdapterManagerException: Unable to parse attributes from the web auth response

Login Testing: The Test Is Successful, but Results Do not Show Useful data (i.e., user name, eMail)

1. Go to settings for your IdP.
2. Add the desired fields to the Scope section. You may have to reference the developers documentation Scope document

Scope	Name	Action
	email	Remove
	user name	Remove

3. Click Save.
4. Repeat the Login test.

Login Testing: Revised User Attributes Are not Shown in the Results of a Successful Test

1. Go to the settings for your IdP.
2. Add the desired fields to the Scope section.
3. Click Save.
4. Repeat the Login test.

You Are not Re-directed Back to Ignition after a Successful IdP Login

1. Verify the Adapter Configuration: Authorization URL.
2. Re-import the IdP's metadata.

Referencing User Information

When we work with Identity Providers, the user signs into the provider, and the provider sends Ignition a response document containing information about the user that successfully logged in. This document is the basis for how Ignition handles the user, including what ID, username, roles, and contact info the user is assigned. When configuring a custom identity provider in Ignition, it will often be necessary to determine the contents of the response document, and make decisions about the logged-in user based on what is found.

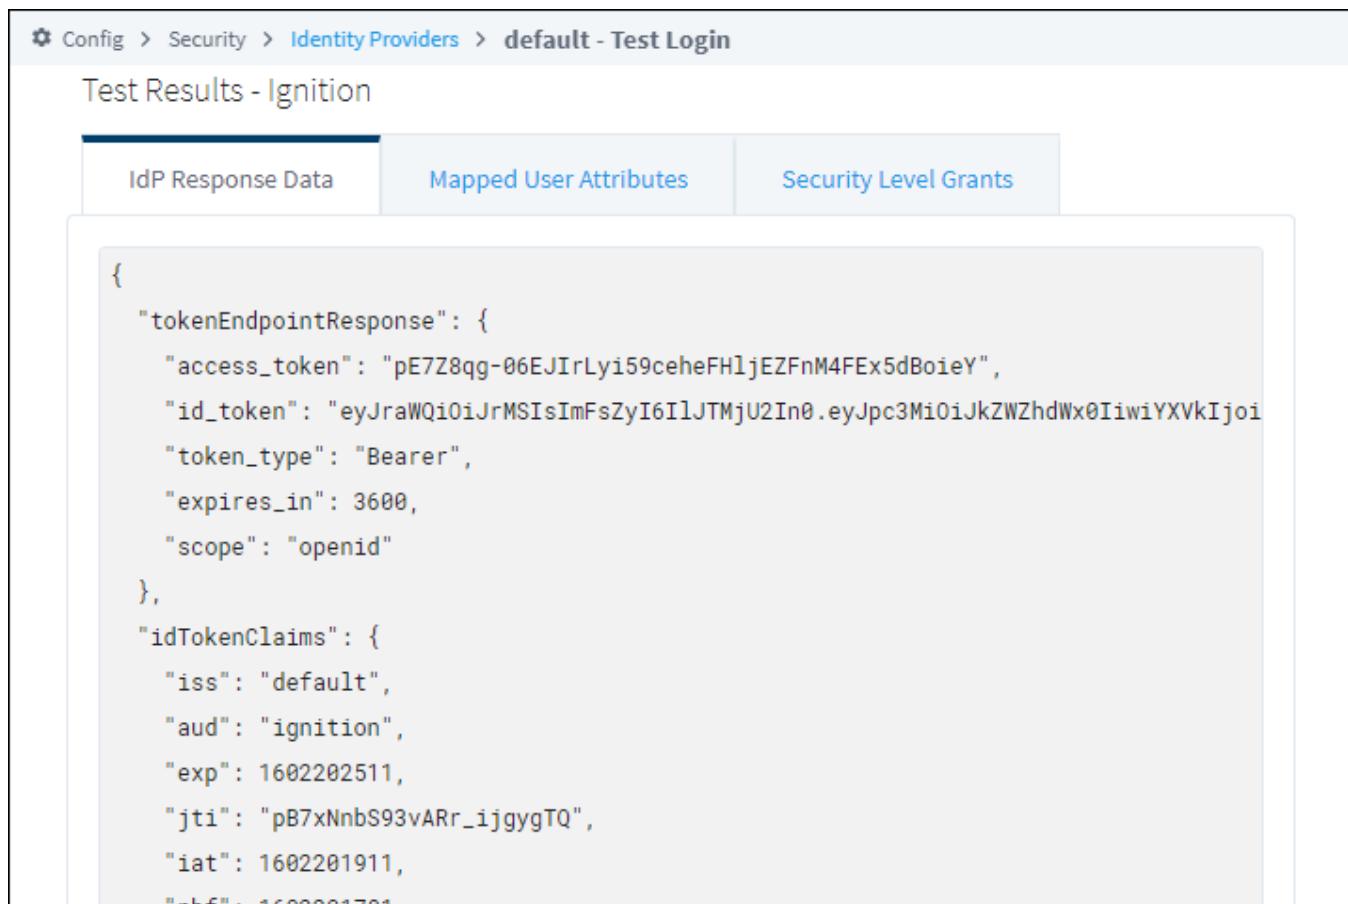
Here, we'll discuss some of the ways of referencing user information in Identity Provider-driven expressions.

On this page ...

- [Accessing a Response Document](#)
- [Referencing Response Elements](#)
 - [Reference Structure](#)
 - [Special Object Reference](#)
 - [Special Function Reference](#)
 - [Evaluating Tag Values](#)
 - [Special Considerations for Rules](#)
- [Referencing Mapped Attributes](#)

Accessing a Response Document

Many of the expressions found in the following sections will require us to know the layout of our Identity Provider's response document, as returned for a user when they authenticate. To view the structure of the response document for a given user, the [Test Login](#) feature can be used for the provider. The **IdP Response Data** tab will contain the returned document.



The screenshot shows the Ignition Test Results interface for a 'default - Test Login' session. The top navigation bar includes 'Config > Security > Identity Providers > default - Test Login'. Below the navigation is a section titled 'Test Results - Ignition' with three tabs: 'IdP Response Data' (selected), 'Mapped User Attributes', and 'Security Level Grants'. The 'IdP Response Data' tab displays a JSON response document:

```
{  
  "tokenEndpointResponse": {  
    "access_token": "pE7Z8qg-06EJIrLyi59ceheFHljEZFnM4FEx5dBoieY",  
    "id_token": "eyJraWQiOjJrMSIsImFsZyI6IlJTMjU2In0.eyJpc3MiOiJkZWZhdx0IiwiYXVkJoi  
    "token_type": "Bearer",  
    "expires_in": 3600,  
    "scope": "openid"  
  },  
  "idTokenClaims": {  
    "iss": "default",  
    "aud": "ignition",  
    "exp": 1602202511,  
    "jti": "pB7xNnbS93vARr_ijgygTQ",  
    "iat": 1602201911,  
    "nbf": 1602201701  
  }  
}
```

Referencing Response Elements

Whenever we work with a response document in an Ignition expression, as from [User Attribute Mappings](#) or [Security Level Rules](#), we may need to reference items in the document directly.

Reference Structure

In IdP contexts, it is possible to reference IdP document elements using a three-part format: {attribute-source:X:Y}

In this format, "X" is the attribute source for the provider, while "Y" is a path to the property inside the attribute source. Possible values for the attribute source will vary by IdP type:

IdP Type	Attribute Source(s)	Expression Path Name
Ignition	ID Token Claims (default)	idTokenClaims
	Token Endpoint Response	tokenEndpointResponse
	User Info Claims	userInfo
OIDC	ID Token Claims (default)	idTokenClaims
	Token Endpoint Response	tokenEndpointResponse
SAML	Authentication Response	authnResponse

In Ignition and OIDC Identity Providers, the Attribute Source is included in the response document already, with member properties inside. For SAML Identity Providers, the authnResponse string will not be shown in the structure, but must be included in the expression.

Special Object Reference

These special objects can be used to reference information gathered from the IdP response document or the Security Zone that the user falls under.

Object Name	Description
{security-zones}	This object gives the collection of security zones that the user currently has. The collection can be handled using one of the unique functions: containsAll or containsAny.
{attribute-source:X:Y}	This object can be used to access values in deeper structures by using colons to delimit each object. Where "X" is the attribute source for the provider, as defined on the User Attribute Mapping page. "Y" is any number of additional attributes along the path to the desired attribute. <div style="border: 1px solid #ccc; padding: 5px;"><p>Editor notes are only visible to logged in users As part of https://youtrack.ia.local/issue/DOC-319, add an extra sentence above. Something like: "The test login page is the best way to find the paths possible".</p></div>

Assuming the following JSON response from an identity provider:

```
{
  "idTokenClaims": {
    "foo" : "bar"
  },
  "userInfo": {
    "email" : "person@place.com"
  }
}
```

{attribute-source:userInfo:email} would retrieve the value of the user's email.

Attribute	Path	Example

ID	sub	{attribute-source:idTokenClaims:sub}
Username	preferred_username	{attribute-source:idTokenClaims:preferred_username}
First Name	given_name	{attribute-source:idTokenClaims:given_name}
Last Name	family_name	{attribute-source:idTokenClaims:family_name}
Email	email	{attribute-source:idTokenClaims:email}
Roles	roles	{attribute-source:idTokenClaims:roles}

Editor notes are only visible to logged in users

Add a new row with the following:

Add the following

ONE

- {user:id}
- {user:userName}
- {user:firstName}
- {user:lastName}
- {user:email}
- {user:roles}

TWO

Change "These special objects can be used to reference information gathered from the IdP response document or the Security Zone that the user falls under."

To "These special objects can be used to reference information gathered from the IdP response document, mapped user attributes, or the Security Zone that the user falls under."

THREE

Mention that the new items are post parsed, not pre-parsed

FOUR

Mention that the previous two rows are pre-mapped ("from the response document directly"

Note: The {attribute-source:x:y} object is only available to Security Levels that fall within the Authenticated Security Level. See below under Special Considerations for Rules.

Special Function Reference

When writing an expression to determine Security Level Rules, there are functions available that are not a part of the normal set available to Expression Bindings. These additional functions are:

Function Name	Description	Example
containsAll(collection, element 0, ..., element N)	Checks to see if all of the listed elements are present in the collection object. The function requires at least two arguments, a collection and an element.	containsAll({security-

		<pre> zones}, 'PlantA', 'Floor1', 'Press Room') </pre>
containsAny(collection, element 0, ..., element N)	Checks to see if any of the listed elements are present in the collection object. The function requires at least two arguments, a collection and an element.	<pre> containsAny({attribute- source:idTokenClaims: roles}, 'Manager', 'Operator') </pre>

Evaluating Tag Values

Tag values can be accessed in the Security Level Rules expression area by encasing the Tag Path (including the Tag Provider) in braces ("{}")

```
{[tagProvider]path/to/tag}
```

The screenshot shows the 'Config > Security > Identity Providers > default - Security Level Rules' interface. On the left, there's a tree view with 'Public' expanded, showing 'Authenticated' and 'CustomLevel'. 'CustomLevel' is selected and highlighted in blue. On the right, under the 'Rules' section, there are two fields: 'Level Name' containing 'CustomLevel' and 'Rule' containing the expression '[[Sample_Tags]Writeable/WriteableBoolean2]'. A vertical scroll bar is visible on the right side of the main content area.

Note: Security Levels are determined on initial login for each session, so if a Security Level is using an expression that references a Tag value, changing the value while the session is running won't change the Security Levels applied to the users already logged in.

Special Considerations for Rules

When defining rules for a Security Level, it is important to notice where in the Security Level tree you are. If you want to access information out of the Identity Provider such as the username, you will need to ensure that the Security Level is located in the Authenticated branch. User information is only captured once a user logs in, so that information will only fall under Security Levels that come from the user being Authenticated. If a Security Level lies outside of the Authenticated branch, then the level will only have access to information such as Tag values and Security Zones.

Referencing Mapped Attributes

In the context of [Security Level Rules](#), it is possible to reference response document elements as discussed above, but it is also possible to use mapped attributes directly.

Service Security

Service Security

After creating some Security Zones, a Security Policy can then be defined for each zone. This can be found by going to the **Config** section of the Gateway Webpage and navigating to **Security > Service Security**. At first, none of the zones will have a policy defined, and the Default zone will be at the top. Selecting **Edit** for any of them will bring up the Security Policy definition page for that zone. The Security Policy has four sections: Alarm Notification, Alarm Status, History Provider Access, and Tag Access. They work together to define how the local Gateway gives access to incoming Gateway connections. All four sections also have the ability to completely block access to specific services with the Service Access setting in each section. Setting that to deny will prevent zone access to that particular information, regardless of what the rest of the options are set to.

Note: It is important to realize that if you have a single Gateway, limiting access of certain clients to certain Tags is still done in the [individual Tags](#).

The following feature is new in Ignition version **8.0.7**

[Click here](#) to check out the other new features

On this page ...

- [Service Security](#)
 - [Default Security Zone](#)
 - [Setting Zone Priority](#)

- **Alarm Journal Access** - Alarm Journal Access has two main settings associated with it. The Default Profile Access is the default access rights for the Alarm Journal service. There will be an additional setting for each Alarm Journal configured on the local Gateway. As an example, the image below shows an "Access Level: 'MyJournal'" setting which corresponds to the configured Alarm Journal named "MyJournal". This setting can be set to **Inherited** which will cause this specific Alarm Journal to inherit the access rights set in the Default Provider Access Level. It can also be set to **No Access** to block query and storage to this specific Alarm Journal. Setting it to **Query Only** will allow users to only query data from this Alarm Journal without any storing capability. In contrast, the **Query and Storage** option allows users to store and query data from this Alarm Journal. It is important to note that every time a new Alarm Journal is created in the local Gateway, a new setting for this journal will be added to this Security Policy and it will automatically default to inherited.

- **Alarm Notification** - The Accessible Pipeline Filter setting is a list of Pipelines in the current Gateway that other connections can use for alarm notification. Pipelines must be entered in the format "project_name/pipeline_name". The list is a comma separated list, and it can make use of the (*) wildcard. This setting is an inclusionary list not an exclusionary list, meaning that if there are no pipelines listed here, then all of them will be available.

- **Alarm Status** - The Allow Acknowledge setting will allow the Gateways that fall within the zone to acknowledge alarms on the local Gateway. The Allow Shelving setting will allow the Gateways that fall within the zone to Shelve alarms on the local Gateway. IE: Other Gateways can shelve alarms on this Gateway. For this Gateway to shelve alarms on others, this must be set on the remote Gateway.

The following feature is new in Ignition version **8.0.7**

[Click here](#) to check out the other new features

- **Audit Log Access** - Audit Log Access has two main settings associated with it. The Default Profile Access is the default access rights for the Audit Profile service. There will be an additional setting for each Audit Profile configured on the local Gateway. Similar to the Alarm Journal Access, the image below will show an "Access Level: 'MyProfile'" setting which in this case corresponds to the configured Audit Profile named "MyProfile". This setting can be set to **Inherited** which will cause this specific Audit Profile to inherit the access rights set in the Default Provider Access Level. It can also be set to **No Access** to block query and storage to this specific Audit Profile. Setting it to **Query Only** will allow users to only query data from this Audit Profile without any storage ability. In contrast, the **Query and Storage** option allows users to store and query data from this Audit Profile. Just like with Alarm Journals, it is important to note that every time a new Audit Profile is created in the local Gateway, a new setting for this profile will be added to this Security Policy and it will automatically default to inherited.

- **History Provider Access** - The History Provider Access has two different settings. First, it has a Default Access Profile. This is the default access rights for Tag History. Second, there will be a setting for each History Provider set up on the local Gateway. In the image below, there is an "Access Level: 'DB'" that can be set that corresponds to the History Provider that was created when a database was connected. It can be set to Query and Storage, which will allow connections in the current zone to both run queries and store Tag History against the Tag History provider, Query Only, which will only allow the zone to query out history data, but not store it, and No Access, which will completely block access to that History Provider. The final setting is Inherited, which will inherit the Default Profile Access rights. Any new history providers will automatically get added to the Security Policy set at inherited so it may be

beneficial to set the Default Profile Access to be either Read Only or No Access so that a recently added history provider does not accidentally get storage rights when it should not.

Note: The Default Access Profile should not be set to Inherited. This also goes for the Default Provider Access Level in the Tag Access section.

- **Tag Access** - The [Tag Access](#) section also has a few different settings. The Default Provider Access Level sets the default access rights for realtime Tag providers.

The Impersonation Role Name field allows you to specify a role name to use when writing to a Tag from an incoming Gateway Network connection (from the selected Zone). Finally, the Tag Access section will then have a setting for each Tag provider configured in the local Gateway, as well as an additional one for system Tags. These can be set to **ReadWriteEdit**, which will allow connections in the current zone to read, write to, and edit the Tags in that provider, **ReadWrite**, which allows the zone to read and write to Tags, and **ReadOnly**, which only allows the zone to read the Tags. It also can be set to **None**, which will prevent the zone from interacting with the Tag Provider altogether, and **Inherited**, which will again inherit the access rights set in the Default Provider Access Level. Any new Tag Providers will automatically get added to the Security Policy with Inherited access rights.

The following feature is new in Ignition version **8.1.2**
[Click here](#) to check out the other new features

As of 8.1.2, the Trust Remote Security Levels setting allows users to opt into trusting the Security Levels of remote Gateway users when remote Gateways read, write, and subscribe to local Tags. If checked, security levels passed from the remote Gateway will be used for determining access to Tags on the local Gateway. If unchecked, or if the remote Gateway is on a version which does not support this feature, the remote Gateway's security zones and the impersonation role will be used as the security levels.

Default Security Zone

While the Default zone may not have a custom Security Policy defined, it does default to not include any notification pipelines, allow alarm acknowledgment, query only history access, and read only Tag access. This means that if a remote Tag Provider is set up on a remote Gateway, and the local Gateway has not changed the default security settings, the remote Gateway will have read only access to the Tag History Provider. This can be changed by editing the Default zone's Security Policy to fit a different preference, or creating new Security Zones with custom security policies. Once a Security Policy has been defined on a zone, it will automatically jump to the top of the list. A new option will also become available that will clear the policy from the zone.

Alarm Journal Access	
Service Access	Allow
Default Profile Access	Query Only
Access Level: 'MyJournal'	Inherited

Alarm Notification	
Service Access	Allow
Accessible Pipeline Filter	A comma separated list of alarm pipeline paths on this gateway (which can include the wildcard *) that will be made available for use in the alarm notifications of other gateways in this security zone.

Alarm Status	
Service Access	Allow
Allow Acknowledge	<input checked="" type="checkbox"/>
Allow Shelving	<input type="checkbox"/>

Audit Log Access	
Service Access	Allow
Default Profile Access	Query Only
Access Level: 'MyProfile'	Inherited

History Provider Access	
Service Access	Allow
Default Profile Access	Query Only
Access Level: 'DB'	Inherited

Tag Access	
Service Access	Allow
Default Provider Access Level	ReadOnly
Trust Remote Security Levels	<input type="checkbox"/> If checked, security levels passed from the remote Gateway will be used for determining access to tags on the local Gateway, if unchecked, or if the remote Gateway is on a version which does not support this feature, the remote Gateway's security zones and the impersonation role will be used as the security levels.
Impersonation Role Name	<input type="text"/> This role name will be injected into the security model of the tag for read/write/subscribe/edit operations.
Access Level: 'default'	Inherited
Access Level: 'System'	Inherited

Setting Zone Priority

Once a Security Policy has been defined for two or more zones, a new option appears on the Service Security page to move the zones up and down the list. This allows a priority to be set on the Security Zones, since a connection can apply to multiple zones. For example, say Zone 2 dictates that all requests coming from a range of IP addresses have query only history access, and read write access to Tags. Zone 1 includes specific Gateways, one of which is also contained in Zone 2, that will have query and storage history access and read write edit access to Tags. When a request comes in from a connection, it first determines which Security Zones it belongs to. The request then starts at the top of the Service Security list and goes down until it finds the first zone that it is in, and uses the access rights of that zone. In our example, we want to make sure Zone 1 is above Zone 2, so that the Gateway that is in both Zone 1 and Zone 2 gets the full access rights afforded to it by the Security Policy of Zone 1 instead of getting the limited access rights from Zone 2.

Service Security

⚠ Security policies are defined based on Security Zones. The highest ranking policy (from the top down) for a connection's zones will be used. If no other policies match, the "Default" policy will be used.

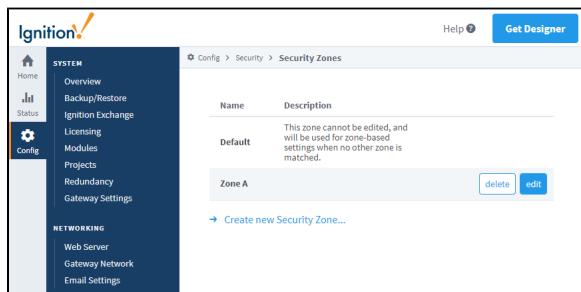
Security Zone	Policy Defined?	
Zone 1	true	Edit Clear Policy Move Down
Zone 2	true	Edit Clear Policy Move Up
Default	false	Edit

Related Topics ...

- [Security in Vision](#)

Security Zones

A Security Zone is a list of Gateways, Computers, or IP addresses that are defined and grouped together. This group now becomes a zone on the [Gateway Network](#), which can have additional policies and restrictions placed on it. While Users and Roles restrict access to specific functions within the Gateway like making certain controls read-only for certain users and read/write for others, Security Zones provide this functionality to the Gateway Network, limiting locations instead of people to be read-only for specific actions. This allows for greater control over the type of information that is passing over the network, improving security and helping to keep different areas of the business separate, while still allowing them to interconnect.



The screenshot shows the Ignition software interface. On the left, there's a navigation sidebar with sections for SYSTEM (Home, Status, Config), NETWORKING (Web Server, Gateway Network, Email Settings), and a central section for Config. Under Config, there are links for Overview, Backup/Restore, Ignition Exchange, Licensing, Modules, Projects, Redundancy, and Gateway Settings. The main content area is titled "Config > Security > Security Zones". It displays a table with two rows: "Default" and "Zone A". The "Default" row has a note: "This zone cannot be edited, and will be used for zone-based settings when no other zone is matched." Below the table are "delete" and "edit" buttons, and a link to "Create new Security Zone...".

Using Security Zones

Sometimes, in addition to knowing who the user is, it is important to know their location. An operator may have permissions to turn on a machine from an HMI, but if the operator is logged into a project on a different Gateway in the network that has remote access to those Tags, it might not be a good idea to let the operator write to those Tags from a remote location. The operator can't see if the physical machine is clear to run.

This is where Security Zones come in. While Security Zones themselves don't define the security, they instead define an area of the Gateway Network, breaking up Gateways and network locations into manageable zones that can then have a Security Policy set on them. Once there are zones defined, a Security Policy can be assigned to each zone, and a priority of zones can be set in the event that more than one zone applies in a given situation.

Caution: When using zone-based security in a project, the project stores the name of the security zone as a string. This means that if you were to modify the name of the zone in the Gateway, the zone-based security in your project will not update to reflect the new name, and instead will try searching for a zone with the original name. Be very careful when modifying the names of security zones.

A connection must pass all of the qualifier checks before being accepted into a Security Zone. So if Require Secure Connection was checked, and Allow Client Scope was not, any requests coming from Clients would be rejected even if they are secure, and the same goes for any non-secure connections coming from sources other than a Client.

Requests can be a part of more than one zone, depending on how the zones are set up. This can be useful for making a whole section of IP addresses read only, but a specific Gateway in that IP address range may be listed specifically in another zone, which can be given read/write access. Any connection which does not fall into one of the zones will be placed in the Default zone.

Define a Security Zone

When setting up a new Security Zone, it is a good idea to set up a [Gateway Network](#) first if you haven't already. While Security Zones can be defined and used without a connected Gateway, they work best when used in conjunction with other Gateways on a Gateway Network. There is a special zone called Default. It is always present and can't be modified, and will be used if an incoming connection does not match any of the other defined zones. Identifiers are how incoming connections are distinguished between different zones. While there are different options to define the incoming connection, it only needs to match one of the Identifiers to match a zone. After first being identified as part of a particular Security Zone, the connection then checks the Qualifiers. With the Qualifiers, the incoming connection needs to fit in with all of the properties before it is fully placed into the Security Zone.

1. Under the **Config** tab of the Gateway Webpage go to **Security > Security Zones**.
2. Select the **Create new Security Zone** link.

On this page ...

- [Using Security Zones](#)
- [Define a Security Zone](#)
- [Settings Table](#)



INDUCTIVE
UNIVERSIT

Security Zones and Service Security

[Watch the Video](#)

The screenshot shows the 'Security Zones' section of a configuration interface. At the top, there's a navigation bar with 'Config > Security > Security Zones'. Below this, a table lists five security zones:

Zone	Description	Actions
Default	This zone cannot be edited, and will be used for zone-based settings when no other zone is matched.	
localhost		delete edit
Zone 1	US Western Sales Region	delete edit
Zone 2	US Central Sales Region	delete edit
Zone 3	US Northeast Sales Region	delete edit

At the bottom of the list, there is a button labeled 'Create new Security Zone...' with a red border around it.

3. Enter a name and description for the new zone.
4. Under **Identifiers**, enter an **IP Address** and a **Gateway Name**.
5. Under **Qualifiers**, select the **Require Secure Connection** option. Leave the remaining options at their default settings.
6. Click **Create New Security Zone**. The page will refresh and you will see a green banner stating that your new Security Zone was successfully created.

Settings Table

Setting	Description
Identifiers	
IP Addresses	This defines an IP address that the connection is coming from. This can be a list of IP addresses by using commas to separate them. It can also make use of the (*) wildcard like '192.168.100.*', or use a range such as '100.100.1-100.0-255'. With IP addresses, virtually all connections can be listed. Use 127.0.0.1 for the local connection.
Host Names	The host name refers to the system name of the machine generating the request such as Joe_Workstation. This can be a list of names separated by commas, and it can also use the (*) wildcard like '*_Workstation'.
Gateway Names	A list of Gateway system names that qualify for this zone. Note: When identifying a Gateway through a proxy Gateway, the IP Address should be using the IP of the proxy, but the Gateway name should use the name of the Gateway we are trying to identify.
Qualifiers	
Require Secure Connection	If this is true, only connections that are made over a secure channel will be accepted.
Direct Connection Required	If this is true, only connections that come from a direct connection will be accepted. The Gateway Network allows you to connect three Gateways in a 1-2-3 configuration, where Gateway 1 can see Gateway 3 through the proxy Gateway 2.
Allow Client Scope	If this is false, any client scoped requests will not be accepted.
Allow Designer Scope	If this is false, any Designer scoped requests will not be accepted.
Allow Gateway Scope	If this is false, any Gateway scoped requests will not be accepted.

Related Topics ...

- [Service Security](#)

Project Security in the Designer

When several users are all working on the same project, managing changes to the project can become cumbersome. By default, all users with Designer access can modify, delete, save, and publish all resources available in the Designer. In some situations, it is desirable to limit what each user can do in the Designer. Ignition has several built-in Designer restriction methods to help in these scenarios.

On this page ...

- [Designer Project Permissions](#)
 - [Controlling Project Edits by Role](#)
 - [Protecting Project Resources](#)

Designer Project Permissions

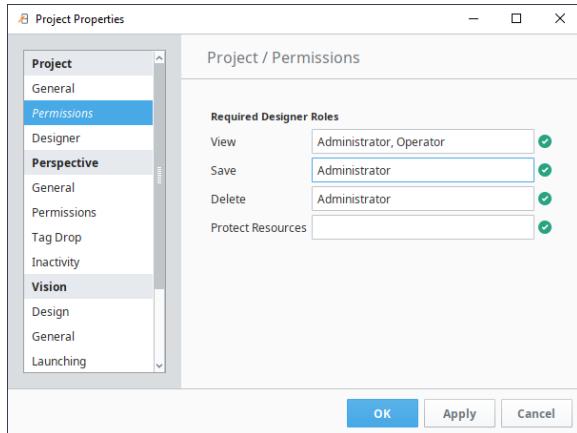
Actions such as, viewing, saving, deleting, and editing of project resources are restricted to users who have sufficient roles to do so. Editing of these required roles is done in the permissions section of the Project Properties dialog in the Designer. If required roles are not set for an action, then all users with Designer access can perform the action.

The Designer does not poll for role changes, so if a user who is currently logged into the Designer has their roles changed, they will need to re-launch the Designer for the new role(s) to take effect.

Controlling Project Edits by Role

You can control who gets to login to a project by assigning roles and giving permissions to those roles in the Required Designer Roles property which you set up in the Designer.

1. In the Designer, from the menubar, choose **Project > Properties**.
2. Go to the **Project > Permissions** area.
3. Under the Required Designer Roles, enter the appropriate roles next to each project-level restriction, as required. You can enter a comma-separated list of role names that are required to access the project. As you start typing, matching role names will pop up.



4. Click **OK** to save the changes.

The following table describes each of these four options:

Option	Affect
View	User must have at least one of these roles to view the project in the Designer.
Save	User must have at least one of these roles to save the project.
Delete	User must have at least one of these roles to delete the project.
Protect Resources	User must have at least one of these roles to access protected resources.

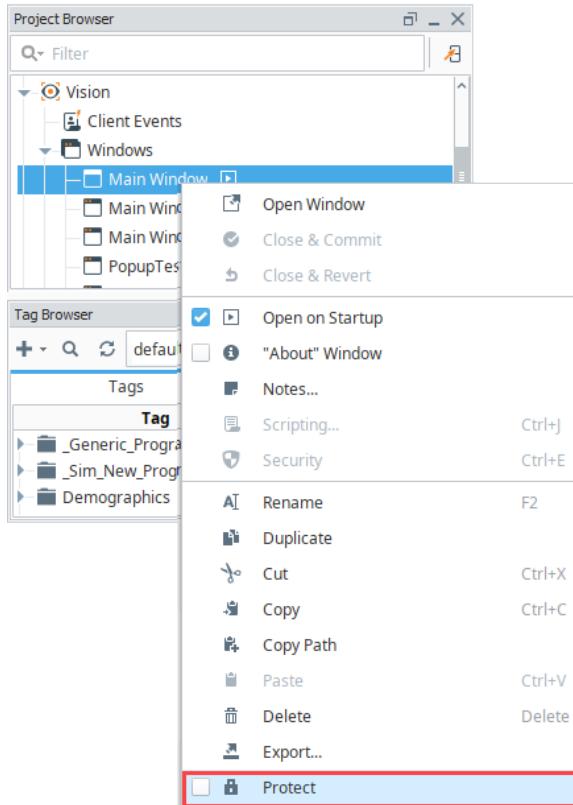


Project Permissions

[Watch the Video](#)

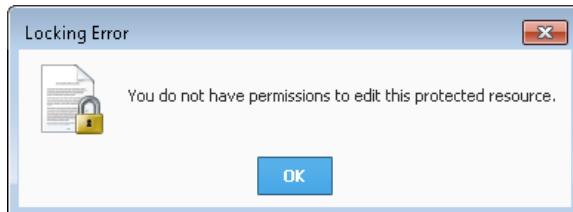
Protecting Project Resources

You can lock individual project resources from inside Designer by opening the Project Browser, and right clicking on any of the objects that you want to lock in. Select the **Protect** option to protect it. Once it's protected, it cannot be changed except by someone that has the permission to unprotect it, and modify it.

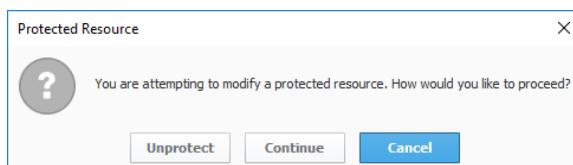


Protected resources are global or project resources that can only be edited by select users with the required roles. These roles are required to protect resources from being edited in the Designer, and do not apply to the clients. This means you can prevent a resource from being edited by other users who have Designer access. It is often used in scenarios where development work is finished on a window or object, and no further changes should be made to it. Other objects like [Vision Templates](#) or [Alarm Pipelines](#) are often protected so they may be used, but not modified.

Users without a required role will see the following message in the Designer when attempting to open a protected resource:



Users with a required role are allowed to modify the resource, but a message will appear informing them that the resource is protected, and will be asked how to proceed:

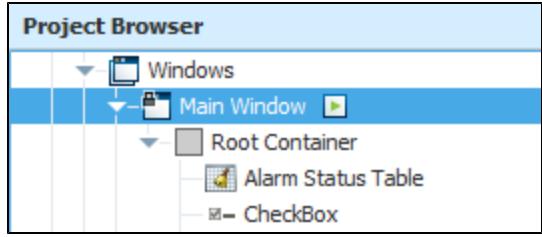


Additionally, a lock icon will appear on the resource informing users that it is protected. An example can be seen on the 'Audit Events' window below:



Locking Project Resources

[Watch the Video](#)



To remove the protection, simply right click the object and select the **Protect** option to unprotect it.

Security Certificates

A security certificate, also known as a digital certificate, is used to provide trust between connections. Trusted certificates establish the identity, authenticity and reliability of incoming and outgoing traffic.

The Ignition platform uses Secure Sockets Layer/Transport Layer Security (SSL/TLS), which requires certificates in multiple features. The purpose and type of certificate determine how the certificate is installed and where it is stored within Ignition. It is important to know where certificates are needed and what their purpose is to make sure all requirements are met. For example, the Gateway Network and OPC UA security both impact client and server connections, but because the Gateway Network connections are between local and remote gateways and OPC UA connections are between devices, the process for adding and trusting certificates is different.

The following is a list of locations where certificates are required, and a link pointing to a page containing the general security purpose, settings, and certificate management properties.

- [Launchers and Clients](#)
- [Gateway:](#)
 - [Acting as a server](#)
 - Acting as a client, see [Adding Security Certificates into KeyStores](#) below.
- [GAN Security](#)
- [OPC UA Security](#)

On this page ...

- [Types of Certificates](#)
 - [SSL Certificates](#)
 - [OPC UA Certificates](#)
- [Adding Security Certificates into KeyStores](#)

Types of Certificates

It may be helpful to understand the different types of certificate Ignition can use if you are new to certificates.

SSL Certificates

SSL certificates allow systems to verify identity and establish an encrypted network connection to another system using SSL/TLS protocols. There are two types of signed SSL Certificates, self-signed certificates and trusted certificate authority (CA) certificates.

- Self signed certificates are generated internally for free
- Trusted CA certificates are signed by a trusted certificate authority
 - Ignition supports CA certificates from your organization's internal CA or any publicly trusted certificate authority.

Both signed certificates offer encryption, but without the signature of a trusted certificate authority, warning messages will appear for self-signed certificates that are not trusted.

Since SSL/TLS requires the installation of a security certificate, both the Gateway Network and the Web Server can use self-signed certificates if CA certificates are not yet available or needed, such as during testing. It is important to note that although the functionality of certificates installed on the Gateway Network and Gateway Web Server are similar, they must be treated separately because settings made on one page on the Gateway do not apply to the other, even the case of shared ports.

OPC UA Certificates

UA security contains authentication and authorization as well as encryption and data integrity by signing. Security is integral to UA and OPC UA protocols are a hybrid variant of TLS, using binary encoding and HTTPS for transport. The Ignition platform inherently offers OPC UA client functionality and the Gateway can connect to any compliant OPC UA server.

Adding Security Certificates into KeyStores

In some cases when the Gateway is acting as a client, you may need to provide supplemental security certificates so the Gateway can communicate with other systems, such as databases or devices elsewhere on the network. These supplemental certificates can be added to a Gateway by simply placing them in the following directory on the Gateway's file system:

```
%gateway installation directory%data/certificates/supplemental
```

Once added, you will need to restart the Gateway before the certificates will be used.

Supported formats are DER encoded binary X.509, and Base-64 encoded X.509 (PEM-encoded ASCII).

OAuth 2.0 Clients

The following feature is new in Ignition version **8.1.24**
[Click here](#) to check out the other new features

What is OAuth 2.0?

OAuth 2.0 is an authorization standard that is developed by [IETF OAuth Working Group](#). Unlike basic authorization which requires a user to input a **fixed password** to log in, OAuth 2.0 requires users to input an **access token** instead. This makes using OAuth 2.0 more secure than basic authorization, since the access token is:

- Unique to the client and the requested scope of access
- Difficult to guess
- Constantly changing

Some well-known authorization protocols, such as OpenID Connect, are built on the OAuth 2.0 standard. See the [official OAuth 2.0 documentation](#) for more information.

On this page ...

- [What is OAuth 2.0?](#)

[Configuring OAuth2 Clients](#)

- [Creating a new OAuth2 Client](#)
- [OAuth2 Client Properties](#)

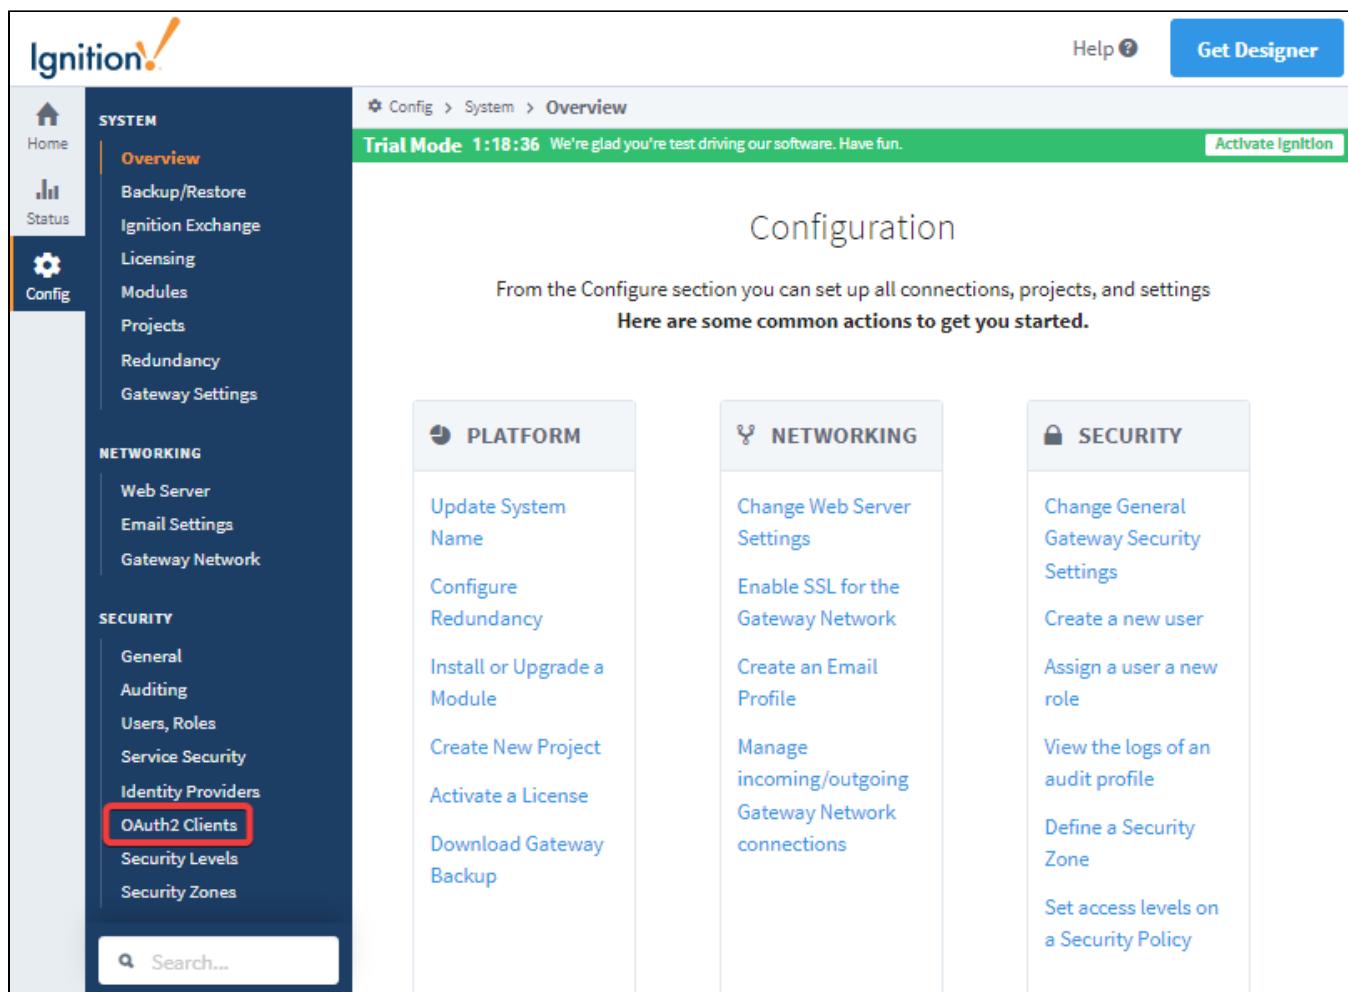
[OAuth2 Client Options](#)

- [Test Authorization](#)
- [Test Token](#)

[Test Authorization and Token Example \(Google Identity OAuth2\)](#)

Configuring OAuth2 Clients

OAuth2 Clients are used by OAuth2 Email Profiles. Additionally, multiple OAuth2 Email Profiles can use the same OAuth2 Client. When configuring an [OAuth2 SMTP Email Profile](#), you can reference the name of your OAuth2 Client to use those properties. To set up or modify OAuth2 Clients, navigate to the OAuth2 Clients section. You can find this section by going to your [Gateway webpage > Config > Security > OAuth2 Clients](#).



The screenshot shows the Ignition Configuration interface. The left sidebar has a dark blue background with white text. It includes links for Home, Status, and Config. Under Config, there are sections for SYSTEM (Overview, Backup/Restore, Ignition Exchange, Licensing, Modules, Projects, Redundancy, Gateway Settings), NETWORKING (Web Server, Email Settings, Gateway Network), and SECURITY (General, Auditing, Users, Roles, Service Security, Identity Providers, OAuth2 Clients, Security Levels, Security Zones). A search bar at the bottom of the sidebar contains the placeholder "Search...". The main content area has a light gray background. At the top, it shows "SYSTEM Overview" and "Trial Mode 1:18:36 We're glad you're test driving our software. Have fun." with a "Activate Ignition" button. Below this is a section titled "Configuration" with the sub-instruction "From the Configuration section you can set up all connections, projects, and settings Here are some common actions to get you started." There are three columns of actions: "PLATFORM" (Update System Name, Configure Redundancy, Install or Upgrade a Module, Create New Project, Activate a License, Download Gateway Backup), "NETWORKING" (Change Web Server Settings, Enable SSL for the Gateway Network, Create an Email Profile, Manage incoming/outgoing Gateway Network connections), and "SECURITY" (Change General Gateway Security Settings, Create a new user, Assign a user a new role, View the logs of an audit profile, Define a Security Zone, Set access levels on a Security Policy).

Creating a new OAuth2 Client

1. Click on the OAuth2 Clients option. This will bring up a window where you can select an existing OAuth2 Client to modify, or create a new OAuth2 Client. Select "Create new OAuth2 Client..."

The screenshot shows the Ignition software interface. On the left, there is a dark sidebar with various navigation options under three main categories: SYSTEM, NETWORKING, and SECURITY. Under the SECURITY category, the 'OAuth2 Clients' option is highlighted with a blue border. The main content area has a header 'Config > Security > OAuth2 Clients'. A green banner at the top says 'Trial Mode 1:01:21 We're glad you're test driving our software. Have fun.' and a 'Activate Ignition' button. Below the banner, there is a table with two columns: 'Name' and 'Description'. The table shows a single entry: 'No OAuth2 Clients'. At the bottom of the table, there is a red rectangular box containing a blue arrow icon and the text 'Create new OAuth2 Client...'.

2. From here, you will be able to configure your OAuth2 Clients by specifying various properties. Fill out the fields on this page and click "Create New OAuth2 Client".

 Most OAuth 2.0 Authorization Servers will require registering Ignition as a client. After the registration process is complete, the authorization server will generate a client ID and secret for Ignition, which is required below. This gives Ignition the ability to communicate securely with the authorization server. Most authorization servers will also require a set of redirect URIs. The redirect URI for this Ignition Gateway is:
(where {name} is the name of the OAuth2 client)

Main

Name	<input type="text" value="OAuth2"/>
Description	<input type="text"/>
Client ID	The client identifier registered within the authorization server. This value is provided by the authorization server.
Client Secret	The client secret registered within the authorization server. This value is provided by the authorization server.
Client Secret	Re-type client secret for verification.
Authorization URL	<input type="text" value="https://accounts.google.com/o/oauth2/v2/auth"/> Optional URL of the authorization server's OAuth 2.0 Authorization Endpoint.
Token URL	<input type="text" value="https://oauth2.googleapis.com/token"/> Required URL of the authorization server's OAuth 2.0 Token Endpoint.

HTTP Client Settings

HTTP Version	<input type="text" value="HTTP/2"/> The maximum HTTP version supported by the authorization server's token endpoint. The Gateway's HTTP client will use this version when making HTTP requests to the authorization server. (default: HTTP_2)
HTTP Connect Timeout	<input type="text" value="30000"/> In the case where the Gateway's HTTP client needs to establish a new HTTP connection with the authorization server, if the connection cannot be established within the given amount of time in milliseconds, then the connection attempt and any pending request requiring the connection will time out. There will be no timeout if the value is less than or equal to zero (the Gateway will wait forever for the connection to be made). (default: 30,000)
HTTP Request	<input type="text" value="120000"/> The timeout in milliseconds for each HTTP request sent from the Gateway's HTTP client to the authorization server. If an HTTP response is not received within the

Timeout	specified timeout after sending an HTTP request then the HTTP request will time out. There will be no timeout if the value is less than or equal to zero (the Gateway will wait forever for a response). (default: 120,000)
Create New OAuth2 Client	

3. Your OAuth 2.0 Client is now set up and you can now test your token, test authorization, or begin using it.

OAuth2 Client Properties

The reference table below lists the available properties to configure and edit for OAuth2 Clients.

Property Name	Property Description
Main	
Name	The name of your OAuth2 Client. This field is important, as redirect URIs will use the name of the OAuth2 Client. See the note at the top of the OAuth2 Client properties.
Description	A description of your OAuth2 Client.
Client ID	The client identifier registered within the authorization server. This value is provided by the authorization server.
Client Secret	The client secret registered within the authorization server. This value is provided by the authorization server.
Client Secret	Re-type the client secret for verification.
Authorization URL	Optional URL of the authorization server's OAuth 2.0 Authorization Endpoint.
Token URL	Required URL of the authorization server's OAuth 2.0 Token Endpoint.
HTTP Client Settings	
HTTP Version	The maximum HTTP version supported by the authorization server's token endpoint. The Gateway's HTTP client will use this version when making HTTP requests to the authorization server. Choose between <code>HTTP_1.1</code> or <code>HTTP_2</code> . Default is <code>HTTP_2</code> .
HTTP Connect Timeout	In the case where the Gateway's HTTP client needs to establish a new HTTP connection with the authorization server, if the connection cannot be established within the given amount of time in milliseconds, then the connection attempt and any pending request requiring the connection will time out. There will be no timeout if the value is less than or equal to zero (the Gateway will wait forever for the connection to be made). Default is 30,000.
HTTP Request Timeout	The timeout in milliseconds for each HTTP request sent from the Gateway's HTTP client to the authorization server. If an HTTP response is not received within the specified timeout after sending an HTTP request then the HTTP request will time out. There will be no timeout if the value is less than or equal to zero (the Gateway will wait forever for a response). Default is 120,000.

OAuth2 Client Options

After you set up your OAuth2 Client, you should have the option to edit your properties, delete the client, test authorization, or test your token.

Test Authorization

Note: The "test authorization" option will be unavailable if the optional **Authorization URL** configuration setting is blank on the OAuth2 Client's configuration settings.

The "test authorization" tool allows you to test the integration between the OAuth2 Client configured in Ignition with the OAuth2 Authorization Server. More specifically, it allows you to walk through a test run of the **Authorization Code Grant** flow. During this process, the authorization server will validate the end user and may prompt the end user to give the Client permissions to perform actions on their behalf, such as sending emails. Once these permissions have been granted, the authorization server will generate an authorization code, which the Client can use to obtain an access token. To see an example of this process work alongside **testing tokens**, go to the **Test Authorization and Token Example**.

If you want to test the authorization, choose the **test authorization** option.

Config > Security > OAuth2 Clients

Trial Mode 1:10:26 Activate Ignition

Successfully created new OAuth2 Client "OAuth2"

Name	Description
OAuth2	<button>More ▾</button> <button>edit</button>

→ Create new OAuth2 Client...

More options menu:

- delete
- test authorization** (highlighted with a red box)
- test token

You will need to add any **scopes** you want to test with, along with any additional **Request Parameters** you want to try. You can find Request Parameters in your authorization server's documentation. Once those fields are filled out, click on "**Test Authorization**".

Config > Security > OAuth2 Clients > OAuth2 > **Test Authorization** Activate Ignition

Test Authorization

Scope	Scope	Action
Scope	No Scopes found	
	Add Scope	

Additional Request Parameters	Key	Value	Action
Additional Request Parameters	No Request Parameters found		
	Add Request Parameter		

Test Authorization (highlighted with a red box)

Test Token

The **test token** option is useful in cases where you want to test the **authorization code** to make sure it successfully gives the Client an access token. It is also useful if you want to test an access token request using a **client credentials grant**, or refreshing an access token using a **refresh token**. During this process, the Gateway will make an access token request against the authorization server using one of three methods and get a response back. To see an example of this process work alongside **testing authorization**, see the [Test Authorization and Token Example](#).

You can find the **test token** option on the OAuth2 Clients page.

The screenshot shows the Ignition configuration interface under the Security section. A green banner at the top indicates "Trial Mode 1:10:26" and "Activate Ignition". Below it, a message box says "Successfully created new OAuth2 Client 'OAuth2'". The main table lists an "OAuth2" client with a "More" dropdown and an "edit" button. A context menu is open over the "OAuth2" row, with options "delete", "test authorization", and "test token". The "test token" option is highlighted with a red box.

The section below details the **Authorization Code**, **Client Credentials**, or a **Refresh Token** grant types, with the main difference between the three being the type of information you will need to enter.

- **Authorization Code**
 - If you use an Authorization Code, you will need to enter an **Authorization Code**, **Redirect URL**, and any additional **Request Parameters**. Once you input that information, click on "**Test Token**".

The screenshot shows the "Test Token" configuration page for the "OAuth2" client. It has three tabs: "Authorization Code", "Client Credentials", and "Refresh Token". The "Authorization Code" tab is selected. It contains fields for "Authorization Code" and "Redirect URL", and a "Test Authorization" button. The "Client Credentials" tab shows "No Request Parameters found" and a "Add Request Parameter" button. The "Refresh Token" tab is not visible. At the bottom is a large blue "Test Token" button, which is also highlighted with a red box.

- **Client Credentials**
 - If you use Client Credentials, you will need to enter a **scope** and any additional **Request Parameters**. Once you input that information, click on "**Test Token**".

Config > Security > OAuth2 Clients > OAuth2 > Test Token

Trial Mode 0:08:30 Activate Ignition

Test Token		
Authorization Code	Client Credentials	
Scope	Scope	Action
	No Scopes found	Add Scope
Additional Request Parameters	Key	Value
	No Request Parameters found	
Add Request Parameter		
Test Token		

- Refresh Token

- If you use a Refresh Token, you will need to enter the **Refresh Token**, **scope**, and any additional **Request Parameters**. Once you input that information, click on "**Test Token**".

Config > Security > OAuth2 Clients > OAuth2 > Test Token

Trial Mode 0:09:24 Activate Ignition

Test Token		
Authorization Code	Refresh Token	
Refresh Token		
Scope	Scope	Action
	No Scopes found	Add Scope
Additional Request Parameters	Key	Value
	No Request Parameters found	
Test Token		

Test Authorization and Token Example (Google Identity OAuth2)

This example walks you through a "test authorization" and "test token" workflow for an OAuth2 Client configured against [Google's OAuth2 Authorization Server](#).

1. Go to your OAuth2 Clients page. This is located at your **Gateway webpage > Config > Security > OAuth2 Clients**.
2. Set up your OAuth2 Client if you haven't already. For this example, we are using the OAuth2 Client we set up earlier called "OAuth2":

Trial Mode 1:10:26

Activate Ignition

Successfully created new OAuth2 Client "OAuth2"

Name	Description
OAuth2	

→ Create new OAuth2 Client...

More ▾ edit

- delete
- test authorization
- test token

- Click on the "test authorization" option. This will bring up the Test Authorization window. Add any Scopes and additional Request Parameters for your OAuth2 Client, and then press "Test Authorization".

Test Authorization

Scope	Action
https://mail.google.com	Remove

Add Scope

Key	Value	Action
access_type	offline	Remove
prompt	consent	Remove

Add Request Parameter

Test Authorization

The additional Request Parameters in the example are specific to Google's Authorization Server. They do two things:

- access_type: Setting this parameter to "offline" will allow your Client to refresh access tokens while the user is not present at the browser.
- prompt: Setting this parameter to "consent" will force the Client to prompt the user for consent.

- Since we are using Google Identity, we will be redirected, and a Google permission prompt will appear. Click "Allow".

 Sign in with Google

[REDACTED] wants to access
your Google Account

D [REDACTED]

This will allow [REDACTED] to:



Read, compose, send, and permanently delete all [\(i\)](#)
your email from Gmail

Make sure you trust [REDACTED]

You may be sharing sensitive info with this site or app. You
can always see or remove access in your [Google Account](#).

Learn how Google helps you [share data safely](#).

See [REDACTED] Privacy Policy and Terms of Service.

[Cancel](#)

[Allow](#)

- Once permissions are granted to the OAuth2 Client, you will be redirected to the Gateway. Information about the authentication, including the authorization code, timestamps, results, and more will be displayed.

Timestamp	
Request Timestamp:	2023-01-12T18:36:45.963Z
Response Timestamp:	2023-01-12T18:36:50.728Z
Elapsed Time:	PT4.765S
Result:	✓ Success
Request	
Response	
Test Token	

Note: Besides the main OAuth2 Client page, you can use the **Test Token** button on the authentication results page to test access tokens. The benefit of testing your token from the authentication results page however, is that some information, such as the Authentication Code and Redirect URL will already be filled out based off the information from the authorization test.

- Click on the "Test Token" button on the authentication results page. This will bring you to the Test Token page, where some fields will already be filled in.

Test Token

Authorization Code	Authorization Code									
Authorization Parameters	Authorization Code <input type="text"/> Redirect URL <input type="text"/> Test Authorization →									
Additional Request Parameters	<table border="1"> <thead> <tr> <th>Key</th> <th>Value</th> <th>Action</th> </tr> </thead> <tbody> <tr> <td colspan="3">No Request Parameters found</td> </tr> <tr> <td colspan="3">Add Request Parameter</td> </tr> </tbody> </table>	Key	Value	Action	No Request Parameters found			Add Request Parameter		
Key	Value	Action								
No Request Parameters found										
Add Request Parameter										
Test Token										

Click on "Test Token" to see what the token test results will be.

7. A "Token Response" dialog box will appear, showing the results of the token test, as well as timestamps, the response, and more. Click "Done" to close out the dialog box.

Token Response

Timestamp	
Request Timestamp:	2023-01-12T18:58:57.426Z
Response Timestamp:	2023-01-12T18:58:57.641Z
Elapsed Time:	PT0.215S
Result:	Success
Request	
<pre>[REDACTED]</pre>	
Response	
<pre>[REDACTED]</pre>	
Done	

Designer

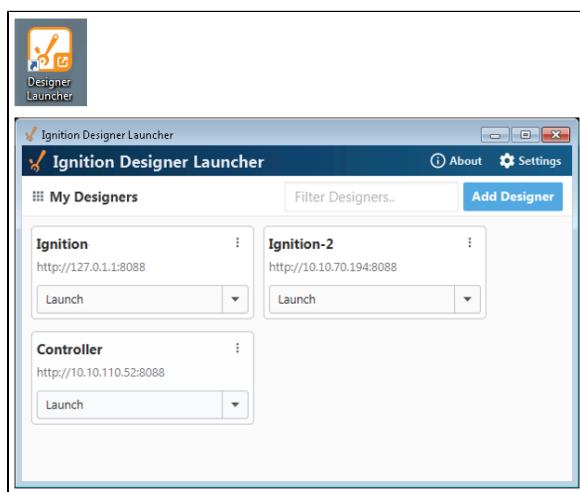
The Ignition Designer is where the majority of configuration and design work is done. The Designer uses web-launch technology to open and edit your projects. This is how you can configure your [Perspective](#) and [Vision](#) projects.

The Designer provides all the firepower to bring your projects to life. It uses a drag-and-drop configuration making screen development quick and intuitive. You can create user interfaces by dragging [Perspective components](#) onto a view, [Vision components](#) onto a window, and Tags onto your components to instantly bind data to tables, charts, and graphs. You can set up Tags and [Transaction Groups](#) to log data to your databases. You can set up [Reports](#) to generate and save data however you'd like. The Designer saves all your projects to the Gateway so everything is controlled in one place.

This page provides some good information about the Designer, Designer workspace, Tools, and how to create a project.

Opening the Designer

Opening the Designer requires running the [Designer Launcher](#). Once the Designer Launcher is installed and configured, you'll have all your projects at your fingertips. If you created a desktop shortcut for the Designer Launcher at install, simply click to open it and select a Designer to open. From there you can create a new project or open an existing project.



On this page ...

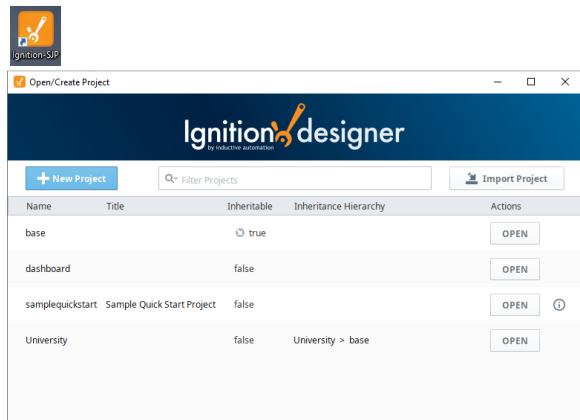
- [Opening the Designer](#)
- [Creating a Project](#)
 - [New Perspective Project](#)
 - [New Vision Project](#)
 - [Project Creation Settings](#)
 - [Open or Create a Project from within the Designer](#)
 - [Updating Project Settings](#)
- [Designer's Workspace](#)
- [Designer Tools](#)
 - [Tools Menu](#)
 - [Status Bar](#)
 - [Previewing the Project](#)
 - [Find and Replace](#)
 - [Keyboard Shortcuts](#)
- [Concurrent Editing and Conflict Resolution](#)
 - [Concurrent Users UI](#)
 - [Conflict Resolution](#)



The Designer Launcher

[Watch the Video](#)

If you have a Designer that has projects that you work on regularly, you can create a shortcut to that Designer and keep it on your desktop. When you click on your shortcut, it opens the Designer, then click on the project you want to edit. Now you can hit the ground running designing your project!

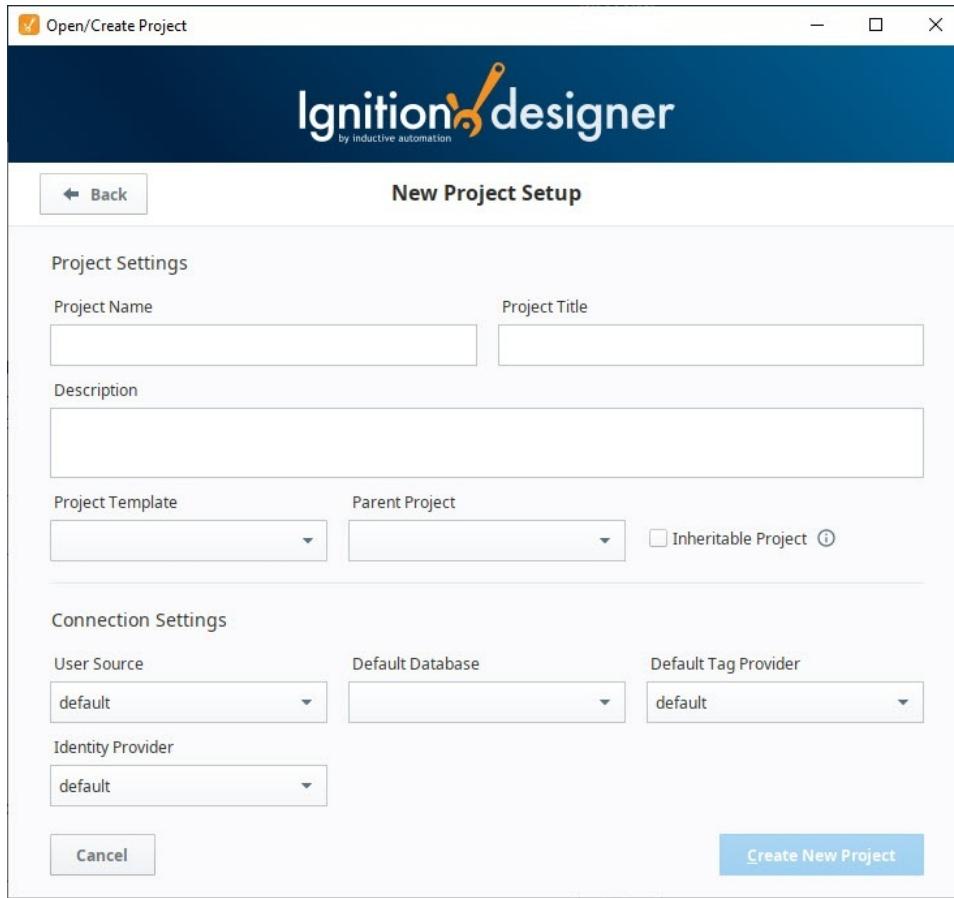


Refer to the [Designer Launcher](#) page for downloading, installing, and configuring the Designer Launcher.

Creating a Project

The first step in working with a project is creating one. Once you launch the Designer, the Open/Create Project window is displayed. Here you have the option to create a new project or open an existing project. Let's create a new project!

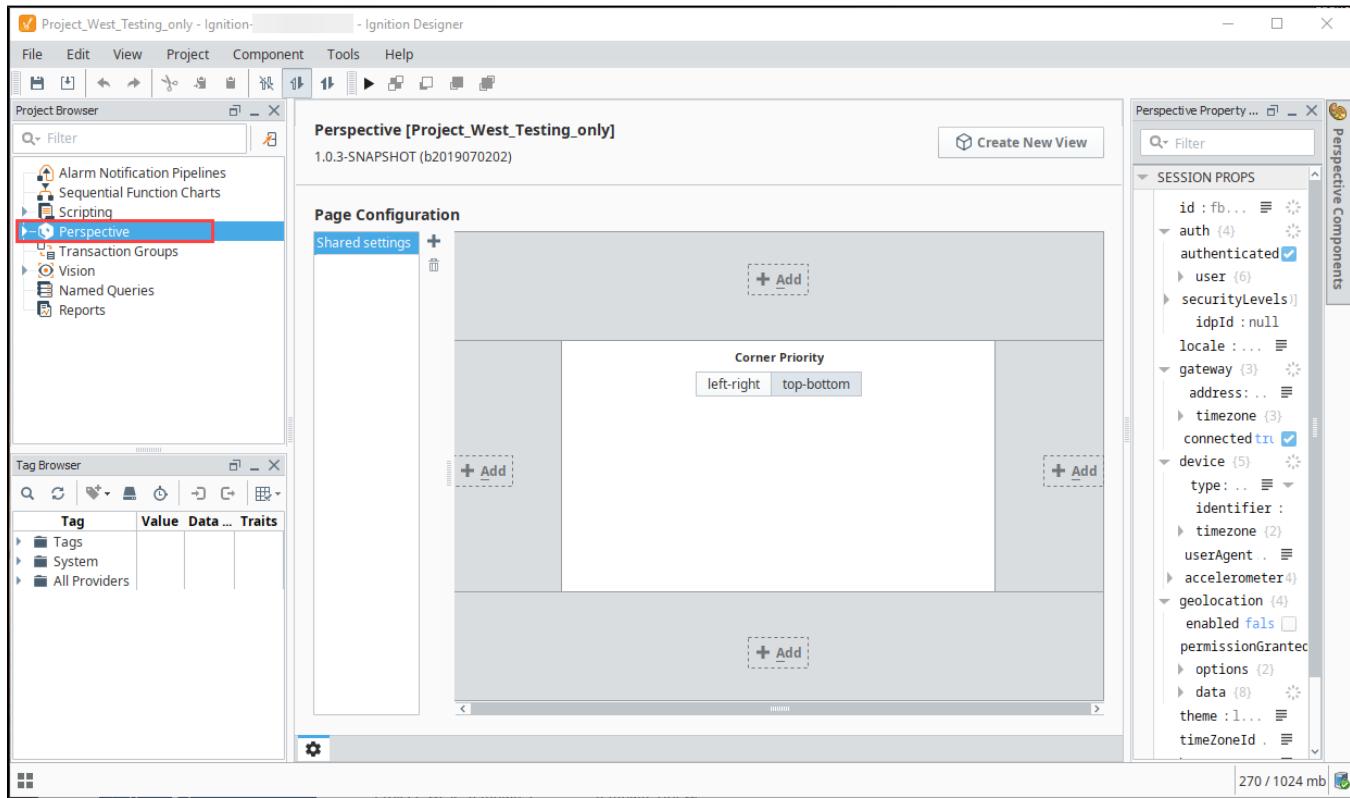
1. After opening the Designer, click on **Create a New Project** button. The following window will be displayed.



2. Enter the **Project Name** (required) and any other configuration settings you need for your project. Most settings are optional. Refer to the [New Project Creation Settings](#) Table below for a description of each property.
3. Click on **Create New Project**.

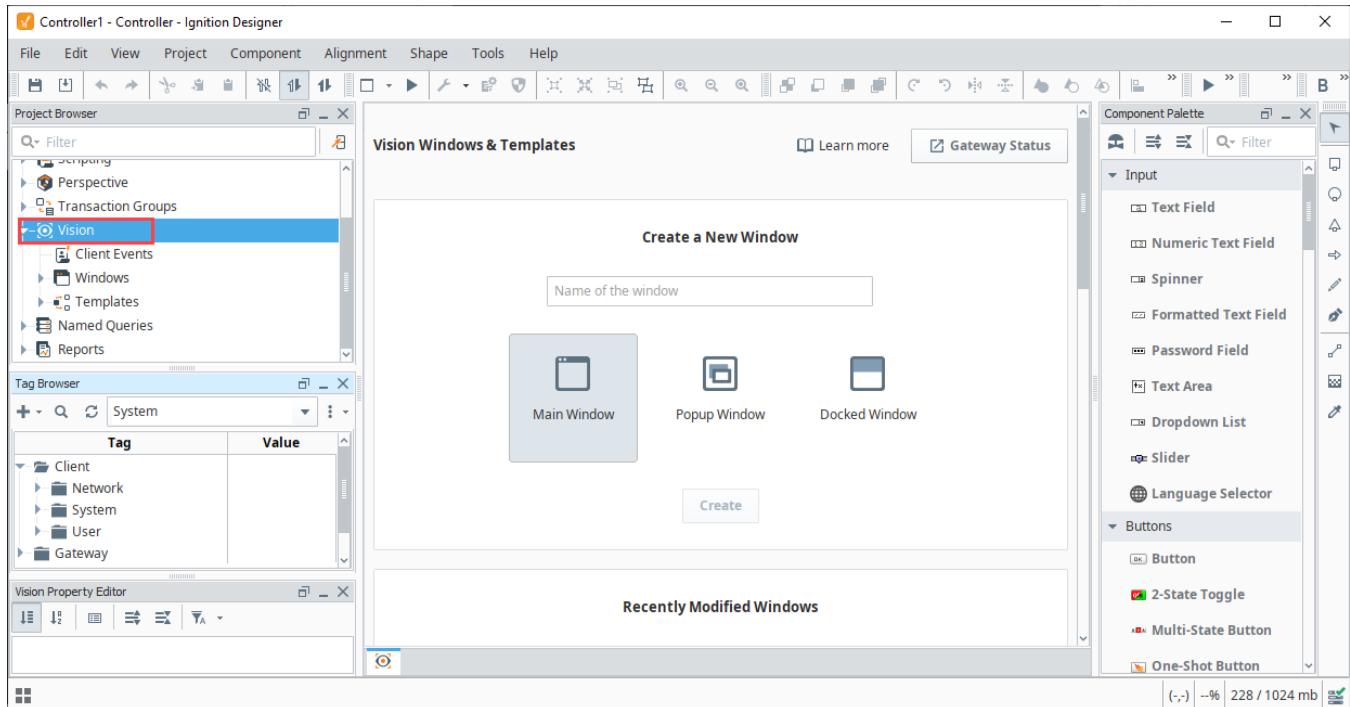
New Perspective Project

If you have the Perspective Module installed, your project will open on the **Perspective Configuration Page** after the project is created. From here, you can begin designing your Perspective project.



New Vision Project

To design your project in Vision, expand the **Vision** module in the Project Browser. The **Vision Getting Started** window will open, allowing you to create windows, add components, and create Client Tags.



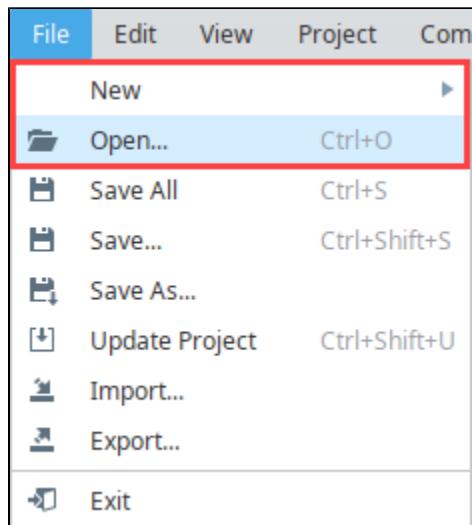
Project Creation Settings

The table below contains descriptions for available settings when creating a new project in the Designer.

Property	Description
Project Name	The Project Name can only consist of alphanumeric characters and the '_' (underscore) character. Spaces and other special characters are not supported. Note: It is not advisable to change the Project Name after it's been created, instead, change the Title property.
Project Title	This is the name that will be displayed on the launch page of the Gateway and in the runtime Client or Session, (optional) There are no restrictions on using special characters or spaces. If you do not specify a title, the project name will be displayed on the Gateway launch page and in the runtime.
Description	Enter a description of the project (optional). Once a project is created, this description can be viewed on the Open/Create Project screen when you hover over the Information  icon.
Project Template	Select a Project Template (optional). There are several pre-built project templates focused on navigation that support either Perspective or Vision. Click the dropdown to see all the available options.
Parent Project	A project may have a parent project, and will inherit all of the resources of that parent project, (optional). Click the dropdown list to see all the available options.
Inheritable Project	Inheritable projects are not runnable as a stand-alone project, but are intended to provide shared resources to one or more child projects.
User Source	Determines the User Source associated with this project (when using the Classic authentication strategy).
Default Database	Select a Default Database (optional). Any queries to the database will use this database connection unless explicitly specified otherwise.
Default Tag Provider	Select a default Tag Provider (optional). If left blank, bindings and references to tags will always need to include the tag provider
Identity Provider	Determines the Identity Provider associated with this project (when using the Identity Provider authentication strategy).

Open or Create a Project from within the Designer

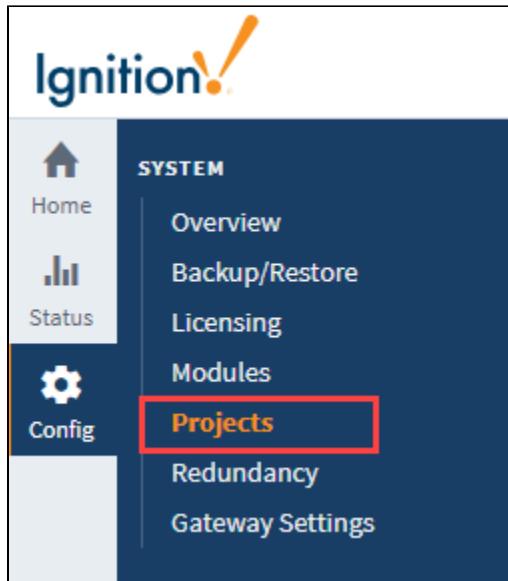
To create a new project or open a different project from within the Designer, use the **File > Open** menu in the top menubar.



The Open/Create Project screen will be displayed. You can choose from existing projects or create a new project.

Updating Project Settings

Project settings such as the title, description, connections, and inheritance are set through the Gateway Webpage **Config Tab**, under **System > Projects**. For more information, see [Projects](#).

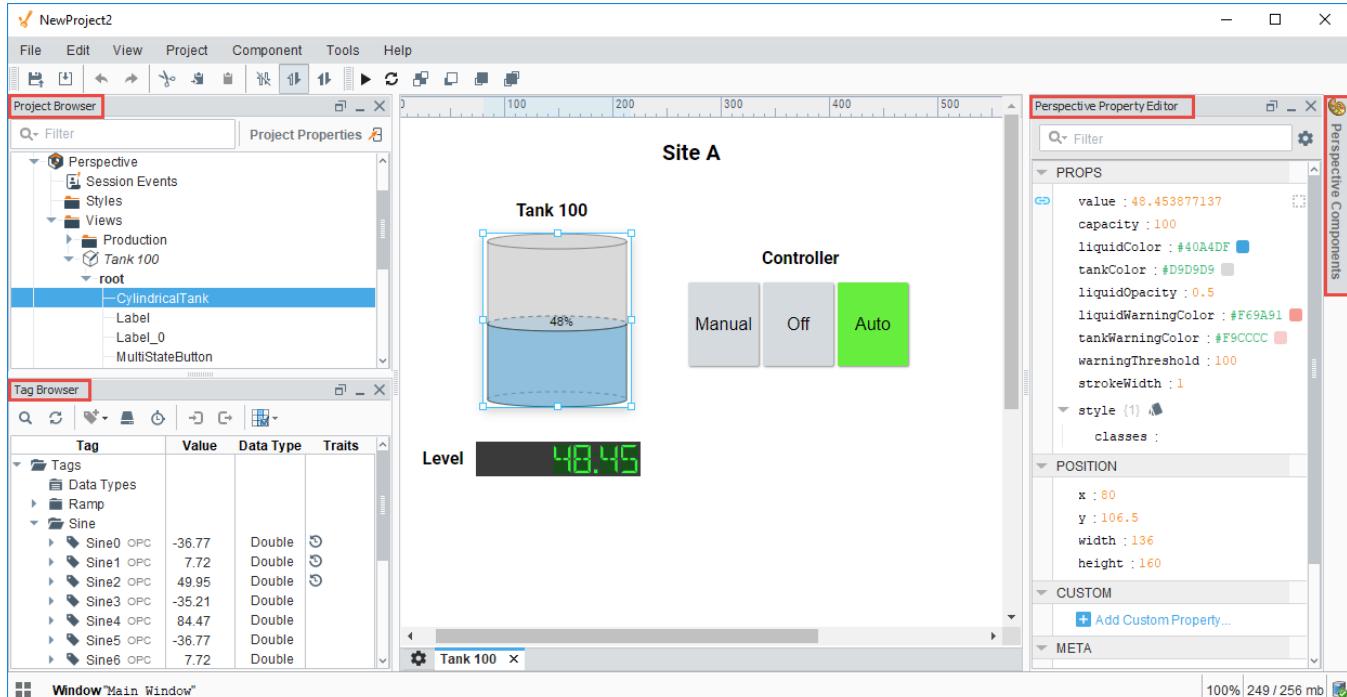


Designer's Workspace

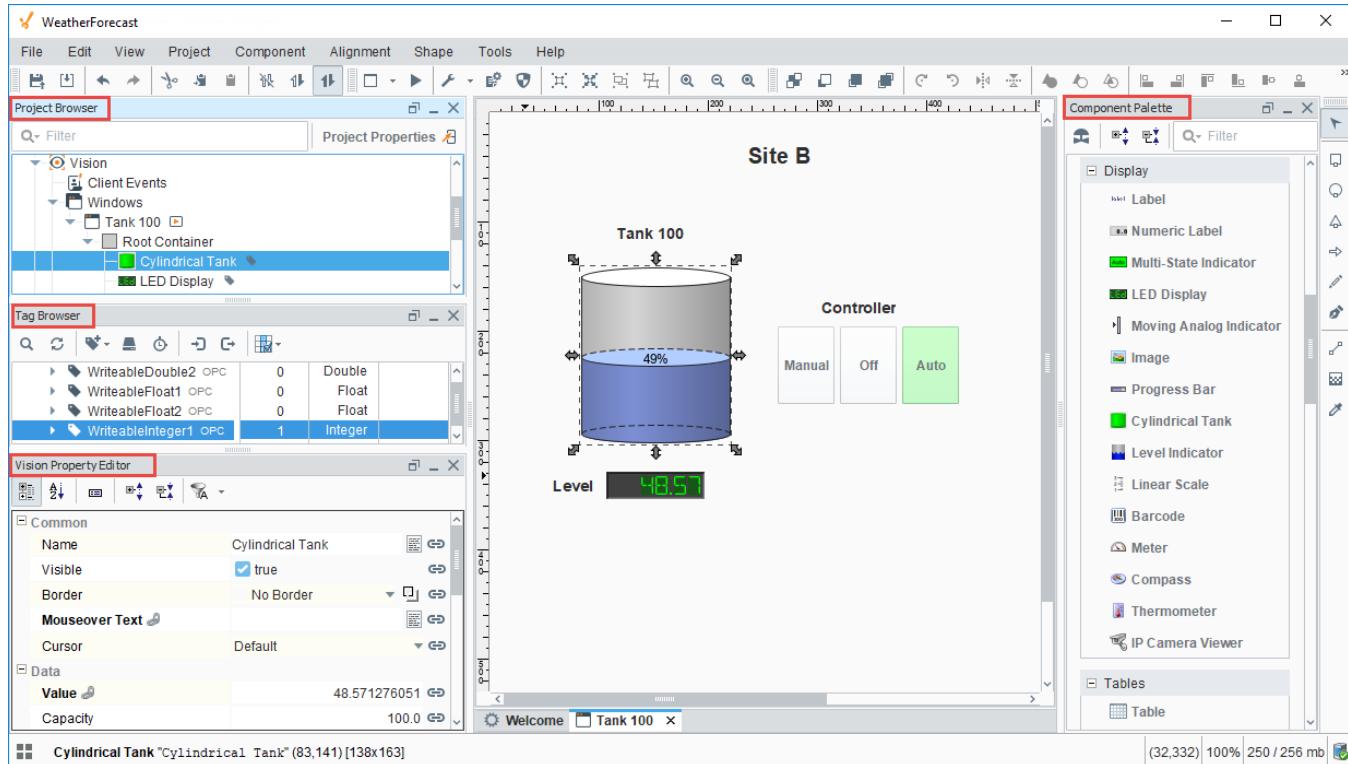
The Designer workspace is centrally located and organized by panels. Some of the panels include a Project Browser, Tag Browser, Component Palette, and Property Editor. These panels can change depending on the type of resource you are currently editing. For example, if you are editing a Perspective view or Vision window, the Designer workspace has Component Palette and Property Editor panels. If you're editing an Alarm Notification Pipeline, your Designer workspace will be the Pipeline Block Editor. If you're editing a Report, your Designer workspace will be the Report Designer. Each type of workspace has panels that are only valid when that workspace is active.

Here are two images showing the Perspective Designer workspace and Vision Designer workspace, and an example displaying the same components. At a glance, they look very similar, but there are some differences, including each having their own Component Palettes. To learn more, refer to the [Perspective Designer Interface](#) and the [Vision Designer Interface](#) pages.

Perspective Designer's Workspace



Vision Designer's Workspace



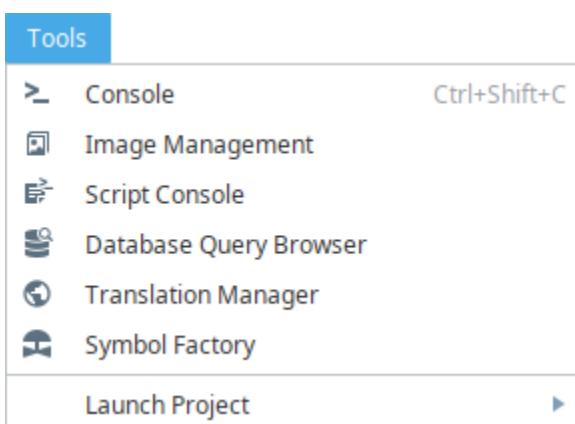
Designer Tools

The Designer has a host of tools to help you accelerate building, testing, and deploying your project. Let's talk about a few here.

Tools Menu

In addition to all the panels available in the Designer workspace, there is also a [Tools menu](#) to help you create your projects.

- **Console** - The [Output Console](#) is the script-writer's best friend, and most frequently used to test and debug Python scripts in Ignition.
- **Image Management** - The [Image Management](#) tool manages and stores images used for your projects.
- **Script Console** - The [Script Console](#) is used to test and debug Python scripts.
- **Database Query Browser** - The [Database Query Browser](#) is a very convenient tool that lets you make simple selects and edits in a database table, and interact with all of the databases that Ignition is connected to (i.e., running queries, browsing tables and schemas). It is very common during the course of project design to inspect the database directly, or to experiment with a SQL query to get it just right.
- **Translation Manager** - Opens up the [Translation Manager](#) panel, and allows you to configure language translations. See also: [Localization and Languages](#)
- **Symbol Factory** - The [Symbol Factory](#) contains a variety of high quality vector graphics symbols that can bring your projects to life.



Status Bar

The Status Bar at the bottom of the Designer workspace allows you to view and adjust certain Designer settings.

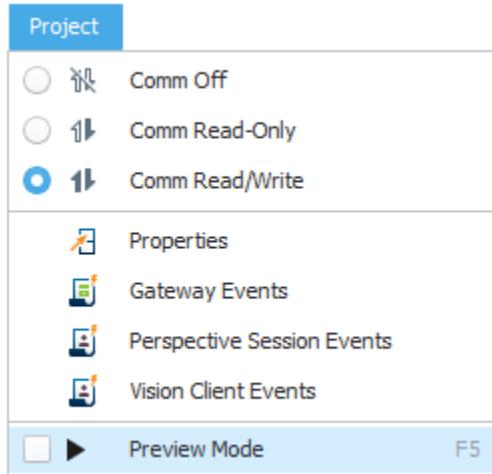


- **Panel Chooser** - Clicking the Panel Chooser icon opens a menu that allows you to open and close Designer panels, as well as to reset all panels to their default positions.
- **Current Window** - Displays the name of the currently open Window. (Vision only)
- **Mouse Coordinates** - Displays the coordinates of your mouse relative to the open Window. (Vision only)
- **Zoom Level** - Changes the zoom level. You can adjust the zoom level between a minimum of 25% and a maximum of 500%.
- **Memory Usage** - Displays your used/maximum available memory. You can change the maximum in the server configuration.
- **Gateway Connection** - This icon shows the current status of your Gateway connection. Clicking on the icon opens the connected Gateway in a browser.

Previewing the Project

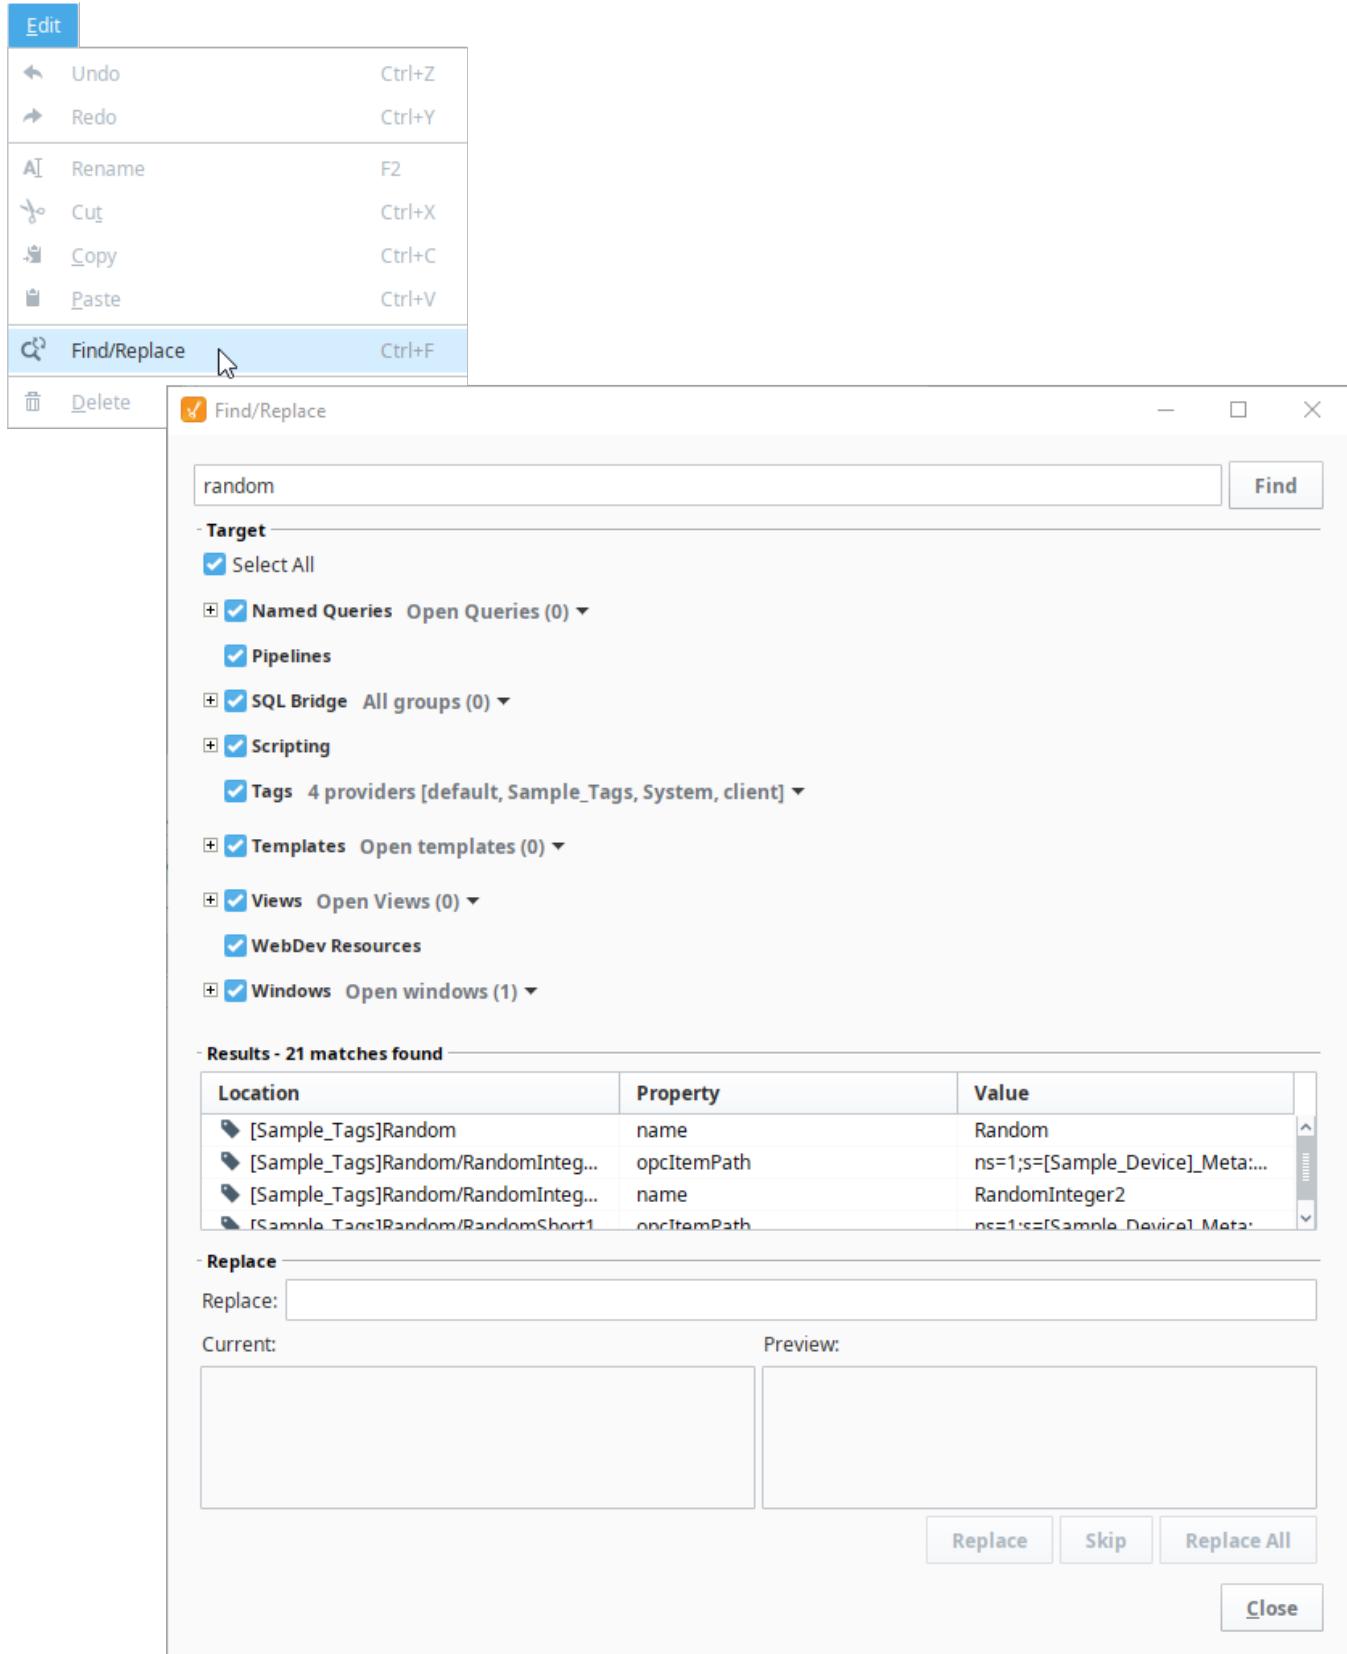
The Designer provides the capability to preview, test, and interact with the screens and functionality before you deploy your project. The Designer workspace operates in two distinct modes: Design mode and [Preview mode](#). Designers can easily switch between these modes to make sure their project is working as expected during the course of development by simply clicking **Preview Mode** icon from the top menubar, or clicking **Project** also from the top menubar, and selecting **Preview Mode**.

Note: The **Preview Mode** icon toggles to be a **Design Mode** icon when it is pressed.



Find and Replace

The [Find and Replace](#) is a handy tool in the Designer workspace. You can search your entire project for specific components, properties, and scripts. You can even use the replace command to make mass changes expeditiously to a project with very little effort.



Keyboard Shortcuts

There are a lot of ways to speed up your development once you are familiar with how Ignition works. There are many [Keyboard shortcuts](#) in Designer that are listed throughout the Designer interface alongside menu options.

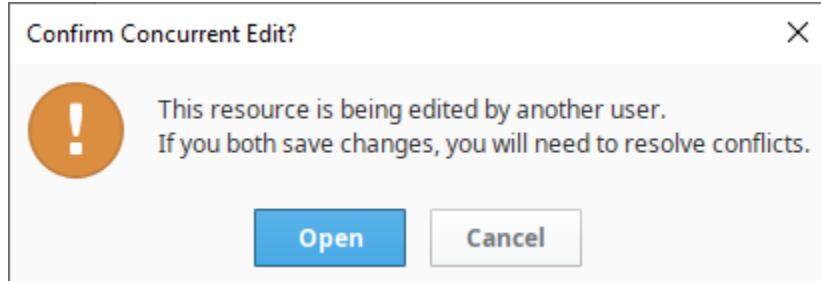
Concurrent Editing and Conflict Resolution

The Ignition Designer uses a lock-free strategy for handling concurrent editing. Multiple people can work together and make changes to a project at the same time. An unlimited number of Designers can be open concurrently, and modifying any resource in the Designer doesn't lock it. The Designer keeps track of the resources that are being edited, and any conflicting edits will be resolved at the time the project is saved. The developer who creates a saving conflict will be prompted to resolve the conflict by choosing whether to use their own changes, accept the other designer's changes, or cancel their save and figure out what to do in another way.

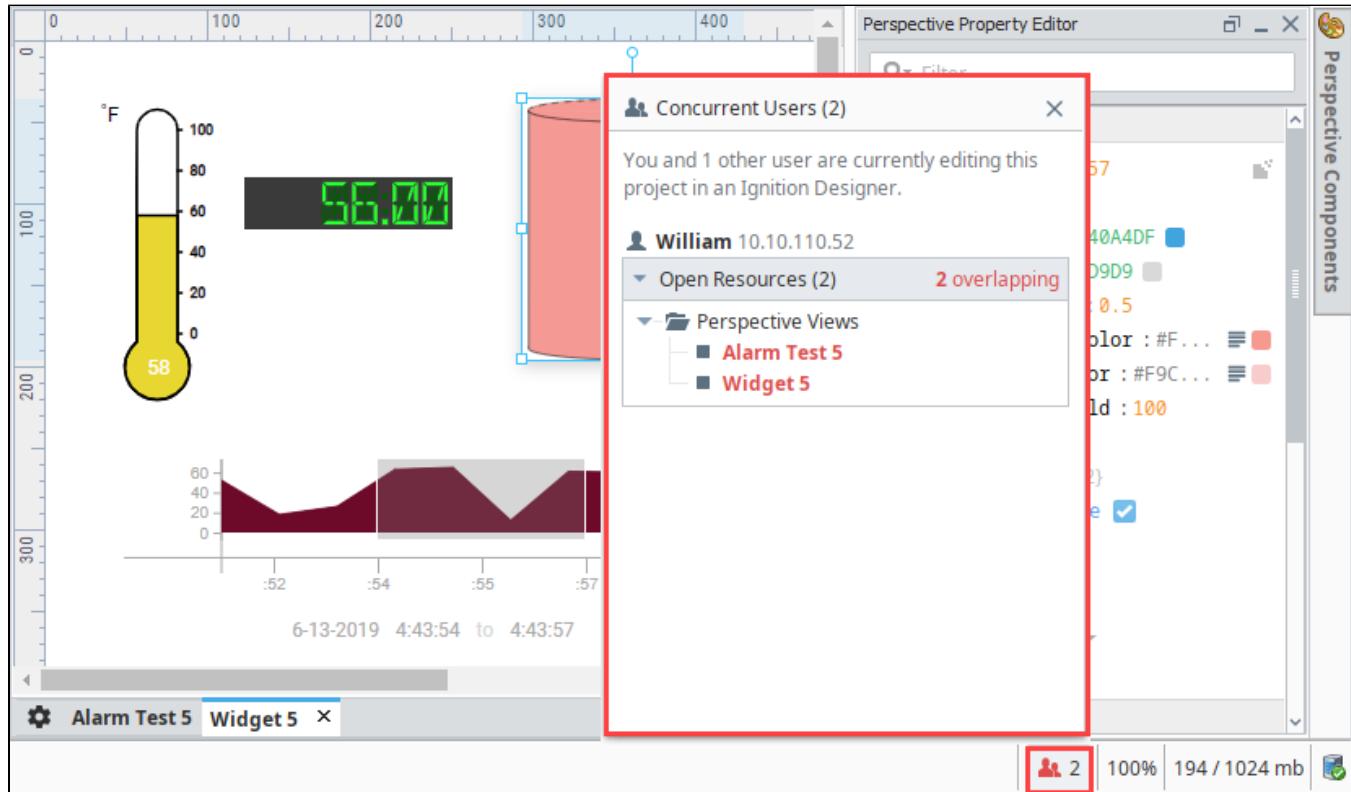
Concurrent Users UI

The Concurrent User Interface allows users to see which project resources are open in other Designer instances, the names of the users that have them open, and when a project update is available. The UI is located in the lower right corner of the Designer.

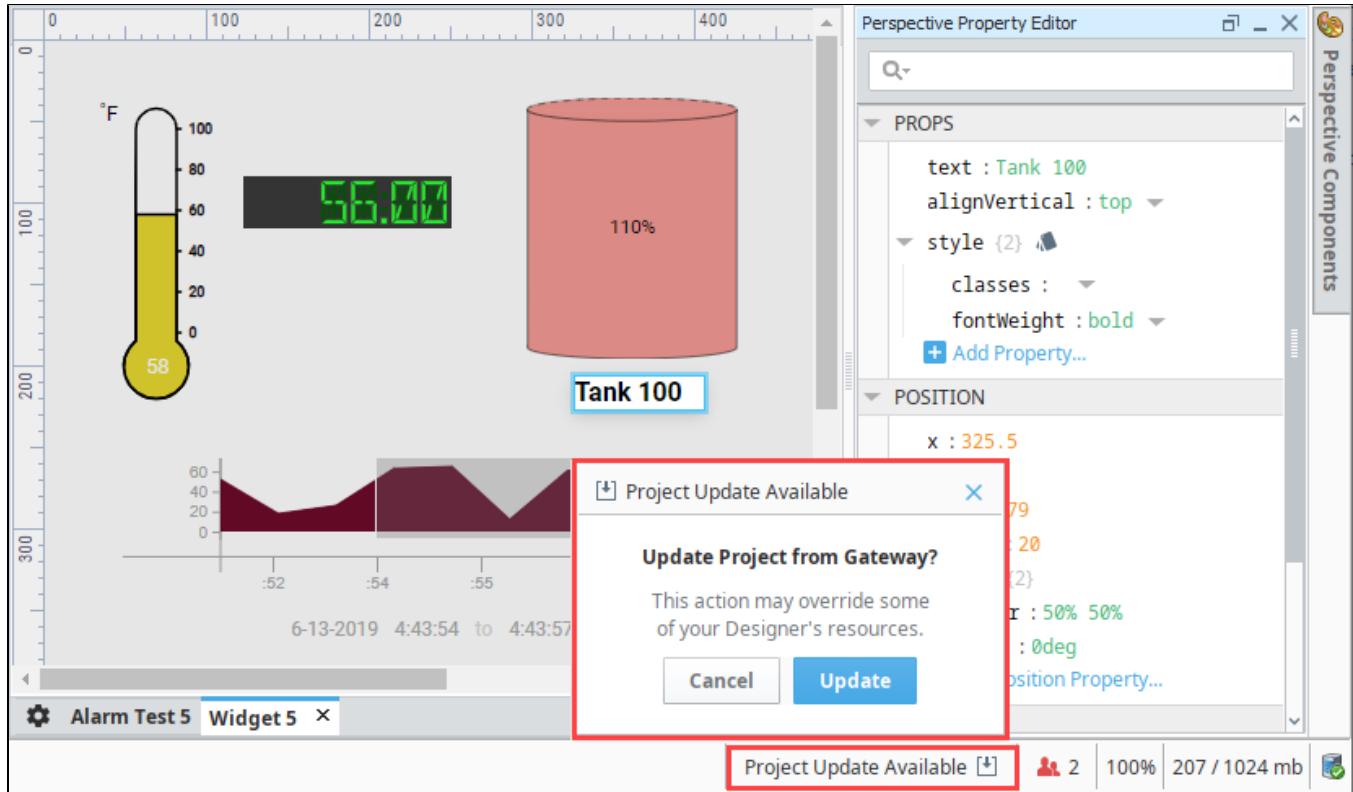
If you open a resource that is already open in another Designer, you will be greeted with a new popup confirming that you want to proceed.



The UI makes use of the color red to denote when there are conflicts with the changes in another Designer instance, or if their users are saving on overlapping resources (multiple users are making changes and saving while the same resource is open).

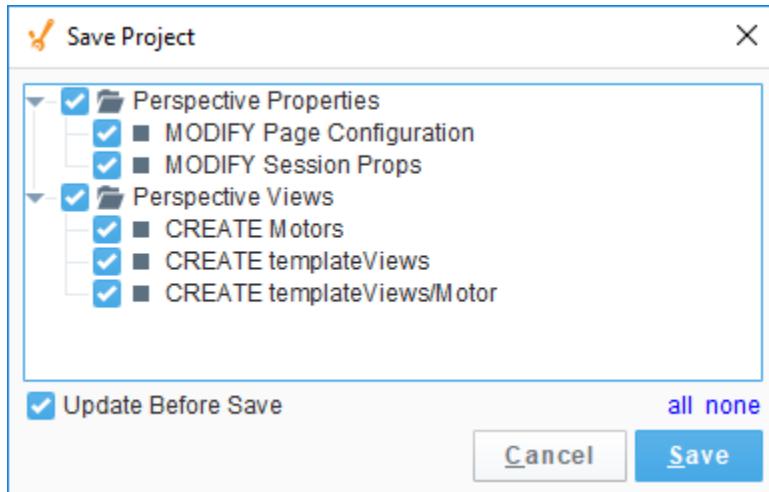


If the user in the other Designer instance saves their project, the UI will display that a project update is available. Click on **Project Update Available** to either update or cancel the project update.



Conflict Resolution

When you're ready to save your project, go to the Menubar and select **File Save**. When the Save Project dialog pops up, select the items you want to save and click the **Save** button.



If there are any conflicts, the Designer will ask you to update your project and the conflict resolution screen will open automatically. You will see a thumbnail image of both screens and the .json code showing the conflict. To resolve the conflict, you will be given a choice; accept your changes, accept the other developer's changes, or cancel the save and contact the other developer. Buttons are provided at the top and bottom of the screen for you to enter how you want to resolve the conflict.

The Resource Tree displays the conflicts and the resolutions. The selected conflict will be highlighted in blue. The image below shows changes were made to the Thermometer. The line number and the code are highlighted at the bottom so you can quickly identify the conflict. Once the conflict is resolved, a checkmark will appear in the Resolution column next to the conflict and how it was resolved. You'll notice that there is second conflict. To resolve the conflict, select it in the Resource Tree or click the arrow to navigate to the next conflict.

Resolve Conflicts

Resolve resource conflicts (2)
Resources that you have changed have also been changed remotely. Review each changed resource and resolve the conflict.

Resource	Resolution
Perspective Views	
Tower Template	<input checked="" type="checkbox"/> This Designer
ViewX	Undecided

Tower Template 2 files

This Designer | Gateway | ← | 1 OF 2 | →

This Designer

```

    graph TD
        FanSpeed[Fan Speed] --> Temp[Temp]
        Temp --> FanSpeed
    
```

Gateway

```

    graph TD
        FanSpeed[Fan Speed] --> Temp[Temp]
        Temp --> FanSpeed
    
```

This Designer

```

        },
        "props": {
            "unit": "F",
            "valueFontColor": "#FFFFFF00"
        },
        "type": "ia.display.thermometer"
    },
    {
        "meta": {
            "hasDelegate": true,
            "name": "Dropdown"
        },
        ...
    }
    }
}
    
```

1 change

Gateway

```

        },
        "props": {
            "valueFontColor": "#FFFFFF00"
        },
        "type": "ia.display.thermometer"
    },
    {
        "meta": {
            "hasDelegate": true,
            "name": "Dropdown"
        },
        ...
    }
    }
}
    
```

Use all from:

This Designer Gateway

Cancel Apply

Editor notes are only visible to logged in users



Related Topics ...

- [Designer Launcher](#)
- [Perspective Designer Interface](#)
- [Vision Designer Interface](#)

In This Section ...

General Designer Interface

Designer Spaces

The Designer has a lot of panels and menus that allow you to build out a project tailored to your needs. However, while some of these like the File Menu are shared throughout the Designer, some of the panels and menu options are specific to certain objects and will typically only be displayed when an object of that type is selected. For example, when editing a Vision Window, the Designer has the Property Editor and Component Palette panels, but when editing an Alarm Pipeline, the Pipeline Blocks and Pipeline Block Editor panels are displayed instead. This creates different Designer Spaces that are used for different areas of a project.

There are many dockable and draggable panels that surround the workspace, as well as the familiar menu bars and toolbars. The dockable panels can be rearranged as you wish and will snap into place as you move them around the screen. Each workspace remembers its layout, which is the docking arrangement of the panels around it.

Docking System

The Designer's docking system provides a very flexible user interface allowing you to customize the layout as you wish. To rearrange the dockable panels, simply drag on their title bars. As you drag the panel, it will try to snap into place and show you a highlighted border. Use the highlighted border that appears to gauge where the panel will be moved to. Hold the CTRL key as you drag these panels to keep them from snapping into place. You can also drag these panels outside of the Designer or onto a second monitor.

Dockable panels can be in one of four modes:

1. **Docked** - The panel is visible, and located somewhere around the perimeter of the workspace. If two panels are docked in the same location, a tab strip will appear to switch between the two panels.
2. **Floating** - A panel can be dragged outside of the workspace perimeter to be floated. The panel can now be positioned anywhere on your desktop.
3. **Pinned** - Pinning a panel makes it minimize to one of the four sides of the Designer, represented by a small tab. Hover over the tab to use the panel.
4. **Hidden** - A hidden panel is not shown. You can open it again by selecting it in the **View > Panels** menu.

Toolbars can also be rearranged and floated to your liking. Simply drag on the "textured" left edge of the toolbar. If you re-arranged your panels into a layout that you don't like, you can quickly revert back to the default by selecting the **View > Reset Panels** option from the menu bar.



Expert Tip

Your docking preferences are stored under %USER_HOME%/.ignition/*/.layout. If you really want to reset your preferences, remove these files and restart the Designer.

Project Browser

The **Project Browser** panel allows you to view the different Designer Spaces, and their component hierarchies at a glance. By clicking **Project Properties**  icon you can view or change many of the properties settings in your project. You can expand the folders and navigate down through each folder to see elements of your project such as windows, views, containers, and components. The Project Browser shows the entire tree structure from the project level folder down to the component level. You can navigate within your project in two ways, by selecting nodes in the Project Browser tree or directly clicking on an element like a component in the project workspace. Regardless of how you select an element, properties for that element

On this page ...

- [Designer Spaces](#)
- [Docking System](#)
- [Project Browser](#)
 - [Project Resources](#)
 - [Right-Click Menu](#)
- [Tag Browser](#)
- [Menubar](#)
 - [File Menu](#)
 - [Project Menu](#)
- [Comm Mode](#)
 - [Previewing the Project](#)
 - [Tools Menu](#)
- [Help Menu](#)

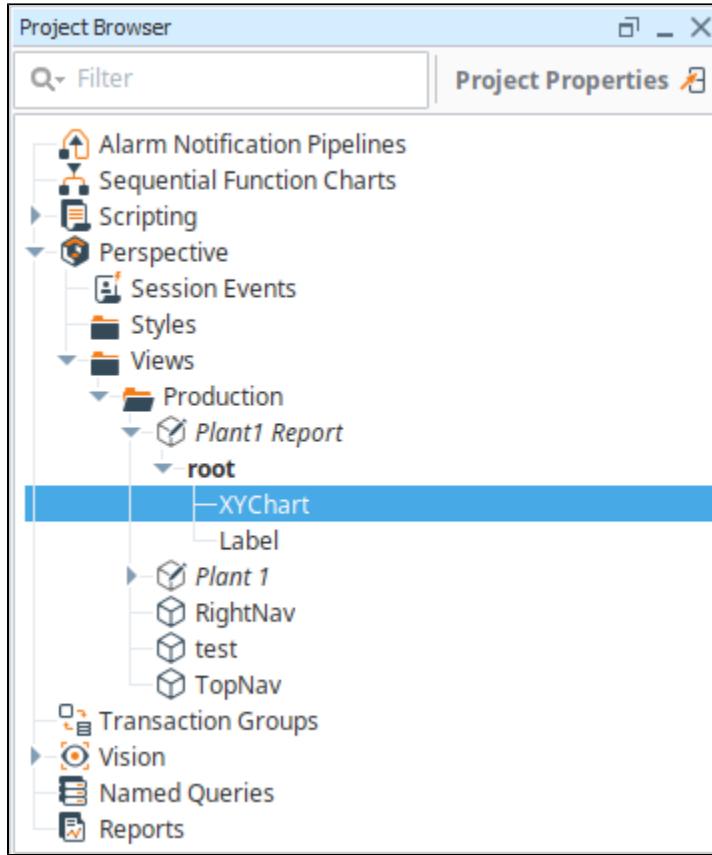


INDUCTIVE
UNIVERSITY

The Designer User Interface

[Watch the Video](#)

are displayed in the Property Editor.



Project Resources

All projects have resources, such as Vision Windows, Vision Templates, Perspective Views, project scripts, reports, and named queries. Project resources can be inherited from one project to another. For more information, see [Project Inheritance](#).

Naming Conventions

Project resource names cannot be blank. They must start with a letter, a numeral, or an underscore. The following characters are reserved and cannot be used in names for project resources.

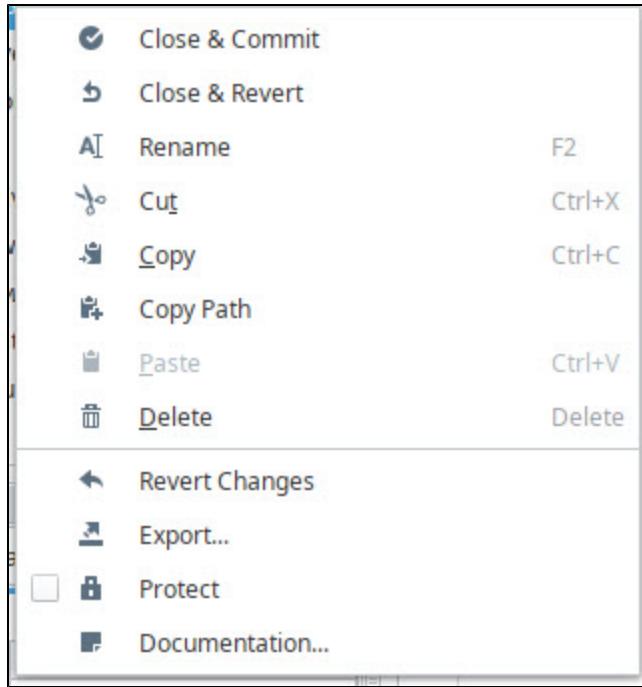
```
< (less than)
> (greater than)
: (colon)
" (double quote)
/ (forward slash)
\ (backslash)
| (vertical bar or pipe)
? (question mark)
* (asterisk)
```

In addition, the following names cannot be used as project resource names:

CON, PRN, AUX, and NUL
COM1, COM2, COM3, COM4, COM5, COM6, COM7, COM8, and COM9
LPT1, LPT2, LPT3, LPT4, LPT5, LPT6, LPT7, LPT8, and LPT9

Right-Click Menu

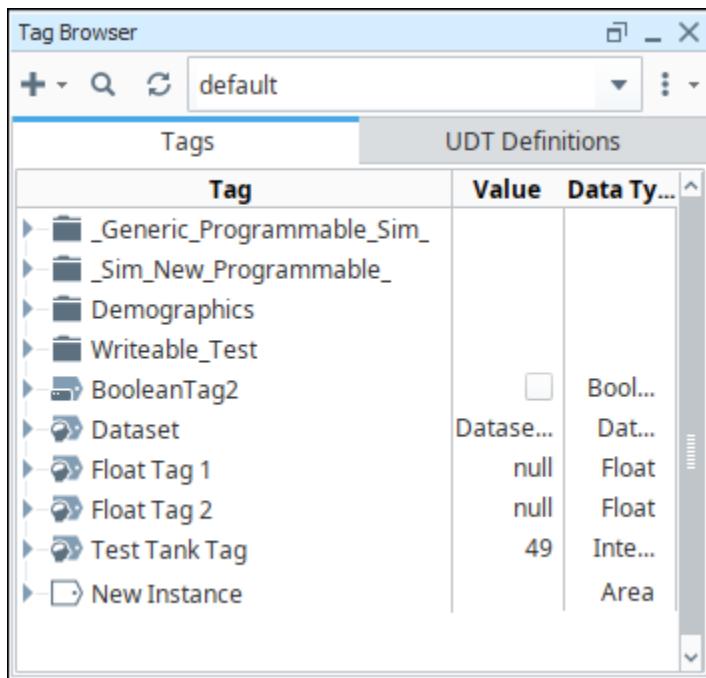
The Project Browser basic functionality is similar to many applications that run on your PC. When you right-click on a project resource in the Project Browser, several options are available like Copy, Paste, and Delete. The presence of certain options will change depending on the type of project resource you are currently editing. For example, [inherited project resources](#) have additional menu options that are only visible in inherited projects. The table below lists and describes common right-click menu options:



Function	Description
Close & Commit	Saves the updates made to the project resource and closes it.
Close & Revert	Reverts the project resource to its last saved state. Reverting an unsaved resource is the same as deleting it.
Configure View Permissions	View permissions determine whether or not a user has access to load the view.
Configure Events	Configures an event or action on a component.
Rename	Change the name of the current selection to a new name.
Cut	Cuts the current selection into the clipboard.
Copy	Copies the current selection into the clipboard.
Copy Path	Copies the path of the current selection into the clipboard.
Paste	Pastes the contents in the clipboard to the selected content.
Delete	Deletes the current selection.
Revert Changes	Reverting a modified resource will revert it back to its last saved state, assuming it's been saved before. Reverting new resources deletes them.
Export	Exports resources from the project.
Protect	Protects the Project Resource. For more information, see Protecting Project Resources .
Documentation...	<p>The following feature is new in Ignition version 8.1.19 Click here to check out the other new features</p> <p>Documentation can be added to all project resources. Previously, documentation was only available on Vision Windows, Vision Templates, and Transaction Groups.</p>

Tag Browser

The **Tag Browser** panel, located on the left side of the Designer workspace under the Project Browser, allows you to browse Tags in the Designer and OPC servers. Tags can be created, edited, displayed, exported, and imported directly from the Tag Browser. See [Tag Browser](#) for more detailed information.



Menubar

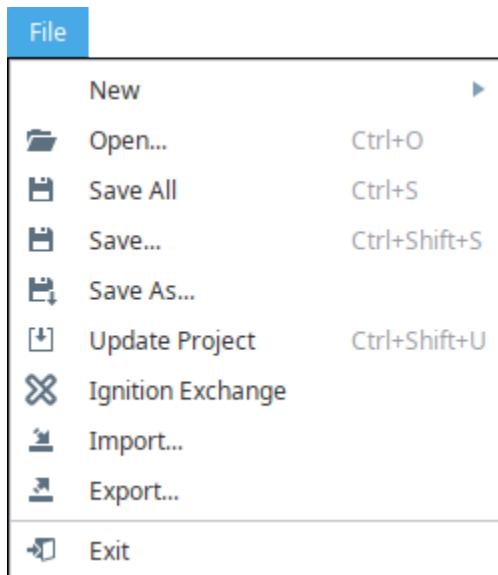
This section addresses only the functionality that the Designer menubar shares between both the Perspective and Vision modules.

For information on the menubar tabs unique to the Vision module, see [Vision Designer Interface](#).

For information on the menubar tabs unique to the Perspective module, see [Perspective Designer Interface](#).

File Menu

The **File Menu** contains most of your basic options related to saving, similar to many other software applications on your PC.

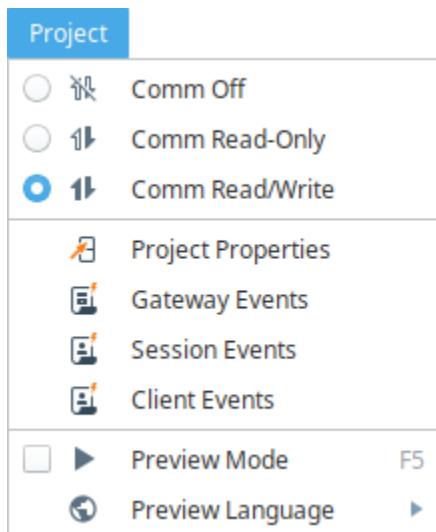


Function	Description
New	The New menu option allows you to create a new Object. That Object could be a new window, a new template, or even a new project. It can vary depending on what modules you have installed, also providing the ability to make new transaction groups,

	reports, Sequential Function Charts, and more.
Open	Allows you to Open a previously created project.
Save All	Saves all project changes, but are not pushed out to the client.
Save	Save brings up a popup menu where you can choose what to save.
Save As	Saves the current project as a new project.
Update Project	Merges any changes from the Gateway to the project, if for example there were multiple people working on the same project.
Ignition Exchange	Provides access to resources, templates, and tools that can be shared with various other industries and used in your own Ignition projects.
Import	Allows you to import or export specific resources into or out of the project such as Tags, scripts, and templates, to name a few.
Export	Exports globally scoped resources, such as Alarm Pipelines.
Exit	Exits the project and allows you to save project changes.

Project Menu

Communication between the Designer and the Gateway is controlled from the Project Menu. The Designer offers three data communication modes for your projects: **Comm Off**, **Comm Read-Only**, or **Comm Read/Write**. Comm Read-Only is the default mode which does not allow writing from the Designer to Tag or Database sources.





**INDUCTIVE
UNIVERSITY**

**About
Communication
Modes**

[Watch the Video](#)

Function	Description
Comm Off, Comm Read-Only, Comm Read/Write	Changes the communication mode for this designer session. See Comm Mode for more details. <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> Note: These settings do not affect the execution of a project's Transaction Groups because Transaction Groups execute on the Gateway, not in the Designer </div>
Properties	Opens up the Project Properties window, allowing project settings to be changed. See also: Project Properties
Event Scripts	Opens up the appropriate event script window, either client, session, or Gateway. These can also be accessed from the Project Browser. See also: Client Event Scripts and Gateway Event Scripts .
Preview Mode	Puts the Designer into Preview Mode, allowing you to interact with it like a client. See also: Previewing the Project .
Preview Language	Determines the language that the Designer will revert to when in Preview Mode. See also Localization and Languages and Localization in Vision .

Comm Mode

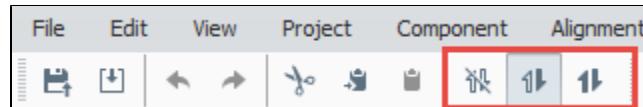
The Designer, and Vision Clients, have a communication mode ("comm mode") that determines if the Designer/Client is able to read and write to tags and databases. The comm mode is determined per Designer/Client instance, and changing the comm mode in one running instance will not impact any other running instance's comm mode setting. The comm mode can be set to the following:

- **Comm Off:** In this mode all database query traffic and Tag subscriptions and writes are blocked. This can be useful to temporarily disable polling component bindings.
- **Comm Read-Only:** Tag subscriptions and SELECT queries work, but Tag writes and UPDATE/INSERT/DELETE queries requested by the Designer/Client are ignored by the Gateway.
- **Comm Read/Write:** The Designer/Client can write to tags. In addition UPDATE/INSERT/DELETE queries requested by the Designer/Client will be processed by the Gateway.

In addition to changing the comm mode from the Project Menu, it can be changed in the following ways.

Setting the Comm Mode from the Main Toolbar

If the Main Toolbar is enabled, go to **Project**, and you'll see a corresponding button for each Comm Mode. The currently selected Comm Mode will have its button highlighted in gray.

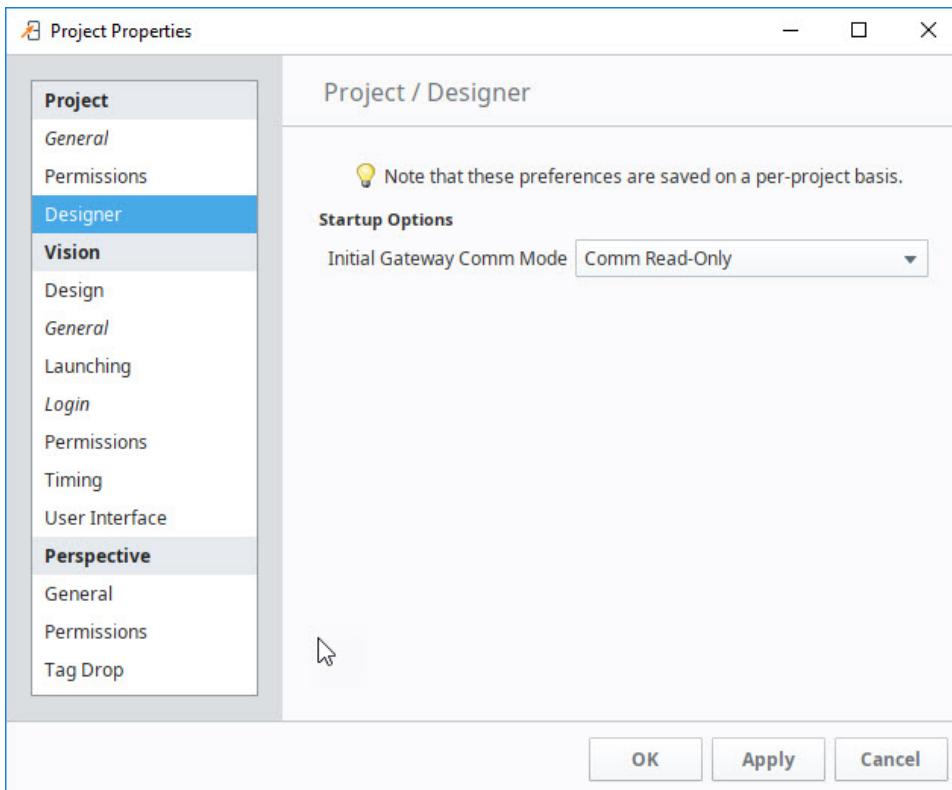


Setting the Comm Mode from the Project Properties Window

This will set the default Comm Mode that the Designer starts up in for the current project.

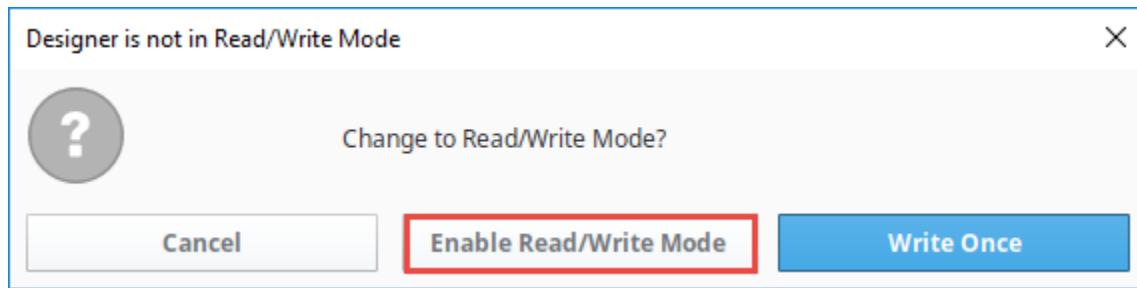
1. In the Designer, either double-click on the **Project > Properties** node in **Project Browser**, or click on the **Project > Properties** command on the top menu.
The **Project Properties** window is displayed.
2. Go to **Project > Designer**.
3. Under **Startup Options**, for the **Initial Gateway Comm Mode**, choose from the dropdown **Comm Off**, **Comm Ready-only**, or **Comm Read /Write**.

Note: These property settings are saved on a per-project basis.



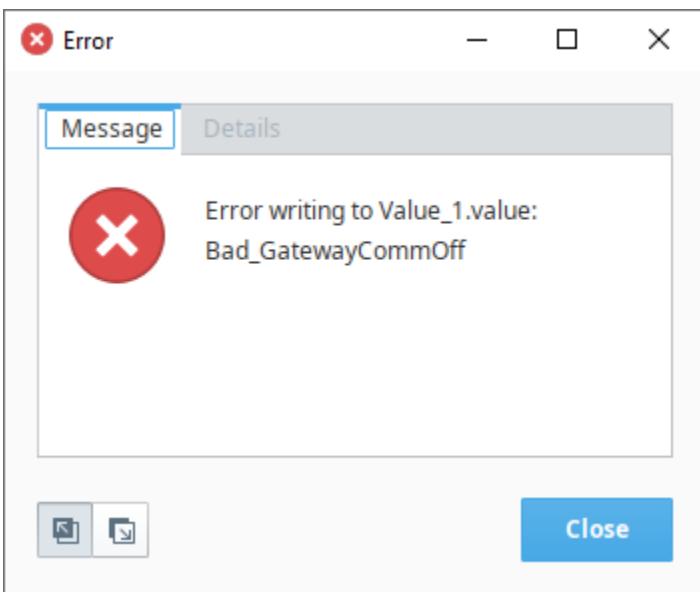
Communication Error Message

If a user is attempting to write to a project in the Designer or from a Client, and the Comm Mode is not enabled for Read/Write, a dialog box will popup stating the Designer is not in Read/Write Mode.



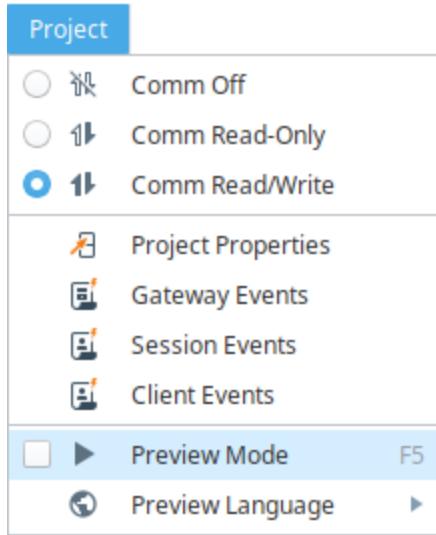
If your Tag is not being written too, there are a couple places to check:

- From Main Toolbar go to **Project**, and make sure the Comm Mode is set to Read/Write.
- If Comm Read/Write is checked, you may have a Client Event Script enabled that could be preventing you from writing to a project or a Tag, as shown in the error message example below. You may need to edit your Client Event Script.



Previewing the Project

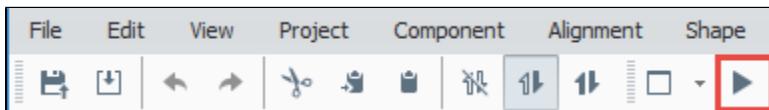
Many times, it is useful to test the components on the screen to ensure that certain bindings or scripts are working the way that was intended. The Designer can go into **Preview Mode** that will allow you to interact with the currently opened window as though you were working in a client. This means that instead of clicking between components and seeing their properties, you will be able to interact directly with the components such as clicking on a button to execute its script, or entering a value into a text field to update a Tag value.



Preview Mode

The window workspace operates in two distinct modes: Design Mode and Preview Mode. There are three different ways you can switch between Design and Preview Mode:

- One of the easiest ways is from the Main Toolbar using the **Preview**  icon or the **Design**  icon to switch between modes.



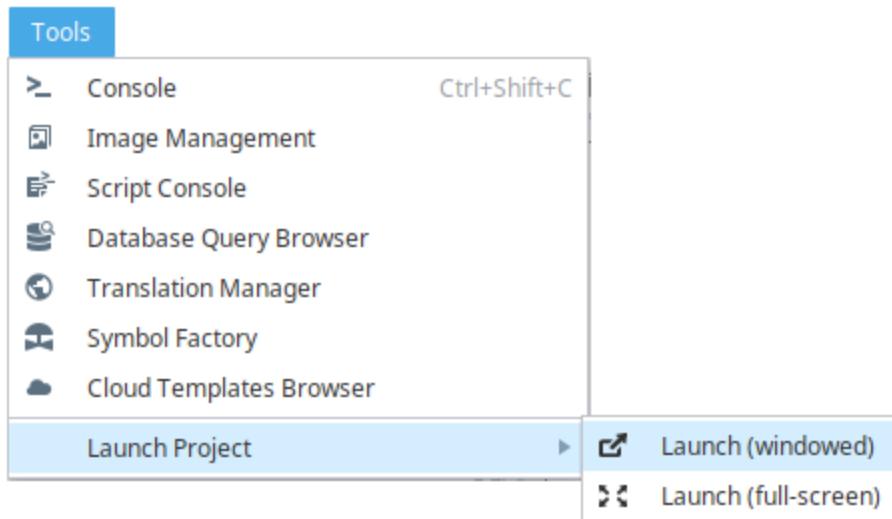
- From the Main Toolbar using the **Project > Preview Mode** menu item.
- Using the **F5** key to toggle between the two modes.

In Design Mode, your mouse is used to manipulate components in a window. You can select, drag, and resize them. You can alter data bindings and event script configuration. Data bindings are active in Design mode, but event handlers are not.

In Preview Mode, you are interacting with a "live" version of the window. Property bindings and event handlers will run, just like in the Client.

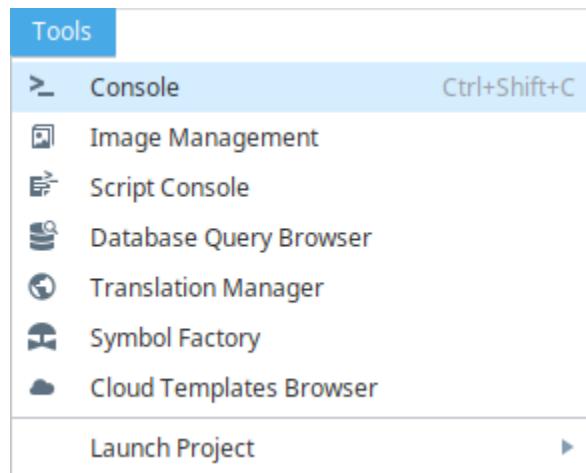
Note: Some of Ignition's functionality will not work in Preview Mode, for example, the retarget and openWindowInstance scripting functions must be tested in a Client.

Preview Mode is useful for a quick check of the operation of a window, but it becomes cumbersome when trying to test a whole project. For that, we recommend having a launched Client up as well, and doing testing in the true Client. You can quickly launch a client in one of the following two launch modes via the **Tools > Launch Project** menu.



Tools Menu

The **Tools Menu** provides some tools to help you when creating projects.



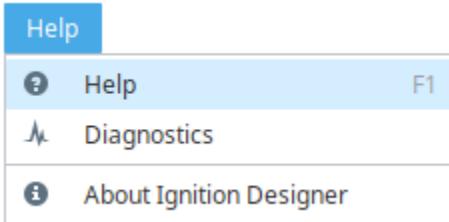
Each of the tools are described here.

Function	Description
Console	The Output Console is a dockable panel prints system messages that are coming from the designer. This can vary from simple info messages that show when things are loaded, to error messages when something goes wrong. The Console is also used frequently to test and debug Python scripts, as print statements on components that are run in the designer are printed here.
Image Management	The Image Manager is available from the Tools > Image Management menu. This tool is a drag-and-drop browser that helps manage the images that are stored on the Gateway. It is important to realize that these images are shared across all projects: they are not stored inside a project itself. Use the Image Management tool to do common tasks like uploading new images and creating folders. You can drag images and folders from your computer's desktop or hard drive into this window to easily upload new images to the Gateway. You can also get to this tool by putting an Component Palette > Display > Image component on a window, and clicking the browse button on the image's Image Path property. See Images and SVGs in Vision and Images and Icons in Perspective .
Script Console	Opens up the Script Console . window. This is where you can test scripts out.
Database Query Browser	Opens up the Database Query Browser panel, which allows you to run SQL queries against your database connections.
Translation Manager	Opens up the Translation Manager panel, which allows you to configure translations. See Localization and Languages .

Symbol Factory	If you have the Symbol Factory module installed, you'll be able to open the Symbol Factory browser under the Tools menu in the Designer. You can browse through the symbols or use the convenient search function to find the symbol you need. Once you find a symbol, you can drag-and-drop it into a window. Each symbol is dropped as a shape group. You will be able to un-group it or double-click into the group just as if you had drawn the symbol yourself using fundamental shapes. This means that you can alter the shape if you need to, or bind any colors inside the shape to a Tag to make the shape dynamic.
Launch Project	Allows you to launch the project directly from the Designer.

Help Menu

The **Help Menu** provides online assistance when looking for information or troubleshooting an issue.



Functions	Description
Help	Opens up your web browser and takes you to this User Manual for quick reference.
Diagnostics	The Help menu in the Designer and the Vision Client has a Diagnostics window that contains a number of tabs each providing a useful troubleshooting feature. You can right-click on any of the tabs to show or hide the other tabs. For more information on these tabs and the troubleshooting features, go to Designer Diagnostics .
About Ignition Designer	Provides information about the Designer such as the versions of the modules, Java version, and Gateway IP Address that the Designer is using.

Related Topics ...

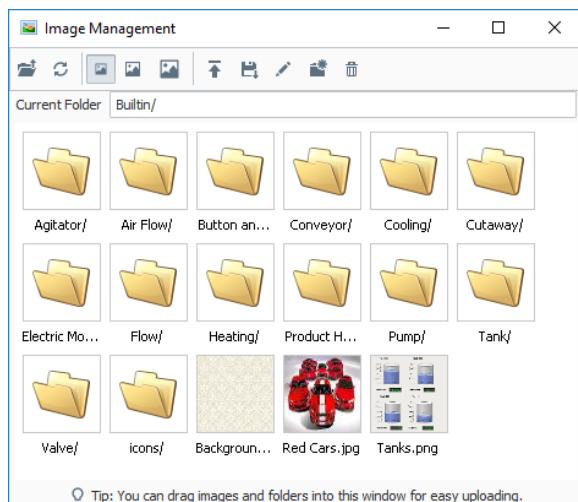
- [Vision Designer Interface](#)
- [Perspective Designer Interface](#)

Designer Tools

The Designer comes with many tools that allow you to manage and test various resources within a project. Each of the tools have their own interface and are accessed within the Tools menu on the menu bar of the Designer.

Image Management Tool

The [Image Management Tool](#) allows you to manage the images that are stored within the Ignition Gateway. The path to the images can be copied out and pasted into a component's property that is expecting an image path.

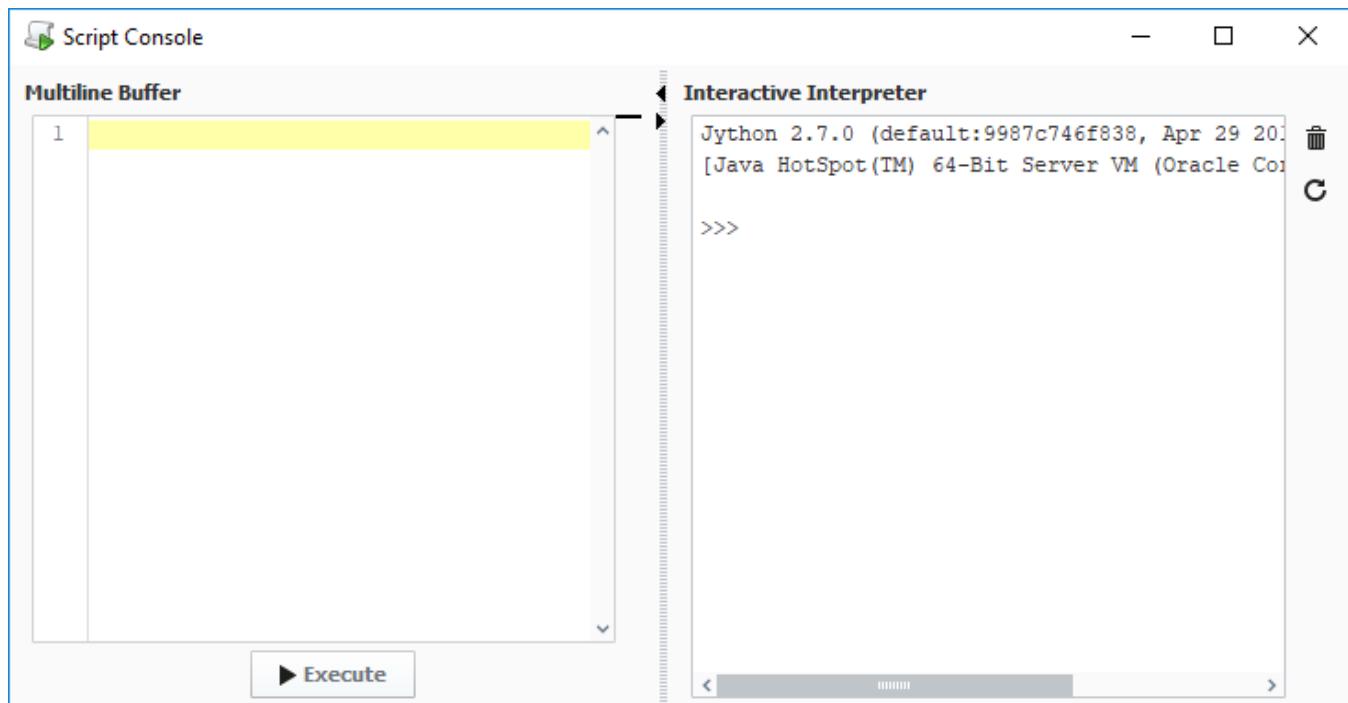


On this page ...

- [Image Management Tool](#)
- [Script Console](#)
- [Database Query Browser](#)
- [Translation Manager](#)
- [Symbol Factory](#)
- [Output Console](#)
- [Keyboard Layouts](#)

Script Console

The [Script Console](#) allows you to test out code snippets, printing the results out in the panel on the right.



Database Query Browser

The [Database Query Browser](#) is like the Script Console, but for SQL and databases. Here, you can test out queries to ensure you are returning the correct data from your database.

The screenshot shows the Database Query Browser interface. At the top, there is a toolbar with a search bar containing the query "SELECT * FROM audit_events". To the right of the search bar is an "Execute" button with a yellow lightning bolt icon. Below the search bar, there is a checkbox labeled "Limit SELECT to: 1000 rows". The main area displays two result sets, "Resultset 1" and "Resultset 2". "Resultset 1" contains 23 rows of data from the "audit_events" table, with columns: AUDIT_EVENTS_ID, EVENT_TIMESTAMP, ACTOR, ACTOR_HOST, and ACTION. The data includes various events like logout, login, project save, etc., occurring between October 22 and 23, 2018. "Resultset 2" is currently empty. At the bottom of the results, it says "23 rows fetched in 0.004s". Below the results are buttons for Auto Refresh, Edit, Apply, and Discard. To the right of the results, there is a sidebar titled "Default" showing the database schema for "audit_events", listing columns such as ACTION, ACTION_TARGET, ACTION_VALUE, ACTOR, ACTOR_HOST, AUDIT_EVENTS_ID, EVENT_TIMESTAMP, ORIGINATING_CONTEXT, ORIGINATING_SYSTEM, and STATUS_CODE.

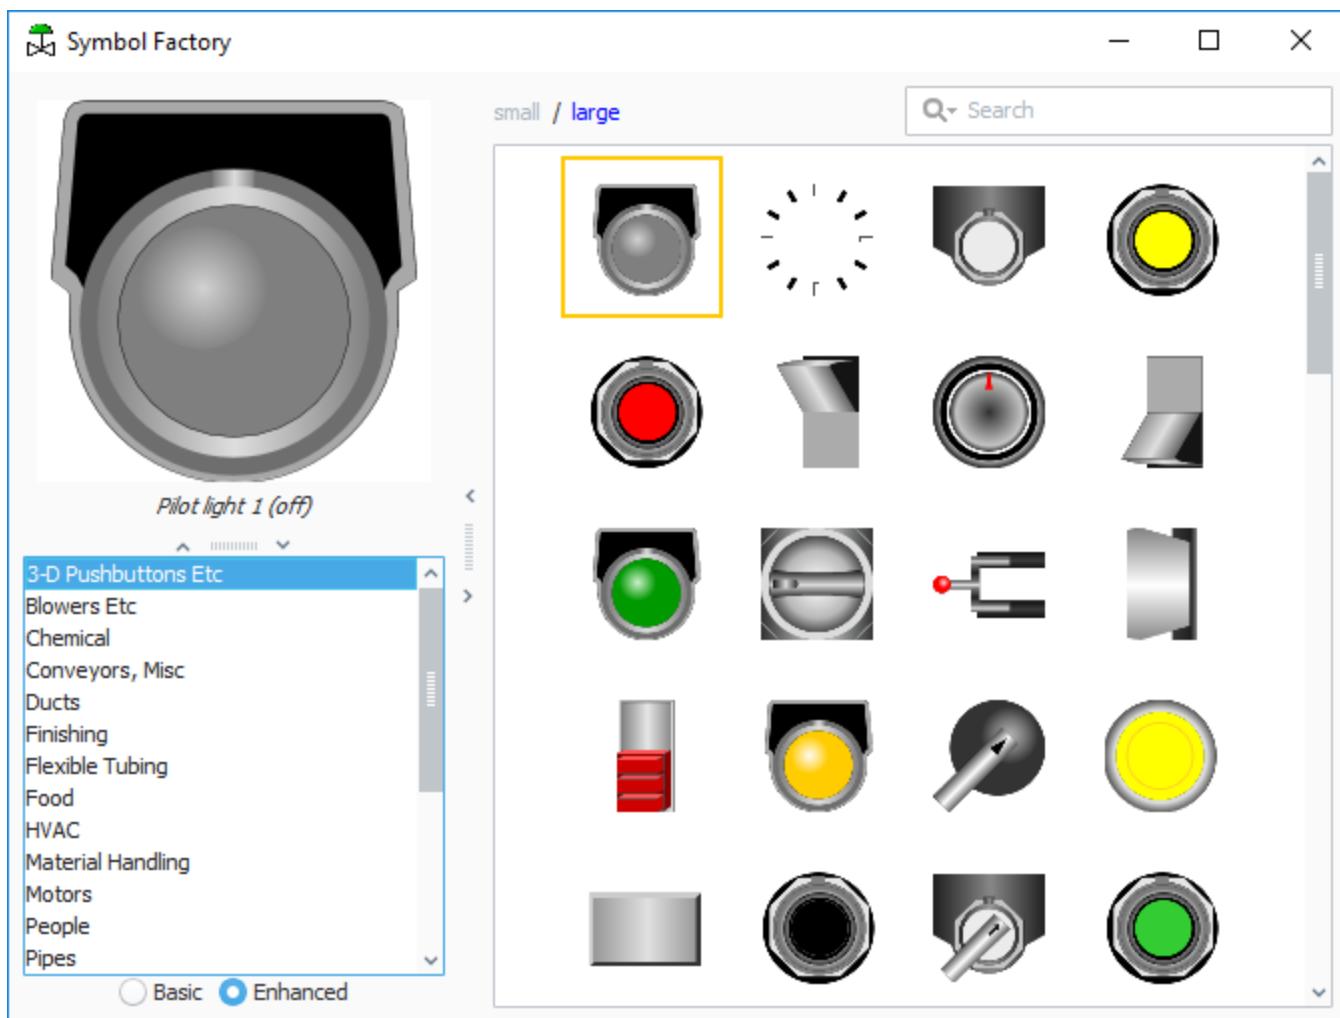
Translation Manager

The [Translation Manager](#) allows you to add, edit, and remove translation mappings to your system. Works in conjunction with the [Localization](#) system.

The screenshot shows the Translation Manager interface. On the left, there is a sidebar titled "Languages" with checkboxes for "(All)", "English", and "Spanish", all of which are checked. In the center, there is a table titled "Translation Terms" with a search icon. The table has three columns: "Key", "English (Alternate)", and "Spanish". It lists several terms with their translations: button (buton), component (componente), industry (industria), level (nivel), tank (tanque), temperature (temperatura), and user (usuario). On the far right, there is a vertical toolbar with icons for adding (+), deleting (x), saving (disk), and canceling (cancel).

Symbol Factory

The [Symbol Factory](#) module is a unique designer tool included with Vision or Perspective.



Output Console

The [Output Console](#) prints system messages coming from the Designer from simple info messages to error messages. The Console is also frequently used to test and debug Python scripts as print statements on components that are run in the Designer.

A screenshot of the Ignition software's Output Console window. The window title is "Output Console". The content area displays several log messages in a monospaced font. The messages are as follows:

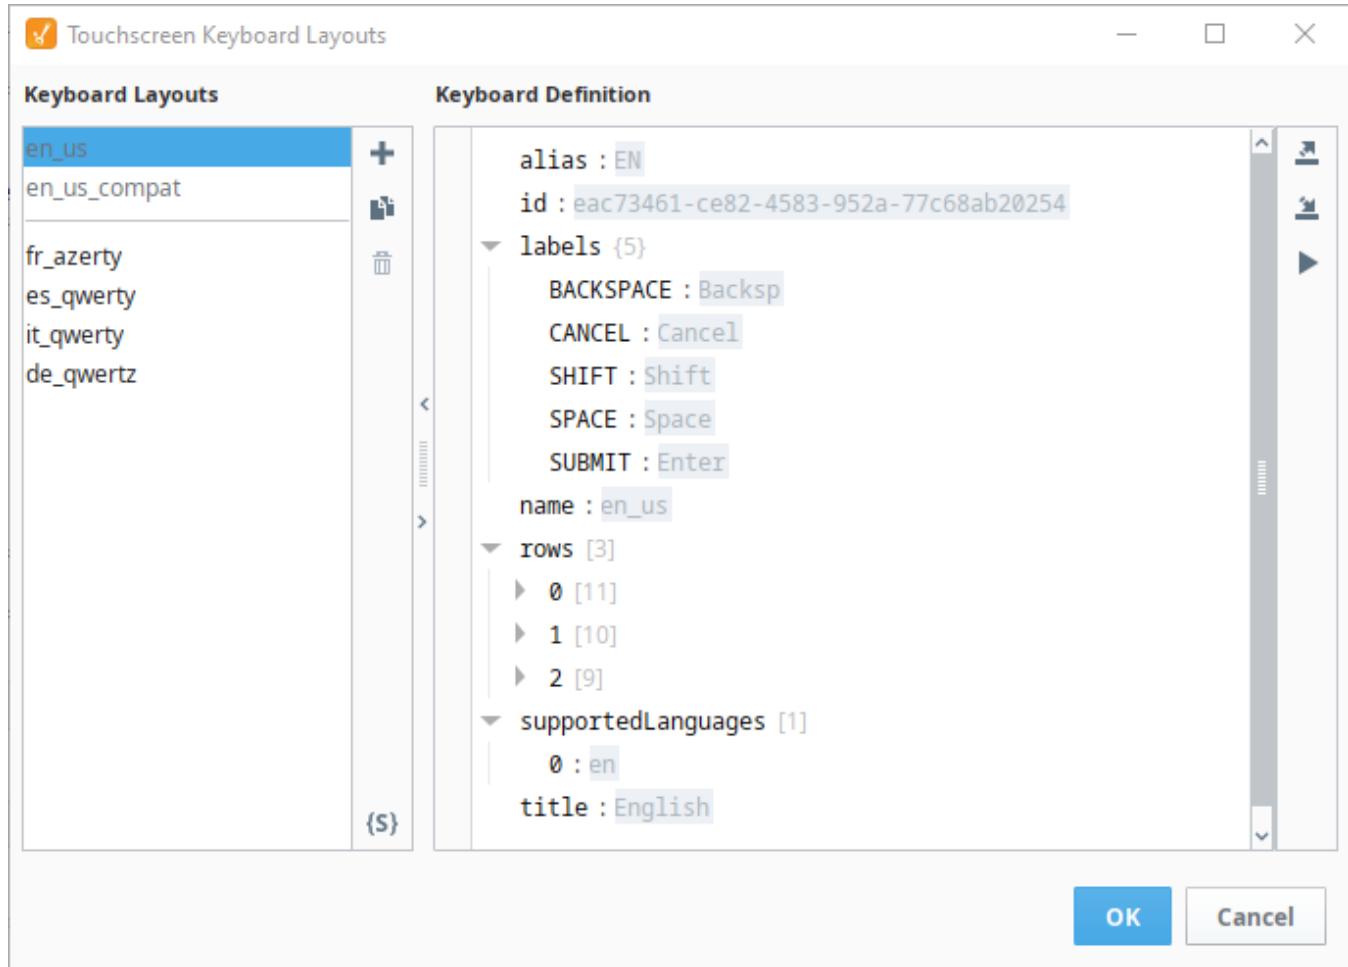
```
11:10:14.337 [Designer-Startup] INFO designer.main - Downloading... [+0]
11:10:47.360 [Designer-Startup] INFO designer.main - Cleanup... [+33023]
11:10:47.368 [Designer-Startup] INFO designer.main - Initializing UI... [+33031]
11:10:47.433 [Designer-Startup] INFO tags.manager.gwinterface - Tag poll rate changed to 250 ms
11:10:47.434 [DesignerExecEngine-1] INFO tags.subscriptions - Changing connected quality to 'Good'
11:10:49.362 [Designer-Startup] INFO designer.main - Starting module: Perspective [+35025]
```

Keyboard Layouts

The following feature is new in Ignition version **8.1.28**
[Click here](#) to check out the other new features

The [Keyboard Layouts](#) editor adds the ability to define custom keyboard layouts, which are then applied on Vision's touchscreen keyboard. Although with this tool keyboard layouts are completely customizable, you can quickly switch between a few preloaded language options to display French,

Spanish, Italian, German, or two variations of an English keyboard. This tool is not available in Perspective.



[Related Topics ...](#)

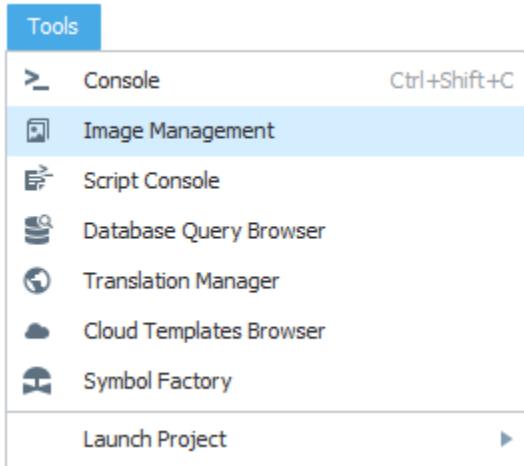
- [Images and SVGs in Vision](#)
- [Localization and Languages](#)

[In This Section ...](#)

Image Management Tool

Images such as PNGs, JPGs, GIFs, and SVGs can be uploaded to the Image Management Tool and used inside of windows in Ignition. Once uploaded, these images may be used on windows and in templates. The Image Manager tool, available from the Tools > Image Management, provides an interface to upload, download, or select images.

Note: The Image Management tool does not support bitmap files.



On this page ...

- [Uploading an Image to the Image Management Tool](#)
- [Downloading Images from the Image Management Tool](#)
- [Exporting and Importing Images in Projects](#)

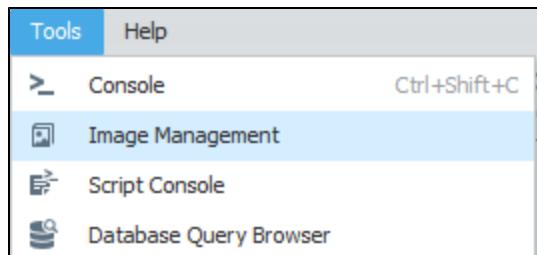


Images (png, jpg, gif)

[Watch the Video](#)

Uploading an Image to the Image Management Tool

There are two ways to upload an image into the Image Management tool. Both ways involve having the Image Management tool open. At the top of the Designer in the Menu Bar, select **Tools > Image Management**.

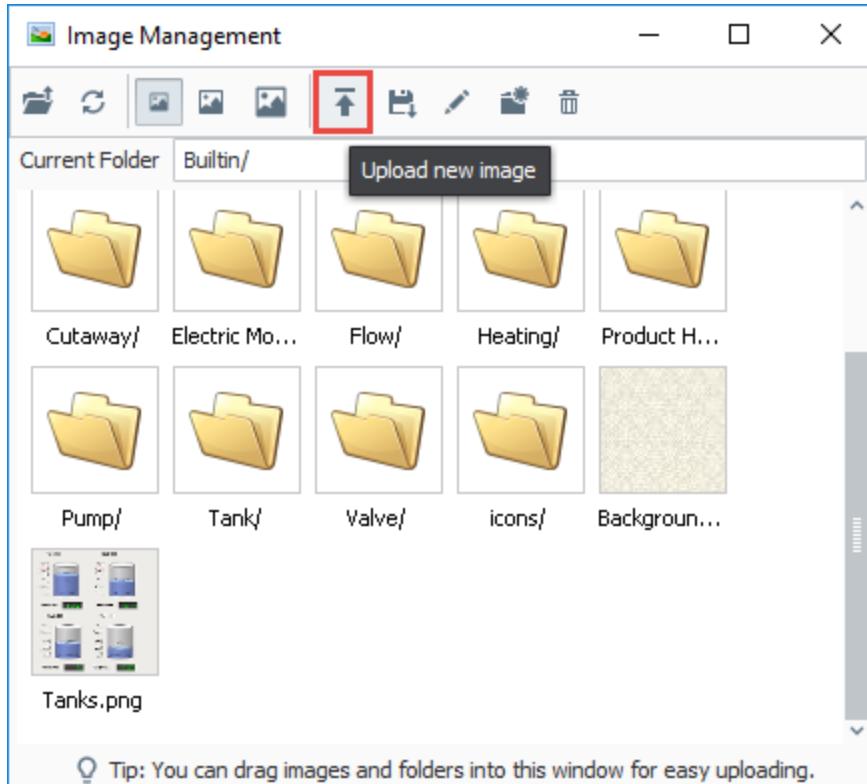


Upload on Drag and Drop

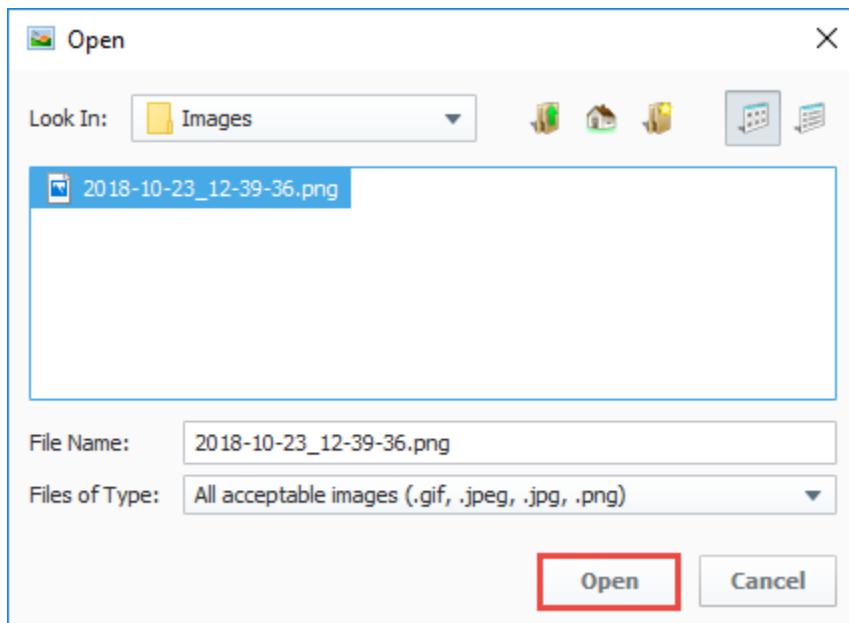
Images can simply be dragged and dropped from the local file system into the **Image Management** window.

Manual Upload

Alternatively, the Image Management window has an Upload button to pass images in. Locate the directory you wish to upload the image. You may use the root folder, or create a new folder to keep your images organized. When ready, click on the **Upload Image** button.



An **Open** dialog window will appear. Simply find your image on the local system, and click **Open** to upload the selected image.



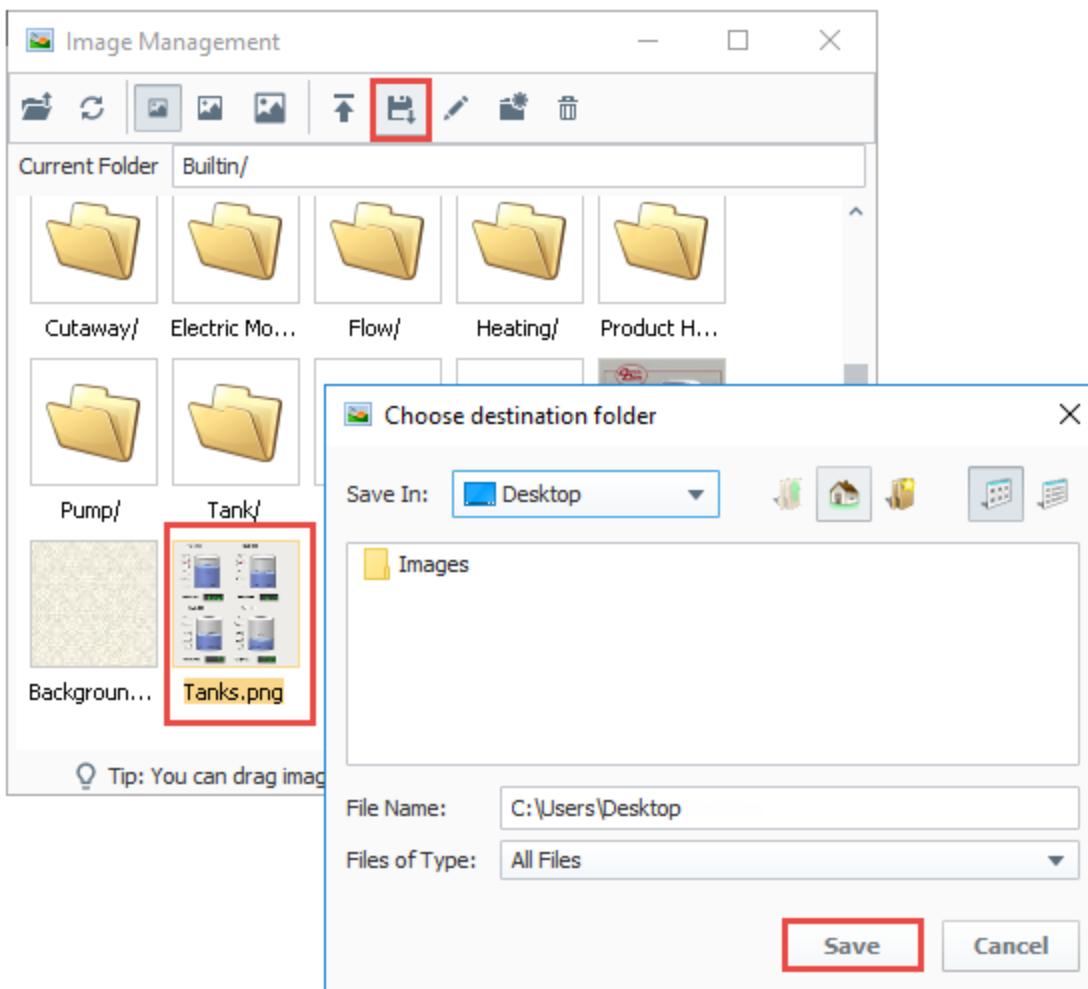
Downloading Images from the Image Management Tool

Single images, as well as entire directories, may be downloaded from the Image Management tool. This is useful when migrating a project to another Gateway.

Image downloads can be taken from either the **Image Management** or **Image Browser** windows. If at least one folder or image is selected, the **Save**



icon will become enabled. Click the **Save** icon, pick a local directory to save the images to, and click the **Save** icon again. All images and subfolders in the selected folder will be copied to the selected directory.



Exporting and Importing Images in Projects

Uploading an image involves storing the file in the Gateway's internal database. This means that project exports do not contain any referenced images.

When exporting a project for use in another Gateway, it is recommended to also export any images that the project uses, and upload them into the new Gateway at the same directory. Examples of uploading and downloading images can be found on this page.

Related Topics ...

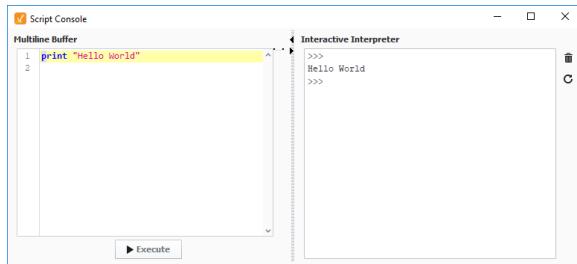
- [Images and Icons in Perspective](#)
- [Images and SVGs in Vision](#)
- [Using Images](#)
- [Symbol Factory](#)

Script Console

The Script Console is a live Python terminal that is only accessible in the Designer. It is a great way to very quickly test a script as it does not rest on a scripting event or specific component. The Script Console can be opened via the **Tools > Script Console** menu. It consists of two parts: a **Multiline Buffer**, and an **Interactive Interpreter**. Code can be typed into both sides.

Due to how "scope" works, the Script Console can not interact with components on a window, but it can call Project and Shared scripts. If a Project or Shared script was recently added, then the console will need to be reset before it can be called. This can be accomplished by clicking on the **Reset C** icon in the upper right.

Note: Gateway-scoped information will not appear in either the Script Console or [Output Console](#). Instead, the output will be sent to the wrapper.log file. Alternatively, `system.util.getLogger` will send messages to the Gateway Console, and is a preferred method of troubleshooting Gateway scoped scripts.



On this page ...

- Features
 - Multiline Buffer
 - Interactive Interpreter

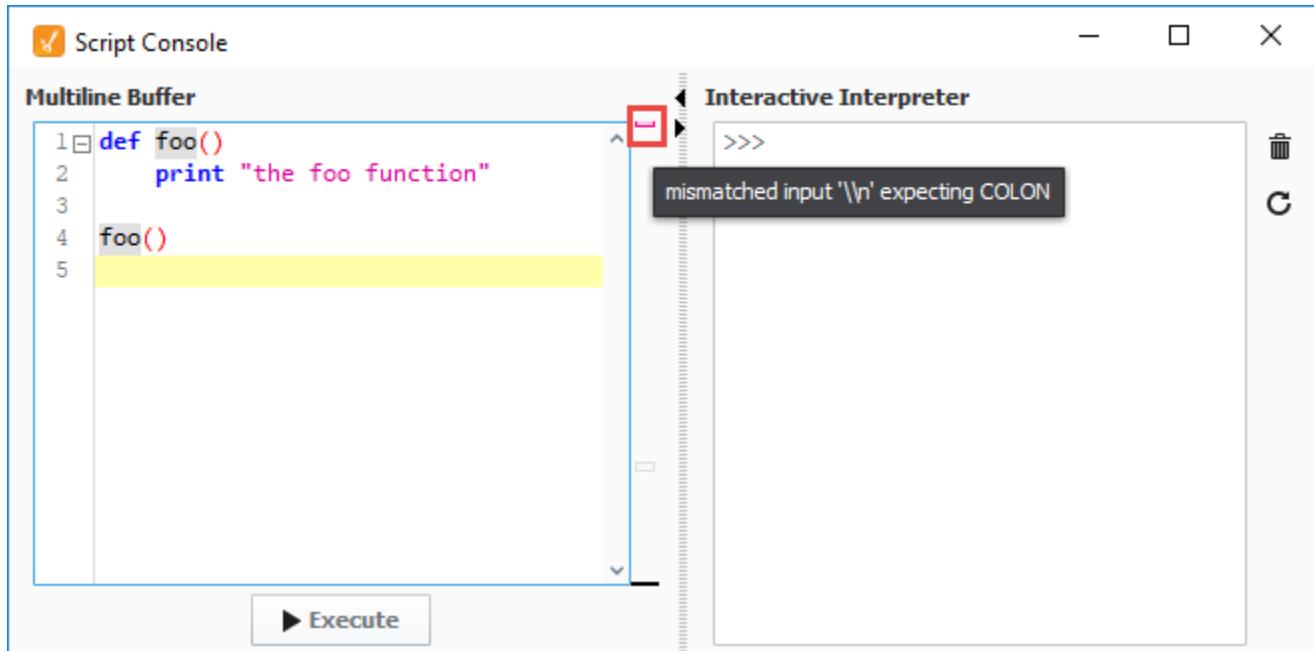


Scripting Console

[Watch the Video](#)

Features

There are several icons and user interface elements located on the Script Console window. The image below highlights a syntax error.



A reference of the icons and descriptions are found in the table below:

Icon	Name	Description
	Clear	Clears the text from the Interactive Interpreter.
	Reset	Clears and resets the text, and deletes all user defined objects (variables and functions) from the Interactive

		Interpreter.
	Expand/Collapse	Expands / Collapses Multiline Buffer and Interactive Interpreter.
	Syntax Error Highlight	Shows up next to a line in the Multiline Buffer identifying an error. Hover over the Error Symbol to see information on the exception.

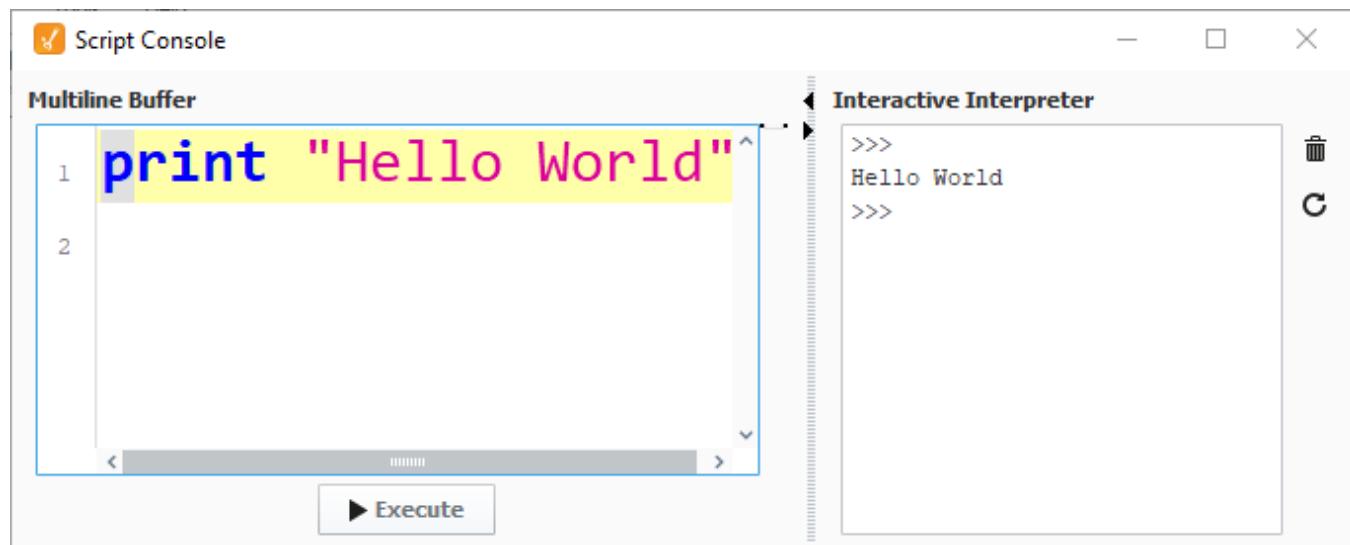
Multiline Buffer

The Multiline Buffer, located on the left side of the Script Console, allows for multiple lines of code to be entered and then executed by clicking on the button. All `print` statements will output to the Interactive Interpreter. It also supports code folding for function definitions and comments.

When executing a script in the console, the button will change to an button. Developers can press the Interrupt button to interrupt / stop a script from executing when testing code with a lot of data, or when the script inadvertently gets stuck in an infinite loop.

Font Size Adjustment

Font size in the Multiline Buffer can be adjusted by holding **Ctrl** and scrolling the mouse wheel.



Right-Click Menu

Right clicking on the Multiline Buffer opens a menu. The menu options are described in the table below.

Script Console

Multiline Buffer

```
1 def foo()
2     print "Hello World"
3
4
```

Execute

Main Menu Options	
Name	Description
Undo	Undoes the last action.
Redo	Gets rid of the last undo action.
Cut	Cuts the selected text.
Copy	Copies the selected text.
Paste	Pastes the selected text.
Delete	Deletes the selected text.
Select All	Selects all text in the window.

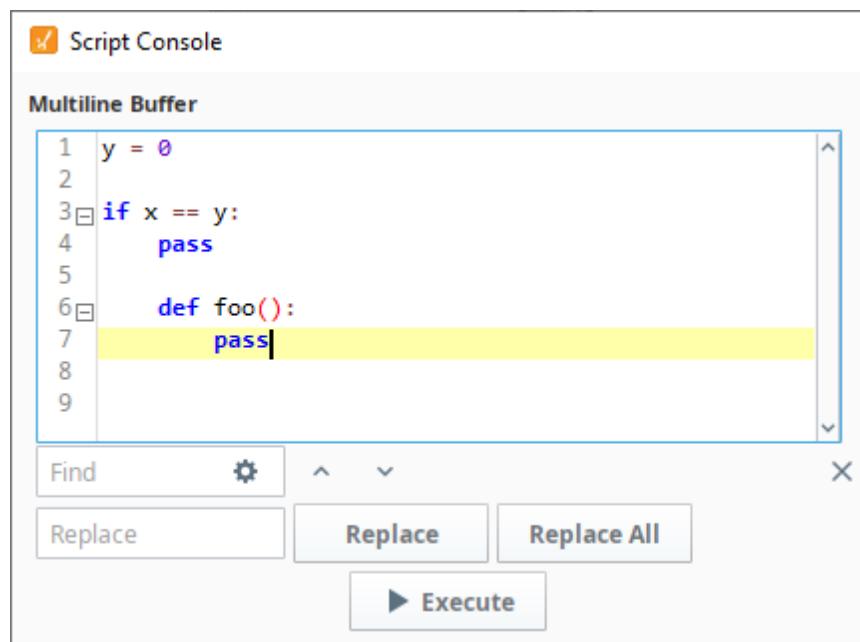
Folding	
Name	Description
Toggle Current Fold	Expand or collapse the fold where the text cursor is located.
Collapse All Comments	Collapse all instances of multi-line comments. Only contiguous comments are collapsible.
Collapse All Folds	Collapse all expanded folds.
Expand All Folds	Expand all collapsed folds.

Autocompletion	
Name	Description

Automatic Activation	Determines access to the Autocompletion window. If set, the window will automatically appear after a second of inactivity occurs when "system." has been typed. If not set, the window can still be accessed manually by pushing Ctrl+Space .
Description Window	Shows or hides the Description pane in the Autocompletion window.
Parameter Assistance	Enables parameter assistance for known function arguments.
Appearance	
Name	Description
Whitespace	Paints an arrow character in whitespace sections (tabs and spaces) for each line.
Tab Lines	Paints lines at the indentation level of each tab stop.

Find/Replace

Pressing **Ctrl+R** while the text cursor is in the Multiline Buffer opens a Find and Replace window. This will search for instances of text throughout the Multiline Buffer, and allows the user to replace all or some instances with new text.



Keyboard Shortcuts

The following shortcuts apply only to the Multiline Buffer.

Key(s)	Description
Ctrl + </> (on the number pad)	Collapse all folds.
Ctrl + <*> (on the number pad)	Expand all folds.
Ctrl + <-> (on the number pad)	Collapse the fold on the same line as the text cursor.
Ctrl + <+-> (on the number pad)	Expand the fold on the same line as the text cursor.
Ctrl + <Space>	Open Autocompletion window. By default, the window will automatically open once "system." has been typed.
Ctrl + <R>	Open Find/Replace window.
Ctrl + <Mouse Wheel Scroll>	Increase and decrease the font size.
Ctrl + <Enter>	Executes the script in the script editor.

Interactive Interpreter

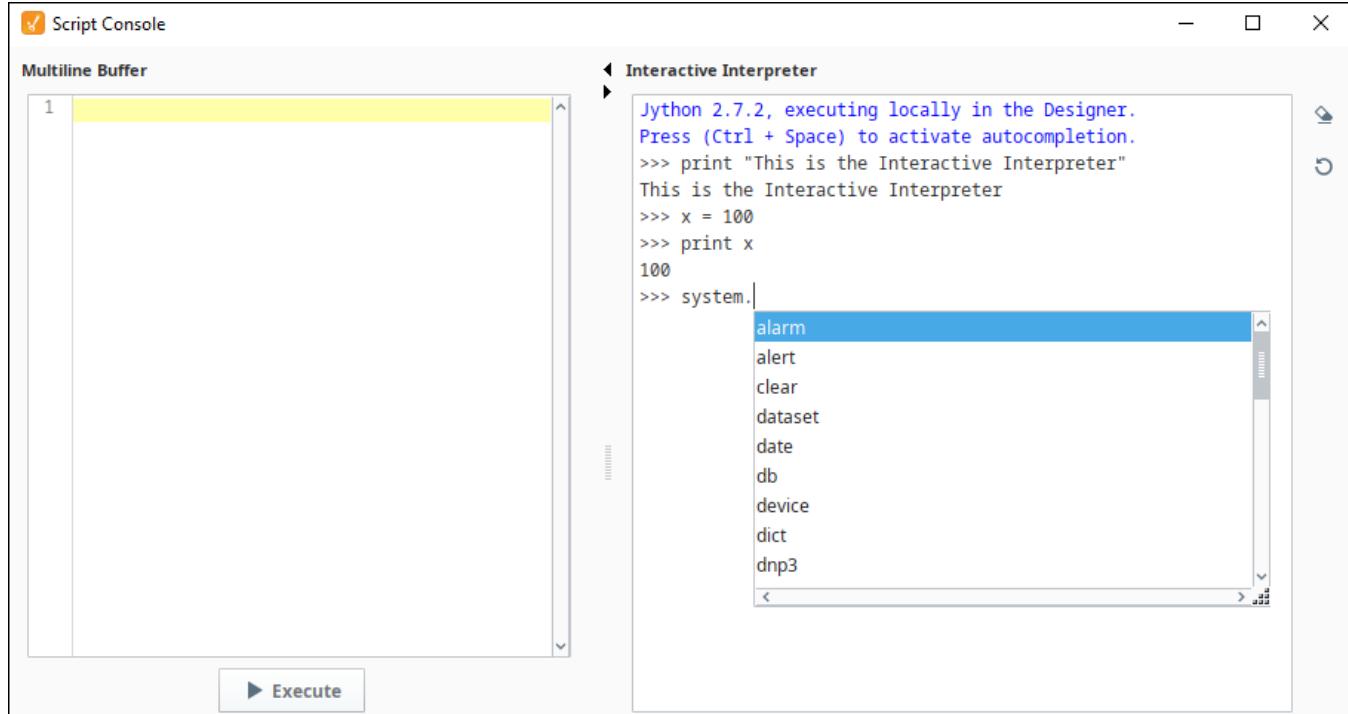
The Interactive Interpreter is located on the right side of the Script Console, and allows you to run a single line of code at a time. Code is executed from the Interactive Interpreter by pressing the **Enter** key. Print statements from both sides of the Script Console will appear in the Interactive Interpreter.

The Autocompletion window, available in the Interactive Interpreter, has access to the current working environment so items such as Project and Shared scripts will automatically appear. They can also be typed in manually.

The following feature is new in Ignition version **8.1.33**

[Click here](#) to check out the other new features

When first opened or refreshed, the Interactive Interpreter default header text has displayed the Jython version. Now, it also includes a reminder that the code executes in the local Designer scope (no access to Gateway methods), and instructions for how to trigger the autocomplete hint popup using the the Ctrl + Space keys.



Keyboard Shortcuts

The following shortcuts apply only to the Interactive Interpreter

Key(s)	Description
Ctrl + <L>	Clear the Interactive Interpreter. Functionally the same as clicking the Clear button.
Ctrl + <R>	Reset the Interactive Interpreter. Functionally the same as clicking the Reset button.
Up arrow	Cycle backward through command history.
Down arrow	Cycle forward through command history.
Ctrl + <C>	Keyboard interrupt.
Ctrl + <Space>	Open Autocompletion window. By default, the window will automatically open once an "object." has been typed such as "system" or "project," and a Project script has already been defined.
Ctrl + <A>	Move the text cursor to the start of the line. Similar to pressing the Home key.

Related Topics ...

- [Output Console](#)

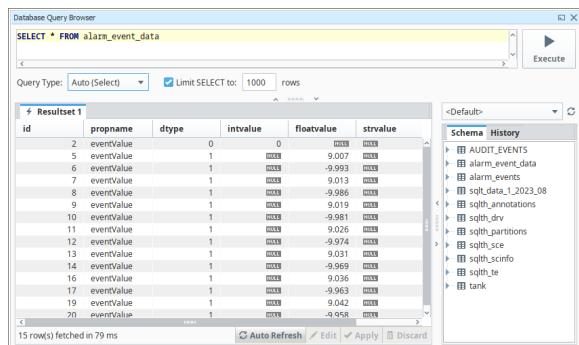
- Python Scripting
- Scripting in Ignition

Database Query Browser

The Database Query Browser Panel

The **Database Query Browser** is a very convenient tool that lets you query any database connected to Ignition, and interact with tables. Because Ignition is so heavily integrated with databases, it is very common in the course of project design the need to inspect the database directly, or to experiment with a SQL query to get it just right.

The Database Query Browser is found in the tools menu and has a few basic parts to it. The first is the query text field at the top of the window where a query can be written and then executed against the selected database connection. This can accept any type of query, but if the query would update the database, the Designer needs to have its Gateway **Communication Mode** set to read/write first. The database connection that the query is executed against can be chosen from the dropdown, below the Execute button. The Result Data space below the query field is where the results of the executed query will appear. If a SELECT query was run, then the table data will be shown. If an UPDATE, INSERT, or DELETE query was run, then the number of rows affected will be displayed instead. Finally, the Table List on the right-hand side of the window will display all of the tables in the specified database connection. The tables can be expanded to show the columns and their data types for that table to assist in writing queries. Additionally, when a table is double-clicked, the query will automatically be populated with a **SELECT * FROM** for that table.



On this page ...

- [The Database Query Browser Panel](#)
- [Features of the Query Browser](#)
 - [Query Type](#)
 - [Select Limit](#)
 - [Multiple Resultsets](#)
 - [Query History](#)
 - [Auto Refresh](#)
 - [Editing the Table in the GUI](#)



Using the Query Browser

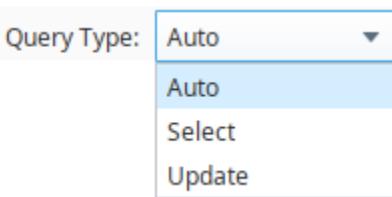
[Watch the Video](#)

Features of the Query Browser

The Database Query Browser has a few features that can help manage and build any SQL query.

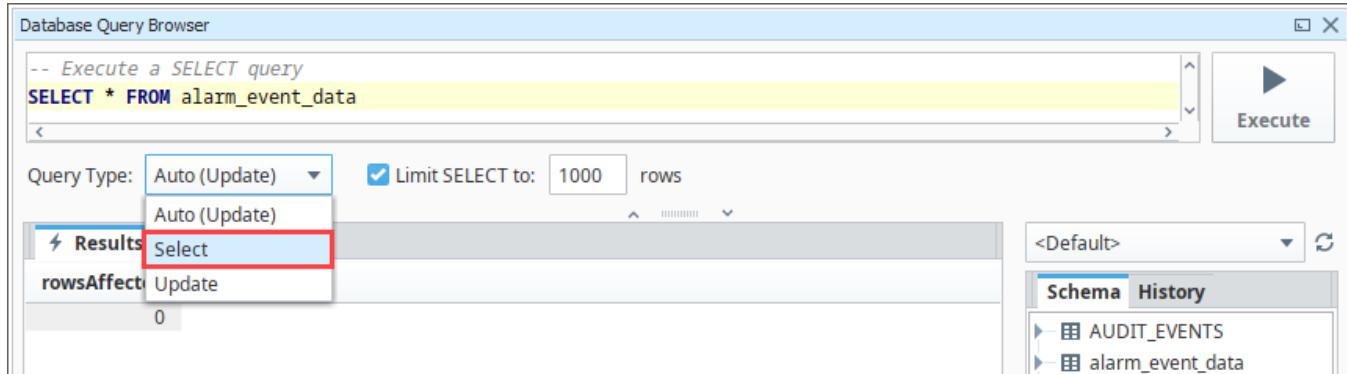
Query Type

The following feature is new in Ignition version **8.1.32**
[Click here](#) to check out the other new features



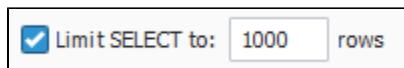
The Query Type field displays the mode that is used when a query is executed. By default, this field will show Auto, which will change to reflect the query type mode as a query is entered to show either Auto (Select) or Auto (Update). This allows confirmation that the correct mode is used when executing the query. If the auto-detected mode is incorrect, the drop down includes Select and Update to set the mode type.

For example, since the check for the query string will assume an update is being run if it doesn't detect SELECT text at the beginning of a query, you may want to use the Query Type dropdown to choose Select as the mode option to confirm the query will execute correctly.



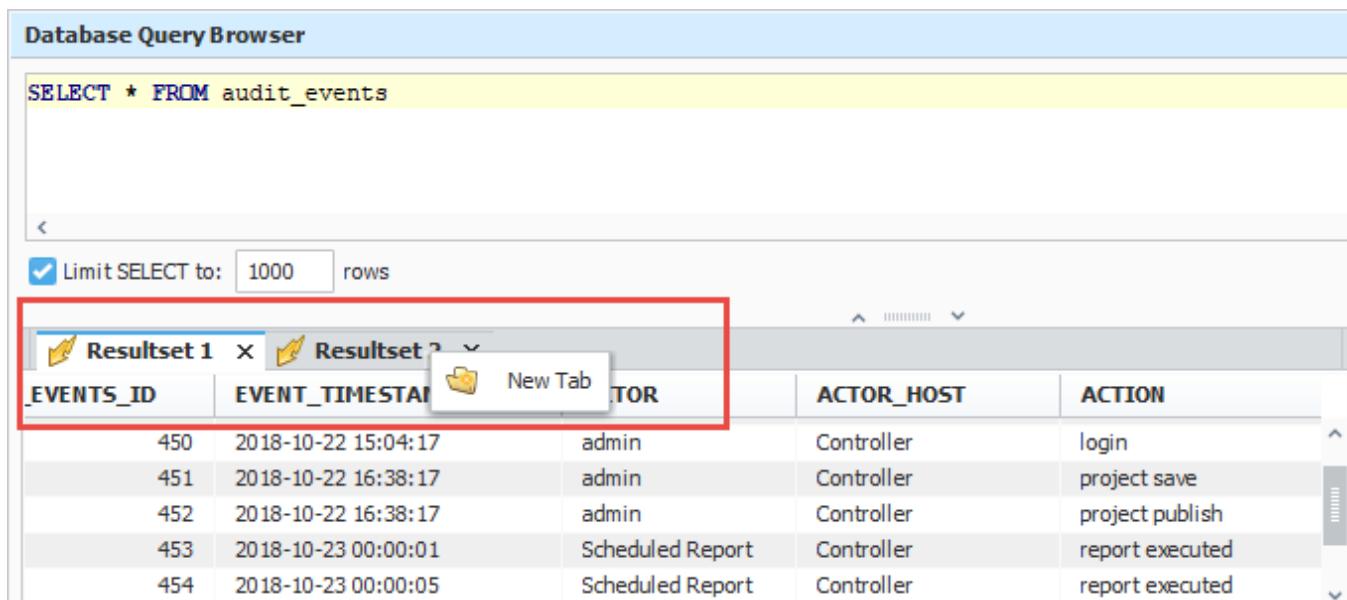
Select Limit

By default, any SELECT statement is limited to 1000 rows. This is to help the queries return quickly, however, it may not always be wanted. This can be turned off or adjusted by either clicking the checkbox or manipulating the row number located underneath the Query Area on the left. It is recommended to leave this on unless you know the result set size. It's better to use a count query than to return all results to see the result count. IE: `SELECT COUNT(*) FROM table...`



Multiple Resultsets

The Query Browser allows you to make multiple tabs of results, so that multiple queries can be run and their results compared to each other. Simply right click on the **Resultset** area and select **New Tab**. Your queries will be saved for each tab.



Query History

At the top of the Table List, there is a second tab labeled **History**. This will switch the Table List to instead show the history of queries run in the query browser. Double clicking one of the entries will push the query into the Query Area (for the selected tab). This allows you see what queries have been run previously to help you tweak your current query.

Auto Refresh

The Query Browser also lets you monitor a database table for changes by using the **Auto Refresh** button. This is often convenient when designing Transaction Groups. As the group runs, you can view the table that it is targeting with Auto Refresh turned ON to watch how the group is altering the table. Simply click the Auto Refresh button at the bottom of the Query Browser to periodically rerun the query in the Query Area. Make sure to include an ORDER BY clause so your results show up in the order you want.



Editing the Table in the GUI

In addition to editing the table data using INSERT, UPDATE, and DELETE statements, the data can also be edited within the Result Data. Simply click the Edit button at the bottom of the Query Browser window.



Editing in this way requires that the data be obtained from a single table. The table must also contain a primary key.

Once the Edit button is clicked, the values in the table can be edited by double clicking on the value and entering a new one.

id	machine_name	area_number
1	Conveyor	1
2	Press	2
3	Tank	1
4	Packer	3
5	Loader	3
6	Oven 1	3
7	Oven 2	2
9	Mixer	3
10	Cold Storage	2
12	Chopper	1

10 rows fetched in 0.006s

Right clicking on a row also provides a few options:

- **Add Row:** Will add a new row to the table for data to be entered into.

- **Clear Field:** Will clear out the value in the selected cell so that it will be NULL. This is different than simply deleting the value out and leaving it empty.
- **Delete Row(s):** Will delete the selected row or rows from the table.
- **Copy Row Values:** Will copy the row values in a comma separated form to the clipboard.
- **New Tab:** Will Create a new Resultset tab for a new query to be run in.

When editing values, cells will highlight depending on what is being done to them. Green cells are new, and typically indicate a new row was added. Red cells are marked for deletion, and will be deleted when the changes are confirmed. Blue cells are cells that have had values changed during editing.

The screenshot shows the Database Query Browser interface. A table named 'machines' is displayed with columns: id, machine_name, and area_number. A context menu is open over the row with id 9, labeled 'New Mixer'. The menu options are: Add Row, Clear Field (highlighted in yellow), Delete Row(s), Copy Row Values, and New Tab. The browser also shows a schema tree on the right and a status bar at the bottom.

id	machine_name	area_number
1	Conveyor	1
2	Press	2
3	Tank	1
4	Packer	3
5	Loader	3
6	Oven 1	3
7	Oven 2	2
9	New Mixer	3
10	Cold Storage	2
12	Chopper	1
0		0

After making edits to the table data, the changes then either need to be applied or discarded. This gives you the opportunity to revert the table to the way it was before the current edit session, or apply the changes and rewrite the table appropriately. Simply click the corresponding button at the bottom of the Query Browser next to the Edit button. Make sure your Gateway **Communication Mode** is set to Read/Write before Applying your changes.



The following feature is new in Ignition version 8.1.25
[Click here](#) to check out the other new features

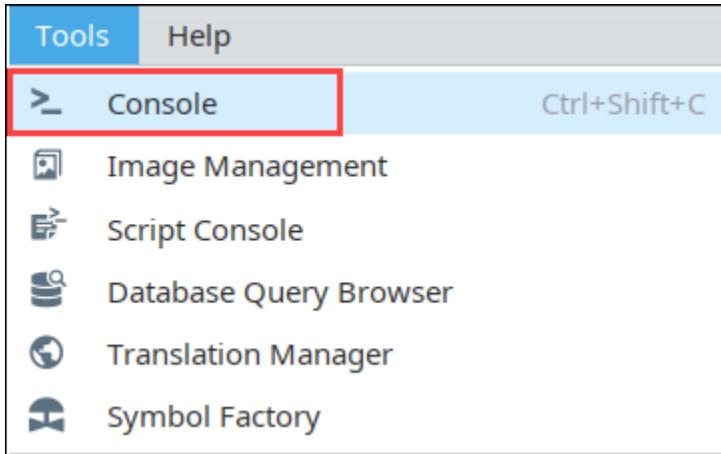
Note: Copied values from the Database Query Browser will paste in an RFC 2339 compatible local datetime format. Additionally, using the shift key while copying will copy the values in a tab-separated format ideal for pasting into spreadsheet tools such as Microsoft Excel.

[Related Topics ...](#)

- [Keyboard Shortcuts](#)

Output Console

The Output Console is a dockable panel that you can open via Tools > Script Console menu or the Ctrl-Shift-C keyboard shortcut. The Output Console prints system messages that are coming from the Designer. This can vary from simple info messages that show when things are loaded, to error messages when something goes wrong. The Console is also frequently used to test and debug [Python scripts](#), as print statements on [components](#) that are run in the Designer are printed here.



For example, by using the `print` function in your script, you can observe the inner workings of your script as it executes. If you executed the following script:

```
# A Python script that demonstrates the print statement.

print "Print me first"
x=10
z=2
print x, y, x/y

#Result looks like this:
#print me first
#10 2 5
```

Note: Gateway-scoped information will not appear in either the Output Console or [Script Console](#). Instead, the output will be sent to the [wrapper.log file](#). Alternatively, [system.util.getLogger\(\)](#) will send messages to the Gateway Console, and is a preferred method of troubleshooting Gateway scoped scripts.

The following feature is new in Ignition version **8.1.24**
[Click here](#) to check out the other new features

Note: You can use certain options, such as copying, in the Edit Menubar at the top of the Designer window.

The Output Console is most frequently used to test and debug Python scripts on components in Ignition. By using the `print` keyword in your script, you can observe the inner workings of your script as it executes. For example, if you executed the following script:

[Python - Using the Output Console to Test and Debug Python Scripts](#)

```
# A function that intercepts tag writes, printing out the previous value first.
def writeToTag(path, value):
    prevValue = system.tag.getTagValue(path)
    print "Writing value '%s' to %s, was previously '%s'" % (value, path, prevValue)
    system.tag.writeToTag(path, value)

writeToTag("Compressor/HOA", 2)
```

It would print the following to the console:

```
Writing value '2' to Compressor/HOA, was previously '0'
```

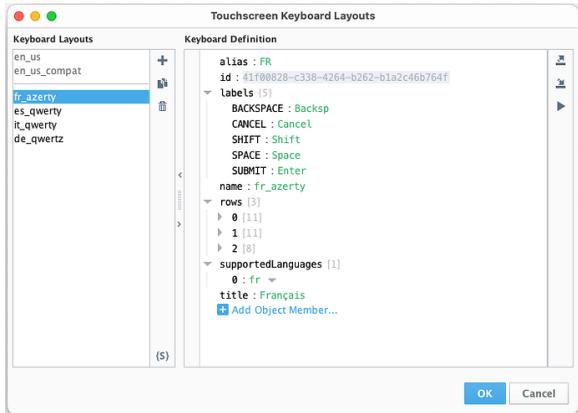
Note: The Output Console is also available in the Vision Client from **Help > Diagnostics** and selecting the **Console** tab.

Related Topics ...

- [Python Scripting](#)
- [Scripting in Ignition](#)

Keyboard Layouts

Keyboard layouts can be added, removed, and customized with the Keyboard Layout editor, which is available in the Tools menu of the Designer Vision workspace. The Keyboard Layouts editor allows you to create, modify, and import/export custom keyboard layouts across your entire Gateway. Keyboards are all configured with a common JSON structure and stored directly in the IDB so that they are shared across all projects.



Each layout can be associated with any number of locales by a language tag. When an end user invokes the touchscreen keyboard, Ignition will attempt to find a keyboard with support for their specific locale, falling back to the English layout if no more appropriate keyboard is found. However, there are a few more ways to assign what keyboard layouts appear. Keyboards can be set individually on touchscreen enabled components using the Touchscreen Keyboard Layout property. Keyboard layouts can also be switched in real time by operators if the keyboard layout includes a language selection key. Furthermore, keyboard layouts can be explicitly requested with the [system.gui.showTouchscreenKeyboard](#) function.

Keyboard Layout Editor

Access the Keyboard Layout Editor in the Designer Vision workspace by selecting **Tools > Keyboard Layouts**. There will already be six keyboard layouts prepopulated under Keyboard Layouts when you first open the editor. The default keyboard layouts are designed to make the best use of screen space and be mobile-friendly. Any created custom keyboard layouts will also be listed here during configuration and after saving. The following keyboard layouts are included by default:

- English (Modern) - A legacy layout utilizing a traditional looking English keyboard.
- English (Compatibility) - A revamped layout of the legacy keyboard with a design geared more toward mobile use.
- Spanish
- French
- German
- Italian

The international keyboards are included to assist in quicker implementation as needed. If your system does not have these keyboards, they can be deleted

in the editor using the Remove icon. The Add icon is used to create any new keyboard layouts that may be required. There are no significant restrictions placed on the number or location of keys. Additionally, if you don't want to build a keyboard layout from scratch, you can duplicate an existing keyboard configuration as a starting point for a new custom keyboard. To do this, select the keyboard you want to duplicate, then

select the Duplicate icon. At the bottom of the Keyboard Layouts section is an Export Schema icon. Use this to export a JSON schema of the selected keyboard layout for use in external JSON editors to create new keyboard layouts.

On the right side of the editor is the Keyboard Definition section. Defining the properties listed within will establish how your keyboard layout will be configured. The Keyboard Definition area also includes Import and Export icons to import and export keyboard layouts. Imported keyboard layouts must be JSON files. The Keyboard Definition also includes a Preview icon to display an instant, live preview in the Designer of the currently selected keyboard layout.

On this page ...

- [Keyboard Layout Editor](#)
 - [Keyboard Definition Properties](#)
 - [Default Keyboard Layouts](#)
 - [Using the Long Press Feature](#)

Touchscreen Keyboard Layouts

Keyboard Layouts

- en_us
- en_us_compat
- fr_azerty**
- es_qwerty
- it_qwerty
- de_qwertz

Keyboard Definition

```

alias : FR
id : 41f00828-c338-4264-b262-b1a2c46b764f
labels {5}
  BACKSPACE : Backsp
  CANCEL : Cancel
  SHIFT : Shift
  SPACE : Space
  SUBMIT : Enter
  name : fr_azerty
rows [3]
  0 [11]
  1 [11]
  2 [8]
  supported_languages [1]

```

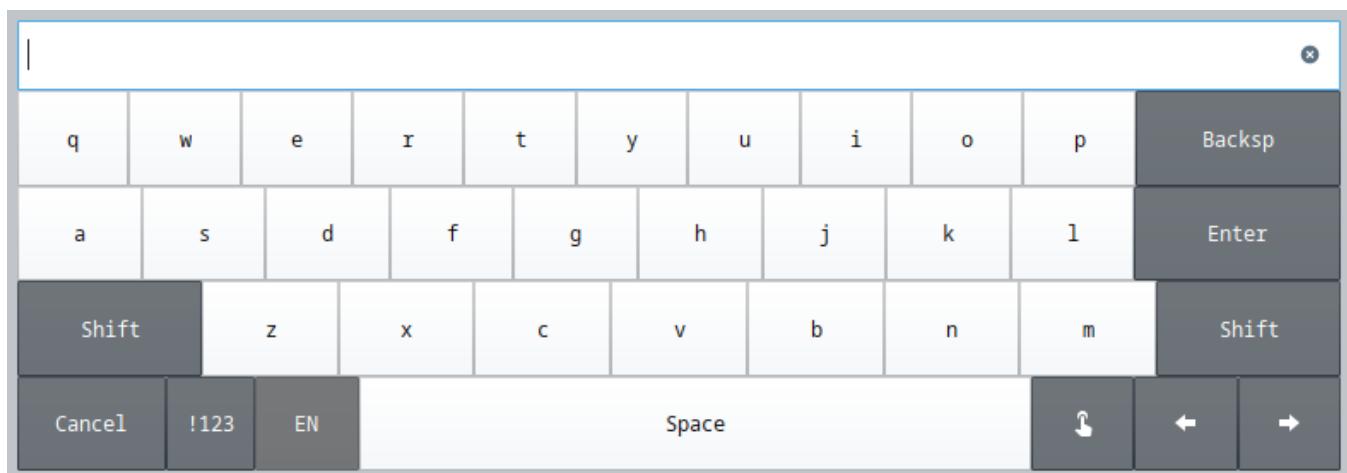
Keyboard Definition Properties

Property	Description
alias	The 2-3 character string that appears on the keyboard language selector key to identify this keyboard.
id	The fixed UUID for the keyboard layout. These are randomly generated and automatically updated to track layouts as they are updated on the Gateway.
labels	Localized strings to change the text on various system keys, such as backspace, shift, and cancel.
name	The name of the keyboard.
rows	<p>The designation of character keys on the keyboard sorted by row. Each character within a row can have uppercase and lowercase fields, along with the ability to add accent options for both fields. For example, row 0 character 2 on the EN keyboard is the letter "e". The uppercase and lowercase selections for "e" include an accents dropdown that contains either the uppercase or lowercase "ê", "é", "è", "ë", "" accent options. Row designation can be structured in any of the following three ways:</p> <ul style="list-style-type: none"> • A string literal member of the enum "BACKSPACE", "SUBMIT", "SHIFT", "CANCEL", "CAPSLOCK", "NUMERIC_LAYOUT" will create a system key at that location. • A simple key can be created with an object with at least one member lowercase: { "lowercase": "a" }. Add an uppercase key to define an alternate form to use when the shift or caps lock keys are active or more complex keys can be defined with a full object pattern: { "lowercase": { "character": "e", "accents": ["ê", "é", "è", "ë", ""] }, "uppercase": { "character": "E", "accents": ["Ê", "É", "È", "Ë", ""] } }. • An alias, allowing the typed character to be different than the label for the key: { "character": "\t", "alias": "Tab" }

supported languages	An array of IETF language tags representing the languages this keyboard language supports. This includes a title field to enter the string that is presented to the user when selecting this keyboard to use. For example, if your language tag was en, your title would be English.

Default Keyboard Layouts

The default keyboard layouts include standard Backspace, Enter, and Shift keys in addition to system keys that switch language selection, toggle layouts between alphabetic and numeric/symbol displays, highlight accent keys, and navigate forward or backward through the preview text field.

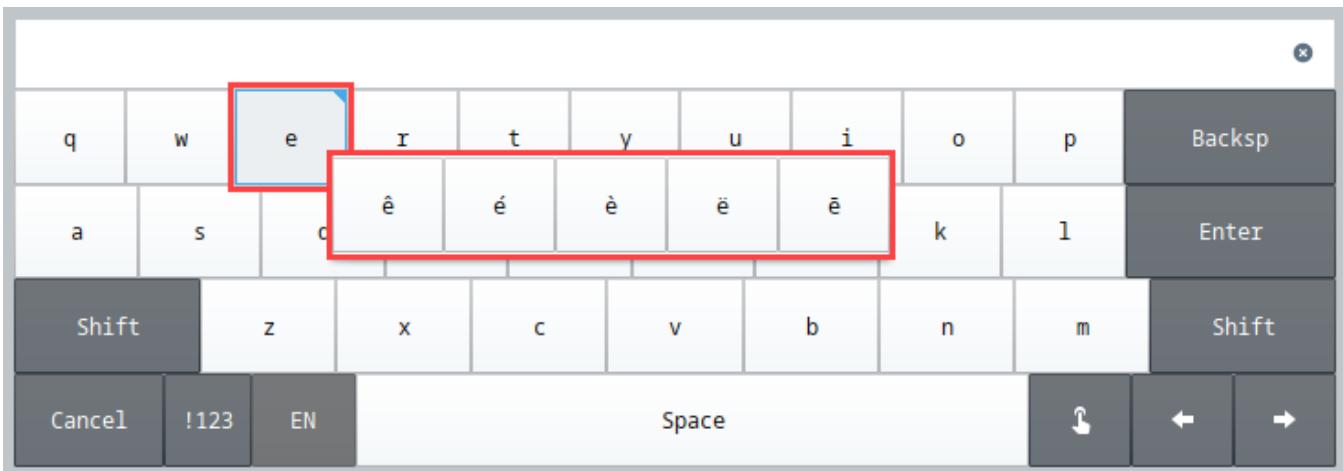


Additionally, there is a clear all icon at the end of the preview text field to make it easier to remove all existing text in the field.

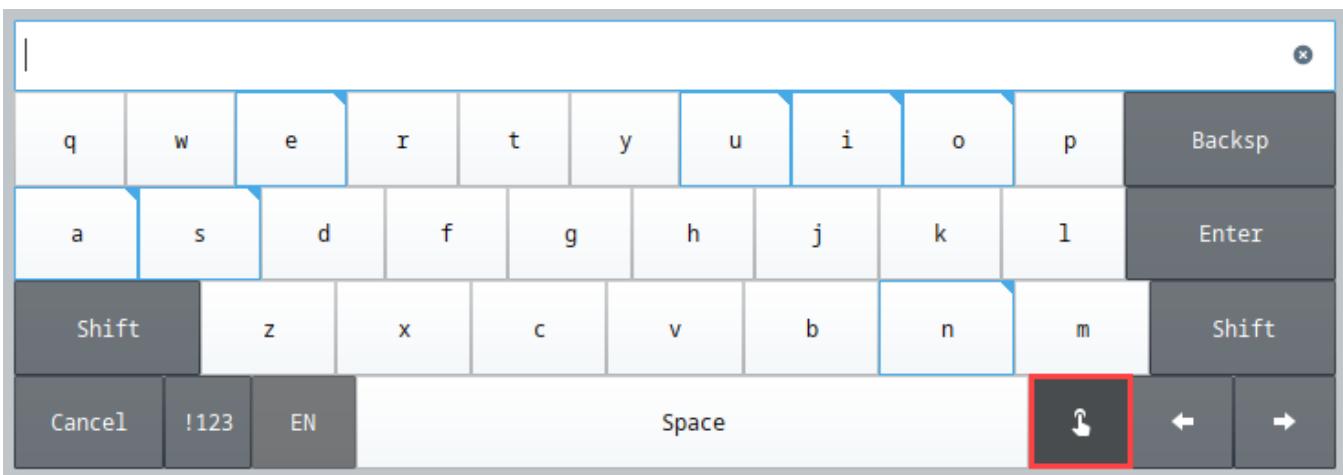
Note: For easy navigation within the text body, use the Shift and Arrow Keys to jump to the start or end of the current text. The arrow keys can be held to repeatedly navigate within the text field. Similarly, the backspace key can be held down to remove multiple characters at a time.

Using the Long Press Feature

Long press on any key with accent characters defined and a pop up window will open allowing you to enter any alternate variations of that key.



Another option to access accent keys is to touch the Finger Push key at the lower right of the keyboard. When this is selected, all keys with accent options will be highlighted in blue and a single touch will invoke the accents pop up.



Note: If even greater configuration of the long press feature is required, the long press delay time can be adjusted using the `ignition.touchescreen.longPressDelay` system property.

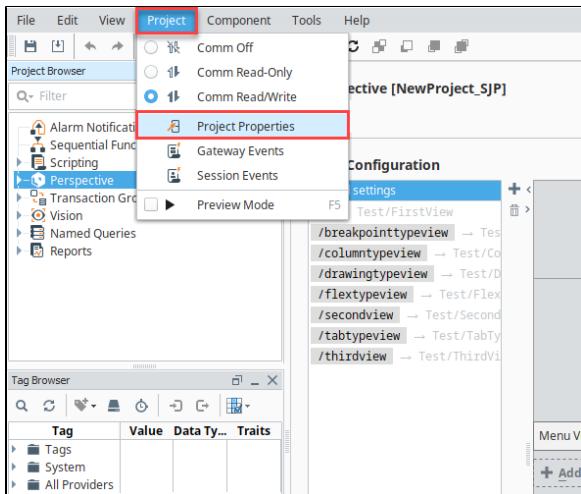
Project Properties

There are a number of properties you can set for your projects within the Designer, Vision Clients, and Perspective Sessions. For example, there are properties for setting the touchscreen mode, customizing a client's auto-login, or configuring how the clients receive updates.

The property settings on the Project Property window apply to the whole project. This page identifies and describes all the available project properties.

Accessing Project Properties

In the Designer, click on Project tab on the menu bar at the top of the Designer, then select **Project Properties**.



The Project Properties screen is displayed. Project Properties span several functional areas each containing settings applicable to that area.

Project General Properties

A project's General properties apply to the project as a whole, across all module functionality. For a new project, there are many default settings that you can use. For example, there are default settings for the Tag provider, database, initial comm mode, window editing, and Client launching.

When properties in a section have been updated but not saved yet, the section heading will change to italicized text.

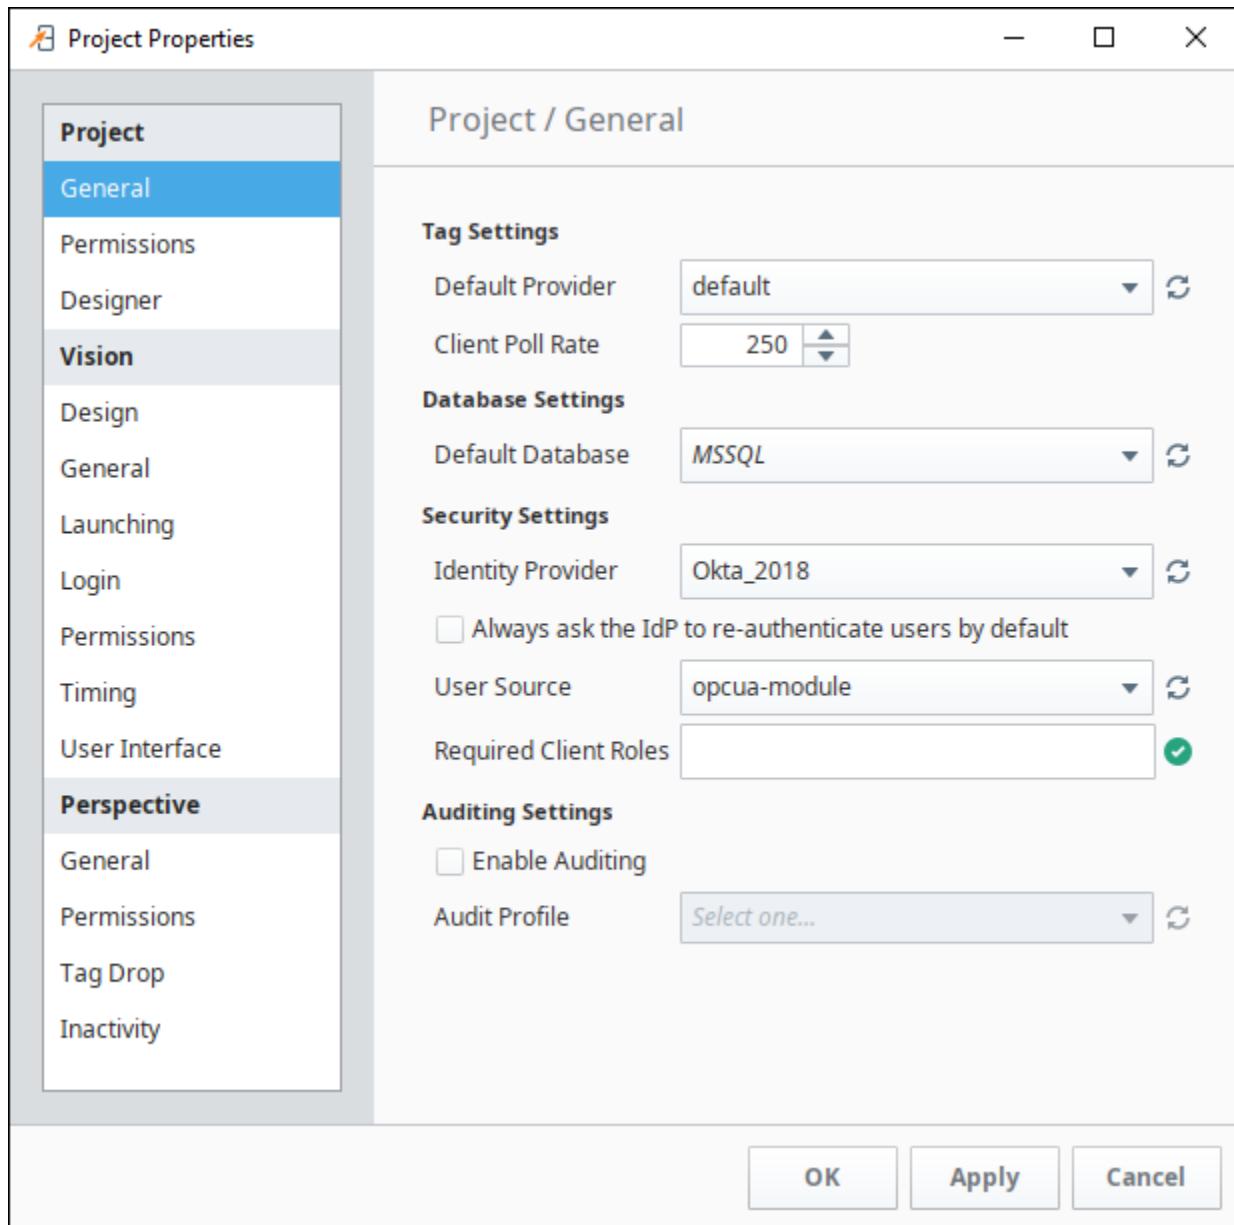
On this page ...

- [Accessing Project Properties](#)
- [Project General Properties](#)
 - [Default Database and Default Tag Provider](#)
 - [General Properties](#)
 - [Permissions Properties](#)
 - [Designer Properties](#)
- [Vision Project Properties](#)
- [Perspective Properties](#)
- [Property Inheritance](#)

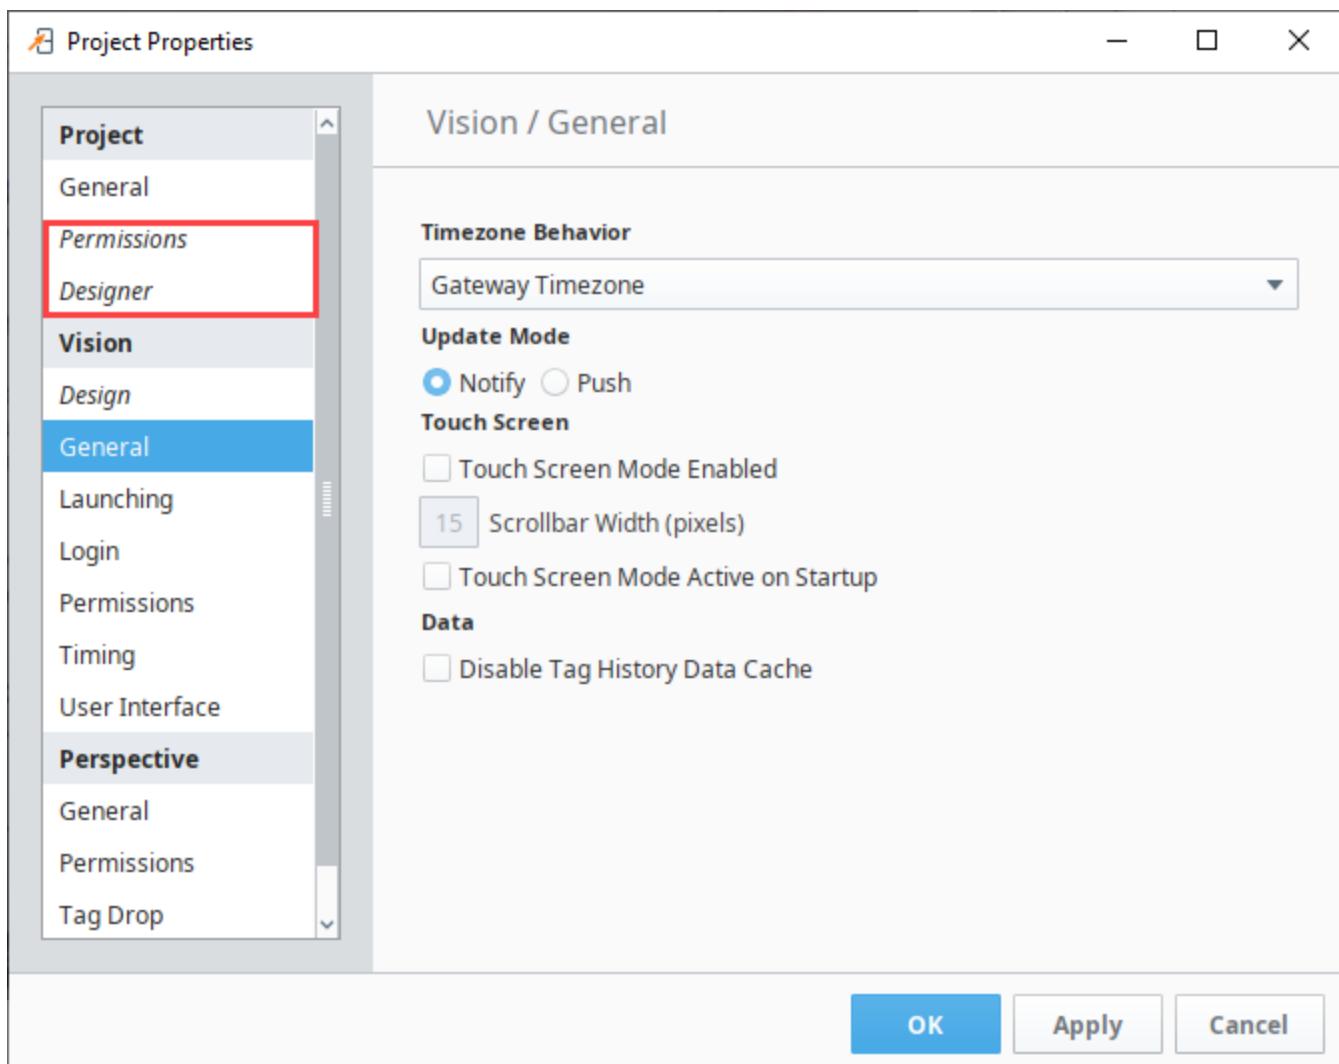


About Project Properties

[Watch the Video](#)



When properties in a section have been updated but not saved yet, the section heading will change to italicized text. In the following example, changes have been made in the the Perspective General and Permissions properties but they have not been saved or applied yet.



Note: There are a few properties of a project, such as its name, description, and title that are not available from the Designer and you need to set them in the Gateway. To do so, go to the [Gateway](#), under the **Config > Projects** section, click on the **edit** link next to the project. You cannot edit these settings while the project is open in the Designer.

Default Database and Default Tag Provider

Under **Project > General** is where you set the project's Default Database and its Default Tag Provider. It is important to understand how to use defaults effectively for proper project design.

Wherever you use a database connection or a Tag in a project, you are always given the option to use the project's default, or an explicitly named connection or provider. If your project is like most typical projects, it primarily uses a single database and a single Tag provider. By consistently using the default option, you make your project more resilient to change.

For example, suppose you have a project, and it has a database connection named `Production_DB`. Now you want to adapt the project to a new, similar plant, while leaving the existing project intact. You copy the project and create a new database connection, named `New_DB`. If your project consistently used its default database connection, the switchover will be as simple as changing the copied project's default database. However, if you used the explicit `Production_DB` connection in your groups and screens, you will need to laboriously switch the bindings over to `New_DB`.

General Properties

Tag Settings	
Property	Description
Default Provider	The Tag provider chosen here will act as the project's default Tag provider.

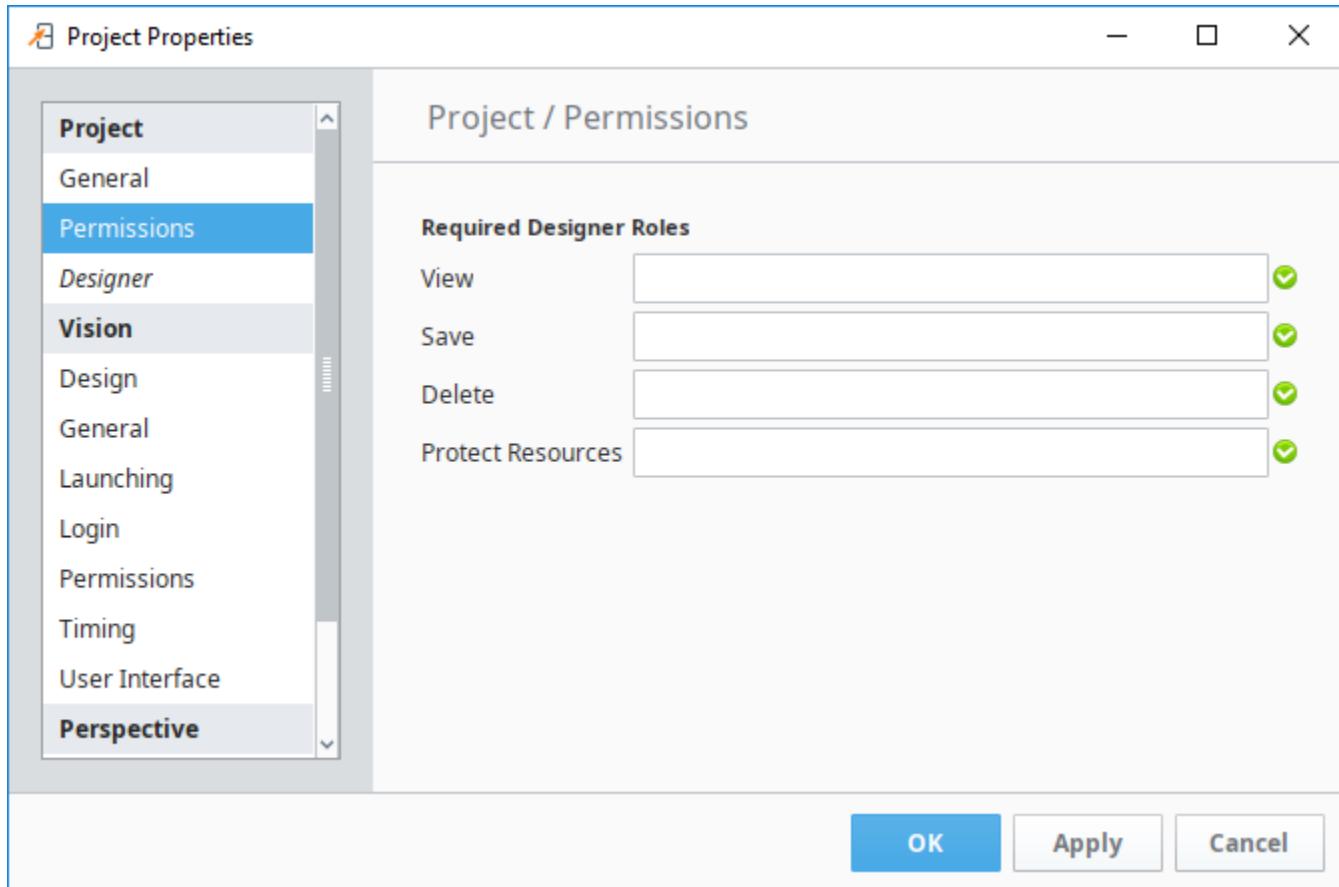
Client Poll Rate	The Client poll rate (in milliseconds) is the rate at which a Vision Client or Ignition Designer polls the Gateway for updates to its subscribed Tags.
Database Settings	
Default Database	This is the default database connection to use for this project. To use the default database connection, use the special <default> connection, or in scripting, the empty-string connection "".
Security Settings	
Identity Provider	<p>The default Identity Provider (IdP) for this project. Choose from a dropdown list of IdPs. You can also choose <None> if the project already had an IdP but you want to remove it.</p> <p>Note: If your Identity Provider was set in an Ignition version prior to 8.0.6, this property may be located in Project Properties > Perspective General instead.</p>
Always ask the IdP to re-authenticate users by default	<p>When enabled, authenticated users will always need to re-enter their credentials on login. When false, if a user is already authenticated, then they will not be required to re-type their credentials when an action in the session triggers a login (such as a call to <code>system.perspective.login</code>).</p> <p>Note: Allowing re-authentication is entirely determined by the Identity Provider. This setting simply makes the Gateway send the re-authentication request. Consult your Identity Provider's documentation for information on re-authentication support.</p>
Automatically redirect users to the IdP if authentication is required	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>If enabled, users will be automatically redirected to the IdP for authentication if the user is not yet authenticated and authentication is required. If disabled, users will be shown a "speed bump" page, with a message explaining how authentication is required and a button allowing the user to redirect to the IdP for authentication.</p> <p>Note: This setting only applies to Perspective Sessions.</p>
User Source	<p>Choose the User Source that governs this project's security: specifically, which group of users (User Sources) is allowed to log into the client. The User Sources are all defined in Gateway's Config section under the Security > User, Roles page.</p> <p>Note: This is for Vision projects only.</p>
Required Client Roles	<p>This property is for Client logins, and determines what role(s) a user must have before they can log into the Client. You can optionally specify a list of roles that are required for a user to log into this project. Use commas to separate the roles. Users must have at least one of the roles in order to log in. If no roles are specified, the user only needs to correctly authenticate with the User Source in order to log in. To define the roles, go to the Gateway's Config section under the Security > User, Roles page. Click the manage users link, and then go to the Roles tab.</p> <p>See also Security.</p>
Audit Settings	
Enable Auditing	If auditing is enabled, audit events that relate to this project will be stored in the chosen audit profile.
Audit Profile	The audit profile stores the audit events when auditing is enabled.

Permissions Properties

When opening the project in the Designer, these properties determine which roles are required when making certain changes to the project. More details can be found on the [Project Security in Designer and Gateway](#) page.

Required Designer Roles	
View	Users must have at least one of these roles to view the project in the Designer.

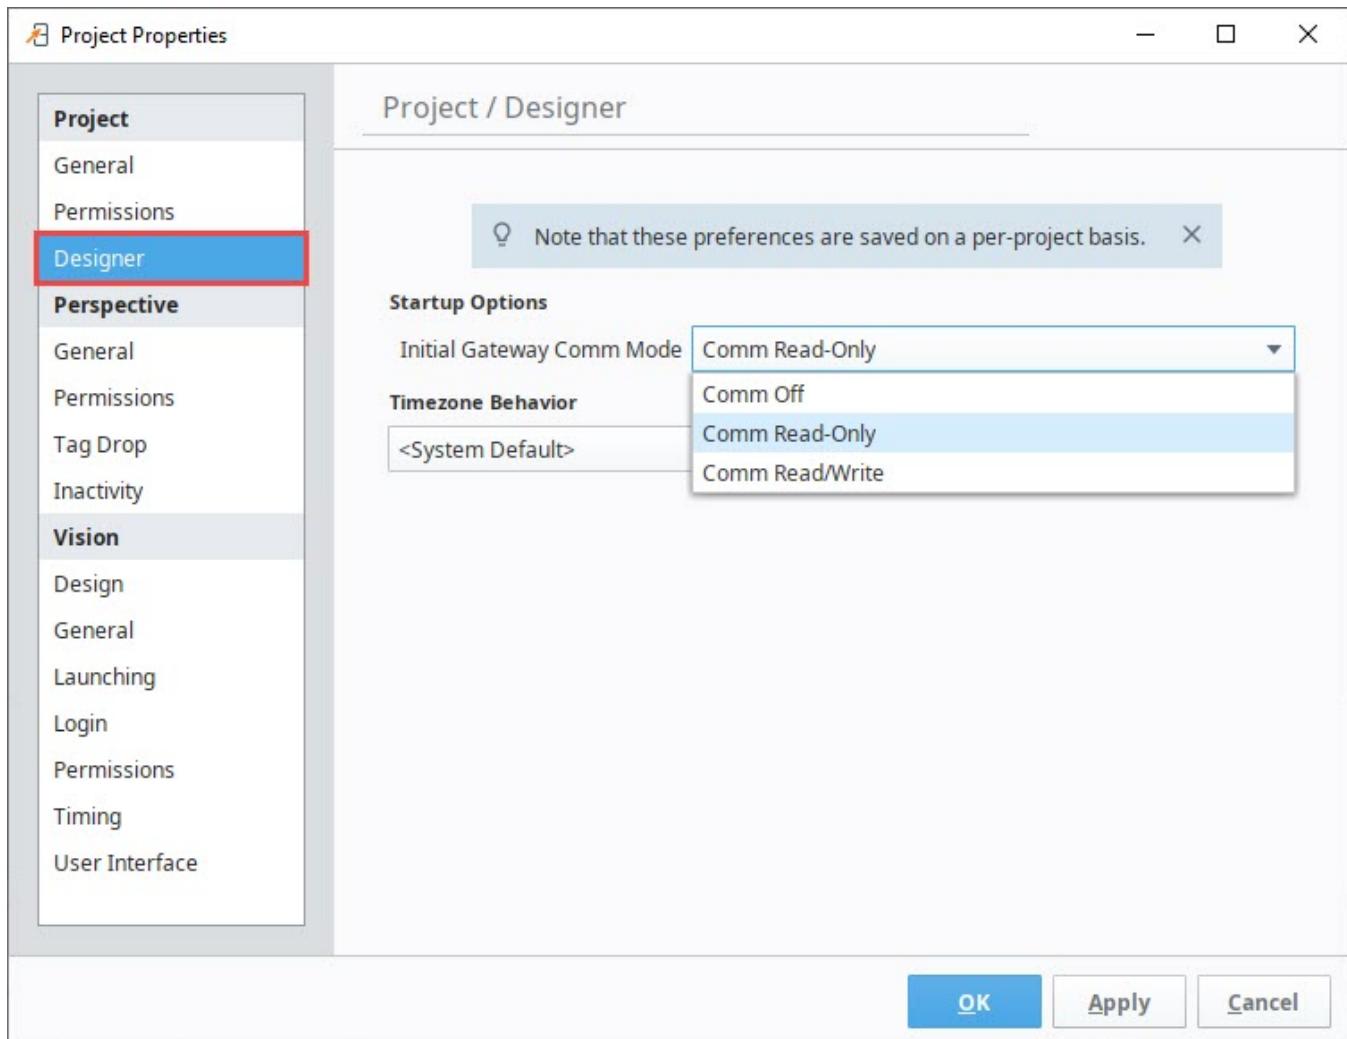
Save	Users must have at least one of these roles to save the project.
Delete	Users must have at least one of these roles to delete the project.
Protected Resources	Users must have at least one of these roles to access protected resources.



Designer Properties

These preferences are saved on a per-project basis.

Startup Options	
Initial Gateway Comm Mode	<p>The Designer starts up in the default Comm Read-Only mode. The property allows you to change the mode the Designer starts in when viewing the project. The options are: Comm Off, Comm Ready-Only, Comm Read/Write.</p> <ul style="list-style-type: none"> • Comm Off - In this mode, all database query traffic and Tag subscriptions and writes are blocked. • Comm Read-Only - Tag subscriptions and SELECT queries work, but Tag writes and UPDATE/INSERT/DELETE queries are blocked. • Comm Read/Write - The Designer may freely request Tag and database values from the Gateway, as well as write or change these values. <p>For more information, see the Communication Modes page.</p>
Timezone Behavior	<p>The following feature is new in Ignition version 8.1.22 Click here to check out the other new features</p> <p>Sets the timezone of the Designer. This takes precedence over Vision's Client Timezone setting If Vision is installed and applies only to the Designer.</p>



Vision Project Properties

There are many project properties that apply specifically to Vision Clients. You can find more information at [Vision Project Properties](#).

Perspective Properties

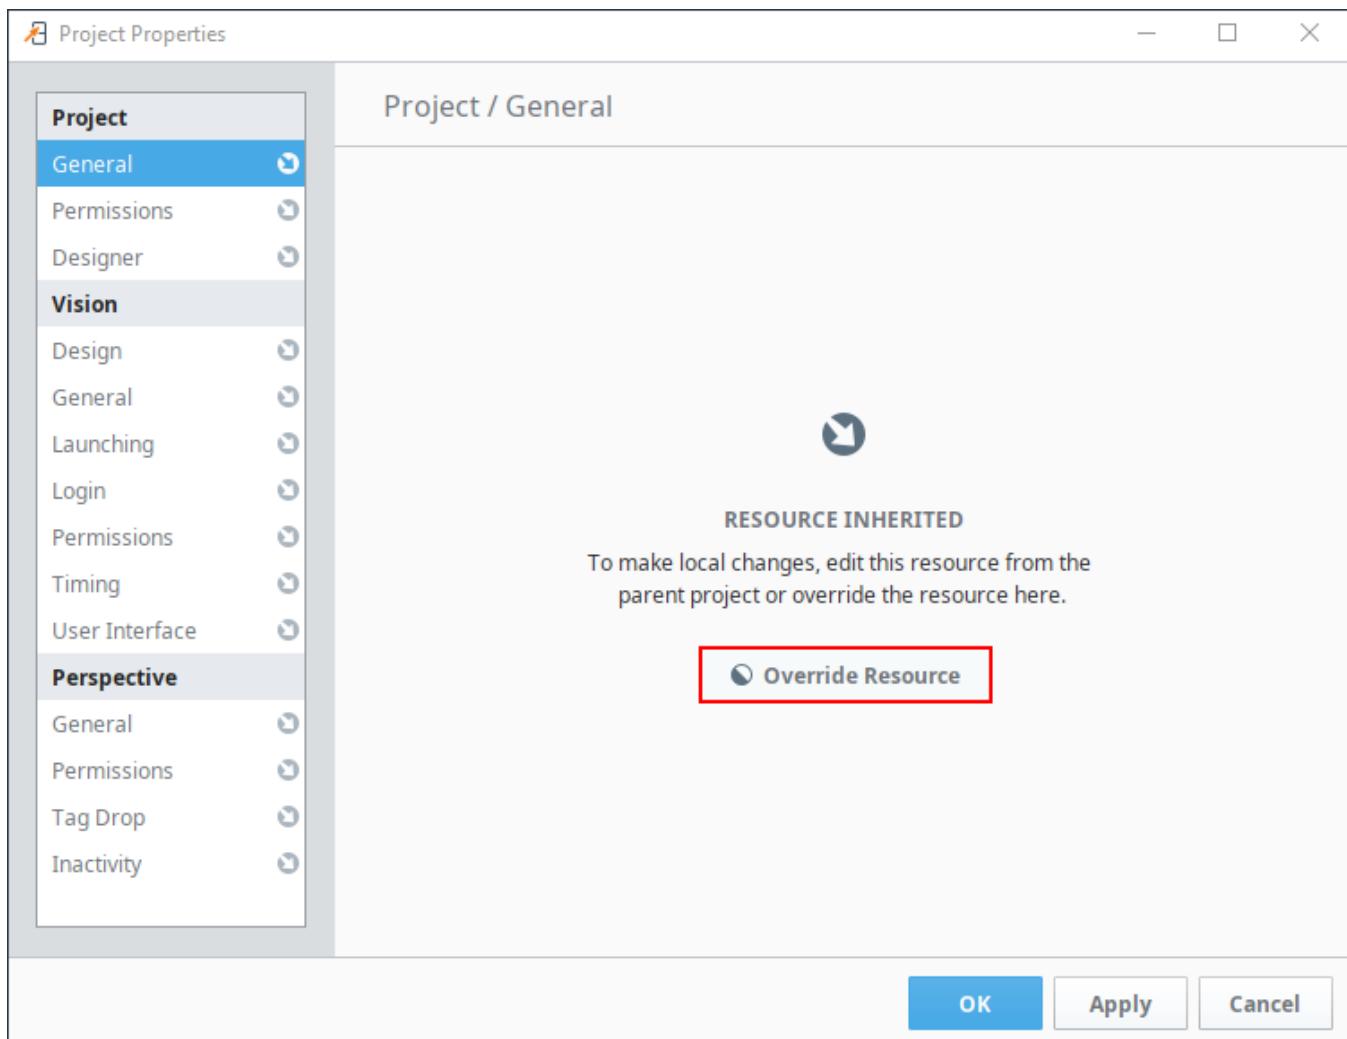
There are many project properties that apply specifically to Perspective Sessions. You can find more information at [Perspective Project Properties](#).

Property Inheritance

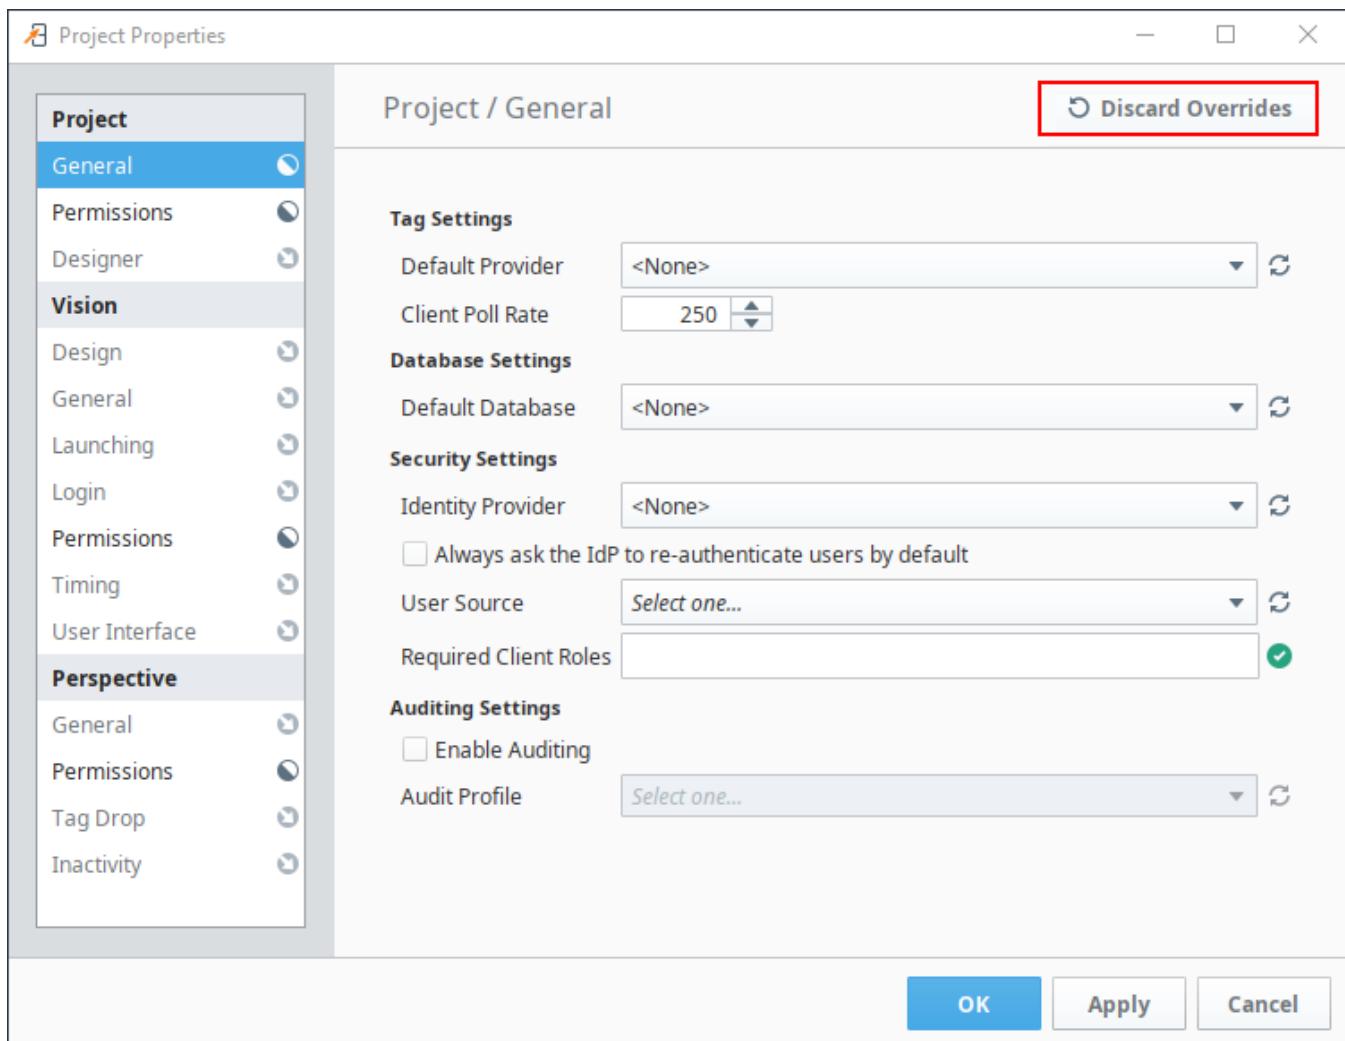
Project properties can be inherited from parent projects. You can find more information at [Project Inheritance](#).

The following feature is new in Ignition version 8.1.2
[Click here](#) to check out the other new features

Inheriting Project Properties results in a **Resource Inherited** overlay on the section. The **Override Resource** button can be used to make changes locally:

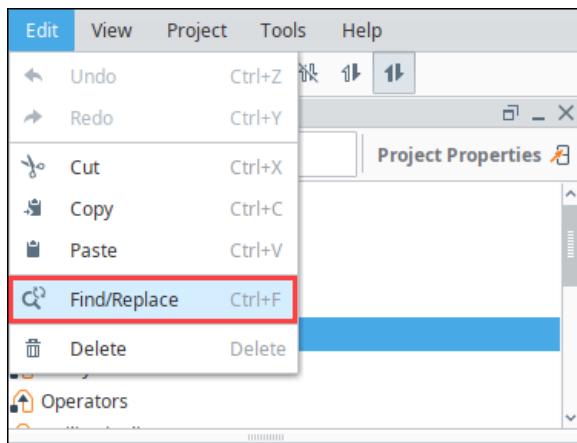


While a resource is overridden, the **Discard Overrides** button can be used to clear overrides for the current section.



Find and Replace

The Find/Replace tool in the Edit menu of the Designer allows you to search your entire project for specific components, properties, scripts, Transaction Groups, and Tags. You can then use the replace feature to make mass changes to a project with the click of a button.



On this page ...

- [Search Options](#)
- [Contextual Find and Replace](#)
- [Wildcards](#)
- [Using Find and Replace](#)



INDUCTIVE
UNIVERSIT

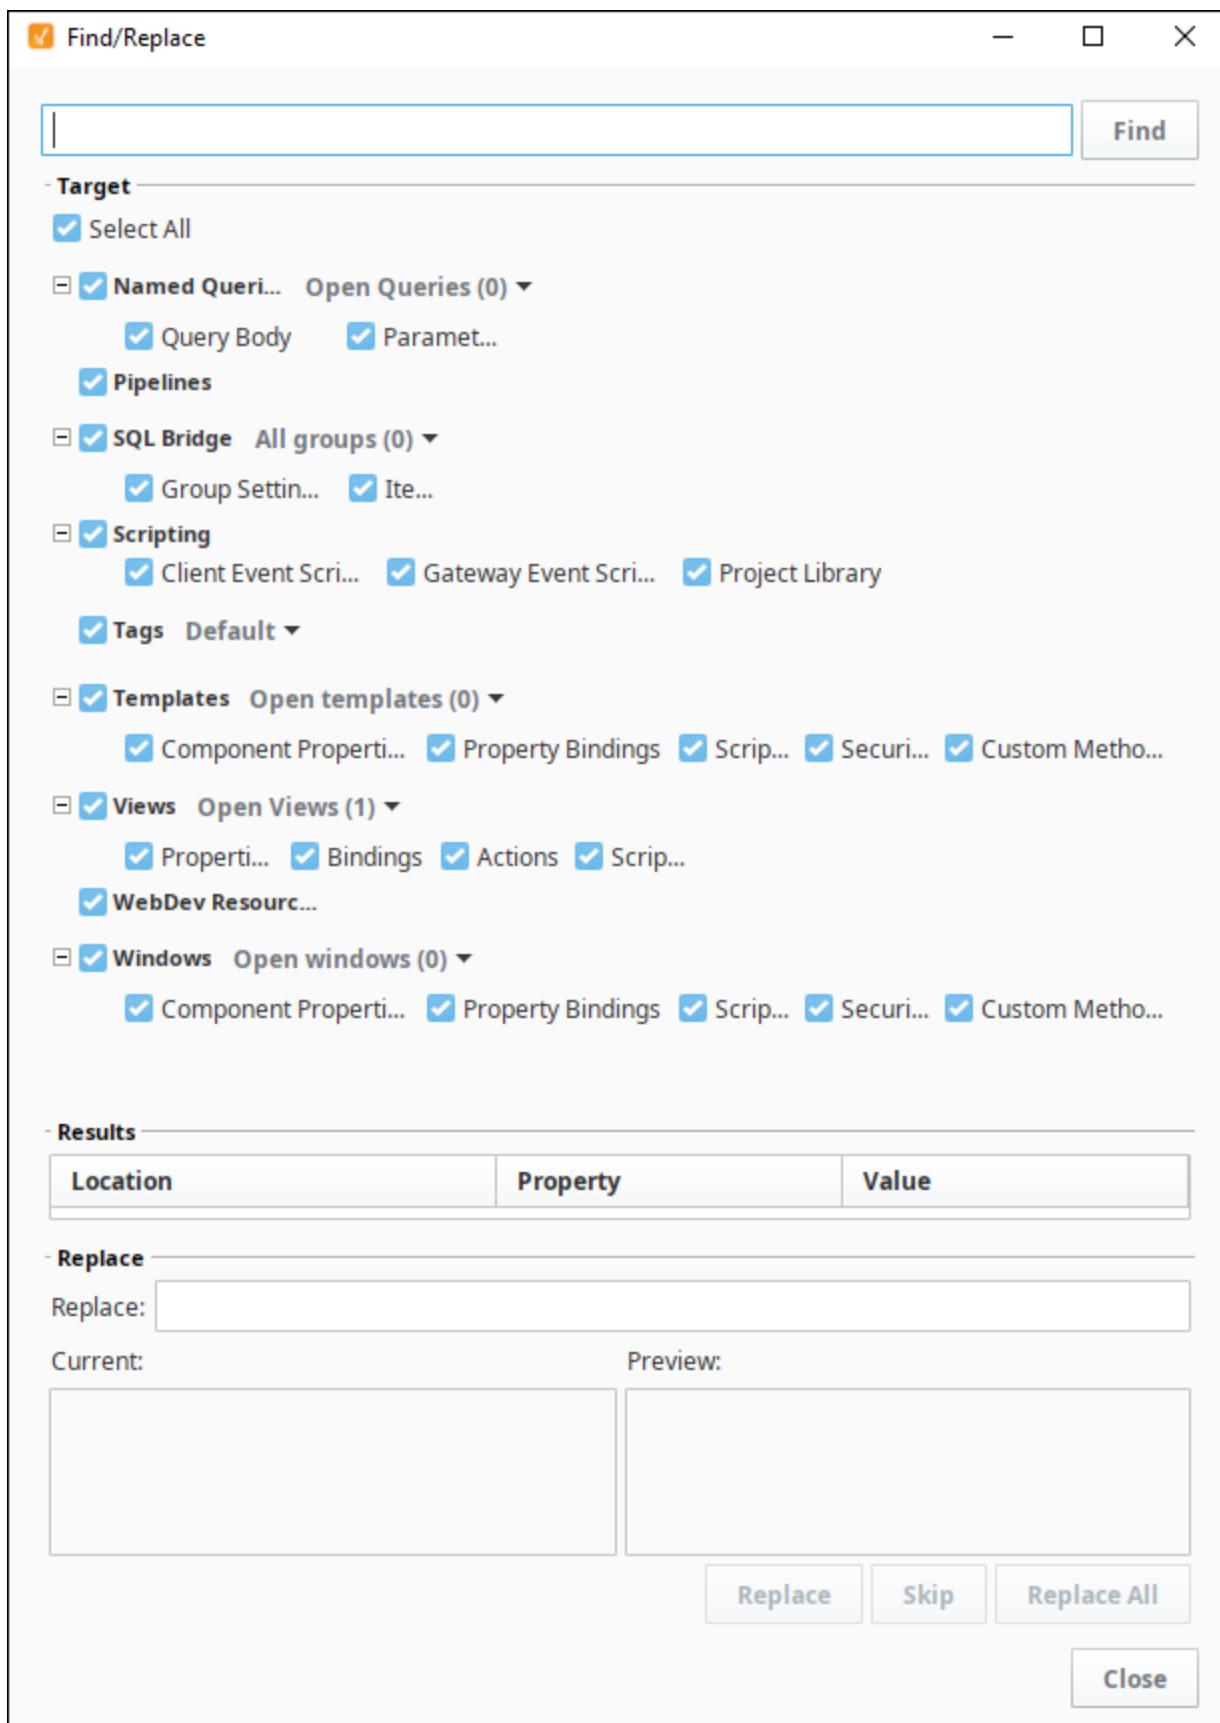
Using Find and Replace

[Watch the Video](#)

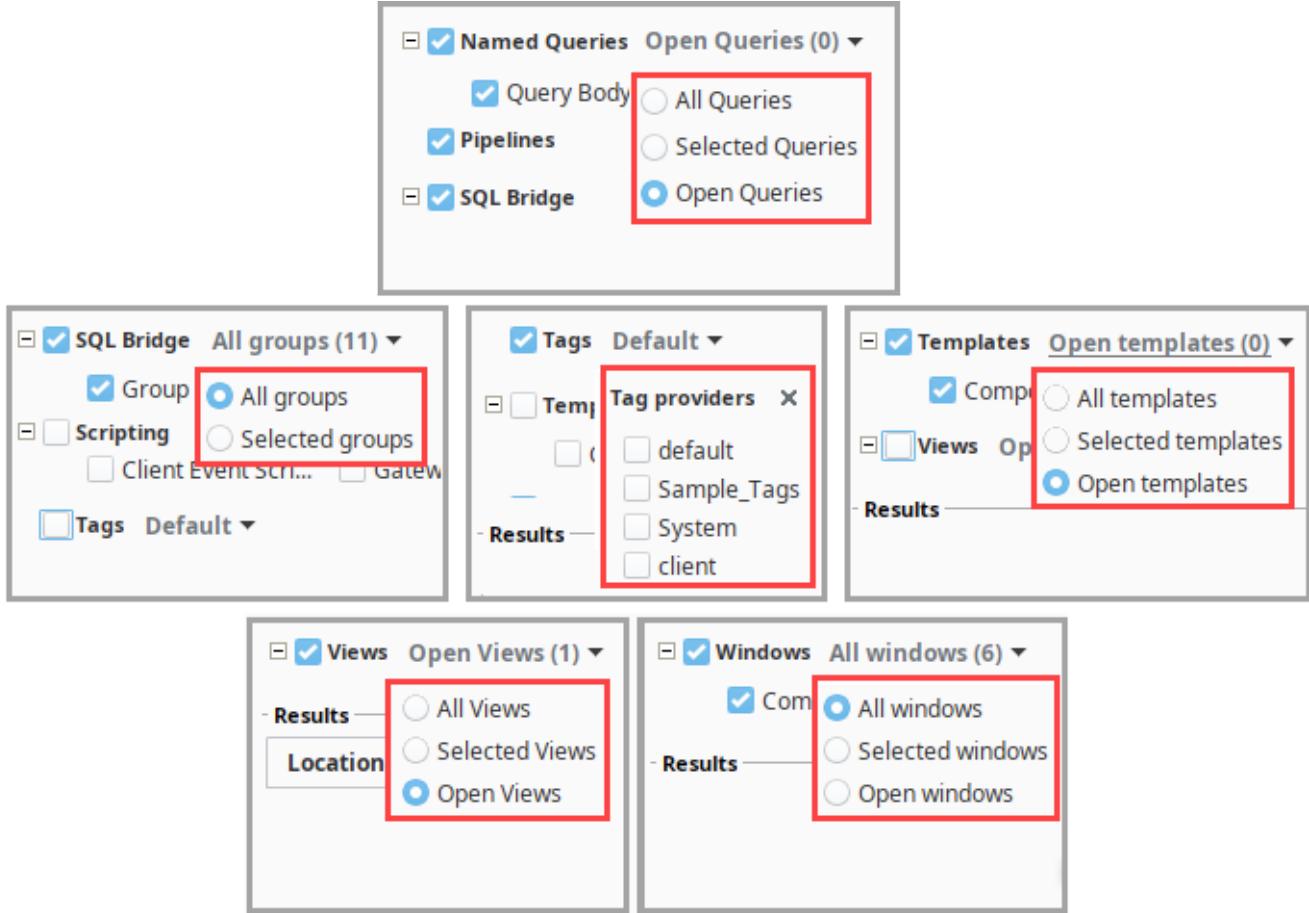
Search Options

In the Target section, the Find/Replace tool has options for searching through many different parts of a project.

- Named Queries (Added in 8.1.3)
- Pipelines
- SQL Bridge Transaction Groups
- Scripting
- Tags
- Templates (Vision)
- Views (Perspective)
- Windows (Vision)
- WebDev Resources (Added in 8.1.1)

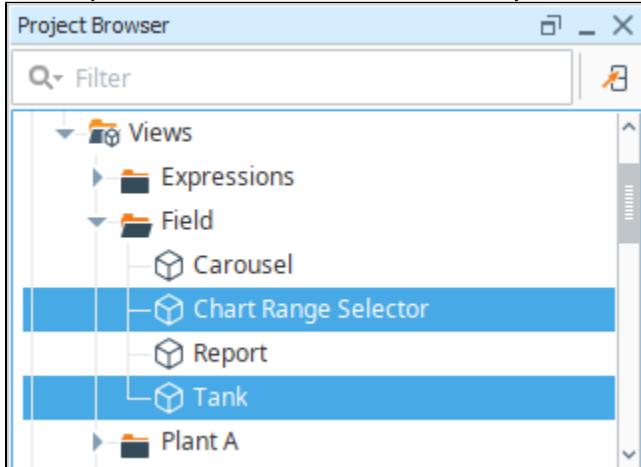


You can narrow down your search by selecting and deselecting categories you include in the search. The SQL Bridge, Tags, Templates, Views, Named Queries, and Windows options also have dropdown options to customize your search.

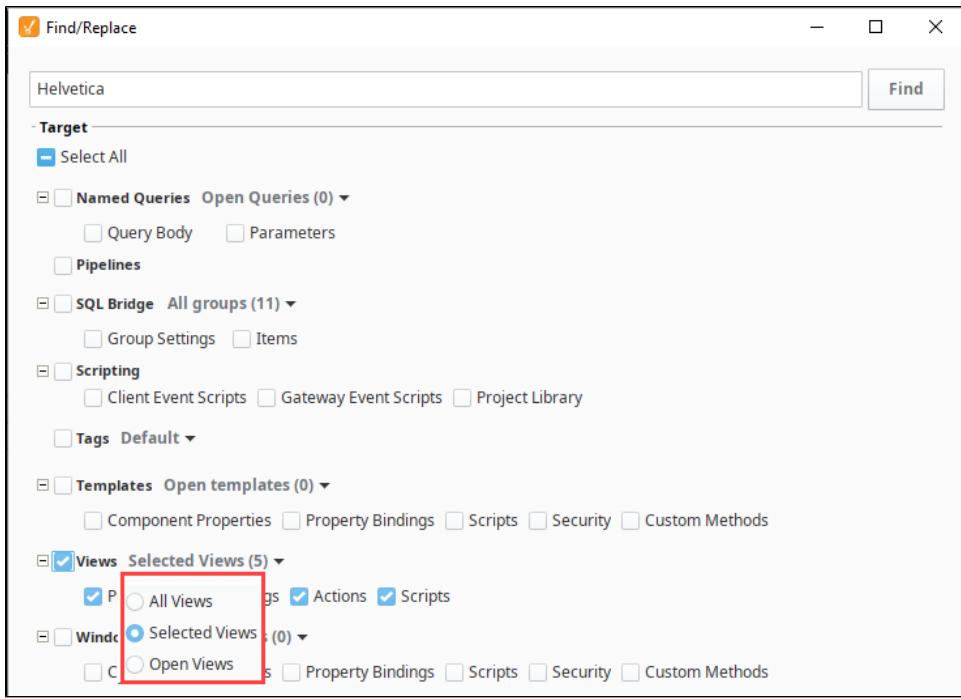


For example, if you want to search only a couple Views, do the following:

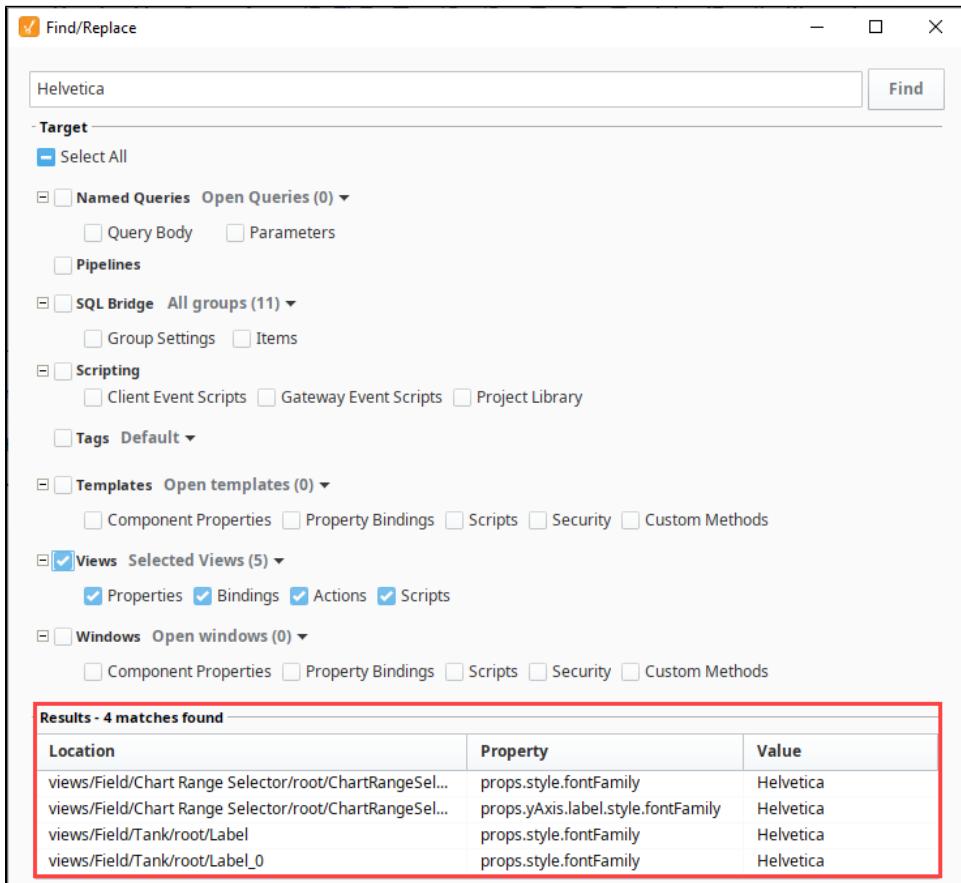
1. In the Project Browser, use **Ctrl-Click** to select the views you want to search.



2. In the Designer toolbar, go to **Edit > Find/Replace** or use the shortcut **Ctrl-F**. The Find/Replace window is displayed.
3. De-select all options by unchecking **Select All**.
4. In Views, click on the expand arrow and choose the **Selected Views** option.



5. Enter the property, action, script, or such that you want to search for. In this example, we searched for property value **Helvetica**.
6. Click **Find**.
7. The search results are displayed in the **Results** box.

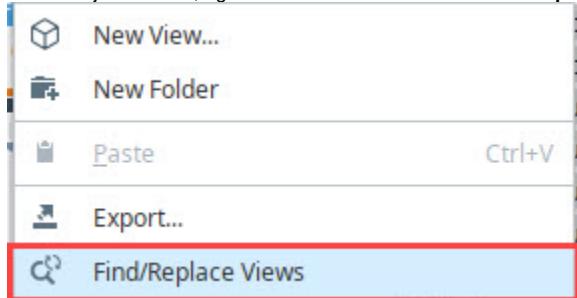


Contextual Find and Replace

The following feature is new in Ignition version **8.1.18**
[Click here](#) to check out the other new features

Starting in 8.1.18, you can right-click on a node to Find/Replace within all Windows, Templates, Views, Transaction Groups, Named Queries, Scripts, and/or selected folders in that node. This bypasses the need to narrow your search criteria with the Search Options detailed in the previous section. For example, to search for all instances of a "Gauge" within all Perspective Views, follow these steps:

1. In the Project Brower, right-click on **Views** and select **Find/Replace Views**.



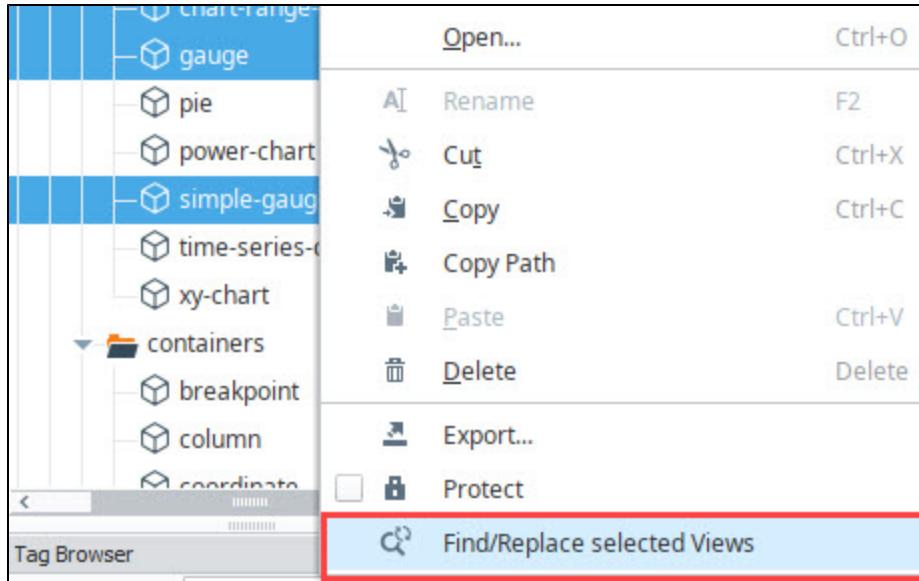
2. This will open the Find/Replace window. Type "gauge" in the search bar and click **Find**.

3. The results displayed will be all instances of "gauge" within the Views node:

The Find/Replace window is titled 'Find/Replace - All Views (154)'. The search bar contains 'gauge'. The results table shows 38 matches found, with columns for Location, Property, and Value. One row is selected, showing 'views/components/component-views/fra...' under Location, 'meta.name' under Property, and 'SimpleGauge' under Value. Below the table is a 'Replace' section with 'Current:' and 'Preview:' fields. The 'Current:' field shows 'SimpleGauge' and the 'Preview:' field shows 'Simple'. At the bottom are 'Replace', 'Skip', 'Replace All', and 'Close' buttons.

Results within All Views (154) - 38 matches found		
Location	Property	Value
views/components/component-views/fra...	View Name	simple-gauge
views/components/component-views/fra...	meta.name	SimpleGauge
views/components/component-views/dis...	props.availableWidgets[0].head...	Gauge
views/components/component-views/dis...	props.availableWidgets[0].name	Gauge
views/components/component-views/dis...	props.availableWidgets[0].viewP...	framework/widgets/gauge
views/components/component-views/cha...	View Name	simple-gauge
views/components/component-views/cha...	props.params.text	Simple Gauge
views/components/component-views/cha...	props.text	The Simple Gauge component p...
views/components/nav/charts/root/Menu...	props.items[1].label.icon.path	sample-components/chart-gauge
views/components/nav/charts/root/Menu...	props.items[1].label.text	Gauge
views/components/nav/charts/root/Menu...	props.items[1].target	/components/charts/gauge

To limit your search within selected views, you may select specific views or folders within the Project Brower. Right-click on any of the selected views and select **Find/Replace selected Views**.



Wildcards

To customize your search further, you can use the wildcard characters * and ?.

Use an asterisk (*) to indicate that any character(s) could be where the asterisk is. For example, to search for Tags that start with the word "Motor" enter "Motor*". This would return Motor 1, Motor2, Motor_East3, and so forth.

Use a question mark (?) to indicate any single character could be where the question mark is. For example, to search for "Valve1Status" through "Valve9Status", you could enter "Valve?Status". Note that this would not return something like Valve44Status because that is more than one character where the wildcard is.

Editor notes are only visible to logged in users

The multi-line wildcard search was implemented as a part of IGN-821, but looks like there was some regressions that prevent it from actually working.

IGN-1113 is the fix for the issue. Let's hold off on the editor note below until it's fixed.

The following feature is new in Ignition version 8.1.1

[Click here](#) to check out the other new features

As of release 8.1.1 Find/Replace can now use wildcards to match items across multiple lines of text.

Using Find and Replace

1. In the Designer toolbar, go to **Edit > Find/Replace** or use the shortcut **Ctrl-F**. The **Find/Replace** window is displayed.
2. Enter what you're searching for in the textbox at the top to search through your project.
3. Click the **Find** button at the top right to execute the search. All matching items appear in the **Results** section. Double-click on an item in the Results table to bring that item into editing focus in the Designer.
4. To replace a value, select an entry in the Results section and you'll see the Value of that entry in the **Current** box of the **Replace** section.
5. Enter the text you want to use as a replacement in the **Replace** textbox and a preview of the new value is shown in the **Preview** box.
6. Click the **Replace** button to execute the replace. This will move your selection down in the Results table so that you can rapidly execute multiple replacements. If you're satisfied and you'd like to make the identical replacement to many items, select them all in the Results table and click the **Replace All** button.
7. Close the Find and Replace window.

Note: The actual replacement on some resources may not take effect until the Find and Replace window is closed.

Related Topics ...

- [Database Query Browser](#)

Windows, Linux, and Mac Keyboard Shortcuts

Using Keyboard Shortcuts in the Designer

You can interact with screens and more by using some of the popular keyboard shortcuts in the Designer.

Action	Windows /Linux Keystrokes	MacOS Keystrokes
Nudge Moves selected component(s) in the direction of the arrow key by the default nudge distance.	, , ,	, , ,
Alt-Nudge Same as Nudge, but uses the “alt-nudge” distance.	Alt + , Alt + , Alt + , Alt +	Option + , Option + , Option + , Option +
Resize Right Moves the right edge of the component left or right. Add Alt to use the alt-nudge distance.	Shift + , Shift +	Shift + , Shift +
Resize Bottom Moves the bottom edge of the component to top or bottom.	Shift + , Shift +	Shift + , Shift +
Resize Left Moves the left edge of the component left or right.	Ctrl + Shift + , Ctrl + Shift +	Control + Shift + , Control + Shift +
Resize Top Moves the top edge of the component to top or bottom.	Ctrl + Shift + , Ctrl + Shift +	Control + Shift + , Control + Shift +
Move Forward (Vision only) Moves the selected components(s) forward in the Z-order	PgUp	Fn +
Move Backward (Vision only) Moves the selected component(s) backwards in the Z-order	PgDn	Fn +
Move to Front (Vision only) Moves the selected component(s) to the front of the Z-order	Home	Fn +
Move to Back (Vision only) Moves the selected component(s) to the back of the Z-order	End	Fn +
Copy-Move Copies the component when holding Ctrl while doing a mousemove.	Ctrl + drag a component	Click and hold on a component + Control + drag
Orthogonal-Move Restricts to only moving straight up, down, left, or right when holding Shift while doing a mousemove.	Shift + drag a component	Shift + drag a component
Selection-Move Drags the components that are currently selected without having to press the mouse button down on the component first.	Hold Alt + drag a component	Hold Option + drag a component
Copy-Axis-Move Combines copy move with axis move.	Hold Ctrl + Alt + drag a component	Click and hold on a component + Control + Option + drag
Proportional Resize Resizes a component while maintaining its aspect ratio.	Hold Ctrl + resize a component	Click and hold on a component + Control + resize
On-Center Resize Resizes the selection using the center as the anchor point.	Hold Shift+ resize a component	Hold Shift + resize a component
Select All (Vision only) Selects all components that are siblings of the selected component.	Ctrl + A	Command + A
Select Same Type (Vision only) Selects all components that are siblings of the selected component and the same component type.	Ctrl + Shift + A	Command + Shift + A
Select Same Type in Window (Vision only) Selects all components in a window that are the same type as the selected component.	Ctrl + Alt + Shift + A	Command + Option + Shift + A
Layout (Vision only) Opens the Layout Constraints window to let you specify layout for the component(s).	Ctrl + L	Command + L

Size & Position (Vision only) Opens Size & Position to let you specify exact size and position for selected component or window.	Ctrl + P	Command + P
Customizer (Vision only) Opens the Custom Properties window to let you configure complex component properties.	Ctrl + U	Control + U
Customizer 2 (Vision only) Some components have a secondary customizer. This command will open that customizer, if available.	Ctrl + 2	Control + 2
Jython Opens Event Configuration page (Perspective) or the Component Scripting page (Vision) to let you configure actions for component or window.	Ctrl + J	Command + J
Script Configuration Opens the Script Configuration for the View or Component selected. (Perspective Only)	Ctrl + K	Command + K
Security (Vision only) Opens Security Settings to let you set security for Component(s) or window.	Ctrl + E	Command + E
Save Saves the project.	Ctrl + S	Command + S
Open Opens the Open/Create Project to let you select a different project.	Ctrl + O	Command + O
Undo Undoes the last action.	Ctrl + Z	Command + Z
Redo Gets rid of the last undo action.	Ctrl + Y	Command + Y
Copy Copies selected component(s) or window.	Ctrl + C	Command + C
Duplicate Duplicates selected component(s) or window.	Ctrl + D	Command + D
Cut Cuts selected component(s) or window.	Ctrl + X	Command + X
Paste Pastes selected component(s) or window in clipboard. Pasted component(s) wait for position before pasting.	Ctrl + V	Command + V
Immediately Paste (Vision only) Places pasted component(s) at the same location where they were copied/cut.	Ctrl + I	Command + I
Comment/Uncomment Lines of Code Quickly comment or uncomment lines of a script or query in Designer.	Ctrl + /	Command + /
Cancel Cancels a pending paste operation, deselects the current row of a table, cancels dragging components onto window.	Esc	Esc
Find/Replace Opens Find/Replace to let you Search and replace the project based on string, pattern, or regex.	Ctrl + F	Command + F
Delete Deletes the current selection.	Delete	Delete
Snap to Grid (Vision only) Toggles whether or not moving and resizing components snaps to the grid.	Ctrl + G	Command + G
Snap to Guides (Vision only) Toggles whether or not moving and resizing components snaps to guides.	Ctrl + Shift + G	Command + Shift + G
Console Opens the Output Console .	Ctrl + Shift + C	Command + Shift + C
Help Launches the Ignition User Manual in a web browser.	F1	Fn + F1
Rename Renames the selected item (Tag, window, transaction group, component, and so on).	F2	Fn + F2
Preview Mode Toggles preview/design mode.	F5	Fn + F5
	F10	Fn + F10

Launch Client/Session (Perspective and Vision only) When in the Vision workspace, launches a Vision Client in windowed mode.		
While in the Perspective workspace, launches a Perspective Session in windowed mode.		
Launch Full Screen Client (Perspective and Vision only) While in the Vision workspace, launches a Vision Client in full screen mode.	F11	Fn + F11
While in the Perspective workspace, launches a Perspective Session in full screen mode.		
Update Project Updates project to server. Receives concurrent edits from other Designers.	Ctrl + Shift + U	Command + Shift + U
Zoom Zooms in/out in the Designer.	Ctrl + Mousewheel	Control + Mousewheel
Touch Selection (Vision only) Draws a line while dragging. Every component in the path of the line will be selected when letting go of the mouse button. To activate, a container must first be selected	Alt + Left Click and Drag	Option + Left Click and Drag
Select Through (Perspective and Vision only) Click on a component that is beneath another component. You can do this multiple times if there are several layers of components.	Alt + Click	Option + Click
Code Folding/Unfolding Select a line of code, and this command will collapse the selection. Press again while the folded code is selected, and the code will unfold. The Script Console has a separate command for code folding.	Ctrl + .	Command + .
Select Through (Perspective only) The following feature is new in Ignition version 8.1.26 Click here to check out the other new features Click on a component that is beneath another component. You can do this multiple times if there are several layers of components.	Ctrl + Shift + Click	Command + Shift + Click
Deep Selection (Perspective only) The following feature is new in Ignition version 8.1.26 Click here to check out the other new features Selects into a container and components within.	Alt + Shift + Click	Option + Shift + Click

Related Topics ...

- [Saving Projects](#)

Saving Projects

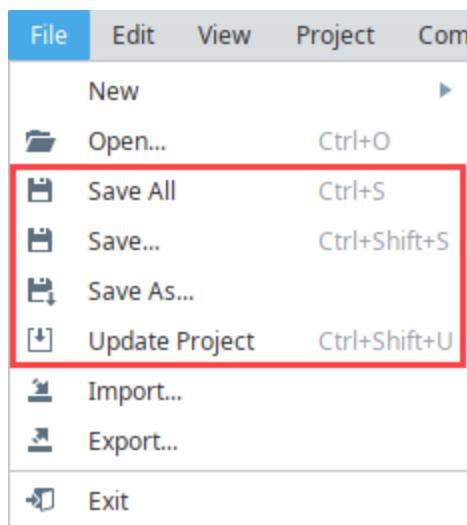
Once you created the Windows, Templates, Tags, Views, and anything else that goes into your project, you need to save your work. Saving your progress means pushing the changes from your Designer into the Gateway. If you are saving viewable resources like Vision Windows or Perspective Views, saving the changes means telling all of the clients or sessions that they can use the new updates to your project.

When you save, it's possible to save individual resources on a project as they are finished instead of saving the entire project. This is particularly helpful if you want to get resource updates to users quickly, or parts of a project into production sooner rather than waiting for the entire project to be completed.

You also have options on how to perform Client and Session updates. You can choose between notifying the operator of an available update, or automatically pushing updates as soon as it is published making it transparent to the operator. This page assumes you are using the default Notify mode, see [Client Update Modes](#) or [Session Project Updates](#) for more information.

Saving a Project

The project save options are all located in the Designer under the **File** menu.



On this page ...

- [Saving a Project](#)
 - [Save All](#)
 - [Save...](#)
 - [Save As](#)
- [Update Project](#)
 - [Project Updates in a Vision Client](#)
 - [Project Updates in a Perspective Session](#)



Project Creation and Publishing

[Watch the Video](#)

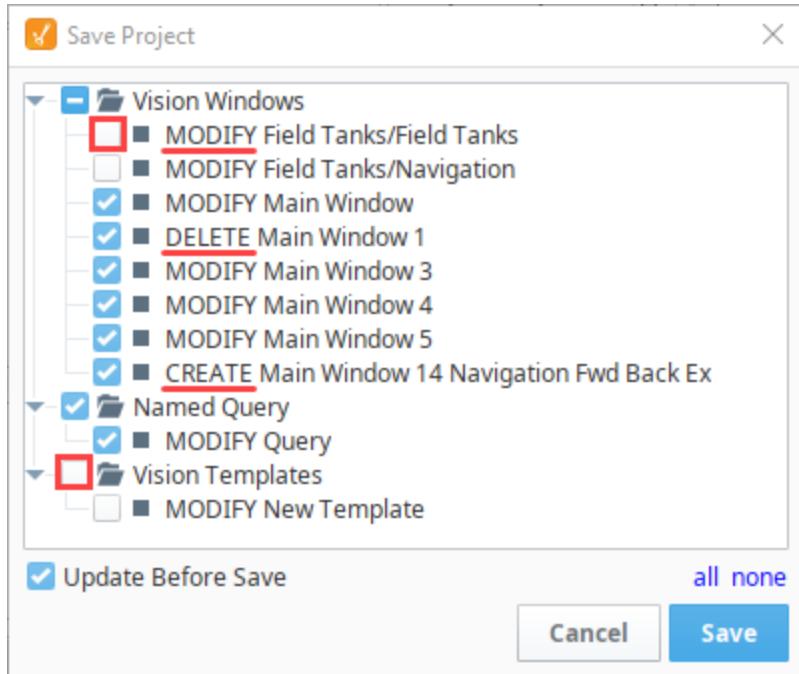
Save All

 **Save All** option saves the entire project. The project will be saved immediately, unless there are any conflicts due to [concurrent editing](#). If there are conflicts, the [Resolve Conflicts](#) screen will open automatically. Once you've resolved the conflicts, the Designer saves the project.

Save...

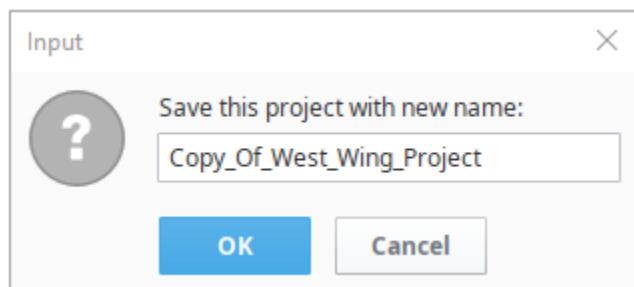
 The **Save...** option displays a list of all items that have been created, modified, or deleted since the last save. On this list you can save (or not save) individual resources on a project rather than of saving the entire project.

- To deselect all items, click "none" in the lower right corner.
- To select all items, click "all" in the lower right corner.
- To select or deselect individual items, click on the Checkbox  icon next to the item.



Save As

The **Save As**  option enables you to save your open project with a different name. When you choose **File > Save as**, the Designer will display an Input window. It will append "Copy_Of_" to the beginning of the current project name. However, you can enter a different name in this Input window. Click **OK** to save the project with a new name.

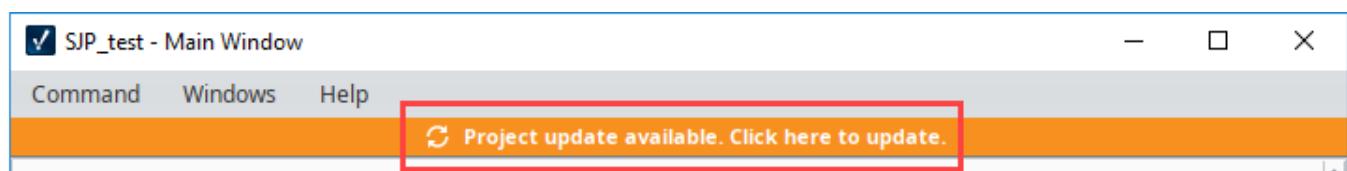


Update Project

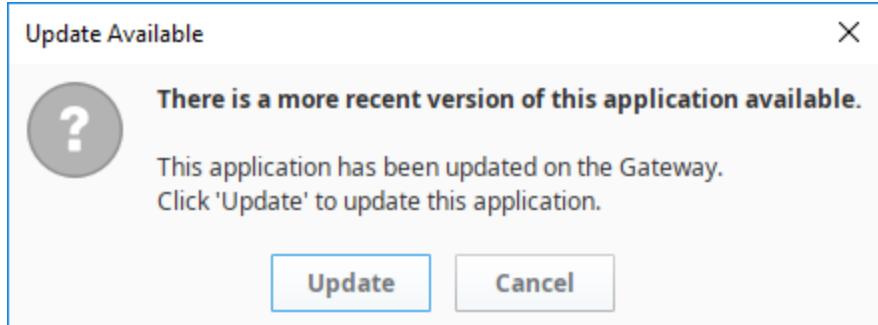
The **Update Project**  option pulls in any changes (saves by other people) that have happened since you opened the project or last updated it. If there are any conflicts due to [concurrent editing](#), the Designer [Resolve Conflicts](#) screen will open automatically. Once you resolve the conflicts, the Designer updates the project.

Project Updates in a Vision Client

Depending on how Clients are set up to receive project updates, operators may be notified with a banner stating that a project update is available. For more information, see [Client Update Modes](#).



When the operator clicks on the banner, a confirmation window appears. The operator can choose to update or cancel and keep working in the previous version of the client.



Project Updates in a Perspective Session

If you have a Perspective Session open and a change was made in the Designer that was saved and published, one of two things may happen. Either the project will silently update, or an Update Notification window will appear in the session. Your session will automatically update in 30 seconds or you can click **Update Now**. For more information, see [Session Project Updates](#).

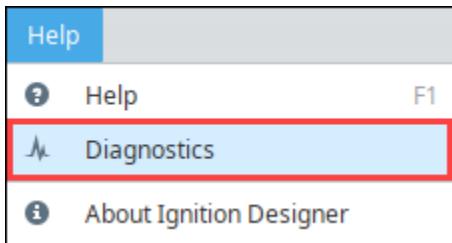
A screenshot showing an update notification in a Perspective Session. On the left is a green icon with a white downward arrow. To its right, the text 'Project Update • a few seconds ago ^' is displayed. Below that, a message reads 'This project has been changed. Please save your work, this session will automatically update in 22 seconds.' At the bottom of the notification area is a blue 'UPDATE NOW' button.

Related Topics ...

- [Vision Client Launcher](#)
- [Client Update Modes](#)
- [Session Project Updates](#)

Designer Diagnostics

The Help menu in the Designer has a **Diagnostics** option, which displays the Diagnostics window with tabs for six different troubleshooting features.



On this page ...

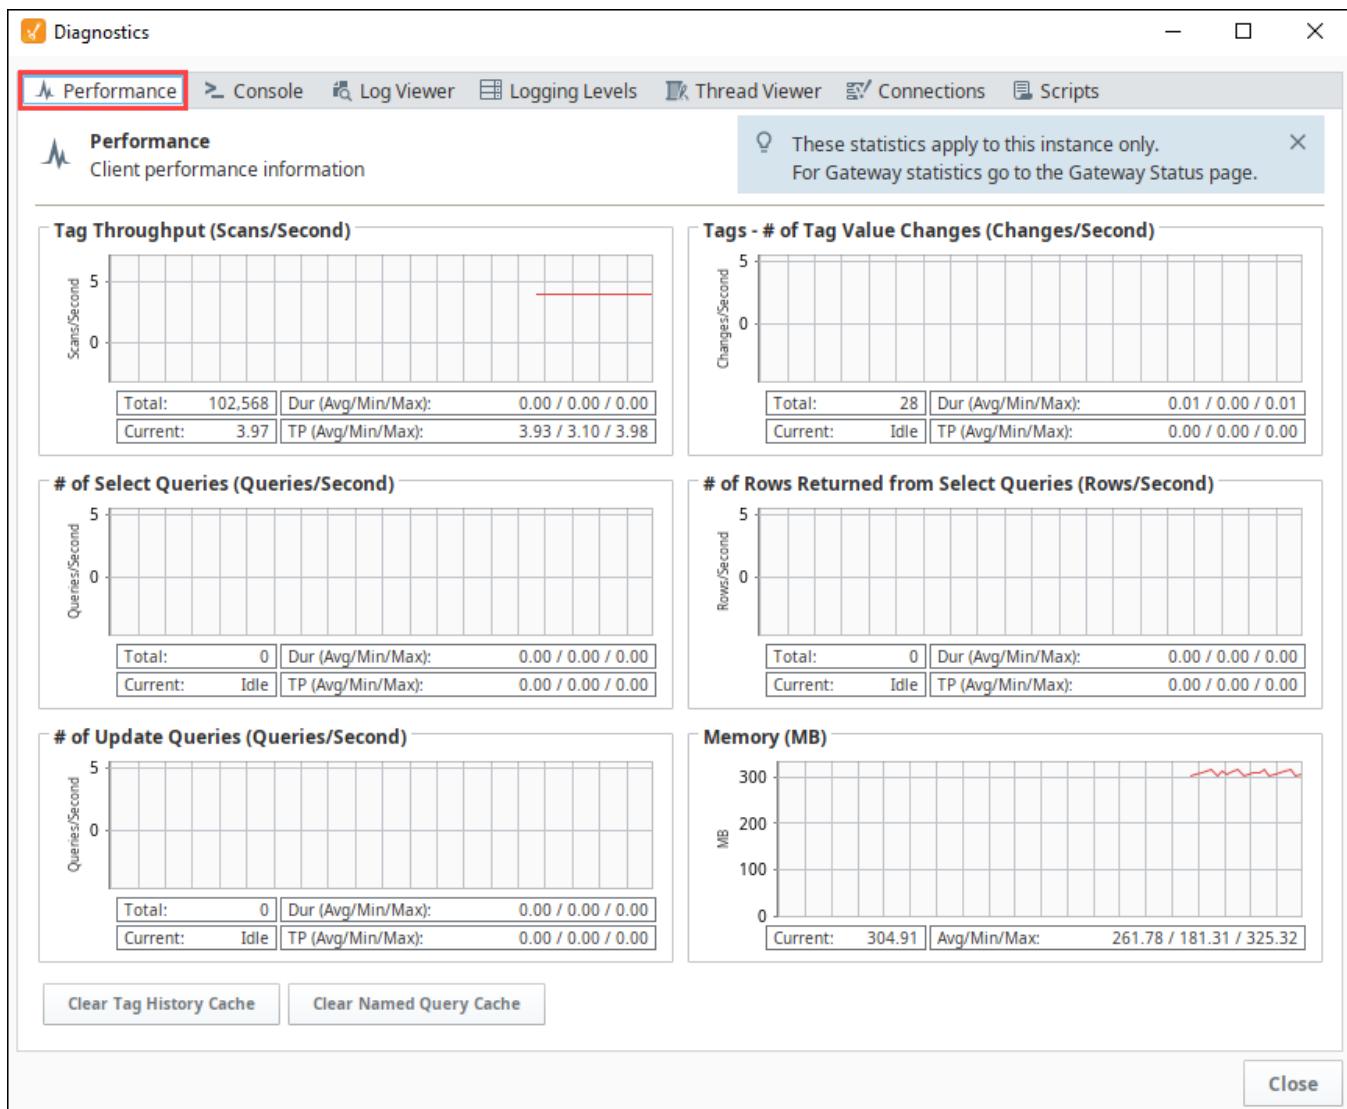
- [Performance Tab](#)
- [Console Tab](#)
- [Log Viewer](#)
- [Logging Levels](#)
- [Thread Viewer](#)
- [Connections](#)
- [Scripts](#)

Note: The information accessed through Designer Diagnostics is specific to the client runtime or the session. For Gateway statistics, see [Diagnostics - Logs](#).

Performance Tab

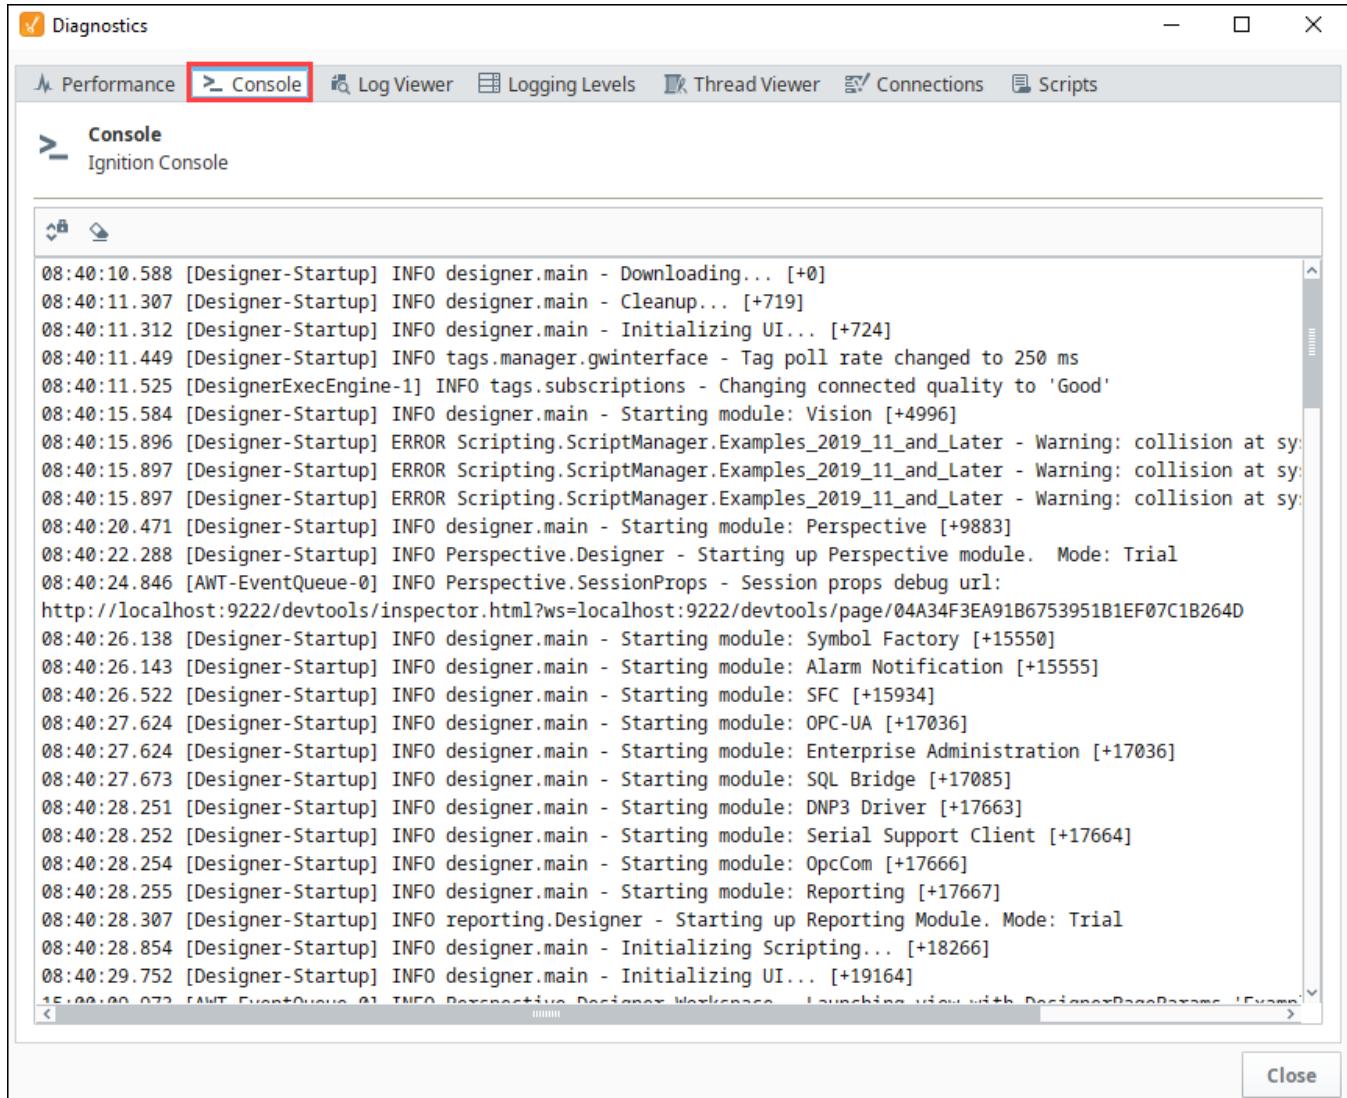
The Performance tab displays six realtime charts, each showing an aspect of the current client's performance. These charts can be very useful to help troubleshoot performance issues, especially slow queries.

- **Tag Throughput (Scans/Second)** - Displays the Tag throughput in scans per second.
- **Tags - # of Tag Value Changes (Changes/Second)** - Displays the number of Tag changes in changes per second.
- **Select Queries (Queries/Second)** - One of the most common causes of query slowdown is simply running too many queries too frequently, and the # of Select Queries (Queries/Second) chart can help identify when this is occurring.
- **Rows Returned from Select Queries (Rows/Second)** - Displays the number of rows returned from selected queries in rows per second.
- **Update Queries (Queries/Second)** - Displays the number of update queries in queries per second.
- **Memory (MB)** - Displays the client's memory usage in megabytes. This will almost always be a saw tooth pattern since memory is used, discarded, and re-acquired on a regular basis.



Console Tab

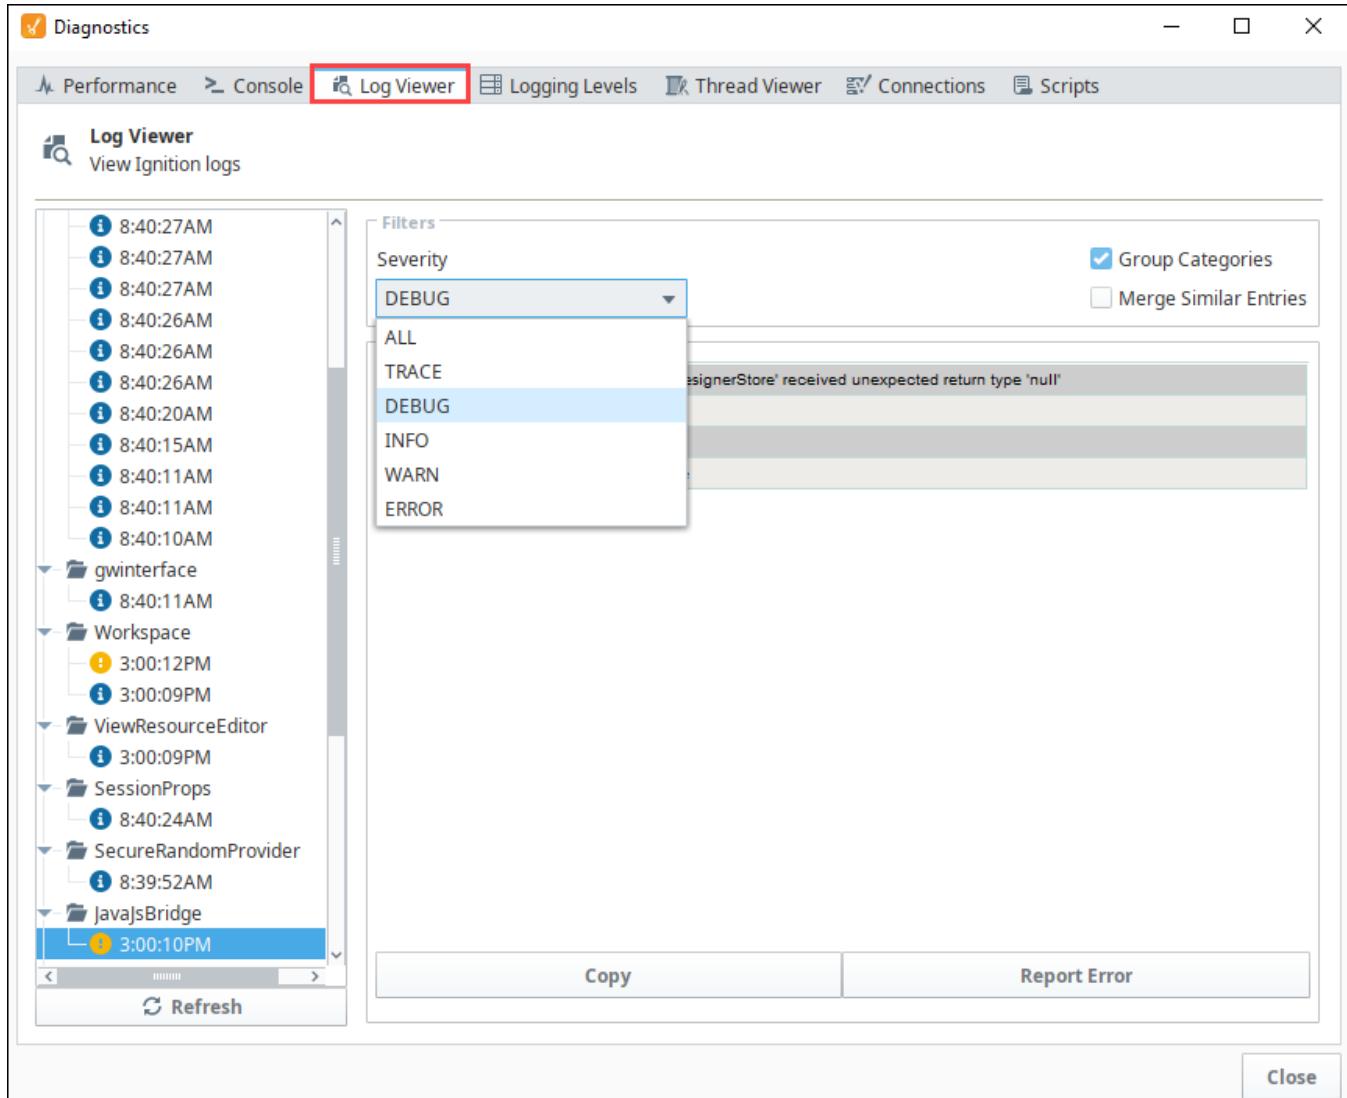
The Console tab displays the Ignition Console. This console displays messages that are generated by the entity that the console was launched from, such as the Designer. Any print statements or errors in your scripts will show up here (except Gateway scope scripts).



Log Viewer

The Log Viewer tab displays the logged events for the current entity, such as the Designer. Whenever messages occur in the console, they are logged and displayed in this tab. Each message has a logging level associated with it. This is a good place to go when troubleshooting an issue, as any errors shown here may illuminate the cause of the problem.

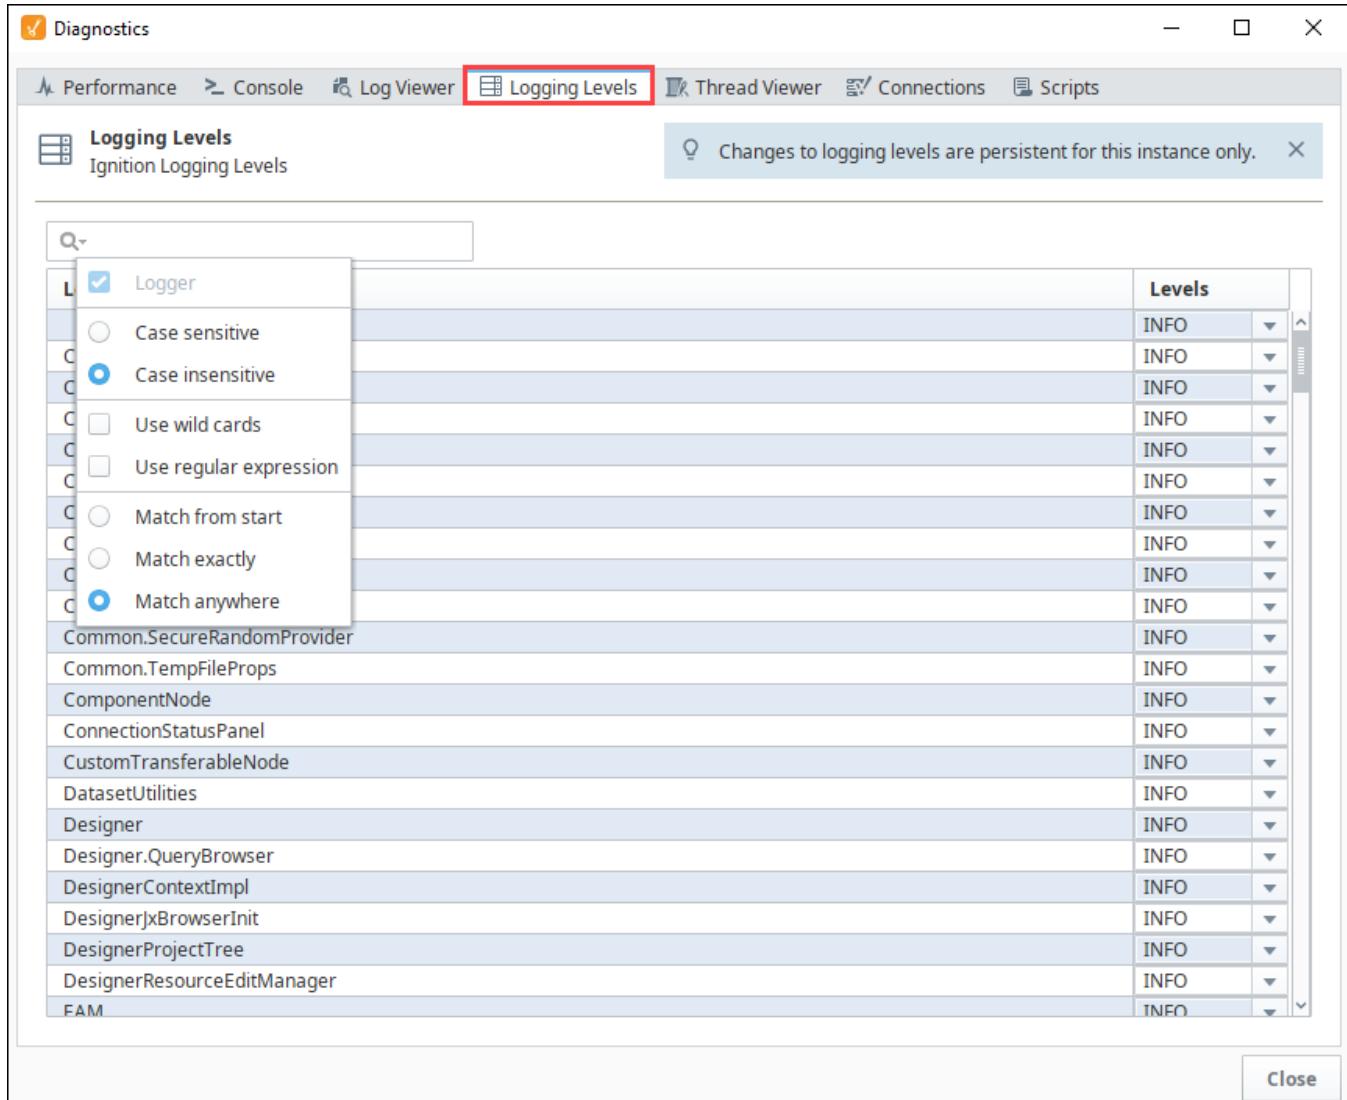
Logs can be filtered by severity by selecting an option in the Filters dropdown. To view entries across all categories chronologically, uncheck the **Group Categories** checkbox.



Logging Levels

The Logging Levels tab displays the list of internal loggers. Most users will not use this tab unless prompted by a technical support representative. Make sure to change logging levels back to info when you are done, otherwise they will flood your console and make it difficult to find any new problems.

A Search box in the upper left enables you to search the loggers. You can also set options to make the search case sensitive, use wildcards, use regular expressions, match from start, match exactly, or match anywhere.



Thread Viewer

The Thread Viewer tab displays information about the currently running threads. Each thread can be expanded by clicking the **Expand** icon or collapsed by clicking the **Collapse** icon. Most users will not use this tab unless prompted by a technical support representative.

The following feature is new in Ignition version **8.1.12**
[Click here](#) to check out the other new features

Now in the Designer and the Vision Client, there is an option to save a thread dump via the Thread Viewer tab. Additionally, automated thread dumps will be saved to the Client Launcher's directory in **.ignition/cache** when a UI thread takes longer than 2-5 seconds. Should the Designer or Vision Client recover, a message will appear in the console stating that a thread dump was taken, along with the full path where the file was saved.

Diagnostics

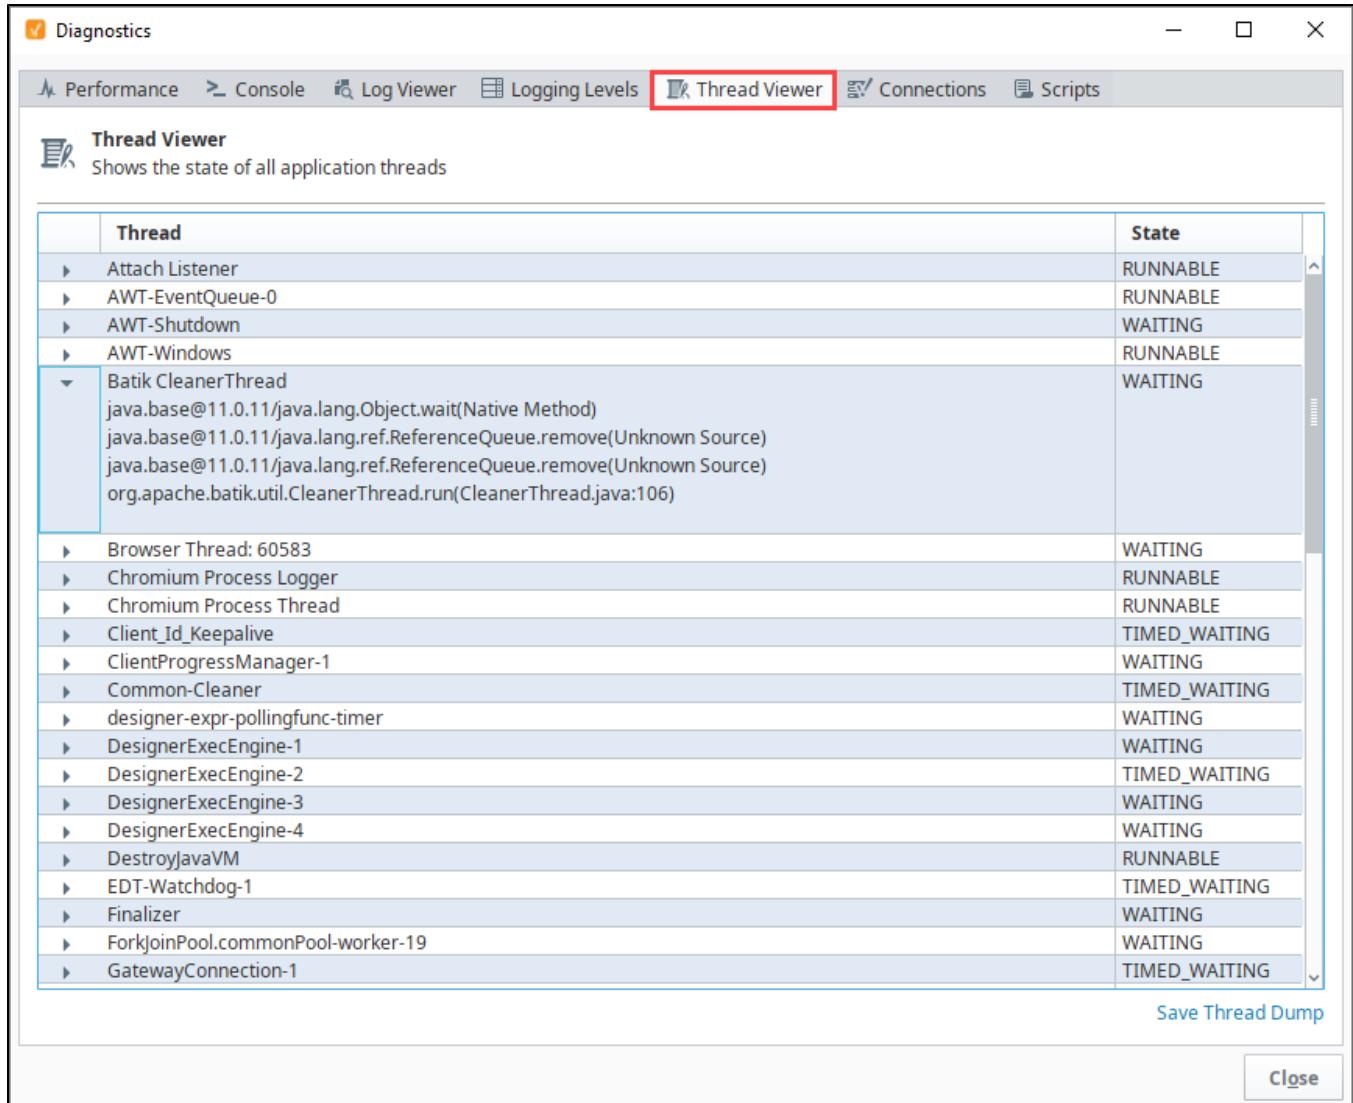
Performance Console Log Viewer Logging Levels Thread Viewer **Connections Scripts**

Thread Viewer
Shows the state of all application threads

Thread	State
Attach Listener	RUNNABLE
AWT-EventQueue-0	RUNNABLE
AWT-Shutdown	WAITING
AWT-Windows	RUNNABLE
Batik CleanerThread java.base@11.0.11/java.lang.Object.wait(Native Method) java.base@11.0.11/java.lang.ref.ReferenceQueue.remove(Unknown Source) java.base@11.0.11/java.lang.ref.ReferenceQueue.remove(Unknown Source) org.apache.batik.util.CleanerThread.run(CleanerThread.java:106)	WAITING
Browser Thread: 60583	WAITING
Chromium Process Logger	RUNNABLE
Chromium Process Thread	RUNNABLE
Client_Id_Keepalive	TIMED_WAITING
ClientProgressManager-1	WAITING
Common-Cleaner	TIMED_WAITING
designer-expr-pollingfunc-timer	WAITING
DesignerExecEngine-1	WAITING
DesignerExecEngine-2	TIMED_WAITING
DesignerExecEngine-3	WAITING
DesignerExecEngine-4	WAITING
DestroyJavaVM	RUNNABLE
EDT-Watchdog-1	TIMED_WAITING
Finalizer	WAITING
ForkJoinPool.commonPool-worker-19	WAITING
GatewayConnection-1	TIMED_WAITING

Save Thread Dump

Close



Connections

The Connections tab displays the Gateway connection status as well as a realtime chart of the Gateway ping time and a summary with the current ping time and average, minimum, and maximum ping times.

The screenshot shows the Diagnostics interface with the 'Connections' tab selected. A red box highlights the 'Connections' tab in the top navigation bar. Below it, the 'Gateway Connection Status' section shows a green checkmark and the connected URL: <http://10.10.110.54:8088>. The main area displays a chart titled 'Gateway Ping Time' with 'ms' on the y-axis. The chart shows a single data series with two small peaks at approximately 0.14 ms and 381.00 ms. Below the chart, a table provides summary statistics: Current: 0.00, Avg/Min/Max: 0.14 / 0.00 / 381.00.

Scripts

Shows running scripts. Use the **Delete** icon to terminate a script.

Tags

What Is a Tag?

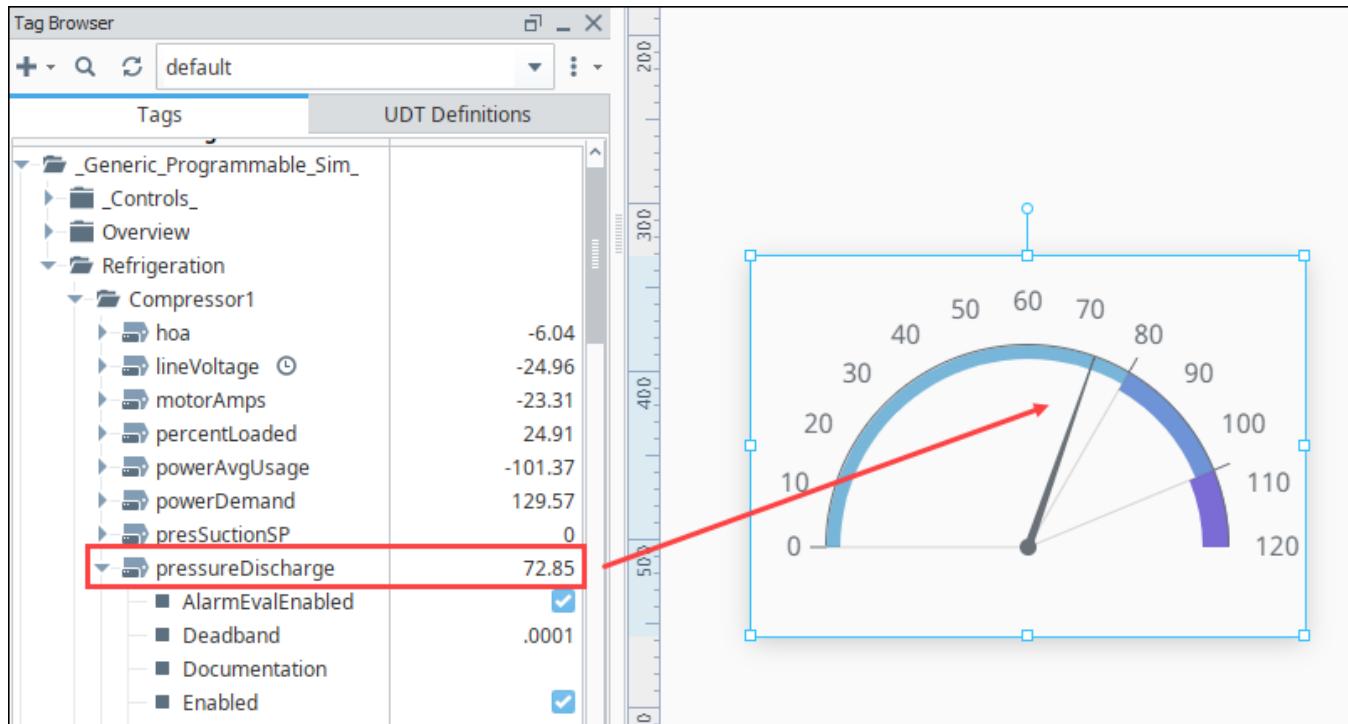
Tags are points of data and may have static values or dynamic values that come from an OPC address, an expression, or a SQL query. The values can be used on screens and in transaction groups.

Tags provide a consistent data model throughout Ignition, and offer the easiest way to get up and running creating realtime status and control systems. Despite their fast initial learning curve, however, Tags offer a great amount of power in system design and configuration. The ability to aggregate Tags from a variety of installations means that you can build widely distributed SCADA systems more easily than ever before with a high level of performance and relatively easy configuration.

While the goal of Tags in Ignition is to create an easy yet powerful model, the variety of options and terminology can sometimes make configuration confusing. Tags are created and controlled using both the Gateway and the Designer for configuration.

- In the Designer, you create or import the Tags. There are several types of Tags such as an OPC Tags and Memory Tags. Each Tag has many properties and other functionality such as alarming, history, etc. Once your Tags are created, you can use them in your windows, views, and reports.
- In the Gateway, you create and modify Tag Providers. You can create these Realtime Providers to store groupings of Tags for use in your projects either locally in Ignition or share them externally with connected Gateways. There are also Historian Providers used to store historical data for the Tags, but these are automatically created for each datasource you have. These Tag Provider configurations in the Gateway apply globally to all your projects.

The following example shows a pressureDischarge Tag in the Tag Browser and a Gauge component in the Designer. The value on the Designer component is bound to the Tag and updates in realtime. This is just a simple example of how Tag values can be represented in your SCADA designs.



Tag Providers

There are two types of Tag providers; Internal and Remote. By default, a fresh Ignition installation will have an internal Tag provider. This can be thought of as a standard internal Tag database, and stored in the Ignition Gateway. Additionally, it is possible to create Remote [Tag Providers](#), linking one installation of Ignition to the Tags on another Ignition. This ability opens up some very flexible architectures.

Tag User Defined Types

Tag [User Defined Types](#) (UDTs) provide an object-oriented approach to Tag building, allowing you to define parameterized data types, extend and override types, and then rapidly generate instances. A change to the type definition is then inherited by all instances, drastically saving time when making routine changes. The UDTs are fully supported by Vision templates, which means you can configure templates for your custom data types and take advantage of drag-and-drop binding to rapidly build complex screens.

On this page ...

- [What Is a Tag?](#)
 - [Tag Providers](#)
 - [Tag User Defined Types](#)
- [Tag Features](#)
- [Importing and Exporting Tags](#)
- [Tag Naming](#)

Tag Features

Tags work naturally and easily with Ignition to offer the following features:

- **Performance and Scalability**

Tags offer great performance on the Gateway, in Perspective Sessions, and in the Vision Client. On the Gateway, the system can support many thousands of value changes per second and millions of Tags. In runtime, Tags improve efficiency with their lightweight subscription architecture. Adding additional Clients creates a nearly negligible effect on the database and the Gateway performance.

- **Object-Oriented Design**

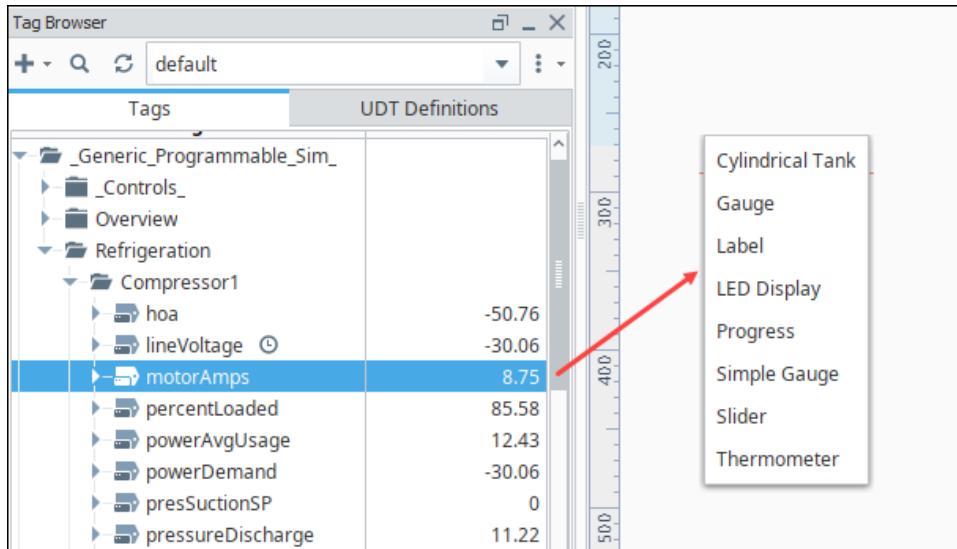
Use [Tag UDTs](#) (User Defined Types) to design re-usable, parameterized, and extendable data types. You can create and configure new instance Tags in seconds, saving a great amount of time over traditional Tag systems.

- **Powerful Alarming Model**

Each Tag can have any number of [alarms](#) configured on it. There are many different [alarm modes](#) accommodating simple digital alarms, analog high/low value alarms, as well as more specialty alarms like bad data quality and bit-packed alarms. The settings for alarms can bound to other Tags, making the alarm configuration dynamic.

- **Drag-and-Drop Screen Design**

You can drag and drop Tags onto a window or view to automatically create new bound components. Drag Tags onto existing components or properties to quickly bind them to the data.

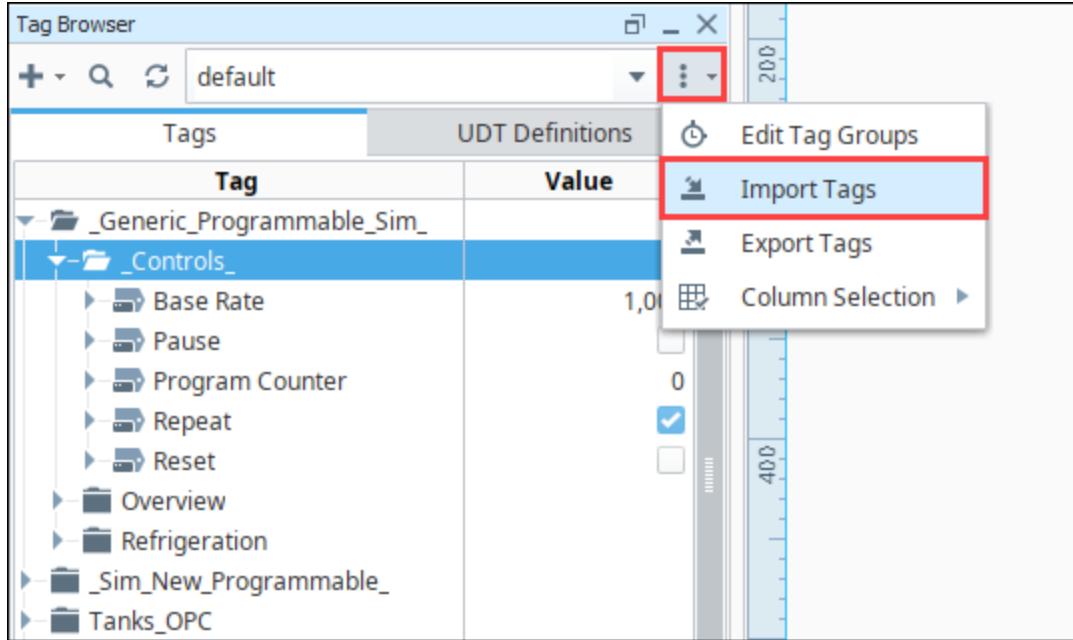


- **Historical Logging**

The [Tag Historian Module](#) makes it easier than ever to store and use [historical data](#). When you simply select a check box on a Tag, historical data is stored in an efficient format in your SQL database. This data is then available for querying through scripting, historical bindings, and reporting. Also, you can drag-and-drop Tags directly onto many components to create trends or display historical values. Tags Historian's robust querying provides you great flexibility in how you retrieve the data.

Importing and Exporting Tags

Ignition Tags can easily be imported and exported from the Designer by selecting the More Options menu, then either the Import Tags or Export Tags icon. You'll be prompted to choose the Tags or folders that you want. See the [Exporting and Importing Tags](#) page for more information.



Tag Naming

Tags names are flexible and do not have to match data source names (like an OPC path) or tag codes (such as N7, F8, etc.). It is not necessary that Tag's name be related at all to its underlying data source (OPC path, for instance). This provides a level of indirection that is convenient for systems with many repeat Tag structures.

It is important to give Tags a meaningful structure and arrange them in hierarchical Tag folders so that they are easy to understand, identify, and locate for all developers. By default, Ignition Tags are named after their OPC Server address when a Tag is dragged into the Tag Browser. You can change this name to just about anything that you want. We recommend using names that mean something to your process, such as "Motor 3 Amps." Alternatively you could create folders in your Tag Browser such as "Motor 3/Amps.". When renaming Tags and folders, there is really only one question to ask: "does this structure make sense?"

Another important concept to consider when naming and organizing your Tags, is to do this early in your project. If you rename or move any of your Tags to another folder, and your Tag is being used in other places, chances are you are going to break the reference to the Tag on your screen. So keeping your Tags organized and defining your Tag structure early on in your project is critical.

When you choose a new name for your Tags and folders, there are some rules that must be followed. The first character of the Tag name must be one of the following:

- Letter - specifically, a letter as recognized by Unicode's Letter (L) category.
- Number
- Underscore

The second character, and every character after that can then be one of the following:

- Letter
- Number
- Underscore
- Space
- Any of the following special characters:

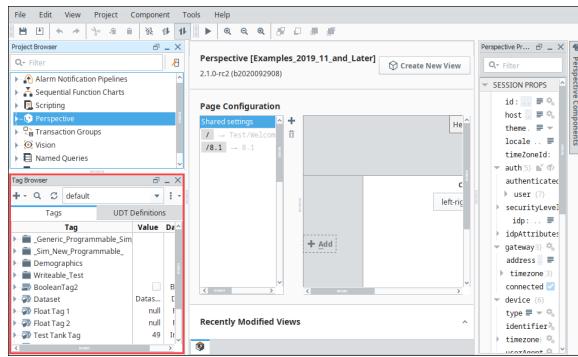


All other special characters are not allowed in a Tag name.

Tag Browser

The Tag Browser is the central location for interaction with all types of Tags on your system. It gives you full view of the Tags including the current value, datatype, and any traits. When panels are in their default configuration in the Designer, the Tag Browser appears on the left side.

As of release 8.1, the Tag Browser has been updated with a new design for ease of use. Tags and [UDT Definitions](#) now have their own tabs. The Tag Provider Selector enables you to view Tags for a specific Tag provider. There is a new interface for [Creating Tags](#). Many other improvements have been integrated, including icons for each [Tag type](#).



On this page ...

- [Tag Browser Tree](#)
- [Tag Browser Toolbar](#)
 - [Add Tag](#)
 - [Browse Devices](#)
 - [Find/Replace](#)
 - [Refresh Providers](#)
 - [Tag Provider Selector](#)
- [More Options Menu](#)
 - [Tag Groups](#)
 - [Import/Export](#)
 - [Column Selector](#)
- [Right-Click Menu](#)
- [Tag Traits](#)



INDUCTIVE
UNIVERSITY

Tags in Ignition

[Watch the Video](#)

Tag Browser Tree

The Tag Browser is set up in an interactive tree structure with folders that can be expanded or collapsed to view more Tags.

Click on the **Expand** ➤ icon to expand any folder or the **Collapse** ▾ icon to collapse the folder. In the example below, the pH Tag for Tower2 was expanded.

Tag Browser

default

Tags UDT Definitions

Tower 2

pH

Property	Value
AlarmEvalEnabled	<input type="checkbox"/>
Deadband	.0001
Documentation	
Enabled	<input checked="" type="checkbox"/>
EngHigh	100
EngLow	0
EngUnit	
FormatString	#,##0.##
HistoryEnabled	<input type="checkbox"/>
Name	pH
Quality	Good
ReadOnly	<input type="checkbox"/>
TagGroup	
Timestamp	
Tooltip	Default
value	98
ValueSource	memory

The following feature is new in Ignition version **8.1.18**
[Click here](#) to check out the other new features

If you create a tooltip on a Tag by filling in the [Tooltip property](#), hovering over the base Tag at the top of the Tag Browser Tree will display the tooltip.

Tag Browser

The screenshot shows the Tag Browser window with the 'Tags' tab selected. A tooltip is displayed over the value '1' of the tag 'Look at me I'm a Tag'. The tooltip contains the text 'Memory Tag: [default]Look at me I'm a Tag' and 'Hi I'm a tooltip'. The 'UDT Definitions' tab is also visible.

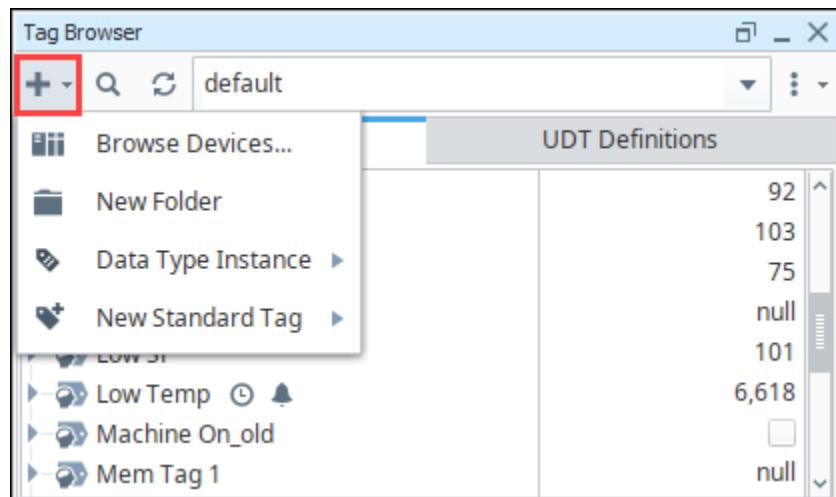
Tag	Value	Data Type
Look at me I'm a Tag	1	Integer
Canwrite		String
Deadband	0	Double
Documentation		String
Enabled	<input checked="" type="checkbox"/>	Boolean
EngHigh	100	Double
EngLow	0	Double
EngUnit		String
FormatString	#,##0.##	String
HistoryEnabled	<input type="checkbox"/>	Boolean
Name	Look at m...	String
Quality	Good	String
ReadOnly	<input type="checkbox"/>	Boolean
TagGroup	Default	String
Timestamp	2022-06-0...	DateTi...

Tag Browser Toolbar

The Tag Browser toolbar has several options for working with Tags.

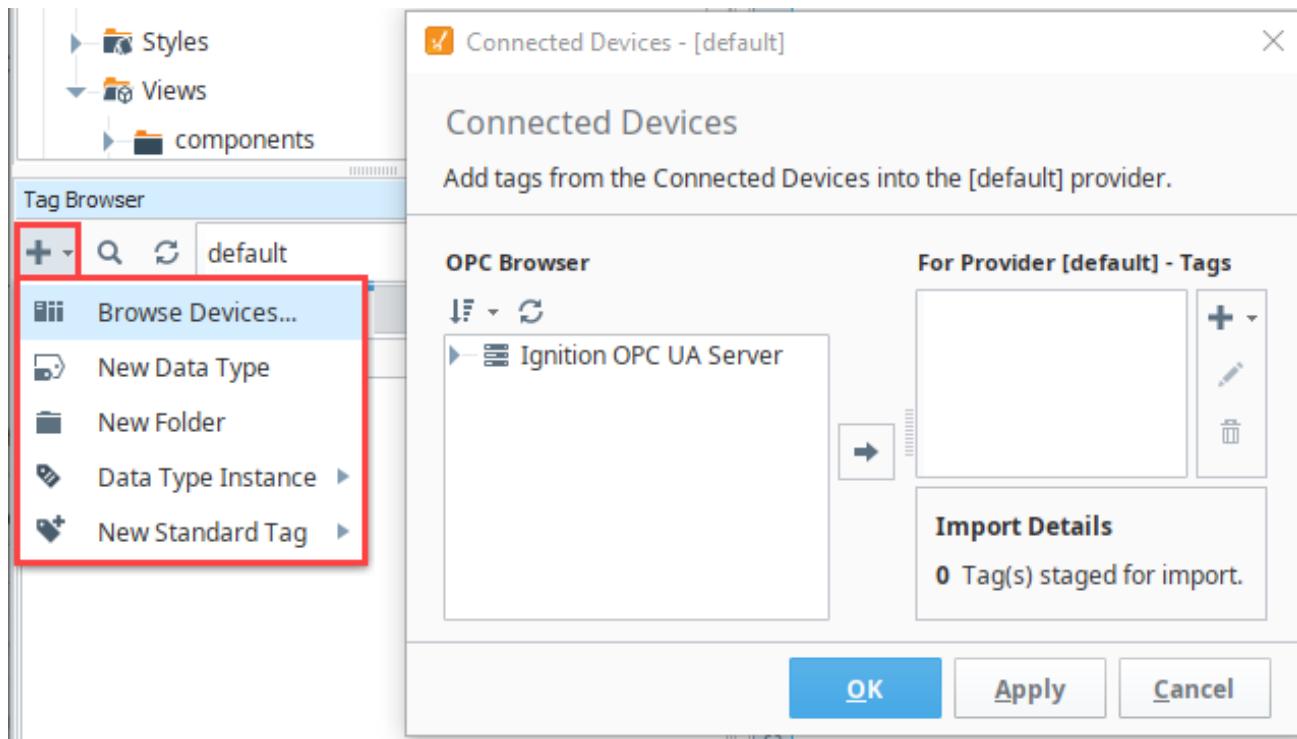
Add Tag

The **Add** icon opens a context menu showing all the options to Browse Devices, add a Tag, Folder, UDT Instance or a UDT Definition. The new object is added under the Folder you selected, or as a sibling to the Tag you selected. This button is disabled if there is no selection.



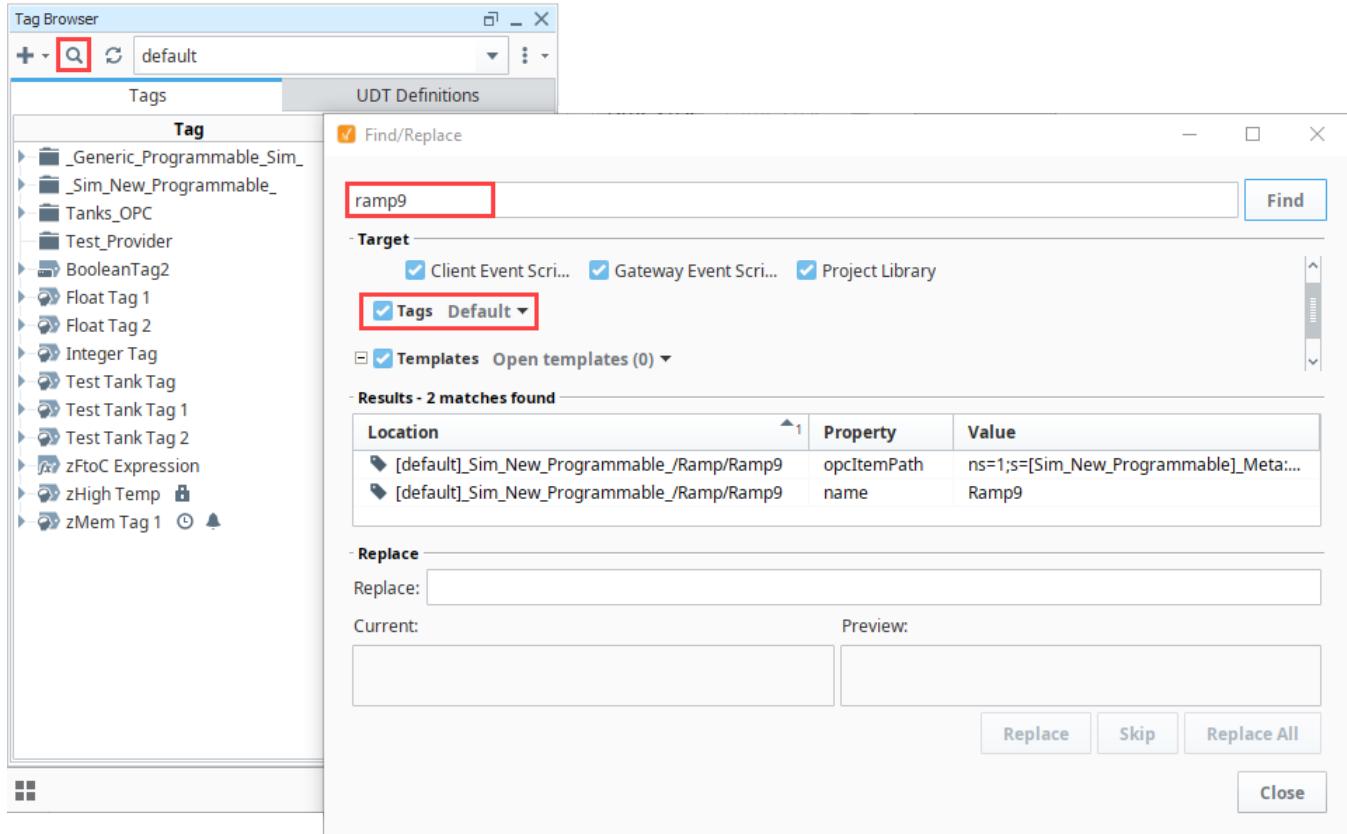
Browse Devices

With the Browse Devices, you can browse to find external PLC or OPC Tags. Click the Add  icon at the top of the Tag Browser to open the Connected Devices window. You can then select Tags and move them to the Tag Browser to be used in the Ignition system. For complete information, see [Creating Tags](#).



Find/Replace

Clicking on the **Search**  icon in the Tag Browser will open up the Designer's global **Find/Replace** screen. In the example below, we searched for the Ramp9 Tag and limited the search to the default Tags. Results are shown at the bottom of the screen. For additional information, see [Find and Replace](#).



Refresh Providers

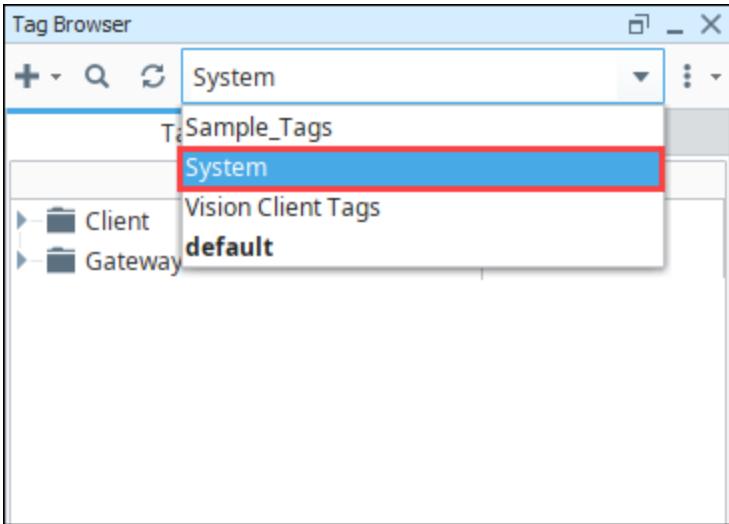
The Refresh Providers  icon refreshes all of the Providers in the Tag Browser. This is useful if you or others have modified Tags and do not see an update. In general, this button is not used very often.

Tag Provider Selector

The Tag Provider Selector is a dropdown menu with a list of available Tag Providers. Use this selector to view the System Tags, Vision Client Tags, default Tags, or other [Tag Providers](#) you have in your project.

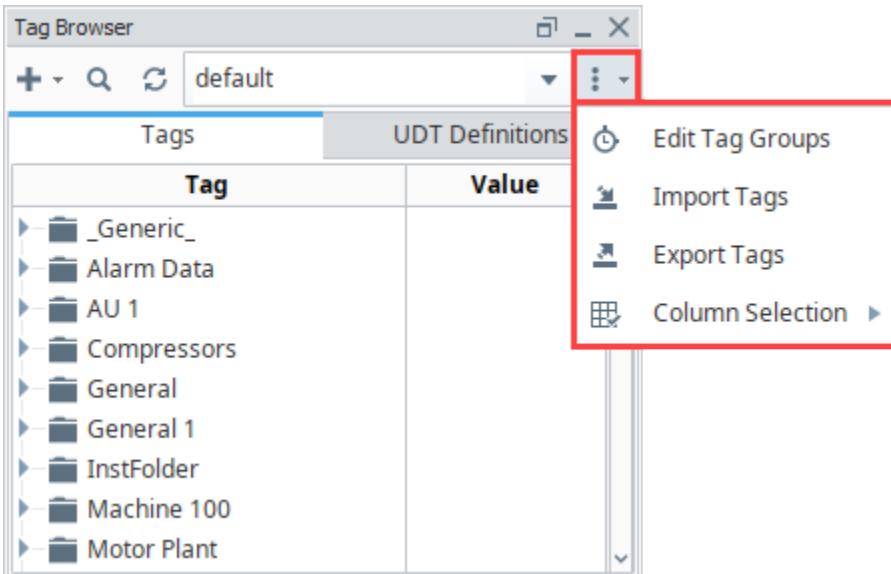
On a new install of Ignition, a single "default" Tag Provider is created for you, but there are no Tags added to the "default" provider initially. You'll notice the "default" provider is in bold. The bold entry in the dropdown list is determined by the Default Provider for the project. You can change the default provider in **Project Properties > Project > General > Tag Settings > Default Provider**. Save and restart the Designer, and the next time you open the Tag Provider in the dropdown of the Tag Browser, a different Default Provider will be displayed in bold.

When a provider is selected to System, the Browsing Devices under the Add  icon is disabled.



More Options Menu

A More options icon on the upper right of the Tag Browser, opens a dropdown menu of additional options.



Tag Groups

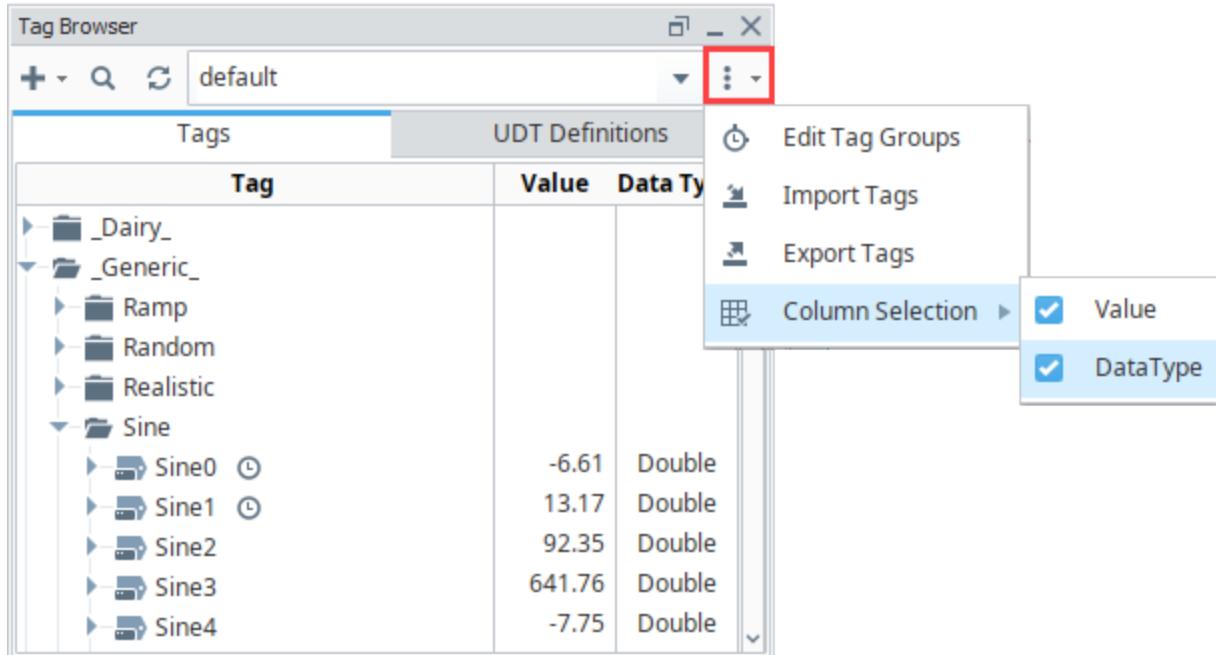
In the Tag Browser, the **Edit Tag Group** opens the **Tag Group Editor** window. Tag Groups dictate the rate of execution of Tags. This is where you set up your Tag Groups and scan rates. See [Tag Groups](#) for more information.

Import/Export

Ignition can export and import Tag configurations to and from the **JSON** (JavaScript Object Notation) file format. Use the **Import** icon or **Export** icon to import and export Tags in this Gateway. See [Exporting and Importing Tags](#) for more information.

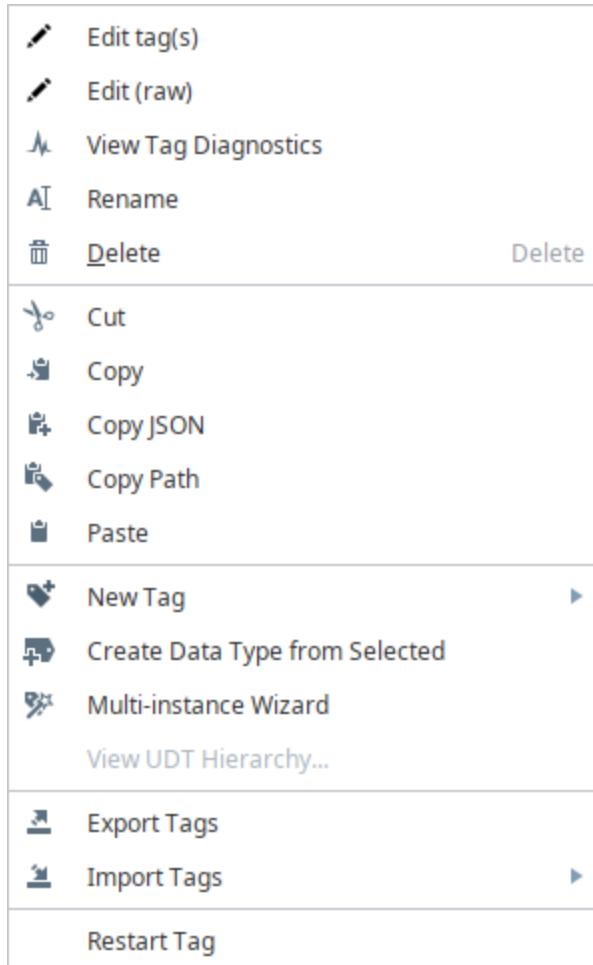
Column Selector

The Tag Browser displays the **Value** and **Data Type**. The Value type is set by default. To toggle any of the options, click on the **Column Selector** icon, then click the checkbox. In the example below, the Tag values are shown next to the Tag names.



Right-Click Menu

Editing Tags is done mostly through the Tag Browser. The **Tag Browser** allows you to right click on a Tag or folder to perform any of the following functions. Different objects will have different options available. The special **Data Types** folder is slightly different than a regular folder and will have even fewer options.

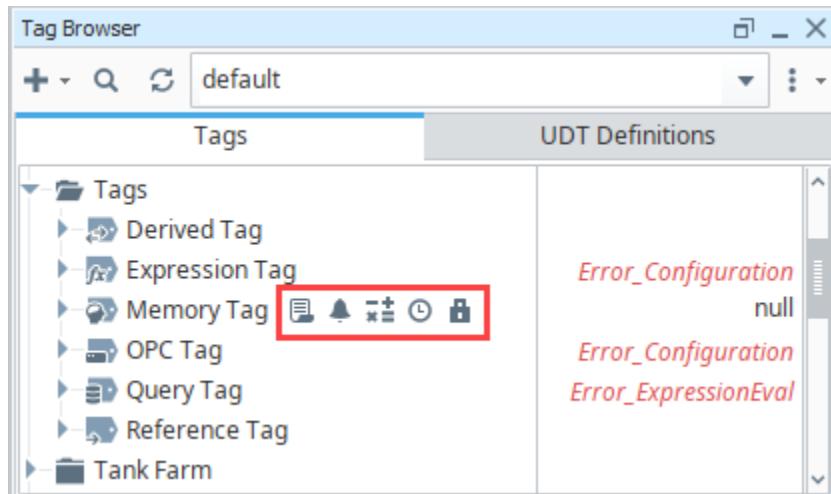


Function	Description
Edit Tag	<i>Disabled when a Folder is selected.</i> Opens the Tag Editor window so the Tag can be edited.
Edit (raw)	<i>Disabled when a Folder is selected.</i> Opens a JSON editor, allowing you to view and edit the underlying JSON that makes up the Tag.
View Tag Diagnostics	The following feature is new in Ignition version 8.1.34 Click here to check out the other new features Opens the Tag Diagnostics window, which also accesses Active Subscriptions and Reference Log data.
Rename	Renames the current selection.
Delete	Deletes the current selection.
Cut	Cuts the current selection into the clipboard.
Copy	Copies the current selection into the clipboard.
Copy JSON	Available for non-client tags only. Copies the JSON for the selected Tags into the system clipboard. In addition, pasting the JSON into a different provider/designer will create or overwrite Tags.
Copy Path	Copies the currently selected Tag path into the clipboard.
Paste	Pastes the content in the clipboard into the selected context.

New Tag	<p><i>Disabled when a Tag is selected.</i></p> <p>For Folders, this option opens a sub-menu to create a Tag or Tags.</p> <table border="1"> <thead> <tr> <th>Function</th><th>Description</th></tr> </thead> <tbody> <tr> <td>New Folder</td><td>Creates a new Tag folder.</td></tr> <tr> <td>Data Type Instance</td><td>Creates a new instance of an existing data type. The instance is linked to the parent type so when the parent changes, the instances are overwritten with the parent type changes. Sub Menu - based on Data Types</td></tr> <tr> <td>New Standard Tag</td><td>Creates different types of Tags such as Derived, Expression, Memory, OPC, Query, and Reference Tags. Sub Menu - Standard Tag Types</td></tr> </tbody> </table>	Function	Description	New Folder	Creates a new Tag folder.	Data Type Instance	Creates a new instance of an existing data type. The instance is linked to the parent type so when the parent changes, the instances are overwritten with the parent type changes. Sub Menu - based on Data Types	New Standard Tag	Creates different types of Tags such as Derived, Expression, Memory, OPC, Query, and Reference Tags. Sub Menu - Standard Tag Types
Function	Description								
New Folder	Creates a new Tag folder.								
Data Type Instance	Creates a new instance of an existing data type. The instance is linked to the parent type so when the parent changes, the instances are overwritten with the parent type changes. Sub Menu - based on Data Types								
New Standard Tag	Creates different types of Tags such as Derived, Expression, Memory, OPC, Query, and Reference Tags. Sub Menu - Standard Tag Types								
Multi-instance Wizard	Creates many instances of a UDT at the same time.								
Export Tags	Exports the selected Tags.								
Import Tags	Imports Tags into the project.								
EAM	Sends Tag to EAM Agents.								
Restart Tag	<p>Attempts to restart the selected Tag.</p> <p>"Restarting" a Tag effectively refreshes the following traits:</p> <ul style="list-style-type: none"> • Value generation (Expressions for expression tags, queries for Query Tags, OPC value for OPC tags, etc). • Scaling • Engineering Limits • Alarms • Deadbands • Event Scripts <p>If a folder is restarted, then all tags under the folder will restart.</p>								

Tag Traits

Certain settings or Tag configurations are visually represented next to the Tag in the Tag Browser.



The following icons enable you to note some important settings on the Tag at a glance. A description of the icons are listed below.

Icon	Setting	Description
	Scaling	The Scale Mode property under the Numeric Tag Properties section of the Tag Editor has been set to a value other than "Off." The value on the Tag will be scaled to some degree.

	Alarming	At least one alarm has been configured on this Tag.
	Tag History	This Tag has been configured to log data into the Tag Historian system.
	Tag Event Script	At least one Tag Event Script has been enabled on this Tag.
	Lock	Shows the Tag has permissions enabled.
	Inheritance	Denotes inheritance. Displays the Parent Type, name of the UDT and instance name (i.e., Parent Types: Motor UDT > Complex Motor). Refer to UDT Inheritance Traits .
	Override	Denotes that the property in the UDT instance overrides the parent property. Refer to Overriding Properties of the Parent UDT .

Types of Tags

There are many different types of Tags in Ignition: standard Tags, [System Tags](#), and [Vision Client Tags](#). All these Tag types are available in the Tag Browser.

Tags executed in the Gateway support all of the primary features of Tags: scaling, alarming, history, and role-based permissions. These Tags run in the Gateway, and the values of the Tags are shared among all running sessions and clients. They are identical in their configurations, apart from defining how the value is generated. As of release 8.1, each Tag Type has its own icon in the Tag browser.

Icon	Tag type
	Memory Tag
	OPC Tag
	Expression Tag
	Query Tag
	Reference Tag
	Derived Tag
	System Client Tags, Vision Client Tags

On this page ...

- [OPC Tags](#)
- [Memory Tags](#)
- [Expression Tags](#)
- [Query Tags](#)
- [Reference Tags](#)
- [Derived Tags](#)
 - [Changing the Source](#)
- [User Defined Types \(UDTs\)](#)
- [System Tags](#)
- [Vision Client Tags](#)

OPC Tags

An OPC  Tag is driven by an OPC Item Path and OPC server. The OPC Item Path is a string path to a particular device connection. The exact path is defined by the driver and OPC server used to communicate with the device. Many drivers support browsing, allowing you to automatically create OPC Tags by dragging-and-dropping from the OPC Browser. However, in cases where browsing isn't supported, OPC Tags can manually be created. In the

In the **Tag Browser**, double click on any existing OPC Tag, to see the the **OPC Server** name and **Item Path**.

Tag Editor

motorAmps

default

Properties

Name	motorAmps
Tag Group	Default
Enabled	true
Value Source	OPC
Data Type	Double
OPC Server	Ignition OPC UA Server
OPC Item Path	ns=1;s=[Generic_Programmable]

Details

- Documentation
- Diagnostics

OK **Apply** **Cancel**

Memory Tags

 Tags are simple Tags, that do not automatically poll or update their value. They hold the same value until some other user-created mechanism (most likely a script or binding) changes their value. They're useful in situations where a value must be stored outside of a PLC or database.

The state of memory tags are stored inside the gateway's internal database. This allows the tag and its value to be retained across gateway restarts.

Tag Editor

Events

default

Properties

Name	Events
Tag Group	Default
Enabled	true
Value Source	Memory
Data Type	Integer
Value	35

Details

- Documentation
- Diagnostics

OK **Apply** **Cancel**

 INDUCTIVE
UNIVERSITY

Memory Tags

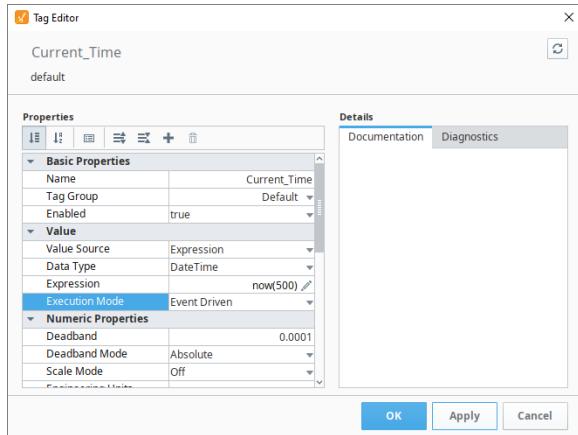
[Watch the Video](#)

Expression Tags

Expression  Tags are driven by an [expression](#), allowing their values to be determined from a calculation.

The **Expression** property on Expression Tags determines their value. The expression can reference values and properties on other Gateway scoped Tag values. However, due to scoping, they can not reference property values on Vision Client and Perspective Session components.

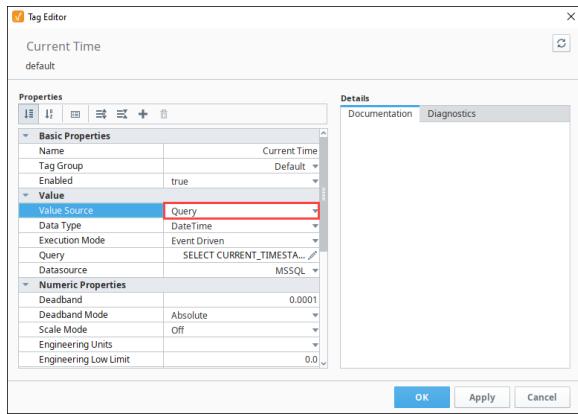
The expression on an Expression Tag executes based off of the **Execution Mode**. More information on Execution Mode can be found on the [Tag Properties](#) page.



Note: It may be helpful when troubleshooting or testing Expression tags to increase the default threadpool count. Refer to the [Gateway Configuration File Reference - Threadpool Counts](#) for more information.

Query Tags

A Query  Tag executes a SQL Query; the result of that query is returned to the value on the Tag. Query Tags can reference other Gateway-scoped Tags to build dynamic queries. The **Query** property dictates the query that will execute, and the **Execution Mode** determines how often the query will run. Furthermore, the **Datasource** property determines which database connection the query will execute against.



Expression Tags

[Watch the Video](#)



Query Tags

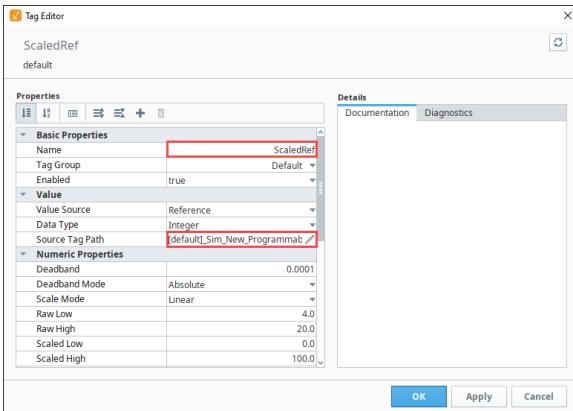
[Watch the Video](#)



Note: It may be helpful when troubleshooting or testing Query tags to increase the default threadpool count. Refer to the [Gateway Configuration File Reference - Threadpool Counts](#) for more information.

Reference Tags

A Reference  Tag simply refers to an existing Tag, using the **Source Tag Path** property to determine which other Tag to reference. Writes targeting the Reference Tag will send a write request to the source Tag.



Reference Tags

[Watch the Video](#)

Derived Tags

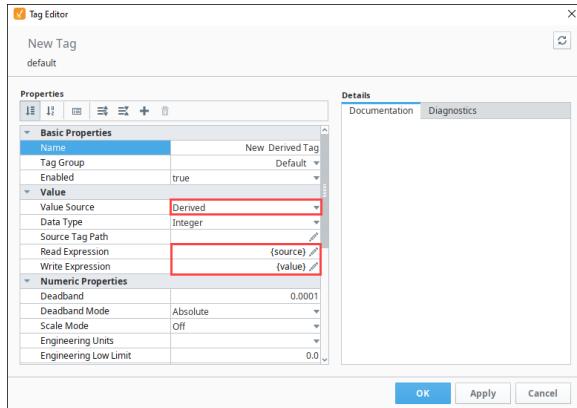
A Derived  Tag is an abstracted Tag that refers to another Tag. They are similar conceptually to Reference Tags in that their value is dependent on the **Source Tag Path** property, but Derived Tags have some additional functionality. Namely, they can apply expressions to the referenced value with the **Read Expression** property, allowing the value on the Derived Tag to differ from the source Tag.

In addition, The **Write Expression** property will apply an expression to the value of any write requests targeting the Derived Tag, allowing the expression to modify the value of the incoming write before it's applied back to the source Tag.



Derived Tags

[Watch the Video](#)



Read Expression	Determines what value should appear on the Derived Tag. The current value of the source Tag may be referenced with the {source} reference.
Write Expression	When writing to the Derived Tag, this expression determines what value should be written to the source Tag. The current value of the source Tag may be referenced with the {source} reference. The raw value passed to the Derived Tag may be referenced with the {value} reference.

This interface provides an opportunity to scale the written and read value. For example, if the source Tag was in Fahrenheit, and the derived Tag should be Celsius, we could use the following expressions:

```
//Read Expression
({source}-32)*(5/9)

//Write Expression
{value}*(9/5) + 32
```

Changing the Source

Aside from editing the Tag from the Designer, the source Tag path on a Derived Tag may be modified by writing directly to the **SourceTagPath** property on the Tag, or using a Tag Binding on a Vision component.

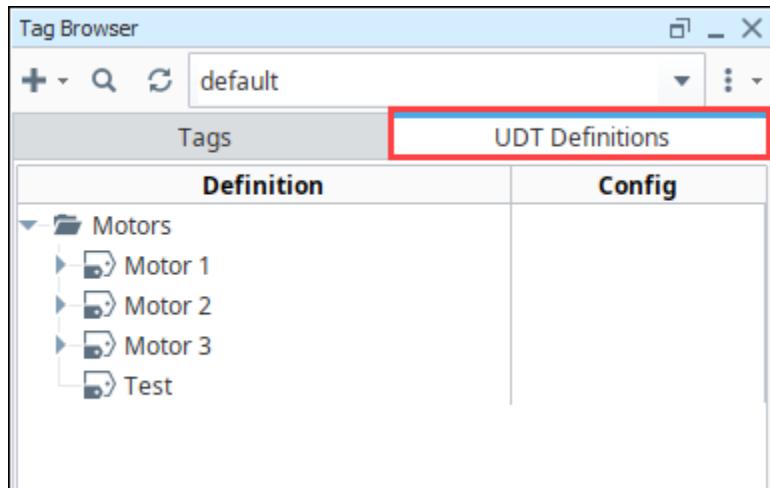
Additionally, the **SourceTagPath** property may be changed through scripting:

```
#Example of writing to the SourceTagPath attribute via Python Script
system.tag.writeBlocking(["Derived Example/Derived Tag.SourceTagPath"], [".]Folder/New Source Tag"])
```

User Defined Types (UDTs)

UDTs are created out of standard Tag types, but they offer a variety of additional features. You can think of them as a way to create "data templates", where a particular structure of Tags is defined, and can then be created as if it were a single Tag. This UDT example shows two Motor instances, the data type Motor, and all the Parameters and Tags that make up the structure (i.e., Amps and HOA). For more information, see [User Defined Types -](#)

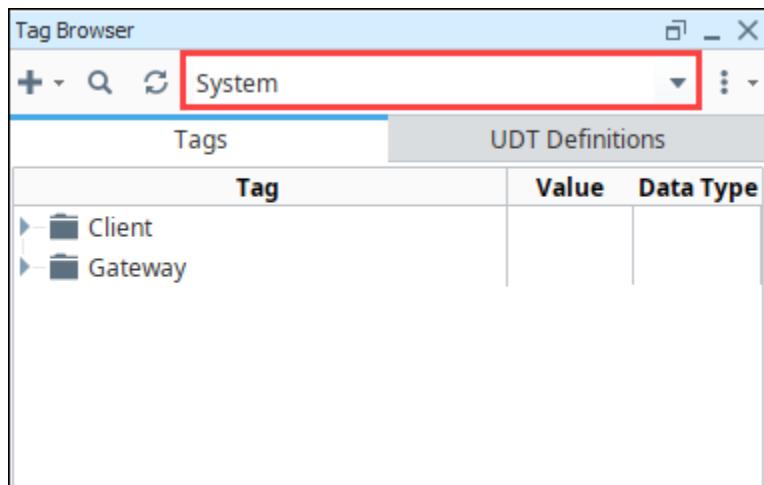
UDTs.



System Tags

System Tags provide status about the Ignition system. They're generally not interactable, but provide use information about how the system is performing.

More information can be found on the [System Tags](#) page.



Vision Client Tags

Within the Vision module, you can also have Vision Client Tags that are specific to a Vision Client. Their values are isolated to a Client runtime. For more information, see [Vision Client Tags](#).

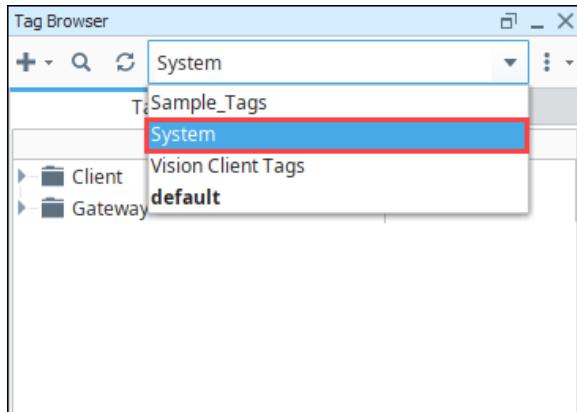
Related Topics ...

- [System Tags](#)
- [User Defined Types - UDTs](#)

In This Section ...

System Tags

System Tags provide status about the Ignition system, such as memory usage, performance metrics, and so on. System Tags cannot be deleted or modified. To view the System Tags in the Tag Browser, go to the Tag Provider Selector and select **System**.



On this page ...

- [System Client Tags \(Vision Only\)](#)
- [Gateway System Tags](#)



System Tags

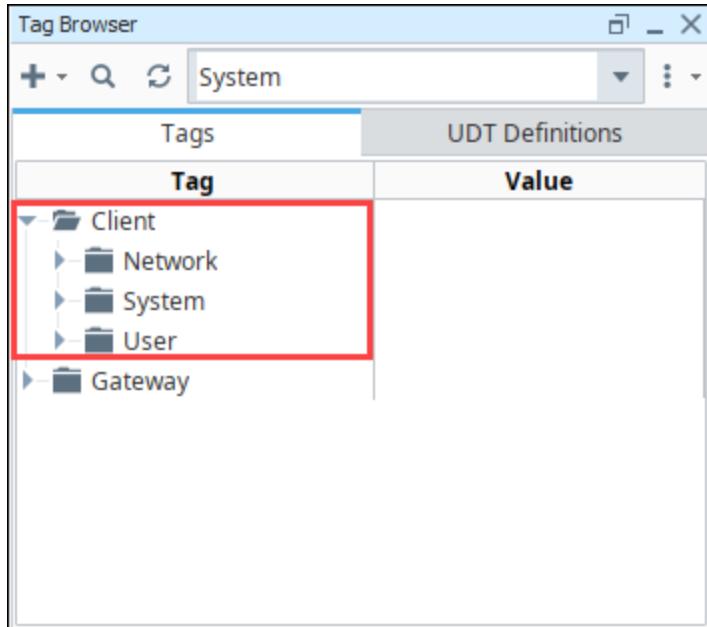
[Watch the Video](#)

The System Tags folders are displayed: **Client** and **Gateway**. The scope for each is slightly different.

Tag		Value
Client		
Network		
System		
User		
Gateway		
Alarming		
Database		
Devices		
EAM		
Gateway Network		
OPC		
Performance		
Redundancy		
Sessions		
CurrentDateTime	2020-10-06 10:18...	
SystemName	Ignition-TR-89MC...	
Timezone	America/Los_Ange...	
UptimeSeconds	579,872	

System Client Tags (Vision Only)

Client-scoped System Tags provide status information about the client's system. They can be used with the Vision module for any Vision Client. Every individual client is going to have their own values, such as IP address, host name, username, and more. There are three folders within the System > Client folder: **Network**, **System**, and **User**. You cannot modify Client System Tags.



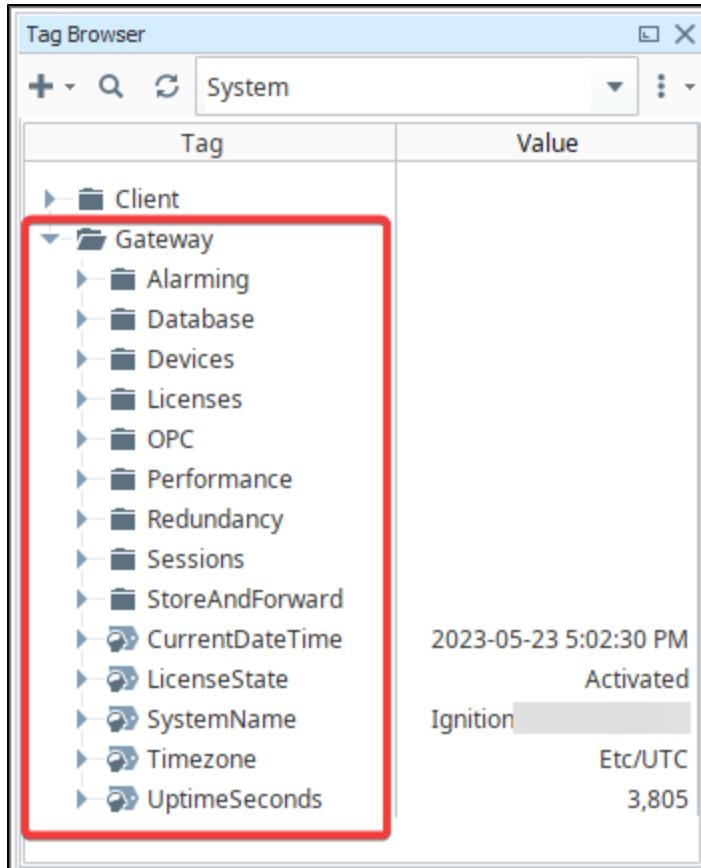
Vision System Client Tags

Tag	Description	Data Type
Network Folder		
GatewayAddress	Gateway URL address.	string
GatewayRedundancyRole	Redundancy State of the Gateway that the client is connected to. Independent, Master, Backup.	string
Hostname	Hostname (name) of the computer that the Client is running on.	string
IPAddress	IP Address of the computer that the Client is running on.	string
MACAddress	MAC Address of the computer that the Client is running on.	string
System Folder		
CurrentDateTime	Current system date and time. Format is yyyy-MM-dd hh:mm:ss a.	DateTime
DefaultDatabase	Name of the default database connection used by the project.	string
DefaultTagProvider	Name of the default Tag Provider used by the project.	string
FPMIVersion	Current Ignition version in use.	string
JavaVersion	Current Java version in use by the client.	string
LastProjectUpdate	Date of the last received update notification.	DateTime
OperatingSystem	Operating system of the computer that the Client is running on.	string
ProjectName	Name field for the current project.	string
ProjectTitle	Title field for the current project.	string
ProjectUpdateAvailable	The following feature is new in Ignition version 8.1.24 Click here to check out the other new features	

	Whether or not there is a project update available for the currently open Vision Client project. Intended to be used alongside the <code>None</code> option in Vision's Project Properties .	
SystemFlags	A byte array of <code>flags</code> for the current state of the Client.	integer
UserSource	Name of the user source for the current Client.	string
User Folder		
Country	Two letter country code according to operating system. for example: US.	string
CurrentWindow	The current main window open in the project (top most Floating, Maximized window).	string
DateFormatFull	Full date format according to the operating system. Format: EEEE, MMMM d, y.	string
DateFormatLong	Long date format: MMMM d, y.	string
DateFormatMedium	Medium date format: MMM d, y.	string
DateFormatShort	Short date format: M/d/yy.	string
DateTimeFormatFull	Full date and time format: EEEE, MMMM d, y 'at' h:mm:ss a zzzz.	string
DateTimeFormatLong	Long date and time format: MMMM d, y 'at' h:mm:ss a.	string
DateTimeFormatMedium	Medium date and time format: MMM d, y, h:mm:ss a zzzz.	string
DateTimeFormatShort	Short date and time format: M/d/y, h:mm a.	string
HomeFolder	Home folder according to OS. For example: C:\Users\psmith.	string
Language	Language according to OS. For example: "en" for English.	string
OSUsername	OS user name, for example: PSmith.	string
RolesDataSet	Dataset with Roles for currently logged in user. For example: Dataset[2R x 1C].	dataset
RolesString	Comma separated string with Roles for currently logged in user. For example: Administrator, Operator.	string
TimeFormatFull	Full time format according to the operating system. Format: h:mm:ss a zzzz.	string
TimeFormatLong	Long time format: h:mm:ss a z.	string
TimeFormatMedium	Medium time format: h:mm:ss a.	string
TimeFormatShort	Sort time format: h:mm a.	string
Timezone	Current timezone, for example, America/Los Angeles.	string
Username	Currently logged in username, for example, PSmith.	string

Gateway System Tags

Gateway System Tags exist in the Gateway scope. There are several folders within the Gateway Tags folder which report metrics on various gateway level systems.



The following Gateway-scoped System Tags are available.

Tag	Description	Type
CurrentDateTime	Current system date and time. Format is yyyy-mm-dd hh:mm:ss a.	Date/Time
LicenseState	<p>The following feature is new in Ignition version 8.1.15 Click here to check out the other new features</p> <p>Returns the license state of the Gateway. If a valid license is applied, the tag will return "Activated". This includes Emergency Activation mode and Maker licensing.</p> <p>For any other license state, the tag will return "Trial".</p>	String
SystemName	Returns the name of the system where the Ignition Gateway is installed.	String
Timezone	Timezone on the Gateway computer. For example, America/Los Angeles.	String
UptimeSeconds	Number of seconds since Ignition was started.	Long
Alarming		
Active and Acked	Number of alarms currently active and acknowledged.	Integer
Active and Unacknowledged	Number of alarms currently active and unacknowledged.	Integer

Unacked		
Clear and Acked	Number of alarms cleared and acknowledged.	Integ
Clear and Unacked	Number of alarms cleared and unacknowledged.	Integ

Database

Note: There will be a subfolder for each database connection, or none if there are no connections. Each subfolder will have the following Tags.

Active Connections	Number of active connections in the pool to this database connection.	Integ
Available	Indicates whether this datasource is available.	Bool
AvailableThroughFailover	Indicates if any database along the failover chain attached to this data source can be reached.	Bool
AvgQueryTime	Average time, in seconds, that it is taking database queries to run.	Integ
ConnectionSaturation	Percentage of possible query throughput that is being used (ratio of currently active connections to maximum possible connections).	Double
QueriesPerSecond	Number of queries running per second.	Integ

Devices

Note: For each device configuration on the gateway, a separate subfolder will exist. The contents of each subfolder are listed below.

Description	Description of the device connection as configured on the device connection on the Gateway.	String
Enabled	Boolean representing whether or not the device connection is enabled.	Bool
Name	Name of the device connection.	String
Status	Status of the connection.	String

EAM

Note: System Tags are created on the Controller for each agent, allowing you to build EAM monitoring clients, set alarms on individual Tags, and many other Tag-related operations. A Tag folder is created for each agent. As with other system Tag values, all EAM system Tag values are read-only.

Metrics Folder	Data table	<p>Note: There will be a subfolder for each database connection, or none if there are no connections. Each subfolder will have the following Tags.</p>
----------------	------------	---

seF old er	Tag		Description	Data Type
	Active Connections		Number of active connections in the pool to this database connection.	Integer
	Active Queries		Number of active queries.	Integer
	Active Query Duration		Average time, in seconds, that it is taking database queries to run.	Integer
Loggi ngF old er	Average Throughput Per Second		Number of records forwarded to the data sink per second. Throughput will be -1 if idle.	Integer
	Tag		Description	Data Type
	Errors Per Hour		Number of errors logged per hour.	Integer
	Errors Per Minute		Number of errors logged per minute.	Integer
Sessi onsF old er	Tag		Description	Data Type
	Clients		Number of currently running Vision clients.	Long
	Designers		Number of currently running Designers.	Long
	Perspective Sessions		The following feature is new in Ignition version 8.1.16 Click here to check out the other new features	Long
Syste mF old er	Number of currently running Perspective sessions.			
	Tag		Description	Data Type
	CPU Usage		CPU Utilization as reported to the Java Virtual Machine.	Float
	Memory Usage		Amount of RAM currently in use by the Gateway, in whole percentages from 0-100.	Float
Agent Group	Name of the group that the Agent belongs to.			
Agent Name	Name of the Agent.			
Brow serU RL	Agent Gateway URL address.			
IsApp roved	Whether this network connection has been approved.			
IsCon nected	Whether the Gateway Network connection is active.			
IsRun ning	Whether the Agent Gateway is considered "healthy". Healthy Agent Gateways may not be connected to the Controller, but not enough time has passed to consider the Agent Gateway as "down".			
LastC omm unicat ion	A timestamp representing the last known communication received from the Gateway.			
Licen seKey	Returns the license key of the Gateway.			
Node	Role of the connected Agent.			

Role	Independent - Redundancy is not enabled and this Ignition system runs as an independent node. Master - This is the Master node, who listens for a connection from the Backup node, and is in charge of managing system synchronization. Backup - This is the Backup node, who will connect to the Master and receive system updates.	
PlatformEdition	Ignition Edition running on the Agent Gateway. Values include standard, edge, and maker.	String
RunningState	String representation of the Running State. Possible values for RunningState are: Disconnected, Running, Warned, Errorred	String
RunningStateInt	Integer representation of the Running State.	Integer
ServerId	ID of the Agent Gateway.	String
Version	Version of Ignition software running on the Agent Gateway.	String

Gateway Network

Note:

Each system connected over the gateway network receives a special folder designated by the following nomenclature:

- Independent Gateways: 0:0
- Redundant Master Gateway: 0:1
- Redundant Backup Gateway: 0:2
- "Virtual" route aliased to either the Master or Backup Gateways: 0:0

Each folder contains the following Tags:

IsAvailable	Whether the remote gateway is available or not.	Bool
LastComm	A timestamp representing the last known communication received from the remote Gateway.	Date/Time

Licenses

The following feature is new in Ignition version **8.1.28**
[Click here](#) to check out the other new features

Note: This Tag folder will only appear if your Gateway is using a [Leased License](#). For each leased license on the Gateway, a separate subfolder will exist. The contents of each subfolder are listed below.

ConsecutiveErrorCount	Rolling counter for how many times the license tried to renew itself, if failed.	Long
ErrorCode	Displays an error code if the license failed to renew.	Integer
Expiration	Time when the session expires.	Date/Time
ExpiresInSeconds	Amount of time left before the current lease period expires, in seconds.	Long
NextCheck	Time when the license will attempt to renew.	Date/Time
Problem	Description of why the license failed to renew or activate.	String

em		
Status	Current status of the leased license.	String
OPC		
Note: There will be a subfolder for each OPC UA Server. Each subfolder will have the following Tags.		
Connected	Whether the OPC UA server is connected to Ignition.	Bool
Enabled	Whether the OPC UA server connection is enabled.	Bool
State	The state name of the connection. For example: Connected, Faulted, Connecting.	String
Performance		
Available Disk Space (MB)	Available disk space on the computer Ignition is installed on, in megabytes.	Long
CPU Usage	CPU Utilization as reported to the Java Virtual Machine.	Double
Disk Utilization	Percentage of hard disk that is in use.	Double
Max Memory	Maximum amount of RAM the Gateway can use, in megabytes.	Long
Memory Usage	Amount of RAM currently in use by the Gateway, in megabytes.	Long
Memory Utilization	Current memory usage/maximum memory usage.	Double
Redundancy		
Connection, Is Connected	Whether this Gateway is connected to another for redundancy.	Bool
Connection, PeerId	The ID of the Gateway connected to, empty string if not connected.	String
ActivityLevel	Indicates where the Gateway is in the redundant state. Can be undecided, cold, warm, or active.	String
IsActive	Whether the Gateway is running.	Bool
IsMaster	Whether the Gateway is the master. False if the backup is in control.	Bool
Role	Named role of the Gateway. Options: Independent, Master, Backup.	String
Sessions		
SessionCount	Number of active sessions on this Gateway. Note this value does not account for Perspective Client Sessions.	Integer

Store and Forward

The following feature is new in Ignition version **8.1.2**
[Click here](#) to check out the other new features

The Store and Forward System Tags were added in version 8.1.2

Note: There will be a subfolder for each database connection, or none if there are no connections. Each subfolder will have the following Tags.

Database Storage

Available	Indicates whether Database Storage is available for this Store and Forward engine.	Bool
isStore	Indicates if the data storage sink is able to store records. This data store is used for the optimization of the records before they are forwarded to the database and no data is technically stored in it. For this reason, this tag's value will always be false.	Bool

Local Cache

Forw ardM etrics Folder	Tag	Description	Data Type
	AverageDuration	The average duration for records to be forwarded to the data sink from the local cache.	Double
	MaxDuration	The maximum duration for records to be forwarded to the data sink from the local cache.	Double
	MinDuration	The minimum duration for records to be forwarded to the data sink from the local cache.	Double
	Throughput	Number of records forwarded to the data sink per second. Throughput will be -1 if idle.	Double
	TimeUnit	Unit of time for this variable. The unit of time for this variable is seconds.	String
	Total	Total number of records forwarded to the data sink from the local cache.	Long
Stora geMe trics Folder	Tag	Description	Data Type
	AverageDuration	The average duration for records to be forwarded to the data sink from the local cache.	Double
	MaxDuration	The maximum duration for records to be forwarded to the data sink from the local cache.	Double
	MinDuration	The minimum duration for records to be forwarded to the data sink from the local cache.	Double
	Throughput	Number of records forwarded to the data sink per second. Throughput will be -1 if idle.	Double
	TimeUnit	Unit of time for this variable. The unit of time for this variable is seconds.	String
	Total	Total number of records forwarded to the data sink from the local cache.	Long
Take Metrics Folder	Tag	Description	Data Type
	AverageDuration	The average duration to read a record from the local cache.	Double
	MaxDuration	The maximum duration to read a record from the local cache.	Double
	MinDuration	The minimum duration to read a record from the local cache.	Double
	Throughput	The number of records read from the local cache per second. Throughput will be -1 if idle.	Double
	TimeUnit	Unit of time for this variable. The unit of time for this variable is seconds.	String
	Total	Total number of records read from the local cache.	Long
Available	Indicates if the local cache for this Store and Forward engine is available.	Bool	
CanQuarantine	Indicates if this local cache can quarantine records. If the local cache cannot quarantine a record, the record will be dropped and lost forever.	Bool	
DroppedRecords	Indicates the number of dropped records for this local cache. A record is considered dropped if it can not be added to one of the buffers, (i.e., when a buffer is full and the Store and Forward engine can no longer accept new records).	Integ	

IsStore	Indicates if the local cache is able to store records.	Bool
MaxRecords	Maximum number of records this local cache can accept.	Integ
PendingRecords	Number of pending records for this local cache.	Integ
QuarantineRecords	Number of quarantined records in this local cache. Quarantined data is data that has erred-out multiple times during attempts to forward it, or data that could not be stored because of some configuration issues.	Integ

Memory Buffer

	Tag	Description	Data Type
ForwardMetrics Folder	AverageDuration	The average duration for records to be forwarded to the data sink from the memory buffer.	Double
	MaxDuration	The maximum duration for records to be forwarded to the data sink from the memory buffer.	Double
	MinDuration	The minimum duration for records to be forwarded to the data sink from the memory buffer.	Double
	Throughput	Number of records forwarded to the data sink per second. Throughput will be -1 if idle.	Double
	TimeUnit	Unit of time for this variable. The unit of time for this variable is seconds.	String
	Total	Total number of records forwarded to the data sink from the memory buffer.	Long
StorageMetrics Folder	AverageDuration	The average duration to store and record into the memory buffer.	Double
	MaxDuration	The maximum duration to store a record into the memory buffer.	Double
	MinDuration	The minimum duration to store a record into the memory buffer.	Double
	Throughput	The number of records that go through the memory buffer per second. Throughput will be -1 if idle.	Double
	TimeUnit	Unit of time for this variable. The unit of time for this variable is seconds.	String
	Total	Total number of records stored in the memory buffer.	Long
TakeMetrics Folder	AverageDuration	The average duration to read a record from the memory buffer.	Double
	MaxDuration	The maximum duration to read a record from the memory buffer.	Double
	MinDuration	The minimum duration to read a record from the memory buffer.	Double
	Throughput	The number of records read from the memory buffer per second. Throughput will be -1 if idle.	Double
	TimeUnit	Unit of time for this variable. The unit of time for this variable is seconds.	String
	Total	Total number of records read from the memory buffer.	Long
Available	Indicates if the memory buffer for this Stored and Forward engine is available.	Bool	
CanQuarantine	Indicates if this memory buffer can quarantine records. The memory buffer cannot quarantine records so this will always be false.	Bool	
DroppedRecords	Indicates the number of dropped records for this memory buffer. A record is considered dropped if it can not be added to one of the buffers, (i.e., when a buffer is full and the Store and Forward engine can no longer accept new records).	Integ	
IsStore	Indicates if the memory buffer is able to store records.	Bool	
MaxR	Maximum number of records this memory buffer can accept.	Integ	

ecords		
Pendi ngRe cords	Number of pending records for this memory buffer.	Integ
Quar antine dRec ords	Number of quarantined records for this memory buffer. Quarantined data is data that has errored-out multiple times during attempts to forward it, or data that could not be stored because of some configuration issues.	Integ
Availa ble	Indicates if this database engine is available.	Bool
Dropp ed	Number of quarantined records for this Store and Forward engine.	Integ
Quar antine	Number of quarantined records for this Store and Forward engine.	Integ

Related Topics ...

- [Vision Client Tags](#)
- [Types of Tags](#)

Creating Tags

Tags are created in the Designer in one of two ways; either using the Connected Devices window or creating Tags manually in the Tag Browser.

The Connected Devices window allows you to browse your connected devices and drag OPC Tags into your Tag Provider, as well as create folders and other Tags manually. You can also make immediate changes to the Tag properties in the Connected Devices window.

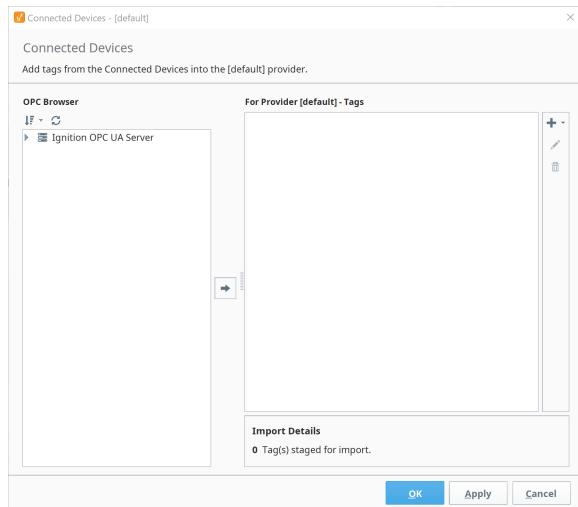
Or you can manually create Tags in the Tag Browser. The Tag Browser allows you to choose from a list of these standard Tag types and change the Tag properties directly in the Tag Editor.

This page describes both methods for creating Tags.

Connected Devices Window

The following feature is new in Ignition version 8.1.0
[Click here](#) to check out the other new features

The Connected Devices window lets you browse all your connected devices and OPC Servers for datapoints. When the Tag you're interested in is not available for browsing, you can manually add a Tag and change any Tag properties directly in the Connected Devices interface. This approach is common for devices like Modbus and Siemens where memory addresses aren't browsable. This interface allows you to create any standard type of Tag. Note that the Connected Devices window is not available when the **System** or **Vision Client Tags** providers are selected.



Tags that are grayed out in the **For Provider** area are Tags that already exist in your Tag Browser. The Tag Provider represented here is dependent on the Provider selected by the Tag Browser before the Connected Devices window was opened.

New Tags added by the Connected Devices window will appear in black font. These new Tags are effectively in a "staging" area, meaning they won't be added to the Tag Provider until you press **OK** or **Apply**. In the Connected Devices window, you can also select the Preview icon to see a summary of new Tags you are adding.

You have the option of editing your new Tags in the Connected Devices interface by selecting the Tag you want to edit in the **For Provider** area and double clicking or clicking on the edit (pencil) icon on the right. The Connected Devices edit window will open. Here you can edit any properties of the Tag. When you're finished, click **Commit** to accept your changes. When you are done adding Tags, click the **OK** or **Apply** buttons in the lower right corner to add your Tags to the Tag Browser.

Creating Tags

The example below describes creating Tags using the Connected Devices browser.

Note: In order to have any items appear under the OPC Browser tree, you must have a [device connection](#).

On this page ...

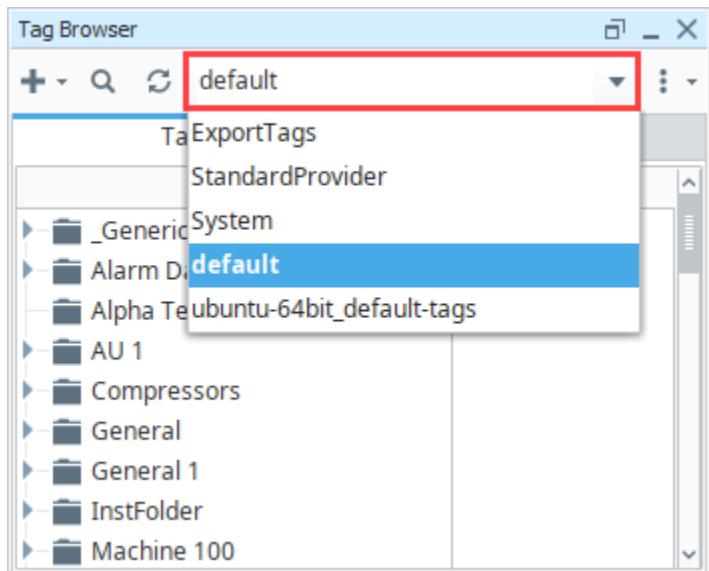
- [Connected Devices Window](#)
 - [Creating Tags](#)
- [Creating Tags in the Tag Browser](#)
- [Editing Tags](#)
 - [Edit a Tag](#)
 - [Rename a Tag](#)
 - [Cut, Paste, or Copy a Tag](#)
- [Addressing Bits](#)
 - [Address an Individual Bit in a Micrologix](#)



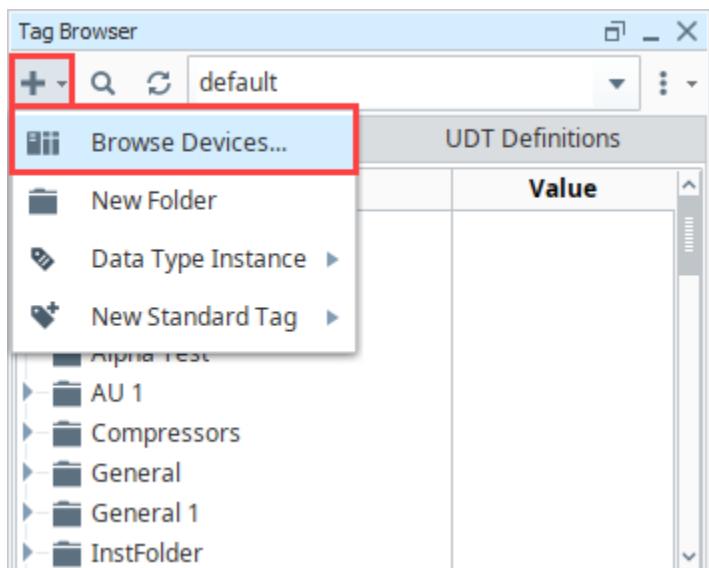
INDUCTIVE
UNIVERSITY

Creating OPC Tags

[Watch the Video](#)

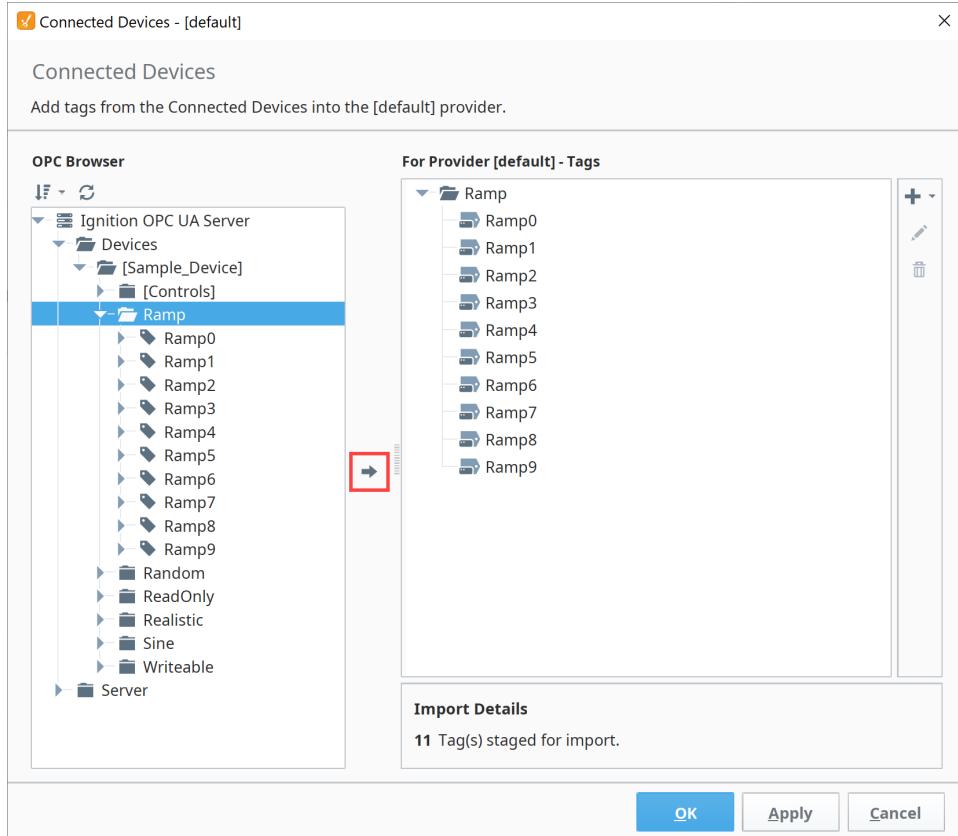


2. In the Tag Browser, click the Add icon to open the context menu. Here you can Browse Devices to add OPC Tags, or simply create a Tag or folder. Select Browse Devices...



3. The Connected Devices window will open where you can either add an OPC Tag or create a new standard Tag. To add an OPC Tag:
- Browse the OPC Tags area to find Tags that you're interested in.
 - Select the folder(s) or Tag(s) you want to add. To unselect a row, click out on a blank area, thus clearing any selected row.

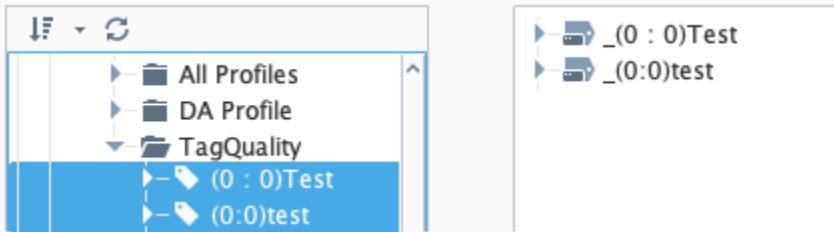
- c. Move the selected OPC Tags into your selected For provider area by clicking the Right Arrow icon or dragging them into the right panel.



Note:

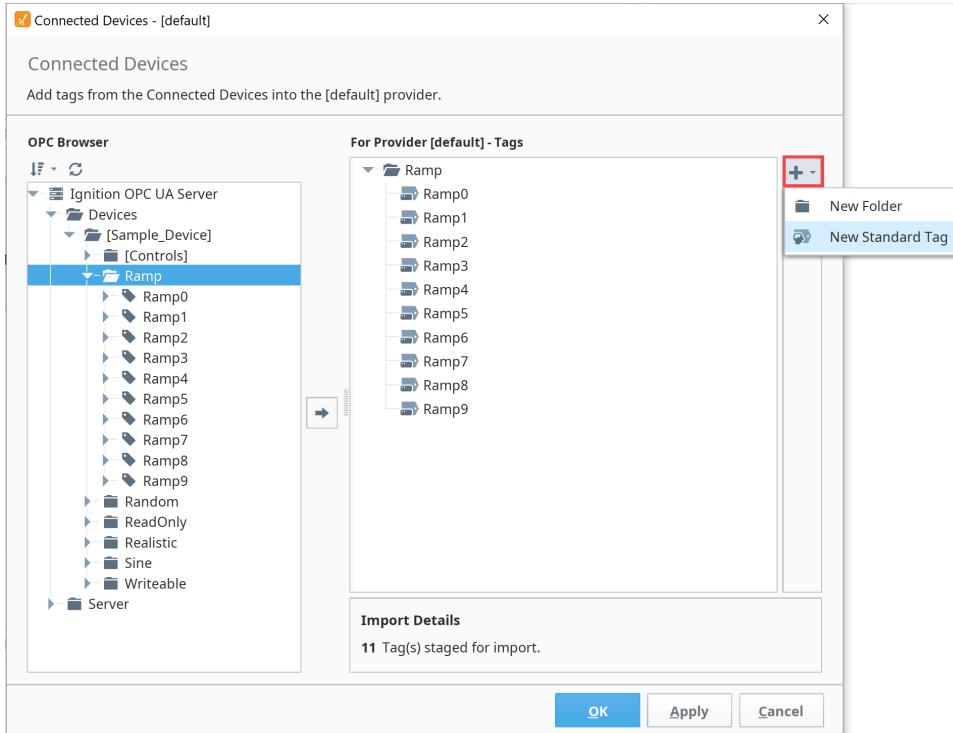
The following characters are invalid and will be replaced by an underscore when dragging and dropping an OPC item into the Tag Browser: ! @ # \$ ^ & * + []

Parentheses are valid characters but are not valid as a first character in a tag name. If an OPC item name starts with a parenthesis, the tag name will be prefixed with an underscore:

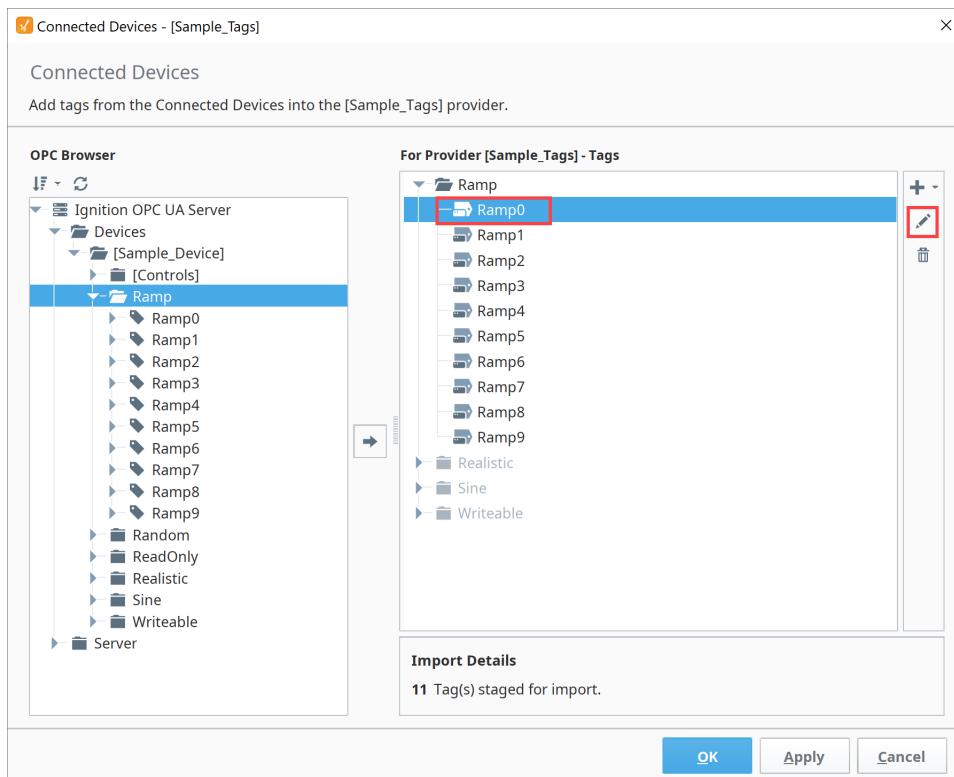


To create a New Standard Tag:

- Click the Add icon.
- Click New Standard Tag. A new blank Tag will appear in the For provider area in the currently selected folder (if one is selected), or at the bottom of the Tag list.



4. You have the choice of editing your Tags in the Connected Devices window or editing them in your Tag Browser. To edit Tags in the Connected Devices window, select a Tag and click on the Edit icon. Here you can change the Data Type, Item Path, set alarms or store history, and anything else that you can do to modify a Tag.
5. Click the Commit button to submit your changes.

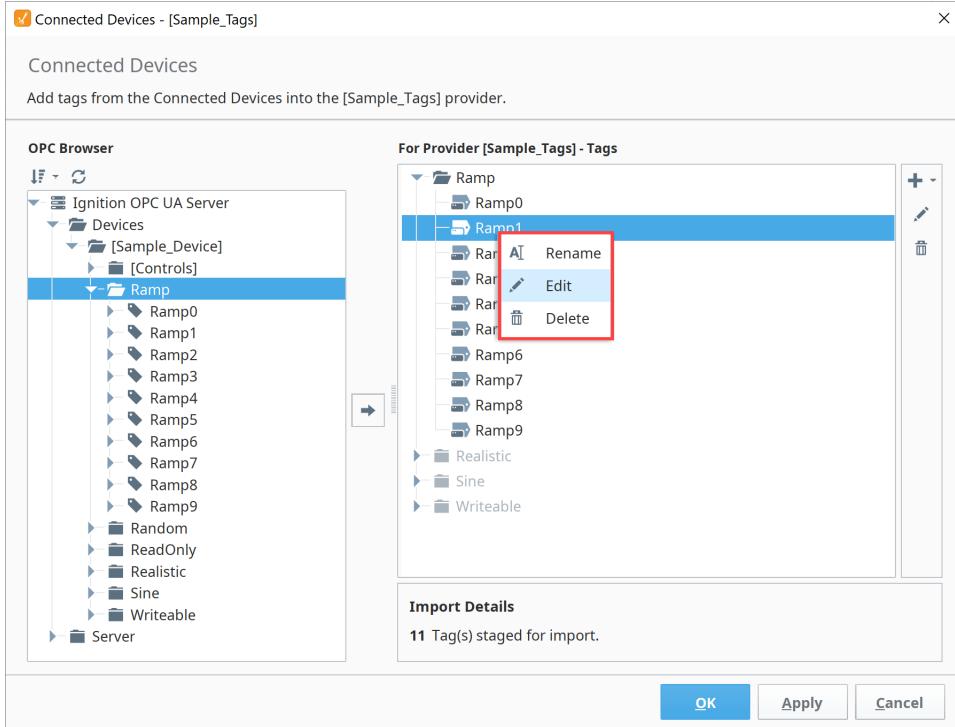


Connected Devices - [Sample_Tags]

Categories	All Properties																																																
<input type="button" value="All Properties"/> Basic Meta Data Value Numeric Security Scripting Alarms History Custom	<h3>All Properties</h3> <div style="display: flex; justify-content: space-between;"> <input type="button" value="List"/> <input type="button" value="Grid"/> <input type="button" value="Table"/> <input type="button" value="Edit"/> <input type="button" value="New"/> <input type="button" value="Delete"/> </div> <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> Basic <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Name</td><td>Ramp0</td></tr> <tr><td>Tag Group</td><td>Default</td></tr> <tr><td>Enabled</td><td>true</td></tr> </table> Meta Data <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Documentation</td><td></td></tr> <tr><td>Tooltip</td><td></td></tr> </table> Value <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Value Source</td><td>OPC</td></tr> <tr><td>Data Type</td><td>Double</td></tr> <tr><td>OPC Server</td><td>Ignition OPC UA Server</td></tr> <tr><td>OPC Item Path</td><td>ns=1;s=[Sample_Device]_Meta:Ramp/Ramp0</td></tr> </table> Numeric <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Deadband Mode</td><td>Absolute</td></tr> <tr><td>Deadband</td><td>0.0001</td></tr> <tr><td>Scale Mode</td><td>Off</td></tr> <tr><td>Engineering Units</td><td></td></tr> <tr><td>Engineering Low Limit</td><td>0.0</td></tr> <tr><td>Engineering High Limit</td><td>100.0</td></tr> <tr><td>Engineering Limit Mode</td><td>No_Clamp</td></tr> <tr><td>Format String</td><td>#,##0.##</td></tr> </table> Security <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Read Permissions</td><td>Public</td></tr> <tr><td>Read Only</td><td>false</td></tr> <tr><td>Write Permissions</td><td>Public</td></tr> </table> Scripting <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Tag Event Scripts</td><td>No scripts</td></tr> </table> Alarms <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Alarms</td><td>No alarms</td></tr> </table> History <table border="1" style="width: 100%; border-collapse: collapse;"> <tr><td>Alarm Eval Enabled</td><td>true</td></tr> <tr><td>History Enabled</td><td>false</td></tr> </table> </div>	Name	Ramp0	Tag Group	Default	Enabled	true	Documentation		Tooltip		Value Source	OPC	Data Type	Double	OPC Server	Ignition OPC UA Server	OPC Item Path	ns=1;s=[Sample_Device]_Meta:Ramp/Ramp0	Deadband Mode	Absolute	Deadband	0.0001	Scale Mode	Off	Engineering Units		Engineering Low Limit	0.0	Engineering High Limit	100.0	Engineering Limit Mode	No_Clamp	Format String	#,##0.##	Read Permissions	Public	Read Only	false	Write Permissions	Public	Tag Event Scripts	No scripts	Alarms	No alarms	Alarm Eval Enabled	true	History Enabled	false
Name	Ramp0																																																
Tag Group	Default																																																
Enabled	true																																																
Documentation																																																	
Tooltip																																																	
Value Source	OPC																																																
Data Type	Double																																																
OPC Server	Ignition OPC UA Server																																																
OPC Item Path	ns=1;s=[Sample_Device]_Meta:Ramp/Ramp0																																																
Deadband Mode	Absolute																																																
Deadband	0.0001																																																
Scale Mode	Off																																																
Engineering Units																																																	
Engineering Low Limit	0.0																																																
Engineering High Limit	100.0																																																
Engineering Limit Mode	No_Clamp																																																
Format String	#,##0.##																																																
Read Permissions	Public																																																
Read Only	false																																																
Write Permissions	Public																																																
Tag Event Scripts	No scripts																																																
Alarms	No alarms																																																
Alarm Eval Enabled	true																																																
History Enabled	false																																																
	<input style="border: 1px solid red; border-radius: 5px; padding: 2px 10px; margin-right: 10px;" type="button" value="Commit"/> <input type="button" value="Revert"/>																																																

The following feature is new in Ignition version **8.1.2**
[Click here](#) to check out the other new features

As of 8.1.2, the For Provider area now features a right-click menu that also has Edit, Rename, and Delete options.



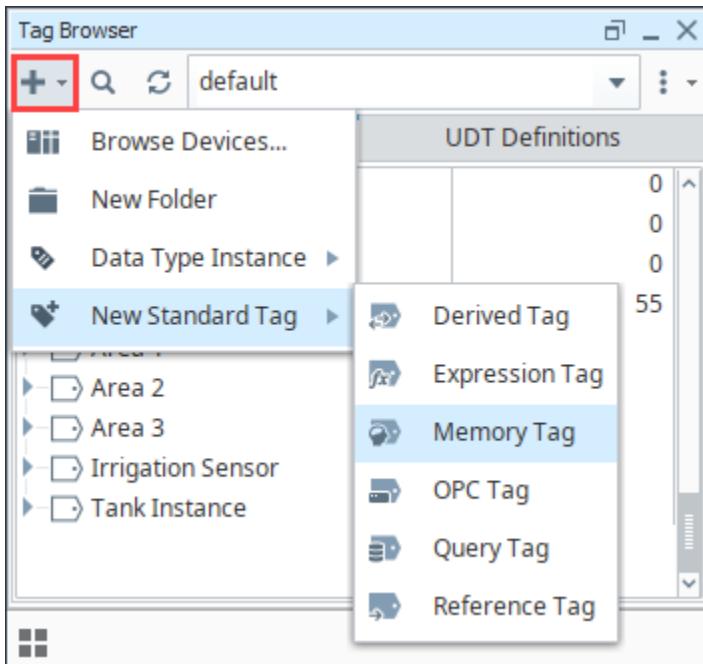
6. When finished, click the Apply button to add your new Tags to the Tag Browser. As soon as the Tags are in the Tag Browser, the values will start updating automatically.

Tag Browser	
Tags	
Tag	Value
Dairy	
Generic	
Ramp	
Ramp0	9.26
Ramp1	62.76
Ramp2	1.85
Ramp3	-3.83
Ramp4	430.86
Ramp5	452.41
Ramp6	253.63
Ramp7	387.78
Ramp8	24.91
Ramp9	1,615.74
Ramp10	9.26
Ramp11	9.26

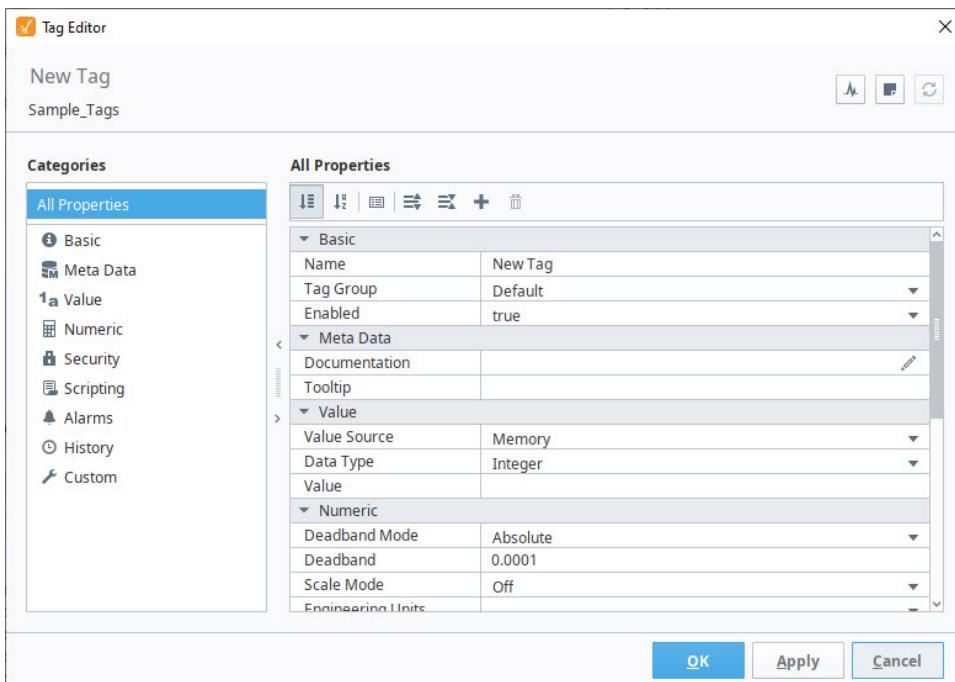
Creating Tags in the Tag Browser

Tags can be created manually in the Tag Browser. You can create any of the Standard [Tag Types](#) or a Data Type Instance. This example creates a Memory Tag.

1. You can create Tags in two different ways: click on the Add  icon to open the context menu and select New Standard Tag and choose the Tag type, or right click anywhere in the Tag Browser and select New Tag and choose the Tag type.

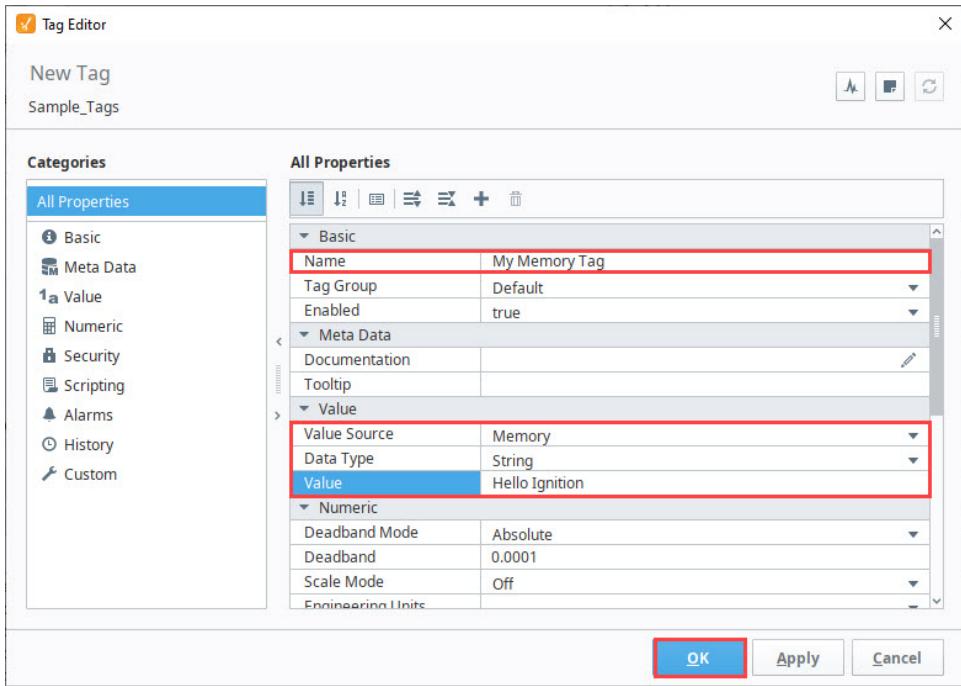


2. The Tag Editor will open to edit the Tag's properties.



3. Set the properties as follows:

Name: **My Memory Tag**
Value Source: **Memory**
Data Type: **String**
Value: **Hello Ignition**



- When finished, click **OK**. Your new Tag will be placed in the Tag column of the Tag Browser. Since this is a Memory Tag, its value will not change unless you write to it.

Editing Tags

The Tag Editor is a powerful tool used when creating Tags and for editing them. The properties displayed in the Tag Editor are custom to the type of Tag you've selected. You can find additional information on Understanding Tags and the Tag right click menu [here](#).



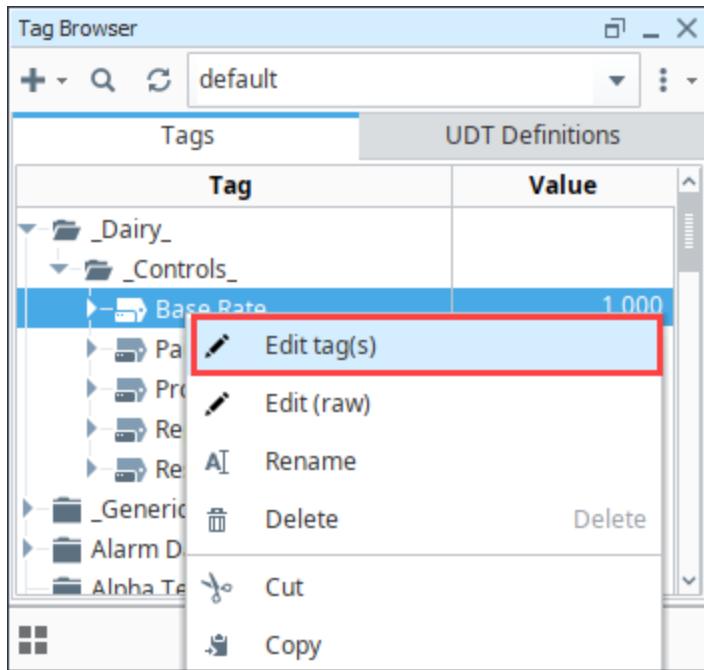
**INDUCTIVE
UNIVERSITY**

The Tag Editor

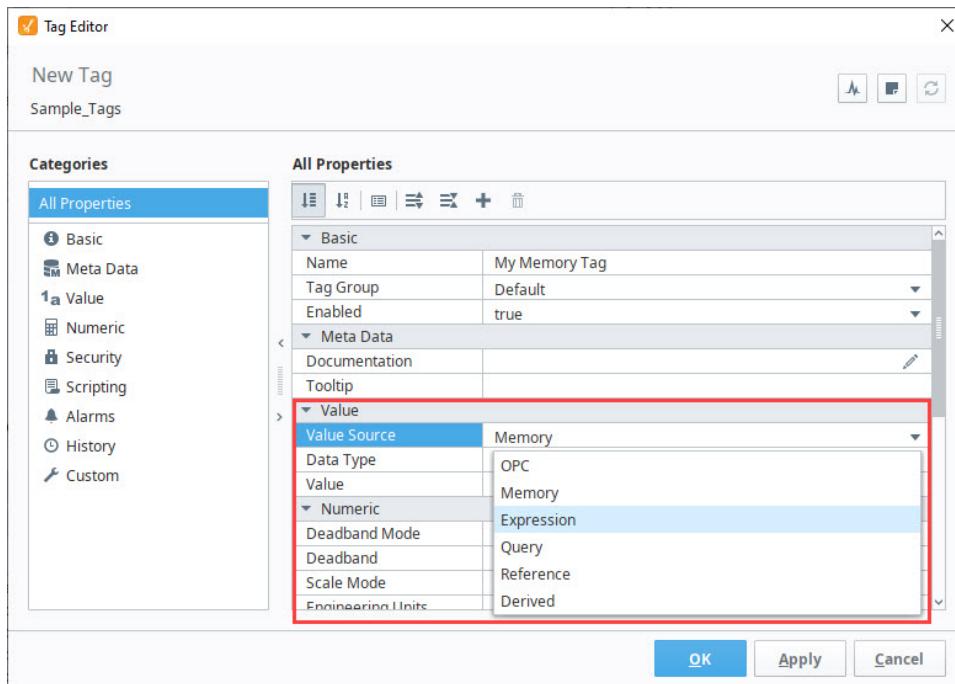
[Watch the Video](#)

Edit a Tag

- To edit an existing Tag, right-click on the Tag, and select the **Edit Tag** icon .



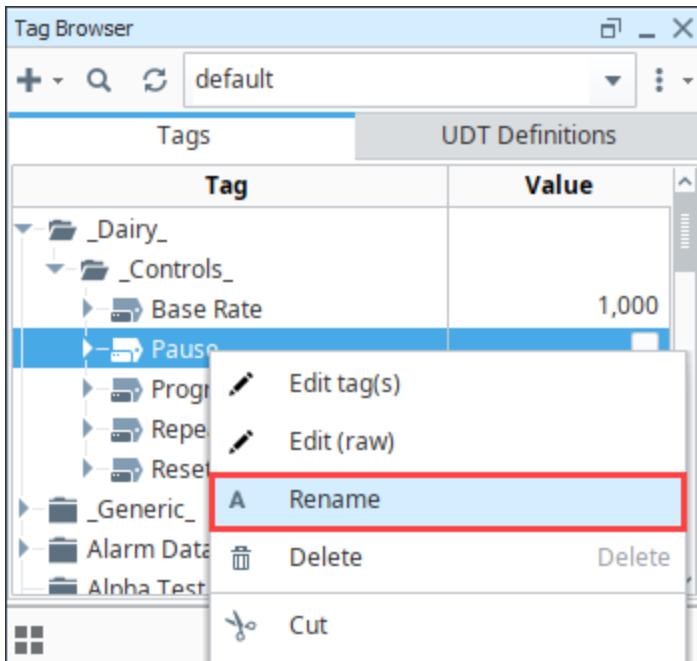
- Once in the Tag Editor, you can update the Tag properties. For example, if you want to change the Tag to a different type – such as from OPC to Expression – go to the Value Source property, click the Expand icon, and choose the type of Tag (OPC, Memory, Expression, Query, Reference, or Derived) that you want.



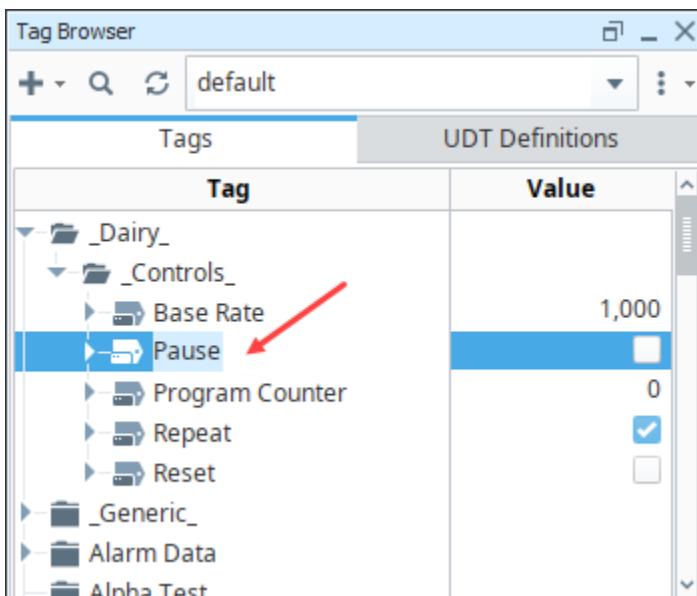
Rename a Tag

Tags names are flexible. For naming conventions, see [Tag Browser](#).

- To rename an existing Tag, right click on the Tag in the Tag Browser and select the **Rename** option.



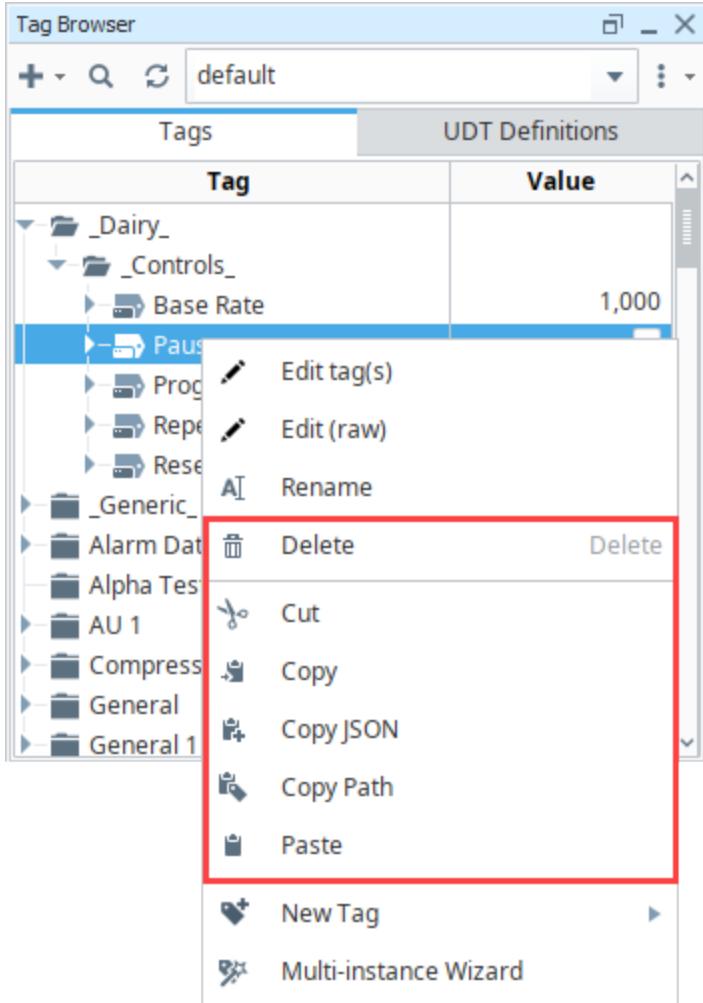
2. The cursor will now blink inside the Tag name and you can type the new name.



Cut, Paste, or Copy a Tag

You can also cut, paste, and copy Tags within the Tag Browser. Right click on the Tag in the Tag Browser. Choose the command you want.

- **Delete:** Completely removes the Tag.
- **Cut:** Delete the Tag from the current location, but leave it in the clipboard to be pasted elsewhere in the browser.
- **Copy:** Make a copy of the Tag and leaves it in the clipboard to be pasted elsewhere in the browser.
- **Paste:** Pastes the Tag you've cut or copied into the currently selected location in the Tag Browser.



Addressing Bits

In bindings and scripting there are ways to read a single bit of a word, but you can also have a Tag to read and write to a single bit. In order to address individual bits in Ignition, you must create a separate OPC Tag pointing directly to the specific bit in the PLC.

When the integer values that come from the OPC Tags are a series of binary bits, it is then possible to address each bit. For example, an integer value can have a 16-bit binary representation as shown here:

Integer	Bit level representation	How it works
4096	0,0,0,0,0,0,0,0,0,0,0,1,0,0,0	$2^{12} = 4096$
1025	1,0,0,0,0,0,0,0,0,1,0,0,0,0,0	$2^0 + 2^{10} = 1025$

Note:

Addressing bits may work differently depending on the type of device you are addressing. Most commonly you will either use /<bit> like /0 or /1, or [<bits>] like [0] or [1], or will use a . like .0 or .1. IE:

- Micrologix: [device]N7:1/0
- ControlLogix: [device]Folder/Tag.0
- Siemens: [device]I0.0

Address an Individual Bit in a Micrologix



Addressing Bits

[Watch the Video](#)

In this example, we want to address bits from a known integer value: 1025. This is represented as above, with the first (0) and eleventh (10) bits true and all others false. Our Gateway is connected to a MicroLogix PLC named MLX. To address an individual bit, do the following:

1. From the **Tag Browser** window, create an OPC Tag to have a specific value:

Data Type: **Integer**

OPC Item Path: **[MLX]B3:0**

2. Once created, set the Tag value to **1025**.

3. Then create another new OPC Tag to show only the value of the first bit of our previous Tag:

Data Type: **Boolean**

OPC Item Path: **[MLX]B3:0/0** (for Micrologic, you can specify the bit as: 0 /0 or 0.0. That is, with a slash/ or a period.)

[MLX]B3:0/0 has a value of "1" or a Boolean value of "True" because the first bit is true (integer value is odd).

4. You can create a Tag for any of the other individual bits. For example, create a new OPC Tag with a Boolean value to the second bit of the original Tag as follows:

Data Type: **Boolean**

OPC Item Path: **[MLX]B3:0/1**

[MLX]B3:0/1 has a value of "0" or a Boolean value of "False" because the second bit is false.

Once you have your Tags created, try changing the boolean values and see the Integer Tag change.

User Defined Types - UDTs

What Is a UDT?

UDTs (User Defined Types), also referred to as Complex Tags, offer the ability to leverage object-oriented data design principles in Ignition. UDTs are extremely important in Ignition. With UDTs, you can dramatically reduce the amount of work necessary to create robust systems by essentially creating parameterized "data templates".

By defining UDTs and using these "data templates", you can generate Tag instances to rapidly build complex screens. A change to the type definition is then inherited by *all instances*, drastically saving time when making routine changes.

The UDT data types are fully supported by [Vision Templates](#), which means you can configure templates for your custom data types and take advantage of drag-and-drop binding to rapidly build complex screens.

Primary UDT Features

- **Object Oriented** - Use small or large groups of Tags to create a single object. Create objects that match your real world devices or the existing structures in your PLCs.
- **Central Definition** - Once you define your data type, you can then create instances of it. If at a later time you want to change some aspect of the type, you can simply edit the type definition, and all instances of it are automatically updated.
- **Parameterized Settings** - Define custom parameters on your data type, and then reference them inside some or all of your member Tags. When it comes time to create instances, you can simply modify their parameter values in order to change where the underlying data comes from.
- **Extendable** - Data types can inherit from other data types in order to add additional members or override settings. Instances can also override settings, allowing for flexibility when dealing with irregularities and corner cases.

On this page ...

- [What Is a UDT?](#)
- [Primary UDT Features](#)
- [UDT Terminology](#)
- [Creating a UDT Definition and Instance](#)
 - [Creating a Definition](#)
 - [Creating a Data Type from Existing Tags](#)
 - [Creating an Instance](#)
 - [Override Instance Properties](#)
 - [Make Changes to the Definition](#)
- [UDT Root Node Properties](#)
- [Assigning Colors to UDTs](#)
- [Binding to UDTs](#)



INDUCTIVE
UNIVERSITY

Understanding UDTs

[Watch the Video](#)

UDT Terminology

Many terms are frequently used when discussing UDTs:

Definition

A Definition represents the structure of a UDT. Definitions don't run, so Tags inside of a Definition won't poll or subscribe to anything. Rather they represent a Tag structure which Instances will inherit from. Changes made to a Definition are automatically applied to any Instances of that Definition.

In the Tag Browser, UDT Definitions are always located under the **UDT Definitions** tab.

Tag Browser	
Tags	
Definition	Config
My Definition	
A Folder	
A Memory Tag	123
Some other Definition	
New Tag	

Instances

Instances are running copies of a Definition. All Instances have a "Parent Type", which is the definition that the Instance is inheriting from. The structure of an Instance is defined by its parent Definition, so you can not add new Members to an Instance. However, you can override the values on properties in any Member.

In the Tag Browser, UDT Instances can be found under the Tags tab, and are signified by either a plain white Tag icon, or a Tag icon with a vertical stripe. Furthermore, you can expand the UDT Instance and find the members (other Tags) in the UDT.

Tag Browser	
Tags	
Tag	Value
A Standard Tag	
Instance With Color	123
Instance Without Color Defined	

Parameters

Parameters are user created properties that can be used to create parameterized data templates. Parameters are configured on Definitions, and their values can be overridden on individual Instances. You can replace values on a member in a UDT with a reference to a parameter, allowing for example, if a data type consists of three OPC Tags that only differ by a number in the path, you can use a parameter for the "base address", allowing instances to be created with only one setting.

Root Node

The top level item in a UDT.

Members

Members are the Tags inside of a data type or instance. Members are always under a Root Node. Members can be standard Tag types or an instance of another UDT.

Override

Instances are copies of a Definition, but in some cases you may wish to change the value of a property on a particular member (Tag) in a UDT instance. This is called Overriding the property, allowing the property to have a value that deviates from the Definition.

Creating a UDT Definition and Instance

In these series of examples, we will create a very simple UDT Definition, make an instance of it, and make some additional configuration changes.

Creating a Definition

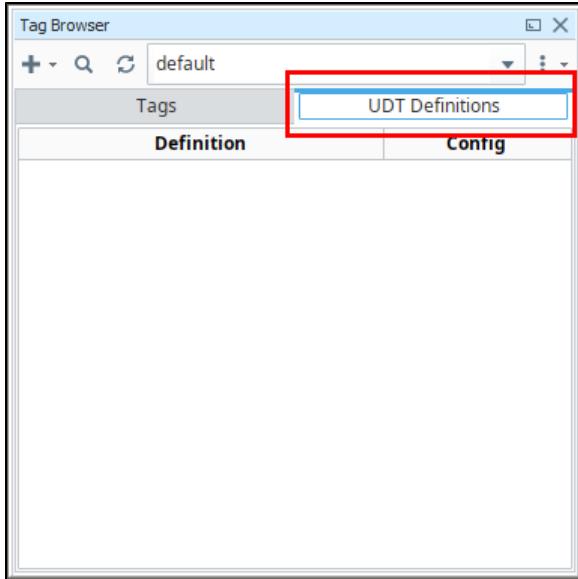
In this example, we'll demonstrate how to create a UDT Definition.

1. To create a UDT Definition, first switch click on the **UDT Definitions** tab in the Tag Browser. This is the only section where you can make UDT Definitions.

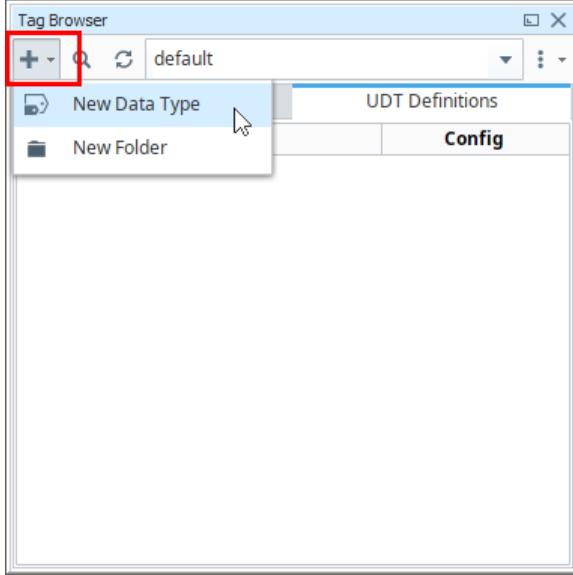
**INDUCTIVE
UNIVERSITY**

Creating UDT Definitions

[Watch the Video](#)

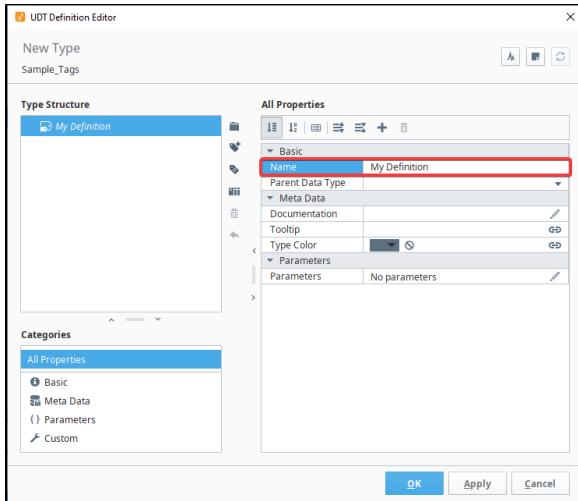


2. Next click the Add icon, and select New Data Type from the dropdown.

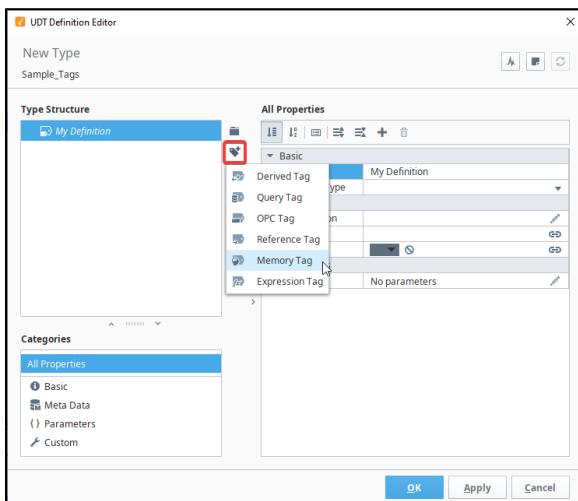


3. The Tag Editor will open, allowing you to create a new UDT Definition. To start, enter a name into the **Name** property. (For this example we used "My Definition".) This name is how the UDT will be identified by the rest of the system.

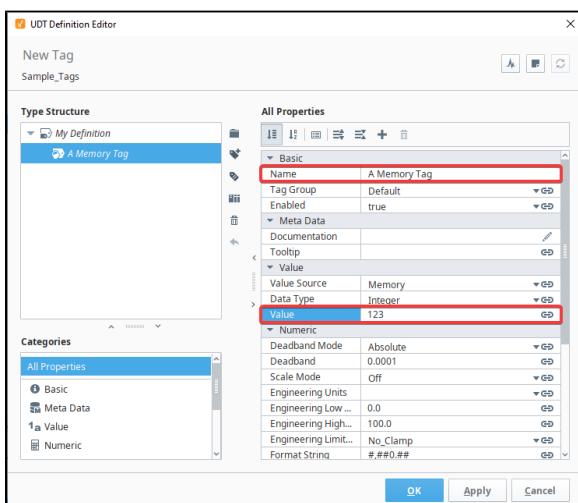
Note: Names can be changed later, but doing so after you've already made instances will create what's called an "orphaned UDT instance:" an instance that is no longer associated with a definition. It's generally advised to avoid changing the name of a definition unless you're also willing to update any of the instances.



- We'll keep this demonstration simple, and only add two members. We'll create a Memory Tag and add a Tag from a connected device. Let's create the Memory Tag first. Click the Add icon, and select **Memory Tag**.



- The Tag Editor will now add a memory Tag to the **Type Structure** tree and select it, allowing the **Properties** table to show settings for the selected Tag.
- Change the Name to "A Memory Tag", and set a Value of **123**.



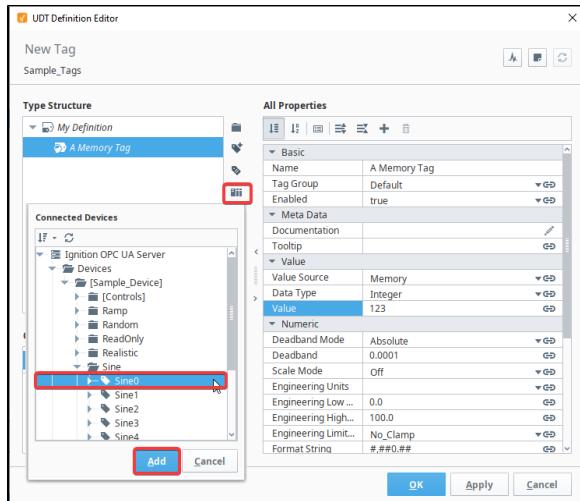
- Now, let's add a Tag from the OPC UA server. Click the Add icon, and a window will open showing you your Connected Devices. Expand **Ignition OPC UA**

Server>Devices>Generic>Sine, and select **Sine0**. Click **Add**.

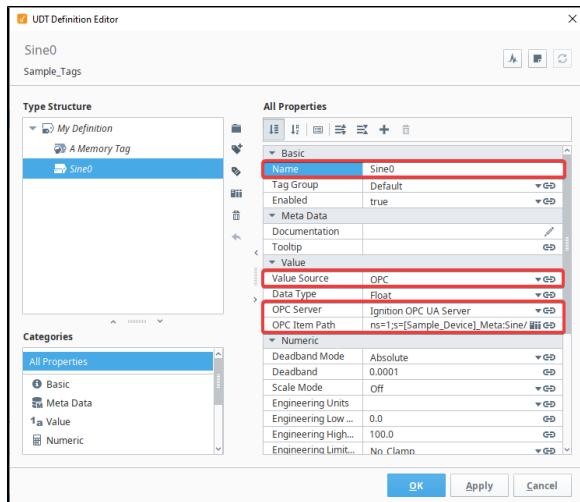
The following feature is new in Ignition version **8.1.1**
[Click here](#) to check out the other new features

As of 8.1.1, you can browse OPC devices from within the Tag Editor in order to add OPC nodes to UDT definitions.

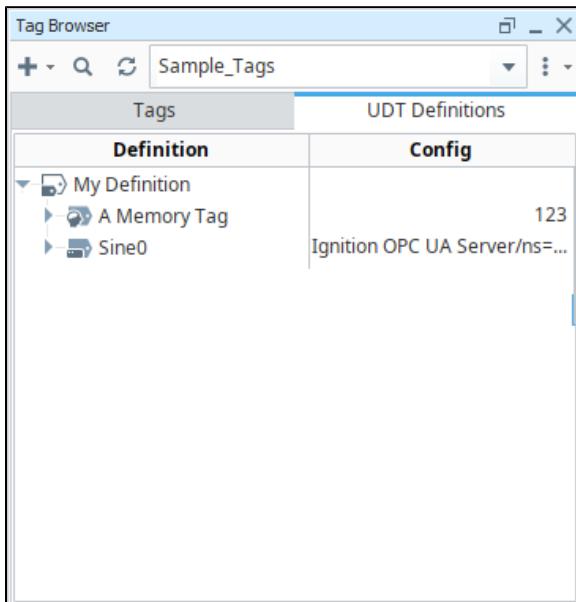
Note: When adding Tag(s) from a device, it requires that you have some connected devices such as the Ignition OPC UA Server as in this example.



8. You'll see that the **Sine0** Tag was added to the **Type Structure**. Select the **Sine0** Tag and you'll see all the property settings for that Tag.



9. Click **OK**. This will close the Tag Editor, and apply your changes. The Tag Browser will now show your UDT Definition. You can expand the "My Definition" item to make the Memory and Sine0 Tags visible.



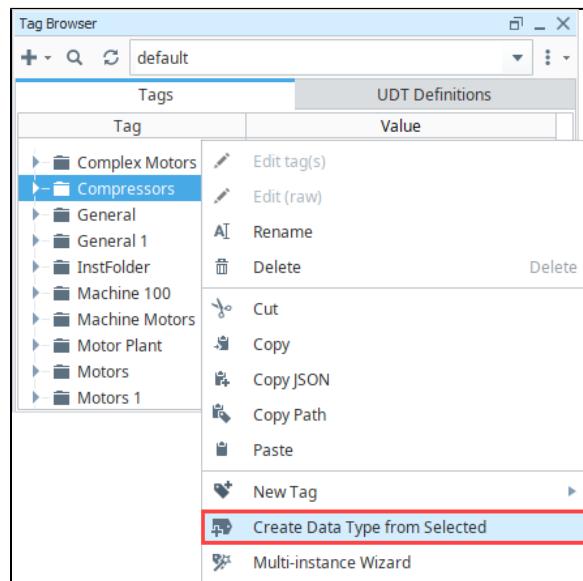
Creating a Data Type from Existing Tags

The following feature is new in Ignition version 8.1.5
[Click here](#) to check out the other new features

In this example we will demonstrate how to create a data type from existing tags.

To create a data type from existing Tags, simply select the Tags or folders you wish to include, right-click and select **Create Data Type from Selected**. The Tag Editor window is displayed with the selected tags pre-populated as members. From here you can modify the Tags, add parameters, and so on. The original Tags will not be affected.

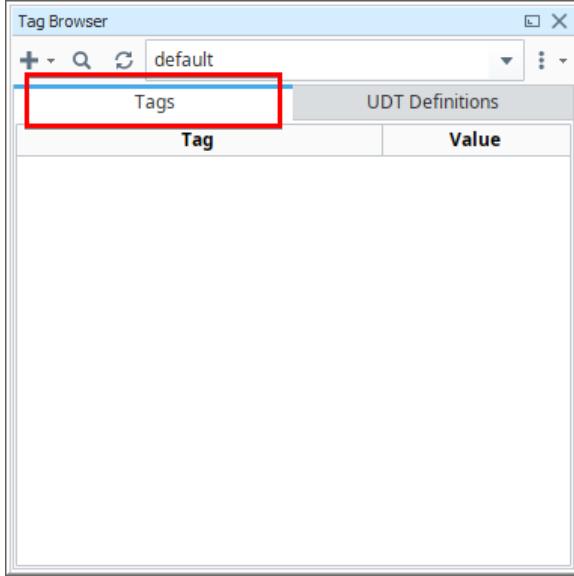
Tip: If you select a single folder as the root to create the type from, its sub-members will be added, and its name will be the basis for the type (that is, the folder itself won't be included in the structure).



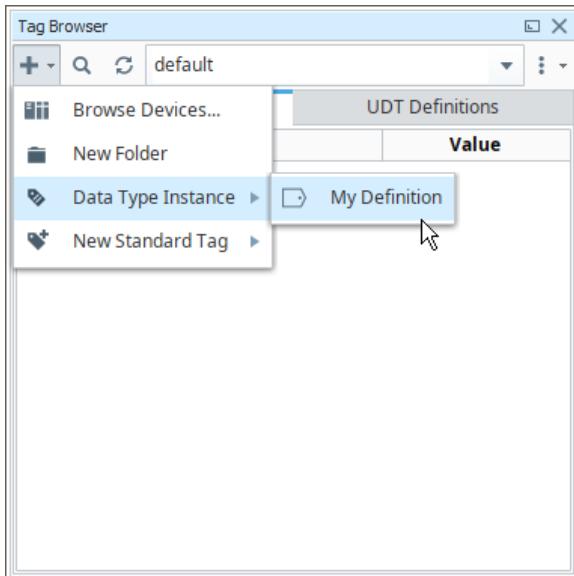
Creating an Instance

Now that we have a UDT Definition, we can create a UDT Instance. You can use the [Multi-Instance Wizard](#) to make many instances quickly, but you can also create a single instance.

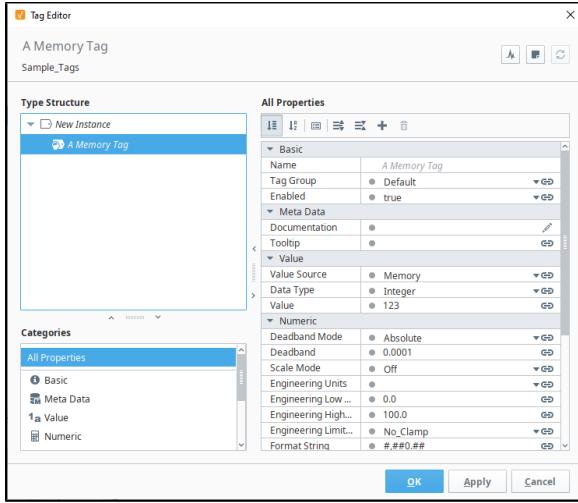
1. In the Tag Browser, switch back to the **Tags** tab. This section of the Tag Browser is where you create UDT instances. Definitions can not be placed in this section.



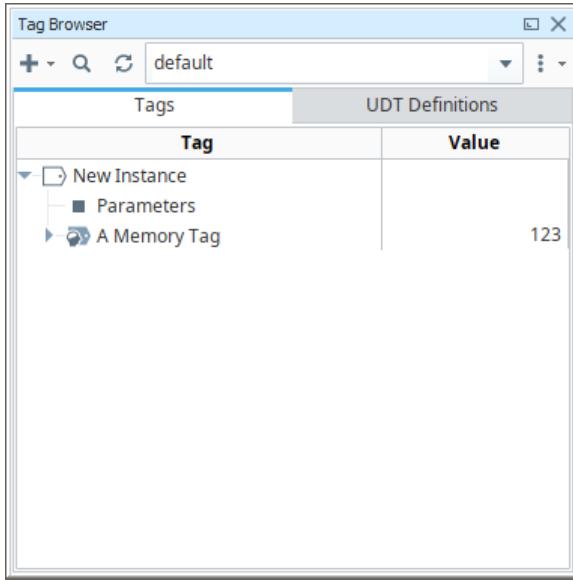
2. Click the Add icon and hover the mouse cursor over **Data Type Instance**. This will expand and show all of the UDT Definitions that exist within the active **Tag Provider**. Select the Definition we created in the previous example.



3. Again, the **Tag Editor** will open. This time allowing you to edit an instance of the UDT. From here you can name the Instance. In this case we'll use the name "New Instance".
4. If you select the member ("A Memory Tag"), you'll notice that many of the properties have a grey dot. This signifies that the properties are inheriting their values from the definition.



5. We won't make any changes here. Just click **OK**, which closes the Tag Editor and creates a UDT instance.



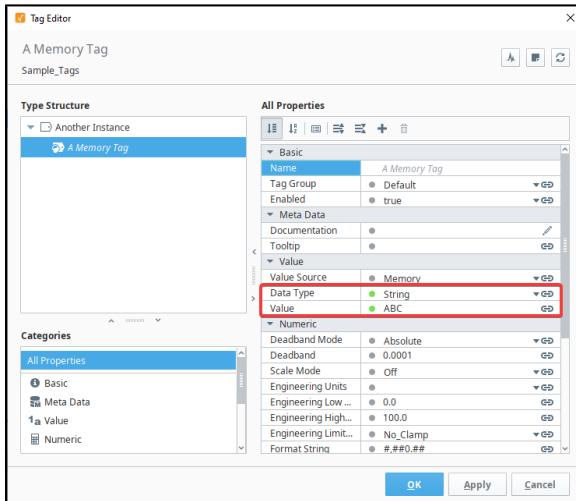
Override Instance Properties

Instances inherit their structure and properties from the Definition, but property values can be overridden on an instance.

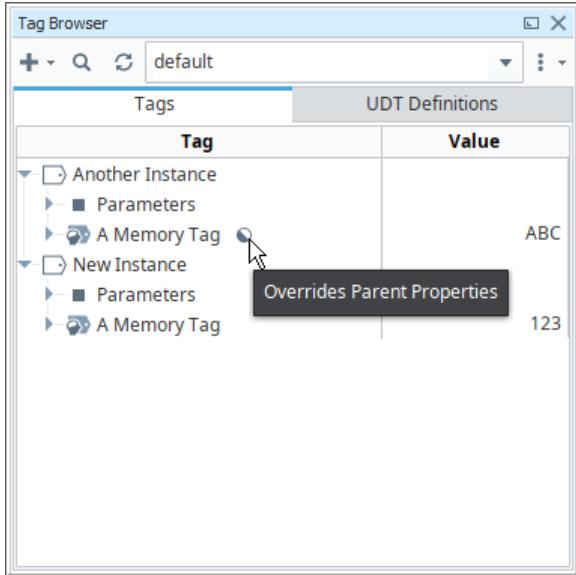
1. Now that we have an instance, let's create another. You can use the same steps in the last example, or just copy and paste **New Instance** from the Tag Browser's right-click menu.
2. Name the new instance "Another Instance".
3. If you need to make changes to a UDT Instance, you can open the **Tag Editor** by double clicking on the Instance, or any of its members.

Tags		UDT Definitions
Tag	Value	
Another Instance		
Parameters		
A Memory Tag	123	
New Instance		
Parameters		
A Memory Tag	123	

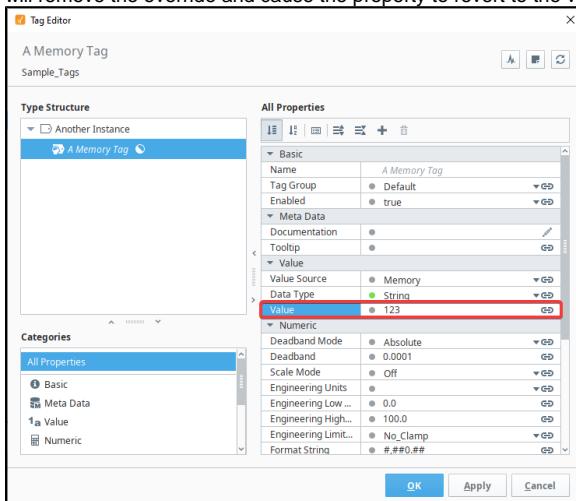
4. With both instances expanded, we can see that aside from the names, they're identical. Let's change that. Double click on **Another Instance** or any of its members to open the Tag Editor.
5. With the Tag Editor open and focused on **Another Instance**, select the memory Tag. Change the Data Type to a **String**, and the Value to "ABC". You'll notice that as you make changes to the properties, the grey dots become green signifying that the property values now differs from the Definition. This concept is known as "Overriding" a UDT member property.



6. Click **OK** to close the Tag Editor and apply your changes. You'll notice that the memory Tag in the instance we edited now has an attribute icon, signifying that one or more properties on the member are overridden.



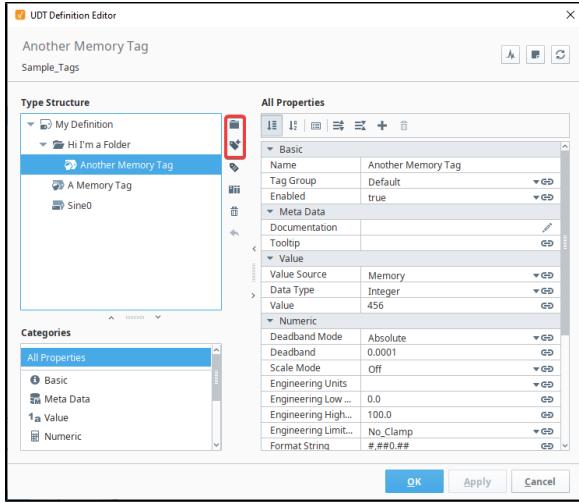
7. If you wish to remove the overrides, simply edit the Tag again, and click on the green dot. This will remove the override and cause the property to revert to the value on the Definition.



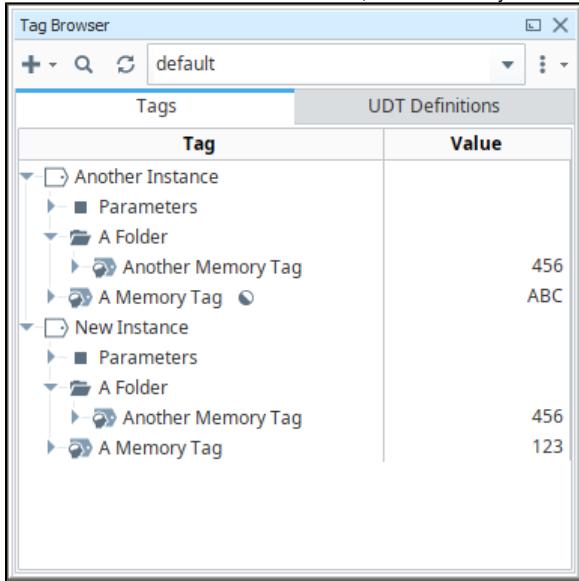
Make Changes to the Definition

Now that we have some instances, we can make a modification to the Definition, which will propagate the change down to the Instances.

1. Switch back to the **UDT Definition** tab in the Tag Browser.
2. Edit the **My Definition UDT**.
3. Use the **Add Folder** and **Add Standard Tag** icons to add a new folder, and a member inside of that folder.

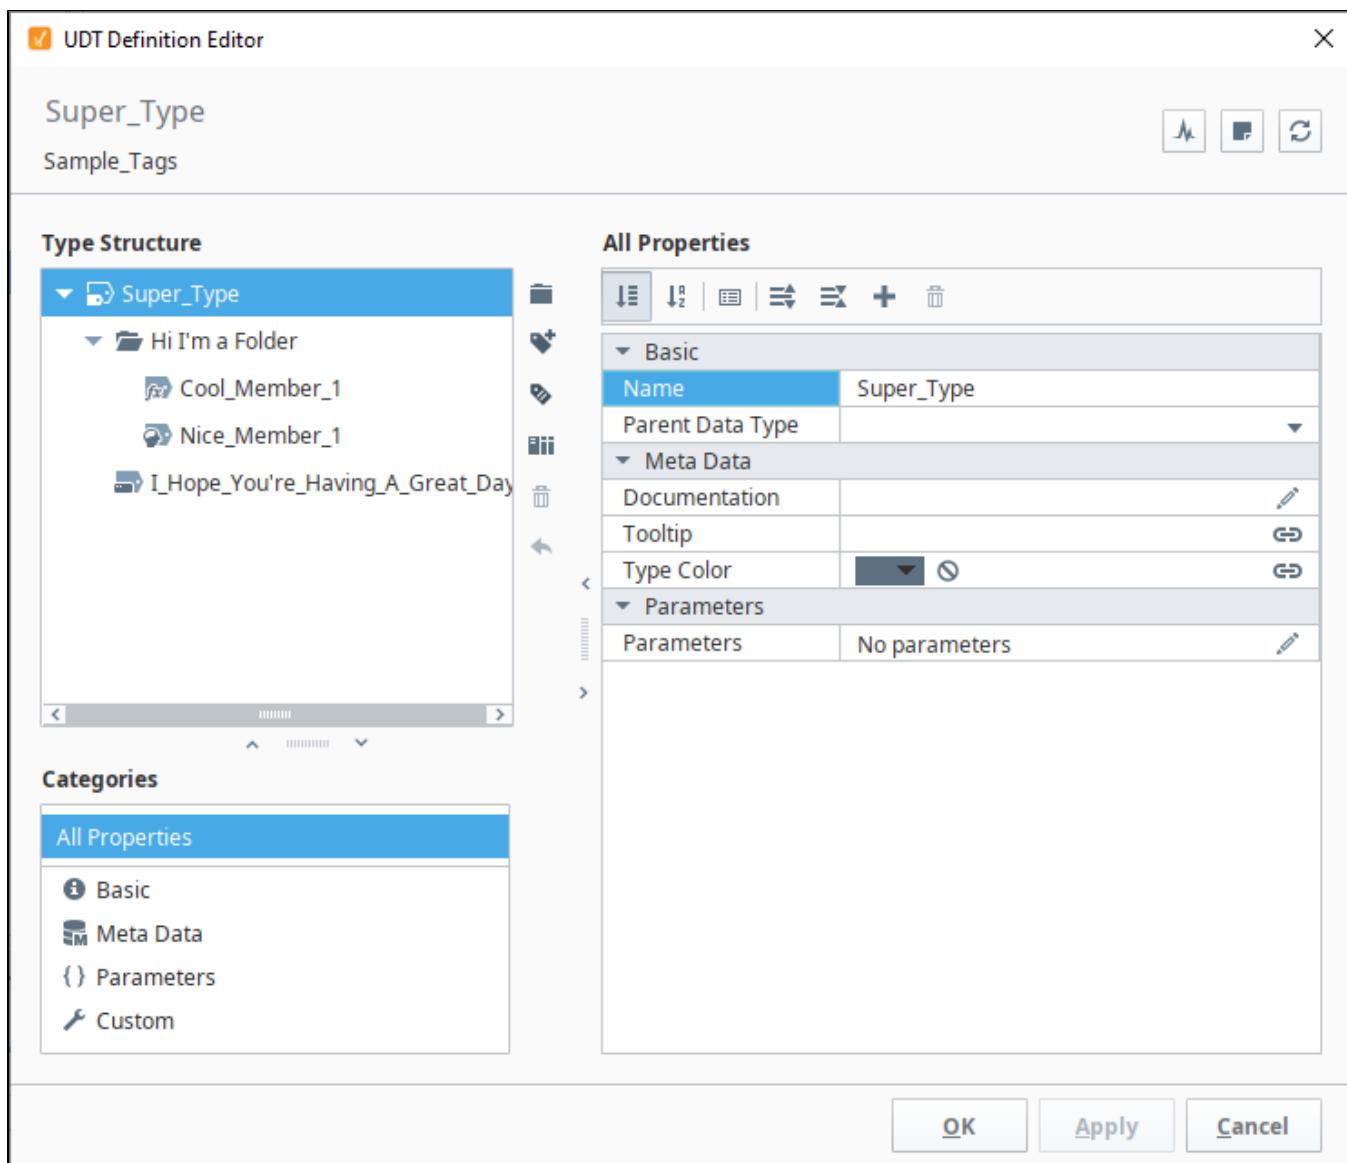


- Click **OK** to switch back to our instances.. You'll see they now both have new members.

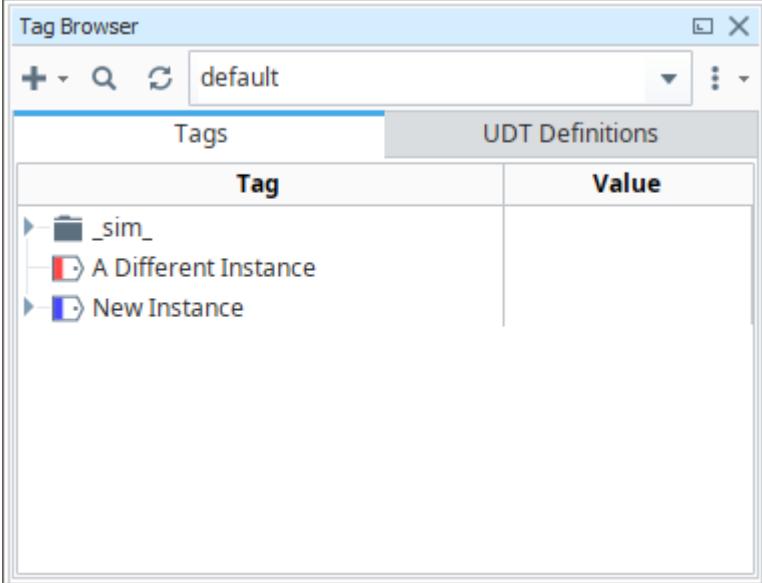


UDT Root Node Properties

While editing a UDT Definition or Instance, the Tag Editor will show some unique properties on the Root Node.



Property	JSON Name	Description
Name	name	The name of the UDT Definition.
Parent Data Type	typeId	<p>Both Instances and Definitions have this property, but the implications of the property are different.</p> <p>On a Definition - The name of the UDT Definition that this Definition is inheriting from. If blank, then the UDT being edited does not inherit from another UDT.</p> <p>On an Instance - The name of the UDT Definition this UDT is an instance of. Changing the Parent Data Type of an instance is not supported.</p>
Documentation	documentation	<p>The following feature is new in Ignition version 8.1.17 Click here to check out the other new features</p> <p>A freeform text property for information about the Tag or UDT.</p>
Tooltip	tooltip	<p>The following feature is new in Ignition version 8.1.17 Click here to check out the other new features</p>

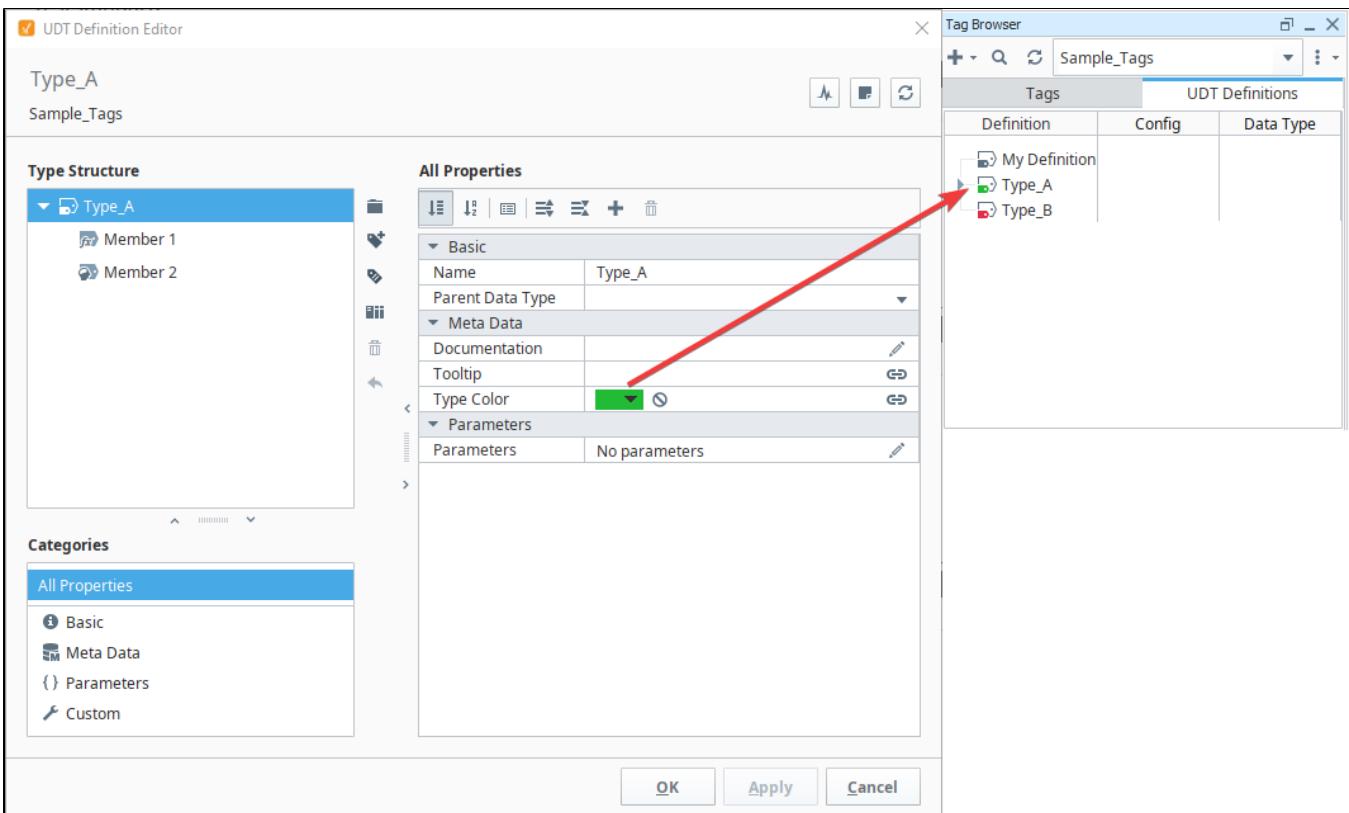
		The tooltip provides a hint to visual components as to what should be displayed when the user hovers their mouse cursor over the component that is being driven by the value of this Tag or UDT.
Parameters	parameters	A collection of parameters configured on the Definition. Note that you can only add or remove parameters on Definitions.
Type Color	typeColor	<p>The following feature is new in Ignition version 8.1.0 Click here to check out the other new features</p> <p>A color that will be applied to the Definition and any Instances. This property is only cosmetic, but can be useful to have certain UDTs stand out from others in the Tag Browser.</p> 

The following feature is new in Ignition version **8.1.0**
[Click here](#) to check out the other new features

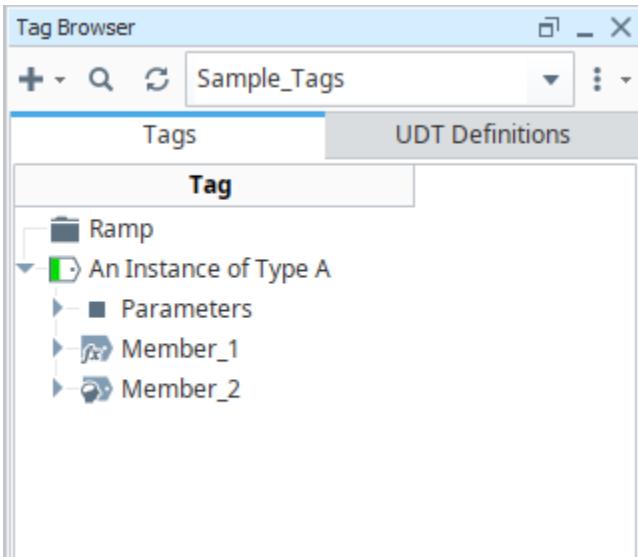
Assigning Colors to UDTs

UDTs can be color coded, which applies a color to the Root Node. This is purely a cosmetic change, but can be helpful in systems with a large number of instances, as the colors can make certain UDTs stand out from one another.

Color is applied to the Definition, via the **Type Color** property.



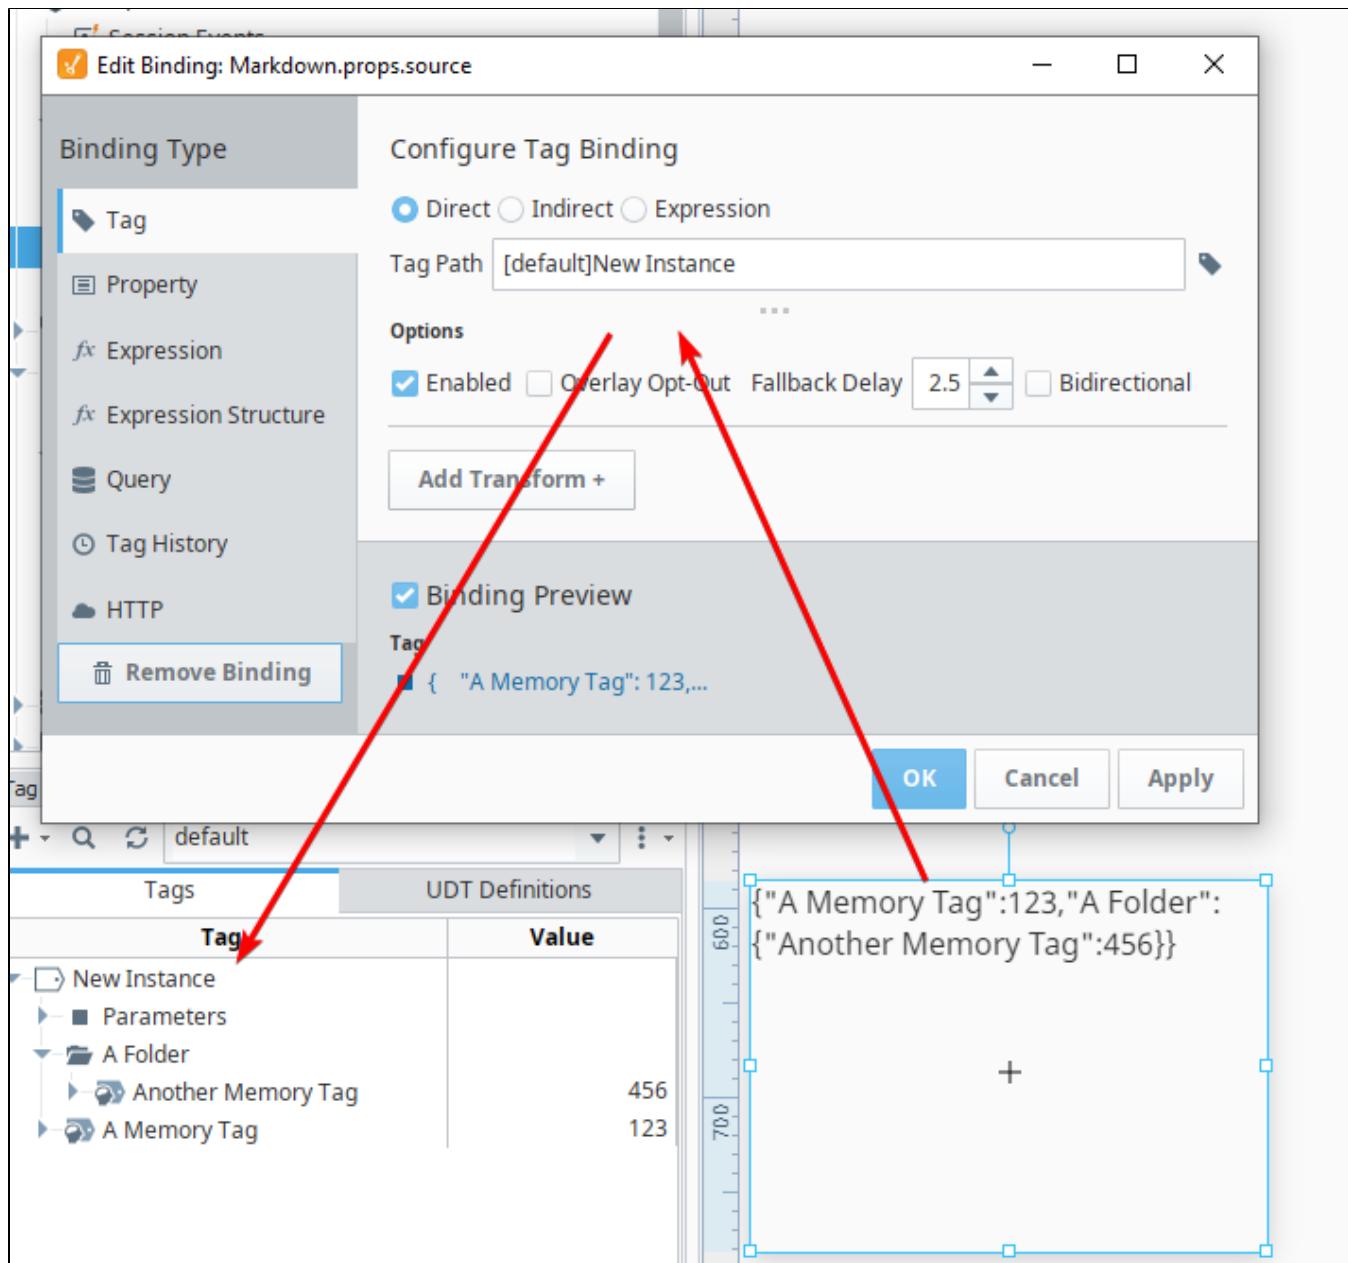
Any instances of that type will apply the color to the Root Node.



Binding to UDTs

Both Perspective and Vision feature a "binding" system, that allows components in those modules to display live values on Tags. In regards to UDT instances, component bindings can bind to members just like any standard Tag, but they can also bind directly to the UDT instance, which results in the binding receiving live live values from all members of the UDT.

In the image below, a Perspective Markdown component has a Tag binding on its **source** property. The binding is leading to the root of a UDT instance. As a result, the member values are shown on the component. They're also live values, so any value changes on any member will appear on the Markdown component.



Related Topics ...

- [UDT Parameters](#)

In This Section ...

UDT Parameters

Parameters effectively act as variables that can be referenced by properties on members. A common use case for parameters in UDTs is to make the OPC Item Path on OPC Tag members dynamic, allowing you to replace parts of the OPC Item Path with the parameter's value. However, parameter values can be referenced by other member properties, such as expressions on Expression Tags.

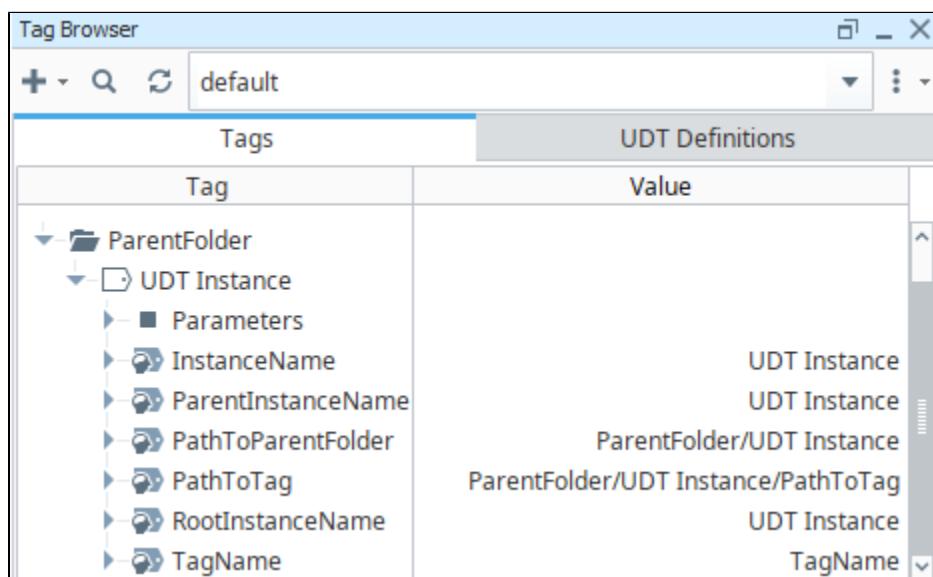
UDT Parameters are configured on UDT Definitions. Instances of a UDT can override the value of a parameter, much like any other property on an Instance.

On this page ...

- Pre-Defined Parameters
- Adding a Parameter to a UDT
 - Referencing Parameters from Member Properties
- Data Type Parameters in Expressions
- Combining Parameters and Tag References
- Attribute Referencing and Parameterized Types
 - Calculations and Numerical Parameter Names
 - Nonexistent UDT Parameters
 - Null Values on Parameters

Pre-Defined Parameters

UDTs have a few parameters already defined to make things easier for you. They give you access to the name and various paths associated with a UDT member Tag. These parameters can be accessed from anywhere in a Tag that a normal parameter can be used. Each of these parameters uses that Tag it is in as a starting point for its path.



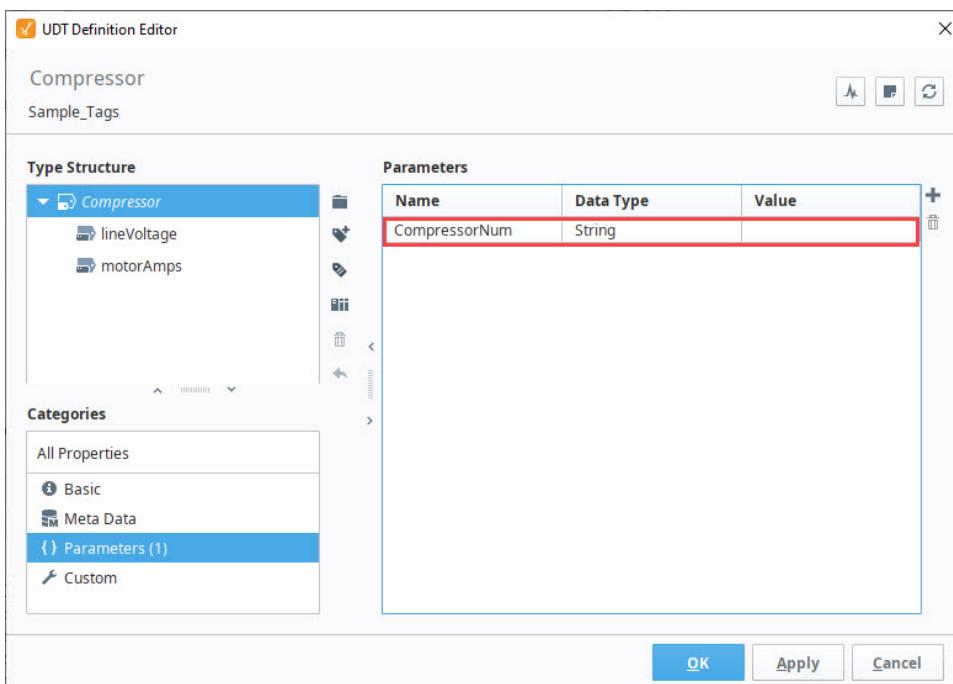
Parameter Name	Description
{InstanceName}	The name of the UDT Instance that this Tag is inside. In cases where UDTs are nested, this parameter will return the name of the UDT the member belongs to. Thus referencing this parameter in a nested UDT will return the name of the nested UDT.
{ParentInstanceName}	<p>The following feature is new in Ignition version 8.1.13 Click here to check out the other new features</p>

	The instance name of the parent UDT Instance. This is similar to InstanceName, except the name of parent UDT Instance will be returned when binding from UDT Instance parameters.
{PathToParentFolder}	The full path to the folder that this Tag is in.
{TagName}	The name of the Tag that is using this parameter.
{PathToTag}	The full path to the Tag using this parameter.
{RootInstanceName}	<p>The following feature is new in Ignition version 8.1.13 Click here to check out the other new features</p> <p>Returns the top-most UDT Instance name.</p>

Adding a Parameter to a UDT

In this example, our plant has multiple compressors. We created a Compressor UDT, but we want each instance of the Compressor UDT to reference a different set of Tags. Our Compressor UDT has two OPC Tags: lineVoltage and motorAmps. These two Tags are pointing to a specific address in the PLC. In order to reference a different set of Tags for each instance, we need to add a parameter to our Compressor UDT that we called "**CompressorNum**".

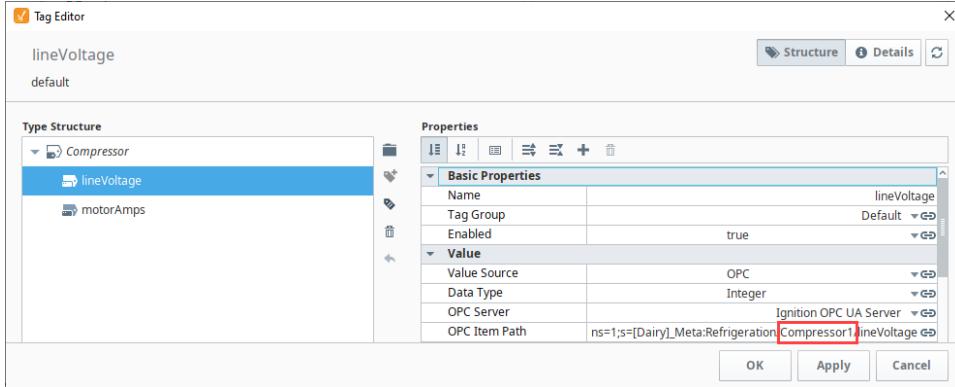
1. In the Tag Browser, we clicked on the **UDT Definitions** tab to find our Compressor UDT.
2. Double click on the Compressor UDT to open it in the Tag Editor and click on the pencil icon next to the **Parameters** property. The Parameters pane will open.
3. Click the Add icon to create a parameter. Enter the parameter **Name** and Data Type of **String**.



Referencing Parameters from Member Properties

Continuing with the example from above, this next example will show you how parameters are referenced from member properties.

1. In the Tag Editor, let's replace the Compressor Number for each Tag with the new parameter, **CompressorNum**. Select the **lineVoltage** Tag and click on the binding icon and click **Edit**.

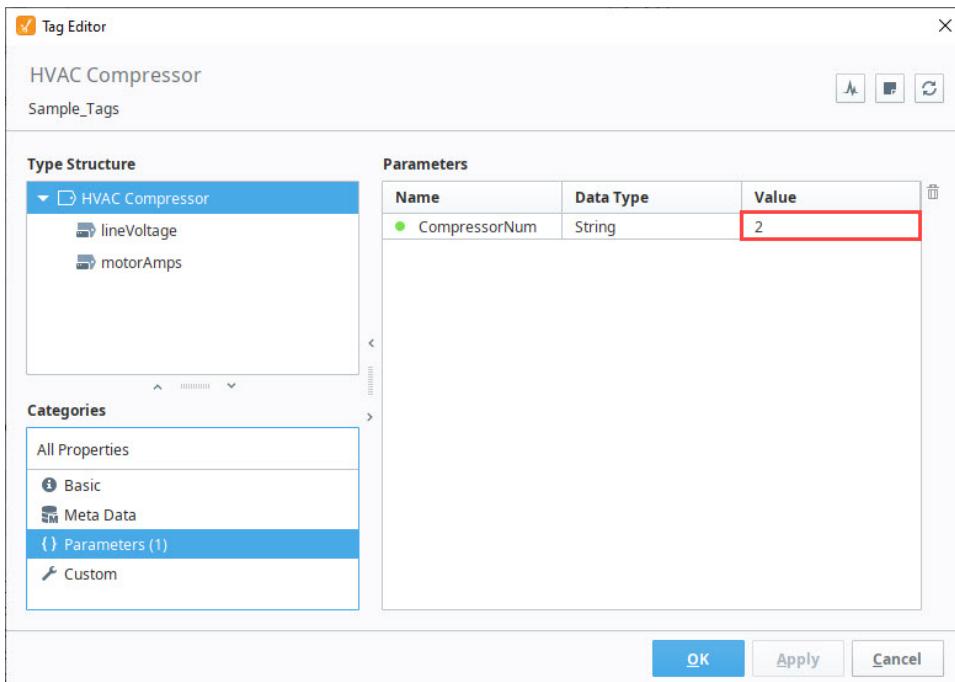


2. This opens the OPC Item Path window for editing the **lineVoltage** Tag. Place your cursor at the end of '**Compressor1**', delete the '**1**', and enter '**{CompressorNum}**'. Click **Commit**.

This feature was changed in Ignition version 8.1.17:

Starting in version 8.1.17, you may also edit tag paths inline in the Tag Editor rather than opening the OPC Item Path window. For this example, you would highlight '1' in the OPC Item Path and replace it with the new parameter, '**{CompressorNum}**'

3. Repeat Step 2 for the **motorAmps** Tag.
4. Both Tags will now show the **CompressorNum** parameter in the OPC Item Path.
5. Now, let's create an instance of the Compressor UDT using the CompressorNum parameter under the **Tags** tab of the Tag Browser. To keep UDT instances organized, we created a Plant Compressors folder.
6. Right click on the Plant Compressors folder and select **New Tag > Data Type Instance > Compressor UDT**.
7. Enter a **Name** for the Instance (i.e., HVAC Compressor), then click on the pencil icon next to the **Parameters** property and enter a value (i.e., **2**). Click **OK** to save the HVAC Compressor instance.



8. Under the **Tags** tab, you'll see the HVAC Compressor was created showing the Parameter that was used and the values for the OPC Tags listed in the OPC Item Path.

Tag Browser

default

UDT Definitions

Tags	UDT Definitions
Motors 1	
Plant Compressors	<ul style="list-style-type: none"> HVAC Compressor <ul style="list-style-type: none"> Parameters <ul style="list-style-type: none"> CompressorNum lineVoltage motorAmps
Plant Motors	

Data Type Parameters in Expressions

It is possible to use the value of data type parameters directly in expression bindings within a UDT. Parameter references can be quickly inserted into an expression.

While a UDT member is selected in the Tag Editor, you can edit bindable properties, such as the Expression on Expression Tags by clicking the **Edit** icon next to the property.

Tag Editor

Member Number

default

Type Structure

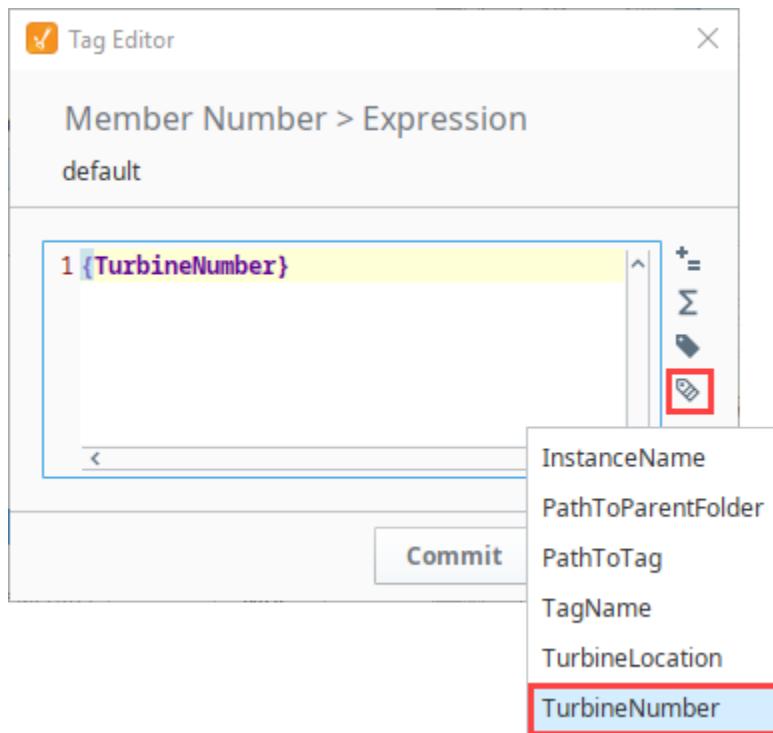
- Turbine
 - Member Location
 - Member Number

Properties

Basic Properties	
Name	Member Number
Tag Group	Default
Enabled	true
Value	
Value Source	Expression
Data Type	Integer
Expression	{TurbineNumber}

OK Apply Cancel

This opens the Expression window. Click on the **UDT Parameters** icon on the right of the expression area, select a parameter, then click **Commit**



Combining Parameters and Tag References

Because parameter and Tag references differ in syntax, some consideration must be made when attempting to use both in the same expression. Tag references must not be placed inside of quotes. After adding a string Tag to the Turbine UDT, a reference to the Tag can be added to Member Location's expression. Single quotes were added to create a space between the Member's Location and the string value.

Tag Editor

Member Location

default

Type Structure

- Turbine
 - Member Location
 - Member Number
 - String Tag
 - Wind Speed

Properties

Basic Properties	
Name	Member Location
Tag Group	Default
Enabled	true
Value	
Value Source	Expression
Data Type	String
Expression	{TurbineLocation} + ' ...'

Tag Editor

Member Location > Expression

default

```
1 {TurbineLocation} + ' ' + {[.]String Tag}
```

Commit Revert

Here's what it looks like in the Tag Browser.

Tag Browser

+ - X

Tags

UDT Definitions

Turbines		Turbine
Turbine10		Turbine
Turbine 100		Document
Parameters		
Member Location	Folsom, CA This is a string value	String
Member Number	100	Integer
String Tag	This is a string value	String
Wind Speed		Boolean

Attribute Referencing and Parameterized Types

As mentioned above, many properties in the member Tag configuration can reference the parameters available in the data type. When instances are created, these references are replaced with the values defined for the type. Parameter references also support basic offsets and numerical formatting, providing a great deal of flexibility. To reference a parameter, use the syntax `{ParameterName}`.

To offset a value, use the form `{ParameterName+offset}`.

To format a value, use the form `{ParameterName|format}`. The format pattern is the same as that used for the `numberFormat` expression function. In short, "0" can be used to require a digit, and "#" can be used for optional digits. ie: `##0`

Example:

For this example, we'll assume that we're parameterizing the **OPC Item Path**, and that the data type has an integer attribute named `BaseAddress` defined. We'll pretend the **OPC Server** provides Tags named like `DataPoint1`.

Standard Referencing

OPC Item Path: `DataPoint{BaseAddress}`

Offset

Imagine that our data type had three fields, and these were laid out sequentially in the device. Instead of specifying each address for each Tag, we can simply offset from the base address:

Member 1: `DataPoint{BaseAddress+0}`
Member 2: `DataPoint{BaseAddress+1}`
Member 3: `DataPoint{BaseAddress+2}`

Formatting

Continuing from the example above, imagine that our OPC server actually provided addresses in the form `DataPoint001`, in order to stay consistent up to "DataPoint999". This can be accommodated using number formatting in the reference:

Member 1: `DataPoint{BaseAddress+0|000}`
Member 2: `DataPoint{BaseAddress+1|000}`
Member 3: `DataPoint{BaseAddress+2|000}`

This format of three zeros means "three required digits". If our instance has a base address of 98, the resulting paths will be `DataPoint098`, `DataPoint099`, `DataPoint100`.

Parameters support more mathematical operators in addition to offsets and formatting. There is a simple expression language available that can be used in conjunction with formatting. The following table shows all available operators in their order of operations (they are evaluated starting at the top of the table).

Operator	Description	Example
<code>()</code>	Parenthesis. These operators are used for grouping any number of values. Also used to change the order of operations.	<code>{Baseaddress*(2+3)}</code>
<code>^</code>	Power. This operator is used to raise a number to a power.	<code>{BaseAddress^2}</code>
<code>-</code>	Negative. Used to create a negative value from one number.	<code>{BaseAddress*-2}</code>
<code>*</code>	Multiplication. Multiply two numbers.	<code>{BaseAddress*2}</code>
<code>/</code>	Division. Dividing the first number by the second number.	<code>{BaseAddress/2}</code>
<code>%</code>	Modulus. This operator returns the remainder of a division operation. IE: $7/3 = 2$ with a remainder of 1, so $7\%3 = 1$	<code>{BaseAddress%2}</code>
<code>+</code>	Addition. Add two numbers.	<code>{BaseAddress+2}</code>
<code>-</code>	Subtraction. Subtract two numbers	<code>{BaseAddress-2}</code>
<code> </code>	Used to define a formatting pattern. Patterns are defined with 0 and # characters.	<code>{BaseAddress##0.00}</code>

Example

```
# This dynamic OPC Item path takes in three parameters to determine the tag path
ns=1;s=[DeviceName]Path/to/tag{BaseAddress+(ParamNum*Multiplier)|0000}
```

```
# The OPC Item path resolves to the following assuming the following values:
```

```
# BaseAddress = 5
# ParamNum = 8
# Multiplier = 2
ns=1;s=[DeviceName]Path/to/tag0021
```

Calculations and Numerical Parameter Names

If the parameter names are purely numerical values (we don't recommend this: it gets confusing), then quotation marks must encase the parameter to run any sort of calculations on the value of the parameter.

For example, if a UDT contains a parameter named 0, and its value is 10:

```
// This will evaluate to 0, because it thinks you mean the integer 0, not the parameter named "0"  
{0 * 1000}  
  
// This will evaluate to 10000, because the quotation marks denote a parameter named "0"  
{"0" * 1000}
```

Nonexistent UDT Parameters

This feature was changed in Ignition version 8.1.8:

Parameter references in bindings that do not evaluate to an existing UDT instance parameter will be returned as a string.

For example, say a binding on a member used the following:

```
SomeText / {myParameter}
```

If the UDT definition does not contain a parameter named "myParameter", then the binding above would be treated as a string literal, returning a value of "SomeText/{myParameter}".

Null Values on Parameters

This feature was changed in Ignition version 8.1.8:

If the value of a UDT parameter is null, then any bindings that reference to the value of that parameter will instead return a string literal, similar to how referencing nonexistent parameters works.

For example, say a UDT definition has the a parameter named "NullParameter". If the value of that parameter is null, then references to it would return the string "{NullParameter}". Thus, if we had a binding that included the parameters:

```
SomeText / {NullParameter}
```

The binding would return a value of:

```
"SomeText / {NullParameter}"
```

UDT Multi-Instance Wizard

The Multi-Instance Wizard provides a powerful, but simple mechanism for rapidly generating many instances of a UDT at the same time by specifying patterns for UDT parameters.

Value Patterns and Tag Names

Value Patterns

In order to define values for parameters (and the Tag names), you can use several different types of patterns (and combinations of patterns):

Range number1-number2[/step]

A numeric range of values, such as 1-10. Optionally, a step parameter can be included, in order to only generate numbers at certain multiples. For example, 0-100/10 would generate 0,10,20, and so on.

Repeat value*count

A value (numerical or string), and the number of times to use it. For example, North Area*10 would use the parameter North Area for 10 items.

List value1, value2, value3

A comma separated list of values (or other patterns) to use.

Examples:

1-10,21-30,31-40 Results in 30 Tags being created, with the specified value ranges (so, for example, there would be no parameter 15).

A,B,C Results in 3 Tags, with each of the values.

0-100/5 Results in 21 Tags (because range is inclusive), with values 0, 5, 10...100.

As mentioned, the size of the pattern will dictate how many Tags will be created. If some patterns are smaller than others, the last value will be repeated for the other Tags.

Tag Names

The names of the generated instances can be specified using a system similar to that of the parameter patterns. If you just want to use sequential names, you don't need to specify a pattern, as values will be generated automatically starting at one. You can also set the pattern to simply be the starting number to generate sequential names from there.

Base Name

A string base for the Tag name. This can also be a list of names, in which case the names will be used directly, and the name pattern won't be used.

Name Pattern

A pattern that will be used to generate values that will be appended to the base name.

At any time, you can use the **Preview** button to view the Tag names and parameters that will be created. Once you are satisfied, click **OK** to generate the Tags under the selected folder in the Tag provider.

How to Make New Instances of a UDT

Once you have a UDT Definition created, you can make multiple instances of it with the Multi-instance Wizard.

1. In the **Tag Browser**, click the Add  icon to create a new folder. We created one called Machine Motors.
You can create a single instance of the motor or use the Multi-instance Wizard to rapidly create many instances at the same time. In this example, we will use the Multi-Instance Wizard.
2. Right-click on the **Machine Motors** folder, and select **Multi-instance Wizard** to open the Instance Creation Wizard window.

On this page ...

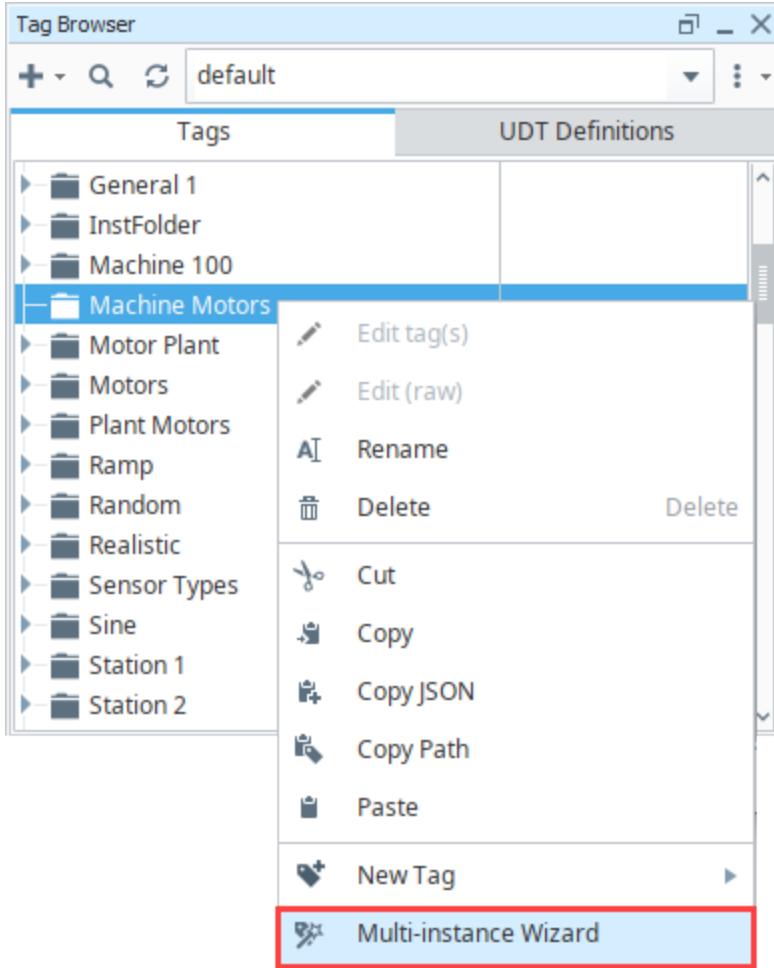
- [Value Patterns and Tag Names](#)
 - [Value Patterns](#)
 - [Tag Names](#)
- [How to Make New Instances of a UDT](#)



INDUCTIVE
UNIVERSITY

UDT Multi-Instance Wizard

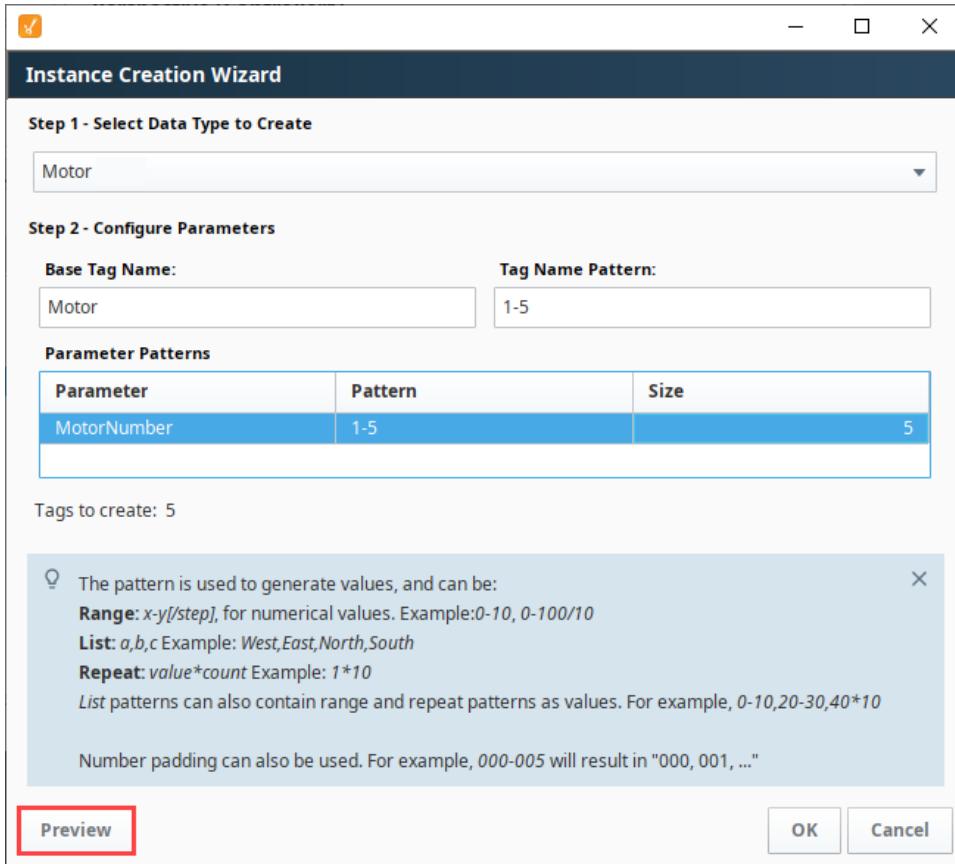
[Watch the Video](#)



3. In **Step 1 - Select Data Type to Create**, select a UDT (i.e., **Motor**) from the dropdown.
4. In **Step 2 - Configure the Parameters**, enter the following for your Motor:

- Base Tag Name: "Motor." Note the space at the end. Without this space your Tag names will look like Motor1, Motor2, etc.
- Tag Name Pattern: 1-5 This creates five Tags **Motor 1**, **Motor 2**, thru **Motor 5**.
- Parameter Patterns: the **MotorNumber** parameter is entered by default when we selected the data type to create in Step 1.
- Pattern: **1-5** is the pattern of the parameter so the Motor 1 Tag will have a parameter of 1, Motor 2 will have a parameter of 2, and so on through Motor 5.

You'll notice that after you enter the Pattern, the number of Tags to create is updated. In this example, five Motor Tags will be created. Click **Preview**.

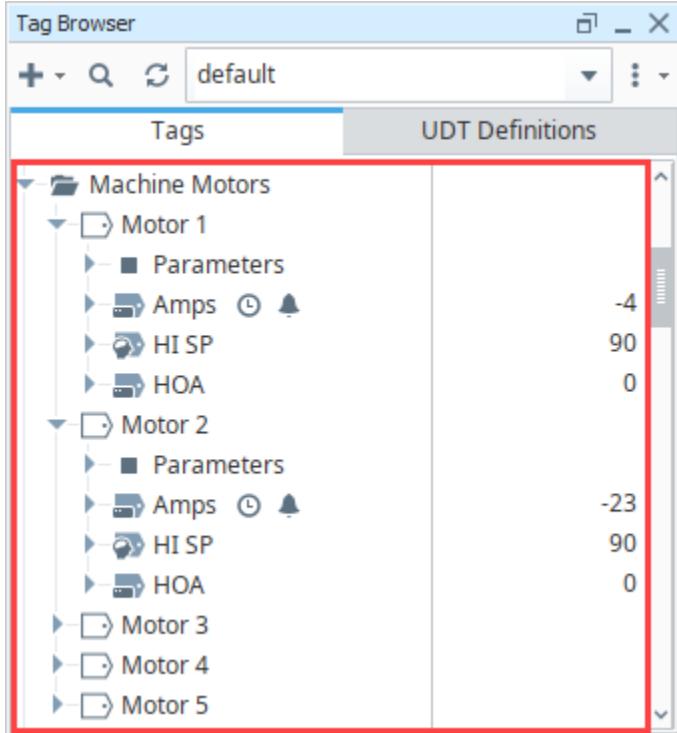


5. In Preview, you will see how the Base Tag Names and Parameter Values get created. Click **Back** to go back to the Instance Creation Wizard window if you want to make an update. If you like what you see on the Preview window, click **OK**.

Tag Name	MotorNumber
Motor 1	1
Motor 2	2
Motor 3	3
Motor 4	4
Motor 5	5

At the bottom are 'Back', 'OK' (highlighted by a red box), and 'Cancel' buttons.

6. In the Tag Browser, expand Motor Tags 1-5 to see if all the members of the Motor UDT were created and are running.



Cannot Edit Existing Instances using the Multi-Instance Wizard

You cannot edit existing instances using the Mult-Instance Wizard. The Mult-Instance Wizard is only used for quickly creating many instances of a UDT at the same time. If you want to make a change to all your instances, refer to [How to Edit an Existing UDT](#).

Related Topics ...

- [UDT Parameters](#)

UDT Inheritance

Once you have a single data type created, it is possible to set up UDT inheritance where data types extend to other data types, to add additional members, or override default values. For example, you can create a new data type and using the inheritance feature it will inherit all Tags from the parent data type including the parameters. Then you can add additional Tags and/or override any settings in your new data type. UDT Inheritance is a way to extend to a class of data types to add more functionality to that class.

For example, you may have a simple motor and a complex motor. The complex motor can inherit from the simple motor, which means all simple motor values will be in the complex motor and you can add more.

Nesting (using one or more UDTs to make up a larger UDT) is different from inheritance and can be found under [UDT Nesting](#).

This feature was changed in Ignition version [8.1.17](#):

In 8.1.17, the Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts.

Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

On this page ...

- [Inheriting Property Values from an Existing UDT](#)
- [Creating the Data Type Instance](#)
- [Overriding Properties of the Parent UDT](#)
 - [UDT Inheritance Traits](#)
 - [UDT Hierarchy Tool](#)



INDUCTIVE
UNIVERSITY

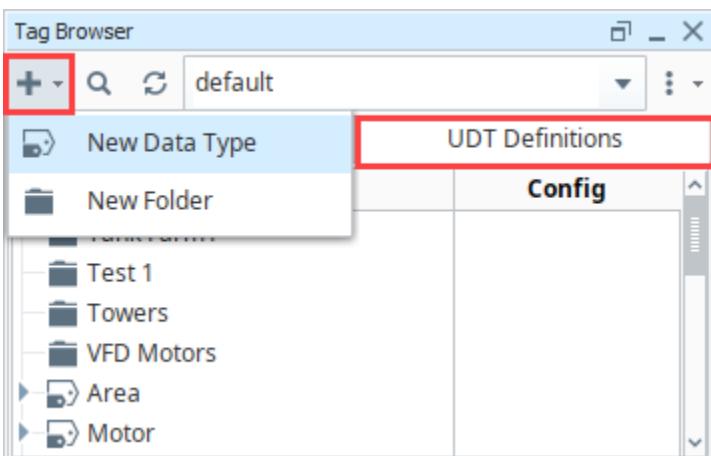
UDT Inheritance

[Watch the Video](#)

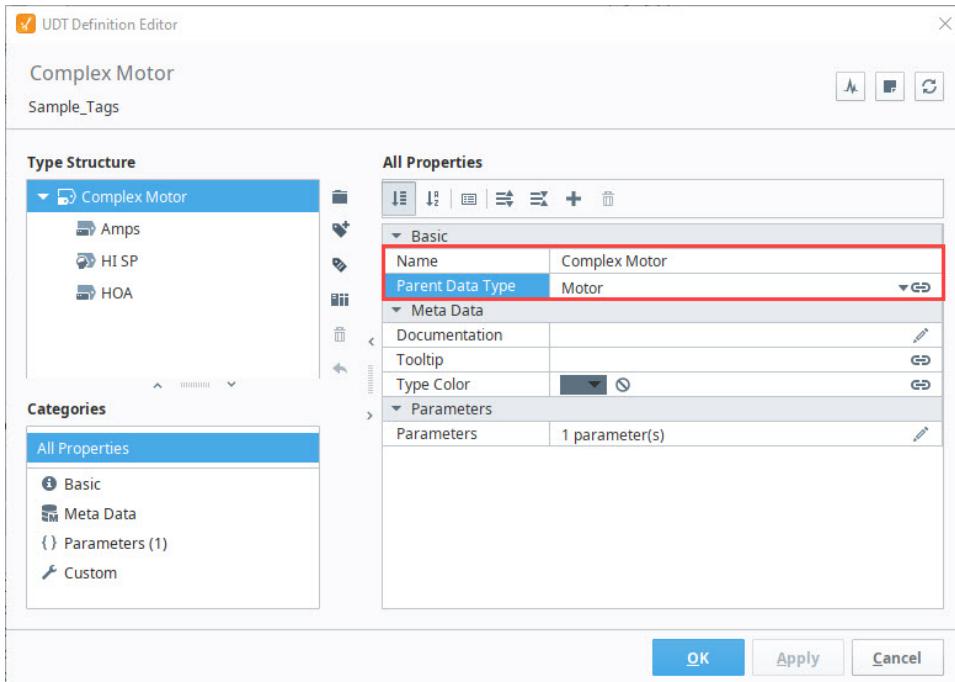
Inheriting Property Values from an Existing UDT

Let's use our data type Motor from the previous sections to create another data type. We'll set the parent to Motor so our new data type automatically inherits all the properties of Motor.

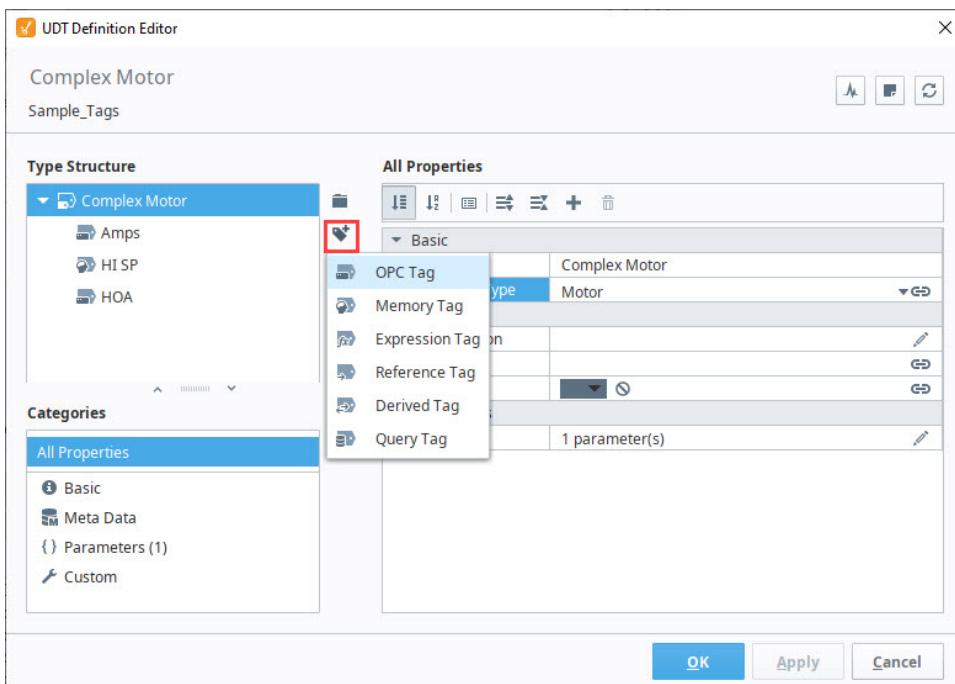
1. In the **Tag Browser**, click on the **UDT Definitions** tab and click the Add **+** icon to create a **New Data Type**. The Tag Editor window will open on the.



2. In the **Name** field, enter name for your new UDT (i.e. Complex Motor). Under **Parent Data Type** property go to the dropdown and select the **Motor** data type and click **Apply**. Now your new Complex Motor UDT is inheriting from all the properties of the parent Motor UDT: Amps, HI SP and HOA.



3. With the Tag Editor still open, let's add an **OPC Tag** to the Complex Motor UDT. Click on the **Add Tag**  and select **OPC Tag**.



4. Enter the following properties for your new Tag and click **Apply**. You will see the new Tag was added to the Complex Motor UDT.

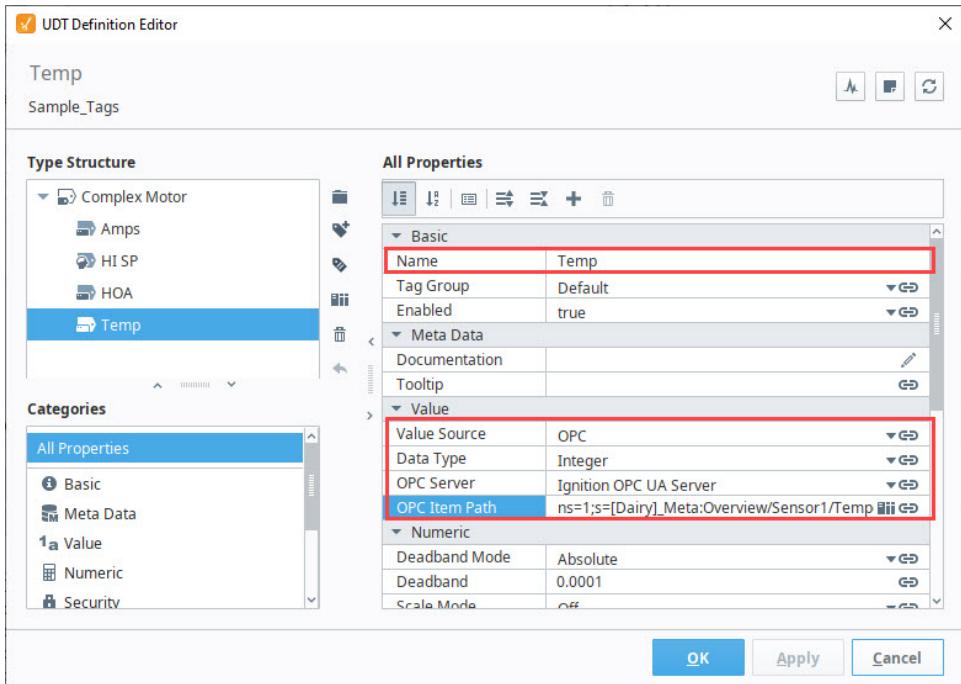
Name: Temp

Value Source: OPC

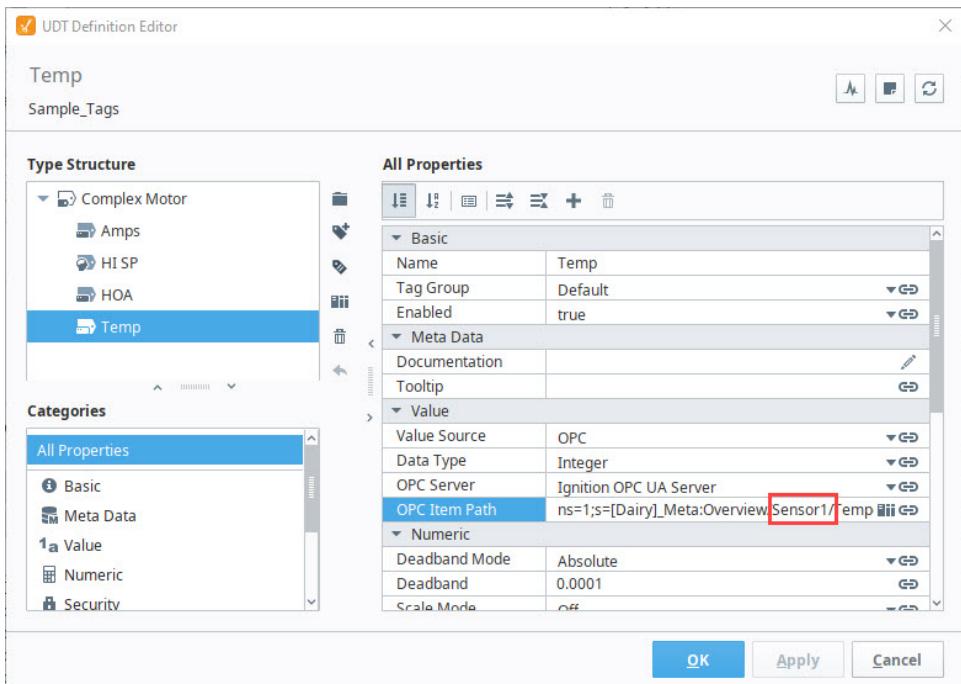
Data Type: Integer

OPC Server: Click on the binding () icon and select **Ignition OPC UA Server**

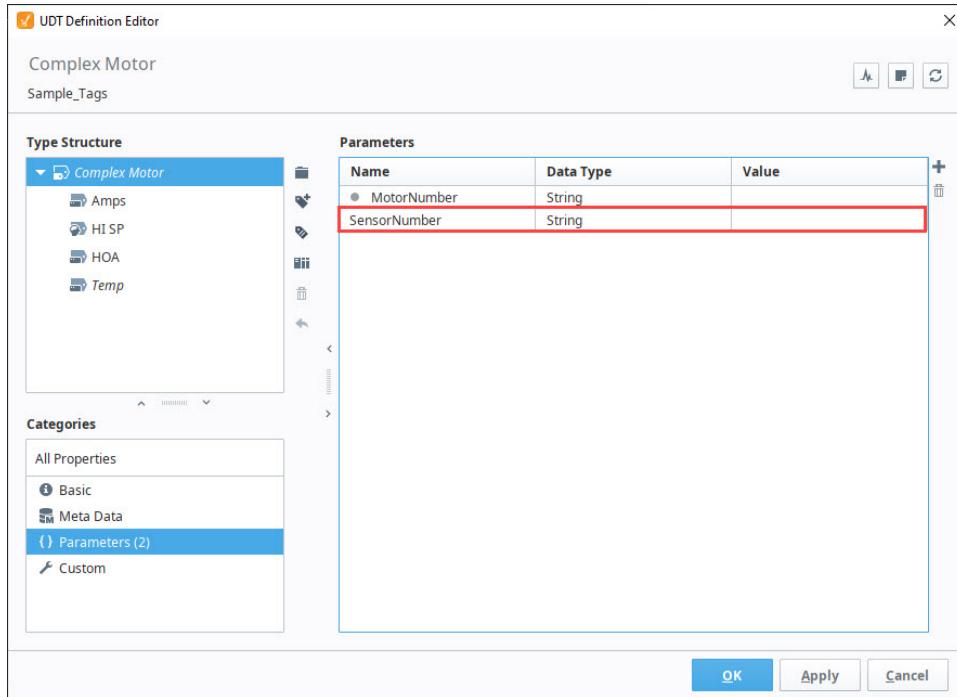
OPC Item Path: Browse the [OPC](#) and find the Tag you want to use. This example uses a **Temperature Tag** from a Sensor in the Dairy program.



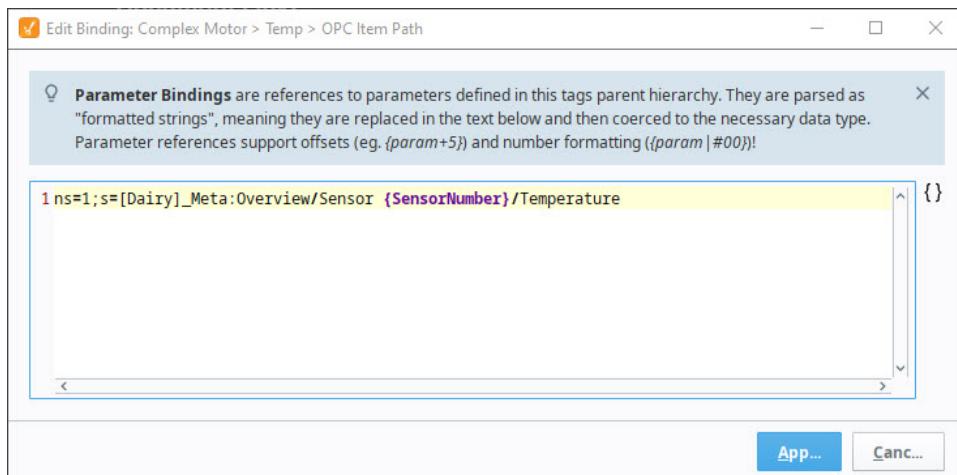
5. In the image below the Temp Tag is pointing to a specific address in the PLC. Because we're creating a new Tag in our UDT, we don't want to point to one specific set of 'Temp' Tags. We want each instance of the Complex Motor UDT to reference a different set of 'Temp' Tags. To do that, we need to add a parameter to the Complex Motor data type that we will call 'SensorNumber'.



6. With the Complex Motor UDT **Tag Editor** open, let's create a new UDT parameter. Right click the **Edit** icon next to the **Parameters** property. The Parameters pane will open.
7. Click the Add **+** icon and add the new parameter, '**SensorNumber**'



8. With the Tag Editor still open, select the "Temp" Tag. In the **OPC Item Path** field, click the **binding** icon, select **Edit**, and the **Temp > OPC Item Path** window will open. Place your cursor at the end of '**Sensor1**', delete the '**1**', add a space, and enter '**{SensorNumber}**'. Don't forget the curly braces. Click **Apply** to save your updates and go back to the previous window.



This feature was changed in Ignition version 8.1.17:

Starting in version 8.1.17, you may also edit tag paths inline in the Tag Editor rather than opening the OPC Item Path window. For this example, you would highlight '1' in the OPC Item Path and replace it with the new parameter, '{SensorNumber}'

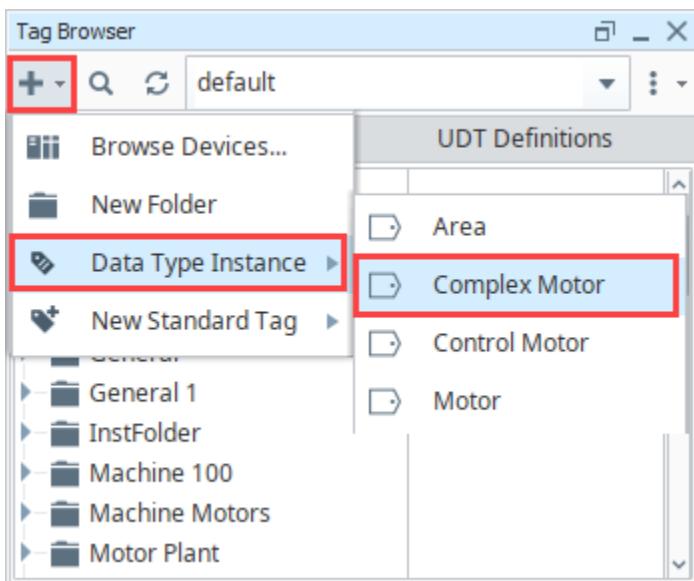
9. In the **Tag Editor** window, click **OK** to save your UDT. In the Tag Browser, the new data type is now visible in the , "Complex Motor." You can see all the inherited Tags from the Motor data type and the newly added Temp Tag.

Tags		UDT Definitions
Definition	Config	
Tank Farm1		
Test 1		
Towers		
VFD Motors		
Area		
Complex Motor		
Amps	Ignition OPC UA Ser...	
HISP	90	
HOA	Ignition OPC UA Ser...	
Temp	Ignition OPC UA Ser...	
Control Motor		
Motor		

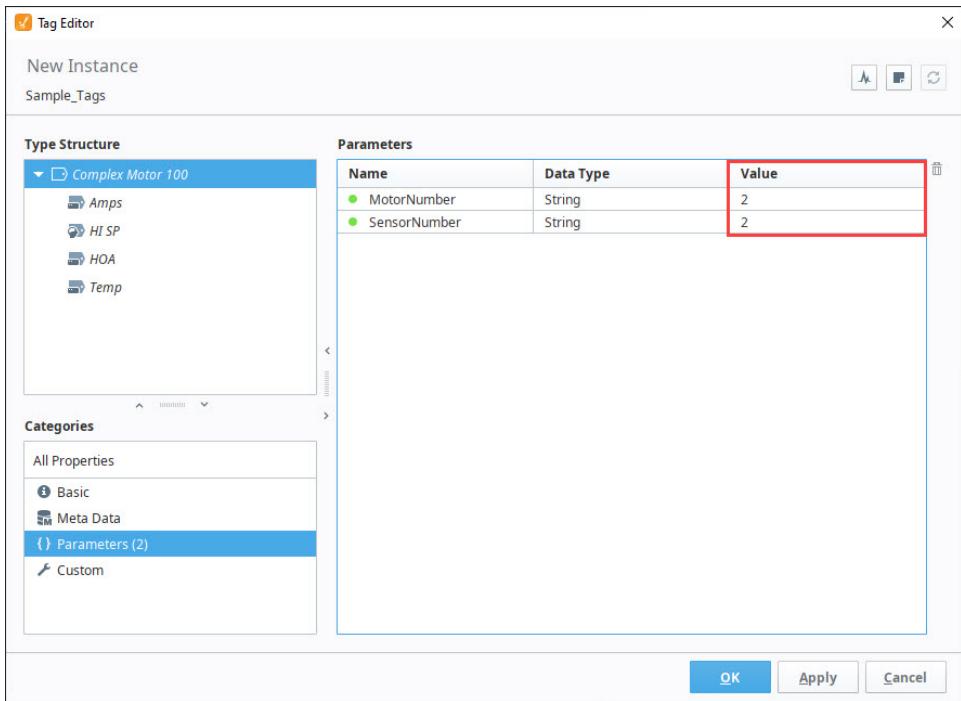
Creating the Data Type Instance

Now that our Complex Motor UDT is set up, let's create a data type instance of the Complex Motor.

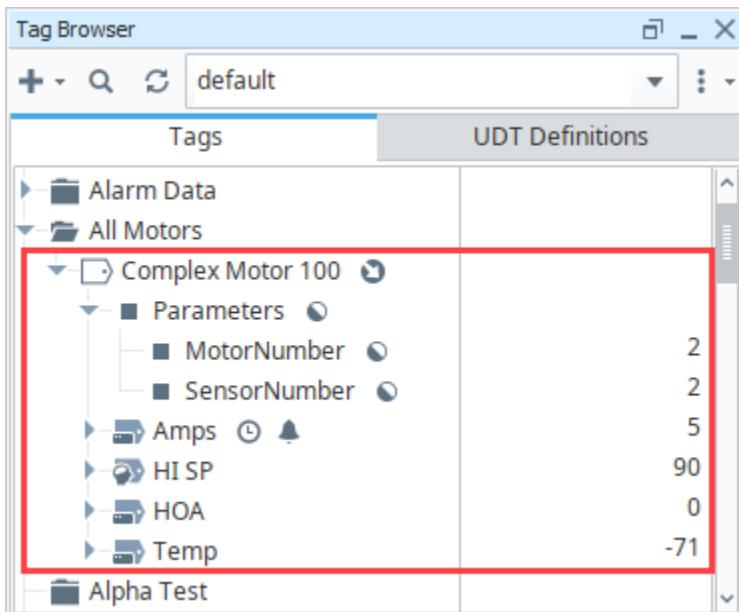
1. Click on the **Tag** tab, and from the context menu, select the **Data Type Instance > Complex Motor**.



2. Enter the **Name** for the instance (i.e., Complex Motor 100). Click the **Edit** icon next to the **Parameters** property and enter the parameter value of '2' for **MotorNumber** and **SensorNumber**.



3. Click **OK** to create the instance. Now, you'll be able to see all the values for Complex Motor 100 including the Temp Tag that was added.

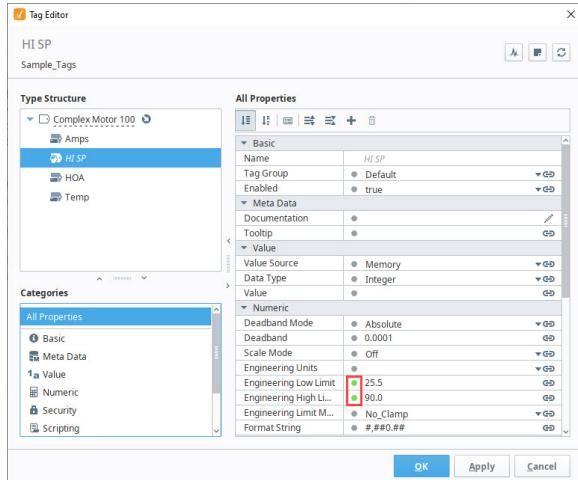


Overriding Properties of the Parent UDT

Another benefit of the UDT inheritance feature is it allows you to override some of the properties of the parent. For example, since the Complex Motor has Motor as the parent, you can go to any of the Tags and override any of the settings of that data type. Click the circle to the right of the property and enter a new value, or change a property's value and the green circle changes to green automatically. This overrides the property inherited from the parent.



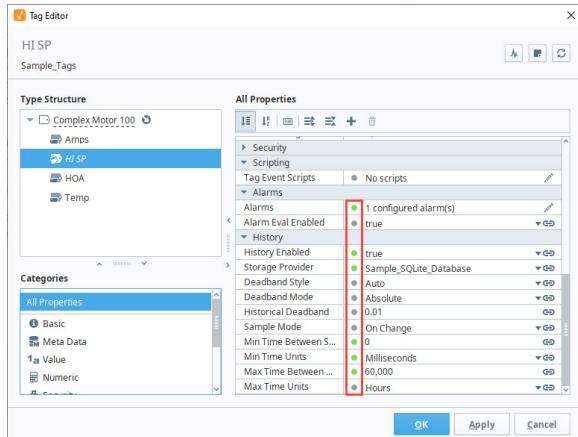
**Overriding
Properties in UDT
Instances**



[Watch the Video](#)

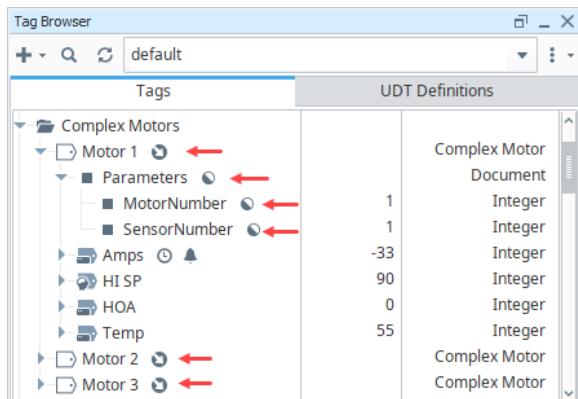
You can also turn on Alarming and History that wasn't initially turned on in the parent UDT by simply using the override feature. Next to the Alarm property, click the green circle to change it to green, and

click the pencil () icon to configure the alarm if it is not already configured. If you want to turn on History, click the green circle or change any of the History properties which will cause any of the green circles to change from gray to green.



UDT Inheritance Traits

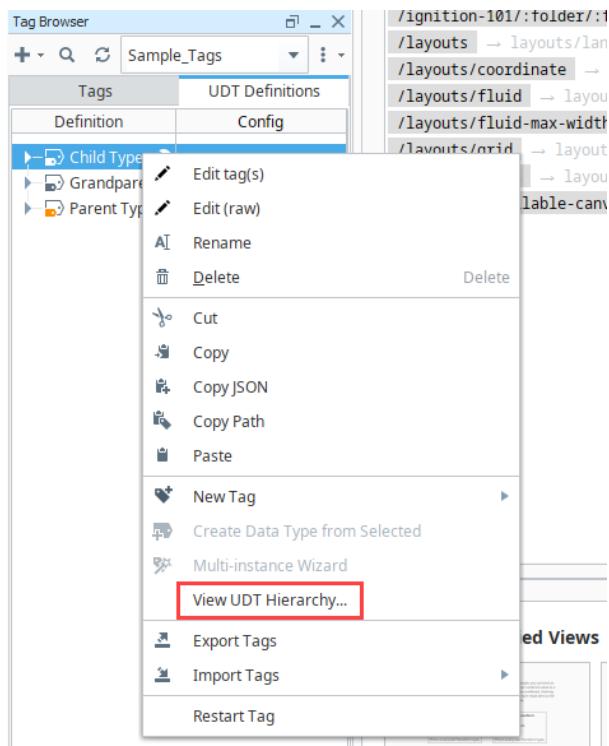
UDT instances that are inherited and have properties that override the parent properties have a visual representation next to them.



UDT Hierarchy Tool

The following feature is new in Ignition version 8.1.23
[Click here](#) to check out the other new features

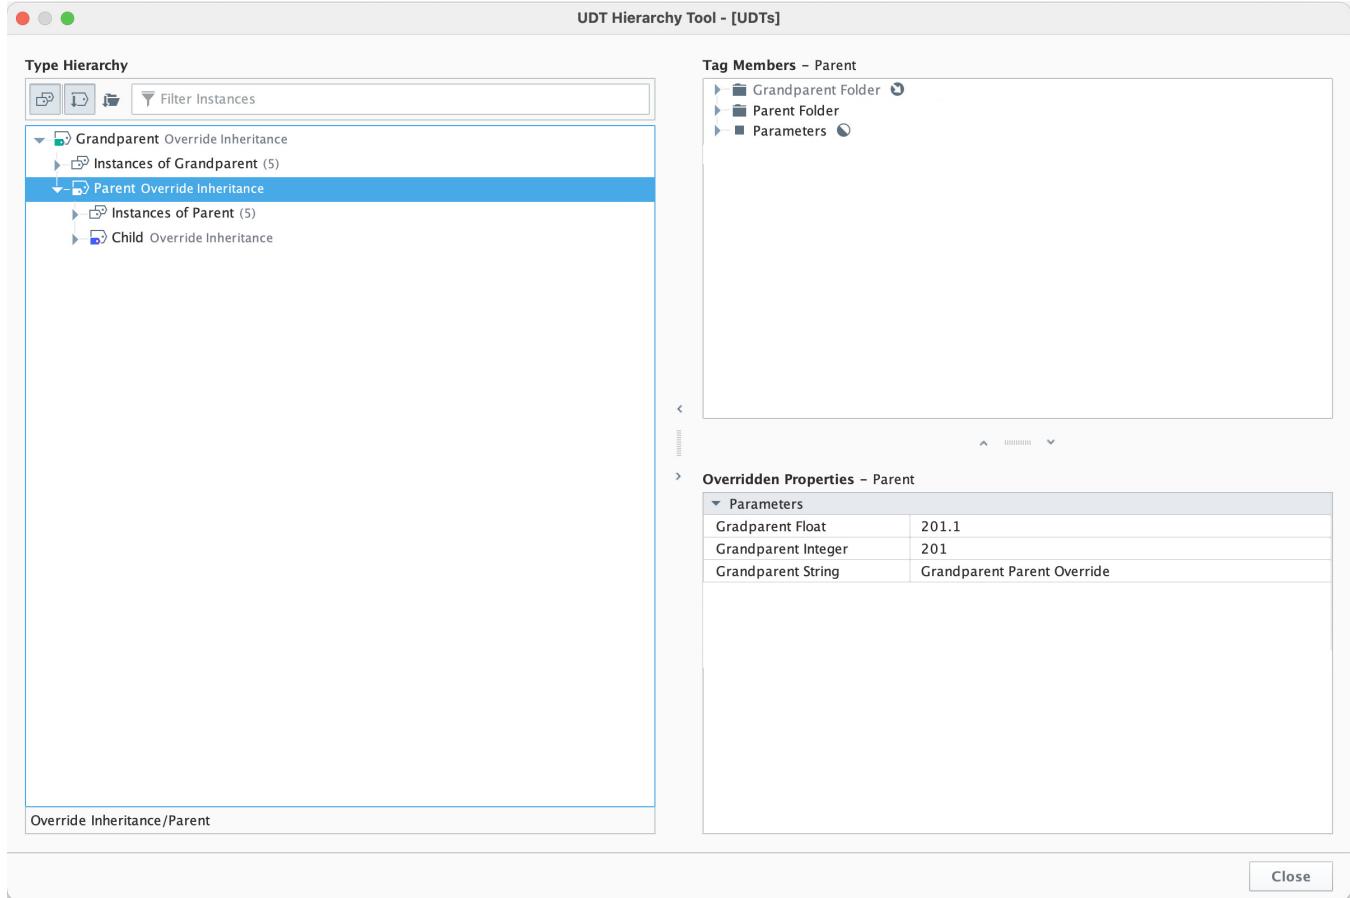
The UDT Hierarchy Tool displays the relationships between connected UDT definitions and instances. It shows inheritance between UDT definitions, overrides of inherited properties, and instances directly related to selected definitions. Right click a UDT Definition or UDT Instance and select **View UDT Hierarchy** to access the UDT Hierarchy Tool.



The Type Hierarchy panel shows the related definitions. When you have the Show Instance icon selected, you can sort instances by typing into the filter bar and choose between grouping by folder path or name with the respective or icons. If you expand the **Instances of** folders, you will see UDT instances that directly inherit the definition as a parent. Note, nested UDT Instances are not visible in this list.

If you need additional filtering for instances, such as specific overrides, right click the **Instances of** text for the specific UDT instances to open the [Tag Report Tool](#) with the search pre-populated for that specific set.

Information about member tags and overridden properties for the selected UDT Definition or UDT Instance are visible on the right side of the tool. As shown in the example below, the inherited tag member types will be grayed out and followed by the Inherited icon.



Related Topics ...

- [UDT Parameters](#)
- [UDT Nesting](#)

UDT Nesting

It's possible to set up UDT nesting in Ignition where you are putting one UDT inside of another UDT. The UDT is nested as an instance within another UDT. It facilitates quicker development of projects since you're able to piece together multiple UDT definitions as needed without having to build everything from the ground up as with each UDT definition. This is particularly useful because it promotes rapid development if you are expanding a plant or facility where all you have to do is make a few Tag changes to existing parameters and property settings.

For example, you may have a production line that is built out of several different machines. You don't need to re-create a motor for each line, instead you can create it once and use it in every line.

Inheritance (having simple and complex version of similar objects) is different from nesting UDTs, and can be found under [UDT Inheritance](#).

On this page ...

- [Set Up UDT Nesting](#)



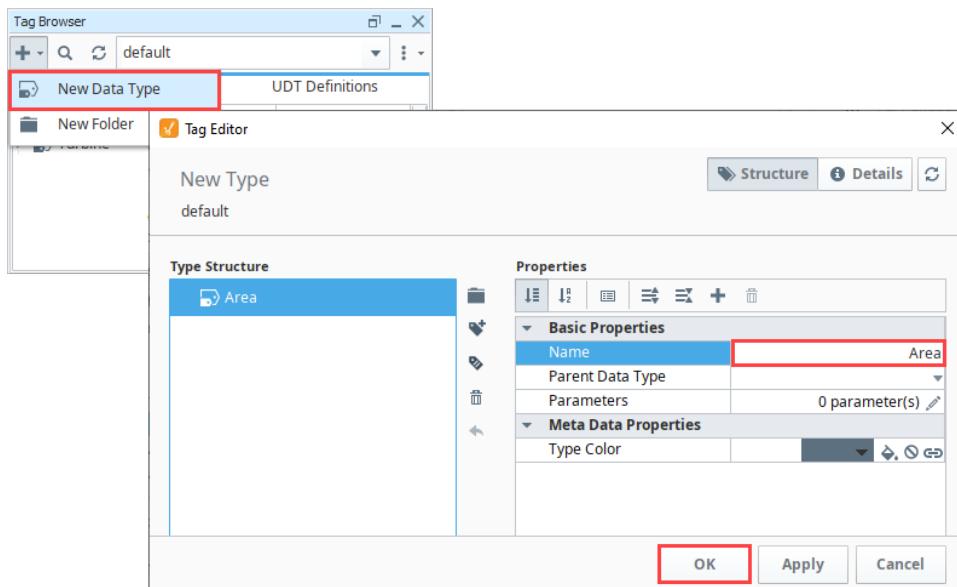
Nested UDTs

[Watch the Video](#)

Set Up UDT Nesting

In this example, let's use our Motor and Sensor data types that were created and used in previous sections of this manual. We are going to create a third UDT called Area that will contain the Motor and Sensor data types inside of it.

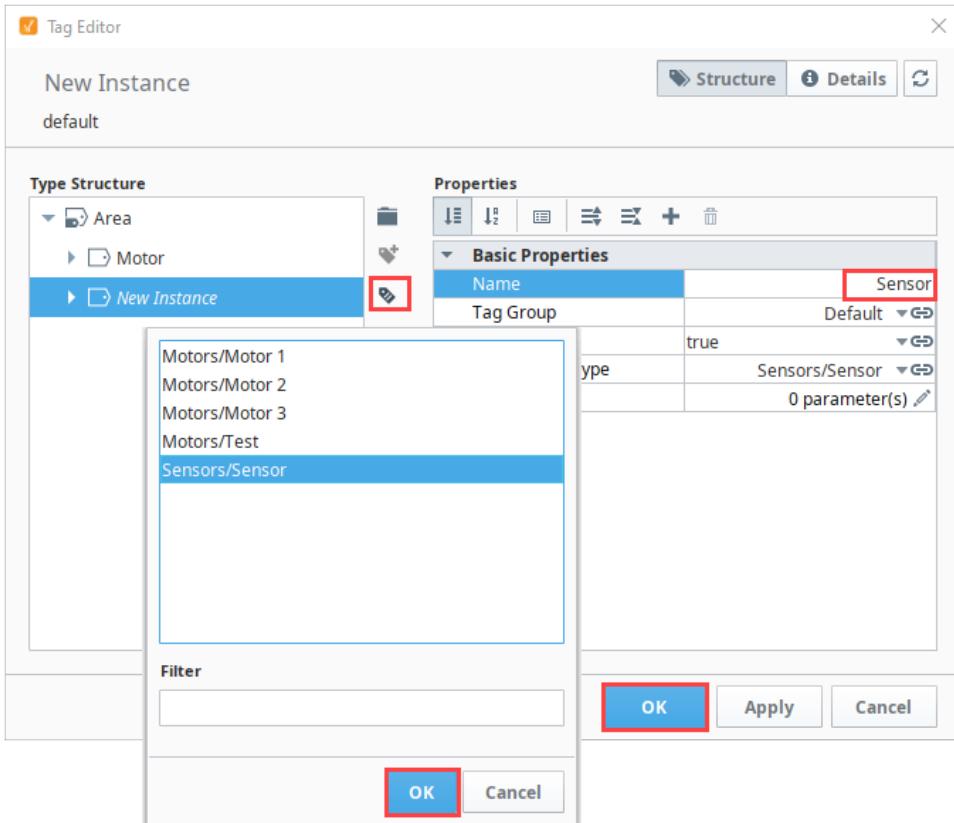
1. In the Tag Browser, right click on **Data Types** and select **New Tag > New Data Type**. The Tag Editor window will open. Assign the new data type a Name called **Area**.



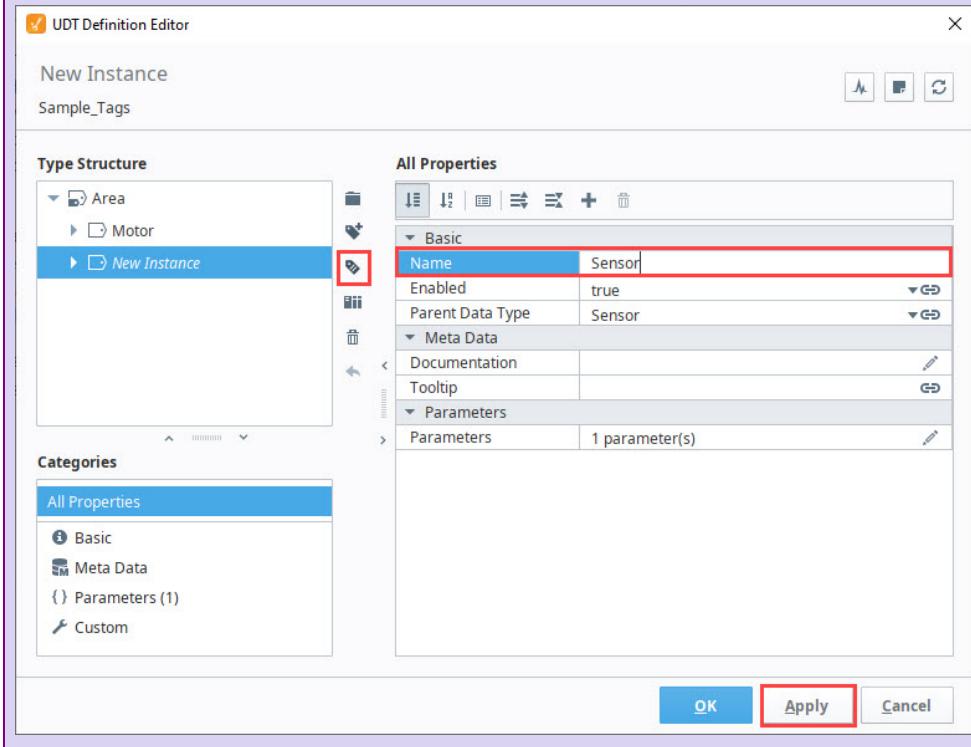
*Editor notes are only visible to logged in users
On 8.1.17 release, delete this screenshot*

2. Inside of the **Area** data type, create two data type instances, one for Motor and the other for Sensor:
 - a. Click on the **New UDT Instance** icon, select **New UDT Instance > Motor/Motor1**. Click **OK**.
 - b. Rename the new instance '**Motor**'. Click **Apply**.
 - c. Click on the **New UDT Instance** icon, select **New UDT Instance > Sensors/Sensor**. Click **OK**.

- d. Rename the new instance 'Sensor'. Click Apply.

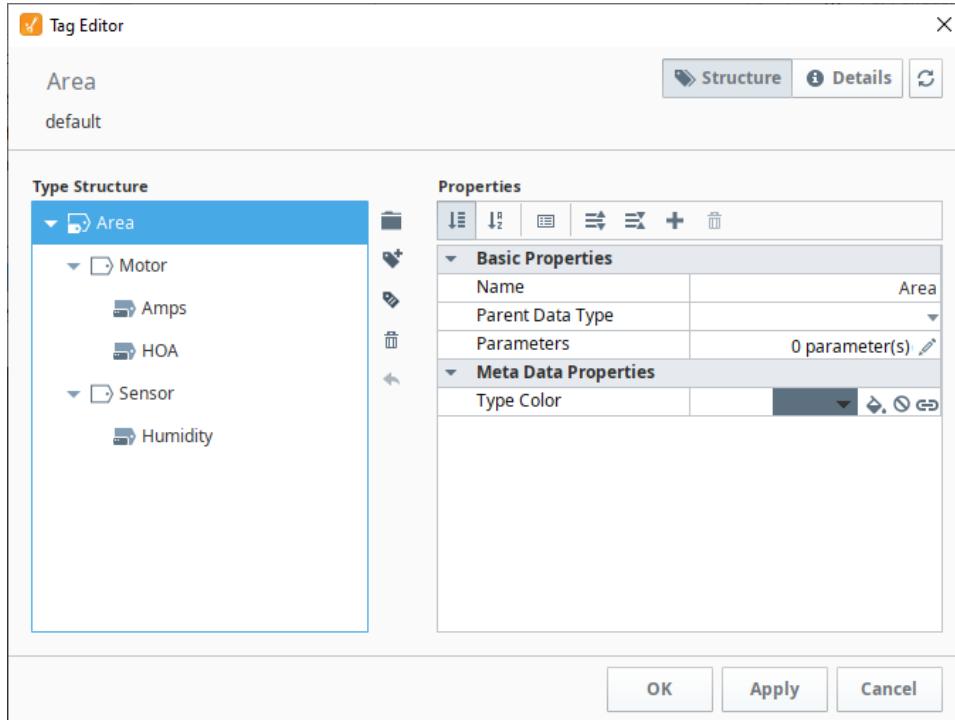


*Editor notes are only visible to logged in users
On 8.1.17 release, replace this screenshot*

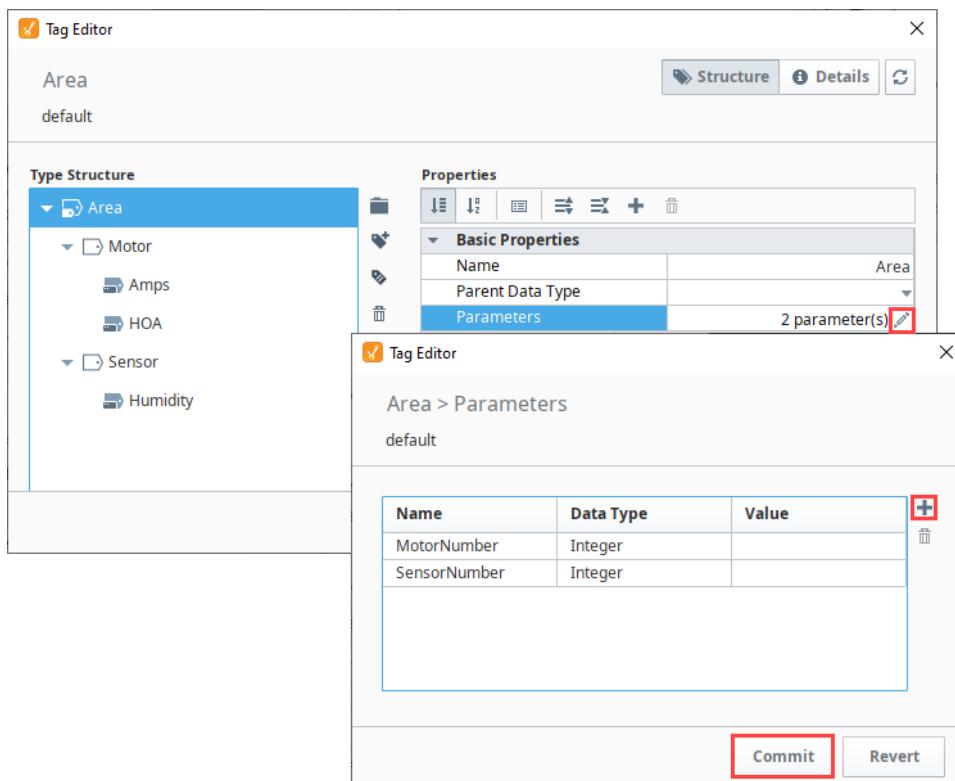


3. With the Tag Editor still open, you'll notice that both the Motor and Sensor UDTs were added. For every UDT that you add inside another UDT, those UDT instances have parameters that need to be specified. In this example, the Motor UDT has the 'MotorNumber' parameter, and the Sensor UDT has the 'SensorNumber' parameter. You must pass a value into the each of these UDTs (Motor and Sensor) from the

parent UDT (Area). To view the parameters for each UDT instance, select each UDT instance (Motor and Sensor) and click the **Edit**  icon.



4. Now that you know what parameters are in each UDT instance, go to the Area UDT, and click the **Edit**  icon next to the **Parameters** property. The Parameters window will open.
5. Click on the **Add**  icon to add the 'MotorNumber' and 'SensorNumber' parameters, then click **Commit**.

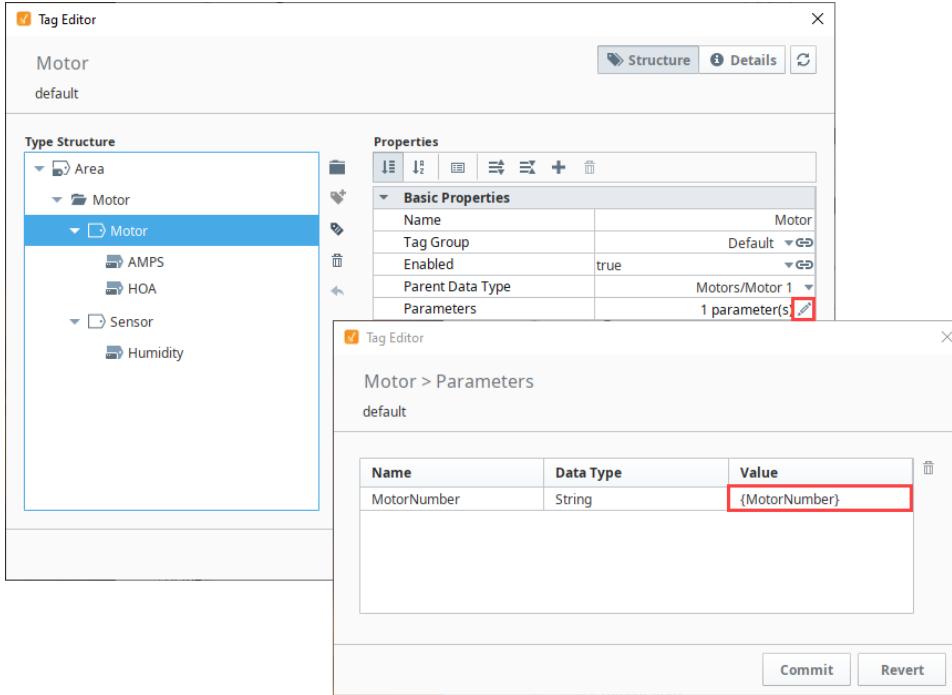


6. Now we need to pass these values into the UDT instances by adding a reference. Select the **Motor** UDT, and click the **Edit**  icon next to the **Parameters** property.

7. Enter the reference for MotorNumber: '{MotorNumber}'. Click **Commit**.

i If Multiple Data Types use the same Parameter

In the event data types use the same parameter, you only need to enter it once in the new data type (i.e., Area).



8. Do the same steps to add a reference to the **Sensor** UDT. Select the Sensor UDT and click the **Edit** icon next to the **Parameters** property. Enter the reference for SensorNumber: '{SensorNumber}'. Click **Commit**.
 9. Click **OK**.

Related Topics ...

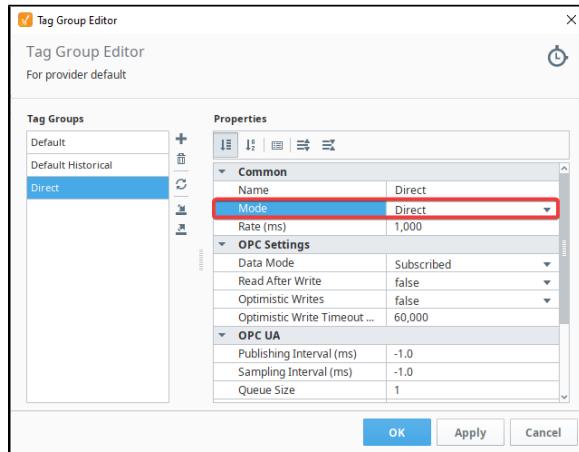
- [UDT Parameters](#)
- [Data Type Parameters in Expressions](#)

Tag Groups

What Is a Tag Group?

Tag Groups dictate the rate of execution of Tags, and therefore play a crucial role in the design of large and high-performance systems. Highly optimized systems generally use multiple Tag Groups which allows for each data point to be polled only as fast as it needs to be. For example some Tags may need to be updated every 500ms, while others may only need to be updated every 30 seconds.

Creating different Tag Groups allows you to organize your Tags into groups that subscribe or poll at different rates.



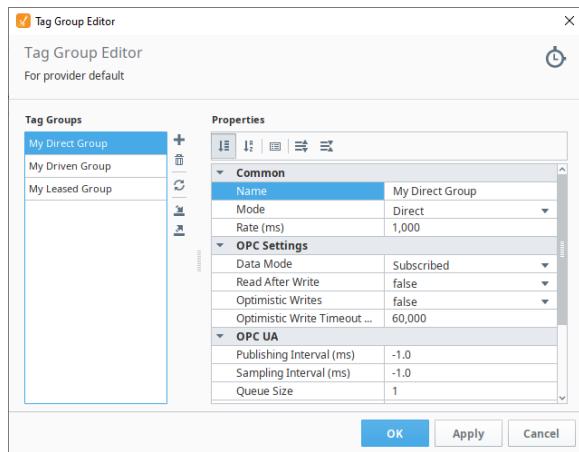
Tag Groups Modes

Tags are assigned to a Tag Group that determine how often they update. For example, how often an OPC value is polled from the PLC, how often an Expression Tag calculates its expression, and how often a Query Tag runs its query. It's easy to create Tag Groups in Ignition for just about any scenario you can think of. Tag Groups are extremely powerful and flexible, and you can create them based entirely on your individual business requirements.

There are three different Tag Group modes in Ignition that you can use. Each mode works a bit differently.

Direct

Executes at one fixed rate, as determined by the **Rate** property in milliseconds. This is the most simplistic option, as the rate of the group doesn't intentionally change between two different rates. The default tag group on newly installed Ignition systems uses this mode. It's the simplest mode in that there is only a single rate for this mode.



See the [Direct Tag Group Example](#) page on steps to add a Direct Tag Group.

On this page ...

- [What Is a Tag Group?](#)
- [Tag Groups Modes](#)
 - [Direct](#)
 - [Driven](#)
 - [Leased](#)
- [Adding and Editing Tag Groups](#)
- [Setting a Tag Group on a Tag](#)
- [Tag Group Properties](#)



INDUCTIVE
UNIVERSIT

Tag Group Overview

[Watch the Video](#)



INDUCTIVE
UNIVERSIT

Driven Tag Group

[Watch the Video](#)



INDUCTIVE
UNIVERSIT

Leased Tag Group

[Watch the Video](#)

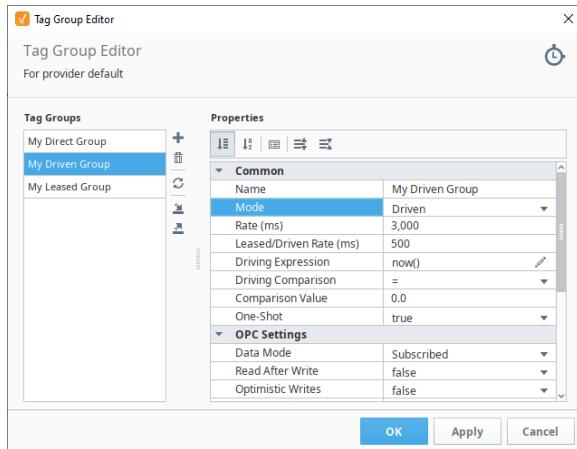
Driven

Driven Tag Groups switch between two different rates based on a condition you configure directly onto the Tag Group. The condition in this case is an **expression** on the **Driving Expression** property. The result of the expression is then compared to the **Comparison Value** property. If the comparison of the two values is equivalent, then the entire tag group executes at the **Leased/Driven Rate** in milliseconds, otherwise it executes at the **Rate**.

Thus, you can have a tag group that switches between a fast and slow rate under certain conditions.

Furthermore, Driven Tag Groups have a **One Shot** property. When set to true, means the group only executes every time the comparison is equivalent. Furthermore, it only executes once every time that comparison transitions from false to true: this is sometimes called "a rising edge".

The One Shot property allows for more interesting tag executions. Instead of running at set rates, you can trigger an execution for all tags on the group by incorporating a tag value as part (or all) of the driving expression, allowing a single tag value change to cause many other tags to update.



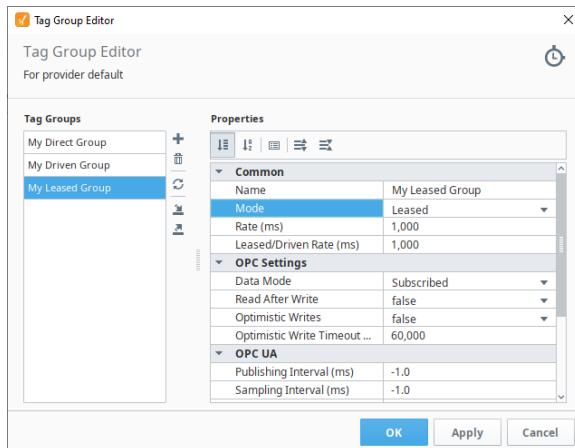
See the [Driven Tag Group Examples](#) page on steps to add a Driven Tag Group.

Leased

Leased groups have two different execution rates, much like a Driven group. However Leased groups don't rely on a comparison to determine rate. Instead, the driving mechanism is whether a Tag is being displayed on an open window or view: generally via a Tag Binding of some sort.

Tags that are displayed (leased) will run at the **Leased/Driven Rate**, while those Tags on the same group that aren't leased will run at the **Rate**.

Leased groups are unique in that Tags on the same group may execute at different rates while being in the same group. For example, if Tags **A** and **B** are both on the same Leased group, and all Designer, Vision Client, and Perspective Sessions are closed, then both tags execute using the **Rate** value. If a user launches a Perspective Session and switches to a view where the value of **A** is displayed on a component binding, then **A** will switch to the **Leased/Driven Rate**. However, since **B** isn't displayed anywhere, it will continue to execute using the **Rate** value.



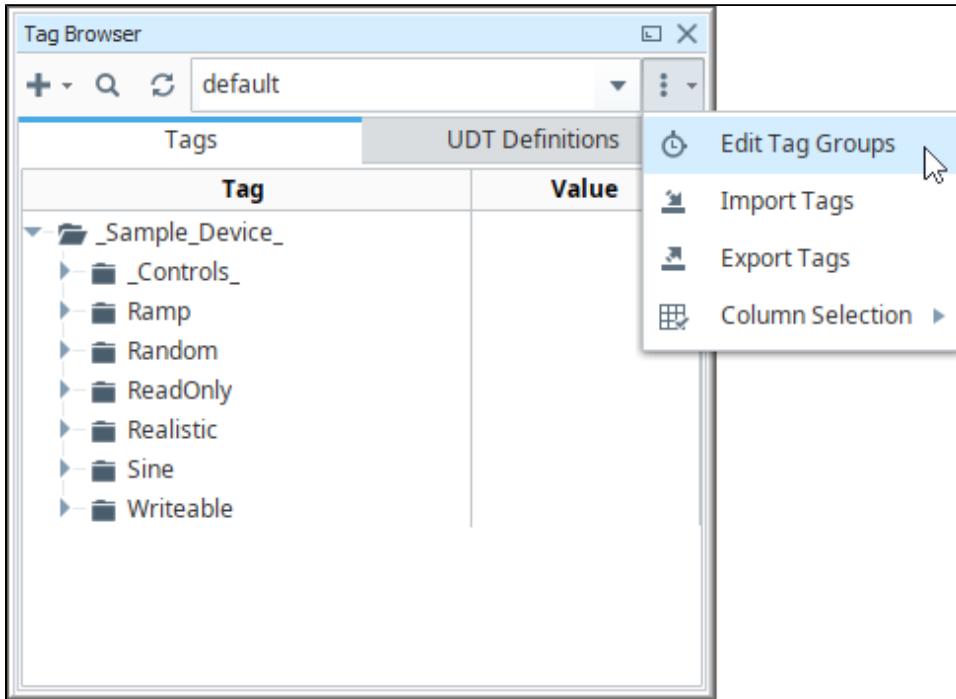
See the [Leased Tag Group Example](#) page on steps to add a Leased Tag Group.

Note: Viewing the Tag only in the Tag Browser will not cause the Tag Group to execute at the Leased Rate.

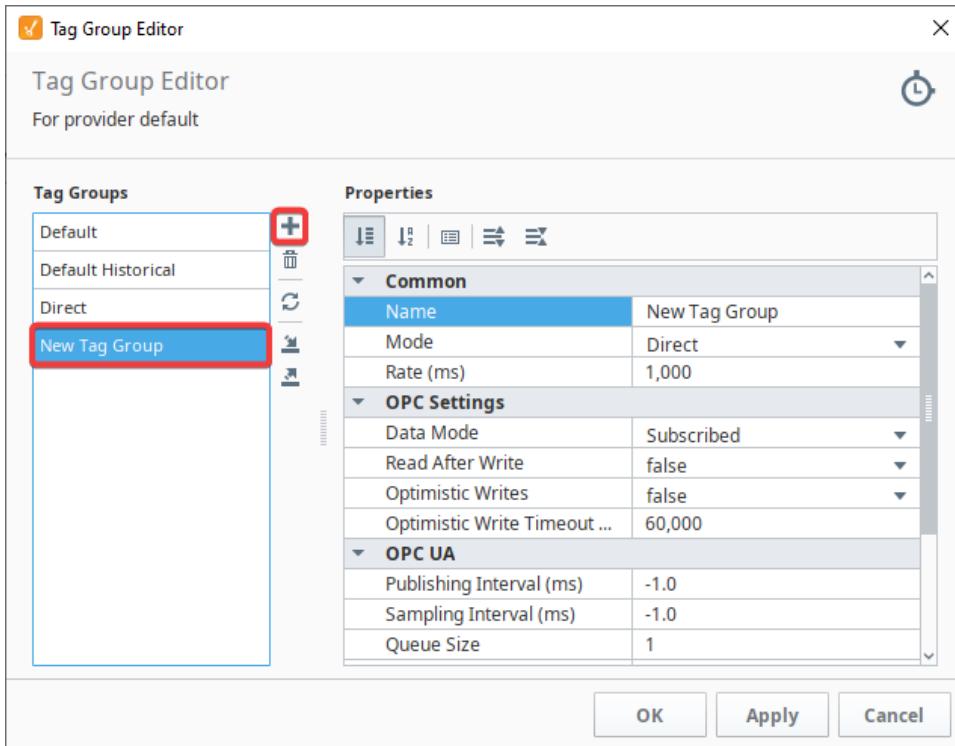
Adding and Editing Tag Groups

Adding and editing Tag Groups is easy in the Designer once you understand how the different Tag Group modes work. It's just a matter of choosing which Tag Group mode you want to use for your Tag, and entering the properties for your Tag Group.

1. In the Tag Browser, click on the **Edit Tag Group**  icon under the Options Menu to open the Tag Group Editor window.



2. A list of already configured Tag Groups appear on the left side of the window and configuration settings on the right. To add a Tag Group, click the **Add**  icon. (Alternatively, you can click on an existing Tag Group to edit it.)

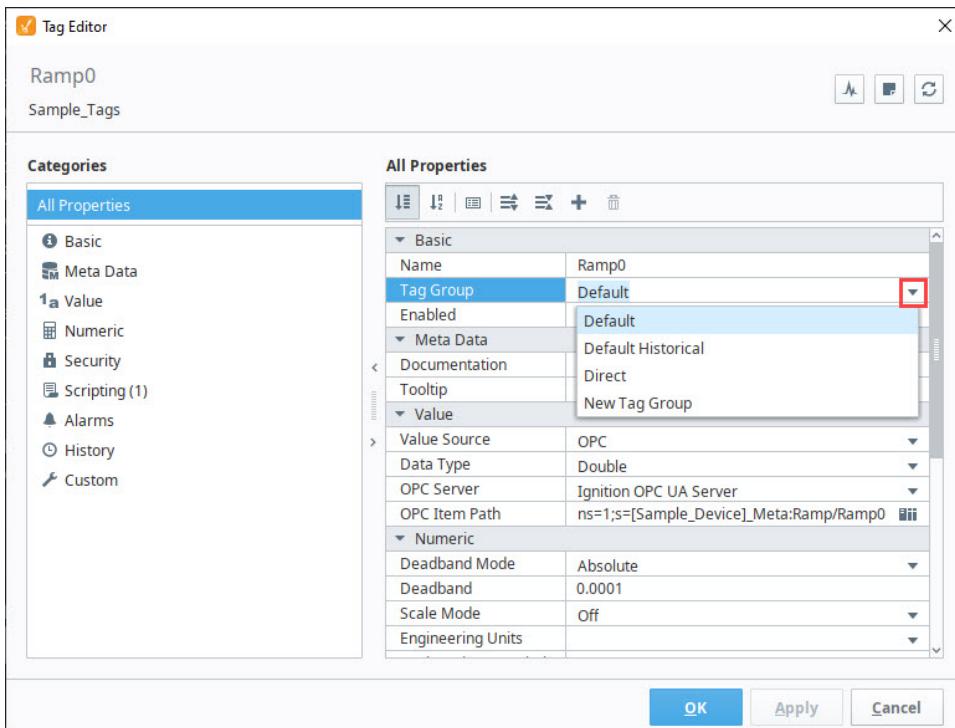


When a Tag Group in the list is selected, the properties for the group will appear on the Properties Table. Each mode will have slightly different settings that will need to be configured. You can find descriptions for those properties in the [Tag Group Properties](#) section below.

Setting a Tag Group on a Tag

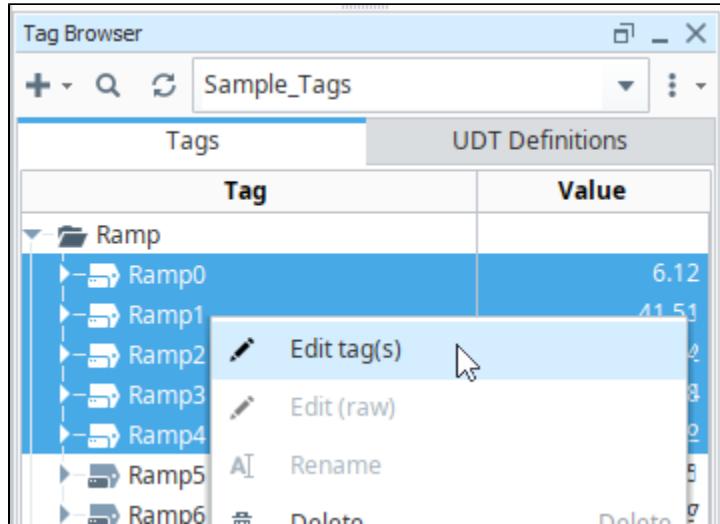
The example below demonstrates how to find the Tag Group on a tag, and change it to something else.

1. In the **Tag Browser**, right-click on any Tag, and click the **Edit tag** icon. The **Tag Editor** window opens. A list of Tag properties is displayed.
2. Under **Basic Properties**, on the right side of the **Tag Group** property, click the dropdown list and a list of available Tag Groups will appear. Again, Tag Groups are configured per Tag Provider, so this list will only contain Tag Groups that are present in the Tag's Tag Provider.



- If you wish to change the Tag Group. Select the new option, and click **OK** or **Apply** to apply the new Tag Group. The Tag will switch rates and start using the new Tag Group.

As a reminder, you can select multiple Tags in the Tag Browser by right clicking to edit the Tags. This opens the Tag Editor, and sets the Tag Group for all the selected Tags at the same time.



Tag Group Properties

The following table lists the properties for Tag Groups. Note that some properties are available only for specific modes.

Property	Description	Applicable Modes
Common		
Name	Unique name of the Tag Group.	All
Driven Mode	The rate of the Tag Group is based on the value of a driving Tag. The condition is a simple comparison between a Tag value and a number. If the condition is true, the Tag Group will execute at the fast rate. If false, it will run at the slow rate. There are two exceptions to this: the Any Change operator, and One-shot mode. Using either of these conditions will not run at a rate. Instead, it will be triggered by a change in the driving Tag's value. Keep in mind that the driving Tag can be an Expression Tag that performs complex calculations and references other Tags. In this way, it's possible to create robust Tag Group triggering.	All
Rate	Base update rate, specified in milliseconds, at which Tags will be executed. Note: If the rate is set to 0, the Tag Group will not execute.	All
Leased /Driven Rate	Used by both the Leased and Driven Modes to determine when the Tag Group should run at the fast rate.	Leased, Driven
Driving Expression	The Tag Group executes based on the condition set on the Driving Expression: Tag or Expression.	Driven
Driving Comparison	How the Comparison Value property should be compared to the Driving Tag's value. If the comparison is true, then the Fast Rate will be used by the Tag Group, otherwise, the Slow Rate will be used. The Any Change operator works differently than the other operators: The Tag Group will execute immediately whenever the driving Tag changes value. Using the Any Change operator means that the Tag Group no longer uses the Slow Rate or Fast Rate properties.	Driven
Comparison Value	Used by the Driving Comparison property to determine if the Tag Group should execute at the slow or fast rate.	Driven
One Shot	One-shot will execute once when the comparison condition is true, and not again until the condition becomes false, and subsequently true.	Driven
OPC Settings		

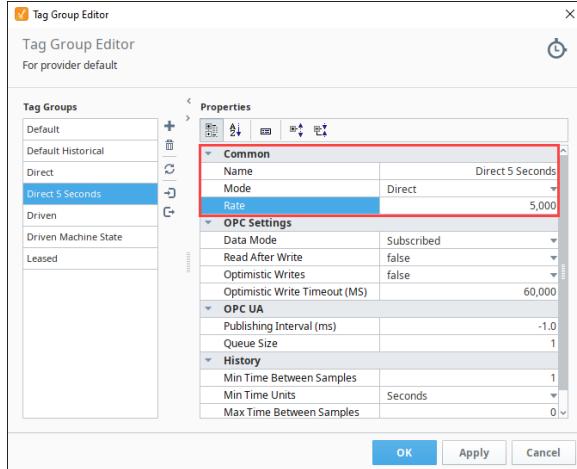
Data Mode	<p>This mode dictates how OPC values are obtained. The default mode, Subscribed, is preferred because it is more efficient than a read.</p> <p>Subscribed All OPC Tags in the Tag Group will be subscribed according to the Tag Group rate. Values will come in asynchronously as they change.</p> <p>Polled Tags will not be subscribed, but will instead be synchronously read each time the Tag Group executes. This operation is less efficient, but allows more precise control over when values are obtained. This mode is particularly useful when collecting data over a slow or expensive connection for display. When combined with the one-shot execution mode above, and a static Tag tied to a momentary button, it's easy to create a manual refresh button on a screen that pulls data on-demand.</p>	All
Read After Write	<p>When enabled, a read request will be sent immediately after a write request. This means that the value on the Tag will be updated much quicker to reflect the latest written value.</p> <p>Enabling this property is less efficient as a single write to a Tag becomes two separate requests. This is especially helpful with slower Tag Groups as the Tags will show the latest value quicker than the normal execution would allow.</p>	All
Optimistic Writes	<p>Optimistic Writes are only valid on OPC Tags. Optimistic Writes set a newly written Tag value in Ignition before receiving confirmation of the write from the PLC. This helps the operators see their newly entered value right away and is useful if you have slow a Tag Group rate. A faster rate (1 second or quicker) will have less need to turn on Optimistic Writes.</p> <p>If enabled, written values will be applied to the Tag in Ignition immediately. Normally, the system must receive confirmation that a write was successful from the device before the Tag in Ignition's value would change. The Optimistic Writes property changes the behavior by assuming the write went through until the next read value or subscription update proves otherwise. Enabling this will make writes appear to execute much quicker.</p> <p>Works in conjunction with the OPC Optimistic Write Timeout property below. If the Tag in Ignition does not receive confirmation that the new write was successful within the timeout, the Tag will change back to the last known value. While in an ambiguous state, the Tag will have a quality of "Good (Provisional)".</p> <p>This setting can be paired with the OPC Read After Write: the Ignition Tag will assume the newly written value, while an asynchronous read request is quickly sent out to confirm the write went through.</p> <p>While the write is pending, values received from subscription activities will override the current value. Assuming an initial value of 0, if a write of 10 is applied to the Ignition Tag, then the Tag will show a value of 10 until the system can confirm the new value. If a subscription update then returns a value of 5, the Ignition Tag will change to 5.</p>	All
Optimistic Write Timeout (ms)	The timeout period for Optimistic Writes. A value of 0 effectively disables the fallback functionality: the new value is maintained on the Tag until the next read or subscription activity.	All
OPC UA		
Publishing Interval (ms)	<p>The rate at which data is delivered to the OPC-UA client.</p> <p>A value of -1 means automatic, allowing the OPC-UA client to determine the rate.</p>	All
Sampling Interval (ms)	<p>The following feature is new in Ignition version 8.1.24 Click here to check out the other new features</p> <p>Allows the Requested Sampling Interval column on the Gateway Status > OPC Connections > Clients > Nodes page to be requested independently of the Tag Group Rate. The default value for this is -1, which indicates that the value will be inherited from the Tag Group Rate value.</p> <p>Note: Setting the Sampling Interval value to 0 against a driver that supports it will show the requested value on the Gateway. Setting the value to 0 against a driver that doesn't support it will show a Revised Sampling Interval of 500. Setting the value below 0 will show the the interval rate of the Tag Group setting in the Requested /Revised Sampling Interval columns.</p>	All
Queue Size	<p>The OPC-UA specifications states that in cases where the sampling interval (the rate at which the server checks the data source for changes) is faster than the publishing interval (rate at which the data is delivered to the client), the samples may be queued or batched together before publishing. This setting determines the maximum size of that queue. When the maximum is reached and a publish has not yet occurred, oldest samples are dropped first.</p> <p>Currently, there are not many features in Ignition that utilize multiple entries in the queue, but 3rd party OPC-UA clients may be able to take advantage of this setting.</p> <p>Note that values on Ignition Tags will only ever show one value at a time, regardless of what this property is set to. (If the value is 0, nothing will be queued.)</p>	All

	<p>Support for this feature is dictated by the driver.</p> <ul style="list-style-type: none"> • DNP3 - See the DNP3 page for details on how buffered data and Sequence of Events works 	
Include Timestamp -Only Changes	<p>The following feature is new in Ignition version 8.1.10 Click here to check out the other new features</p> <p>If true, allows the OPC-UA client to receive timestamp-only changes for tags in this group. By default, the OPC-UA client only receives value and quality changes.</p>	All
History		
Min Time Between Samples	Minimum time between samples (integer).	All
Min Time Units	Minimum time in units is defined as: Milliseconds, Seconds, Minutes, Hours, Days, Weeks, Months, and Years.	All
Max Time Between Samples	<p>Maximum time between samples (integer).</p> <p>When a Tag uses the "Tag Group" sample mode, and is set to this Tag Group, it will ignore its own Max Time Between Settings setting (configured directly on the Tag), but only if this Max Time Between Samples setting (on the Tag Group) is using non-default values. Defaults to 0.</p>	All
Max Time Units	Maximum time in units is defined as: Milliseconds, Seconds, Minutes, Hours, Days, Weeks, Months, and Years.	All

Direct Tag Group Example

Add a Direct Tag Group

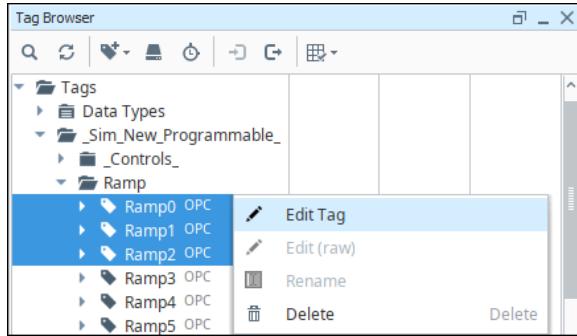
1. In the Tag Browser, click on the **Tag Groups**  icon to open the **Tag Group Editor**.
2. On the bottom left side, click the **Add**  icon to create a new Tag Group, and enter the values for the following properties:
 - a. Name - Enter a unique name for the Tag Group: **Direct 5 Seconds**
 - b. Select the Mode: **Direct**
 - c. Enter Rate: **5,000**
3. Click **OK** to save.



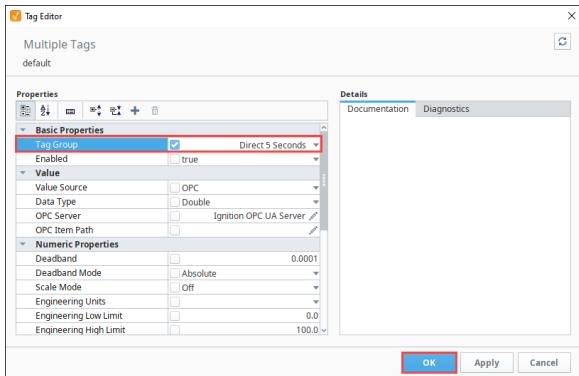
On this page ...

- [Add a Direct Tag Group](#)

4. Now that you have your Tag Group created, let's add multiple Tags to the Tag Group. Go to your **Tag Browser**, find some Tags you want to add to the Tag Group. This example uses several Ramp Tags. Right click on the selected Tags, and click on **Edit Tag**  icon.



5. This opens the **Tag Editor** window. It also shows you that you have multiple Tags selected. Select the **Direct 5 Seconds** Tag Group from the dropdown list.
6. Click **OK**.

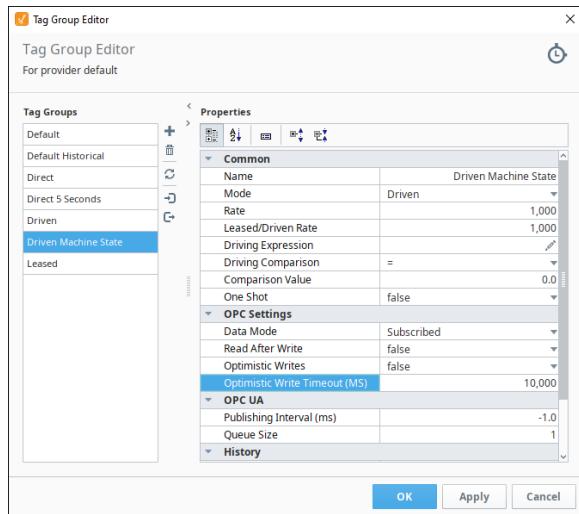


You created a new Direct 5 Second Tag Group and added your Tags. Just make sure you want to poll the 5 second values all the time (24/7) when you use the Direct 5 Seconds Tag Group.

Driven Tag Group Examples

Driven Tag Group - Machine State

Driven Tag Groups that are based on a machine can be very important when you only want to poll values differently based on machine state. A machine state Tag Group is when the Tags are polled at one rate with the machine ON, and a different rate when the machine is OFF. You can easily see the rate change in the the Tag Browser using this type of condition.



On this page ...

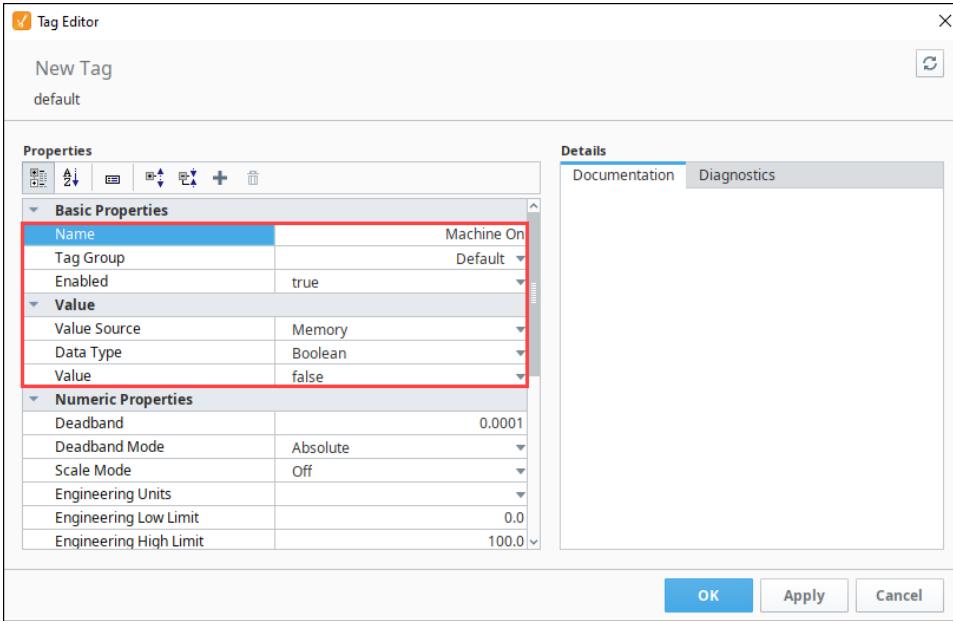
- [Driven Tag Group - Machine State](#)
 - [Create a Machine State Tag](#)
 - [Add a Driven Tag Group Based on the Machine State](#)
- [Driven Tag Group - Time of Day](#)
 - [Add a Driven Tag Group Based on the Time of Day](#)

Create a Machine State Tag

First, you need to create a driving Tag if you don't already have one.

1. In the **Tag Browser**, right-click on the **Tags** folder, then go to **New Tag > Memory Tag** to [create a memory Tag](#). The Tag Editor window is displayed.
2. In **Tag Editor**, enter the following:
Name: **Machine On**
Tag Group: **Default**
Enabled: **true**
Data Type: **Boolean**
Value: **false**

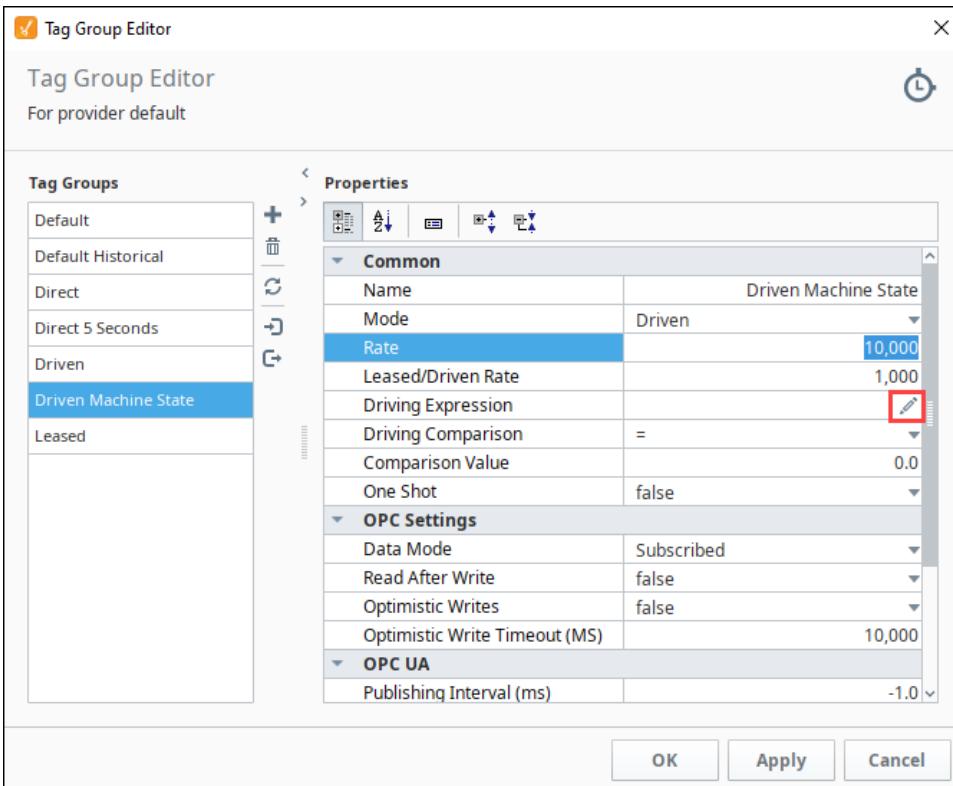
3. Click **OK** to add the **Machine On** Tag to the Tag Browser.



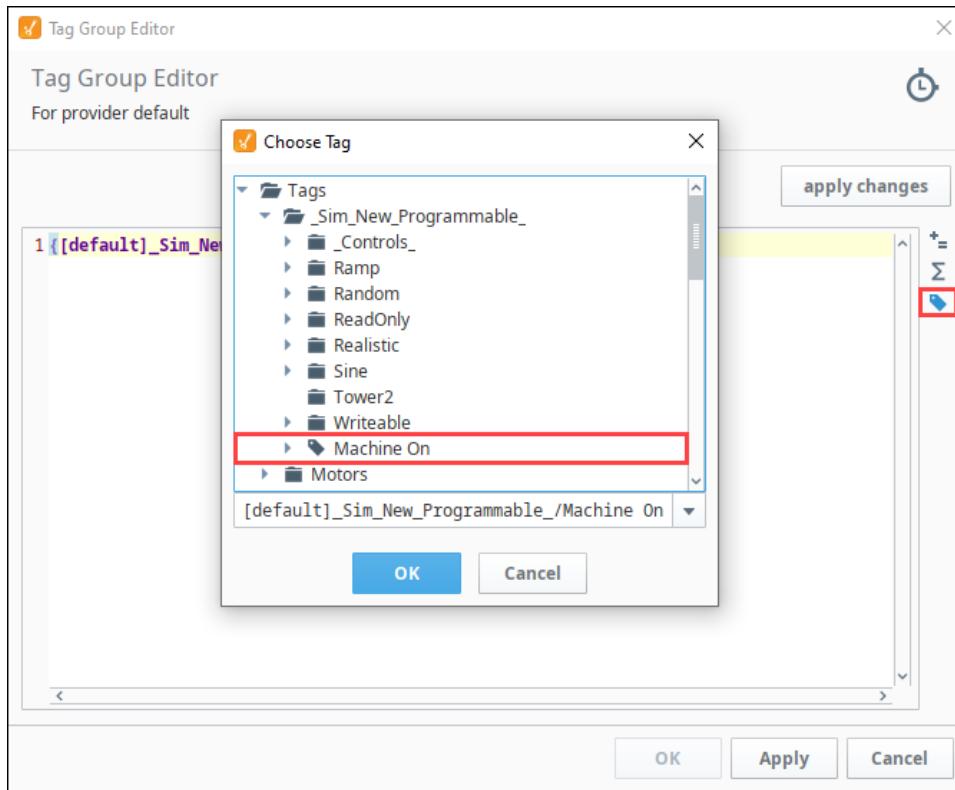
Add a Driven Tag Group Based on the Machine State

Once you have your driving Tag created, add a Driven Tag Group that updates the Tags based on when the machine is ON.

1. In the **Tag Browser**, click on the **Edit Tag Groups** icon to open the Tag Group Editor.
2. To create a new Tag Group, click **Add** icon.
3. Enter the name of the Tag Group. For example, you can name it **Driven Machine State**, and set the **Mode** to **Driven**.
4. Set the **Rate** to 10,000 (10 seconds) and the **Leased/Driven Rate** to 1000 (1 second).
5. This Tag Group executes based on the condition you set on the **Driving Expression**. Set the **Driving Expression** by clicking the **Edit** icon.



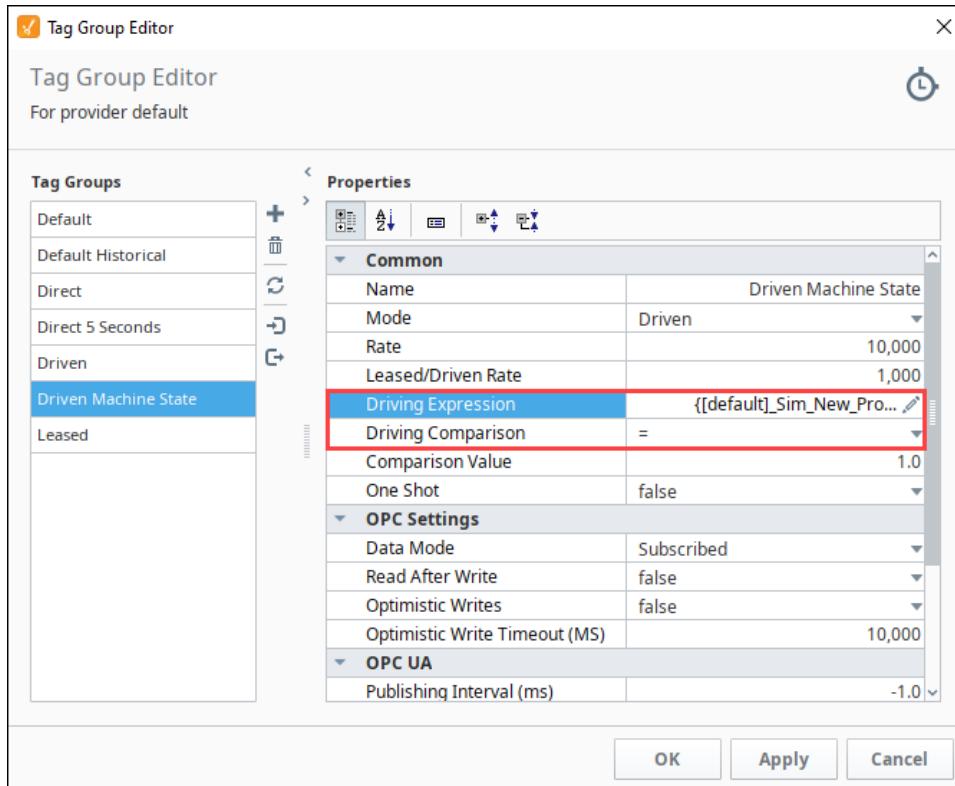
6. This Tag Group executes based on the condition you set on the **Driving Expression**. Set the **Driving Expression** by clicking the **Edit** icon.
7. Click the **Tag** icon and choose the Tag you want to set the condition on (i.e., Machine On).
8. Click **OK** on the popup and **apply changes** in the upper right of the Tag Group Editor.



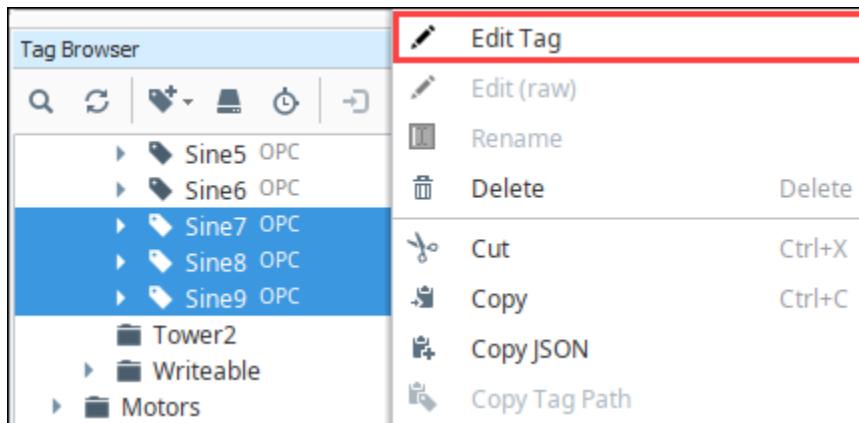
9. While still in the Tag Group Editor, set the following values:

- Driving Comparison: =
- Comparison Value: 1

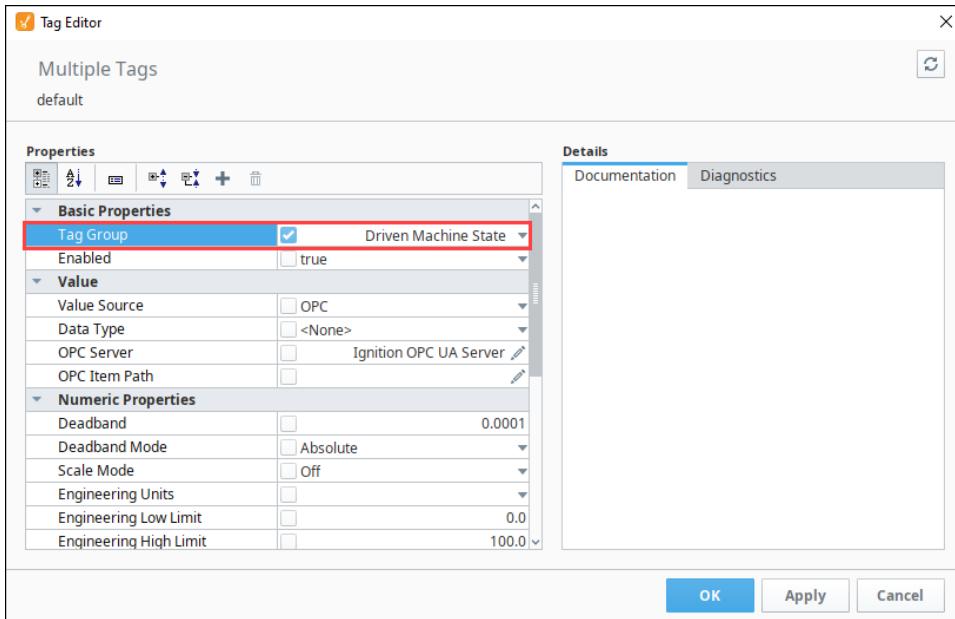
10. Click **OK** in the lower right of the Tag Group Editor.



11. When the machine is on, it will use the Leased/Driven Rate (1 second). When the machine is off, it will use the Rate (10 seconds)
12. Next, select all the Tags that you want to use for the Drive Machine State Tag Group. This example uses **Sine7**, **Sine8**, and **Sine9** Tags.
13. Right click on the selected Tags and choose **Edit Tag**.



14. The Tag Editor will open showing you are editing multiple Tags. From the Tag Group dropdown list, choose **Drive Machine State**, and click **OK**.



15. Now let's turn the machine on. In the Tag Browser, mark the checkbox next to the **Machine On** Tag that we created earlier. Now, when the machine is on, polling will be at the 1 second rate for the Sine Tags. When the machine is off, polling will be at the 10 second rate for the Sine Tags. Test it by toggling the Machine On tag on and off.

Tag Browser			
Search Refresh New Open			
Tags Servers Databases Groups			
▼ Sine <ul style="list-style-type: none"> ▶ ◆ Sine0 OPC -45.22 Float ▶ ◆ Sine1 OPC 10.21 Float ▶ ◆ Sine2 OPC -93.46 Double ▶ ◆ Sine3 OPC 26.88 Double ▶ ◆ Sine4 OPC 5 Integer ▶ ◆ Sine5 OPC -20 Integer ▶ ◆ Sine6 OPC -25 Long ▶ ◆ Sine7 OPC 19 Long ▶ ◆ Sine8 OPC 66.39 Double ▶ ◆ Sine9 OPC 3.73 Double 			
▼ Tower2			
▼ Writeable			
▶ ◆ Machine On Memory	<input checked="" type="checkbox"/>	Boolean	
▼ Motors			

You can also set the **Rate** (slow rate) to a value of 0, which means polling is stopped when the condition is false. So for example, when the machine is off, you will see the Tags stay at the last known value.

Tag Browser

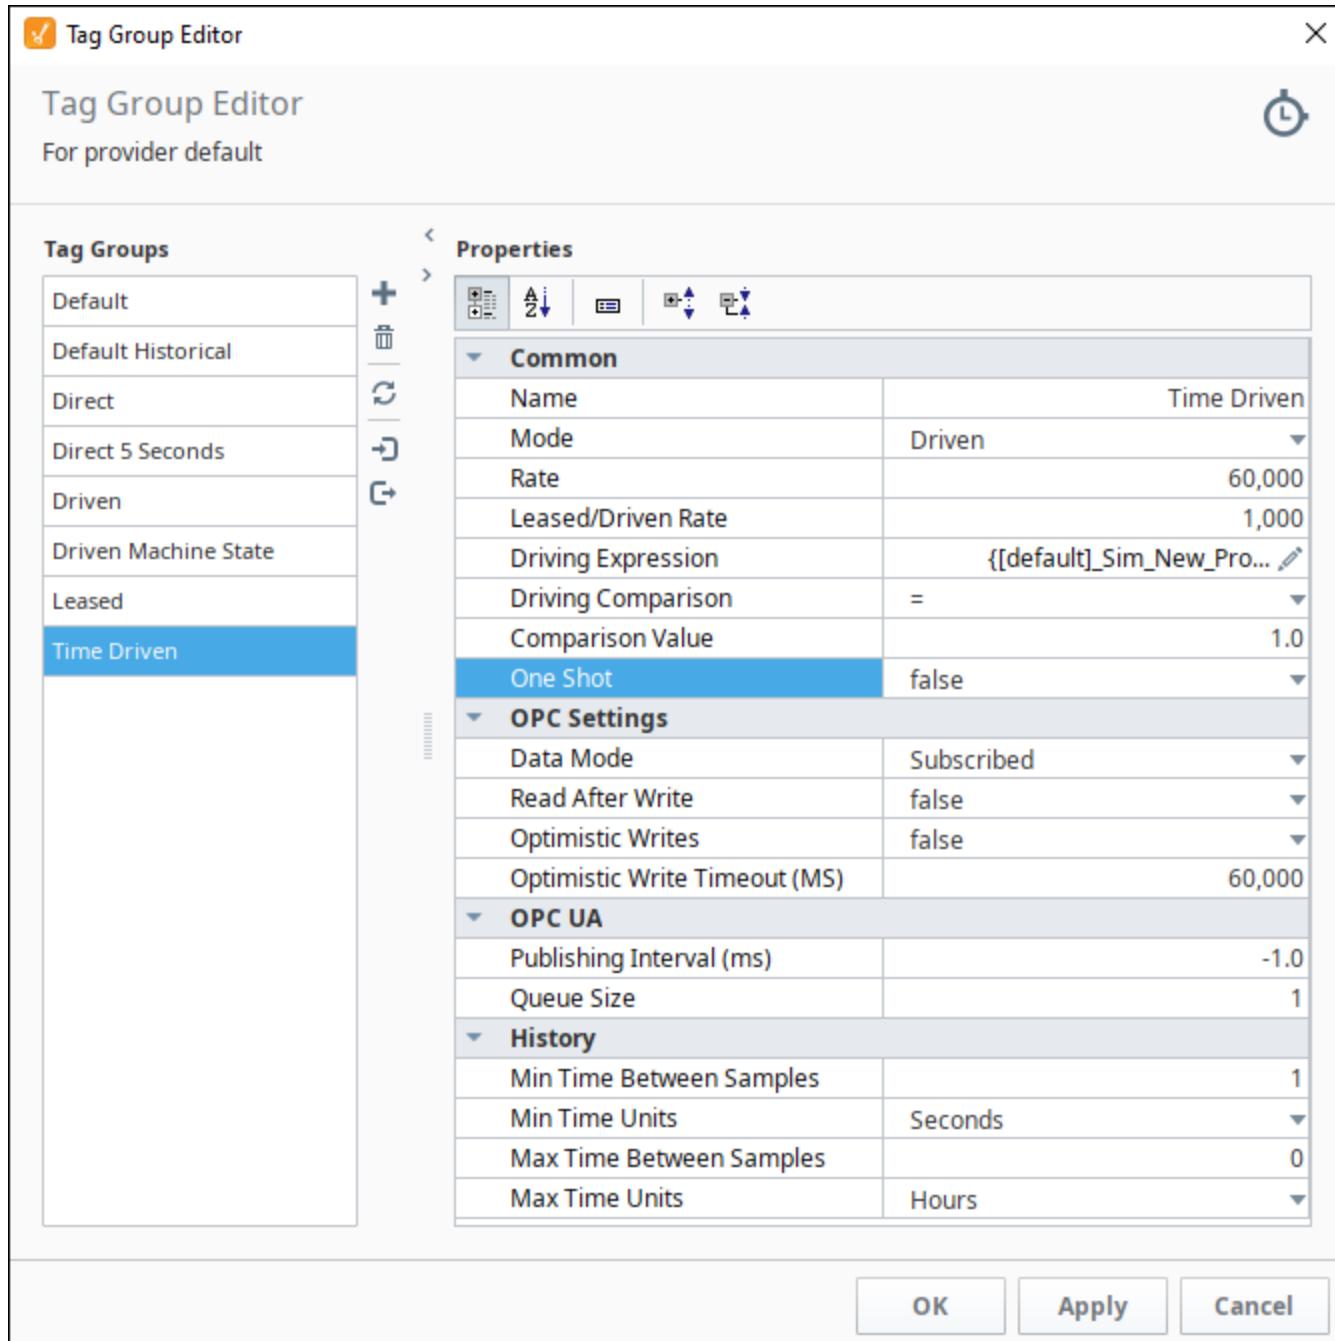
▼	Sine		
▶	Sine0 OPC	48.85	Float
▶	Sine1 OPC	242.11	Float
▶	Sine2 OPC	-18.42	Double
▶	Sine3 OPC	-244.69	Double
▶	Sine4 OPC	-16	Integer
▶	Sine5 OPC	9	Integer
▶	Sine6 OPC	17	Long
▶	Sine7 OPC	14	Long
▶	Sine8 OPC	84.38	Double
▶	Sine9 OPC	2.65	Double
▶	Tower2		
▶	Writeable		
▶	Machine On Memory	<input type="checkbox"/>	Boolean
▶	Motors		

Once the machine is turned back on, the Tags resume polling.

Driven Tag Group - Time of Day

Driven Tag Groups can be used as the polling rate for Tags to trigger at different rates at different times of the day. They can also be used as a one-shot event at a specific time of the day. You can accomplish this by setting the condition of the Driven Tag Group to be true at certain times of the day, and false at other times.

Keep in mind, that if you use a Driven Tag Group based on the time of day with a 0 Slow Rate for history, the Tag will only store history during those specified times.



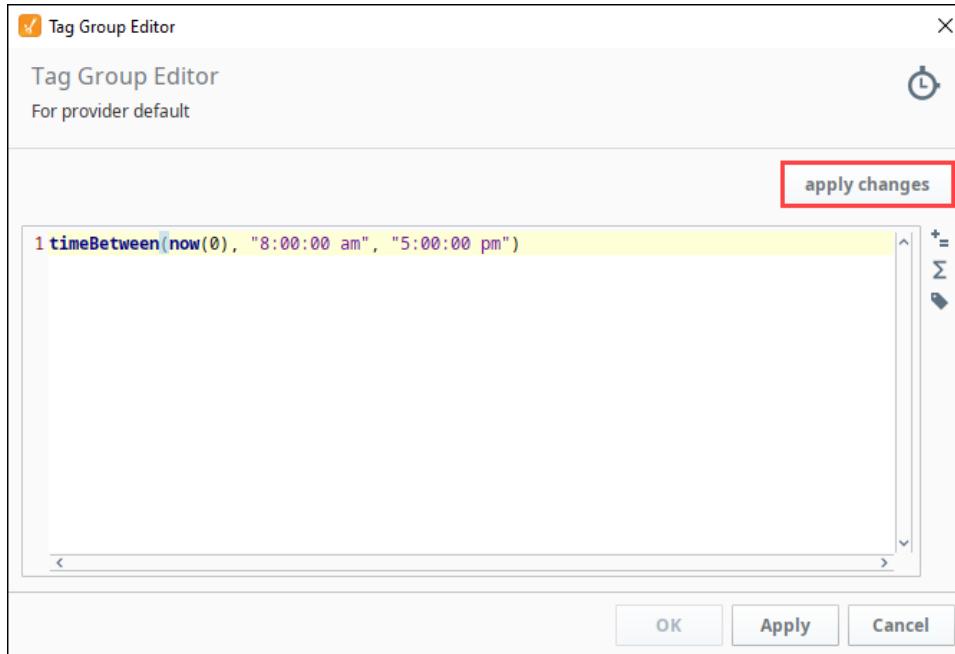
Add a Driven Tag Group Based on the Time of Day

Let's add a Driven Tag Group that updates the Tags based on a time of day. We will use an Expression to drive the Tag Group. You can use the functions that are in the expression language to poll the PLC during the hours of 8am to 5pm.

1. In the Tag Browser, click on the **Tag Groups** icon to open the Tag Group Editor.
2. To create a new Tag Group, click on the **Add** icon.
3. Enter the name for the Tag Group. In this example, we named it **Time Driven**, and set the Mode to **Driven**. This Tag Group executes based on the condition of the **Driving Expression**. In this example, we'll use the Poll Time Tag we created above.
4. Set the **Rate** to 60,000ms so it polls at a slow rate, and the **Leased/Driven Rate** to 1,000ms so it polls at a faster rate.
5. Enter the **Driving Expression**. Click on the **Edit** icon, and copy and paste the following expression in the expression box, then click **Apply Changes**.

timeBetween and Machine On expression

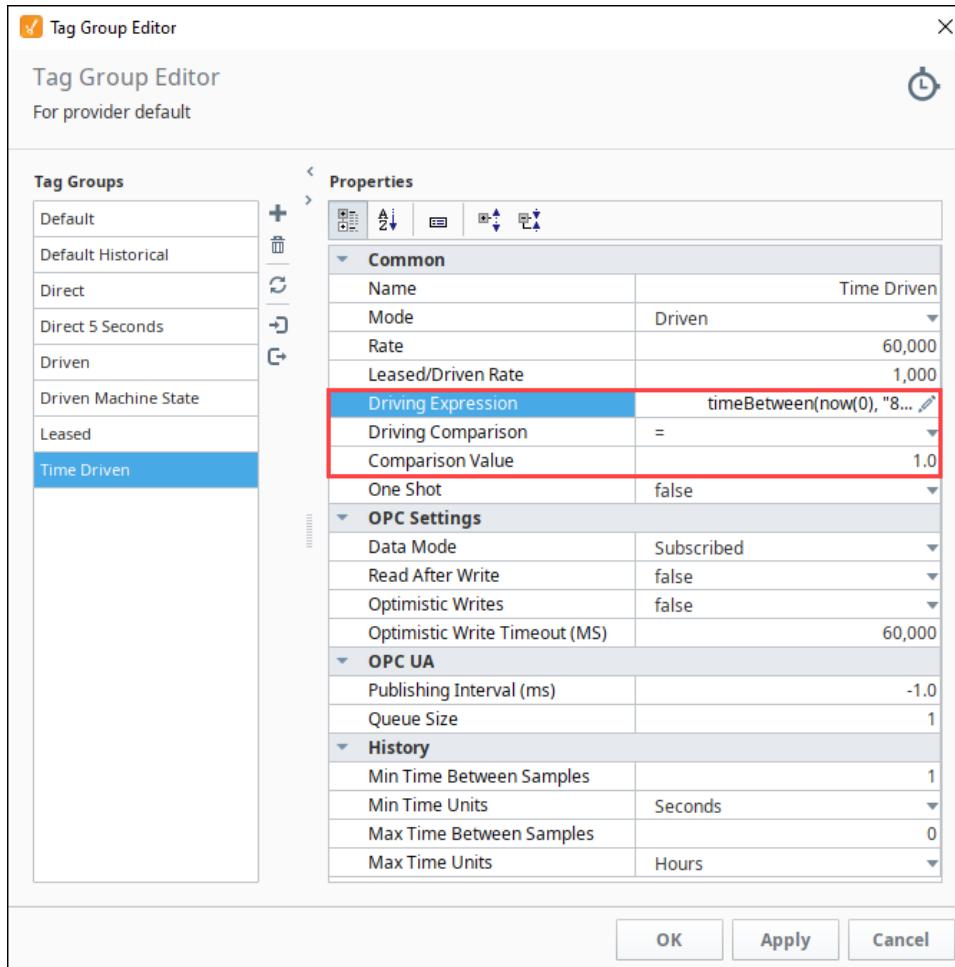
```
timeBetween(now(0), "8:00:00 am", "5:00:00 pm")
```



6. While still in the Tag Group Editor, set the following values:

- Driving Comparison: =
- Comparison Value: 1

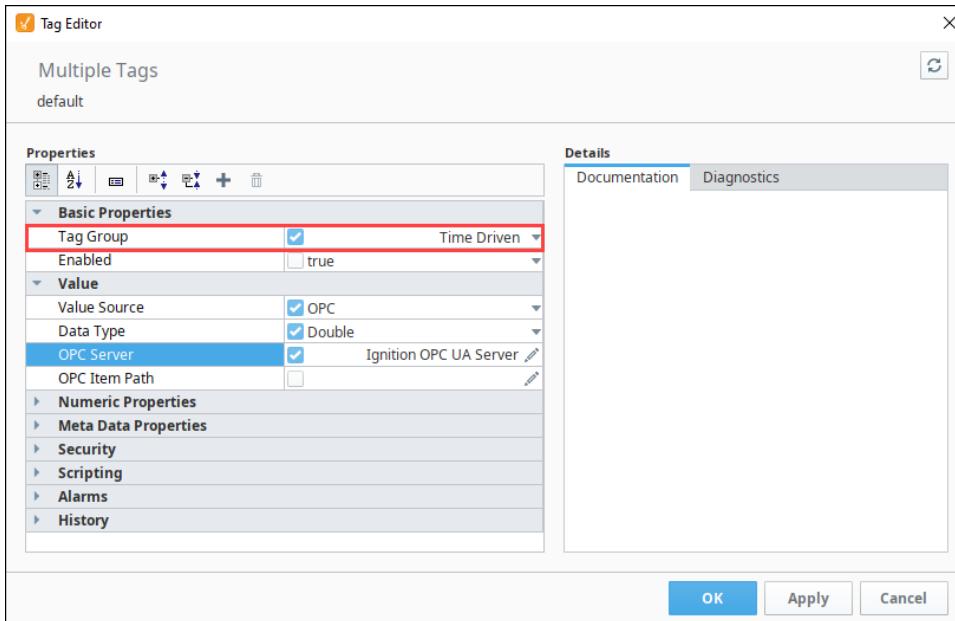
7. Click **OK**.



8. Next, select all the Tags that you want to specify in the Tag Group (i.e., **Pressure** and **Temperature**). Right click on the selected Tags, and choose **Edit Tag**.

▶ Test			
▶ Writeable			
▶ Machine On Memory	<input type="checkbox"/>	Boolean	
▶ Poll Time Expression	<input type="checkbox"/>	Boolean	
▶ Pressure OPC	68.56	Double	
▶ Temperature OPC	98.54	Double	
▶ System			
▶ Vision Client Tags			

9. The Tag Editor will open showing you are editing multiple Tags. To change the Tag Group, choose **Time Driven** from the dropdown, and click **OK**.



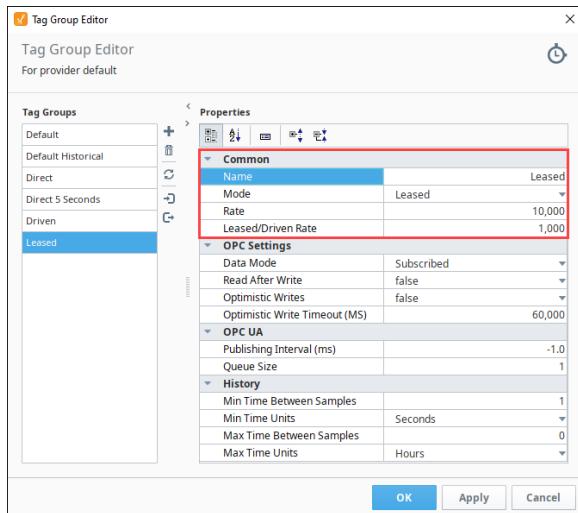
10. In the Tag Browser you can look at the Tags to see they are updating at the correct rate. Try adjusting the time range in your expression to change the rate of polling. When the driving conditions are true, that is the time between hours of 8am and 5pm, polling is at the Leased /Driven Rate of 1000 ms. When not, polling is stopped, and the last known value is displayed.

Leased Tag Group Example

Add a Leased Tag Group

Let's add a new Leased Tag Group that polls a Tag from the PLC at a 1 second rate when someone needs to view that Tag in the Client. If the Client is closed the tags will not poll at all.

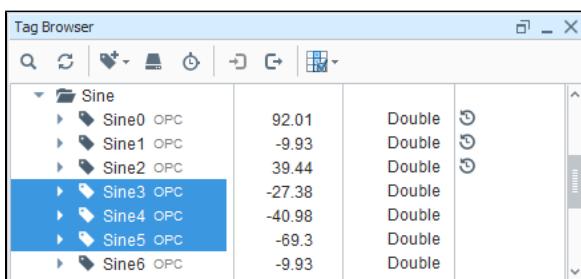
1. In the **Tag Browser**, click on the **Tag Group** icon to open the **Tag Group Editor**.
2. On the left side of the Tag Group Editor window, you can see all existing Tag Groups. Click on the **Add +** icon on the lower left side of the window to add a new Tag Group.
3. Enter the name of the Tag Group. For example, you can name this new Tag Group **Leased**, and set the **Mode** to **Leased**.
4. Set the **Rate** to 10,000ms (10 seconds) and the **Leased/Driven Rate** to 1,000ms (1 second).
5. Click **OK**.



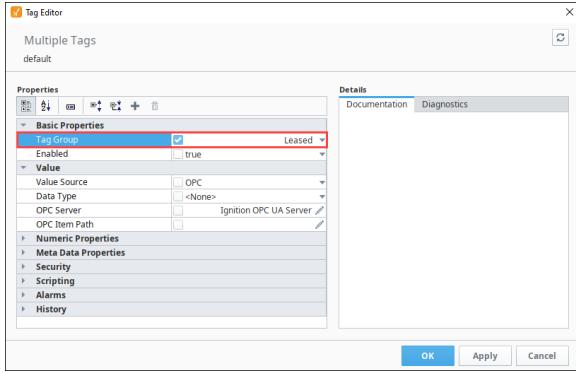
On this page ...

- [Add a Leased Tag Group](#)

6. Next, select the Tags you want to use on the new Leased Tag Group.
 - a. Go to your Tag Browser, and select some Tags you want to add to the Leased Tag Group. This example uses three Sine Tags: Sine3, Sine4, and Sine5.
 - b. Right click on the selected Tags, and click on **Edit tags**.



7. The **Tag Editor** will open showing you are editing multiple Tags. From the Tag Group dropdown list, choose **Leased**, and click **OK**.



You just created a new Leased Tag Group and added some Tags. These Tags are only polled quickly when a component that is bound to it is showing in a Client, Session, or when the Tags are showing in the Designer. For example, these three Sine Tags that are part of the Leased Tag Group are each bound to an LED component.

Tag	Value	Type
Sine0 OPC	-44.17	Float
Sine1 OPC	10.29	Float
Sine2 OPC	-60.5	Double
Sine3 OPC	135.45	Double
Sine4 OPC	2	Integer
Sine5 OPC	20	Integer
Sine6 OPC	-25	Long
Sine7 OPC	22	Long
Sine8 OPC	88.78	Double

Note: Remember that Tag History is stored based on the current Ignition Tag value. If your Leased Tag Group has a **0** slow rate, you will not get updated data to store for history. Always make sure your Tag Group slow speed is at least as fast as your History Tag Group speed.

Tag Providers

Tag Providers

At the highest level of Tag configuration is the Tag Provider. A provider is a Tag database (a collection of Tags) and a name. An Ignition Gateway can have any number of Tag Providers, and therefore the name is used to distinguish which provider a Tag comes from. Tag Providers can be set up with security or even disabled independent of each other.

Every copy of Ignition has its own Tags, spread across one or more Tag Providers. With the remote Tag Provider, Ignition can also see the Tags on another Gateway, as long as the two Gateways are connected through a Gateway network.

All Tags reside in a Tag Provider and have realtime values. Additionally, there is the concept of [Tag historian providers](#), which can store and query historical data for Tags. Each Tag can optionally have a historian provider assigned to it to whom it will report value changes for historical storage.

Realtime Provider Types

There are two types of Realtime Tag Providers that you can choose from:

Standard Tag Provider

Standard Tag Providers store all configuration and do any execution (read, write, history, alarms) through the local Ignition Gateway. Every new Ignition installation automatically creates a Standard Tag Provider named "default." You can add as many Standard Tag Providers as you want. This provider can be exposed or hidden from other Gateways on the network through the Gateway's OPC UA settings.

Remote Tag Provider

Remote Tag Providers connect a remote installation of Ignition and access those Tags. The remote Tag Provider works by creating a link from the local Gateway to a Tag provider on a remote Gateway using a Gateway Network connection. The local Ignition may be allowed to read and write to the remote Tags, but any execution is handled by the remote Gateway. So, things like writing to a PLC, alarms, and history will still be handled by the remote Ignition.

Put another way, "tags" in a remote tag provider are simply representations of the tags as they exist on the remote system. By default, a Remote Tag Provider will fall under the [Default Security Zone](#) and be read only.

Note: Due to iterative changes in the platform, UDT Definitions on 8.1+ Gateways can only be viewed or edited remotely from 8.1+ Gateways configured with a remote tag provider.

Configuring Realtime Providers

Realtime Tags providers are configured in the Gateway's Config section under [Tags > Realtime](#). After installation, the Ignition Gateway will start with a standard provider defined. You can edit its name and settings by selecting [edit](#) to the right of its entry in the table, or create new providers by selecting [Create new Realtime Tag Provider](#) below the table.

On this page ...

- [Tag Providers](#)
- [Realtime Provider Types](#)
 - [Standard Tag Provider](#)
 - [Remote Tag Provider](#)
- [Configuring Realtime Providers](#)
 - [Standard Tag Provider](#)
 - [Remote Tag Provider](#)



INDUCTIVE
UNIVERSIT

Tag Providers

[Watch the Video](#)



INDUCTIVE
UNIVERSIT

Remote Tag Provider

[Watch the Video](#)

The screenshot shows the Ignition software interface. The top navigation bar includes the Ignition logo, Help, and Get Designer buttons. The left sidebar, titled 'SYSTEM', contains links for Home, Status, Config (which is selected), Overview, Backup/Restore, Licensing, Modules, Projects, Redundancy, and Gateway Settings. Under 'NETWORKING', there are links for Gateway Network and Email Settings. Under 'SECURITY', there is a link for Security. The main content area is titled 'Realtime Tag Providers'. It displays a table with two rows. The first row has columns for Name ('default'), Description ('Default tag provider'), Enabled ('true'), Type ('Standard Tag Provider'), and actions ('delete' and 'edit'). The second row has columns for Name ('Test_Provider'), Description (''), Enabled ('true'), Type ('Remote Tag Provider (Gateway Network)'), and actions ('delete' and 'edit'). Below the table is a link to 'Create new Realtime Tag Provider...'.

Note: When setting up your Realtime Tag Provider, naming your Provider the same name as the database connection is **not** recommended. Doing so may cause no values to be returned when querying Tag data.

Standard Tag Provider

Tags are stored inside of Ignition and executed by the system.

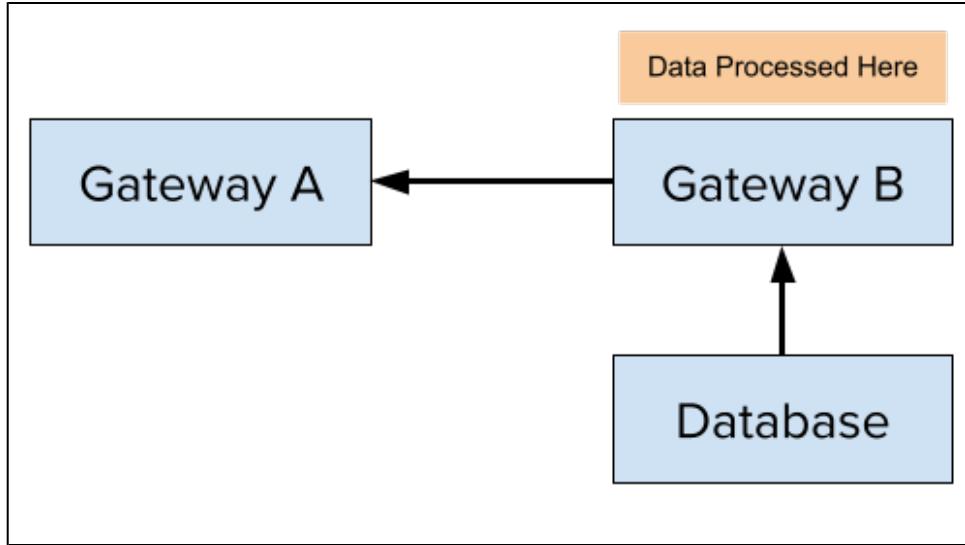
Setting	Description
Name	The name of the provider.
Description	The description of the provider.
Enabled	If true, Tag provider is enabled. Default is true.
Default Database	The default database connection to use for expression Tags that run SQL queries. All query Tags with default database providers selected will run their queries against this database source.
Tag Read Permissions	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>Any actor attempting to read any tag within this tag provider must have these permissions in addition to the target tag's permissions. Multiple security level paths can be specified by separating them with commas. If blank, tag reads for this provider will not be restricted by default.</p> <p>For example:</p> <pre>Authenticated/Roles/Administrator, SecurityZones/MyZone</pre> <p>When multiple security levels are provided, the radio buttons determine if the user needs all of the listed security levels, or at least one.</p>
Read Only	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>If enabled, writes to any tag within this tag provider will be rejected, regardless of the target tag's permissions. Disabled by default.</p>
Tag Write Permissions	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p>

	<p>Any actor attempting to write to any tag within this tag provider must have these permissions in addition to the target tag's permissions. Multiple security level paths can be specified by separating them with commas. If blank, tag writes for this provider will not be restricted by default.</p> <p>See the example for Tag Read Permissions.</p>
Tag Editing Permissions	<p>Determines the roles required to edit, create, or delete Tags in the provider. Expects a path to a Security Level. Multiple levels can be specified, separated by a comma. For example:</p> <pre>Authenticated/Roles/Administrator, SecurityZones/MyZone</pre> <p>When multiple security levels are provided, the radio buttons determine if the user needs all of the listed security levels, or at least one.</p> <div style="border: 1px solid orange; padding: 5px; margin-top: 10px;"> <p>The following feature is new in Ignition version 8.1.2 Click here to check out the other new features</p> </div> <p>As of version 8.1.2, Edge Gateways now have access to the Tag Editing Permissions setting.</p>
Advanced Properties	
Allow Back-fill Data	<p>The following feature is new in Ignition version 8.1.4 Click here to check out the other new features</p> <p>If enabled, data will be allowed to arrive out of order from the source. Data from the past will be stored to history, but will not be used for alarms, scripts, or subscriptions. If false (default behavior), each value will be processed fully as it arrives. Default is false.</p>
Enable Tag Reference Tracker Store	<p>The following feature is new in Ignition version 8.1.34 Click here to check out the other new features</p> <p>Enables the storing of Tag Reference entries to a database on the local Gateway for analysis in a Designer. Default is true.</p>

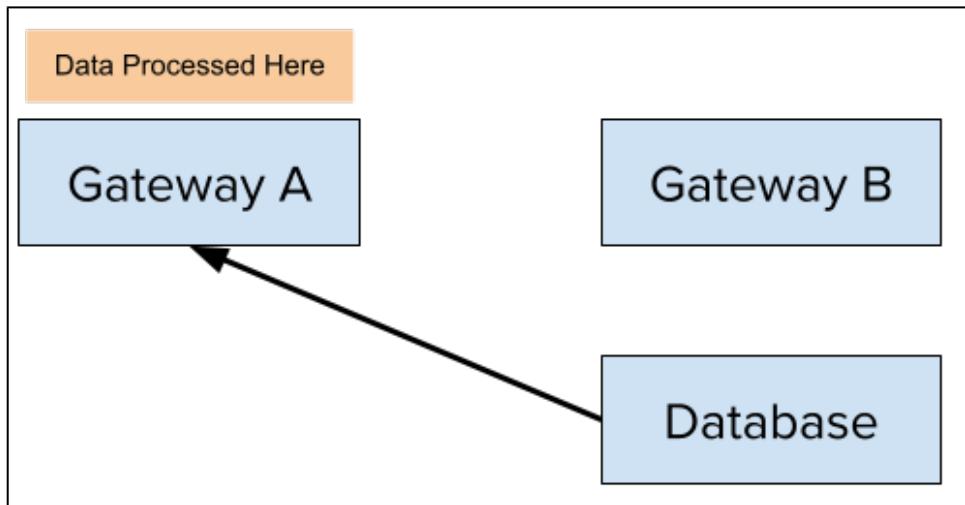
Remote Tag Provider

Tag Provider from one Gateway is brought in to another Gateway.

Setting	Description
Name	The name of the provider.
Description	The description of the provider
Gateway	The name of the Gateway on the Gateway Network that this provider is coming from.
Provider	The name of the provider as it is on the remote Gateway. This does not have to be the same as its name on the new Gateway.
History Access Mode	<p>This setting dictates how Tag history is queried for remote Tags. Ultimately determines where records are processed.</p> <p>When set to GatewayNetwork, history requests will go through the Gateway Network. Meaning the remote gateway will query records from its database, process the records, then send the results over to the requesting gateway.</p>



When set to **Database** history requests will run against the database directly. In this scenario, the requesting gateway then processes the data.



History Datasource	The datasource to query when History Access Mode is set to Database.
History Driver	If querying the database directly, this is the Gateway name of the remote system. It is used to identify data from that system in the database.
History Provider	If querying the database directly, this is the name of the Tag provider on the remote system. It is used along with driver name to identify the correct Tags in the database.
Alarms Enabled	If true, alarms configured on the remote Gateway will be enabled on the new Gateway.
Alarm Mode	How alarm state should be monitored. In 'queried', state will be queried through the Gateway network when necessary. In 'subscribed', the state will be subscribed, and updates will be sent asynchronously. Subscribed provides better performance, but uses more memory.
Advanced Properties	
Allow Back-fill Data	<p>The following feature is new in Ignition version 8.1.4 Click here to check out the other new features</p> <p>If enabled, data will be allowed to arrive out of order from the source. Data from the past will be stored to history, but will not be used for alarms, scripts, or subscriptions. If false (default behavior), each value will be processed fully as it arrives. Default is false.</p>
Enable Tag Reference	

Tracker Store

The following feature is new in Ignition version **8.1.34**
[Click here](#) to check out the other new features

Enables the storing of Tag Reference entries to a database on the local Gateway for analysis in a Designer. Default is true.

Tag Event Scripts

Scripts can be attached to Tags. When you edit a Tag, you can navigate to the Tag Events section and click on the **Edit**  icon to see a list of all of the available events. Those events are

- Value Changed
- Quality Changed
- Alarm Active
- Alarm Cleared
- Alarm Acknowledged

Because Tags are stored in the Gateway, they aren't scoped to a specific project. In these cases, some system functions may require that you specify a project or other resource, such as a database for the `syst em.db.runPrepUpdate` function.

Note: Because these scripts are Gateway scoped, certain things like print statements will not print to the Designer console, but will print instead to the wrapper log file in Ignition's installation directory. The prefer approach to adding logging statements to tag event scripts is to use the `system.util.getLogger` function, which will send the messages to the [Gateway's Logs](#) page.

This feature was changed in Ignition version **8.1.17**:

The Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts.

Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

Note: It is not recommended to have long running scripts on Tag Events due to limited thread pool availability.

On this page ...

- Event Options
 - Value Changed
 - Quality Changed
 - Alarm Active
 - Alarm Cleared
 - Alarm Acknowledged
- Using the Project Library in a Tag Event Script
- UDT Parameters in Tag Event Scripts
 - Modern Approach
 - Legacy Approach
- Tag Script Examples
- Troubleshooting



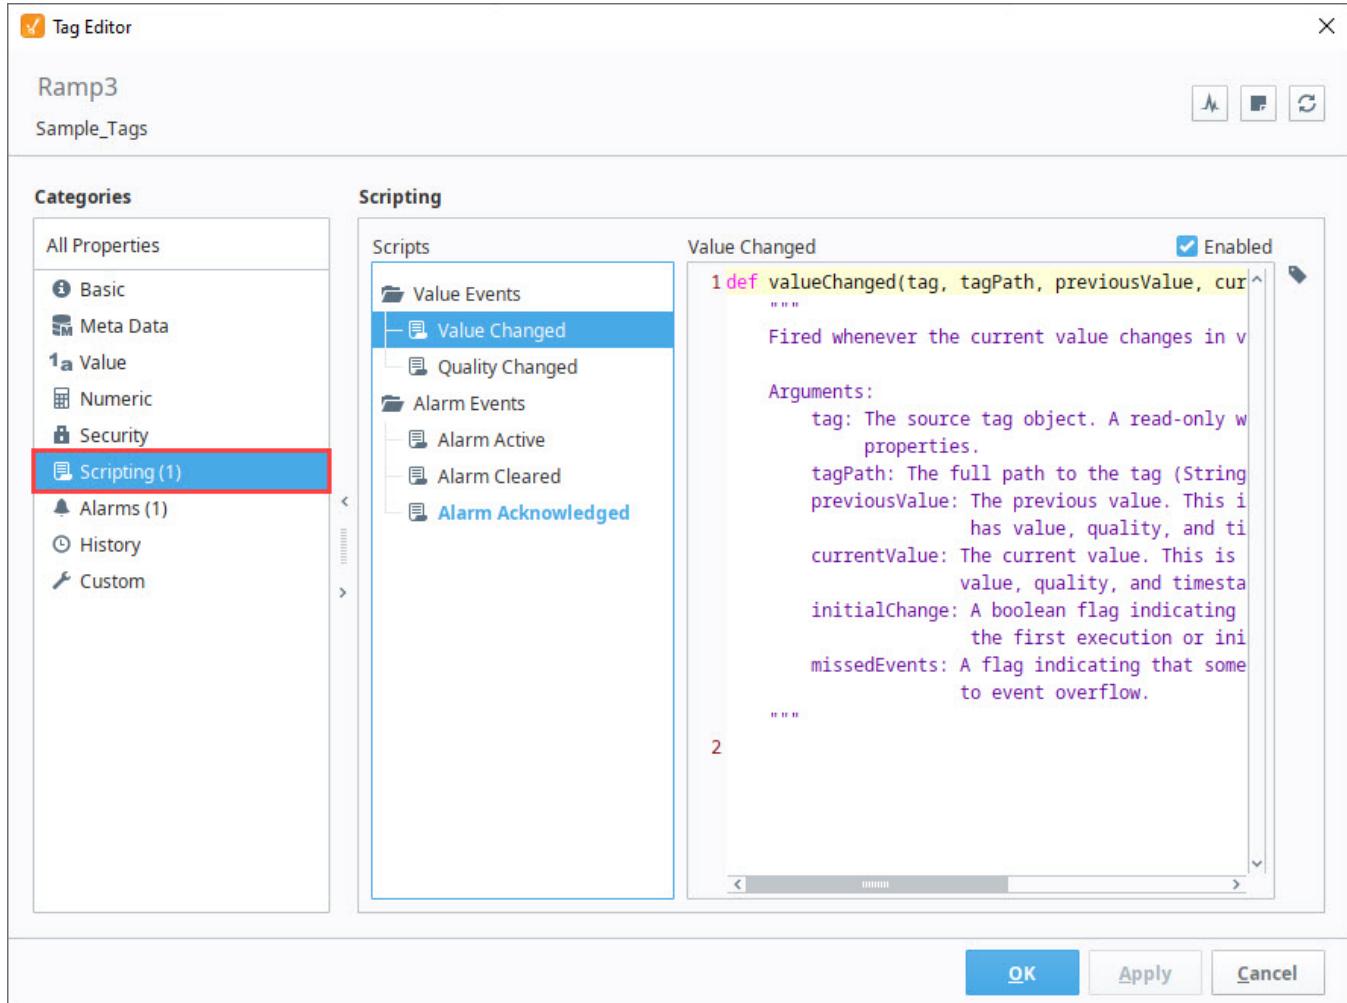
INDUCTIVE
UNIVERSITY

Tag Event Scripts

[Watch the Video](#)

Event Options

Once an event has been selected, note that the top of the text area is [defining a function](#). As a result, all code that should execute when this event occurs should be indented at least once to be included in the definition.



Value Changed

The Value Changed event is fired whenever the value of this particular Tag changes. Since Tags use a QualifiedValue, which include a value, quality, and timestamp, this script will fire whenever any of those change. Refer to [Scripting Object Reference](#) for more information.

This event has a variety of arguments available for use in the script:

- String tagPath - The full path to the Tag. Example: [tagProvider]Folder/Folder/Tag
- Object previousValue - The previous value. This is a qualified value, so it has value, quality, and timestamp properties.
- Object currentValue - The current value. This is a qualified value, so it has value, quality, and timestamp properties.
- Boolean initialChange - A boolean flag indicating whether this event is due to the initial subscription or the first execution after a Tag update.
- Boolean missedEvents - A flag indicating that some events have been skipped due to an event overflow.



The currentValue and previousValue arguments are qualified values: objects that contain a value, timestamp, and quality. This means that to get to the value of the currentValue, your script would need to access currentValue.value

Quality Changed

The Quality Changed event is fired whenever the quality of this particular Tag changes. This event has a variety of arguments available for use in the script:

- String tagPath - The full path to the Tag. Example: [tagProvider]Folder/Folder/Tag
- Object previousValue - The previous value. This is a qualified value, so it has value, quality, and timestamp properties.
- Object currentValue - The current value. This is a qualified value, so it has value, quality, and timestamp properties.
- Boolean initialChange - A boolean flag indicating whether this event is due to the initial subscription or the first execution after a Tag update.
- Boolean missedEvents - A flag indicating that some events have been skipped due to an event overflow.

Alarm Active

The Alarm Active event fires whenever a new active alarm event occurs. This event has a variety of arguments available for use in the script:

- String tagPath - The full path to the Tag. Example: [tagProvider]Folder/Folder/Tag
- String alarmName - The name of the alarm. This does not include the full alarm path.
- Object alarmEvent - The full alarm event object. The properties available to this object are:
 - eventId
 - source
 - name
 - priority
 - displayPath
 - displayPathOrSource (the display path if it is set, otherwise the source)
 - state (the current state, active/cleared + acked/unacked)
 - eventState (the last transition, active, clear, acknowledged)
 - isClear (a boolean if the alarm is currently cleared)
 - isAcked (a boolean if the alarm is currently acknowledged)
 - isShelved (a boolean if the alarm is currently shelved)
 - notes
- String alarmPath - The full alarm path.
- Boolean missedEvents - A flag indicating that some events have been skipped due to an event overflow.

Alarm Cleared

The Alarm Cleared event fires whenever an alarm becomes cleared on the Tag. This event has a variety of arguments available for use in the script:

- String tagPath - The full path to the Tag. Example: [tagProvider]Folder/Folder/Tag
- String alarmName - The name of the alarm. This does not include the full alarm path.
- Object alarmEvent - The full alarm event object. See the Alarm Active alarmEvent object for the full list of available properties.
- String alarmPath - The full alarm path.
- Boolean missedEvents - A flag indicating that some events have been skipped due to an event overflow.

Alarm Acknowledged

The Alarm Acknowledged event fires whenever an alarm event becomes acknowledged on the Tag. This event has a variety of arguments available for use in the script:

- String tagPath - The full path to the Tag. Example: [tagProvider]Folder/Folder/Tag
- String alarmName - The name of the alarm. This does not include the full alarm path.
- Object alarmEvent - The full alarm event object. See the Alarm Active alarmEvent object for the full list of available properties.
- String alarmPath - The full alarm path.
- String ackedBy - The full path to the user that acknowledged the alarm.
- Boolean missedEvents - A flag indicating that some events have been skipped due to an event overflow.

Using the Project Library in a Tag Event Script

Scripts defined in a project script can be called from a Tag Event Script. However, only scripts defined in the Gateway Script Project may be used. For more information on configuring the Gateway Scripting Project, please see the [Project Library](#) page.

UDT Parameters in Tag Event Scripts

Parameters on UDTs can be interacted with from Tag Event Scripts. There are two main approaches.

Modern Approach

Parameters values can be accessed as dictionary values. The benefit of this approach is that value changes to UDT parameters will be reflected in subsequent calls. In most cases, this modern approach is preferable. Thus, if the script needs to access the most recent parameter values on a UDT, and the parameters can change through a means other than editing the UDT (which would restart the tag), then this approach should be used.

If trying to access the value of a parameter named "myParam" from a Tag Event Script within the UDT, the script would look like:

```
# Reminder: "tag" is a built-in argument on all Tag Event Scripts. Accessing the "parameters" key on the tag
# argument will
# provide access to all UDT parameters.
paramValue = tag['parameters']['myParam']
```

Legacy Approach

In Ignition release 7.9 and prior, UDT parameters could be accessed in Tag Event Scripts with the familiar curly brackets approach. Thus, if trying to access the value of a parameter named "myParam" from a Tag Event Script within the UDT, the script would look like:

```
paramValue = {myParam}
```

A large caveat with this approach is that value changes made to the parameter ("myParam", in the example above) would not be reflected in scripts until the UDT was restarted. UDT parameters are pre-compiled, which in this case means value changes are mostly ignored until the UDT is restarted.

In all cases, the modern approach above is preferred.

Tag Script Examples

Printing all parameters

```
# This script will fetch all of the possible parameters in the Tag Changed Script.  
# Note that this will not print out to the console, but it will print to the wrapper logs located on the  
Gateway.  
  
path = tagPath  
prev = previousValue  
cur = currentValue  
init = initialChange  
missed = missedEvents  
print path, prev.value, cur.quality, init, missed
```

Automatic Reset

```
# This code can be used on a Value Changed script, and will automatically reset the value of the tag to 0  
after it goes above 3000.  
# This can be useful for counter tags.  
if currentValue.value > 3000:  
    system.tag.writeBlocking(["[default]IntegerCounterTag"], [0])
```

Copy to another Tag

```
# This code can be used on a Value Changed script, and will record the highest value of the current tag onto  
another memory tag.  
# This can be useful for recording the highest temperature.  
highestRecordedTemp = system.tag.readBlocking(["[default]HighestTempTag"])[0].value  
if currentValue.value > highestRecordedTemp:  
    system.tag.writeBlocking(["[default]HighestTempTag"], [currentValue.value])
```

Troubleshooting

It may be helpful when troubleshooting or testing Tag Event Scripts to increase the default threadpool count. Refer to the [Gateway Configuration File Reference - Threadpool Counts](#) for more information.

Tag Properties

Tags are points of data and may have static values or dynamic values that come from an OPC address, an Expression, or a SQL query. The values can be used on screens and in Transaction Groups. Additionally, Tags offer a core set of features above and beyond simple values, such as scaling, alarming, and history logging. Depending on the specific [type of Tag](#), even more options are available. In general, Tags provide a common interface for tying together many types of data in Ignition.

This feature was changed in Ignition version [8.1.17](#):

In 8.1.17, the Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts.

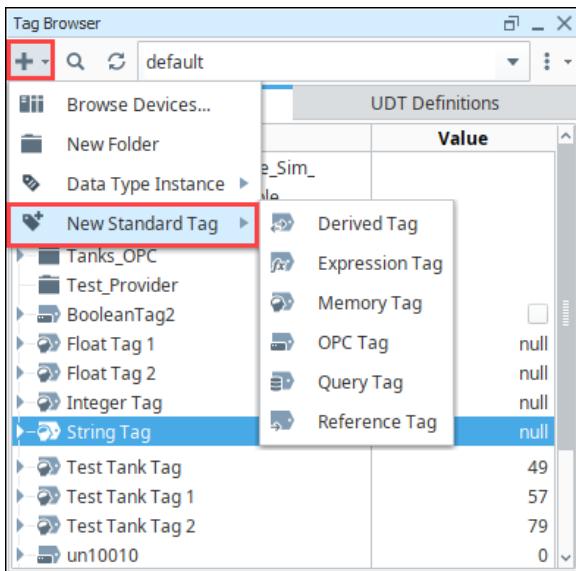
Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

On this page ...

- [Tag Configuration in the Designer](#)
- [Tag Object Types](#)
- [Standard Tag Properties Table](#)
 - [Runtime Properties](#)
 - [Custom Tag Properties](#)
- [Vision Client Tags](#)

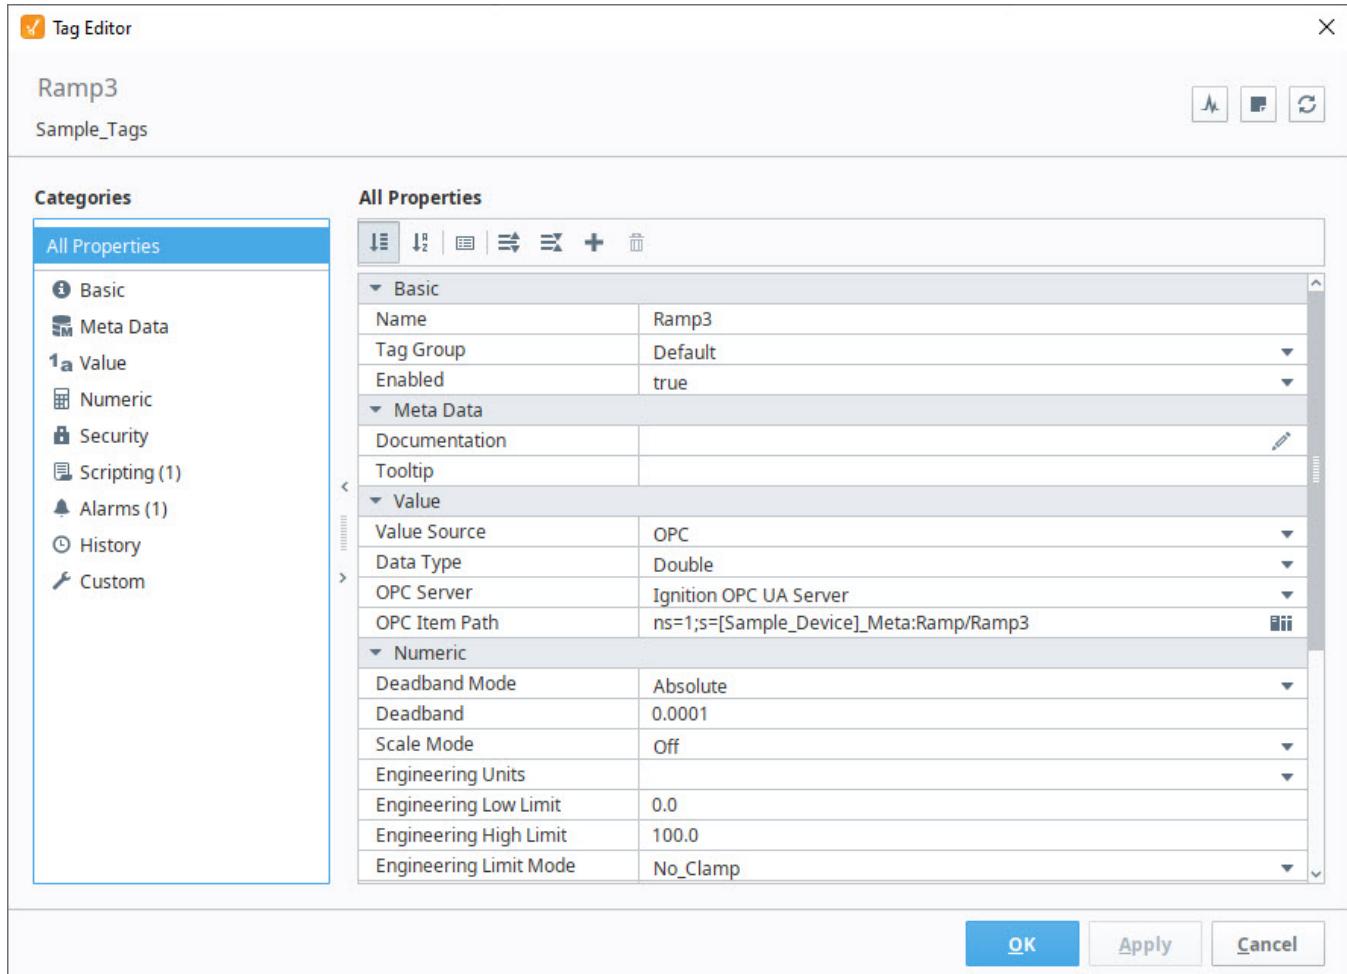
Tag Configuration in the Designer

Tags are managed in the Tag Editor. To configure a Tag, right-click on it and select **Edit Tag**. Or create a new Tag by right-clicking on the Tags folder in the Tag Browser and use the **New Tag** option to select a new Tag type.



Once the **Tag Editor** window is displayed you can set the properties for that Tag. The **Tag Editor** window has the following sections depending on the type of Tag you are creating:

- Basic Properties
- Value
- Numeric Properties
- Meta Data Properties
- Security
- Scripting
- Alarms
- History



Tag Object Types

Some features, such as `system.tag.browse`, can access the Object Type of the tag (sometimes called "tagType"). Below is a table representing the possible types.

Tag Object	Type
Property	A single value underneath an node.
Node	An entity that may have a value and may have children. "Node" is a generic term for other objects in this table, such as a Folder or AtomicTag.
Folder	Represented by a folder in the Tag Browser. Folders generally have child nodes, but don't have values or other properties that make up a tag.
AtomicTag	A "normal" type of tag. Objects with this type can be one of the following (based on the Value Source property): <ul style="list-style-type: none"> • OPC Tag • Query Tag • Expression Tag • Derived Tag • Reference Tag • Memory Tag
UdtInstance	An instance of a complex tag instance (UDT Instance). It's important to note that UdtInstances contain other nodes, so this type is generally only seen at the root of a UDT instance. <p>Thus, nodes under a UdtInstance are not considered to have a type of "UdtInstance", unless the child node is actually a UdtInstance: in other words, a nested UDT instance.</p>

UdtType	Represents the root of a complex tag definition (UDT Definition). Similar to UdtInstance, nodes under a UdtType have their own object type, so a UdtType represents the root of a complex tag.
Provider	Represents a Tag Provider.

Standard Tag Properties Table

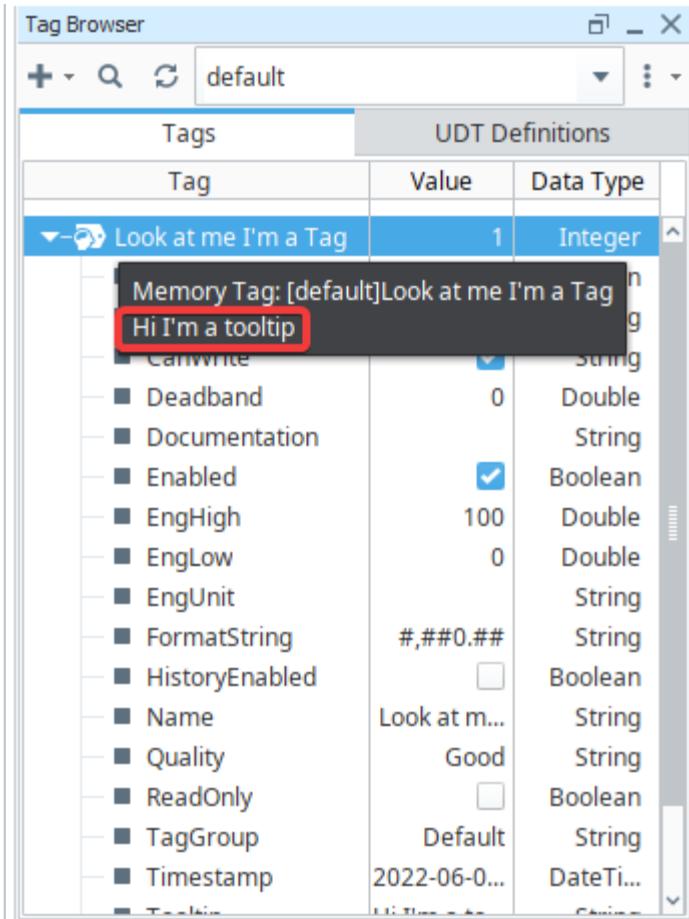
This following table provides detail on each Tag Property, including the binding name, description, data type, and the Tag Types that use the property.

Property	JSON /Scripting Name	Description	Data Type																								
Basic Properties																											
Name	name	How the Tag will be presented and referenced in the system. The Tag path will be the provider, the folder structure, and this name.	String																								
Tag Group	tagGroup	The Tag Group that will execute the Tag. The Tag Group dictates the rate and conditions on which the Tag will be evaluated. For more details, see Tag Groups .	String																								
Enabled	enabled	Whether the Tag will be evaluated by the Tag Group. If false, the Tag will still be present, but will have a bad quality.	Boolean																								
Value																											
Tag Type (unlisted)	tagType	<p>The type of the node. See the Tag Object Types table for more information.</p> <p>Note: This property does not appear in the Tag Editor, but is accessible via scripting.</p>	String																								
Type ID (unlisted)	typeId	<p>Returns a path representing which UDT this node is derived from. If the node is not a UDT, then this property will return a <code>None</code> object.</p> <p>Note: This property does not appear in the Tag Editor, but is accessible via scripting.</p>	String																								
Value Source	valueSource	<p>Specifies how the Tag determines its value. In other words, sets the type of the Tag (Memory, OPC, Expression, etc).</p> <table border="1"> <thead> <tr> <th>Value Source</th><th>JSON Name</th></tr> </thead> <tbody> <tr> <td>Derived</td><td>derived</td></tr> <tr> <td>Expression</td><td>expr</td></tr> <tr> <td>Memory</td><td>memory</td></tr> <tr> <td>OPC</td><td>opc</td></tr> <tr> <td>Query</td><td>db</td></tr> <tr> <td>Reference</td><td>reference</td></tr> </tbody> </table>	Value Source	JSON Name	Derived	derived	Expression	expr	Memory	memory	OPC	opc	Query	db	Reference	reference	String										
Value Source	JSON Name																										
Derived	derived																										
Expression	expr																										
Memory	memory																										
OPC	opc																										
Query	db																										
Reference	reference																										
Data Type	dataType	<p>The data type of the Tag. It is important that this be set as correctly as possible with regards to the Tag's underlying data source. The Tag system will attempt to coerce any raw incoming value (for example, from OPC or a SQL query) into the desired type. For detailed information, see Tag Data Types.</p> <table border="1"> <thead> <tr> <th>Data Type</th><th>String Value</th><th>Integer Value</th></tr> </thead> <tbody> <tr> <td>Byte</td><td>Int1</td><td>0</td></tr> <tr> <td>Short</td><td>Int2</td><td>1</td></tr> <tr> <td>Integer</td><td>Int4</td><td>2</td></tr> <tr> <td>Long</td><td>Int8</td><td>3</td></tr> <tr> <td>Float</td><td>Float4</td><td>4</td></tr> <tr> <td>Double</td><td>Float8</td><td>5</td></tr> <tr> <td>Boolean</td><td>Boolean</td><td>6</td></tr> </tbody> </table>	Data Type	String Value	Integer Value	Byte	Int1	0	Short	Int2	1	Integer	Int4	2	Long	Int8	3	Float	Float4	4	Double	Float8	5	Boolean	Boolean	6	String
Data Type	String Value	Integer Value																									
Byte	Int1	0																									
Short	Int2	1																									
Integer	Int4	2																									
Long	Int8	3																									
Float	Float4	4																									
Double	Float8	5																									
Boolean	Boolean	6																									

		<table border="1"> <tr><td>String</td><td>String</td><td>7</td></tr> <tr><td>DateTime</td><td>DateTime</td><td>8</td></tr> <tr><td>Text</td><td>Text</td><td>10</td></tr> <tr><td>Byte Array</td><td>Int1Array</td><td>17</td></tr> <tr><td>Short Array</td><td>Int2Array</td><td>18</td></tr> <tr><td>Integer Array</td><td>Int4Array</td><td>11</td></tr> <tr><td>Long Array</td><td>Int8Array</td><td>12</td></tr> <tr><td>Float Array</td><td>Float4Array</td><td>19</td></tr> <tr><td>Double Array</td><td>Float8Array</td><td>13</td></tr> <tr><td>Boolean Array</td><td>BooleanArray</td><td>14</td></tr> <tr><td>String Array</td><td>StringArray</td><td>15</td></tr> <tr><td>DateTime Array</td><td>DateTimeArray</td><td>16</td></tr> <tr><td>Binary Data</td><td>ByteArray</td><td>20</td></tr> <tr><td>Dataset</td><td>DataSet</td><td>9</td></tr> <tr><td>Document</td><td>Document</td><td>29</td></tr> </table>	String	String	7	DateTime	DateTime	8	Text	Text	10	Byte Array	Int1Array	17	Short Array	Int2Array	18	Integer Array	Int4Array	11	Long Array	Int8Array	12	Float Array	Float4Array	19	Double Array	Float8Array	13	Boolean Array	BooleanArray	14	String Array	StringArray	15	DateTime Array	DateTimeArray	16	Binary Data	ByteArray	20	Dataset	DataSet	9	Document	Document	29	
String	String	7																																														
DateTime	DateTime	8																																														
Text	Text	10																																														
Byte Array	Int1Array	17																																														
Short Array	Int2Array	18																																														
Integer Array	Int4Array	11																																														
Long Array	Int8Array	12																																														
Float Array	Float4Array	19																																														
Double Array	Float8Array	13																																														
Boolean Array	BooleanArray	14																																														
String Array	StringArray	15																																														
DateTime Array	DateTimeArray	16																																														
Binary Data	ByteArray	20																																														
Dataset	DataSet	9																																														
Document	Document	29																																														
		<p>Note: Regarding Array data types, Alarming, Scaling, and Historical settings applied to an array Tag are propagated down to elements in the array.</p>																																														
Value	value	The value of the Tag. Can only be modified if the Tag allows value writing and the user has sufficient privileges.	Object (depends on the data type of the Tag)																																													
OPC Server	opcServer	Only present for OPC Tags. The server against which to subscribe the data point.	String																																													
OPC Item Path	opcItemPath	Only present for OPC Tags. The path to subscribe to on the server. The point will be subscribed at the rate dictated by the Tag Group. <div style="border: 1px solid orange; padding: 5px; margin-top: 10px;"> The following feature is new in Ignition version 8.1.5 Click here to check out the other new features </div> <p>It's possible to escape curly braces {} in the item path by using additional curly braces. For example: {{device_name}} would evaluate to {<device_name value>}, allowing you to include braces in the Item Path.</p>	String																																													
Source Tag Path	sourceTagPath	The path to the Tag that this Tag is referencing. Only present for Reference and Derived Tags.	String																																													
Execution Mode	executionMode	Only present for Query and Expression Tags . Determines how when the Tag executes. <ul style="list-style-type: none"> • Event Driven - Updates when something happens (i.e., value event or alarm event) within the expression. • Fixed Rate - Tag will be executed at the set or fixed rate. Adds the Execution Rate property, which determines how often the Tag executes in milliseconds. • Tag Group - Tags are executed by Tag Groups, which dictate the rate of execution. <table border="1" style="margin-top: 10px;"> <thead> <tr> <th>Execution Mode</th><th>JSON Name</th></tr> </thead> <tbody> <tr><td>Event Driven</td><td>EventDriven</td></tr> <tr><td>Fixed Rate</td><td>FixedRate</td></tr> <tr><td>Tag Group</td><td>TagGroupRate</td></tr> </tbody> </table>	Execution Mode	JSON Name	Event Driven	EventDriven	Fixed Rate	FixedRate	Tag Group	TagGroupRate	String																																					
Execution Mode	JSON Name																																															
Event Driven	EventDriven																																															
Fixed Rate	FixedRate																																															
Tag Group	TagGroupRate																																															
Expression	expression	The expression the Tag will use to determine its value.	String																																													
Read Expression	deriveExpressionGetter	The expression that determines how the value on the Derived Tag is displayed.	String																																													
Query	query	The SQL query to be run, which drives the tag's value. Queries doing database reads and writes are possible, see the Query Type property description for details.																																														
Write Expression		The expression that determines how the value on the Derived Tag is displayed.	String																																													

	deriveExpressionSetter																				
Datasource	datasource	The database connection that the query Tag will execute against.	String																		
Query Type	queryType	<p>The following feature is new in Ignition version 8.1.3 Click here to check out the other new features</p> <p>Defines whether the query is executing a database read or a database write. Important for determining the value behavior of the Tag.</p> <p>Possible values are:</p> <ul style="list-style-type: none"> • AutoDetect - Query type is determined from the query itself. • Select - Dictates that the query is reading data from the database. The query result set will be stored on the Tag's value. • Update - Dictates that the query is writing data to the database (but does not require an UPDATE statement in the query, specifically). The value on the query Tag will be the number of affected rows. 	String																		
Numeric Properties																					
Deadband	deadband	A numerical value used to prevent unnecessary updates for Tags whose values change by small amounts.	Numeric																		
Deadband Mode	deadbandMode	<p>Defines how the deadband value is used.</p> <ul style="list-style-type: none"> • Absolute - The deadband setting is considered to be an absolute value. • Percent - The actual deadband is calculated as a percent of the Tag's engineering unit span. <p>The following feature is new in Ignition version 8.1.2 Click here to check out the other new features</p> <ul style="list-style-type: none"> • Off - The deadband setting is the equivalent to a value of 0.0, so that all values pass through if their timestamp has changed. <p>Valid values are as follows:</p> <table border="1"> <thead> <tr> <th>Deadband Mode</th><th>JSON Name</th></tr> </thead> <tbody> <tr> <td>Absolute</td><td>Absolute</td></tr> <tr> <td>Percent</td><td>Percent</td></tr> <tr> <td>Off</td><td>Off</td></tr> </tbody> </table>	Deadband Mode	JSON Name	Absolute	Absolute	Percent	Percent	Off	Off	String										
Deadband Mode	JSON Name																				
Absolute	Absolute																				
Percent	Percent																				
Off	Off																				
Scale Mode	scaleMode	If and how the Tag value will be scaled between the source, and what is reported for the Tag.	String																		
		Valid values are as follows:																			
		<table border="1"> <thead> <tr> <th>Mode</th><th>JSON Name</th><th>Int Value</th></tr> </thead> <tbody> <tr> <td>Off</td><td>Off</td><td>0</td></tr> <tr> <td>Linear</td><td>Linear</td><td>1</td></tr> <tr> <td>Square Root</td><td>SquareRoot</td><td>2</td></tr> <tr> <td>Exponential Filter</td><td>ExponentialFilter</td><td>3</td></tr> <tr> <td>Bit Inversion</td><td>BitInversion</td><td>4</td></tr> </tbody> </table>	Mode	JSON Name	Int Value	Off	Off	0	Linear	Linear	1	Square Root	SquareRoot	2	Exponential Filter	ExponentialFilter	3	Bit Inversion	BitInversion	4	
Mode	JSON Name	Int Value																			
Off	Off	0																			
Linear	Linear	1																			
Square Root	SquareRoot	2																			
Exponential Filter	ExponentialFilter	3																			
Bit Inversion	BitInversion	4																			
Raw Low	rawLow	Start of the "raw" value range. Only present if Scale Mode is set to Linear or Square Root .	Numeric																		
Raw High	rawHigh	End of the "raw" value range. Only present if Scale Mode is set to Linear or Square Root .	Numeric																		
Scaled Low	scaledLow	Start of "scaled" value range. Raw low will map to Scaled low for the Tag. Only present if Scale Mode is set to Linear or Square Root .	Numeric																		
Scaled High	scaledHigh	End of "scaled" value range. Raw high will map to Scaled high for the Tag. Only present if Scale Mode is set to Linear or Square Root .	Numeric																		
Clamp Mode	clampMode	How values that fall outside of the ranges will be treated. "Clamped" values will be adjusted to the low/high scaled value as appropriate. Only present if Scale Mode is set to Linear or Square Root .	String																		
		Valid values are as follows:																			
		<table border="1"> <thead> <tr> <th>Clamp Mode</th><th>JSON Name</th><th>Int Value</th></tr> </thead> </table>	Clamp Mode	JSON Name	Int Value																
Clamp Mode	JSON Name	Int Value																			

		<table border="1"> <tr><td>No_Clamp</td><td>No_Clamp</td><td>0</td></tr> <tr><td>Clamp_Low</td><td>Clamp_Low</td><td>1</td></tr> <tr><td>Clamp_High</td><td>Clamp_High</td><td>2</td></tr> <tr><td>Clamp_Both</td><td>Clamp_Both</td><td>3</td></tr> </table>	No_Clamp	No_Clamp	0	Clamp_Low	Clamp_Low	1	Clamp_High	Clamp_High	2	Clamp_Both	Clamp_Both	3				
No_Clamp	No_Clamp	0																
Clamp_Low	Clamp_Low	1																
Clamp_High	Clamp_High	2																
Clamp_Both	Clamp_Both	3																
Scale Factor	scaleFactor	For single parameter modes (currently only Exponential Filter), the factor parameter for the equation. Used when the Scale Mode property is set to Exponential Filter	Numeric															
Engineering Units	engUnit	The engineering units of the value.	String															
Engineering Low Limit	engLow	The lowest expected value of the Tag.	Numeric															
Engineering High Limit	engHigh	The highest expected value of the Tag.	Numeric															
Engineering Limit Mode	engLimitMode	<p>Dictates how the engineering range should be enforced on the Tag. If not "Off", the Tag will change to bad quality ("limit exceeded"), when the value falls outside the specified range.</p> <p>Valid values are as follows:</p> <table border="1"> <thead> <tr><th>Limit Enforcement</th><th>JSON Name</th><th>Int Value</th></tr> </thead> <tbody> <tr><td>No_Clamp</td><td>No_Clamp</td><td>0</td></tr> <tr><td>Clamp_Low</td><td>Clamp_Low</td><td>1</td></tr> <tr><td>Clamp_High</td><td>Clamp_High</td><td>2</td></tr> <tr><td>Clamp_Both</td><td>Clamp_Both</td><td>3</td></tr> </tbody> </table>	Limit Enforcement	JSON Name	Int Value	No_Clamp	No_Clamp	0	Clamp_Low	Clamp_Low	1	Clamp_High	Clamp_High	2	Clamp_Both	Clamp_Both	3	Numeric
Limit Enforcement	JSON Name	Int Value																
No_Clamp	No_Clamp	0																
Clamp_Low	Clamp_Low	1																
Clamp_High	Clamp_High	2																
Clamp_Both	Clamp_Both	3																
Format String	formatString	<p>How the value should be formatted when converted to a string (only applies to numerical data types). Uses # and 0 characters to describe the format.</p> <p># : If the number in this position is non-zero, then do not show the position. Otherwise, show the number. Useful when you only want to show a decimal position if the value is non-zero.</p> <p>0 : If the number in this position is non-zero, then show that number. Otherwise, show a zero. Useful to add leading and trailing zeros to a value.</p> <p>See Data Type Formatting Reference.</p>	String															
Meta Data Properties																		
Tooltip	tooltip	<p>The tooltip provides a hint to visual components as to what should be displayed when the user hovers their mouse cursor over the component that is being driven by the value of this Tag.</p> <div style="border: 1px solid orange; padding: 5px; margin-top: 10px;"> <p>The following feature is new in Ignition version 8.1.18 Click here to check out the other new features</p> </div> <p>Hovering over the Tag itself in the Tag Browser will also display this hint.</p>	String															



Documentation	documentation	A freeform text property for information about the Tag.	String
---------------	---------------	---	--------

Security			
Read Permissions	readPermissions	Defines the security levels required in order to read values from a Tag. For more information, see Tag Security Properties . Contains the following elements:	JSON Object

Name	JSON Name	Description
Type	type	Represents the selected radio button on the security level UI, determining if all of the elements in the securityLevels array are required, or if any of the elements are allowed. Possible values are: AnyOF AllOF
Security Levels	securityLevels	Represents allowed security levels for this permission. Each level is represented as a JSON object, containing a "name" value that represents the name of a security level, and a "children" array which represents any levels under the current. The actual "selected" levels are any levels that have an empty "children" object. See the example below for more information.

The JSON in this example uses the configuration shown in the image below. Permission is granted if the security levels on the request are from either an "Administrator" user, or if the request originated from the "Zone A" Security Zone.

Security Levels (including Roles) may be added by going to the [Config > Security > Security Levels](#) page of the Gateway Web Interface

The security levels of the user must match **all** of the required security levels
At least one of the security levels of the user must match **any** of the required security levels

```
"readPermissions": {
    "type": "AnyOf",
    "securityLevels": [
        {
            "name": "Authenticated",
            "children": [
                {
                    "name": "Roles",
                    "children": [
                        {
                            "name": "Administrator",
                            "children": []
                        }
                    ]
                }
            ]
        },
        {
            "name": "SecurityZones",
            "children": [
                {
                    "name": "Zone A",
                    "children": []
                }
            ]
        }
    ]
}
```

Read Only	readOnly	Defines whether a Tag is read-only or writeable. For more information, see Tag Security Properties .	value: boolean
Write Permissions	writePermissions	Defines the security levels required in order to read values from a Tag . For more information, see Tag Security Properties . Contains the following elements:	JSON Object

Name	JSON Name	Description
Type	type	

		Represents the selected radio button on the security level UI, determining if all of the elements in the securityLevels array are required, or if any of the elements are allowed. Possible values are: AnyOf AllOf
Security Levels	securityLevels	Represents allowed security levels for this permission. Each level is represented as a JSON object, containing a "name" value that represents the name of a security level, and a "children" array which represents any levels under the current. The actual "selected" levels are any levels that have an empty "children" object. See the example below for more information.

The JSON in this example uses the configuration shown in the image below. Permission is granted if the security levels on the request are from either an "Administrator" user, or if the request originated from the "Zone A" Security Zone.

 Security Levels (including Roles) may be added by going to the [Config > Security > Security Levels](#) page of the Gateway Web Interface

The security levels of the user must match **all of** the required security levels

At least one of the security levels of the user must match **any of** the required security levels

```
"writePermissions": {
    "type": "AnyOf",
    "securityLevels": [
        {
            "name": "Authenticated",
            "children": [
                {
                    "name": "Roles",
                    "children": [
                        {
                            "name": "Administrator",
                            "children": []
                        }
                    ]
                }
            ]
        },
        {
            "name": "SecurityZones",
            "children": [
                {
                    "name": "Zone A",
                    "children": []
                }
            ]
        }
    ]
}
```

Scripting																					
Tag Event Scripts	eventScripts	<p>Each Tag has the option to have Tag Event Scripts on it. When you edit a Tag, you can navigate to the Tag Events screen to see a list of all of the Tag scripts. You can then select which event you would like to write a script for. You can even write a script for multiple events if you like. For detailed information, see Tag Event Scripts.</p> <p>When interacting with a Tag from a script, the Tag Event Scripts are represented as an array of JSON objects. Each JSON object is described in the expandable area below:</p> <p>Key Description</p> <table border="1"> <thead> <tr> <th>Key</th><th>Description</th></tr> </thead> <tbody> <tr> <td>eventid</td><td>A value representing the type of event script</td></tr> <tr> <td>script</td><td>A value representing the content of the script</td></tr> </tbody> </table> <p>Possible eventid values</p> <table border="1"> <thead> <tr> <th>Event Script</th><th>eventid value</th></tr> </thead> <tbody> <tr> <td>Quality Changed</td><td>qualityChanged</td></tr> <tr> <td>Value Changed</td><td>valueChanged</td></tr> <tr> <td>Alarm Active</td><td>alarmActive</td></tr> <tr> <td>Alarm Cleared</td><td>alarmCleared</td></tr> <tr> <td>Alarm Acknowledged</td><td>alarmAcked</td></tr> </tbody> </table>	Key	Description	eventid	A value representing the type of event script	script	A value representing the content of the script	Event Script	eventid value	Quality Changed	qualityChanged	Value Changed	valueChanged	Alarm Active	alarmActive	Alarm Cleared	alarmCleared	Alarm Acknowledged	alarmAcked	JSON Array
Key	Description																				
eventid	A value representing the type of event script																				
script	A value representing the content of the script																				
Event Script	eventid value																				
Quality Changed	qualityChanged																				
Value Changed	valueChanged																				
Alarm Active	alarmActive																				
Alarm Cleared	alarmCleared																				
Alarm Acknowledged	alarmAcked																				
Alarms																					
Alarms	alarms	<p>Tags have the ability to define any number of alarms. Each alarm is a condition that will be evaluated when the value of the Tag changes. When the condition becomes true, the alarm is said to be active. When it becomes false, the alarm is said to be cleared.</p> <p>For detailed information, see Tag Alarm Properties.</p>	JSON Array of JSON objects. For detailed information see Tag Alarm Properties																		
Alarm Eval Enabled	alarmEvalEnabled	Determines if alarms will be evaluated on this tag.	Boolean																		
History																					
History Enabled	historyEnabled	Whether the Tag will report its history to the Tags Historian system.	Boolean																		
Storage Provider	historyProvider	<p>Which Tag Historian data store the Tag will target. A particular Tag can only target one history store. For more information, refer to History Providers on the Tag History Gateway Settings page.</p> <p>Note: The Storage Provider dropdown displays the provider names as they are written at the time of tag configuration. If the Storage Provider name is updated later, this setting will need to be adjusted to match the new Storage Provider name.</p>	String																		
Deadband Style	historicalDeadbandStyle	<p>There are three styles to choose from: Auto, Analog, or Discrete.</p> <p>When set to Auto, this setting will automatically pick from Analog or Discrete, based on the data type of the Tag.</p> <ul style="list-style-type: none"> If the data type of the Tag is set to a float or double, then Auto will use the Analog Style If the data type of the Tag is any other type, then the Discrete style will be used. <p>More information on the Analog and Discrete types can be found on the Configuring Tag History page.</p> <p>Valid values are as follows:</p> <table border="1"> <thead> <tr> <th>Deadband Style</th><th>JSON Name</th></tr> </thead> <tbody> <tr> <td>Auto</td><td>Auto</td></tr> <tr> <td>Analog</td><td>Analog_Compressed</td></tr> <tr> <td>Discrete</td><td>Discrete</td></tr> </tbody> </table>	Deadband Style	JSON Name	Auto	Auto	Analog	Analog_Compressed	Discrete	Discrete	String										
Deadband Style	JSON Name																				
Auto	Auto																				
Analog	Analog_Compressed																				
Discrete	Discrete																				
Deadband Mode	historicalDeadbandMode	Defines how the deadband value is used.	String																		

- **Absolute** - The deadband setting is considered to be an absolute value.
- **Percent** - The actual deadband is calculated as a percent of the Tag's engineering unit span.

The following feature is new in Ignition version **8.1.2**
[Click here](#) to check out the other new features

- **Off** - The deadband setting is the equivalent to a value of 0.0.

Deadband Mode	JSON Name
Absolute	Absolute
Percent	Percent
Off	Off

Historical Deadband	historicalDeadband	A deadband that applies only to historical evaluation.	Numeric																		
Sample Mode	sampleMode	<p>Determines how often a historical record should be collected.</p> <ul style="list-style-type: none"> • On Change - Collects a record whenever the value on the Tag changes. • Periodic - Collects a record based on the Sample Rate and Sample Rate Units properties. • Tag Group - Collects a record based on the Tag Group specified under the Historical Tag Group property. <p>Valid values are as follows:</p> <table border="1"> <thead> <tr> <th>Max Time Between Records Mode</th> <th>JSON Name</th> </tr> </thead> <tbody> <tr> <td>On Change</td><td>OnChange</td></tr> <tr> <td>Periodic</td><td>Periodic</td></tr> <tr> <td>Tag Group</td><td>TagGroup</td></tr> </tbody> </table>	Max Time Between Records Mode	JSON Name	On Change	OnChange	Periodic	Periodic	Tag Group	TagGroup	String										
Max Time Between Records Mode	JSON Name																				
On Change	OnChange																				
Periodic	Periodic																				
Tag Group	TagGroup																				
Sample Rate	historySampleRate	When the Sample Mode property is set to "Periodic", this property (working in conjunction with the Sample Rate Units property) determines how often a record should be collected.	Numeric																		
Sample Rate Units	historySampleRateUnits	When the Sample Mode property is set to "Periodic", this property (working in conjunction with the Sample Rate property) determines the unit of time that will be used in record collection.	String																		
Historical Tag Group	historyTagGroup	When the Sample Mode property is set to "Tag Group", this property determines which Tag Group will be used to collect records.	String																		
Min Time Between Samples	historyTimeDeadband	Minimum time between records. Useful in restricting the number of records collected when the Sample Mode is set to "Tag Change". Prevents multiple consecutive Tag changes from triggering consecutive record collections. Works in conjunctions with the Min Time Units property. The Value is calculated off of the value timestamp.	Integer																		
Min Time Units	historyTimeDeadbandUnits	<p>Units of time to use with the Min Time Between Samples property.</p> <table border="1"> <thead> <tr> <th>Unit of Time</th> <th>JSON Name</th> </tr> </thead> <tbody> <tr> <td>Milliseconds</td><td>MS</td></tr> <tr> <td>Seconds</td><td>SEC</td></tr> <tr> <td>Minutes</td><td>MIN</td></tr> <tr> <td>Hour</td><td>HOUR</td></tr> <tr> <td>Day</td><td>DAY</td></tr> <tr> <td>Week</td><td>WEEK</td></tr> <tr> <td>Month</td><td>MONTH</td></tr> <tr> <td>Year</td><td>YEAR</td></tr> </tbody> </table>	Unit of Time	JSON Name	Milliseconds	MS	Seconds	SEC	Minutes	MIN	Hour	HOUR	Day	DAY	Week	WEEK	Month	MONTH	Year	YEAR	String
Unit of Time	JSON Name																				
Milliseconds	MS																				
Seconds	SEC																				
Minutes	MIN																				
Hour	HOUR																				
Day	DAY																				
Week	WEEK																				
Month	MONTH																				
Year	YEAR																				

		<table border="1"> <tr><td>Day</td><td>DAY</td></tr> <tr><td>Week</td><td>WEEK</td></tr> <tr><td>Month</td><td>MONTH</td></tr> <tr><td>Year</td><td>YEAR</td></tr> </table>	Day	DAY	Week	WEEK	Month	MONTH	Year	YEAR											
Day	DAY																				
Week	WEEK																				
Month	MONTH																				
Year	YEAR																				
Max Time Between Samples	historyMaxAge	<p>Maximum time between samples. Works in conjunction with the Max Time Units property. If a sample has not been collected by the time range specified by these two properties, then a record will be collected on the next sample interval. Setting this value to 0 will disable automatic record collection. Default is 0.</p> <p>Note: This setting will be ignored if the Sample Mode is set to Tag Group, and the targeted Tag Group is using non-default values for its Max Time Between Samples setting. The implication being that non-default values on the Tag Group settings take precedence over this setting.</p>	Integer																		
Max Time Units	historyMaxAgeUnits	<p>Maximum time in units is defined as: Milliseconds, Seconds, Minutes, Hours, Days, Weeks, Months, and Years.</p> <table border="1"> <thead> <tr> <th>Unit of Time</th><th>JSON Name</th></tr> </thead> <tbody> <tr><td>Milliseconds</td><td>MS</td></tr> <tr><td>Seconds</td><td>SEC</td></tr> <tr><td>Minutes</td><td>MIN</td></tr> <tr><td>Hour</td><td>HOUR</td></tr> <tr><td>Day</td><td>DAY</td></tr> <tr><td>Week</td><td>WEEK</td></tr> <tr><td>Month</td><td>MONTH</td></tr> <tr><td>Year</td><td>YEAR</td></tr> </tbody> </table>	Unit of Time	JSON Name	Milliseconds	MS	Seconds	SEC	Minutes	MIN	Hour	HOUR	Day	DAY	Week	WEEK	Month	MONTH	Year	YEAR	String
Unit of Time	JSON Name																				
Milliseconds	MS																				
Seconds	SEC																				
Minutes	MIN																				
Hour	HOUR																				
Day	DAY																				
Week	WEEK																				
Month	MONTH																				
Year	YEAR																				

Runtime Properties

In addition to properties listed in the Tag Editor, some properties are exposed as runtime properties in the Tag Browser. These properties are valid targets for component bindings and tag reads/writes.

My Tag	10
AlarmEvalEnabled	<input checked="" type="checkbox"/>
CanRead	<input checked="" type="checkbox"/>
CanWrite	<input checked="" type="checkbox"/>
Deadband	0
Documentation	
Enabled	<input checked="" type="checkbox"/>
EngHigh	100
EngLow	0
EngUnit	
FormatString	#,##0.##
HistoryEnabled	<input type="checkbox"/>
Name	My Tag
Quality	Good
ReadOnly	<input type="checkbox"/>
TagGroup	Default
Timestamp	2021-07-23 1:24:42 PM
Tooltip	
value	10
ValueSource	memory
Alarms	

Most runtime properties are representations of properties that can be configured in the Tag Editor. However there are some properties only listed in the Tag Browser:

Property	Description
CanRead	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>A read-only property that represents whether or not this tag can be read from the current security context. This is determined by looking at the read permission settings on the tag and the tag provider's permission settings.</p>
CanWrite	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>A read-only property that represents whether or not this tag can be written to from the current security context. This is determined by looking at the write permission settings on the tag, the Read Only property, and the tag provider's permission settings.</p>

Custom Tag Properties

Custom Tag properties allow application designers to configure their own properties on Tags to store unique values on any Tag. Once added, a custom property can be referenced like any other Tag property via bindings, expressions, and scripts.

Both the Perspective and Vision visualization systems can bind to custom properties. In the following example, we already have Array Tag. Now let's add a custom property.

1. Open an existing Tag and select the Custom properties category.
2. Click on the Add  icon in the Tag Editor. This will open the Custom Property dialog box.

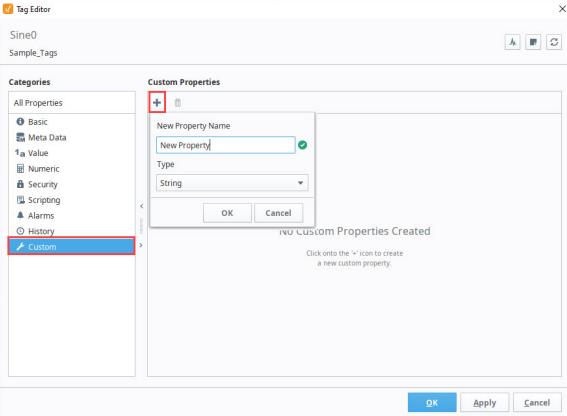


**INDUCTIVE
UNIVERSITY**

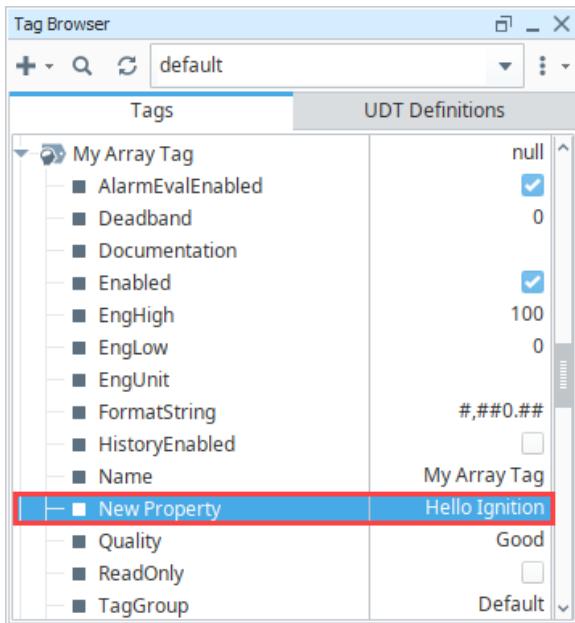
**Custom Tag
Properties**

[Watch the Video](#)

3. Enter a **Name** for your property, select the **Data Type**, and click **OK**.



4. When you open your Tag Browser and expand the My Array Tag, you'll see your Custom Property.



Vision Client Tags

Client Tags have the ability to be used as either Expression or SQL Query Tags. There is an **Expression/SQL** page in the Tag editor that allows you to select what type it is.

Query/Expression Attributes						
Property	Binding /Scripting Name	Description	Data Type	Applicable Tag Type		
OPC Server	OPCWriteBackServer	The server against which to subscribe the data point.	String	Query, Expression		
OPC Item Path	OPCWriteBackItemPath	The path to subscribe to on the server.	String	Query, Expression		
Query	Expression	Text area to build your query or expression.	String	Query, Expression		
Query Type	QueryType	When the TagType property is set to 1, this property determines if the Tag should be a Memory, Expression, or Query Tag.	Integer <table border="1" style="margin-left: auto; margin-right: auto;"> <tr> <td>Tag Type</td> <td>Integer Value</td> </tr> </table>	Tag Type	Integer Value	Query, Expression, Memory
Tag Type	Integer Value					

			<table border="1"> <tr> <td>Memory Tag</td><td>0</td></tr> <tr> <td>Expression Tag</td><td>1</td></tr> <tr> <td>Query Tag</td><td>2</td></tr> </table>	Memory Tag	0	Expression Tag	1	Query Tag	2	
Memory Tag	0									
Expression Tag	1									
Query Tag	2									
Datasource	SQLBindingDataSource	The default data source of the Tag provider.	String	Query						

Tag Alarm Properties

This feature was changed in Ignition version 8.1.17:

In 8.1.17, the Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts.

Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

Alarm Property Reference

The table in this section provides a description of alarm properties.

Aarms in Scripting

When interacting with the tags system in via scripting, such as with the `system.tag.configure` function, alarms are represented as a JSON array of JSON objects, where each object contains the configurations for a single alarm. Thus, any scripting names here are assumed to exist under the alarms array.

Reference Table

The descriptions of these properties can also be viewed in the Tag Editor > Alarms popup by selecting the Show/Hide Description Area  icon available above the property list.

Property Name	JSON /Scripting Name	Description	Data						
Main									
Name	name	Identifier of the alarm. Combined with the location of the alarm, this will be the unique alarm ID. For dynamic values, used Label or Display Path.	String						
Enabled	enabled	This boolean determines whether or not the alarm will be evaluated. A disabled alarm's condition will not be evaluated, and thus will not generate any alarm events.	Boolean						
Priority	priority	The priority or severity of this alarm. Alarm priorities can be examined by many other systems in Ignition, including the visualization modules, pipelines, and even scripting.	<table border="1"><tr><td>Priority</td></tr><tr><td>Diagnostic</td></tr><tr><td>Low</td></tr><tr><td>Medium</td></tr><tr><td>High</td></tr><tr><td>Critical</td></tr></table>	Priority	Diagnostic	Low	Medium	High	Critical
Priority									
Diagnostic									
Low									
Medium									
High									
Critical									
Timestamp Source	timestampSource	Indicates where the timestamp for the alarm event should come from: the system time of when the event was generated (i.e., the Gateway's time), or the timestamp of the value that triggered the event (i.e., the timestamp of the value from the OPC server).	<table border="1"><tr><td>Timestamp Source</td></tr><tr><td>System</td></tr><tr><td>Value</td></tr></table>	Timestamp Source	System	Value			
Timestamp Source									
System									
Value									
Label	label	An optional name that will be used for display purposes. Provides a dynamic alternative to the static <code>name</code> property. If left blank, the <code>name</code> will be used.	String						
Display Path	displayPath	Optional text that will be used for display and browsing purposes. If this is blank, this property will show the path to the Tag and the name of the alarm instead.	String						
Ack Mode	ackMode	Dictates how acknowledgement works for the alarm.	String						

On this page ...

- [Alarm Property Reference](#)
 - [Alarms in Scripting](#)
 - [Reference Table](#)
- [Runtime Alarm Metric Properties](#)
 - [Runtime Alarm Metrics for Individual Alarms](#)
- [Binding](#)

		<ul style="list-style-type: none"> Unused - Any alarm event that is generated will automatically be marked as acknowledged. Auto - The alarm is acknowledged automatically when the alarm becomes cleared. Manual - The alarm is never set to be acknowledged by the system, and it is up to the user to manually acknowledge alarms. 	Ack Mode
Notes	notes	A place for any free-form documentation about the alarm that can be displayed to users.	String
Ack Notes Required	ackNotesReqd	If this setting is true, the operators will be required to provide some explanation when the alarm is acknowledged.	Boolean
Shelving Allowed	shelvingAllowed	If this setting is false, the shelving feature will be unavailable for this alarm.	Boolean

Alarm Mode Settings

Mode	mode	<p>This setting controls what condition this alarm is evaluating. The available modes are as follows:</p> <ul style="list-style-type: none"> Equal - Active when the Tag's value equals the alarm's setpoint. Not Equal - Active when the Tag's value does not equal the alarm's setpoint. Above Setpoint - Active when the Tag's value is above the alarm's setpoint. Below Setpoint - Active when the Tag's value is below the alarm's setpoint. Between Setpoints - Active when the Tag's value is between the low and high setpoints. If any change is true, an event will be generated for each value change between the setpoints. Outside Setpoints - Active when the Tag's value falls outside the low and high setpoints. If any change is true, an event will be generated for each value change outside the setpoints. Out of Range - The same as Outside Setpoints, but uses the Tag's Engineering High and Engineering Low as the high and low setpoints. Bad Quality - Active if the Tag value becomes a bad quality, for example, on communication loss. Any Change - An alarm event is generated every time the Tag value changes. <p>Note: This alarm will never be "active" because each active event is paired with a matching clear event, instantly.</p> <ul style="list-style-type: none"> Bit State - This alarm mode is used to alarm when a specific bit out of an integer Tag becomes high. You must specify which bit position to use, with zero being the least significant bit. The On Zero property is used to invert the logic and alarm when the bit is low. On Condition - This free-form alarm mode is used for when you want to specify the condition using an expression or another Tag. To do this, bind the "Is Active" property to an appropriate expression or Tag. 	String
Setpoint/Low Setpoint	setpointA	<p>Used to determine if the alarm is active by comparing this value to the tag value.</p> <p>Some modes under the Mode property allow for multiple setpoints (i.e., a low setpoint and a high setpoint). In these cases, this property is considered to be the Low setpoint.</p> <p>Note: Setpoints must be numerical values, not string representations of numbers. If a string is passed to a setpoint then it will be converted to a numeric value of 0.</p>	Numeric
Inclusive	inclusiveA	If true, the Setpoint value is used inclusively for the condition to alarm.	Boolean
High Setpoint	setpointB	The high value used to initiate an alarm when the alarm mode calls for two setpoints. Available for modes: Between Setpoint, Outside Setpoints.	Numeric
High Inclusive	inclusiveB	If true, the High Setpoint value is used inclusively for the condition to alarm. Available for modes: Between Setpoint, Outside Setpoints.	Boolean
Any Change	anyChange	If true, will alarm on each value change given the alarm mode conditions are met.	Boolean
On Zero	bitOnZero	<p>If true, will alarm when the specified bit is not high (when the bit is 0).</p>	Boolean
Bit Position	bitPosition	The position of the bit, starting at 0 that will be watched. Available for modes: Bit State.	Numeric

Is Active	activeCondition	When this property is active, the alarm will be active. Typically has a binding of some sort that will be used to determine when the alarm goes active. If the expression evaluates to True, the alarm is active. If the expression evaluates to False, the alarm is not active.	Boolean
-----------	-----------------	--	---------

Deadbands and Time Delays

Deadband	deadband	The value for the deadband, interpreted according to the Deadband mode. Note: All alarms are only evaluated after the Tag's value changes, which means that the Tag's own deadband will be considered first. When the deadband is positive, an active alarm condition needs to clear its setpoint(s) by the amount of the deadband for the alarm to clear. For example, suppose you had a Between Setpoints alarm with a low setpoint of 50 and a high setpoint of 70, and with a deadband of 2. The alarm will go active if the value is between 50 and 70, but will only clear if the value falls below 48 or rises above 72.	Numeric
Deadband Mode	deadbandMode	Defines how the deadband value is used. <ul style="list-style-type: none">• Absolute - The deadband setting is considered to be an absolute value.• Percent - The actual deadband is calculated as a percent of the Tag's engineering unit span. <div style="background-color: #ffccbc; padding: 10px; border: 1px solid #ff9800; border-radius: 5px; width: fit-content; margin-left: auto; margin-right: 0;"><p>The following feature is new in Ignition version 8.1.2 Click here to check out the other new features</p></div> <ul style="list-style-type: none">• Off - The deadband setting is the equivalent to a value of 0.0.	<div style="display: flex; align-items: center; justify-content: space-between;"> Numeric <div style="flex-grow: 1; text-align: right;"> Alarmin Mode Absolute Percent Off </div> </div>
Active Delay	timeOnDelaySeconds	The time, in seconds, before the alarm will be considered active after the alarm's condition becomes true. Also known as a "rising edge time deadband."	Numeric
Clear Delay	timeOffDelaySeconds	The time, in seconds, before an active alarm will be considered clear after the alarm's condition becomes false. Also known as a "falling edge time deadband."	Numeric

Notification Properties

Active Pipeline	activePipeline	The name of an alarm notification pipeline to put this alarm into when it becomes active in order to send out active alarm messages. Many alarms may share a single pipeline.	String
Clear Pipeline	clearPipeline	The name of an alarm notification pipeline to put this alarm into when it becomes cleared in order to send out cleared messages.	String
Ack Pipeline	ackPipeline	The name of the alarm notification pipeline to put this alarm into when the alarm has been acknowledged.	String

Phone Call Settings

Custom Message	voip.customMessage	The voice message you want to play when sending the Alarm. See Configuring Messages for more information on custom phone messages.	String
----------------	--------------------	--	--------

Email Notification Properties

Custom Subject	CustomEmailSubject	A string that will be used as the subject line of an email notification message. If blank, the message settings defined on the notification block that sent the email out will be used instead.	String
Custom Message	CustomEmailMessage	A string that will be used as the body of this alarm's email notification message. If blank, the message settings defined on the notification block that sent the email out will be used instead. Supports HTML tags such as <html>, <body>, <h1>, etc.	String

SMS Notification Properties

Custom Message	CustomSmsMessage	If specified, will be used for the SMS message. If blank, the message defined in the notification block will be used.	String
----------------	------------------	---	--------

Associated Data

User Defined Data		Associated Data are custom alarm properties that can be added to any alarm. These properties will often be bound to other Tags that represent associated contextual data that may be related to the alarm. A snapshot of the values of these properties will be taken when the alarm becomes active. These values will be attached to the alarm event as it moves through the rest of the alarming system, meaning that the values will be available from the alarm status system, the alarm journal system, and in the alarm notification system.	String
-------------------	--	--	--------

Runtime Alarm Metric Properties

The Alarms property under a tag in the Tag Browser contains properties that report the current state of alarms on the tag.

ValueSource	memory
Alarms	
ActiveAckCount	0
ActiveUnackCount	1
ClearUnackCount	0
HasActive	<input checked="" type="checkbox"/>
HasUnacknowledged	<input checked="" type="checkbox"/>
HighestAckedName	null
HighestAckedPriority	null
HighestActiveName	High Setpoint Alarm
HighestActivePriority	1
HighestUnackedName	High Setpoint Alarm
HighestUnackedPriority	1
LastActiveTime	2021-07-02 2:19:52 PM
ShelvedCount	0

Property Name	Description
ActiveAckCount	The number of alarms on the Tag that are both active and acknowledged.
ActiveUnackCount	The number of alarms on the Tag that are both active and unacknowledged.
ClearUnackCount	The number of alarms on the Tag that are both clear and unacknowledged.
HasActive	True, if the Tag has at least one active alarm. False, if there are zero alarms.
HasUnacknowledged	True, if the Tag has at least one unacknowledged alarm. False, if there are zero unacknowledged alarms.
HighestAckedName	The Name of the highest acknowledged alarm, ranked by Priority.
HighestAckedPriority	The highest Priority of all acknowledged alarms on the Tag.
HighestActiveName	The Name of the highest active alarm, ranked by Priority.
HighestActivePriority	The highest Priority of all active alarms on the Tag.
HighestUnackedName	The Name of the highest unacknowledged alarm, ranked by Priority.
HighestUnackedPriority	The highest Priority of all unacknowledged alarms on the Tag.
LastActiveTime	A timestamp representing the last time an alarm went active on the Tag.
ShelvedCount	The number of currently shelved alarms on the Tag.

Runtime Alarm Metrics for Individual Alarms

In addition to the metrics above, each alarm configured on a tag has further properties that are available. These properties are located after **ShelvedCount**. For each alarm on a tag, additional expandable items will be visible in the property editor. The name on these items will match the name of their associated alarm.

In the image below, the parent tag has two alarms configured, one named **High Setpoint Alarm**, the other named **Low Setpoint Alarm**. Because the tag has two alarms, we see two expandable entries towards the bottom of the image, with names that match the names of the alarms.

Alarms	
■ ActiveAckCount	0
■ ActiveUnackCount	1
■ ClearUnackCount	0
■ HasActive	<input checked="" type="checkbox"/>
■ HasUnacknowledged	<input checked="" type="checkbox"/>
■ HighestAckedName	null
■ HighestAckedPriority	null
■ HighestActiveName	High Setpoint Alarm
■ HighestActivePriority	1
■ HighestUnackedName	High Setpoint Alarm
■ HighestUnackedPriority	1
■ LastActiveTime	2021-07-02 2:19:52 PM
■ ShelvedCount	0
▶ ■ High Setpoint Alarm	
▶ ■ Low Setpoint Alarm	

Expanding these items will reveal additional properties, as listed in the table below.

Property Name	Description
AckTime	If the most recent alarm event for this alarm has been acknowledged, then this property will show a timestamp representing when acknowledgement occurred. If the most recent alarm event has not yet been acknowledged, the value of this property will be null.
AckUser	If the most recent alarm event for this alarm has been acknowledged, then this property will show the username that acknowledged the alarm event. Otherwise the value of this property will be null. If the alarm event was acknowledged by the system, such as when the Alarm Mode is set to Auto, then the property will show a null value. If the user that acknowledged the alarm event was authenticated via a User Source, then the value will follow the pattern below, where X is the name of the User Source, and Y is the name of the user. usr-prov:X:/usr:Y
	If the user that acknowledged the alarm event was authenticate via an Identity Provider, then the value will follow the pattern below, where Y is the name of the user. usr:Y
AckUserName	If the most recent alarm event for this alarm has been acknowledged, then this property will show the username that acknowledged the alarm event. Otherwise the value of this property will be an empty string. If the alarm event was acknowledged by the system, such as when the Alarm Mode is set to Auto, then the property will show an empty string. This property is effectively a simplified version of AckUser, containing only the username.
ActiveTime	Shows a timestamp representing when the most recent alarm event became active.
ClearTime	Shows a timestamp representing when the most recent alarm event transitioned from an active state to a clear state.
DisplayPath	Represents the value of the Display Path property on the alarm.
DisplayPathOrSource	Shows the display path if one has been configured. Otherwise shows the source path.

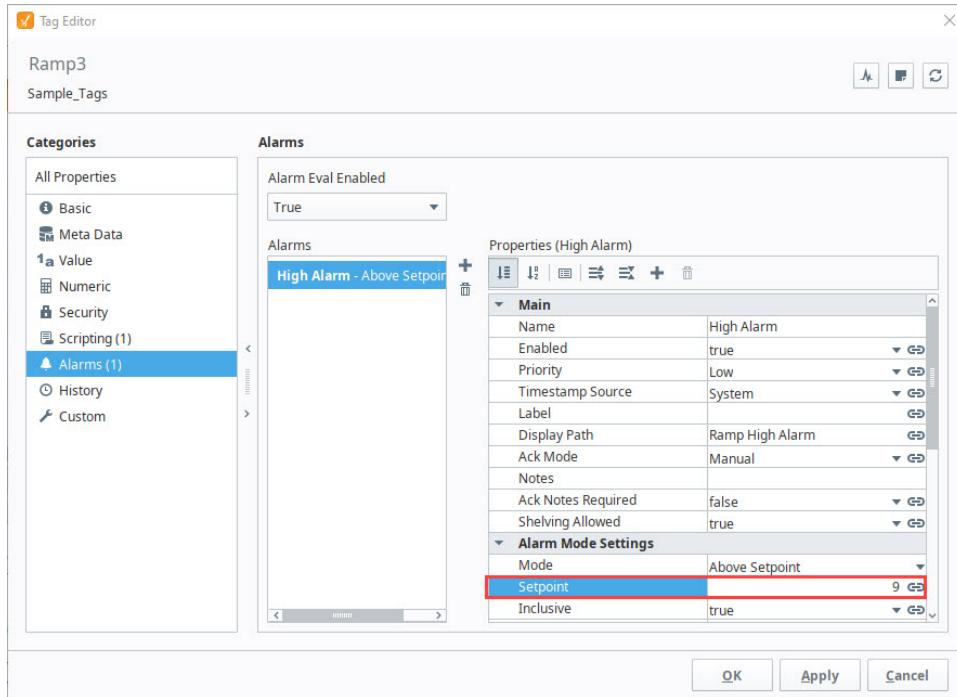
Enabled	Represent whether or not this alarm is enabled.												
EventState	Represents the most recent state transition. Values and their numerical state are listed below: <table border="1"> <thead> <tr> <th>Value</th><th>State</th></tr> </thead> <tbody> <tr> <td>0</td><td>Active</td></tr> <tr> <td>1</td><td>Clear</td></tr> <tr> <td>2</td><td>Acknowledged</td></tr> </tbody> </table>	Value	State	0	Active	1	Clear	2	Acknowledged				
Value	State												
0	Active												
1	Clear												
2	Acknowledged												
EventTime	A timestamp representing when the event last changed state.												
EventValue	Represents the last value that made the alarm event transition to either active or cleared.												
IsAcked	A boolean representing whether or not the most recent alarm event has been acknowledged.												
IsActive	A boolean representing whether or not the most recent alarm event is active.												
IsClear	A boolean representing whether or not the most recent alarm event is cleared.												
IsShelved	A boolean representing whether or not the alarm has been shelved.												
Label	A string representing the Label property on the alarm. Shows a null value if a Label wasn't defined.												
Name	The name of the alarm.												
Priority	The priority of the alarm. <table border="1"> <thead> <tr> <th>Value</th><th>Priority</th></tr> </thead> <tbody> <tr> <td>0</td><td>Diagnostic</td></tr> <tr> <td>1</td><td>Low</td></tr> <tr> <td>2</td><td>Medium</td></tr> <tr> <td>3</td><td>High</td></tr> <tr> <td>4</td><td>Critical</td></tr> </tbody> </table>	Value	Priority	0	Diagnostic	1	Low	2	Medium	3	High	4	Critical
Value	Priority												
0	Diagnostic												
1	Low												
2	Medium												
3	High												
4	Critical												
SetpointA	When the Mode property on the alarm has a single setpoint, this property shows the setpoint value. When the Mode property has multiple setpoints, this represents the Low Setpoint property.												
SetpointB	When the Mode property on the alarm has a single setpoint, this property is hidden. When the Mode property has multiple setpoints, this represents the High Setpoint property.												
Source	The source path of the alarm.												
State	The state of the alarm. <table border="1"> <thead> <tr> <th>Value</th><th>State</th></tr> </thead> <tbody> <tr> <td>0</td><td>Clear and Unacked</td></tr> <tr> <td>1</td><td>Clear and Acked</td></tr> <tr> <td>2</td><td>Active and Unacked</td></tr> <tr> <td>3</td><td>Active and Acked</td></tr> </tbody> </table>	Value	State	0	Clear and Unacked	1	Clear and Acked	2	Active and Unacked	3	Active and Acked		
Value	State												
0	Clear and Unacked												
1	Clear and Acked												
2	Active and Unacked												
3	Active and Acked												

Binding

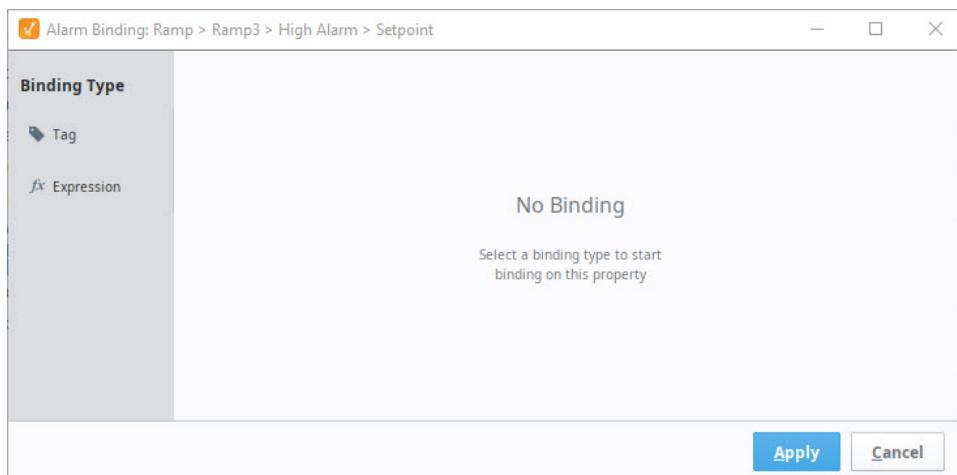
Many alarm properties are bindable, which means they can be bound to other tags in the system, or expressions. For example, you might bind the enabled property to another tag which represents whether or not your process is running, thereby, disabling the alarm when production is stopped. Another example is you might bind the setpoint of an alarm to a tag that operators can manipulate, thereby, letting the setpoint be changed at runtime. For more information, see [Configuring Alarms](#).

To bind an alarm property of a tag:

1. Click on the binding icon. An Alarm Binding popup window will appear.

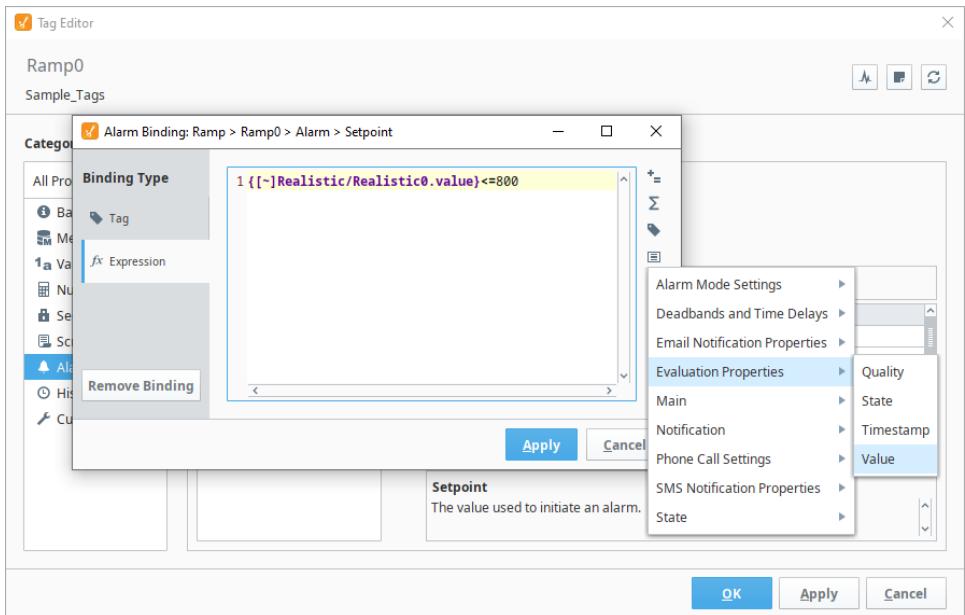


- From this window, you can select a Binding Type to configure a Tag or an Expression binding. For this example, let's select Expression. Binding to an Expression can reference many useful values such as the tag's value and other settings of the alarm.



- Now, enter your expression, using the Operators , Functions , Tags , and Alarm Properties selections as needed.

Note: The alarm properties that can be selected using the Alarm Properties icon are the properties included in the **Tag Editor > Alarms** window or under the tag in the Tag Browser. See the [Reference Table](#) and the [Runtime Alarm Metric Properties](#) for descriptions of these properties. The State properties will only be available to view after alarming is enabled on the tag.



4. When you configured the binding to your liking, click **Apply**.

Tag Scaling Properties

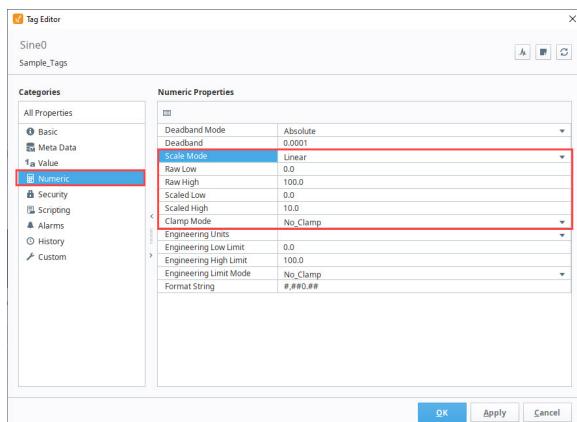
Configuring a Tag's scaling will condition the data for use within the Ignition Designer. Scaling will take the raw PLC value driving a Tag, do some math, and use the resulting value as the value of that Tag. Scaling works both ways. When you write to that Tag, Ignition will scale it in the opposite direction before writing to the PLC.

For example, if the capacity of a tank is 5250 gallons, but the tank's fill level is better represented in the Designer as percentage of 0 through 100, configuring the Tag's scaling property, will result in the Tag displaying 0 through 100, while the actual tank fill moving is between 0 and 5250 gallons. For this example, you can double-click on the Tag to open the **Tag Editor**, and expand **Numeric Properties** to configure the scaling of the Tag. When scaling between a Raw Low and High, and Scaled Low and High, select the **Linear** Scale Mode. So what Ignition is actually doing, is setting up the calculations behind the scenes to scale the value appropriately.

This feature was changed in Ignition version 8.1.17:

In 8.1.17, the Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts.

Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).



On this page ...

- [Linear Scaling](#)
- [Square Root Scaling](#)
- [Exponential Filter Scaling](#)
- [Bit Inversion Scaling](#)
- [Scaling Examples](#)
 - [Simple Divide by Ten](#)
 - [4-20 Milliamp Signal to Percent](#)
 - [4-20 Milliamp Signal to Gallons](#)



Tag Scaling

[Watch the Video](#)



If you are using scaling...

The numbers and units don't have to match-up. Scaling is straight from one number to another number. It doesn't matter what the units are, and it doesn't matter what the conversion is. What is important, is that the data type of the Tag must match the data type of the scale value (i.e., dividing an integer in the PLC by 10 probably means your Ignition Tag should be a float).

Scaling is not available on Memory Tags: Memory Tags are not driven by an external source such as a PLC or SQL query, so scaling will never be applied. In these scenarios, it is recommended to scale the mechanism that is writing to the Memory Tag instead. Numerical properties of Tags can be scaled allowing automatic bi-directional conversion outside of the PLC. Scaling types include **Linear** scaling, **Square Root** scaling, and **Exponential Filter** scaling. The numerical properties are available to **OPC**, **Expression**, **Database**, and **Client** Tags whose data types are numeric. For a complete list of all of the Tag Scaling properties, see [Tag Properties](#).

Linear Scaling

The value will be scaled linearly between the low and high values, and clamped as appropriate.
The linear equation is:

```
ScaledValue = S * (Value - RL) / R + SL
```

Where:

S = (ScaledHigh-ScaledLow)
R = (RawHigh - RawLow)
RL = RawLow
SL = ScaledLow

Square Root Scaling

The equation for square root scaling is:

```
ScaledValue = S * ((Value-RL) / R) + SL
```

Where:

S = (ScaledHigh-ScaledLow)
R = (RawHigh - RawLow)
RL = RawLow
SL = ScaledLow

Exponential Filter Scaling

This mode implements a simple linear recursive filter to smooth values. The scale factor corresponds to the weight of the smoothing effect, and is a value between 0.0 and 1.0. The smaller the factor, the greater the degree of smoothing.

The equation for the filter is:

```
y(t) = (1-f)*y(t-1)+f*x(t)
```

Where:

y(t) = the output at time t
y(t-1) = the previous output
x(t) = the input value (current value)
f = the scale factor, with $0.0 \leq f \leq 1.0$

Note: Only **good** quality values are considered for the filter. **Bad** quality values are ignored.

Bit Inversion Scaling

This simple scaling mode will generate the complement of a binary value. If the current value is coming in as 0001_0101 (21), this will return a 1110_1010 (-22) instead. A popular use for this scale mode is that it can be used to invert modbus values if your device stores them in reverse bit order. Note that Bit Inversion Scaling uses a little-endian format.

Scaling Examples

Simple Divide by Ten

This is common when storing a single decimal point of precision as an Integer in the PLC. This is to save space by not using a Float type.

Raw Low: 0.0
Raw High: 100.0
Scaled Low: 0.0
Scaled High: 10.0

4-20 Milliamp Signal to Percent

This is common when using a simple pressure sensor. The sensor is calibrated to send 4 millamps (minimum value) when the tank is empty, and 20 millamps (maximum value) when the tank is full.

Raw Low: 4.0
Raw High: 20.0
Scaled Low: 0.0
Scaled High: 100.0

4-20 Milliamp Signal to Gallons

This is common when using a simple pressure sensor. The sensor is calibrated to send 4 millamps (minimum value) when the tank is empty, and 20 millamps (maximum value) when the tank is full (5000 gallons).

Note: There is no direct conversion between amps and gallons. In scaling it doesn't matter.

Raw Low: 4.0

Raw High: 20.0

Scaled Low: 0.0

Scaled High: 5000.0

Related Topics ...

- [Exporting and Importing Tags](#)

Tag Security Properties

This feature was changed in Ignition version 8.1.17:

In 8.1.17, the Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts. Additional improvements include the following:

- Only expression bindings now require opening the dialog to change. Tag and Parameter bindings can be edited in-line within the tag editor and are not shortened.
- Sorting in UDT Editor now sorts folders first, then tags.
- Users can add tags in the UDT Editor without having to click the root of the UDT/folder.
- UDT Parameters now show when they are overridden.
- The Tag Editor uses the entire width allowed to display content.
- UDT properties are added in-line to bindings now rather than clearing out the content.
- Property values are now all left-aligned.
- The Expression editor now accepts drag and drops from the UDT Editor or Tag Browser.

Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

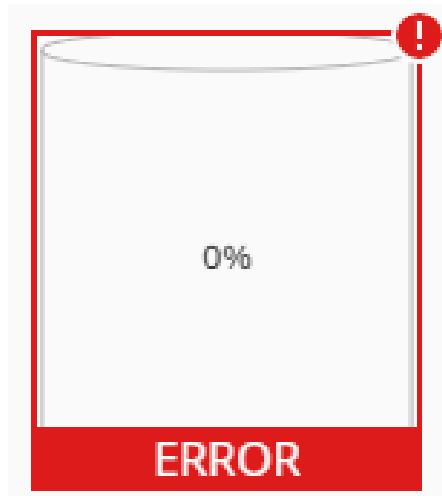
Tag security is often the best way to configure security for data access. By defining security on a Tag, you affect the Tag across wherever it is used, as opposed to configuring component security on each component that displays or controls that Tag.

There are three properties on Tags that can restrict access.

- Read Permissions: Defines the security levels required in order to read values from a Tag
- Read Only: Defines whether a Tag is read-only or writable
- Write Permissions: Defines the security levels required in order to write values to a Tag

Users with specific roles and zones can be given read/write access to a Tag, while other users with other roles are excluded from modifying the Tag.

If a user opens a Perspective view or a Vision client window that has components that are bound to a Tag they do not have permissions for, the user will see an overlay on top of the component. For more information, see [Quality Codes and Overlays](#). The following example shows a tank displayed in a session, but the user does not have read permission for the Tag it is bound to.

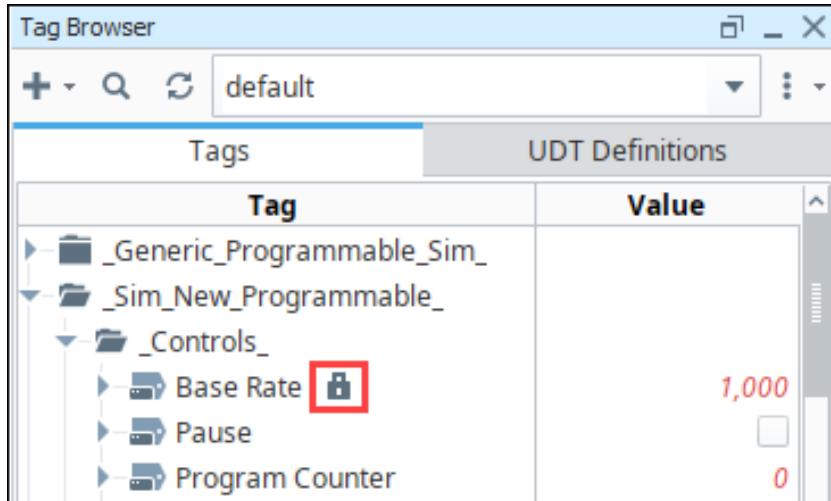


Read Only Security

When a Tag is set to read only, a Lock  icon is displayed next to the Tag in the Tag Browser.

On this page ...

- [Read Only Security](#)
- [Read and Write Permissions](#)
 - [Read Permissions](#)
 - [Write Permissions](#)
- [Using Security Zones](#)



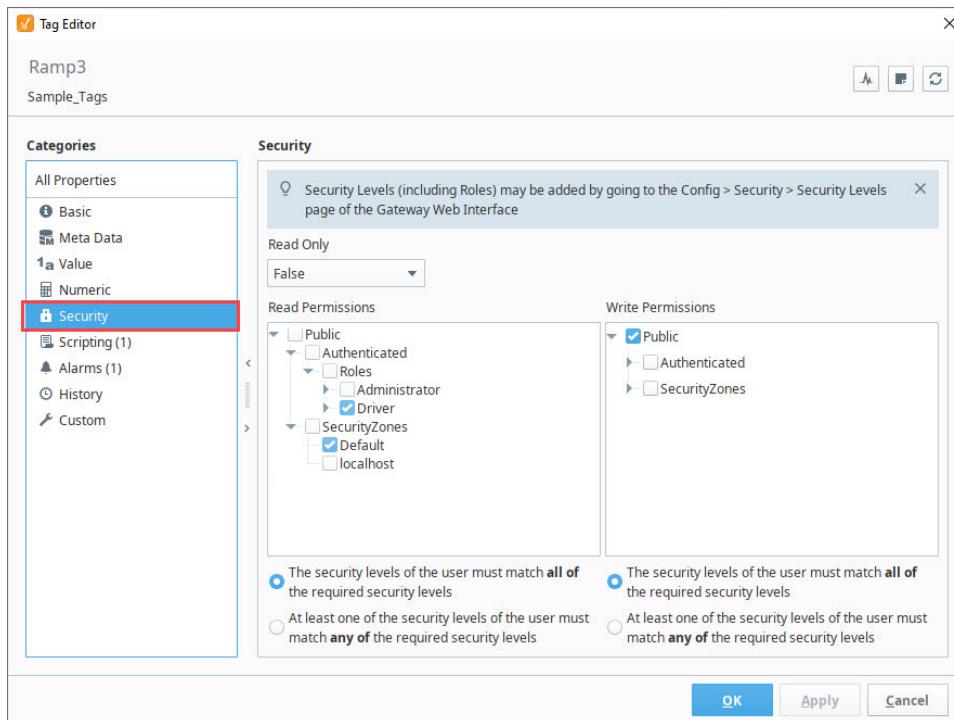
Read and Write Permissions

Instead of making a tag Read Only for all users, you can conditionally provide read and write access based on [Security Levels](#). Doing so involves a adjusting the security settings on the tag in question. The checkbox tree you are presented with will show you all of the security levels configured in the Gateway Config > Security > [Security Levels](#) page.

Read Permissions

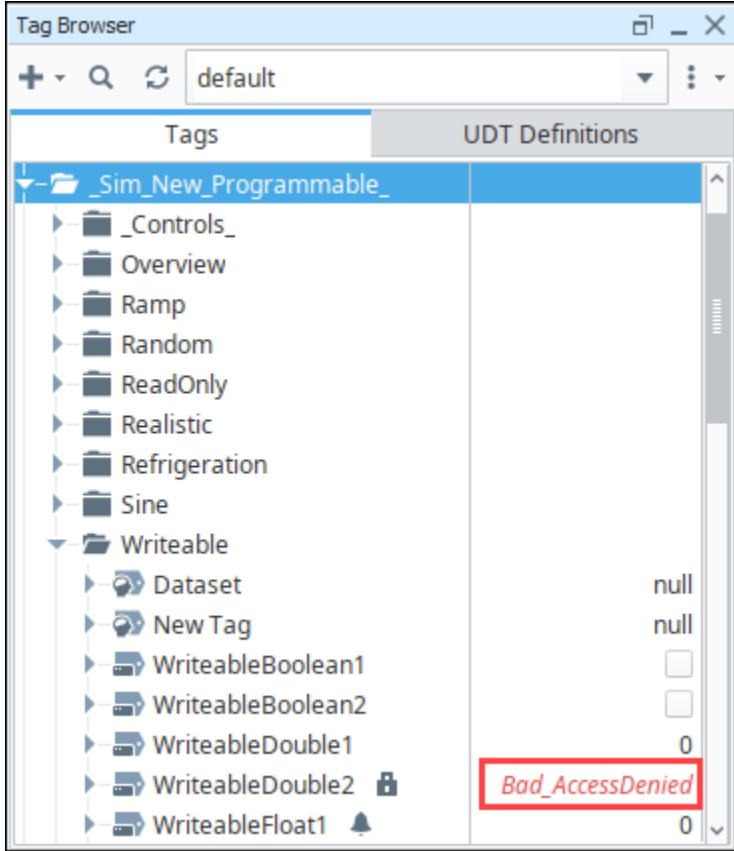
Read permissions define the security levels required in order to read values from a Tag. By default, Tags have Read Permissions set to "Public". You can change the Read security using the Tag Browser in the Designer.

1. In the Tag Browser, right-click on the **Tag**, and select the **Edit** icon.
2. Select the Security section.
3. On the screen, choose the security levels you want to have Read permissions for this Tag. In this example, only users with role of "Driver" will be able to see the Tag value.



4. Click **Commit** to accept the settings.
5. Click **OK** to save the changes to the Tag.

6. If you are logged in as a user other than Driver, you will now see the "Bad_AccessDenied" in the Tag Browser instead of the Tag's value.

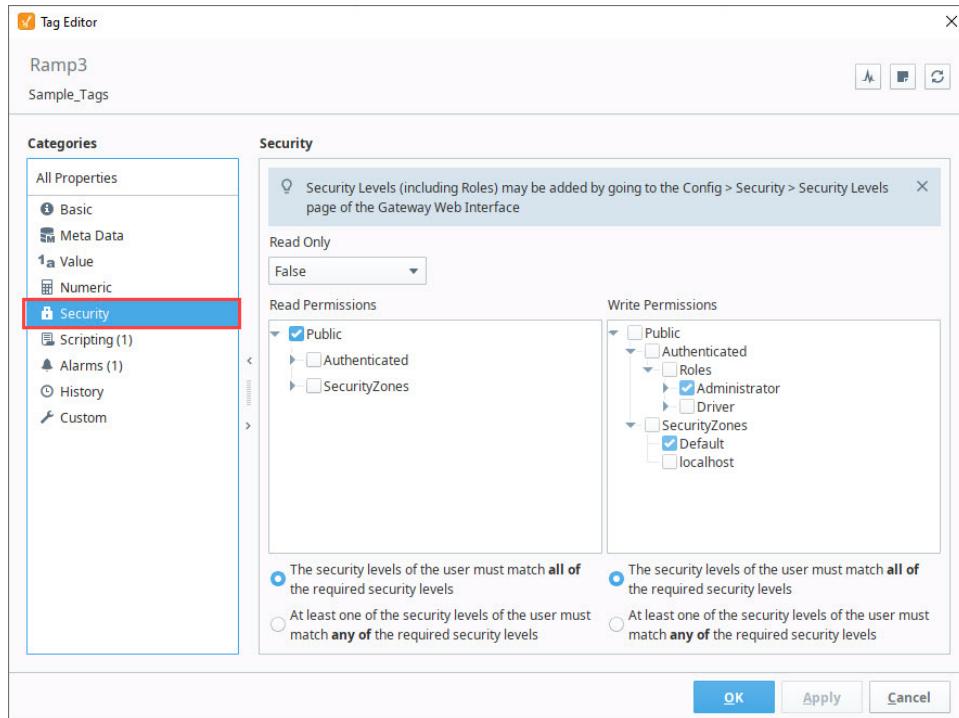


Write Permissions

Write permissions define the security levels required in order to write values to a Tag. By default, Tags have Write Permissions set to "Public". You can change the Write security using the Tag Browser in the Designer.

1. In the Tag Browser, right-click on the **Tag**, and select the **Edit** icon.
2. Select the Security section.

3. On the screen, choose the security levels you want to have write permissions for this Tag. In this example, only users with role of Administrator will be able to write to the Tag value.



Note:

The following feature is new in Ignition version **8.1.25**
[Click here](#) to check out the other new features

If any security levels under the Read Permissions or Write Permissions sections are deleted from the Gateway, they will appear grayed-out with a red warning underline. A warning indicator icon will also appear in the upper right-hand corner with the number of selected security levels that no longer exist. If you follow the security levels tree up to the parent levels, you'll notice affected levels now include a dotted underline.

The deleted security level and all warning indications are removed when the deleted security level is unchecked and new settings are saved.

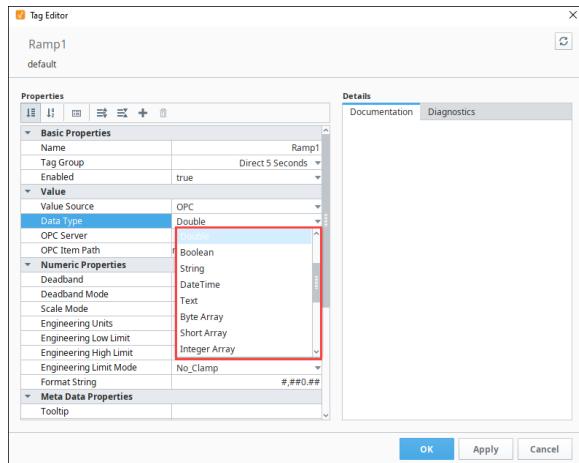
Using Security Zones

In addition to setting up security on individual Tags, you can set up security policies specific to Security Zones. This is useful in cases where you want to make all tags in a provider Read Only from network locations. See the [Security Zones](#) page for more details about Tag Access options.

Tag Data Types

This page details the different data types that can be applied to standard tags.

The data type of a tag is determined by the Data Type property, which is accessible from the Tag Editor. The tags system will attempt to coerce any raw incoming value (for example, from OPC or a SQL query) into the desired type.



On this page ...

- [Array and Dataset Data Types](#)
 - [Array Tags](#)
 - [Dataset Tags](#)
 - [Dataset Tag Example](#)
- [Document Type](#)



INDUCTIVE
UNIVERSIT

Array and Dataset Tags

[Watch the Video](#)

The following table lists all the data types available for tags in Ignition.

Data Type	String Value	Integer Value
Byte	Int1	0
Short	Int2	1
Integer	Int4	2
Long	Int8	3
Float	Float4	4
Double	Float8	5
Boolean	Boolean	6
String	String	7
DateTime	Datetime	8

Note:

The Siemens driver does not support Ignition tag's DateTime type. In these cases it is recommended to extract each byte in the Siemens DATE_AND_TIME value item, storing each into a separate Ignition tag. Then use an expression tag to combine each byte into a human readable datetime.

Example

```
// The following assumes each byte of the DATE_AND_TIME value is under a tag named  
// like "BASIC_DATA_AND_TIME", and that each tag is a String type.
```

```

"20" + //First part of the year, assuming in 21st century
toHex({[.]BASIC_DATE_AND_TIME[0]}) + "-" +
//Second part of the year, reading from an Ignition tag.
numberFormat(toHex({[.]BASIC_DATE_AND_TIME[1]}), "00") + "-" + //Month
numberFormat(toHex({[.]BASIC_DATE_AND_TIME[2]}), "00") + "-" + //Day
numberFormat(toHex({[.]BASIC_DATE_AND_TIME[3]}), "00") + ":" + //Hours
numberFormat(toHex({[.]BASIC_DATE_AND_TIME[4]}), "00") + ":" + //Minutes
numberFormat(toHex({[.]BASIC_DATE_AND_TIME[5]}), "00") + ":" + //Seconds
numberFormat(fromBinary(left(numberFormat(toBinary({[.]BASIC_DATE_AND_TIME[6]} + {[.]BASIC_DATE_AND_TIME[7]}), "0000000000000000"), 12)), "000") + " " + //Milliseconds, it
includes the first 4 bits of byte 7
case( //Day of the week, last 4 bits of byte 7
right(numberFormat(toBinary({[.]BASIC_DATE_AND_TIME[7]}), "00000000"), 4),
"0000", "Sunday",
"0001", "Monday",
"0010", "Tuesday",
"0011", "Wednesday",
"0100", "Thursday",
"0101", "Friday",
"0111", "Saturday",
"Error - No Day Found")

```

This feature was changed in Ignition version 8.1.29:

The Siemens driver supports DateTime data types for the following Siemens devices:

- S7-300
- S7-400
- S7-1500

Text	Text	10
This feature was changed in Ignition version 8.1.8: The Text data type on tags was deprecated		
Byte Array	Int1Array	17
Short Array	Int2Array	18
Integer Array	Int4Array	11
Long Array	Int8Array	12
Float Array	Float4Array	19
Double Array	Float8Array	13
Boolean Array	BooleanArray	14
String Array	StringArray	15
DateTime Array	DateTim eArray	16

Binary Data	ByteArray	20
Dataset	DataSet	9
Document (JSON Document)	Document	29

Array and Dataset Data Types

The Array and Dataset data types available on tags allow for multiple data points to be stored in a single tag. Configuring a tag as an array or dataset is as easy as changing the data type in the Tag Editor.

Note: Most OPC device drivers do not support array and dataset data types. These types work best when used on query tags or memory tags. When working with arrays via our OPC UA device drivers, it is recommended to create Ignition tags for each array element instead of using an array data type tag.

Array Tags

For OPC servers and drivers that support array type tags, each element in the array can easily be represented with the array data types in Ignition. Because the core data type of each element in the array is the same, it is possible to apply Tag History, Alarming, or Scaling configurations onto the array, and these configurations will be inherited by each element.

The following feature is new in Ignition version **8.1.33**

[Click here](#) to check out the other new features

The array tag type will now accept and coerce JSON array types, allowing methods such as `jsonGet`.

Array Tag Write-Back

OPC Array tags support writing back to the device. How this is done can vary, depending on the type of OPC Server in use. Some OPC Servers support writes to individual array elements, where a write would occur just like any other tag write. However, some OPC Servers do not support individual element writes, which means the whole array will need to be written back to the array tag, even if only a single element is changing.

Dataset Tags

Dataset tags allow multiple rows and columns worth of data to be stored in a tag. Each column is exposed as a separate folder in the tag (i.e., the "name" folder in the image below). Dataset tags can be driven by a query, so it's possible to query for multiple columns on a row in a single tag. This is more efficient than using multiple query tags (and thus multiple queries) to retrieve the same data.

While dataset tags are convenient, note that the Tag History system and Alarm system do not support tags with dataset types.

The screenshot shows the Ignition Tag Browser window titled "Tag Browser". The search bar at the top contains the text "default". Below the search bar, there are two tabs: "Tags" and "UDT Definitions", with "Tags" being the active tab. The left pane displays a tree structure of tags under a folder named "Demographics". One tag, "City_Dataset", is expanded, showing its properties. The right pane shows the configuration details for "City_Dataset".

Property	Value
Type	Dataset [5R x 4C]
Value	.0001
Format	#,##0.##
Unit	City_Dataset
Quality	Good
Default Value	Default
Source	Dataset [5R x 4C] memory

Valid data types that can be stored in dataset tags include the following:

- Float
- Short
- Long
- Date
- Integer
- Boolean
- String
- Color
- Double
- Timestamp

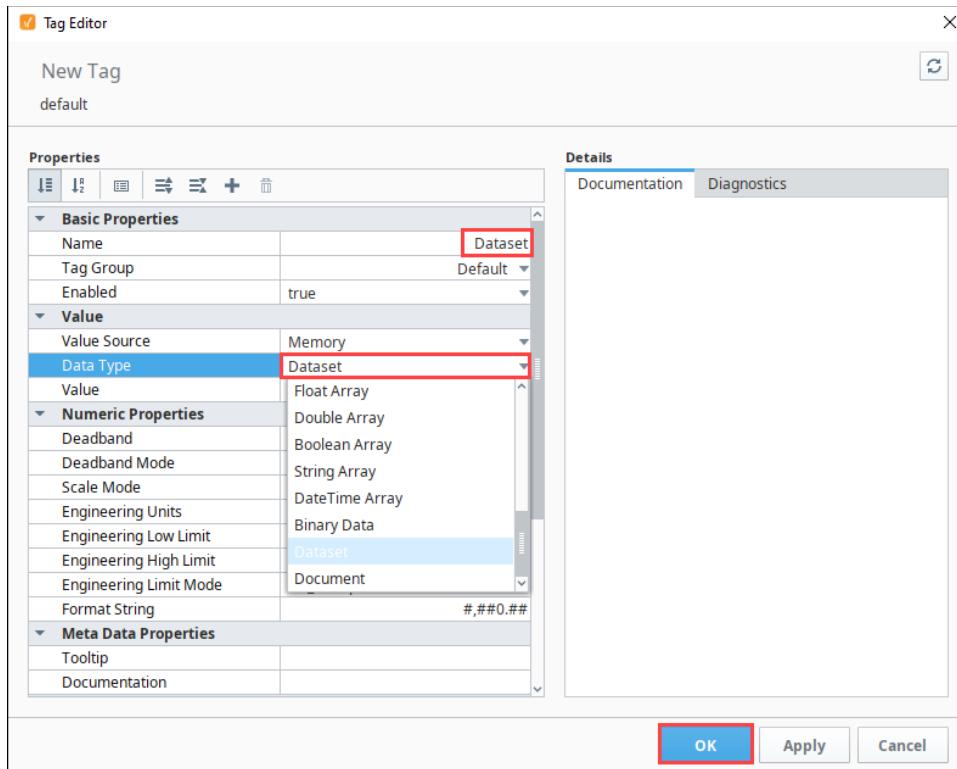
The following feature is new in Ignition version **8.1.20**
[Click here](#) to check out the other new features

- Byte Arrays are now supported as of 8.1.20

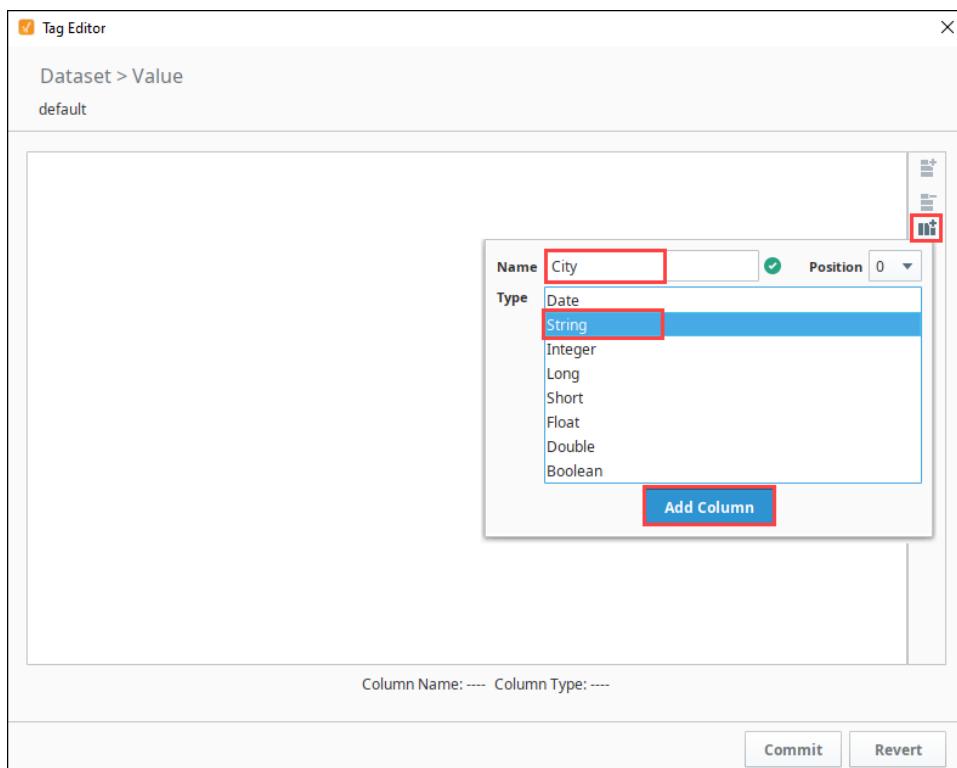
Dataset Tag Example

The following example will create a dataset memory tag and display the contents in a Table component.

1. Create a new **Memory Tag**. Name it **Dataset**, and change the data type to **Dataset**. The Dataset will be empty by default.



2. Click the **Edit** icon next to **Value**. The Value screen is displayed. For this example, we created a simple dataset with four columns and five rows.
3. Click the **Add Column** icon. Name the first column **City** and set type to be **String**.
4. Click **Add Column**.



5. Repeat adding columns as follows:

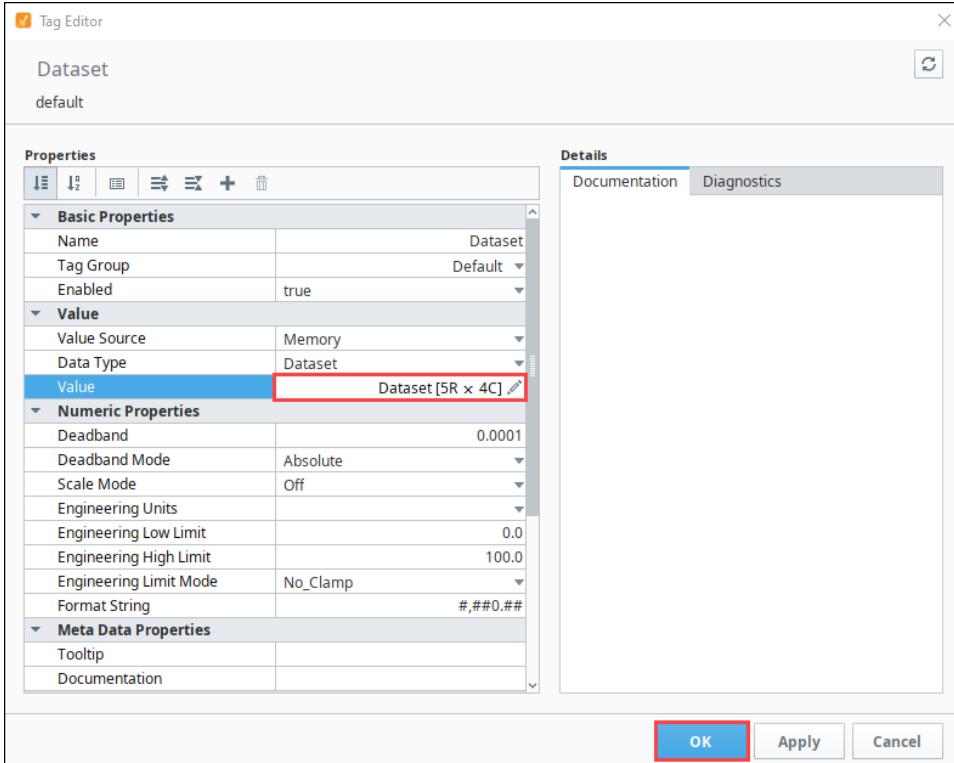
Column Name: **Population** Type: **Integer**
 Column Name: **TimeZone** Type: **String**

Column Name: **GMTOffset** Type: **Integer**

6. Click the **Add Row**  icon. Add the row information as follows:

New York	8368710	EST	-5
Los Angeles	3833995	PST	-8
Chicago	2853114	CST	-6
Houston	2242193	CST	-6
Phoenix	1567924	MST	-7

7. Click the **Commit** button.
8. Click **OK** to save the tag.



The tag will now contain rows, columns, and values based on the configurations you made earlier in this example. Now you have a tag with a dataset value that can be bound to by components in Vision and Perspective.

Document Type

The document type allows a tag to have a JSON document as a value. Note that the Tag History system does not support tags with a document type.

Note: Most OPC device drivers do not support the document data type. This type works best when used on query tags or memory tags.

The following feature is new in Ignition version **8.1.27**
[Click here](#) to check out the other new features

The document tag type now supports writing to OPC tags containing DocumentArray values.

Setting a tag to this type will cause the value field to display an icon.

New Tag

default

Properties

Basic Properties	
Name	New Tag
Tag Group	Default
Enabled	true

Value

Value Source	Memory
Data Type	Document
Value	{document}

Numeric Properties

A screenshot of the Tag Editor interface. It shows a 'Basic Properties' table with columns for Name, Tag Group, and Enabled. Below it is a 'Value' section with Value Source set to Memory and Data Type set to Document. The Value field contains '{document}' with a pencil icon. A red box highlights the pencil icon. At the bottom is a 'Numeric Properties' section.

Clicking the icon will transition the Tag Editor into a JSON Editor, allowing you to manually write JSON directly to the value.

Tag Editor X

New Tag > Value

default

Filter

object {0} Add Object Member...

Commit Revert

A screenshot of the Tag Editor's JSON Editor view. It shows a tree structure with 'object {0}' expanded, and a '+ Add Object Member...' button. At the bottom are 'Commit' and 'Revert' buttons.

Click the + icon to add new members to the object. When finished, press **Commit**.

When saved, the value of the tag will render as a JSON document in the Tag Browser.

```
 New Tag { "key": "value", "array": [ "value", "value" ] }
```

The following feature is new in Ignition version **8.1.32**
[Click here](#) to check out the other new features

The document tag type will now accept and coerce PyDictionary types.

Tag Paths

Tags and their properties can be referenced by a string-based path in many areas of Ignition, such as [expressions](#) and [scripts](#). Each Tag has a unique absolute path and often has many equivalent relative paths when referenced from other Tags. In most cases these paths are generated automatically via helper buttons. However, it's a good idea to understand how Tag paths work, particularly if you need to configure an Indirect Tag Binding, or access a Tag from an expression, or script.

A Tag path looks something like this: [Tag Provider]folder/path/tag.property

The folder/path/tag.property portion of the path may contain the following:

- A Tag
- Any number of nested folders followed by a Tag, separated by forward slashes (/)
- A period (.) followed by a property name after the Tag. Omitting this is equivalent to using the . value property

The [Tag Provider] portion surrounded by square braces can have the following options:

Source Option	Meaning
[Tag Provider Name]	The name of the Tag provider that hosts the Tag.
[] or not specified	The default Tag provider for the current project. If used in the Gateway scope, this notation can (generally) result in an invalid path, as the Gateway doesn't have a default Tag provider.
[.]	Relative to the folder of the Tag that is being bound. This is especially useful in UDT definitions.
[~]	Relative to the Tag provider of the Tag that is being bound (root node).
[Client]	Refers to the Vision Client Tag provider, which contains only Vision Client Tags.
[System]	Refers to a System Tag.

On this page ...

- [Using Relative Tag Paths](#)
- [Tag Path Manipulation](#)
- [Array Type Tag Paths](#)
- [Document Type Tag Paths](#)
 - [Writing to Document Type tags](#)
- [Using the {this} Keyword](#)
 - [Using {this} in Alarms](#)

Using Relative Tag Paths

Tag paths that begin with [.] or [~] are known as *relative paths*. They are used inside Tags that bind to other Tags, and they are relative to the host Tag's path. Using the relative path syntax helps to avoid problems caused by moving Tags and renaming providers.

[.] refers to the Tag's current folder. By using [.], Tags can be moved from folder to folder without problem (provided that all of the applicable Tags are moved together). Additionally, you can use ".." (two periods) to go back one folder from the current relative position, for example [..]../../tag allows you to reference a Tag that is two folders up.

[~] refers to the Tag's provider root. It replaces an explicit provider name and thus protects the Tag path from "breaking" if the provider is renamed or if the Tag is imported/exported/moved between different providers.

Tag Path Manipulation

Ignition provides a great deal of flexibility for Tag addressing since Tag paths and Tag properties are string-based. The underlying strings that compose a valid Tag path can be assembled from many different parts in which the eventual construction results in a valid Tag path.

The following scripting demonstrates this concept. Suppose there was a Tag path to a level indicator in a tank. In this case it is the default Tag provider, Tanks folder, Tank 1 Folder, and the Level Tag.

```
tagPath = "[default]Tanks/Tank 1/Level"
```

But suppose that there was more than just Tank 1 and instead there was Tank 2, Tank 3, Tank 4, etc. Dynamically changing the Tag paths is simple because Ignition's Tag paths are string representations. The following takes the tank number and inserts it into a new Tag path. The tankNumber variable changes the eventual creation of the tagPath. Using this method in scripting or in an expression binding will look slightly different.

Python Dynamic Tag Path

```
tankNumber = 2  
tagPath = "[default]Tanks/Tank %i/Level" % tankNumber
```

Expression Dynamic Tag Path

```
tag("[default]Tanks/Tank "+{Root Container.tankNumber}+"/Level")
```

The result of the tagPath variable will be **[default]Tanks/Tank 2/Level** which is a valid Tag path to the the level sensor for Tank 2.

Array Type Tag Paths

When a path leads to an array type tag, individual elements can be accessed using square brackets, and the index offset.

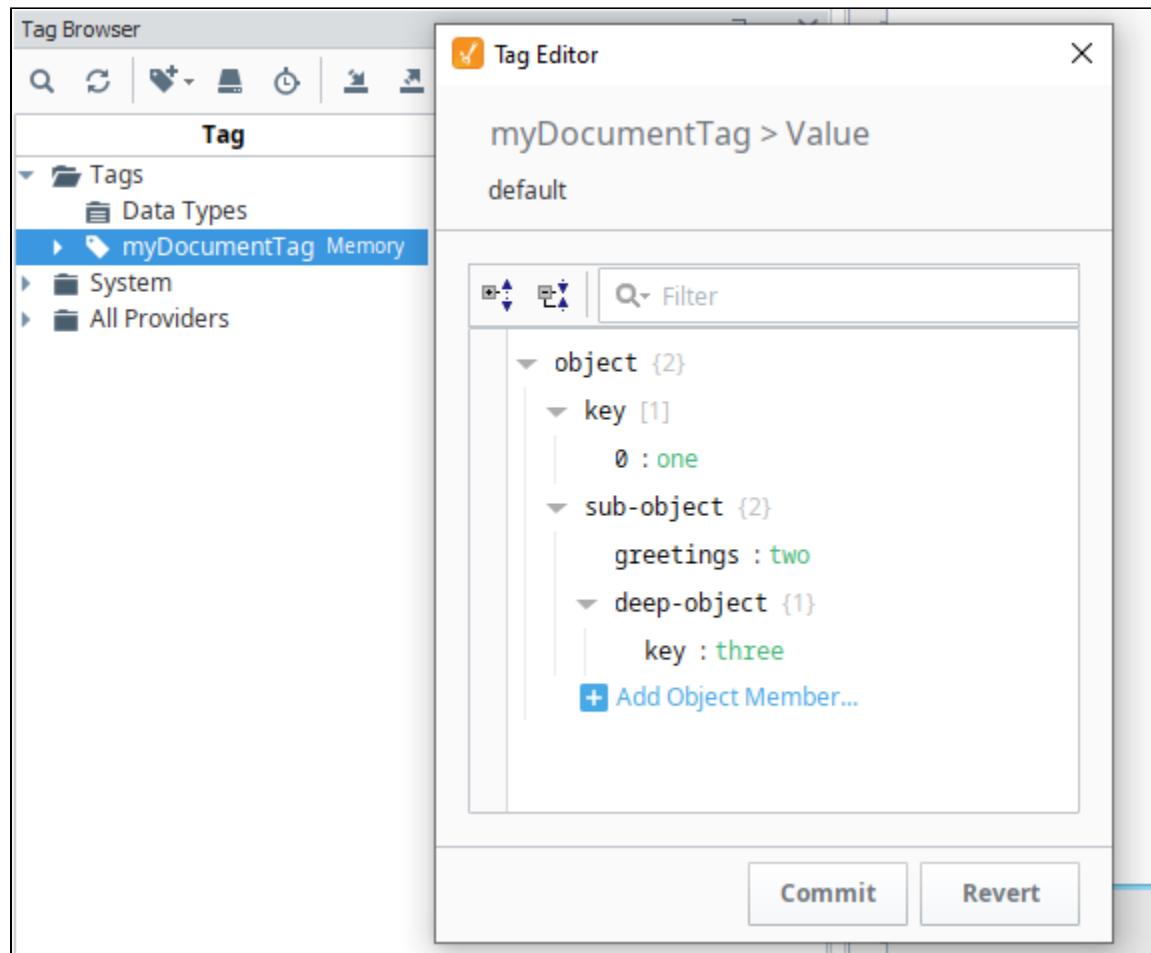
```
[default]Folder/myArrayTag[0]
```

Document Type Tag Paths

The following feature is new in Ignition version **8.1.0**
[Click here](#) to check out the other new features

When a path leads to a document type tag, individual objects within the tag's value can be accessed by using square brackets wrapped around a set of quotations marks that lead to the object.

For example, say we have a documentation tag with the following value:



The following would return "**one**", which is in the first element of the "**key**" array:

```
[default]myDocumentTag[ 'key[ 0 ]' ]
```

JSON Values are references with a '.' character, so the following would retrieve the value of "**two**" inside of the "**sub-object**":

```
[default]myDocumentTag[ 'sub-object.greetings' ]
```

For more complex structures, you can continue adding to the JSON string. The following would return "**three**" from the key in "**deep-object**":

```
[default]myDocumentTag[ 'sub-object.deep-object.key' ]
```

Writing to Document Type tags

The paths above aren't only for reading. Using the same addressing mentioned above on a bi-directional component binding would allow the binding to write back to the specific JSON object in the document.

Using the {this} Keyword

Tags have a built-in "this" keyword, that can be used as a reference to the Tag. The keyword is useful in cases where an expression is being configured on a Tag property and you want to reference the value of another property.

```
// The expression below always returns the name of the Tag  
{this.name}
```

Using {this} in Alarms

The "this" keyword is also available for use on expressions on alarms. However when used like this, the keyword still refers to the Tag, not the alarm. Thus, "["this.name"](#)" on an alarm property expression would return the name of the Tag, not the name of the alarm.

For more information on Ignition's Expression language, see [Expression Overview and Syntax](#).

Related Topics ...

- [Tag Event Scripts](#)

Quality Codes and Overlays

A Quality Code represents a level of confidence in a value. Quality codes are made up of a level (Good, Bad, Error, or Uncertain) and an integer. If a value's quality is not Good, the value generally should *not* be trusted.

Tag Quality in the Designer

In the Tag Browser, find your Tag, expand it, and scroll down to the meta property called **Quality**. Here, you can check the quality of the Tag. This example shows a Good Quality Tag, meaning the Tag can be trusted.

The screenshot shows the Tag Browser interface with the 'Tags' tab selected. A tree view on the left shows a 'Tank Level' node expanded, containing various tags like AlarmEvalEnabled, Deadband, Documentation, etc. On the right, a table lists these tags with their corresponding values and data types. The 'Quality' column is highlighted with a red border. For the 'Quality' tag under 'Tank Level', the value is 'Good' and the data type is 'String'. Other tags like 'AlarmEvalEnabled' have error messages in the Quality column.

On this page ...

- [Tag Quality in the Designer](#)
- [Component Overlays](#)
- [Quality Code Reference Table](#)
- [Perspective Component Overlays](#)
- [Vision Component Overlays](#)
- [Template Overlay](#)
- [Tag Quality and Referenced Tags](#)
- [Overlay Opt-Out](#)



Tag Quality and Overlays

[Watch the Video](#)

One obvious indicator if the Tag is of bad quality is if there is a red error message next to the Tag in the Value column. Hover over the error message to see if there is any additional information about the error. You can also expand the Tag to see the quality issue. This example shows the Query Tag with with a Error Expression Evaluation which helps you resolve the issue promptly.

The screenshot shows the Tag Browser with a 'Query Tag' expanded. Under 'Error_ExpressionEval', the value is 'Error_ExpressionEval("java.lang.NullPointerException")'. The 'Value' column for this tag is highlighted with a red border and contains the error message. The 'Quality' column for this tag is also highlighted with a red border and shows 'Error'.

Component Overlays

HMI screens allow users to quickly gauge the health and accuracy of what is displayed at a glance. In a highly distributed system like Ignition, it is especially important as the client may be located at quite a distance (maybe across the world) from the physical process it is monitoring and controlling.

For these reasons, Perspective and Vision components display visual overlays for various reasons to indicate that the data they are displaying is not good, or pending a reply from the device. Each data binding that drives a component is evaluated for quality. If any of these qualities becomes poor, the Perspective or Vision component will show an overlay. The different overlays can mean different things, denoting their underlying cause. What they indicate is based on the Quality properties of Tags.

Component overlays appear in the Designer workspace, Perspective Session, and Vision Client to let designers and operators know when there is a problem with one of the bindings on a component. What is cool about component overlays is that they not only tell you that there is a problem, but they also help diagnose the problem. Vision and Perspective overlay systems are similar, but each look a little bit different.

The sections below describe in detail Perspective and Vision overlays. Each module has its own Tag Quality Code Reference Table displaying the error codes and what they mean.

Quality Code Reference Table

There are four primary quality codes which quickly inform the user of the quality of the Tag: Good, Uncertain, Bad, and Error.

Note: For users who have upgraded from 7.9, be aware that old code values have been adapted to match the quality codes below. Because the quality of the old values are adapted, tag history quality results and bindings can still be trusted in your upgraded versions. However, it is important to check areas where manual entry may prevent adapted values from populating correctly such as scripts for live checks of tag quality.

The following feature is new in Ignition version **8.1.27**
[Click here](#) to check out the other new features

Note: OPC UA clients are now allowed to write to exposed Tags, which include a StatusCode. This StatusCode is converted to its Qualitycode equivalent, while the write value is passed to the Tag system as a QualifiedValue.

Each quality code has a range of subcodes that provide more specific information about the value of the Tag. The following tables outline the primary data qualities. Quality codes are broken up into several ranges. Note that not every value within a range has a code: the space is there for future codes to be added.

- 0-225: Good quality. The value is generally considered reliable.
- 256 - 511 : Uncertain quality. Generally represents a value that was good, but the reliability is somewhat questionable. These are generally used when the system hasn't received a new value in a timely manner.
- 512 - 767: Bad quality. There's a problem with the value, but it's an "expected" or well recognized type of the problem is. For example, the trial expired or read access was explicitly denied.
- 768 - 1023: Error quality. There is a problem with the value, and the problem was completely unexpected. For example, a query on a query tag couldn't execute, but the Tag doesn't know why. More details on such errors are usually sent to a console somewhere, such as the Gateway's console.

Good Quality	Subcodes 0-225	Meaning
Good_Unspecified	0	A generic "good" code. Generally used in conjunction with a matching good quality subcode, (1,2, or 192).
Good_WritePending	2	Used when a write is in progress. Generally, values use this code until the system knows the write through successfully, which would then result in a 192 code.
Good	192	This data has met all criteria for being considered reliable.
Good_Provisional	200	Good data that should not be considered valid long-term.
Good_Initial	201	Indicates that the value is an initial/seed value for a system that is starting up.
Good_Overload	202	<p>The following feature is new in Ignition version 8.1.2 Click here to check out the other new features</p> <p>Represents good data that is being sampled slower than requested due to a resource limitation.</p>
Good_Backfill	203	Used to indicate good quality values that have arrived out of order. Different systems can choose to process them



Component Overlays

[Watch the Video](#)

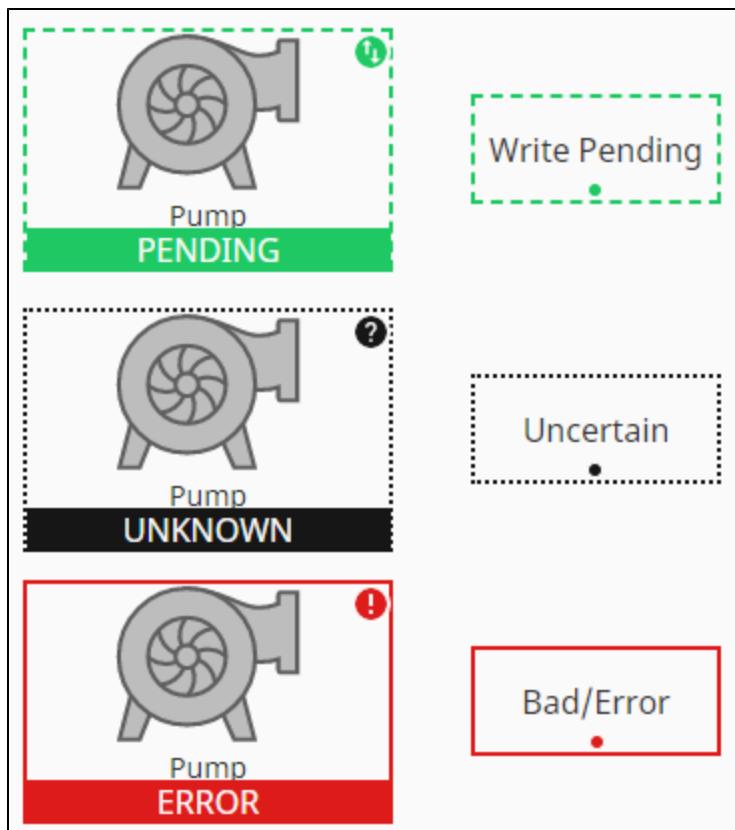
		accordingly.
Uncertain Quality	Subcodes 256 - 511	Meaning
Uncertain	256	An unspecified degree of uncertainty exists in this value.
Uncertain_LastKnownValue	257	The current value is unavailable and represents the last known value.
Uncertain_InitialValue	258	Indicates that a subscription has been made and a good value should be arriving shortly.
Uncertain_DataSubNormal	259	Insufficient good-quality sources required for the derivation of this value.
Uncertain_EngineeringUnitsExceeded	260	Indicates that a value has gone beyond its configured engineering units.
Uncertain_IncompleteOperation	261	An asynchronous operation is currently pending and its result is unknown.
Bad Quality	Subcodes 512 - 767	Meaning
Bad	512	General code for a bad value.
Bad_Unauthorized	513	An unauthorized request was made for data that requires authorization.
Bad_AccessDenied	514	Data requested that requires credentials not held by the requesting user.
Bad_Disabled	515	Data source is currently not enabled.
Bad_Stale	516	Data is out-of-date based upon the requested refresh interval.
Bad_TrialExpired	517	The Trial Mode's timer expired.
Bad_LicenseExceeded	518	The license limit has been exceeded.
Bad_NotFound	519	Object requested was not found.
Bad_ReferenceNotFound	520	Derived or referenced value required an object which was not found.
Bad_AggregateNotFound	521	Requested aggregate was not found.
Bad_NotConnected	522	A connection required for this value is not currently connected.
Bad_GatewayCommOff	523	Connection to the Ignition Gateway is currently turned off. See General Designer Interface .
Bad_OutofRange	524	This value exceeded its allowed range.
Bad_DatabaseNotConnected	525	A database connection required for this value is not connected.
Bad_ReadOnly	526	A write was attempted on a read only target.
Bad_Failure	527	A "failure" code was received from the underlying system. Additional details may be in the diagnostic message. This generally does not indicate an exception, which would be handled by Error_Exception, but instead a simple failure from a system that can return success or failure.
Bad_Unsupported	528	The operation is not supported by the target.
Error Quality	Subcodes 768 - 1023	Meaning
Error	768	An unexpected error occurred while retrieving or calculating this value.

Error_Configuration	769	The source of this value is not configured correctly.
Error_ExpressionEval	770	The source expression was unable to be executed.
Error_TagExecution	771	The source Tag could not be executed.
Error_TypeConversion	772	The actual value was not able to be coerced into the configured data type for the source of this value.
Error_DatabaseQuery	773	A database query required for this value caused an error upon execution.
Error_IO	774	An input/output error occurred while attempting to retrieve or calculate this value.
Error_TimeoutExpired	775	An asynchronous operation failed due to a timeout.
Error_Exception	776	An exception was caught, and logged in the relevant system.
Error_InvalidPathSyntax	777	A path (i.e., Tag path, property path, etc..) was not able to be parsed because the syntax is invalid.
Error_Formatting	778	Attempted formatting (i.e., numeric, date formatting) failed.
Error_ScriptEval	779	A script needed to create this value failed to execute.
Error_CycleDetected	780	Calculating the value involved an execution cycle.

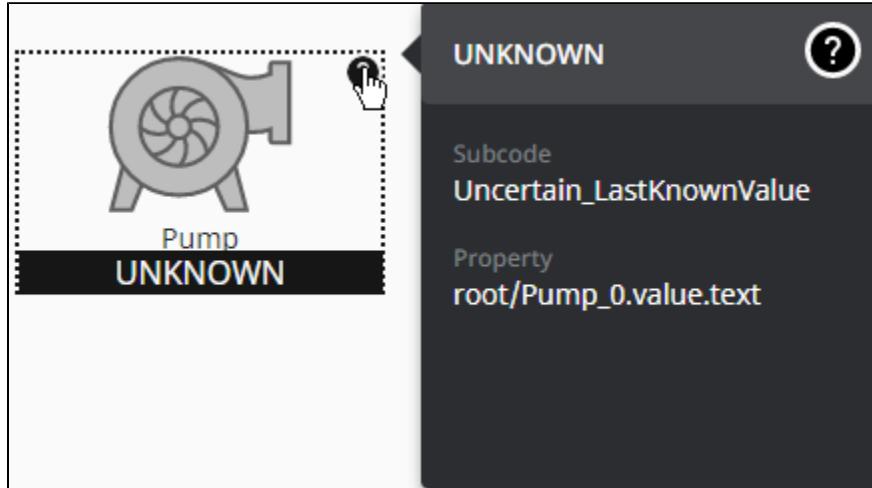
Perspective Component Overlays

Component properties may show a quality overlay if the source of data has a certain code. In Perspective, there are three types of quality overlays, each with a large and small variant.

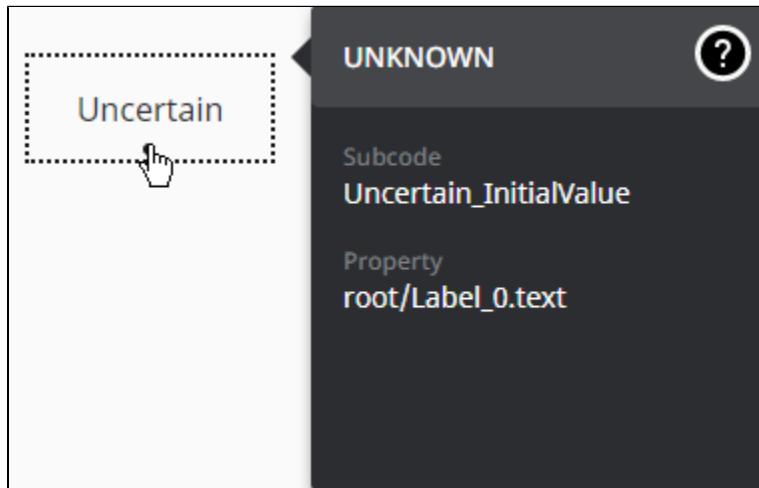
- The Pending overlay is displayed for Good quality, subcode 2 only
- The Unknown overlay is displayed for any Uncertain quality subcode
- The Error overlay is shown for any Bad and Error quality, regardless of subcode



For any large overlay, clicking the icon in the upper right will yield diagnostic information.



For the smaller overlays, clicking on the dot will show the diagnostic information.



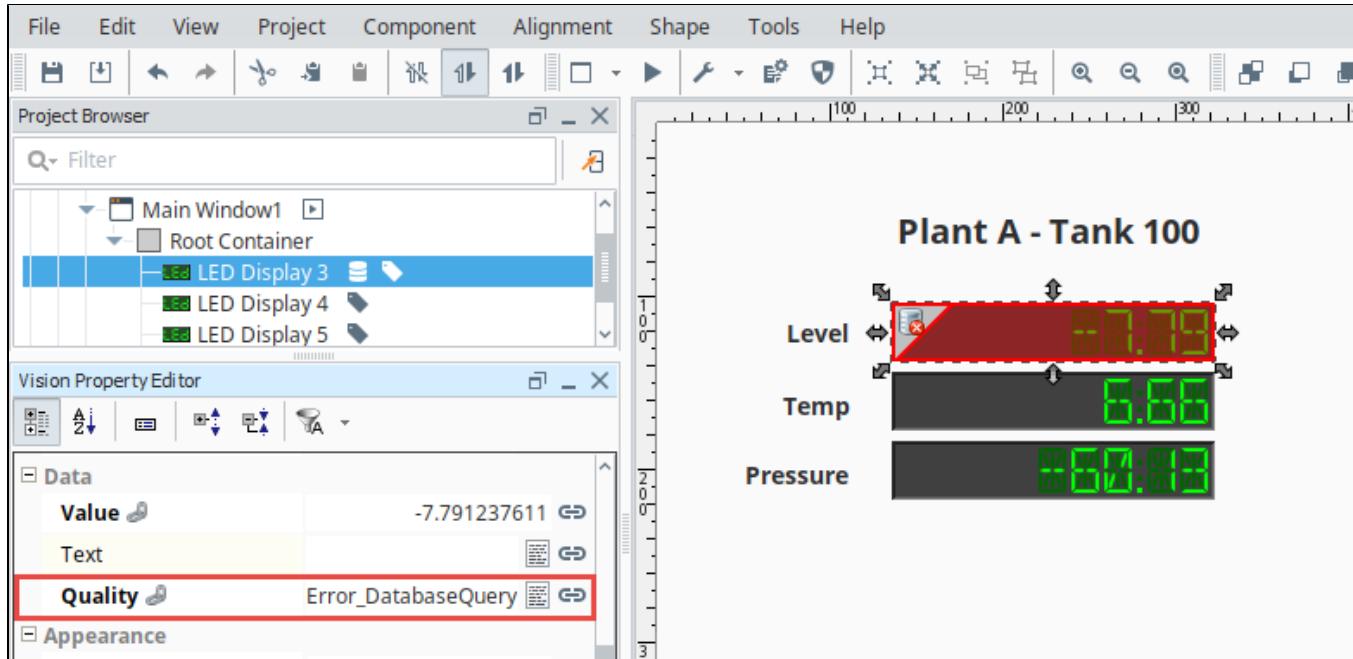
Vision Component Overlays

An overlay on a Vision component lets the operator know that they could be looking at a bad value for that Tag. When the overlay goes away and the values start coming in again, the operator knows that it's a valid Tag, and the values can be trusted.

Component Overlays in Designer Mode

In the following example, you see a red overlay with an icon in the top left corner of the selected LED Display component. The icon gives you a clue to the source of the problem. In this example, it is an SQL Database error. In the Vision Property Editor, the Quality property is highlighted and you'll notice there is a "Error_DatabaseQuery" error message.

The overlays table in the next section show all the possible Vision overlays and what they mean.



Component Overlays in Preview Mode

Let's switch from the Designer mode to the Preview Mode. To put your active view in **Preview Mode**, press the Preview / Designer mode icon ► in the top menubar. Components that have a problem will have a red overlay and an icon in the top left of the component overlay to indicate the problem. The overlay is identical to the overlay that is displayed in the Designer, but the component cannot be selected.

Component Overlays in the Vision Client

Component overlays in a Vision Client work the same way as they do in Preview Mode of the Designer. You have to look at the icon on the overlay to help you diagnose the problem. Go back to the Designer to correct the problem.

Vision Component Overlay Chart

For legacy reasons, Vision has different subcodes. Possible subcodes are listed below.

Good_WritePending		Good_Provisional
Good_WritePending (2)		Good_Provisional (200)
Good_Initial		Good_Overload
Good_Initial (201)		Good_Overload (202)
Good_Backfill		Uncertain
Good_Backfill (203)		Uncertain (1073742080)
Uncertain_LastKnownValue		Uncertain_InitialValue
Uncertain_LastKnownValue (1073742081)		Uncertain_InitialValue (1073742082)
Uncertain_DataSubNormal		Uncertain_EngineeringUnitsExceeded
Uncertain_DataSubNormal (1073742083)		Uncertain_EngineeringUnitsExceeded (1073742084)
Uncertain_IncompleteOperation		Bad
Uncertain_IncompleteOperation (1073742085)		Bad (-2147483136)
Bad_Unauthorized		Bad_AccessDenied
Bad_Unauthorized (-2147483135)		Bad_AccessDenied (-2147483134)
Bad_Disabled		Bad_Stale
Bad_Disabled (-2147483133)		Bad_Stale (-2147483132)
Bad_TrialExpired		Bad_NotFound
Bad_TrialExpired (-2147483131)		Bad_NotFound (-2147483129)
Bad_ReferenceNotFound		Bad_AggregateNotFound
Bad_ReferenceNotFound (-2147483128)		Bad_AggregateNotFound (-2147483127)
Bad_NotConnected		Bad_GatewayCommOff
Bad_NotConnected (-2147483126)		Bad_GatewayCommOff (-2147483125)
Bad_OutOfRange		Bad_DatabaseNotConnected
Bad_OutOfRange (-2147483124)		Bad_DatabaseNotConnected (-2147483123)
Bad_ReadOnly		Bad_Failure
Bad_ReadOnly (-2147483122)		Bad_Failure (-2147483121)
Bad_Unsupported		Error
Bad_Unsupported (-2147483120)		Error (-1073741056)
Error_Configuration		Error_ExpressionEval
Error_Configuration (-1073741055)		Error_ExpressionEval (-1073741054)
Error_TagExecution		Error_TypeConversion
Error_TagExecution (-1073741053)		Error_TypeConversion (-1073741052)
Error_DatabaseQuery		Error_IO
Error_DatabaseQuery (-1073741051)		Error_IO (-1073741050)
Error_TimeoutExpired		Error_Exception
Error_TimeoutExpired (-1073741049)		Error_Exception (-1073741048)
Error_InvalidPathSyntax		Error_Formatting
Error_InvalidPathSyntax (-1073741047)		Error_Formatting (-1073741046)
Error_ScriptEval		Error_CycleDetected
Error_ScriptEval (-1073741045)		Error_CycleDetected (-1073741044)

Template Overlay

In addition, there is a quality overlay for an unknown template. This can occur when the Template Path property on a template instance does not resolve to an existing template. This can happen due to a name change on the template definition, or if a project import file was recently imported, and included usages of a template instance but not the definition.

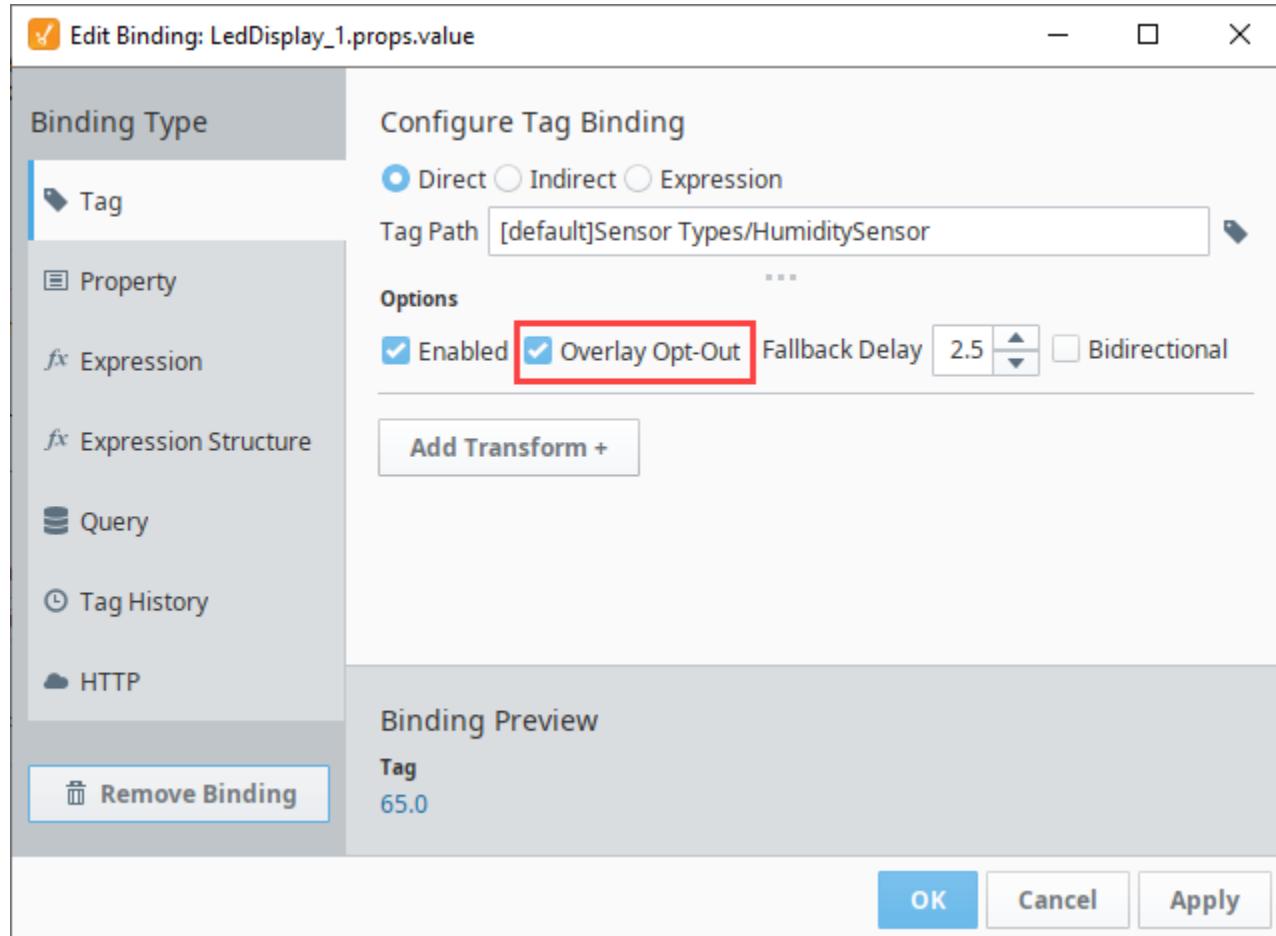


Tag Quality and Referenced Tags

When Tags reference other Tags, such as in expressions, they will often pass the worst sub-quality up as their own. For example, even though a particular Tag's expression executes without problem, if the expression references a Tag whose quality is "Bad", the expression Tag will also report "Bad."

Overlay Opt-Out

Choosing the Overlay Opt-Out option will ignore the quality of the chosen Tag, making it have no effect on the component's quality overlay. The Overlay Opt-Out option is located in the Tag bindings for both Perspective and Vision components. If this option is enabled, the operator will not see any overlays and will have no indication that the underlying Tag quality is something other than good. A word of caution when you use the Opt-Out option because you always want to give the operator some indication that the values they are seeing on the screen can be trusted, and by opting out, you are removing that indicator for the operator.



Related Topics ...

- [Tag Scaling Properties](#)
- [Tag Properties](#)
- [Bindings in Perspective](#)
- [Tag Bindings in Vision](#)
- [Indirect Tag Bindings in Vision](#)

Exporting and Importing Tags

Ignition can export and import tag configurations to and from the **JSON** (JavaScript Object Notation) format. You can import **XML** (Extensible Markup Language) or **CSV** (Comma Separated Value) file formats as well, but Ignition will convert them to **JSON** format while in a Tag Provider. Tag export files can be edited directly in any text editor, allowing you to make bulk edits to tags before importing them back into a Tag Provider.

Export Tags

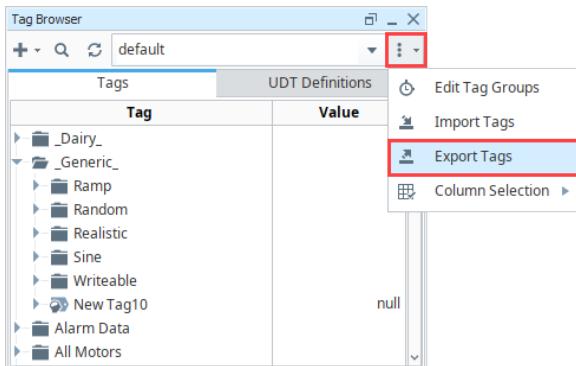
1. In the **Designer**, go to the **Tag Browser**, and select the tags tab to export all your tags. You can also select an individual folder that contains tags you want to export. You can even export individual tags as long the individual tags are in the same folder.

Exporting UDTs and UDT Instances

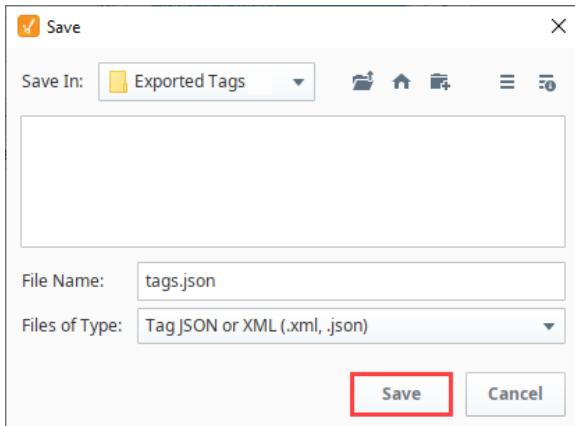
If you select a UDT instance to export, the UDT definition is not automatically included. You must export the definition as a separate file by clicking on the UDT Definitions tab.

When later importing these UDTs, it is recommended to import UDT definitions before importing any instances.

2. On the **Tag Browser** toolbar, click the More Options menu  to open the dropdown. Select **Export Tags**.



3. The **Save** window will open. Specify the folder where you want to save your exported tag files, and then click **Save**. Ignition will export tags by default to a **.json** file.



Import Tags

You can import tags to an individual folder or under the Tags tab. To import tags under the Tags tab, you can click the Tags tab or the empty space at the bottom of the Tag Browser below all your tags. When importing tags, you need to tell Ignition how to handle duplicate tags. If any of the tags being imported already exist in the folder you specify, Ignition can abort the import, overwrite the tags, rename them, ignore them, or merge them. This is called a collision policy.

On this page ...

- [Export Tags](#)
- [Import Tags](#)
 - [Advanced Tag Import](#)
- [Tag File Formats](#)
 - [CSV Format](#)
 - [Property Values in the CSV Import](#)
 - [Tag Properties](#)
- [JSON Example](#)
- [XML Example](#)



Importing and Exporting Tags

[Watch the Video](#)

Collision Policy Options Table

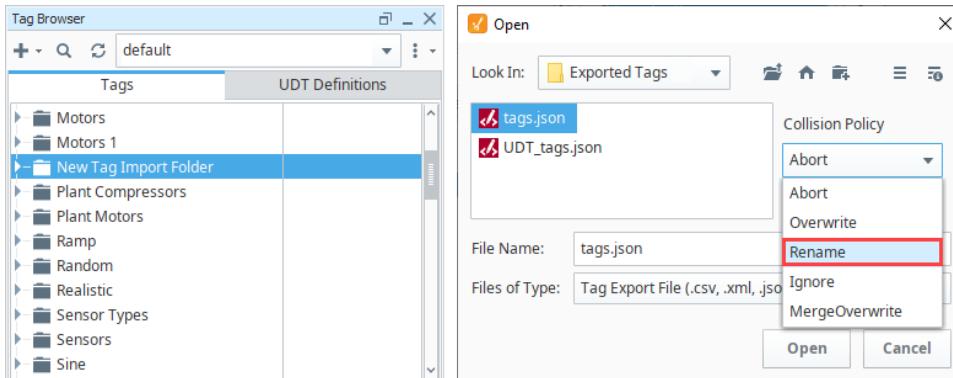
Policy	Description
Abort	Aborts the import if duplicate tags are found.
Overwrite	Overwrites any tags in the folder that have the same name as tags being imported. Note this is a complete overwrite of the tag. This feature was changed in Ignition version 8.1.8 : As of 8.1.8, importing UDT definitions with this collision policy will remove any members that are not included in the import file.
Rename	Renames any duplicate tags.
Ignore	Ignores duplicate tags and imports only those that are unique.
MergeOverwrite	Overwrites the tag with the exception of any properties that aren't defined in the import folder. Those properties will be merged.

To import tags, do the steps that follow.

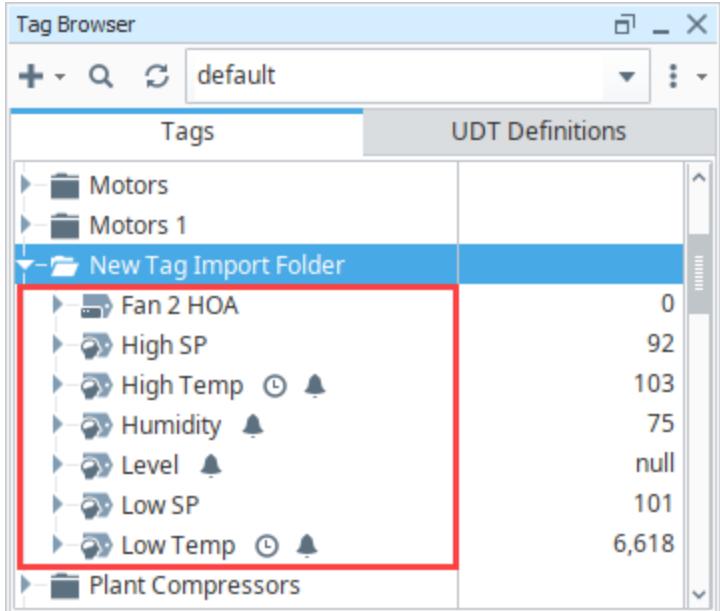
1. In the Tag Browser toolbar, right click on your folder and select **Import Tags**.

This feature was changed in Ignition version **8.1.15**:
In version 8.1.15 and newer, right click on your folder and select **Import Tags > Direct**.

2. Specify the folder you want to import your Tags from, and choose a previously exported file either **.json**, **.xml**, or **.csv** file type.
3. Choose a Collision Policy, which indicates how Ignition will deal with duplicate tags.
4. Click **Open** to import the tags.



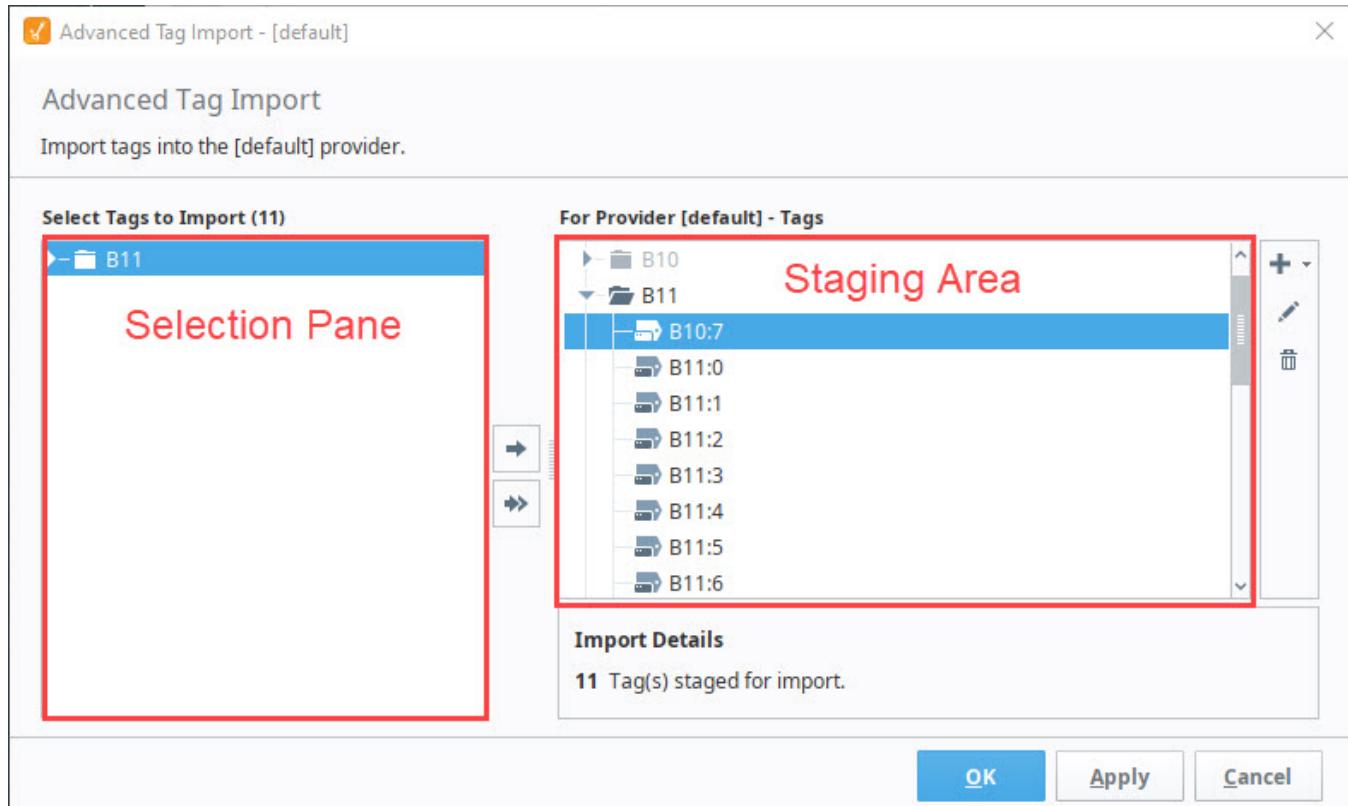
5. The tags now appear in the Tag Browser in the folder you indicated for the import.



Advanced Tag Import

The following feature is new in Ignition version 8.1.15
[Click here](#) to check out the other new features

The Advanced Tag Import tool allows users to edit tag properties in the Designer before import. It combines the functionality of the basic Import Tags tool with the Tag Editor window.



Interface Element	Description
	Adds the selected tag, folder, or UDT to the Staging Area.

	Adds all tags, folders, and UDTs in the Selection Pane to the Staging Area.
	Creates a new folder or tag at the selected node in the Staging Area, or the root if a node is not selected.
	Opens a Tag Editor window for the selected tag or UDT.
	Deletes the selected tag or UDT from the Staging Area.

1. From the Tag Browser, click the **More Options** menu and select **Import Tags > Interactive**.
2. Select the **.json**, **.xml**, or **.csv** file to import.

3. Move the desired tags and UDTs from the Selection Pane to the Staging Area using the **Add** or **Add All** buttons.
4. Edit tags in the Staging Area by selecting a tag and clicking **Edit** .
5. If any collisions are detected, you will be prompted to select a [Collision Policy](#) from the dropdown menu.

Import Details

10 collision(s) detected. Set collision action to: **Overwrite**

11 Tag(s) staged for import.

Note: Unlike the basic Import Tags tool, you may only choose to **Overwrite** or **Ignore** duplicate tags imported via the Advanced Tag Import tool.

6. Click **OK** to import the tags and UDTs.
7. The tags now appear in the Tag Browser.

Note: If your tags are not appearing as expected for a large tag import, the Designer's memory allocation may need to be increased. Access your Gateway and navigate to the **Config > Gateway Settings > Designer Memory** to adjust memory limitations. The default size is 1.0 GB, with available dropdown options from MB128 to 4.0 GB.

Tag File Formats

Tags can be imported from CSV, JSON, and XML. tags can only be exported in XML or JSON. There are many configuration settings for tags than what is displayed in a JSON or XML export file. The tag export feature only exports the configuration properties that have been edited in at least one of the tags in the selected export folder. Therefore, to ensure the desired configuration setting is available in the export file, at least one tag within the selected export folder must have that configuration property changed

CSV Format

Importing

Ignition supports importing tags from a CSV format. Details of the format are below, if you expand "CSV Example Format". This format can contain tag types, OPC paths, and most tag properties. One difference between the CSV format and the XML and JSON format is that the CSV format does not include support for alarm configurations. Alarms can certainly be added to tags in the Ignition Designer after tags have been imported from CSV, but alarms cannot be defined directly in the CSV.

Below is an example of the legacy CSV format, which contains a couple of different tag types, purely for demonstrative purposes. It contains:

- An OPC Tag
- A Folder
- An OPC Tag located in a folder
- A Derived Tag
- An Expression Tag
- A Memory Tag
- A Query Tag

```

Path,Name,Owner,TagType,DataType,Value,Enabled,AccessRights,OPCServer,OPCItemPath,ScanClass,DriverName,
ScaleMode,RawLow,RawHigh,ScaledLow,ScaledHigh,ClampMode,ScaleFactor,Deadband,DeadbandMode,FormatString,
EngUnit,EngLow,EngHigh,EngLimitMode,Tooltip,Documentation,ExpressionType,Expression,OPCWriteBackServer,
OPCWriteBackItemPath,SQLBindingDatasource,HistoryEnabled,PrimaryHistoryProvider,HistoricalScanclass,
HistoricalDeadband,HistoricalDeadbandMode,InterpolationMode,HistoryMaxAgeMode,HistoryMaxAge,
HistoryTimestampSource,UDTParentType,PersistValue,SourceDataType,SourceTagPath,SQLBindingPollRate,Permissions
# version=1,,,
,_types_,,6,2,,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,0,,,,FALSE,,,
Default Historical,0.01,0,3,0,1,0,,FALSE,,,
,A Folder,,6,2,,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,0,,,,FALSE,,,
Default Historical,0.01,0,3,0,1,0,,FALSE,,,
A Folder/,OPC in a folder,,0,2,,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,0,,,,FALSE,,,
Default Historical,0.01,0,3,0,1,0,,FALSE,,,
,Derived Tag,,13,2,100,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,0,,,,FALSE,,,
Default Historical,0.01,0,3,0,1,0,,FALSE,,,[~]Expression Tag,,,
,Expression Tag,,1,2,100,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,1,"//This is an expression
100",,,FALSE,,Default Historical,0.01,0,3,0,1,0,,FALSE,,,
,Memory Tag,,1,7,I'm a memory Tag,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,0,,,,FALSE,,,
Default Historical,0.01,0,3,0,1,0,,FALSE,,,
,OPC Tag,,0,2,,TRUE,Read_Write,Ignition OPC-UA Server,[devicename]folder/path,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,0,,,,FALSE,,,
Default Historical,0.01,0,3,0,1,0,,FALSE,,,
,Query Tag,,1,2,,TRUE,Read_Write,,,Default,,0,0,100,0,10,0,0,1.00E-04,0,"#,##0.##",,0,100,0,,,2,"/*Some
Query*/
SELECT 100",,,FALSE,,Default Historical,0.01,0,3,0,1,0,,FALSE,,,

```

Property Values in the CSV Import

The following table shows the configuration property names and values contained in legacy CSV tag import files. Tags were overhauled in Ignition 8.0, so the properties listed here are not the same as tags export from an Ignition 8.0+ system. For modern tag properties, see the [Tag Properties](#) page.

Tag Properties			
Property Name	Type	Values (if applicable)	Description
Value			The value of the tag, dependent on the data type.
Data Type	Int	0 - Int1 1 - Int2 2 - Int4 3 - Int8 4 - Float4 5 - Float8 6 - Boolean 7 - String 8 - DateTime 9 - DataSet	
Enabled	Boolean	true/false	
Tagtype	Int	0 - OPC Tag 1 - DB Tag (see ExpressionType) 2 - Client Tag 6 - Folder 10 - UDT Instance 13 - Derived Tag	Determines the type of the tag. A value of one is a "DB Tag", which is either a Memory Tag, Query Tag, or Expression Tag, depending on the value of the ExpressionType field.
ExpressionType	Int	0 - None 1 - Expression 2 - SQL Query	Used in conjunction when the TagType is set to 1 (DB Tag), otherwise this field is ignored.
AccessRights	Int	0 - Read Only 1 - Read/Write 2 - Custom	If custom, will be defined by a Permissions Tag.
OPCServer	String		
OPCItemPath	String		
OPCWriteBackServer	String		Write back target for expression Tags.
OPCWriteBackItemPath	String		

ScaleMode	Int	0 - Off 1 - Linear 2 - Square Root 3 - Exponential Filter	
ScaleFactor	Float		For exponential filter.
RawLow	Float		Defines scale range.
RawHigh	Float		
ScaledLow	Float		
ScaledHigh	Float		
ClampMode	Int	0 - None 1 - Low 2 - High 3 - Both	
Deadband	Float		
DeadbandMode	Int	0 - Absolute 1 - Percentage	
FormatString	String		
EngUnit	String		
EngLow	Float		
EngHigh	Float		
EngLimitMode	Int	0 - None 1 - Low 2 - High 3 - Both	
Tooltip	String		
Documentation	String		
DriverName	String		Used for external tags.
ScanClass	String		The export will only include the name of the Scanclass, not the configuration of the Scanclass itself. A Scanclass with the same name needs to already exist on the Gateway that the tags are being imported to, prior to importing them.
HistoryEnabled	Boolean	true/false	
PrimaryHistory Provider	String		The history provider to use if storing history.
HistoricalDeadband	Float		
HistoricalDeadbandMode	Int	0 - Absolute 1 - Percentage	
HistoricalScanclass	String		
InterpolationMode	Int	0 - Discrete 2 - Analog (deadband) 3 - Analog (compressed)	How values are interpolated. 2 exists for backwards compatibility (and is equivalent to 1), but only 0 or 3 should be used in the future.
HistoryTimestampSource	Int	0 - System 1 - Value	
HistoryMaxAge Mode	Int	0 - Unlimited 1 - Limited	
HistoryMaxAge	Int		Max cycles between storage.
UDTParentType	String		The path to the parent UDT type. Used by sub-types and instances.

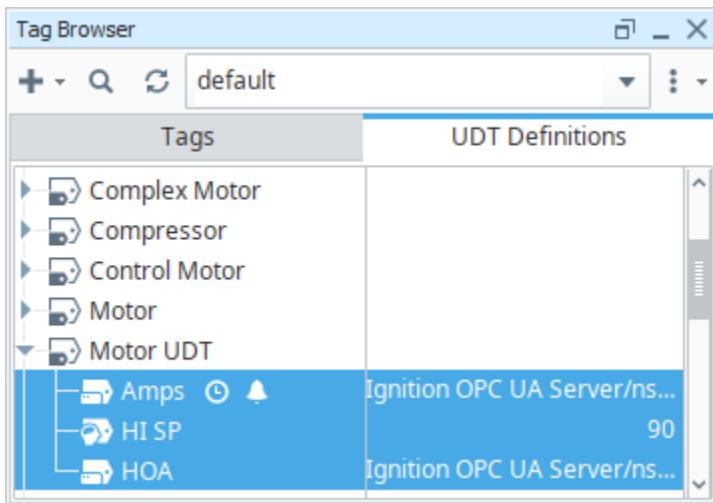
Exporting

Although Ignition can import tags from a CSV format, Ignition does not export tags to a CSV format. Since multiple alarms can be configured per tag, the XML or JSON formats provide a much better format to allow for the tree structures needed to fully represent Ignition tags.

JSON Example

In this example, we exported three tags from the Motor UDT in JSON format.

- Amps - Expression tag with an Alarm and History enabled
- HI SP - Memory tag which is bound to a parameter to an OPC tag
- HOA - OPC tag



The following exported tag file is in JSON format. As you browse through the JSON file, you will see the tag properties and configuration settings for each of the three tags listed above.

Below is an example of a JSON format tag export. Descriptions on the various properties can be found on the [Tag Properties](#) page.

Exported Tags in JSON Format

```
{  
    "tags": [  
        {  
            "opcItemPath": {  
                "bindType": "parameter",  
                "binding": "ns\u003d1;s\u003d[Dairy]_Meta:Overview/Motor {MotorNumber}/Amps"  
            },  
            "valueSource": "opc",  
            "historyProvider": "MySQL",  
            "alarms": [  
                {  
                    "mode": "BelowValue",  
                    "setpointA": 25.0,  
                    "name": "Low Amps",  
                    "priority": "Critical",  
                    "displayPath": {  
                        "bindType": "Expression",  
                        "value": "Motor{MotorNumber}"  
                    }  
                },  
                {  
                    "mode": "AboveValue",  
                    "name": "High Amps",  
                    "priority": "Critical",  
                    "setpointA": {  
                        "bindType": "Expression",  
                        "value": "{[.]HI SP}"  
                    }  
                }  
            ],  
            "name": "Amps",  
            "historyEnabled": true,  
            "tagType": "AtomicTag",  
            "value": 90  
        },  
        {  
            "opcItemPath": {  
                "bindType": "parameter",  
                "binding": "ns\u003d1;s\u003d[Dairy]_Meta:Overview/Motor {MotorNumber}/HI SP"  
            },  
            "valueSource": "opc",  
            "historyProvider": "MySQL",  
            "alarms": [  
                {  
                    "mode": "BelowValue",  
                    "setpointA": 25.0,  
                    "name": "Low HI SP",  
                    "priority": "Critical",  
                    "displayPath": {  
                        "bindType": "Expression",  
                        "value": "Motor{MotorNumber}"  
                    }  
                },  
                {  
                    "mode": "AboveValue",  
                    "name": "High HI SP",  
                    "priority": "Critical",  
                    "setpointA": {  
                        "bindType": "Expression",  
                        "value": "{[.]HOA}"  
                    }  
                }  
            ],  
            "name": "HI SP",  
            "historyEnabled": true,  
            "tagType": "AtomicTag",  
            "value": 90  
        },  
        {  
            "opcItemPath": {  
                "bindType": "parameter",  
                "binding": "ns\u003d1;s\u003d[Dairy]_Meta:Overview/Motor {MotorNumber}/HOA"  
            },  
            "valueSource": "opc",  
            "historyProvider": "MySQL",  
            "alarms": [  
                {  
                    "mode": "BelowValue",  
                    "setpointA": 25.0,  
                    "name": "Low HOA",  
                    "priority": "Critical",  
                    "displayPath": {  
                        "bindType": "Expression",  
                        "value": "Motor{MotorNumber}"  
                    }  
                },  
                {  
                    "mode": "AboveValue",  
                    "name": "High HOA",  
                    "priority": "Critical",  
                    "setpointA": {  
                        "bindType": "Expression",  
                        "value": "{[.]Amps}"  
                    }  
                }  
            ],  
            "name": "HOA",  
            "historyEnabled": true,  
            "tagType": "AtomicTag",  
            "value": 90  
        }  
    ]  
}
```

```

        "opcServer": "Ignition OPC UA Server"
    },
{
    "valueSource": "memory",
    "name": "HI SP",
    "value": 90,
    "tagType": "AtomicTag"
},
{
    "opcItemPath": {
        "bindType": "parameter",
        "binding": "ns\u003d1;s\u003d[Dairy]_Meta:Overview/Motor {MotorNumber}/HOA"
    },
    "valueSource": "opc",
    "name": "HOA",
    "tagType": "AtomicTag",
    "opcServer": "Ignition OPC UA Server"
}
]
}

```

XML Example

In this example, we exported the same three tags from our Tag Browser, that were also used in the JSON example, in XML format. As you browse through the XML file, you will see the tag properties and configuration settings for each of the same three tags.

Below is an example of an XML format tag export. Descriptions on the various properties can be found on the [Tag Properties](#) page.

Exported Tags in XML Format

```

<Tags MinVersion="8.0.0" locale="en_US">
    <Tag name="Amps" type="AtomicTag">
        <Property name="opcItemPath" boundValueType="parameter">ns=1;s=[Dairy]_Meta:Overview/Motor
{MotorNumber}/Amps</Property>
        <Property name="valueSource">opc</Property>
        <Property name="historyProvider" datatype="String">MySQL</Property>
        <CompoundProperty name="alarms">
            <PropertySet>
                <Property name="mode">3</Property>
                <Property name="setpointA">25</Property>
                <Property name="name">Low Amps</Property>
                <Property name="priority">4</Property>
                <Property name="displayPath" bindtype="Expression">Motor{MotorNumber}</Property>
            </PropertySet>
            <PropertySet>
                <Property name="mode">2</Property>
                <Property name="name">High Amps</Property>
                <Property name="priority">4</Property>
                <Property name="setpointA" bindtype="Expression">{[.]HI SP}</Property>
            </PropertySet>
        </CompoundProperty>
        <Property name="historyEnabled" datatype="Boolean">true</Property>
        <Property name="opcServer">Ignition OPC UA Server</Property>
    </Tag>
    <Tag name="HI SP" type="AtomicTag">
        <Property name="valueSource">memory</Property>
        <Property name="value">90</Property>
    </Tag>
    <Tag name="HOA" type="AtomicTag">
        <Property name="opcItemPath" boundValueType="parameter">ns=1;s=[Dairy]_Meta:Overview/Motor
{MotorNumber}/HOA</Property>
        <Property name="valueSource">opc</Property>
        <Property name="opcServer">Ignition OPC UA Server</Property>
    </Tag>

```

</Tags>

Related Topics ...

- [Tag Data Types](#)
- [Tag Properties](#)

Tag Editor

This feature was changed in Ignition version 8.1.17:

In 8.1.17, the Tag Editor was redesigned to improve usability. The new Tag Editor now requires fewer clicks and keeps relevant tag information visible while modifying bindings, alarms, and event scripts. Additional improvements include the following:

- Only expression bindings now require opening the dialog to change. Tag and Parameter bindings can be edited in-line within the tag editor and are not shortened.
- Sorting in UDT Editor now sorts folders first, then tags.
- Users can add tags in the UDT Editor without having to click the root of the UDT/folder.
- UDT Parameters now show when they are overridden.
- The Tag Editor uses the entire width allowed to display content.
- UDT properties are added in-line to bindings now rather than clearing out the content.
- Property values are now all left-aligned.
- The Expression editor now accepts drag and drops from the UDT Editor or Tag Browser.

This feature was changed in Ignition version 8.1.34:

In 8.1.34, [Tag Diagnostics](#) data was removed from the Tag Editor and now is displayed it's own window.

Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

The Tag Editor is a robust interface that contains all the properties that can be configured for Tags. In the Tag Editor, you set the Tag's name, value, numeric and meta data properties, security, alarming, history, and more. For information an explanation on all possible tag configurations, see [Tag Properties](#).

Categories

- All Properties
- Basic
- Meta Data
- Value**
- Numeric
- Security
- Scripting (1)
- Alarms (1)
- History
- Custom

All Properties

Basic	
Name	Ramp3
Tag Group	Default
Enabled	true
Meta Data	
Documentation	
Tooltip	
Value	
Value Source	OPC
Data Type	Double
OPC Server	Ignition OPC UA Server
OPC Item Path	ns=1;s=[Sample_Device]_Meta:Ramp/Ramp3
Numeric	
Deadband Mode	Absolute
Deadband	0.0001
Scale Mode	Off
Engineering Units	
Engineering Low Limit	0.0
Engineering High Limit	100.0
Engineering Limit Mode	No_Clamp

Buttons: OK, Apply, Cancel

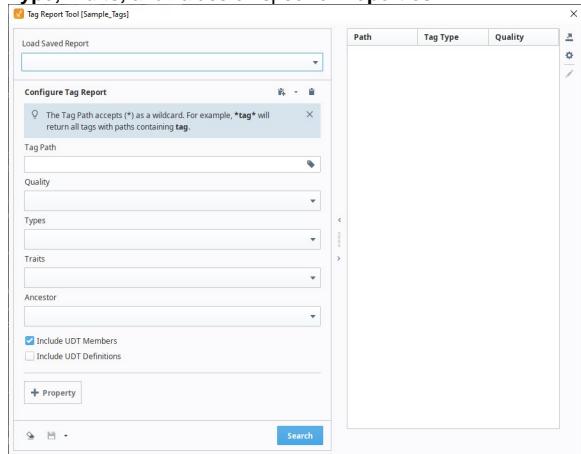
Diagnostics, Documentation, and Refresh icons are included on the upper right-hand side of the Tag Editor to access more tag information:

- The **Diagnostics**  icon opens the [Tag Diagnostics](#) window.
- The **Documentation**  icon displays a Documentation pane at the bottom of the Tag Editor.
- The **Refresh**  icon refreshes the editor with the current definition of the tag from the Gateway.

Tag Report Tool

The following feature is new in Ignition version 8.1.19
[Click here](#) to check out the other new features

The Tag Report Tool allows users to search for tags by multiple criteria, including **Tag Path**, **Quality**, **Type**, **Traits**, and values on specific **Properties**.



On this page ...

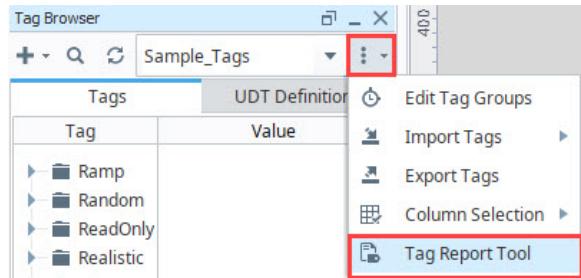
- [Configure, Save, and Copy Reports](#)
- [Export Reports](#)
- [View Tags in Context](#)

Interface Elements

Configure Tag Report

- [Tag Path](#)
- [Quality](#)
- [Types](#)
- [Traits](#)
- [Ancestor](#)
- [UDT Members and Definitions](#)
- [Properties](#)

The Tag Report Tool is accessible through the **More Options**  menu on the Tag Browser. Select the Tag Provider you wish to search, then click the  icon and select **Tag Report Tool** to begin configuring your report.



Note: The Tag Report Tool can only search within one Tag Provider at a time. To search within a different Tag Provider, exit the Tag Report Tool and select a new Tag Provider from the drop-down menu.

Configure, Save, and Copy Reports

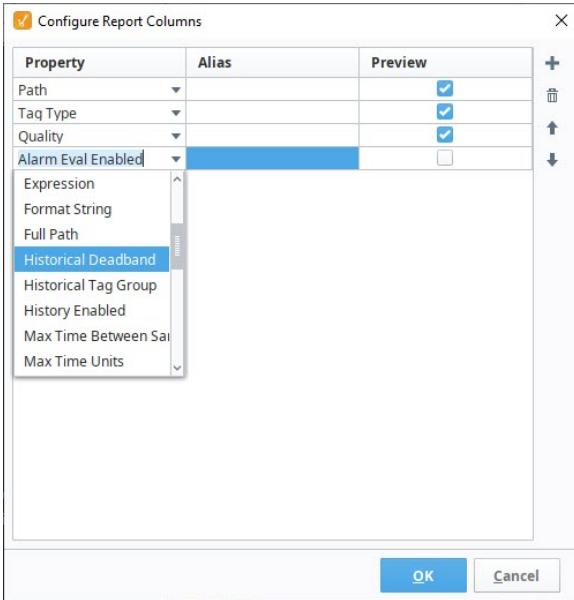
Save your query as a Saved Report to easily repeat searches with complex criteria. Saved reports are accessible to all users and all designers on the same Gateway, and can be reused to generate reports for different Tag Providers.

You can also **Copy**  a report as either a JSON object or as a script. Copying a report as a

JSON object allows you to **Paste**  the same query into the Tag Report Tool on a separate Gateway. Copying a report as a Script will allow you to paste a `system.tag.query()` script anywhere in Ignition where scripting can be used.

Export Reports

Export reports as CSV files to analyze data in Excel or other platforms outside of Ignition. Configure Report Columns to add, remove, reorder, or assign aliases to columns displayed in the report.



The following feature is new in Ignition version **8.1.21**
[Click here](#) to check out the other new features

New columns are now available for export:

Property	Description
Overrides	Lists all overridden properties on the tag.
Timestamp	A timestamp representing when the stored value was seen by the system. This property can also be used to filter reports .

View Tags in Context

Right-click on any tag in the Selection Pane to view additional options, such as **Edit**, **Delete**, or **Copy Path**. Choose **Select in Tag Browser** to highlight the selected tag in the Tag Browser.



Changes made within the Tag Report Tool will edit tags directly. Deleting a tag in the Tag Report Tool will also delete the tag from the Tag Browser, not just from the generated report.

The screenshot shows the Tag Report Tool [default] window and the Tag Browser window side-by-side. In the Tag Report Tool, a context menu is open over a row in the table. The menu items are: Edit, Delete, Copy, Copy Def as JSON, Copy Path, and Select in Tag Browser. A red arrow points from the 'Select in Tag Browser' option to the 'Select in Tag Browser' button in the bottom right corner of the Tag Report Tool window.

Path	Tag Type	Quality
Ramp/Ramp7	AtomicTag	Good
Ramp/Ramp0	AtomicTag	Good
Ramp/Ramp4	AtomicTag	Good
Ramp/Ramp2	AtomicTag	Good
Ramp/Ramp9	AtomicTag	Good
Ramp/Ramp1	AtomicTag	Good
Ramp/Ramp5	AtomicTag	Good
Ramp/Ramp6	AtomicTag	Good
Ramp/Ramp3	AtomicTag	Good
Ramp/Ramp8	AtomicTag	Good

Interface Elements

Icon	Interface Element	Description
	Copy Query	<p>Allows the user to copy the configured query:</p> <ul style="list-style-type: none"> Copy query as JSON - Copies the query as a JSON object that can be pasted into the Tag Report Tool in another Gateway. Copy as Script - Copies the query as a script (<code>system.tag.query</code>)
	Paste JSON Query	Paste a JSON query into the Tag Report Tool.
	Download Tag Report	Exports the Tag Report as a CSV file.
	Configure Report Columns	Add, remove, reorder, or assign aliases to columns displayed in the report.
	Edit tag	<p>Opens the Tag Editor for the selected tag.</p> <p>Warning: Changes made within the Tag Report Tool will edit tags directly. Deleting a tag in the Tag Report Tool will also delete the tag from the Tag Browser, not just from the generated report.</p>
	Clear	Clears the configured query.
	Save	Save options for the configured query:

		<ul style="list-style-type: none"> Save - Saves changes to an existing saved report. Save As - Save as a new saved report. Manage saved reports - Opens a new window
	Load Saved Report	Select and load a previously saved report. Saved reports are accessible to all users and all designers on the same Gateway. 

Configure Tag Report

Configure a search query using the criteria on the left-hand panel and click **Search** to generate a new report.

Tag Path

Search for tags by full or partial Tag Path. The Tag Path accepts (*) as a wildcard anywhere in the search string.

 Tag Report Tool [Sample_Tags]

Load Saved Report

Configure Tag Report

The Tag Path accepts (*) as a wildcard. For example, *tag* will return all tags with paths containing tag.

Tag Path

Quality

Types

Traits

Ancestor

Include UDT Members

Include UDT Definitions

+ Property

Path	Tag Type	Quality
Ramp/Ramp1	AtomicTag	Bad_TrialExpir...
Ramp/Ramp0	AtomicTag	Bad_TrialExpir...
Ramp/Ramp3	AtomicTag	Bad_TrialExpir...
Ramp/Ramp7	AtomicTag	Bad_TrialExpir...
Ramp/Ramp8	AtomicTag	Bad_TrialExpir...
Ramp/Ramp2	AtomicTag	Bad_TrialExpir...
Ramp/Ramp9	AtomicTag	Bad_TrialExpir...
Ramp/Ramp4	AtomicTag	Bad_TrialExpir...
Ramp/Ramp5	AtomicTag	Bad_TrialExpir...
Ramp/Ramp6	AtomicTag	Bad_TrialExpir...

Complete - 10 results

Search

Quality

Search for tags by Quality. You may search by **Good**, **Bad**, **Error**, **Uncertain** or narrow your search down to specific qualities like **Good_Initial** or **Bad_Stale**. For more information on tag qualities, see [Quality Codes and Overlays](#).

Tag Report Tool [Sample_Tags]

Load Saved Report

Configure Tag Report

The Tag Path accepts (*) as a wildcard. For example, *tag* will return all tags with paths containing tag.

Tag Path

Quality

Good

- Good
- Good_Unspecified
- Good_WritePending
- Good_Provisional
- Good_Initial
- Good_Overload
- Good_Backfill

Bad

Error

Uncertain

Clear

Search

Path	Tag Type	Quality
ReadOnly/Rea...	AtomicTag	Good
Ramp/Ramp5	AtomicTag	Good
Ramp/Ramp7	AtomicTag	Good
Ramp/Ramp1	AtomicTag	Good
Ramp/Ramp0	AtomicTag	Good
Ramp/Ramp2	AtomicTag	Good
Ramp/Ramp4	AtomicTag	Good
Ramp/Ramp6	AtomicTag	Good
Ramp/Ramp9	AtomicTag	Good
Ramp/Ramp3	AtomicTag	Good
Ramp/Ramp8	AtomicTag	Good

Complete - 74 results

Types

Search for tags by their value source.

Tag Report Tool [Sample_Tags]

Load Saved Report

Configure Tag Report

The Tag Path accepts (*) as a wildcard. For example, *tag* will return all tags with paths containing tag.

Tag Path

Quality

Types

OPC

- OPC
- Memory
- Expression
- Query
- Reference
- Derived
- UDT Type
- UDT Instance

Search

Path	Tag Type	Quality
ReadOnly/Rea...	AtomicTag	Good
Ramp/Ramp5	AtomicTag	Good
Ramp/Ramp7	AtomicTag	Good
Ramp/Ramp1	AtomicTag	Good
Ramp/Ramp0	AtomicTag	Good
Ramp/Ramp2	AtomicTag	Good
Ramp/Ramp4	AtomicTag	Good
Ramp/Ramp6	AtomicTag	Good
Ramp/Ramp9	AtomicTag	Good
Ramp/Ramp3	AtomicTag	Good
Ramp/Ramp8	AtomicTag	Good

Complete - 74 results

Traits

Search for tags by specific traits. Searchable traits include the following:

- Event Script Configured
- Value Scaling Applied
- Alarm(s) Configured
- History Enabled
- Overrides Parent Properties
- Custom Security Permissions
- Tag Disabled

Note: Trait searches configured in the Tag Report Tool are AND queries. Selecting multiple criteria will return a report including only tags that match all selected criteria.

Tag Report Tool [Sample_Tags]

Load Saved Report

Configure Tag Report

The Tag Path accepts (*) as a wildcard. For example, *tag* will return all tags with paths containing tag.

Tag Path

Quality

Types

Traits

Alarm(s) Configured

- Event Script Configured
- Value Scaling Applied
- Alarm(s) Configured
- History Enabled
- Overrides Parent Properties
- Custom Security Permissions
- Tag Disabled

Path Tag Type Quality

Path	Tag Type	Quality
Ramp/Ramp3	AtomicTag	Good
Writeable/Writ...	AtomicTag	Good
Writeable/Writ...	AtomicTag	Good

Complete - 3 results

This screenshot shows the 'Tag Report Tool' interface for a database named 'Sample_Tags'. On the left, there's a sidebar with sections for 'Load Saved Report', 'Configure Tag Report' (with a note about wildcards), 'Tag Path', 'Quality', 'Types', and 'Traits'. The 'Traits' section is expanded and highlighted with a red box, showing various configuration options like 'Event Script Configured', 'Value Scaling Applied', and 'Alarm(s) Configured'. An 'Alarm(s) Configured' checkbox is checked. On the right, a main panel displays a table of search results with columns for 'Path', 'Tag Type', and 'Quality'. Three entries are listed: 'Ramp/Ramp3' (AtomicTag, Good), 'Writeable/Writ...' (AtomicTag, Good), and 'Writeable/Writ...' (AtomicTag, Good). The bottom right corner of the main panel says 'Complete - 3 results'.

Ancestor

Search for tags by UDT Parent types.

Tag Report Tool [default]

Load Saved Report

Configure Tag Report

The Tag Path accepts (*) as a wildcard. For example, *tag* will return all tags with paths containing tag.

Tag Path

Quality

Types

Traits

Ancestor

Motor

— ✓ Motor

Path Tag Type Quality

Path	Tag Type	Quality
Motor 2	UdtInstance	Good
Motor 1	UdtInstance	Good

Complete - 2 results

This screenshot shows the Ignition Tag Report Tool interface. On the left, there's a configuration panel with sections for 'Tag Path', 'Quality', 'Types', 'Traits', and 'Ancestor'. The 'Ancestor' section is highlighted with a red box and contains a dropdown set to 'Motor' with an option 'Motor' selected. On the right, a results table displays two entries: 'Motor 2' and 'Motor 1', both categorized as 'UdtInstance' with 'Good' quality. A status message at the bottom indicates 'Complete - 2 results'.

UDT Members and Definitions

The following feature is new in Ignition version **8.1.21**
[Click here](#) to check out the other new features

The Tag Report Tool provides the option to include UDT Members or UDT Definitions in the query's search results.

The screenshot shows the 'Tag Report Tool [Sample_Tags]' interface. On the left, there's a configuration panel with fields for 'Tag Path', 'Quality', 'Types', 'Traits', 'Ancestor', and two checkboxes: 'Include UDT Members' (checked) and 'Include UDT Definitions'. A red box highlights the 'Include UDT Members' checkbox. Below these are buttons for '+ Property' and search operators. On the right is a results grid titled 'Path' with columns for 'Path', 'Tag Type', and 'Quality'. The top right corner of the interface has icons for search, settings, and edit.

Properties

Search for Tags by values on specific properties. All tag properties are available to query on using the following operators:

Note:

Properties contained within a set or properties that are a set require wildcards around search terms when using **Like** and **Not Like**. These properties are listed below.

- Alarms (Only returns results when searching for alarm names)
- Parameters
- Qualified Value
- Read Permissions
- Tag Event Scripts
- Value (If the value data type is a set)
- Write Permissions

Operator	Description
Has	Has been configured to a value other than the default value.
Has Not	Has not been configured.
Overridden	Value of this property has been overridden from the parent property.
Like	Is like. This argument accepts a wildcard (*) in the search string.
Not Like	

The following feature is new in Ignition version **8.1.21**
[Click here](#) to check out the other new features

Is not like. This argument accepts a wildcard (*) in the search string.

> Greater than

>= Greater than or equal to

< Less than

<= Less than or equal to

= Equals

!= Does not equal

Tag Report Tool [Sample_Tags]

Load Saved Report

Configure Tag Report

?

The Tag Path accepts (*) as a wildcard. For example, *tag* will return all tags with paths containing tag.

X

Tag Path

Quality

Types

Traits

Ancestor

+ Property

Property 1

Alarms Has <input type="text"/>

Property 2

Value > <input type="text" value="123"/>

Alarm Eval Enabled
Alarms
Clamp Mode
Data Type
Datasource
Deadband
Deadband Mode
Deadband Mode

Search

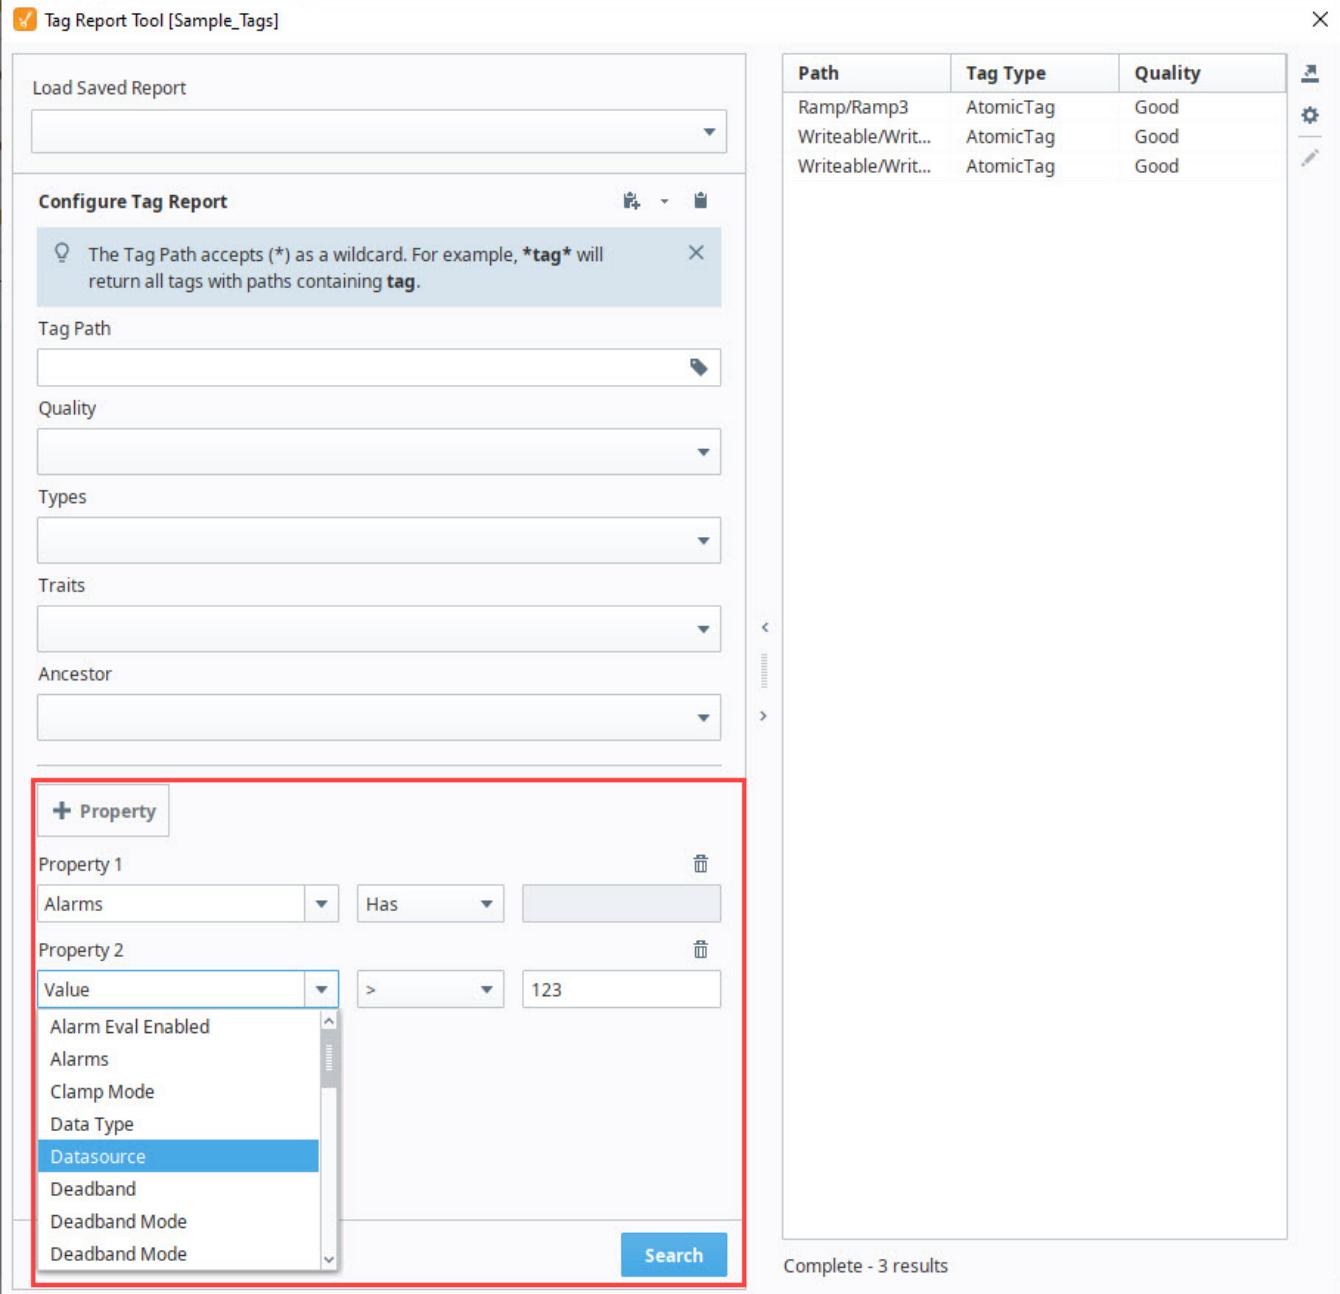
Path Tag Type Quality

Ramp/Ramp3 AtomicTag Good

Writeable/Writ... AtomicTag Good

Writeable/Writ... AtomicTag Good

Complete - 3 results



The following feature is new in Ignition version **8.1.28**
[Click here](#) to check out the other new features

You can search for multiple Tags with different properties using the `Or` operator. Additionally, you can nest multiple `And` operators within an `Or` operator.

The screenshot shows the Ignition search interface. At the top, there is a search bar with dropdown menus for 'Name' and 'Like', and a text input field containing '*tag*'. Below this is a section labeled 'And' with a '+' button. In the middle, there is a section labeled 'Or' with a '+' button. Below the 'Or' section is another section labeled 'And' with a '+' button. At the bottom right is a large blue 'Search' button, and at the bottom left are icons for refresh, save, and cancel.

Note: You will not be able to use the `Or` operator to query a remote Tag Provider that is from an older Gateway, since Gateways before version 8.1.28 will not have `Or` operator functionality.

You can also use the `Or` operator for specific system functions that use JSON queries. The following is a list of system functions that can use `Or`:

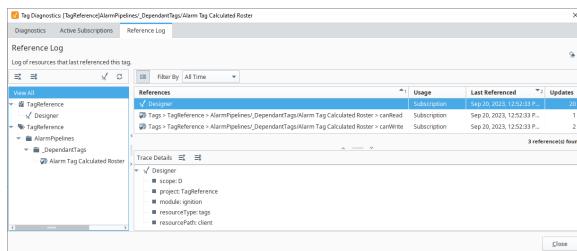
- [system.alarm.queryJournal](#)
- [system.alarm.queryStatus](#)
- [system.tag.query](#)

See the system function pages above on examples of how to use the `Or` operator.

Tag Diagnostics

This feature was changed in Ignition version [8.1.34](#):

In 8.1.34, the Tag Diagnostics popup window was added to maintain quick access to relevant tag information and introduce the Tag Reference Tracker. Note that with this change the Tag Editor no longer displays tag diagnostics information. Pages detailing features of the previous Tag Editor can be found in [Deprecated Ignition Features](#).

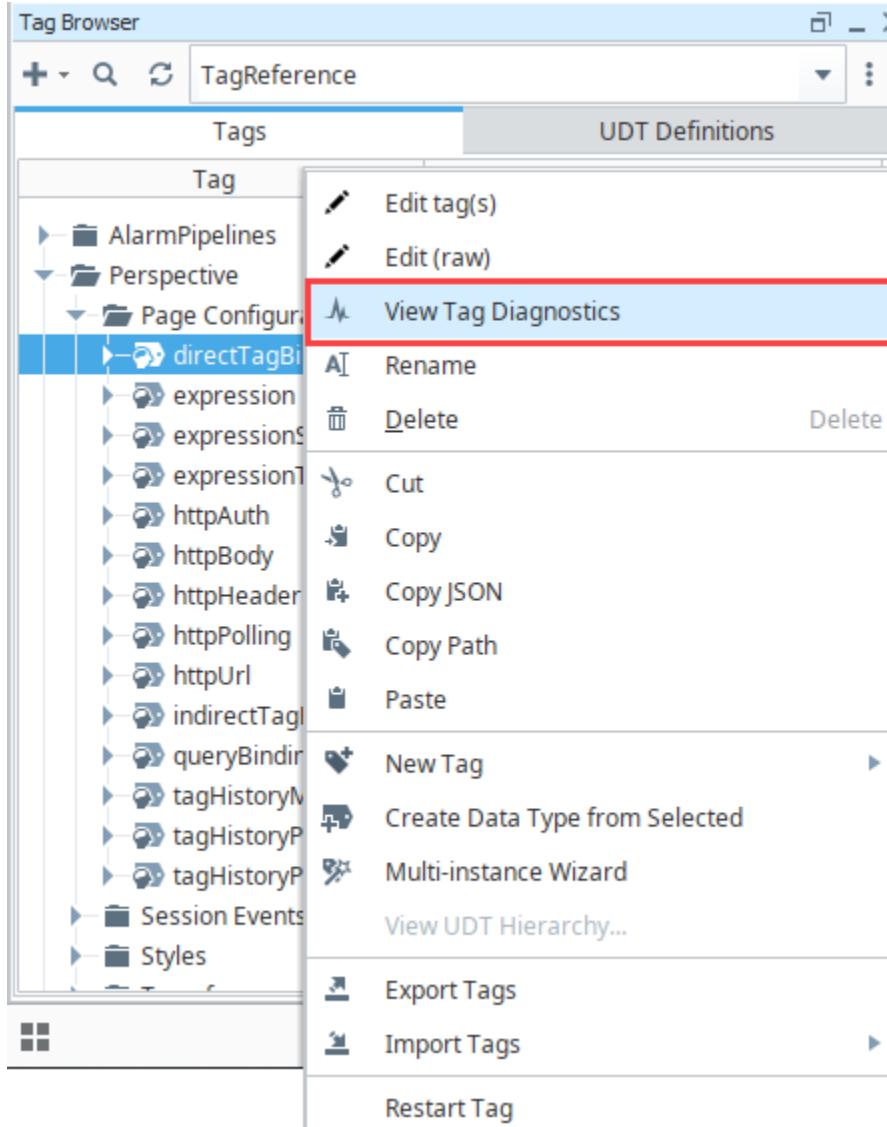


On this page ...

- [Accessing the Tag Diagnostics Window](#)
- [Tag Diagnostic Values](#)
- [Active Subscriptions](#)
- [Tag Reference Tracker](#)
 - [Enabling/Disabling the Tag Reference Store](#)
 - [Collecting Data Backups](#)

Accessing the Tag Diagnostics Window

To access Tag Diagnostics for a specified tag, right-click on the tag in the Tag Browser and select **View Tag Diagnostics**. This will open the Tag Diagnostics window, which contains a Diagnostics, Active Subscriptions, and Reference Logs tabs.



Tag Diagnostic Values

The Diagnostics tab provides detailed metrics about the specified tag including:

- Current value, quality and timestamp information
- Error messages
- Status of dependent datasources

Note: The diagnostic values do not update while the Tag Diagnostics window is open. The displayed information is collected when opened and will remain static until the window is closed and opened again.

Tag Diagnostics: [TagReference]Perspective/Page Configurations/directTagBinding

Diagnostics Active Subscriptions Reference Log

Name	Value
Deadband	
Is Leased	true
Last published value	[directTag, Good, Tue Sep 19 10:39:22 PDT 2023 (1695145162996)]
Tag Definition Valid?	true
Tag id	baf54a67-3d4b-4407-afeb-cbcda3dcc852

Close

Name	Description						
Alarm*	A subcategory for details pertaining to alarms configured on the tag. Each alarm on a tag will receive its own subcategory. <table border="1"><thead><tr><th>Name</th><th>Description</th></tr></thead><tbody><tr><td>Last State</td><td>Displays the last state of the alarm.</td></tr><tr><td>Live Event Count</td><td>The number of active or unacknowledged events for the alarm. The maximum count is determined by the Gateway's Live Event Limit setting.</td></tr></tbody></table>	Name	Description	Last State	Displays the last state of the alarm.	Live Event Count	The number of active or unacknowledged events for the alarm. The maximum count is determined by the Gateway's Live Event Limit setting .
Name	Description						
Last State	Displays the last state of the alarm.						
Live Event Count	The number of active or unacknowledged events for the alarm. The maximum count is determined by the Gateway's Live Event Limit setting .						
Deadband	A subcategory for details pertaining to the deadband settings on the tag. These metrics are only available when the Numeric category's deadband is set to a value larger than 0. <table border="1"><thead><tr><th>Name</th><th>Description</th></tr></thead><tbody><tr><td>Last value</td><td>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features The last value on that tag, as recognized by the Numeric deadband.</td></tr><tr><td>Limit</td><td>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features The Numeric deadband value on the tag.</td></tr></tbody></table>	Name	Description	Last value	The following feature is new in Ignition version 8.1.11 Click here to check out the other new features The last value on that tag, as recognized by the Numeric deadband.	Limit	The following feature is new in Ignition version 8.1.11 Click here to check out the other new features The Numeric deadband value on the tag.
Name	Description						
Last value	The following feature is new in Ignition version 8.1.11 Click here to check out the other new features The last value on that tag, as recognized by the Numeric deadband.						
Limit	The following feature is new in Ignition version 8.1.11 Click here to check out the other new features The Numeric deadband value on the tag.						
History	A subcategory for details pertaining to the history settings on the tag.						

Name	Description														
Deadband (Analog)	<p>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features</p> <p>A subcategory of diagnostic metrics that are only available when the History category's Deadband Style is set to Analog.</p>														
	<table border="1"> <thead> <tr> <th>Name</th><th>Description</th></tr> </thead> <tbody> <tr> <td>Current Lower Slope and Current Upper Slope</td><td> <p>These properties represent the slope for each window.</p> <p>These values are calculated every time the tag evaluates, based upon the tag's sample mode. The resulting numbers are then compared to these current slopes. If the new upper slope is less than the Current Lower Slope, or the new lower slope exceeds the Current Upper Slope, then the new value is stored. In either case, the new slope values replace the value on these Current Slope properties.</p> <p>See Deadband and Analog Compression for more information.</p> </td></tr> <tr> <td>Last Stored Time</td><td> <p>Represents the previously stored timestamp, after the most recent timestamp. Works in a similar manner to Last Stored Value except this stores the time of the storage, instead of the value.</p> </td></tr> <tr> <td>Last Stored Value</td><td> <p>Represents the value stored before the most recently stored value.</p> <p>For example, say a historian system stores the following for a single tag, in order:</p> <ul style="list-style-type: none"> • 1000 • 1100 • 1200 <p>The most recently stored value would be 1200. The value of this Last Stored Value metric would be 1100, since that was the previously stored value.</p> </td></tr> <tr> <td>Last value</td><td> <p>Represents the last QualifiedValue on the tag, regardless of whether or not that value was stored. In other words, the value most recently written to the tag.</p> </td></tr> <tr> <td>Mode</td><td> <p>The Historic category's Deadband Style used by the tag. Useful in cases where a style of "Auto" is used, as this diagnostic metric will display which option was selected by the system.</p> </td></tr> <tr> <td>Tolerance</td><td> <p>The tolerance used to calculate each new slope. The value of this metric is determined by the History category's Deadband setting on the tag.</p> </td></tr> </tbody> </table>	Name	Description	Current Lower Slope and Current Upper Slope	<p>These properties represent the slope for each window.</p> <p>These values are calculated every time the tag evaluates, based upon the tag's sample mode. The resulting numbers are then compared to these current slopes. If the new upper slope is less than the Current Lower Slope, or the new lower slope exceeds the Current Upper Slope, then the new value is stored. In either case, the new slope values replace the value on these Current Slope properties.</p> <p>See Deadband and Analog Compression for more information.</p>	Last Stored Time	<p>Represents the previously stored timestamp, after the most recent timestamp. Works in a similar manner to Last Stored Value except this stores the time of the storage, instead of the value.</p>	Last Stored Value	<p>Represents the value stored before the most recently stored value.</p> <p>For example, say a historian system stores the following for a single tag, in order:</p> <ul style="list-style-type: none"> • 1000 • 1100 • 1200 <p>The most recently stored value would be 1200. The value of this Last Stored Value metric would be 1100, since that was the previously stored value.</p>	Last value	<p>Represents the last QualifiedValue on the tag, regardless of whether or not that value was stored. In other words, the value most recently written to the tag.</p>	Mode	<p>The Historic category's Deadband Style used by the tag. Useful in cases where a style of "Auto" is used, as this diagnostic metric will display which option was selected by the system.</p>	Tolerance	<p>The tolerance used to calculate each new slope. The value of this metric is determined by the History category's Deadband setting on the tag.</p>
Name	Description														
Current Lower Slope and Current Upper Slope	<p>These properties represent the slope for each window.</p> <p>These values are calculated every time the tag evaluates, based upon the tag's sample mode. The resulting numbers are then compared to these current slopes. If the new upper slope is less than the Current Lower Slope, or the new lower slope exceeds the Current Upper Slope, then the new value is stored. In either case, the new slope values replace the value on these Current Slope properties.</p> <p>See Deadband and Analog Compression for more information.</p>														
Last Stored Time	<p>Represents the previously stored timestamp, after the most recent timestamp. Works in a similar manner to Last Stored Value except this stores the time of the storage, instead of the value.</p>														
Last Stored Value	<p>Represents the value stored before the most recently stored value.</p> <p>For example, say a historian system stores the following for a single tag, in order:</p> <ul style="list-style-type: none"> • 1000 • 1100 • 1200 <p>The most recently stored value would be 1200. The value of this Last Stored Value metric would be 1100, since that was the previously stored value.</p>														
Last value	<p>Represents the last QualifiedValue on the tag, regardless of whether or not that value was stored. In other words, the value most recently written to the tag.</p>														
Mode	<p>The Historic category's Deadband Style used by the tag. Useful in cases where a style of "Auto" is used, as this diagnostic metric will display which option was selected by the system.</p>														
Tolerance	<p>The tolerance used to calculate each new slope. The value of this metric is determined by the History category's Deadband setting on the tag.</p>														
Deadband (Discrete)	<p>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features</p> <p>A subcategory of diagnostic metrics that are only available when the History category's Deadband Style is set to Discrete.</p>														
	<table border="1"> <thead> <tr> <th>Name</th><th>Description</th></tr> </thead> <tbody> <tr> <td>Last Value</td><td> <p>The last value on that tag, as recognized by the History category's Deadband property.</p> </td></tr> <tr> <td>Limit</td><td> <p>The Historical Deadband value on the tag.</p> </td></tr> <tr> <td>Mode</td><td> <p>The Historic category's Deadband Style used by the tag. Useful in cases where a style of "Auto" is used, as this diagnostic value will display which option was selected by the system.</p> </td></tr> </tbody> </table>	Name	Description	Last Value	<p>The last value on that tag, as recognized by the History category's Deadband property.</p>	Limit	<p>The Historical Deadband value on the tag.</p>	Mode	<p>The Historic category's Deadband Style used by the tag. Useful in cases where a style of "Auto" is used, as this diagnostic value will display which option was selected by the system.</p>						
Name	Description														
Last Value	<p>The last value on that tag, as recognized by the History category's Deadband property.</p>														
Limit	<p>The Historical Deadband value on the tag.</p>														
Mode	<p>The Historic category's Deadband Style used by the tag. Useful in cases where a style of "Auto" is used, as this diagnostic value will display which option was selected by the system.</p>														
Min /Max Age	<p>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features</p>														
	<table border="1"> <thead> <tr> <th>Name</th><th>Description</th></tr> </thead> <tbody> <tr> <td>Max Age Limit (ms)</td><td> <p>A milliseconds representation of the Max Time Between Samples and Max Time Units properties.</p> </td></tr> </tbody> </table>	Name	Description	Max Age Limit (ms)	<p>A milliseconds representation of the Max Time Between Samples and Max Time Units properties.</p>										
Name	Description														
Max Age Limit (ms)	<p>A milliseconds representation of the Max Time Between Samples and Max Time Units properties.</p>														

		<table border="1"> <tr> <td>Minimum Age Limit (ms)</td><td>A milliseconds representation of the Min Time Between Samples and Min Time Units properties.</td></tr> <tr> <td>Pending Value</td><td>Represents one or more pending values that can't yet be written to the tag due to the Minimum Age Limit.</td></tr> <tr> <td>System Time of Last Storage</td><td>A date representing the last time a historical record value was stored for this tag.</td></tr> <tr> <td>Timestamp of Last Stored Value</td><td>A date representing the timestamp associated with the last collected historical record.</td></tr> </table>	Minimum Age Limit (ms)	A milliseconds representation of the Min Time Between Samples and Min Time Units properties.	Pending Value	Represents one or more pending values that can't yet be written to the tag due to the Minimum Age Limit .	System Time of Last Storage	A date representing the last time a historical record value was stored for this tag.	Timestamp of Last Stored Value	A date representing the timestamp associated with the last collected historical record.		
Minimum Age Limit (ms)	A milliseconds representation of the Min Time Between Samples and Min Time Units properties.											
Pending Value	Represents one or more pending values that can't yet be written to the tag due to the Minimum Age Limit .											
System Time of Last Storage	A date representing the last time a historical record value was stored for this tag.											
Timestamp of Last Stored Value	A date representing the timestamp associated with the last collected historical record.											
	Stored Value Count (since startup)	<p>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features</p>										
		The total number of historical records collected for the tag since history on the tag was enabled.										
OPC	A subcategory for details pertaining to OPC details on the tag. These diagnostic values are only available when Value Source is set to OPC .											
	<table border="1"> <thead> <tr> <th>Name</th><th>Description</th></tr> </thead> <tbody> <tr> <td>Last Subscription Value</td><td>A qualified value representing the last subscribed value. The value on this setting ignores the deadband.</td></tr> <tr> <td>Server Name</td><td>The name of the OPC server the tag is subscribing to.</td></tr> <tr> <td>Subscription Name</td><td>The name of the subscription, used by the OPC server.</td></tr> <tr> <td>Valid Tag Group?</td><td>States whether or not the tag group the tag is configured to use is valid.</td></tr> </tbody> </table>		Name	Description	Last Subscription Value	A qualified value representing the last subscribed value. The value on this setting ignores the deadband.	Server Name	The name of the OPC server the tag is subscribing to.	Subscription Name	The name of the subscription, used by the OPC server.	Valid Tag Group?	States whether or not the tag group the tag is configured to use is valid.
Name	Description											
Last Subscription Value	A qualified value representing the last subscribed value. The value on this setting ignores the deadband.											
Server Name	The name of the OPC server the tag is subscribing to.											
Subscription Name	The name of the subscription, used by the OPC server.											
Valid Tag Group?	States whether or not the tag group the tag is configured to use is valid.											
Alarm Acknowledged	Displays any errors from the Alarm Acknowledged tag event script. Only visible if a script is defined on the corresponding tag event .											
Alarm Active	Displays any errors from the Alarm Active tag event script. Only visible if a script is defined on the corresponding tag event .											
Alarm Cleared	Displays any errors from the Alarm Cleared tag event script. Only visible if a script is defined on the corresponding tag event .											
Datasource Status	States which datasource is configured for the tag. Only available when Value Source is set to Query .											
Is Leased	If true, indicates that the tag is using a leased tag group , and is currently executing at the leased rate.											
Is Scaled	If true, indicates that Scale Mode has been enabled.											
Last published value	Represents the last value on the tag when the diagnostics tab was last updated.											
Range	<p>The following feature is new in Ignition version 8.1.11 Click here to check out the other new features</p> <p>When Deadband Mode is set to Percent, indicates the range of values between the Engineering Low Limit and Engineering High Limit.</p>											
Quality Changed	Displays any errors from the Quality Changed tag event script. Only visible if a script is defined on the corresponding tag event .											
Quality Diagnostics	Reports any quality errors about the tag. This is typically only present when there is a configuration issue with the tag, such as when Value Source is set to Query and a Datasource value has not been set.											
Tag Definition Valid?	A flag that indicates whether or not a tag configuration is valid.											

Tag id	A unique identifier for the tag. Used by the tag system to distinguish one tag from another.
Value Changed	Displays any errors from the Value Changed tag event script. Only visible if a script is defined on the corresponding tag event.

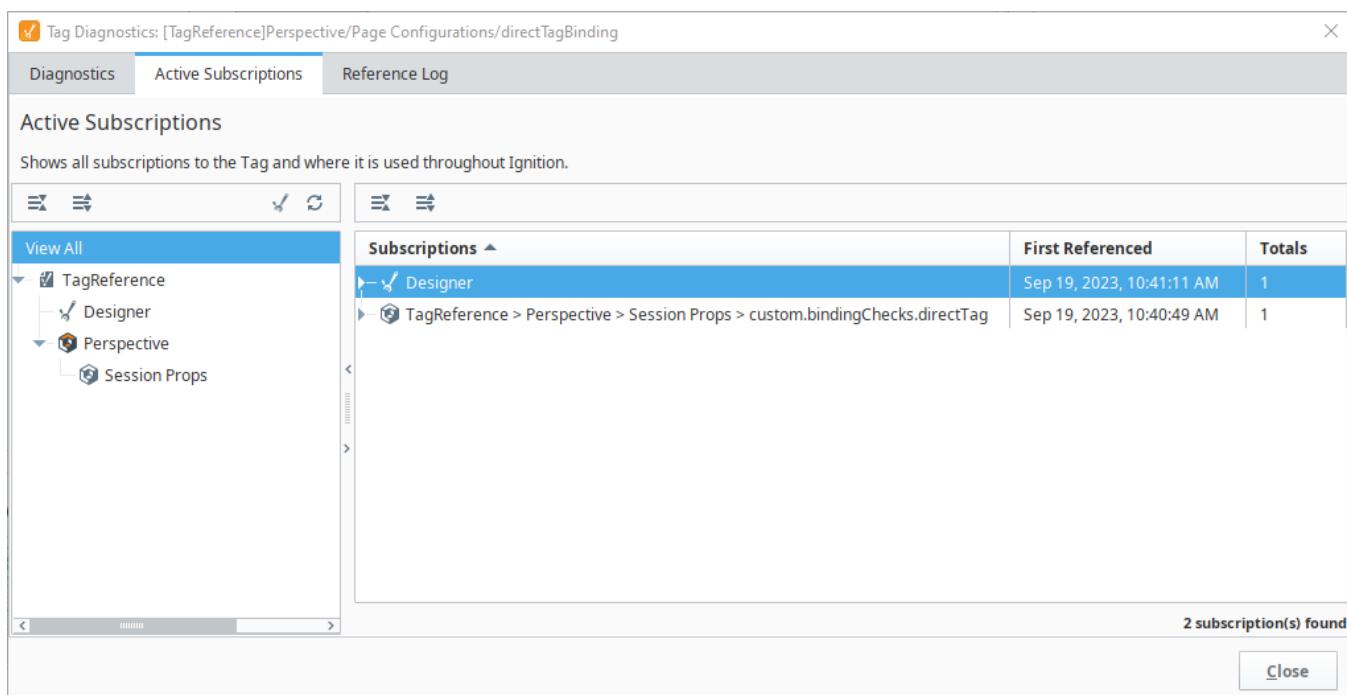
Active Subscriptions

The Active Subscriptions tab displays real-time tracking of a tag subscription by sharing the following information:

- **Subscriptions:** Displays the path to a resource that is currently subscribed to the selected tag.
- **First Referenced:** Displays the first logged time a tag subscription at the given location was created. The logged time is reflective of the first subscription for the current grouping of a specific subscription.
- **Totals:** Displays the number of currently active subscriptions from the referenced location. For example, if there are multiple Perspective Sessions with the same View bound to the selected tag loaded, this number would increase.

Additionally, the Reference Browser on the left-hand side lists all locations where the tag has an active subscription. Use the Collapse All  and

Expand All  icons to adjust how much information is displayed. You can also update the active subscription information by clicking the Refresh  icon. The Designer  icon is available to show/hide references to the specified tag from the Designer, including in the tag browser and Vision components when in Preview mode.



The screenshot shows the Tag Diagnostics dialog with the "Active Subscriptions" tab selected. The left pane shows a tree view under "View All" with nodes for "TagReference", "Designer", and "Perspective". The "Perspective" node has a "Session Props" child. The right pane displays a table titled "Subscriptions" with two entries:

Subscriptions	First Referenced	Totals
Designer	Sep 19, 2023, 10:41:11 AM	1
TagReference > Perspective > Session Props > custom.bindingChecks.directTag	Sep 19, 2023, 10:40:49 AM	1

At the bottom right of the dialog, it says "2 subscription(s) found" and has a "Close" button.

Tag Reference Tracker

The following feature is new in Ignition version **8.1.34**
[Click here](#) to check out the other new features

The Reference Log tab shows the data stored in the Tag Reference Tracker Store as a way to find where tags are being used in a project. The Tag Reference Tracker Store provides visibility to the areas that rely on any given tag as it displays what resources have subscribed, read, written to or configured specified tags.



The Reference Log can only show references that are in use or have been in use since the reference tracking began. It will miss potential references that haven't been active since tracking started. The longer it is in use, the more reliable the tracker becomes as the history will grow and continue to log with each new session that is opened.



Tag Diagnostics - Tag Reference Tracker

[Watch the Video](#)

The bulk of the Tag Reference Tracker information includes:

- **References:** The tag path to display where the tag was used.
- **Usage:** The type of reference to this tag. The usage type listed will match the tracking type recorded in the Ignition database file. The four tracking types for tag usage include reads, writes, subscriptions, and configuration changes.
 - **Configuration:** Resource made a configuration change to the tag definition.
 - **Subscription:** Resource subscribed to the tag to display information on a particular property.
 - **Read:** A read event occurred to a property on the tag.
 - **Write:** A write event occurred to a property on the tag.
- **Last Referenced:** The timestamp for when the reference was last updated.
- **Updates:** The number of times a reference with a specific usage type occurred.

Note: Double-clicking on a reference line will open the tag reference root location or the Tag Editor depending on the reference type. If it is in a Tag Value Change Script or an Expression used within the Tag definition, it will open the Tag Editor.

Like the Active Subscriptions tab, the Reference Browser on the left-hand side will show where the tag is used in any project. You can use the

reference tree or Collapse All and Expand All icons to locate where in the Designer a tag is used for each project, the Designer icon to show/hide Designer references, and the Refresh icon to refresh the browser list. Refer to the Trace Details section at the bottom for more information that might be helpful in understanding tag usage. Raw tracking logs can also be found by right-clicking on the Reference Log tab and selecting **Show Raw**. The Raw Tag Logs Browser will now be visible on the Tag Diagnostics window. Clicking the Export icon to download the raw logs.

References	Usage	Last Referenced	Updates
Designer	Subscription	Sep 19, 2023, 10:41:11 AM	1
TagReference > Perspective > Session Props > custom.bindingChecks.directTag	Subscription	Sep 19, 2023, 10:40:49 AM	1

Enabling/Disabling the Tag Reference Store

The Tag Reference Store is enabled by default on new installations and disabled by default on upgrades from Ignition versions 8.1.33 and earlier. There are two options to enable or disable the Tag Reference Store depending on the provider type.

- For Realtime Tag Providers, standard and remote: Navigate to the **Config > Tags > Realtime** on the Gateway to access the **Realtime Tag Providers** page. Select **edit** for the provider you want to enable/disable and expand the **Show advanced properties** section. If the **Enable Tag Reference Tracker Store** property is checked, it is enabled.

Config > Tags > Realtime Tag Providers

Any actor attempting to write to any tag within this tag provider must have these permissions in addition to the target tag's permissions. Multiple security level paths can be specified by separating them with commas. If blank, tag writes for this provider will not be restricted by default. Example: Authenticated/Roles/Administrator,SecurityZones/localhost

Tag Editing Permissions

- Users must belong to all of these security levels in order to edit this provider's tags.
- Users must belong to at least one of these security levels in order to edit this provider's tags.

Any actor attempting to edit any tag within this tag provider must have these permissions. Multiple security level paths can be specified by separating them with commas. If blank, tag editing for this provider will not be restricted. Example: Authenticated/Roles/Administrator,SecurityZones/localhost

Show advanced properties

Advanced

Allow Back-fill Data	<input type="checkbox"/> If enabled, data will be allowed to arrive out of order from the source. Data from the past will be stored to history, but will not be used for alarms, scripts, or subscriptions. If false (default behavior), each value will be processed fully as it arrives. (default: false)
Enable Tag Reference Tracker Store	<input checked="" type="checkbox"/> Enables the storing of Tag Reference entries to a database on the local Gateway for analysis in a Designer (default: true)

Create New Realtime Tag Provider

- For managed providers, like third party modules and the system provider: Navigate to the **Config > System > Gateway Settings** on the Gateway to access the **System Settings** page. If the **Enable Tag Reference Tracker Store** property is checked, it is enabled. Be prepared that this setting impacts all managed providers throughout the Gateway, so a large number of third party modules can result in large data collections.

Config > System

Gateway Settings

System Name	TagReferenceTracker
The name of this Ignition system, used to differentiate this system from others in a larger architecture. (default:)	
Persist Alarms	<input checked="" type="checkbox"/> Whether or not alarm properties such as acknowledgment should be persisted across Gateway restarts. (default: true)
Homepage Redirect URL	/web/home
The URL this gateway will redirect to when http://ipport/ is visited. Can either be a relative path (e.g., /web/home), or fully qualified (e.g., https://inductiveautomation.com). (default: /web/home)	
Gateway Scripting Project	TagReference
The name of the Project that gateway-scoped scripts with no project affiliation can access user script libraries in. (default:)	
Enable Tag Reference Tracker Store	<input type="checkbox"/> Enables the storing of Tag Reference entries to a database on the local Gateway for analysis in a Designer for 3rd party modules. (default: true)

Note: Disabling the Tag Reference Tracker will not impact the Diagnostics or Active Subscriptions tabs, but the Reference Log will no longer track tag information.

Collecting Data Backups

Gateway backups do not include Tag Reference Store data. This information is saved instead in a database that creates a new Diagnostics folder within the Data directory of an Ignition installation. Make sure to collect and add the reference data from this folder to any backups so no information is lost.

The default folder location for each operating system are listed below:

Windows

```
C:\Program Files\Inductive Automation\Ignition\data\diagnostics\tags
```

MacOS

```
/usr/local/ignition/data/diagnostics/tags
```

Linux

```
/usr/local/bin/ignition/data/diagnostics/tags
```

Alarming

Alarms are always configured on a Tag, excluding Vision client tags. Every time a Tag with a configured alarm receives an updated value, the configured alarm will examine the new value. If the new value on the Tag meets some criteria, as defined by the **alarms mode settings**, then the alarm generates a new **alarm event**.

Alarms are monitored by the gateway, and are considered gateway scoped. However, some module features and functions can be used to [monitor the live status of alarm events](#).

Alarm Events

Each alarm event contains several attributes that store metadata about the source Tag at the time the alarm event was created. Alarm events are effectively a "snapshot" of several parameters when the alarm went active. Tags can have multiple alarms configured on them, meaning that it's possible for multiple alarm events to be created from a single value change on a Tag.

Alarm Events and Their States

The state of an alarm event is determined by two conditions

Active Condition

Determines if the alarm is active, or "live".

All alarm events start with an Active condition, and then move to a Cleared condition. Alarm events never transition from Cleared to Active. Instead, a new alarm event would be created should the value on a tag meet the setpoint criteria.

- **Active** - The alarm event is still active. Meaning, the value on the source tag still meets the setpoint configured on the alarm. If the alarm is used to signal that there's a problem, then the active condition means the problem is still active.
- **Cleared** - The alarm event is no longer active. The value on the source tag no longer meets the setpoint configured on the alarm.

Acknowledged Condition

Alarms feature an Acknowledgement functionality. Acknowledgement is simply a flag that can be marked on an alarm event. Commonly the acknowledged condition is used as a way for users to signal to other users that an alarm is being handled.

Imagine a scenario where multiple operators are monitoring any alarm events that occur. Acknowledgement allows an operator to "claim" an alarm, signaling to the others that the event is being handled by someone.

Alarm event acknowledgement is generally handled by either the Vision or Perspective Alarm Status Table components, or the [system.alarm.acknowledge](#) function.

- **Unacknowledged (Unacked)** - The alarm event has not yet been flagged as "acknowledged".
- **Acknowledged (Acked)** - The alarm event has been flagged as "acknowledged".

Alarm Event States

The two conditions above are combined form the state of an event. The four possible states are:

- Active and Unacknowledged
- Active and Acknowledged
- Cleared and Unacknowledged
- Cleared and Acknowledged

Alarm Shelving

On this page ...

- [Alarm Events](#)
 - [Alarm Events and Their States](#)
 - [Alarm Shelving](#)
 - [Alarms in Transaction Groups](#)
 - [Configuring Alarms](#)
 - [Monitoring Alarm Status](#)
 - [Alarm Count Tags](#)
 - [Viewing Alarm History](#)
 - [Alarm Notifications](#)
 - [On-Call Rosters](#)
 - [Alarm Notification Pipelines](#)

Shelving an alarm tells the alarm system that it should ignore the alarm, preventing new alarm events from being created. Shelving always has a duration associated with it. At the end of the shelving period, the alarm will evaluate its source tag, potentially creating a new alarm event. Shelving is useful to temporarily ignore an alarm for a period of time, such as maintenance periods.

Shelved alarms will not generate alarm notifications. In addition, shelving an alarm will hide all alarm event for that alarm from the Vision and Perspective Alarm Status Table components.

Alarms in Transaction Groups

OPC items in Transaction groups can also have alarms configured on them, similar to tags. This functionality exists mostly for legacy reasons, and offers no real benefit over configuring alarms on a Tag instead.

Configuring Alarms

Alarm configuration in Ignition is flexible and highly customizable to your needs. You can configure alarms with one alarm on a Tag or multiple alarms on a Tag. You can add alarms in UDTs so every instance of that Tag will automatically have alarms configured when a new instance of your Tag is created. You can use the alarm settings to create alarms that equal or don't equal a setpoint, above or below a setpoint, between setpoints, outside setpoints, dynamic setpoints, out of range, bad quality, etc. Alarms can be configured for any alarm condition imaginable.

More information about configuring alarms can be found in [Alarming Properties](#) and [Configuring Alarms](#) pages.

The image below shows an alarm configured on an OPC Tag. You can see that an alarm has quite a few properties including alarm mode settings where you can set specific alarm attribute values.



Ramp8 > Alarms

OPC Alarm - Outside Setpoints, Critical

< OPC Alarm

Main	
Name	OPC Alarm
Enabled	true
Priority	Critical
Timestamp Source	System
Label	
Display Path	
Ack Mode	Manual
Notes	
Ack Notes Required	false
Shelving Allowed	true
Alarm Mode Settings	
Mode	Outside Setpoints
Low Setpoint	150
Low Inclusive	true
High Setpoint	300
High Inclusive	true
Any Change	false
Deadbands and Time Delays	
Deadband	0
Deadband Mode	Absolute
Active delay (seconds)	0
Clear delay (seconds)	0
Notification	
Ack Pipeline	
Active Pipeline	
Clear Pipeline	
Email Notification Properties	
Custom Message	
Custom Subject	
SMS Notification Properties	
Custom Message	

Monitoring Alarm Status

Alarms can be monitored from [Vision](#) and [Perspective](#) using their respective Alarm Status Table components. In addition, the [system.alarm.queryStatus](#) function can be used to retrieve realtime alarm status.

Active Time	Display Path	Priority	Current State
<input checked="" type="checkbox"/> 5/30/19, 11:23 AM	Motors/Motor 4/Amps/Low Amps	Critical	Active, Unacknowledged
<input type="checkbox"/> 5/30/19, 11:25 AM	Motors/Motor 2/Amps/Low Amps	Critical	Active, Unacknowledged
<input type="checkbox"/> 5/30/19, 11:26 AM	Ramp/Ramp8/OPC Alarm	Critical	Active, Unacknowledged
<input type="checkbox"/> 5/30/19, 10:33 AM	Turbine Number 200 located at Livermore, CA	High	Active, Unacknowledged
<input type="checkbox"/> 5/30/19, 10:33 AM	Turbine Number 300 located at Fresno	High	Active, Unacknowledged
<input type="checkbox"/> 5/30/19, 11:13 AM	Motors/Motor 2/Amps/Low Amps	Critical	Cleared, Unacknowledged
<input type="checkbox"/> 5/30/19, 11:14 AM	Motors/Motor 2/Amps/Low Amps	Critical	Cleared, Unacknowledged

Details Notes

Config Properties

- + **On Active**
- mode: Below Setpoint
- setpointA: 30
- Event Value: 28
- name: Low Amps
- Event Time: 5/30/19, 11:23 AM
- priority: Critical

Acknowledge Shelve

Alarm Count Tags

Ignition provides a set of [System Tags](#) to view information about the Ignition server which includes four Tags that count the number of alarms in each state. A quick way to see if any alarms are currently active and get an alarm count is to add a Label component on the Navigation window. The four system alarm states are:

- Active and Unacknowledged
- Active and Acknowledged
- Clear and Acknowledged
- Clear and Unacknowledged

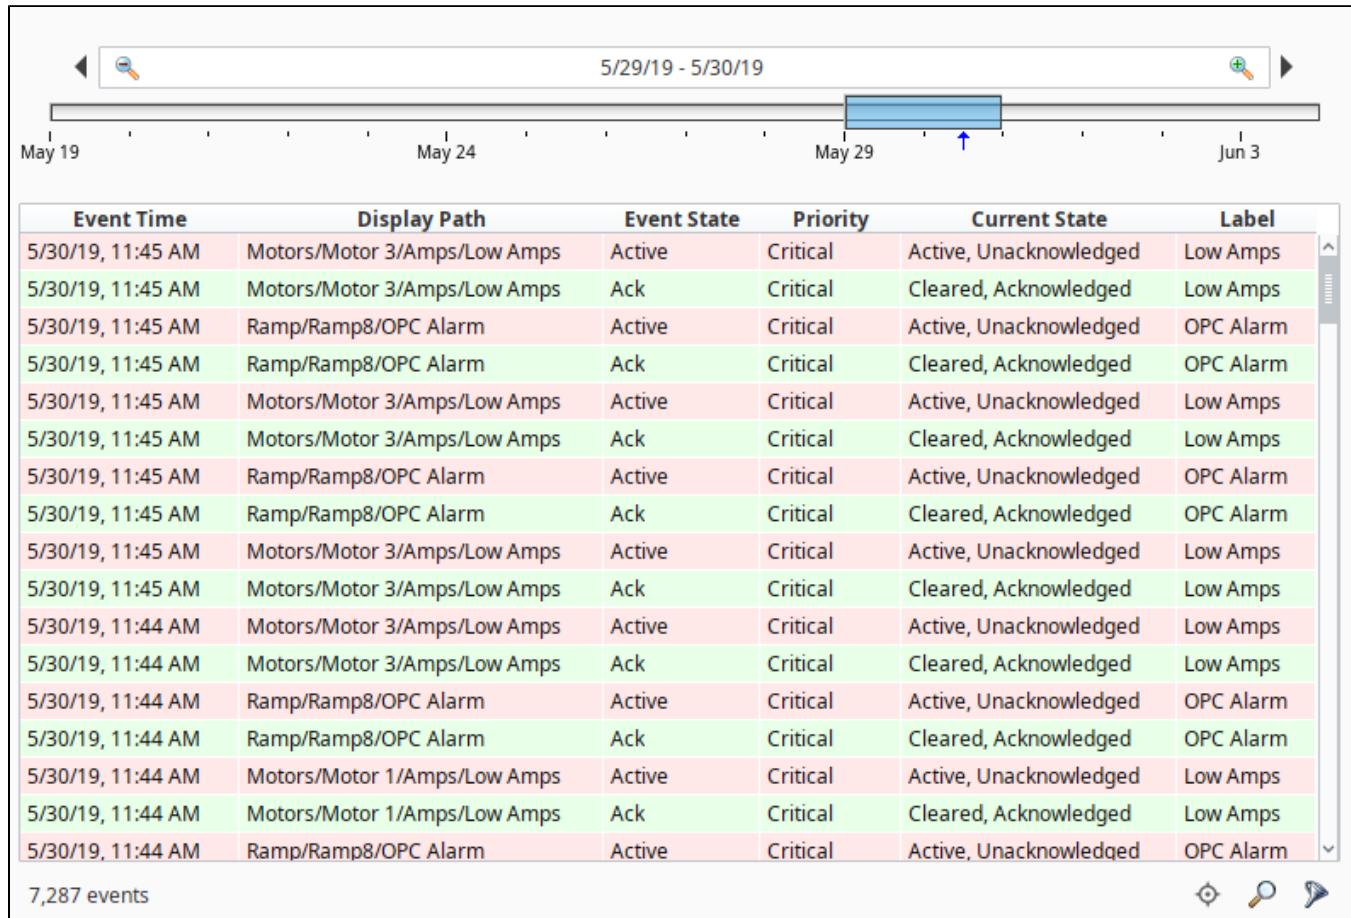
In the **Tag Browser** of the Designer, scroll down to the **System > Gateway > Alarming** folder. These system Tags can easily be used to visualize all alarms in the system.

Tag	Value	Data Type	Traits
Active and Acked	3	Integer	
Active and Unacked	4	Integer	
Clear and Acked	6	Integer	
Clear and Unacked	40	Integer	

Viewing Alarm History

The [Alarm Journal](#) stores historical information about alarms in a database. It stores basic data about alarms that have occurred, such as their source and timestamp, associated data on an alarm, and the values of an alarm's properties at the time the event occurred. It captures all status changes for each alarm to an external SQL database of your choosing. To begin viewing alarm history, all you need to do is create an [Alarm Journal Profile](#) in the Gateway webpage.

Like the Alarm Status Table, both [Vision](#) and [Perspective](#) have components that can be used to view entries in an Alarm Journal.



The screenshot shows the Alarm Journal component in Ignition. At the top, there is a search bar with a magnifying glass icon and a date range selector showing "5/29/19 - 5/30/19". Below the search bar is a horizontal timeline with major ticks for May 19, May 24, May 29, and Jun 3. A blue rectangular selection box highlights the period from May 29 to May 30. To the right of the timeline is a small upward-pointing arrow. The main area contains a table with 15 rows of alarm data. The columns are labeled: Event Time, Display Path, Event State, Priority, Current State, and Label. The data shows various alarms for different paths like Motors/Motor 3/Amps/Low Amps and Ramp/Ramp8/OPC Alarm, with states ranging from Active to Cleared, Acknowledged, and priorities from Critical to None. The table has a light gray background with alternating row colors. At the bottom left of the table area, it says "7,287 events". On the far right of the table, there are three small icons: a magnifying glass, a wrench, and a gear.

Event Time	Display Path	Event State	Priority	Current State	Label
5/30/19, 11:45 AM	Motors/Motor 3/Amps/Low Amps	Active	Critical	Active, Unacknowledged	Low Amps
5/30/19, 11:45 AM	Motors/Motor 3/Amps/Low Amps	Ack	Critical	Cleared, Acknowledged	Low Amps
5/30/19, 11:45 AM	Ramp/Ramp8/OPC Alarm	Active	Critical	Active, Unacknowledged	OPC Alarm
5/30/19, 11:45 AM	Ramp/Ramp8/OPC Alarm	Ack	Critical	Cleared, Acknowledged	OPC Alarm
5/30/19, 11:45 AM	Motors/Motor 3/Amps/Low Amps	Active	Critical	Active, Unacknowledged	Low Amps
5/30/19, 11:45 AM	Motors/Motor 3/Amps/Low Amps	Ack	Critical	Cleared, Acknowledged	Low Amps
5/30/19, 11:45 AM	Ramp/Ramp8/OPC Alarm	Active	Critical	Active, Unacknowledged	OPC Alarm
5/30/19, 11:45 AM	Ramp/Ramp8/OPC Alarm	Ack	Critical	Cleared, Acknowledged	OPC Alarm
5/30/19, 11:45 AM	Motors/Motor 3/Amps/Low Amps	Active	Critical	Active, Unacknowledged	Low Amps
5/30/19, 11:45 AM	Motors/Motor 3/Amps/Low Amps	Ack	Critical	Cleared, Acknowledged	Low Amps
5/30/19, 11:44 AM	Motors/Motor 3/Amps/Low Amps	Active	Critical	Active, Unacknowledged	Low Amps
5/30/19, 11:44 AM	Motors/Motor 3/Amps/Low Amps	Ack	Critical	Cleared, Acknowledged	Low Amps
5/30/19, 11:44 AM	Ramp/Ramp8/OPC Alarm	Active	Critical	Active, Unacknowledged	OPC Alarm
5/30/19, 11:44 AM	Ramp/Ramp8/OPC Alarm	Ack	Critical	Cleared, Acknowledged	OPC Alarm
5/30/19, 11:44 AM	Motors/Motor 1/Amps/Low Amps	Active	Critical	Active, Unacknowledged	Low Amps
5/30/19, 11:44 AM	Motors/Motor 1/Amps/Low Amps	Ack	Critical	Cleared, Acknowledged	Low Amps
5/30/19, 11:44 AM	Ramp/Ramp8/OPC Alarm	Active	Critical	Active, Unacknowledged	OPC Alarm

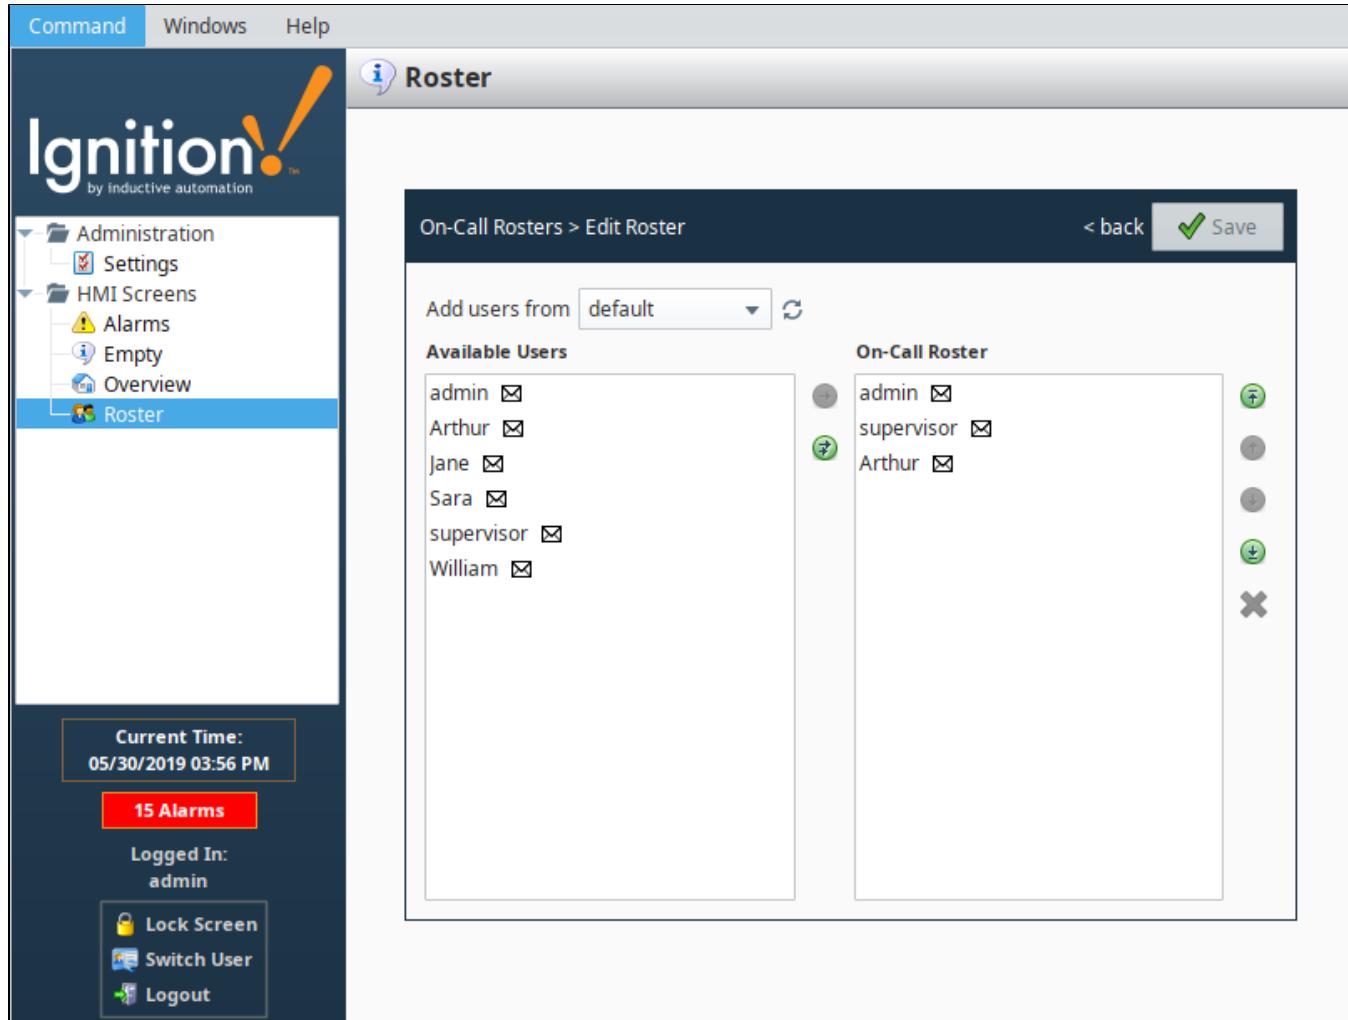
Alarm Notifications

Alarms can also generate notifications that are delivered to users allowing Ignition to immediately communicate events and problems to your users. Alarm Notification Pipelines control how and when notifications are sent to users. You can select the delivery channel for how alarms are sent: [Email](#), [SMS](#), or [Voice](#). The notification system has access to Ignition's Authentication Profiles so users can easily be added to notification [On-Call Rosters](#). [Schedules](#) can be created allowing users to receive notifications only when on-schedule, so there is no need to worry about notifying a supervisor or manager when they are not on-site, or if it's in the middle of the night. Pipelines coupled with on-call rosters and schedules allow you to build your own custom alarm notification process.

On-Call Rosters

The [On-Call Roster](#) is a collection of users that are notified when an alarm occurs. When an alarm is triggered, a notification is sent to a designated On-Call Roster where it evaluates the users schedules, and only notifies those users that have an active schedule.

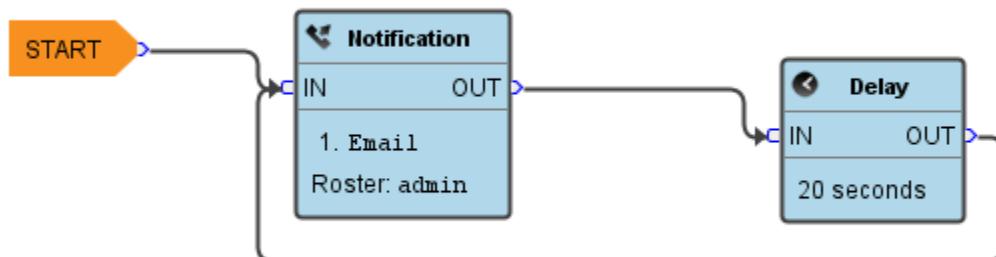
Roster Management from the Vision Client Window



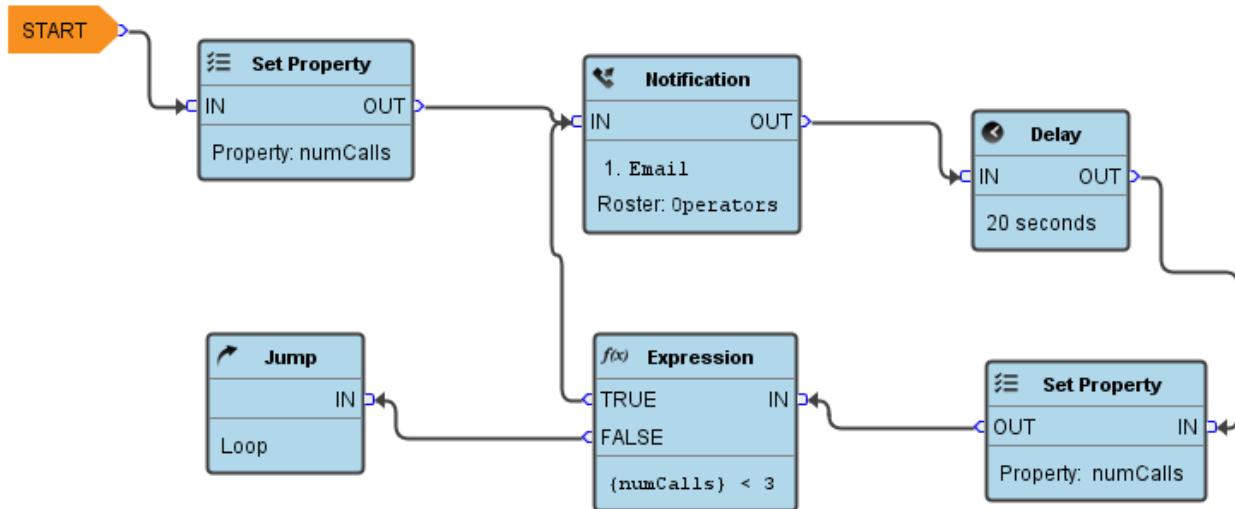
Alarm Notification Pipelines

The [Alarm Notification Pipeline](#) feature is an innovative tool that lets you easily create routes for your alarms. By designing your own alarm notification routing, you have control of what happens when an alarm goes active, when an alarm is sent out, and who receives the alarm notification.

The alarm notification pipeline has a simple drag-and-drop interface so you can build various types of alarm logic. Alarm notification pipelines can be very simple to very complex. In this simple notification pipeline, when an alarm is triggered, the people listed in the [On-call Roster](#) are notified via Email. If no one acknowledges the alarm in 20 seconds, the alarm notification is routed back to the same users listed in the On-Call Roster.



In this more complex alarm notification pipeline, if an operator doesn't respond to the alarm after three attempts, the pipeline jumps it to another pipeline (possibly an escalation pipeline).



To learn more about building your own pipelines, go to [Alarm Notification Pipelines](#).

Alarm Journal

By default, current alarm data is only stored in memory, and a finite number of events are retained for each alarm. Fortunately, Ignition can easily be configured to store alarm data into a SQL database with an Alarm Journal Profile. The journal can store basic data about alarms that have occurred, such as their source and timestamp, associated data on the alarm, and the values of the alarm's properties at the time the event occurred. The Alarm Journal is used by the **Alarm Journal Table** component, and can be accessed through scripting functions and direct database queries.

The Gateway can have more than one Alarm Journal. Alarm Journals have options to filter which Alarms are stored in the journal, therefore, by having more than one alarm journal configured on the Gateway, it is possible to store some alarms in one journal, and different alarms in another journal. Once configured, the journal can be accessed by the **Alarm Journal Table** component, scripting functions, or direct database queries.

Alarm Journals can store data in one of three ways, and store data indefinitely unless a Data Pruning value is set:

- In a database, using an existing database connection from the Gateway
- Remotely, using another Ignition Gateway's Alarm Journal profile
- Internally, storing alarm information into the Ignition install directory

It is strongly encouraged to set a Data Pruning value for Internal Alarm Journal Profiles. Otherwise, it could cause your computer to run out of hard drive space.

Note: You must have an Alarm Journal Profile created and have a valid [database connection](#) to use the Alarm Journal Table.

Creating an Alarm Journal Profile

In Ignition, an Alarm Journal stores historical information about alarms in a database. It can store basic data about alarms that have occurred, such as source and timestamp, along with associated data on the alarm, and the values of the alarm's properties at the time the event occurred. You can choose to store your alarm data on an external database or automatically send it to a remote gateway's Alarm Journal to be logged. Both options are described in the sections below.

Create a single Alarm Journal Profile to store all of your alarms, or create multiple journals to store alarms across multiple databases. Each journal stores alarms based on the filters you set up and can prune data automatically after a set time limit.

Create an Alarm Journal to Log Events to an External Database

1. Go to the **Config** section of the Gateway webpage.
2. Choose **Alarming > Journal** from the menu on the left.
3. Look for the blue arrow and click on **Create new Alarm Journal Profile....** The Alarm Journal Profiles screen will be displayed.

On this page ...

- [Creating an Alarm Journal Profile](#)
 - [Create an Alarm Journal to Log Events to an External Database](#)
- [Remote Alarm Journal Profile](#)
 - [Creating a Remote Alarm Journal Profile](#)
 - [Create an Internal Alarm Journal to Log Events Locally](#)
- [Alarm Journal Component](#)
- [Journal Properties](#)
 - [Table Definitions](#)

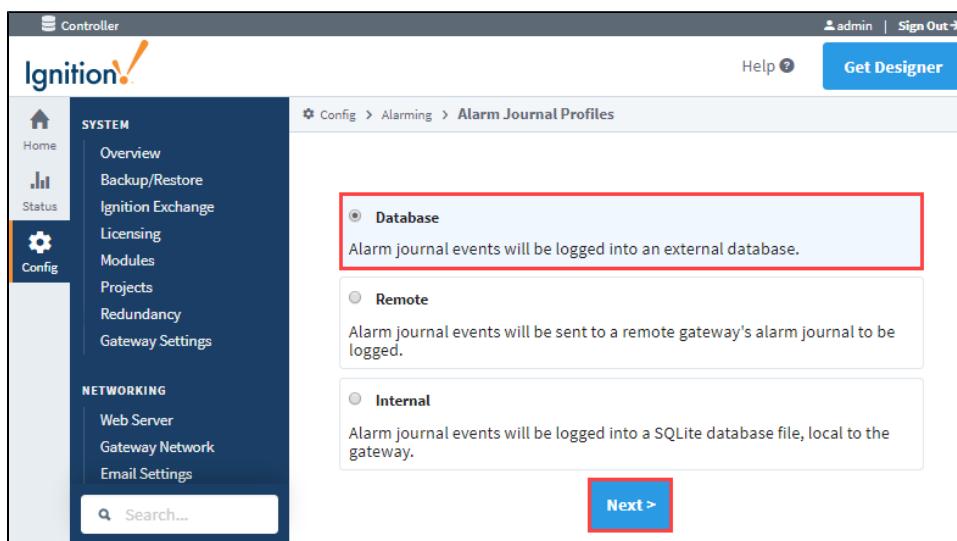


INDUCTIVE
UNIVERSITY

Create Alarm Journal Profile

[Watch the Video](#)

4. You have the option of logging alarm journal events to an external database, logging locally, or sending them to a remote gateway's Alarm Journal. In this example, select **Database**, and click **Next**.



5. Enter the **Name** of your alarm journal profile. The default name is 'Journal.' Most of the fields have default settings. Refer to the [journal properties table](#) below for setting descriptions, and update as necessary. Click the **Create New Alarm Journal Profile** button at the bottom of the page. Once completed, the tables will be created for you once an alarm event occurs.

Config > Alarming > Alarm Journal Profiles

Main	
Name	Packaging_Area
Description	
Enabled	<input checked="" type="checkbox"/> (default: true)

Main	
Datasource	MySQL
	Events will be stored to this datasource.
Minimum Priority	Low
	Only events equal to or greater than the specified priority will be stored. (default: Low)
Store Shelved Events	<input type="checkbox"/> If enabled, events generated by "shelved" alarms will still be written to the journal system. (default: false)
Use Store and Forward	<input checked="" type="checkbox"/> If enabled, journaled events will be stored through the store and forward system. If not enabled, they will be stored directly against the database. (default: true)

Note: If you only have one alarm journal specified on your Gateway, then you do not need to specify the journal name on the Journal Name property. Ignition will set this for you. If you have more than one alarm journal created, then you need to provide the name of the journal you'd like to query in the Journal Name component property of the Property Editor.

Remote Alarm Journal Profile

Utilizing the gateway network, remote alarm journal profiles allow one Ignition gateway to send local alarm events to a remote gateway for journal logging. This type of profile is useful in cases where alarm events need to be recorded by multiple alarm journal profiles. In addition, this type of profile is useful in [Hub and Spoke architectures](#) as it allows the hub to record the alarm events from each spoke.

Creating a Remote Alarm Journal Profile

Just like configuring alarm journal events to be logged into a database, it is done from the Gateway Webpage, **Config > Alarming > Journal**.

- To have your alarm journal events automatically sent to a remote gateway's alarm journal, select **Remote**, and click **Next**.

The screenshot shows the Ignition configuration interface. On the left, there's a sidebar with 'Config' selected. The main area is titled 'Config > Alarming > Alarm Journal Profiles'. It displays three options: 'Database', 'Remote', and 'Internal'. The 'Remote' option is selected and highlighted with a red box. Below each option is a brief description. At the bottom right of the main area is a blue 'Next >' button.

- A list of known Gateways will be displayed. If you don't see a gateway that you expected to see, check your Gateway Network settings to verify that the connections are valid. You also have the option to specify a gateway manually. This example selects a valid gateway. Click **Next**.

The screenshot shows the same configuration page as before, but now it lists known gateways. One gateway, 'Ignition-Test', is selected and highlighted with a red box. Its connection details are shown: 'Connection: ignition-test|https://10.10.110.59:8060/system, State: Connected'. Below this is another section for specifying a gateway manually. At the bottom right is a blue 'Next >' button.

- If an Alarm Journal exists on the remote gateway, the fields will automatically populate. The name of the gateway and the Alarm Journal Profile name will appear in the **Name** field prefaced with the alarm journal profile name,(i.e., Ignition_Test_Journal), as shown in the following example. Click **Create New Alarm Journal Profile**.

- You will receive a successful message stating your new Alarm Journal Profile was created.

The screenshot shows the final step where a success message is displayed: 'Successfully created new Alarm Journal Profile "Ignition_Test_Journal"'. Below this is a table of alarm journal profiles. The first row, 'Ignition_Test_Journal', is selected and highlighted with a red box. It shows the 'Name' (Ignition_Test_Journal), 'Type' (Remote), and two buttons ('delete' and 'edit'). Other rows show 'Journal' (Database) and 'Packaging_Area' (Database). At the bottom left is a link to 'Create new Alarm Journal Profile...'. The table has columns for Name, Type, and Description.

Name	Type	Description
Ignition_Test_Journal	Remote	<button>delete</button> <button>edit</button>
Journal	Database	<button>delete</button> <button>edit</button>
Packaging_Area	Database	<button>delete</button> <button>edit</button>

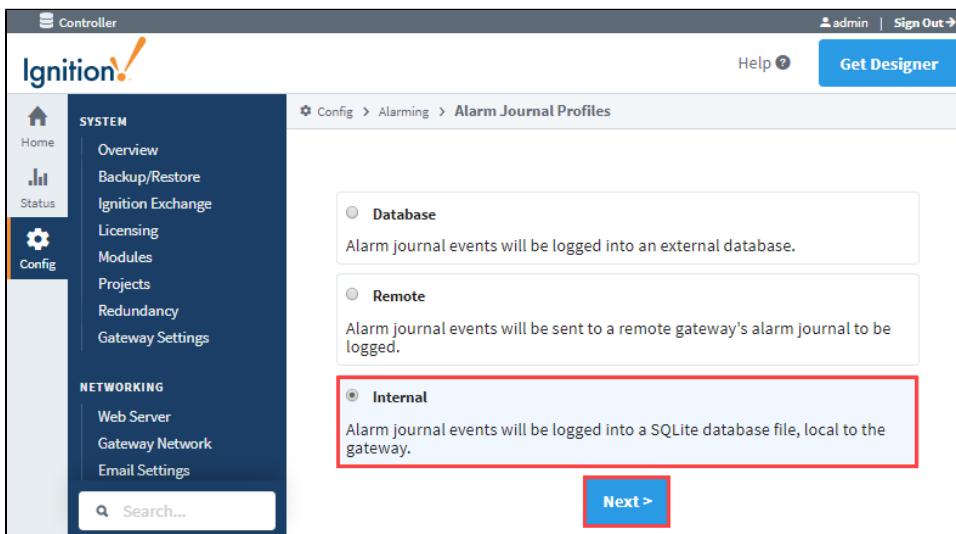
Remote Gateway Alarm Journal Properties Table

Main	
Name	The default name, is the name of the Remote Gateway and Alarm Journal.
Enabled	By default. The journal profile is enabled.
Description	Description of the journal profile. Optional
Query Only	<p>The following feature is new in Ignition version 8.1.5 Click here to check out the other new features</p> <p>If set to true, allows the alarm journal to opt out of being used for storage. When set to true, all alarm events will be discarded by the given journal.</p>
Remote Gateway	
Gateway Name	Name of the Remote Gateway.
Alarm Journal	Name of the Alarm Journal on the Remote Gateway.
Advanced	
Use Store and Forward	<p>The following feature is new in Ignition version 8.1.23 Click here to check out the other new features</p> <p>If enabled, alarm journal events will be stored through the Store and Forward system. If not enabled, they will be stored directly against the remote Gateway. Default is true.</p>

Create an Internal Alarm Journal to Log Events Locally

Ignition Gateways can now create an Internal Alarm Journal Profile that stores Journal entries locally. Go to the Gateway webpage, **Config > Alarming > Journal** to create the Internal alarm journal profile.

1. Click on the the **Create new Alarm Journal Profile...** link.
2. Select **Internal** to have your alarm journal events logged locally.



3. Enter the name of your alarm journal profile and update any settings as required, then click **Next**.

Config > Alarming > Alarm Journal Profiles

Main	
Name	Ignition_Internal_Journal <input type="button" value="Edit"/>
Minimum Priority	Low <input type="button" value="▼"/> Only events equal to or greater than the specified priority will be stored. (default: Low)
Description	<input type="text"/>
Store Shelved Events	<input type="checkbox"/> If enabled, events generated by "shelved" alarms will still be written to the journal system. (default: false)
Enabled	<input checked="" type="checkbox"/> (default: true)

4. You will receive a successful message stating your new Alarm Journal Profile was created.

Internal Alarm Journal Properties Table

Main	
Name	The default name, is the name of the Remote Gateway and Alarm Journal.
Minimum Priority	Only events equal to or greater than the specified priority will be stored.
Description	Description of the journal profile. Optional
Stored Shelved Events	If enabled, events generated by "shelved" alarms will still be written to the journal system. Default is false.
Enabled	By default. The journal profile is enabled.

Alarm Journal Component

The Vision and Perspective modules feature built-in components that can automatically retrieve events recorded in an Alarm Journal. See [Vision - Alarm Journal Table](#) and [Perspective - Alarm Journal Table](#) for more information.

Event Time	Event Id	Display Path	Event State	Priority	System E...	Ack'd By	Event Value	Current S...	Label
2/10/22, 7:06 PM	fa15c4a...	Ramp High Al...	Active	Low	False		9.0515	Active, U...	High Ala...
2/10/22, 7:01 PM	5d53c80...	Ramp High Al...	Clear	Low	False		-9.9588	Cleared, ...	High Ala...
2/10/22, 7:01 PM	5d53c80...	Ramp High Al...	Active	Low	False		9.0407	Active, U...	High Ala...
2/10/22, 6:56 PM	4a5c6fa...	Ramp High Al...	Clear	Low	False		-9.9705	Cleared, ...	High Ala...
2/10/22, 6:56 PM	4a5c6fa...	Ramp High Al...	Active	Low	False		9.0289	Active, U...	High Ala...
2/10/22, 6:51 PM	1590678...	Ramp High Al...	Clear	Low	False		-9.9831	Cleared, ...	High Ala...
2/10/22, 6:51 PM	1590678...	Ramp High Al...	Active	Low	False		9.0161	Active, U...	High Ala...
2/10/22, 6:46 PM	c83fc9d...	Alarm Fault	Clear	Low	False	Auto-Ack		Cleared, ...	Fault
2/10/22, 6:46 PM	00e91bd...	Ramp High Al...	Clear	Low	False	Auto-Ack		Cleared, ...	High Ala...
2/10/22, 6:46 PM	fd32ba2...	Level Lo Alarm	Active	Medium	False		0	Active, U...	Lo
2/10/22, 6:46 PM	ch2d7hd	Level Hi Alarm	Clear	Critical	False	Auto-Ack		Cleared	Hi

12 events   

Journal Properties

Main	
Name	The default name is Journal.
Datasource	Events are stored to this datasource. (Only available on Database type profiles)
Enabled	By default, the journal profile is enabled.
Description	Description of the journal profile.
Query Only	When enabled, the alarm journal will not store alarm events.
Use Store and Forward	Enabled by default, which means the alarm journaled events will be stored through the Store and Forward system . If not enabled, they will be stored directly against the database. This system protects data from being lost due to temporary database connectivity issues. (Only available on Database type profiles)
Events	
Minimum Priority	Only events equal to or greater than the specified priority will be stored. The default is Low. You can set the priority to be: Diagnostic, Low, Medium, High, and Critical.
Store Shelved Events	Not enabled by default. If enabled, events generated by "shelved" alarms will still be written to the journal system.
Store Enabled & Disabled Events	<p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>When enabled, events generated by enabling or disabling alarms will be stored in the journal. This includes cases where the Enabled property on an alarm is toggled, as well as cases where a Tag's Alarm Eval Enabled property is changed. This property additionally relies on setting the Perspective/Vision Alarm Journal Table properties for the enabled and disabled events.</p>
Event Data	
Static Config	By default, it is not selected. If selected, will store the values of static alarm configuration. That is, the alarm properties that are <i>not</i> bound. These do not change during evaluation, only when a user modifies them in the Designer, and so they are not stored by default.
Dynamic Config	If selected, will store the values of dynamically bound alarm configuration properties. The value of these properties can change at any time, and the values at the time of the alarm are captured on the alarm event.
Static Associated Data	If selected, will store the values of non-bound associated data (properties created by the user) properties on alarm that do not change during execution.
Dynamic Associated Data	If selected, will store the values of dynamically bound associated data (properties created by the user) properties.
Data Filters	
<p>Note:</p> <p>The following three properties interact via logical AND, meaning an alarm must meet the criteria for all three for it to be logged in the Journal. Thus, if you supply values for all three Data Filter properties, then an alarm must match the Filter by Alarm Source, Filter by Display Path, and Filter by Display Path or Source properties.</p> <p>Example: If a journal has values for all three properties, and an alarm only meets the requirements for Filter by Alarm Source and Filter by Display Path or Source, but not Filter by Display Path, then the alarm will not be logged to the Journal.</p> <p>If you want to filter on both the Display Path and Source Path, you would want to use only one of the two following approaches:</p> <ul style="list-style-type: none"> • Filter by Alarm Source and Filter by Display Path • Only use Filter by Display Path or Source <p>It is recommended to avoid using all three properties simultaneously on the same Journal.</p>	

Filter by Alarm Source	Only events matching the source will be stored. Multiple sources to match can be comma separated. Leave blank to store events from all sources.	
Filter by Display Path	Only events matching the display path will be stored. Multiple display paths to match can be comma separated. Leave blank to store events from all display paths.	
Filter by Display Path or Source	Only events matching the display path, if defined, will be stored. Multiple matches can be comma separated. If no display path is defined, only events matching the source will be stored. Leave blank to store all events.	

Data Pruning

Enable Data Pruning	If selected, data will be deleted after the specified time period as set by the Prune Age and Units below. Default is false. Note: Since the data is stored directly in a database, an administrator is free to manually delete data at any time.	
Prune Age	The number of Prune Age Units to store data for. i.e., 1 year, 5 hours, etc. The default is 1.	
Prune Age Units	The type of Prune Age Unit. Default is Years. You can choose the unit to be Milliseconds, Seconds, Minutes, Hours, Days, Weeks, Months, or Years.	

Advanced

These settings let you specify your own table names. This is especially useful when trying to use multiple alarm profiles within a single database (not common, but can happen, especially with multiple systems sharing a single database).

Table Name	The table name for the core event table. The default is <code>alarm_events</code> .	
Event Data Table Name	The table name for event data associated with alarms. The default is <code>alarm_event_data</code> .	

Table Definitions

The Alarm Journal system will automatically create the necessary tables for you, and scripting functions can be used to query the system without having to know about the table structure. However, understanding the structure of the Alarm Journal tables can be useful for accessing the data in situations where SQL queries are more convenient.

Alarm Events (`alarm_events`)

This table stores the core data for each event that occurs. An event is a transition for an alarm between active, cleared, or acknowledged. Additionally, other events may be stored in this table that aren't directly related to an alarm, such as a system shutdown event. This table defines a primary key "id", that is used as a foreign key by the Alarm Event Data table, which stores additional information for each event.

Column Name	Data Type	Description
<code>id</code>	integer	A unique integer id for each event entry event
<code>eventid</code>	string	The UUID of the alarm event.
<code>source</code>	string	The qualified path of the entity that generated the alarm event. See below for more information about qualified paths.
<code>display path</code>	string	The value set for the "Display Path" of the alarm. Generally a user defined, friendlier version of the source.
<code>priority</code>	integer	The priority or severity of the alarm: 0: Diagnostic 1: Low 2: Medium 3: High 4: Critical

eventtype	integer	<p>The type of transition represented by this event:</p> <p>0: Active 1: Clear 2: Acknowledgement</p> <p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>The following values were added in 8.1.8</p> <p>4: An alarm was enabled 5: An alarm was disabled</p>
eventflags	integer	<p>A numeric bitmask flag field providing additional information about the event.</p> <p>Bit 0: System Event - One of the designated system events. (System Startup, System Shutdown) Bit 1: Shelved Event - The alarm was "shelved" at the time that the event occurred. Shelving alarms does not prevent execution, so if the journal is configured to store shelved events, they will be stored even if they're not sent to the notification system, or shown to users. Bit 2: System Acknowledgement - When the "live event limit" (defined in general alarm settings) is exceeded, the system will automatically acknowledge overflow events, and the acknowledgment event will have this flag set. Bit 3: Acknowledge Event - The event was acknowledged at the time of the event. For events that are cleared after being acknowledged. Bit 4: Cleared Event - The event was cleared at the time of the event. For alarms that are acknowledged after being cleared.</p> <p>The following feature is new in Ignition version 8.1.8 Click here to check out the other new features</p> <p>The following bit was added in version 8.1.8:</p> <p>Bit 5: Enabled - Signifies that the enabled state on the alarm was changed.</p>
eventtime	datetime	The time of the event.

Alarm Event Data (alarm_event_data)

This table stores the properties associated with an alarm event. The individual event is referenced through the ID column, against the alarm event table.

Column Name	Data Type	Description
id	integer	The id that corresponds to the alarm event in the alarm_events table.
propname	string	The name of the property. May be one of the common alarm properties (a configuration setting), or the name of an associated data property.
dtype	integer	The data type of the property, indicating which data column should be used: 0: Integer 1: Float 2: String
intvalue, floatvalue, strvalue	integer, float (double), string	The corresponding value columns for the property. Unused columns will receive "null" values.

Qualified Paths

A qualified path in Ignition is a path to an object, described by various annotated components. Each component has a type identifier and a value, separated by a colon (:), and each component is separated by colon-forward slash (:/). For example, an alarm is identified by alm:Alarm Name. It usually exists under a Tag, in which case, its fuller path would be tag:Path/To/Tag:/alm:Alarm Name. Paths can be built up further depending on the level of specificity required by the situation.

Configuring Alarms

Alarms are conditions that are evaluated when the value of the Tag changes. When the condition becomes true, the alarm is said to be **active**. When it becomes false, the alarm is said to be **clear**. Alarms may also be **acknowledged**. This flags the alarm in Ignition so the acknowledgement state of each alarm can be made visible throughout Ignition.

Alarms can be configured on Tags or OPC items in [SQL Bridge \(Transaction Groups\)](#). You can also put alarms on [System Tags](#) that Ignition inherently provides such as the Gateway Performance and CPU Usage.

Note: Dataset type tags are not supported by the Tag History system.

Alarm Names and Forward Slashes

Alarm names can make use of forward slashes ("/"). However the alarm name property throughout Ignition is designed to only show characters after the rightmost forward slash. Assume you name an alarm with the following:

one/two/three

The alarm name property (as seen on the various alarm table components, Tag Browser, and various alarm scripting functions) will omit "one" and "two", leaving the alarm name as:

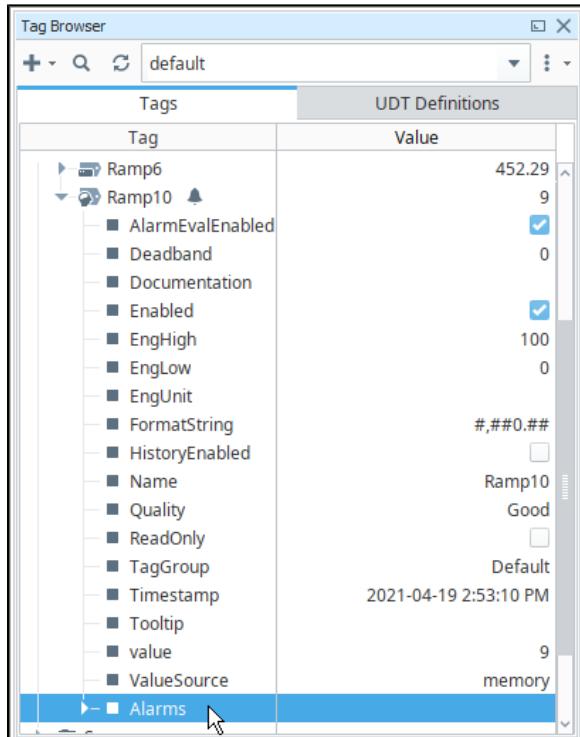
three

This behavior can be confusing, but is easily avoided by not using forward slashes in alarm names.

Alarm Properties

Similar in concept to properties on Vision components, alarm settings, also known as alarm properties, allow you to modify the behavior of each alarm. Each alarm will have its own alarm properties such as a unique Name, Priority, Display Path, Notes, and many other properties, some of which are optional. Descriptions of each alarm property can be found on the [Tag Alarm Properties](#) page.

Aside from viewing alarm properties from the Tag Editor, you can also examine them from the Tag Browser. If you expand a tag that has an alarm on it, you'll find an **Alarms** item:



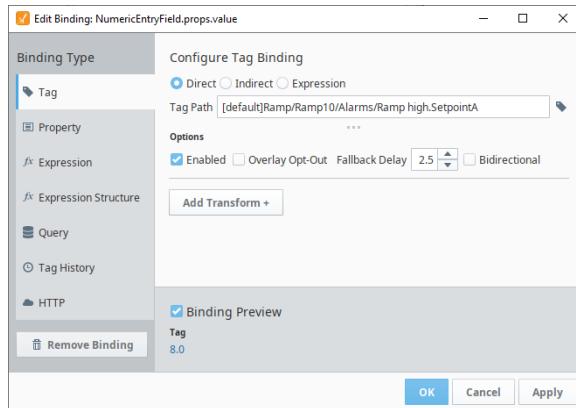
On this page ...

- [Alarm Names and Forward Slashes](#)
- [Alarm Properties](#)
 - [Alarm Modes and Setpoints](#)
 - [Alarming on Individual Bits](#)
 - [Deadband and Time Delays](#)
 - [Associated Data](#)
- [Configuring Alarms](#)
 - [Configuring an Alarm on a Memory Tag](#)
 - [Configuring Alarms on Tags in a UDT](#)
 - [Configuring Multiple Alarms on a Single Tag](#)
- [Alarm Property Bindings](#)
- [Alarms in Transaction Groups](#)

From here we can expand the **Alarms** item. We'll find several data points that represent the state of any alarms on the Tag. In addition, can find runtime information for each alarm on the Tag. In the image below, the Tag has an alarm named "**Ramp high**". We can expand that tag to learn more about that alarm.

The screenshot shows a hierarchical tree view under the 'Alarms' node. The tree includes items like ActiveAckCount, ActiveUnackCount, ClearUnackCount, HasActive, HasUnacknowledged, HighestAckedName, HighestAckedPriority, HighestActiveName, HighestActivePriority, HighestUnackedName, HighestUnackedPriority, LastActiveTime, and ShelvesCount. Below this tree, a specific alarm named 'Ramp high' is selected, highlighted with a blue background. This selection reveals its sub-properties: AckTime, AckUser, and AckUserName. The 'AckTime' property is set to null. The 'AckUser' and 'AckUserName' properties are also set to null. The 'LastActiveTime' property is displayed as 2021-04-19 2:53:10 PM.

These sub properties are simply attributes on a tag, so Tag Bindings can be used to retrieve their realtime value. In the image below, we see a Perspective Tag Binding pointing to the **State** attribute of an alarm.

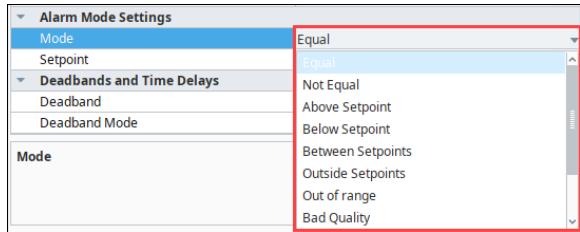


The following feature is new in Ignition version **8.1.5**
[Click here](#) to check out the other new features

Alarm setpoint properties and the Enabled property can now be modified directly from their attributes on the tag. This means component property bindings and scripts can change the setpoint or enabled state of an alarm during runtime.

Alarm Modes and Setpoints

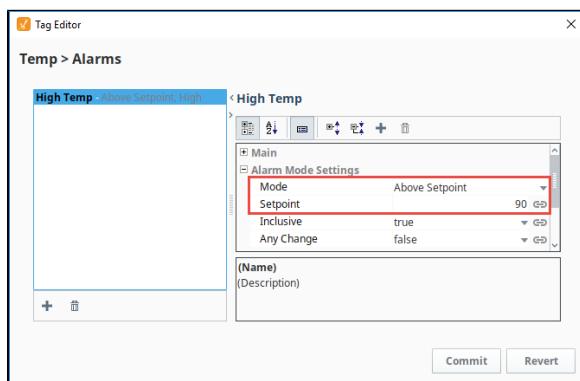
There are many different conditions that can be set up on an alarm. Alarm Mode Settings is where you define the actual conditions when the alarm goes 'true.' Some of the various Alarm Modes are shown in the image below.



Each alarm is configured with one mode, and usually one or more Setpoints. **Modes** determine the method in which alarm activity is evaluated, while Setpoints are the thresholds or limits that determine when a Tag is within the alarm state. Both properties work in together to determine when alarms become active, as well as when they are cleared.

For example, when Mode is set to "Above Setpoint", the Setpoint attribute is compared to the Tag's value. When the Tag value is above the the Setpoint value, the alarm becomes active. Once the value of the Tag is less than the Setpoint value, the alarm will transition to a cleared state.

For a complete list of Alarm Modes and their descriptions, refer to the [Reference Table](#) on the Tag Alarm Properties page.



Alarming on Individual Bits

Additionally, alarms can become active by bit state. A common practice for PLC programming is to store alarm conditions as series of bits inside the PLC and expose them to the world as an integer value. Each bit is essentially a binary value. A series of 8 bits, for example, could be represented as an 8-bit integer. This integer value can be monitored by Ignition's OPC-UA server as an integer value as a Tag. This Tag can then have as many alarms as the integer has bits by monitoring the integer's bit state. Using the Bit Position Mode, multiple alarms in Ignition can be configured to monitor the state of each bit.

The following table shows binary and decimal equivalents:

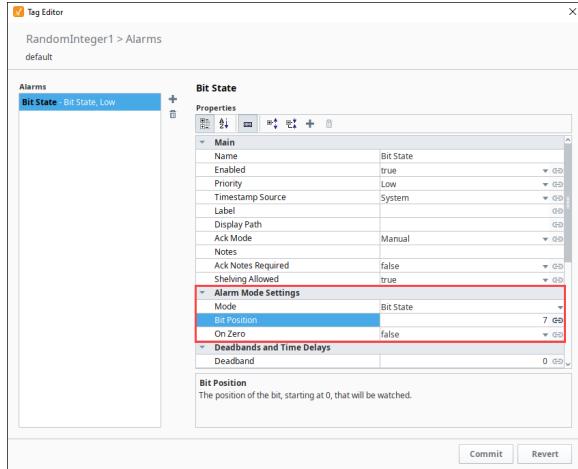
Binary	Decimal
10001110	142

In this case, the Tag would be an integer with a value of 142. This Tag would have eight alarms. Each one of these alarms become active in accordance with its Boolean value.

Create Alarms on an Integer Tag Value

1. Edit a Tag that has an integer value.
2. Scroll down to Alarms, and next to 'No alarms,' click on the **Edit** icon to create an alarm.
3. Click the **Add** icon to add a new alarm to the Tag.
4. Fill in the appropriate properties, for example, enter the **Name**, **Priority**, and other properties as appropriate.
5. From the **Mode** section dropdown, select **Bit State** and enter a value for the **Bit Position**.

- Click **Commit** to save the alarm. Repeat the same process for the remainder of the bits.

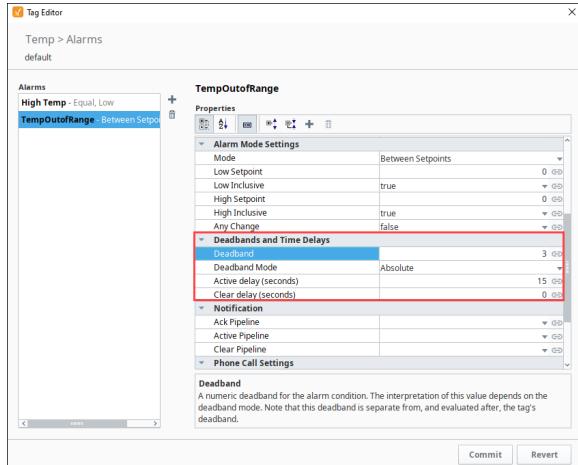


Deadband and Time Delays

The value for the deadband is interpreted by the Deadband Mode. All alarms are only evaluated **after** the Tag's value changes, which means that the Tag's own deadband will be considered first. When the deadband is positive, an active alarm condition needs to clear its setpoint(s) by the amount of the deadband for the alarm to clear.

For example, suppose we have a "Between Setpoints" alarm with a Low Setpoint of 45 and a High Setpoint of 65, and with a Deadband of 3. The alarm will go active if the value is between 45 and 65, but will only clear if the value falls below 42 or rises above 68.

In other situations, a Tag may frequently enter and leave an alarm state, but only for a brief moment. Normally an alarm would be generated each time the alarm condition was met, but the **Active Delay** attribute will prevent any alarms from being generated until the value on the Tag has stayed in an alarm state for a set period of time.



To learn more about deadband and time delays, refer to the [Reference Table](#) on the Tag Alarm Properties page.

Associated Data

You can extend the list of alarm properties by adding your own **Associated Data** or custom properties to an alarm you already have configured. These values can be static or dynamic. Static properties are excellent for filtering alarms. Dynamic properties can be driven by another Tag, or evaluated with an expression. Regardless of the type, the Associated Data property values will be attached to the alarm event, viewed in [real-time](#), and recorded in the [Alarm Journal](#) system.

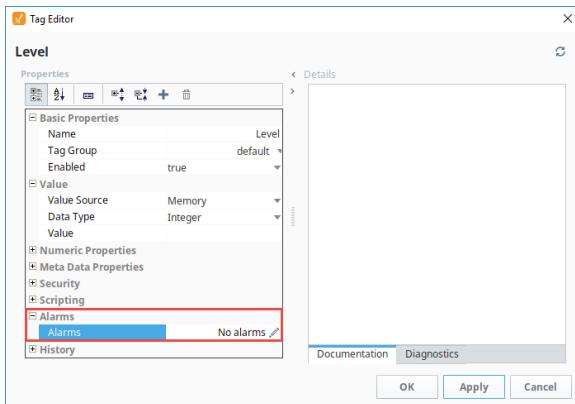
Configuring Alarms

Alarms can be configured on any Tag type: Memory Tag, Query Tag, Expression Tag, as well as Tags inside of a UDT.

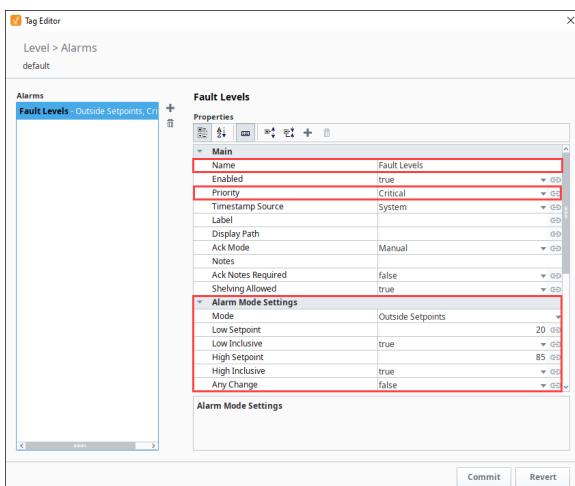
Configuring an Alarm on a Memory Tag

Let's use a simple Memory Tag with an Integer data type and create an alarm.

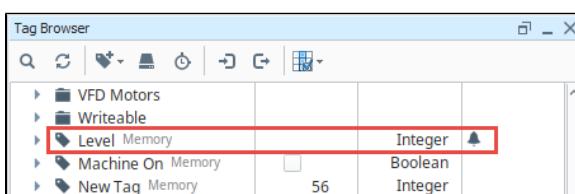
1. In the Tag Browser, double click a memory Tag (i.e. Level) to open the **Tag Editor**.
2. Scroll down to Alarms, and next to 'No alarms,' click on the **Edit** icon to create an alarm, and the alarm UI will slide in from the right.



3. Click the **Add** icon in the bottom left corner of the window, or double click **New Alarm** item to add a new alarm to the Tag.
4. Configure the following alarm settings:
Name: **Fault Levels**
Priority: **Critical**
Alarm Mode: **Outside Setpoints**
Low Setpoint: **20** (when the Low Setpoint is below 20, an alarm is triggered)
Low Inclusive: **true** (alarm is triggered when the low setpoint is equal to 20 or below)
High Setpoint: **85** (when the High Setpoint is above 85, an alarm is triggered)
High Inclusive: **true** (when the High Setpoint is equal to or above 85, an alarm is triggered)
5. Click **Commit** to save your alarm, or **Revert** to cancel.



6. Click **OK** to save your Tag edits.
7. Ignition will start monitoring the alarm immediately. Notice that the Alarm icon will appear in the **Traits** column of the Tag Browser next to the name of the Tag. This means that at least one alarm is configured on this Tag.



Configure an Alarm

[Watch the Video](#)

Configuring Alarms on Tags in a UDT

Alarms can be configured on Tags inside a UDT so when you create instances of that UDT, the alarms will automatically be configured.

In this example, we have a Motor UDT that contains two Tags: Amps and HOA. Let's configure an alarm on the Amps Tag.

1. In the **Tag Browser**, double click on a UDT (i.e., Motor) to open the **Tag Editor**.
2. In the Type Structure area, click on a Tag (i.e., Amps).
3. Scroll down to Alarms, and next to **No alarms**, click on the **Edit** icon to create an alarm.

4. Click the **Add** icon.

5. For this example, we entered the following:

Name: **Low Amps**

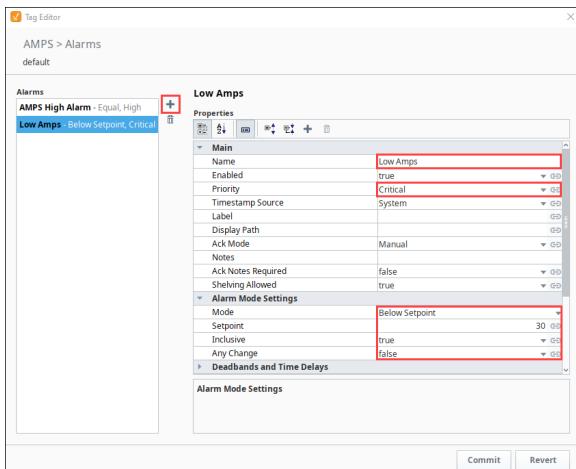
Priority: **High**

Alarm Mode: **Below Setpoint**

Setpoint: **30**

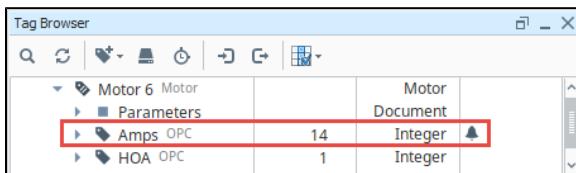
Inclusive: **true** (an alarm is triggered when the low setpoint is equal to 30 or below)

6. Click **Commit** to save your alarm.



7. Click **OK** to save the UDT. Now you're ready to create instances of the UDT. Once you create your UDT instances, Ignition will start monitoring the alarm immediately.

Note: The Alarm icon will appear in the **Traits** column of the Tag Browser next to the name of the Tag in the UDT instance. (Notice how you see the name of the Motor UDT next to the Amps Tag of the Motor 6 instance.)



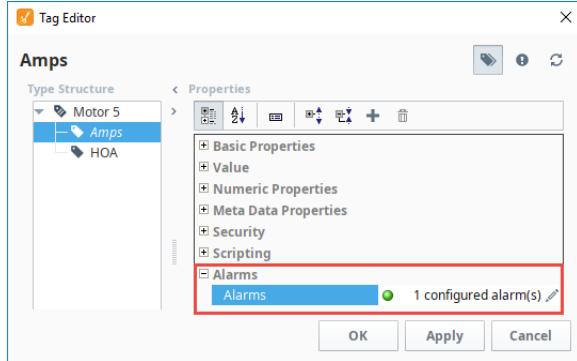
Updates to the UDT and Overriding Alarm Settings

If the alarm in the UDT is updated, the instance will automatically receive the updates and be refreshed. It's not uncommon that you may need to make alarm attribute values unique for particular instances, in which case, you can override alarm settings. Click the override button on the alarm UI of the Tag Editor making it green and edit the alarm properties and/or their values To learn more, refer to Overriding Properties in UDT Instances.



Configure Alarm in UDT

[Watch the Video](#)



Configuring Multiple Alarms on a Single Tag

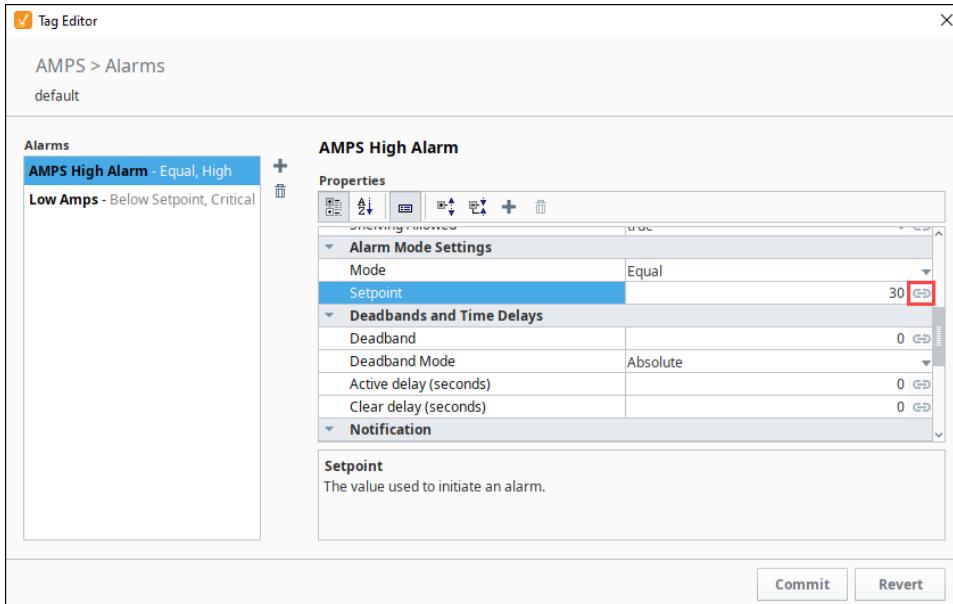
Tags can have multiple alarms configured, each with unique setpoints and other configurations. The number of alarms that you can configure is virtually unlimited and will never exceed the bounds of reasonable design consideration. Adding additional alarms is just as easy as adding the first alarm. Simply click on the **Add** icon in the alarm UI and start configuring the new alarm. This way separate alarms can monitor different setpoints, so Tags can have alarms for high setpoints and low setpoints.

Alarm Property Bindings

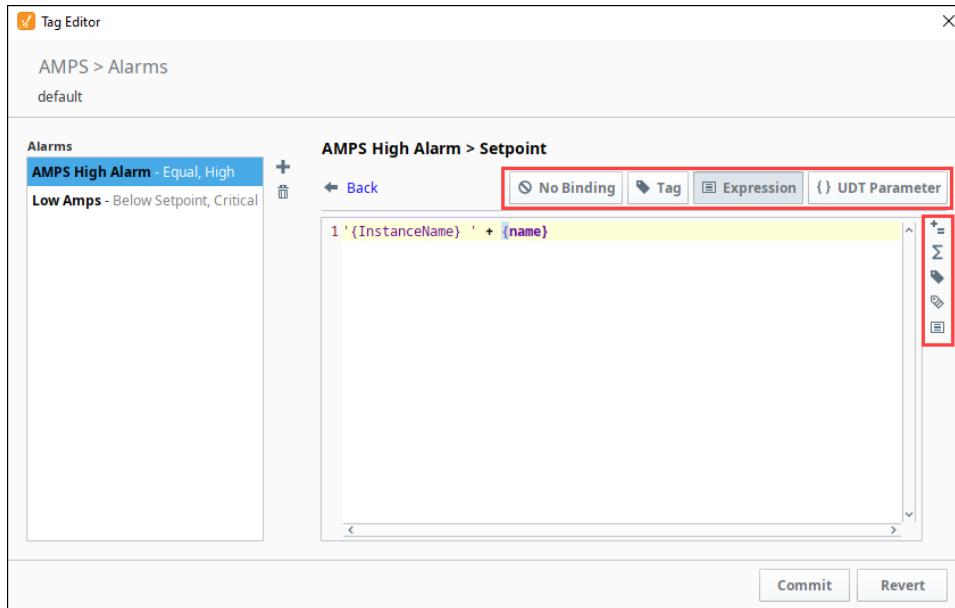
Many alarm properties are bindable, which means they can be bound to other Tags in the system, expressions and even a UDT parameter. For example, you might bind the Enabled property to another Tag which represents whether or not your process is running, thereby disabling the alarm when production is stopped. Or, you might bind the setpoint of an alarm to a Tag that operators can manipulate, thereby letting the setpoint be changed at runtime.

This example shows how to bind an alarm property from the Tag Browser in Perspective.

1. Double click on your Tag to open the **Tag Editor**.
2. Click on the **Edit** icon next to **Configured Alarms** to see your configured alarms. (If you don't have any alarms, create an alarm using the steps in the [Configuring an Alarm on a Tag](#).)
3. Select the alarm and the screen will refresh with all the alarm properties.
4. From here, find the alarm property you want to add a binding to and click on the binding icon. The binding UI will slide in from the right.



5. Select the binding type (**No Binding**, **Tag**, **Expression**, or **UDT Parameter**, if applicable). The image below shows an example of an Expression binding. Notice that the expression can reference many useful values such as the Tag's value and other settings of the alarm. Enter your expression.
6. Once you configured the binding to your liking, click **Commit**, or **Revert** if you decide to cancel.



7. Click **OK** to save the changes to the Tag.

For more information on property bindings see, [Property Bindings in Perspective](#) and [Property Bindings in Vision](#).

Alarms in Transaction Groups

Alarms can also be added to OPC items in Transaction Groups. This means alarms can be used without ever creating a Tag in Ignition. Simply edit an OPC item, and an Alarming section will appear in the Tag Editor window. From here, adding an alarm to the item is similar to adding an alarm to an Ignition Tag.

Value Source	OPC
Data Type	Integer
OPC Server	Ignition OPC UA Server
OPC Item Path	ns=1;s=[Dairy]_Meta:Overview/AU 1/Fan 2 HOA

Dynamic Alarm Attributes

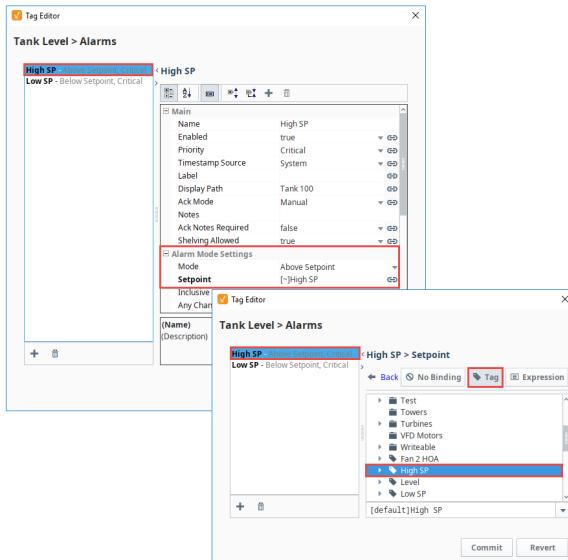
Dynamic Setpoints

You can configure alarms with dynamic setpoints inside of a Tag definition. It is similar to configuring an alarm, however, configuring an alarm with a dynamic setpoint requires additional Tags to serve as these setpoints. A good example of this is when an operator changes a high or low setpoint, it also changes how the alarm is evaluated.

To Configure Dynamic Setpoints on an Alarm

Suppose you want to alarm a Tank Level if it goes above or below a certain setpoint. Rather than hardcoding a value in the high and low Setpoint properties, you can bind them to either a Tag in the system or an Expression to make them dynamic. This example uses an OPC Tag called Tank Level, and two Memory Tags called High SP and Low SP. The Tank Level will use the values in the Memory Tags as the high and low setpoints. Let's alarm the tank level when the setpoint goes above 85 and below 20.

1. Create two Memory Tags calling one **High SP** with value of 85, and another one called **Low SP** with a value of 20.
2. Use an OPC Tag and rename it to **Tank Level**. Click on the Tag to open the **Tag Editor**.
3. Scroll down to Alarms, and click on the **Edit** icon next to your alarm to open the Alarm UI.
4. If you don't have a configured alarm, you will need to create one. Select the configured alarm that you want associate a high setpoint with. (This examples uses the **High SP** alarm).
5. Under Alarm Mode Settings, set the Mode to **Above Setpoint**.
6. Click on the **binding** icon for the **Setpoint** property. You can select a **Tag** or **Expression**. (This examples binds the **Setpoint** property to a Memory Tag that was created called **High SP**). Click the **Commit** to commit your changes.



7. Now let's configure the low setpoint. Create another alarm to associate with the low setpoint if you don't have one.
8. Under Alarm Mode Settings, set the Mode to **Below Setpoint**.
9. Click on the **binding** icon for the **Setpoint** property. You can select a **Tag** or **Expression**. (This examples binds the **Setpoint** property to a Memory Tag that was created called **Low SP**). Click the **Commit** to commit your changes.
10. Click **OK**. Now, the tank alarm will be evaluated based upon the high and low setpoints of each Tag.

On this page ...

- [Dynamic Setpoints](#)
 - [To Configure Dynamic Setpoints on an Alarm](#)
- [Dynamic Enabling and Disabling](#)
 - [Enabling Based on Machine State](#)
 - [Enabling Based on Time of Day](#)



INDUCTIVE
UNIVERSITY

UDT Alarm Dynamic Setpoints

[Watch the Video](#)

11. To test it, change the value of the Tank Level to 99, and you'll notice it activates the alarm.

Tag Browser			
-> Tank Level OPC	99	Float	▲
Enabled	<input checked="" type="checkbox"/>	Boolean	
OpcItemPath	ns=1;s=[Generic]_Meta:Writeable...	String	
OpcServer	Ignition OPC UA Server	String	
Quality	Good	String	
TagGroup	Default	String	
Timestamp	2019-06-03 12:19:55 PM	DateTi...	
value	99	Float	
Alarms		String	
High SP		DateTi...	
AckTime		String	
AckUser		String	
AckUserName		DateTi...	
ActiveTime	2019-06-03 12:19:56 PM	DateTi...	
ClearTime		DateTi...	
DisplayPath	Tank 100	String	
DisplayPathOrSource	Tank 100	String	
EventState	Active	Integer	
EventTime	2019-06-03 12:19:56 PM	DateTi...	
EventValue	99.0	String	
IsAcked	<input type="checkbox"/>	Boolean	
IsActive	<input checked="" type="checkbox"/>	Boolean	
IsClear	<input type="checkbox"/>	Boolean	
IsShelved		Boolean	
Label		String	
Name	High SP	String	
Priority	Critical	Integer	
Source	prov:default/tag:Tank Level/al...	String	
State	Active, Unacknowledged	Integer	
Low SP		String	



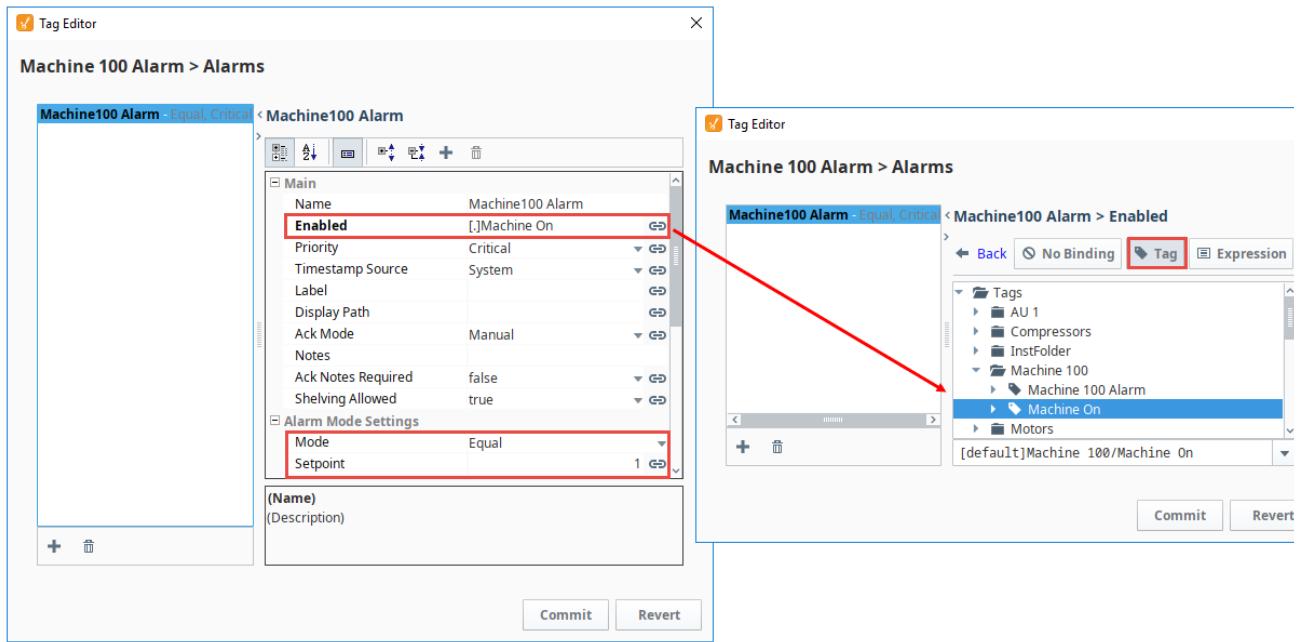
While you can configure dynamic values on any property that is showing the binding icon on the right, some other properties (like Name) may also accept dynamic properties using the {myParam} notation. We recommend against using Dynamic setpoints on these static alarm properties because they will only evaluate on startup, not while the Gateway is running.

Dynamic Enabling and Disabling

Enabling Based on Machine State

Allowing a dynamic condition to determine if an alarm is enabled or disabled is possible inside the alarm's properties, like whether or not a machine is on. This example uses two Tags: one for the Alarm, and one to determine the running state of the machine. Two conditions must be true for the alarm to become active. The value on the alarm Tag must match the setpoint ('true' in the example below), and the **Enabled** property must evaluate to 'true' driven by the value of the 'Machine On' Tag.

1. Select a **Tag** that has the alarm you want to configure.
2. To access the alarm properties, click on the **Edit** icon next to your alarm to open the Alarm UI. Select an alarm or create a alarm if one doesn't exist. The alarm in this example is called Machine 100 Alarm.
3. Click the **binding** icon for the **Enabled** property and bind it to a **Boolean Tag** or an **Expression** that evaluates a 'true' or 'false' condition. This example uses a Tag called Machine On. Enter the **Mode** to '**Equal**' and the **Setpoint** to '**1**'. Click **Commit**.



4. Click **Commit** to save your changes. Now, the alarm will only be evaluated based upon the state of the Machine On Tag. Set the Machine to **On** or 'true.'

5. In the Tag Browser, expand your Tag (i.e., Machine 100 Alarm), and you'll notice the **'IsAlert'** is now active. As you can see, making the Enable property dynamic based on another Tag in the system is extremely easy.

Tag Browser			
	Machine 100		
	Machine 100 Alarm OPC		
	Enabled	<input checked="" type="checkbox"/>	Boolean
	OpcItemPath	<input checked="" type="checkbox"/>	String
	OpcServer	<input checked="" type="checkbox"/>	String
	Quality	<input checked="" type="checkbox"/>	String
	TagGroup	<input checked="" type="checkbox"/>	String
	Timestamp	<input checked="" type="checkbox"/>	DateTime
	value	<input checked="" type="checkbox"/>	Boolean
	Alarms		
	Machine100 Alarm		
	AckTime	<input type="checkbox"/>	DateTime
	AckUser	<input type="checkbox"/>	String
	AckUserName	<input type="checkbox"/>	String
	ActiveTime	<input checked="" type="checkbox"/>	DateTime
	ClearTime	<input type="checkbox"/>	DateTime
	DisplayPath	<input type="checkbox"/>	String
	DisplayPathOrSource	Machine 100/Machine 100 Alarm/Machine...	String
	EventState	Active	Integer
	EventTime	2019-06-03 3:43:18 PM	DateTime
	EventValue	true	String
	IsAcked	<input type="checkbox"/>	Boolean
	IsActive	<input checked="" type="checkbox"/>	Boolean
	IsClear	<input type="checkbox"/>	Boolean
	IsShelved	<input type="checkbox"/>	Boolean
	Label	<input type="checkbox"/>	String
	Name	<input type="checkbox"/>	String
	Priority	Critical	Integer
	Source	prov:default:/tag:Machine 100/Machine 10...	String
	State	Active, Unacknowledged	Integer
	Machine On Memory	<input checked="" type="checkbox"/>	Boolean

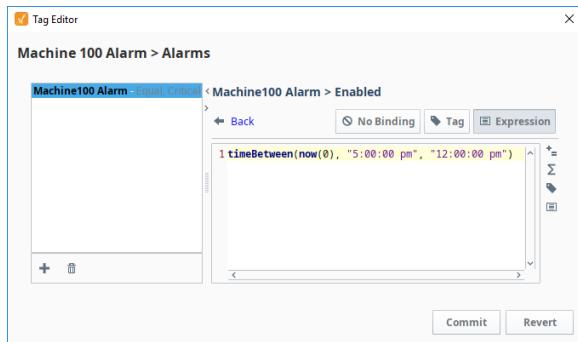
Enabling Based on Time of Day

There are occasions when you only want alarms to be evaluated at certain times of the day. In Ignition, you can automatically enable and disable alarms for specific times of the day. This is typically achieved by binding the alarm's Enabled property to a Tag or an Expression. (This example uses an Expression).

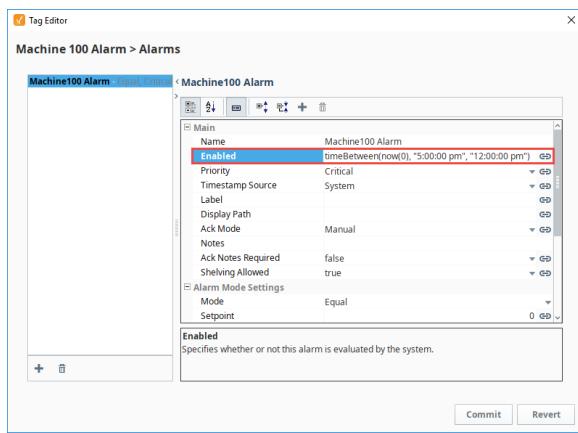
- From the **Tag Browser**, select the **Tag** on which you want to configure the alarm.
- To access the alarm properties, click on the icon next to your alarm to open the Alarm UI.
- Select an alarm you want to place the binding on, or create a alarm if one doesn't exist. The alarm in this example is called Machine 100 Alarm.
- Bind the **Enabled** property to an expression that evaluates a 'true' or 'false' condition based upon the current time. Click the binding button to the right of the Enabled property, and click on **Expression** tab on the top right of the screen, and enter your expression or copy and paste from the code block below, then click **Commit**.

In the expression language, there are various functions for dates that can be used. For example, the following expression would return 'true' if the time is between the hours of 5pm and 12pm, and return 'false' if it is not.

```
timeBetween(now(0), "5:00:00 pm", "12:00:00 pm")
```



5. Click **Commit** again to save your edits. The alarm will only be evaluated when the system clock falls between the specified time.



Alarms in UDTs

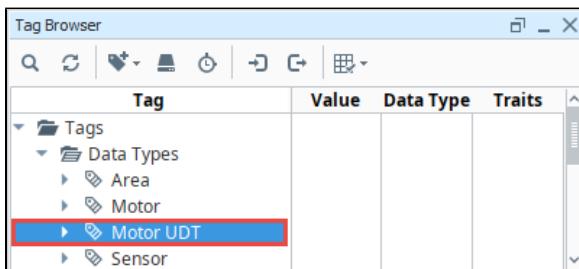
Configuring an Alarm on a UDT Member

The great thing about UDTs is that you configure it in one place, inside of the UDT definition, and every instance of that UDT will automatically inherit that same configuration. The same concept works for alarms on a UDT. If an alarm is configured inside a UDT, every instance of that UDT will automatically have that same alarm configuration. Even if a new instance is created, it will automatically get that same alarm configuration.

Configuring an Alarm on a UDT

This example uses a Motor UDT. The Motor UDT contains two OPC Tags: Amps and HOA. Let's configure an alarm on the Motor UDT when the Amps setpoint goes under 25.

1. In the **Tag Browser**, go to the **Data Types** folder, and double click on a UDT (i.e., Motor UDT) to edit the definition.



2. Select a Tag (i.e., Amps) and scroll down to Alarms, and click on the **Edit** icon next to your alarm to open the Alarm UI.
3. Click on the **Add** icon on the lower left side of the screen to create a new alarm. Enter an **Alarm Name**, **Display Path**, **Priority**, **Mode**, and **Setpoint**. UDT configurations are set up the same way as the normal [alarm configuration on a Tag](#).

Display Path

Since you have multiple instances of a Motor, as in this example, you need to make sure that the Display Path is unique for every instance, otherwise, every instance of the UDT will have the same Display Path name, and the operator won't know which Motor alarm to respond to. There are a couple of ways to address this. You can leave the Display Path blank, in which case, Ignition will use the full Display Path to the instance of that Tag, or you can bind a property to an expression involving a UDT parameter (i.e., "Motor" + {MotorNumber}).

4. Click **Commit** to save your alarm edits.
5. Click **OK** to save your Motor UDT edits. Now, all instances of the UDT will have a similar alarm setup, but each alarm can be unique to that UDT instance by changing the **Below Setpoint** value.

On this page ...

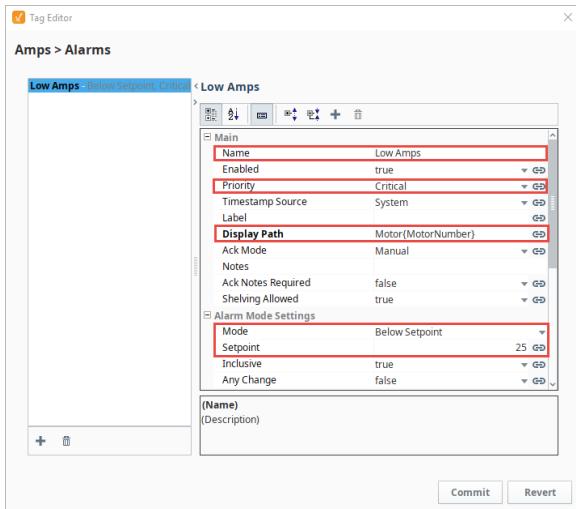
- Configuring an Alarm on a UDT Member
 - Configuring an Alarm on a UDT
- Dynamic Setpoints in UDTs
 - Configuring Dynamic Setpoints Inside a UDT



INDUCTIVE
UNIVERSITY

Configure Alarm in UDT

[Watch the Video](#)



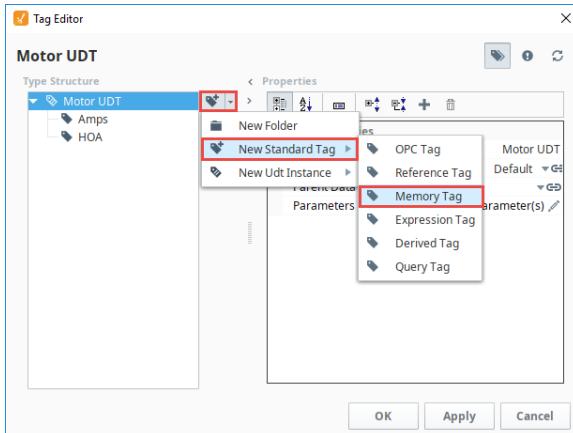
Dynamic Setpoints in UDTs

Instead of using a hardcoded setpoint, you can configure alarms with dynamic setpoints inside of a UDT definition. It is similar to configuring a UDT alarm, however, configuring a UDT alarm with a dynamic setpoint requires additional Tags to serve as these setpoints in all the deployed UDTs.

Configuring Dynamic Setpoints Inside a UDT

In this example, we'll use the Motor UDT in the above section to create a **Memory Tag** inside the UDT definition to serve as the setpoint.

1. In the **Tag Browser**, go to the **Data Types** folder, and click on your UDT to edit the definition.
2. Click the **Add Tag** icon and select **New Standard Tag > Memory Tag** from the dropdown.



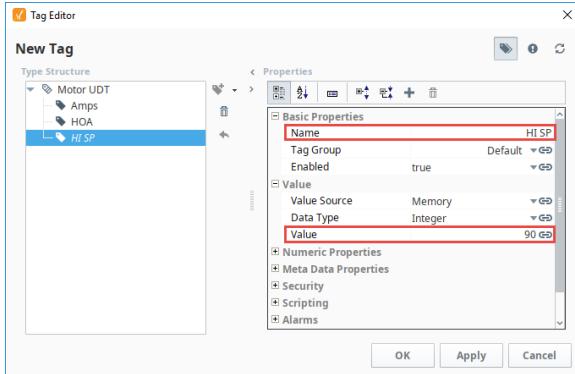
3. Enter the **Name** of the Tag (i.e., HI SP) and a **Value** (i.e., 90). Click **Apply**.



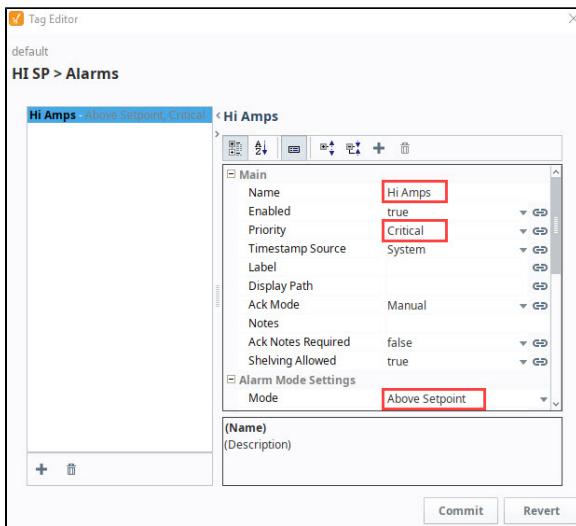
**INDUCTIVE
UNIVERSITY**

**UDT Alarm
Dynamic Setpoints**

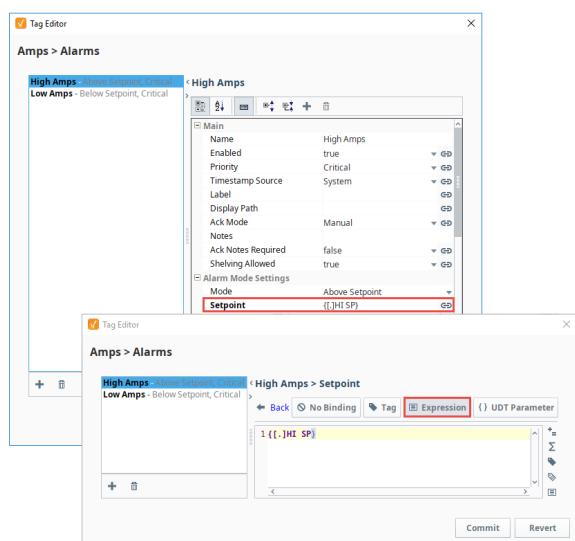
[Watch the Video](#)



4. In the **Type Structure** area, select the Tag you want to associate with the dynamic setpoint (i.e., Amps)
5. Under Properties, scroll down to Alarms. Click on the **Edit** icon next to Alarms.
6. Click on the Add icon to create a new alarm for the high setpoint and call it "Hi Amps".
7. Set the **Priority** as **Critical**, and the **Alarm Mode** as **Above Setpoint**. Click **Commit**.



8. Bind the **Setpoint** property to the new UDT Memory Tag (i.e., HI SP). Click on the **binding** icon and you can either create an expression or use the UDT Tag Tab to select the Memory Tag (i.e., HI SP). This example uses an expression as shown in the image below.
9. Click **Commit** to save the expression.
10. Click **Commit** again to save your alarm edits.



11. Click **OK** to save all your UDT updates. Now you're ready to create instances of your UDT. All of the UDT instances will now have Memory Tags that serve as alarm setpoints.
12. Test it out by dragging an Alarm Status Table component in a window and locating your alarm.

Related Topics ...

- [Configuring Alarms](#)
- [Dynamic Alarm Attributes](#)
- [UDT Multi-Instance Wizard](#)

Alarm Associated Data

What is Alarm Associated Data

Every alarm in Ignition has alarm properties associated with it as it moves throughout the system, like active time, clear time, who acknowledged the alarm, time the alarm was acknowledged, priority, and display path. You can extend the list of alarm properties by adding your own associated data or custom properties to an alarm you already configured. Values can either be static or dynamic. Static properties are great for filtering alarms. Dynamic properties will often be bound to other Tags or an expression. Regardless of the whether a value is static or dynamic, these values will be attached to the alarm event as it moves through the system, and the values will be available from the [Alarm Status](#) system, the [Alarm Journal](#) system, and the [Alarm Notification](#) system.

Note: Associated Data values are always saved as strings. You can typecast these values to any other data type using [Expression](#) bindings or through [Scripting](#).

On this page ...

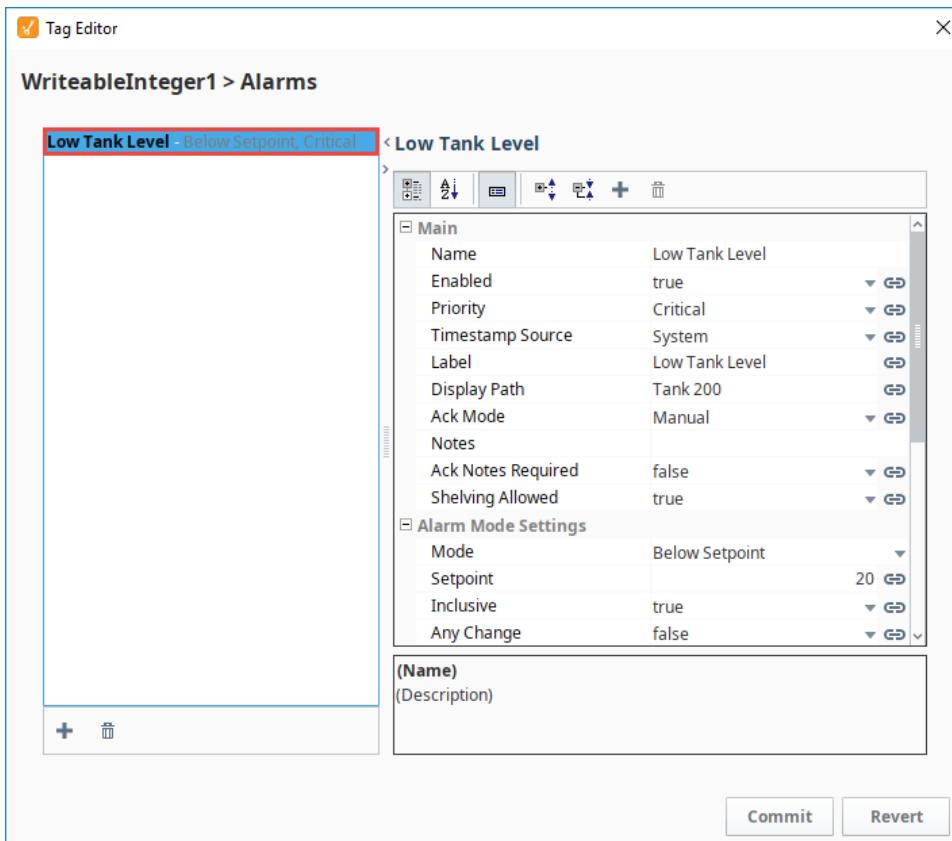
- [What is Alarm Associated Data](#)
 - [Creating Associated Data for both Static and Dynamic Properties](#)
- [Alarm Grouping](#)
 - [Creating an Alarm Grouping](#)

It's easy to go back to an existing alarm and add associated data.

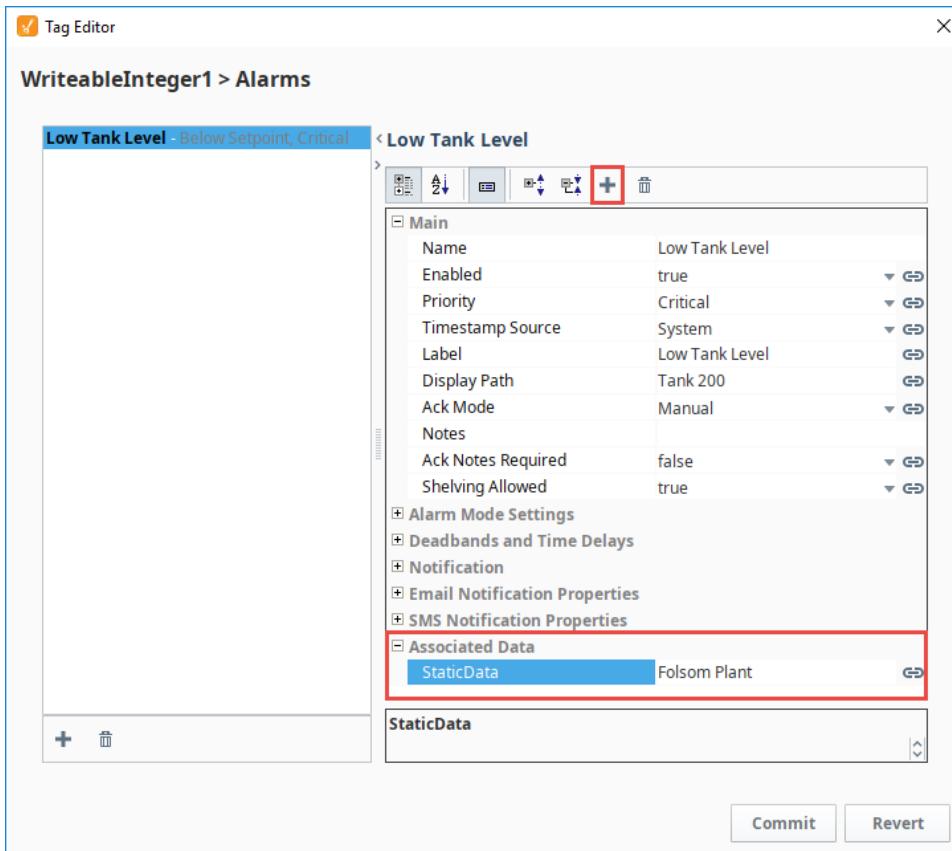
Creating Associated Data for both Static and Dynamic Properties

For this example you can use an existing alarm you already configured or [create a new alarm](#) if you don't have one.

1. Right-click on your Tag and select **Edit Tag**. This example uses a WriteableInteger OPC Tag, but any alarm-capable Tag will suffice.
2. In the **Tag Editor**, scroll down to Alarms, and next to 'No alarms,' click on the **Edit** icon and click on your alarm. The Tag below already has an alarm named ' Low Tank Level.' Select the alarm to open the Alarm UI.



3. Click the **Add** icon above the alarm properties to add new associated data for the alarm. Scroll to the bottom of the property list, and you'll see a new associated data property was added. By default, the new property name is called '**New Data**'.
- Double-click the property name to rename it to something more meaningful. In this example, make the associated data property a static value. We renamed ours to '**StaticData**' with a value of '**Folsom Plant**' and press enter to submit.



4. Now, let's create another Associated Data property with a dynamic value. Click the **Add** icon above the alarm properties.
5. Scroll down to the bottom of the property list and rename your 'New Data' property to '**DynamicData**' and bind it to a Tag using the **binding** icon on the right side of the window. This example uses a Ramp0 OPC Tag.
6. Click the '**Commit**' to submit the Tag selection, and then click **OK** to save your Tag edits.

The screenshot shows two Tag Editor windows. The top window displays the 'Associated Data' list for the 'Low Tank Level' alarm, with 'DynamicData' added and its binding set to '[~]Ramp/Ramp0'. The bottom window shows the 'DynamicData' binding configuration, where 'Ramp0' is selected from a list of tags under the 'Ramp' category. The 'Tag' tab is selected in the binding dialog.

7. When the alarm goes active, Ignition stores the values of the associated data with the alarm. You can view the details of the alarm in the [Alarm Status Component](#) along with the new associated data properties. To view the **Details** of the alarm, check the box of the alarm you want to see, and click on the magnifying glass at the bottom of the window to open the Details window.

The screenshot shows the Ignition Alarm Status Component interface. At the top, there is a table with columns: Active Time, Display Path, Priority, Current State, and Label. Three rows of data are shown:

Active Time	Display Path	Priority	Current State	Label
6/7/19, 11:46 AM	Tank 200	Critical	Active, Unacknowledged	Low Tank Level
6/7/19, 11:40 AM	Tank 200	Critical	Cleared, Unacknowledged	Low Tank Level
6/7/19, 11:41 AM	Tank 200	Critical	Cleared, Unacknowledged	Low Tank Level

Below the table is a details panel titled "Details" which contains a list of alarm properties. The "DynamicData" and "Group" properties are highlighted with red boxes. The "Group" property is set to "Production".

Details

- On Active mode Below Setpoint
- setpointA 20
- Event Value 3
- DynamicData** 995.093333333332
- name Low Tank Level
- Event Time 6/7/19, 11:46 AM
- label Low Tank Level
- StaticData** Folsom Plant
- priority Critical
- displayPath Tank 200

Acknowledge **Shelve**

Alarm Grouping

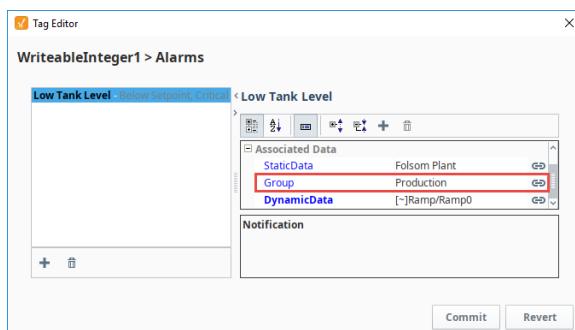
Grouping alarms in Ignition is an important concept. Instead of seeing every alarm in a single view, often times, you may want to view alarms for a particular area of your plant, or for a specific set of alarms.

There are several ways to group alarms. One way is to use a folder structure which requires you to organize your Tags into a hierarchy. Another way is to use the Display Path field in the alarm configuration. The third way, and the most recommended way, is to use associated data. It's a common design practice to associate alarm groupings on the associated data of the alarm.

Creating an Alarm Grouping

This example uses the WriteableInteger Tag that we used in the example above.

1. Use an existing alarm that you already configured. If you don't have one, [create a new alarm](#).
2. Right-click on your **Tag** and select **Edit Tags**. In the **Tag Editor**, scroll down to Alarms, and next to 'No alarms,' click on the **Edit** icon and click on your alarm. Click the **Add** icon above the alarm properties to add new associated data for the alarm.
3. Scroll to the bottom of the property list, and you'll see a new associated data property was added. In this example, make the associated data property a static value, and rename '**New Data**' to '**Group**' with a value that represents your Group (i.e., Production).



4. Click **Commit** to save your alarm edits, and then press **OK**. The alarm will now always have the '**Group**' designation associated with it. This can be used for filtering or part of your alarm pipeline notifications.
5. Before you can use the Alarm Journal to filter on your new associated data property, edit the '**filterAlarm**' Extension Function to include your new associated data property by selecting the [Alarm Journal component](#), and right clicking on **Scripting**. Under the **Extension Functions** folder



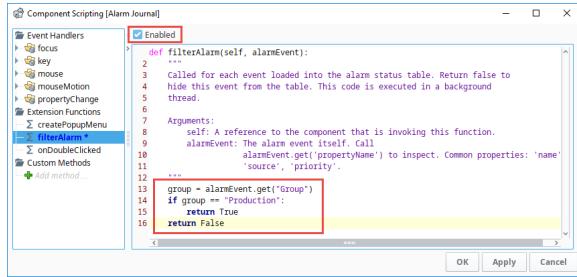
Alarm Grouping

[Watch the Video](#)

, click on **filterAlarm**, then click on the **Enabled** checkbox to enable the script. Edit the script to add your new associated property and property value. You can copy the new code below to add to your filterAlarm script.

filterAlarm scripting function

```
group = alarmEvent.get("Group")
if group == "Production":
    return True
return False
```



- Now, you're good to go! Open the [Alarm Journal](#) component and filter on your associated data. The Alarm Journal Table shown below displays all the alarms that are associated with the group, Production.

Event Time	Display Path	Event State	Priority	Event Value	Current State	Label
6/7/19, 2:35 PM	Pressure/HighPressure	Active	Critical	99	Active, Unacknowledged	HighPressure
6/7/19, 2:35 PM	Pressure/HighPressure	Active	Critical	99	Active, Unacknowledged	HighPressure
6/7/19, 2:30 PM	Tank 200	Active	Critical	7	Active, Unacknowledged	LowTank Level
6/7/19, 2:30 PM	Tank 200	Active	Critical	7	Active, Unacknowledged	LowTank Level
6/7/19, 2:30 PM	Tank 200	Clear	Critical	55	Cleared, Unacknowledged	Low Tank Level
6/7/19, 2:30 PM	Tank 200	Clear	Critical	55	Cleared, Unacknowledged	Low Tank Level
6/7/19, 12:42 PM	Tank 200	Active	Critical	3	Active, Unacknowledged	LowTank Level
6/7/19, 12:42 PM	Tank 200	Active	Critical	3	Active, Unacknowledged	LowTank Level

8 events

Related Topics ...

- [Configuring Alarms](#)
- [Extension Functions](#)
- [Alarm Status Table](#)
- [Alarm Journal Table](#)

Gateway General Alarm Properties

The alarm system has several settings available on the Gateway's web interface. These properties can be found under the **Config > Alarming > General**

Note: The General Alarm Settings page is not available on Edge gateways.

The screenshot shows the Ignition web interface with the following details:

- Header:** Ignition! (with a gear icon), Help ? (blue button), Get Designer (blue button).
- Left Sidebar (Config section):**
 - Home
 - Status
 - Config** (selected)
 - Overview
 - Backup/Restore
 - Ignition Exchange
 - Licensing
 - Modules
 - Projects
 - Redundancy
 - Gateway Settings
- Header:** Config > Alarming > General Alarm Settings
- Alarm Evaluation:**
 - Live Event Limit:** Value: 5. Description: "The number of "live" events (active or unacknowledged) that can exist for a single alarm at a given time. When surpassed, older events will be acknowledged automatically by the system. (default: 5)"
- Event Suppression:**
 - Continuous Event Detection Window (min):** Value: 10. Description: "How long to store events for on shutdown, to prevent new events on startup. Prevents unacknowledged active events from being generated due to reboot. If set to 0, will not be used. (default: 10)"
 - Notify Initial Events:** Description: "If false, active alarms caused by the "initial state" (that is, the first value checked after being created, or after the enabled state changes) won't be sent to the notification system. (default: false)"

Alarm Evaluation	
Live Event Limit	Default is 5. The number of "live" events (active or unacknowledged) that can exist for a single alarm at a given time. When surpassed, older events will be acknowledged automatically by the system. This means as an alarm cycles on and off, Ignition will keep track of the last five times the alarm event happened until the user acknowledges them. This does not store history for those events.
Event Suppression	
Continuous Event Detection Window (min)	Default is 10. The amount of time to store events before shutdown to prevent new duplicate events from being created on startup. This setting prevents unacknowledged active events from being generated due to reboot. If set to 0, will not be used.
Notify Initial Events	Default is false. If false, active alarms caused by the "initial state" (that is, the first value checked after being created, or after the enabled state changes) won't be sent to the notification system. This means if you add an alarm to a Tag, a notification won't be immediately sent when the new state is created.

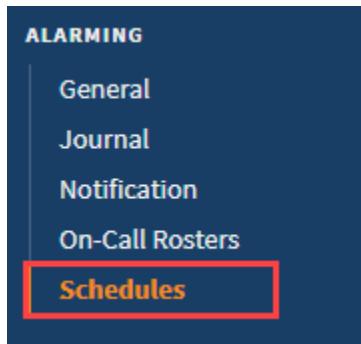
Alarming Schedules

The alarm notification system uses schedules to determine which users should be notified about an active alarm by looking at the users defined in the [notification block](#). This means that notification messages are sent only to active users based on the defined schedules.

You can set a schedule for each user in the alarm notification system and utilize [on-call rosters](#) to create user groups. For example, suppose you have alarms that should be sent to all supervisors. You can put all of the supervisors in one on-call roster and the scheduling system will automatically only notify those supervisors who are on-shift (based on their individually set schedules) when the alarm goes active.

Define a New Schedule

1. Go to the **Config** tab of the Gateway Webpage and scroll down to **Alarming > Schedules**.



2. The Schedule Management page is displayed. Here you can see an Always and an Example schedule.

The **Always** Schedule is a built-in schedule that is always available: 24/7/365.
The **Example** Schedule is an example of a M-F 8am-5pm schedule with a lunch break. Clicking **edit** will access the detailed settings.

A screenshot of the 'Schedule Management' page. At the top, there are tabs for 'Schedules' (which is selected) and 'Holidays'. Below the tabs is a table with columns: Name, Type, Description, and Active Now?. There are two rows: 'Always' (Standard Schedule, Built-in schedule that is always available: 24x7x365, Yes) and 'Example' (Standard Schedule, An example of a M-F 8am-5pm schedule with a lunch break, No). To the right of the table are 'delete' and 'edit' buttons. At the bottom left is a link 'Create new Schedule...'. The URL in the address bar is 'Config > Alarming > Schedule Management'.

3. Click on **Create new Schedule**.
4. For our example, we'll set up a new Standard schedule. Enter a schedule name, description, and set the hours:
Name: **Weekend Basic**
Description: **Regular Weekend schedule, no holidays**
Observe Holidays: **No** (Unselect this option)
All days: **No** (Unselect this option)
Weekdays: **No** (Unselect this option)
Saturday: **Yes** (Select this option)
Sunday: **Yes** (Select this option)

On this page ...

- [Define a New Schedule](#)
- [Manage User Schedules from the Vision Client](#)



User Schedules

[Watch the Video](#)

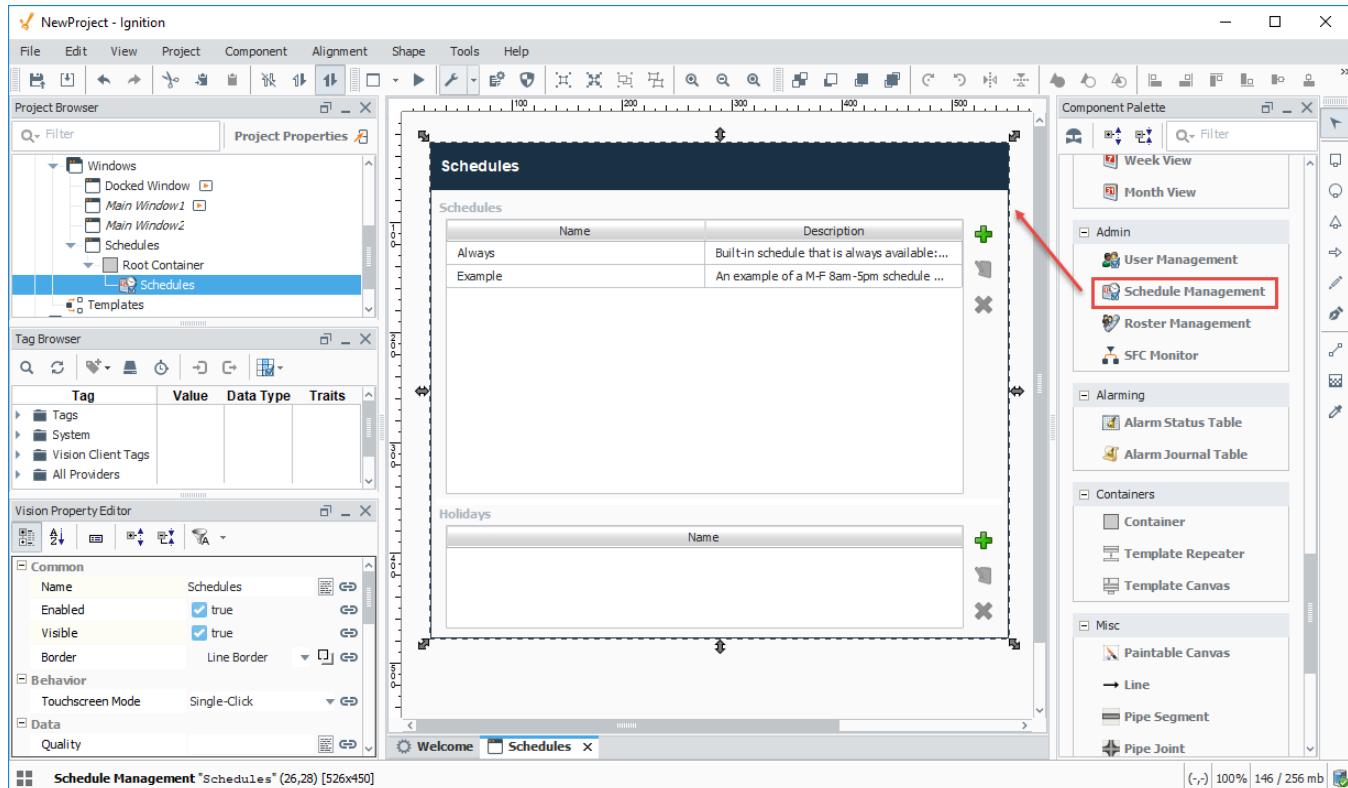
5. Click **Add Schedule**.

The following feature is new in Ignition version **8.1.25**
[Click here](#) to check out the other new features

Note: If you want to ensure users are unable to log in when their assigned schedule is not active, you can select the **Schedule Restricted** option for the user's corresponding User Source. The user will then be denied any login access for attempts made outside active hours.

Manage User Schedules from the Vision Client

There are a few ways to manage user schedules from the Vision client. The first is to use the **Schedule Management** component on a window. This component allows you to quickly and easily manage the schedules from the Vision client.



For more granular control, you may instead want to use scripting to manage the schedules. This may offer a more granular control at the click of a button. There are a number of system functions that allow you to create, read, edit, and delete schedules or holidays from a user source using

scripting. (i.e., [system.user.addSchedule](#), [system.user.getHoliday](#), etc.). To learn what system functions are available for user scheduling and holiday scripting, refer to the [System Functions](#) in the Appendix.

```
# This code creates a new schedule by using an old schedule but setting observe holidays to true.  
mySchedule = system.user.getSchedule("WeeklySchedule")  
if mySchedule != None and mySchedule.getType() == "basic schedule":  
    mySchedule.setObserveHolidays(True)  
    mySchedule.setName("NewWeeklySchedule")  
    system.user.addSchedule(mySchedule)
```

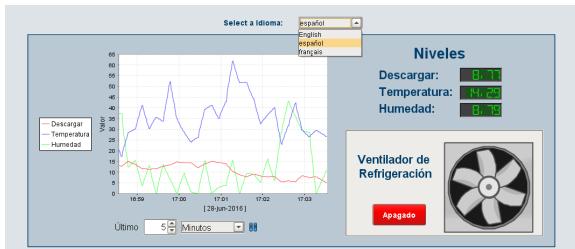
Related Topics ...

- [On-Call Rosters](#)
- [User Notifications](#)

Localization and Languages

Translating Ignition

Localization allows you to translate text into multiple languages in a project for display on client screens. The localization feature allows users located in different countries to set their default language so client screens can be displayed in their native language. The user can easily choose which language their Client displays with a click of a button! Text contained in components, that have their properties correctly mapped with localization, will change to reflect the language the user has selected. There is a special [Language Selector](#) dropdown list that displays the available translation mapping options for these translations.

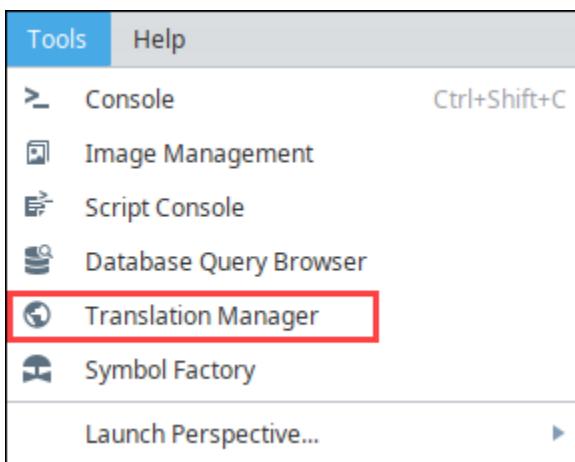


On this page ...

- Translating Ignition
- Translation Manager
- Translation Database and Term Lookup
 - Quickly Identify Fields to Translatable Terms
- Previewing Translations

Translation Manager

All translations are stored centrally in the Gateway and are distributed to each Client and Designer. All projects share the same translations, and those translations can be used in other locations, such as Gateway scripts, and alarm messages. The full translation database can be viewed and edited in the **Translation Manager** tool located on the Tools menu the Designer.



You can add new terms and phrases to the Translation Manager. Make your phrases as long or short as you like with single words or whole paragraphs. It allows users to easily [import and export](#) translatable components and share them across projects and with other users. It also has the ability to define new languages.

A central term database and support are automatically built in to all component text properties, as well as other text-based properties, such as alarm messages. Aside from defining translations for terms, there is usually no other work needed to take advantage of the translation system.

Translation Manager			
Languages	Translation Terms		
(All)	Key	English (Alternate)	Spanish
<input checked="" type="checkbox"/> English	+ Auto		Automatico
<input checked="" type="checkbox"/> Spanish	+ brake		freno
	+ Brake Motors		Motores de Freno
	+ Days		Dias
	+ Hand		Manual
	+ Humidity		Humedad
	+ Last		Ultimo
	+ Minutes		Minutos
	+ Off		Apagado
	+ Seconds		Segundos
	+ Tank		Deposito
	+ Temperature		Temperatura
	+ Value		Valor

Translation Database and Term Lookup

Terms are referenced in the translation database using direct string comparison. Ignition has special term **Keys** that are used to identify when to translate text. The base term may be a user-readable string such as "Start", or any special code you want such as "START_COMMAND" or "#start", which would have an alternate translation for English defined as "Start." This way you can quickly identify which terms will be translated when working in the Designer. To keep things simple, the English key is automatically used if no match is found for the requested language.

Quickly Identify Fields to Translatable Terms

You can quickly and easily identify, and translate all terms by selecting a container from your Project Browser and opening the Translatable Terms panel.

Note: If you don't see the Translatable Terms panel, go to the menubar and select **View > Panels** and check **Translatable Terms**. Here, you can see all terms for each component in the container, and if each term has an associated translation for the requested language.

Translatable Terms

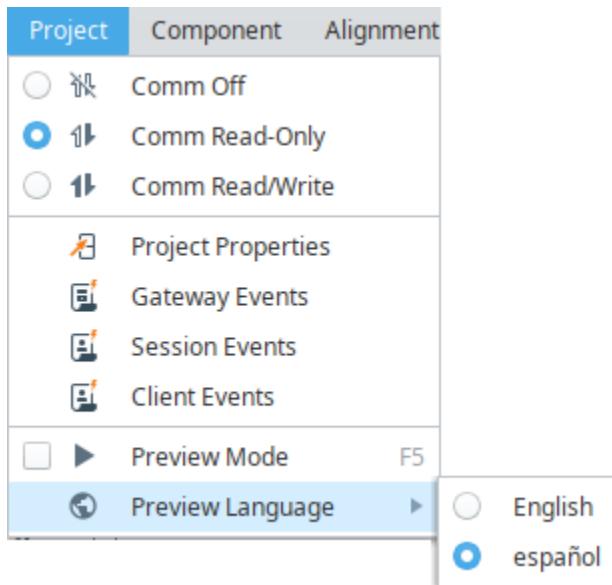
The dialog box has a toolbar with icons for file operations and search. Below the toolbar is a section titled "Component Terms". A table lists terms with their descriptions and translations:

Description	Key	Spanish (Com...)	Spanish (Glob...
Confirm Text	Are you sure?		
selectedText.1	Hand		Manual
unselectedTex...	Hand		Manual
selectedText.2	Off		Apagado
unselectedTex...	Off		Apagado
selectedText.3	Auto		Automatico
unselectedTex...	Auto		Automatico
Confirm	Confirm		

Selected Component: Multi-State Button

Previewing Translations

You can easily preview your translations in the Designer. To change the preferred language for Preview Mode, use the **Preview Language** menu item under the **Project** menu in the Designer. When you put the Designer into Preview Mode, everything on your windows will be automatically translated. Using the Language Selector Component or scripting can temporarily change the current language, and the language will revert back any time Preview Mode is disabled. The Designer will remember your choice, even when the Designer is restarted.



In This Section ...

Creating Translation Lists

Project Translation Overview

Project translation is also known as Client localization. This is how the user can choose their language-of-choice for the project they are working with. When the user selects a language, the Client will change its text in accordance with the project's translation settings. Translations work with all components (i.e., Labels, Charts, or Multi-State Buttons).



Translation Notes

- Text within Vision Table components (those under the "Table" group) can not be translated. However the Vision Alarm Status Table and Vision Alarm Journal Table components do support translations.
- For Text elements that have bindings, the translation occurs after the binding is evaluated.

The Translation Manager allows you to view all the terms from all your projects in the global translation database. You can add, edit, and remove terms, and provide translations, as well as allow users to easily import and export translatable components. The translation database shares terms across all projects and with other users.

Setting up New Languages, Terms, and Translations

Setting up new languages and adding terms for translation is pretty easy. For every element of text on the screen, you can provide a translation in the desired language. When a user logs into a project, they can choose their preferred language, and all the text will be displayed in their preferred language. They can also have their preferred language defined in [user settings](#) so when a user logs into the Client, their preferred language is automatically enabled.

Component vs Global Terms

There are two types of terms used in the translation system: Component and Global. Both behave a little differently, and it's important to know the difference.

- **Component Terms** - Component level terms are specific to that one component, and translations are added to the translation system using the Translatable Terms Panel. Any component level translation will automatically take precedence in the event there is a matching global term.
- **Global Terms** - Global terms and translations are available in both the Translatable Terms Panel and in the Translation Manager's global database. The global translation database provides language translations on all components, in all windows, and on all projects. You can set up global translations so that anywhere a term is used within your project, it will get replaced with that global translation. Although, there are times, when a component might have a different context than the global translation. For this reason, you may want to create a component level translation to take precedence over the global translation. For example, some users may work in a specific window of a project, and what is typically called a 'Tank,' this group of users refer to it as a 'Barrel.' You can create a component level translation for 'Barrel' to take precedence on the root container of a window. This way the users see the term in the context that is familiar to them.

Add a New Language

1. In Designer, right click in your workspace and select **Translations** from the dropdown. The Translatable Terms window will open.
2. Select the icon, and click **Add Language**. By default, the English language is defined.

On this page ...

- [Project Translation Overview](#)
- [Setting up New Languages, Terms, and Translations](#)
 - [Component vs Global Terms](#)
 - [Add a New Language](#)
 - [Add Global Translation Terms](#)
 - [Add Component Level Translation Terms](#)
- [Translation List Import and Export](#)
 - [Export Terms from a Window](#)
 - [Import Terms](#)



INDUCTIVE
UNIVERSITY

Introduction to Project Translation

[Watch the Video](#)

The screenshot shows the 'Translatable Terms' window. At the top, there's a toolbar with icons for file operations and search. Below it is a component selection dropdown where 'Spanish' is selected. A red box highlights the 'Add Language' button next to the dropdown. The main area is a grid table with columns: Component, Description, Spanish (Component), and Spanish (Global). The table contains several entries, such as 'Label.Text' with 'Motors' and 'Motores de Freno', and 'Multi-State Button' with various states like 'Hand', 'Off', 'Auto', and 'Confirm'. The bottom of the window shows the selected component as 'Root Container' and has tabs for 'Properties' and 'Translations'.

3. Select a new language from the dropdown.

The screenshot shows the 'Add new language' dialog box. It has a dropdown menu at the top with 'French' selected. A red box highlights the 'French' option in the list. Other languages listed include Finnish, Friulian, Fulah, Galician, Ganda, and Georgian.

Add Global Translation Terms

Creating global translations is always a good first choice over the component level translation because most of your translated terms will be used across all projects unless you want one particular term to be different from the global translated term.

There are a couple of ways to add terms to the Translation Manager. You can add new terms directly to the Translation Manager or by selecting terms from the Translation Terms Panel.

Let's add some terms directly to the Translation Manager.

1. From the menu bar at the top, go to **Tools > Translation Manager**.
2. The Translation Manager window will open, and you can view all the translatable terms. This example has terms that were already added and translated.
3. On the right side of the window, there is a list of icons where you can add, delete, import, export, and edit the Translation settings.

Translation Manager

Languages: (All) English Spanish

Translation Terms

Key	English (Alternate)	Spanish
+ Auto		Automatico
+ brake		freno
+ Brake Motors		Motores de Freno
+ Days		Dias
+ Hand		Manual
+ Humidity		Humedad
+ Last		Ultimo
+ Minutes		Minutos
+ Off		Apagado
+ Seconds		Segundos
+ Tank		Deposito
+ Temperature		Temperatura
+ Value		Valor

4. To add a new global term, click on the Add  icon. Enter the term (Tanks), and click **Save**.
5. Double click on the row of the new term, and under Spanish translation, enter the word "**Deposito.**" Click **Save**.

Translation Manager

Languages: (All) English Spanish

Translation Terms

Key	English (Alternate)	Spanish
+ Last		Ultimo
+ Minutes		Minutos
+ Off		Apagado
+ Seconds		Segundos
+ Tank		Deposito

Original Term
Tank

English (Alternate)

Spanish translation
Deposito

Save

6. To view all the global terms added to a project, go to **Tools > Translation Manager**. If you want to provide a translation to a term that has not been translated, double click on the row of the term, enter the translation, and click **Save**.

Translation Manager

Translation Terms

Key	English (Alternate)	Spanish
+ Auto		Automatico
+ brake		freno
+ Brake Motors		Motores de Freno
+ Days		Dias
+ Hand		Manual
+ Humidity		Humedad
+ Last		Ultimo
+ Minutes		Minutos
+ Off		Apagado
+ Seconds		Segundos
+ Tank		Deposito
+ Temperature		Temperatura
+ Value		Valor

7. Now test it out. Drag a **Language Selector** component to your window. Go into **Preview Mode**, and toggle the languages. You will see the languages for the text switch between English and Spanish for these two components.

The image consists of two vertically stacked screenshots of a user interface. The top screenshot shows a Language Selector component with the dropdown menu open, displaying 'español' at the top, followed by 'English' and 'español' again. Below the dropdown are three buttons: 'MANUAL' (gray), 'APAGADO' (red), and 'EN' (gray). At the bottom is a button labeled 'Motores'. The bottom screenshot shows the same interface after switching to English. The Language Selector dropdown now displays 'English' at the top, followed by 'español'. The buttons below have changed to 'Manual' (gray), 'Apagado' (red), and 'Automatico' (gray). A red box highlights the Language Selector dropdown in the bottom screenshot.

Add Component Level Translation Terms

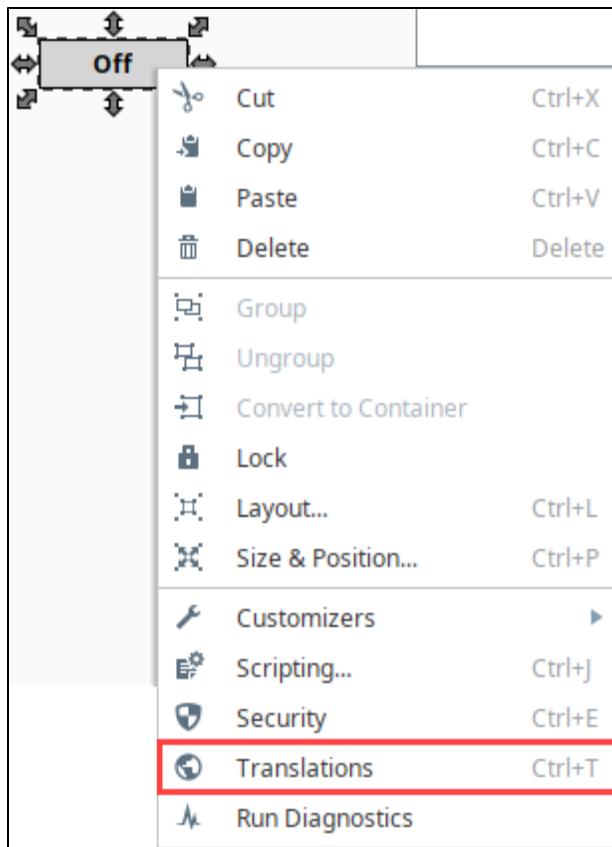




Component level translation terms are only available in Vision.

What's nice about adding component translation terms is that the Translatable Terms panel provides all the terms for that component on the window. All you have to do is add the translations. The Translatable Terms Panel is a central place to manage all the translations associated with a specific component. One important reason to use a component level term is if it is going to be different from the global term. Just remember, that because a term is translated at the component level doesn't mean that other components get to use it. Typically, it is used for that component only.

1. In Designer, drag a **Multi-State Indicator** from the Component Palette in to your workspace.
2. Right click on the component and scroll down to **Translations**. The Translatable Terms panel will open and display all the terms for the specific component. Alternatively, you can right-click on the Root Container of a window and select **Translations**. The **Translatable Terms Panel** will display all the terms for all the components on the window.



Translatable Term Options

In the Translatable Terms Panel, you have the option of adding your translation to the Component, Translation Manager global database, or marking it for translation at a later time. You'll notice under the Component Terms section of the Translatable Terms window that there are four columns: Component Description, Key, Component translation, and Global translation.

You'll see below, all the terms for the Multi-State Button, and none of the terms have translations added. If you don't see the language you want your terms translated to, click the Earth icon , and select the language (i.e., Spanish). You will then see a column for your component translation and global translation.

Translatable Terms

The screenshot shows the 'Translatable Terms' panel with the title 'Component Terms'. It contains a table with four columns: 'Description', 'Key', 'Spanish (Co...)', and 'Spanish (Gl...)'. The rows represent different text elements and their states:

Description	Key	Spanish (Co...)	Spanish (Gl...)
Text	Off	Pending	
Style State 0.Text	Off	Pending	
Style State 1.Text	Auto	Pending	
Style State 2.Text	Manual		
Style State 3.Text	Fault		
Style State 4.Text	Fault		
Style State 5.Text	Pending		

Selected Component: Multi-State Indicator

- To add translations for your component terms, double click the row of your term (i.e. ON). In the **Spanish Component** box, enter the translation (i.e. EN). You also have the option of creating a global translation on the Translatable Terms Panel as well. Then, click the **Back** link, and repeat this step for all your terms.

The screenshot illustrates the translation process. On the left, the 'Translatable Terms' panel shows a list of component terms. The row for 'Style State 1.Text' is highlighted with a red box. On the right, a secondary panel titled 'Back' displays two text fields: 'Original Term' containing 'Auto' and 'Spanish (Component)' containing 'Automatico'. The 'Automatico' entry is also highlighted with a red box.

- Once you entered the translations for all your terms, you will see them on the Translatable Terms Panel in the Spanish (Component) column. If a term doesn't have a translation, you can select the term and mark it for translation by clicking the **Earth with a green plus icon**. This adds the term(s) to the translation database for someone else to add the translation.

In this example, you'll notice several things about each of the terms so look closely at them. You'll find examples of both component and global level terms, a component level term, a global term, a component term that overrides a global term, and terms marked for translation.

Translatable Terms

Component Terms Mark the selected term(s) for translation (adds them to the translation database)

Description	Key	Spanish (Com...)	Spanish (Glob...
Text	Off		Pending
Style State 0.T...	Off		Pending
Style State 1.T...	Auto	Automatico	Pending
Style State 2.T...	Manual		
Style State 3.T...	Fault		Pending
Style State 4.T...	Fault		Pending
Style State 5.T...	Pending		

Selected Component: Multi-State Indicator

- To test it out, drag a **Language Selector** component from the Component Palette to your workspace. Go to **Preview Mode**, select Spanish from the dropdown to see the text change from English to Spanish for these two components.



Translation List Import and Export

The ability to import and export terms and translations can be extremely important particularly to a large project. It allows you disseminate terms and translations rapidly, send them to a third party for translation, and keeps all related projects up-to-date.

Export Terms from a Window

Let's export some terms from the Translation Manager. In this example, you can see all the terms that were added to the Translation Manager. You can export selected terms, or all terms from the Translation Manager.

- Go to Tools > Translation Manager. Select all terms or only specific terms to export. Click on the **Export Terms** icon on the right side of the **Translation Manager** window.



Import and Export Translations

[Watch the Video](#)

Translation Manager

Languages	Translation Terms		
	Key	English (Alternate)	Spanish
<input checked="" type="checkbox"/> (All)	#HAND	Hand	Manual
<input checked="" type="checkbox"/> English	#MyText	<html>This is some long text and I...</html>	<html>Este es un texto largo yquier...
<input checked="" type="checkbox"/> Spanish	#OFF	Off	Apagado
	#ON	On	en
	#Tank1Name	Tank	Tanque 1
	#Tank2Name	Tank	Tanque 2
	#Auto		Automatico
	#Barrel	Barrel	Barril
	#Button		Tanques
	#Check Box		Recharzar
	#Confirm		
	#Controllers		Controladores
	#Goodbye		Adios
	#HAND		MANUAL
	#Hand		Manual
	#High		
	#Label		Tanque
	#Motors		Motores
	#OFF		APAGADO
	#Off	Off	Apagado
	#ON		EN
	#Sine0		Nivel
	#Sine1		Temperatura
	#Tab 1	Home	Casa
	#Tab 2	Office	Oficina
	#Tab 3	School	Estudiar
	#Tank		Tanque
	#Tanks		Tanques
	#Temperature	Temperature	Temperaturas
	#Toggle Button	Motors	Motores

2. The Exported Terms window will open. Enter the following information:

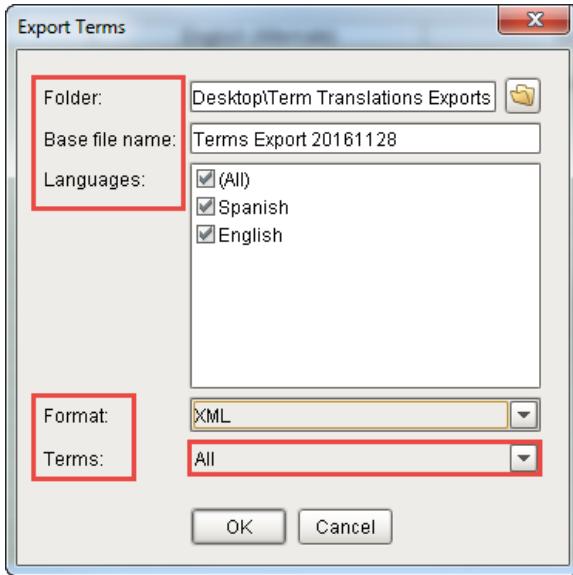
- a. Select the **Folder** location where you want to put your exported terms.
- b. Enter the **Base file name**. (File name of your exported terms).
- c. Select the appropriate **Languages**. You can have multiple languages selected depending on the number of language translations you are using (i.e., English, Spanish, and French). This example exports two files: one for the English translation, and one for the Spanish translation.
- d. From the **Format** dropdown list, select either **PROPERTIES** or **XML**. Both formats behave the same, the only difference is that the XML format supports UTF-8 encoding directly. A person can now translate the XML file directly and import it back into Ignition's Translation Manager.
- e. From the **Terms** dropdown list, choose either **Selected** or **All**.



For Specific Terms

If you want to select specific terms to export, choose '**Selected**', otherwise, '**All**' terms will be exported from the dropdown list at the bottom of the window.

- f. Click **OK**.



```

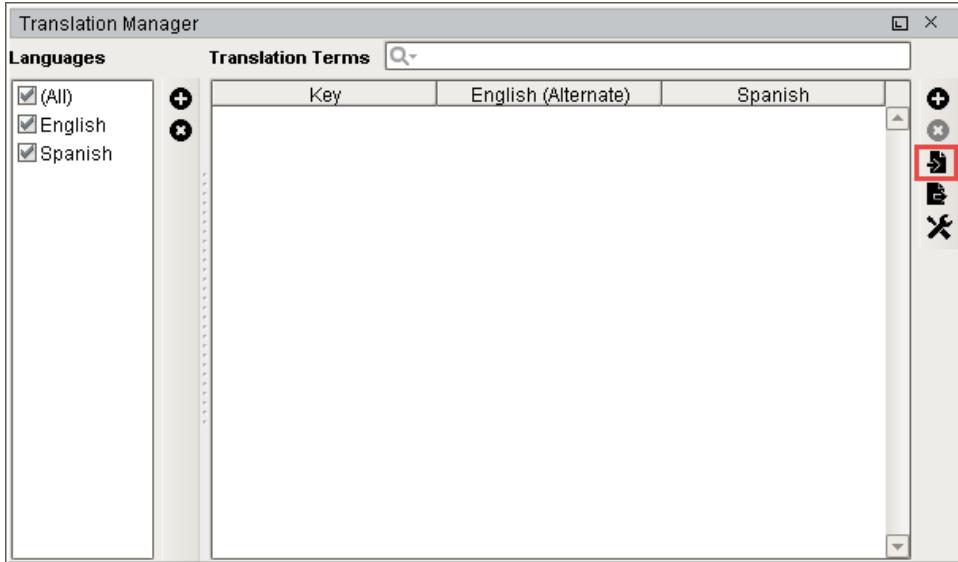
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE properties SYSTEM "http://java.sun.com/dtd/properties.dtd">
- <properties>
  <comment>Locale: en</comment>
  <entry key="#Sine1">Temperature</entry>
  <entry key="OFF"/>
  <entry key="Sine0"/>
  <entry key="Tanks"/>
  <entry key="#Tank1Name">Tank</entry>
  <entry key="Goodbye"/>
  <entry key="Toggle Button">Motors</entry>
  <entry key="Controllers"/>
  <entry key="#OFF>Off</entry>
  <entry key="Off>Off</entry>
  <entry key="Tab 3">School</entry>
  <entry key="#mytext"><html>This is some long text and I want it
    to wrap</entry>
  <entry key="Tab 2">Office</entry>
  <entry key="Tank"/>
  <entry key="Tab 1">Home</entry>
  <entry key="#HAND">Hand</entry>
  <entry key="Barrel">Barrel</entry>
  <entry key="Welcome"/>
  <entry key="Button"/>
  <entry key="Auto"/>
  <entry key="#ON">On</entry>
  <entry key="Hello"/>
  <entry key="Label"/>
  <entry key="High"/>
  <entry key="Hand"/>
  <entry key="Check Box"/>
  <entry key="#Tank2Name">Tank</entry>
  <entry key="HAND"/>
  <entry key="Motors"/>
  <entry key="Confirm"/>
</properties>

```

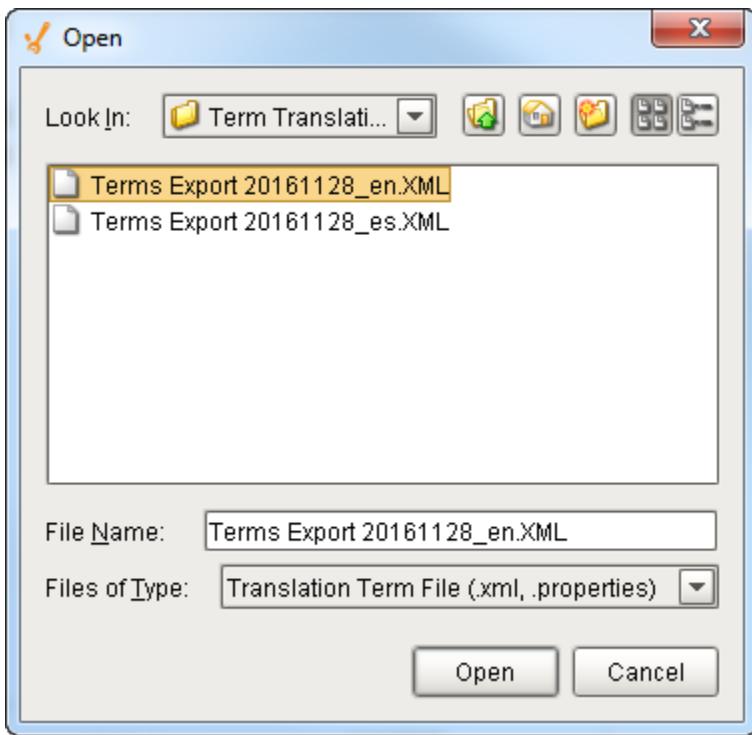
3. Go to the folder location and check to see if your exported files are there. You will see **one file exported file for each language**. Open the English exported file and verify that all the terms you intended to export are there.

Import Terms

1. Go to Tools > Translation Manager, and select the **Import Terms** icon from the right side of the Translation Manager window.



2. Navigate to the folder where your exported files reside, and choose the English file from the **Open** window, and click **Open**.

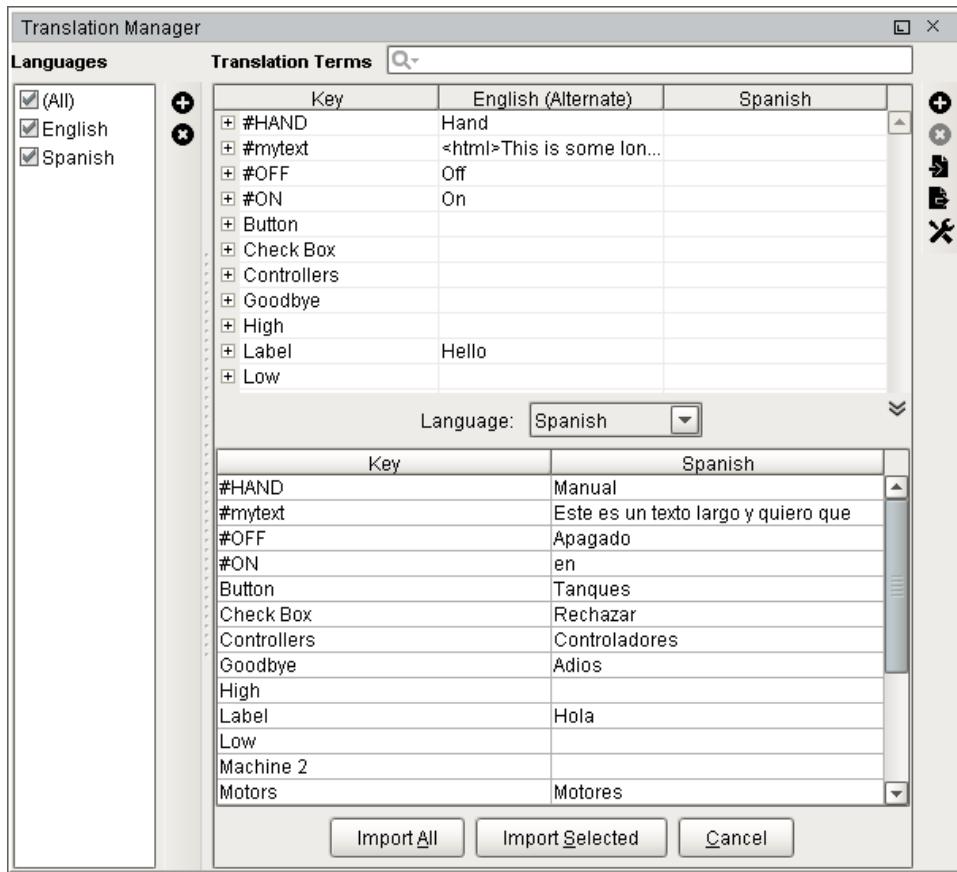


3. You have the option of importing selected terms or all terms. Click either **Import All** or **Import Selected** depending on what terms you want to import.

Repeat this step to import each translation file. This example imports the Spanish translation file.

Imported Terms Overwrite Existing Terms

Be aware that existing terms will be overwritten by the imported term.



4. Check to see that all your terms were imported into the Translation Manager by going to **Tools > Translation Manager**.

Related Topics ...

- [Switching the Current Language](#)
- [Localization Best Practices](#)

Switching the Current Language

Language Selection

Because Ignition has multiple visualization systems, each system has separate ways to switch the current language.

Vision

In Vision, the language can be selected using the language selector dropdown before logging in to the client, or can be added to the client in the form of the language selector component.

Perspective

In Perspective, the language is selected by changing the locale of the session by modifying the session.props.locale. This can be written to using a script when the session first opens, or something can be bound to it that the user can change.

Expression and Scripting Functions

You can look up translations using the following functions:

- Expression Function
translate()
- Scripting Functions
system.util.translate()
system.util.modifyTranslation() scripting functions.

Translations are matched by looking for the base language value in the translation database. This is especially useful for message boxes and other warnings or errors that you show in your scripts.

System Considerations During Translation

The inherent operating system may affect Ignition's ability to provide a complete translation in certain circumstances.

For example, suppose a user selected the Spanish option from their Language Selector component. An English to Spanish translation term exist for the word "Information" as well as "Start." The message box script takes two parameters: one for the content while the other parameter is optional. If the second parameter is not included in the function then the resulting title of the message box will include the English word "Information." The following code is executed on a button clicked event handler.

```
message = system.util.translate("Start")
system.gui.messageBox(message)
```

What results is a message box with a translated content while the title remains in English. In addition, the accept button is translated as well even though the English word for "OK" does not exist in the Translation manager. This is due to the system level translations that exist for inherent language support.

A similar occurrence exist with the file open dialog window. The following code is executed on a button clicked event handler.

```
system.file.openFile()
```

The result is a translation for the title as well as the open and cancel buttons while the references to the "File Name" and the "Files of Type" remain in the operating system's language and thus are beyond the scope of Ignition's ability to translate.

Similar occurrences will appear in print and error dialog boxes.

On this page ...

- [Language Selection](#)
 - [Vision](#)
 - [Perspective](#)
- [Expression and Scripting Functions](#)
- [System Considerations During Translation](#)

Localization Best Practices

Best Practices

Before you begin using multiple languages in your projects, we thought we would start you off with a few best practices for using project terms and translations.

As you already know, there are global level translations which are available to all components and text elements in both the Translatable Terms Panel and the global database, and component level translations that are only specific to the component and only available in the Translatable Terms Panel. It's good to understand a little about their behavior before investing a lot of manpower setting up your translations only to find out down the road that you might set them up differently after you become more familiar with them.

Global vs Component Terms

We recommend using global terms over component level terms because this way you will only have to add your term once since global terms are shared across all windows and projects. This will save you a little time from having to add the same term again if you discover that a component level term should extend beyond the specific component and shared across all projects. The only time you would want to use a component level term, is if you want it to be different from the global level term. This ensures that the component level term will override the global term.



Where the component level term takes precedence

The [component level term](#) always takes precedence when there is a matching global term.

Using Codes

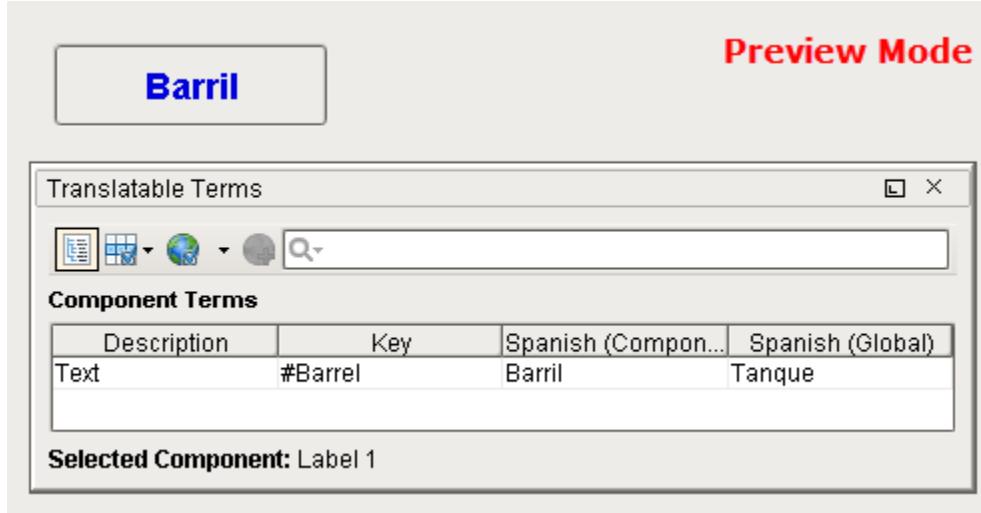
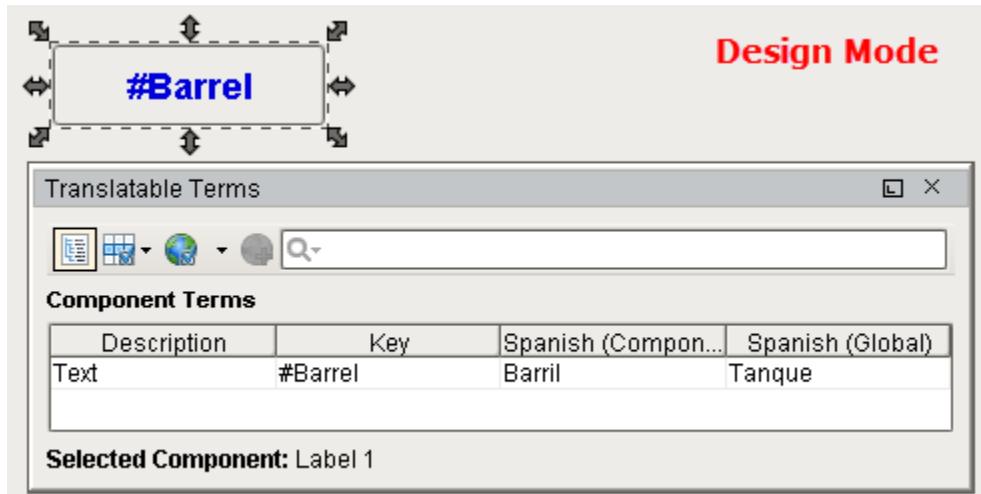
It's a good idea to use codes for the Key field of your Terms for any descriptions that you want to translate.(i.e: #introduction or #welcome_screen_info) so your global term won't ever accidentally translate on another window or component. (The '#' is just for convention, and is not a special character). Using codes works particularly well for long text strings such as introduction paragraphs. One thing to be very careful of is if you edit the codes in a minor way, it could potentially break the translation.

Here are a couple of examples:

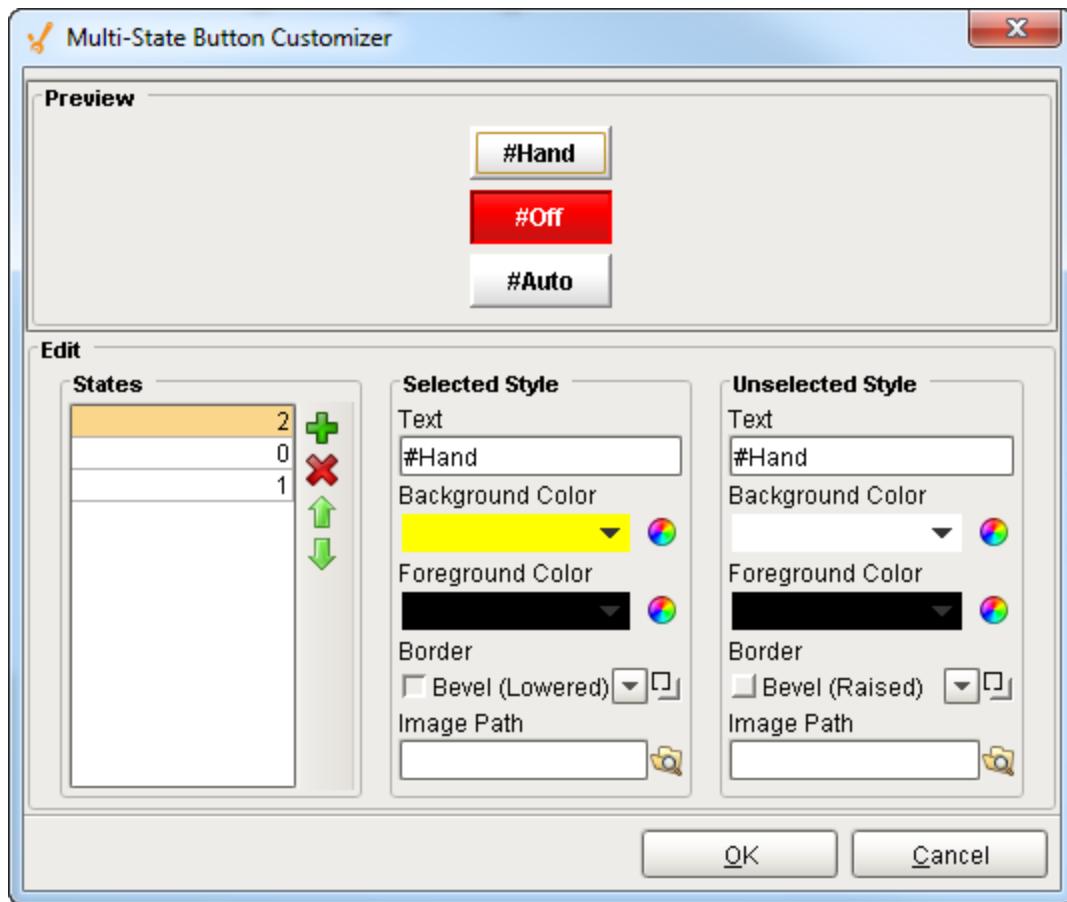
The first example is of a Label component with a Key field of "**#Barrel**." You'll notice the Label has global translation of "**Tanque**" and a component level translation of "**Barril**." By using the "#," the component level translation of "Barril" was not overwritten by the global translation of "Tanque" by using the "#." Here is what the translation looks like in **Design Mode** and **Preview Mode**.

On this page ...

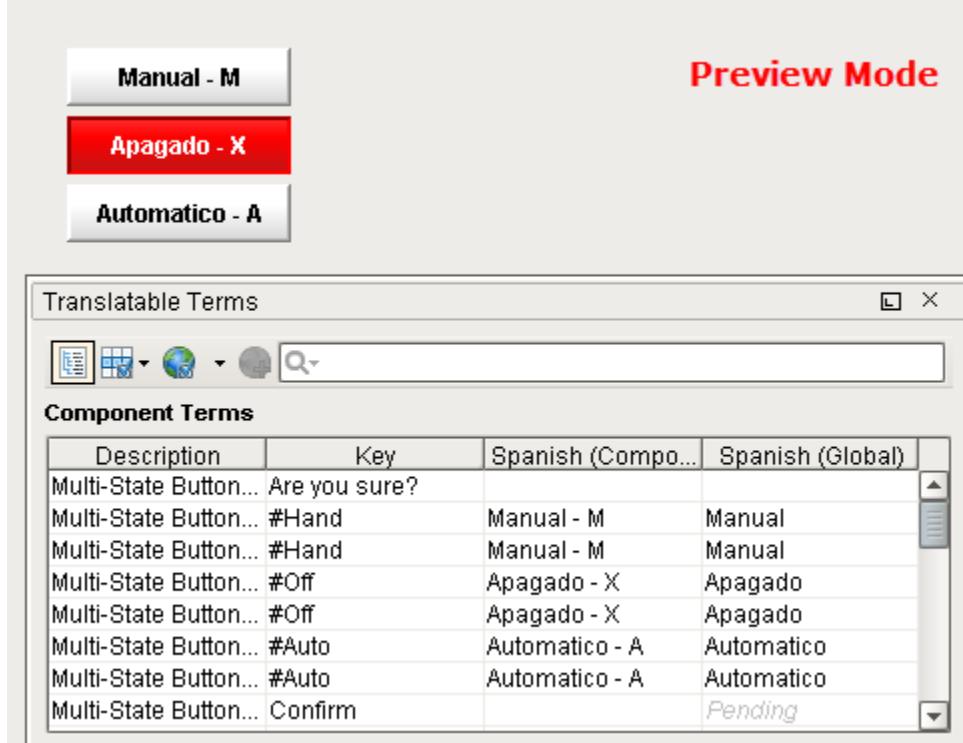
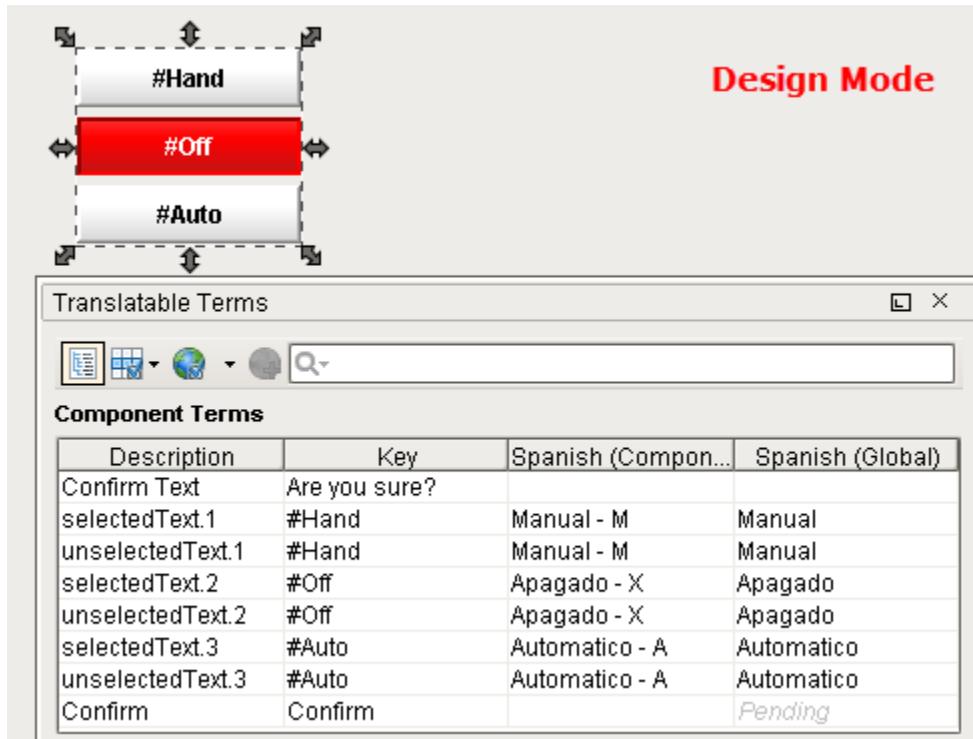
- [Best Practices](#)
 - [Global vs Component Terms](#)
 - [Using Codes](#)
 - [Using HTML for Text Wrapping](#)



The "#" can also be used on components for text that you don't want accidentally translated. This example uses the Multi-State Button component. The only way to change the states / words on the buttons (i.e., Hand, Off, Auto) is to use the Multi-State Button Customizer. Add the "#" to the text on each of the three Keys so they don't get translated by global terms that have already been set up.



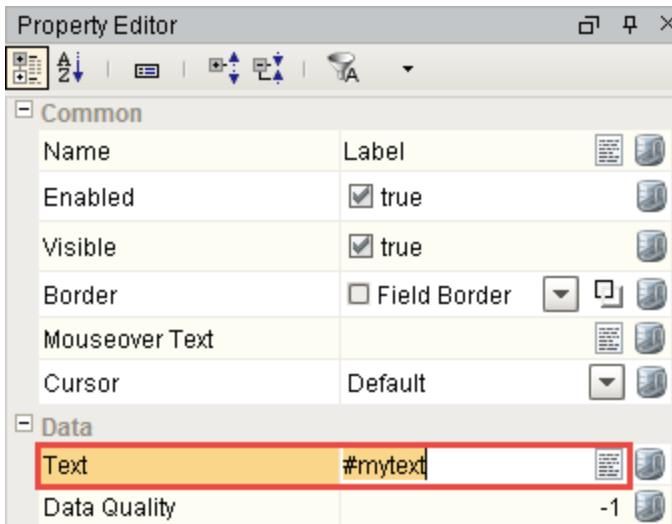
Here's what the translations look like in **Design Mode** and **Preview Mode**.



Using HTML for Text Wrapping

Another good practice is to use HTML to wrap text within your component's buttons, labels, etc.,. You can use it with translating terms that contain lots of text. HTML is just a good way to wrap text to fit within a label and button components.

1. In the **Property Editor** of a Label component, go to the **Text** property and give the Text a unique name (i.e., #Mytext).



2. Right click on the **Label** component, and click on **Translations**. Double click on the row for Mytext. Enter the text as shown below for both the **Spanish Global** and English **Global translations**. Make sure you precede your text with "<HTML>." Click the **Back** link.

```
Mytext

#English
<html>This is some long text and I want to wrap it

#Spanish
<html>Este es un texto largo y quiero que se envuelva
```

Design Mode

The screenshot shows the 'Translatable Terms' dialog in 'Design Mode'. At the top, there's a preview area with a dashed border containing the text '#mytext'. Below it, the title 'Design Mode' is displayed in red. The main area is titled 'Component Terms' and contains a table:

Description	Key	Spanish (Comp...)	Spanish (Global)
Text	#mytext	<html>Este es un texto lar...	<html>Este es un texto largo y quiero que se envuelva

Below the table, there's a tree view under 'Translatable Terms' with a 'Back' link. The 'Original Term' is '#mytext'. The 'Spanish (Global)' section and the 'English (Global)' section both contain the same long HTML string, which is highlighted with a red box.

3. Once in **Preview Mode**, your long text will be wrapped!

The screenshot shows the 'Preview Mode' interface. On the left, there's a box containing the text 'Este es un texto largo y quiero que se envuelva'. To the right of the box, the words 'Preview Mode' are displayed in red. The text inside the box is wrapped at the word 'quiero'.

Translating Built-in Terms

There are several built-in terms used by Ignition that are translatable. This page contains the Keys that can be used to translate these terms.

Built-in Access Denied

When [Security Settings](#) are configured on a window, the only restriction is **Do Not Open**, which prevents the window from opening based on User Roles. The text in this message box is not normally configurable, but can be altered via translations using the following steps.

On this page ...

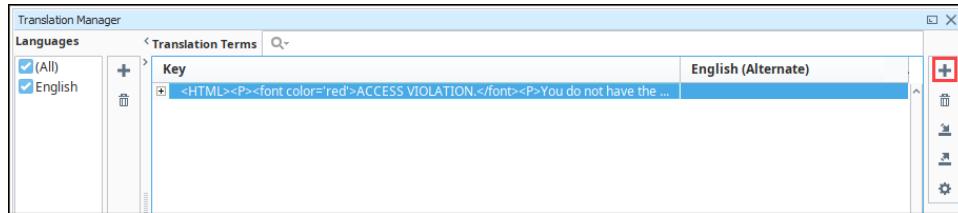
- [Built-in Access Denied](#)
- [Screen Locked Message](#)

1. Go to the top menubar in the Designer and select **Tools > Translation Manager**. This opens the Translation Manager window.
2. Copy the **Key** from the code block below, click the plus  icon and paste it in the **Add Item** field. Click **OK** and you'll notice the code will be displayed in the **Key** field.

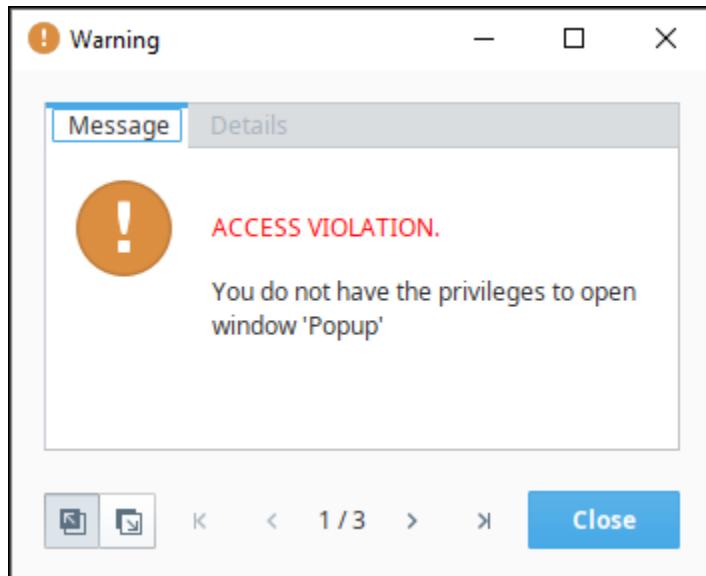
The line below should be used as the Key.

Access Violation - Key

```
<HTML><P><font color='red'>ACCESS VIOLATION.</font><P>You do not have the privileges to open window '%s'
```



3. Save your project.
4. When a user without the proper permission attempts to access the popup from the client, they will get the following message.



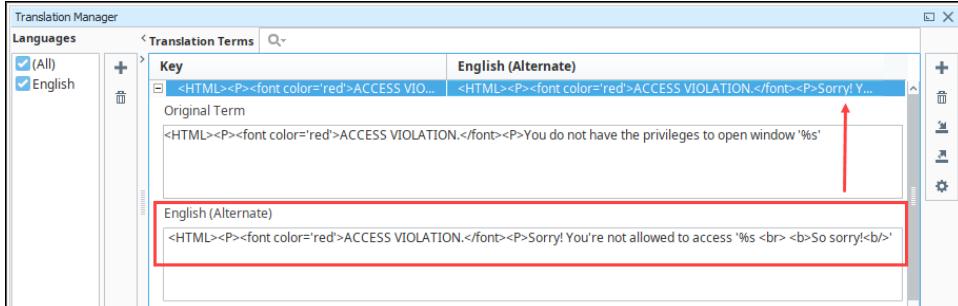
Using the same example from above, we modified a version of the text that appeared in the message box. We used the **English (Alternate)** as the language where the name of the window is passed into the message box via the '%s' substring.

1. Copy the new message from the code block below.

Access Violation - Translated Term

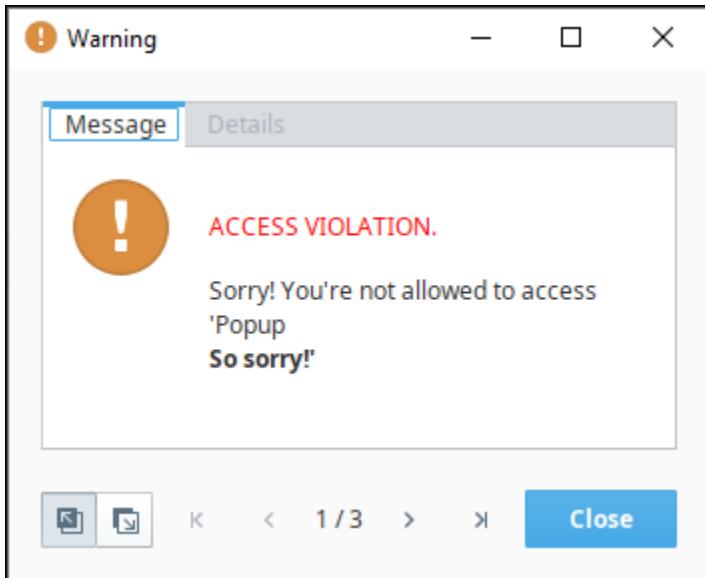
```
<HTML><P><font color='red'>ACCESS VIOLATION.</font><P>Sorry! You're not allowed to access '%s <br> <b>So sorry!<b/>'
```

2. In the Transaction Manager window, double click the row for the **Key** and paste it in the **English (Alternate)** field. Click **Save**.



3. Save your project.

Now, our message box appears as the following when a user does not have the proper access.



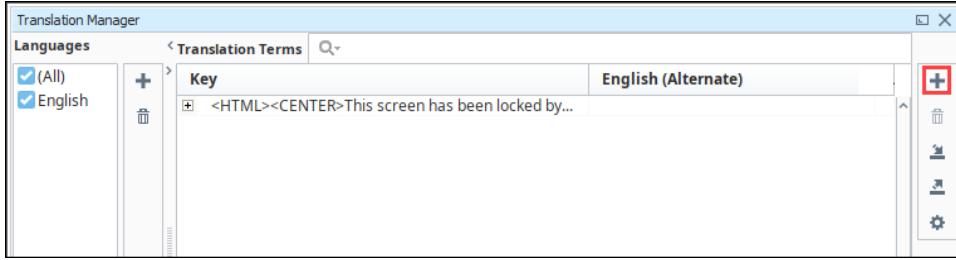
Screen Locked Message

The built-in [system.security.lockScreen](#) function in a [client timer script](#) can be used to lock the screen. In some cases, you may wish to change the text on the Screen Locked window. Use the following steps to change the text.

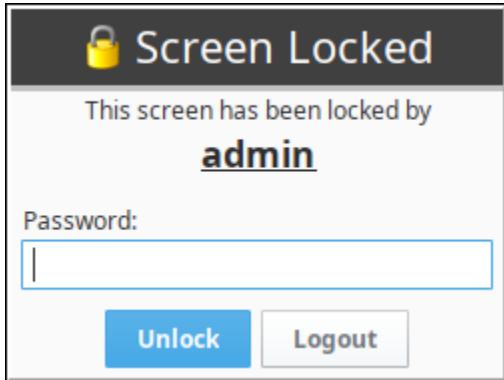
1. Open the Transaction Manger window.
2. Copy the **Key** from the code block below. Click the plus icon and paste it in the **Add Item** field. Click **OK** and you'll notice the code will be displayed in the **Key** field.

Screen Locked - Key

```
<HTML><CENTER>This screen has been locked by<BR><font size=+1><b><u>%s</u></b></font>
```



- Save your project. Now the Screen Locked window will appear as follows.



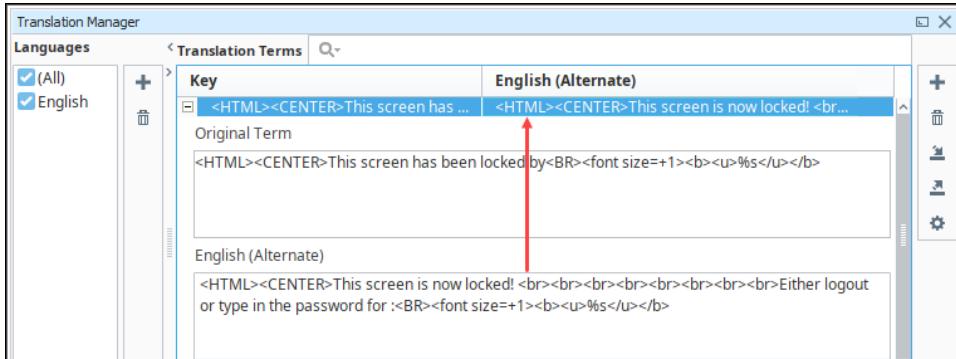
Building on the previous example, we can use the following key to translate the message on the Screen Locked window. The '%s' substring will reference the username of the user currently logged in. In this example, we added a few line breaks, which will make the Screen Locked window appear taller.

- Copy the updated text from the code block below:

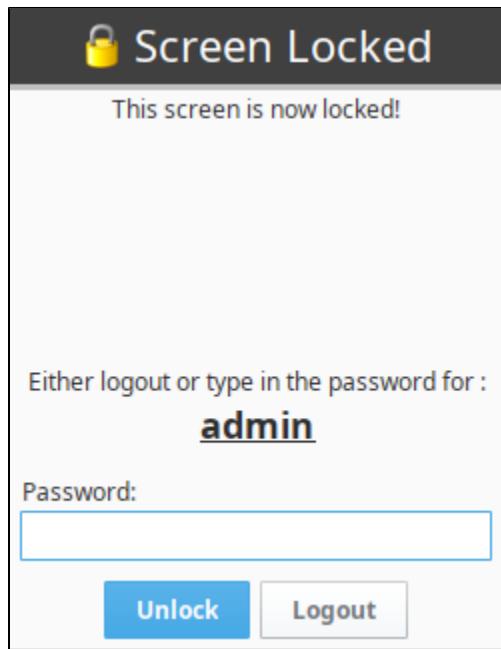
```
Screen Locked - Translated Term

<HTML><CENTER>This screen is now locked! <br><br><br><br><br><br>Either logout or type in the password for :<BR><font size=+1><b><u>%s</u></b></font>
```

- In the Translation Manager window, click the row for the **Key** and paste it in the **English (Alternate)** field. Click **Save**.



3. Save your project. Now, the message on the Screen Locked window appears as follows.



Expression Language and Syntax

Expression Language

The expression language is used to define dynamic values for component properties and [expression tags](#). Expressions often involve one or more other values that are used to calculate a final value. In most cases, expressions only return a value.

The classic example for an expression is to change a temperature that is stored in Celsius to Fahrenheit in order to display it. Suppose you had a Temperature tag that was in Celsius. If you wanted to display that tag in Fahrenheit on a Label, you would use an expression binding on the label's text property using the following expression:

```
1.8 * {My Temperature Tag} + 32
```

Every time that the Temperature tag changes, the expression will re-calculate the value and push it into the Label's text property. Now lets say that you wanted to append °F to the end of the label so that the user knew the units of the temperature. You could simply use some string concatenation in your expression, like this:

```
(1.8 * {Some Folder/Some Tag} + 32) + " °F"
```

For more information see [Expression Bindings in Vision](#) or [Expression Bindings in Perspective](#).

On this page ...

- [Expression Language](#)
- [Syntax](#)
 - [Literal Values](#)
 - [Operators](#)
 - [Bound Values](#)
 - [Dataset Access](#)
 - [Collection and Mapping access](#)
 - [Expression Functions](#)
 - [Whitespace and Comments](#)
 - [Tag Paths](#)
- [Additional Examples](#)
 - [String Concatenation](#)
 - [Celsius to Fahrenheit](#)
 - [Format Date](#)
 - [Date Manipulations](#)
 - [Bit Functions](#)
 - [Switch](#)
 - [Checking Conditions](#)
 - [Tag History](#)

Syntax

As its name suggests, everything in the expression language is an "expression". This means that everything returns a value: 5 is an expression, so is 5 +1, and so are {MyTags/TankLevel} and {MyTags/TankLevel}+1. Expressions can be combined in many powerful ways. Lets take a look at how expressions are written.

More formally, an expression is any one of the following:

- Number
- Boolean
- String
- Bound Tag
- Bound property
- Function call
- Dataset access
- Equation involving any of these

Literal Values

Literal values are things like numbers, booleans, and strings that are represented directly in the language. In the expression language, numbers can be typed in directly as integers, floating point values, or using hexadecimal notation with a 0x prefix. Examples:

```
42  
8.456  
0xFFC2
```

Strings are represented by surrounding them with double or single quotes. You can use the backslash character to escape quotes that you want to be included in the string. Examples:

```
"This is a regular string"  
'This one uses single quotes'  
"This string uses \"escaping\" to include quotes inside the string"
```

In addition, the following escape characters are available:

Character	Description
\n	New line
\t	Tab
\r	Carriage return

```
// The words "Hello" and "User" will be placed on separate lines
"Hello\nUser"

// Each "\t" will inject a tab
"Lots\tOf\tSpace"
```

Boolean values are represented with `True` and `False` (without quotation marks). Note that boolean values are case insensitive, so `true` is the same as `True`. For simplicity, we recommend writing `True` and `False`, since Python uses similar casing for boolean values.

```
// Multiple ways to specify a true boolean value.
True
true
tRuE
TRUE
```

Operators

You can use these arithmetic, logical, and bit-shifting operators to combine expressions.

Operator	Name	Description
//	Comments	Allows for comments following this operator.
-	Unary Minus or Subtraction	If both preceded by a number, then returns a value by subtracting the operand right from the operand on the left of the operator: If preceded by anything else (or nothing, such as the start of the expression) and is followed by a number, this operator will negate the number. <pre>// This will return 6. 10 - 4 // This will return -10. - (15 - 5)</pre>
!	Not	Logical opposite of a boolean.
^	Power	Raises a number to the power of another number.
%	Modulus	Modulus or remainder of two numbers. <code>a%b</code> is the remainder of $a \div b$.
*	Multiply	Multiplies the number on the left of the operator by the number on the right of the operator.
/	Divide	Divides the number on the left of the operator by the number on the right of the operator.
+	Add or Concatenation	If both operands are numbers, this will add them together. Otherwise treats arguments as strings and performs concatenation. <pre>// This will return 10. 4 + 6 // This will return 'FirstSecond'. 'First' + 'Second'</pre>

		<pre>// This will return '2Alarms'. 2 + 'Alarms'</pre>
-	Subtraction	Subtracts the number on the right of the operator from the number on the left of the operator.
~	Bitwise NOT	Examines the bits of an operand, and performs logical negation: bits with a value of 0 become 1, and vice versa.
&	Bitwise AND	<p>Examines the bits of two operands, and performs a logical AND to each set, comparing the bits in each position from both sets. Returns true for any position where the bits being compared are true.</p> <pre>// 0101 // AND 0011 // = 0001</pre> <p>// Performs the bitwise operation above, resulting in the decimal 1. 5 & 3</p>
	Bitwise OR	<p>Examines the bits of two operands, and performs a logical OR to each set, comparing the bits in each position from both sets. Returns true for any position where either bit being compared is true.</p> <pre>// 0101 // OR 0011 // = 0111</pre> <p>// Performs the bitwise operation above, resulting in the decimal 7. 5 3</p>
xor	Bitwise XOR	<p>Examines the bits of two operands, and performs a logical exclusive OR to each set, comparing the bits in each position from both sets. Returns true for any position where only one of the bits are true.</p> <pre>// 0101 // XOR 0011 // = 0110</pre> <p>// Performs the bitwise operation above, resulting in the decimal 6. 5 xor 3</p>
<<	Left Shift	A signed bitwise left shift.
>>	Right Shift	A signed bitwise right shift.
>	Greater Than	Logical greater-than test between two numbers. Returns a boolean.
<	Less Than	Logical less-than test between two numbers. Returns a boolean.
>=	Greater Than or Equal To	Tests if the operand on the left is greater or equal to the operand on the right. Returns a boolean.
<=	Less Than or Equal To	Tests if the operand on the left is less than or equal to the operand on the right. Returns a boolean.
=	Equal	Tests for equality between two operands.
!=	Not Equal	Tests for equality, returning true when not equal.
&&	Logical AND	Returns true when both operands are true. Anything non-zero is considered true.
	Logical OR	Returns true when either operand is true. Anything non-zero is considered true.
like	Fuzzy String Matching	Compares the left-hand value with the pattern on the right side. The pattern may consist of %, *, and ? wildcards.

Bound Values

Bound values are paths to other values enclosed in braces. These will appear red in the expression editor. When you are writing an expression for a Expression Binding in Vision, you can reference Tag values and property values using the brace notation. When you are writing an expression for an expression Tag, you can only reference other Tag values. You can use the Insert Property  icon and Insert Tag  icon to build these references for you.

Dataset Access

If you have an expression that returns a dataset, you can pull a value out of the dataset using the dataset access notation, which takes one of these forms:

```
Dataset_Expression ["Column_Name"]      //Returns the value from the first row at the given column name.  
Dataset_Expression [Column_Index] //Returns the value from the given column at the first row.  
Dataset_Expression [Row_Index, "Column_Name"] //Returns the value from the given row at the given column name.  
Dataset_Expression [Row_Index, Column_Index] //Returns the value from the given row at the given column index.
```

For example, this expression would pull a value out of a [Table](#) at row 6 for column "ProductCode":

```
{Root Container.Table.data}[6, "ProductCode"]
```

Note that you'll often have to convince the expression system that what you're doing is safe. The expression language can't tell what the datatype will be for a given column, so you may have to use a type-casting function to convince the expression language to accept your expression, like this:

```
toInt({Root Container.Table.data}[6, "ProductCode"])
```

Collection and Mapping access

The following feature is new in Ignition version **8.1.8**
[Click here](#) to check out the other new features

When working with collections such as sequences or arrays, individual elements can be addressed with square brackets and the index of the desired element.

```
// Returns the third elements in the array  
{An Array}[2]
```

Keys in maps ("dictionaries and JSON objects") can be accessed with square brackets and string name of the desired key:

```
// Returns the value of "myKey"  
{A Map}["myKey"]
```

Expression Functions

The expression language's functions are where much of the real power lies. A function may take various arguments, all of which can themselves be any arbitrary expression. This means that you can use the results of one function as the argument to another function. In general, the syntax for a function call is:

```
functionName(expression1, expression2, ...)
```

Whitespace and Comments

Whitespace, such as spaces, tabs and newlines, are largely ignored in the expression language. It is often helpful to break your expression up onto multiple lines for clarity. Comments are delimited by two forward slashes. This will make the rest of that line be ignored. This example shows an if function spread over four lines with comments annotating the arguments.

```

if( {Root Container.UseTagValueOption.selected},
    {MyTags/SomeValue}, // Use the Tag value.
    "Not Selected"      // Use default value if the user doesn't check the box.
)

```

Tag Paths

While referencing a tag path in expressions, you may see some special notation, such as "~" and "[.]". More information on this notation can be found on the [Tag Paths](#) page.

Additional Examples

The following headings demonstrate many simple examples that better demonstrate the expression language's syntax.

String Concatenation

You can use an expression binding to concatenate strings resulting in a new string that reflects the concatenation of different strings.

The function starts with `CONCAT()` and inside the function is a list of strings. They can be manually typed like "42" or they can come from Tags or properties.

Example 1

```
concat("The answer is: ", "42") //returns "The answer is: 42"
```

Example 2

You have a date, and need to extract the year, and concatenate the word "Vintage" to the end for a label display. Bind a label's text property to:

```
dateExtract({Root Container.VintageDate}, 'year') + ' Vintage'
```

Celsius to Fahrenheit

Use an expression binding to convert a temperature from Celsius to Fahrenheit. This is an example of how an expression binding can handle calculations.

```
{celsiusTemp} * 9/5 + 32
```

The reference to the Celsius temperature can come from a property or a Tag. As the property or the Tag changes, so does the expression binding.



Expression Binding – Concat Strings

[Watch the Video](#)


Expression Binding - Celsius to Farenheit

[Watch the Video](#)

Format Date

You can format a date in an expression binding by using the `dateFormat` and `now` functions.

To make a label that updates to show the current time:

1. Drag a **Label** component onto the window.
2. Select the label's **Text** property binding icon and select **Expression** binding.
3. Enter the following code into the expression, and click **OK**.

```
dateFormat(now(), "MMM d, yyyy hh:mm:ss a")
```



Expression Binding – Format Date

The `dateFormat` function takes two arguments. The first argument is any date type variable. This can include another function (like `now`) that returns a date type. The second argument refers to the date format that you want returned. The `now` function returns the current time and in this case it will update every second.

For more information on the date formatting expression, see the [Appendix](#).

[Watch the Video](#)

Date Manipulations

You can manipulate dates in expression bindings such as a date addition/subtraction with the `dateArithmetic` function. This is important when you want to use the expression bindings to select a date that is offset by a certain amount.

The following example will return the time 15 minutes ago by using the `dateArithmetic` expression function:

```
dateArithmetic(now(), -15, "minute")
```



Expression Binding – Date Manipulations

[Watch the Video](#)

Bit Functions

You can use various bit functions in expression bindings like `getBit` to return individual bits of a word.

Example 1

Assuming a Tag path 'Folder/BitTag', the following would return the binary representation of the Tag's value at the 0 position

```
getBit({Folder/BitTag}, 0)
```

For more details on the `getBit` function, please see the [getBit page](#) in the Appendix.



Expression Binding - Bit Functions

[Watch the Video](#)

Example 2

You have 3 bits in a PLC, only one of which will be on at a time. You want to turn these 3 bits into a single integer (0,1,2) to drive a component's Styles. Bind a custom integer property to:

```
binEnum({MyTags/Bit1}, {MyTags/Bit2}, {MyTags/Bit3})
```

Switch

You can use the `switch` function in expression bindings to evaluate conditional statements. This function acts like the `switch` statement in C-like programming languages. It takes the `value` argument and compares it to each of the `case1` through `caseN` expressions.

The following example returns the string "Running" when it is given the value of 1. Its options are 0,1, and 2. And when comparing the value to the options the switch statement returns one of the corresponding results. If a result cannot be found, a fail-over option is returned.

```
switch(1, 0, 1, 2, "Off", "Running", "Fault", forceQuality("!BAD STATE!", 0))
```



Expression Binding – Switch

[Watch the Video](#)

Checking Conditions

You can use expression bindings to return true or false based on different conditions.

Example 1

Consider the following expression that references the Tag with a path of 'Folder/Machine State':

```
{Folder/Machine State} = 0
```

The above expression simply tests the value of the Machine State Tag. If the value of the Machine State Tag is ever equal to 0 then the above expression would return true. In every case where Machine State is not equal to 0, then the expression would return false.

Example 2

It is possible to check for multiple conditions in the same expression. If you have two boolean Tags and you only want the expression to return true if both Tags are true then the binding would look like:

```
{boolTag1}=True && {boolTag2}=True
```

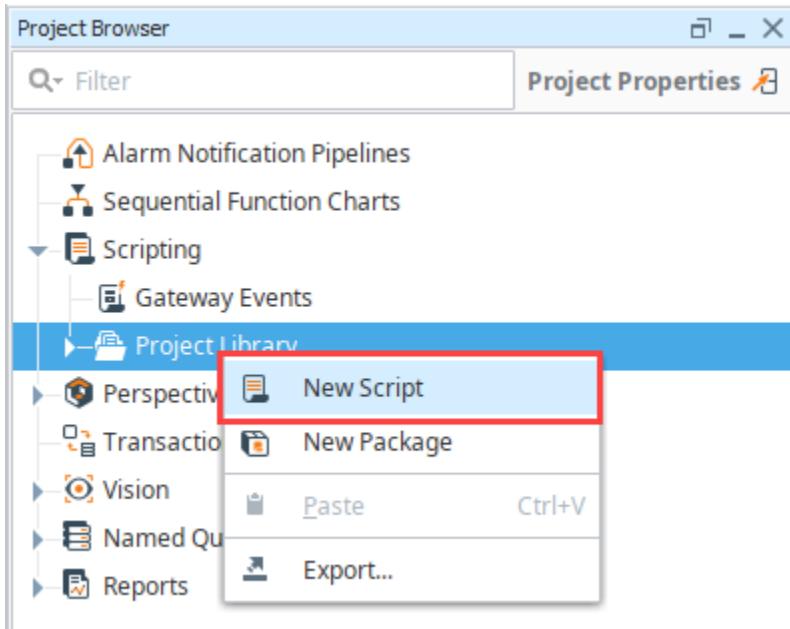
Tag History

You can use expression tags to display tag history data that is stored in a database. The example below demonstrates how to pull a single tag's average history on an expression tag using a Project Library script and the runScript function. The expression used here can be altered to incorporate multiple tag values if desired. The tag history used for this example is from the Sine0 tag in the Quick Start configuration. This was enabled in our project by accessing the **default** Tag Provider in Tag Browser dropdown, selecting **Browse Devices** to import the Sine0 tag, and selecting **History > History Enabled > true** and a **Storage Provider** in the Sine0 Tag Editor.



Since this method uses both project library scripting and requires access to tag data, make sure your project name is defined in the [Gateway Scripting Project](#) property on your Gateway Webpage.

1. In your Project Browser, expand the **Scripting** dropdown.
2. Right-click **Project Library** and select **New Script**.



3. In the name field, enter **TagHistory**.
4. Select **Create Script**.
5. Paste the following script:



Expression Binding - Checking Conditions

[Watch the Video](#)

```

def tagHistory(tagpath):
    results = system.tag.queryTagCalculations(paths=[tagpath], calculations=
    ['Average'], noInterpolation=False)
    return results.getValueAt(0,1)

```

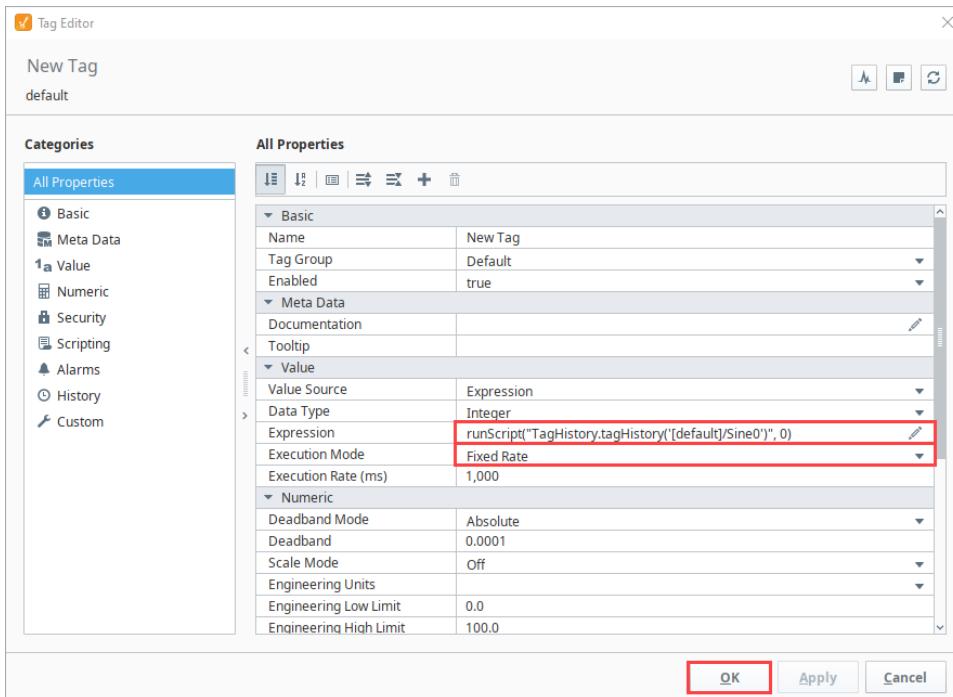
Note: See the [system.tag.queryTagCalculations](#) page for more information on the function.

6. Save your project.
7. Navigate to the Tag Browser and select the Add Tag  icon.
8. Select **New Standard Tag > Expression Tag**.
9. Click the Expression Edit  icon and paste the following expression into the pop-up field:

```
runScript("TagHistory.tagHistory('[default]/Sine0')", 0)
```

Note: See the [runScript](#) page for more information on the function.

10. Click **Apply**.
11. Change the Execution Mode to **Fixed Rate**.
12. Click **OK**.



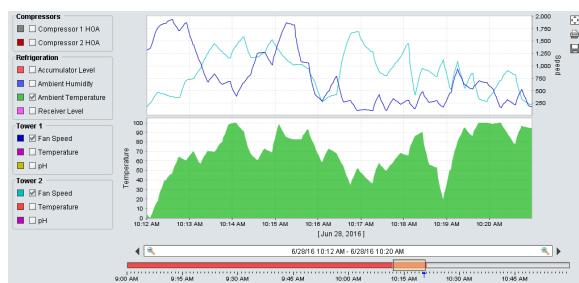
SQL in Ignition

Leveraging Databases

In addition to all the normal HMI functionality you would expect, Ignition has the ability to connect to databases, and this can greatly increase the functionality available to you! You can use databases to store history, create easy to search lists and configurations, and retrieve data from ERP or other systems. When you start using SQL, you can expand your project from a simple HMI to a project that brings your whole process together. The best part is that Ignition connects to as many databases as you want, all from one central location. Because the database lives outside of Ignition, we don't maintain any control or rules over what you can do with your data.

Displaying Data

You can easily display information from your databases on a window along with anything else in Ignition. You can show parts lists, step sequences, [realtime or historical charts](#), add the ability to search for inventory, or anything else you can think of.



On this page ...

- [Leveraging Databases](#)
 - [Displaying Data](#)
 - [What is SQL?](#)
 - [SQL Queries](#)
 - [Database Connections](#)
 - [Using SQL in Ignition](#)
 - [Queries in Bindings](#)
 - [Queries in Scripting](#)
 - [Queries in Tags](#)
 - [Queries in Reports](#)
 - [Queries in Transaction Groups](#)
 - [Database Query Browser](#)
 - [Auto Generated Queries](#)

What is SQL?

SQL stands for "Structured Query Language" and is the backbone of most modern relational databases. It's often referred to as "S.Q.L." or "Sequel," but both are correct and widely recognized. This language (different than the [Python Scripting Language](#)) allows you to write out requests or "queries" against the existing data to [view, add, edit, or remove](#) the information you want in a simple format.

Everything in a relational database is based around tables. Tables store the basic information for any system and can be combined together to make very efficient queries to retrieve your data.

	id	name	description	location	qty
▶	576	Mold 1	Mold for part # 1523	Warehouse A	23
	577	Mold 2	Mold for part # 8825	Warehouse A	12
	578	Bin 1	Finished Goods Bin 1	Warehouse B	9
	579	Bin 2	Scrap Bin 1	Warehouse B	14

```
SELECT *
FROM inventory
JOIN location ON location.name = inventory.location;
```

SQL Queries

SQL queries are crucial to Ignition's database-centric model. Queries can [show what is available](#) or [alter data in the databases](#), and some companies have positions just dedicated to running databases and creating queries. Anywhere Ignition is fetching data, you can choose to use your own custom queries to get exactly what you want out of the database. You can make your queries as simple or complex as you like. If your database is large, you might have a whole team dedicated to creating these queries for you and Ignition will happily execute them.

SQL - Select Statement - Select Data from a Table

```
SELECT * FROM mytable
```

SQL - Inner Join Statement - Selects Records that have Matching Values

```
SELECT users.id, users.firstname, users.lastname, roles.name as 'rolename'  
FROM users  
INNER JOIN mapping ON users.id = mapping.userid  
INNER JOIN roles ON mapping.roleid = roles.id  
WHERE roles.name = 'Administrator'
```

Database Connections

Any SQL query you use needs a [Database connection](#), but Ignition simplifies all that by creating database connections in the Gateway instead of in the clients. This means from one central location you can manage all your database connections, and you don't have to worry about planning around adding clients in the future. Any special rules or connection restrictions are taken care of in the Gateway.

Using SQL in Ignition

There are many types of queries, and many ways to use them in Ignition. Some provide an easy to use builder to automatically store or fetch data, and some allow you to completely customize your queries.

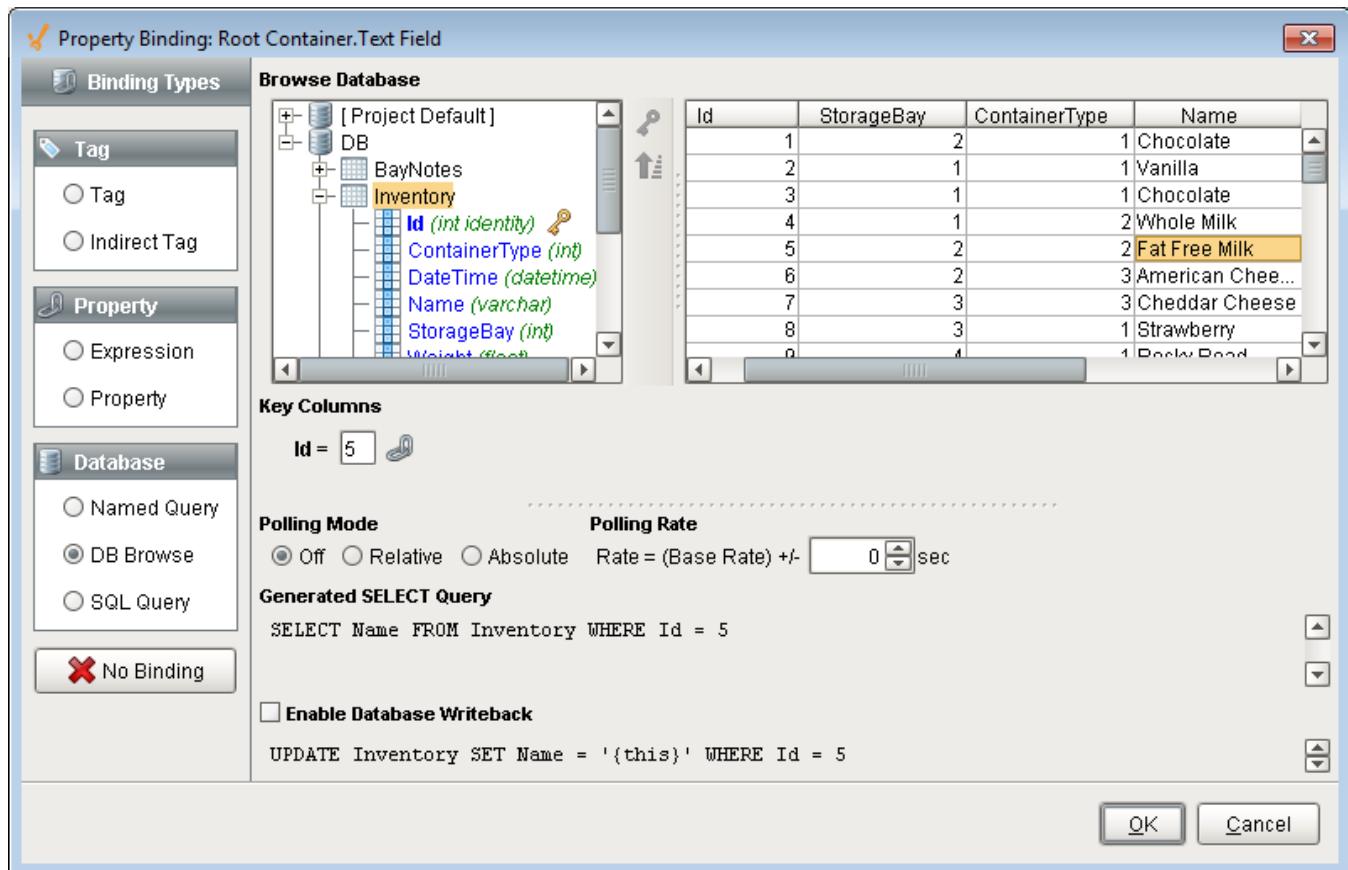
Queries in Bindings

Ignition's binding system offers a lot of flexibility in how database data can be used in a binding. The [Named Query](#) binding allows you to select one of the Named Queries that were previously built for that project, offering a very secure method of pulling data from the database. The [DB Browse](#) builder provides an interface that will build the query based on the data in the table that was selected. This allows even users with little SQL knowledge to pull data from the database. Finally, the [SQL Query](#) option will accept a straight query, so that a query specific to that binding can be written directly in the binding. When binding to a basic (non-dataset) data type, you can use the Writeback feature directly to send any changes back to the database.



Querying Data from Database

[Watch the Video](#)



Queries in Scripting

Ignition offers a number of built-in scripting functions for you to use to query your databases. This makes it very simple to view data, create dynamic scripts that use real data, and more. You can pull individual pieces of information, return whole tables of data, or update your database directly. Depending on the type of query and the sort of results you want, you will use different functions. The following functions are the ones you will use most, and all of them can use a special placeholder (?) to allow for dynamic query building.

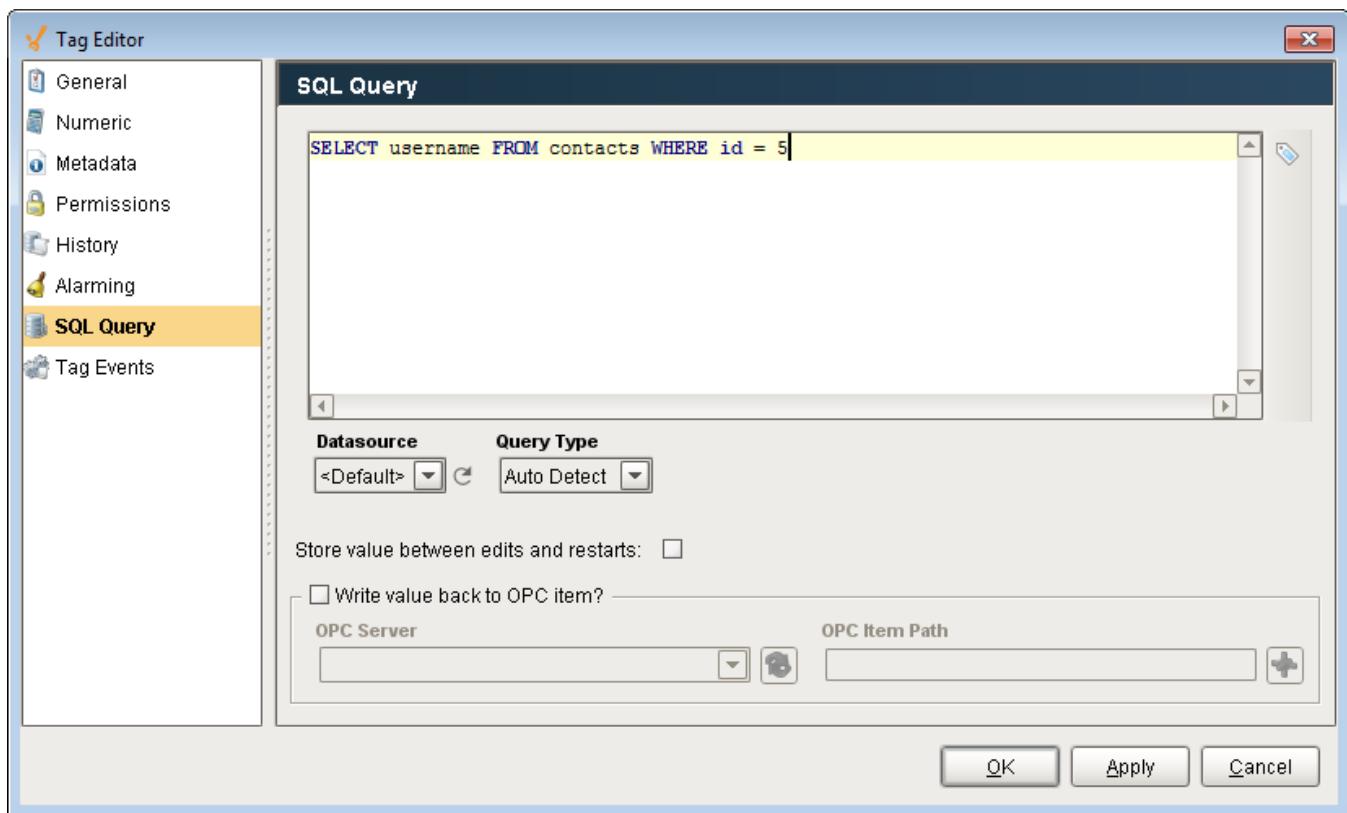
Scripting Function	Description
<code>system.db.runNamedQuery()</code>	Used to run a previously set up Named Query from within a script.
<code>system.db.runPrepQuery()</code>	Used to run basic SELECT queries to fetch whole datasets. This can be used to populate tables, or to sift through the data to do your own calculations.
<code>system.db.runPrepUpdate()</code>	Used to run queries that change the data in the database. Usually used on input form windows to update your database records.
<code>system.db.runScalarPrepQuery()</code>	Used when you want only one value from your results. Perfect for fetching a single value like the highest ID, or first timestamp of a result set.

Each of the different functions takes in different arguments (values) and provides slightly different options and functionality. For example, the `runPrepUpdate()` can return the auto-generated key from insert queries. This can be extremely helpful and eliminate the need to hit the database multiple times if you are using linked tables.

You can find examples of each of these and all the other database functions in the [system.db](#) section of the appendix.

Queries in Tags

Ignition offers Query Tags, which can run queries and return the result as a Tag value, giving all of the projects in the Gateway access to the same Database values.



Queries in Reports

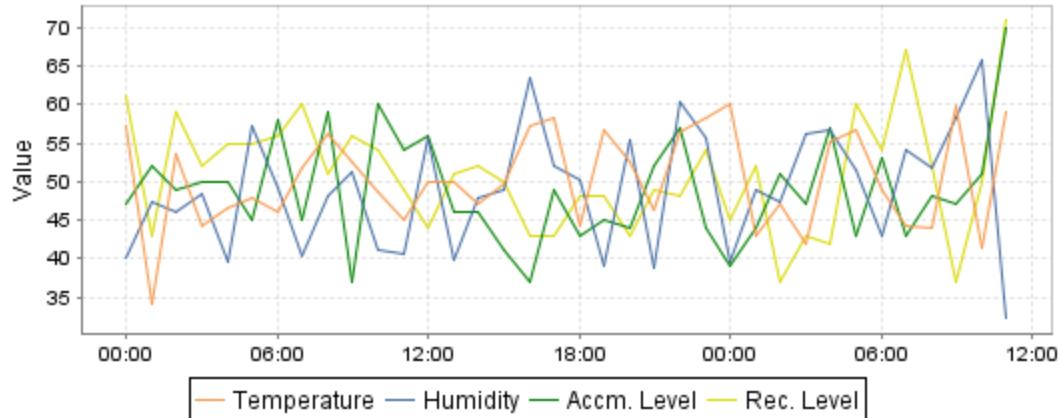
You can leverage queries to access data from all database connections to [create rich reports](#), from simple data logging to complex reports with grouped charts and datasets.

Tag History Report

A report that queries tag history and displays a graph and table

Report Start: 6/27/2016 12:00 AM

Report End: 6/28/2016 11:00 AM



Date / Time	Temperature	Humidity	Accm. Level	Rec. Level
6/27/2016 12:00:00 AM	57.34	40.14	47	61
6/27/2016 1:00:00 AM	33.98	47.33	52	43
6/27/2016 2:00:00 AM	53.61	46.15	49	59
6/27/2016 3:00:00 AM	44.3	48.5	50	52
6/27/2016 4:00:00 AM	46.57	39.59	50	55
6/27/2016 5:00:00 AM	47.9	57.16	45	55
6/27/2016 6:00:00 AM	46.04	49.16	58	56
6/27/2016 7:00:00 AM	51.87	40.42	45	60
6/27/2016 8:00:00 AM	56.24	48.07	59	51
6/27/2016 9:00:00 AM	52.42	51.17	37	56
6/27/2016 10:00:00 AM	48.74	41.01	60	54
6/27/2016 11:00:00 AM	44.89	40.67	54	49
6/27/2016 12:00:00 PM	49.85	55.73	56	44
6/27/2016 1:00:00 PM	49.82	50.94	46	51

Queries in Transaction Groups

While [Transaction Groups](#) are great at storing Tag data to a Database automatically, the built-in [Expression Items](#) can execute a SQL Query within the Transaction Group.

Group

Execution Disabled **Save project to apply changes** Enabled **Disabled** Pause

Edit group tag

Expression/SQL

Expression Type *: SQL Query

```

SELECT
    COUNT(target_number)
FROM
    workorders
WHERE
    id = { [~]currentWoID}
  
```

Datasource: <Default> **Query Type**: Auto Detect

Trigger **Action**

Execution Scheduling:

- Timer Schedule
- 1 second(s)

Update mode: OPC to DB

Data source: <Default>

Table name: group_table

Automatically create

Use custom index co

Store timestamp to:

Store quality code to

Delete records older

1 day(s)

Table action

Execution Disabled

Last execution: N/A

Total executions: N/A

Database Query Browser

The [Database Query Browser](#) offers an easy to use environment to run queries in for testing. Here, queries can be tested to figure out what values get returned, or data can be updated through raw queries, or the Database Query Browsers easy to use GUI editor.

The screenshot shows the Ignition Database Query Browser interface. In the top-left, the query `SELECT * FROM machines` is entered. To the right is an **Execute** button with a lightning bolt icon. Below the query is a checkbox for **Limit SELECT to:** followed by a text input set to 1000 rows. The main area displays a **Resultset 1** table with columns `id`, `machine_name`, and `area_number`. The data is as follows:

<code>id</code>	<code>machine_name</code>	<code>area_number</code>
1	Conveyor	1
2	Press	2
3	Tank	1
4	Packer	3
5	Loader	3
6	Oven 1	3
7	Oven 2	2
9	Mixer	3
10	Cold Storage	2
12	Chopper	1

At the bottom left, it says "10 rows fetched in 0.032s". To the right of the table is a **Schema** tree view showing the database structure. The `machines` table is selected, revealing its columns: `id (INT)`, `area_number (IN)`, and `machine_name (C)`. Other tables listed in the schema include `agent_events`, `alarm_event_data`, `alarm_events`, `containers`, `files`, `scada_roles`, `scada_user_ci`, `scada_user_ex`, `scada_user_rl`, and `scada_user_sa`.

Auto Generated Queries

Many systems within Ignition utilize a database connection, but the queries that are executed are constructed automatically by the system and do not require you to build the queries manually. These systems such as the [Tag Historian](#), the [Alarm Journal](#), or the [Database User Source](#) are very easy to set up and use since each system will automatically generate the necessary tables in the database, insert the relevant data, and even has prebuilt tools to extract the data. However, it is important to note that while these systems can automatically generate queries for you using the various components, these systems are simply storing data in a database which you can manually query out by building your own SQL queries.

Related Topics ...

- [Database Connections](#)
- [Connecting to Databases](#)
- [Scripting](#)

In This Section ...

Writing SQL Queries

SQL Tables



For information about databases and how to get connected, see the [Database Connections](#) section.

The foundation of every database system is a table. Every database consists of one or more tables, which store the database's data/information. Each table is identified by a name (for example `CUSTOMERS` or `ORDERS`), and consists of column definitions and rows of data.

The database table columns have their own unique names and have pre-defined data types. Table columns can have various attributes defining the column functionality (such as the primary key, index, default value, and so on).

While table columns describe the data types, the table rows contain the actual data for the columns.

ID	Name	Address
1	Safeway	123 Safeway Lane
2	Costco	456 Costco Way
3	Target	789 Target Ave

Current common databases are almost all Relational Databases. This means that their tables can relate to each other by including an ID in one table that matches the key of another. These are called foreign keys.

Primary Key

A primary key is a way to uniquely identify each row in a table. While it is possible to create a database table without a primary key, it is highly recommended to configure one for each table. A primary key is comprised of either a single column, or set of columns. When multiple columns are specified as a primary key, this is known as a **composite primary key**. No two distinct rows in a table can have the same value (or combination of values) in those columns.

While Primary Keys can be configured in several ways, they *typically* meet the following criteria:

- **Integer Data Type:** The data type of the key column is typically an integer, and not a varchar. The primary key is only an identifier to a specific row in a table, so an integer data type can easily be used. Some databases support a UID or a UUID (Universally Unique Identifier) that looks like a character string, but is something specially made for primary keys.
- **Automatically Incrementing:** The value of the primary key increments as rows are added. The key is usually configured to automatically increment in the database so that external applications (such as Ignition) don't have to figure out the next available value when inserting a new row.
- **Statically Defined:** Any row that is inserted must fill in these value(s) without creating duplicates. Configuring the primary key as automatically incrementing means that the database will automatically handle this criteria.
- **Non-NUL:** NULL (empty) values should not be present in the primary key. This column (or columns) will usually not allow NULL values.

Index

Indexes speed up the querying process by providing swift access to rows in the data tables, similarly to the way a book's index helps you find information quickly within that book. Indexes are extremely important when querying large sets of data. You should create an index for the set of columns you use commonly in a `WHERE` clause. For example, you should add an index on the timestamp column of a historical table when querying the table by a start and end date. Ignition does this automatically when it creates tables for Tag History or Transaction Groups.

Foreign Key

A Foreign Key is a referential constraint between two tables. The foreign key identifies a column or a set of columns in one (referencing) table that refers to a column or set of columns in another (referenced) table. The columns in the referencing table must be the primary key in the referenced table. For example, you might have a list of suppliers with an integer ID column. Then your invoices would use the supplier's ID instead of the name. These linked tables can save a lot of space because you don't have to include everything about the supplier in each invoice.

Example

Supplier (**SupplierNumber**, Name, Address, Type)
Invoices (InvoiceNumber, **SupplierNumber**, Text, Cost)

On this page ...

- [SQL Tables](#)
 - [Primary Key](#)
 - [Index](#)
 - [Foreign Key](#)
 - [Null Value](#)
 - [Comments](#)
- [SQL Queries](#)
 - [Select Command](#)
 - [Select Examples](#)
 - [Where Clause](#)
 - [Where Examples](#)
 - [Insert Into Command](#)
 - [Insert Examples](#)
 - [Update Command](#)
 - [Update Examples](#)
 - [Delete Command](#)
 - [Delete Examples](#)

Null Value

NULL is a special marker used in SQL to indicate that a data value does not exist in the database. This way it is clear that there is no data, instead of guessing if a value of 0 is correct or just missing data. By definition, NULL is not equal to anything, even other NULL values. Every database has a function to identify if a value is NULL, usually called `isNULL()` or something similar.

Comments

Comments can be added to any SQL query just like with scripting. Single line comments are done with two dashes and a space: '-- '

SQL - Single Line Comment

```
-- This is a single line comment in a SQL query.  
SELECT * FROM my_table
```

You can also do multi line comments by wrapping text within a forward slash and an asterisk: /* text */

SQL - Multi Line Comment

```
/* A multi line comment  
can span multiple lines. The  
comment will stop when it is closed  
with another asterisk and forward slash.*/
```

SQL Queries

SQL queries (or statements) are used to create, maintain, and view relational databases like MySQL, SQLServer, Oracle, etc. They follow a specific format and use just key words to determine the structure of the query. Unlike most coding languages, SQL does not rely on newlines or start/end markers for its format, each query is a single line of code. You will often see SQL queries split over several lines, but that is just to make them easier to read.



You might notice a lot of CAPITALIZED words in SQL queries. While these key words are not case sensitive, it is still common practice for people to capitalize them in a query. Things like `SELECT`, `FROM`, `WHERE` (and a few others) are almost always capitalized because they have a special meaning. Try not to have table or column names that use any of these special words. You will see this capitalization format in examples throughout this user manual and other online references.

Select Command

The `SELECT` statement is used to select data from a database. The result is returned as a data set, called the result set. This is true even if there is only one value returned. The syntax for a `SELECT` statement is as follows:

SQL - The Select Command

```
SELECT myColumn FROM myTable
```

Select Examples

The "*" character can be used to specify all columns from a database table. While this is the easiest way to retrieve results from a table, this is not the recommended approach.

SQL - Selecting All Columns

```
SELECT * FROM Customers
```

The recommended approach is to instead specify only the columns that are required for a query. There are several reasons for this, but performance would be the main one: less columns in a statement means less work for the database, and the resulting data set in Ignition will use less memory.

SQL - Selecting a Single Column

```
SELECT Name FROM Customers
```

SQL - Selecting Multiple Columns

```
SELECT Name, Address FROM Customers
```

Where Clause

The WHERE clause is used in conjunction with other commands to extract only those records that fulfill a specified criterion. The WHERE clause usually goes near the end of the query followed by a condition that the values must meet to be returned.

Pseudocode - Where Clause

```
SELECT myColumn FROM myTable WHERE condition
```

The WHERE clause can use various operators for its condition, with the basic operators being:

Operator	Description
=	Equal to.
<>	Not equal to.
>	Greater than.
<	Less than.
>=	Greater than or equal to.
<=	Less than or equal to.

Where Examples

Only return customers from CA.

SQL - Simple Where Clause

```
SELECT * FROM Customers WHERE State = 'CA'
```

Only return users over a specified age.

SQL - Select Users over 25

```
SELECT * FROM Users WHERE Age > 25
```

Insert Into Command

The INSERT INTO statement is used to insert a new row in a table. If any columns have default values or are auto-incrementing, they can be omitted from the INSERT query.

Pseudocode - The Insert Command

```
INSERT INTO myTable (column1, column2)
VALUES ('Value1', 'Value2')
```

Insert Examples

SQL - Insert Values into Columns

```
INSERT INTO Customers (Name, Address, City, State, Zip, Country, Phone)
VALUES ('Inductive Automation', '90 Blue Ravine', 'Folsom', 'CA', '95630', United States, '1-800-266-7798')
```

If inserting a value into every column of the table, the columns do not need to be listed on the INSERT INTO statement. The values just need to be listed in the same order as the columns in the table. The table in the query below has four columns: id, first name, last name, and title.

SQL - Inserting to all Columns

```
INSERT INTO Users
VALUES (5628, 'Bob', 'Smith', 'Project Manager')
```

Update Command

The UPDATE statement is used to update existing records in a table. If a WHERE clause is not used, **all rows in the table will be updated**. As a result, the UPDATE statement should be used in conjunction with a WHERE clause in most cases. Many official management tools like SQLServer's Management studio will not allow UPDATE commands without a WHERE clause.

Pseudocode - The Update Command

```
UPDATE myTable SET myColumn = 'myValue'
```

Update Examples

SQL - Updating All Rows in a Single Column

```
UPDATE Customers SET Name = 'Inductive Automation'
```

SQL - Updating a Single Column for a Single Row

```
UPDATE Customers SET Address = '2110 21st Street' WHERE ID = 1
```

Delete Command

The DELETE statement is used to delete records in a table. **NEVER** run a delete command without a WHERE clause. It will delete **ALL** records from that table. Many official management tools like SQLServer's Management studio will not allow DELETE commands without a WHERE clause.

Pseudocode - The Delete Command

```
DELETE FROM myTable WHERE myColumn = value
```

Delete Examples

SQL - Deleting Based on Column Value

```
DELETE FROM Customers WHERE Name = 'Inductive Automation'
```

SQL - Deleting Rows Based on an ID Column

```
DELETE FROM Customers WHERE id < 538
```

Related Topics ...

- [Inserting Data into a Database](#)
- [Named Queries](#)
- [Query Builder](#)
- [Database Query Browser](#)

In This Section ...

SQL Select Statements

While the SELECT command in its basic form can be very simple to use, the SELECT statement can be used with other statements or in certain ways that allow you to bring in exactly the data you need.

Selecting Static Values

Static values can be inserted into a resultset returned from a SELECT query as another column. Simply use the static value as a column to select, and the query will return a column where the name of the column is the static value, and every row in that column will return that same static value.

Pseudocode - Selecting Two Columns From a Table, and a Third Column with a Value of 10 for every row

```
SELECT column1, column2, 10 FROM table
```

On this page ...

- [Selecting Static Values](#)
- [Select Distinct Command](#)
- [Order By Clause](#)
- [Limiting Rows Returned](#)
- [Aliases](#)
- [Union Command](#)

Select Distinct Command

The SELECT DISTINCT statement works much like a SELECT statement works, in that it selects data from a database. However, SELECT DISTINCT will only return distinct or different values, not duplicates.

Pseudocode - The Select Distinct Command

```
SELECT DISTINCT column FROM table
```

Examples

This can be useful for getting a better idea of the range of values in a particular column.

SQL - Select Distinct Countries

```
SELECT DISTINCT country FROM Customers
```

Order By Clause

The ORDER BY keyword is used to sort the result-set by a specified column set of column. The ORDER BY keyword sorts the records in ascending (ASC) order by default. If you want to sort the records in a descending order, you can use the DESC keyword.

Pseudocode - Order By Clause

```
SELECT column1, column2 FROM table ORDER BY column2 DESC
```

Examples

SQL - Ordering by One Column

```
SELECT * FROM Customers ORDER BY Name ASC
```

You can use multiple columns to sort, this will sort by state first, and for each state the rows will be sorted by name.

SQL - Ordering by Multiple Columns

```
SELECT * FROM Customers ORDER BY State ASC, Name DESC
```

Limiting Rows Returned

SELECT commands can have the number of rows that the query returns limited using a special keyword. The keyword differs between database providers but the effect is the same, limiting the number of rows returned to a value that you specify.

Database	Keyword	Example
MS SQL Server/ MS Access	SELECT TOP value/percent	<p>Note that the SELECT TOP command is unique in that you can also specify a percentage value instead of an exact number of records.</p> <pre>SELECT TOP 200 column1, column2 FROM table</pre> <pre>SELECT TOP 10 PERCENT column1, column2 FROM table</pre>
MySQL	Limit value	<pre>SELECT column1, column2 FROM table LIMIT 200</pre>
Oracle DB	ROWNUM <= value	<p>Note that the ROWNUM command is unique in that you can use it to identify the row number for any reason, not just limiting the number of rows returned.</p> <pre>SELECT column1, column2 FROM table WHERE ROWNUM <= 200</pre>

Aliases

In a SQL query, aliases are used to give columns or even tables a temporary name for that query. Simply place the keyword AS after a column or table, followed by the alias name. If the alias is two words, it needs to be encapsulated in single quotes.

Pseudocode - Aliasing Columns and the Table

```
SELECT column1 AS a, column2 AS 'b c' FROM table AS t
```

Example

This can be really useful when the table has complex column names.

SQL - Aliasing Complex Column Names

```
SELECT id AS 'Badge Number', name AS 'Employee Name', dob AS Birthday FROM employees
```

This can also be useful when using multiple tables in a query, such as with a [JOIN](#).

SQL - Joining Columns from Two Tables with Aliases

```
SELECT * FROM Contacts AS co JOIN Customers AS cu ON cu.ID = co.CustomerID
```

Union Command

The UNION operator is used to combine the results of two different SELECT statements. This differs from a JOIN in that there does not have to be a relationship between columns. However, both SELECT statements need to select the same number of columns with similar data types in a similar order. So if my first statement selects an int column and then a string column, the second statement needs to do the same. The name of the columns in the resultset will take the name of the columns from the first SELECT in the UNION.

Pseudocode - Union Two Tables

```
SELECT stringCol, intCol FROM table1
UNION
SELECT stringCol, intCol FROM table2
```

Examples

By default, the UNION operator will only select distinct values between the two tables.

SQL - Union of Users and Customers

```
SELECT username FROM users
UNION
SELECT name FROM customers
```

To select all values from both tables, we can use UNION ALL instead.

SQL - Union All

```
SELECT jobTitle FROM jobs
UNION ALL
SELECT position FROM employees
```

Static values can be used in a UNION to help differentiate the rows from each table.

SQL - Differentiating Between Users and Customers

```
SELECT 'User' AS Type, username FROM users
UNION
SELECT 'Customer', name FROM customers
```

Related Topics ...

- [Writing SQL Queries](#)

SQL Where Clauses

Overview

The SQL WHERE clause is utilized to restrict the number of rows impacted by a query. A WHERE clause is commonly utilized in two scenarios:

- In conjunction with a SELECT statement to filter the number of rows returned.
- As part of an UPDATE or DELETE statements to restrict which rows are manipulated by the query.

In either scenario, the syntax of a WHERE clause is used the same, and can have multiple predicates:

Pseudocode - Where Clause Syntax

```
SELECT * FROM table WHERE column = value
```

On this page ...

- Overview
 - Where-Clauses and Data Manipulation
- And Operator
- Or Operator
- Not Condition
- Between Operator
- Like Condition
- In Condition
- Combining Multiple Operators

Where-Clauses and Data Manipulation

In most cases, statements that modify the content of a database table via UPDATE or DELETE **should** include a WHERE clause: otherwise the manipulation will be applied to every row. To demonstrate, you typically want to avoid queries that look like the following

Pseudocode - Never do this!

```
UPDATE table SET column = 'This was a horrible mistake'
```

When manually modifying a database table, a good habit to develop involves first writing a SELECT statement. If you can successfully write a WHERE clause that only returns the results you need to modify, then you can simply change the rest of your query to manipulate the table. Thus, we could write a query like the following:

Pseudocode - Select-Statement with a Where Clause

```
SELECT * FROM table WHERE id in (100,101,150,174)
```

If we receive only the results we need to modify in the query above, we can make a simple modification to our query to now delete just the rows we want.

Pseudocode - Delete-Statement with a Where Clause

```
DELETE FROM table WHERE id in (100,101,150,174)
```

Another common use of the WHERE clause is to search through records and return the ones during a particular time frame. Timestamp columns can use the < and > operators to compare to each other, and certain string formats can be compared to timestamps like this query:

Pseudocode - Where Clause with a Timestamp

```
SELECT * FROM table WHERE t_stamp > '1984-01-25 16:35:55'
```

And Operator

The AND operator allows you to specify two or more conditions in a WHERE clause, where each condition must be true for the row to be returned.

Pseudocode - Using And to Specify Multiple Conditions

```
SELECT column1, column2, column3 FROM table WHERE column1 > value AND column2 < value AND column3 = value
```

Example

This helps to narrow down the result set even further by adding in additional conditions that must be met. This will only return rows for customers from Germany who are also over 20 years old.

SQL - Customers from Germany Over 20

```
SELECT * FROM customers WHERE country = 'Germany' AND age > 20
```

Or Operator

The OR operator allows you to specify two or more conditions in a WHERE clause, but only one of the conditions need to be true for the row to be returned.

Pseudocode - Using Or to Specify Multiple Conditions

```
SELECT column1, column2, column3 FROM table WHERE column1 > value OR column2 < value OR column3 = value
```

Example

The OR operator can help pull in data from two different subsets in the table. This will only return rows for customers from Germany or customers who are over 20 years old.

SQL - Customers from Germany or Customers Over 20

```
SELECT * FROM customers WHERE country = 'Germany' OR age > 20
```

Not Condition

The NOT condition allows you to specify a condition that must not be met for the row to be returned.

Pseudocode - Using Not to Specify a Condition that Shouldn't be met

```
SELECT column1, column2, column3 FROM table WHERE NOT column1 = value
```

Example

This can be useful for finding all data other than a certain subset. This will return all customers who are not from Germany.

SQL - Customers that are not from Germany

```
SELECT * FROM customers WHERE NOT country = 'Germany'
```

Between Operator

The BETWEEN condition allows you to specify a range of values separated by an AND that the value must be in for a condition to be true. The value can be numbers, text or dates and is inclusive of the first and last values in the range.

Pseudocode - Using Between to Specify a Range of Values

```
SELECT column1, column2, column3 FROM table WHERE column1 BETWEEN value1 AND value2
```

Examples

SQL - Customers that are Between the Ages of 20 and 40

```
SELECT * FROM customers WHERE age BETWEEN 20 AND 40
```

Note that the BETWEEN operator would work similarly to using a greater than or equal condition and a less than or equal condition.

SQL - Customers that are Between the Ages of 20 and 40 with no Between Operator

```
SELECT * FROM customers WHERE age >= 20 AND age <= 40
```

Timestamps can also use the BETWEEN operator to check for a given start time and end time.

SQL - Customers that are Between the Ages of 20 and 40 with no Between Operator

```
SELECT * FROM customers WHERE start_time BETWEEN '1984-01-25 00:00:00' AND '1984-01-25 16:35:55'
```

Like Condition

The LIKE condition allows you to specify a condition that must meet a certain pattern. Typically used to compare to string values, the pattern can be built using a combination of characters and the two wildcard values.

- % - Used to specify any number of any characters including zero characters.
- _ - Used to specify exactly one character.

Pattern Examples

Pattern	Meaning	Possible Matches
'%a%	Values that have an 'a' in them.	'a', 'Inductive Automation', 'almost', 'create'
'_a_'	Values that have an 'a' with exactly one character before and after the 'a'.	'bat', 'cat', 'can'
'_a%	Values that have an 'a' as the second character.	'da', 'saw', 'catcher'
'a%t'	Values that start with 'a' and end with 't'.	'about', 'at'
'%a%_%_%_'	Values that contain an 'a' with at least 3 other characters after it.	'trains', 'airplane', 'canteen'
'%a%a%	Values that contain at least two 'a' characters in them.	'Inductive Automation', 'separate', 'apart'

Once the pattern has been constructed, it can be used with the LIKE operator to find values that match the specified pattern.

Pseudocode - Using Like to Specify a Pattern of Values

```
SELECT column1, column2, column3 FROM table WHERE column1 LIKE '%a%'
```

Example

The LIKE operator can be used to find all values that match a criteria, such as all countries with 'land' in that name

SQL - Customers that are not from Countries with 'land' in the Name

```
SELECT * FROM customers WHERE country LIKE '%land%'
```

In Condition

The IN operator allows you to specify a subset of values, with the condition that the return match at least one of them. Using an IN operator is similar to using multiple OR operators for the same column.

Pseudocode - Using IN to Specify Multiple Values

```
SELECT column1, column2, column3 FROM table WHERE column1 IN (value1, value2, value3)
```

Examples

The IN can be used as a shorthand way of writing out multiple conditions for the same column separated by OR operators. This would select all values where the country is either Germany, France, or USA.

SQL - Customers from Germany or France or USA

```
SELECT * FROM customers WHERE country IN ('Germany', 'France', 'USA')
```

This would be similar to doing something like this.

SQL - Customers from Germany or France or USA

```
SELECT * FROM customers WHERE country = 'Germany' OR country = 'France' OR country = 'USA'
```

The real power of the IN operator is that instead of specifying static values, an entirely new query can be run to compare values against.

SQL - Customers from Countries that users are also in

```
SELECT * FROM customers WHERE country IN (SELECT country FROM users)
```

Combining Multiple Operators

Multiple AND and OR operators can be combined to specify multiple different conditions that need to be met in order for a particular row to be returned. Additionally, each condition can be simple using the mathematical operators or complex using the conditions listed above. When using AND and OR operators in a WHERE clause, the AND will take precedence, evaluating first before the OR. In the pseudocode below, the row will be returned if either both the first and second conditions are met, or the third condition is met.

Pseudocode - Using Multiple Where Clause Operators

```
SELECT column1, column2, column3 FROM table WHERE column1 > value AND column2 < value OR column3 = value
```

However, the order at which the operators get evaluated can change by placing parentheses around the conditions which should be evaluated first. In the pseudocode below, the row will be returned if both the first condition is met, and either the second or third condition is met.

Pseudocode - Using Multiple Where Clause Operators with Parentheses

```
SELECT column1, column2, column3 FROM table WHERE column1 > value AND (column2 < value OR column3 = value)
```

Examples

We can use complex conditions with different operators to find all customers who are over the age of 50 in a country that has 'land' in the name, or any customers in Germany or France.

SQL - Multiple Complex Conditions

```
SELECT * FROM customers WHERE country LIKE '%land%' AND age > 50 OR country IN ('Germany', 'France')
```

Using parentheses in the same query can drastically change what valid return conditions are. Here, the customer must both be from a country with 'land' in the name, as well as either over 50 or from Germany or France.

SQL - Multiple Complex Conditions

```
SELECT * FROM customers WHERE country LIKE '%land%' AND (age > 50 OR country IN ('Germany', 'France'))
```

Related Topics ...

- [Writing SQL Queries](#)
- [system.db.runNamedQuery](#)

SQL Table Joins

Overview

The SQL JOIN allows you to run a single SELECT statement that references multiple tables. This can be used for more advanced filtering, as well as combining data from multiple tables in a single result set. The process of joining two tables involves stating both tables in the query, and then specifying that a column from one table relates to another in some way.

Joins may look imposing at first, but they are simply SELECT statements that utilize columns from multiple tables.

The JOIN keyword works in conjunction with a SELECT statement. However, there are some key concepts that must be addressed when attempting to use the JOIN keyword.

Specifying Each Column

When listing column in a statement that uses the JOIN keyword, you must denote which table each column is being retrieved from. You can do this by using an [Alias](#), or with the fully qualified column name. This prevents ambiguous columns in the context of the query, and makes it easier for you to use other keywords in the statement: i.e., adding WHERE clauses that apply to multiple tables.

Pseudocode - Fully Qualified Column Name

```
table_name.column_name
```

On this page ...

- Overview
 - Specifying Each Column
 - Declare the Relation
 - Combining the Concepts
 - Joins with Three or More Tables
- Join
 - Join in Action
- Left Join
 - Left Join in Action
- Right Join
 - Right Join in Action
- Full Join
 - Full Join in Action
 - Full Joins in MySQL

Declare the Relation

After the JOIN keyword, you must state which columns from each table relate to each other. This is accomplished by stating the name of the table, using the ON keyword, and then stating that a column on the first table is equal to a column on the second table. The columns specified are typically [primary keys](#) for their respective tables.

In the example below, we're stating that the values in **some_column** on **tableA** should be associated with matching values in **some_other_column** on **tableB**.

Pseudocode - Using the Join Keyword

```
FROM
    TableA
JOIN tableB ON tableA.some_column = tableB.some_other_column
```

Note that the order you present the columns **after** the ON clause does not matter: A = B is equivalent to saying B = A, so we could switch the columns listed with no distinguishable impact on the resulting query.

Combining the Concepts

Altogether, a basic JOIN looks like the following:

Pseudocode - Join Syntax

```
SELECT
    tableA.column,
    tableB.column
FROM
    tableA
JOIN tableB ON tableA.identity_column = tableB.identity_column
```

Joins with Three or More Tables

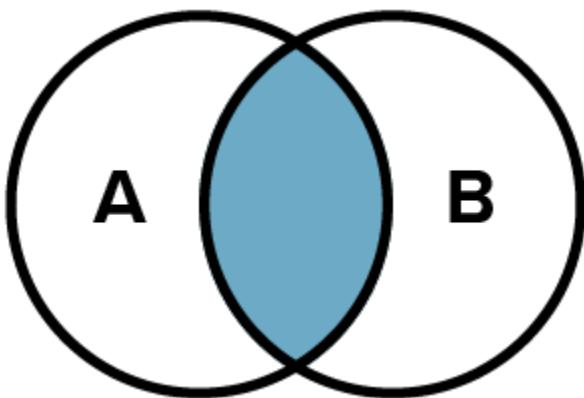
Joins can even be done between three or more tables. The syntax is similar, with each new table relation declared below the first.

Pseudocode - Join Syntax

```
SELECT
    tableA.column,
    tableB.column,
    tableC.column
FROM
    tableA
JOIN tableB ON tableA.identity_column = tableB.identity_column
JOIN tableC ON tableA.identity_column = tableC.identity_column
```

Join

The standard JOIN, also referred to as INNER JOIN, will only return rows where the joined columns contain matching values. If one of the joined columns contains a value that is not present in the other, then the row is not represented in the result set. You would use JOIN when you only want results that are represented in both tables.



This section will demonstrate the various uses of the JOIN keyword. For the sake of clarity, the queries will run against tables that look like the following:

Products Table	
id	product_name
1	Apples
2	Oranges
3	Grapes
4	Plums

Inventory Table			
id	product_id	quantity	product_vendor
1	1	15	Apple Corp
2	2	25	Orange Ya-Glad
3	3	56	Grape Escape
4	5	45	Banana Solutions

Join in Action

In this demonstration, only rows that pertain to Apples, Oranges, and Grapes are being returned. We're using a JOIN between **products.id** and **inventory.product_id**, so our results will only contain rows that have matching values from both of those columns. Our result set does not contain any information on products with **products.id** values of 4 or **inventory.product_id** values of 5, because those values are not present in **both** of the joined columns

SQL - Joining Products and Inventory

```

SELECT
    products.id
    ,inventory.product_id
    ,products.product_name
    ,inventory.product_vendor
    ,inventory.quantity
FROM
    products
JOIN inventory ON products.id = inventory.product_id

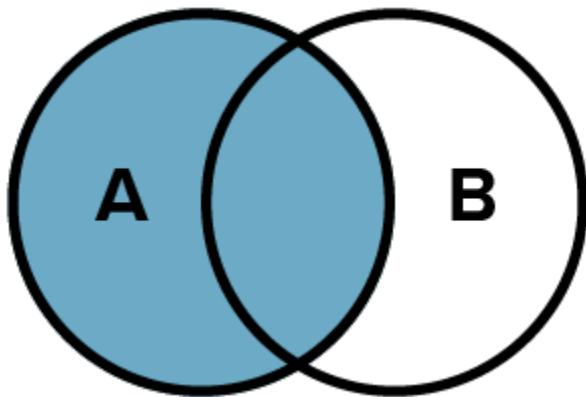
```

Example Results

id	product_id	product_name	product_vendor	quantity
1	1	Apples	Apple Corp	15
2	2	Orange	Orange Ya-Glad	25
3	3	Grapes	Grape Escape	56

Left Join

Return all rows from the left-most table (**table A** in the diagram), even if there are no matches on the right-most table (**table B**). If there isn't a matching record in the right table, then NULL values are returned.



Left Join in Action

Here we see all rows returned from our products table (since it is the left-most table in our query). In row 4, columns that are being populated via the inventory table (**product_id**, **product_vendor**, and **quantity**) contain NULL values, because there isn't a row on the inventory table that matches with a **product_id** value of 4. The query must return something in this case, so it returns NULL for these columns.

SQL - Left Joining Products and Inventory

```

SELECT
    products.id
    ,inventory.product_id
    ,products.product_name
    ,inventory.product_vendor
    ,inventory.quantity
FROM
    products
LEFT JOIN inventory ON products.id = inventory.product_id

```

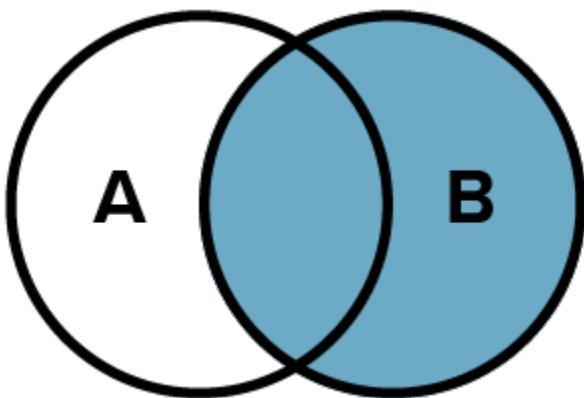
Example Results

id	product_id	product_name	product_vendor	quantity
1	1	Apples	Apple Corp	15
2	2	Orange	Orange Ya-Glad	25
3	3	Grapes	Grape Escape	56
4	4	Cherries	NULL	NULL

1	1	Apples	Apple Corp	15
2	2	Orange	Orange Ya-Glad	25
3	3	Grapes	Grape Escape	56
4	NULL	Plums	NULL	NULL

Right Join

Return all rows from the right-most table (**table B**), even if there are no matches on the left-most table (**table A**). If there isn't a matching record on the left table, then NULL values are returned.



Right Join in Action

When using a RIGHT JOIN, all rows will be returned from the inventory table. The products table does not have a row that contains an **id** value of 5, so the **id** and **product_name** columns will show NULL values in our result set.

SQL - Right Joining Products and Inventory

```

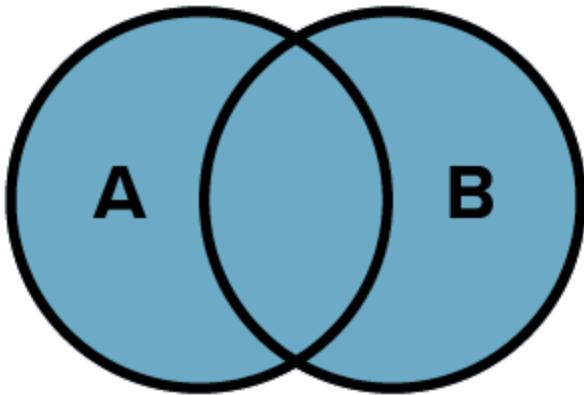
SELECT
    products.id
    ,inventory.product_id
    ,products.product_name
    ,inventory.product_vendor
    ,inventory.quantity
FROM
    products
RIGHT JOIN inventory ON products.id = inventory.product_id
  
```

Example Results

id	product_id	product_name	product_vendor	quantity
1	1	Apples	Apple Corp	15
2	2	Orange	Orange Ya-Glad	25
3	3	Grapes	Grape Escape	56
NULL	5	NULL	Banana Solutions	45

Full Join

The FULL JOIN returns all rows from both tables, regardless if there are matching values in the joined columns. You would use a FULL JOIN in cases where you want to show all applicable records from both tables, and synchronize the data across both tables via the joining columns where possible.



Full Join in Action

Note that we're using the same query as the standard JOIN, but we've prepended "FULL" to the last line of the query. Note the NULL values returned in cases where our product ID is not fully represented on both tables.

SQL - Full Joining Products and Inventory

```

SELECT
    products.id
    ,inventory.product_id
    ,products.product_name
    ,inventory.product_vendor
    ,inventory.quantity
FROM
    products
FULL JOIN inventory ON products.id = inventory.product_id
  
```

Example Results

id	product_id	product_name	product_vendor	quantity
1	1	Apples	Apple Corp	15
2	2	Orange	Orange Ya-Glad	25
3	3	Grapes	Grape Escape	56
4	NULL	Plums	NULL	NULL
NULL	5	NULL	Banana Solutions	45

Full Joins in MySQL

MySQL does not have an equivalent FULL JOIN. However, you can emulate one by utilizing a LEFT JOIN, RIGHT JOIN, and the UNION keyword. For the sake of simplicity, we will return all columns in the following example, but you would still want to specify individual columns in both SELECT query.

SQL - Full Join in MySQL using Left Join, Right Join, and Union Keyword

```

SELECT * FROM products
LEFT JOIN inventory ON products.id = inventory.product_id

UNION ALL

SELECT * FROM products
RIGHT JOIN inventory ON products.id = inventory.product_id
WHERE products.id IS NULL
  
```

Related Topics ...

- [Writing SQL Queries](#)
- [SQL Query Data Source](#)

SQL Common Functions

Functions are available in most SQL databases, and can provide some helpful utility to any queries you may be calling in Ignition.

This page contains some commonly used SQL functions that some databases contain. The exact functions available and usage depends on the database, so always check your database's documentation for a more complete list of available functions.

Using Column Values

It is important to understand that when calling these functions, you generally use a column name instead of a static number. For the sake of simplicity, the tables below demonstrate how to use the functions with static values, but they usually are switched for column names, ie:

Pseudocode - Passing a Column to a Function

```
SELECT SUM(downtime_duration) FROM downtime_events
```

On this page ...

- Using Column Values
 - Example Table
- Numeric Functions
- String Functions
- Date Functions
- Logic Functions
- Group By Clause

Example Table

Some of the functions on this table are better demonstrated when used in conjunction with a table (i.e., using the AVG() function with a single value isn't too interesting). Thus, the following table contains sample data that the functions on this page will utilize if necessary.

Products Table

id	product_quantity	product_name	date_added	date_updated	origin_state
1	100	apples	Mon Jan 29 00:00:00 PST 2018	Mon Jan 29 12:00:00 PST 2018	California
2	24	oranges	Mon Feb 13 00:00:00 PST 2017	Mon Feb 13 09:00:00 PST 2017	Florida
3	56	grapes	Mon Mar 07 00:00:00 PST 2016	Mon Mar 07 05:00:00 PST 2016	California

Numeric Functions

Function	Description	Example	Output
ABS(value)	Returns the absolute value of the passed number or column.	<pre>SELECT ABS (-3.5)</pre>	3.5
AVG(value)	Takes the values of a single numeric column, and returns an average. A WHERE clause may be used in the same statement to filter out some of the rows on the table.	<pre>SELECT AVG (product_quantity) FROM products</pre>	60
CEILING(value)	Returns the next greatest integer value based on the argument provided. Thus, CEILING(10.1) would return 11.	<pre>SELECT CEILING(10.1)</pre>	11

COUNT(value)	Returns a row count. Typically takes either a single column, *, or 1. Regardless of which row is passed, the function will return the number of rows on the table that meet the criteria of any WHERE clauses.	<pre>SELECT COUNT(*) FROM products</pre>	
FLOOR(value)	Returns the next smallest integer value based on the argument provided. Thus, FLOOR(10.9) would return 10.	<pre>SELECT FLOOR(10.9)</pre>	10
MAX(value)	Returns the largest value from the specified column.	<pre>SELECT MAX(product_quantity) FROM products</pre>	100
MIN(value)	Returns the smallest value from the specified column.	<pre>SELECT MIN(product_quantity) FROM products</pre>	24
ROUND(value, decimal_places)	Returns a number rounded to a certain number of decimal places. Takes two parameters. The first is the number to round to, and the second is the number of decimal places to round to.	<pre>SELECT ROUND(1.234, 1)</pre>	1.2
SUM(value)	Takes the value of a single numeric column, and returns the sum. A WHERE clause may be used in the same statement to filter out some of the rows on the table.	<pre>SELECT SUM(product_quantity) FROM products</pre>	180

String Functions

String Functions		Example	Output
CONCAT (value1, value2,..., valueN)	Concatenates multiple strings or values. Some databases may require you to convert each value to a string before concatenating.	<pre>SELECT CONCAT(product_name, ':', product_quantity) FROM products</pre>	apples:100 oranges:24 grapes:56

LOWER (value)	Converts a string to lowercase.	<pre>SELECT LOWER('MAKE Me smALL')</pre>	make me small
LTRIM(value)	Removes leading space from a string.	<pre>SELECT LTRIM(' Take a little off the left')</pre>	Take a little off the left
REPLACE (original_string, target_string, replacement_string)	Searches a string for a substring (target_string), and replaces the substring with the replacement_string.	<pre>SELECT REPLACE('Who is awesome', 'Who is', 'You are')</pre>	You are awesome
RTRIM(value)	Removes trailing space from a string	<pre>SELECT RTRIM('Take a little off the right ')</pre>	Take a little off the right
SUBSTRING (original_string, character_index, [length])	Extracts a substring from another string based on character index. Takes two parameters: the original string, and the character index to start at. An optional third parameter can specify the number of characters to extract. Character index is one-based, so the first character in the string rests at index 1.	<pre>SELECT SUBSTRING('This is my string!', 9, 9)</pre>	my string
TRIM(value)	Removes both leading and trailing space from a string.	<pre>SELECT TRIM(' Trim Both Sides ')</pre>	Trim Both Sides
UPPER(value)	Converts a string to uppercase.	<pre>SELECT UPPER('super size me')</pre>	SUPER SIZE ME

Date Functions

There are many date and time functions for each database (MySQL, MSSQL, Oracle, etc), but they all vary wildly. These examples work in most databases:

Function	Description	Example	Output
CURRENT_TIMESTAMP()	Returns the current date and time, as reported by the database.	<pre>SELECT CURRENT_TIMESTAMP()</pre>	Returns the current time
TIMEDIFF(date1, date2)	Returns a difference between two dates. Assumes that date1 is the most recent datetime.	<pre>SELECT TIMEDIFF (date_updated, date_added) FROM products WHERE id =1</pre>	Thu Jan 01 12:00: 00 PST 1970

Logic Functions

Function	Description	Example	Output
COALESCE (value1, value2,... valueN)	Returns the first non-null expression.	<pre>SELECT COALESCE (NULL, 'Pick me!')</pre>	Pick me!
ISNULL (expression)	Returns true if an expression is NULL.	<pre>SELECT ISNULL (NULL)</pre> <pre>SELECT ISNULL (14)</pre>	True False
NULLIF (expression1, expression2)	Compares two expressions. If they are equal to each other, then the function returns a NULL. If the two expressions are not equal, the first expression passed to NULLIF() is returned.	<pre>SELECT NULLIF (100, 100)</pre> <pre>SELECT NULLIF (100, 3)</pre>	NULL 100

Group By Clause

The GROUP BY statement is used in conjunction with the aggregate functions to group the result-set by one or more columns. This way you can find the MIN, MAX, Average, COUNT, etc., for each group of rows. Grouping can also be done for multiple columns, with precedence going in the order that they are listed.

Pseudocode - Passing a Column to a Function

```
SELECT SUM(column1) FROM table GROUP BY column2
```

Example

Here, we are grabbing the sum of product quantity for each origin state.

SQL - Simple Group By

```
SELECT SUM(product_quantity), origin_state, FROM products_table GROUP BY origin_state
```

Related Topics ...

- [Writing SQL Queries](#)

- SQL Query Data Source
- system.db.runNamedQuery

SQL Stored Procedures

For the uninitiated, Stored Procedures are a series of predefined SQL statements that are configured and stored in the database, and can be easily called from another application, such as Ignition, with a single statement. Conceptually, they are very similar to a scripting function: parameters may be passed to them, they can utilize looping and conditional logic, and can return a value.

Stored Procedure Syntax

Stored procedures are created and maintained in your database. As a result, the creation of a stored procedure is outside the scope of the Ignition User Manual: the commands used to create a stored procedure vary per database.

In regard to calling a Stored Procedure, the syntax can also differ.

SQL Server

The EXEC command is utilized to execute a stored procedure.

Pseudocode - Executing a Stored Procedure in SQL Server

```
EXEC dbo.myStoredProcedure
```

Parameters may be passed to the Stored Procedure. SQL Server's documentation has more details on utilizing parameters with Stored Procedures.

Pseudocode - Executing a Stored Procedure in SQL Server with Parameters

```
EXEC dbo.myStoredProcedure @myParam = 100, @AnotherParameters = 'Hello'
```

MySQL

MySQL uses the CALL command to execute a Stored Procedure. Note the parentheses characters at the end, as they must be present even when the Stored Procedure is not utilizing parameters.

Pseudocode - Executing a Stored Procedure in MySQL

```
CALL myStoredProcedure()
```

If parameters are defined, they can be passed via the parenthesis, similar to how functions work in Python.

Pseudocode - Executing a Stored Procedure in MySQL with Multiple Parameters

```
CALL myStoredProcedure(100, 'Hello')
```

For information on the creation of a stored procedure, as well as proper SQL syntax to call a Stored Procedure, reference your database's documentation. Alternatively, if you have a database administrator, they can typically help with the creation and execution of a Stored Procedure.

Calling Stored Procedures in Ignition

There are several locations in Ignition where Stored Procedures may be utilized from.

SQL Query Bindings

Instead of typing a query directly into a [SQL Query binding](#), a Stored Procedure may be executed instead. Assuming a MySQL database contains a Stored Procedure named 'return_all_bays', we can call the procedure on a binding with the following:

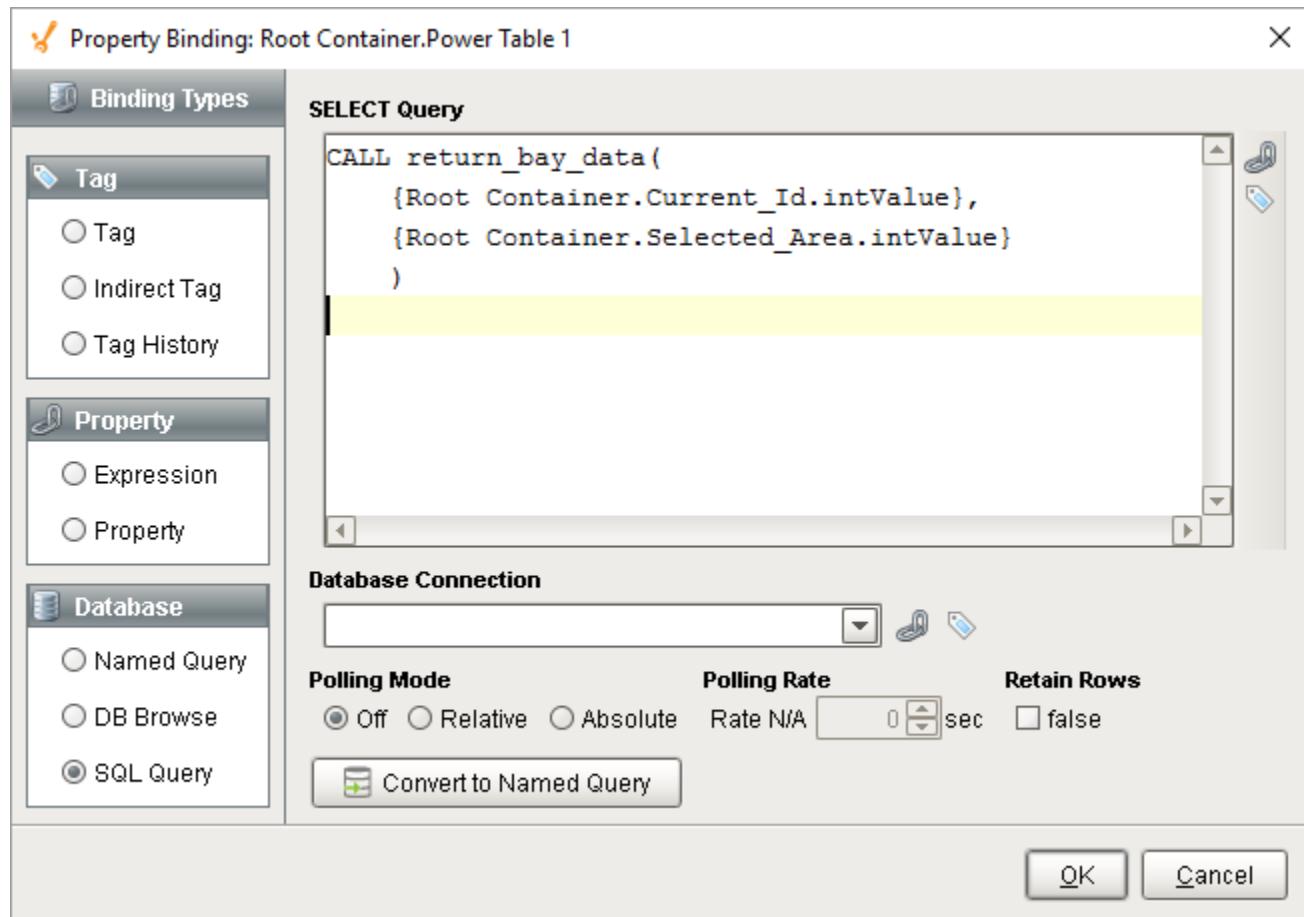
On this page ...

- [Stored Procedure Syntax](#)
 - [SQL Server](#)
 - [MySQL](#)
- [Calling Stored Procedures in Ignition](#)
 - [SQL Query Bindings](#)
 - [Named Queries](#)
- [Stored Procedure Groups](#)
- [Stored Procedures in Python Scripts](#)
 - [Using createSProcCall and execSProcCall](#)
 - [Other System Functions](#)

SQL - Calling a Stored Procedure in MySQL

```
CALL return_all_bays()
```

As with all bindings, Tag  and Property  references may be utilized by clicking the appropriate icons in the binding window.



Named Queries

Named Queries support Stored Procedure calls. As mentioned under [Stored Procedure Syntax](#), the syntax is based entirely on how your database expects a Stored Procedure to be called. Be mindful of the [Query Type setting](#), as it has to match what the stored procedure is doing: if it is returning a result set, leave it set to Query, if it is modifying a record in the database, then set the type to Insert Query.

The screenshot shows the Ignition Project Browser interface. On the left, the Project Browser tree includes nodes like Alarm Notification Pipelines, Sequential Function Charts, Scripting, Perspective, Transaction Groups, Vision, Named Queries (with sub-nodes paramProblems, Inventory, MyNamedQuery, OEEValues, Query), Stored Procedure, Update Query, and Reports. The 'Stored Procedure' node is selected. The main workspace is titled 'Stored Procedure' and has tabs for Settings, Authoring (which is active), and Testing. Under 'Authoring', there are sections for Database Connection (set to <Default>) and Query Type (set to Query). Below these are 'Parameters' and a 'Query' section. The 'Parameters' table has two rows: 'Value' for 'firstParam' (Data Type Int4) and 'Value' for 'secondParam' (Data Type String). The 'Query' section contains the SQL command:

```

1 CALL myStoredProcedure( :firstParam, :secondParam)
2

```

Stored Procedure Groups

One of the easiest ways to utilize Tags with a Stored Procedure is to use the [Stored Procedure Group](#). Parameters can be easily assigned to each item in the group, and utilize all of the features of a [Transaction Group](#), such as scheduled execution and triggers.

Each item in the group is linked to a specific parameter in the Stored Procedure. Any IN or INOUT parameters can write directly to the Tags, while new values can be fed into OUT and INOUT parameters allowing you to easily move data from Tags into the database with the Stored Procedure.

The screenshot shows the configuration interface for a Stored Procedure Group. The left side lists three types of items: 'Basic OPC/Group Items (1)', 'Run-Always Expression Items (ignore trigger) (0)', and 'Triggered Expression Items (0)'. The 'Basic OPC/Group Items' table has one row with the following data:

Item Name	Value	Target Name	Output	Data Type	Properties
Daily Total ...	100	passBack		Int4	

The 'Run-Always Expression Items' and 'Triggered Expression Items' tables are currently empty. The right side of the interface contains configuration options:

- Action:** Enabled (green play button)
- Trigger:** Trigger button
- Options:** Options button
- Execution Scheduling:**
 - Timer:** Selected (radio button)
 - Schedule:** Unselected (radio button)
 - Interval: 1 second(s)
- Data source:** <Default>
- Procedure name:** retrieve_daily_total

Stored Procedures in Python Scripts

There are a few ways to call a Stored Procedure from a script in Ignition.

Using `createSProcCall` and `execSProcCall`

The recommended approach to calling a Stored Procedure from a Python script in Ignition typically involves two main steps:

1. Calling [system.db.createSProcCall](#) to create a call context, or object that effectively represents the impending stored procedure call. This object can be used to specify parameter values that will be passed to the Stored Procedure.
2. Using [system.db.execSProcCall](#) to execute the Stored Procedure.

Once the Stored Procedure has been executed, the call context generated in step #1 can be used to reference any values that were returned.

Other System Functions

Technically, most other system functions in the "db" library, such as [system.db.runPrepQuery](#), can be used to call a Stored Procedure. We generally recommend against this, as [system.db.createSProcCall](#) and [system.db.execSProcCall](#) are better suited to work with Stored Procedures and have some additional functionality not found in the other db functions.

Related Topics ...

- [SQL Query Bindings in Vision](#)
- [Understanding Transaction Groups](#)

Query Builder

Crafting Queries with the Query Builder

Many places in Ignition that allow for SQL queries have a link to the SQL Query Builder tool. The Query Builder is a powerful Drag-and-Drop query building GUI that allows you to make complex queries from your connected databases. While a basic understanding of SQL helps make the most of the Query Builder tool, most people will have no problem creating effective queries after a brief tutorial. Additionally, the Query Builder does go over many advanced features of SQL that may be unfamiliar. We suggest looking up how these work in your favorite SQL resource guide as this covers how to use them, not what they do.

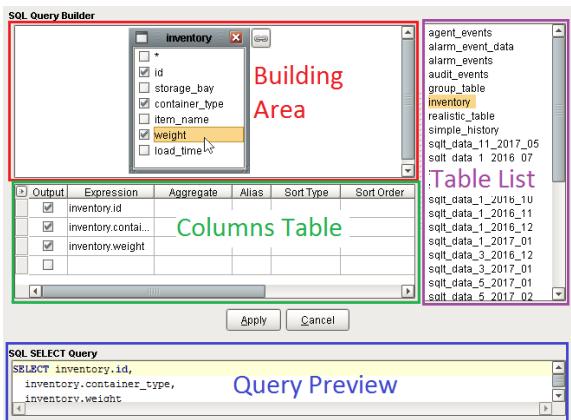
Builder Anatomy

Once opened, the Query Builder has the following items:

- **Building Area:** Visualizes the table relationships of all tables included in the query. Tables are typically dragged in from the **Table List**.
- **Columns Table:** Shows which columns from all tables are referenced in the query.
- **Table List:** Shows all database tables in the selected Database Connection.
- **Query Preview:** Shows a preview of the query that will be created once the **Apply** button has been pressed.

On this page ...

- [Crafting Queries with the Query Builder](#)
 - [Builder Anatomy](#)
 - [Opening the Query Builder](#)
 - [Using the Syntax Parser](#)
 - [Using the Builder](#)
 - [Query Properties](#)
- [Joining Tables in the Query Builder](#)
 - [Joins Right-Click Menu](#)
 - [Join Properties](#)
- [Columns Table](#)
 - [Right-Clicking on an Item](#)
 - [Field Description](#)



Opening the Query Builder

The Builder can be accessed from several different resources in Ignition, notably when using [Named Queries](#) and setting up [Report Data](#). To use the Builder, click the **Show Builder** button.

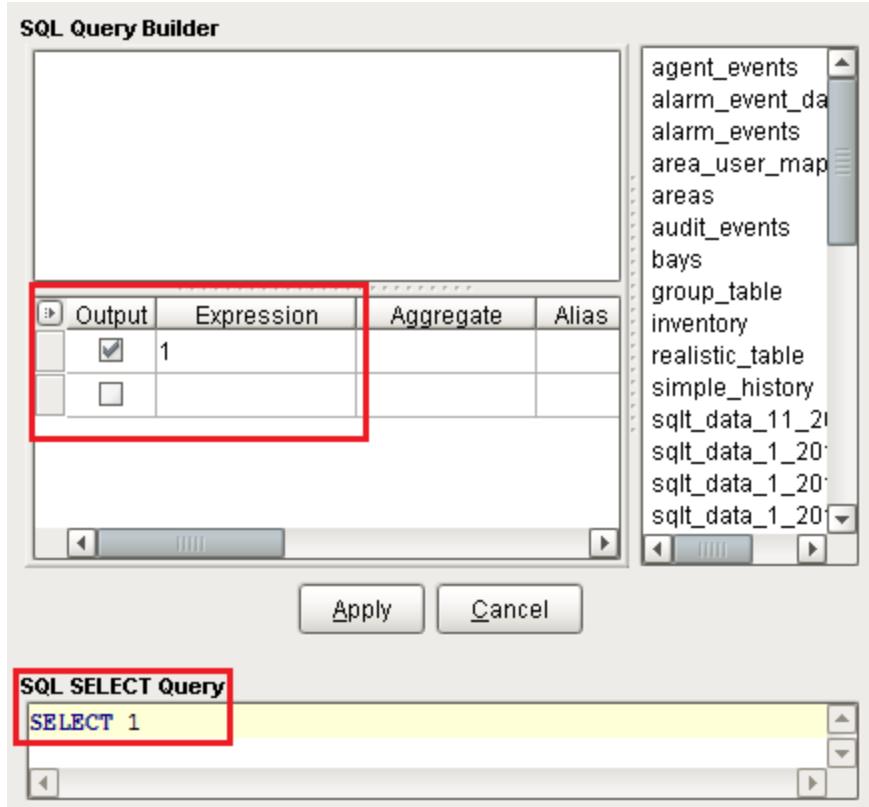


Clicking this button will open the Builder workspace with an empty query. If a query was already written before the button was pressed, then the Builder will attempt to load the query into the Builder's interface. For example, if the following query was typed before pressing the button:

SQL - Select Statement - Loading into the Builder's Workspace

```
SELECT 1
```

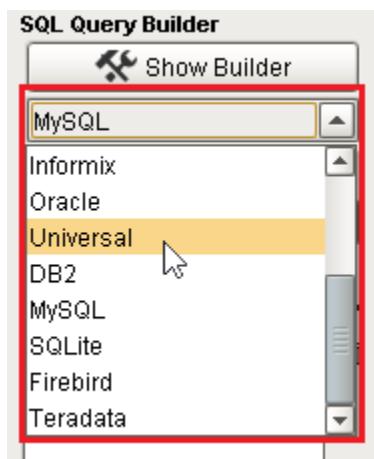
The following would be loaded into the Builder's workspace. Note that the **Building Area** is empty, because we're not querying from a table.



The Builder attempts to parse your query based on the **Syntax Parser**.

Using the Syntax Parser

When starting with a pre-existing query, the Syntax Parser tells the Builder which implementation of SQL your query's syntax is using. In most cases, this can be left with the default value of **Universal**. When set to a specific implementation of SQL, this allows the Builder to understand/accept implementation-specific keywords and syntax.



Syntax Parser in Action

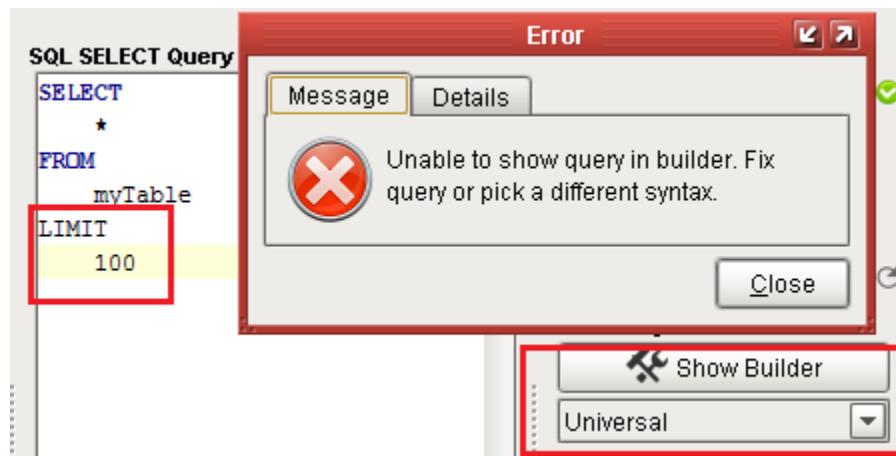
In MySQL, query results may be limited with the **LIMIT** keyword:

```
SQL - Limit Statement

SELECT
    *
FROM
    myTable
```

```
LIMIT  
    100
```

Attempting to open the Builder while using **LIMIT**, and the Syntax Parser is configured to a syntax that does not recognize the **LIMIT** keyword, will result in an error:

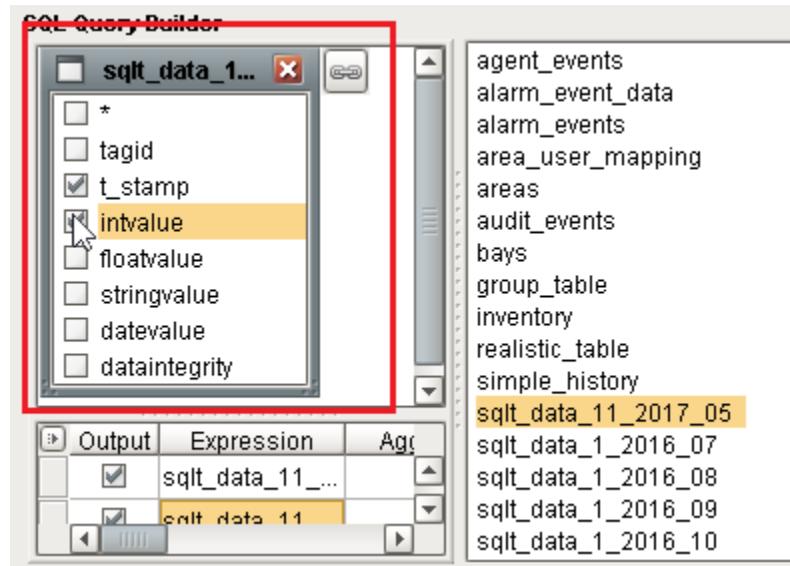


In these scenarios, you can switch to a syntax that supports the query, or remove the offending lines.

Using the Builder

Once the Builder is open, (and if you have tables in your database connection), you can start building a query by dragging and dropping a table from the Table List into the Building Area. From here, you can select which columns on the Table Object you want to bring into the query, or drag and drop them into the Columns Table below.

The * option at the top of the table is special. If the * character is selected, then all columns from the table will be included in the results just like a "SELECT * " query.



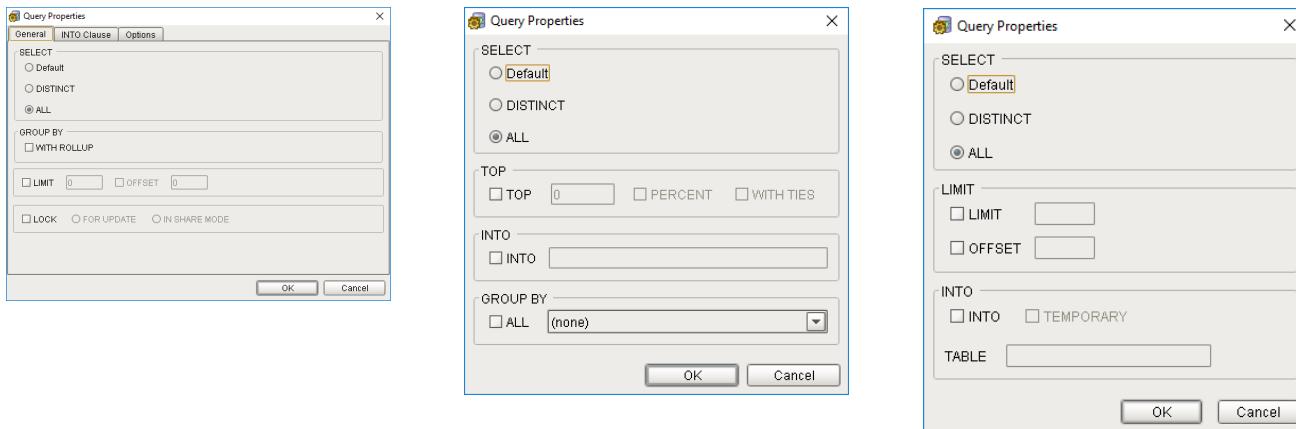
Query Properties

Right clicking on the empty space in the Building Area and selecting properties will bring up the Query Properties window that allows you to customize how the query works. How it looks and what it contains can vary, depending on what syntax parser you have selected.

[MySQL Query Properties](#)

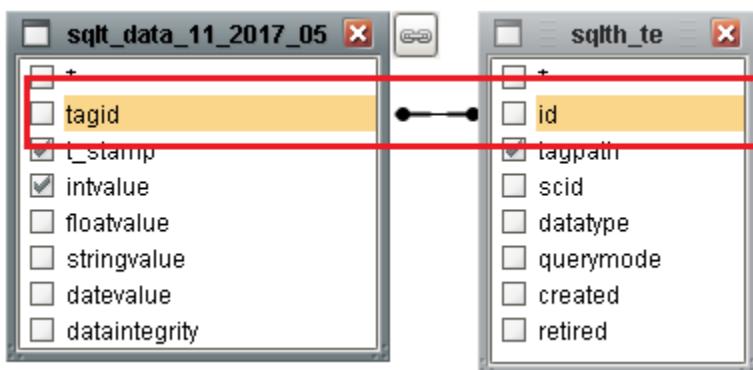
[MSSQL Query Properties](#)

[PostgreSQL Query Properties](#)



Joining Tables in the Query Builder

Multiple tables may be added to the Building Area. Once there, the Builder can JOIN the two tables by dragging from one column in a table on to another column on a different table.

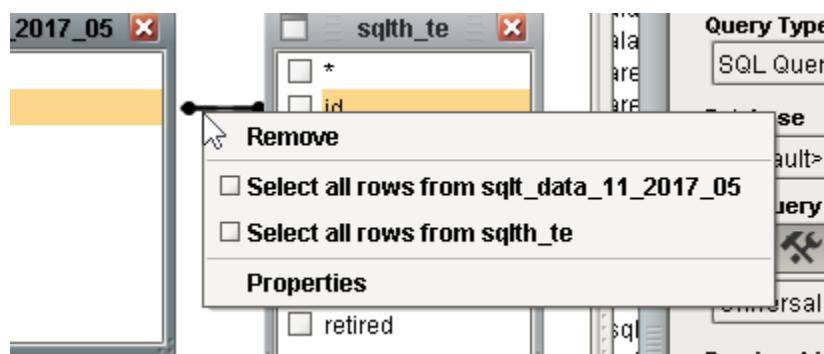


Notice that "tagid" and "id" are now linked together. This means the resulting query will JOIN the two tables based on the linked columns. Much like a SQL JOIN, the joining columns do not need to be selected to be utilized in the JOIN.

Joins Right-Click Menu

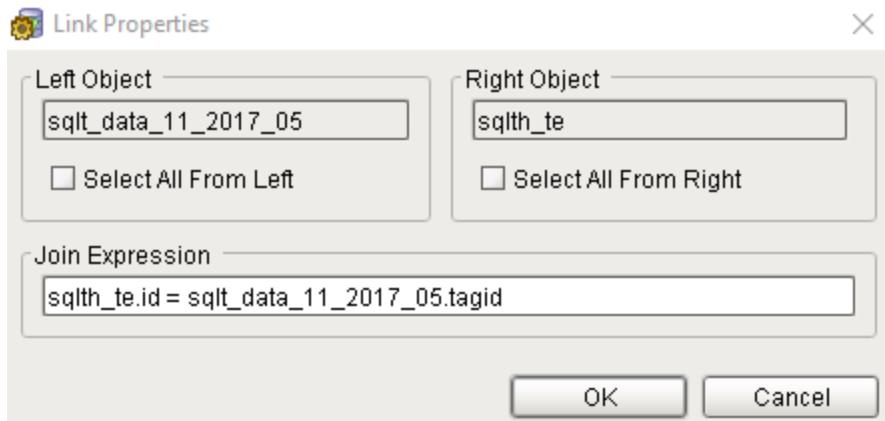
Right-clicking on the link will create a menu that allows you to remove the JOIN, as well as access properties of the JOIN.

Additionally, you can select all rows from one of the joining tables, which effectively creates a **RIGHT OUTER JOIN** or **LEFT OUTER JOIN**, depending on which table you select.



Join Properties

Clicking on **Properties** on the Right-Click Menu allows you to view the relation between each table, as well as apply **LEFT** or **RIGHT JOINS**, as mentioned above.



Columns Table

The Columns Table allows you to modify individual columns in the query. Typically, this allows you to add aggregates or aliases to each column from the Builder. Each row represents an expression, or combination of functions, columns, variables, and constants that will ultimately be a single column in the resulting query.

Output	Expression	Aggregate	Alias	Sort Type	Sort Order
<input checked="" type="checkbox"/>	sqlth_te.tagpath				
<input checked="" type="checkbox"/>	sqlt_data_11_2...				
<input checked="" type="checkbox"/>	sqlt_data_11_2...				
<input type="checkbox"/>					

Right-Clicking on an Item

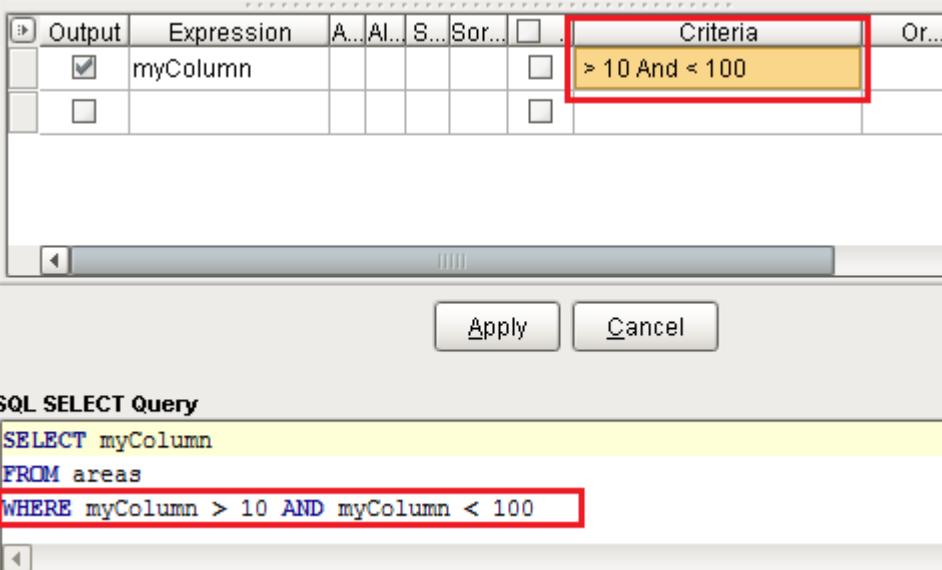
Right-clicking on a row of the Columns Table brings up a menu with the following options:

Option	Description
Move up	Moves the selected row up in the table. This means the column will appear in the query results prior to the other listed columns.
Move down	Similar to move up, but moves the row down, or towards the end of the listed columns.
Delete item	Removes the row from the Columns Table, effectively removing it from the resulting query.
Insert empty item	Adds a new row to the Columns Table with blank values.

Field Description

Each column in the Columns Table allows you to modify the resulting query in some way. The following is a description of each column:

Name	Description
Output	Specifies if the row should appear in the resulting query. Disabling a row on the table removes the column from the results.
Expression	The column from one of the tables in the Building Area that this row represents. Left-clicking on this field will create a dropdown list of possible rows.
Aggreg	

ate	<p>Allows you to aggregate the expression. Blank by default (no aggregate), configuring this column allows you to do things like sum column, or return the average. Possible values are:</p> <ul style="list-style-type: none"> • Blank (no aggregation will be performed) • Avg • Count • Max • Min • Sum <p>You can also add the Distinct keyword to any non-blank option. For example, you could enter "Sum Distinct" or "Count Distinct".</p>
Alias	Creates an Alias, or alternative name for the column. Effectively adds the SQL "AS" keyword to the column, which allows you to rename the resulting column.
Sort Type	Sorts the results based on this column. Enabling a Sort Type is similar to added an ORDER BY to your query. Possible values are: <ul style="list-style-type: none"> • Blank (no sorting on this column will be performed) • Ascending • Descending
Sort Order	When a Sort Type has been specified on multiple rows, this determines which row the query will sort on first.
Grouping	Allows you to group the results. Adds a GROUP BY statement to the resulting query. Enabling grouping on one of the columns makes the Criteria for column appear.
Criteria for	<p>A dropdown list that sets how you want the Criteria column to apply to the grouping. Does not apply if both Grouping and the WHERE clause Criteria are not being used.</p> <ul style="list-style-type: none"> • For Values will place the WHERE clause specified in the Criteria column first, filtering out rows that don't apply, and then applying the grouping. • For Groups will instead change the WHERE clause into a HAVING clause that is used to filter after the grouping has taken place.
Criteria	<p>Allows you to add a WHERE clause. Supports the use of both the OR and AND keywords for multiple conditions. Example: assuming a column named "myColumn", we could limit the results of our query to rows where myColumn has a value greater than 10 AND less than 100:</p>  <p>The screenshot shows the Ignition Query Builder interface. At the top, there's a toolbar with icons for Output, Expression, Alias, Sort, and Criteria. Below the toolbar is a table with columns: Output, Expression, Alias, Sort, Criteria, and Or... The 'Criteria' column contains the value '> 10 And < 100', which is highlighted with a red box. At the bottom of the dialog, there are 'Apply' and 'Cancel' buttons. Below the dialog, the 'SQL SELECT Query' pane displays the generated SQL code:</p> <pre> SQL SELECT Query SELECT myColumn FROM areas WHERE myColumn > 10 AND myColumn < 100 </pre> <p>The 'WHERE' clause is highlighted with a red box.</p>
Or...	Allows you to add an additional condition to a WHERE clause. Will separate each grouped condition with "(" characters to maintain logic.

Note: The Query Builder is a third party tool that we brought into Ignition and while we go over how to use it here, you can also check out the [Active Query Builder's documentation](#) for more information on how this feature works.

Related Topics ...

- [Named Queries](#)

- Named Query Data Source

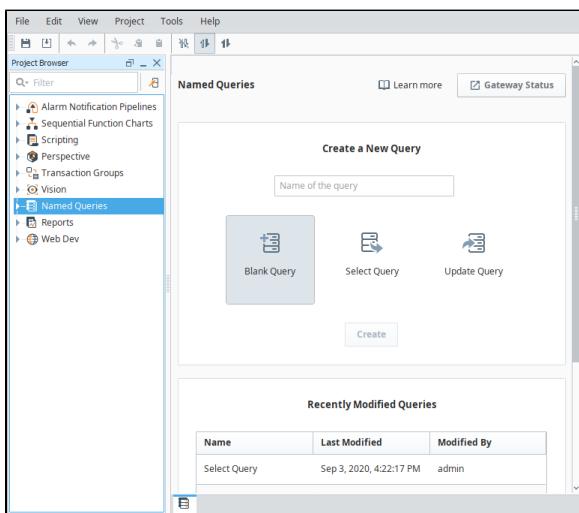
Named Queries

What Are Named Queries?

Named Queries are preconfigured queries that may be invoked elsewhere in a project. Named Queries are conceptually very similar to Project Scripts; defined in a single location, and then referenced in multiple places throughout the project. When executing a Named Query, parameters may be passed to return a dynamic result set. This way, a query may be written once, and then called from multiple locations throughout the project.

The Named Queries Welcome tab allows you to create three types of named queries. Each one of the named query types is basically a template to help you get started creating your own query. Once you select a Named Query type, enter a name, and press 'create', and the specific named query template will open. The Select Query and the Update Query will have some sample parameters and queries to help you get started. The Named Queries Welcome tab will show you any recently modified named queries along with the date it was modified and who modified it. You can even double click on a recently modified query and open it.

The Named Queries Welcome tab provides a quick way to create a new query and update existing ones.



Name	Last Modified	Modified By
Select Query	Sep 3, 2020, 4:22:17 PM	admin

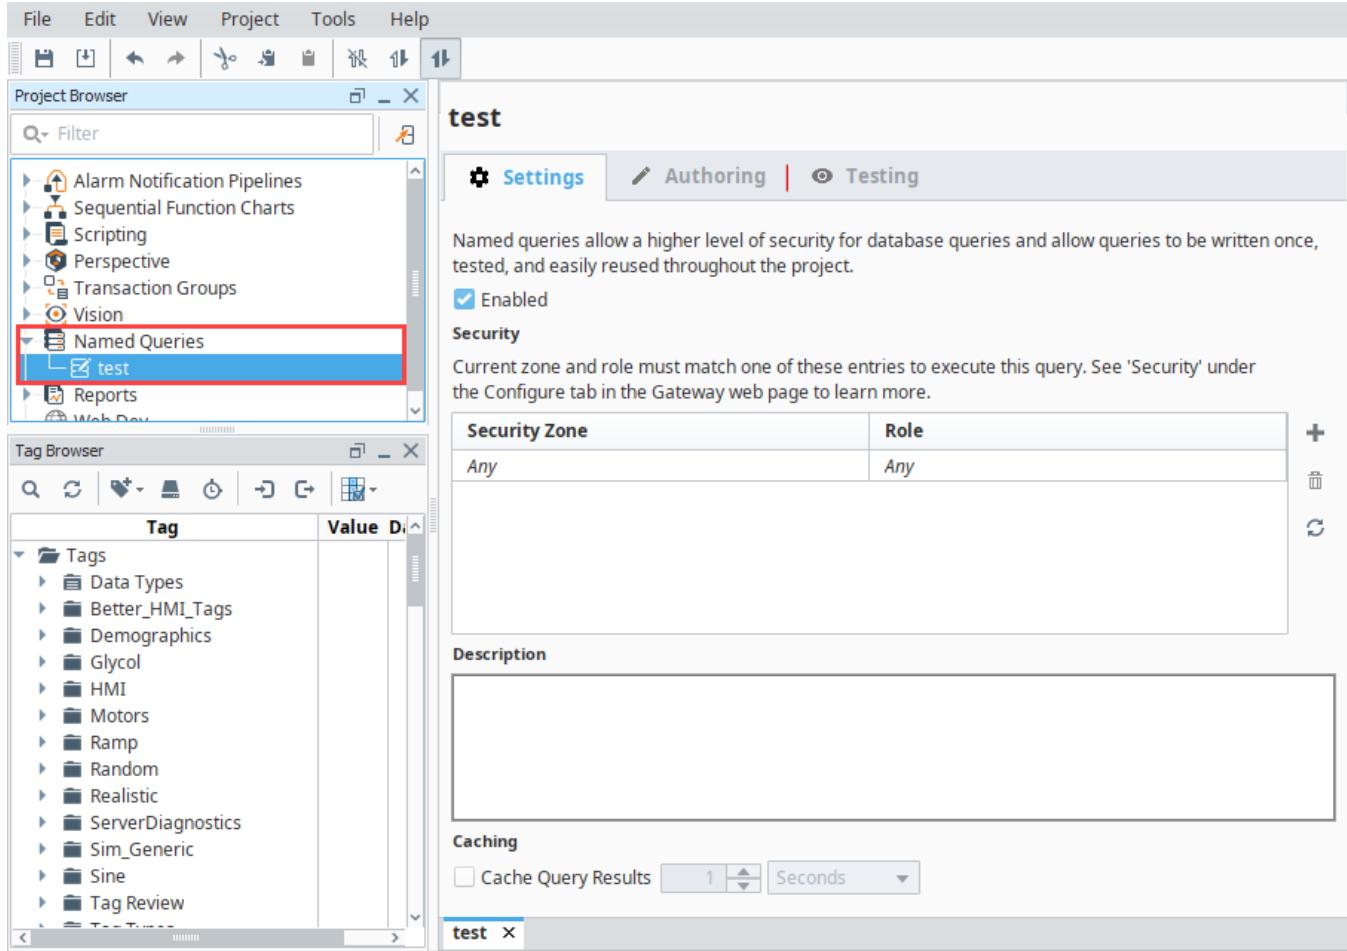
On this page ...

- [What Are Named Queries?](#)
- [Secure Query Execution](#)
- [Where Are Named Queries Used?](#)
 - [Binding](#)
 - [Reporting](#)
 - [Scripting](#)
 - [Named Query File Location](#)



Named Query Overview

[Watch the Video](#)



Secure Query Execution

While clients may request data from a Named Query, the actual execution of the query always takes place on the Gateway. Clients simply specify which query should run, and pass parameters that the Gateway will use. Additionally, the Gateway has an opportunity to [restrict access](#) to the query based on Security Zone and/or User Role. This provides a single interface to restrict access to the queries, and better protect your data. Additionally, queries cannot be modified by a Client other than by passing variables into it. This creates a very secure method to control what queries are being run against your database.

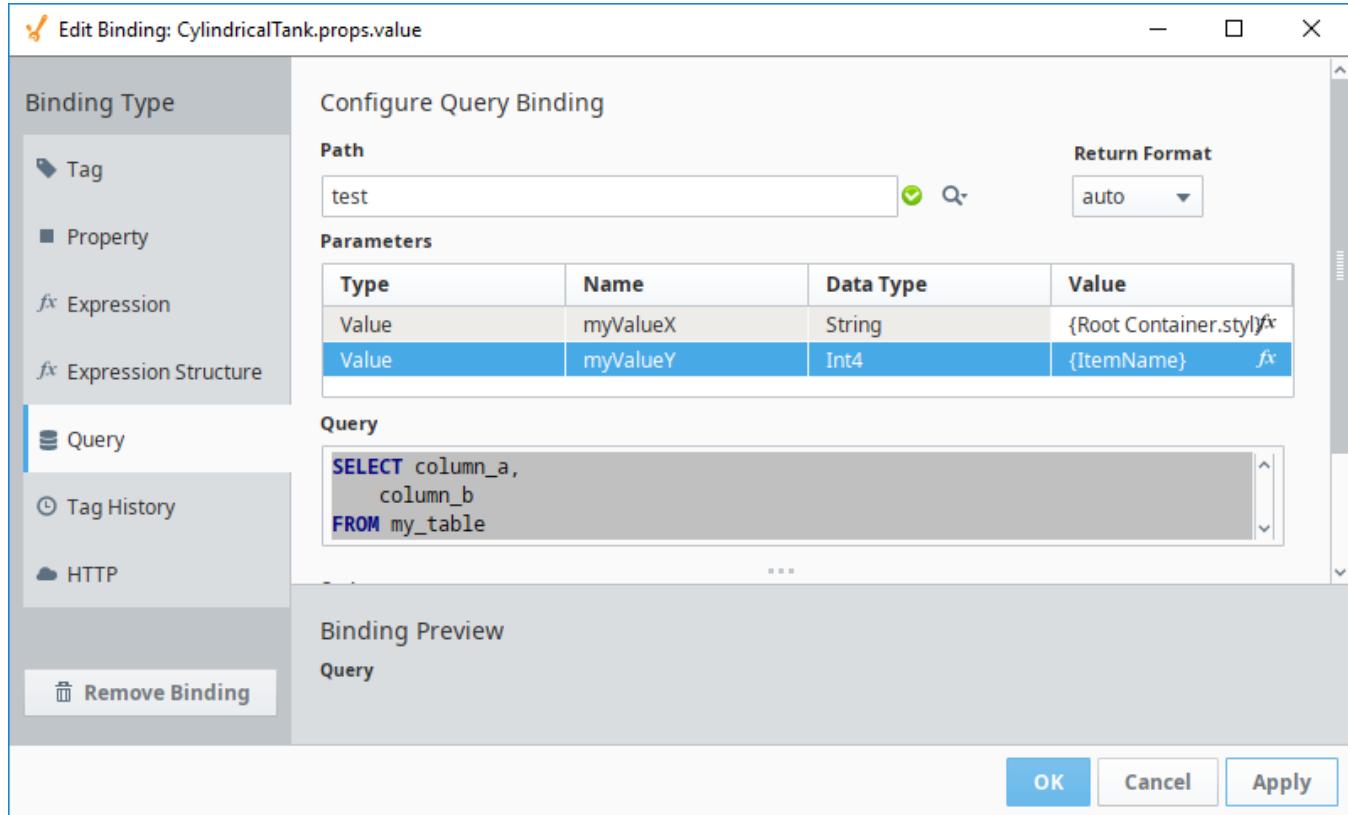
Where Are Named Queries Used?

Named Queries may be used in multiple locations in Ignition. They are used in all the same places that a normal SQL query would be used. Here are a few examples of locations in which named queries can be used. A more detailed example of a named query is provided [here](#).

Binding

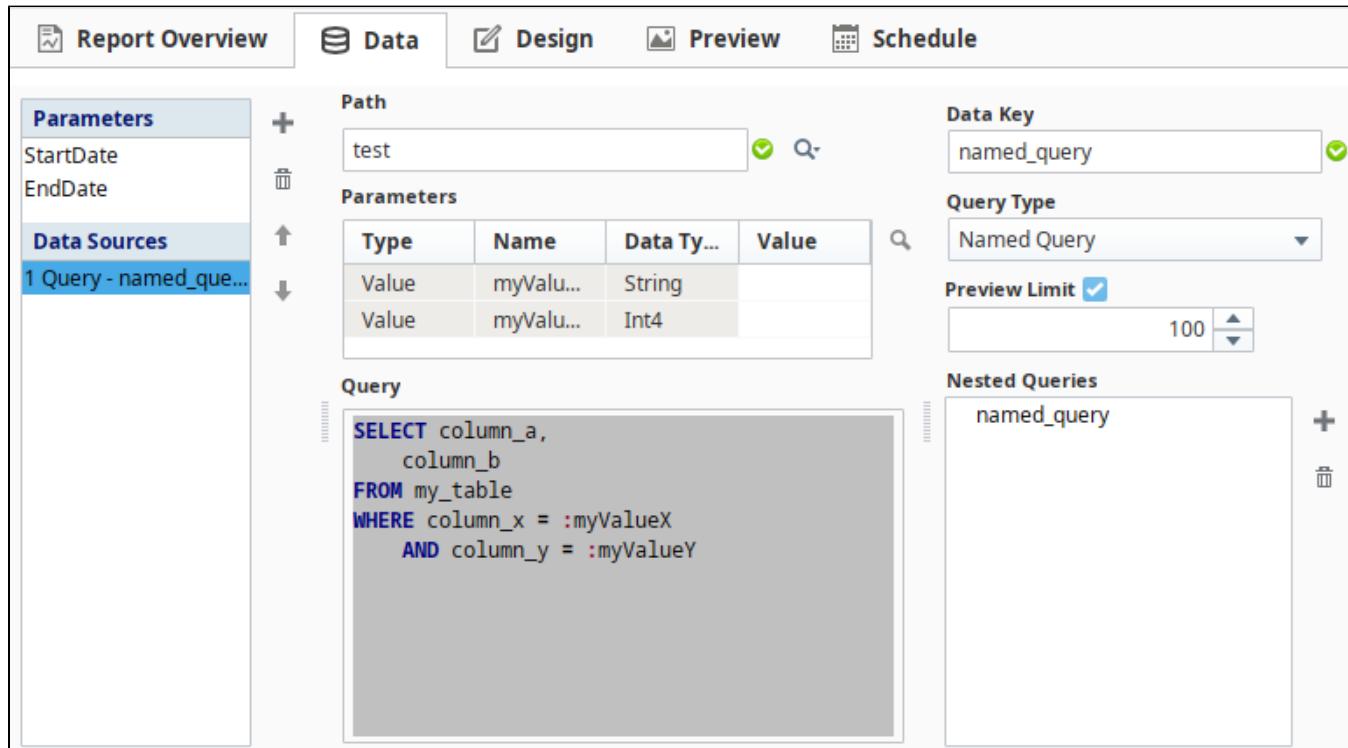
A [Query Binding](#) type has been added to leverage named queries for any component property bindings. [Named Query Parameters](#) may be bound to Tags or other properties on the same window, allowing your users to modify the resulting dataset by manipulating other components similar to the

original SQL Query binding type. You can also select and update a query to simulate a bi-directional binding to the database.



Reporting

A [Named Query Data Source](#) has been added to Reports. Report Parameters may be used by the Named Query to generate dynamic reports.



Scripting

The [system.db.runNamedQuery](#) function may be used to called a Named Query from any Python script. This provides a secure method to execute a query from any location in Ignition.

Named Query File Location

The following feature is new in Ignition version **8.1.6**
[Click here](#) to check out the other new features

The queries and metadata used by named queries are stored in SQL and JSON files respectively on the Gateway's file system.

```
%installationDirectory%\data\projects\PROJECTNAME\ignition\named-query\QUERYNAME
```

Where PROJECTNAME is the name of the project the Named Query resides in, and QUERYNAME is the name of the Named Query.

Named Query Workspace

Named Query Workspace Overview

Named Queries have a dedicated workspace inside of the Project Section of the Designer. This workspace allows for the creation and testing of Named Queries. Once created, the Named Query may be called from another resource, such as a [datasource in a report](#) or a component using a [Named Query Bindings](#).

Named Queries are created by right-clicking the Named Query item in the Project Browser. Like other resources in the Project Browser, Named Queries can be organized in folders, which creates a unique path to the query and helps keep your queries organized.

Note: Named Queries are referenced by path, so renaming the Named Query or any parent folders will require you to update the path on any other resources that are using it.

Also like other resources, multiple Named Queries may be opened in the same Designer session. Tabs at the bottom of the Designer allow for easy swapping between Named Queries.

On this page ...

- [Named Query Workspace Overview](#)
- [Workspace Sections](#)
 - [Settings](#)
 - [Authoring](#)
 - [Testing](#)

Workspace Sections

The Named Query workspace contains three tabs: Settings, Authoring, and Testing. A description of each section follows.

Settings

The Settings  tab contains configuration properties and security for the selected Named Query.

Test Query

 Settings

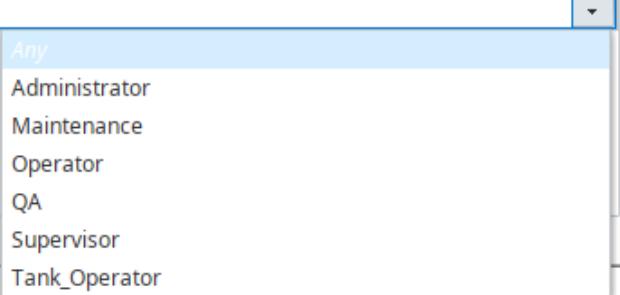
 Authoring |  Testing

Named queries allow a higher level of security for database queries and allow queries to be written once, tested, and easily reused throughout the project.

Enabled

Security

Current zone and role must match one of these entries to execute this query. See 'Security' under the Configure tab in the Gateway web page to learn more.

Security Zone	Role
Ridgefield East	 <ul style="list-style-type: none">AnyAdministratorMaintenanceOperatorQASupervisorTank_Operator   
Description	
Caching	<input type="checkbox"/> Cache Query Results  1 Seconds
Test Query 	

The following properties are available:

Item Name	Description
Enabled	Determines if the Named Query is enabled or disabled. A disabled Named Query may not be executed.
Security	Specifies a combination of Security Zones and Roles that may call the Named Query. Only roles in the projects user source will be available in the Role dropdown. Multiple rows may be configured to account for granular access restrictions (i.e., requests from Administrator roles originating from the Office security zone area could be allowed, while requests from same users in the plant floor zone could be denied). If the request does not match any of the specified zone and role combinations, then the query will not run. Additionally, if a ScalarQuery type has a Fallback value configured, that Fallback value will not be returned either: the query will not execute due to security settings, so there is never a chance for other errors to occur. Either the Security Zone or the Role (but not both) may be left blank. This means it is available to all objects of that type (i.e., with a blank Role and the "office" Security Zone, all roles in the "office" zone are valid).
Description	Allows you to give the Named Query a description.
Caching	Allows the Gateway to cache the results of the query. See the Named Query Caching page for more details.

Authoring

The Authoring  tab is where the query and parameters are created. There is also a Table Browser and [Query Builder](#) that can be used to help you to create your query.

Test Query

Settings

Authoring

Testing

Database Connection

<Parameter>



Query Type

Query

Parameters

Type	Name	Data Type
Value	myValueX	String
Value	myValueY	Int4

Table Browser

--

Query

```
1 SELECT column_a,  
2     column_b  
3 FROM my_table  
4 WHERE column_x = :myValueX  
5     AND column_y = :myValueY
```

Query Builder

Open Query Builder

Builder Syntax

Universal

Test Query X

The following feature is new in Ignition version **8.1.24**
[Click here](#) to check out the other new features

Note: Starting in Ignition 8.1.24, switching to another tab from the Authoring tab will automatically commit any changes to the Named Query's parameters.

Item Name	Description
Database Connection	<p>The database connection the Named Query should run against. In addition to a list of the database connections configured in the Gateway, this dropdown contains two unique values: <Default> and <Parameter>.</p> <ul style="list-style-type: none">• <Default>: The query will execute against the project's default database connection.• <Parameter>: The query expects the database connection name to be passed in as a parameter when called. This allows you to use a dynamic database connection. <p>Note:</p> <p>A parameter for the database connection does not need to be manually created in the Parameters section. Instead, the Named Query will have a special "Database" type parameter available when called. Below is an image of a Named Query binding that is utilizing the <Parameter> connection type.</p>

SELECT Query

Path

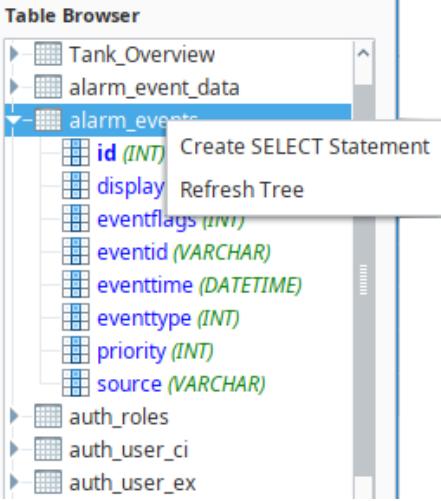
Test Query

Parameters

Type	Name	Data Type	Value
Database	database	String	
Value	myValueX	String	
Value	myValueY	Total	

Query Type	<p>SELECT Query</p> <p>Path</p> <p>Test Query </p> <p>Parameters</p> <table border="1" style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th>Type</th><th>Name</th><th>Data Type</th><th>Value</th></tr> </thead> <tbody> <tr> <td>Database</td><td>database</td><td>String</td><td></td></tr> <tr> <td>Value</td><td>myValueX</td><td>String</td><td></td></tr> <tr> <td>Value</td><td>myValueY</td><td>Total</td><td></td></tr> </tbody> </table> <p>The type of query to execute. The following options are available:</p> <ul style="list-style-type: none"> Query: Allows SELECT Queries and returns a full dataset. This type should be selected when running a SELECT statement that returns multiple rows or columns. ScalarQuery: Allows SELECT queries and returns a single value. This type should be used when running SELECT statements that only return a single value. The very first cell returned will be the only output. This type is special in that a Fallback value may be defined. The Fallback value will be returned if the Named Query would return an error. Note, the Fallback parameter will not be returned if the request does not meet the security requirements. UpdateQuery: Allows all UPDATE types of queries (queries that mutate or otherwise modify rows on a table) and returns the number of rows affected by the query. This type should be selected when modifying the database in some way, such as when running an INSERT, UPDATE, or DELETE query. 	Type	Name	Data Type	Value	Database	database	String		Value	myValueX	String		Value	myValueY	Total	
Type	Name	Data Type	Value														
Database	database	String															
Value	myValueX	String															
Value	myValueY	Total															
Parameters	A table of the parameter names and types that will be used in the query. These parameters have three types, Value, QueryString, and Database. Most commonly, the Value type is used and can be accessed by using the <code>:paramName</code> notation. More details on this field may be found on the Named Query Parameters page.																
Query	<p>The query that will execute when the Named Query is called. You can type directly into this field or use the Table Browser on the right to get started. Right-Clicking inside this field will cause a popup menu to appear:</p> <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;"> <p>Query</p> <pre>1 SELECT column_a, 2 column_b 3 FROM my_table 4 WHERE column_x = :myValueX 5 AND column_y = :myValueY</pre> </div> <div style="border: 1px solid #ccc; padding: 5px; background-color: #f9f9f9;"> <p>Parameterize ></p> <p>Insert Parameter > Make Value</p> <p>Undo</p> <p>Can't Redo</p> <p>Cut</p> <p>Copy</p> <p>Paste</p> <p>Delete</p> <p>Select All</p> <p>Themes ></p> </div>																

Most of the items on this menu are self-explanatory, but a few require special mention:

	<p>Parameterize: Contains two sub-menu items, which are detailed below. Note, that the sub-items will be disabled unless you right-click on some text that does not reference a parameter, table name, or column name in the query.</p> <ul style="list-style-type: none"> ◦ Make Value: Turns the selected text into a value-type parameter. The new value-type parameter will appear in the Parameter table above the query. ◦ Make QueryString: Turns the selected text into a QueryString. The new QueryString will appear in the Parameter table above the query. Note that QueryStrings are susceptible to SQL injection attacks. Because of this, the Make Value option is recommended over this option. <p>Insert Parameter: Quickly creates a reference to the selected parameter. This menu is an alternative to dragging-and-dropping from the Parameters table or typing the name of the parameter.</p>
Table Browser	<p>Provides a list of the tables in the selected Database Connection. Tables may be dragged into the Query field to quickly insert the name of the table. Additionally, right-clicking on a table in the list will cause a popup menu to appear:</p>  <ul style="list-style-type: none"> • Create SELECT Statement will populate the Query field with a SELECT statement targeting the selected table or selected row of the table. • Refresh Tree will refresh the Table Browser.
Query Builder	<p>Opens the Query Builder, which provides an easy way to create SQL queries using a drag-and-drop interface. This button will be disabled if the Database Connection property is set to <Parameter>. This is the same Query Builder used in other places like the reporting data page.</p>
Builder Syntax	<p>Specifies the syntax the Query Builder should use. Contains syntax for many popular databases, and has a Universal selection that should work in most scenarios.</p>

Testing

In the Testing  tab, you can test your query without leaving the workspace. Fill in your values and click the **Execute Query** button to see your results.

Test Query

Settings | Authoring

Testing

Test Parameters

Type	Name	Data Type	Value
Database	database	String	
Value	myValueX	String	
Value	myValueY	Int4	

Use Sample Size Limit

100

Execute Query

Export to CSV

Results

result

1 row fetched in 0.071s

Test Query

The following feature is new in Ignition version 8.1.28
[Click here](#) to check out the other new features

If your Named Query is using a DateTime parameter, you can expand the DateTime picker to easily set a DateTime value.

Item Name	Description
Test Parameters	Allows you to manually supply test values to the parameters to the queries. The table is populated by the Parameters field on the Authoring section.
Use Sample Size Limit	When checked, allows you to set the maximum number of rows the query will return while testing the Named Query. In addition this setting also limits the results of that named query called elsewhere in the Designer. This property is only enabled when the Query Type on the Authoring section is set to Query .
Execute Query	Runs the Named Query using the parameter values listed above. The Results area will display any results returned by the query.
Export to CSV	Exports the results of the query to a CSV file. The button becomes available after results are returned. Note that this will only return the values shown, check the Sample Size Limit when using this button.
Results	The results returned by the Named Query when testing. Populated by the Execute Query button.

Related Topics ...

- [Named Query Parameters](#)
- [Named Query Conversions](#)
- [Named Query Caching](#)

Named Query Parameters

Parameters allow you to make Named Queries dynamic. They act as placeholders you can pass values into when requesting the query to execute. Other resources in Ignition can then pass arguments into the parameters. The exact implementation depends on what resource is requesting the Named Query, such as a [Named Query Bindings](#), a [Named Query Data Source](#), or the `system.db.runNamedQuery` function. Check out the [Named Query Example](#) page for a complete example on passing parameters into a named query.

Types of Named Query Parameters

There are three types of Parameters in Named queries. Each varies in usage.

Value

The Value type should be used whenever a Named Query needs a dynamic WHERE clause. These act like values passed to a prepared statement, meaning they can never be used to parameterize column or table names. However, they are resilient to SQL injection attacks.

SQL Query - Using a Parameter

```
SELECT * FROM mytable  
WHERE name = :myParam
```

On this page ...

- [Types of Named Query Parameters](#)
 - [Value](#)
 - [QueryString](#)
 - [Database](#)
- [Parameters while Authoring a Named Query](#)
 - [Creating Parameters](#)
 - [Using Parameters in the Query](#)

QueryString

QueryStrings are more flexible than the Value type in that they can be used to parameterize column and table names. However, their values are never sanitized, which causes them to be more susceptible to SQL injection attacks. When using QueryStrings, it is best to avoid situations where the user can manually type in the value that will be passed to the Named Query. Additionally, if you are using a QueryString for a string in the where clause, you would need to provide quotation marks.

SQL Query - A Using QueryString

```
SELECT {myColumnName} FROM mytable  
WHERE name = '{myName}'
```

Database

Database type parameters cannot be created manually. Instead, it is automatically created when the **Database Connection** dropdown on the **Authoring** section is set to **<Parameter>**. Additionally, this parameter is not used in the body of the query. This type allows you to parameterize the database connection when the Named Query is called. This way the Named Query can run against multiple database connections specified by the resource that made the request.

Parameters while Authoring a Named Query

Creating Parameters

New parameters can be created in the Authoring section of a Named Query by clicking the Add  icon next to the Parameters table.

test

Settings Authoring Testing

Database Connection Query Type

<Default> Query

Parameters

Type	Name	Data Type
Value	myValueX	String
Value	myValueY	Int4
QueryString		
Value		

Query

```
1 | SELECT column_a,
```

i In Named Queries, Parameters are referenced by their name, so renaming the Parameter will require you to update it on any other resources that are using it.

Each parameter has three properties represented by different columns in the table. These may be edited by double-clicking on the cell you wish to modify:

- **Type:** Changes the type between Value and QueryString.
- **Name:** Determines the name of the parameter, and how it will appear in the query. Names are **not** case-sensitive and must be unique. Additionally, they may only use letters, numbers, dashes and underscores.
- **Data Type:** Specifies the datatype of the parameter. The Type of the parameter determines which data types are available. Note, that QueryStrings may only be configured as strings, where as Value-type parameters have more types available.

Using Parameters in the Query

Once created, parameters can be inserted into the Query field by **drag-and-drop** from the Parameter table onto the Query field, or by using the **right-click menu** in the Query field.

Query

```
1 | SELECT column_a,
2   column_b
3 FROM my_table
4 WHERE column_x = :myValueX
5 AND column_y =
```

Parameterize ▶

Insert Parameter ▶

myValueX
myValueY

Undo

Can't Redo

Additionally, the parameters may be typed in manually, but the correct syntax must be used.

Related Topics ...

- [Named Query Bindings](#)
- [system.db.runNamedQuery](#)
- [Named Query Conversions](#)
- [Using Named Queries - Example](#)

Named Query Caching

Overview

Named queries can opt-in to caching the results on the Gateway. This means if another request to the same Named Query comes in, the Gateway can return the cached result instead of having the database executing the query again. This will use more memory on the Gateway (to maintain the results), but could result in less queries running against the database.

Named Query caching is disabled by default, but can be enabled on the **Settings** section of each Named Query. If caching is enabled, the spinner and dropdown fields set the **lifespan** of the cache. Once the lifespan expires, the cache is invalidated. After the cache is invalidated, the next request for the Named Query will force the Gateway to re-execute the query and build a new cache.

Caching is especially useful for tables that are not updated often like recipe or inventory tables. Tables that update often like historical storage tables are bad candidates.

The screenshot shows the 'Settings' tab of a Named Query configuration. It includes sections for 'Enabled' (checkbox checked), 'Security' (checkbox checked), 'Description' (empty), and 'Caching'. Under 'Caching', there is a checkbox labeled 'Cache Query Results' which is checked, and a dropdown menu set to '5 Seconds'.

On this page ...

- Overview
 - When Caching is Disabled
 - Scenario
 - Considerations
 - Gateway Memory
 - How often the Database Values are Updated
 - How often the Named Query is Executed
 - Cached Query Updates and Designer Values
 - Scripting Functions to Clear Cache from a Named Query

When Caching is Disabled

Update queries are not allowed to cache their results. With UPDATE or DELETE statements, the work will already have been performed after the first execution, and INSERT statements typically utilize different parameters each execution. Thus if the Caching setting is disabled, then this typically means that the **Query Type** in the **Authoring** tab is set to an **Update Query**.

The screenshot shows the 'Authoring' tab of a Named Query configuration. It includes sections for 'Caching' (checkbox unchecked), 'Database Connection' (set to '<Default>'), and 'Query Type' (set to 'Update Query'). The 'Query Type' dropdown is highlighted with a red box.

Scenario

Consider the following:

- A Named Query is created. This runs a query to fetch data from an arbitrary table.
- A component in a project is configured to call the Named Query on a window.
- **Client_A** navigates to the window, which triggers a request for the Named Query be executed.
- Several seconds later, **Client_B** opens the same window, and needs the same results.

In this scenario, if caching was enabled on the Named Query, then **Client_B** would not cause another query execution (assuming both **A** and **B** passed the same values to the parameters). This would result in less network traffic between the Gateway and the Database, and less work for the Database. An example with two clients isn't exciting, but the same scenario with ~50 clients would mean a huge potential performance boost, especially if those requests were polling at a 5 second rate, and the cache period was configured to 5 seconds.

Considerations

The following are considerations that should determine whether or not caching the results of a Named Query is helpful.

Gateway Memory

Each time the Named Query is called, if the arguments passed differ from those used to create the current cache, a new cache is created. This means Named Queries that are frequently called with varying parameter values will create multiple caches. If the results are large datasets, this can result in a large amount of the Gateway Memory being tied up maintaining these caches. In this scenario, you will want to monitor the memory usage of the Gateway.

This is especially important to consider when dealing with queries that accept a timestamp parameter that uses an expression like now() that will return time to the current second. If 5 clients are opened with each one 1 second apart, they will all have different timestamps and create separate cache entries.

How often the Database Values are Updated

Once a cache is created, the Named Query will not look for any changes made to the database table until the query next executes. This means changes, such as adding a new row to the table, will not appear in the Named Query's results until the cache is invalidated. Data that changes often may not be a good candidate for caching. In this case, the lifespan of the cache should be set to a short amount of time, depending on how often the data may get updated.

How often the Named Query is Executed

If a Named Query is called frequently and there are not many updates to the table data, then there can be a huge performance benefit to caching the results. Fewer calls to the database result in less network traffic and better overall performance.

Cached Query Updates and Designer Values

Cached query results have some notable interactions in the Designer. Specifically:

- **Named Queries results executed in the Designer never cache:** This is because the Designer uses the live version of the query in the Designer instead of the saved Gateway version. This also prevents bad results being stored in the cache, and then appearing in the Clients.
- **Making changes to a Named Query and saving will invalidate all caches for that Named Query:** When a change is made to a Named Query and saved, the new query is pushed to the Gateway, which means all current caches are immediately outdated.

Scripting Functions to Clear Cache from a Named Query

You can also clear the Named Query cache using the following scripting functions.

- [system.db.clearNamedQueryCache](#) - Clears the cache of a single Named Query.
- [system.db.clearAllNamedQueryCaches](#) - Clears the caches of all Named Queries.

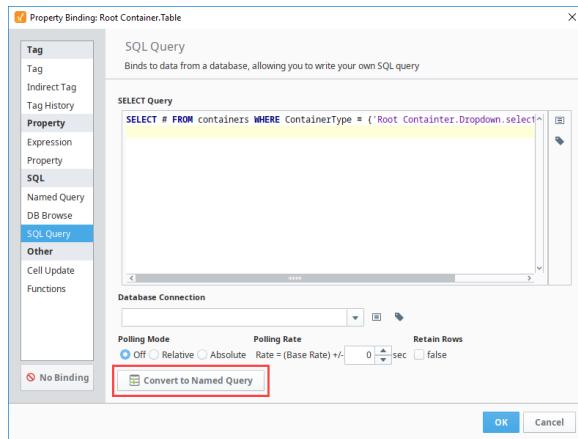
Related Topics ...

- [Named Query Conversions](#)
- [Using Named Queries - Example](#)

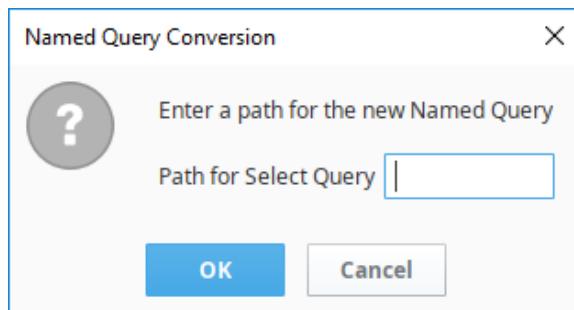
Named Query Conversions

Converting SQL Queries to Named Queries

If you have a SQL Query binding on a component, it is easy to convert it to a Named Query. Simply open up the binding and click the **Convert to Named Query** button.



A popup appears where you can enter a path to the new Named Query.



On this page ...

- [Converting SQL Queries to Named Queries](#)
- [Converting with an Update Query](#)
- [Modifying Converted Query to Use Parameters](#)



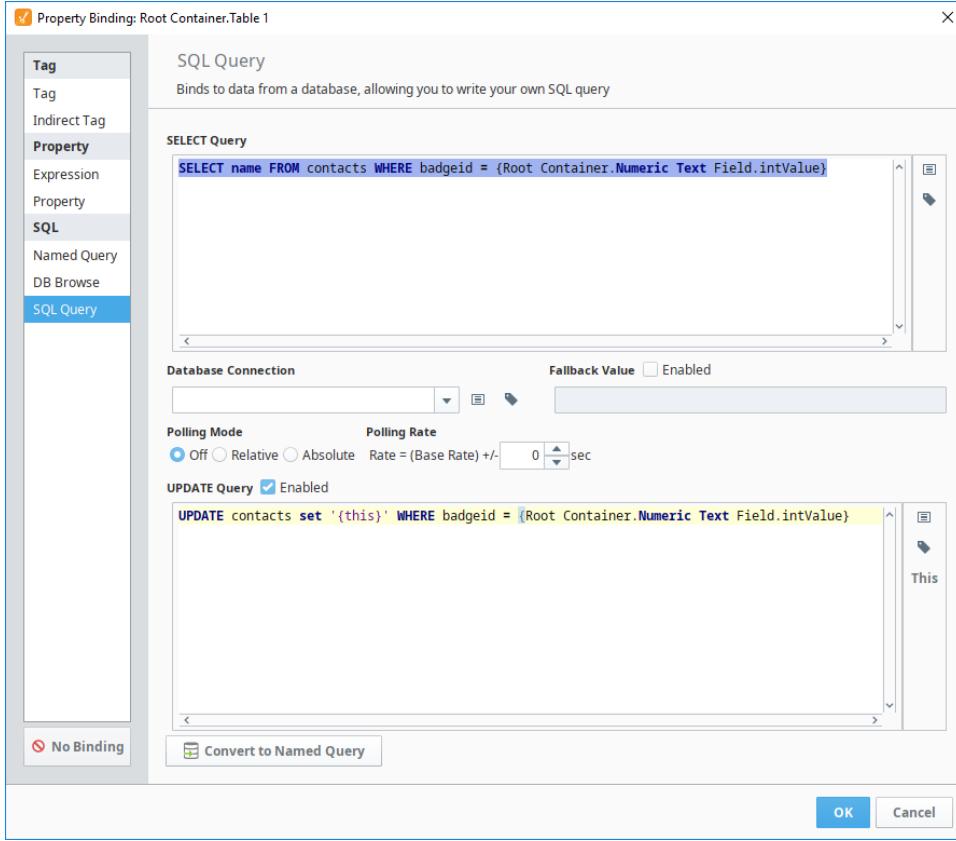
Named Query Conversion Example

[Watch the Video](#)

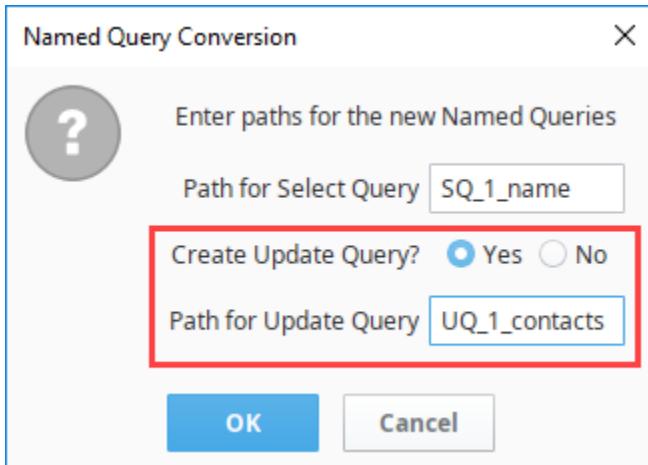
Converting with an Update Query

If your SQL Query binding is on an component property whose value can be updated and written back to the database (for example, input components and button components), you have the option to create two Named Queries: one for the Select Query in the binding and a second for the Update Query.

1. Enter the Select Query.
2. Click the Enabled checkbox for the Update Query.
3. Enter the Update Query.
4. Click OK.



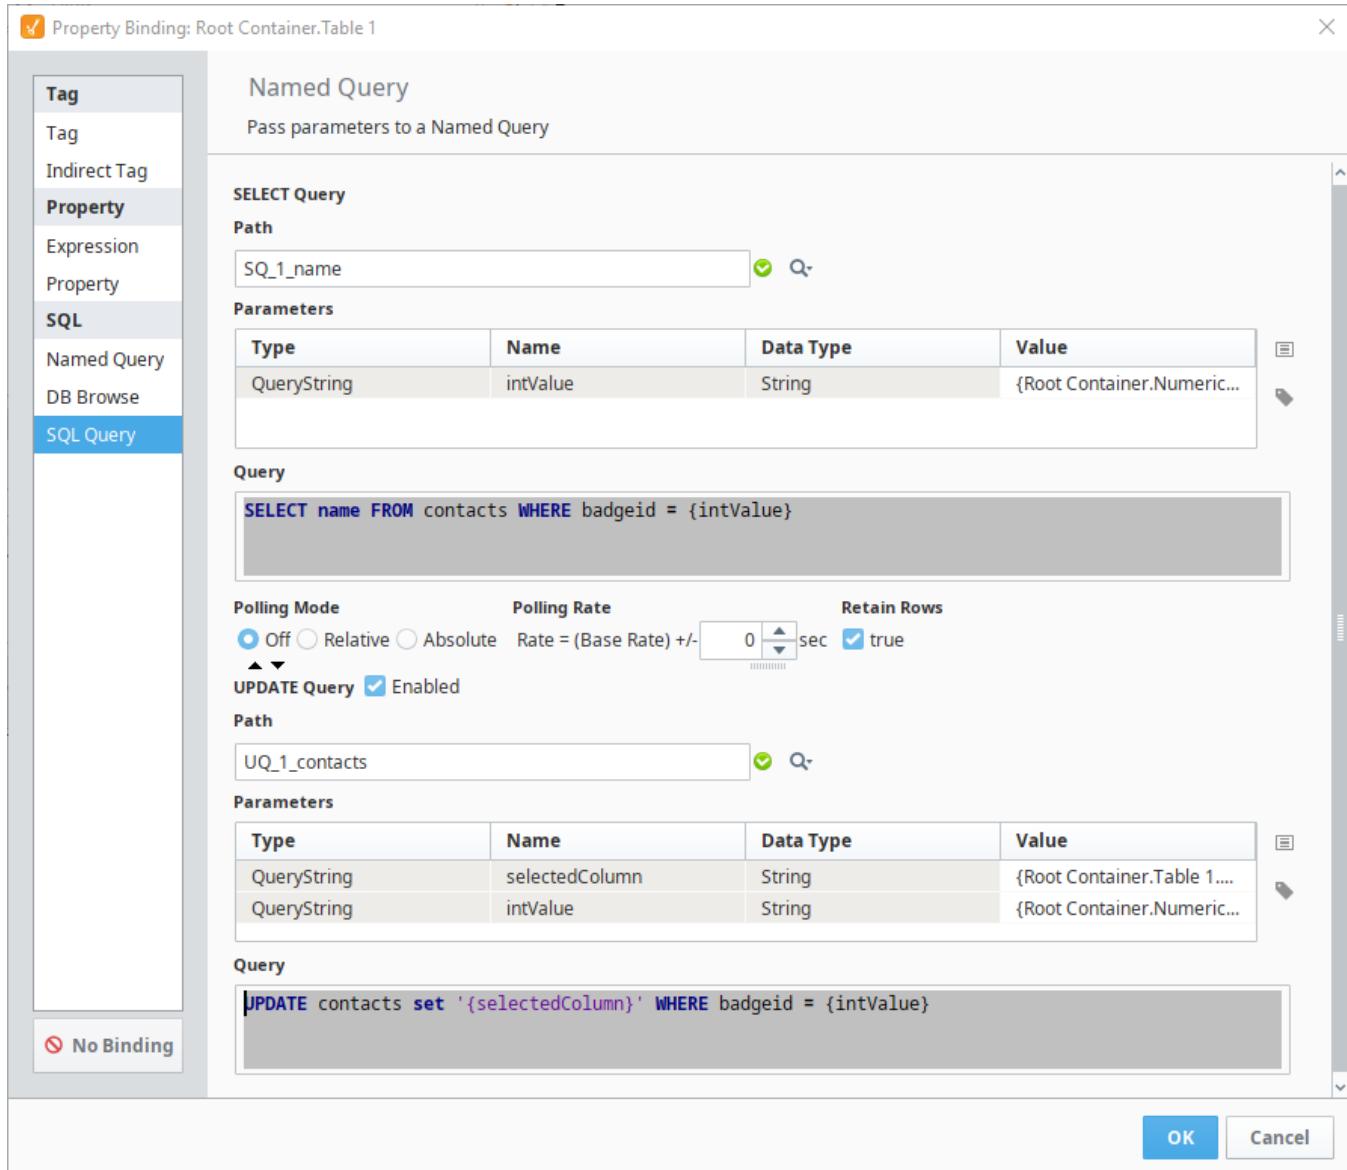
5. In the popup, enter the path for the Select Query.
6. Click the Yes button next to Create Update Query,
7. Enter the path for the Update Query,
8. Click Save.



Warning for 7.9.5 and Prior

When converting a SQL Query Binding that contains an UPDATE Query in 7.9.5 and prior, an Update query will not be generated. Thus, it was possible to lose the Update Query upon conversion.
If using these versions of Ignition, it is highly recommended to manually create the Update Named Query before pressing the Conversion button. As of 7.9.6, this is no longer an issue, and the **Convert to Named Query** button may freely be used.

After confirming the conversion, the Window will show the query/queries. Click the **OK** button to save this change.



When converting in this manner, all parameters will be created at **QueryStrings**. It is highly recommended that you modify your new Named Queries so that these values become **Value Parameters** instead, as mentioned later on this page.

Modifying Converted Query to Use Parameters

When converting a query to a Named Query, it is **strongly** recommended that you go back into the query and convert the parameter from a **QueryString** type to a **Parameter** type. To do this, you first need to change all **QueryString** type parameters to the **Parameter** type. This is as simple as selecting **Parameter** from the dropdown under **Type**.

Now that your parameter(s) types have been modified, the second thing you need to do is modify syntax in the query to use the parameter(s). The **Parameter** and **QueryString** types are referenced differently in the query, so you will need to ensure that you modify how the parameter is referenced. In addition because the **Parameter** type works like a prepared statement, it does not need any quotation marks around any string type parameters like a **QueryString** would, so all quotation marks around parameters should be removed from the query. See the images below for an example.

From this:

To this:

MyConvertedQuery

Settings Authoring Testing

Database Connection: <Default> Query Type: Scalar Query Fallback: Enabled

Parameters:

Type	Name	Data Type
QueryString	selectedStringValue	String

Query:

```
1 SELECT * FROM containers WHERE ContainerType = '{selectedStringValue}'
```

MyConvertedQuery

Settings Authoring Testing

Database Connection: <Default> Query Type: Scalar Query Fallback: Enabled

Parameters:

Type	Name	Data Type
Value	selectedStringValue	String

Query:

```
1 SELECT * FROM containers WHERE ContainerType = :selectedStringValue
```

Once the syntax in the query has been modified, test it out in the **Testing** section at the top to make sure everything works correctly. If so, then you have successfully converted to a Named Query.

Related Topics ...

- [Named Queries](#)
- [Named Query Bindings](#)
- [SQL Query Bindings in Vision](#)

Using Named Queries - Example

A Named Query Example

Here we will go over the steps necessary to put together a basic named query. This example will create a Named Query that uses a single parameter to run a select query, and then add a Table to a window and create a binding that uses our new Named Query.

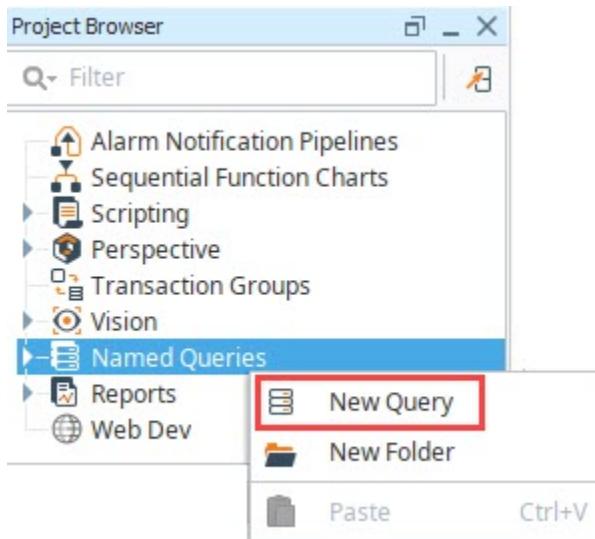
We are going to be querying a table in the database that holds information about products stored in storage bays. The database table is named "containers" and has the following structure.

id	StorageBay	ContainerType	ItemName	Weight	Time
1	1	Jug	Vanilla	25.2	2017-06-25 15:58:47
6	1	Mason Jar	Chocolate	12.3	2017-06-26 16:05:27
18	2	Tray	Swiss Cheese	88.8	2017-06-25 01:21:31
22	3	Tray	Cheddar Cheese	54.7	2017-06-25 03:52:16
23	3	Basket	Strawberry	36.8	2017-06-25 15:56:03
31	3	Jug	Whole Milk	80.1	2017-06-27 09:51:31
32	3	Jug	Fat Free Milk	76.9	2017-06-27 09:52:52

If you want to follow along with the example, feel free to make a database table that looks similar to this one and add as many rows of data as you would like, otherwise, you can use your own and substitute in the proper column names from your database.

Creating a Named Query and Adding Security

1. Start by opening up the Designer and [loading a project](#).
2. Locate the Named Query section of the project browser, right click on it, and select the **New Query** option.



3. You should now have a fresh Named Query that you can rename whatever you want. We used **FirstNamedQuery** in the example.

On this page ...

- [A Named Query Example](#)
 - [Creating a Named Query and Adding Security](#)
 - [Building the Query](#)
 - [Using the Query](#)

The screenshot shows the Project Browser interface with the 'Named Queries' section expanded. The 'FirstNamedQuery' item is selected and highlighted with a red box. The 'Settings' tab is selected, showing the following configuration:

- Enabled:** Checked
- Security:** Current zone and role must match one of these entries to execute this query. The table shows:

Security Zone	Role
Any	Any
- Description:** (Empty)
- Caching:** Cache Query Results: 1 Seconds

- To set up security on the Named Query, set required Security Zone and Role combinations in the Security table of the Settings tab. You can leave this blank if you don't have **roles** or **zones** set up yet.
- If multiple security combinations are required, use the Add  icon to add additional rows.

The screenshot shows the 'FirstNamedQuery' settings page. The 'Security' table is displayed, showing the following data:

Security Zone	Role
Any	Administrator
Production West	Maintenance
	Any
	Administrator
	Maintenance
	Operator
	QA
	Supervisor
	Tank_Operator

A red box highlights the 'Role' column for the selected row 'Production West'. A dropdown menu is open over the 'Role' field, showing a list of roles: Any, Administrator, Maintenance, Operator, QA, Supervisor, and Tank_Operator. The 'Maintenance' option is selected.

Building the Query

- Click on the **Authoring** tab. Here is where we do most of the work.
- Under **Database Connection**, we need to select a database connection that this named query will use. We selected **<Default>**.
- For the **Query Type**, we can decide what type of query this will be. For this example, we are running a select query that will return a dataset, so we chose **Query**.

4. In the Parameters section, we can decide on a list of parameters that will be used in this query. This query is fairly simple and will only use a single parameter.

Click the Add  icon to add a new parameter and set the following values:

Type: Parameter

Name: BayNum

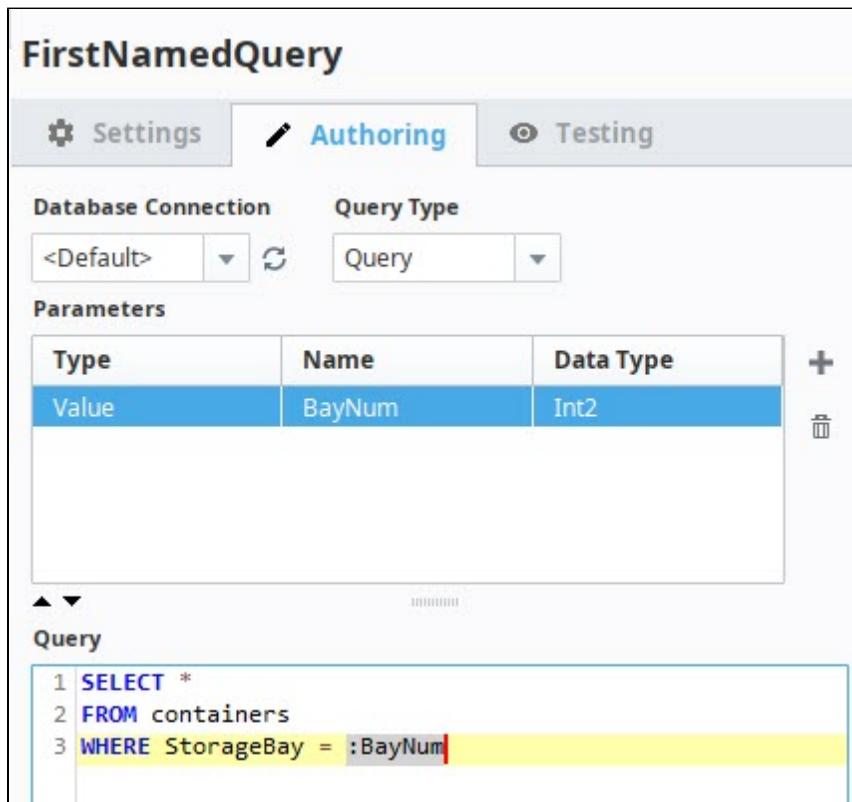
Data Type: Integer

5. The Query section below is where we construct our query using the Table Browser.

a. Right click on the containers table in the Table Browser on the right, and click on **Create SELECT Statement**. This will populate our query field with a basic select all statement.

b. Type into the **Query** field and add the following WHERE clause: "**WHERE StorageBay =**"

c. Now drag the **BayNum** parameter from the Parameters table to the end of the query you just typed. Notice "**:BayNum**" will be added at the end of the query.

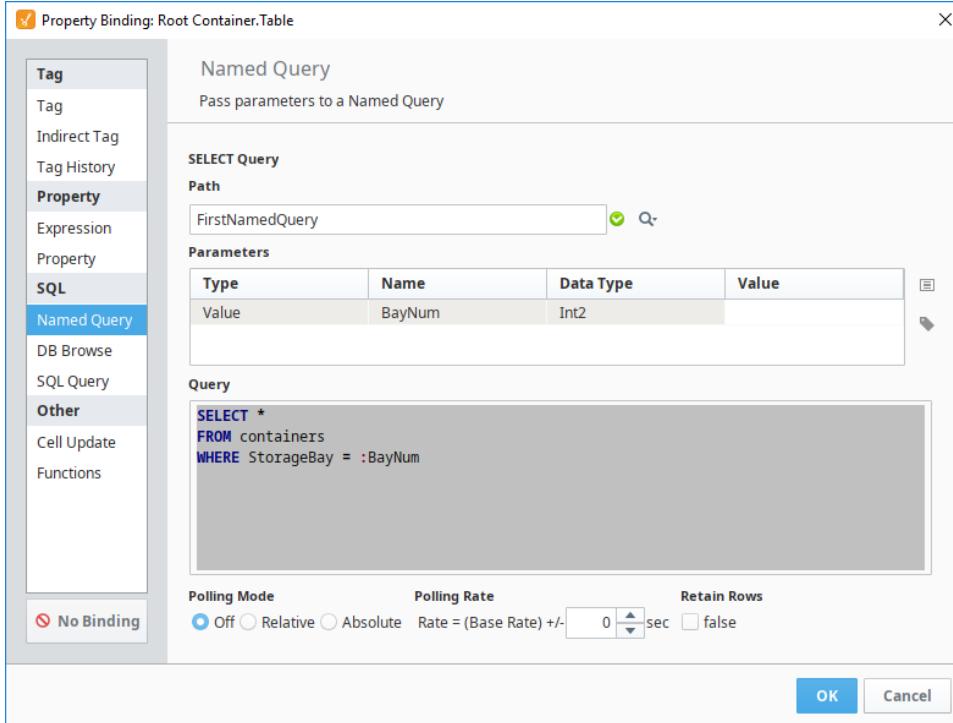


Type	Name	Data Type
Value	BayNum	Int2

Using the Query

- In the Project Browser, create or open a [Main Window](#).
- Drag a **Dropdown List** component onto the window.
- The Dropdown List is where we will be able to select a Bay Number to use as our Named Query's BayNum Parameter. Use the [Dataset Viewer](#) to set the **Data** property of the Dropdown to look like this:

Value	Label
1	Bay 1
2	Bay 2
3	Bay 3



4. Now drag a **Table** component onto the window. We can set up a Named Query binding on the Data property.
5. Click on the **Binding** icon for the **data** property and select the **Named Query** binding type.
6. Set the Path property by clicking on the **Select Resource Path** icon and selecting your new Named Query from the list. For this example, it is **FirstNamedQuery**. Alternately, you can type the name in.
7. Highlight your **BayNum** Parameter (in the Parameters table) and click on the Insert Property icon. We want to select the **Selected Value** property of our Dropdown List.
8. Finally, we want to ensure the Polling Mode is set to **Off**. This means the query will not run continuously, but will only run when it changes such as when a new bay is selected from the dropdown.
9. Click the **OK** button to save your binding, and put your Designer into **Preview Mode** to test it. We can then make a selection (i.e., Bay 3) with our Dropdown, and see the table populate with data.

id	StorageBay	ContainerType	ItemName	Weight	Time
22	3	Tray	Cheddar Cheese	54.7	Jun 25, 2017 3:5..
23	3	Basket	Strawberry	36.8	Jun 25, 2017 3:5..
31	3	Jug	Whole Milk	80.1	Jun 27, 2017 9:5..
32	3	Jug	Fat Free Milk	76.9	Jun 27, 2017 9:5..

Bay 3

Related Topics ...

- [Named Queries](#)

Queries in Scripting

Overview

In addition to using a binding, queries can be called from a Python script. This typically involves using one of Ignition's built-in [Scripting functions](#). This page presents several approaches to interacting with a database from a Python Script. See the [Scripting](#) section for more information about using Python.

Named Queries

Named Queries may be called from a Python script using the [system.db.runNamedQuery](#) function. Named Queries can execute as a Prepared Statement when not utilizing [Query String parameters](#), so they can offer a secure approach to calling a query from a script.

Additionally, Named Queries are incredibly easy to call from a script, since you only need to know their name, and don't have to bother with SQL syntax in the script. This is another way to protect your database from other users in the Designer, as you can [Protect Named Queries](#).

Python - Calling a Named Query

```
namedQuery = "Add New Order"
parameters = { "accountId":123, "productName": "Bananas" }

system.db.runNamedQuery(namedQuery, parameters)
```

For more information, refer to the [Named Queries](#) section.

On this page ...

- [Overview](#)
- [Named Queries](#)
- [Prepared Statements vs Standard Statements](#)
 - [Standard Statements vs Prepared Statements](#)
 - [How do I Call a Prepared Statement?](#)
 - [Standard Statements](#)
- [Stored Procedures](#)
- [Transactions](#)

Prepared Statements vs Standard Statements

A prepared statement is a feature in a database that executes a parameterized query. One of the main reasons to utilize a prepared statement is because they are resilient to SQL Injection Attacks. Because of this, we **highly recommend** you utilize Prepared Statements over Standard Statements, but Named Queries are the most secure. Especially in cases where the users can type values that will be included in the query.

Prepared Statements typically involve passing two pieces of information to the database:

- A query string that contains placeholders to represent where arguments will be passed in later. These are represented in the query as question mark characters ("?")
- A series of arguments to replace the placeholder characters in the static query.

Standard Statements vs Prepared Statements

Typical SQL insert queries look like the following:

SQL - Standard Statement

```
INSERT INTO orders (account_id, product_name) VALUES (123, 'Bananas')
```

A Prepared Statement instead looks like the following. Note, that the placeholders do not require quotation marks even when a string will be passed in.

SQL - Prepared Statement

```
INSERT INTO orders (account_id, product_name) VALUES (?, ?)
```

How do I Call a Prepared Statement?

Prepared Statements can be called from a script using specific functions. Typically, they contain "Prep" in the name such as [system.db.runPrepQuery](#), or [system.db.runPrepUpdate](#). When in doubt, take a look at the sub pages in the [system.db](#) section.

There are typically two required parameters with these functions: a string literal that represents the query, and a list of parameters.

Python - Calling an Insert Query as a Prepared Statement

```
query = "INSERT INTO orders (account_id, product_name) VALUES (?, ?)"  
args = [123, "Bananas"]  
system.db.runPrepUpdate(query, args)
```

Notable Prepared Statement Functions

- [system.db.runPrepQuery](#): Should be used when SELECT statements are used, as it returns a dataset containing the results. This function is effectively "read-only", as it does not manipulate data in the database.
- [system.db.runPrepUpdate](#): Should be used when manipulating the database in some way. When executing a statement that utilize INSERT, UPDATE, or DELETE, the runPrepUpdate function should be called. Note that it does return the number of rows affected, so the return value can be used to keep track of how much of an impact the query had.
- [system.db.runScalarPrepQuery](#): Like runPrepQuery above, but only returns the first column of the first row: i.e. a single value is returned instead of a full dataset. This is useful when using an aggregate function of some sort to return a count or total, as it saves your script the work of extracting the value from the full dataset that runPrepQuery normally returns.

Standard Statements

Queries can be called as a Standard Statement (a statement that that isn't a Prepared Statement) by using the [system.db.runQuery](#) and [system.db.runUpdateQuery](#) functions. However, these are susceptible to [SQL Injection attacks](#), and should be avoided where possible: especially when users have access to a keyboard and can directly type values that will be used in the query.

Calling a Standard Statement involves building the entire query as a single string, and passing the string on to our Standard Statement functions.

Python - Calling an Insert Query as a Standard Statement

```
query = "INSERT INTO orders (account_id, product_name) VALUES (%i, '%s')" % (123, "Bananas")  
system.db.run(query)
```

Notable Standard Statement Functions

- [system.db.runQuery](#): Executes a SELECT statement, returning a result set as a dataset.
- [system.db.runUpdateQuery](#): Executes a statement that manipulates the database in someway. Should be used with INSERT, UPDATE, and DELETE statements.
- [system.db.runScalarPrepQuery](#): Similar to runQuery, except only a single value is returned: the first column of the first row. Generally used in conjunction with SELECT statements that contain an aggregate function.

Stored Procedures

If your database administrator has already configured Stored Procedures for you to use, then they can easily be called from a Python Script. Using Stored Procedures in a script typically involves two main steps:

1. A SProcCall object is created with the [system.db.createSProcCall](#) function. The SProcCall object contains several functions that can be used to register parameters, and access the results set returned by the Stored Procedure after it has been executed.
2. The [system.db.execSProcCall](#) function must be used to execute the Stored Procedure.

Python - Creating and Executing a Stored Procedure

```
# Create a SProcCall object, which will be used to configure parameters on the Stored Procedure, and then  
# executed.  
myCall = system.db.createSProcCall("insert_new_order")  
  
# Register parameters on the SProcCall object.  
myCall.registerInParam(1, system.db.INTEGER, 123)  
myCall.registerInParam(2, system.db.VARCHAR, "Bananas")  
  
# Execute the Stored Procedure.  
system.db.execSProcCall(myCall)
```

Take a look at the [SQL Stored Procedures](#) page for more details.

Transactions

A SQL Transaction can also be executed from a script. For the unfamiliar, a Transaction is a batch of statements that will be executed together, and either succeed or fail as a group. Note, that the statements executed in the Transaction are not visible by other connections in the database until you commit them.

Transactions typically involve several steps:

1. Call `system.db.beginTransaction`. This returns a **transaction identifier** that can be used with other statements. Using this identifier is how you specify that a statement should be included in the transaction.
2. Start calling other statements with other functions, such as `system.db.runPrepUpdate`. The function's "tx" parameter will be passed the **transaction identifier**.
3. Commit or Rollback the transaction. Use `system.db.commitTransaction` to commit, and `system.db.rollbackTransaction` to rollback. These options are essentially the same as applying or canceling the results of the queries. Committing will make the updated results available to other database connections.
4. Close the Transaction once you're done, which can be accomplished with the `system.db.closeTransaction` function. This invalidates the **transaction identifier**.

Python - The SQL Transaction Workflow

```
# 1) Begin the transaction. This returns a transaction identifier that we can use with other statements.
transactionId = system.db.beginTransaction(timeout = 5000)

# 2) Now we can execute statements. Because we want these to run as part of the transaction, we need to
# include our identifier.
query = "INSERT INTO orders (account_id, product_name) VALUES (?, ?)"
args = [123, "Bananas"]
system.db.runPrepUpdate(query, args, tx = transactionId)

# 3) We can continue to add statements, but in this case we'll commit them. We could instead rollback if
# there was an issue with our previous statement.
system.db.commitTransaction(transactionId)

# 4) We're done, so close the Transaction.
system.db.closeTransaction(transactionId)
```

Related Topics ...

- [Writing Basic SQL Queries](#)
- [system.db.runPrepUpdate](#)
- [Named Queries](#)
- [SQL Stored Procedures](#)

Common SQL Tasks

This section contains examples of items we identified as "common" tasks. When first starting out with [SQL in Ignition](#), many users are looking for examples in order to grasp concepts, or examples to possibly use in their projects. Additionally, this section aims to demystify some of the more complex or abstract tasks that our users may encounter when working with SQL in Ignition.

The examples in this section document several types of SQL tasks that may also touch many other areas of Ignition. While these examples are typically focused on a single concept or end result, they can easily be expanded or modified. In essence, they serve as a great starting point for users new to Ignition, as well as experienced users that need to get acquainted with a new or unfamiliar feature.

Below is a list of common SQL tasks related to using SQL in Ignition.

Filter Data in a Table

The [Filter Data in a Table](#) section explains how using a simple Dropdown box and making some modifications to the SQL query can allow users to dynamically filter data coming into the table to only show a specific subset of data, or all of the data. This section also includes variations of the SQL query, thus returning different filtered results to the table.

On this page ...

- [Filter Data in a Table](#)
- [Add Data to a Database](#)
- [Edit Data in a Database](#)
- [Refreshing SQL Data on a Component](#)
- [Editing Multi-Selected Rows from Table](#)
- [Storing PDF Files in a Database](#)
- [Simple Database Editor](#)

Add Data to a Database

A common method of [adding data to a database](#) from within a client is using a Button component that executes an SQL query. This section provides an SQL script showing how to collect relevant properties and insert them into a database, as well as how to set up your components on a window.

Edit Data in a Database

[Editing data in a database](#) is very similar to how we add data to a database. Users can edit data within a table on the screen in realtime, and the changes are then pushed back to the database table. This section goes over the different ways to edit data in the database depending if you're using a Table component or Power Table component.

Refreshing SQL Data on a Component

Sometimes, it may be unnecessary to have a table constantly updating and querying the database for data. By using the [system.db.refresh](#) function, we can turn Polling "Off" on our query, and have a button that allows us to manually update the table with new data when the Button is pressed. Alternatively, we can add the [system.db.refresh](#) function at the end of a script to refresh the newly entered data automatically without pressing a button.

Editing Multi-Selected Rows from Table

Tables have the ability to select and [edit multiple rows](#) at a time. Oftentimes, this can be used to manipulate multiple rows simultaneously, such as deleting all of the selected rows at once.

Storing PDF Files in a Database

The database is a powerful tool and can allow you to store files such as PDFs in a database table. The [Storing PDF Files in a Database](#) section explores how we can take a PDF file and store it in the database table, as well as how to pull it back out and display it in the [Vision - PDF Viewer](#) component.

Simple Database Editor

It may be beneficial to have a frontend to a database table built right into an Ignition project that allows you to control the data in the table without having to go into the database's built-in frontend, such as MySQL workbench. In the [Simple Database Editor](#) section, you will learn how to build a simple database table editor on a window that can add, edit, and delete data from the database.

Related Topics ...

- [Writing SQL Queries](#)
- [Scripting Examples](#)
- [Common Reporting Tasks](#)

In This Section ...

Filter Rows in a Table

Creating the Components

Filtering table data by using a Dropdown box is possible if the **Table** component's **Data** property is bound to a SQL query. We can create a dynamic WHERE clause that will allow us to select all the data or a specific subset of it. To implement this solution, you must have a **Table** component and a **Dropdown List** component on a window. With a Dropdown component, we can create an option for each way of filtering the data, each with a unique value associated with it. An additional "all" entry can also be added, with its own unique value.

We can then use the value of the Dropdown to drive the query. When you select a different option from the Dropdown component the binding gets re-evaluated on the table's data binding resulting in the query executing the WHERE clause with the new parameters. You may want to turn Polling Mode off on the tables Data property binding in order to limit the periodic querying of the database.

Here we can use a value of 1 to show only entries which match that filter column value.

Pseudocode - Filtering for Area 1

```
SELECT * FROM table WHERE filterColumn = 1
```

While the following query will return everything, but will still evaluate the query's WHERE clause.

Pseudocode - Return All Rows

```
SELECT * FROM table WHERE 0=0
```

In this case 0=0 will always evaluate as true, therefore, the query will return every row.

We can combine the logic of each WHERE with an OR into a single query, and substitute in our Dropdown value. The SQL Query binding would then look something like this:

SQL - Where Clause Combined with OR Condition

```
SELECT * FROM table WHERE filterColumn = {dropdown value} OR 0 = {dropdown value}
```

This way, the WHERE clause will only be true when the filter column matches with our selected option, or it will return all rows, if we set up our "all" option to have a value of 0.

Example - Filtering Data on an Area Number

We can put together an easy example on how this might with data that stores what machines are in which areas. Our data table should look something like this:

id	machine_name	area_number
1	Conveyor	1
2	Press	2
3	Tank	1
4	Packer	3
5	Loader	3
6	Oven 1	3
7	Oven 2	2
8	Wrapper	1
9	Mixer	3

On this page ...

- [Creating the Components](#)
- [Example - Filtering Data on an Area Number](#)



INDUCTIVE
UNIVERSIT

Filter Data in Table

[Watch the Video](#)

10	Cold Storage	2
11	Dryer	2

We can then put together the query and components necessary to get this working:

1. Create a new Named Query. Set up security to fit your needs, and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - a. Create a single Value type Parameter that is an Int2 data type. I called mine dropdownValue
 - b. Add in the combined query that we went over above, but use the machine table name and column names.

SQL - Selecting Values from a Table and Filtering on a Dropdown

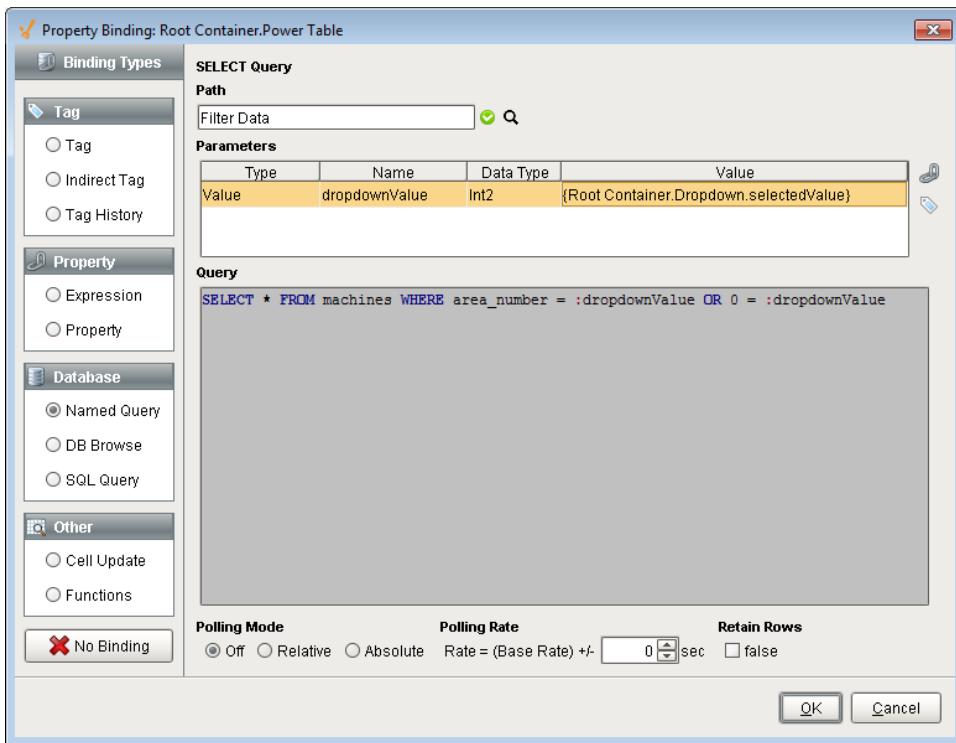
```
SELECT * FROM machines WHERE area_number = :dropdownValue OR 0 = :dropdownValue
```

2. Create a new Main Window and add a Power Table component and a Dropdown component to the window.
 - a. On the Dropdown's **Data** property, create a dataset that looks like this:

Value	Label
1	Area 1
2	Area 2
3	Area 3
0	All Areas

3. On the Power Table component, create a Named Query Binding on the **Data** property.

- a. Select the Named Query that was created in step 1.
- b. For the parameter, bind its value to the **Selected Value** property of the Dropdown component, and click the **OK** button.



4. Put the Designer into **Preview** mode and try out the Dropdown Filter. You will see that selecting a different value in the dropdown filters the data coming back so that only certain rows are shown.

id	machine_name	area_number
1	Conveyor	1
2	Press	2
3	Tank	1
4	Packer	3
5	Loader	3
6	Oven 1	3
7	Oven 2	2
8	Wrapper	1
9	Mixer	3
10	Cold Storage	2
11	Dryer	2

Related Topics ...

- [Writing SQL Queries](#)

Inserting Data into a Database

Inserting Data on a Button Press

A common way to insert information into a database is to execute a SQL query after the user presses a Button. The button can run a script to collect the information it needs and then execute a SQL `INSERT` statement to push that data into the correct Database table. A script executed on the Button's `actionPerformed` event handler would collect the relevant properties and insert them into a database. The script would look like this:

Pseudocode - Collects Data and Insert into the Database

```
value1 = {component property reference}  
value2 = {tag value}  
value3 = {static value}  
  
query = "INSERT INTO table (col1, col2, col3) VALUES (?,?,?)"  
args = [value1, value2, value3]  
system.db.runPrepUpdate(query, args)
```

Notice that the script isn't limited from taking values from any one place, grabbing tag values and property values and inserting them into a table. It is important to note that using the `runPrepUpdate` function will require that certain [Client Permissions](#) have been enabled.

On this page ...

- [Inserting Data on a Button Press](#)
- [Example - Inserting Values](#)



INDUCTIVE
UNIVERSITY

Add Data to Database

[Watch the Video](#)

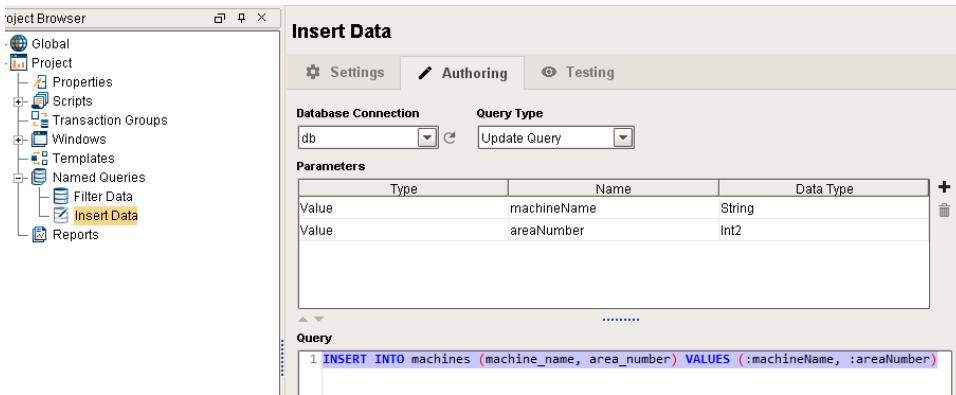
Example - Inserting Values

Say we have a table called `machines` in a database with three columns: an `id` column, a `machine_name` column, and an `area_number` column. We can build a query and a script that will insert the data into the database:

1. Create a new Named Query. Set up security to fit your needs and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - a. Set the Query Type to Update Query.
 - b. Create two Value type Parameters.
 - i. The first will be a string data type and will be for the name of the machine so it can be called `machineName`.
 - ii. The second will be an Int2 data type and will be for the area number so it can be called `areaNumber`.
 - c. Create the insert query.

SQL - Inserting Values from Components on the Window

```
INSERT INTO machines (machine_name, area_number) VALUES (:machineName, :areaNumber)
```



2. On a Main Window, add a Button component, a Text Field component, and a Dropdown Component.

- a. On the Dropdown's Data property, create a dataset that looks like this:

Value	Label
1	Area 1
2	Area 2
3	Area 3

- b. On the Button's Text property, change the value to something like "Submit"



3. Right Click on the Button and select Scripting. Navigate to the Script Editor tab on the actionPerformed Event Handler.

- a. Copy in this script, which will pull in the value from the Dropdown and the Text Field and insert them into the table using the Named Query we built in step 1.

```
Python - Insert Values into a Database Table

# Grab the area number and machine name from the components we added to the window.
areaNum = event.source.parent.getComponent('Dropdown').selectedValue
machineName = event.source.parent.getComponent('Text Field').text

# A call to our Named Query, inserting the two parameters using dictionary syntax.
system.db.runNamedQuery("Insert Data", {"machineName":machineName, "areaNumber":areaNum})
```

4. Test it out by selecting an area, entering in a machine name, and clicking the submit button. You can check out the new rows using the [Database Query Browser](#) to select the values from that table.

Related Topics ...

- [Simple Database Editor](#)

Updating the Database through the Power Table

Updating Table Data

Table components can do more than show data from a database. A properly configured Table can make the data of the Table accessible to the Client and allow the user to edit the data in realtime. This page shows two examples using the Table and Power Table components editing data in a database in realtime.

Suppose your database had a Table like this:

id	UserName	FirstName	LastName	Notes
1	JS	John	Smith	Likes bikes
2	LJ	Luke	Johnson	Lives in town
3	PB	Peter	Burke	Enjoys cooking

On this page ...

- [Updating Table Data](#)
- [Using a Power Table Component](#)
- [Using a Table Component](#)



INDUCTIVE
UNIVERSITY

Edit Data in Database

[Watch the Video](#)

Using a Power Table Component

Let's create a script on a Power Table component and the **onCellEdited** extension function that will write to the cell in the database.



This function will not work using default [Client Permission](#) settings. The Legacy Database Access will need to be enabled for this to work.

1. Drag a Power Table component to your window, and bind the **Data** property of the Power Table using a SQL Query Binding Type and the query below.

SQL - Selects All Records

```
SELECT * FROM users
```

2. Right-click on the **Power Table** and select **Customizers > Table Customizer**.
3. Select the **Editable** checkbox for each of the columns that we want to edit (i.e., same as the above example). Note, the **id** column should not be editable. Click the **OK** button.
4. Right-click the **Power Table** and select **Scripting**.
5. Under the Extension Functions folder, select the **onCellEdited** function, and check the **Enable** box. We can put a script in here that will grab the **id** column value of the row that we modify, and then use that and the new value that we entered to update the Database table. The extension function makes this easy, because it provides variables for the column name of the row we are editing, the row number that was edited, as well as the new value. Paste the following script in the scripting area.

Python - Writes to a Cell in the Database

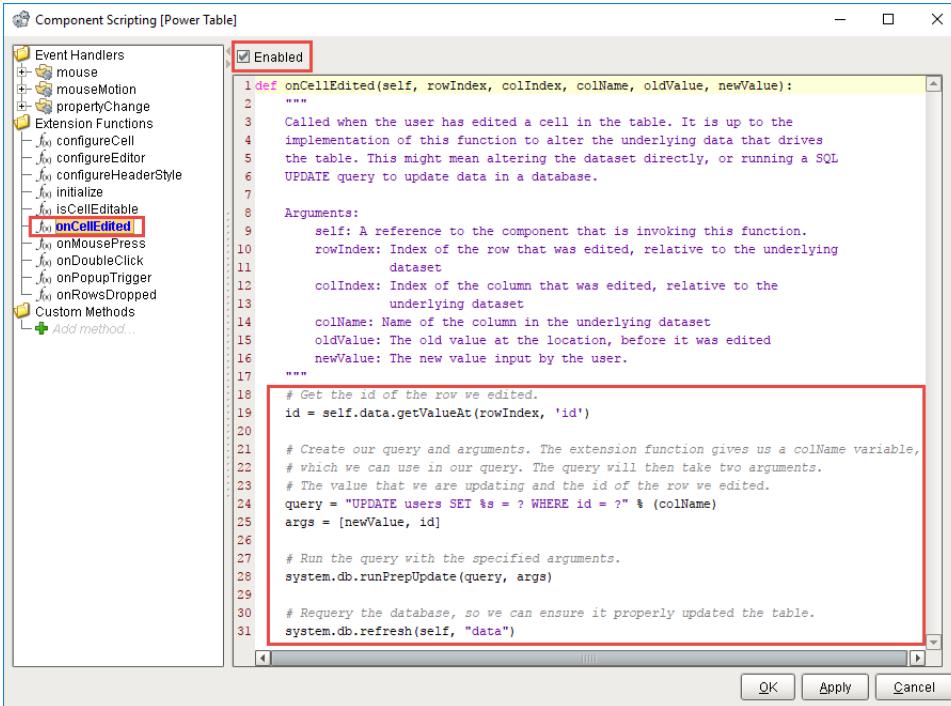
```
# Get the id of the row we edited.  
id = self.data.getValueAt(rowIndex, 'id')  
  
# Create our query and arguments. The extension function gives us a colName variable,  
# which we can use in our query. The query will then take two arguments.  
# The value that we are updating and the id of the row we edited.  
query = "UPDATE users SET %s = ? WHERE id = ?" % (colName)  
args = [newValue, id]
```

```

# Run the query with the specified arguments.
system.db.runPrepUpdate(query, args)

# Requery the database, so we can ensure it properly updated the table.
system.db.refresh(self, "data")

```



- Now, let's test it out. Put the Designer in **Preview Mode**, select the cell you want to edit, hit enter to commit the change, or tab to the next cell to make additional edits.

id	UserName	FirstName	LastName	Notes
1	JS	John	Smith	Likes bikes
2	LJ	Luke	Johnson	Lives in town
3	CP	Casey	Peters	Loves Ignition

Using a Table Component

Let's create another script like the example above, but this time we'll use a Table component's **cellEdited** event that will write to the cell in the database when we update it on the component. The script will be slightly different.



This function will not work using default [Client Permission](#) settings. The Legacy Database Access will need to be enabled for this to work.

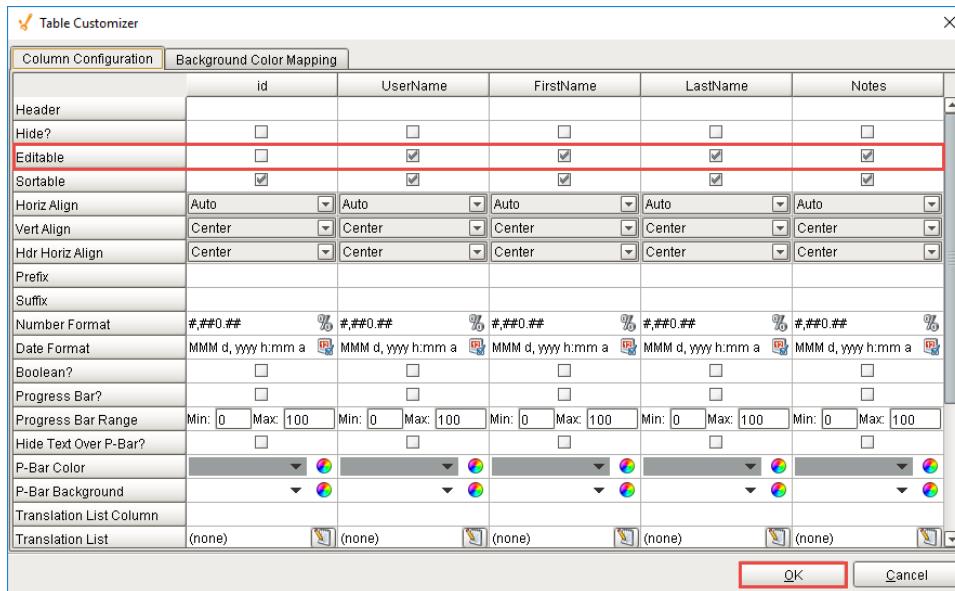
1. Drag a Table component to your window. With the Table component selected, bind the **Data** property of the Table using a SQL Query Binding Type and the query below.

```
SQL - Select All Records
SELECT * FROM users
```

id	UserName	FirstName	LastName	Notes
1	JS	John	Smith	Likes bikes
2	LJ	Luke	Johnson	Lives in town
3	PB	Peter	Burke	Enjoys cooking

2. Right-click on the **Table** and select **Customizers > Table Customizer**.

3. Select the **Editable** checkbox for each of the columns that we will want to edit. Note, the **id** column should not be editable. Click the **OK** button.



4. Right-click the **Table** and select **Scripting**.

5. Create a script in the **cell > cellEdited** event handler using the script provided below. In this script, we have variables that contain the row number that was edited, the column number that was edited, and the new value. Because we do not have the name of the column that was edited, we must first grab the list of columns using the **getColumnHeaders()** system function. We can then use the column number that was edited to find the name of the column that changed and use it in the query.

Python - Writes to a Cell in the Database When Component is Updated

```
# Get the id of the row we edited.
id = event.source.data.getValueAt(event.row, 'id')

# Get the header names of the dataset.
headers = system.dataset.getColumnHeaders(event.source.data)

# Build our Update query. The column name is substituted in from the column that was edited.
```

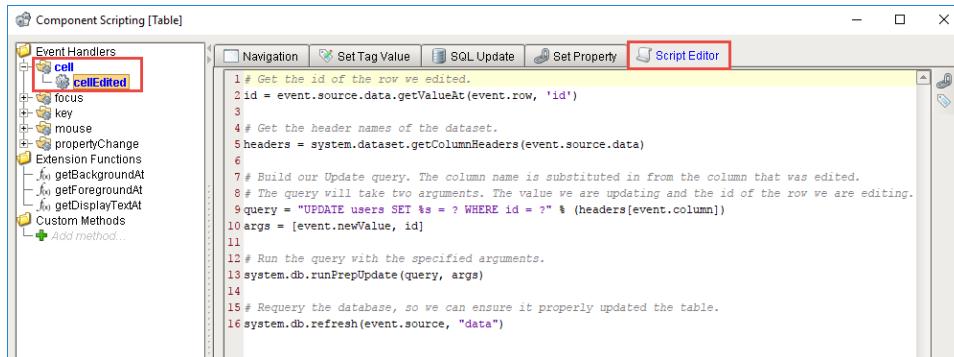
```

# The query will take two arguments. The value we are updating and the id of the row we are editing.
query = "UPDATE users SET %s = ? WHERE id = ?" % (headers[event.column])
args = [event.newValue, id]

# Run the query with the specified arguments.
system.db.runPrepUpdate(query, args)

# Requery the database, so we can ensure it properly updated the table.
system.db.refresh(event.source, "data")

```



- Now, you're ready to test it out! Put the Designer in **Preview Mode**, select the cell you want to edit, hit enter to commit the change, or tab to the next cell to make additional edits.

id	UserName	FirstName	LastName	Notes
1	JS	John	Smith	Likes bikes
2	LJ	Luke	Johnson	Lives in town
3	EF	Ellen	Frank	Loves Ignition

Related Topics ...

- [Simple Database Editor](#)

Refreshing a SQL Query

Refreshing a Query in Vision

The SQL query that populates a property on a component will refresh its data periodically if the [Polling Mode](#) is set to either Relative or Absolute. However, there are times when you want the data to query the database once when the window is opened, and then retain explicit control over each subsequent refresh. The Table data can be requeried using `system.db.refresh`. There are two main ways of doing this: either placing the script on a button, or at the end of some code.

Refresh After Edits

If edits are being made to the Table data, or additional rows are being added, it can be a good idea to call the `system.db.refresh()` function after manipulating data. This way, the data in the Table will automatically refresh with the newly entered data, saving the user the hassle of clicking a refresh button.

Pseudocode - Refreshing a Table after an Edit

```
doWork()

system.db.runPrepUpdate("INSERT INTO table")

system.db.refresh(component, "propertyName")
```

On this page ...

- [Refreshing a Query in Vision](#)
 - [Refresh After Edits](#)
 - [Refresh on a Button](#)
- [Refreshing a Query in Perspective](#)



INDUCTIVE
UNIVERSITY

Refreshing SQL Data on Component

[Watch the Video](#)

Refresh on a Button

We can use a Button component with a script on it to refresh the component with the query binding on it. Typically the button is placed close by the query bound component, and will say something like "Refresh Data". The script is actually fairly simple, and can be placed on the `actionPerformed` event of the Button.

1. Drag a **Button** component onto a window that has a Power Table or Table component querying a Database table.
 - a. Change the Button's Text property to say "Refresh."
2. On the Power Table component, open the **data** property binding and set the polling mode to **off**.
3. Right click on the Button and select **Scripting**.
 - a. Select the **actionPerformed** Event Handler and navigate to the Script Editor tab.
 - b. Add in this script, which will force the Table's data property to refresh.

Python - Refreshing the Database when the Button is Pressed

```
# Will force the Power Table's Data property to run the query again.
system.db.refresh(event.source.parent.getComponent('Power Table'), "data")
```

4. Try it out by adding new data to the Database table and then clicking the Refresh Button.

Refreshing a Query in Perspective

The `system.db.refresh()` function does not work in Perspective. However, an additional component method has been added to handle the more general task of refreshing a binding. The method is called `refreshBinding()`, and can be called on any Perspective component:

```
self.refreshBinding("props.data")
```

The function takes a string as a parameter, the path to the property to be refreshed. More info on the `refreshBinding()` function can be found at [Perspective Component Methods](#).

Related Topics ...

- [system.db.refresh](#)
- [SQL Query Binding](#)
- [DB Browse Binding](#)
- [Named Query Bindings](#)

Editing Multi-Selected Rows from a Table

Editing Multiple Table Rows

A common user experience is to select multiple rows of a table and then edit all of those rows from the database at once, such as deleting all of them. When a user selects more than one row in a table, there is a special function called `getSelectedRows()` available on both the Table and Power Table components that returns the row indexes of the table as a Python list. This list can be iterated through in order to delete the selected rows from the database.

Example - Deleting Selected Rows

We can use the list of selected rows to delete them from the database. Start with a table in the Database table that looks like this:

id	machine_name	area_number
1	Conveyor	1
2	Press	2
3	Tank	1
4	Packer	3
5	Loader	3
6	Oven 1	3
7	Oven 2	2
8	Wrapper	1
9	Mixer	3
10	Cold Storage	2
11	Dryer	2

On this page ...

- [Editing Multiple Table Rows](#)
- [Example - Deleting Selected Rows](#)



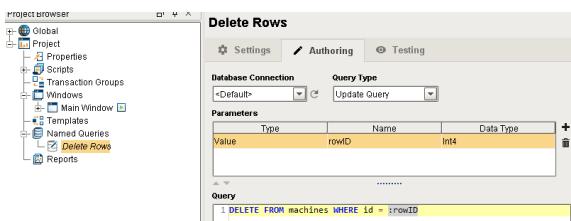
Deleting Multi-Selected Rows from Table

[Watch the Video](#)

1. Create a new Named Query that will be used to delete rows of data. [Set up Security](#) to fit your needs, and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - a. Set the Query Type to **Update Query**.
 - b. Create a single Value type Parameter that is an Int4 data type. This will hold the id of the row, so we can name it `rowID`.
 - c. Create the Delete Query.

SQL - Delete Row with Matching ID

```
DELETE FROM machines WHERE id = :rowID
```



2. Create a second Named Query as in step 1, but this one will be used to select the data into a table.
 - a. Set the Query Type to Query, no parameters, and add a basic select query.

SQL - Selecting from a Table

```
SELECT * FROM machine
```

3. On a new Main Window, add a Power Table component and a Button component.
4. On the Power Table's **Data** property, set up a Named Query binding to the Select Query that was made in step 2.
 - a. Ensure that the Selection Mode property is set to **Multiple Interval**.
5. On the Button, change the Text property to say "Delete".
 - a. Right click on the Button and select Scripting.
 - b. Select the actionPerformed Event Handler and navigate to the Script Editor Tab.
 - c. Here we need to call the **getSelectedRows()** function on the Power Table to determine what rows are selected, and then loop through those to grab the value of the id column in each row and delete the row based on that id.

Python - Looping Through the Selected Rows and Deleting Them

```
# Get the data from the table and assign it to the variable
# called data.
data = event.source.parent.getComponent('Power Table').data

# Get the rows of the data that the user has currently
# highlighted.
rows = event.source.parent.getComponent('Power Table').
getSelectedRows()

# Iterate through each row of the list that is associated
# with the rows variable.
for row in rows:

    # Get the value associated with the current row and
    # the column called "id".
    id = data.getValueAt(row, "id")

    # Run the query to delete from the database, where
    # tableName is the name of the database table.
    system.db.runNamedQuery("Delete Machine Rows",
    {"rowID":id})

    # Refresh the table data.
    system.db.refresh(event.source.parent.getComponent('Power
    Table'), "data")
```

6. You can test it out by putting the Designer into Preview Mode, selecting a few rows, and then clicking the Delete Button.

Related Topics ...

- [Updating the Database through the Power Table](#)
- [Refreshing a SQL Query](#)

Storing Files in a Database

Storing and Displaying Files in a Database

Ignition can store different types of files into a database by storing the raw file bytes into a special database column. Ignition can also pull these file bytes out and display certain files within a Client.



Each database has different column types that are used to store files, so it is important to check with your database documentation to see which data type the column would need to be set to for it to accept file bytes. For example, in MySQL, the data type that accepts PDF bytes is a **LongBlob** data type, so you will need to set the PDF Data column to the **LongBlob** data type. MS SQL accepts the **Varbinary** data type, so you'll need to set the PDF Data column to a **Varbinary** data type.

On this page ...

- [Storing and Displaying Files in a Database](#)
- [Example - PDF File](#)
 - [Uploading PDF Files to the Database](#)
 - [Displaying PDF Files from the Database](#)

Example - PDF File

One of the most common file types that is stored is PDF files. This allows you to store each PDF file within the central database where each client will have access to it, instead of placing the file in a shared drive that all Client computers have access to.



INDUCTIVE
UNIVERSITY

Storing Files in a Database

[Watch the Video](#)

Uploading PDF Files to the Database

With a simple script on a Button component, we can store a PDF file into the database so that any user can view it later. This part does not use a Named Query, because the Named Query Parameters do not have a data type that allows us to pass in the raw file bytes. We can instead use [system.db.runPrepUpdate](#) to call a query from the script.

This example requires that you have a table with a byte array column in it. For example: MySQL uses the BLOB data type and MSSQL uses the varbinary() data type.



This function will not work using default [Client Permission](#) settings. The Legacy Database Access will need to be enabled for this to work.

1. Add a Button component to a new Main Window.
2. Change the Text property to say Add File.
3. Right click on the Button and select Scripting.
 - a. Navigate to the Script Editor Tab of the actionPerformed Event Handler. Here we can put a script that will grab the file bytes using the file path and the [system.file.readFileAsBytes](#) function, and then insert that into the database, along with a user selected file name.

Python - Uploads PDF Files to a Database Using a Button

```
# Find the path to the PDF file.
path = system.file.openFile("pdf")

# Check to ensure that the user actually selected a filepath.
if path != None:

    # Read the file as bytes.
    data = system.file.readFileAsBytes(path)

    # Ask the user to enter in a filename.
    # Will grab just the filename and extension from the path as a suggestion.
```

```

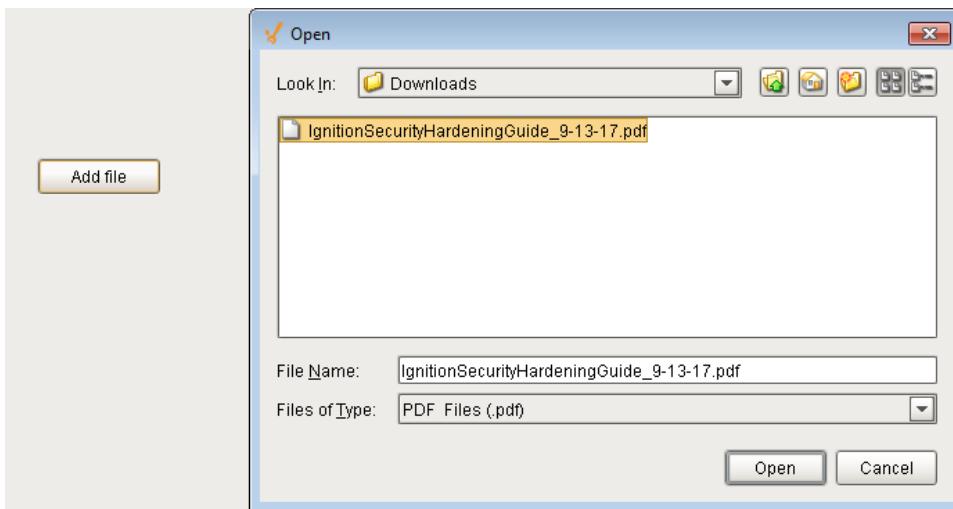
name = system.gui.inputBox("Enter a name for the file", path.split('\\')[-1])

# Check to ensure that the user entered a name for the file.
if name != None:

    # Insert the data and name into the database.
    system.db.runPrepUpdate("INSERT INTO files (fileName, fileBytes) VALUES
(?,?)", [name,data])

```

4. Test out the script by putting the Designer into Preview Mode and clicking the Button.
 a. First, select the file to load.



- b. Then enter in a file name.



Displaying PDF Files from the Database

Ignition can render a PDF document inside the **Vision - PDF Viewer** component, which is a part of the **Reporting Module**. To view PDF files in the Client, your Ignition server must have the Reporting Module installed. Once the module is installed, you can load the bytes from the database into the PDF Viewer component.

1. Create a new Named Query that will be used to select the file names and ids for a Dropdown selection. Set up security to fit your needs, and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - a. With no Parameters, add a query to select all the files in our files table.

SQL - Selecting all Files
<pre>SELECT id, fileName FROM files</pre>

2. Create a second Named Query, same as in step 1. This will be used to grab the file name and bytes after the user has chosen a file.
 - a. Add a single Value type Parameter, "fileID" that will be an Int4 data type.
 - b. Add the query to select the file name and bytes based on the selected ID.

--

SQL - Selecting a File based on an ID

```
SELECT fileName, fileBytes FROM files WHERE id = :fileID
```

3. On the Main Window, add a Dropdown List component and a PDF Viewer component.
4. On the Dropdown List component, add a Named Query binding to the Data property.
 - a. Set the binding to the Named Query that was made in step 1.
 - b. Place a small refresh Button next to the Dropdown that will refresh this query. See [Refreshing a SQL Query](#) for more information for how to refresh on a button.
5. Right click on the Dropdown List component and select Scripting.
 - a. Select propertyChange and navigate to the Script Editor tab.
 - b. This script will take any new selected value and use it in the Named Query we made in step 2 to get the file name and bytes. We can then load the bytes into the PDF Viewer.

Python - Displays PDF Files from a Database Using the PDF Viewer

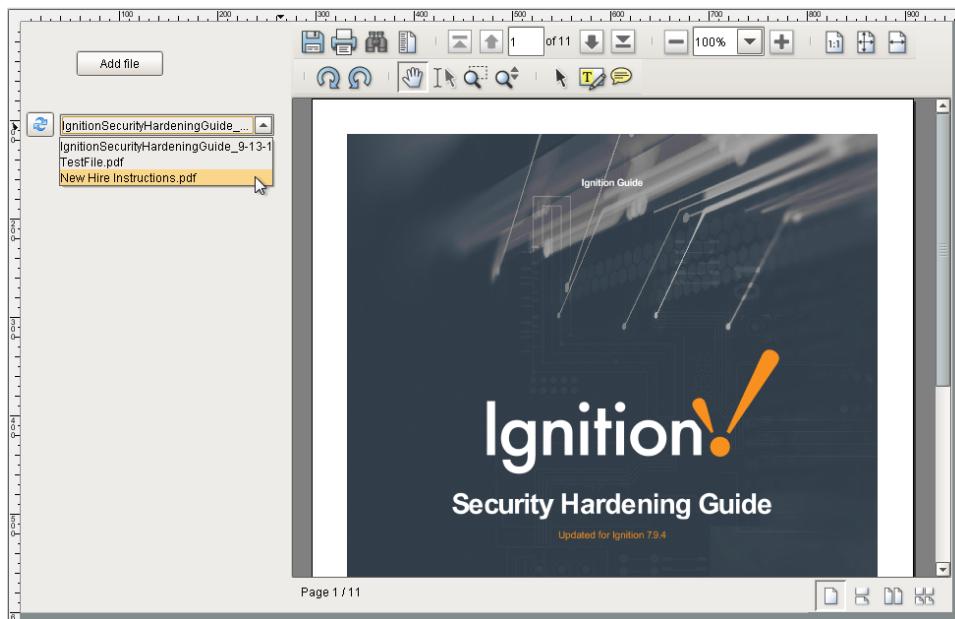
```
# Check to see if the property that changed was the Selected Value property.
if event.propertyName == "selectedRow":

    # Run the query to grab the file name and bytes using the new selected ID value.
    data = system.db.runNamedQuery("Read File", {"fileID":event.newValue})

    # Grab the file bytes and name from the same row.
    bytes = data.getValueAt(0, "fileBytes")
    name = data.getValueAt(0, "fileName")

    # Load the bytes into the PDF Viewer component.
    event.source.parent.getComponent('PDF Viewer').loadPDFBytes(bytes, name)
```

6. Place the Designer into Preview Mode, and try selecting one of the stored files.



Related Topics ...

- [Vision - PDF Viewer](#)

Simple Database Editor

Building a Database Table Editor

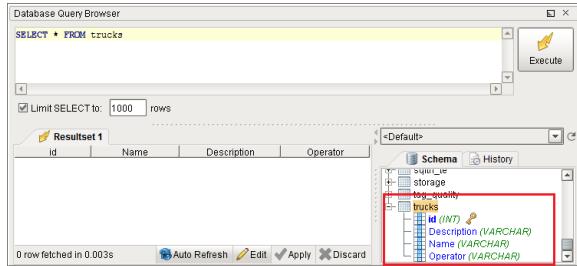
Building a Simple Database Table Editor in Ignition is actually quite easy. Using a few simple components and some scripting, we can easily make a window that will allow users to add, edit, and delete the data from a particular database table. Before we get started, there is a table that has already been created in a database that is used for this example. You can either create this table and follow along, or alter the example to fit your table.

The table is called **trucks**, and it has four columns.

id	Name	Description	Operator
----	------	-------------	----------

- The **id** column is our primary key, it is non null, and it auto increments
- All other columns are varchar(45)

That is it. There is no need to add any data, as we can add data when testing our tool!



Adding the Components

This example consists of a Table component that will display the data, Button components to add and delete rows from the database table, Text Fields where we can enter in new values, and an Update Details Button that will push those values to the database table.

Let's start off by going to the **Designer** and adding the necessary components to the window.

1. Add a **Power Table** component to the window.
 - a. We want to query the table in the database that we made earlier, so let's add a **SQL Query Binding** on the **Data** property of the Table.
 - b. Add the following code under the **SELECT Query** area.

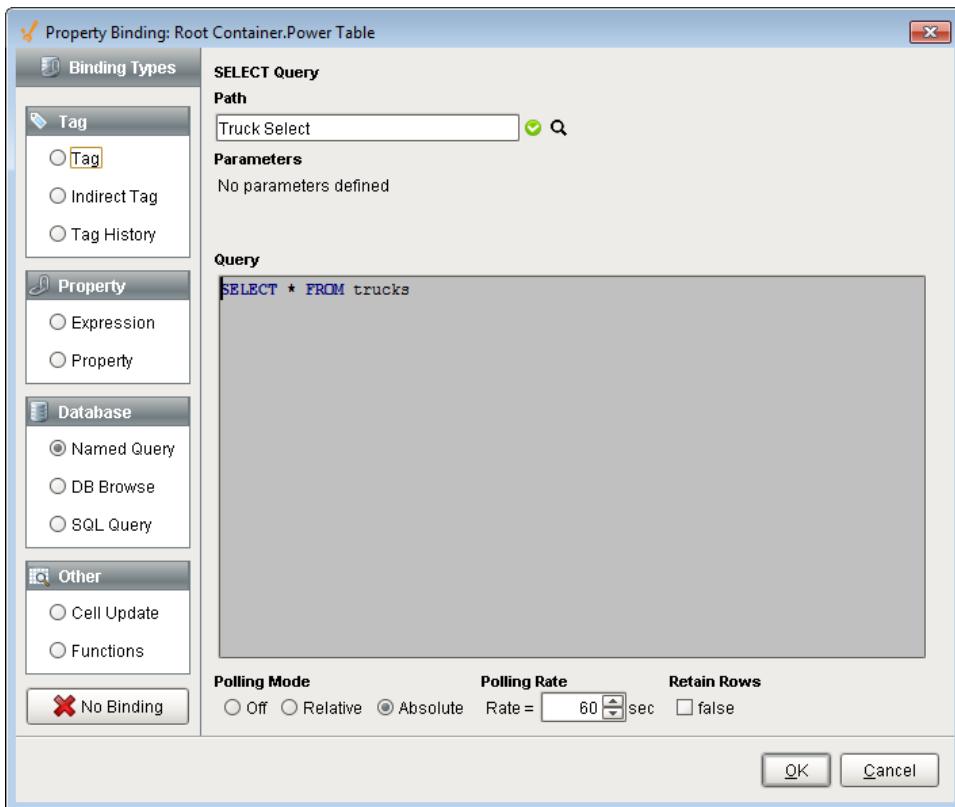
SQL - Selects all Trucks

```
SELECT * FROM trucks
```

- c. Set the **Polling Mode** to **Absolute** and the **Polling Rate** to **60 seconds**, so that it periodically updates if left open.
- d. Since we want to use the secure Named Query system, click the convert to Named Query button under the polling mode, and call the query "Truck Select"

On this page ...

- [Building a Database Table Editor](#)
 - [Adding the Components](#)
 - [Add a Truck](#)
 - [Delete a Truck](#)
 - [Populate Text Fields](#)
 - [Update Data](#)

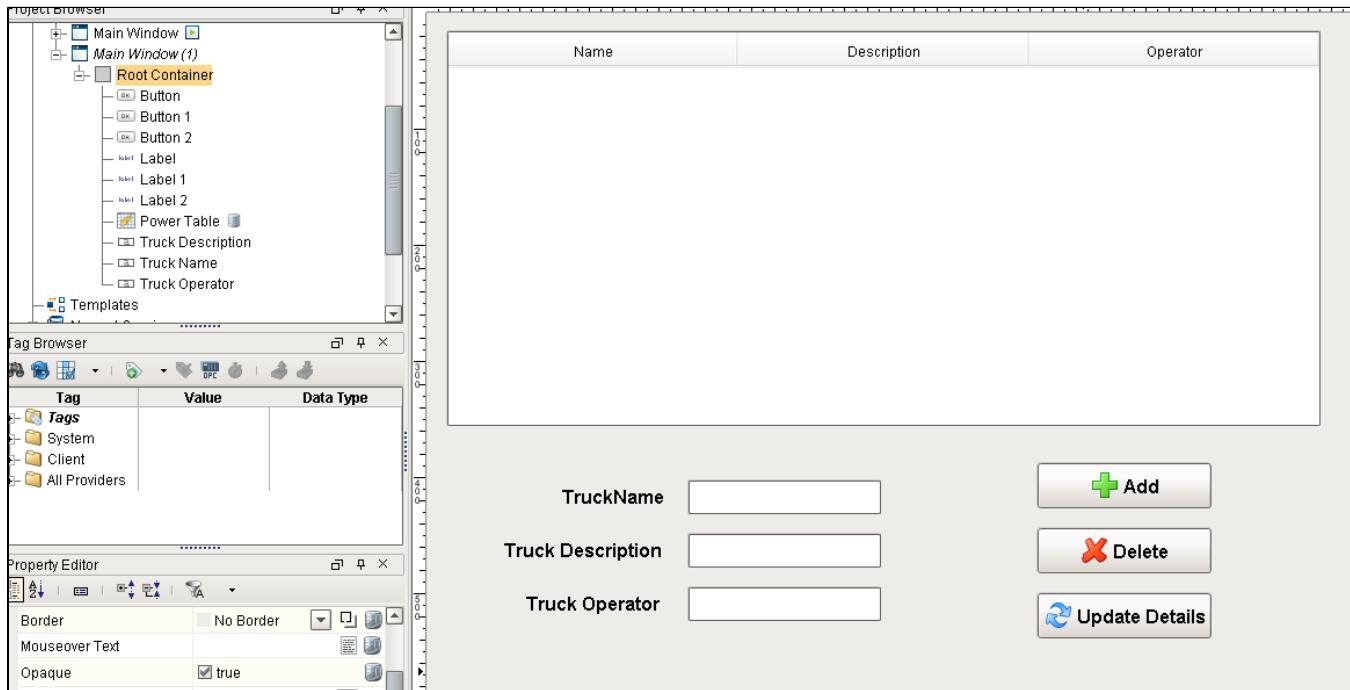


- e. With the Table still selected, open the **Table Customizer** by right clicking on the Power Table and going to **Customizers > Table Customizer**.
- f. Set the **Hide** property of the **id** column to **true**. We don't need to show that column to our users. Click **OK** to close the Table Customizer.
- g. In the **Property Editor**, make sure the **Selection Mode** property of the Power Table is set to **Multiple Interval**.

- 2. Add three **Button** components to the right of the **Table**.
 - a. Change the **Text** property on one Button to say "**Add**".
 - b. Change the **Text** property of the second Button to say "**Delete**".
 - c. Change the **Text** property of the final Button to say "**Update Details**".
 - d. Optionally, you can also add an image to the **Image Path** of each button such as a green plus and a red X, and resize them as necessary.

- 3. Add three **Label** components and three **Text Field** components under the **Table**.
 - a. Pair up a Label with each Text Field.
 - b. Rename the Text Fields:
Truck Name, Truck Description, Truck Operator.
 - c. Change the **Text** property of each Label to match the name of their corresponding Text Fields.

Now we have our window configured with all of the components we need. Next we can start adding scripts that will alter the components.



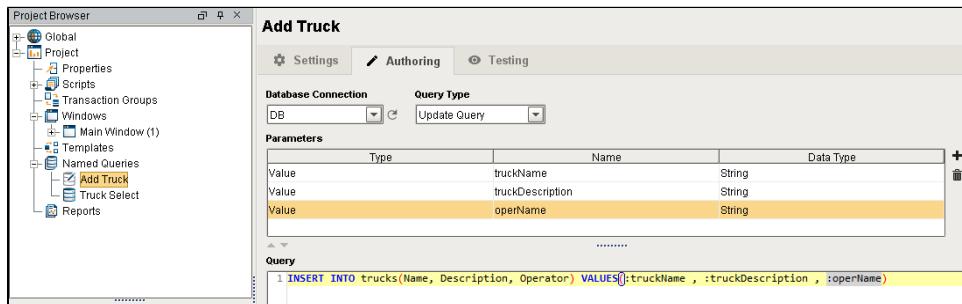
Add a Truck

First, we can add a script to our **Add** button that will allow us to add a row to the Table. We can use the Text Fields that we added so the user can enter in values to insert into the table.

1. Create a new Named Query which we will use to insert new data. Set up **security** to fit your needs and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - a. Set the Query Type to Update Query.
 - b. Create three Value type Parameters.
 - i. The first will be a string data type and will be used for the name of the truck so it can be called "truckName".
 - ii. The second will be a string data type and will be used for the description of the truck so it can be called "truckDescription".
 - iii. The third will also be a string data type and will be used for the operators name so it can be called "operName".
 - c. Create the insert query.

SQL - Adding a Truck

```
INSERT INTO trucks(Name, Description, Operator) VALUES(:truckName , :truckDescription , :operName)
```



2. Back on our window with components, right click on the **Add Button** and select Scripting. Select the **actionPerformed** Event Handler and navigate to the **Script Editor** tab.
 - a. Here we can add some code that will pull the values from the three Text Field components and then use those values in the Named Query we made in step 1.

Python - Adds a New Blank Row and Refreshes the Table

```
1 Python - Adds a New Blank Row and Refreshes the Table
```

```

# Grab the values from the Text Fields.
name = event.source.parent.getComponent('Truck Name').text
description = event.source.parent.getComponent('Truck Description').text
operator = event.source.parent.getComponent('Truck Operator').text

# Use those values in the Add Truck Named Query.
system.db.runNamedQuery("Add Truck", {"truckName":name, "truckDescription":description,
"operator":operator})

# Refresh the table to immediately bring in the new row.
system.db.refresh(event.source.parent.getComponent("Power Table"), "data")

```

3. Now we can test out our script and query by putting the Designer into **Preview Mode**, adding some data to the **Text Fields**, and clicking the **Add Button**. You should see a row populate the table. Try adding a few rows to fill our table with data. Make sure you have the designer Communication Mode set to Read/Write to test. Go back to **Design Mode** before moving on.

Name	Description	Operator
Green Monster	It is so green!	Sarah
Little Truck	The smallest truck.	Joe
Big Truck	This truck is big!	James
Fire Truck	Puts out fires. Does not transport fire.	Dale
Ice Cream Truck	Sells ice cream and various other treats.	Amy
Toy Truck	Transports toys, not a truck that is a toy.	Jane
Blue Truck	A truck that is painted blue.	Anne

TruckName + Add

Truck Description ✖ Delete

Truck Operator ⟳ Update Details

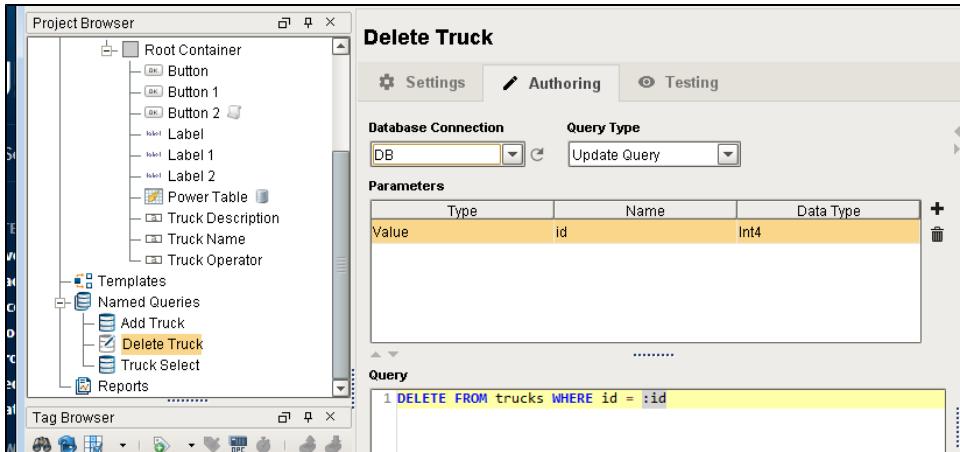
Delete a Truck

Now that we can add rows, let's add a script to our **Delete button** that will delete rows. Since our users can select multiple rows, our script needs to take that into account, and delete all of the rows that the user selected. It will also have a **Message Box** popup if there are no rows selected, informing the user that they need to select at least one row. We also want to add a confirmation before the rows actually get deleted.

- Create a new Named Query which we will use to delete a row of data. Set up **security** to fit your needs and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - Set the Query Type to Update Query.
 - Create one Value type Parameters.
 - It will be a Int4 data type and will be used for the id of row so it can be called "id".
 - Create the delete query.

SQL - Deleting a Truck

```
DELETE FROM trucks WHERE id = :id
```



2. Back on our window with components, right click on the **Delete Button** and select Scripting. Navigate to the **Script Editor** tab on the **action Performed** Event Handler.

a. Here we can add some code that will delete the user's selected rows. First we grab the Power Table's selected rows, and check to make sure rows are actually selected. If there are selected rows, then we have a popup that confirms the user wants to delete rows, and it shows the number of rows that will be deleted. Once the user confirms, the script will then loop through the selected rows and delete each one.

```
Python - Deletes the Row or Rows

# Get the selected row or rows.
selRows = event.source.parent.getComponent('Power Table').getSelectedRows()

# Check to see that a row is actually selected.
if len(selRows) > 0:

    # If a row is selected, ask for confirmation before deleting the rows.
    if system.gui.confirm("Are you sure you want to delete " + str(len(selRows)) + " row(s)?", "Are You Sure?", 0):

        # If the user gave confirmation, we need to loop through all of them.
        for row in selRows:

            # For each selected row, we need to get the value of the id column in that row.
            id = event.source.parent.getComponent('Power Table').data.getValueAt(row, "id")

            # Use the id of the row to delete it from the database table.
            system.db.runNamedQuery("Delete Truck", {"id":id})

            # Refresh the table after deleting all selected rows
            # to immediately remove the selected rows from the Power Table.
            system.db.refresh(event.source.parent.getComponent('Power Table'), "data")

        # If the user said no to the delete.
    else:
        system.gui.messageBox("User canceled the delete.", "Delete Canceled")

    # If no row is selected, have a popup message that asks the user to select a row.
else:
    system.gui.messageBox("Please select at least one row.", "Select A Row")
```

3. Now we can test our **Delete button** by putting the Designer into **Preview Mode** and clicking the **Delete button**. When one or more rows are selected and the **Delete button** is pressed, a popup will confirm we want to delete the rows before deleting them. (You can select multiple rows with **Shift-Click** or **Control-Click**). When no rows are selected and the **Delete button** is pressed, a **Message Box** will pop up, informing the user to select a row first. You can deselect rows in the Table by clicking on the Table and pressing the **Escape** key. Be sure to go back to **Design Mode** before moving on.

Name	Description	Operator
Green Monster	It is so green!	Sarah
Little Truck	The smallest truck.	Joe
Big Truck	This truck is big!	James
Fire Truck	Puts out fires. Does not transport fire.	Dale
Ice Cream Truck	Sells ice cream and various other treats.	Amy
Toy Truck	Transports toys, not a truck that is a toy.	Jane
Blue Truck	A truck that is painted blue.	Anne

Are You Sure?



Are you sure you want to delete 2 row(s)?

Yes **No**

TruckName  Add

Truck Description  Delete

Truck Operator  Update Details

Name	Description	Operator
Green Monster	It is so green!	Sarah
Little Truck	The smallest truck.	Joe
Fire Truck	Puts out fires. Does not transport fire.	Dale
Toy Truck	Transports toys, not a truck that is a toy.	Jane
Blue Truck	A truck that is painted blue.	Anne

Are You Sure?



Are you sure you want to delete 2 row(s)?

Yes **No**

TruckName  Add

Truck Description  Delete

Truck Operator  Update Details

Populate Text Fields

Now that we can add and delete rows from our Table, we need a way to edit the data contained in them. However, before we can create a script on our **Update Details button**, we first need to populate the Text Fields with the current values of the selected row. We can do this by adding a simple expression to each of the text fields.

1. On the Truck Name Text Field, add an expression binding to the Text property.
 - a. The expression will grab the dataset and will grab the value at the selected row of the Name column. The expression is wrapped in a try so that if no rows are selected, it will return an empty string.

Expression - Grabbing the Selected Row's Name Value

```
try({Root Container.Power Table.data}[{Root Container.Power Table.selectedRow}, "Name"], "")
```

2. On the Truck Description Text Field, add an expression binding to the Text property.

- a. The expression will grab the dataset and will grab the value at the selected row of the Description column. The expression is wrapped in a try so that if no rows are selected, it will return an empty string.

Expression - Grabbing the Selected Row's Description Value

```
try({Root Container.Power Table.data}[{Root Container.Power Table.selectedRow},
"Description"] , "")
```

3. On the Truck Operator Text Field, add an expression binding to the Text property.

- a. The expression will grab the dataset and will grab the value at the selected row of the Operator column. The expression is wrapped in a try so that if no rows are selected, it will return an empty string.

Expression - Grabbing the Selected Row's Operator Value

```
try({Root Container.Power Table.data}[{Root Container.Power Table.selectedRow}, "Operator"],
"")
```

4. Click on a row to see the fields below populate. Press the ESC key to empty the selection.

Name	Description	Operator
Green Monster	It is so green!	Sarah
Little Truck	The smallest truck.	Joe
Fire Truck	Puts out fires. Does not transport fire.	Dale
Toy Truck	Transports toys, not a truck that is a toy.	Jane
Blue Truck	A truck that is painted blue.	Anne

TruckName

Truck Description

Truck Operator

+ **Add**

✖ **Delete**

🕒 **Update Details**

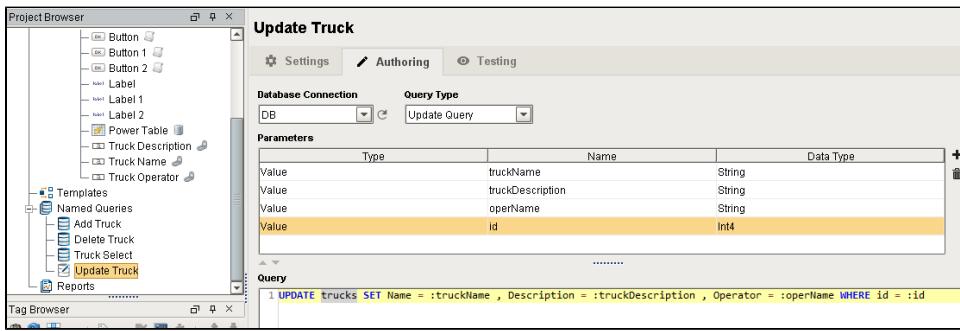
Update Data

The last part of setting up this window is to add a script to the **Update Details** button that will pull the data from the Text Fields and insert it into the selected row of the database Table. Additionally, the script will check to ensure that a row selected, and that only one row is selected. We wouldn't want to confuse users who may have accidentally selected multiple rows at a time, since our button will only update a single row.

- Create a new Named Query which we will use to update a row of data. Set up **security** to fit your needs and name it appropriately. For more information on creating Named Queries, see [Using Named Queries - Example](#).
 - Set the Query Type to Update Query.
 - Create four Value type Parameters.
 - The first will be a string data type and will be used for the truck name of row so it can be called "truckName".
 - The second will be a string data type and will be used for the truck description of row so it can be called "truckDescription".
 - The third will be a string data type and will be used for the operator name of row so it can be called "operName".
 - The fourth will be a Int4 data type and will be used for the id of row so it can be called "id".
 - Create the update query.

SQL - Deleting a Truck

```
UPDATE trucks SET Name = :truckName , Description = :truckDescription , Operator = :operName
WHERE id = :id
```



2. Back on our window with components, right click on the **Update Button** and select Scripting. Navigate to the **Script Editor** tab on the **action Performed** Event Handler.

a. Here we can add some code that will update the users selected row with data pulled from the Text Fields. First we check to ensure that only one row is selected. Next the script asks them to confirm that the user wants to update the selected row. If the user confirms, the script then pulls in the values of all three Text Fields and uses those values along with the selected rows id to update the database table with new data. It will then refresh the table after updating so that the new data can be brought into the table.

```
Python - Deletes the Row or Rows

# Check to make sure only one row is selected.
if len(event.source.parent.getComponent('Power Table').getSelectedRows()) == 1:

    # If a row is selected, ask for confirmation before updating the row.
    if system.gui.confirm("Are you sure you want to update the selected row?", "Are You Sure?", 0):

        # Grab the values from the text fields.
        name = event.source.parent.getComponent('Truck Name').text
        description = event.source.parent.getComponent('Truck Description').text
        operator = event.source.parent.getComponent('Truck Operator').text

        # Grab the selected row.
        selRow = event.source.parent.getComponent('Power Table').selectedRow

        # Using the selected row, we need to get the value of the id column in that row.
        id = event.source.parent.getComponent('Power Table').data.getValueAt(selRow, "id")

        # Run a query that will update the values of the row matching the id.
        system.db.runNamedQuery("Update Truck", {"truckName":name, "truckDescription":description, "openName":operator, "id":id})

        # Refresh the table to immediately show the updated data.
        system.db.refresh(event.source.parent.getComponent('Power Table'), "data")

        # If the user said no to the update.
        else:
            system.gui.messageBox("User canceled the update.", "Update Canceled")

    # If there more than one row selected, or no rows selected.
    else:

        # Have a popup asking the user to select one row.
        system.gui.messageBox("Please select one row to edit.", "Select One Row")
```

3. Now that we added this last bit of code, we can test it by selecting a row, adding data to each **Text Field**, and clicking the **Update Details** button. You should see the data populate the row of the Table, since the table in the database has been updated. Additionally, you should notice how the Text Fields will automatically update when selecting a new row. Try adding a few rows of valid data to your Table.

Name	Description	Operator
Green Monster	It is so green!	Sarah
Little Truck	The smallest truck.	Joe
Fire Truck	Puts out fires. Does not transport fire.	Dale
Toy Truck	Transports toys, not a truck that is a toy.	Jane
Yellow Truck	A truck that is painted yellow.	Abby

TruckName	<input type="text" value="Yellow Truck"/>	 Add
Truck Description	<input type="text" value="A truck that is painted yell"/>	 Delete
Truck Operator	<input type="text" value="Abby"/>	 Update Details

Related Topics ...

- [Inserting Data into a Database](#)
- [Updating the Database through the Power Table](#)
- [Editing Multi-Selected Rows from Table](#)
- [Named Query Caching](#)

Basic SQL Troubleshooting

Learning when something is wrong with your database queries and learning how to resolve the problem is key to getting the best possible use out of the database connection to Ignition. Typically, the biggest problem that users face with queries is the data taking too long to load on the window, but an error in a query can also be difficult to track down. This section details what to look for when the query fails to execute, as well as what to do when facing slow queries and how to optimize them within the project to help return data to the window as quickly as possible.

Error Message Box

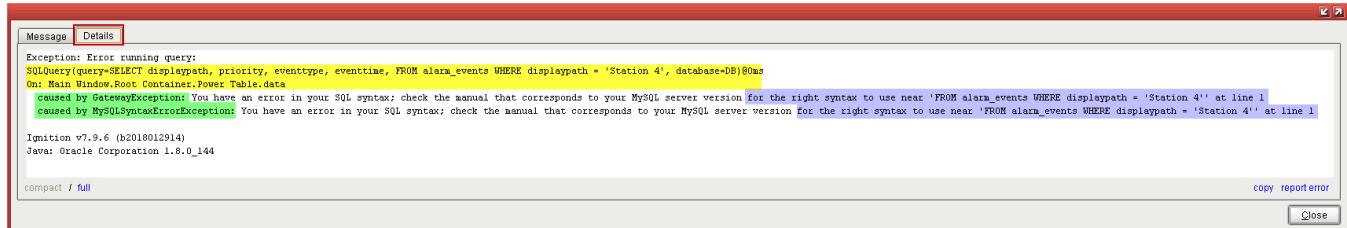
If a query fails to execute, chances are there was an error box that popped up. Resist the urge to close it! The Error Box usually contains a lot of useful info that can help troubleshoot why the query is failing. When looking at the error message, there are a few things to look for. First, you will want to click on the **Details** tab in the upper left corner of the Error Box. This will bring up the details of the error, which contains the information we need to track down the error.

In the image below, you can see the yellow highlighted part shows us what our query is, as well as where it is executing from. This helps to ensure we are looking at the correct error message for the query we are trying to fix. Next, you want to look for the words '**caused by**' in the text of the error details, which tell us the reason for the error. You can see in the image below, there are two of '**caused by**' messages highlighted in green. The first one is from the Gateway, and the second one is from the database (in this case, a MySQL database).

After locating the '**caused by**', the following message will help to pinpoint what the error is. In this example, we have an error with our syntax in our SQL statement. The most helpful portion of the error is at the end of the message highlighted in purple. This lets us know where the syntax error is in the SQL statement. The trick is to look immediately before the quoted query. In the image below, it starts to quote the original query with '**FROM alarm_events WHERE.....**', so in my original query, I need to look at what was right before the '**FROM**'. You can see in my original query highlighted in yellow, there is a comma right before the '**FROM**', which is incorrect, as the last column in the **SELECT** statement should not have a comma after it.

On this page ...

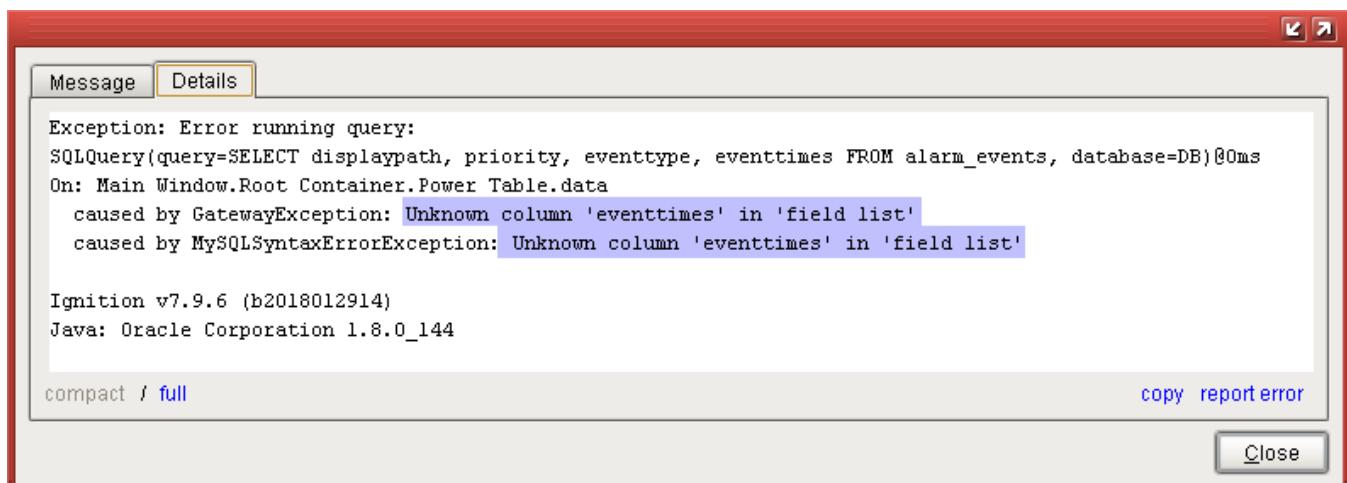
- Error Message Box
 - Other Common SQL Errors
- Checking the Database Connection
- Testing Query Results
- Checking for Slow Queries



Other Common SQL Errors

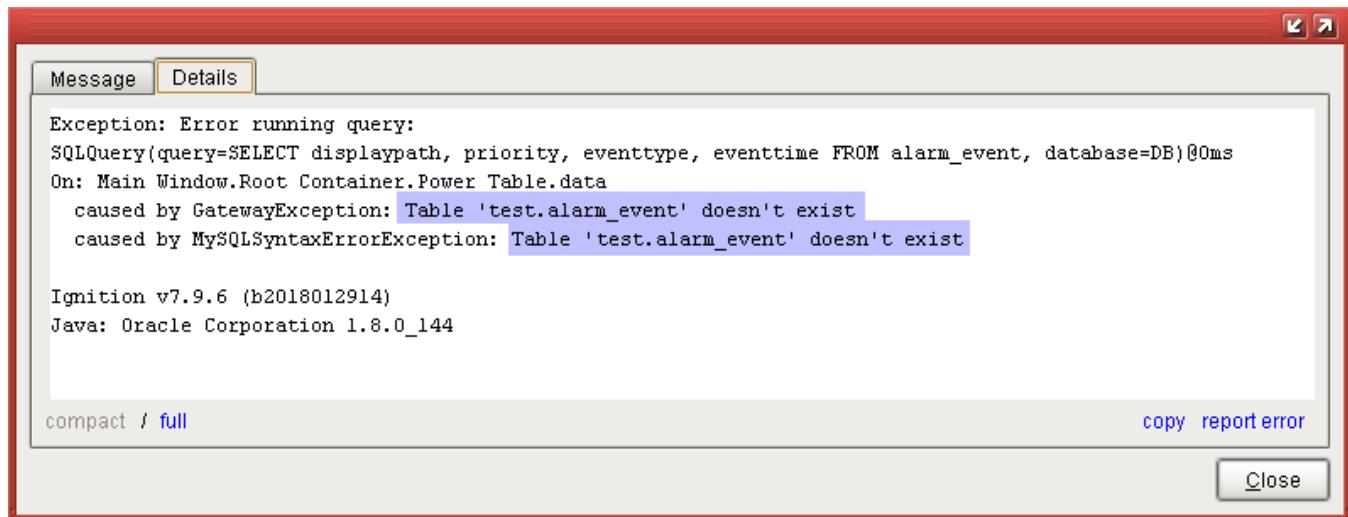
Unknown Column

Typically, the name of the column is wrong because of a misspelling. In this case, the correct spelling was 'eventtime', but there was an accidental 's' added to the name.



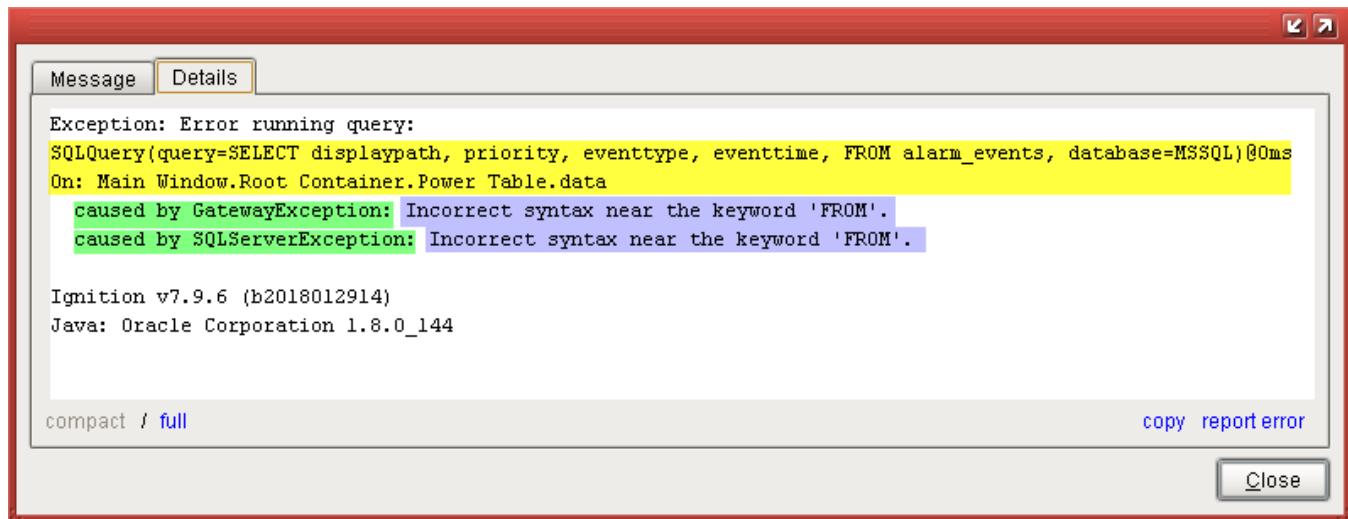
Unknown Table

If the table is not found in the database, it may also be because of a misspelling. Here, the correct table name should be 'alarm_events'.



Any Database Error

The previous error examples have all been from a MySQL database, but the same principles apply to any database. Simply locate the 'caused by', and look at the message afterwards. In this MSSQL syntax error, instead of displaying everything after the syntax error, it only displays the part immediately after. The problem in this query is an extra comma before the word 'FROM'.



Checking the Database Connection

Checking the Database Connection in the Gateway can also be useful to ensure there is a valid connection and there is nothing blocking the execution of queries. The [Status](#) page of the Gateway can be used to determine if queries are taking too long, while the Database Connections page can be used to alter the settings of the database to better handle the number of queries running from the Gateway to the database. See the [Slow Queries](#) page for more information on changing these settings.

Testing Query Results

Using the [Database Query Browser](#) is a good way to test out a query before actually running it on a component. This can help you see what results will be returned so that the query can be modified to narrow down how the query should be formatted.

Checking for Slow Queries

There are several ways that slow running queries can cause improper behavior. See [Slow Queries](#) for more details.

[Related Topics ...](#)

- [SQL in Ignition](#)

[In This Section ...](#)

Slow Queries

Slow running queries can be a big problem. Not only can data take a long time to display on the screen, but it can end up slowing down the whole client. Here, we take a look at some of the things that you can do when your project has a slow running query.

Identify the Slow Query

The first step in dealing with slow queries is identifying which query is actually running slowly. It is typically obvious when a window is opened and a component takes time before database data is displayed. Chances are, the query is on that component somewhere. However, it may also be a good idea to check the queries that are being run against that [database connection in the Gateway](#). From the Gateway webpage, navigate to **Status Connections Databases** to see a list of all database connections. Clicking on the **Details** button to the right of the database connection will show all of the currently running queries, the most recent long running queries, as well as some basic metrics for that connection. Here, you can get a good idea of any queries that may be a little slower than the others.

On this page ...

- [Identify the Slow Query](#)
- [Execute Against the Database Management Software](#)
- [Fast Query in the Database Management Software](#)
 - [Check the Database Connection](#)
 - [Check Currently Executing Queries](#)
- [Slow Query in the Database Management Software](#)

The screenshot shows the Ignition software interface. The top navigation bar includes 'HOME', 'STATUS' (which is highlighted in orange), and 'CONFIGURE'. On the far right, there are user account and sign-out options, and a 'Launch Designer' button. The main content area is titled 'Connections' and 'Databases / Connection [ACME_DB]'. It features a 'Database Stats' section with two gauges: 'Connections' (1/8) and 'Queries / Sec' (0). To the right is a histogram showing the distribution of query execution times. Below this is a table for 'Active Queries' with one row: 'SELECT sleep(60)' started 'a few seconds ago' with a 'Cancel' button. At the bottom is a table for 'Longest Recent Queries' with three rows of data. The left sidebar contains navigation menus for 'SYSTEMS' (Overview, Performance, Alarm Pipelines, etc.), 'CONNECTIONS' (Databases, Designers, Devices, etc.), and 'DIAGNOSTICS' (Execution, Logs, Threads).

Execute Against the Database Management Software

The next step in identifying the type of slow query we are dealing with is to run the query directly within the database management software. By cutting out Ignition, we can determine if the query is actually running slowly, or if there is a problem within Ignition that is making the query run slow.

Fast Query in the Database Management Software

If the query runs quickly in the database management software, then there are a few things we can take a look at that may help out.

Check the Database Connection

First, check to make sure the database connection is valid, and there are no warnings associated with it. You may want to go into the connection settings for that database and take note of the value of the **Max Active** property, which determines the maximum number of active connections to the database as well as the **Max Wait** property, which is the number of milliseconds to wait for a connection to come available. To get to the database connections settings, go to **Configure** page of the Gateway webpage and select **Databases Connections**, locate your database, and hit the **edit** button. Here, you can check the database settings. Open **Advanced Properties** at the bottom of the screen, and you'll find the **Max Active** property as well as the **Max Wait** property.

Check Currently Executing Queries

Back in the **Status** section of the Gateway Webpage on the **Database Connection** page mentioned above, we can see a list of currently running queries, as well as how many of the active connections are being used. If the max number of connections is being used, it may be that there are so many queries running that each query needs to wait for an active connection to open up. If this is the case, you may want to increase the amount of active connections to the database or take steps to [reduce the query load on the database](#).

Slow Query in the Database Management Software

If the query also runs slow within the database management software, then the query is just a slow query. Unfortunately, nothing within Ignition can speed up the execution of that query, so you would want to instead take a look at what the query is doing. If the query is pulling in lots of data, you can try breaking the query down into smaller queries, or writing the query in a more efficient way.

For particularly large tables, it may also be helpful to add an index to one of the table's columns. Indexes are something that the user can't see, but help the database speed up data retrieval. However, adding an index to a table will increase the amount of time that an update to the table takes, because the index also needs to be updated. For this reason, it is recommended to only make indexes on columns that are frequently searched against.

To make an index, most database management software have built in interfaces that allow you to customize the index on each table. An index can be made for a single column, or a combination of columns in the table. Talk to your Database Administrator about adding or updating table indexes.

Alternately, you may need to take a look at the database system as a whole. As the size of the database grows, you may need to update the hardware resources available to it. If the database is installed on a server with another system like Ignition, be aware that although the two systems are now sitting next to each other, they now have to share the hardware resources available to them, which may cause issues for both systems. In many cases, it is often better to have the database run on a separate server, which gives it ample room to grow.

Related Topics ...

- [SQL Query Volume Optimization](#)
- [Connections - Databases](#)

SQL Query Volume Optimization

Overview

Leveraging a SQL Database can drastically improve the quality of a project, but improper database consideration while designing can lead to poor performance later on. This page contains some best practices and considerations when incorporating SQL queries.

Optimizing Individual Queries

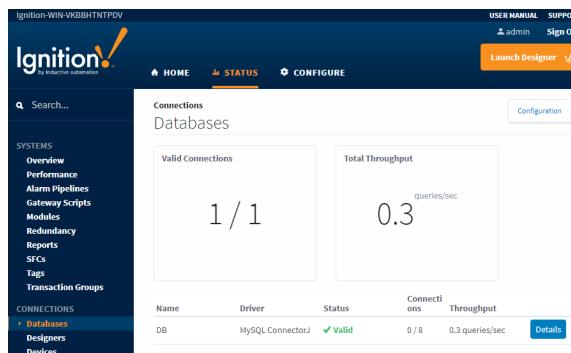
Optimizing individual queries to run faster and be more efficient is very difficult to do properly and can vary widely depending on a number of factors such as the way the tables are set up in the database and what data is being pulled out. Because each database can vary widely from another, there isn't any general way of improving the efficiency of your queries. Instead we recommend becoming more familiar with the SQL language as well as how the data is set up in your system and exactly what data you want to retrieve. This knowledge can help you build better queries.

Your company may also have a database administrator who would manage the database system for your company and would be familiar with figuring out the best way of retrieving data. They may be able to help you retrieve the data that you need.

Always Think Large

Consider how many instances of a query may be running at any given point in time. A single [SQL Query Binding](#) will be called for each instance of the window that is open, so if 50 clients are all looking at the same window, 50 separate queries will be called. If the binding is configured to poll, then 50 queries will poll at the rate specified for this single binding. This is already a fair amount of work without factoring in other systems, such as Tag Historian.

To provide context, you can always check the [Status Section](#) of the Gateway Webpage to check current throughput of each database connection.



Use Cached Named Queries When Possible

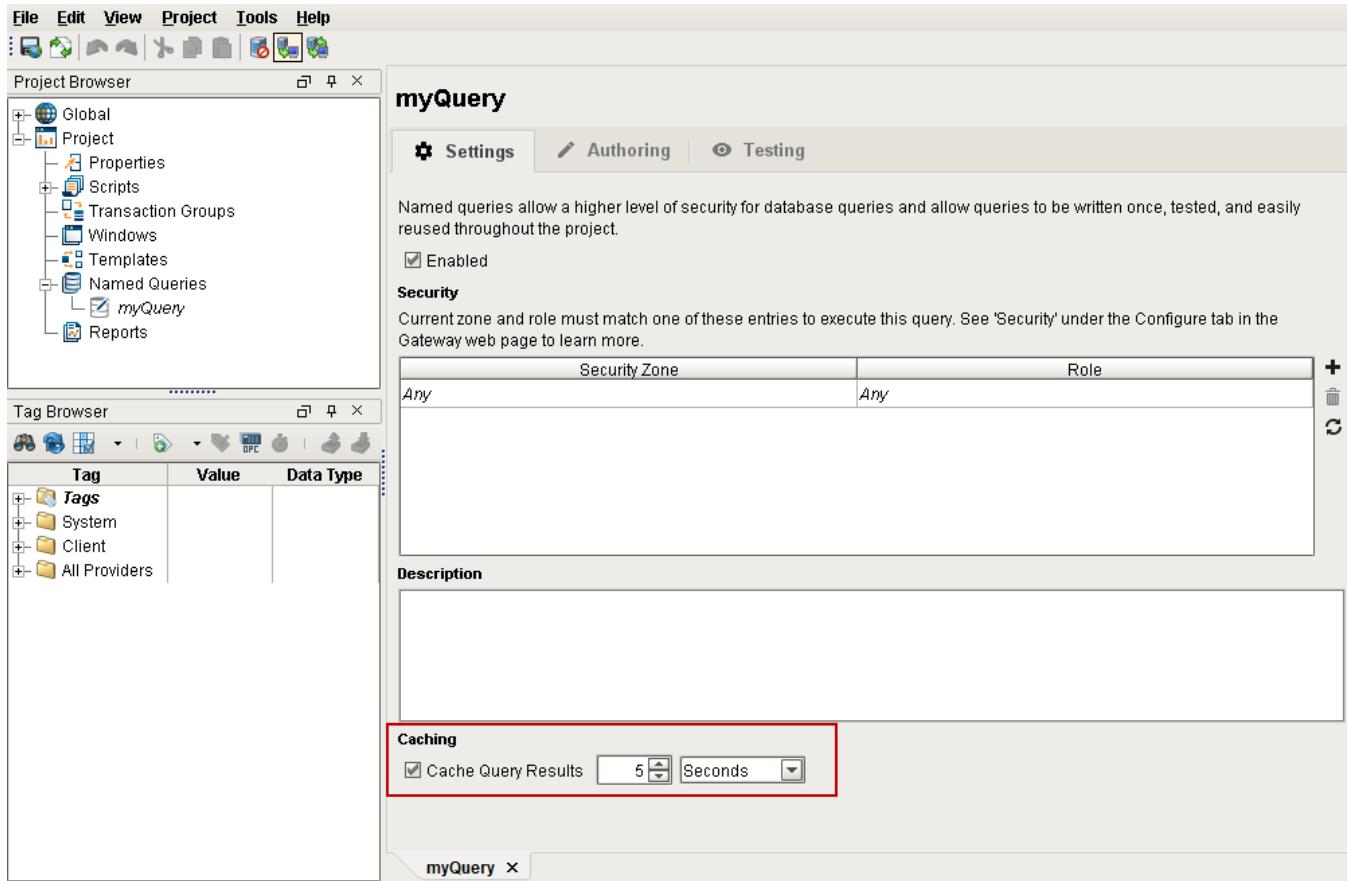
Most resources in Ignition that can request a query will not cache the results for use by other resources: a SQL Query Binding that returns a result set will only do so for the one component, and can't be utilized by other resources in the same project. Thus, if two clients have the same window, the same query must be fired twice for both bindings to receive information from the database which is wasteful.

[Named Queries](#) are the exception. They can [cache the resulting dataset](#) for use by other resources in the project, as well as other instances of the same project. In the previously mentioned scenario, one client would trigger the Named Query to execute, and the other client would simply utilize the cached result set, reducing the number of queries running against the database.

If your project contains queries that poll slowly, or results sets that aren't frequently modified, then a Named Query with caching enabled is an efficient alternative to a SQL Query or DB Browse binding.

On this page ...

- Overview
- Optimizing Individual Queries
- Always Think Large
- Use Cached Named Queries When Possible
- Use the Expression Language to Consolidate Multiple Queries
- Restrict the Number of Query Tags in a UDT
 - Single Database Tables
 - Multiple Database Tables



Use the Expression Language to Consolidate Multiple Queries

If multiple resources (such as multiple Tags, or multiple components) need separate values from the same database table, or a window contains multiple components that are all querying data from the same table, such as multiple Numeric Labels, it may be more efficient to have a single query run and fetch the large portions of the table, and then distribute the individual values to each component. This typically involves having some property or Tag query for several rows of data from a database table, and then using expression bindings or Expression Tags to retrieve individual values from the query.

Individual values may be retrieved from a dataset via the Expression Language: either an Expression binding on a property, or an Expression Tag. Here are two commonly used approaches to extracting a single value from a dataset in the Expression Language:

- The [Expression Language's Dataset Access](#) syntax. You may want to use the `try()` function in case the dataset is empty:

Pseudocode - Expression Language Dataset Access Syntax

```
{dataset}[rowIndex,columnIndex]
```

- The `lookup` function:

Pseudocode - Expression Language Lookup Function

```
lookup({dataset}, lookupValue, noMatchValue)
```

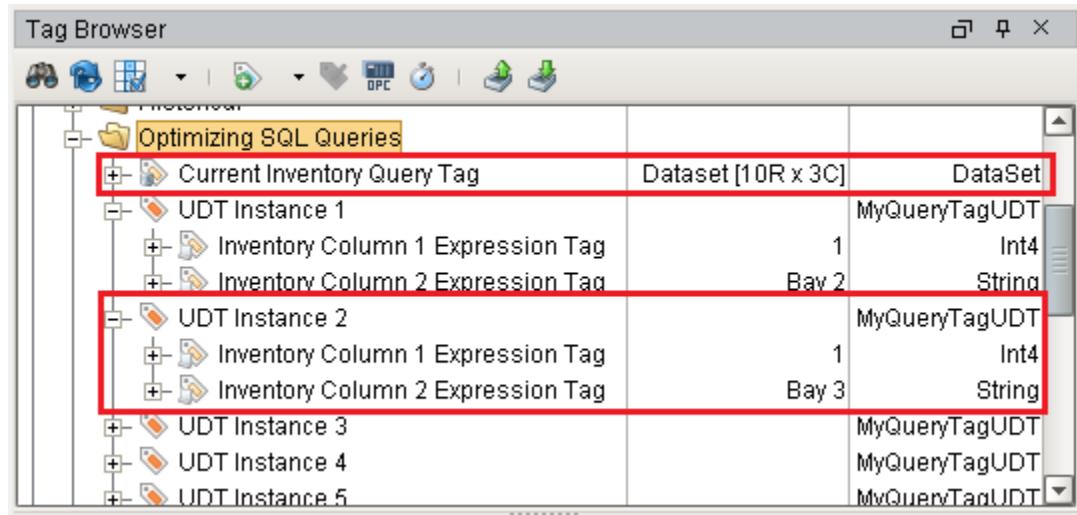
Restrict the Number of Query Tags in a UDT

Each Query Tag in a UDT will run a separate query per instance of the UDT. Assuming a scan class of one second, if a UDT definition contains 5 Query Tags, and there are 5 instances of that UDT, then there will be 25 queries executing every second.

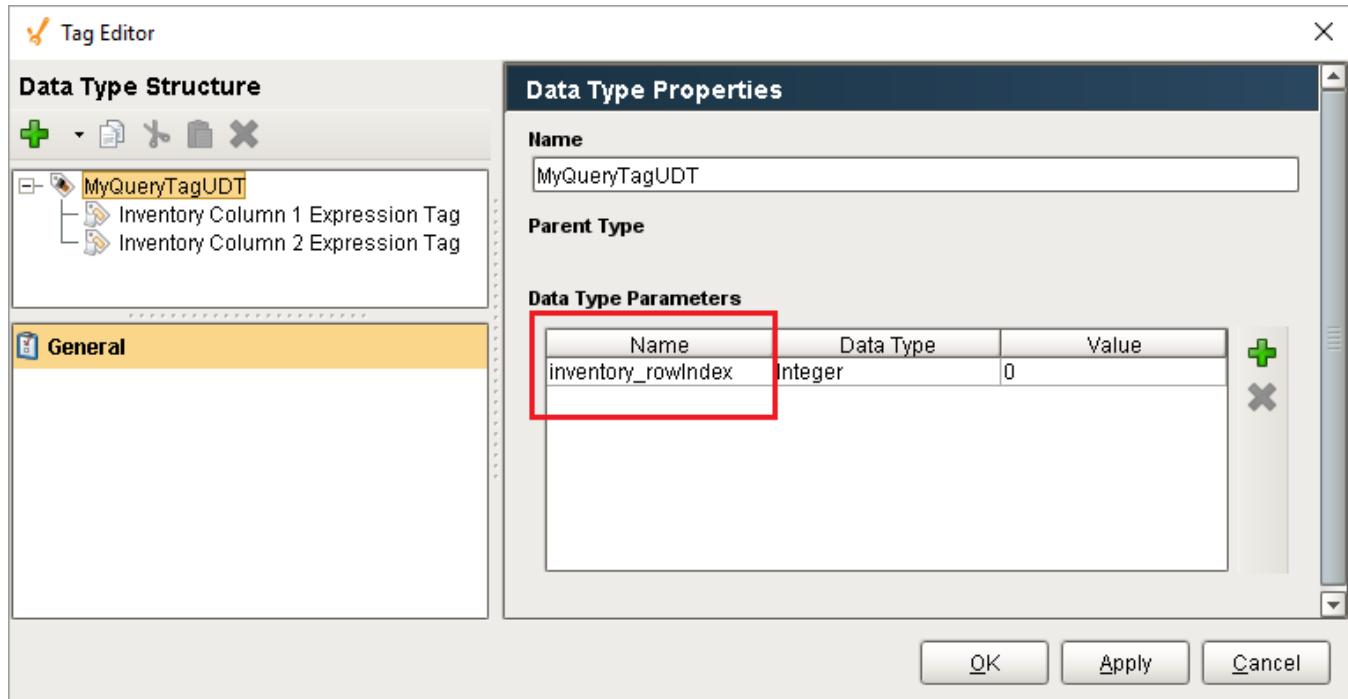
As mentioned on this page, the Expression Language can be used to reduce load on the database if multiple Query Tags are retrieving data from the same database table. Furthermore, UDT parameters can be utilized in the Expression Tags, so new UDT Instances can easily be configured to look up the appropriate values.

Single Database Tables

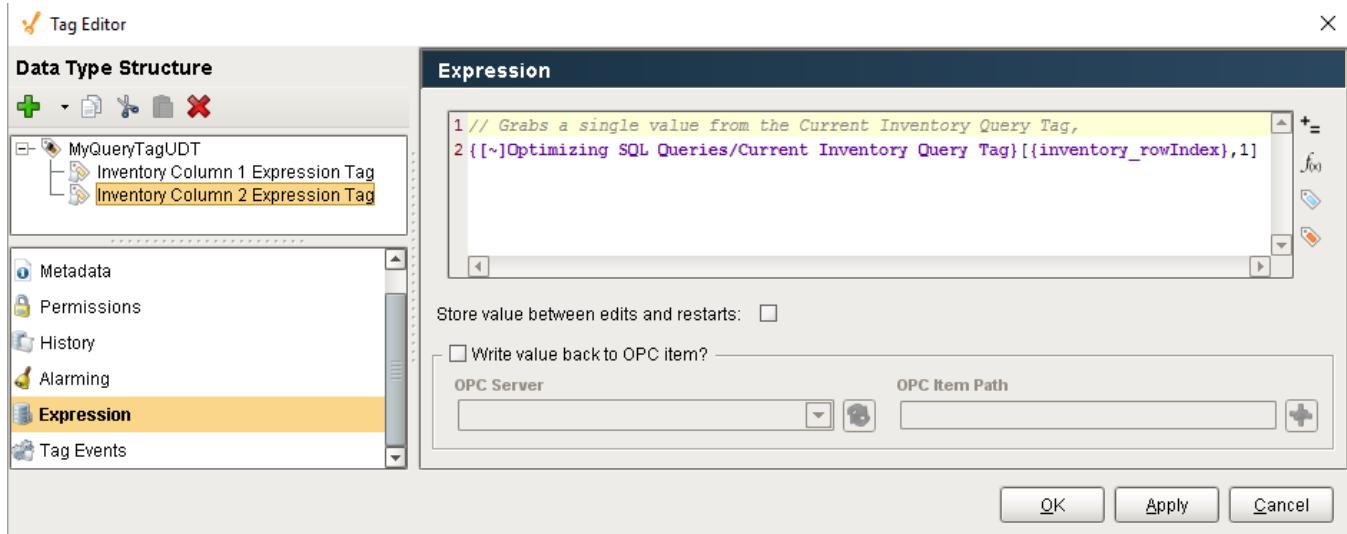
Below, the Tag named **Current Inventory Query Tag**, as the name implies, is a Query Tag retrieving multiple rows of data from a database table. We see that the highlighted **UDT Instance 2** contains two members: **Inventory Column 1 Expression Tag** and **Inventory Column 2 Expression Tag**, which are simply Expression Tags that are referencing individual cells from the **Current Inventory Query Tag**.



The UDT definition can use a parameter to specify an individual row in the Query Tag that each instance should focus on.



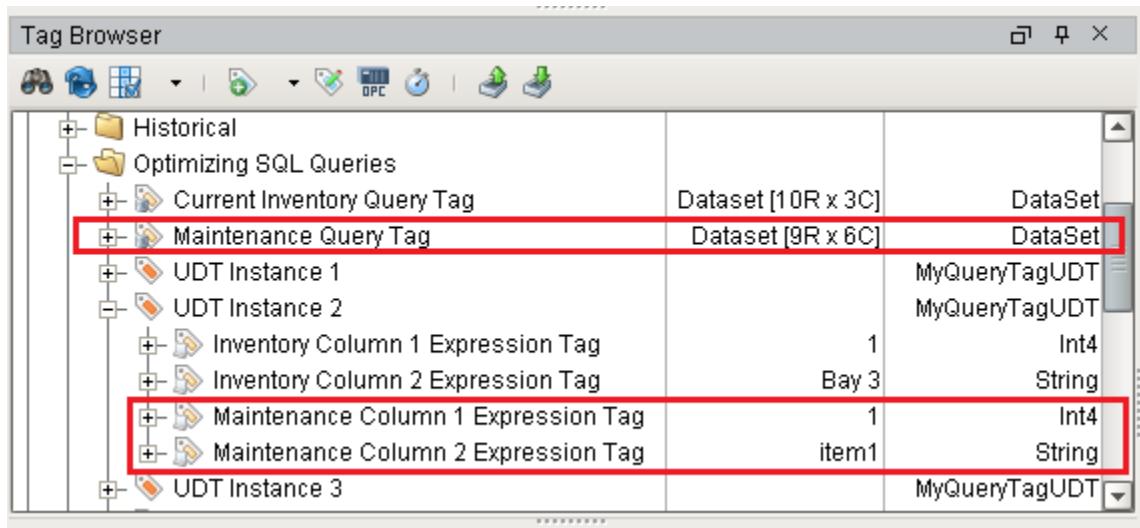
Each Expression Tag could use an expression like the following to look up individual values:



Multiple Database Tables

To add values from a separate database table, we simply need a separate Query Tag. In the image below, a new Tag named **Maintenance Query Tag** has been added, which is querying from a separate Database table. To incorporate this new data into our UDT instances, new Expression Tags have been added (**Maintenance Column 1 Expression Tag**, and **Maintenance Column 2 Expression Tag**) that simply reference specific values in the new Query Tag. Now, regardless of how many UDT instances exist in the Tag provider, we only have two Tags that are executing queries against the database.

For each separate table, we need to incorporate a single new Query Tag to collect all of the rows we want to show, add index parameters to the UDT definition, and add Expression Tags to our UDTs.



[Related Topics ...](#)

- [Named Query Caching](#)
- [Slow Queries](#)

Scripting

What Is Scripting?

Most of the time when we talk about "scripting" in Ignition we are talking about Python scripting, or writing code in the Python language. Python is a general purpose programming language that was developed in the early 1990s and has gained significant popularity in the 2000s. It is extremely readable, elegant, powerful, and easy to learn. As an added bonus, it gracefully interacts with Java, giving programmers an extremely powerful tool when paired with Ignition, which is written in Java.

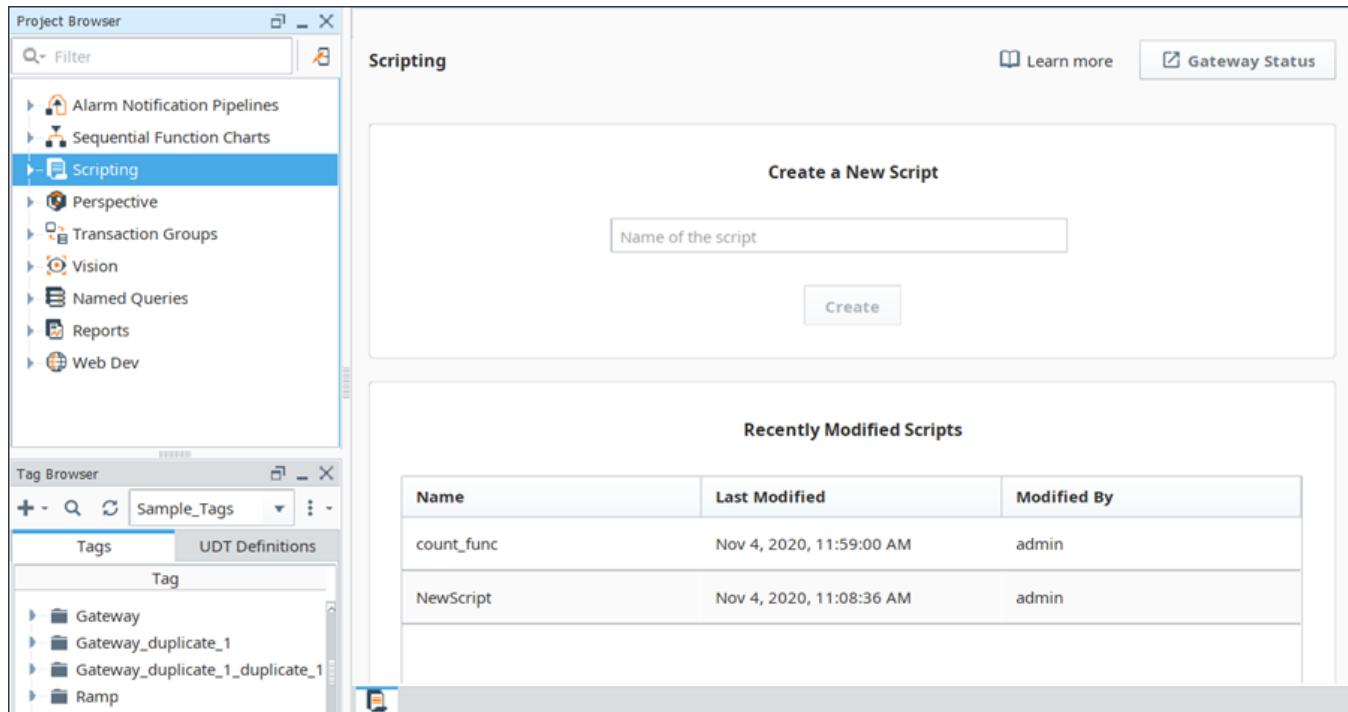
The Scripting Welcome tab will appear when you have either the Scripting window or one of its children open. It allows you to create a script, and once you click Create, it immediately opens the window so you can start writing your script. It's a quick and efficient way to get right to your scripting tasks. At a glance, the Scripting Welcome tab will also show you any recently modified scripts along with the date it was modified and who modified it. You can even double click on a recently modified script and open it.

The Scripting Welcome tab provides a quick way to create a new script and update existing ones. The **Welcome Tab** should appear when the user has either Scripting, or one of its children items selected.

In Ignition, you will be mixing the core Python language with references to other components and a variety of our built-in [system functions](#).

On this page ...

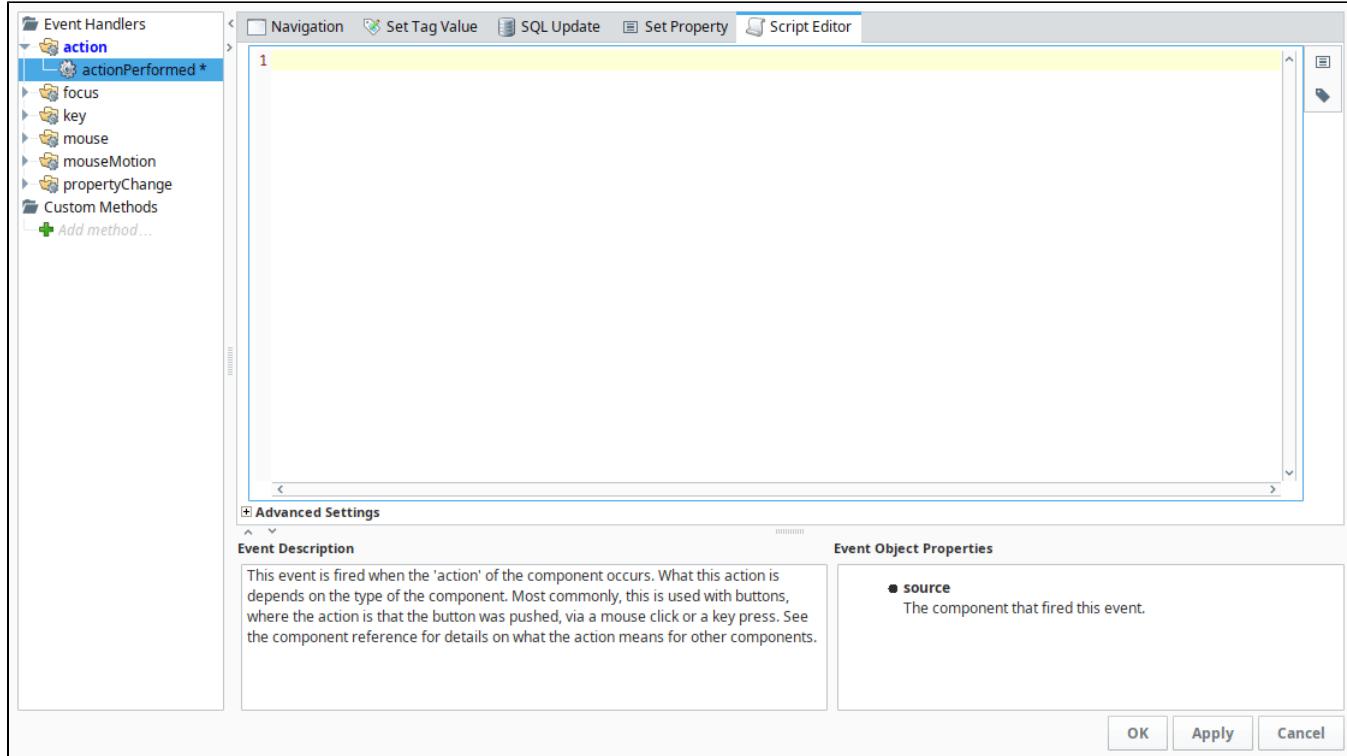
- [What Is Scripting?](#)
- [Is it Easy to Use?](#)
- [Where Is Scripting Used?](#)



Is it Easy to Use?

Luckily, Python is a simple language to get started with. Using it in an event-driven system takes away a lot of the extra code that normally makes programming time consuming. For those that are already familiar with scripting (and those of you that are learning), we have a huge list of functions inside Ignition to do some of the common tasks in a single line of code. These [System Functions](#) are available while typing. Just start with "**system**"

and press "ctrl-space" to see a list of available functions. The list will filter itself as you continue typing.



Where Is Scripting Used?

Python is used in many places in Ignition. Each location has its own events that trigger your scripts to run, and add functionality to your projects in different ways.

- **Components** - Add actions to components like buttons, customize the look and feel of charts and tables, and even set up a custom navigation schema.
- **Tags** - Create a script that runs on a Tag change and when an alarm goes active!
- **Reports** - Use scripting to create a customized datasource, or create your own unique action to use with the scheduling system.
- **Alarm Notification** - Create custom rosters using scripting to dynamically change who gets notified with each new alarm event.
- **Client, Session, and Gateway** - Add a script that will run when certain events happen, such as when the Client or Gateway starts up, or on certain keystrokes.

Users that are new to Ignition focus mainly on the component binding system in [Vision](#) or [Perspective](#), and for good reason. It's simple, flexible, and generally easy to understand without much of a computer background. However, Ignition has a complete scripting system built into every place you can think of. Using it is not a requirement, but it can add a significant degree of flexibility and customization to your projects. It allows you to create exactly what you need, giving you total control where pre-canned options fall short.

The majority of your scripting will be done in Event Handlers inside of components. This system makes it very easy to get started scripting with little to no experience. With the script builders, basic scripting like Navigation and setting Tag values takes just a few clicks. See also [Script Builders in Vision](#) and [Component Events and Actions](#).

Related Topics ...

- [Getting Started with Scripting in Ignition](#)
- [Gateway Event Scripts](#)
- [Client Event Scripts](#)
- [Tag Event Scripts](#)
- [Scripting in the Report Module](#)

In This Section ...

Python Scripting

About Python

While it is entirely possible to create a complete and powerful project in Ignition without writing a line of script, many designers will find that in order to complete projects with specific requirements, they need to learn at least a little Python. In our experience, most industrial projects involve lots of very complex and specific requirements.

This section is a short tutorial specifically for Python, which should help get you started. It goes over all of the core concepts you will need for scripting in Ignition but then next section ([Scripting in Ignition](#)) goes over using Python directly inside Ignition.

- [Python Variables, Data Types, and Operators](#): Learn what a variable is and how to create it, what the various data types are, and what operators you can use on them.
- [Conditions and Loops](#): Learn the common if/else type of statements, as well as loops.
- Functions: Learn about the [built-in functions](#) that can do complex work with a simple command, as well as [user defined functions](#), which help compartmentalize the code.
- [Libraries](#): Learn about our built-in system functions, as well as pulling in outside libraries.

Python or Jython?

You'll often hear Python referred to as "Jython" by advanced users of Ignition. Python is the language, Jython is the implementation of the language that we use. Most users of Python use the implementation called "CPython" - they just don't realize it. See [http://en.wikipedia.org/wiki/Python_\(programming_language\)#Implementations](http://en.wikipedia.org/wiki/Python_(programming_language)#Implementations).

One of the powerful things about using Jython is that your script has access to the entire Java standard library. For more information, see [Accessing Java](#).

Many scripting users are blown away by their script's speed. We can't take credit for this - the Jython engine compiles the code when it is run. Your Jython code is converted to Java bytecode, which means it runs natively in the JVM, which in turn can compile it to machine code. It's fast.

Which Version of Python Is Used?

Ignition uses Jython 2.7. Jython is the Python programming language implemented over the Java Virtual Machine. When looking at outside documentation, such as on www.python.org, verify that you are looking at the correct version.

Jython 2.7 allows us to use the standard functions and tools in Python 2.7, so if you want to look up something in the Python docs, make sure to use version 2.7 (<https://docs.python.org/2/>).

Scripting Basics

Python is easy to learn, and with some understanding of its basic syntax, you can get started making your own scripts.

Hello World

Let's get right to everyone's favorite example, "Hello World." The following script will print out "Hello World" to the [Output Console](#).

The `print` keyword is a handy tool in Python, allowing you to write text to the Output Console. This is useful for debugging your scripts. You can print multiple things by separating them with commas.

Variables

Variables are created by simply assigning a value to them. Variables do not need to be declared, because Python has a dynamic type system. That means Python figures out the type of the variable on the fly when the script is executed.

The following script would print out: 15

Python - Using Variables

```
x=5  
y=3  
print x*y
```

On this page ...

- [About Python](#)
- [Python or Jython?](#)
 - [Which Version of Python Is Used?](#)
- [Scripting Basics](#)
 - [Hello World](#)
 - [Variables](#)
 - [Strings](#)
 - [Whitespace](#)
- [Comments](#)
 - [Individual Lines](#)
 - [Blocks of Lines](#)
 - [Comment Many Lines with a Keyboard Shortcut](#)
- [Control Flow](#)
 - [If and Else](#)
 - [For and While](#)



Basic Python - Variables and Comments

[Watch the Video](#)

Strings

Strings are defined in Python with a matching pair of 1 or 3 single or double quotes. There are few times when the type of quotation mark you use matters - but one common reason to choose one or the other is for 'escaping' other quotes inside your content. Some of the rules are shown here:

Python - Using Single and Double Quotes

```
print "This is my text"          # Using double quotation marks
print 'This is my text'         # Using single quotation marks
print "This is my text"         # This will not work because Python
does not allow mixing the single and double quotation marks
print "My name is 'David'"      # This will print: My name is 'David'
print 'My name is "David"'      # This will print: My name is "David"
print 'My name is Seamus O\'Malley' # This will print: My name is Seamus
O'Malley
```

Triple quotes (single or double) can be used to make 'escaping' both single and double quotes inside your string easier, or to write multi-line comments:

Python - Multiple Lines of Comments Using a Triple Quote

```
'''
This is a lot of text
that you want to show as multiple lines of
comments.
Script written by Professor X.
Jan 5, 1990
'''
print 'Hello world'
```

Strings can also be prefixed with certain characters to change how they are interpreted - for instance, a leading `u` character marks a string as Unicode, allowing for characters outside of the ASCII range to be used.

Python - Unicode Prefix on String

```
print u"äöü"
```

Whitespace

Perhaps Python's most unique feature is logical blocks which are defined by an indentation in Python. A colon (`:`) starts a new block, and the next line must be indented (typically using a tab or 4 spaces). The block ends when the indentation level returns to the previous level. For example, the following will print out `"5 4 3 2 1 Blast-off!"` with each value on a new line. The final `print` is not part of the while loop because it isn't indented.

Python - Logical Blocks / Indentation

```
countdown = 5
while countdown > 0:
    print countdown
    countdown = countdown - 1
print "Blast-off!"
```

Comments

Comments are a way to document your Python script. There are several ways to use comments, but the best advice we can give is to use them as much as possible! There are a few ways to make a comment in Python.

Individual Lines

You can start a line with a pound/hash (#) sign, or put one anywhere in a normal line of code.

Python - Document Scripts Using Comments

```
# this is a comment
print 'Hello world' # this is also a valid comment
```

Blocks of Lines

While Python doesn't explicitly have a way to block comment (comment out multiple lines), [multi-line strings](#) are functionally similar, and a common convention.

Comment Many Lines with a Keyboard Shortcut

In Ignition, you can use the Ctrl-/ keyboard shortcut to comment several lines of code at once. Just highlight one or more lines of code and hold the Ctrl key and press "/". This will prepend all of the selected lines of code with the pound/hash (#) sign. Press Ctrl-/ again to remove the pound/hash sign.

Control Flow

Control Flow statements, that is the ifs and loops, make the language do things differently based on the various conditions. Python has all of the basic control flow statements that you'd expect from a programming language.

If and Else

An if statement allows you to check if a condition is true or not true. Depending on the condition, you can either execute a block of code, or do something else. Many of these can be chained together to determine under what conditions should certain code execute.

Pseudocode - If Statement

```
if condition == True:
    print value1
```



Basic Python - Flow Control

[Watch the Video](#)

For and While

Looping can be done with a for, which executes a block of code a set number of times, or a while, which executes a block of code until a certain condition is met. Both can be incredibly useful.

Pseudocode - For Statement

```
for item in myList:
    print item
```

Related Topics ...

- [Scripting in Ignition](#)
- [Getting Started with Scripting in Ignition](#)

In This Section ...

Variables, Data Types, and Objects

Without incorporating any Ignition-specific objects, it is important to understand the basics of Python. This section seeks to introduce three main Python principles: Variables, Data Types, and Operators. Additionally, this page and the sub-pages herein will provide plenty of examples to get you started.

This section of the manual attempts to introduce these principles so that they may later be used in conjunction with Ignition. The information in this section is far from comprehensive, as [Python's official documentation](#) is a better reference for all things Python. However, this page serves as a great way to jump right into the action.

Variables

Variables are created by simply assigning a value to them. Variables do not need to be declared, because Python has a dynamic type system. That means, Python figures out the type of the variable on the fly when the script is executed.

The following script would print out: 15

Python - Declaring and Assigning Variables

```
x=5  
y=3  
print x*y
```

On this page ...

- [Variables](#)
- [Built-in Data Types](#)
 - [None](#)
 - [Booleans](#)
 - [Strings](#)
 - [Numeric Types](#)
 - [Colors](#)
 - [Lists and Tuples](#)
 - [Dictionaries](#)
 - [JSON](#)
 - [Datasets](#)
 - [Dates](#)
- [Using the In-Keyword](#)
- [Basic Operators](#)
 - [Operator Reference](#)
 - [Order of Operations](#)

A space may be included on either side of the assignment operator ('='), but is not required. Thus, the following example would be functionality identical as the example above.

Python - Declaring and Assigning Variables with Spaces

```
x = 5  
y = 3  
print x * y
```



Python Variables

[Watch the Video](#)

Built-in Data Types

Python features several built-in data types. Below is an overview and links to appropriate pages where applicable.

None

There is a special value in Python called `None` (with a capital N). This is simply a special value that means no value. This value is equivalent to Java's `null` value. `None` can be used to initialize a variable, but is best when checking to see if something exists before doing extra work:

Python - None

```
# If some value is not equal to None...  
if something != None:  
    # ...then do some work  
    doWork()
```



Python Datatypes

[Watch the Video](#)

Booleans

Python has two built-in values to represent true and false values: `True` and `False`, respectively (note the capital letters). These can be used when testing for truth, and are implicitly returned when the comparison operator is used:



Python - Booleans

```
# Prints True  
print 1 == 1  
  
# Prints False  
print 1 == 0
```

When using Booleans, suchs as predicates in an if-statement, you typically don't have to use True or False directly. Instead, many other values are considered True or False. The following values are also considered False:

- None
- Numeric values of 0, such as 0
- Empty Sequences and Dictionaries, like [] or {}

Values aside from the ones mentioned above are considered true, so you can easily utilize the existence of a non-zero value as a True. Note that for integers both positive and negative numbers are True, only a value of 0 is False.

Strings

Literal [strings](#) can be typed in using either double quotes or single quotes. This can be handy when your string contains one quote or the other. You can also use the backslash character to escape special characters including these quotes. See the [Strings](#) page for more information.

Numeric Types

Numbers can just be typed in normally, like 42 or 3.14159 . Adding a decimal point differentiates an Integer from a Float. More information on Numeric types can be found on the [Numeric Types](#) page.

Colors

Working with colors in Python is remarkably easy. You can simply use any tuple of 3 or 4 integers to represent a color in RGB or RGBA. For example, to set a label's text color to red, you can simple do something like this:

Python - Tuple as Color

```
label = event.source  
label.foreground = (255,0,0) #(red,green,blue)
```

Additionally, the [system.gui.color](#) function allows you to pick a color in a similar fashion:

Python - System Function as Color

```
newColor = system.gui.color(255,0,0)
```

Lists and Tuples

Python offers a variety of sequence types: most notably [Lists and Tuples](#). These are ordered collections, meaning they are indexed and the sorted order is maintained. More information on these types can be found on the [Lists and Tuples](#) page.

Python - System Function as Color

```
 newList = [1,2]  
newTuple = (2,5,7)
```

Dictionaries

Working with Different Datatypes

[Watch the Video](#)

A [Dictionary](#) is a very useful type that holds a set of key-value pairs. Unlike sequences, they are not ordered, so there is no need to sort them. Instead you give each Value in the dictionary a Key, which handles as a reference to Value. You may have used these in other languages and know them as hashmaps, maps, associative memories, or associative arrays. More information on Dictionaries can be found on the [Dictionaries](#) page.

Python - System Function as Color

```
newDictionary = {"itemName":5}
```

JSON

JSON stands for JavaScript Object Notation. While it is not a data type, it is a way of defining data in a human readable format, and is commonly used in many applications. It comes from the Javascript programming language, but it is language independent. Each JSON object contains lists and objects. A [list](#) is a series of ordered values separated by commas that is commonly used within Python. The object works like a [dictionary](#), using any number of name/value pairs, where each value could be any basic data type, a list, or even another object with its own name/value pairs. Because JSON is just a way of defining data, it can be used in many different programming languages, including Python. This makes it a useful tool for defining data in a way that humans can easily read.

Ignition also has two scripting functions that allow you to convert between a JSON string and a native Python object: [system.util.jsonEncode](#) and [system.util.jsonDecode](#).

Official Documentation

For more information on JSON, see <http://www.json.org>.

Datasets

A [Dataset](#) is a multidimensional collection of values, stored in a manner similar to how values on a spreadsheet appear. Python does **not** natively have a Dataset type. Instead these datasets were created for use inside of Ignition. There are two types of datasets:

- Dataset: Sometimes called a "Standard Dataset", this type is commonly used on many Vision Components, such as the [Power Table](#) or [Chart](#), to display multiple values simultaneously.
- PyDataset: Short for "Python Dataset", these datasets act in a manner very similar to a Python Sequence when it comes to accessing specific values, or iteration (see [Lists](#) and [Tuples](#)). Ignition's built-in system functions that interact with the database typically return a PyDataset.

While there are two types of datasets, you can easily convert one type of dataset to the other. Additionally, you can easily create a dataset from a script.

Python - Creating a Dataset

```
header = ["The Only Column"]
rows = [[1],[2],[3]]

myDataset = system.dataset.toDataSet(header, rows)
```

Dates

Dates and times can be created in Python with the `datetime` and `time` libraries. However, the Ignition's built-in `system.date` functions can also be used instead without having to import either library.

Python - Show the Current Datetime

```
currentTime = system.date.now()

print "The current time is: %s" % currentTime
```

Using the In-Keyword

Python's `in` keyword can be used to check the contents of something for a specific instance of another object. One use is to check the contents of a string for a certain substring:

Python - Using In Keyword to Look for a Substring

```

myString = "Hello World"
subString = "World"

# Here we are using an if-statement to look for the substring. If the substring exists inside of myString,
# the expression returns True.
if subString in myString:
    # This will print out, because the substring "World" exists in the string "Hello World"
    print "The word 'World' appears in myString"

```

The `in` keyword can be used to look for a certain object inside of a sequence:

Python - Check for the Presence of an Object

```

myList = [1,2,3,4]

# If the integer 3 exists in myList, then this will print True, otherwise it will print False.
print 3 in myList

```

Python - Check for the Presence of an Object

```

# Check for the Administrator role.
if "Administrator" in system.security.getRoles():
    print "Administrator found"

```

Additionally, the `in` keyword can be used in an expression on a `while` loop:

Python - Check for the Presence of a Keyword in an Expression

```

myList = [1,2,3,4]

# As long as the integer 4 is in myList, the while loop will continue to iterate.
while 4 in myList:
    # pop(0) will remove the first element in the list
    myList.pop(0)
else:
    print "all done!"

```

Basic Operators

Python has many common operators as you would expect, or are at least familiar with if you've worked with other scripting languages.

These are just the basics. There are other operators, like bit shift operators and more. Read about them at: <http://docs.python.org/library/stdtypes.html>

Operator Reference

Arithmetic Operators



In regards to Arithmetic Operators, the precision of the returned value depends on the data types of the arguments: including at least one `float` will return a `float`, otherwise, an `integer` is returned.

Operator	Meaning	Example	Output
<code>+</code>	Addition. Note that the data type of the returned object depends on the data type of the arguments being used:	<pre> print 5 + 9 print 5.0 + 9 </pre>	14 14.0

-	Subtraction	<pre>print 5 - 9</pre>	-4
*	Multiplication	<pre>print 5 * 9</pre>	45
/	Division	<pre>print 10 / 4.5</pre>	2.22222222222 22223
//	Floored Division: will divide and return just the nearest integer value, even if dividing floats.	<pre>print 10 // 4.5</pre>	2.0
%	Modulo: returns just the remainder of the dividend (the left argument) after being divided by the right argument. In the example on the right, the 4.5 divides evenly twice, leaving 1.0 as a remainder.	<pre>print 10 % 4.5</pre>	1.0
**	Power: raise the number on the left to the power of the number on the right.	<pre>print 5 ** 3</pre>	125

Boolean Operators

Operator	Meaning	Example	Output
or	Returns True if either argument is True. If both are False, then returns False.	<pre>x = False y = True print x or y</pre>	True
and	Returns True only if both arguments are True. Otherwise, returns False.	<pre>x = False y = True print x and y</pre>	False
not	Returns a boolean value that represents the opposite value of the trailing expression: False becomes True, and True becomes False.	<pre>x = True print not x</pre>	False

Comparison Operators

Operator	Meaning	Example	Output
<	Less than	<pre>print 10 < 5 print 10 < 10</pre>	False False
<=	Less than or equal to	<pre>print 10 <= 5 print 10 <= 10</pre>	False True

>	Greater than	<pre>print 10 > 5 print 10 > 10</pre>	True False
>=	Greater than or equal to	<pre>print 10 >= 5 print 10 >= 10</pre>	True True
==	Equal	<pre>print 10 == 10 print 10 == 0</pre>	True False
!=	Not equal	<pre>print 10 ! = 10 print 10 ! = 0</pre>	False True
is	Returns True if both arguments are referring to the same object, otherwise, returns False.	<pre>x = 10 y = 5 print x is y</pre>	False
is not	Returns True if both arguments are referring to different objects.	<pre>x = 10 y = 5 print x is not y</pre>	True

Order of Operations

Note that there is an order of operations for Arithmetic and Boolean operators that can be modified with parenthesis. For arithmetic they are grouped in tiers and evaluated left-to-right within a given tier:

- Parenthesis
- Exponents
- Multiplication, Division, Modulo
- Addition, Subtraction

Python - System Function as Color

```
6+10*2    # produces 6+(10*2) = 26
6*10+2    # produces (6*10)+2 = 62

10*6%2    # produces (10*6)%2 = 0
10%6*2    # produces (10%6)*2 = 8
```

For Boolean operators:

- Not
- And
- Or

Related Topics ...

- [Conditions and Loops](#)

In This Section ...

Numeric Types

The Differences Between the Types

Within Python, there are a few numeric types that we will often use within the scope of Ignition: integers, floats and booleans. **Integers** have at least 32 bits of precision and consist of non decimal values (1, 2, 50, 246). When working with larger numbers, Python also supports a **long** type that has unlimited precision. **Floats** consist of numbers with decimal values (1.0, 2.4, 50.7, 246.8734). **Booleans** are a unique subset of integers. They are either **True** or **False**, but can also be represented using a 1 or 0, respectively.

On this page ...

- The Differences Between the Types
 - Numbers in Arithmetic
 - Operators and Functions

Numbers in Arithmetic

When using numeric values in expressions or equations, Python is very flexible in how it handles multiple types in the same operation. When using two of the same type in an expression, the outcome will always be of that type. Adding two integers together yields an integer, while multiplying two floats will yield a float. With division, using integers will always result in a floor division so that the answer remains an integer. Floor division will simply remove the decimal, and will not round up at any point.

When using different types, the more precise type will be used for the outcome. Doing an operation with one integer and one float will produce an outcome that is a float. Adding a float and an integer together will yield a float.

Below, we have some examples that show off some of the different ways that Python handles expressions. You can try these out in the interactive interpreter of the [Script Console](#).

Type of Arithmetic	Description	Example	Output
Addition and Subtraction	Adding/Subtracting integers will return an Integer. If at least one of the numbers is a float, then the resulting number will also be a float. Otherwise the result will be an integer	<pre>4 + 5 # int + int = int 4.0 + 5 # float + int = float 4.0 + 5.3 # float + float = float</pre>	9 9.0 9.3
Multiplication	Similar to Addition and Subtraction, multiplication will only return a float if at least one of the numbers is a float	<pre>4 * 5 # int * int = int 4 * 5.0 # int * float = float</pre>	20 20.0
Division	Dividing an integer by another integer results in floor division: meaning the remainder of division is never returned. Note that this is not the same as rounding to the nearest integer. To include the remainder, at least one of the numbers needs to be a float	<pre>4 / 5 # int / int = int 4 / 5.0 # int / float = float</pre>	0 0.8

Operators and Functions

Python offers a wide range of operators that can be used to perform calculations on numeric values. Additionally, there are many functions that can perform specific operations to numeric values. The tables below contain some of the most common operations and functions, but this is not an exhaustive list.

Operators			
Operator	Description	Example	Types that use the Operator
+	Used to find the sum of two numeric values Can also be used to concatenate string values	<pre>5 + 1 = 6 "Hello " + "World"</pre>	Numeric Types, Strings

-	Used to find the difference of two values	5 - 1 = 4	Numeric Types
*	Used to find the product of two values	5 * 2 = 10	Numeric Types
/	Used to find the quotient of two values	4 / 2 = 2	Numeric Types
%	Modulus operator. When used with numeric types, returns the remainder of the quotient of two values The operator is also seen in string formatting .	5 % 2 = 1	Numeric Types, Strings
**	Used to find a value to the power of another value	5 ** 2 = 25	Numeric Types

Functions

Function	Description	Example
int(x)	Converts x to an integer	int(4.8) = 4
long(x)	Converts x to a long	long(5.7) = 5
float(x)	Converts x to a float	float(4) = 4.0
abs(x)	Absolute value of x	abs(-4.6)=4.6
round(x[, n])	Rounds x to n digits. If n is not specified, default is 0	round(5.674)=6.0 round(5.674, 2)=5.67

Related Topics ...

- [Strings](#)
- [Lists and Tuples](#)
- [Dictionaries](#)
- [Datasets](#)
- [Dates](#)

Strings

What are Strings?

In Python, as well as most other programming languages, a string is a grouping or string of characters that are used literally instead of as names of variables or other objects. In Python, they can be enclosed in either single quotes or double quotes:

Python - Print Statements

```
# These two print statements will both print out a sentence that looks the same as the other.  
print 'These will both print the same!'  
  
print "These will both print the same!"
```

Strings are also considered sequences, which means they can be [iterated through](#) much like a list.

On this page ...

- [What are Strings?](#)
 - [String Escape Character](#)
 - [Multi-Line Strings](#)
 - [Raw Strings](#)
 - [Unicode Strings](#)
- [Combining Strings](#)
- [String Indexing](#)
 - [Slicing Strings](#)
- [String Formatting](#)
- [String Search Example](#)
- [String Functions](#)

String Escape Character

Strings use a special character called an "escape" character to denote when something different should be happening within the string. In Python, the escape character is the backslash (\). The escape character is mainly used in a few different ways:

The Escape Character and Quotation Marks - "" and ''	Example	Output
<p>Comments</p> <p>Here, we use the escape character to let Python know we want to ignore the second single quote and directly print it.</p> <p>This works with either single or double quotes.</p>	<pre>print 'doesn\'t'</pre>	doesn't
<p>Alternately, we can enclose the string in the opposite quotes that we are using in the string.</p> <p>Because we enclose the string in double quotes here, the single quote prints normally.</p>	<pre>print "doesn't"</pre>	doesn't
<p>This also works in reverse, where we can enclose the string in single quotes and use double quotes within.</p>	<pre>print '"Test", he said.'</pre>	"Test", he said.
<p>Of course, we can also use the escape character with double quotes enclosing the string.</p>	<pre>print "\"Test\\\"", he said."</pre>	"Test", he said.
<h3>New Lines and Tab Spacing</h3> <p>A new line may be specified with the escape character and "n".</p> <p>Alternatively, a multi-line string may be used to add line breaks to a string literal.</p>	<pre>print "First line\nSecond line"</pre>	First line Second line
<p>Tab-spacing may be added to a string literal with the escape character and "t"</p>	<pre>print "Hello \tWorld"</pre>	Hello World
<h3>Specifying a Backslash</h3>		

If you need to place a backslash inside of a string literal, simply use the escape character twice.

```
print "Folder\\file.txt"
```

```
Folder\file.txt
```

Multi-Line Strings

If a string literal should contain multiple lines, you can use a multi-line string, which is created by using three quotation marks.

Example	Output
<pre>print '''This is on multiple lines'''</pre>	This is on multiple lines

This is useful when using the [system.db](#) functions, as full SQL queries are typically passed to these functions, so using a multi-line string allows you to write the query in an easy to read format.

Python - Multi-line String

```
query = '''      SELECT * FROM myTable
                  WHERE id < 10'''

system.db.runPrepQuery(query, [])
```

Raw Strings

Sometimes, it is necessary to print the raw string without allowing escaped characters. This is done by placing the letter "r" in front of the string.

Python - Raw String

```
print r'This string \n will print out directly as written.'
```

This is especially useful in cases where a file path needs to be hard coded as a string literal:

Python - File Path as a Raw String

```
# Specifying a Windows file path
myPath = r"C:\Folder\Another Folder\file.txt"
print myPath

# Specifying a Linux file path
myPath = r"/home/Directory/Another Directory/file.txt"
print myPath
```

Unicode Strings

Strings that contain characters beyond 7-bit ASCII, such as é or ? need to be marked as unicode strings by placing the letter **u** in front of the string. Implementing unicode can be additionally useful when working with [system functions](#) as some functions will fail if tags paths include special characters.

Python - Unicode String

```
print u'été'
```

Combining Strings

Two different string type variables can actually be combined or concatenated using the plus (+) sign. It is important to understand that they are concatenated exactly.

Python - Concatenated String Type Variables

```
a = "this"  
b = "that"  
  
# Will print 'thisthat' because there was no space at the end of a or the beginning of b.  
print a + b  
  
# Will print 'this that' because we added a space between them.  
print a + " " + b
```

String Indexing

Strings in Python are actually indexed, with the first character having an index of 0. To grab a value at a specific index, you place a value within square brackets ([]) after the string variable.

Python - Indexed Strings

```
a = "Ignition"  
  
# Will print out 'I', since it is in the zero position.  
print a[0]  
  
# Will print out 't', since it is in the fourth position.  
print a[4]
```

You can also use a negative index, which will start from the right side of the string with the last character having a index value of -1. This is useful for getting the last character when you aren't sure how long the string is.

Python - Negative Indexed String

```
a = "Ignition"  
  
# Will print out 'o', since it is the second to last character.  
print a[-2]
```

Slicing Strings

Using the string index values, we can actually grab slices or parts of the string. Similar to grabbing an individual character, we place two values separated by a colon within square brackets. You can think of these numbers as the slices between characters, with 0 before the first character, 1 before the second, and so on. It should look like this: string[4:7] The first value is the start location, while the second value is the end location.

Python - Sliced Strings

```
a = "Inductive Automation"  
  
# Will print out 'ctive Auto'. Note how the space is counted as a character,  
# and the 14th location is before the m, which is not included.  
print a[4:14]
```

If left blank, the first value will default to 0, while the second value will default to the length of the string. Additionally, negative values can be used as well, just like with the index.

Python - Negative Value Sliced Strings

```
a = "Inductive Automation"

# From index 0 to index 12.
print a[:12] # Will print out: Inductive Au

# From index 12 to the end of the string.
print a[12:] # Will print out: tomation

# From the 4th character from the end to the 5th index. Note this prints nothing because the start character
# is after the end character.
print a[-4:5] # Will print out:

# From the 8th character from the end to the 2nd character from the end.
print a[-8:-2] # Will print out: tomati

# From index 5 to the 5th character from the end.
print a[5:-5] # Will print out: tive Autom
```

You can use all of these together (and with integer variables or functions) to have a very flexible way to format strings.

String Formatting

Also known as string substitution, it allows you to enter in values into a string, similar to having a variable inside of the string. This is useful when doing a database query where you can build the query as a string, and then add in the query parameters when executing the query.

Doing this requires the use of the percent (%) sign followed by the parameter type within the string. After the string, you then use the percent sign followed by the parameter values within parentheses in order.

Python - String Substitution

```
a = "I have %i apples."

# Will print exactly as written, because we are not substituting any values in.
print a

# Will print with an 8 substituted in for the %i. I have 8 apples.
print a % (8)
```

Instead of using literal values, you can also use variables as parameter values instead. This is useful when pulling the value from a function. Each string can also have multiple values substituted, as long as the parameters are in the order that they are in the string.

Python - String Substitution using Variables as Parameters

```
apples = 10
oranges = 14
peaches = 5

a = "I have %i apples, %i oranges, and %i peaches."

# Will print with the variable values substituted in for the %i. I have 10 apples, 14 oranges, and 5 peaches.
print a % (apples, oranges, peaches)
```

There are a few different types used with string formatting, the most common of which are listed here. The character values are what go after the percent sign.

Type	Character
Signed Integer	i
Floating Point Decimal Format	f
Floating Point Exponential Format	e

String

s

String Search Example

You can bring all this together for a simple way to print out just the last word in a string.

Python - Find the Last Word of a String

```
# Find just the last word of a string.
myString = "Inductive Automation"

lastSpaceIndex = myString.rfind(" ")

# Add one to not include the space.
print myString[lastSpaceIndex+1:] # Will print out: Automation
```

String Functions

Python strings have many functions available that can manipulate the string or give information on the string. The most common are in the table below.

Function	Description	Example	Output
<code>len(string)</code>	Returns the length of the string .	<pre>a = "Inductive Automation" print len(a)</pre>	20
<code>x in string</code>	Will return True if x is within the string, False if not. Can also be used to iterate through the string.	<pre>string = "Inductive Automation" if 'd' in string: print "There is a d in the string" for letter in string: print letter</pre>	There is a d in the string l n d u c t i ...
<code>string.find(x, [start, end])</code>	Returns the first location of the substring x from the string. Returns -1 if the substring is not found.	<pre>string = "Inductive Automation" print string. find(" ")</pre>	9
<code>string.rfind(x, [start, end])</code>	Similar to find, but returns the last location of the substring x . Returns -1 if not found.	<pre>string = "Inductive Automation" print string. find("i")</pre>	17

string.upper()	Returns a copy of the string with all characters uppercase. This is useful when comparing user input to a string value, as the user may use a different case for certain letters.	<pre>userInput = "administration" print userInput.upper()</pre>	ADMINISTRATION
string.lower()	Returns a copy of the string with all characters lowercase. This is useful when comparing user input to a string value, as the user may use a different case for certain letters.	<pre>userInput = "ADminIStration" print userInput.lower()</pre>	administration
string.capitalize()	Returns a copy of the string with the first character capitalized and all other characters lowercase.	<pre>string = "here is my sentence." print string.capitalize()</pre>	Here is my sentence.
string.title()	Returns a copy of the string with the first letter of each word capitalized and all other characters lowercase.	<pre>string = "here is my sentence." print string.title()</pre>	Here Is My Sentence.
string.strip([x])	Returns a copy of the string with leading and trailing characters removed, where x is the string of characters.. If x is omitted, then removes whitespace from the leading and trailing edges of the string. lstrip() and rstrip() may instead be used to strip characters from the leading or trailing edge	<pre>string = " This string has some empty space This string has some empty space My String My String_____ _____ ____My String fencedString = " ____My String____ " print fencedString.strip(" _") print fencedString.lstrip(" _") print fencedString.rstrip(" _")</pre>	This string has some empty space This string has some empty space My String My String_____ _____ ____My String

string.count([x[, start[, end]]])	Returns the number of occurrences of x in the string. A start and end can be specified that will limit the count to that area.	<pre>string = "Inductive Automation" print string.count('i') print string.count('i', 4, 9)</pre>	2 1
string.split([delimiter[, maxsplit]])	Returns a list of the words in the string. Optionally, specifying a delimiter will split the string on the delimiter string. Specifying a maxsplit will split the string a maximum number of times, with the remainder of the string as the final list object. The number of items in the list will not be more than maxsplit + 1	<pre>sentence = "This is an example of split" print sentence.split() dashedSentence = "Yet-another-sentence-here" print dashedSentence.split("-", 2)</pre>	['This', 'is', 'an', 'example', 'of', 'split'] ['Yet', 'another', 'sentence', 'here']
string.rsplit([delimiter[, maxsplit]])	Similar to <code>split</code> , but splitting is performed from right to left. Thus if maxsplit is smaller than the total number of delimiter in the string, only the rightmost words will be split off as separate items.	<pre>sentence = "This is an example of split" print sentence.rsplit() dashedSentence = "Yet-another-sentence-here" print dashedSentence.rsplit("-", 2)</pre>	['This', 'is', 'an', 'example', 'of', 'split'] ['Yet', 'another', 'sentence', 'here']
string.join(x)	Returns a copy of the string that is the concatenation of the strings in the iterable x . The string calling the <code>join</code> is what will separate the values of the iterable.	<pre>a = '-' b = 'abcdef' print a.join(b)</pre>	a-b-c-d-e-f
string.replace(old, new[, count])	Returns a copy of the string where occurrences of the old substring are replaced with the new substring. Optionally, if count is specified, will only replace the first count number of occurrences.	<pre>string = "Ignition is good!" Ignition is awesome!</pre>	Ignition is awesome!

```
print string.  
replace  
( "good" ,  
"awesome" )
```

Related Topics ...

- [Numeric Types](#)
- [Lists and Tuples](#)
- [Dictionaries](#)
- [Datasets](#)
- [Dates](#)

Lists and Tuples

Sequences

Like [Strings](#), Python has two other common sequence types: Lists and Tuples. While lists and tuples are similar to strings in that they share many of the same functionality, they are unique in that they are used to group other values together. Both lists and tuples define a number of values separated by commas, but lists are enclosed in square brackets [] and are mutable, meaning their contents can change, while tuples are enclosed in parentheses () and are immutable, meaning their contents can't change.

Lists

As stated above, lists are groups of comma separated values enclosed in square brackets, but can utilize many of the features available to other sequences like strings.

Python - Lists

```
# Lists are very simple to create.  
myList = [1, 2, 3]  
  
print myList # Will print out: [1, 2, 3]  
  
# Empty lists can also be created, to have items added to them later.  
myList = []  
  
# Lists are not confined to hold values of a single data type either.  
myList = [1, "hello", 3.3]  
  
# Lists can even hold other lists! In this case, myList would actually  
hold 5 elements:  
# 1, a list, 4, 'that', and another list. The last list also contains a  
list as well.  
myList = [1, ["this", 3.3], 4, 'that', [6, [7.7, 'other'], 9]]
```

On this page ...

- [Sequences](#)
- [Lists](#)
 - [List Concatenation](#)
 - [List Indexing](#)
 - [Appending to Lists](#)
- [List Functions](#)
- [Tuples](#)
- [Tuple Functions](#)



INDUCTIVE
UNIVERSITY

Basic Python - Lists and Dictionaries

[Watch the Video](#)

List Concatenation

Like strings, lists also support concatenation and can be combined to form a single larger list.

Python - List Concatenation

```
a = [1, 2, 3]  
b = ['four', 'five', 'six']  
  
print a + b # Will print out [1, 2, 3, 'four', 'five', 'six']
```

List Indexing

Like strings, lists are also indexed, which allows you to grab single values or splice them to get ranges of values. Just like everything else in Python, lists start with 0.

Python - List Indexing

```
myList = ['a', 'b', 'c', 'd', 'e', 'f', 'g']  
  
print myList[3] # Will print out: d  
  
print myList[-2] # Will print out: f  
  
print myList[2:5] # Will print out: ['c', 'd', 'e']
```

When trying to index a nested list, we simply need to add a second index after the first.

Python - Index a Nested List

```
myList = [1, 2, [3, [4, 5, 6, 7], 8], 9, 10]

print myList[2][1][3] # Will print out: 7
```

Appending to Lists

You can add values directly to the end of a list with the `append()` function. You can add anything that is normally allowed in a list.

Python - Appending to a List

```
myList = []

myList.append('Hello')
myList.append('World')

print myList # Will print out ['Hello', 'World']
```

List Functions

Below is a list of common list functions. Some of them are similar to other sequences like strings, while others are unique to lists because lists are mutable.

Function	Description	Example	Output
<code>len(list)</code>	Returns the length of the <code>list</code> .	<pre>list = [1, 2, 3, 4, 5] print len(list)</pre>	5
<code>x in list</code>	Will return <code>True</code> if <code>x</code> is within the list, <code>False</code> if not. Can also be used to iterate through the list.	<pre>list = [1, 2, 3, 4, 5] if 4 in list: print "There is a 4 in the list" for value in list: print value</pre>	There is a 4 in the list 1 2 3 4 5
<code>list.index(x)</code>	Will return the index number of item <code>x</code> . Throws an error if item is not found.	<pre>list = [1, 2, 3, 4, 5] print list.index(3)</pre>	2
<code>min(list)</code>	Returns the smallest item of <code>list</code> .	<pre>list = [1, 2, 3, 4, 5] print min(list)</pre>	1
<code>max(list)</code>	Returns the largest item of <code>list</code> .	<pre>list = [1, 2, 3, 4, 5] print max(list)</pre>	5
<code>list.count(x)</code>	Will return the number of times <code>x</code> appears in the list.		2

		<pre>list = [1, 2, 3, 4, 5, 4] print list.count(4)</pre>	
list.append(x)	Will add x to the end of the list.	<pre>list = [1, 2, 3, 4, 5] list.append(6) print list</pre>	[1, 2, 3, 4, 5, 6]
list.insert(i, x)	Will insert x at position i .	<pre>list = [1, 2, 3, 4, 5] list.insert(1, 5) print list</pre>	[1, 5, 2, 3, 4, 5]
list.remove(x)	Will remove the first x from the list.	<pre>list = [1, 2, 3, 2, 4, 5] list.remove(2) print list</pre>	[1, 3, 2, 4, 5]
list.pop(i)	Will remove the item at index i , and return it. If no index is specified, it will remove and return the last item in the list.	<pre>list = [1, 2, 3, 4, 5] list.pop(2) print list list.pop() print list</pre>	3 [1, 2, 4, 5] 5 [1, 2, 4]
list.reverse()	Will reverse the items in the list.	<pre>list = [1, 2, 3, 4, 5] list.reverse() print list</pre>	[5, 4, 3, 2, 1]

Tuples

Tuples look similar to lists in that they are defined as a group of comma separated values, but they are enclosed by parenthesis and they are immutable like strings, meaning they can't be altered. Besides that, they are sequences, like lists and strings. This means that tuples have the functionality that other sequences have, such as concatenation, indexing and slicing, and even nesting.

Python - Tuples

```
# Tuples are very simple to create.
myTuple = (1, 2, 3)

print myTuple # Will print out: (1, 2, 3)

# Empty tuples can also be created, to have items added to them later.
myTuple = ()

# Like lists, tuples are not confined to hold values of a single data type either.
```

```

myTuple = (1, "two", 3.3)

# Tuples can even hold other tuples!
myTuple = (1, ("two", 3.3), 4, 'five', (6, (7.7, 'eight')), 9))

a = (1, 2, 3)
b = (4, 5, 6)

# Combine two tuples to make a new tuple
print a + b # Will print out: (1, 2, 3, 4, 5, 6)

myTuple = ('a', 'b', 'c', 'd', 'e', 'f', 'g')

print myTuple[2:5] # Will print out: ('c', 'd', 'e')

```



Lists and tuples can also be nested within each other!

Tuple Functions

All of the functions that work on tuples work on lists, so these should look familiar. However, because they are immutable, not all of the list functions work on tuples. Check out the common tuple functions below.

Function	Description	Example	Output
<code>len(tuple)</code>	Returns the length of the tuple .	<pre> tuple = [1, 2, 3, 4, 5] print len(tuple) </pre>	5
<code>x in tuple</code>	Will return True if x is within the tuple, False if not. Can also be used to iterate through the tuple.	<pre> tuple = [1, 2, 3, 4, 5] if 4 in tuple: print "There is a 4 in the tuple" for value in tuple: print value </pre>	There is a 4 in the tuple 1 2 3 4 5
<code>min(tuple)</code>	Returns the smallest item of tuple .	<pre> tuple = [1, 2, 3, 4, 5] print min(tuple) </pre>	1
<code>max(tuple)</code>	Returns the largest item of tuple .	<pre> tuple = [1, 2, 3, 4, 5] print max(tuple) </pre>	5

Related Topics ...

- [Numeric Types](#)
- [Strings](#)
- [Dictionaries](#)
- [Datasets](#)
- [Dates](#)

Dictionaries

Mapping Type

A Dictionary is a mapping object. Where sequences are indexed with a numeric index value, dictionaries are indexed using keys. These keys then have a matching value pair that is associated with a particular key. For example, with a list we can extract the object at index 0, which we may have decided is the name, whereas with dictionaries, I can instead extract the a value using a key "name". Because of how they work, dictionaries are sometimes known as associative arrays in other programming languages.

Dictionaries are created using braces { }, where each key/value pair is separated by a comma (,) and keys are separated from their values using a colon (:). In the example below, I created a dictionary with two keys: name, id.

Python - Creating a Dictionary

```
# In this dictionary, I associated the name John Smith to the key "name",
# and the id number 12345 to the key "id".
myDictionary = {"name": "John Smith", "id": 12345}
```

On this page ...

- [Mapping Type](#)
 - [Using a Dictionary](#)
 - [Dictionary Functions](#)



INDUCTIVE
UNIVERSIT

Basic Python - Lists and Dictionaries

[Watch the Video](#)

Using a Dictionary

The keys in a dictionary can be numbers, strings, or tuples, but typically a string is used to make a key that best describes the value. Any given key may only appear once in a dictionary, so trying to set another value for a key that already exists will overwrite the previous value for that key. Alternately, attempting to access the value of a key that does not exist will throw an error, while setting a value to a key that does not exist will create a new key/value pair within the dictionary.

To access a value in a dictionary works much like accessing a value in a list; simply place brackets containing the key after the dictionary object.

Python - Accessing Values in a Dictionary

```
# Creates a dictionary with three key/value pairs.
myDictionary = {'Bob': 89.9, 'Joe': 188.72, 'Sally': 21.44}

print myDictionary['Joe'] # Will print out: 188.72

# Adds a key for 'Amir', and alters the value associated with the key 'Sally'.
myDictionary['Amir'] = 45.89
myDictionary['Sally'] = 146.23

print myDictionary # Will print out the whole dictionary: {'Joe': 188.72, 'Amir': 45.89, 'Bob': 89.9,
'Sally': 146.23}
```

It is also easy to loop through all of the values of a dictionary using the keys() function. For example:

Python - Keys Function

```
# The keys() function provides us with a list of keys, which we can iterate through and print out in
addition to using in the value lookup.
for key in myDict.keys():
    print key, myDict[key]
```

There are many use cases for dictionaries, but they are commonly used in Ignition when passing values into a Message Handler or [creating a dynamic roster](#) for alarms.

Dictionary Functions

Dictionaries have a few functions that allow for greater control over the dictionary object and the values contained within.

Function	Description	Example	Output
len(dictionary)	Returns the number of items in the dictionary.	<pre>myDictionary = {"name": "John Smith", "id": 12345} print len (myDictionary)</pre>	2
del dictionary[key]	Will remove the named key.	<pre>myDictionary = {"name": "John Smith", "id": 12345} del myDictionary["id"] print myDictionary</pre>	{'name': 'John Smith'}
key in dictionary	Will return True if the dictionary has that key. Can also use "key not in dictionary"	<pre>myDictionary = {"name": "John Smith", "id": 12345} if "name" in myDictionary: print myDictionary["name"]</pre>	John Smith
dictionary.clear()	Remove all of the items in the dictionary.	<pre>myDictionary = {"name": "John Smith", "id": 12345} myDictionary.clear() print myDictionary</pre>	{}
dictionary.keys()	Returns a list of the dictionary's keys.	<pre>myDictionary = {"name": "John Smith", "id": 12345} print myDictionary.keys()</pre>	['name', 'id']
dictionary.values()	Returns a list of the dictionary's values.	<pre>myDictionary = {"name": "John Smith", "id": 12345} print myDictionary.values()</pre>	['John Smith', 12345]

Related Topics ...

- [Client Message Handler](#)
- [Gateway Message Handler](#)
- [Numeric Types](#)
- [Strings](#)
- [Lists and Tuples](#)
- [Datasets](#)
- [Dates](#)

Datasets

Datasets and PyDatasets

A dataset can be thought of as a two dimensional list, or rather a list where each object is another list of objects. Datasets are not normally native to Python, but are built into Ignition because of their usefulness when dealing with data from a database. It is very common to deal with datasets in scripting, as datasets power many of the interesting features in Ignition, like charts and tables.

The main confusion when dealing with datasets is the difference between the dataset object and the PyDataset object. Dataset is the kind of object that Ignition uses internally to represent datasets. When you get the data property out of a component like a Table, you will get a dataset. The PyDataset is a wrapper type that you can use to make datasets more accessible in Python. The biggest differences are seen in how we access the data in the two different objects. However, you can easily convert between the two with `system.dataset.toDataSet` and `system.dataset.toPyDataSet`, making it simple to use the object that you find easier to use.

Creating Datasets

Because datasets are not native to Python, there is no way to naturally create them within scripting. Instead they must be created using the `system.dataset.toDataSet` function, which also allows you to convert a PyDataset to a Dataset. It requires a list of headers and a list of each row's data. Each data row list must be the same length as the length of the headers list.

Python - Creating a Dataset

```
# First create a list that contains the headers, in this case there are 4
# headers.
headers = ["City", "Population", "Timezone", "GMTOffset"]

# Then create an empty list, this will house our data.
data = []

# Then add each row to the list. Note that each row is also a list object.
data.append(["New York", 8363710, "EST", -5])
data.append(["Los Angeles", 3833995, "PST", -8])
data.append(["Chicago", 2853114, "CST", -6])
data.append(["Houston", 2242193, "CST", -6])
data.append(["Phoenix", 1567924, "MST", -7])

# Finally, both the headers and data lists are used in the function to
# create a Dataset object
cities = system.dataset.toDataSet(headers, data)
```

On this page ...

- [Datasets and PyDatasets](#)
- [Creating Datasets](#)
- [Accessing Data in a Dataset](#)
 - [Looping Through a Dataset](#)
- [Accessing Data in a PyDataset](#)
 - [Looping Through a PyDataset](#)
 - [PyRow](#)
- [Altering a Dataset](#)



INDUCTIVE
UNIVERSITY

Working with Datasets

[Watch the Video](#)

Note: All code snippets on this page will reference the cities dataset we created above, so place that code at the beginning of every code snippet.

Accessing Data in a Dataset

To access the data inside of a dataset, each dataset has a few functions that can be called on to access different parts of the dataset. These are listed in the table below.

Function	Description	Example	Output
<code>data.getColumnAsList(colIndex)</code>	Returns the column at the specified index as a list.	<pre>print cities. getColumnAsList(0)</pre>	[New York, Los Angeles, Chicago, Houston, Phoenix]
<code>data.getColumnCount()</code>	Returns the number of columns in the dataset.	<pre>print cities. getColumnCount()</pre>	4

data. getColumnIndex (colName)	Returns the index of the column with the name colName.	print cities. getColumnIndex ("Timezone")	2
data. getColumnName (colIndex)	Returns the name of the column at the index colIndex.	print cities. getColumnName(1)	Population
data. getColumnNames()	Returns a list with the names of all the columns.	print cities. getColumnNames()	[City, Population, Timezone, GMTOffset]
data. getColumnType (colIndex)	Returns the type of the column at the index.	print cities. getColumnType(3)	<type 'java.lang.Integer'>
data. getColumnTypes()	Returns a list with the types of all the columns.	print cities. getColumnTypes()	[class java.lang.String, class java.lang.Integer, class java.lang.String, class java.lang.Integer]
data. getRowCount()	Returns the number of rows in the dataset.	print cities. getRowCount()	5
data.getValueAt (rowIndex, colIndex)	Returns the value at the specified row and column indexes.	print cities. getValueAt(1, 2)	PST
data.getValueAt (rowIndex, colName)	Returns the value at the specified row index and column name.	print cities. getValueAt(2, "Population")	2853114

Looping Through a Dataset

Oftentimes you need to loop through the items in a dataset similar to how you would loop through a list of items. You can use the functions above to do this.

Python - Looping Through a Dataset

```
# We use the same cities dataset from above. Using the range function, we can come up with a range of values that represents the number of columns.
for row in range(cities.getRowCount()):
    for col in range(cities.getColumnCount()):
        print cities.getValueAt(row, col) # Will print out every item in our cities dataset, starting on the first row and moving left to right.
```

Accessing Data in a PyDataset

Note: PyDatasets can be accessed in the same ways that Datasets can. This means that all of the above functions (getColumnCount(), getValueAt(), etc) can be used with PyDatasets too.

PyDatasets are special in that they can be handled similarly to other Python sequences. Any dataset object can be converted to a PyDataset using the function `system.dataset.toPyDataSet`. All of the functions listed above can be used on a PyDataset, but the data can also be accessed much easier, similar to how you would a list.

Python - Accessing Data in a PyDataset

```
# First convert the cities dataset to a PyDataset.  
pyData = system.dataset.toPyDataSet(cities)  
  
# The data can then be accessed using two brackets at the end with row and column indexes. This will print  
"PST"  
print pyData[1][2]
```

Looping Through a PyDataset

Looping through a PyDataset is also a bit easier to do, working similar to other sequences. The first for loop will pull out each row, which acts like a list and can be used in a second for loop to extract the values.

Python - Looping Through a PyDataset

```
# Convert to a PyDataset  
pyData = system.dataset.toPyDataSet(cities)  
  
# The for loop pulls out the whole row, so typically the variable row is used.  
for row in pyData:  
    # Now that we have a single row, we can loop through the columns just like a list.  
    for value in row:  
        print value
```

Additionally, a single column of data can be extracted by looping through the PyDataset.

Python - Extract a Column of Data by Looping Through a PyDataset

```
# Convert to a PyDataset  
pyData = system.dataset.toPyDataSet(cities)  
  
# Use a for loop to extract out a single row at a time  
for row in pyData:  
    # Use either the column index or the column name to extract a single value from that row.  
    city = row[0]  
    population = row["Population"]  
    print city, population
```

PyRow

A PyRow is a row in a PyDataset. It works similarly to a Python list.

The examples and outputs are based on the results in the table below. In addition, "print" commands are used, but should be replaced by appropriate logging methods (such as `system.util.getLogger`) depending on the scope of the script.

A	B
Apple	Orange
Banana	Orange
Apple	Apple

Method	Description	Syntax	Example	Output
index()	Returns the index of first occurrence of the element. Returns a ValueError if the element isn't present in the list.	index(element)	for row in pyDataset: try:	0 No apples

		<pre> print row. index("Apple") except: print "No apples in this row" </pre>	in this row 0
count()	Calculates total occurrence of given element in the row.	count (element)	<pre> for row in pyDataset: print row.count ("Apple") </pre>

Repeating Elements

You can also have repeating elements in a row:

Example	Output
<pre> for row in PyDataset print row * 2 </pre>	<pre> [u'Apple', u'Orange', u'Apple', u'Orange'] [u'Banana', u'Orange', u'Banana', u'Orange'] [u'Apple', u'Apple', u'Apple'] </pre>

Altering a Dataset

Technically, you cannot alter a dataset. Datasets are immutable, meaning they cannot change. You can, however, create new datasets. To change a dataset, you really create a new one and then replace the old one with the new one. There are system functions that are available that can alter or manipulate datasets in other ways. Any of the functions in the [system.dataset](#) section can be used on datasets, the most common ones have been listed below:

- [system.dataset.addRow](#)
- [system.dataset.deleteRow](#)
- [system.dataset.setValue](#)
- [system.dataset.updateRow](#)

The important thing to realize about all of these datasets is that, again, they do not actually alter the input dataset. They return a new dataset. You need to actually use that returned dataset to do anything useful.

For example, the following code is an example of the `setValue` function, and would change the population value for Los Angeles.

Python - Altering a Dataset Using the `setValue` Function

```

# Create a new dataset with the new value.
newData = system.dataset.setValue(cities, 1, "Population", 5000000)

# The cities dataset remains unchanged, and we can see this by looping through both datasets.
for row in range(cities.getRowCount()):
    for col in range(cities.getColumnCount()):
        print cities.getValueAt(row, col)

for row in range(newData.getRowCount()):
    for col in range(newData.getColumnCount()):
        print newData.getValueAt(row, col)

```

[Related Topics ...](#)

- Numeric Types
- Strings
- Lists and Tuples
- Dictionaries
- Dates

Dates

Dates can normally be tricky since they generally require very specific formats. Furthermore, some functions/objects require a date object instead of a string. Fortunately, there are several ways to create and alter date objects with scripting in Ignition.

Python has some [built-in libraries](#) to create and manipulate dates and times. However, most users find both Ignition's built-in [system functions](#) and even Java's Calendar class easier to use. Regardless, this section will demonstrate some examples from each approach.

Ignition's System Functions

Ignition's [system.date](#) library has a large number of functions that provide easy access to datetime creation and manipulation. This page has just a few simple examples. Additional examples and functions can be found in the [scripting appendix](#).

Creating Dates

New datetimes can be created by using either the [system.date.now](#) or [system.date.getDate](#). The [system.date.getDate](#) function returns a datetime, but the time is set to midnight. However, we can use [system.date.setTime](#) to change the time.

Python - System Functions - Creating Dates

```
# Get the current datetime.  
print system.date.now()  
  
# Create a date. The time will be set to midnight.  
newDate = system.date.getDate(2018, 10, 28)  
print newDate  
  
# Change the time on the new date to 11:30 am.  
print system.date.setTime(newDate, 11, 30, 0)
```

On this page ...

- [Ignition's System Functions](#)
 - [Creating Dates](#)
 - [Formatting Dates](#)
 - [Date Arithmetic](#)
 - [Date Formatting Characters](#)
- [Java's Calendar Class](#)
 - [Creating Dates](#)
 - [Date Arithmetic](#)
- [Python's Time and Datetime Libraries](#)
 - [Creating Dates - Python's Time Library](#)
 - [Creating Dates - Python's Datetime Library](#)
 - [Date Arithmetic](#)



Basic Python - Dates, Colors, and JSON Strings

[Watch the Video](#)

Formatting Dates

When printed, datetimes default to a format like the following: **Sun Jan 1 00:00:00 TZ 2018**. However, this can be manipulated by using special characters in the [system.date.format](#) function:

Python - System Functions - Date Formatting

```
rightNow = system.date.now()  
  
# Demonstrating the standard format.  
print rightNow  
  
# Demonstrating the modified format.  
print system.date.format(rightNow, "YYYY-MM-dd HH:mm:ss")
```

Date Arithmetic

The [system.date.add*](#) functions can be used to add to or subtract some amount of time from a date. See the [system.date.add*](#) functions for more information.

Python - System Functions - Date Arithmetic

```
# Get the current datetime.  
newDate = system.date.now()
```

```
# Change the time on the new date to 30 minutes ago.
print system.date.addMinutes(newDate, -30)
```

Date Formatting Characters

The following is a reference of date formatting characters that can be used by `system.date.format` or Java's `DateFormat` class. Additionally, there are many other non-scripting uses in Ignition (such as the [Vision - Calendar](#) component's Format String property) that can utilize this reference.

Symbol	Description	Presentation	Example	Other Notes
G	Era designator	Text	G=AD	
y	Year	Year	yyyy=1996; yy=96	Lowercase y is the most commonly used year symbol
Y	Week year	Year	YYYY=2009; YY=09	Capital Y gives the year based on weeks (ie. changes to the new year up to a week early)
M	Month in year	Month	MMMM=July; MMM=Jul; MM=07	
w	Week in year	Number	27	If Dec31 is mid-week, it will be in week 1 of the next year
W	Week in month	Number	2	
d	Day in year	Number	189	
d	Day in month	Number	10	
F	Day of week in month	Number	2	2nd Sunday of the month
E	Day name in week	Text	EEEE=Tuesday; E=Tue	
u	Day number of week	Number	1	(1 = Monday, ..., 7 = Sunday)
a	Am/Pm marker	Text	PM	
H	Hour in day (0-23)	Number	0	
h	Hour in am/pm (1-12)	Number	12	
k	Hour in day (1-24)	Number	24	
K	Hour in am/pm (0-11)	Number	0	
m	Minute in hour	Number	30	
s	Second in minute	Number	55	
S	Millisecond	Number	978	
z	Time zone	General time zone	zzzz=Pacific Standard Time ; z=PST	
Z	Time zone	RFC 822 time zone	Z=-0800	
X	Time zone	ISO 8601 time zone	X=-08; XX=-0800; XXX=-08:00	

Java's Calendar Class

While Java's `Calendar` class is useful, in many cases Ignition's built-in `system.date` functions are simpler to use. Furthermore, the `system.date` functions typically use the `Calendar` class to retrieve the current time, so you are not losing any functionality by using the `system` functions.

Ignition's System Functions vs Java's Calendar Class

It is highly advisable to use [Ignition's system functions](#) to generate and manipulate dates. The information on this page pertaining to the `Calendar` class is maintained in the interest for legacy installations.

Creating Dates

To create an arbitrary date, you can use the `java.util.Calendar` class. It has various functions to alter the calendar fields, like `Calendar.HOUR`, `Calendar.MONTH`, and so on. After you're done manipulating the `Calendar`, you can use its `getTime()` function to retrieve the `Date` represented by the calendar. It also has a handy `set()` function that takes the common parameters of a `Date`. The one major "gotcha" here is that January is month zero, not month one. For example:

Python - Calendar Class - Creating Dates

```
from java.util import Calendar
cal = Calendar.getInstance()

# set year, month, day, hour, minute, second in one call
# This sets it to Feb 25th, 1:05:00 PM, 2010
cal.set(2010, 1, 25, 13, 5, 0)
myDate = cal.getTime()
```

Date Arithmetic

Often you'll have a `Date` object from a component like the [Popup Calendar](#) and want to alter it programmatically. Say, subtracting 8 hours from it, or something like that. The `java.util.Calendar` class is used for this as well. Following the example above, this code would subtract 8 hours from the variable `myDate`.

Python - Calendar Class - Date Arithmetic

```
from java.util import Calendar
cal = Calendar.getInstance()
cal.setTime(myDate)
cal.add(Calendar.HOUR, -8)
myNewDate = cal.getTime()
```

Python's Time and Datetime Libraries

Many components in Ignition that contain a `Date` property actually expect a Java calendar object. Creating a `datetime` object using Python's built-in libraries and passing them to a date property on a component will result in an exception.

Ignition's System Functions vs Python's Libraries

It is highly recommended to use [Ignition's built-in system.date](#) functions.

Creating Dates - Python's Time Library

The time library can be used to return dates as well as time. Times are created as a tuple of integers. The integers represent the following values: year, month, day of the month, hour, minute, second, weekday, day of the year, daylight savings time)

Check out [Python's time library documentation](#) for more information.

Python - Python Library - time

```
import time

# Finds the current local time. The time is returned as a tuple of integers.
myTime = time.localtime()

# Print the time into a 24-character string with the following format: Sun Nov 20 12:00:00 2017
print time.asctime(myTime)

# Alternatively, we can reformat the time in a custom manner, then print it
print time.strftime('%H:%M:%S %b %d %Y', myTime)
```

Creating Dates - Python's Datetime Library

Python's datetime library offers a bit more flexibility since arithmetic can easily be applied. Additional information on the datetime library can be found in [Python's official documentation](#).

Notice the double use of 'datetime' in the example below. This is because the 'datetime' library has a class named 'datetime.'

Python - Python Library - datetime

```
import datetime

# Returns the current datetime.
print datetime.datetime.now()
```

However, we can clean up the above by importing the datetime class from the library:

Python - Python Library - datetime

```
# Imports the class named 'datetime' from the 'datetime' library, so we don't have to state it twice.
from datetime import datetime

# Returns the current datetime.
print datetime.now()
```

If you need to create a specific datetime, instead of just using the current, you can pass in the values directly when creating an instance of datetime:

Python - Python Library - Creating a New Time

```
from datetime import datetime

# Prints out the following datetime: 2018-01-02 03:04:05.000006
print datetime(2018,1,2,3,4,5,6)
```

Finding the difference between two datetime objects can easily be accomplished by using the '-' character

Python - Python Library - Date Difference

```
from datetime import datetime

rightNow = datetime.now()
someTime = datetime(2018,1,1,1,1,1,1)

# Find the difference between the two dates.
print someTime - rightNow
```

Date Arithmetic

With Python's built-in libraries, the timedelta class provides the simplest way to perform arithmetic on a date: It simply creates an object that effectively represents a duration. The duration can then be applied to a datetime.

Python - Python Library - Date Arithmetic

```
# We're including the timedelta class here
from datetime import datetime, timedelta

rightNow = datetime.now()

# Creating a timedelta object, and setting the hours to 8
offset = timedelta(hours = 8)

# Print the current time, and then print the time minus the offset.
print rightNow
print rightNow - offset
```

Related Topics ...

- [Numeric Types](#)
- [Strings](#)
- [Lists and Tuples](#)
- [Dictionaries](#)
- [Datasets](#)

Conditions and Loops

If-Statements

The `if` statement should be familiar to anyone with a passing knowledge of programming. The idea of an `if` is that you want your script to execute a block of statements only when a certain condition is true. Python's `if` is simple to use, and has some additional keywords to provide more flexibility.

Simple If-Statement Example

The syntax for `if` is as follows:

Pseudocode - If Statement

```
# Note that 'if' uses lowercase characters.  
# Additionally, a colon is placed after the expression.  
if expression:  
  
    # The statements that should execute when the expression is true  
    # MUST be indented.  
    statement
```

Example

```
x = 5  
z = 15  
if x < 10:  
    # Since the condition "x < 10" is true,  
    # the following line will execute  
    print "'x' is less than 10"  
if z < 10:  
    # This condition "z < 10" is false,  
    # so the following line will not execute  
    print "this will never show"
```

Output

'x' is less than 10

On this page ...

- If-Statements
 - Simple If-Statement Example
 - If and Else
 - Elif (Else If)
- For-Loops and While-Loops
 - For-Loop
 - While-Loop
 - The Break and Continue Statements
 - The Pass Keyword



INDUCTIVE
UNIVERSITY

Control Flow Logic

[Watch the Video](#)

If and Else

You can use the `if...else` form of an `if` statement to do one thing if a condition is true, and something else if the condition is false.

Example	Output
<pre>x = 15 if x < 10: print "x is less than 10" else: print "x is not less than 10"</pre>	x is not less than 10

Elif (Else If)

Lastly, you can use the `if...elif` form. This form combines multiple condition checks. `elif` stands for "else if". This form can optionally have a catch-all `else` clause at the end. For example, this script will print out `three`:

Example	Output
<pre>x = 3 if x == 1: print "one" elif x == 2: print "two"</pre>	three

```

    elif x == 3:
        print "three"
    else:
        print "not 1-3"

```

You can use as many `elif` items as you want, and the `else` is not required at the end.

For-Loops and While-Loops

For-Loop

Python's `for` loop may be a bit different than what you're used to if you've programmed in C. The `for` loop is specialized to iterate over the elements of any sequence, like a list. A `for` loop uses an iterator variable to reference each item as it steps through the sequence. This means it's very simple to write a loop!

Note that the syntax of the for-loop requires use of the `in`-keyword.

Pseudocode - For Loop

```

# In this example, "item" is a variable created specifically by the "for"
# loop to act as an iterator.
# The name "item" is not a keyword, and a different variable name may be
# used.
# Additionally, note that "for" and "in" are lowercase, and a colon is
# present at the end of the line.
for item in sequence:
    # All statements that should execute each iteration must be
    # indented after the "for" statement
    statement

```



Control Flow Loops

[Watch the Video](#)

Example	Output
<pre> listOfFruit = ['Apples', 'Oranges', 'Bananas'] for fruit in listOfFruit: print fruit </pre>	Apples Oranges Bananas

You don't need to manually create a sequence to repeat a task several times in a for loop. Instead, the built-in function `range()` function can generate a variable-size list of integers starting at zero. For example, calling `range(4)` will return the list [0, 1, 2, 3].

Example	Output
<pre> # Even though this example isn't using the value of "x", # the print statement will still be executed once for each item # in the list returned by range(). for x in range(4): print "this will print 4 times" </pre>	this will print 4 times this will print 4 times this will print 4 times this will print 4 times

While-Loop

A while loop will repeat a block of statements as long as a condition is true. This code will print out the contents of the items in the list.

Pseudocode - While Loop

```
# A while loop simply needs the keyword "while", the condition that
# determines when we should stop iterating, and a colon at the end of the line.
while condition:
    # All statements that should be repeated each iteration must be indented after the "while" statement
    # statement.
```

This code uses a function called **len()** , which is a built-in function that returns the length of a sequence.

Example	Output
<pre>listOfFruit = ['Apples', 'Oranges', 'Bananas'] x = 0 while x < len(listOfFruit): print listOfFruit[x] x = x + 1</pre>	Apples Oranges Bananas

The Break and Continue Statements

You can stop a loop from repeating in its tracks by using the **break** statement. This code will print out " Loop " exactly two times, and then print " Finished ".

Example	Output
<pre>for x in range(10): if x >= 2: break print "Loop" print "Finished"</pre>	Loop Loop Finished

You can use the **continue** statement to make a loop stop executing its current iteration and skip to the beginning of the next iteration. The following code will print out the numbers 0-9, skipping 4

Example	Output
<pre>for x in range(10): if x == 4: continue print x</pre>	0 1 2 3 5 6 7 8 9

Infinite Loops

It is incredibly easy to create an infinite loop when using a **while** statement. Depending where the infinite loop was created, it could cause you to lose your work in the Designer, or create a large amount of overhead on the Gateway.

Python - Infinite Loop Created by While Statement

```
x = 0
while x < 10:
    x += 1 # Forgetting to add a way to increment "x" will cause an infinite loop
    print x
```

In many cases, a **for** loop could be used instead of a **while**, but this is not always possible. When using **while**, the best way to avoid an infinite loop is to make sure you always have a way to exit the loop: a simple approach involves using a counter that can eventually trigger a **break** statement, or add the counter as an additional condition to the **while**.

Python - Preventing Infinite Loops Using the Break Keyword

```
####Example 1: using the break keyword
# The counter variable will be used as a guaranteed way out of the While.
counter = 0

# Normally, using True as a condition in a While would be a quick
# way to generate an infinite loop, but the counter helps prevent that.
while (True):

    # Increase the counter
    counter += 1

    # Check the value of the counter. If it's at the point where we can assume we're going to be looping
    indefinitely...
    if counter >= 1000:

        # Break out of the loop
        break
```

Python - Preventing Infinite Loops Using an Additional Condition

```
####Example 2: using an additional condition
# Again, the counter variable will be used as a guaranteed way out of the While.
counter = 0

# Instead of using nested logic, we can simply add counter's value as an additional condition with "and"
while (True and counter < 1000):

    # Increase the counter. Once counter >= 1000, the while loop will be forced to end.
    counter += 1
```

The Pass Keyword

When using conditional statements and loops, the **pass** keyword can be especially useful when writing a new script. When called, the **pass** keyword does nothing, which may seem useless. However it is great when you find yourself in a situation where you need a line of code to meet a syntax requirement, but don't want the code to do any additional work.

Python - Pass Keyword

```
myVar = system.tag.read(tagPath).value

if myVar == 0:
    firstFunction()
elif myVar == 1:
    secondFunction()
elif myVar == 2:
    # I haven't implemented the thirdFunction() yet, so I can use pass here as a placeholder
    pass.
```

Related Topics ...

- [Error Handling](#)
- [Getting Started with Scripting in Ignition](#)

Error Handling

What is Error Handling

The concept of error handling is recognizing when an error might occur in a block of code, and instead of throwing the error, handling it gracefully. This can involve giving the user a more distinct error message, letting the user know that their attempt to run the code failed, or even undoing something that your code set in motion so that it can be started again from the same starting point.

Error Handling in Python, Java, and Jython

Within Python, we can use the `try` and `except` blocks to handle errors. We would first use `try:` and write the code we would like to try indented underneath it. We then must have an `except:` with code that will run if there is an error.

On this page ...

- [What is Error Handling](#)
 - [Error Handling in Python, Java, and Jython](#)
 - [The Pass Keyword](#)
 - [Error Handling Examples](#)
 - [Exception-Specific Failover](#)
 - [Displaying Error Text](#)
 - [Determining the Error Object](#)

Pseudocode - Error Handling (Python)

```
# With try, we can attempt any amount of code
try:
    some code
    more code
    even more code

# If any of lines in the try block were to throw an error, then we move down and run the block under except.
# The except statement is NOT optional: you must define what your code should do in the event an exception
# occurs.
except:
    failover code
```

When running the code above, there is a specific order to the way it executes.

1. The try block is run, and the code within is executed line by line.
 - a. If an error occurs, the code in the try block will immediately stop and the except block will begin to execute.
 - b. If no error occurs after all code in the try block has been executed, the try block will be finished and the code in the except block will be skipped.
2. After either outcome, any code after the except will then execute.

Because of the way the try and except blocks work, it is very useful on situations that require user input, where something may have been incorrectly entered which would usually result in an error.

Much of the scripting in Ignition is done in Python; however, because Ignition is written in Java, many internal system calls may throw Java exceptions that Python won't catch without modifications.

Within Java, we can use the try and catch blocks to handle errors:

Pseudocode - Error Handling (Java)

```
// Use try to test our code
try {
    testing code goes here
}

// Handle any errors the code throws
catch(Exception e) {
    code to handle exceptions
}
```

However, since we are running Jython, we can use Python's `try` and `except` blocks for Java error handling by using Java's `Exception` class.

Pseudocode - Error Handling (Jython)

```

# Import the Exception class
import java.lang.Exception

# With try, we can attempt any amount of code
try:
    code you want to test
    more code
    code that will throw an exception

# If any of lines in the try block were to throw an error, then we move down and run the block under except.
# The except statement is NOT optional: you must define what your code should do in the event an exception occurs.
except java.lang.Exception, e:
    failover code

```

The Pass Keyword

The **pass** keyword is unique in that it does nothing except to fill in a spot where code is required. This is useful if we want to handle the exception so that it doesn't throw an error message, but we don't actually want to do anything with it. In this case, we can use pass to tell the script to continue.

Pseudocode - Pass Keyword

```

try:
    some code
    more code
    even more code

except:
    # The error will bring the code to the exception, and then the exception will simply do nothing.
    pass

```

Error Handling Examples

An easy way to demonstrate how error handling works is with a division example, since it is easy to cause an error by dividing by 0. Take the code below:

When we run it, we get a printed value of 50. There was no error in the division, and the try block finished successfully. However, if we were to change the value of x to 0, we can see that "An error occurred!" is printed.

Python - Error Handling Division

```

# We start with a value, which could represent some user input.
x = 2

# We use that value in our try block, and have a simple print statement if there is an error.
try:
    value = 100/x
    print value
except:
    print "An error occurred!"

```

Continuation of [DOC-1021](#)

The following code block is also with a division example, but in Java:

Java - Error Handling Division

```

// We start with a value, which could represent some user input.
int x = 2;

// We use that value in our try block, and have a simple print statement if there is an error.
try {
    int value = 100/x;
    System.out.println(value);
}

```

```

    catch (Exception e) {
        System.out.println("An error occurred!");
    }
}

```

Finally, the following code block is with the same division example, but in Jython:

Jython - Error Handling Division

```

# Import the Exception class
import java.lang.Exception

# We start with a value, which could represent some user input.
x = 2

# We use that value in our try block, and have a simple print statement if there is an error.
try:
    value = 100/x
    print value
except java.lang.Exception, e:
    print "An error occurred!"

```

Exception-Specific Failover

While each try block must be followed by at least one except block, there can be multiple except blocks to handle different types of errors. This is done by listing the error object after the except keyword and before the colon. Looking back at the example above, I know that my most common error is going to be a divide by zero error. So I can make an exception that is specific to divide by zero errors.

Python - Exception-Specific Failover

```

# We start with a value, which could represent some user input.
x = 0

# Use the user input in division in our try block.
try:
    value = 100/x
    print value

# If the exception that would occur is of type ZeroDivisionError, we can run a specific block of code
except ZeroDivisionError:
    print "Dividing by zero is not allowed. Please stop trying to divide by zero"
# We can then have a second exception without a specific error. This except acts as a catch-all;
# if the user caused an exception we didn't account for above, we can failover to the block of code below.
except:
    print "An error occurred!"

```

The except blocks in the code above cover all possible errors as represented in the table below. Now, we have a more tailored approach to how we handle errors while still catching all of them that may occur.

Inputs (x value)	Output
2	50
0	Dividing by zero is not allowed. Please stop trying to divide by zero
'a'	An error occurred!

Each try block can have many except blocks, and each except block can also name multiple types of errors as shown below. However, an error that happens within a try block will only ever trigger a single except block.

Pseudocode - Try Block with Except Blocks

```

try:
    some code
except (ZeroDivisionError, RuntimeError, TypeError):
    failover code

```

Displaying Error Text

Sometimes you want to get the actual text from an error in addition to protecting your script. There's an easy way to fetch that information that's built into Python. When you are inside an except section of code, `sys.exc_info()` gives you access to the error text as a list of values. You can use this to print out your message, display it on the screen, send it to the database, or anything else.

This example should be put on a button, and will write the error text to a Label component that is a sibling to the button. This is useful in Perspective to get error messages out of views and event actions.

```

try:
    # cause an error
    x=[1,2]
    val = x[5]
except:
    # push the error text to a sibling label
    self.getSibling("Label").props.text = sys.exc_info()

```

You can also use the Ignition loggers to push these error messages out to the Gateway console. You can find these in the Gateway Webpage under the Status section, on the **Logs** page.

```

# This code would log an error to the gateway
try:
    # cause an error
    100/0
except:
    # push the error text to the logger
    logger = system.util.getLogger("myLogger")
    # convert the sys.exc_info() to a string and log it
    logger.info(str(sys.exc_info()))

```

Determining the Error Object

To determine the name of the error object that will be thrown from certain errors, we can take a look at the error to figure that out. We already mentioned that dividing by zero gives a `ZeroDivisionError`, but what about when we divide by a string? If I divide 100 by the letter a without error handling, this is the error I get:

```

Traceback (most recent call last):
File "<buffer>", line 3, in <module>
TypeError: unsupported operand type(s) for /: 'int' and 'str'

```

The last line of that error gives the name of the error object "`TypeError`" followed by the cause of that error. This makes sense, because the string '`a`' is the wrong type to be using in division. However, not all errors are so simple and may require a bit more to find the name of the error. For a list of python error names check out this page in the python docs: <https://docs.python.org/2.7/library/exceptions.html#Exception>

Additionally, some exceptions may be returned by Java. In these cases, Oracle's documentation on the `Exception` class is more useful: <https://docs.oracle.com/javase/8/docs/api/java/lang/Exception.html>

A list of common exceptions are listed below.

Exception	Description	Exception Demonstration
<code>ArrayIndexOutOfBoundsException</code>	This typically occurs when a line of code attempts to access an index in a collection, but the index specified doesn't exist. This exception is unique to datasets in Ignition. Lists, and other built-in Python objects will return the <code>IndexError</code> below.	<pre> myDataset = system.dataset.toDataSet (["colName"],[[0]]) # This will fail because the dataset only has a single row, </pre>

		<pre># so trying to read the value at row index # 5 means we're trying to # read something that doesn't exist. print myDataset.getValueAt(0,5)</pre>
AttributeError	<p>Typically seen when a script tries to access a property that doesn't exist.</p> <p>Example: a script tries to get the value of a property on a component, but the property name is misspelled.</p>	<pre>myVar = 10 # integers do not natively have a name # variable, so this will fail with an # AttributeError: # there isn't an 'attribute' on an integer # by the name of 'name' print myVar.name</pre>
IndexError	Similar to <code>ArrayIndexOutOfBoundsException</code> above, but occurs when Python specific object, such as a list.	<pre>myList = [1,2,3] # There isn't an element in the list at # index 4, so this will fail. print myList[4]</pre>
NameError	Occurs when the local or global name is not found. Typically happens when referencing a variable that hasn't been assigned a value.	<pre># We haven't given a value to the variable myValue, so it will fail. print "The value of the variable is: " , myValue</pre>
TypeError	A <code>TypeError</code> occurs when a function or similar operation is applied to another object of an incorrect type.	<pre>myList = [1,2,3] # The first argument for pop() expects an # integer, so passing a string value is # inappropriate. # Passing a string to pop() will return a # TypeError. print myList.pop("0")</pre>
ValueError	A <code>ValueError</code> is returned when the value of an argument is inappropriate. Typically the exact issue is returned in the returned exception.	<pre>myVar = "Hello" # Strings can be passed to the int() # function, but the value needs to be # something that # can be coerced into an integer, like "0". # Because "Hello" can't be easily # converted, the line below will fail. print int(myVar)</pre>

Related Topics ...

- [Conditions and Loops](#)

Built-In Functions

Python Built-In Functions

Functions are code that can be called repeatedly from other places. Functions can have parameters passed into them and may return a resulting value. Some functions, like `len()`, are built-in. Some functions, like `system.gui.messageBox()`, are part of the [scripting libraries](#) provided by Ignition. Some functions, like `math.sqrt()`, are provided by the Python Standard Library.

Functions are invoked by using their name followed by an argument list surrounded in parentheses. If there are no arguments, you still need an open and close parenthesis.

This section details several useful Built-in Functions, along with some simple examples. See the [official docs](#) for more information.

On this page ...

- [Python Built-In Functions](#)
 - [Type Casting Functions](#)
 - [Checking an Object's Type](#)
 - [Generating a Range of Values](#)
 - [Rounding Numbers](#)

Type Casting Functions

Python has many functions to convert between data types. Some common type casting functions are listed below

Function	Notes	Example	Output
<code>bool()</code>	When casting a numeric value to a boolean, a zero value is false, while all non-zero numbers are True When casting a String or Unicode value to a boolean, an empty string is False, any other string is True.	<pre># Results in False print bool("") # Results in True print bool("Test")</pre>	False True
<code>int()</code> and <code>long()</code>	When casting a float, rounding will not occur automatically to the decimal value. Instead, the round() function should be called. When casting a String or Unicode value, the string literal needs to be a valid integer or long: decimal values contained in the string will result in a <code>ValueError</code> Integers have at least 32 bits of precision, while Longs have unlimited precision.	<pre># Float to Integer print int(123.8) # Float to Long print long(321.8) # String to Integer print int("400") # ValueError: the value # is not base 10 print int("400.5")</pre>	123 321 400 ValueError
<code>float()</code>	When casting a string literal as a float, non-numeric characters in the string will result in an exception, except for a decimal point (".").	<pre># Integer to Float print float(123) # String to Float print float("400.5")</pre>	123.0 400.5
<code>str()</code> and <code>unicode()</code>	Most objects can be cast as a string representation of some sort, including sequences.	<pre>print "First line:" + str(80) # Even sequences can be # cast as a string, # making for easy concatenation myList = [1,2,3] print str(myList)</pre>	80 [1, 2, 3]

Checking an Object's Type

Checking the data type of an object can easily be done with both the `type()` and `isinstance()` functions.

Function	Description	Example	Output
<code>type(object)</code>	When passed a single parameter, this function returns the type of the object.	<pre>var = 10 print type(var) print type(str(var))</pre>	<pre>type 'int'> <type 'str'></pre>
<code>isinstance(object, classinfo)</code>	Returns True if the <code>object</code> is an instance or subclass of <code>classinfo</code> , otherwise, returns false. If checking for a string or unicode type, a <code>classinfo</code> of "basestring", which is the base class for both strings and unicode types, would return True.	<pre>var = 10 print isinstance(var,int) strVar = "string" print isinstance(strVar,basestring)</pre>	<pre>True True</pre>

Python - Type Validation: `type` vs `isinstance`

```
## type() Example
# This example attempts to validate the type of a variable. As written, this will evaluate as True, and thus
the print statement would execute.
var = "My String"
if type(var) == type(""):
    print "Variable 'var' is a string"

## isinstance() Example
# The isinstance() function can offer the same functionality as above.
var = "My String"
if isinstance(var, str): # Note the lack of quotation marks around the classinfo parameter. We want to
reference the class str, not the string "str".
    print "Variable 'var' is a string"
```

Generating a Range of Values

In some cases, it is useful to generate a range of integers for iteration. Python's `range()` function will return a list of integers.

Function	Description	Example	Output
<code>range([start,] stop[, step])</code>	<p>Returns a list of progressively greater integers.</p> <p><code>start</code> - Integer value denoting the initial value in the list. If omitted, defaults to 0. This parameter is inclusive.</p> <p><code>stop</code> - Integer value, determines when to cease generating integers. This parameter is exclusive.</p> <p><code>step</code> - Integer value to increment each new integer by. If omitted, step defaults to 1.</p> <p>If <code>step</code> is positive, integers will be generated as long as (<code>start + i * step < stop</code>) is true.</p> <p>If <code>step</code> is negative, integers will be generated as long as (<code>start + i * step > stop</code>) is true.</p>	<pre>print range(5) print range(1, 5) print range(1, 10, 3) print range(15, 0, -3)</pre>	<pre>[0, 1, 2, 3, 4] [1, 2, 3, 4] [1, 4, 7] [15, 12, 9, 6, 3]</pre>

Assume we need to read from five separate Tags with a nearly identical Tag path in a single script:

Pseudocode - Tag Path

```
[Provider]Folder/Sub_Folder_1/Tag  
[Provider]Folder/Sub_Folder_2/Tag  
[Provider]Folder/Sub_Folder_3/Tag  
[Provider]Folder/Sub_Folder_4/Tag  
[Provider]Folder/Sub_Folder_5/Tag
```

Instead of manually typing each path, we could use `range()` in a for loop that would write the paths automatically.

Python - Range in a For Loop

```
# Initialize an empty list that will ultimately hold all the Tag paths.  
tagPaths = []  
  
# Use range to repeatedly append 5 tag paths to the tagPaths list: starting with a value of 1, and ending  
# with a value of 5.  
for num in range(1,6):  
    # Use String Formatting to create a Tag path with the iterator's (num) value.  
    tagPaths.append("[Provider]Folder/Sub_Folder_%i/Tag" % num)  
  
# Now that tagPaths contains all our tag paths, we can use the list to interact with the tag, such as by  
# reading their values simultaneously.  
tagValues = system.tag.readAll(tagPaths).value
```

Rounding Numbers

You can round numbers inside Python with a few simple functions.

Function	Description	Example	Output
<code>round(number [, digits])</code>	When passed a single parameter, this function returns a rounded integer value of a number. If the decimal is greater than or equal to .5, the it rounds up, less than .5 rounds down. If the optional digits argument is used, then it rounds to that many decimal places.	<pre>var = 10.236981 print round (var) print round (var,3)</pre>	10 10.237
<code>math.floor(number)</code>	Returns a truncated integer from the number given. The largest integer value less than or equal to <i>number</i> . Note that the example needs to import that math library before being able to call <code>floor()</code> .	<pre>import math var = 100.938 print math. floor(var)</pre>	100.0
<code>math.ceil(number)</code>	Returns the ceiling integer from the number given. The smallest integer value greater than or equal to <i>number</i> . Note that the example needs to import that math library before being able to call <code>ceil()</code> .	<pre>import math var = 100.138 print math. ceil(var)</pre>	101.0

Python - Simple Casting

```
stringVar = "40"

# Without using int(), this line would cause an exception. However int() is able to cast
# the type of stringVar's value to an integer.
print 20 + int(stringVar)

# Type casting is also useful in cases where a string parameter is required, but a numerical value
# should be given, such as the message parameter in system.gui.messageBox().
intVar = 60

# Note that this could also be accomplished with String Formatting instead of using str().
system.gui.messageBox(str(intVar))
```

Related Topics ...

- [User Defined Functions](#)
- [Libraries](#)

Libraries

The System Library

Ignition comes with a group of system functions, called the System Library. Using a system function is simple. For example, the following code will access the value of a Tag.

Pseudocode - Reading a Tag Value

```
value = system.tag.read("tagPath").value
```

The [scripting appendix](#) is full of built-in functions such as this.

On this page ...

- [The System Library](#)
- [Python Libraries](#)
- [Python Standard Library](#)
 - [Importing 3rd Party Libraries](#)
- [Accessing Java](#)
 - [Subclassing Java](#)

Python Libraries

Python Libraries are packages of extra functions that expand the functionality of the code and can be imported into a script. We do this by using the `import` keyword:

Pseudocode - Import a Library

```
# This pseudo code will import a library, and then call a function of that
library.
import myLibrary

myLibrary.specialFunction()
```

The `import` keyword imports that entire library and allows you to use all of the functions inside of it by calling them off of the imported library. You can also import a piece of a library:

Pseudocode - Import a Function of a Library

```
# This pseudo code will import a function from a library, and then call
that function.
from myLibrary import specialFunction

specialFunction()
```

Note, that since we are directly importing in the function, we can directly call it instead of having to call it off of the library.

Python Standard Library

Python has an extensive standard library that provides a host of new functionality to the scripting language. The python documentation goes over all of the libraries in its standard library as well as how to use them here: <https://docs.python.org/2/library/index.html>

Let's take a look at an example of using a common library:

Python - Accessing Files in a Python Standard Library

```
# The csv library provides an easy way to read csv files, regardless of how they are formatted.
import csv

# We first grab our filepath, and feed it into the open function, which opens the file.
filepath = "C:\\\\test.csv"
csvFile = open(filepath, 'r')

# We then pass our opened csv file object into the csv.reader function, which will read the file.
# This can be looped through in a for loop to print every row of the csv.
```



System Library

[Watch the Video](#)

```
reader = csv.reader(csvFile)
for row in reader:
    print row
```

Importing 3rd Party Libraries

In addition to the standard libraries, 3rd party libraries can also be imported into Ignition's scripting environment. A Python Library or individual module file will consist of a python file (.py) that contains the code that implements the functions of the library. You can often find python libraries built by other users on the web, or can even create your own. These files can then be placed into a folder within your Ignition server.

- **Windows folder:** C:\Program Files\Inductive Automation\Ignition\user-lib\pylib
- **Linux folder:** /usr/local/bin/ignition/user-lib/pylib
- **Mac OS X folder:** /usr/local/ignition/user-lib/pylib

Once the python file is in that folder, you can then import the library into a script just like any of the standard libraries.



Ignition uses Python version 2.7. This means that any imported libraries must be compatible with Python 2.7.

Accessing Java

Scripting in Ignition executes in the java based implementation of Python called Jython. (See [Python or Jython?](#)). While this doesn't have any great effect on the Python language itself, one of the great side benefits is that your Python code can seamlessly interact with Java code as if it were Python code. This means that your Python code has access to the entire Java standard library.

To use Java classes, you simply import them as if they were Python modules. For example, the following code will print out all of the files in the user's home directory. This code uses the Java classes `java.lang.System` and `java.io.File` to look up the user's home directory and to list the files. Notice that we can even use the Python-style for loop to iterate over a Java sequence.

Python - Accessing Java

```
# Importing the appropriate java libraries.
from java.lang import System
from java.io import File

# Used to look up the files in the users home directory.
homePath = System.getProperty("user.home")
homeDir = File(homePath)

# Loops through the list of files and prints them.
for filename in homeDir.list():
    print filename
```

You can find the reference documentation for the Java standard class library (also known as, the "JavaDocs") at: <http://docs.oracle.com/javase/8/docs/api/>

Subclassing Java

You can also create Python classes that implement Java interfaces. You do need some understanding of Java and object-oriented programming concepts, which are outside the scope of this manual. To create a Python class that implements a Java interface, you simply use the interface as a superclass for your Python class. For example, we could augment the example above to use the overload `java.io.File.list(FilenameFilter)`. To do this, we'll need to create a `FilenameFilter`, which is an interface in Java that defines a single function:

```
boolean accept(File dir, String name)
```

To implement this interface, we create a Python class that has `java.io.FilenameFilter` as its superclass, and implements that Java-style function in a Python-esque way.

Python - Implementing Java Interfaces

```
# Importing the appropriate java libraries.
from java.lang import System
from java.io import *
```

```
# This sets up an extension filter that can check the file extension. Txt is the default.
class ExtensionFilter(FilenameFilter):
    def __init__(self, extension=".txt"):
        self.extension=extension.lower()

    def accept(self, directory, name):
        # make sure that the filename ends in the right extension
        return name.lower().endswith(self.extension)

# Used to look up the files in the users home directory.
homePath = System.getProperty("user.home")
homeDir = File(homePath)

# Prints out all .txt files. Txt is provided if nothing is specified.
for filename in homeDir.list(ExtensionFilter()):
    print filename

# Prints out all .pdf files.
for filename in homeDir.list(ExtensionFilter(".pdf")):
    print filename
```

Related Topics ...

- [Built-in Functions](#)

User Defined Functions

Functions

A function is code that can be called repeatedly from other places. Functions can have parameters passed into them, and may return a resulting value. Some functions, like `len`, are built-in. Some functions, like `system.gui.messageBox()`, are part of the [scripting libraries](#) provided by Ignition. Some functions, like `math.sqrt()`, are provided by the [Python standard library](#). However, functions can also be defined in a script that can be used later on in the script. In these user defined functions, you give the function a name and some code that will run when the function is called. Then later on in the script, you can call the function by its name and it will run the code specified in the function. This is useful, because it allows you to run a segment of code many times without having to repeat it within the script.

Functions are invoked by using their name followed by an argument list surrounded in parentheses. If there are no arguments, you still need an open and close parenthesis.

Defining Functions

Functions are defined using the `def` keyword. A function needs a name and a list of the arguments that it can be passed. For example, this code defines a function that prints "Hello World!".

Python - Defining a Function

```
# First we define our function.  
def printHW():  
    print "Hello World!"  
  
# We can then call our function.  
printHW()
```

On this page ...

- [Functions](#)
 - [Defining Functions](#)
 - [Functions Arguments](#)
 - [Keyword Arguments](#)
 - [Functions Are Objects](#)
 - [Where Can Functions Be Defined](#)
 - [Function Scope](#)



INDUCTIVE
UNIVERSITY

Functions

[Watch the Video](#)

Functions Arguments

When a function accepts arguments, the names of those arguments become variables in the function's namespace. Whatever value was passed to the function when it was invoked becomes the value of those variables. Arguments can have default values, which makes them optional. If an argument is omitted, then its default value will be used. The following code defines a function called `cap`, which will check if a number is within an upper and lower limit. The default lower limit is 0, and the default upper limit is 100.

Python - Defining Arguments

```
# We first define our function. Notice that we have 3 different arguments.  
# x, min, and max. The min and the max are set to equal 0 and 100 respectively.  
def cap(x, min=0, max=100):  
  
    # Check if x is less than the min, return the min if true.  
    if x < min:  
        return min  
  
    # Check if x is greater than the max, return the max if true.  
    elif x > max:  
        return max  
  
    # Return the value if it is within the bounds.  
    else:  
        return x  
  
# We can then see the outcome by running our function with a few different parameters.  
# This will print out a 40, since it is within the bounds.  
print cap(40)  
  
# This will print out "0", since it is less than the min of 0.  
print cap(-1)  
  
# This will print out "100", since it is greater than the max of 100.  
print cap(150)
```

```
# This will print out "150", because it uses a max of 200 instead of the default 100.  
print cap(150, 0, 200)
```

Keyword Arguments

In Ignition, some complicated script functions are designed to take keyword arguments instead of normal parameters. In the description for those functions, you may see the following info box in this User Manual:

 This function accepts keyword arguments.

Arguments can also be specified by keyword instead of by position. In the example above, the only way someone would know that the 200 in the last call to `cap` specified the `max` is by its position. This can lead to hard-to-read function invocations for functions with lots of optional arguments. You can use keyword-style invocation to improve readability. The following code is equivalent to the last line above, using 200 for the `max` and the default for the `min`.

Python - Using Keyword Arguments

```
print cap(150, max=200)
```

Because we used a keyword to specify that 200 was the `max`, we were able to omit the `min` argument altogether, using its default. However, using a keyword argument before a non-keyword or positional argument is not allowed.

Python - Non-keyword Argument

```
# This would fail, because the function isn't sure what 150 is being used for.  
print cap(max=200, 150)
```

Functions Are Objects

Perhaps one of the most foreign concepts for new Python users is that in Python, functions are first-class objects. This means that functions can be passed around to other functions (this concept is similar to the idea of function pointers in C or C++).

Suppose we wanted a general way to filter a list. Maybe sometimes we want the odd entries, while other times we want even ones. We can define a function called `extract` that takes a list and another function, and returns only entries that "pass" through the other function.

Python - Functions Passed to Other Functions

```
# We define a function that checks if the value passed in is odd.  
def isOdd(num):  
    return num % 2 == 1  
  
# We define a function that checks if the value passed in is even.  
def isEven(num):  
    return num % 2 == 0  
  
# We define a function that inserts our list into the appropriate function and returns valid values.  
def extract(filterFunction, list):  
    newList = []  
    for entry in list:  
        if filterFunction(entry):  
            newList.append(entry)  
    return newList  
  
# Prints out [0, 2, 4, 6, 8]  
# Notice that isEven is not invoked, but passed to the filter function.  
print extract(isEven, range(10))
```

Where Can Functions Be Defined

User Defined Functions can be defined anywhere that a script is used. As stated before, they are useful to run segments of code multiple times without having to repeat it. They are also used extensively in [project scripts](#) where multiple functions can be defined in a single script module. Finally, some special Ignition System functions like [system.gui.createPopupMenu](#) or the [runScript](#) Expression function use functions as arguments.

Function Scope

The concept of scope is very important in all programming, and Python is no exception. Scope defines what names are directly accessible without any qualifiers. Another way to put this is that the scope determines what variables are defined. In Python, a variable is defined at the time that it is assigned. What scope it belongs to is also defined by where the assignment occurs.

Pseudocode - Defining a Function for Scope

```
# On this line, there is no variable 'x' in scope.  
doSomeWork()  
  
# Now 'x' is defined in our scope, because we've assigned a value to it  
x = 5  
  
# This will work because x is in scope.  
print x
```

When you define a function, that function gets its own scope. Variables that are assigned within that function body will not be available outside of the function.

Python - Variables Defined within a Function Not Available Outside Scope

```
# x is local to myFunction() because this is where it is defined.  
def myFunction():  
    x = 15  
    print x  
  
# This will fail, because x is not available in the outer scope  
y = x + 10
```

Related Topics ...

- [Built-in Functions](#)
- [Libraries](#)

Scripting in Ignition

Where Is Scripting Used?

Python is used in many places in Ignition. Each location has its own events that trigger your scripts to run, and add functionality to your projects in different ways. The most apparent place is in [event handlers](#) on components and other objects in Vision Clients and Perspective Sessions.

Script Scope

One important thing to keep in mind before scripting in Ignition, is to understand the concept of scope.

Within Ignition, there are different scopes:

- **Gateway Scope** - The script runs on the gateway. Scripts running in this scope cannot interact with components in the other two scopes.
- **Perspective Session Scope** - The script runs as a part of a Perspective Session. Note that scripts in Perspective execute on the gateway, not in the browser, but this scope is still distinct from the Gateway Scope.
- **the Vision Client Scope** - The script runs inside of an instance of a Vision Client.

Where a script was written determines which scope it executes in. For example, Tags are in the Gateway Scope, so [Tag Event Scripts](#) execute in the Gateway Scope.

This means that the script will not be able to access any client level resources such as windows or components that you may have open in the Client. Additionally, some of the system functions like `system.gui.errorBox` only work in the "Client Scope," so you will not be able to use them in the script on the Tag.

System Functions, Hints, and Autocomplete

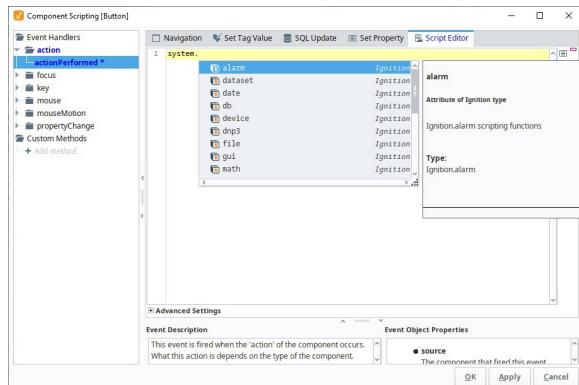
Ignition comes with a group of system functions, which are built-in functions that interact with Ignition features.

Python - Simple Script Using a System Function

```
value = system.tag.readBlocking(["tagPath"]).value
```

A complete list of these functions (with their definitions) is available from the autocomplete popup. Wherever you can add a script, type `system.` and then press **Ctrl+Space** to get a list of all the functions available. If you keep typing, the list will be automatically narrowed down for you. Additionally, the [System Functions](#) page in the appendix contains complete documentation for the built-in system functions.

Note: The autocomplete popup always shows all system functions scoped to the current script. If a system function does not appear in the list, that means the function is not available in the current scope, or has been deprecated.



The following feature is new in Ignition version **8.1.18**
[Click here](#) to check out the other new features

Starting in 8.1.18, the autocomplete popup is enabled by default and will appear after typing "."

On this page ...

- [Where Is Scripting Used?](#)
 - [Script Scope](#)
 - [System Functions, Hints, and Autocomplete](#)
- [Components](#)
- [Client, Gateway, and Session Event Scripts](#)
- [Project Scripts](#)
- [Tag Scripts](#)
- [Reporting](#)
- [Alarming](#)
- [Sequential Function Charts](#)



INDUCTIVE
UNIVERSITY

Scripting in Ignition

[Watch the Video](#)

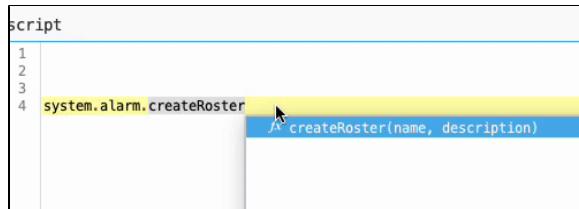


INDUCTIVE
UNIVERSITY

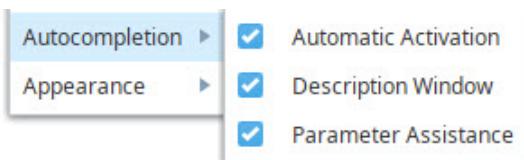
System Library

[Watch the Video](#)

The new editor also offers parameter completion assistance; if you auto-complete a method with multiple required parameters, you'll automatically enter a "parameter assistance" mode, where you can tab through the parameters and enter them one at a time:

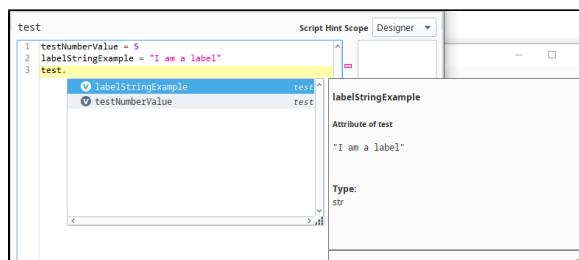


To disable these features, right-click anywhere within the Script Editor window and deselect **Automatic Activation** and/or **Parameter Assistance**.



The following feature is new in Ignition version **8.1.19**
[Click here](#) to check out the other new features

Autocomplete hints are now also displayed for code other than Ignition's system functions. If Ignition detects a function or project script, a popup will automatically appear, from which you can select which function or project script you are trying to reference. You can also bring up this popup by pressing "ctrl-space".



The following feature is new in Ignition version **8.1.32**
[Click here](#) to check out the other new features

Autocomplete hints will extract method parameters, return information, and limited type awareness for project script functions and class docstrings written in [Google Python Style Guide's](#) docstring format. An example of the expected format is as follows:

Code Snippet

```
def setMode(mode):
    """Changes the mode of the running system

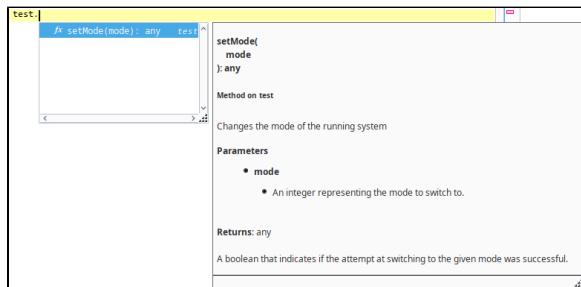
    Args:
        mode: An integer representing the mode to switch to.

    Returns:
        A boolean that indicates if the attempt at switching to the given
        mode was successful.

    Raises:
        Error: If communications are down, an exception will be thrown.

    """
    # Function code goes here
```

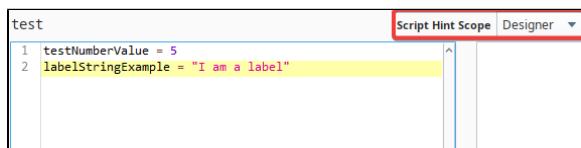
The code above will display the description, parameter, and return information when autocomplete hints render the docstring:



Script Hint Scope

The following feature is new in Ignition version **8.1.19**
[Click here](#) to check out the other new features

You can choose the scope of your scripting hints and suggestions come from by selecting the dropdown menu in your Project Library's scripts:

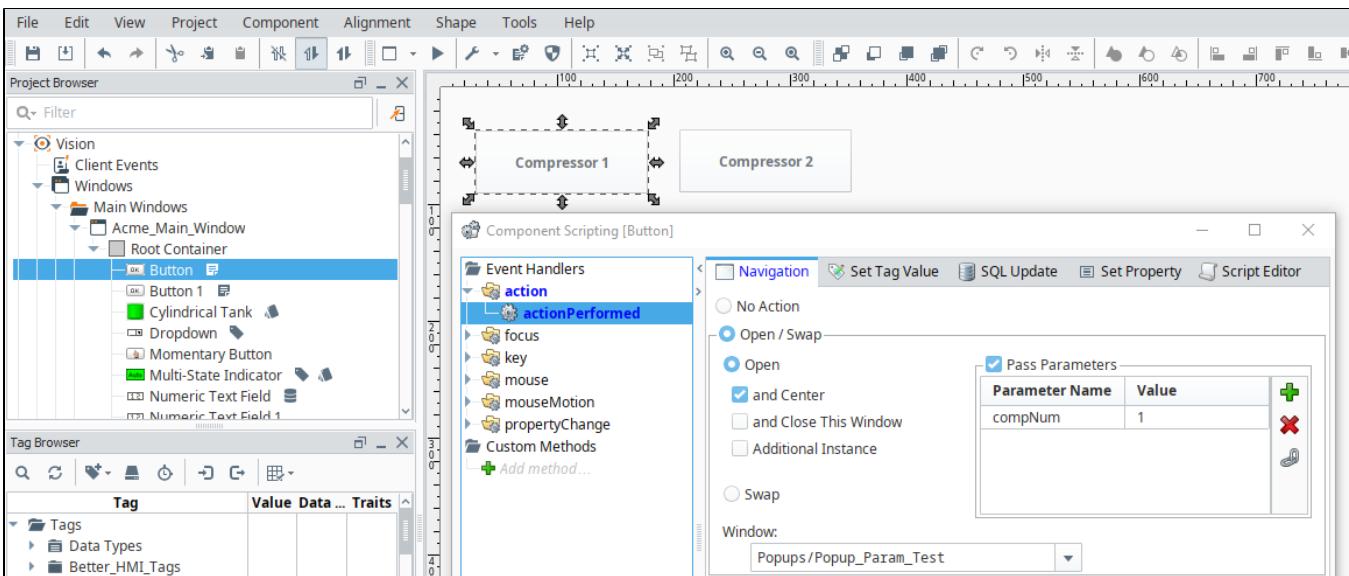


The table below lists possible values you can choose from in the dropdown menu:

Value	Description
None	The project library will not use the Designer nor the Gateway to populate scripting hints.
Designer	The project library will use the Designer scope to populate scripting hints.
Gateway	The project library will use the Gateway scope to populate scripting hints.
All	The project library will use both the Designer and Gateway scope to populate scripting hints.

Components

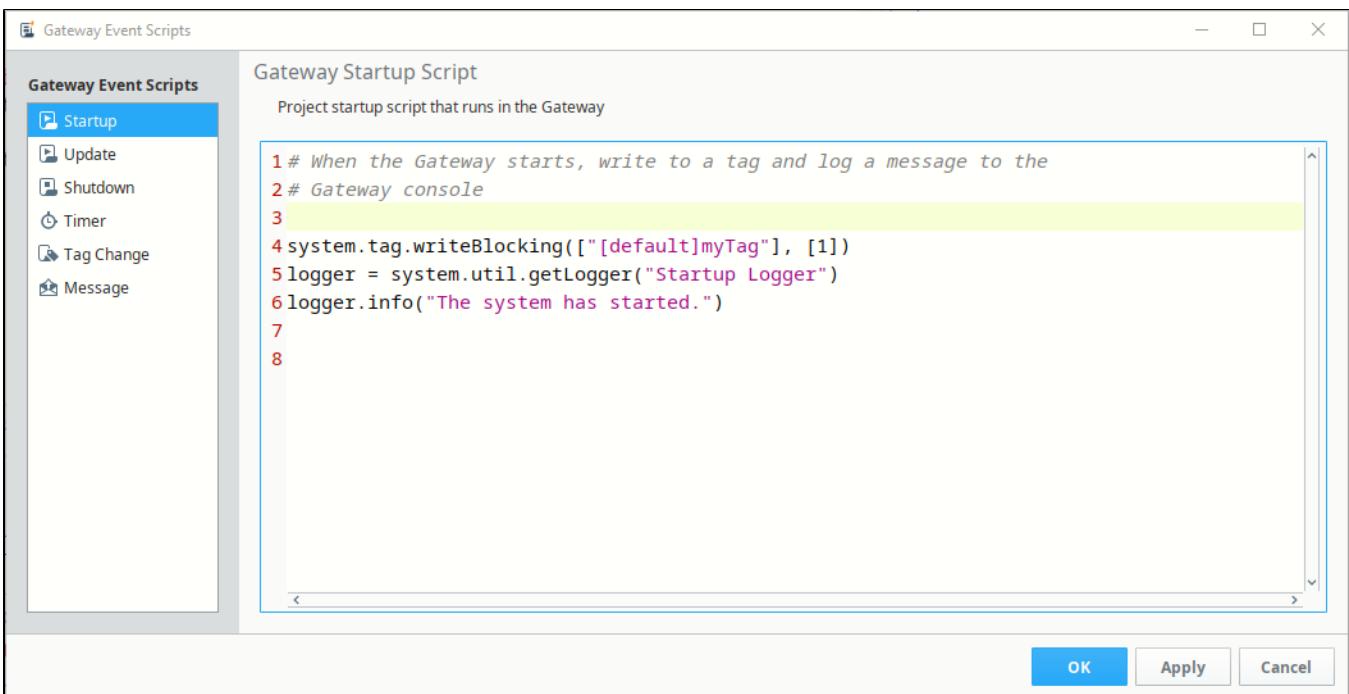
Both Perspective and Vision offer component based scripting triggers, providing a means to execute a script under a number of different situations, such as a user interacting with a component or a component property value changing. For more information on how both module handle component based scripts, take a look at the [Scripting in Perspective](#) and [Scripting in Vision](#) sections.



Client, Gateway, and Session Event Scripts

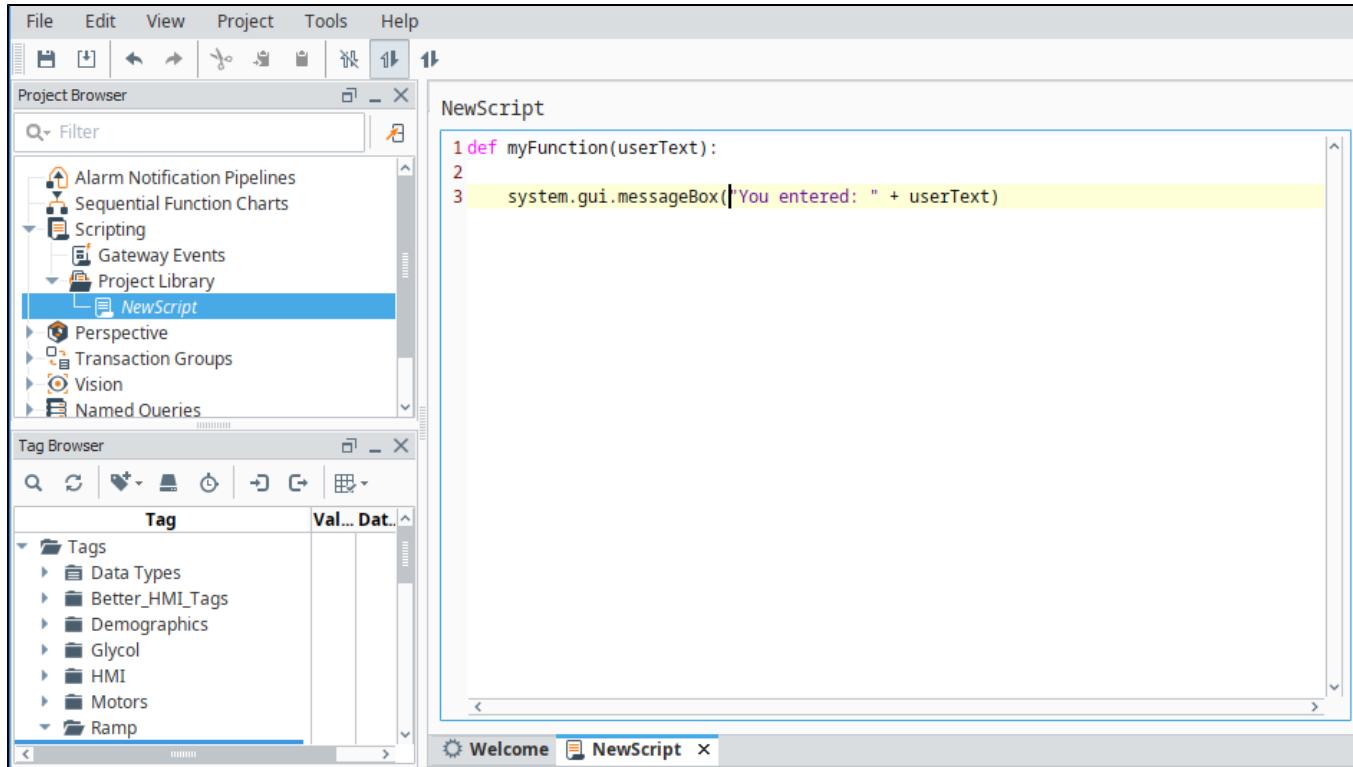
Scripts can be set to activate on specific events that occur during runtime. For example, you can trigger a script to run when a vision client starts, or on certain time intervals.

More information on these events can be found on the [Client Event Scripts](#), [Gateway Event Scripts](#), and [Perspective Session Event Scripts](#) pages.



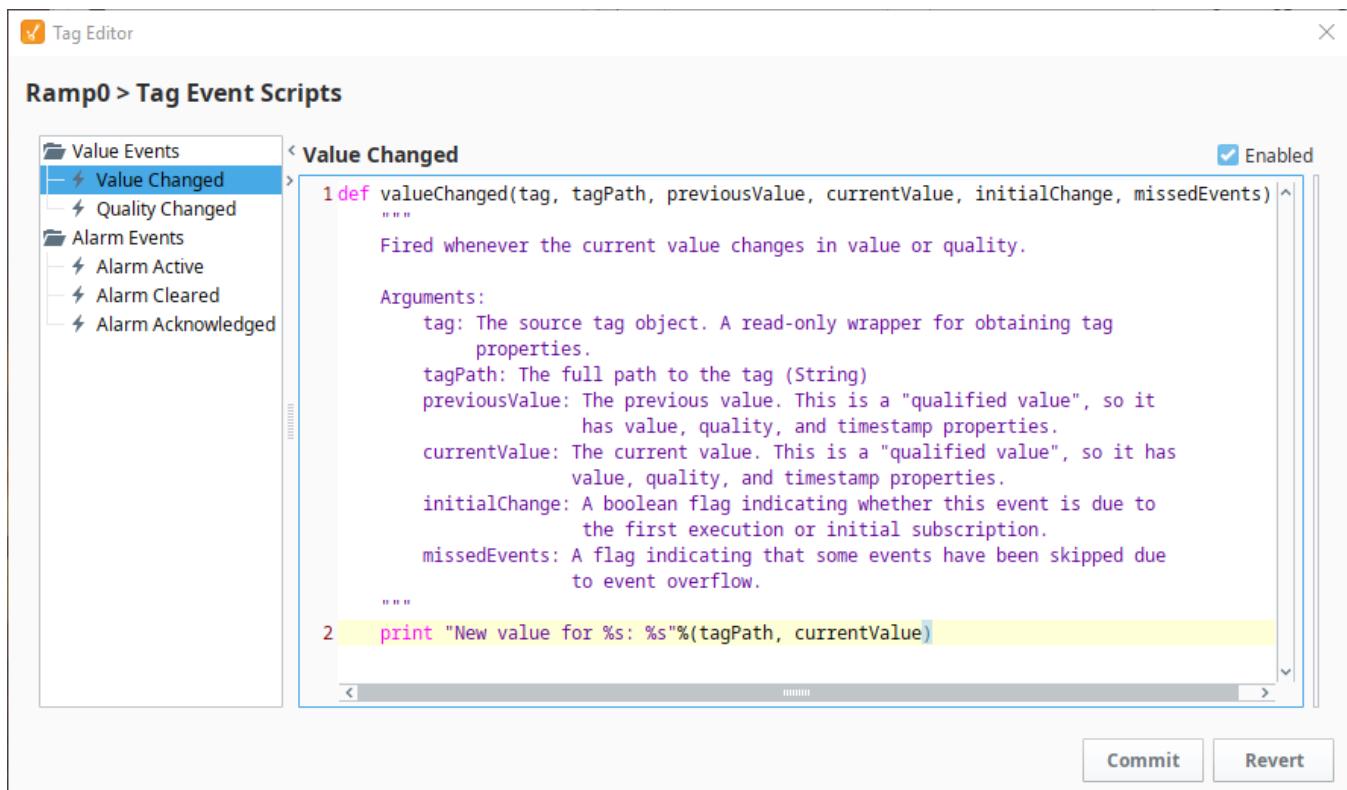
Project Scripts

You can create your own reusable blocks of code in the [Project Library](#). Once configured, these functions can be called from anywhere in a project, just like our `system.*` functions.



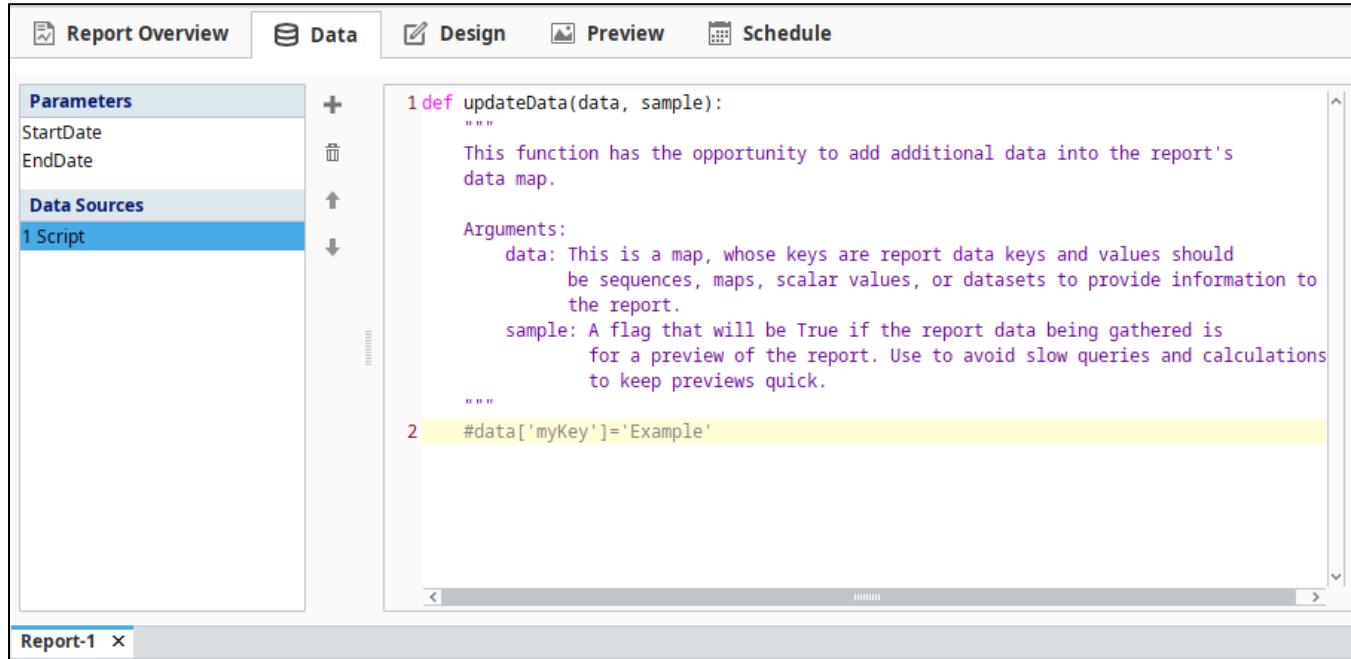
Tag Scripts

Once Enabled, these scripts are fired whenever a [Tag value changes](#) or an alarm event happens. You can use them for additional diagnostics, to set additional Tags, or to react to an alarm event. Because these events are on Tags, they are Gateway Scoped.



Reporting

Reporting uses scripting in many different ways to help increase the effectiveness of the report. Scripting in Reports is used to [create and modify data sources](#), [manipulate charts](#), and set up a script as a [scheduled report action](#).



The screenshot shows the 'Report Overview' interface with several tabs at the top: Report Overview, Data, Design, Preview, and Schedule. The 'Data' tab is selected. On the left, there are sections for 'Parameters' (StartDate, EndDate) and 'Data Sources' (1 Script). The 'Script' item under 'Data Sources' is selected and highlighted in blue. The main pane displays a Python script:

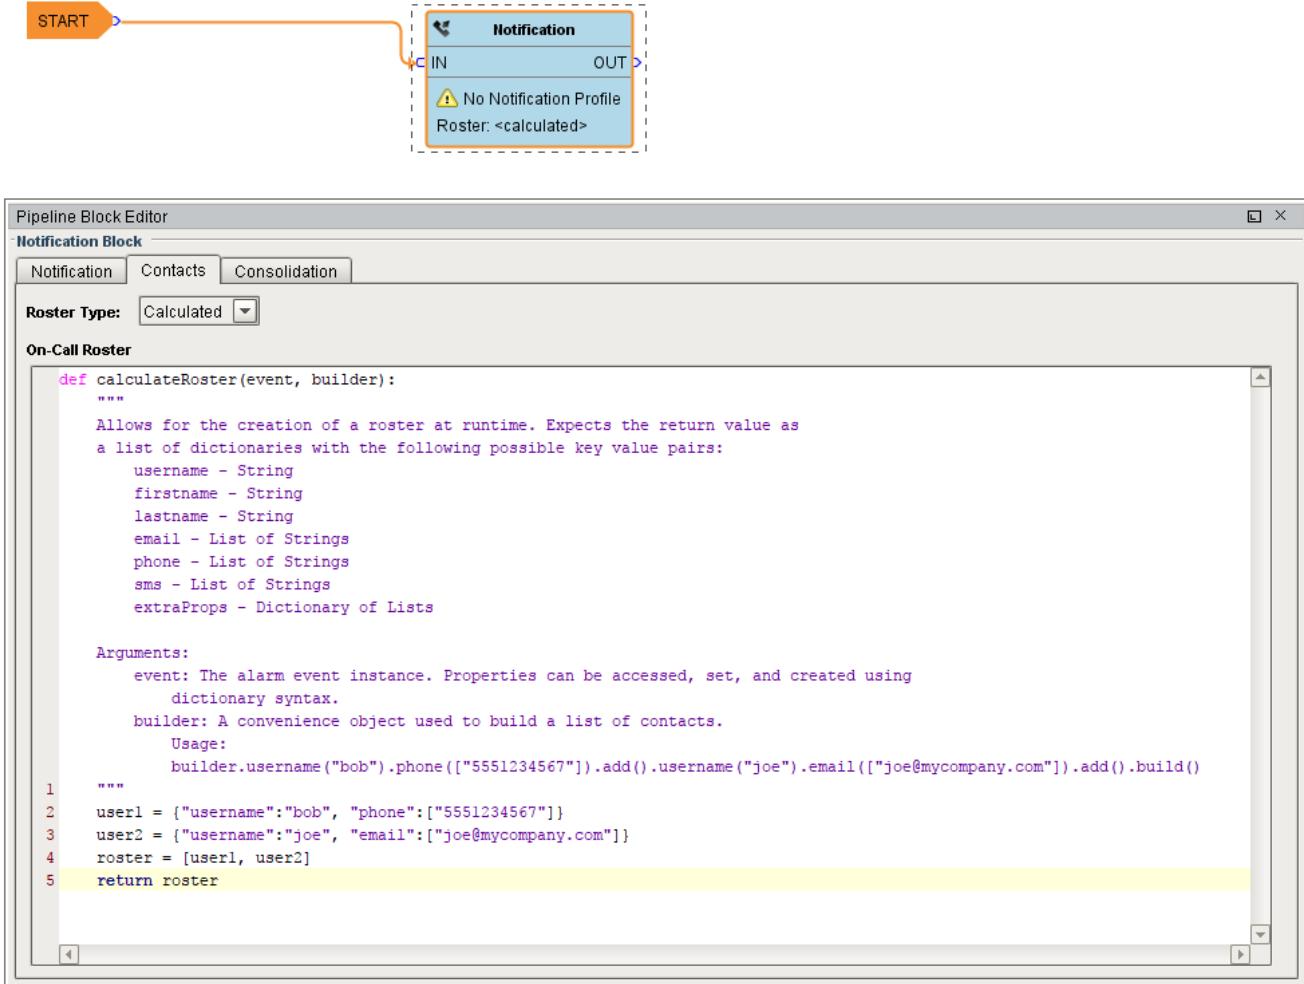
```
1 def updateData(data, sample):
    """
    This function has the opportunity to add additional data into the report's
    data map.

    Arguments:
        data: This is a map, whose keys are report data keys and values should
              be sequences, maps, scalar values, or datasets to provide information to
              the report.
        sample: A flag that will be True if the report data being gathered is
                for a preview of the report. Use to avoid slow queries and calculations
                to keep previews quick.

    """
2     #data['myKey']='Example'
```

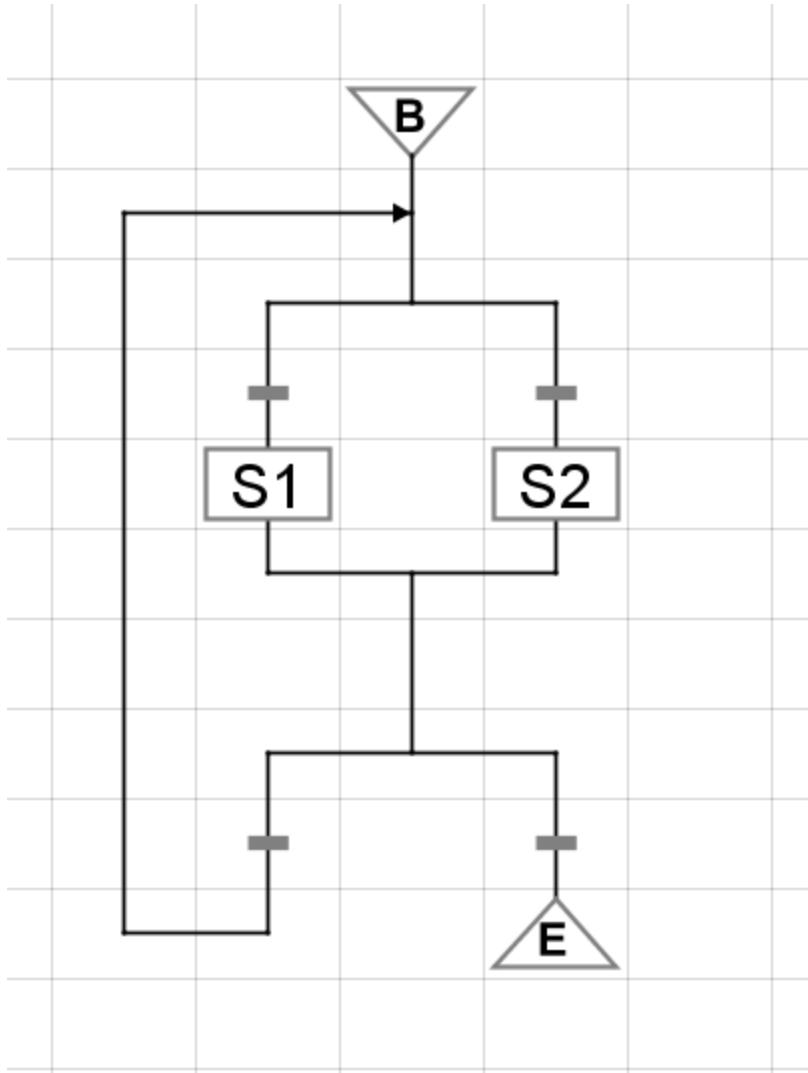
Alarming

The Alarm Notification system can also use scripting to great effect. A [script block](#) allows a script to be run within the pipeline, allowing data to be manipulated as the alarm event travels through the pipeline. Additionally, scripting can be used to generate a [custom roster](#) of users at runtime, giving full customization to who gets notified by the alarm event.



Sequential Function Charts

Sequential Function Charts (SFCs) are a flowchart of blocks that run scripts. They are executed in a specific sequential order along with some logic to potentially loop or call other charts. The scripts here can interact with the Gateway, and provide greater control when each step needs to complete before the next one can begin in multi-step processes.



Related Topics ...

- Scripting Data Source
- Sequential Function Charts
- Notification Block
- Tag Event Scripts
- Client Event Scripts
- Gateway Event Scripts
- Perspective Component Methods
- Perspective Session Event Scripts
- Project Library

In This Section ...

Getting Started with Scripting in Ignition

Overview

The best way to get started with scripting in Ignition is to write some very simple scripts and print them within Ignition. Let's take the phrase, "Hello World," and print it within Ignition. To show how flexible Ignition is, there are actually a few ways to easily do this. One way is to use the Scripting Console, and another way is to use a Button component. Both examples are shown below.

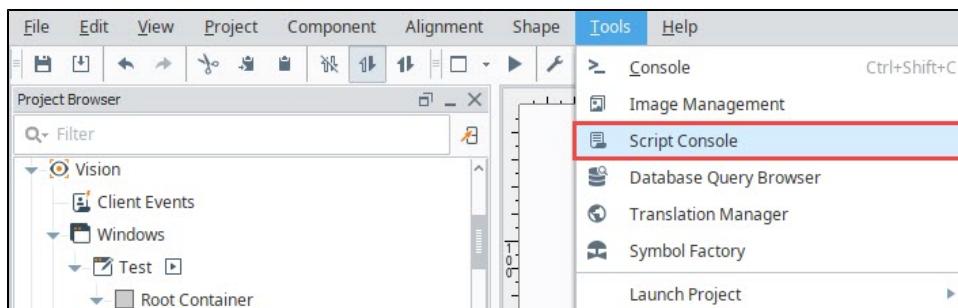
Example - "Hello World" Using the Script Console

When testing or writing a new script, the [Scripting Console](#) is very useful since you can immediately get some feedback on the results of the script. When learning Python, it is a great place to start since you don't have to create a window or component before you begin writing your code.

On this page ...

- [Overview](#)
- [Example - "Hello World" Using the Script Console](#)
- [Example - "Hello World" on a Button Component](#)
- [Example - Using a Message Box](#)

1. In the **Designer**, open the Script Console from the Menu Bar: **Tools > Script Console**.

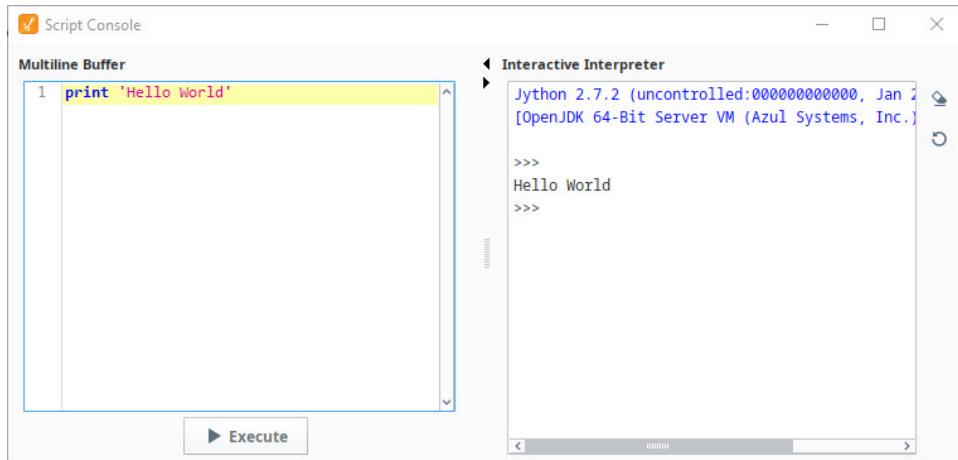


2. The Script Console will appear. Type the following code, or simply copy and paste it into the text area in the Multiline Buffer on the left side of the Script Console:

Python - Simple Print

```
print 'Hello World'
```

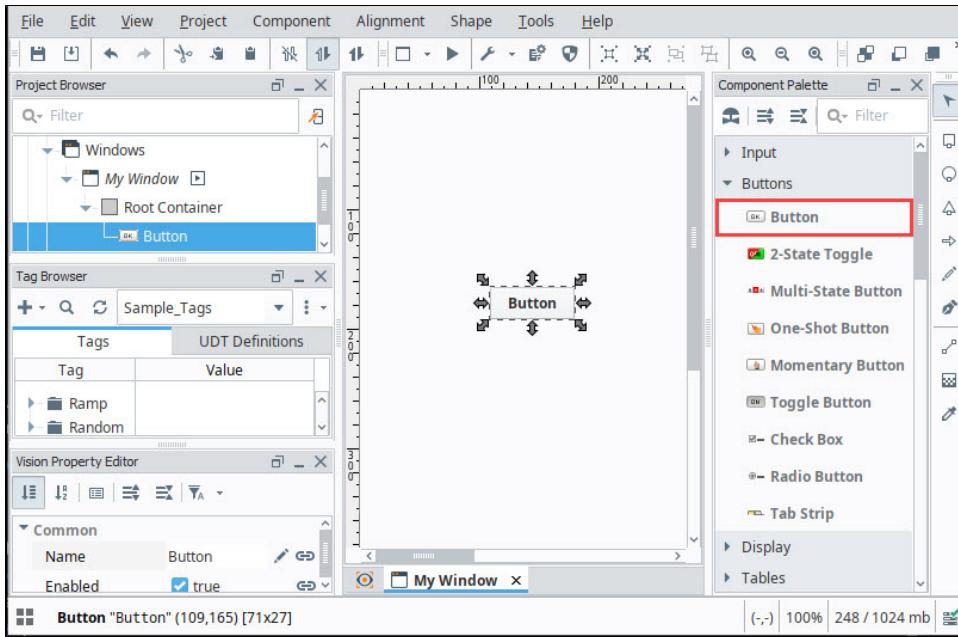
3. Click the **Execute** button at the bottom of the Script Console. You should see the message "Hello World" appear in the Interactive Interpreter on the right side of the Script Console.



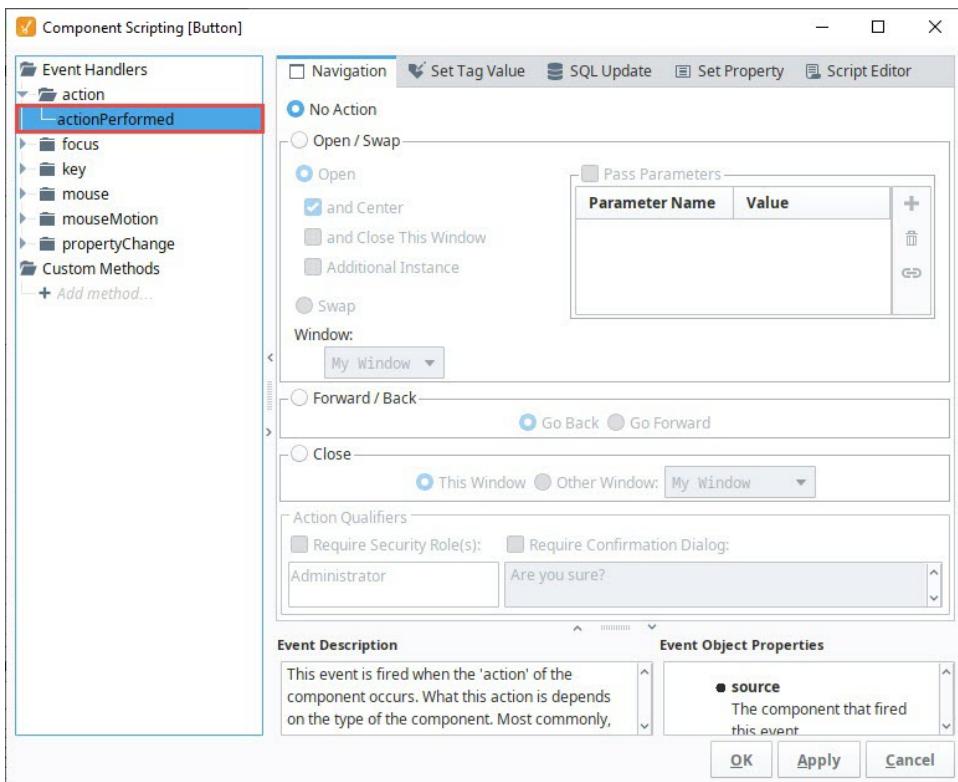
Example - "Hello World" on a Button Component

Scripts are commonly located on components in a window. In this example, let's add a script on a Button component, and print out "Hello World" when the Button is pressed.

1. In the **Designer**, create a new [Main Window](#). Give it a meaningful name if you like. (We won't need to reference the name of the window in any of these examples).
2. Drag a standard **Button** component to your window.

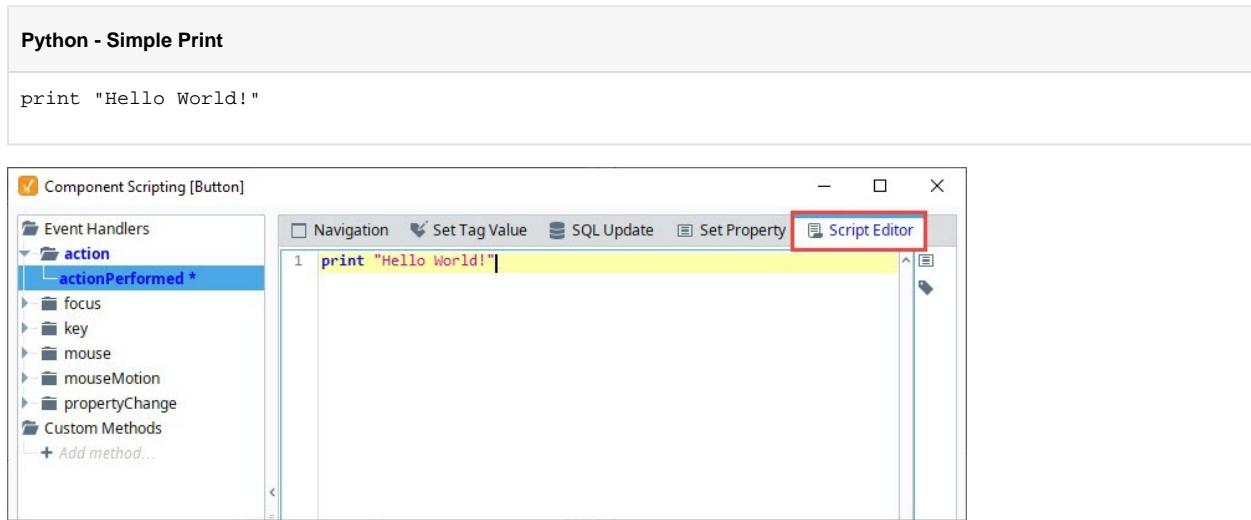


3. Let's add a script that triggers when the Button is pressed. With the Button component selected, right click on **Scripting** to open the **Component Scripting** window.
4. Under Event Handlers, select the **action > actionPerformed** option, and click the **Script Editor** tab. Make sure the **actionPerformed** event is highlighted. If an event is not selected, all of the remaining features in this window will be disabled, so we will not be able to write a script. Additionally, if the wrong event is selected, then our script will not trigger when we expect it to.



5. In the Script Editor, you will see a large editable Text Area. This is where we will type our script.

6. Let's generate a message that shows "Hello World!". Use the following code to get started. Make sure the word "print" lines up exactly to the left edge. Indention in Python means something, so we need to avoid starting our lines with a space or tab unless we're denoting a block of code under something like an **if-statement** or **function definition**. Click **OK** to close the **Component Scripting** window.

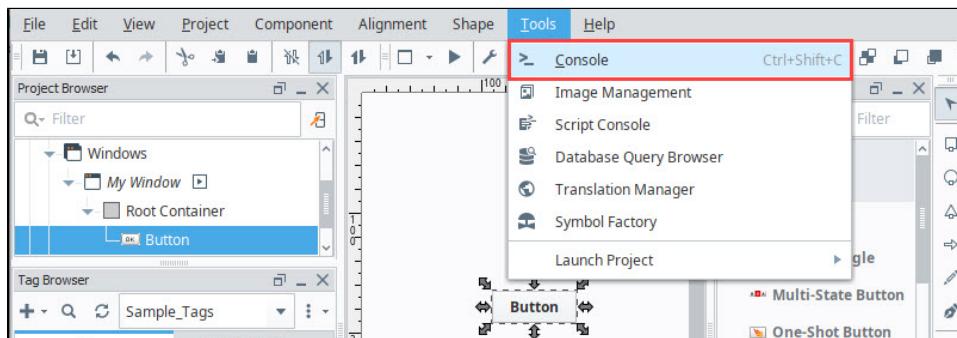


Notice the **actionPerformed** event is blue and bold. This means there is a script on this event. This is useful to know in situations where a component has scripts on multiple events.

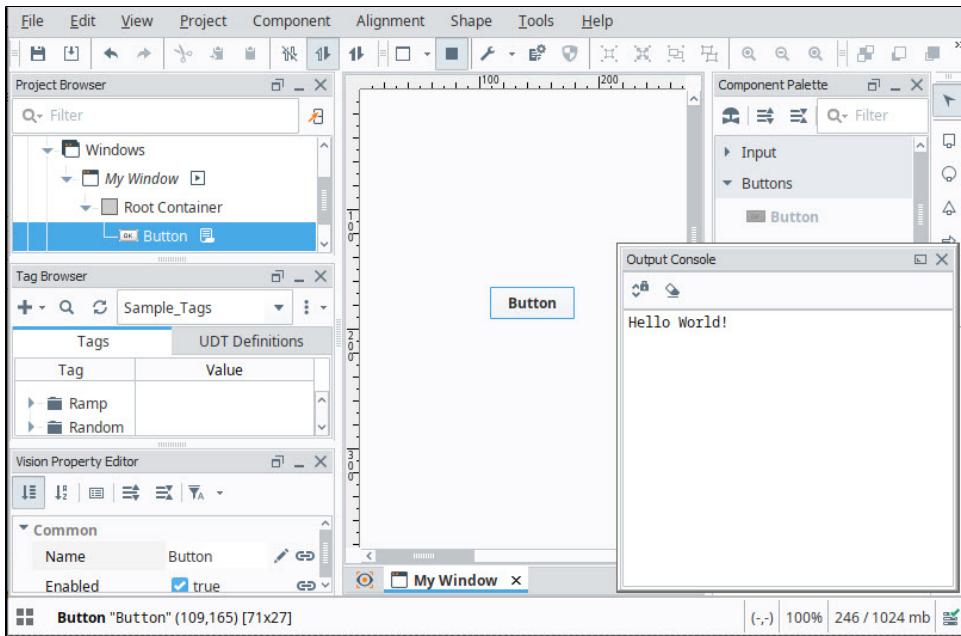
Additionally, an asterisk character (*) is next to the event. This means you have not applied/saved the changes to the script. The asterisk will disappear when you press either **OK** or **Apply**, and reappear whenever you make a change to the script. If you see this, then it means you may want to save any changes you made.

The **Script Editor** tab also has a blue color, denoting where the script is. An event will only ever have a script located on a single tab at any time. If a new tab is selected and configured, it will wipe out the work on the prior tab, for example, writing a script on **Script Editor** and then configuring the **Navigation** tab will erase the script on the **Script Editor** tab.

7. Now we can test the script. Place the **Designer** into Preview Mode, and press the Button. If everything is working as intended, then it should appear as if nothing happened. This is because we used `print` in our script, which always outputs to a console, as opposed to popping up at the user. This means we need to open the Designer's **Console** to see the results of our script. At the Designer's menu bar, select **Tools > Console**. This will make the console appear.



8. There may be a large amount of text in the Output Console. The Designer logs many different types of activities and events here, including polling events from components on other windows that are currently open in the Designer. However, the most recent events should be towards the bottom. As a quick tip, you can also click the eraser icon to clear out the console, and then press your **Button** again to generate a new entry like we did in this example.



We can now see where print statements go when called from a component, but this isn't too useful, as we don't want our users to open the console to see messages. Print statements are very useful when troubleshooting a problematic script, or even when testing a new script. Keep this in mind as you start to delve into scripting more.

Example - Using a Message Box

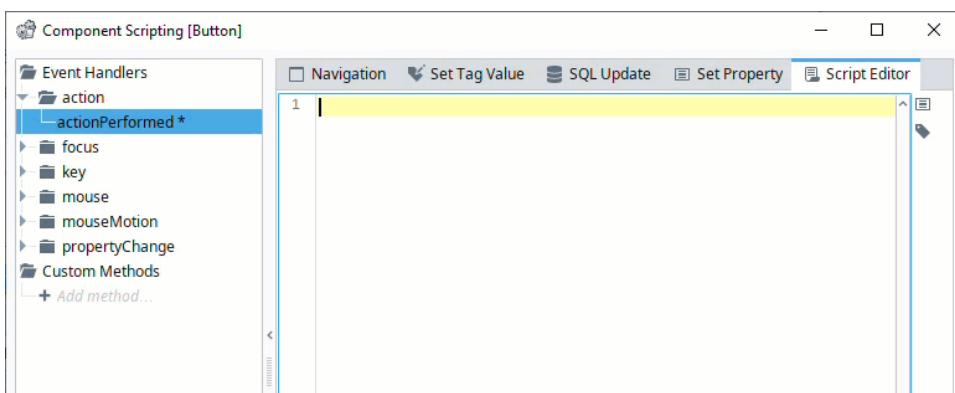
Now that we've seen how to print to the Output Console, let's make our message appear when we call it. This time, we will modify the script on the Button component to bring up a window that the user will see.

1. Open the **Component Scripting** window again, and find the script on the Button.
2. We will use one of Ignition's built-in functions, called `system.gui.messageBox`, to display the message. This will make a message box appear, which is a modal window that we can pass a string to. Remove your old code on the Button, as we will be replacing it. Start by typing the following:

Pseudocode - Built-in Functions Begin with `system`.

```
system.
```

3. With the Text Cursor just to the right of the ".", hold The **Ctrl/Cmd** key and press the **Spacebar**. This will make the Autocompletion popup appear. This lists all of the available system functions. Start by clicking on `gui`, and then `messageBox`. Note, that you can still type while the popup is open to filter the results in the list.



The following feature is new in Ignition version **8.1.18**
[Click here](#) to check out the other new features

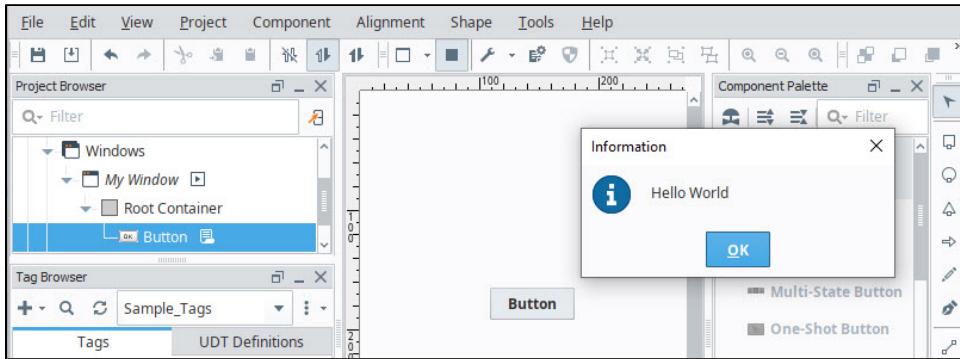
Starting in 8.1.18, the Autocompletion popup appears by default after typing "."

4. Once you've selected **messageBox**, click on it, and it should start some of the code for you automatically.
5. Complete the code by placing some parentheses and a message as a string. Alternatively, you can copy the example below.

Python - Using a System Function to Write to a Message Box

```
system.gui.messageBox('Hello World')
```

6. Place the **Designer** into **Preview Mode**, and click your Button. You will see a Message Box appear displaying the message we passed to the function.



Related Topics ...

- [Python Scripting](#)
- [System Functions](#)
- [Component Events](#)

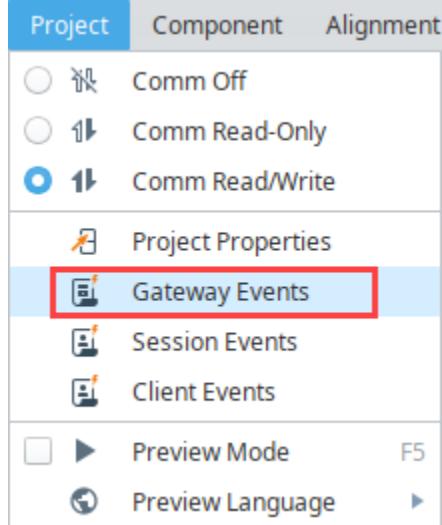
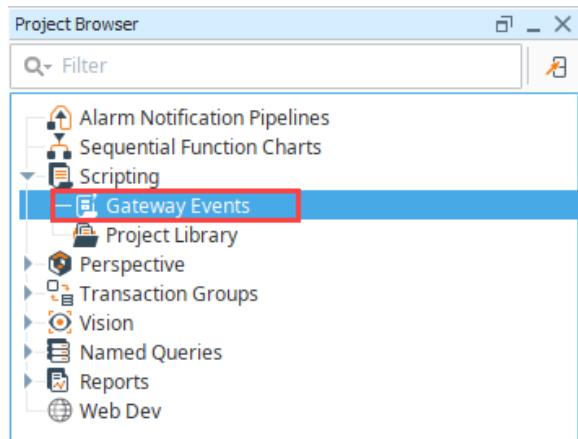
Gateway Event Scripts

Gateway Event Scripts Overview

Gateway Event Scripts are scripts that run directly on the Gateway. They are useful because they always run, regardless if any sessions or clients are open. They also provide a guaranteed way to make sure only a single execution of a particular script occurs at a time, as opposed to placing a script in a window, as there could be multiple instances of the window open at a given point of time.

Note that even though Gateway Event Scripts run on the Gateway, they're still considered a project resource. Project backups will include any Gateway Event Scripts.

The Gateway Event scripting workspace is located in the Scripting menu of the Designer or in the Project Browser under **Scripting > Gateway Events**.



Other Event Scripts

The content on this page will focus primarily on Gateway Event Scripts. However, there is some overlap with Client Event Scripts, as they have similar events. More information can be found on the [Client Event Scripts](#) and [Perspective Session Event Scripts](#) pages.

Note: System functions are available for both Client Event Scripts and Gateway Event Scripts, but some system functions are specific to either one or the other. When you're writing event scripts, it's important to remember the scope of where you're writing the script: Client or Gateway. You can check [System Functions](#) in the Appendix to see list of all system functions, their descriptions, and what scope they run in.

On this page ...

- [Gateway Event Scripts Overview](#)
 - [Other Event Scripts](#)
- [Startup Script](#)
 - [Gateway Startup Behavior](#)
- [Update Script](#)
- [Shutdown Script](#)
 - [Gateway Shutdown Behavior](#)
- [Timer Scripts](#)
 - [Timer Script Settings](#)
- [Tag Change Scripts](#)
 - [Tag Change Script Interface](#)
 - [Tag List](#)
 - [Tag Change Objects](#)
- [Message Scripts](#)
 - [Gateway Message Handler Settings](#)
 - [The Payload](#)
 - [Calling Message Handlers](#)
- [Gateway Scheduled Scripts](#)
 - [Settings Tab](#)
 - [Script Tab](#)
- [Troubleshooting Gateway Scripts](#)

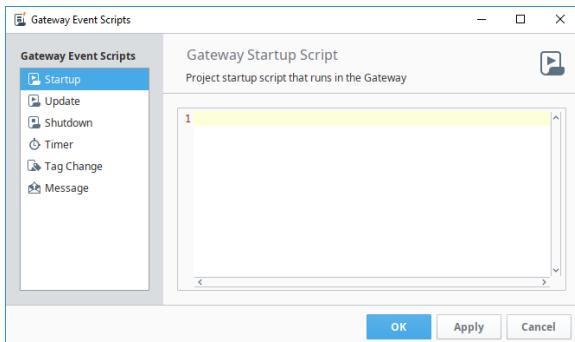


Gateway vs Client Event Scripts

[Watch the Video](#)

Startup Script

The Startup Script event runs at the startup of the Gateway. Additionally, if the project is restarted in someway, such as by making a change to a Gateway Event Script and saving, then the Startup Script will be called. This means that while editing scripts frequently in the Designer, the startup and shutdown events may happen frequently.



**INDUCTIVE
UNIVERSIT**

Startup Scripts

[Watch the Video](#)

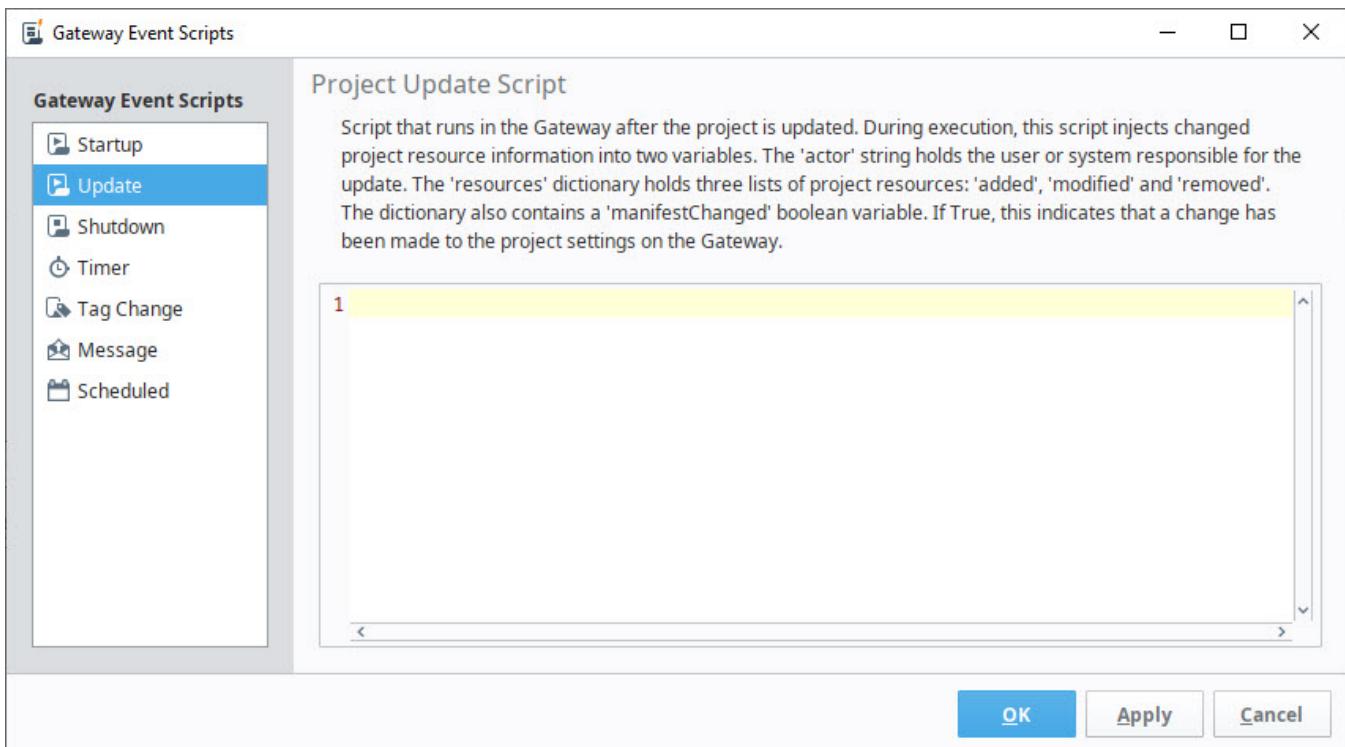
Gateway Startup Behavior

There is a specific order to when the various startup scripts are run. When troubleshooting your Gateway startup times, consider the following:

1. **Gateway starts** - The Gateway will start as an OS service, and start the context. No startup scripts can run before this is complete.
2. **Projects are started** - This includes all of the Gateway scoped items in the projects such as Transaction Groups, SFCs, etc. This does not refer to launching clients, and no clients can be automatically launched at this time. All **Gateway Startup Scripts** are run at this time for each project. Note: if you copied a project, always check for Gateway scoped events such as these. You generally don't want a Gateway Startup Script to run twice because it is in two projects.

Update Script

The Update Script event runs after a project is saved or updated on the Gateway. This enables you to insert a script that will run every time a project is saved.



The following feature is new in Ignition version 8.1.14
[Click here](#) to check out the other new features

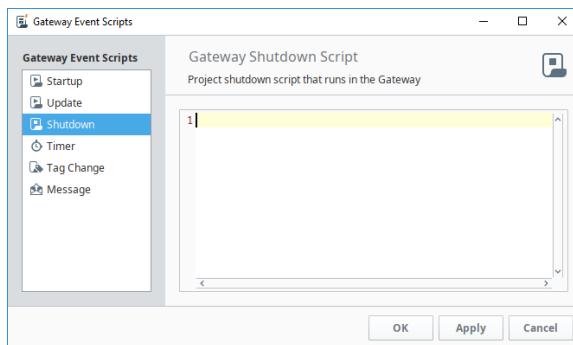
This script now reports which resources were modified during a project save. The following parameters have been added:

Parameter	Description	Type										
actor	The user or system responsible for the project update.	String										
resources	A dictionary that holds the following keys: <table border="1"> <thead> <tr> <th>Key</th> <th>Description</th> </tr> </thead> <tbody> <tr> <td>added</td> <td>List of dictionaries containing information about resources that were added to the project.</td> </tr> <tr> <td>removed</td> <td>List of dictionaries containing information about resources that were removed from the project.</td> </tr> <tr> <td>modified</td> <td>List of dictionaries containing information about resources that were modified.</td> </tr> <tr> <td>manifestChanged</td> <td>A boolean variable indicating whether or not a change has been made to the project settings on the Gateway.</td> </tr> </tbody> </table>	Key	Description	added	List of dictionaries containing information about resources that were added to the project.	removed	List of dictionaries containing information about resources that were removed from the project.	modified	List of dictionaries containing information about resources that were modified.	manifestChanged	A boolean variable indicating whether or not a change has been made to the project settings on the Gateway.	Dictionary
Key	Description											
added	List of dictionaries containing information about resources that were added to the project.											
removed	List of dictionaries containing information about resources that were removed from the project.											
modified	List of dictionaries containing information about resources that were modified.											
manifestChanged	A boolean variable indicating whether or not a change has been made to the project settings on the Gateway.											

Shutdown Script

The Shutdown Script event runs at the shutdown of the project, which means it can be used as a way to trigger a script when the Gateway has to be restarted. It allows you to run a piece of code as the shutdown is occurring. After the script completes, the shutdown will finish.

Note that the Shutdown Script event only gets called if the Gateway is requested to shut down: if the computer power is lost abruptly (power outage, hard restart, etc.) this shutdown script will not run.





**INDUCTIVE
UNIVERSIT**

Shutdown Scripts

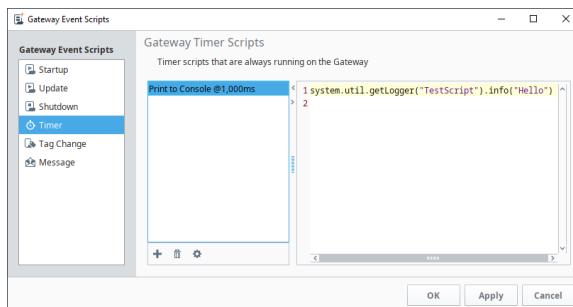
[Watch the Video](#)

Gateway Shutdown Behavior

Similar to how a [Gateway Event Scripts#Startup Script](#) behaves, "Shutdown" in the context of this event means "project shutdown", so shutting down the Gateway would trigger this event, as well as disabling the project containing a script on this event. Additionally, making a change to a Gateway Event Script in the Designer, and then saving the project will cause the project to restart, which means this event can get called by simply making changes in the Designer and saving.

Timer Scripts

The Timer Scripts execute periodically on a timer at a fixed delay or rate. This allows you to set up a sort of heartbeat that can run on the Gateway. This is the ideal event to use if you need the Gateway to periodically perform some scripting task.





**INDUCTIVE
UNIVERSIT**

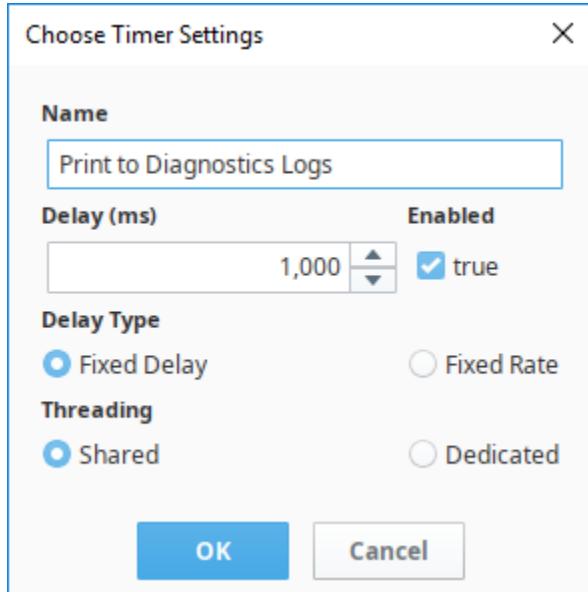
Timer Scripts

[Watch the Video](#)

Since multiple Timer Scripts can be added, there are separate buttons that allow you to manage each Timer Script.

-  **Add Timer Script** - Adds a new Timer Script.
-  **Remove Timer Script** - Will delete the selected Timer Script.
-  **Modify Timer Script** - Will modify the settings for the selected Timer Script.

Timer Script Settings



Below is an overview of the settings for a Timer Script.

- **Name:** The name of the Timer script. Names must be unique per project, so two timer scripts in the same project cannot have the same name.
- **Delay:** The delay period in milliseconds. The meaning of this setting is dependent on the **Delay Type** setting.
- **Enabled:** Allows you disable the Timer Script when set to false.
- **Delay Type:** Determines how the **Delay** setting is utilized.
 - A **Fixed Delay** timer script (the default) waits for the given **Delay** between each script invocation. This means that the script's rate will actually be the delay plus the amount of time it takes to execute the script. This is the safest option since it prevents a script from mistakenly running continuously because it takes longer to execute the script than the delay.
 - **Fixed Rate** scripts attempt to run the script at a fixed rate relative to the first execution. If the script takes too long, or there is too much background process, this may not be possible. See the documentation for [java.util.Timer.scheduleAtFixedRate\(\)](#) for more details.
- **Threading:** Determines which thread this script should run in. In other words, this setting allows you to specify if you want this timer script to share execution resources or not. The rule of thumb here is that quick-running tasks should run in the shared thread, and long-running tasks should get their own dedicated thread.
 - The **Shared** setting means that all timer scripts will share a thread. This is usually desirable, as it prevents creating lots of unnecessary threads: threads have some overhead, so a small amount of resources are used per thread. However, if your script takes a long time to run, it will block other timer tasks on the shared thread.
 - The **Dedicated** setting will create a separate thread specifically for the timer script to use. This setting is desirable when your scripts executions must be as consistent as possible, as other timer scripts can't slowdown or otherwise impact the execution of a script in a separate thread.

Tag Change Scripts

The Tag Change Script event allows you to specify any number of Tags, and trigger a script when one of them change. Since these execute based on a Tag changing value, Tag Change Scripts are ideal when you need a script to run based on some signal from a PLC.

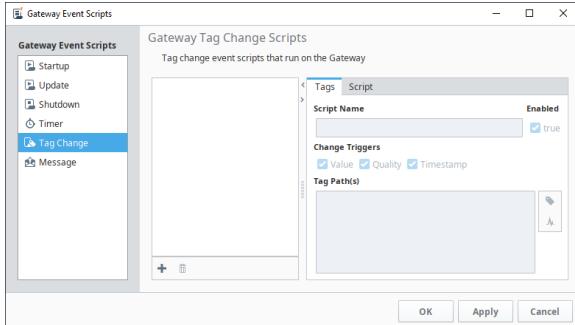
Having the Tag Change Scripts run in the Gateway Scope means that the scripts are active as long as the Gateway is running. Thus, you do not need a client or session to be open for a Gateway Tag Change Script to execute. When executing, each Tag Change Script runs in a separate thread. This prevents long running scripts from blocking the execution of other Tag Change Scripts.



Tag Change Scripts

[Watch the Video](#)

Note: Due to their nature, Client Tags will not trigger Gateway Tag Change Scripts. However, there are other similar events, such as a [Client Event Script](#), that enable you to trigger a script based on a Client Tag changing value.



Tag Change Script Interface

Tag List

Lists all of the available Tag Change scripts in the project. Two icons are available below the list:

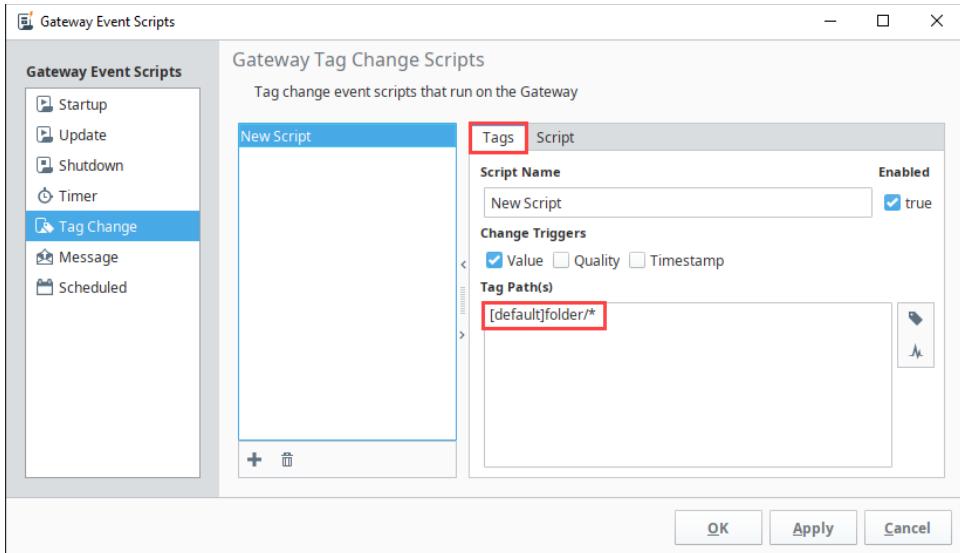
- **Add Script** - Adds a new Tag Change Script to the list.
- **Remove Script** - Removes the currently selected script from the list.

Tags Tab

The **Tags** tab contains settings for the script. See the **Script Settings** description below

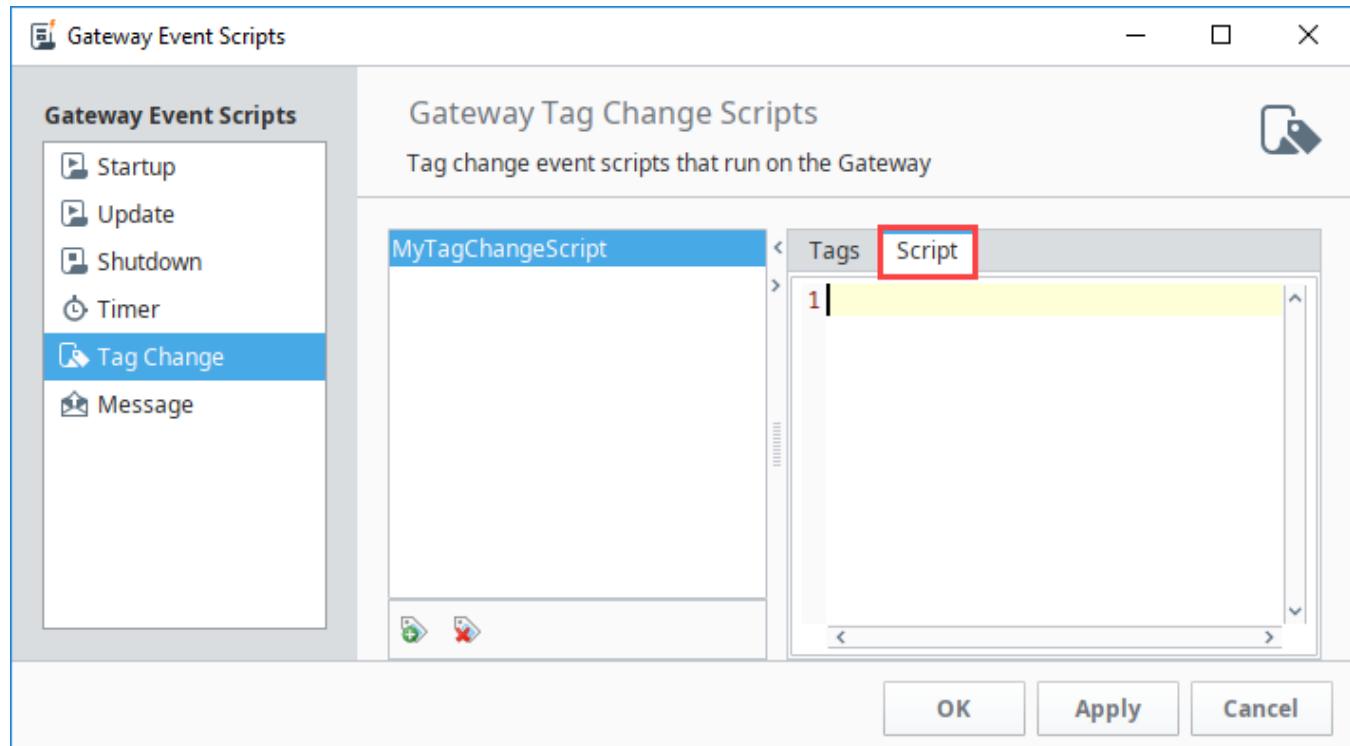
- **Script Name** - The name of the script. Script names must be unique per project.
- **Enabled** - Determines if the script is active or not. Set to false to disable the script.
- **Change Triggers** - When the Tag changes, the script can trigger based on the **Value**, **Quality**, and/or **Timestamp**. Note that regardless of how many triggers changed, the script only executes once per tag, so leaving all triggers enabled will not trigger three executions each time the Tag changes.
- **Tag Paths** - A list of Tag paths to monitor. When any of the Tags listed in this area change, the script will trigger. Note that the list is not comma separated: new paths are specified each line.
 - **Tag Browser** - Opens a Tag Browser window, allowing you to quickly lookup and add Tag paths to the Tag Change Script.
 - **Path Diagnostic** - Click this icon to verify the paths specified under the **Tag Paths** list. This is useful for checking for typos in the list.
 - Wildcards can be used at the end of configured tag paths. Using a wildcard will execute the script for all tags within the same folder. In the below example, the script "New Script" will run when the value of any tag in the "folder" changes.

Note: Wildcards can only be used at the folder level and cannot be used at the tag level. For example, configuring a tag path like [default]folder/* will execute a script on all tags within the folder, but [default]folder/ramp* will not.



Script Tab

The **Script** tab is where the Python script associated with this event will be placed.



Tag Change Objects

Tag Change Scripts contain several built-in objects that are useful for inspecting the event, such as seeing what value the Tag changed to. These objects are listed below.

The initialChange Value

The boolean `initialChange` variable indicates if an event is due to an initial subscription. This is useful as you can filter out the event that is the initial subscription, preventing a script from running when the values haven't actually changed.

```
if not initialChange:
    # Do something useful here
```

The executionCount Value

An integer representing a number of event executions since gateway scripts were restarted (typically by applying a change to a script and saving the project).

```
if executionCount <= 1:  
    pass
```

The newValue Object

A [QualifiedValue](#) object that represents the current value on the tag.

Method	Returns	Description	Usage Example
getValue()	Varies	Returns the new value on the tag	newValue.getValue()
getQuality()	QualityCode	Returns the new quality on the tag	newValue.getQuality()
getTimestamp()	Date	Returns the new timestamp value	newValue.getTimestamp()

The previousValue Object

A [QualifiedValue](#) object that represents the previous value of the tag.

Method	Returns	Description	Usage Example
getValue()	Varies	Returns the previous value on the tag	previousValue.getValue()
getQuality()	QualityCode	Returns the previous quality on the tag	previousValue.getQuality()
getTimestamp()	Date	Returns the previous timestamp value	previousValue.getTimestamp()

The Event Object

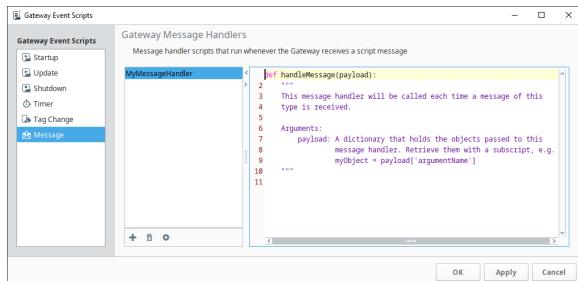
This object offers some additional utility, such as accessing the previous values on the tag.

Method /Attribute	Returns	Description	Usage Example
getCurrentValue()	QualifiedValue	Returns a QualifiedValue object (similar to the <code>newValue</code> object), representing the current value on the tag.	event.getCurrentValue().getValue()
getPreviousValue()	QualifiedValue	Returns a QualifiedValue object, representing the values on the tag before the change.	event.getPreviousValue().getValue()
getTagPath()	TagPath	Returns a TagPath object, that can be further examined for details about the path on the tag that changed. See the TagPath Object table below. Additionally, the TagPath object can easily be turned into a string, providing quick access to the path of the tag that changed value.	event.getTagPath()
getValue()	QualifiedValue	Returns a QualifiedValue object (similar to the <code>newValue</code> object), representing the current value on the tag. This method is functionally identical to <code>getCurrentValue()</code> , and maintained mostly for backwards compatibility reasons.	event.getValue().getValue()
changes	TagChangeType	A RegularEnumSet typed attribute that describes what changed: the value, quality, or timestamp. Values in the set are TagChangeType objects, which can be converted to string with either Python's built-in <code>str()</code> or Java's <code>toString()</code> .	<pre>for i in event.changes: if i.toString() == "ValueChange": foo()</pre>
tagPath		An attribute that describes the tag path on the tag that changed.	event.tagPath.getItemName()

	Method /Attribute	Returns	Description	
	<code>getItemName()</code>	string	Returns the name of the item at the end of the tag path, which can be used to get the name of the tag that changed.	
	<code>getParentPath()</code>	BasicTagPath	Returns a BasicTagPath of the tag's parent folder. Python's str() method can be used to convert the path to a string.	<code>event.tagPath. getParentPath()</code>

Message Scripts

Message Handlers allow you to write a script that will run in the scope they are located in, but they can be invoked by making a call from other projects or even other Gateways. They can be called using three different scripting functions: `system.util.sendMessage`, `system.util.sendRequest`, and `system.util.sendRequestAsync`.




**INDUCTIVE
UNIVERSITY**

Script Messaging

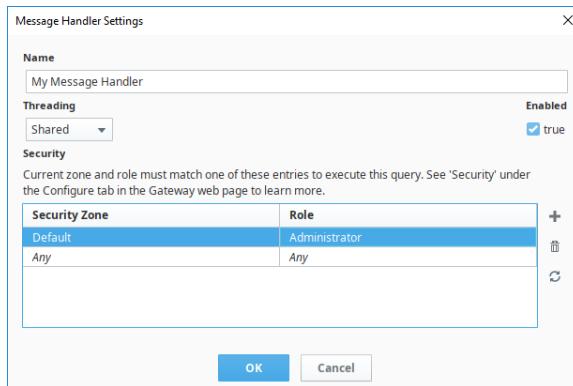
[Watch the Video](#)

Under the list of handlers, three small buttons allow you to add, remove and manage your handlers.

-  **Add Message Handler** - Will add a message handler.
-  **Remove Message Handler** - Will delete the highlighted message handler.
-  **Modify Message Handler** - Will modify the settings for the highlighted message handler.

Gateway Message Handler Settings

When adding or modifying a message handler, a Message Handler settings window will popup.



The following settings are available:

- **Name** - The name of the message handler. Each message handler must have a unique name per project.
- **Threading** - Determines the threading for the message handler. Contains the following options:
 - **Shared** - The default way of running a message handler. Will execute the handler on a shared pool of threads in the order that they are invoked. If too many message handlers are called all at once and they take long periods of time to execute, there may be delays before each message handler gets to execute.
 - **Dedicated** - The message handler will run on its own dedicated thread. This is useful when a message handler will take a long time to execute, so that it does not hinder the execution of other message handlers. Threads have a bit of overhead, so this option uses more of the Gateway's resources, but is desirable if you want the message handler to not be impeded by the execution of other message handlers.

- **Security** - Allows you to specify security zone and role combinations that are allowed to request this message handler.

The Payload

Inside the message handler is your script. The script will have a single object available to it, the **payload**. The payload is a dictionary containing the objects that were passed into it. In essence, the payload is the mechanism that allows you to pass the message handler values.

The payload is simply a python dictionary, so extracting values involves specifying the key:

Pseudocode - Payload Values

```
value1 = payload["MyFirstValue"] # "MyFirstValue" is the key that is  
associated with a value. We are taking the value associated with  
MyFirstValue, and assigning it to value1.  
value2 = payload["MySecondValue"] # Similarly, we are taking the value  
associated with MySecondValue and assigning it to value2.
```

Calling Message Handlers

Once you have your message handlers created, you can then call them from a script using one of three scripting functions: `system.util.sendMessage`, `system.util.sendRequest`, and `system.util.sendRequestAsync`. These functions allow you to call a message handler in any project, even if the project that the message handler resides on is different from the one you are calling it from. The message handler will then execute in the scope in which it was created, and will use any parameters that you pass in through the payload.

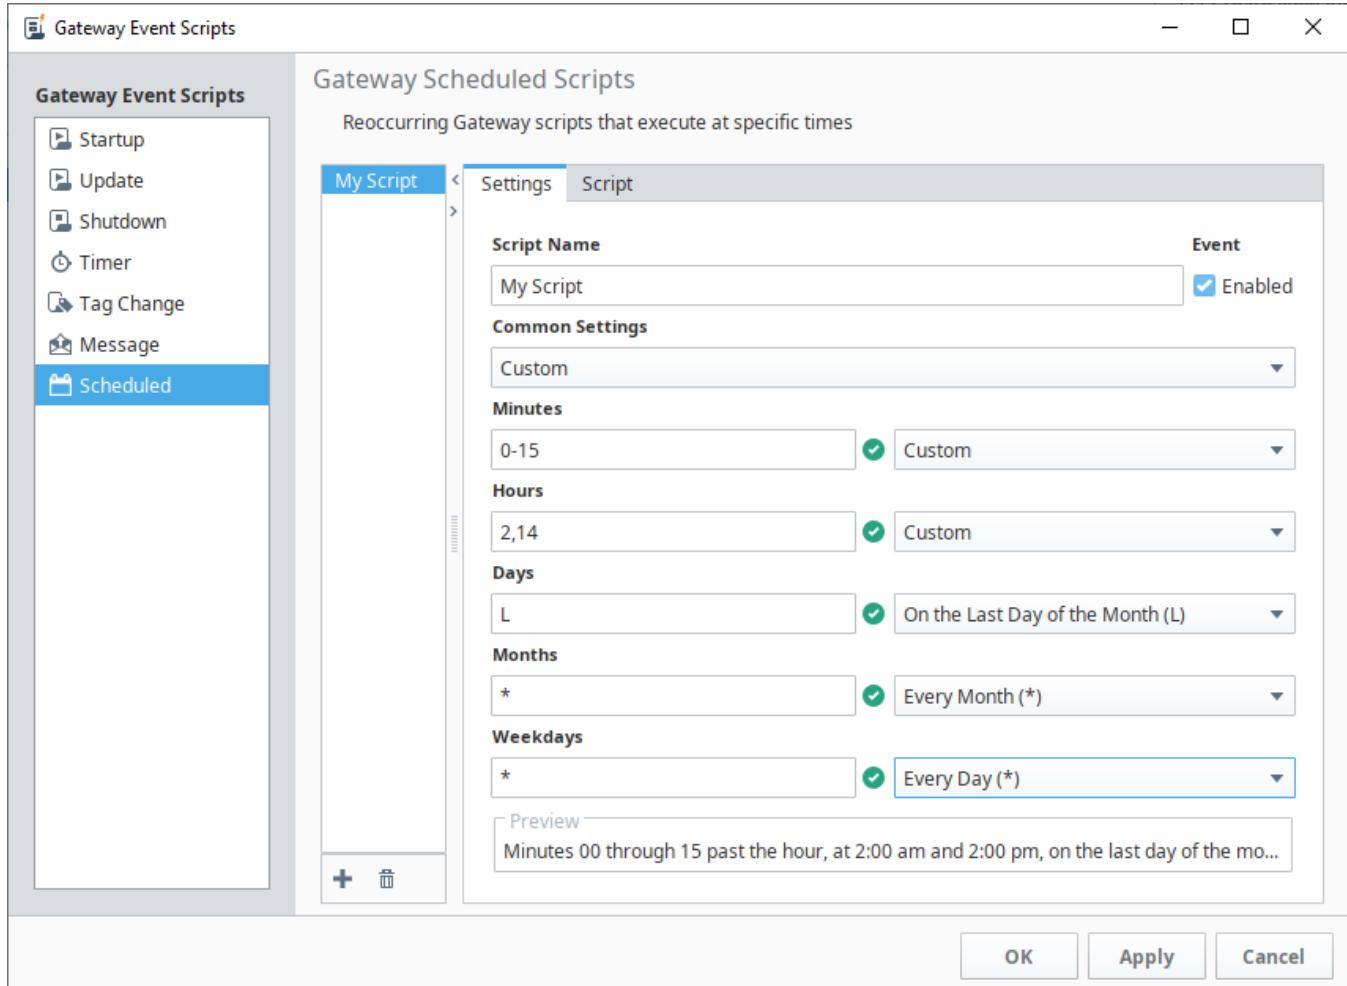
Pseudocode - Calling a Message Handler

```
project="test"  
messageHandler="My Message Handler"  
myDict = {'MyFirstValue': "Hello", 'MySecondValue': "World"}  
results=system.util.sendMessage(project, messageHandler, myDict)
```

Gateway Scheduled Scripts

The following feature is new in Ignition version **8.1.6**
[Click here](#) to check out the other new features

Scheduled scripts are events that execute at fixed times of the day, based off of the Gateway's system time. Configuration for the event is split between two tabs: **Settings** and **Script**.



Settings Tab

The settings tab allows you to give the event a unique name and determine how often it executes. Schedules are driven by cron job scheduling.

The **Common Settings** dropdown contains several common selections, which can be further modified with the fields and dropdowns under each unit of time.

Each unit of time consists of a field, and corresponding dropdown. The dropdown is filled with suggestions and simple options, but more complex values can be provided in the field on the left. See [Crontab Formatting Reference](#).

Script Tab

The script tab houses the Python script that will execute when the scheduled event executes.

Troubleshooting Gateway Scripts

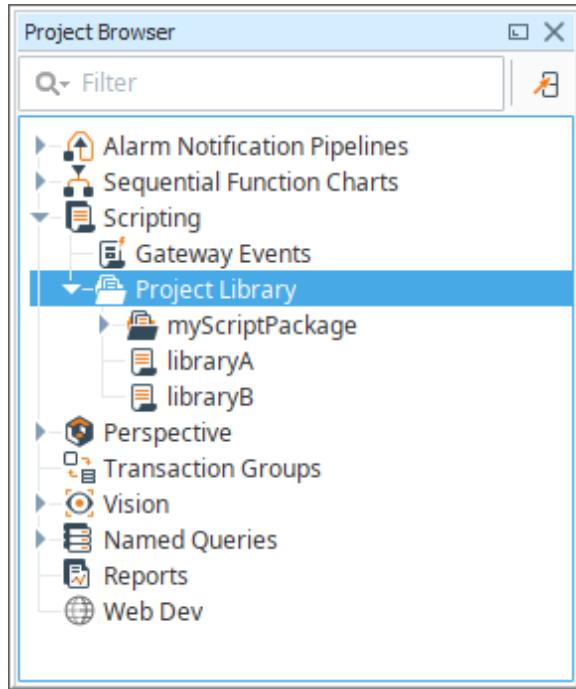
While they are technically project resources, remember that Gateway Event Scripts technically run on the Gateway. Thus the [Status section of the Gateway](#) is useful for diagnosing issues with Gateway Event Scripts.

Related Topics ...

- [Client Event Scripts](#)
- [Perspective Session Event Scripts](#)

Project Library

Scripts under the Project Library are a project-based resource that allows user created Python scripts to be configured. Objects and functions created in a Project Library script can be called from anywhere in the project. Project Library scripts are accessible from the Project Browser, under the **Scripting** item.



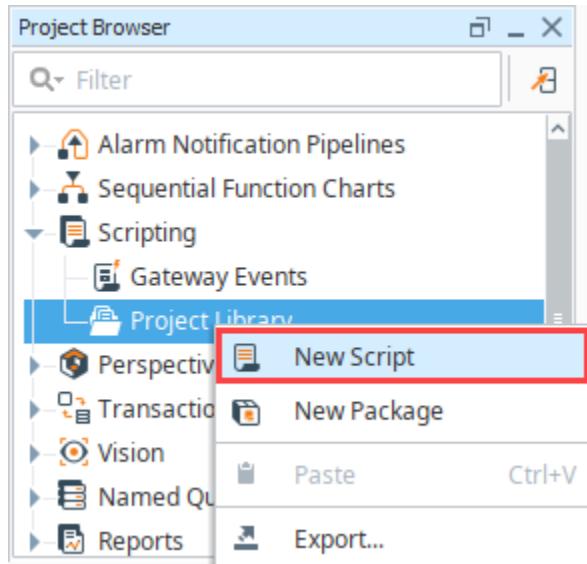
On this page ...

- Add a Script
 - Scripts and Packages
 - Usage Example
 - Project Libraries and Execution
- Gateway Scripting Project
 - Example

Additionally, a single project can be designated as the **Gateway Scripting Project**, meaning that scripts defined in the stated project can be called from the Gateway scope.

Add a Script

To add a Project Library script, right click the **Project Library** item and click the **New Script** option.



Scripts and Packages

There are two main types of resources under the Project Library.

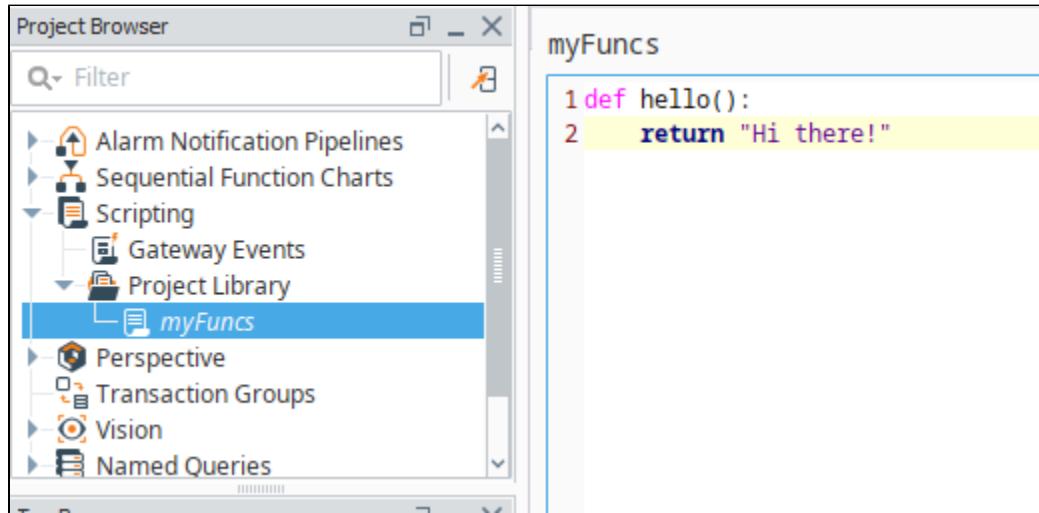
- **Scripts** - Each script resource can contain many **functions** and objects.
- **Packages** - Each package effectively acts as a folder, allowing you to better organize each script resource.

Usage Example

For example, let's suppose you added the following script module named `myFuncs`, whose body is shown below.

Python

```
def hello():
    return "Hi there!"
```



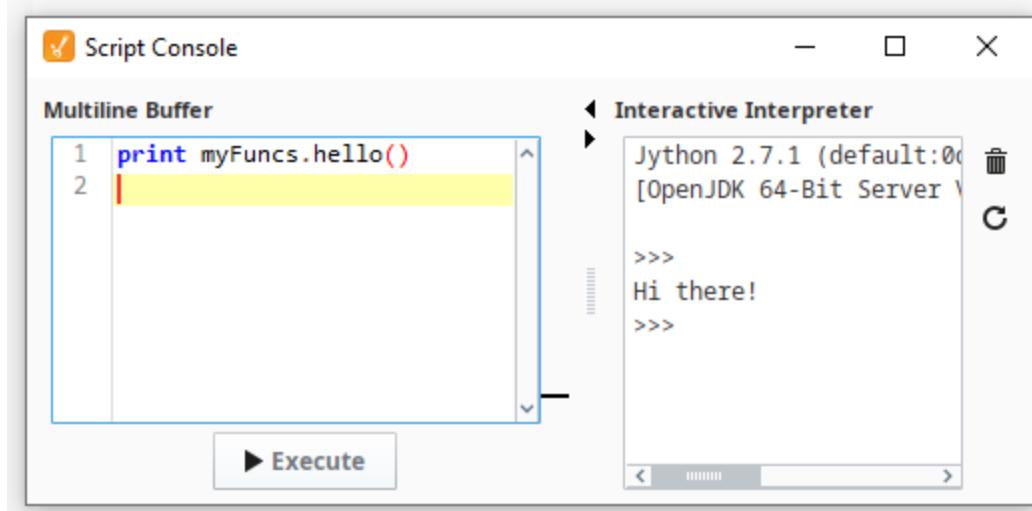
Once we **save** our project, we can now call this function from anywhere within the project using the following syntax

Python - Calling the Project Script

```
myFuncs.hello()
```

Note: Project Library scripts are not accessible to the other resources until the project is saved.

For example, we could open the Script Console (Tools menu > Script Console), write the following, and execute the script.



Each script resource can contain multiple functions and objects. As you add new function definitions, the list on the right will populate giving you a quick way to navigate through long scripts.

The screenshot shows a script resource named 'myFuncs' with the following code:

```

1 myGatewayName = "Inductive_University"
2
3 def hello():
4     return "Hi there!"
5
6 def getGatewayName():
7     return myGatewayName
8
9 def echo(message):
10    return "You said: " + str(message)
11

```

To the right, a sidebar lists the defined functions:

- hello()
- getGatewayName()
- echo(message)

Project Libraries and Execution

Because Python is a dynamic language, any code inside of a project library must be run to build the function and class definitions. This is a common behavior across interpreters. Within Ignition, these project libraries will run under certain conditions. For example, such as when the Script Console in the Designer starts up, if changes are made to third-party libraries inside of the gateway's installation directory, when saving changes to these project libraries, and several other conditions.

When this occurs, code inside of a project library is executed. Meaning classes and functions are defined, and any code that is not contained within either a class or function will execute.

Because of this process, it's generally recommended that all code within a project library is wrapped inside of a function or class definition.

Gateway Scripting Project

Project Library scripts are normally only accessible from the project they were defined in. Thus objects that exist in other scopes, such as Tags that exist in the Gateway scope, are unable to call Project Library scripts. Attempting to do so will result in Gateway log errors stating that "global name 'yourScript' is not defined".

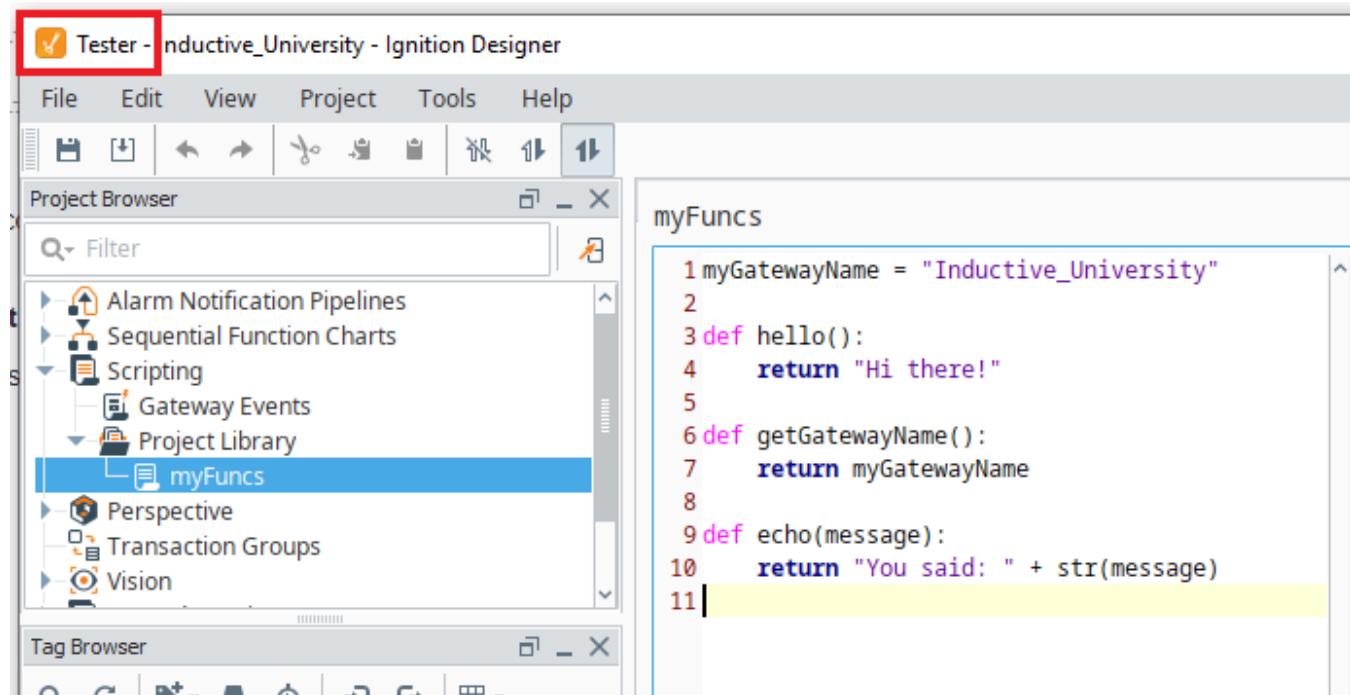
The exception to this rule is the Gateway Scripting Project. This project is specified by the **Gateway Scripting Project** property, which is set in the Config section of the Gateway Webpage under [Gateway Settings](#). Entering the name of a project under this property allows the Gateway access to Project Library scripts configured in the specified project.

Gateway Scripting

Gateway Scripting Project

The name of the Project that gateway-scoped scripts with no project affiliation can access user script libraries in.
(default:)

Thus, if the **myFuncs** library in the prior section was configured in a project named Tester, we would enter Tester in the Gateway Scripting Project field.



Once the name of the project is entered, select **Save Changes**.

Gateway Scripting

Gateway Scripting Project

Tester

The name of the Project that gateway-scoped scripts with no project affiliation can access user script libraries in.
(default:)

Save Changes

Now, you can start calling any of the global scripting events, such as tag event scripts or runScript calls, from the Tester Project Library.

Note: The Gateway Scripting Project has no impact to Gateway event scripts defined in the Designer (startup, shutdown, update, message handler, etc.) as those are Project Resources and thus only function for the project they run under.

Example

This example demonstrates how to call scripts in your Gateway Scripting Project from a tag.

1. Set up your Gateway Scripting Project under **Config > Gateway Settings**:

2. Open your Gateway Scripting Project and add a library called **myFuncs**.
3. Define a new function in this library called **thankYou**:

```
def thankYou(user):
    system.gui.messageBox("Thanks for noticing me, " + user)
```

4. Save your Gateway Scripting Project and exit.
5. Open a new project.
6. Choose a tag and configure a [Tag Event Script](#) on the Alarm Acknowledged event.
7. You can call any function within the **myFuncs** library in your Gateway Scripting Project:

```
myFuncs.thankYou(user = ackedBy)
```

Note: You do not need to specify the name of your Gateway Scripting Project when calling any of its functions.

Web Services, SUDS, and REST

Web Services Overview

Web services are software solutions that allow for interacting with machines residing on a network. In short, web services are nothing more than web pages for machines. They provide a standard way for a third party to request and receive data from a piece of hardware on the network without having to know anything about how that machine works.

Protocols

There are two common approaches to Web Services in Ignition: making a HTTP (HyperText Transfer Protocol) method call, or making a SOAP (Simple Object Access Protocol) call.

- HTTP methods, such as GET and POST, are called using the built-in system functions, such as `system.net.httpGet()`. More information on HTTP calls can be found on the [HTTP Methods](#) page.
- SOAP is XML based, and typically requires a third party Python library. There are [many third party Python libraries that utilize SOAP](#), however, documenting them goes outside the scope of this manual (as do all third party libraries).

The approach you choose typically depends on the server you're trying to make calls to: more specifically, the protocol(s) it supports.



What About the SUDS Library?

The SUDS library, a library that used to come included with the Python Standard Library, offered SOAP based functionality. However, SUDS development has been halted, and is no longer included in the standard library.

In the interest of posterity, the legacy SUDS documentation has been condensed and can be found on the [SUDS - Library Overview](#) page. Note that the legacy documentation should be considered deprecated.

Can Ignition Make RESTful Calls?

Yes it can! However, it is important to understand that REST is an architecture, **NOT** a protocol. Instead, REST utilizes and describes how a protocol should be used. Thus, a RESTful architecture could use both of the protocols mentioned on this page, although HTTP is far more common.

Common Web Services Workflow

While all Web Services follow the same standards, they all do different things. They wouldn't be worth anything if you didn't get the information you need, or if they contained a lot of excess data. If you are unfamiliar with a particular Web Service, there are a few things that you can do to figure out what data is available and how to get it.

1. Identify a Web Service that you will be using. Usually the Web Service has an API somewhere documenting how requests should be made.
2. Write a script to pull some information from the Web Service. If using HTTP, this could mean starting with a GET call, whereas SOAP would involve retrieving the WSDL (Web Services Description Language). In both cases, you may need to find a way to authenticate against the server (usually with some user credentials or an auth token, the API for the service would have more details).
3. Once you have the results from the GET/WSDL, identify the information or functions you want to use.
4. Write a script to use that function and return your values.
5. Parse the results and use them. This can be for display, saving to a database, or anything else you need.

Note: Web Services sometimes take a lot of time to return results, especially the first time they are called. If you put your Web Services script in a button, the client will freeze until the call is complete (this is because the event handlers are run on the GUI thread). It's a good idea to use `system.util.invokeAsynchronous()` or add a waiting image to your screen to let the user know Ignition is working as expected.

Related Topics ...

- [WebDev Module](#)

In This Section ...

On this page ...

- [Web Services Overview](#)
 - [Protocols](#)
 - [Can Ignition Make RESTful Calls?](#)
- [Common Web Services Workflow](#)

HTTP Methods

Overview

Web services calls typically require some protocol to make requests. HTTP is an incredibly common protocol, so this page will introduce how to incorporate these calls in a Python script. Note that all of the examples on this page can be easily called with the Script Console, but can be utilized through some other means (like the actionPerformed event on a Button).

Finding an Endpoint

Ignition doesn't natively expose an endpoint for web services calls, so we'll have to utilize a service of some sort for the examples on this page. Fortunately, there are many public services we can utilize, such as [OpenWeather](#). From here, we can generate an endpoint to use. At the time of this writing, we're using the following endpoint URL template:

Endpoint URL Template

```
https://api.openweathermap.org/data/2.5/weather?lat={lat}&lon={lon}&appid={API key}
```

The following table describes the various parameters in the endpoint URL:

Parameter Key	Description
lat, lon	Geographical coordinates (latitude, longitude).
appid	Your unique API key. Note: The code snippets on this page use "Your API key" as a placeholder for unique API keys. You will need to substitute the placeholder with your own API key for the code snippets to function as expected.

The following code snippets are using the geographical coordinates for the Inductive Automation corporate headquarters located in Folsom, CA.

Making the Call

To retrieve the results of this information in Ignition, we can use [system.net.httpClient\(\)](#) to fetch the results of this call. We can try the following script in the [Scripting Console](#):

Python - Creates a variable to Store the Endpoint and Retrieves Results

```
#Set the endpoint URL
url = "https://api.openweathermap.org/data/2.5/weather"

#Declare a variable for system function we are using
myClient = system.net.httpClient()

#Declare a variable and set the parameters for the endpoint URL
#Instead of specifying the API call parameters for the endpoint URL in the beginning
response = myClient.get(url, {"lat":38.652330, "lon":-121.189773, "appid":"Your API key"})

#print the output
print response.getText()
```

Printing this results in the following Python string:

Python - Results

On this page ...

- [Overview](#)
- [Finding an Endpoint](#)
- [Making the Call](#)
- [Parsing the Results](#)
- [Make the Results Human Readable](#)
- [Troubleshooting HTTP Methods](#)
- [HTTP Response Codes](#)

```
{"coord": {"lon": -121.1898, "lat": 38.6523}, "weather": [{"id": 800, "main": "Clear", "description": "clear sky", "icon": "01d"}], "base": "stations", "main": {"temp": 310.34, "feels_like": 310.11, "temp_min": 308.88, "temp_max": 312.74, "pressure": 1013, "humidity": 26}, "visibility": 10000, "wind": {"speed": 4.12, "deg": 290}, "clouds": {"all": 0}, "dt": 1654893775, "sys": {"type": 2, "id": 2006213, "country": "US", "sunrise": 1654864794, "sunset": 1654918128}, "timezone": -25200, "id": 5349705, "name": "Folsom", "cod": 200}
```

Parsing the Results

If we wanted to extract a single value out of the results, we have a number of approaches. One useful approach would be to turn this JSON string into a Python Dictionary. This way we can single out a key instead of using regex or looking for substrings (both valid approaches in their own regard).

When presented with a JSON string, we can call [system.util.jsonDecode\(\)](#) to turn a JSON string into a native Python Object. Thus, we can modify our code to the following:

Python - Parsing the Results Code

```
#Set the endpoint URL
url = "https://api.openweathermap.org/data/2.5/weather"

#Declare a variable for system function we are using
myClient = system.net.httpClient()

#Declare a variable and set the parameters for the endpoint URL
response = myClient.get(url, {"lat": 38.652330, "lon": -121.189773, "appid": "Your API key"})

#Set a variable for the contents of the API call
results = response.getText()

# Convert the JSON string into a Python object. In this case, it results in a Dictionary.
decodedDict = system.util.jsonDecode(results)

# Now we can treat the results like a nested dictionary, thus we can specify the "weather" key,
# and then the nested "description" key to return a description of the current weather conditions.
# "[0]" is needed to specify the index of the list before specifying the "description" key.
print decodedDict.get("weather")[0].get("description")
```

Now we can easily retrieve a single value by specifying key names on the results. Printing this results in the following Python string (at the current time):

Python - Results

```
clear sky
```

Make the Results Human Readable

Now that we know how to extract the results, we should clean up the output of the GET call. The JSON string returned by the endpoint could potentially be long and cumbersome to read through for a human, but we can use Python's built-in **pprint** library to pretty print the results.

Python - Now with Pretty Print

```
# Import the pprint library
import pprint

# We'll instantiate an instance of PrettyPrinter, and store it in a variable named pp.
pp = pprint.PrettyPrinter(indent=4)

#Set the endpoint URL
url = "https://api.openweathermap.org/data/2.5/weather"

#Declare a variable for system function we are using
myClient = system.net.httpClient()
```

```

#Declare a variable and set the parameters for the endpoint URL
response = myClient.get(url, {"lat":38.652330, "lon":-121.189773, "appid":"Your API key"})

#Set a variable for the contents of the API call
results = response.getText()

# Convert the JSON string into a Python object. In this case, it results in a Dictionary.
decodedDict = system.util.jsonDecode(results)

# Print out the dictionary in an easy to read format.
print pp.pprint(decodedDict)

```

The resulting output, which is much easier to read, looks like the following:

Python - Results

```

{
    u'base': 'stations',
    u'clouds': {   u'all': 0},
    u'cod': 200,
    u'coord': {   u'lat': 38.6523, u'lon': -121.1898},
    u'dt': 1654896566,
    u'id': 5349705,
    u'main': {   u'feels_like': 310.59,
                 u'humidity': 25,
                 u'pressure': 1013,
                 u'temp': 310.84,
                 u'temp_max': 313.46,
                 u'temp_min': 309.19},
    u'name': 'Folsom',
    u'sys': {   u'country': 'US',
                u'id': 2006213,
                u'sunrise': 1654864794,
                u'sunset': 1654918128,
                u'type': 2},
    u'timezone': -25200,
    u'vesibility': 10000,
    u'weather': [   {   u'description': 'clear sky',
                      u'icon': '01d',
                      u'id': 800,
                      u'main': 'Clear'}],
    u'wind': {   u'deg': 330, u'speed': 4.12}
}
None

```

From here we can see all of the keys that lead to our final value:

```

◀ Interactive Interpreter
▶ Jython 2.7.2 (uncontrolled:000000000000, Jan 25 2022, 14:39:15)
[OpenJDK 64-Bit Server VM (Azul Systems, Inc.)] on java11.0.15

>>>
{
  u'base': 'stations',
  u'clouds': {  u'all': 0},
  u'cod': 200,
  u'coord': {  u'lat': 38.6523, u'lon': -121.1898},
  u'dt': 1654896566,
  u'id': 5349705,
  u'main': {  u'feels_like': 310.59,
              u'humidity': 25,
              u'pressure': 1013,
              u'temp': 310.84,
              u'temp_max': 313.46,
              u'temp_min': 309.19},
  u'name': 'Folsom',
  u'sys': {  u'country': 'US',
             u'id': 2006213,
             u'sunrise': 1654864794,
             u'sunset': 1654918128,
             u'type': 2},
  u'timezone': -25200,
  u'vesibility': 10000,
  u'weather': [  {  u'description': 'clear sky',
                  u'icon': '01d',
                  u'id': 800,
                  u'main': 'Clear'}],
  u'wind': {  u'deg': 330, u'speed': 4.12}}
None
>>>

```

To get to the value for the description key, we simply need to address each key along the way:

Python - To Get Value for the Sunset Key

```
# "[0]" is needed to specify the index of the list before specifying the "description" key.
print decodedDict.get("weather")[0].get("description")
```

Troubleshooting HTTP Methods

When making HTTP calls, it is helpful to be familiar with the status codes that are returned by errors. To demonstrate, we could modify an earlier example:

Python - Returns Error Status Codes

```
#Set the endpoint URL
url = "https://api.openweathermap.org/data/2.5/weather"

#Declare a variable for system function we are using
myClient = system.net.httpClient()
```

```

### Note that instead of replacing the "appid" parameter with a unique API key, we will leave it as the
#"Your API key" placeholder.
#Declare a variable and set the parameters for the endpoint URL
response = myClient.get(url, {"lat":38.652330, "lon":-121.189773, "appid":"Your API key"})

#print the output
print response.getText()

```

This will return an error, which looks like the following:

Python - IOError: Response Code 400

```
{"cod":401, "message": "Invalid API key. Please see http://openweathermap.org/faq#error401 for more info."}
```

Note that HTTP response code **401**, which means **unauthorized**, was referenced. This error code is correct because we intentionally used an incorrect API key!

HTTP Response Codes

The World Wide Web Consortium has a [page dedicated to HTTP response codes](#), which details all possible error codes. However, several common codes are listed below:

Response Code	Description
400	Bad Request - The server could not understand the request due to malformed syntax.
401	Unauthorized - The request requires some sort of authentication. Potentially some user credentials or an auth token of some kind.
403	Forbidden - The server understood what you requested, but is intentionally refusing the request. In some cases, the error message may include a reason why the request was not fulfilled (but not always). Typically, if the server doesn't include a reason, they'll use a 404 error code instead
404	Not Found - Your syntax was correct, but the server could not find the resource you were asking for. This could mean a typo, or missing portion of the URL you are using. In this case, double check the address you're specifying. Depending on the configuration, this could also mean that the server does actually have the resource you requested, but doesn't want to confirm its existence (typically due to a security policy. See error code 403 above).

Related Topics ...

- [system.net](#)

SUDS - Library Overview

The SUDS Library

The SUDS library is a SOAP-based web services client developed for Python. It is extremely simple to use and practically eliminates the need for the user to understand or even view the WSDL of a web service.

Disclaimer

The SUDS library used to be included in the Python Standard Library. However it has since been removed, meaning you may not have access to it when performing a fresh install of Ignition. Additionally, development on the library has mostly ceased, so any copies you find online may be drastically outdated.

The information on this page will be maintained for legacy users that need to be familiar with the old SUDS library. As a result, this page and its contents should be considered deprecated.

The SUDS library interprets the WSDL file for you and through a couple simple function calls allows you to get a list of the available methods provided to you by the web service. You can then invoke these methods in Ignition through event scripting to send and receive data in your projects. You will have to familiarize yourself with the SUDS library in order to make use of it.

Simple Example

If you read through the SUDS documentation, you'll see that the Client object is the primary interface for most users. It is extremely simple using this object and a few print statements to view a list of available methods provided by the web service you are connecting to. This example will illustrate how to make an initial connect to a web service, print out the list of available methods, and then call one of these methods and display the resulting value in a label on an Ignition window at the push of a button. The below example uses a public web service and is available for anyone to test against.

1. First, we can use the script playground to test out some scripting calls and see the output. The below example shows how to get a reference to a client object. By printing this client object to the console we get an output of all the available methods, types, and other information about the web service as defined in the WSDL file.

Python - W3Schools WSDL

```
from suds.client import Client
client = Client("http://www.w3schools.com/xml/tempconvert.asmx?WSDL")
print client
```

The screenshot shows the Ignition Script Console and Interactive Interpreter. The Script Console window on the left contains the Python code. The Interactive Interpreter window on the right shows the output of the code execution. The output includes the Java version, the Suds library version, and the WSDL service definition, which lists two methods: CelsiusToFahrenheit and FahrenheitToCelsius.

This WSDL defines two functions: CelsiusToFahrenheit(string tempCelsius) and FahrenheitToCelsius(string tempFahrenheit). These are the functions that this web service makes available to you. Don't worry about the fact that the methods are listed twice. This is just because the WSDL has two definitions of the functions that are formatted for different SOAP version standards. To call these functions in Ignition

On this page ...

- The SUDS Library
 - Simple Example
 - Beyond the Example
 - Complex Arguments

scripting, you have to make use of the "service" member of the client object. You can see printing the returned results shows the conversion.

The screenshot shows a Python script console interface. On the left, a 'Multiline Buffer' window contains the following code:

```
from suds.client import Client
client = Client("http://www.w3schools.com/xml/tempconvert.asmx?WSDL")
print client.service.FahrenheitToCelsius("85")
```

A yellow box highlights the line `print client.service.FahrenheitToCelsius("85")`. Below this is a 'Execute' button. To the right is an 'Interactive Interpreter' window titled 'Jython 2.5.3 (v2.5.3:3d2, Java HotSpot(TM) 64-Bit)'. It shows the output:

```
>>>
29.44444444444444
>>>
```

2. To make a simple conversion window in an Ignition project you can add a button, a numeric textbox, and a label to a window. Then on the button to calculate a Fahrenheit to Celsius calculation, you would place something like the following:

Python - Fahrenheit to Celsius

```
from suds.client import Client
client = Client("http://www.w3schools.com/xml/tempconvert.asmx?WSDL")

far_value = event.source.parent.getComponent('Numeric Text Field').floatValue
cels_value = client.service.FahrenheitToCelsius(far_value)

event.source.parent.getComponent('Label').text = cels_value
```



3. Then you can make a second button to do the opposite: calculate Celsius to Fahrenheit.

Python - Celsius to Fahrenheit

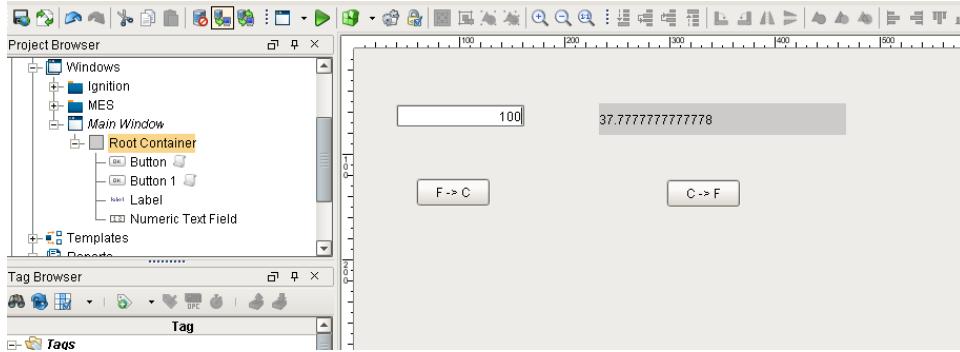
```
from suds.client import Client
client = Client("http://www.w3schools.com/xml/tempconvert.asmx?WSDL")

cels_value = event.source.parent.getComponent('Numeric Text Field').floatValue
far_value = client.service.CelsiusToFahrenheit(cels_value)

event.source.parent.getComponent('Label').text = far_value
```



4. With your scripts in place your window should now function as a simple temperature conversion tool!



Beyond the Example

While the example is relatively simple, it can easily be expanded upon. However, always keep the general workflow in mind when using the SUDS library:

Pseudocode - WSDL Workflow

```
#Import the SUDS Client object
from suds.client import Client

#Instantiate a new Client Object
client = Client("url_to_your_wsdl")

#Call the desired method using the service instance variable
client.service.MyMethod(myArgument)
```

Complex Arguments

In the overview, the methods provided by the web service were very simple and took simple argument types. Sometimes however, the web service will describe complex types and allow you create instances of these types that can then be added to the system/machine that the web service is providing an interface for.

A simple, hypothetical example of this would be a system that stores contact information of clients and can be used as an address book of sorts by other systems on the network. It may provide not only a way to pull contact information for a certain individual out, but also a way to insert new contacts. We'll keep the example simple and say that contacts have only a name and a phone number.

Note: This example is completely hypothetical. It is intended to give insight into complex types. It does not make use of an actual functional web service.

For example, say we create and print the client object we associated with our web service in the following manner:

Pseudocode - Client Object

```
from suds.client import Client
url = 'http://localhost:7575/webservices/hypothetical_webservice?wsdl'
client = Client(url)
print client
```

And the resulting output is the following:

Python - Results

```
Suds ( https://fedorahosted.org/suds/ ) version: 0.4 GA build: R699-20100913

Service (hypothetical_webservice)
Prefixes (0):
Ports (1):
```

```

(Soap)
Methods:
    addContact(Contact contact, )
    getContactList(xs:string str, xs:int length, )
    getContactByName(Name name, )

Types (3):
    Contact
    Name
    Phone

```

Here you can see that, while not too complicated, the web service defines more than just methods that take simple type arguments and return the same. Under the Types section, you can see there are three "complex" types. These are basically just objects like one creates in an object oriented programming language like java. The SUDS Client object has an instance variable called "factory" that allows you to create these complex types so you can use them to invoke methods defined by your web service that take complex arguments.

If we wanted to add a contact using the addContact() method, we have to create a contact object first:

Pseudocode - Using a Method

```

contact = client.factory.create('Contact')
print contact

```

The create function creates a new contact object that knows its own structure. We can view this structure by calling print on this new object and see that it prints the following:

Python - Structure

```

(Contact)=
{
    phone = []
    age = NONE
    name(Name) =
    {
        last = NONE
        first = NONE
    }
}

```

By examining the Contact type object, we can see its structure and know what we need to create in order to have a valid Contact to add to the address book. We could then do the following to supply the necessary information for the Contact object and then call our addContact function.

Pseudocode - Adding a Contact

```

contact = client.factory.create('Contact')

phone= client.factory.create('Phone')
phone.areacode = '916'
phone.number = '5557777'

name = client.factory.create('Name')
name.first = 'John'
name.last = 'Doe'

contact.name = name
contact.phone = phone
contact.age = 30

client.service.addContact(contact)

```

After execution a new contact will have been added via the web service!

Steps to remember when using complex types:

Pseudocode - Complex Type Reminders

```
#Create a new type object using the factory instance variable of the Client object  
my_type = client.factory.create('MyType')  
  
#If you don't know the structure of the newly created object then print it to the console  
print my_type
```

Related Topics ...

- [Web Services, SUDS, and REST](#)

JSON Format

About JSON

JavaScript Object Notation (JSON) is a language-independent data format. It's a lightweight format for storing and transporting data (i.e., when data is sent from a server to a web page). JSON is easy for users to read and write and for machines to parse and generate.

JSON Rules

JSON has a few simple rules:

- Data is in name/value pairs.
- Data is separated by commas.
- Curly braces hold objects.
- Square brackets hold arrays.

On this page ...

- [About JSON](#)
- [JSON Rules](#)
- [How JSON Works](#)
 - [JSON Data - Name and a Value](#)
 - [JSON Objects](#)
 - [JSON Arrays](#)
- [Where JSON Is Used in Ignition](#)
 - [Perspective Component Properties](#)
 - [Tags](#)
 - [JSON in Tag UDTs](#)
 - [Looping through JSON Objects with Scripting](#)

How JSON Works

JSON Data - Name and a Value

JSON data is written as name/value pairs. A name/value pair consists of a field name in double quotes, followed by a colon, followed by a value.

```
{  
    "companyName": "Inductive Automation"  
}
```

JSON Objects

JSON objects are written inside curly braces.

Note: The properties of a JSON object have no defined order. If you need a defined order, use a [dataset](#) or an array instead.

```
{  
    "firstName": "Sally",  
    "lastName": "Smith"  
}
```

JSON Arrays

JSON Arrays are written inside square brackets. An array can contain objects. In the following example, the object "companies" is an array and contains three objects.

```
{  
    "companies": [  
        {  
            "companyName": "Inductive Automation",  
            "cityName": "Folsom",  
            "stateName": "CA"  
        },  
        {  
            "companyName": "Hewlett Packard",  
            "cityName": "Palo Alto",  
            "stateName": "CA"  
        }  
    ]  
}
```

```

        "stateName" : "CA"
    },
    {
        "companyName" : "Apple",
        "cityName" : "Cupertino",
        "stateName" : "CA"
    }
]
}

```

Where JSON Is Used in Ignition

Ignition uses the JSON format to store much of its data internally, including Tags and Perspective component properties.

Perspective Component Properties

Components have properties (props), which are simply named values. These properties are arranged in a tree structure following the structure and data model of the common JSON document format. Component properties are defined as a JSON structure, and are variable according to the type of component that the config object represents. All components registered in the module have a set of default properties included. These defaults are provided to the instantiated component at runtime, and so default props are not saved. Instead, only those which have a value that differs from the default are stored, serialized and sent to the client during loading.

Example - Sample Data from Table Component

```
[ { "city": "Helsinki", "country": "Finland", "population": 635591 }, { "city": "Jakarta", "country": "Indonesia", "population": 10187595 }, { "city": "Madrid", "country": "Spain", "population": 3233527 }, { "city": "Prague", "country": "Czech Republic", "population": 1241664 }, { "city": "San Diego", "country": "United States", "population": 1406630 }, { "city": "Tunis", "country": "Tunisia", "population": 1056247 } ]
```

Tags

Ignition exports and imports Tag configurations to and from JSON. Tags are defined as JSON objects, which consist of properties, arrays, and sub-objects. The system.tag.configure function can take either a String document definition, or a JSON object that defines one or more Tags. Overrides for UDTs are created by simple redefinition of properties, and complex structures like Event Scripts and Alarm configurations will be merged with inherited definitions.

You can copy the JSON or one or more Tags in the Tag Browser. This copies them into the system clipboard. In addition, pasting the JSON into a different provider/designer will create or overwrite tags.

Example 1 - Tag Export

```
{
    "name": "Tank Instance",
    "typeId": "Tank UDT",
    "tagType": "UdtInstance",
    "tags": [
        {
            "value": "80",
            "name": "Tank Level",
            "tagType": "AtomicTag"
        },
        {
            "value": 80,
            "name": "sliderValue",
            "tagType": "AtomicTag"
        }
    ]
}
```

Example 2 - Tag Export

```
{
  "tags": [
    {
      "valueSource": "memory",
      "dataType": "Boolean",
      "alarms": [
        {
          "setpointA": 1,
          "name": "Above Normal"
        }
      ],
      "name": "Boolean Tag",
      "value": false,
      "tagType": "AtomicTag"
    },
    {
      "valueSource": "memory",
      "dataType": "Boolean",
      "name": "One Shot Trigger",
      "tagGroup": "Driven One Shot",
      "value": true,
      "tagType": "AtomicTag",
      "enabled": true
    },
    {
      "valueSource": "opc",
      "opcItemPath": "ns\u003d1;s\u003d[Generic]_Meta:Random/RandomDouble1",
      "dataType": "Float8",
      "name": "Pressure3",
      "tagGroup": "Driven One Shot",
      "tagType": "AtomicTag",
      "enabled": true,
      "opcServer": "Ignition OPC UA Server"
    },
    {
      "valueSource": "opc",
      "opcItemPath": "ns\u003d1;s\u003d[Generic]_Meta:Random/RandomDouble2",
      "dataType": "Float8",
      "name": "Thickness3",
      "tagGroup": "Driven One Shot",
      "tagType": "AtomicTag",
      "enabled": true,
      "opcServer": "Ignition OPC UA Server"
    }
  ]
}
```

JSON in Tag UDTs

You can also set JSON strings as properties in an Ignition Tag. Any properties of a Tag can be set to a string that represents a JSON object.

Note: In a UDT, the {} braces are used to denote a property reference. Make sure you don't have any properties with names that look like JSON objects so there is no overlap. Strings will appear as black text, parameter references will appear as grey and italicized.

Looping through JSON Objects with Scripting

Traversing a JSON object in scripting is simple as long as the structure is known. If there are objects within objects, you can use multiple loops to get through it all.

Let's use the JSON Array object above for a simple example. You can loop through the list to get the repeating items by name.

```
# loop through the JSON data
# fetch the data, this will change depending on where the script is in relation to the table
json = self.items
# access the companies object (which is a list)
companies = json["companies"]
```

```
# loop through the companies list
for company in companies :
    # get each item out of the row object
    name = company[ "companyName" ]
    city = company[ "cityName" ]
    state= company[ "stateName" ]
    # now do something with the data
```

Related Topics ...

- [JSON Functions](#)

Basic Python Troubleshooting

When learning how to code in Python, most Ignition users tend to place most of their learning efforts on memorizing syntax or other aspects of the language. While being comfortable with the language is useful, there are plenty of references available: countless books and websites that describe syntax and usage already exist.

In truth, the best thing you can do to make yourself a better programmer is to learn some basic troubleshooting behaviors. While syntax examples are all over the Internet, examples detailing exactly what you want your code to do, to the extent you want it to, will be difficult if not impossible to find.

This section details how to troubleshoot a script in Ignition. You won't walk away from this section having exact answers to specific problems, but rather examples and concepts that you can apply to your own scripts.

Coding Best Practices

The following are some general and helpful best practices that can help minimize the amount of time you spend troubleshooting, as they can help better direct you to a problem.

On this page ...

- Coding Best Practices
 - Look for Errors
 - If You Find an Error, Read It!
 - Use Print Statements
 - Add Comments
 - Test Early, Test Often
 - Avoid Hard-Coding
 - Arguments: Use Variables Instead
 - Decide on a Naming Convention
 - The Simplest Approach Really Is the Best Approach

Look for Errors

When code fails mid-execution, it always generates an error message. Where the message appears depends on where the script executed:

Scope	Print Command	Print Output Location
Gateway <i>Note: anything that isn't in the Client or Session scopes listed below uses this scope</i>	<code>system.util.getLogger</code>	Scripts on Gateway-scoped resources (Tags, Alarm Pipelines, SFCs, etc.) will appear on the Logs page of the Status section on the Gateway . Additionally, the wrapper.log file in Ignition's installation directory will have these messages. Here are the default wrapper.log file paths for each operating system: <ul style="list-style-type: none">• Windows: Program Files/Inductive Automation/Ignition/logs• Linux: /var/log/ignition• Mac OS X: /Users/UserName/Documents/Ignition-osx-x.x.x/logs
Vision	Python's <code>print</code> command, or <code>system.util.getLogger</code>	The Client Console will contain any errors generated in the client: press Ctrl + Shift + F7 to open the console, or using the menubar in the client to go to Help > Diagnostics then click on the Console Tab . Additionally, a red Error Box should appear with details on the error if a Component Event Handler threw the exception: Extension Functions do not generate the Error Box. If you don't see the error box then it might be minimized, or open in the background (behind the Client).
Perspective	<code>system.perspective.print</code>	Web browsers generally offer a way to inspect a page, which usually contains a console of some sort.
Designer	Depends. See the output location.	The Designer reports errors in a similar manner to Clients: errors appear in the Designer's Console (Ctrl + Shift + c) . Component Event Handlers will generate a red Error Message. Note that events some events, such as Client Startup Scripts, will not trigger in the Designer, so get in the habit of launching a Client when testing non-component scripting events. While in the Designer, and working on Perspective resources, calls to <code>system.perspective.print</code> will appear in the Designer's Console as well.

What happens if I don't see any errors?

If there truly isn't an error message somewhere, but your code isn't doing what you expected, then ask yourself the following questions:

1. What is your script supposed to do?
2. What is it actually doing?

Question #2 is harder to answer: if you knew what it was doing, you wouldn't be stuck! The best way to answer this question is by adding `print` statements to your code.

When a script doesn't perform to expectations, it can suggest a problem with the script's workflow. Some of the [pages in this section](#) can offer some suggestions on what to do.

If You Find an Error, Read It!

It is common for users that are new to scripting to see an error message, immediately close it without much thought, and then stare at their code as if the problem will politely make it self known. While looking over your code line-by-line will eventually lead you to the problem, the error messages can provide you with a shortcut to the issue. [Check out some of the pages in this section for more information on reading an error message.](#)

Use Print Statements

When testing your scripts, the print command can help you verify that your code is behaving the way it should. This allows you to reconstruct what your code did when it executed. Don't be afraid to add helpful print statements to your code.

Once your code works as expected, remember to either remove or comment out the print statements, so they don't flood the console during normal use as this makes troubleshooting other issue more difficult. Just be mindful of the scope of the script, as that determines how you generate print messages, as well as where the console is.

Python - Using print to Troubleshoot

```
myVar = system.tag.read("folder/tag")

# printing out variables you are going to use in an if-statement later allows you to confirm that the values
# are what you expect them to be.
print "myVar is set to: " + str(myVar)

# Sometimes viewing the data type of the variable can prove helpful.
print "myVar is a: " + type(myVar)

# Should your code have multiple if-statements, adding a print statement before and after can show you where
# the flow of the script went.
print "Starting if-statement"

if myVar > 100:

    # If you don't see this print statement, then your if-statement evaluated to False.
    print "Inside if-statement"
    doWork()

# Printing the end of your script doesn't give you any useful troubleshooting information, but if you need
# to trigger the script multiple times, it helps delineate each execution.
print "Script Ended"
```

Add Comments

While comments are useful to remind you how your code works, you can also use them to plan your script before you write any code. Break down what you want the code to do into several smaller steps, and then leave comments describing those steps in order. This provides you a chance to review the script's workflow before worrying about syntax. It also provides natural points to stop and test your code, to make sure it is doing what you think it should do.

```

1 ### When this script triggers, it needs to read the value of a Tag,
2 ### and retrieve a record from the database. It should then parse
3 ### the results, and write the matching value to the Label.
4
5 # Read from the Tag
6
7
8 # Write the Query using the Tag value
9
10
11 # Send off the Query using a system.db.runPrepQuery
12
13
14 # Check the results for the matching record
15
16
17 # Write the record to the Label on the Window.

```

Test Early, Test Often

Unless the script is very simple, avoid writing the entire script and then testing at the end. You may have missed an important line early on, so now you have to adjust all of the code below that line to make the script run. As mentioned above, add print statements, and run your script to make sure it is doing what you think it should.

Additionally, stopping to test your code provides an excellent opportunity to **save your project**. Get into the habit of saving before you execute any new code.

Avoid Hard-Coding Arguments: Use Variables Instead

Instead of doing this:

Python - Hard Coded Message Box

```
system.gui.messageBox("Hello you. Glad to see you")
```

Try to get into the habit of doing this:

Python - Message Box Variable

```
message = "Hello you. Glad to see you"
system.gui.messageBox(message)
```

Simple examples like the above don't make for the best use case, but when you have a large script that references a value multiple times, it is easier to declare the value once in a variable, and then just reference the variable throughout your code. If you need to change the value later, then you can simply change it once where you initialize the variable, and you don't have to search every line of your code looking for the value.

Decide on a Naming Convention

When creating variables, try to adopt a naming convention that comes natural to you, and stick to it. If you consistently write variables using the same conventions, you will be less likely to end up with a typo when referencing that variable later in your code. Remember that code is case-sensitive, so something as simple as forgetting to capitalize a letter will cause an error.

This is especially important if you are working in a group on the same project. It's better to get together with your colleagues and agree upon some naming conventions before you write any code.

Ignition's **system.* functions** use the camelCase naming convention. That is, the first letter of each word is capitalized except for the very first letter. We recommend that you use it for variable names because it is easy to remember to use camelCase for both functions and variable names instead of for just one.

The Simplest Approach Really Is the Best Approach

When learning how to code, it's not uncommon to run into multiple issues that require you to find a workaround. However, these workarounds can cascade into other issues and make your script more complicated. Consider the following:

```
"The goal of my script is to access A and then output B...
Shoot, B requires C, so I'll add C...
Wait, C requires that D exists, so I'll create that...
Oh, D needs interfaces E, F, and G, so let's add those in...
Hmm. E needs H and I, F needs J and K, while G needs Y, Z and...A again?!?"
```

Take a step back and ask yourself "what is this script doing"? If you can't do that in a sentence or two, you may want to rethink the script. If the scope of the script is too large, then it considerably increases the complexity of the code, which in turn could add a plethora of problems later.

If you keep adding workarounds, but your code is not getting any closer to achieving its end goal, there may be an easier way to accomplish what you're trying to do with a different approach.

Related Topics ...

- [Python Scripting](#)

In This Section ...

Reading Error Messages

When an error occurs in the execution of a script, an error message box will pop up. The popup box appears in front of any open Designer windows, and will remain in view until you close it or click on something behind it.

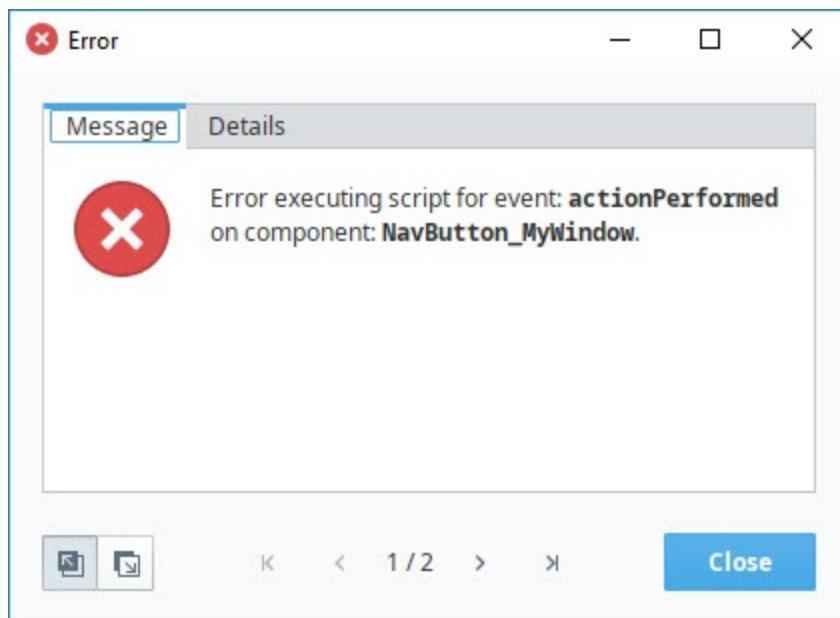
There are two modes for the Error box selectable by the **Send to Front**  and **Send to Back**  icons.

- **Send to Front**  means additional error messages will cause the popup to reappear on top.
- **Send to Back**  will cause the errors to remain hidden below the Designer.

Error Message Box Overview

Exceptions usually include a line number. Take note of the number in the Details tab, and start your search for the problem there. Be aware, however, that the line reported may not be the cause of the issue. The actual problem may be higher up in the code due to a faulty initialization, or some other issue. When troubleshooting, always start looking at the line reported, and work your way back towards the top.

When testing a script, you may come across an Error Message Box like the following. Your first inclination may be to close the error without closer examination: resist it!



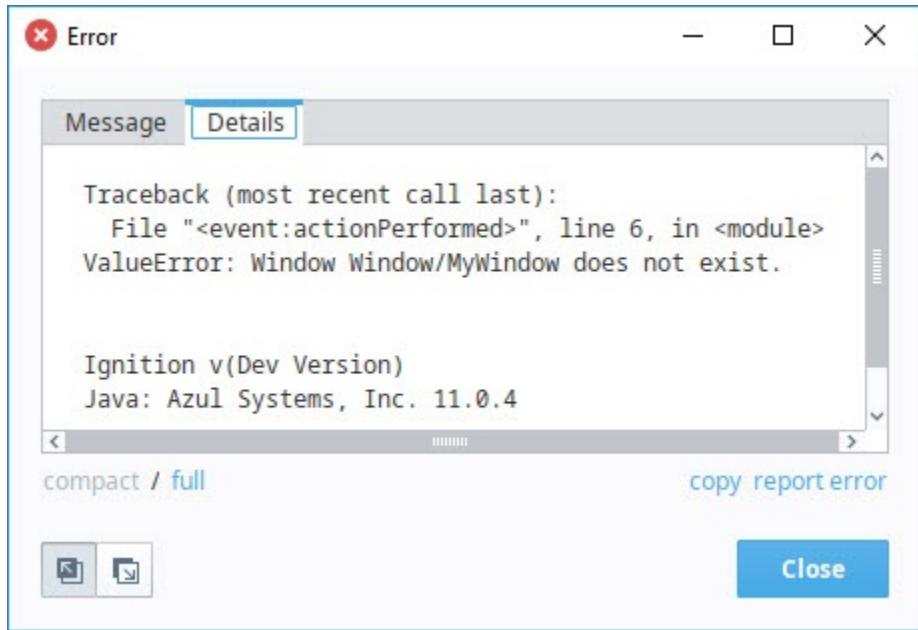
The error messages generated from failed script executions are incredibly helpful. The **Message** tab on the error describes the [Event Handler](#) that encountered the exception; in this case, the actionPerformed event handler. This is important in cases where multiple Event Handlers on the same component have scripts. Without knowing which Event Handler generated the script, we could waste time trying to troubleshoot the wrong script.

Additionally, it shows the name of the component, which is "NavButton_MyWindow". Again, this is helpful if multiple scripts from multiple components are triggering in quick succession. The **Message** tab will clearly point you towards the source of the exception. This is another reason to give meaningful names to your components: doing so makes the process of tracking an error much easier.

The Details tab has even more information, specifically the line number that the exception occurred on, as well as the error message. The message here states that the window at path "MyWindow" does not exist, so we can check the Project Browser to see if we simply mistyped the name of the window.

On this page ...

- [Error Message Box Overview](#)
- [Troubleshooting Errors Using the Error Message Box](#)
 - [Broken Example - Incorrect Attribute](#)
 - [Broken Example - Undefined](#)
 - [Broken Example - Type Error](#)



Troubleshooting Errors Using the Error Message Box

When we don't get our expected results from our script, as in the following examples, always read both the Message and Details tabs in the Error Message Box. They are pretty good about pointing you to the root cause of your error. From there, you can easily find and fix any errors in your script.

The following examples show how to troubleshoot some of these error messages from the Error Message Box. Keep in mind that every script is unique, but at least you'll become familiar with what some of the error messages mean, and gain a little insight of the troubleshooting process.

Broken Example - Incorrect Attribute

In the following script, when a Button is pressed a Message Box is supposed to open on a window and display "Hello World".

Python - Broken: Incorrect Attribute

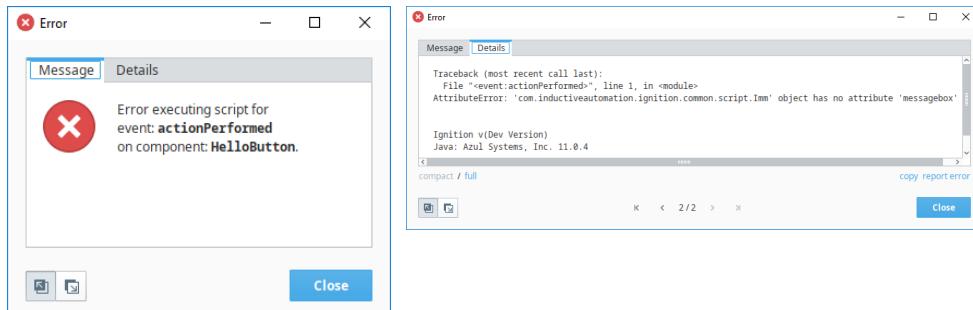
```
system.gui.messagebox( "Hello World" )
```

From the error message box, you know two important things right away:

- The script is using the actionPerformed event on the Hello Button.
- Line 1 displays the description of the error.

Check to be sure you used the correct attribute. Then check the spelling, case sensitive letters, spacing, and operators in your script. If you are familiar with Ignition's [Built-in Scripting Functions](#), you can probably spot the error immediately. If not, you might want to use the [autocomplete popup feature](#) to retype your 'messageBox' scripting function. This will automatically fix any syntax errors.

The error in this example is in the attribute name. It uses case sensitive letters (i.e., **messageBox**).



In this case, "messagebox" must be spelled with a capital "B" since we're trying to use Ignition's `system.gui.messageBox` function. We corrected the script by changing our code to the following:

Python - Corrected: Incorrect Attribute

```
system.gui.messageBox("Hello World")
```

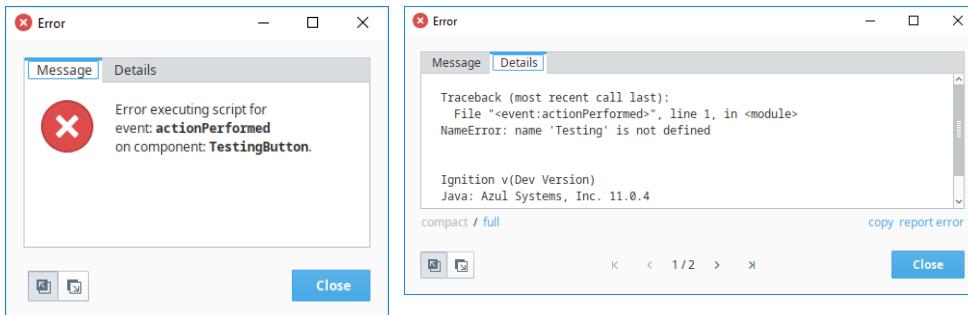
Broken Example - Undefined

Here is a similar script. When a Button is pressed, it's supposed to open a Message Box on a window and display "Testing".

Python - Broken: Undefined

```
system.gui.messageBox(Testing)
```

The Error Message Box displays an error on the actionPerformed script on the button component in line 1 with an **undefined** name. The Details tab is simply telling us that something is not defined in the script. You can either define a variable or create a string by putting "Testing" in quotes to correct the error.



For this example, we corrected the code by turning the argument passed to system.gui.messageBox into a string literal with quotation marks:

Python - Corrected: Undefined

```
system.gui.messageBox("Testing")
```

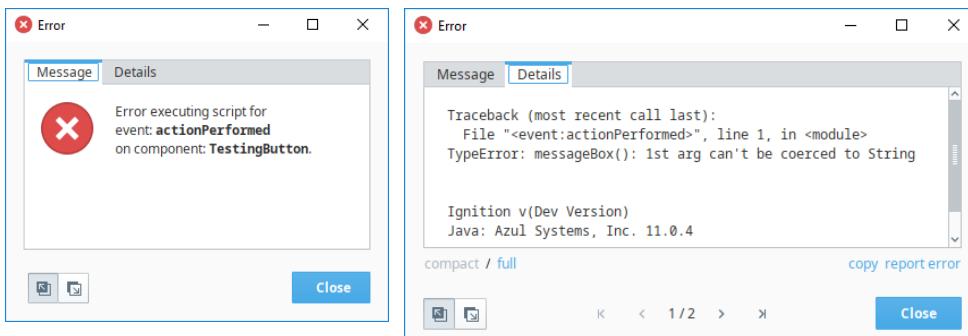
Broken Example - Type Error

Similar to the other scripts, this script opens a Message Box and is supposed to display "100" when a Button is pressed.

Python - Broken: TypeError

```
system.gui.messageBox(100)
```

Since this script failed to execute successfully, the Error Message Box popped up and displayed an error on the actionPerformed script on the button component in line 1 with a **TypeError**. This error tells us that the first parameter is expecting a string, not an integer.



For this example, we corrected the code by changing the argument passed to system.gui.messageBox into a string literal with quotation marks.

Python - Corrected: TypeError

```
system.gui.messageBox( "100" )
```

You can see from the examples above, that the Error Message Box provides quite a bit of information that points to the root cause of an error on an Event Handler script. As you learn from the Error Message Boxes about what these messages mean, you'll be able to quickly spot them in your script and quickly correct any errors you encounter.

Related Topics ...

- [Basic Python Troubleshooting](#)
- [Scripting in Ignition](#)
- [Getting Started with Scripting in Ignition](#)

Troubleshooting - Nothing Happened

Did it Work?

When testing a script, you may find yourself in a situation where the script appears to be running, but doesn't seem to work. Furthermore, you may not see any error messages stating there is a problem.

When a script doesn't perform to expectations and doesn't throw an exception, it can suggest a problem with the script's workflow. This page describes a couple of common scenarios.

Before You Begin...

Make **absolutely** certain that there isn't an Error Message Box hiding somewhere, otherwise you'll waste time applying the troubleshooting tips outlined below, when the real error message was hiding in the background behind another window the whole time.



If a button component on a window is running a script, the output will be displayed on the Designer/Client Output Console. You can check for any error messages by navigating to the Console. From the **Client** menu bar, go to **Help > Diagnostics** and select the **Console tab**. From the **Designer** menu bar, select **Tools > Console**.

Common Scenarios when Nothing Happens

An Important Line Is Missing

Some critical part of your code is missing or commented out. It could be something simple, like your code is supposed to increment a value, but you forgot to write the line that increments the value.

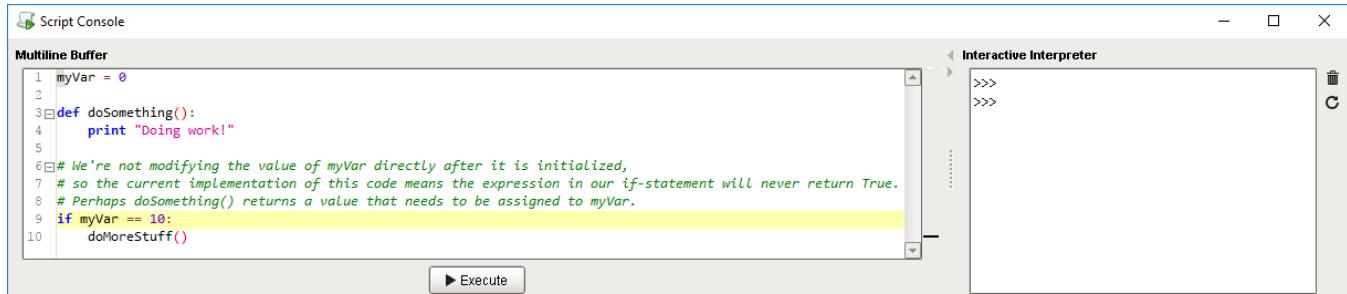
Pseudocode - Missing Code

```
myVar = 0

def doSomething():
    print "Doing work!"

# We're not modifying the value of myVar directly after it is initialized,
# so the current implementation of this code means the expression in our if-statement will never return True.
# Perhaps doSomething() returns a value that needs to be assigned to myVar.
if myVar == 10:
    doMoreStuff()
```

A quick way to test your code is by running your script in the **Script Console** before attaching it to a scripting event or specific component. You can see that running the code above in the Script Console that nothing happened, thus, suggesting an error.



On this page ...

- Did it Work?
 - Before You Begin...
- Common Scenarios when Nothing Happens
 - An Important Line Is Missing
 - The Script Is not Being Called
 - Important Lines Are Being Skipped
- Determining the Cause

The Script Is not Being Called

The script doesn't appear to be working because the mechanism that is supposed to be triggered hasn't been called. This can be caused by using the wrong event: (i.e., perhaps you placed the code on a Button's **propertyChange** event, when you meant to place it on **actionPerformed**).

Alternatively, perhaps you defined a function in your script, but you forgot to call the function.

Additionally, if you're testing the script in the Designer, make sure it is in **Preview Mode**. Event based scripts will not run in the Designer unless it is in Preview Mode.

Important Lines Are Being Skipped

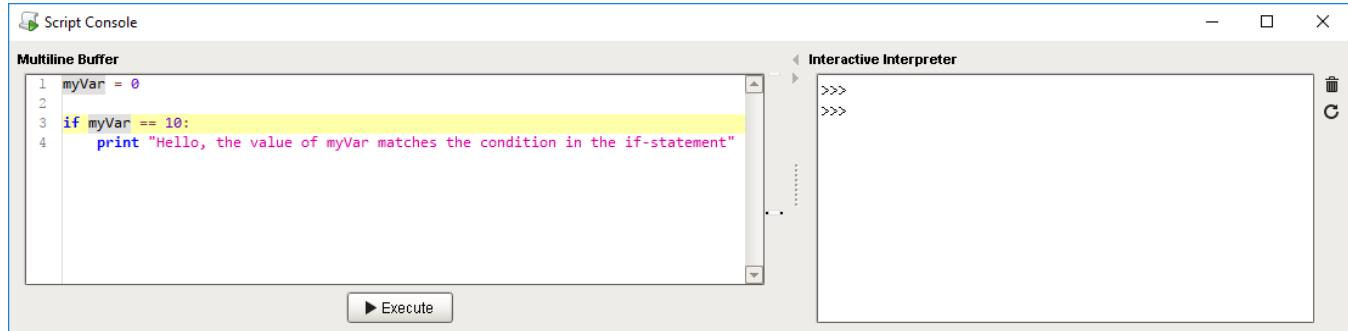
This can be caused by incorrect indentation, or a misconfigured condition in an if-statement's expression. For example, you used "==" when you meant to use "!=". Verify that you are using the correct operators in if-statements.

Pseudocode - Skipping Code

```
myVar = 0

if myVar == 10:
    print "Hello, the value of myVar matches the condition in the if-statement"
```

Test your code for any errors by running your script in the Script Console. You can see nothing happened by running the code above in the Script Editor, once again suggesting an error.



Determining the Cause

The easier step to take when troubleshooting your script is to start adding print statements to your code. From here, you can start piecing together what the code is doing:

- Print the value of variables to make sure they are coming in the way you expect.
- Place print statements at the start and end of your code. This allows you to determine when your script is being called, and when it finishes.
- Adding a print statement before and after an if-statement can show you the flow of your script went as expected, or was filtered out by the if-statement.

Related Topics ...

- [Script Console](#)
- [Basic Python Troubleshooting](#)
- [Reading Error Messages](#)

Troubleshooting Workflow

Troubleshooting a script is an iterative process. Since the script stops executing at the first error, we won't see if there are any other errors on later lines unless, of course, we spot them ourselves while writing the code. As a result, we will have to keep trying our script until it executes successfully.

It is important to understand that scripts always execute from top to bottom, and each line must complete before the next line may move on. In the event, our code returns an exception, we can assume that our troubleshooting process should always start at the line reported in the error, and then work up until we find the problem.

On this page, we will take an indepth look at a script with multiple problems, and work through each one. Note, that this is not a comprehensive list of all possible types of exceptions that could occur when writing a script, but instead, this page attempts to demonstrate the troubleshooting process.

To learn more about troubleshooting specific script errors, refer to the following sections:

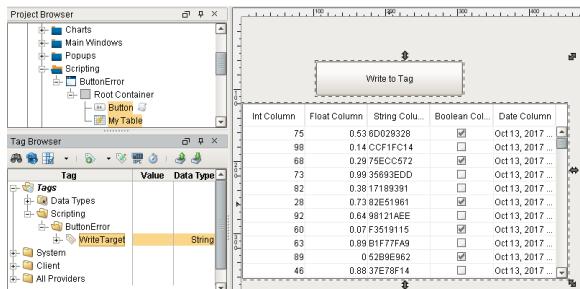
- [Reading Error Messages](#)
- [Troubleshooting - Nothing Happened](#)

Scenario Overview

Here we have a window with two components: a **Button** and **Power Table**. The purpose of the button is to find which cell in the Power Table the user selected, and write the value to a Tag.

On this page ...

- [Scenario Overview](#)
- [First Error - NoneType Object](#)
 - [Exception Error Explained](#)
 - [What Should We Look For?](#)
 - [Solution](#)
- [Second Error - Checking Attributes](#)
 - [Exception Error Explained](#)
 - [What Should We Look For?](#)
 - [Solution](#)
- [Third Error - ArrayIndexOutOfBoundsException](#)
 - [Exception Error Explained](#)
 - [What Should We Look For?](#)
 - [Solution](#)



The code on the button is listed below.

Python - Sample Code Block

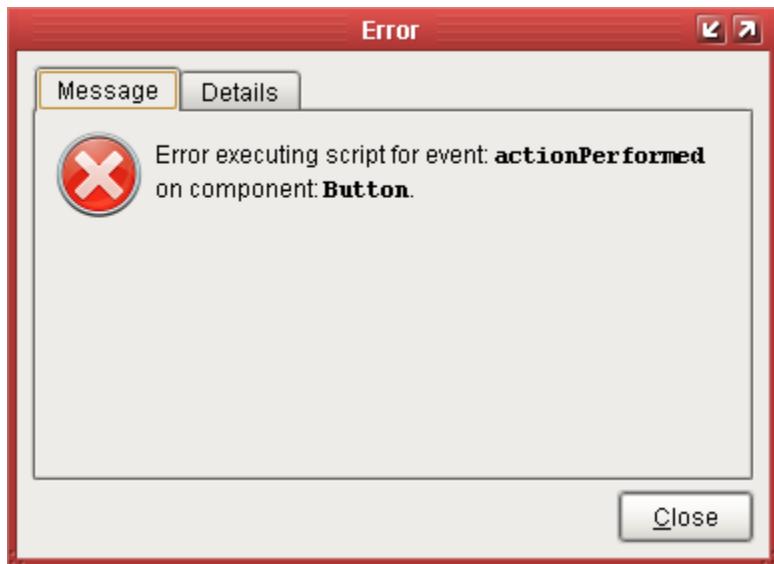
```
# Create a variable that references the Power Table
table = event.source.parent.getComponent('Power Table')

# Find the cell the user currently has selected, and store the value in a
variable
userSelectedValue = table.data.getValueAt(table.selectedrow, table.
selectedColumn)

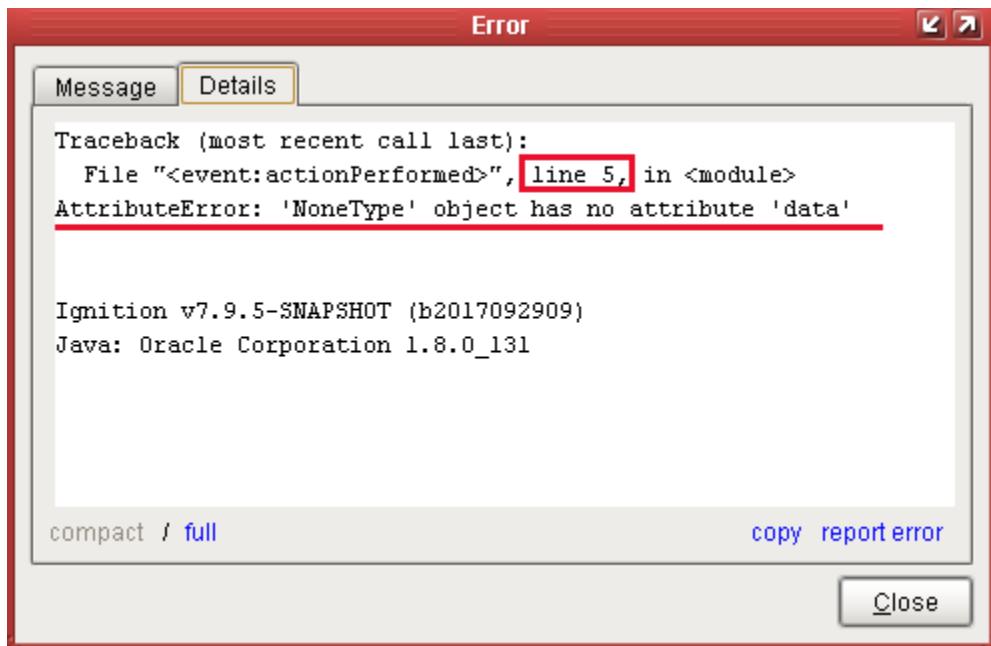
# Write the User Selected Value to a tag
system.tag.write("Scripting/ButtonError/WriteTarget", userSelectedValue)
```

First Error - NoneType Object

When the button is pressed, we are presented with an error. The **Message** tab in the Error Box describes the [Event Handler](#) that encountered the exception.



The **Details** tab tells us where to start looking, specifically the line number where the exception occurred as well as the error message. Line 5 is referenced, so the troubleshooting process should start there.



Exception Error Explained

When an error message refers to a **NoneType** object, that simply means a null, or None in Python. The word '**attribute**' is used to reference a property on an object. In this case, this means the script was trying to access the '**data**' property on nothing. The message is telling us that it couldn't find a property named '**data**' on a nothing, which is correct since **NoneType** objects don't have any properties named '**data**'.

If we were to reword this message to something a bit more straightforward, it would say the following: "I tried to access the 'data' property on 'null', but it doesn't have a property by that name."

What Should We Look For?

If your exception is referring to a **NoneType** object, then some line of code probably tried to reference something else (i.e., Tag value, property, another variable, etc.), but couldn't find anything at the location you specified. It is also helpful to note what attribute was mentioned in the error, which was **data**. If we look at our code, we see the following on line 5.

Python - Line 5

```
userSelectedValue = table.data.getValueAt(table.selectedrow, table.selectedColumn)
```

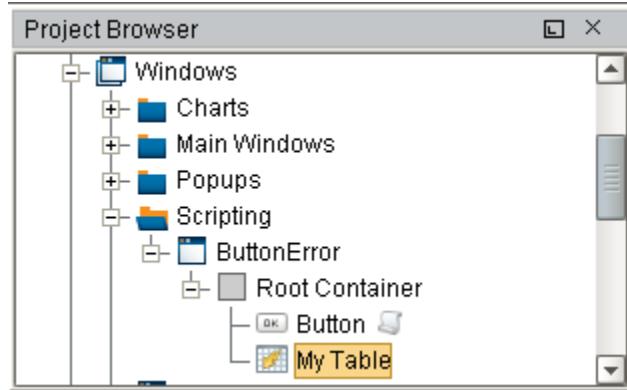
Based on this information, we should take a closer look at the table variable since our code is specifying '**data**' as an attribute of the table. The error states that it could not find the data attribute on a **NoneType** object, but our code only references data in regards to the table object, so maybe there is something wrong with how that table object was referenced or initialized. If we look further up in our code, we see that **table** was initialized on line 2:

Python - Line 2

```
table = event.source.parent.getComponent('Power Table')
```

Since this script was placed on the Button component's **actionPerformed** event, we can trace back to the source of the issue:

event = actionPerformed event
source = Button this script is placed on
parent = Container that the Button was placed in. Based on the Button's position in the window, this appears to be the Root Container.
getComponent('Power Table') = Should return a reference to a component named '**Power Table**' directly inside of the Root Container, assuming one exists. If we take a look at the Project Browser, we see the following:



There is no component named '**Power Table**' in the Root Container. However, there is a Power Table component named '**My Table**', so it appears someone renamed the component, which caused our script to initialize the variable **table** as a '**NoneType**' instead of a reference to a component.

Solution

To fix this issue, we can simply update our code to use the new name of the Power Table, which is '**My Table**'. We can type this in manually, but since it is case sensitive, it must match exactly. It may be easier to find the **Name** property on the component, copy the name to the system clipboard (**Ctrl-C**), and paste it (**Ctrl-V**) into the script. Our code now looks like the following:

Python - Updated Code Block

```
# Create a variable that references the Power Table
table = event.source.parent.getComponent('My Table')

# Find the cell the user currently has selected, and store the value in a variable
userSelectedValue = table.data.getValueAt(table.selectedrow, table.selectedColumn)

# Write the User Selected Value to a tag
system.tag.write("Scripting/ButtonError/WriteTarget", userSelectedValue)
```

Second Error - Checking Attributes

If we try our button again, we get the following exception error.



Exception Error Explained

Like the last exception, the problem has to do with our code trying to access an attribute (property) on an object, but the given attribute doesn't exist on the object.

What Should We Look For?

Again, the exception refers to line 5, but note that the error mentions a different object and attribute this time. Without knowing what a 'com.inductiveautomation.factorypmi.application.com' object is, we can tell that our script is trying to reference an attribute named 'selectedrow'. This gives us a starting point for troubleshooting the error. If we check line 5, we see the following:

Python - Line 5

```
userSelectedValue = table.data.getValueAt(table.selectedrow, table.selectedColumn)
```

We see `table.selectedrow` as the first parameter being passed to the `getValueAt()` function. We checked into the `table` variable earlier, so this doesn't necessarily mean that the variable is the problem, although there is no harm in double checking, especially if you have not personally verified the variable. Assuming that the `table` variable is correct, let's check the attribute. We can check the Property Editor in the Designer to find a matching property, or head over to the Power Table component page in the manual and check the property reference. We can see that there is a **Selected Row** property on the Power Table, but it should be spelled "selectedRow" with a capital "R".

Property Editor

Non-Contiguous Selection	<input type="checkbox"/> false
Row Dragging Enabled	<input type="checkbox"/> false
Row Height	16
Row Selection Allowed	<input checked="" type="checkbox"/> true
Selected Column	-1
Selected Row	-1
Selection Background	250,214,138

Selected Row
The index of the first selected row, or -1 if none.
Property Type: int
Property Scripting Name: 'selectedRow'

We also know that the `getValueAt()` function takes a row index as the first parameter, so it would make sense that our code has a typo.

Solution

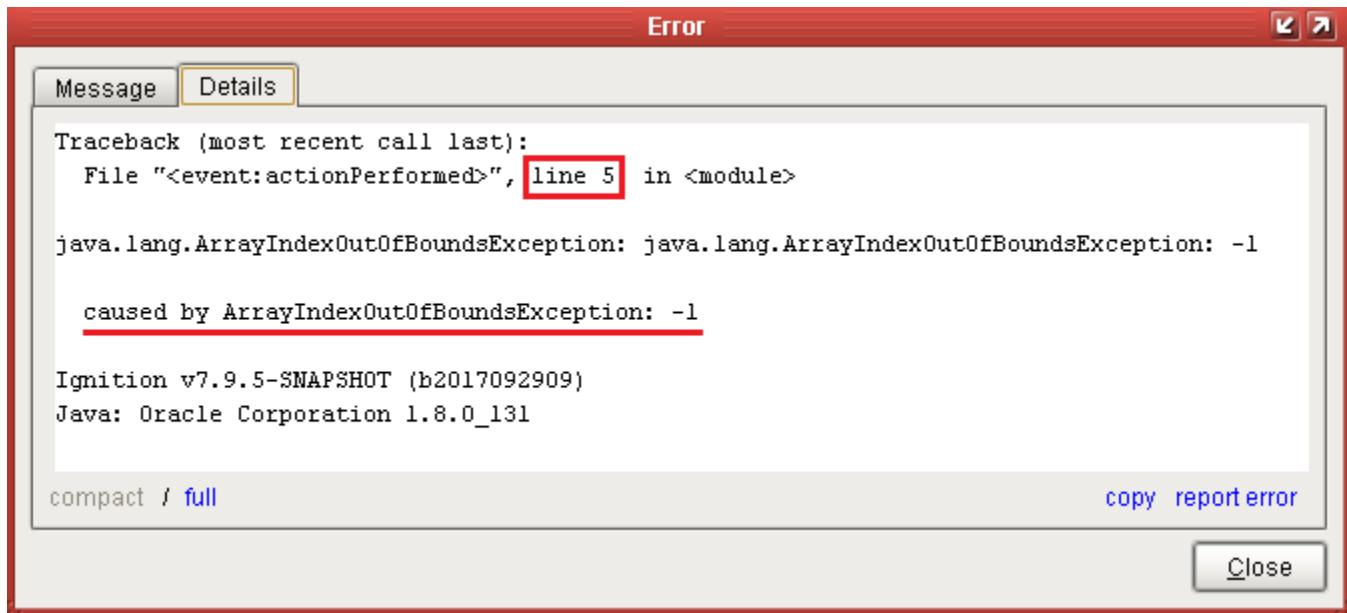
We should fix the typo in the attribute name.

Python - Line 5 Updated

```
userSelectedValue = table.data.getValueAt(table.selectedRow, table.selectedColumn)
```

Third Error - ArrayIndexOutOfBoundsException

After making our most recent change, the button works great! Users are able to select a cell, and write to the Tag. The one exception, is if the user does not have a cell selected before clicking the button. In this case, the following error occurs:



Exception Error Explained

Our code attempted to look up something in a collection of some sort (i.e., sequence, dataset, etc.) by index, but was told to check index -1, which is out of bounds. Indexes in Ignition are typically zero-based, meaning they start at 0, and increment from there, so an index of -1 doesn't exist. Commonly, when a property in Ignition refers to index -1, that means nothing is selected, so our code failed because nothing was selected.

What Should We Look For?

Fortunately, the exception told us what the issue is: the user did not have anything selected before pressing the button, so we don't have to search further.

Solution

The solution to this issue is more open ended, as there are many ways to suggest to our user that they need to select a cell in the table first. The **Selected Row** and **Selected Column** properties on the Power Table will have -1 values if none of the cells are selected, so we could use an 'if' statement checking the values of one of those properties. Our new code could look like the following:

Python - Code Block Update #3

```
# Create a variable that references the Power Table
table = event.source.parent.getComponent('My Table')

# Make sure a cell in the Power Table is selected first
if table.selectedRow != -1:
    # Find the cell the user currently has selected, and store the value in a variable
    userSelectedValue = table.data.getValueAt(table.selectedRow, table.selectedColumn)
```

```
# Write the User Selected Value to a tag
system.tag.write("Scripting/ButtonError/WriteTarget", userSelectedValue)

# If a cell isn't selected, then let the user know
else:
    system.gui.messageBox("Please select a cell in the table first!")
```

Related Topics ...

- [Reading Error Messages](#)
- [Troubleshooting - Nothing Happened](#)

Scripting Vs. SQL Vs. Expressions

Expression Language & SQL Queries vs Scripting

There are three major languages in Ignition, the [Expression language](#), the [SQL Queries](#), and [Python Scripting](#). It is important to understand the differences between the three and to know where each is used. Scripting is used in the event handlers that are available all over Ignition, but Expressions and SQL are in the [Property Binding](#) locations shown here.

Spot the Difference - Comments

When starting out with Ignition, it can be difficult to know which syntax you should be using for a particular area of text. One little trick that may help is to be familiar with how each language handles comments, and then utilize the **Ctrl + /** command which automatically adds the characters that comments out a line of code in the language you are typing in. Once you remember to use **Ctrl + /**, you simply need to be familiar with the characters that each language uses.

Expression Language - Comment

```
// The Expression Language uses double forward slashes.
```

SQL - Comment

```
-- Areas in Ignition that accept SQL syntax use double dashes.
```

Python - Comment

```
# Python uses the pound/number/hash character.
```

In addition to comments, interfaces for each of the languages usually contain other signs or reminders about the language. These will be covered in their respective sections on this page.

Python Scripts

Overview

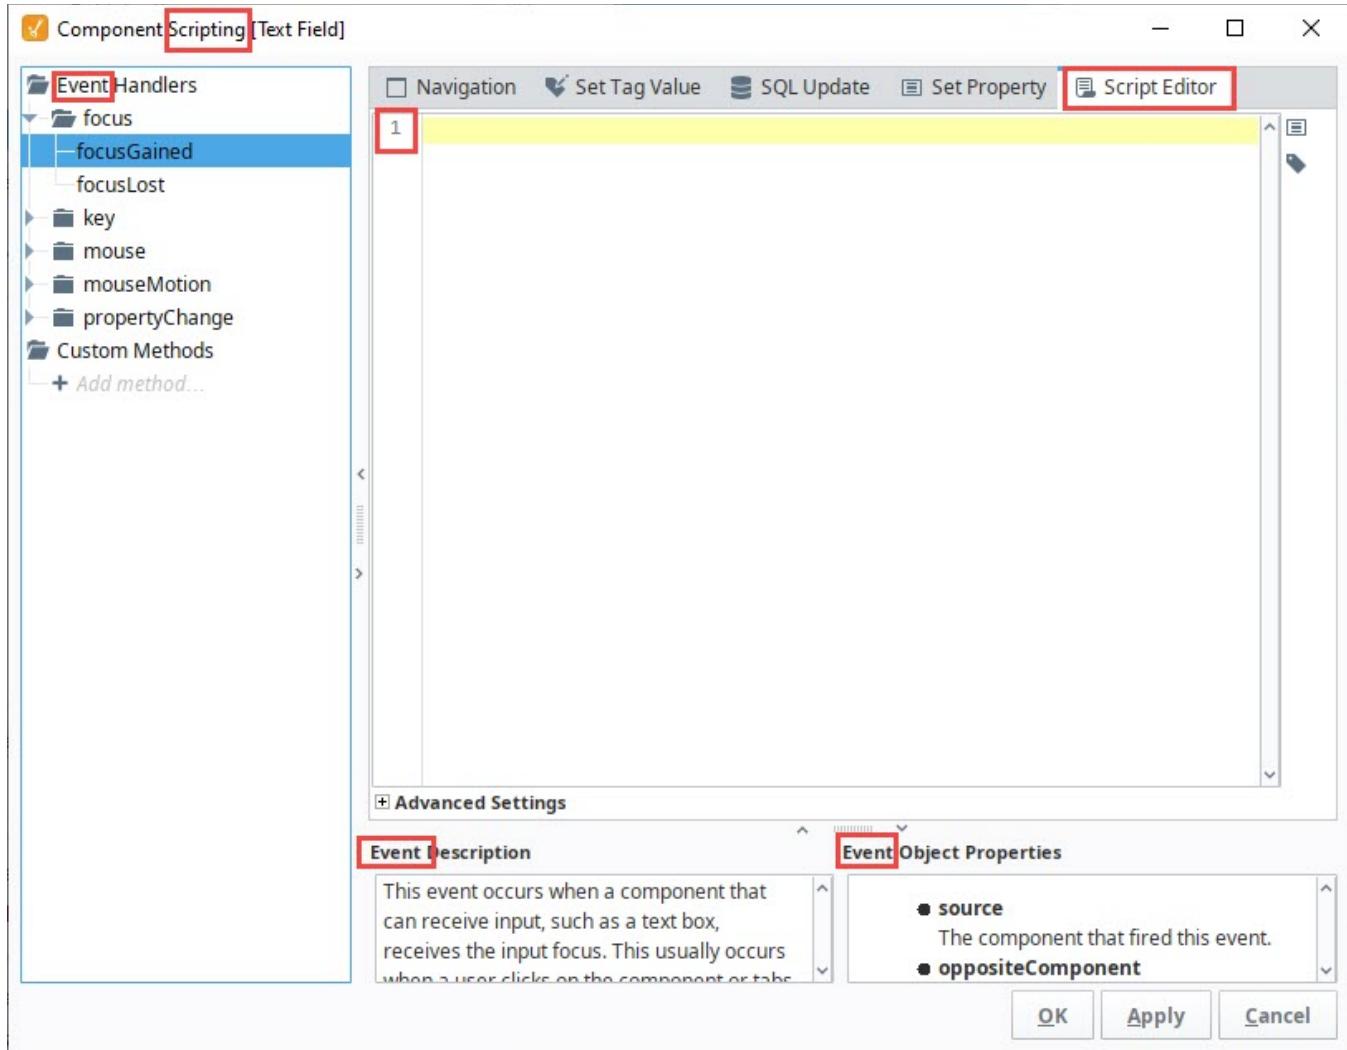
Python is featured prominently throughout Ignition and many different resources can contain a Python Script.

How Can I Tell If I'm Writing a Python Script?

Python Scripts typically use the words **Script** or **Event** in the interface. Additionally, Python requires a particular Event to be selected so if you see event Handlers on the left, you know you are looking at a Python script.

On this page ...

- [Expression Language & SQL Queries vs Scripting](#)
- [Spot the Difference - Comments](#)
- [Python Scripts](#)
 - [Overview](#)
 - [How Can I Tell If I'm Writing a Python Script?](#)
 - [Where Are Python Scripts Used?](#)
- [SQL Queries](#)
 - [Overview](#)
 - [How Can I Tell If I Should Be Using SQL Syntax?](#)
 - [Where Is SQL Used in Ignition?](#)
 - [SQL in Python](#)
- [Expression Language](#)
 - [Overview](#)
 - [How Can I Tell If I'm Writing an Expression?](#)
 - [Where Is The Expression Language Used?](#)



Where Are Python Scripts Used?

Python Scripts are used all throughout Ignition. Some resources, such as components, even have multiple places to write scripts! Below are some common locations for scripts:

- Scripting in Vision
- Extension Functions
- Tag Event Scripts
- Client Event Scripts
- Gateway Event Scripts

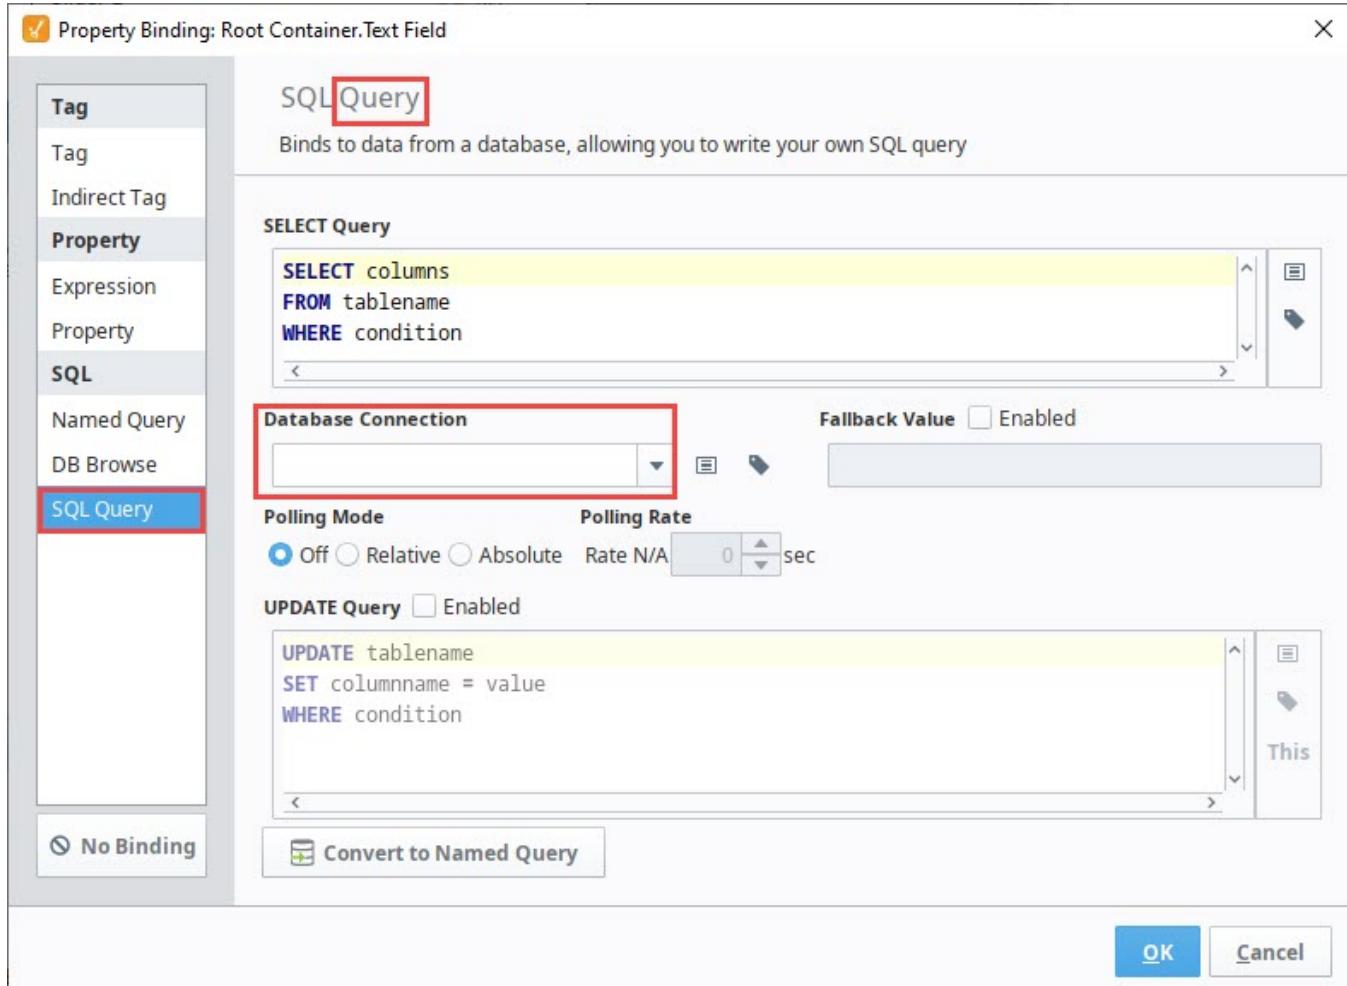
SQL Queries

Overview

The SQL language is used for selecting information from a database. It can be used in a variety of ways, but most of them will either make modifications to the database or return a set of values (with a few notable exceptions like Stored Procedures). The majority of users will be returning large chunks of data into a Table or Report in Ignition. This means a complete dataset will be returned based on a user selection, a time range, or any combination of factors.

How Can I Tell If I Should Be Using SQL Syntax?

Typically, the words **Query** and **Database** appear in the interface somewhere. Additionally, there is usually a way to specify a **Database Connection**.



Where Is SQL Used in Ignition?

Below is a reference of the most common areas in Ignition where SQL queries may be used.

- [SQL Query Bindings](#)
- [Named Queries](#)
- [Database Query Browser](#)
- Reporting Data Source: [SQL Query Data Source](#) and [Basic SQL Query](#)
- [Query Tags](#)
- Python Scripts - See the SQL in Python header below for more details

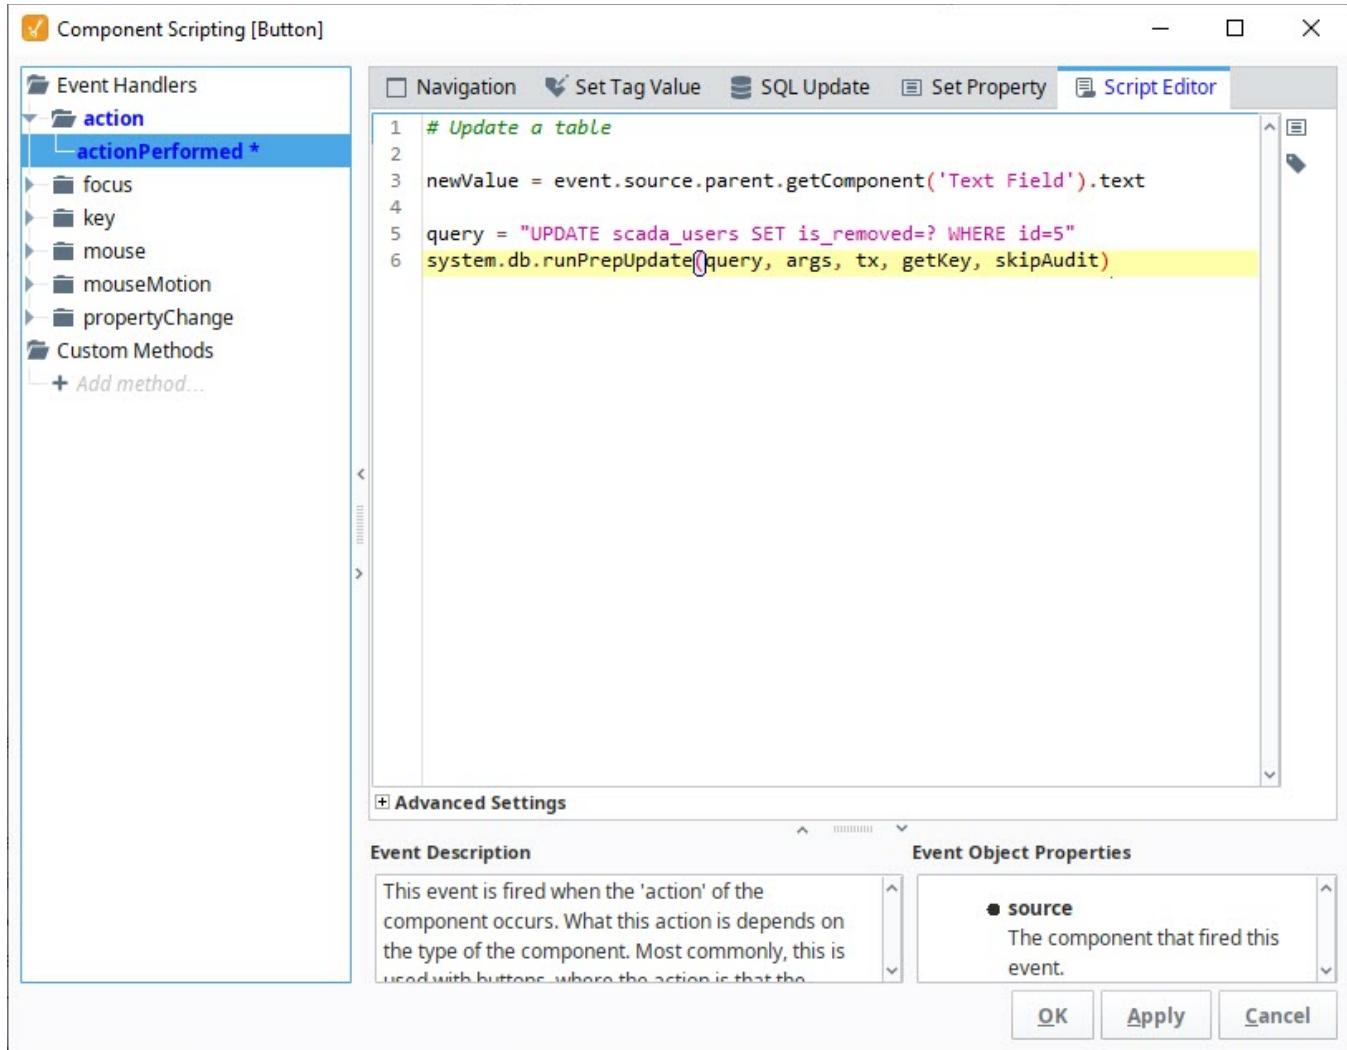
SQL in Python

SQL queries can be called from a **Python** script. There are several system functions in Ignition that allow a script to run a query against the database, such as `system.db.runPrepQuery`. This is a more advanced technique, as you need to adhere to both language's syntax. Furthermore, when typing a SQL query in a Python script interface, the syntax highlighting can not help with the SQL portions. The syntax highlighting in a Scripting Window is only looking for Python syntax, not SQL.

In cases where you plan on calling a SQL query from a Python script, it is highly recommended to write the query in the **Database Query Browser** first (substituting parameters with static values for testing purposes), and then move the query over to the script once the query executes successfully on its own. This approach can save you some time troubleshooting, as there will be less ambiguity when an error occurs since you know the query runs.

Below we see an example of calling a SQL query in a script. Line 5 creates a variable called "query", and assigns it a string consisting of a prepared statement (using SQL). The query is then executed with the `system.db.runPrepUpdate` function.

For more examples of using a query in a Python Script, check out the `system.db.*` functions.



Expression Language

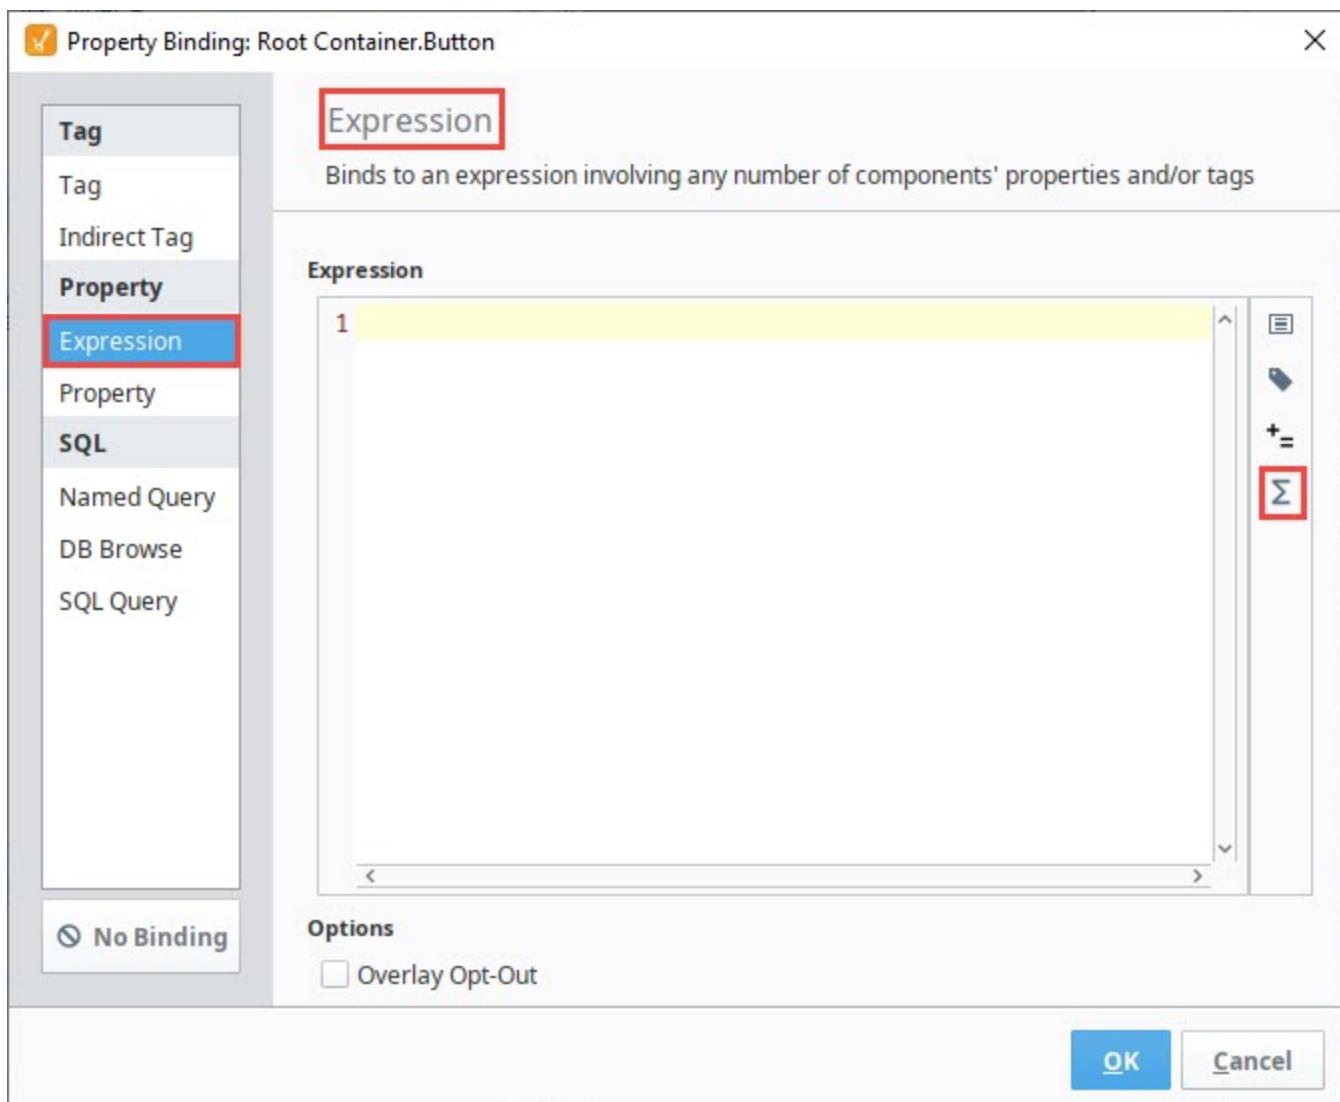
Overview

The Expression Language is a simple programming language that we invented (very similar to many other existing expression languages), and is different from the Python scripting you will find in Ignition. The expression language is a very simple kind of language where everything is an expression - which is a single piece of code that returns a value. This means that there are no statements, and no variables, just operators, literals, and functions.

The most common expression language that most people are familiar with is the one found in Microsoft Excel. You can have Excel calculate a cell's value dynamically by typing an expression like `=SUM(C5:C10)`. Our expression language has similar functionality, but different syntax. It is mainly used to define dynamic values for **Tags** and **component properties**. In Ignition's Expression Language, you will use component properties and functions like `if({Root Container.type}="Type C",True,False)`.

How Can I Tell If I'm Writing an Expression?

Typically interfaces that expect the Expression Language use the word **Expression**. Additionally, you'll commonly find the **Expression Function** button  towards the right side of the text area.



Where Is The Expression Language Used?

Below is a list of resources that commonly utilize the Expression Language

- [Expression Bindings](#) on Property Bindings and [Alarm Bindings](#)
- [Expression Tags and Derived Tags](#)
- Several [Alarm Pipeline Blocks](#), including the [Expression Block](#), [Switch Block](#), and [Notification Blocks](#).
- [Expression Items](#) in Transaction Groups
- [Parameters on Reports](#)
- [SFC Transitions](#)

Related Topics ...

- [Expression Language and Syntax](#)
- [SQL in Ignition](#)
- [Scripting](#)
- [Python Scripting](#)

Scripting Examples

This section contains examples for items we've identified as "common" tasks: undertakings that many users are looking to utilize when first starting out with a specific module or feature in Ignition. Additionally, this section aims to demystify some of the more complex or abstract tasks that our users may encounter.

The examples in this section are self-contained explanations that may touch upon many other areas of Ignition. While they are typically focused on a single goal or end result, they can easily be expanded or modified after the fact. In essence, they serve as a great starting point for users new to Ignition, as well as experienced users that need to get acquainted with a new or unfamiliar feature.

Below is a list of common tasks related to this section of the manual.

On this page ...

- [Reading and Writing to Tags](#)
- [Importing and Exporting a CSV](#)
- [Reading a Cell from a Table](#)
- [Export Tag Historian to CSV](#)

Reading and Writing to Tags

There are simple interfaces in Ignition that allow you to easily write to a Tag on some Event, and Reading can be as simple as creating a Tag binding. Sometimes, however, the built-in approaches can be too simplistic or limiting. The [Reading and Writing to Tags](#) page details how to better interact with Tags from scripting.

Importing and Exporting a CSV

CSV files are used by many software programs to export data so that other systems may utilize the information contained within. The [Importing and Exporting a CSV](#) page demonstrates how to both import a CSV into Ignition, as well as export data from Ignition into a CSV.

Reading a Cell from a Table

Once data is populated into a Table component, it's useful to know how to read and extract a data from a cell in a Table, particularly if users can select a row in a Table. The [Read a Cell from a Table](#) page has some good examples for retrieving data from a single cell and multiple cells in a Table.

Export Tag Historian to CSV

The Tag Historian Data can be great, but it's sometimes difficult to view it outside of Ignition. The [Export Tag Historian to CSV](#) page details how to pull out a subset of Tag history data and export it to a CSV file.

Related Topics ...

- [Basic Python Troubleshooting](#)

In This Section ...

Location Based Vision Startup Scripts

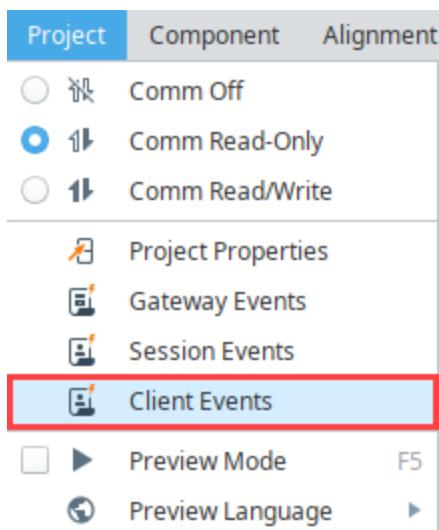
Client Startup Scripts

Security can be created on a Vision Client using a Startup Script that checks certain user information and then customizes the level of access within the project based on that information. We can pull user information from places like [System Tags](#) for Vision Clients or [Session Properties](#) for Perspective Sessions. Using that information, the script can then customize which windows or views the user sees, provide some additional information to the user, or even prevent the user from accessing the project completely. An additional example is located [here](#).

Location Based Vision Client Restrictions

Each Client has access to distinct hostname and IP addresses. These can be used in Client startup scripts that evaluate the Client's information and compares it to a list of acceptable host names or IP addresses. This information can come from a database or a set of Tags. In this example, we will prevent access to the client if the user is logging in from an incorrect location.

1. On the Project tab, choose Client Events.



2. Select the Startup icon.
3. Add the following script:

```
# Prevent access if user is not logging in from the correct location

# Grab the hostname that the user is logging in from
hostname = system.tag.readBlocking(["[System]Client/Network/Hostname"])[0].value

if hostname != "Machine A Computer":
    # If the user logs in on a computer that is not called Machine A Computer,
    # inform them that the project can only be accessed from the Machine A computer,
    # and then log them out.
    system.gui.messageBox("This project can only be accessed from the 'Machine A Computer'.")
    system.util.exit(1)
```

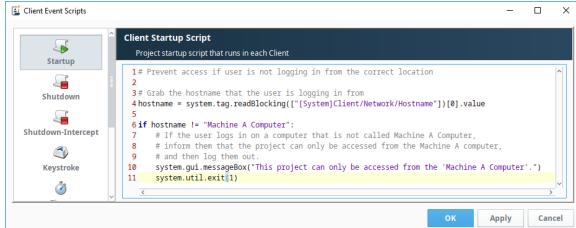
On this page ...

- [Client Startup Scripts](#)
- [Location Based Vision Client Restrictions](#)
- [Location Based Startup Display](#)

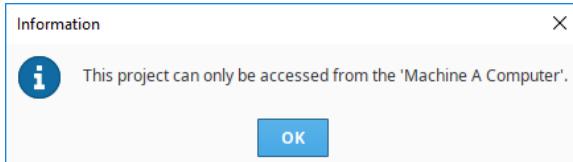


Custom Security

[Watch the Video](#)



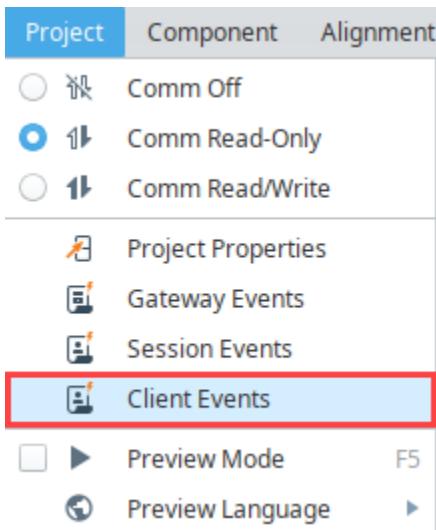
4. Click **OK** to save the script.
5. Save your project and launch the client. You will get a popup that informs you that the project must be launched using the Machine A Computer.



Location Based Startup Display

Another common use for startup scripts on Vision Clients is to open a specific window depending on the location of the log in. In this example, if the user logs in on Machine A Computer, then the Machine A Details window will be displayed. If they log in on a different computer, the Overview window is displayed.

1. On the Project tab, choose Client Events.



2. Select the icon.
3. Add the following script:

```

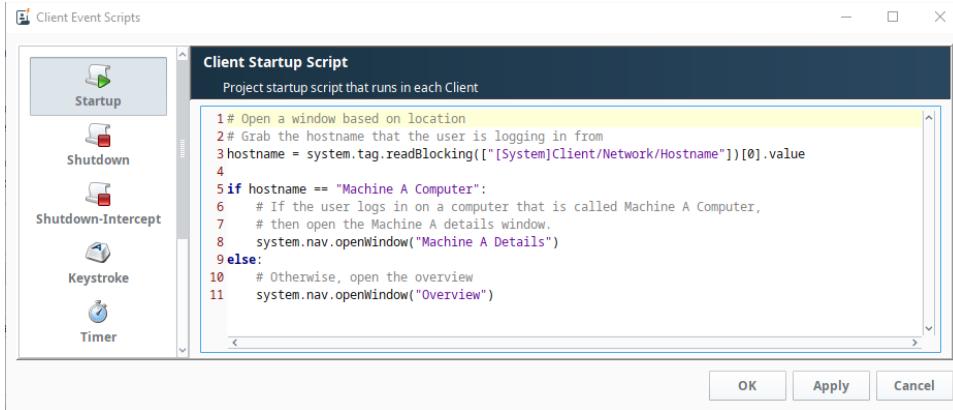
# Open a window based on location

# Grab the hostname that the user is logging in from
hostname = system.tag.readBlocking("[System]Client/Network/Hostname")[0].value

if hostname == "Machine A Computer":
    # If the user logs in on a computer that is called Machine A Computer,
    # then open the Machine A details window.
    system.nav.openWindow("Machine A Details")

else:
    # Otherwise, open the overview
    system.nav.openWindow("Overview")

```



4. Click **OK** to save the script.

Note: Any windows that are set to [Open on Startup](#) will still open in addition to the window specified in the startup script. You should disable any main windows from opening on startup if using this method.

5. Save your project.

Related Topics ...

- [Client Tags for Indirection](#)

Reading and Writing to Tags

In many cases, the binding system is the most appropriate way to display a Tag value on the screen. However, you may wish to access a Tag's value in a script. Using the system functions, you can read from a Tag and write to a Tag in Ignition.

Script Builder

If you simply need to write to value to a Tag from a script, you can use the Set Tag Value tab of the Script Editor. To learn more, refer to the sections on [Script Builders](#) for more information. However, if you need to do more than just send a single write, see the section below.

Know Your Scope

Whenever a script from a Shared resource (such an Alarm Pipeline) attempts to interact with a Tag, the script must specify the Tag Provider, otherwise, the script could return an exception. Tag Providers are always included at the start of the Tag Path and look like the following: [tagProvider]

When interacting with a Tag from the Project scope (such as a script on a component) you may optionally include the Tag provider. If omitted, the Project default Tag provider will be used.

Pseudocode - Tag Read Scope

```
# Reading a Tag without the Tag Provider. You would not want to use this
# from the Shared scope.
system.tag.readBlocking([ "My/Tag/Path" ])

# Reading from the same Tag, but specifying the Tag Provider. This format
# may be used safely in either scope.
system.tag.readBlocking([ "[default]My/Tag/Path" ])
```

On this page ...

- [Script Builder](#)
- [Know Your Scope](#)
- [Manual Tag Reads](#)
 - [Reading from a Single Tag](#)
 - [Reading from Multiple Tags](#)
- [Relative Tag Paths](#)
- [Manual Tag Writes](#)
 - [Writing to a Single Tag](#)
 - [Writing to Multiple Tags](#)



INDUCTIVE
UNIVERSITY

Reading and Writing Tags

[Watch the Video](#)

Manual Tag Reads

Reading from a Single Tag

Reading a Tag from a script is accomplished with the `system.tag.readBlocking()` function, which requires the Tag path you wish to read from. This function returns a 'Qualified Value'; this object is more than just the value. A Qualified Value is the Tag value that has three attributes; Value, Quality, and TimeStamp.

The function normally expects a list of Tag paths, but can be used with a single Tag: simply provide a list of only a single Tag path.

When handling the results of the read, the results will be returned in a list of qualified values, so you'll need to specify which qualified value you're interested in via Python slicing, even if there is only a single qualified value in the list.

Pseudocode - Reading Tag Attributes

```
# get the Tag value
value = system.tag.readBlocking([ "tagPath" ])[0].value

# get all three attributes
tag = system.tag.readBlocking([ "tagPath" ])[0]
value = tag.value
quality = tag.quality
timestamp= tag.timestamp
```

Reading from Multiple Tags

The `system.tag.readBlocking()` function can easily be used to read multiple Tags in a single call:

Python - Reading Multiple Tags and Printing the Values

```

# Create a List of Tag Paths to read
paths = ["Scripting/Tags/Alarm_1", "Scripting/Tags/Alarm_2", "Scripting/Tags/Alarm_3"]

# Read the Tags, and store the complex results in a variable
values = system.tag.readBlocking(paths)

# For each Tag Path, iterate through our results...
for index in range(len(paths)):

    # ...and do something with the individual values
    print values[index].value

```

Relative Tag Paths

Like elsewhere in Ignition, relative paths may be used from within a script. This is especially useful when writing a Tag Event script inside of a UDT, as you can specify relative members in the same UDT with "[.]".

Pseudocode - Reading with Relative Tag Paths

```

# Assuming a UDT with two sub-members, a script from one can read from one member to the other using the
# following code:
system.tag.readBlocking(["[.]otherMember"]).value

```

Manual Tag Writes

Writing to a Single Tag

Much like reading, there is a function you can use to write to Tags: `system.tag.writeBlocking`. It requires a list of Tag paths, as well as a list of values to write to those Tags.

Pseudocode - Tag Write

```
system.tag.writeBlocking(["tagPath"], ["Hello World"])
```

Writing to Multiple Tags

The `system.tag.writeBlocking` function can also write to multiple Tags, again by providing multiple paths and values.

Python - Multiple Tag Writes

```

# Create a List of Tag Paths to write to
paths = ["Scripting/Tags/Alarm_Setpoint_1", "Scripting/Tags/Alarm_Setpoint_2"]

# Create a List of values to write
values = [72, 72]

# Send off the write requests
system.tag.writeBlocking(paths, values)

```

Related Topics ...

- [Tags](#)

Exporting and Importing a CSV

A CSV file, (comma separated values) is one of the most simple structured formats used for exporting and importing datasets. It is a convenient and flexible way to edit and share data across applications. Ignition has a built-in function to convert a dataset to CSV data called `system.dataset.toCSV`. You can even convert the contents of a CSV to a script and move it to an Ignition component, such as a Power Table.

This section contains examples for exporting and importing data to a CSV as well as converting the contents of a CSV to a script.



The examples on this page make use of Vision components, and run in the Vision Client scope.

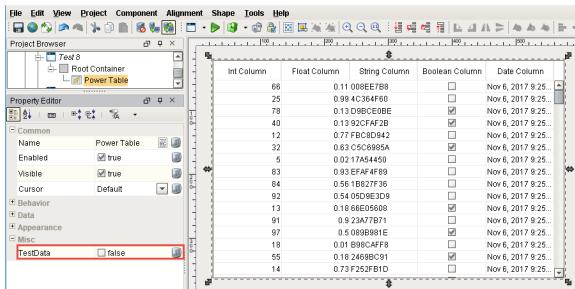
On this page ...

- [Exporting Data to a CSV](#)
- [Importing Data from a CSV](#)
- [Converting the Data into a Dataset](#)
 - [Calling the `system.dataset.fromCSV` Function](#)
 - [Calling the `csv.reader` Function](#)

Exporting Data to a CSV

You can export a dataset from a query or table to a CSV file.

1. Identify the dataset that you want to export to a CSV file. In this case, we can generate some data on a Power Table and export that data. Drag a **Power Table** component on to your window and toggle its **TestData** property to generate some data.



2. Drag a **Button** component on the window, and double click on the Button to open the **Component Scripting** window.
3. Next, let's add our script on the Button component's actionPerformed event. Select the **actionPerformed** event, and click on the **Script Editor** tab.
4. Copy the contents from one of the examples below, and paste the contents to the **Script Editor**. Notice the `system.dataset.exportCSV` scripting function is used in the first example to export the dataset to a CSV file:

Python - Hard Coded Filepath

```
# Create a variable that references our Power Table. You could
# modify this part
# of the example to point to a different component in the window.
component = event.source.parent.getComponent('Power Table')

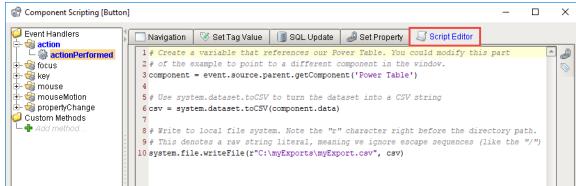
# Use system.dataset.toCSV to turn the dataset into a CSV string
csv = system.dataset.toCSV(component.data)

# Write to local file system. Note the "r" character right before
# the directory path.
# This denotes a raw string literal, meaning we ignore escape
# sequences (like the "/").
system.writeFile(r"C:\myExports\myExport.csv", csv)
```



Exporting Data to CSV

[Watch the Video](#)



5. Instead of hardcoding the path as we did in the above example, we could ask the user to select a directory on the local system with [system.file.saveFile](#):

Python - User Selected Directory

```

# Create a variable that references our Power Table. You could
modify this part
# of the example to point to a different component in the window.
component = event.source.parent.getComponent('Power Table')

# Use system.dataset.toCSV to turn the dataset into a CSV string.
csv = system.dataset.toCSV(component.data)

# Use system.file.writeFile to have the user find a directory to
write to.
filePath = system.file.writeFile("myExport.csv", "csv", "Comma
Separated Values")

# We can check the value of filePath to make sure the user picked a
path before
# attempting to write.
if filePath:
    system.file.writeFile(filePath, csv)

```

6. To test your scripts, put the Designer in **Preview Mode**, and press the **Button**. Open your **myExport.csv** file and check your data.

	A	B	C	D	E	F
10	84	0.56352293	1B827F36	FALSE	11/6/2017 9:25	
11	92	0.53782463	05D9E3D9	FALSE	11/6/2017 9:25	
12	13	0.17691666	66E05608	TRUE	11/6/2017 9:25	
13	91	0.90138406	23A77B71	FALSE	11/6/2017 9:25	
14	97	0.5005872	0898981E	TRUE	11/6/2017 9:25	
15	18	0.010637939	B98CAF8	FALSE	11/6/2017 9:25	
16	55	0.18030751	2469BC91	TRUE	11/6/2017 9:25	
17	14	0.7256514	F252FB1D	FALSE	11/6/2017 9:25	
18	15	0.2902785	197496C1	TRUE	11/6/2017 9:25	
19	99	0.8383941	916B8310	FALSE	11/6/2017 9:25	
20	38	0.03875667	C8081AD1	FALSE	11/6/2017 9:25	

Importing Data from a CSV

There are several ways to import data from a CSV file. First, we could use [system.file.readFileAsString](#) to read the entire file as a string. Note, that this will read the file as is, meaning "`\n`" can be used to denote new lines.

Python - Using [system.file.readFileAsString\(\)](#)

```

# Ask the user to find the CSV in the local file system.
path = system.file.openFile("csv")

# Use readFileAsString to read the contents of the file as a string.
# This string will be the parameter we pass to fromCSV below
stringData = system.file.readFileAsString(path)

```

```

# Split stringData into a List of strings, delimited by the new line character
stringData = stringData.split("\n")

# Iterate through the list, and do something with each line
for i in range(len(stringData)):

    # We're printing the row here, but you could do something more useful with the data.
    print stringData[i]

```

Alternatively, Python's CSV Library could be used to read in the contents of a CSV. In some cases, this is the easier approach, as the reader object is ready to be iterated over. Note, that this approach does read in each row as a List of strings:

Python - Using csv.reader()

```

# Import Python's built-in csv library
import csv

# Ask the user to find the CSV in the local file system.
path = system.file.openFile("csv")

# Create a reader object that will iterate over the lines of a CSV.
# We're using Python's built-in open() function to open the file.
file = open(path)
csvData = csv.reader(file)

# Iterate through the reader object, and do something with each line
for row in csvData:

    # We're printing the row here, but you could do something more useful with the data.
    print row

# Close the file once we're done with it
file.close()

```

Converting the Data into a Dataset

Once you've read in the contents of a CSV into a script, you may wish to move it elsewhere. It is not uncommon to move the data to a **Power Table**, or other components on the screen. The main difficulty with this is converting the CSV data into a dataset so that it fits into the **Power Table** component's **Data** property. We have a couple of approaches listed below.

Calling the system.dataset.fromCSV Function

The [system.dataset.fromCSV](#) function can take a string and convert it to a dataset. Note, that the function expects a very specific format:

CSV File Content

```

#NAMES
Col 1,Col 2,Col 3
#TYPES
I,str,D
#ROWS,6
44,Test Row 2,1.8713151369491254
86,Test Row 3,97.4913421614675
0,Test Row 8,20.39722542161364
25,Test Row 4,20.3972254244222
33,Test Row 5,20.39722542232323
62,Test Row 6,20.39722542111999

```

Example

1. Create a Text file on your local system named "**example.csv**".

2. Open the Text file, copy the contents of the "CSV file Content" box above, and paste the contents to the file. Save the changes to the **example.csv** file.

```

Example.csv - Notepad
File Edit Format View Help
#NAMES
Col 1,Col 2,Col 3
#TYPES
I,str,D
#ROWS,6
44,Test Row 2,1.8713151369491254
86,Test Row 3,97.4913421614675
0,Test Row 8,20.39722542161364
25,Test Row 4,20.39722542444222
0,Test Row 5,20.39722542232323
0,Test Row 6,20.39722542111999

```

3. Let's move the contents of a CSV file to a Power Table. Add a **Power Table** and a **Button** component to your window. Double click on the **Button** component, and paste the following code into the **Script Editor** of the **actionPerformed** event:

```

Python - Using system.dataset.fromCSV()

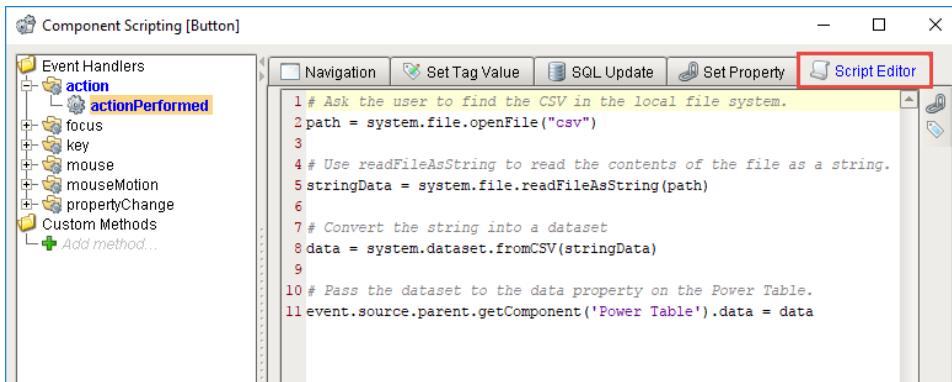
# Ask the user to find the CSV in the local file system.
path = system.file.openFile("csv")

# Use readFileAsString to read the contents of the file as a string.
stringData = system.file.readFileAsString(path)

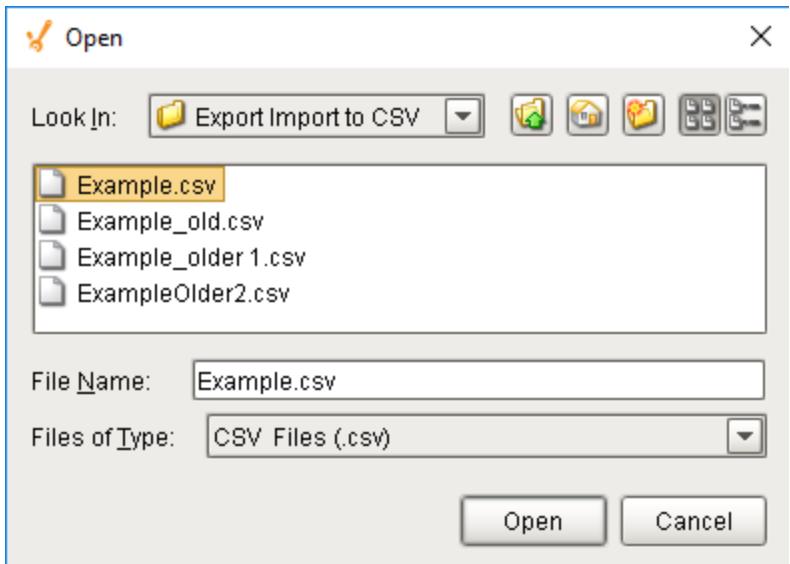
# Convert the string into a dataset
data = system.dataset.fromCSV(stringData)

# Pass the dataset to the data property on the Power Table.
event.source.parent.getComponent('Power Table').data = data

```



4. To execute your script, put the Designer into **Preview Mode**, and press the **Button**. A window will open for you to navigate and choose your CSV file, then click Open.



5. Your data will be displayed in the Power Table as shown below.

Col 1	Col 2	Col 3
44	Test Row 2	1.87
86	Test Row 3	97.49
0	Test Row 8	20.4
25	Test Row 4	20.4
33	Test Row 5	20.4
62	Test Row 6	20.4

Button

Calling the csv.reader Function

As mentioned, `system.dataset.fromCSV()` requires a specific format, which may not match the format of your file. In this case, we can use Python's CSV Library to parse the file and convert it to a dataset.

CSV File Content

```
Col 1,Col 2,Col 3
44,Test Row 2,1.8713151369491254
86,Test Row 3,97.4913421614675
0,Test Row 8,20.39722542161364
25,Test Row 4,20.39722542444222
33,Test Row 5,20.39722542232323
62,Test Row 6,20.39722542111999
```

Here is the code to import the above CSV data.

Python - Using Python's csv Library

```
# Import Python's built-in csv library.
import csv

# Ask the user to find the CSV in the local file system.
path = system.file.openFile("csv")

# Create a reader object that will iterate over the lines of a CSV.
# We're using Python's built-in open() function to open the file.
file = open(path)
csvData = csv.reader(file)

# Create a List of strings to use as a header for the dataset. Note that the number
# of headers must match the number of columns in the CSV, otherwise an error will occur.
# The simplest approach would be to use next() to read the first line in the file, and
# store that at the header.
header = csvData.next()

# Create a dataset with the header and the rest of our CSV.
dataset = system.dataset.toDataSet(header ,list(csvData))

# Store it into the table.
event.source.parent.getComponent('Power Table').data = dataset

# Close the file
file.close()
```

Related Topics ...

- [system.dataset.exportCSV](#)

Adding a Delay to a Script

Overview

In some cases, having a script execute after a delay is preferable. A common use case is waiting for some event elsewhere in the system to finish: a Tag change script executes that needs to wait for a new value from a separate Tag. One approach to this is to trigger our script, and then hold or wait until the other event occurs. On this page, we'll take a look at a couple of different approaches to this problem.

It is important to note that pausing a script can cause your client to lock. It is often preferred to look for another event to trigger the script you need. For the example above where we are waiting on a Tag to change, you might be able to use the Tag change to fire your script instead of waiting in the original script. It is up to you to determine the best trigger based on what exactly your script does.

Using the `system.util.invokeLaterLater` Function

The `system.util.invokeLaterLater` function is a great way to add a delay mid-way through the script. Simply create a function that represents all of the work that should occur after the delay, and pass the function to `invokeLater`, along with a delay period.

The example below calls two Message Boxes: once initially when pressed, and the other after a three second delay.

Python - Two Message Boxes

```
message = "All Done!"  
  
# Create a function that will be called by invokeLater.  
def runThisLater():  
    system.gui.messageBox(message)  
  
# Call invokeLater with a 3000ms delay. Note that our function will not  
run immediately  
# because invokeLater always executes once the rest of this script is  
complete.  
system.util.invokeLaterLater(runThisLater, 3000)  
  
# Bring up another Message Box. This will appear before the "All Done!"  
message, because of invokeLater.  
system.gui.messageBox("Waiting...")
```

One of the main limitations with `invokeLater` is that you can not pass parameters to the function that will be called. Parameters need to be initialized and determined elsewhere in your script, usually in the function definition.

Using a Timer from Python's Threading Library

The Threading Library has a Timer function that works in a very similar fashion to `invokeLater`. The main difference is that you can pass parameters to the function parameter when calling the Timer. Take a look at [Python's official documentation](#) for more details on the Timer object.

The example below will again call two Message Boxes with a three second delay between them. However, the text that is defined in the second Message Box is specified when starting the Timer.

Python - Python Threading Timer

```
from threading import Timer  
  
# Create a function that will be called by the Timer.  
def runThisLater(param):  
    system.gui.messageBox(str(param))  
  
# Constructs a Timer object that runs a function after a specified interval of time has passed.  
# Note the start() at the end: this is required to start the Timer. Don't forget this part!  
Timer(3.0, runThisLater, ["Stop Waiting, I'm done"]).start()
```

On this page ...

- Overview
- Using the `system.util.invokeLater` Function
- Using a Timer from Python's Threading Library
- Calling `time.sleep` from Python's Time Library
- Using a While Loop
- Approaches to Avoid - Locking the Client
 - Reasons to Avoid `time.sleep` and While Loops
 - Recommended Alternative
- Demonstration - Executing a Delay Between User Keyboard Input
 - Workflow for a Delay Using Two Scripts

```
# Bring up another Message Box. This will appear before the "All Done!" message, because of the Timer
system.gui.messageBox("Waiting")
```

Another benefit to using the Timer, is you can cancel the execution of the Timer's action using the cancel() function, but it only works if the Timer is still in its waiting stage.

Calling time.sleep from Python's Time Library

The simplest approach to pausing a script is to use the sleep() function in Python's Time Library:

Python - Sleeping the Code

```
from time import sleep

# This will pause execution of the script for 3 seconds. After that time, the script will continue.
sleep(3)

print "I'm awake!"
```

Using a While Loop

The [While Loop](#) is another simple approach to adding a delay: simply keep looping until the other event occurs. As always, you will want to take steps to ensure that an infinite loop never occurs: easiest by initializing a counter variable before iterating in the loop, and breaking out if the counter reaches a certain value.

Python - While Loop Safeguard

```
counter = 0

while not otherEvent:

    checkOtherEvent()
    counter += 1

    # Use this to break out if the event takes too long. You'll need the rest of your script to account
    # for this possibility.
    if counter >= 10000:
        print "Took too long. Leaving the While Loop"
        break
```

Approaches to Avoid - Locking the Client

While the sleep() and while loop functions are simple to use, they can cause problems by locking up the client and because of this, we generally recommend avoiding them if possible. Typically, [system.util.invokeLater](#) on a delay can accomplish the same task.

Note, that there are use cases for both approaches outlined below especially so in regard to the While Loop. However, neither function should be used to force a delay in a Client.



Using either of these methods to pause or delay a script can lock up the entire client, which may be very dangerous.

Reasons to Avoid time.sleep and While Loops

If used on a component, these approaches could lock up the Client or Designer. Scripts called from a component run in the same thread that the Client and Designer use to refresh the screen. Whenever a script on a component triggers, the screen is unable to refresh until the script finishes. In most cases, this is so fast that no one notices. If your script is intentionally calling sleep() or using a long running While Loop, then the Client will lock up for the duration.

Additionally, these approaches run an extended period of time, and block other component based scripts from executing. They do not yield to other scripts while waiting.

Because of these two reasons, the sleep() and While Loop functions can cause your clients to appear unresponsive, when really they are just running a script that prevents the screen from refreshing.

Recommended Alternative

As mentioned earlier, the [system.util.invokeLater](#) function (also mentioned on this page) can provide the same functionality. If a sleep() or While Loop function must be used, then they should be called from [system.util.invokeAsynchronous](#) to prevent blocking any other scripts.

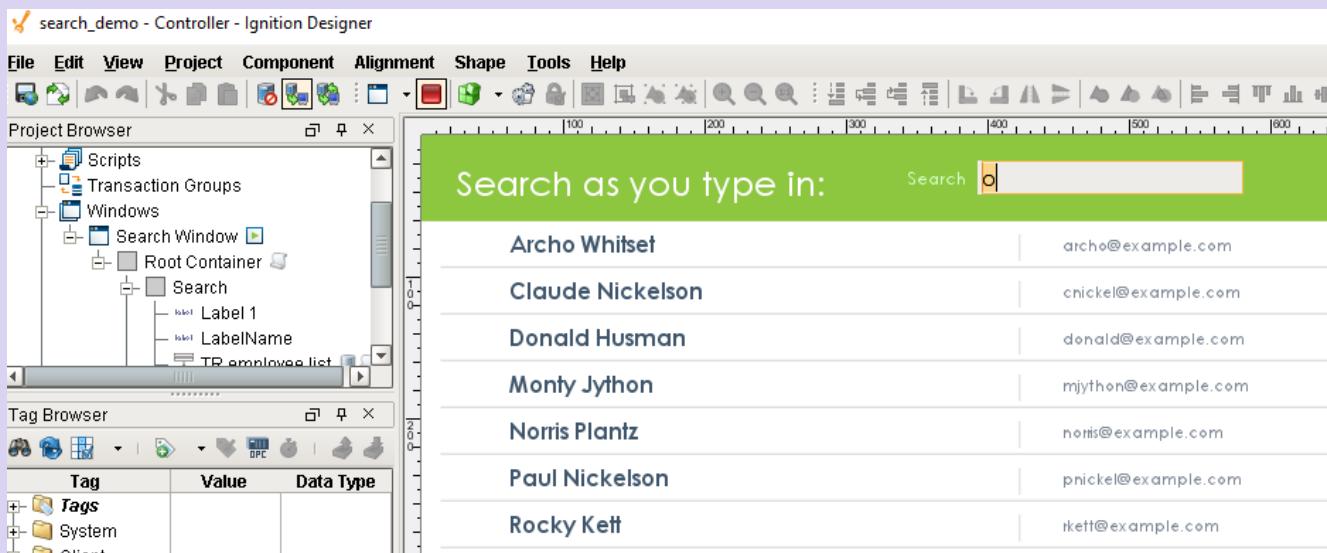
Editor notes are only visible to logged in users

This entire section is in an Editor Note macro because the final script in this section could potentially open people to SQL injection attacks. Leaving in Editor Note macro for now in case we want to revise the example in the future.

Demonstration - Executing a Delay Between User Keyboard Input

In many cases, you may need to show your users a large number of entries on a Table or Template Repeater. As more entries are added to the system, adding a search field that can filter results becomes more appealing. Furthermore, being able to filter as the user types (as opposed to forcing them to hit enter every time) adds some polish to the window. However, if the entries are backed by the database, you will not want to run a query every time a user presses a key. Instead, it would be preferable to wait for a delay in input, and run one query to limit strain on the database.

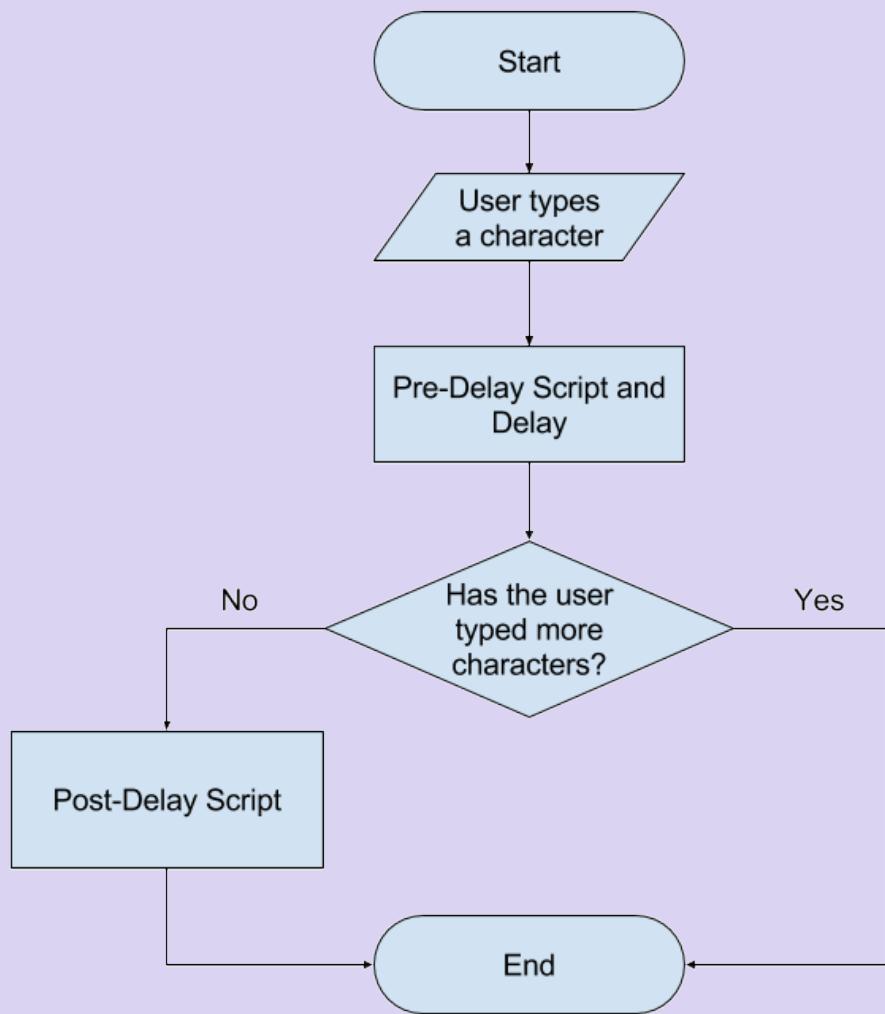
The following is not a traditional example that you would find here in the User Manual, and assumes you are comfortable with many concepts in Ignition.



Workflow for a Delay Using Two Scripts

One approach to adding a delay involves two scripts:

- **Pre-Delay Script:** This script fires before the delay. In this case, we will use a script on the Text Field's keyReleased event. This will call the script everytime a new key is entered into the Text Field, and calls the Post-Delay Script on a delay. It also passes the text that is currently in the Text Field, and increments an edit count (could be configured as a custom property on the same component).
- **Post-Delay Script:** This script runs after the delay, so it needs to check if there were any new keystrokes that occurred during the delay. If there were no new keystrokes, then it can run the query. In this demonstration, we will use a Custom Method not the Text Field component.



Additionally, we will need some criteria that our script can use to determine if it is safe to run our query. You can have multiple criteria here, but for the sake of simplicity we will use the **Text Property on the Text Field**. This way we will know what the user typed before the delay, and then can compare it back to the Text property after the delay. If the Text on the component after the delay is different, then we don't want to run the query as the user may still be typing.

The keyReleased script is shown below. It is using a .05 second delay, but could of course be modified. We're creating a new Timer object every time a key is released, so there will be many of these scripts firing as the user types. However, they are fairly lightweight, and don't directly interact with the database, so having many executions of this script isn't taxing on the system.

Python - Pre-Delay Script

```

# We're using the Timer object for this because we want to pass parameters to the function that will run after the delay.
from threading import Timer

# Ignoring arrow keys.
if (event.keyCode < 37 or event.keyCode > 40):

    # Grab the text in the component currently.
    currentText = event.source.text

    # Calls the function after 0.5 seconds, and pass the currentText.
    Timer (0.5, event.source.sendingQuery,[currentText]).start()
  
```

Below is the sendingQuery script. As mentioned, this script will determine if the user ceased keyboard activity before running a query against the database. This example is using a Custom Method on the Text Field, so it uses the `self` argument to reference the source component.

Python - Post-Delay Script

```
def sendingQuery (self, oldText):

    # Read the value of the text property, since it may have changed during the delay.
    currentSearchText = self.text

    # If the text before the delay is the same as the after the delay, there has been a pause in
    # keyboard activity, so we should run the query.
    if (currentSearchText == oldText):

        # If the text field isn't empty, then we need to filter the results with a WHERE clause
        # and our criteria.
        if(currentSearchText != ""):
            newQuery = "SELECT * FROM employees WHERE CONCAT(firstname, ' ', lastname) like
            '%" + currentSearchText + "%'"

            # If the Text Field Is blank, then we should run the query again, but this time without a
            # filter, so all results appear.
            else:
                newQuery = "SELECT * FROM employees"
```

We can now run our query in between user keyboard activity.

Related Topics ...

- [system.util.invokeLater](#)
- [Timer Scripts](#)

Export Tag Historian to CSV

Obtaining Historian Data

Sometimes, it may be useful to export the values from the Tag Historian. There are many ways that this can be done, but the easiest way is to use the built in [system.tag.queryTagHistory](#) function. This allows you to specify a list of Tag paths to query history for.

Python - Grabbing History Data

```
# First we start by creating a start and end time that we will be querying history for.  
endTime = system.date.now()  
startTime = system.date.addMinutes(endTime, -30)  
  
# Next we call our queryTagHistory function, using the start and end dates as well as a Tag path.  
# The other parameters listed for this function can be altered to fit your need.  
data = system.tag.queryTagHistory(paths=[ "[default]myTag" ],  
startDate=startTime, endDate=endTime, returnSize=10, aggregationMode="Average", returnFormat='Wide')
```

This simple script will query Tag History for a Tag called myTag. It will query the last 30 mins of data, return 10 records, each record will be an average of the time slice, and the return format will be wide. For more information on the function parameters and the different options available, see the [queryTagHistory](#) function page. The data returned is then being stored in a variable **data**, and we can use it later in the script in anyway we want.

Exporting to CSV

Now that we have our data, it is really easy to export it to a CSV file. Ignition has a built-in function to convert a dataset to CSV data called [system.dataset.toCSV](#). This CSV data can then be written to an actual file using [system.file.writeFile](#).

Python - Export History to CSV

```
# First we start by creating a start and end time that we will be querying history for.  
endTime = system.date.now()  
startTime = system.date.addMinutes(endTime, -30)  
  
# Next we call our queryTagHistory function, using the start and end dates as well as a Tag path.  
# The other parameters listed for this function can be altered to fit your need.  
data = system.tag.queryTagHistory(paths=[ "[default]myTag" ],  
startDate=startTime, endDate=endTime, returnSize=10, aggregationMode="Average", returnFormat='Wide')  
  
# Turn that history data into a CSV.  
csv = system.dataset.toCSV(data)  
  
# Export that CSV to a specific file path. The r forces it to use the raw path and not require double  
backslashes.  
system.file.writeFile(r"C:\myExports\myExport.csv", csv)
```

On this page ...

- [Obtaining Historian Data](#)
- [Exporting to CSV](#)
- [History Tag Search and Export](#)

History Tag Search and Export

If you have a lot of Tags storing history, it can be difficult to list them all out in the function above. This example makes things easier by specifying a folder of Tags instead of individual Tags. The script first grabs all of the historical Tags from a particular folder and gets the Tag paths using the [system.tag.browseHistoricalTags](#) function, runs those through the [system.tag.queryTagHistory](#) function, and then exports the results to a CSV file. This script is great, because it can be run from the [scripting console](#), or from a button on the client. However, everything in the script is hardcoded to work with a specific system, though that can be easily changed to fit any system. Below are the list of parts that can be changed with what each of them is used for.

Part of Script	Description
----------------	-------------

<pre>path='histprov: myDB:/drv: myIgnitionGateway: myTagProvider:/tag: myFolderOfTags'</pre>	<p>The path to the parent folder that we are searching through for history Tags. Each part of the string corresponds to a different part of your system.</p> <ul style="list-style-type: none"> • histprov:myDB:/ - This is the name of your Historical Tag Provider that you are searching in. Replace myDB with the name of your historical Tag Provider. • drv:myIgnitionGateway:myTagProvider:/ - The first thing is the name of your Ignition Gateway. Replace myIgnition Gateway with the name of your Ignition Gateway, which can be found by going to the Gateway Configure Webpage and navigating to System > Gateway Settings. It is the System Name. The second thing is the name of the Realtime Tag Provider where the Tags reside. Replace myTagProvider with the name of your Tag Provider. • tag:myFolderOfTags - This is the name of the folder you want to search for historical Tags. Replace myFolderOfTags with the name of a folder in your Tag Browser to search through. <p>Note: Paths to Internal Historian providers (such as those used by Edge Gateways) do not need to include the Gateway name.</p>
<pre>system.tag. queryTagHistory (...)</pre>	<p>This is what will actually query for historical data. While the path parameter should be kept the same, the other parameters can be altered to fit your needs. Currently, the function is querying the last 30 mins of history, has a return size of 10, an aggregation mode of Average, and a return format of wide. Of course, all of this can be changed. See system.tag.queryTagHistory for more information on the parameters that can be used here.</p>
<pre>r"C: \myExports\myExport .csv"</pre>	<p>This is the path that the CSV file will save at. This can be changed to be any location you want. You can even use system.file.saveFile to first ask the user for a save location and filename, and then passing the path returned from that function into the writeFile function. The r at the beginning of the string is important as it allows us to not have to use double backslashes.</p>

Caution:

This script can be dangerous! It recursively looks through a folder to find all historical Tags, and then uses all of those Tags to look up Tag history. Looking through a folder that has too many history Tags, or querying history for too large a period of time can potentially result in locking up your system. Use caution when using this function, and never search through too large a set of Tags or query too much history.

Python - Recursively Browse for History and Export to CSV

```
# Our browse function that will browse for historical tags.  
# By setting this up as a function, it allows us to recursively call it to dig down through the specified folder.  
# Pass in an empty list that we can add historical paths to, and the path to the top level folder.  
def browse(t, path):  
  
    # Loop through the results of the historical tag browse, and append the path to the empty list.  
    for result in system.tag/browseHistoricalTags(path).getResults():  
        t.append(result.getPath())  
  
        # If the result is a folder, run it through the browse function as well.  
        # This will continue until we are as deep as possible.  
        if result.hasChildren():  
            browse(t, result.getPath())  
  
# Start with an empty list to store our historical paths in.  
historyPaths = []  
  
# Call the browse function, passing in an empty list, and the folder that we want to browse for historical tags.  
# This path is a placeholder. It should be replace with your valid path.  
browse(historyPaths, path='histprov:myDB:/drv:myIgnitionGateway:myTagProvider:/tag:myFolderOfTags')  
  
# Create another empty list to store our tag paths that we will pull out of the historical paths.  
tagPaths = []  
  
# Loop through the list of historical tag paths, split out just the tag path part,  
# and push it into our tag path list.  
for tag in historyPaths:  
    tagPaths.append("[myTagProvider]" + str(tag).split("tag:")[1])
```

```
# Now that we have a list of tag paths, we need to grab the historical data from them.  
# Start by creating a start and end time.  
endTime = system.date.now()  
startTime = system.date.addMinutes(endTime, -30)  
  
# Then we can make our query to tag history, specifying the various parameters.  
# The parameters listed for this function can be altered to fit your need.  
data = system.tag.queryTagHistory(paths=tagPaths, startDate=startTime, endDate=endTime, returnSize=10,  
aggregationMode="Average", returnFormat='Wide')  
  
# Turn that history data into a CSV.  
csv = system.dataset.toCSV(data)  
  
# Export that CSV to a specific file path. The r forces it to use the raw path with backslashes.  
system.file.writeFile(r"C:\myExports\myExport.csv", csv)
```

Related Topics ...

- [Exporting and Importing a CSV](#)

Parsing XML with the Etree Library

What is the xml.etree Library?

The etree (ElementTree) library is a part of the python standard library, and contains many tools that make it simple to parse through and pull information out of an XML document. There are other libraries that can parse through XML documents, but etree is commonly used and very easy to get started with. The etree library will break up the XML into easily accessible elements, each representing a single node in the entire XML tree. For more information on using the etree library beyond the scope of this page, see the [python documentation](#).

On this page ...

- [What is the xml.etree Library?](#)
 - [Using the xml.etree Library](#)
 - [A Simple Book Example](#)

Using the xml.etree Library

There are a couple of different ways to import the XML data from etree, depending on how it is being stored. It can pull the data in from an XML file using the filepath, or it can read a string. Notice how regardless of how we import the XML, we end up with a root object.

Python - Reading a File

```
# The library must first be imported no matter how we pull in the data.
import xml.etree.ElementTree as ET

# Here we can grab the filepath using Ignition's built in openFile function, parse that into a tree, then
# grab the root element.
filepath = system.file.openFile()
tree = ET.parse(filepath)
root = tree.getroot()
```

Python - Reading from a String

```
# The library must first be imported no matter how we pull in the data.
import xml.etree.ElementTree as ET

# Alternately, we can start with a string of the xml data.
xmlString = """
<employee id="1234">
    <name>John Smith</name>
    <start_date>2010-11-26</start_date>
    <department>IT</department>
    <title>Tech Support</title>
</employee>
"""

# Then parse through the string using a different function that takes us straight to the root element.
root = ET.fromstring(xmlString)
```

Each Tag is considered an element object. In the example above, the root element would be the employee Tag. Elements can also have attributes, which are within the Tag itself. In the example above, the employee element has an attribute id with a value 1234. Finally, each element can also have additional data, typically in the form of a string. This additional data is usually placed in between the element's start and end Tags. In the example above, the employee element has no additional data, but its children do. The name element would have an additional data value of John Smith. All of this data can be accessed using the Element object's built-in functionality. The major functions are listed below, and each example uses the reading from a string root XML example from above.

Function	Description	Example	Output
Element.tag	Returns the name of the Element's Tag.	<pre>print root.tag</pre>	employee
Element.attrib	Returns a dictionary of the Element's attributes.		{'id':'1234'}

		print root. attrib	
Element. text	Returns the additional data of the Element. The example here will return nothing because the root does not have any text. The next example uses children which do have text.	print root. text	
for child in Element	Will iterate through the Element's children. Each child is then its own element, complete with Tag, attrib, and text properties.	for child in root: print child. tag, child. text	name John Smith start_date 2010-11- 26 departme nt IT title Tech Support
Element [index]	Allows direct reference to an Element's children by index. Since Tags can be nested many times, further nested children can be accessed by adding an additional index in square brackets as many times as necessary: Element[1][4] [0] From the original element, we would go to the child located in the first position, that child's fourth position child, and that child's zero position child. When direct referencing child elements in this way, they still have access to the Tag, attrib, and text properties.	root[2]. tag root[3]. text	departme nt Tech Support

A Simple Book Example

Using the functions above, we can now easily parse through an XML file and use the results for something. Lets keep it simple, and parse through a document and then place the values into a table. First we need to start with an XML document. We have one below for you to test with in a string form, which would need to be pasted at the top of the script.

XML String

```
document = """
<catalog>
    <book id="bk101">
        <author>Gambardella, Matthew</author>
        <title>XML Developer's Guide</title>
        <genre>Computer</genre>
        <price>44.95</price>
        <publish_date>2000-10-01</publish_date>
        <description>An in-depth look at creating applications
with XML.</description>
    </book>
    <book id="bk102">
        <author>Ralls, Kim</author>
        <title>Midnight Rain</title>
        <genre>Fantasy</genre>
        <price>5.95</price>
        <publish_date>2000-12-16</publish_date>
        <description>A former architect battles corporate zombies,
an evil sorceress, and her own childhood to become queen
of the world.</description>
    </book>
    <book id="bk103">
        <author>Corets, Eva</author>
        <title>Maeve Ascendant</title>
        <genre>Fantasy</genre>
        <price>5.95</price>
        <publish_date>2000-11-17</publish_date>
```

```
<description>After the collapse of a nanotechnology  
society in England, the young survivors lay the  
foundation for a new society.</description>  
</book>  
<book id="bk104">  
    <author>Corets, Eva</author>  
    <title>Oberon's Legacy</title>  
    <genre>Fantasy</genre>  
    <price>5.95</price>  
    <publish_date>2001-03-10</publish_date>  
    <description>In post-apocalypse England, the mysterious  
agent known only as Oberon helps to create a new life  
for the inhabitants of London. Sequel to Maeve  
Ascendant.</description>  
</book>  
<book id="bk105">  
    <author>Corets, Eva</author>  
    <title>The Sundered Grail</title>  
    <genre>Fantasy</genre>  
    <price>5.95</price>  
    <publish_date>2001-09-10</publish_date>  
    <description>The two daughters of Maeve, half-sisters,  
battle one another for control of England. Sequel to  
Oberon's Legacy.</description>  
</book>  
<book id="bk106">  
    <author>Randall, Cynthia</author>  
    <title>Lover Birds</title>  
    <genre>Romance</genre>  
    <price>4.95</price>  
    <publish_date>2000-09-02</publish_date>  
    <description>When Carla meets Paul at an ornithology  
conference, tempers fly as feathers get ruffled.</description>  
</book>  
<book id="bk107">  
    <author>Thurman, Paula</author>  
    <title>Splish Splash</title>  
    <genre>Romance</genre>  
    <price>4.95</price>  
    <publish_date>2000-11-02</publish_date>  
    <description>A deep sea diver finds true love twenty  
thousand leagues beneath the sea.</description>  
</book>  
<book id="bk108">  
    <author>Knorr, Stefan</author>  
    <title>Creepy Crawlies</title>  
    <genre>Horror</genre>  
    <price>4.95</price>  
    <publish_date>2000-12-06</publish_date>  
    <description>An anthology of horror stories about roaches,  
centipedes, scorpions and other insects.</description>  
</book>  
<book id="bk109">  
    <author>Kress, Peter</author>  
    <title>Paradox Lost</title>  
    <genre>Science Fiction</genre>  
    <price>6.95</price>  
    <publish_date>2000-11-02</publish_date>  
    <description>After an inadvertant trip through a Heisenberg  
Uncertainty Device, James Salway discovers the problems  
of being quantum.</description>  
</book>  
<book id="bk110">  
    <author>O'Brien, Tim</author>  
    <title>Microsoft .NET: The Programming Bible</title>  
    <genre>Computer</genre>  
    <price>36.95</price>  
    <publish_date>2000-12-09</publish_date>  
    <description>Microsoft's .NET initiative is explored in  
detail in this deep programmer's reference.</description>  
</book>
```

```

<book id="bk111">
    <author>O'Brien, Tim</author>
    <title>MSXML3: A Comprehensive Guide</title>
    <genre>Computer</genre>
    <price>36.95</price>
    <publish_date>2000-12-01</publish_date>
    <description>The Microsoft MSXML3 parser is covered in
    detail, with attention to XML DOM interfaces, XSLT processing,
    SAX and more.</description>
</book>
<book id="bk112">
    <author>Galos, Mike</author>
    <title>Visual Studio 7: A Comprehensive Guide</title>
    <genre>Computer</genre>
    <price>49.95</price>
    <publish_date>2001-04-16</publish_date>
    <description>Microsoft Visual Studio 7 is explored in depth,
    looking at how Visual Basic, Visual C++, C#, and ASP+ are
    integrated into a comprehensive development
    environment.</description>
</book>
</catalog>
"""

```

We can then place a Table component and a Button component on the window, and place this script on the Button's actionPerformed event.

Python - Complete XML Parsing

```

# Start by importing the library
import xml.etree.ElementTree as ET

#####
# Here is where you would paste in the document string.
# Simply remove this comment, and paste in the document string.
#####

# We can then parse the string into useable elements.
root = ET.fromstring(document)

# This creates empty header and row lists that we will add to later.
# These are used to create the dataset that will go into the Table.
# We could fill in the names of the headers beforehand, since we know what each will be.
# However, this allows us to add or remove children keys, and the script will automatically adjust.
headers = []
rows = []

# Now we can loop through each child of the root.
# Since the root is catalog, each child element is an individual book.
# We also create a single row empty list. We can add all of the data for a single book to this list.
for child in root:
    oneRow = []

    # Check if the book has any attributes.
    if child.attrib != {}:

        # If it does contain attributes, we want to loop through all of them.
        for key in child.attrib:

            # Since we only want to add the attributes to our header list once, first check if
            it is there.
            # If it isn't add it.
            if key not in headers:
                headers.append(key)

            # Add the attribute value to the oneRow list.
            oneRow.append(child.attrib[key])

    # Loop through the children of the book.
    for child2 in child:

```

```
# Similar to above, we check if the tag is present in the header list before adding it.
if child2.tag not in headers:
    headers.append(child2.tag)

# We can then add the text of the Element to the oneRow list.
oneRow.append(child2.text)

# Finally, we want to add the oneRow list to our list of rows.
rows.append(oneRow)

# Once the loop is complete, this will print out the headers and rows list so we can manually check them in
# the console.
print headers
print rows

# Convert to a dataset, and insert into the Table.
data = system.dataset.toDataSet(headers, rows)
event.source.parent.getComponent('Table').data = data
```

Related Topics ...

- [system.file.openFile](#)
- [Libraries](#)

Audit Log and Profiles

Ignition's built-in auditing system automatically records certain actions that occur in the system, such as a Tag writes or User Source authentication, into a SQL database table. Utilizing the system involves creating an Audit Profile, followed by enabling auditing in a project. Once both prerequisites have been met, the Gateway will automatically create a database table named AUDIT_EVENTS, and use the table to start tracking user actions.

The Remote Audit Log configuration option allows audit events to be automatically sent to a remote Gateway's audit log. The remote Gateway you plan to connect to must have a Audit Profile created. To learn more about sending audit events to a remote Gateway, refer to section [Creating a Remote Gateway Audit Profile](#) on this page.

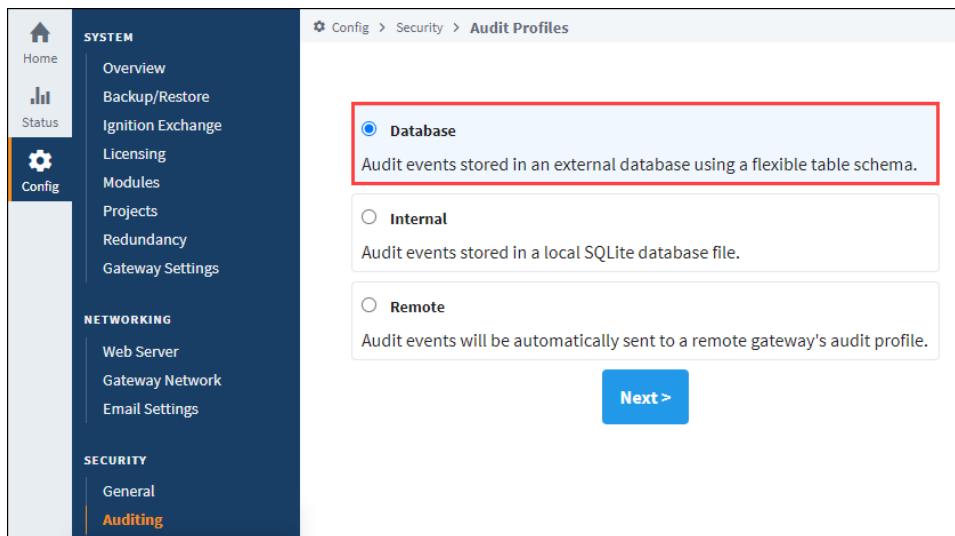
Note: You can use Audit Profiles for Gateway events and project events. See the [Enabling Auditing for Gateway-Spaced Actions](#) or the [Enabling Auditing in a Project](#) sections for more details.

Auditing Actions

For a list of actions that are recorded by an audit profile, see the [Auditing Actions Reference](#) page.

Create a Database Audit Profile

1. Go to the **Config** section of the Gateway Webpage.
2. Scroll down to the **Security > Auditing** from the menu on the left. The Audit Profiles page is displayed.
3. Click the **Create a new Audit Profile** link.
4. You have the option of storing audit logs into an external database or sending them to a remote Gateway. For this example, select **Database**. (Configuring audit events to be sent to a remote Gateway's audit log is addressed in [Creating a Remote Gateway Audit Profile](#) section on this page).



5. Enter the **Name** of the audit log and **Description** (optional).
6. In the Retention field, set a value in days for how long you want audit records kept. (The default is 90 days.)
7. Under the **Database Settings**, select the **Database** where the table will be stored, select the **Auto Create** check box, and enter the desired **Table Name**.
8. Click **Create New Audit Profile**.

On this page ...

- [Create a Database Audit Profile](#)
- [Create an Internal Audit Profile](#)
- [Creating a Remote Gateway Audit Profile](#)
- [Enabling Auditing for Gateway-Spaced Actions](#)
- [Enabling Auditing in a Project](#)
- [Viewing Information in an Audit Log](#)
- [Audit Log Table Descriptions](#)

Config > Security > Audit Profiles

Main

Name	<input type="text"/>
Description	<input type="text"/>
Retention	90 How long (in days) should audit records be kept? Values less than or equal to 0 will disable pruning. (default: 90)

Database Settings

Database	MySQL57 The database connection to use to store audit events.
Auto Create	<input checked="" type="checkbox"/> If true, the table schema specified here will be automatically verified and created if necessary. (default: true)
Pruning Enabled	<input type="checkbox"/> If false, this audit profile will never prune records, regardless of the retention field. Otherwise, the retention field will be followed. (default: false)
Table Name	AUDIT_EVENTS The name of the table to store audit events. (default: AUDIT_EVENTS)

Show advanced properties

Create New Audit Profile

Once some changes have been made to a Tag or a Database table, Ignition will begin recording.

AUDIT_EVENTS_ID	EVENT_TIMESTAMP	ACTOR	ACTOR_HOST	ACTION	ACTION_TARGET	ACTION_VALUE
1	2016-07-25 17:50:09	admin	IU-WorkStation	tag write	B Tags/B3:1	1.0
2	2016-07-25 17:50:51	admin	IU-WorkStation	tag write	B Tags/B3:1	100.0
3	2016-07-25 17:50:53	admin	IU-WorkStation	tag write	B Tags/B3:1	2.0
4	2016-07-25 17:50:56	admin	IU-WorkStation	tag write	B Tags/B3:1	8.0
5	2016-07-25 17:51:20	admin	IU-WorkStation	query	update audit_events set acto...	4
6	2016-07-25 17:51:51	admin	IU-WorkStation	query	UPDATE audit_events SET `A...	1

Database Audit Profile Properties Table

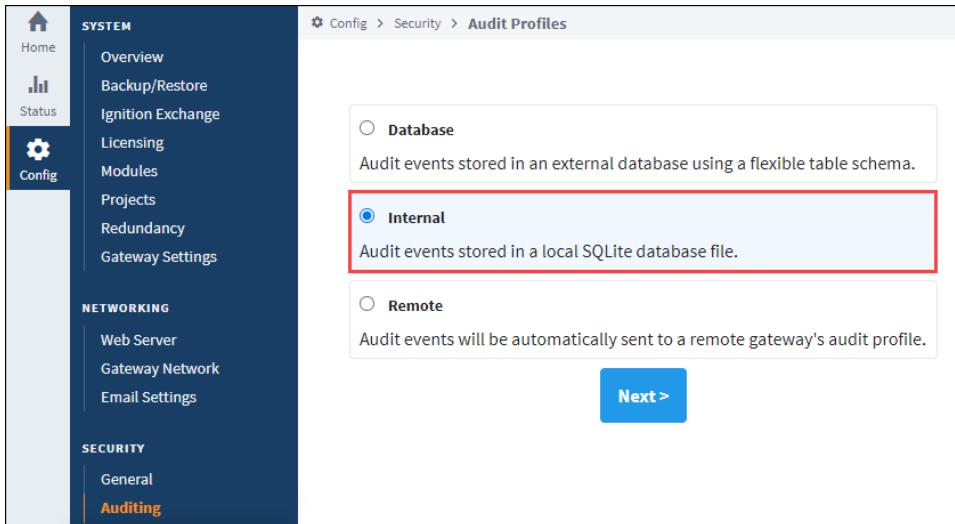
Main	
Name	The default name, is the name of the Audit Profile.
Description	Description of the audit profile. Optional.

Retention	<p>The following feature is new in Ignition version 8.1.1 Click here to check out the other new features</p> <p>How long (in days) should audit records be kept? Values less than or equal to 0 will disable pruning. Default is 90 days.</p>
Database Settings	
Database	The database connection to use to store audit events.
Auto Create	If true (selected), the table schema specified here will be automatically verified and created if necessary. Default is true.
Pruning Enabled	<p>The following feature is new in Ignition version 8.1.3 Click here to check out the other new features</p> <p>If false, this audit profile will never prune records, regardless of the retention field. Otherwise, the retention field will be followed. Default is false.</p>
Table Name	The name of the table to store audit events. Default is AUDIT_EVENTS.

Create an Internal Audit Profile

The Internal Audit Profile option allows an Ignition Gateway to store audit records without an external SQL database. The only way to interact with the Internal Audit Profile is via the Status page of the Gateway webpage.

1. Go to the **Config** section of the Gateway Webpage.
2. Scroll down to the **Security > Auditing** from the menu on the left. The Audit Profiles page is displayed.
3. Click the **Create a new Audit Profile** link.
4. Select **Internal**.



5. Enter a name for the audit log and a description (optional).
6. In the Retention field, set a value in days for how long you want audit records kept. (The default is 90 days.)
7. Click **Create New Audit Profile**.

Internal Audit Profile Properties Table

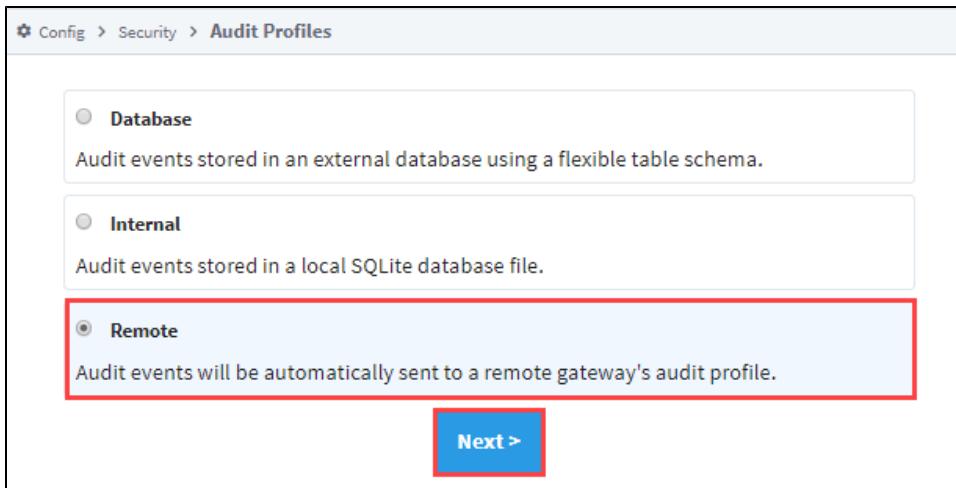
Main	
Name	The default name, is the name of the Audit Profile.
Description	Description of the audit profile. Optional.
Retention	<p>The following feature is new in Ignition version 8.1.1 Click here to check out the other new features</p>

Value in days for how long you want audit records kept. (The default is 90 days.)

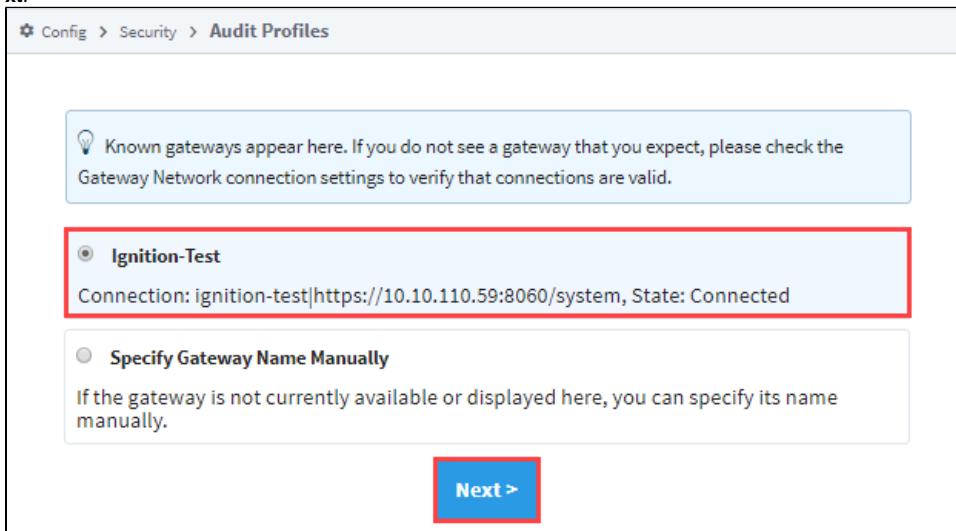
Creating a Remote Gateway Audit Profile

Just like configuring audit events to be logged into an external database, it is done from the Gateway Webpage, **Config > Security > Auditing**.

1. To have your audit events automatically sent to a remote Gateway's audit profile, select **Remote**, and click **Next**.



2. A list of known Gateways will be displayed. If you don't see a Gateway that you expected to see, check your Gateway Network settings to verify that the connections are valid. You also have the option to specify a Gateway manually. This example selects a valid Gateway. Click **Next**.



3. If an Audit profile exists, the fields will auto-populate. The name of the Gateway will appear in the **Name** field prefaced with the audit profile name (i.e., Ignition_Test_Auditing), as shown in the following example. Click **Create New Audit Profile**.

The screenshot shows the 'Audit Profiles' configuration screen. At the top, there's a navigation bar: 'Config > Security > Audit Profiles'. Below it is a 'Main' section with a 'Name' field containing 'Ignition_Test_Auditing' and a 'Description' field which is empty. The next section is 'Remote Settings', containing 'Target System' (set to 'Ignition-Test') and 'Target Profile' (set to 'Auditing'). Underneath is a checked checkbox for 'Show advanced properties', followed by an 'Advanced' section with a 'Use Store and Forward' option (which has a checked checkbox) described as enabling audit events to be stored through a store and forward system or directly against the remote gateway (default: true). At the bottom right of the main form area is a prominent blue button with white text that says 'Create New Audit Profile'.

4. You will receive a successful message stating your new Audit Profile was created.

The screenshot shows the same 'Audit Profiles' page after a profile has been created. A green success message at the top states 'Successfully created new Audit Profile "Ignition_Test_Auditing"'. Below this, a table lists existing audit profiles. The first profile, 'Audit', is of type 'Database' and has a red 'edit' button. The second profile, 'Ignition_Test_Auditing', is of type 'Remote' and is highlighted with a red border around its entire row, indicating it is the newly created profile. It also has a red 'edit' button. At the bottom left of the table area is a blue link that says 'Create new Audit Profile...'

Remote Gateway Audit Profile Properties Table

Main	
Name	The default name, is the name of the Remote Gateway and Audit Profile.
Description	Description of the audit profile. Optional.

Enabled	By default, the journal profile is enabled.
Remote Settings	
Target System	The remote system to send audit events to over the Gateway network.
Target Profile	The audit profile on the remote system to log events into.
Advanced	
Use Store and Forward	<p>The following feature is new in Ignition version 8.1.23 Click here to check out the other new features</p> <p>If enabled, audit events will be stored through the Store and Forward system. If not enabled, they will be stored directly against the remote Gateway. Default is true.</p>

Enabling Auditing for Gateway-Spaced Actions

After setting up an Audit Profile, you can have the audit log record Gateway events.

1. Navigate to the [Gateway General Security Settings](#) page. This is located at your **Gateway's Config page > Security > General**.



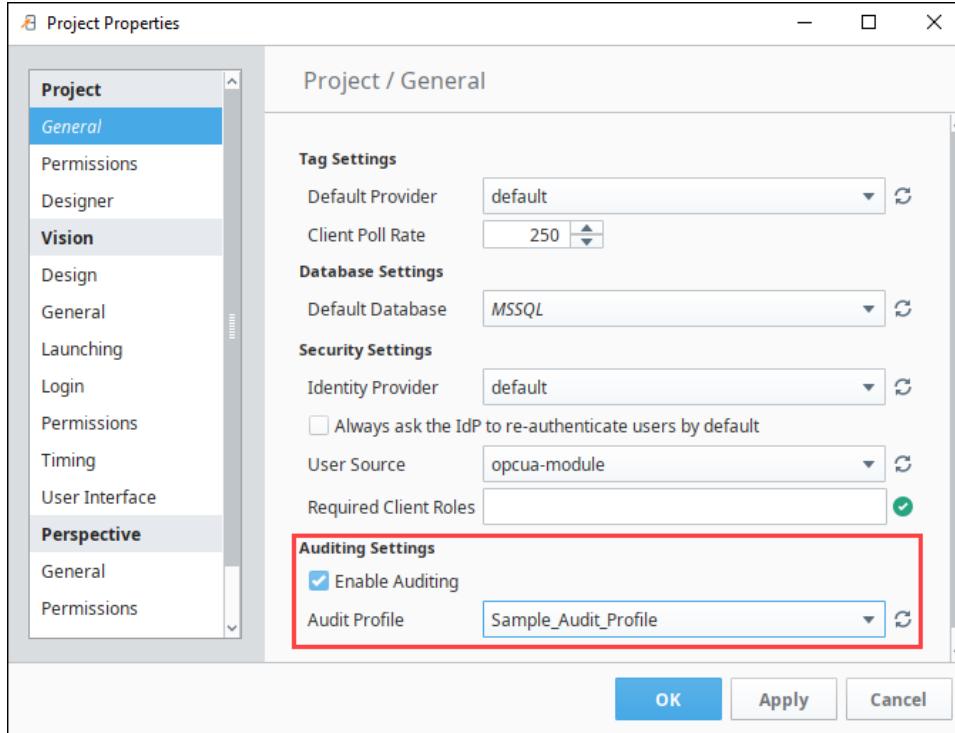
2. Set the **Gateway Audit Profile** setting to the audit profile you want to record with, then save your settings.
3. Once you press save, a confirmation banner will appear at the top of the page and your Gateway will begin recording events to the specified audit profile.

Successfully saved Gateway Security Settings

Enabling Auditing in a Project

1. Go to the Designer, open the project that you want to enable auditing on, then go to **Project > Properties**.
2. Go to the General section, select the **Enable Auditing** check box, and select your **Audit Profile** from the drop-down menu. The audit profile is used to record audit actions for your project. If the new audit profile does not show up, click **Refresh**.
3. Click **OK**.

4. Save your Project.



Viewing Information in an Audit Log

There are a few ways to [view audit information](#): using a Table component, interface on the Gateway, or the Database Query Browser. Here is one example of viewing an Audit Log using the Database Query Browser.

1. In the Designer, go to **Tools > Database Query Browser**.
2. Under the **Schema** area, double click on a **table**, and it will expand the query in the Database Query Browser area.
3. Click **Execute**. All the audit log data will be displayed in the Resultset1 area.

The screenshot shows the 'Database Query Browser' window. At the top, a query is entered in the text area: 'SELECT * FROM audit_events'. Below the query, there is a checkbox 'Limit SELECT to:' followed by a text input field containing '1000 rows'. To the right of the query area is a large 'Execute' button with a yellow arrow icon. The main area displays a table titled 'Resultset 1'. The table has columns: AUDIT_EVENT_ID, EVENT_TIMESTAMP, ACTOR, and ACTOR_HOST. The data shows four rows of audit events. At the bottom of the table, it says '438 rows fetched in 0.011s'. At the very bottom of the window are buttons for 'Auto Refresh', 'Edit', 'Apply', and 'Discard'. On the right side of the window, there is a 'Schema' pane showing the database structure, including tables like 'agent_events' with columns 'agent_name', 'agent_role', 'event_category', and 'event_level'.

Audit Log Table Descriptions

A description of the columns used by the audit log's database table can be found on the [Ignition Database Table Reference](#) page.

Related Topics ...

- [Database Query Browser](#)
- [Audit Log Display](#)

In This Section ...

Alarm Notification Auditing

Alarm Notification profiles can be set to store information in an Audit Profile.

How to Store Alarm Notifications in an Audit Profile

Once you have an [audit profile](#) created in Ignition, you can configure your [Alarm Notification Profile](#) to start using it.

1. Go to the **Config** section of the Gateway.
2. Choose **Alarming > Notification** from the menu on the left.
The Alarm Notification Profiles page is displayed.
3. Edit the appropriate notification profile by clicking the **edit** link.

Alarm Notification Profiles				
Name	Description	Enabled	Type	Status
Email	true	Email Notification	Running	delete edit
Email 1	true	Email Notification	Running	delete edit
SMS	true	SMS Notification	Running	delete edit

4. The Edit Alarm Notification Profile page is displayed.
Scroll down to the **Auditing** section, and select an **Audit Profile** from the drop-down menu.
Click **Save Changes**, and Ignition will automatically begin storing information.

Auditing	
Audit Profile	<input type="text" value="Audit"/> ▼
If an audit profile is selected, events such as emails and acknowledgements will be stored to the audit system. Note that alarm acknowledgements are also stored to the alarm journal.	
<input type="checkbox"/> Show advanced properties	
Save Changes	

Now that the Alarm Notification Profile is storing data into the Audit system, you have a complete log of all alarm emails and acknowledgements that you can review.

See [Audit Log Display](#) for more info on how to retrieve information from the Audit Log.

Related Topics ...

- [Alarm Notification](#)
- [Database Connections](#)
- [Audit Log Display](#)

On this page ...

- [How to Store Alarm Notifications in an Audit Profile](#)



Alarm Notification Auditing

[Watch the Video](#)

Audit Log Display

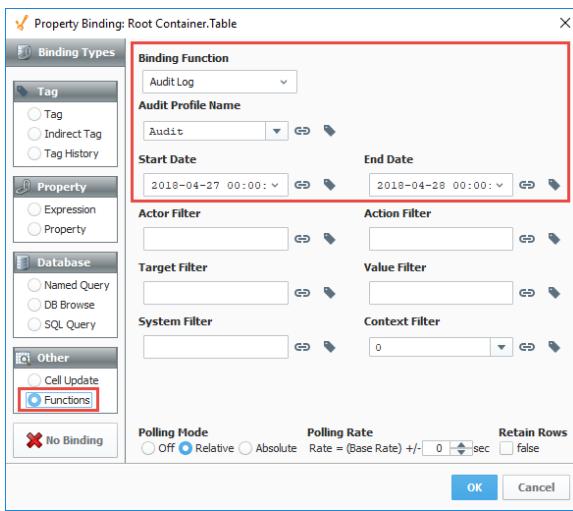
This page documents several ways to view results from the Audit Log System. For more information on how the Audit Log works, see the [Audit Log and Profiles](#) page.

Access the Audit Log with Table Functions

Since Ignition makes accessing data from databases seamless, it is possible to bind a data property of a table directly to database table. Alternatively, it is possible to access the contents of the audit log with table functions.

Once you have an [Audit Log set up and attached to a project](#), you can go back to that project and see what information is in that audit log. This example uses the Table Functions to extract data from the Audit Log.

1. In **Designer**, drag a **Table** component on to a window.
2. In the **Property Editor**, click on the binding icon from the table's **Data** property. The Property Binding window is displayed.
3. Select the **Functions** Binding Type.
4. Select **Functions** from the various binding options.
 - a. **Binding Function**, select the **Audit Log** from the drop-down menu.
 - b. **Audit Profile Name**, select **Audit** or the name of your Audit Profile.
 - c. **Start Date**, enter the appropriate start date.
 - d. **End Date**, enter the appropriate end date.
 - e. **Polling Mode**, select **Relative**.
 - f. Select any other appropriate function options from the menu, and click **OK**.



5. The table will populate with all the information stored in the audit log based on the Table Functions you selected. You can use the [Table Customizer](#) to configure how you want the table to look by reorganizing and hiding columns, making columns sortable, and assigning meaningful headers.

Timestamp	Actor	Action	Action Target	Action Value	Result	System	Context	Context
Apr 27, 2018 11:39 AM	admin	login			0	AuditStatus...	project->parameters	4:Client
Apr 27, 2018 3:48 PM	admin	logout			0	AuditStatus...	project->parameters	4:Client
Apr 27, 2018 3:48 PM	admin	project save	Parameters		0	AuditStatus...	project->parameters	2:Designer
Apr 27, 2018 3:49 PM	admin	project save	Parameters		0	AuditStatus...	project->parameters	2:Designer
Apr 27, 2018 3:49 PM	admin	logout			0	AuditStatus...	project->parameters	2:Designer
Apr 28, 2018 12:00 AM	Scheduled Report	report executed	WorkOrder Report	Failed Action ...	-2,147,-...,-2,147,-...	AuditStatus...	project->parameters	1:Gateway
Apr 28, 2018 12:00 AM	Scheduled Report	report executed	Test Report	Failed Action ...	-2,147,-...,-2,147,-...	AuditStatus...	project->parameters	1:Gateway
Apr 28, 2018 12:00 AM	Scheduled Report	report executed	Veggie Report 2		0	AuditStatus...	project->parameters	1:Gateway
Apr 28, 2018 12:00 AM	Scheduled Report	report executed	Car Inventory Report	Failed Action ...	-2,147,-...,-2,147,-...	AuditStatus...	project->parameters	1:Gateway

On this page ...

- [Access the Audit Log with Table Functions](#)
- [Access the Audit Log Using the Database Query Browser](#)
- [Access the Audit Log on the Gateway](#)



View Audit Information

[Watch the Video](#)

Access the Audit Log Using the Database Query Browser

The Database Query Browser makes it easy to search your database tables to view audit information.

1. In the **Designer** under the **Tools** menubar, select **Database Query Browser**. The Database Query Browser opens.

2. On the right side of Browser window under **Schema**, will be a list of tables from the currently selected database. Double click on the **audit_events** table, and click **Execute**.

The data from the audit_events table will appear in the Resultset 1 tab.

The screenshot shows the Database Query Browser interface. In the top left, there is a code editor with the query: "SELECT * FROM audit_events". Below it is a row limit input field set to 1000 rows. To the right is a toolbar with an "Execute" button. The main area displays a table titled "Resultset 1" with columns: AUDIT_EVENT_ID, EVENT_TIMESTAMP, ACTOR, ACTOR_HOST, ACTION, ACTION_TARGET, and ACTION_VALUE. The table contains 438 rows of audit log data. On the far right, a schema browser pane shows the structure of the audit_events table, including columns like ACTION (VARCHAR), ACTION_TARGET (TEXT), and ACTOR (VARCHAR). At the bottom of the browser are buttons for Auto Refresh, Edit, Apply, and Discard.

Refer to the [Database Query Browser](#) to learn more about how it works.

Access the Audit Log on the Gateway

Ignition provides a simple interface to view Audit Logs on the Gateway.

1. On the Gateway webpage in the **Config** section, scroll down to **Security > Auditing**.
2. The Audit Profile page will be displayed. Select the Audit Profile where your information is stored, and click **More > view log**.

The screenshot shows the "Audit Profiles" page in the Ignition Config interface. It lists two profiles: "Audit" and "Audit1". For each profile, there are "Type" (Database), "Description" (empty), and three buttons: "view log" (highlighted with a red box), "More", and "edit".

3. Choose the parameter settings if you're looking for something specific, otherwise, enter a **Start Date** and **End Date**, and click **Search**.

The screenshot shows the search results for audit logs. The search parameters are highlighted with a red box: "Start Date" (3/31/20, 6:25 AM) and "End Date" (4/1/20, 11:59 PM). Below the search bar is a table with columns: Timestamp, Actor, Host, Action, Target, Value, Result, System, and Context. The table contains three rows of audit log data corresponding to the search criteria.

Timestamp	Actor	Host	Action	Target	Value	Result	System	Context
3/31/20, 2:24 PM	admin	Unknown	resources modified	Production	ignition/global-props	AuditStatus[0x00000000, Severity=Good, Subcode=NotSpecified]	project:Production	Unknown (0)
3/31/20, 2:24 PM	admin	ws3	tag write	[default]Speed	78	AuditStatus[0x00000000, Severity=Good, Subcode=NotSpecified]	sys:Controller/project:Production	Designer
3/31/20, 2:24 PM	admin	ws3	tag write	[default]Speed	104	AuditStatus[0x00000000, Severity=Good, Subcode=NotSpecified]	sys:Controller/project:Production	Designer

Related Topics ...

- [Scripting](#)
- [Database Query Browser](#)
- [Table Customizer](#)

- Security

Auditing Actions Reference

The auditing system in Ignition records actions originating from the Gateway, Perspective and Vision projects. This page lists which actions are logged by the auditing system.

A description of the audit table can be found on the [Ignition Database Table Reference](#) page.

Gateway Audit Actions

The following actions are recorded in an audit log when the Gateway has a [Gateway Audit Profile](#) is configured.

Project System

The following project-based actions are tracked by the auditing system.

- Project Property changes made from the Designer.
- Project setting changes made from the gateway's web interface.
- Creating or deleting a project.
- Saving a project (action recorded as "project update").

Gateway Systems

In addition, project files on the Gateway's file system are closely monitored. If a user or third-party system modifies any of the project files, an entry will be recorded in the auditing system. The following Gateway-level actions are recorded in an audit log when the Gateway has a [Gateway Audit Profile](#) is configured.

Modules

- Installing modules on the Gateway.
- Restarting a module.
- Deleting a module.

Gateway - General

- Gateway startup
- Gateway shutdown (assuming the gateway was requested to shutdown: unintended shutdowns from power failures and such will not be recorded).

The following feature is new in Ignition version **8.1.17**
[Click here](#) to check out the other new features

- Gateway Login
- Gateway Logout

Gateway - Restore

- Restoring the Gateway from a Gateway backup. Specifically, the Gateway will log that it was asked to before a restore, then perform the restoration.

Licensing Changes

- Activating a license.
- Unactivating a license.
- Updating a license.

Redundancy

- Saving after making any changes to the Redundancy Settings page.

Web Server Page

- Installing or removing a security certificate.
- Making changes to the Web Server Settings page.

Gateway Network

On this page ...

- [Gateway Audit Actions](#)
 - [Project System](#)
 - [Gateway Systems](#)
 - [Remote Gateway Tag Writes](#)
- [Perspective Auditing Actions](#)
- [Vision Auditing Actions](#)
 - [Tags](#)
 - [Vision Tag Writes](#)
 - [Vision Component Database Writes](#)
 - [Vision User Login/Logout](#)
 - [Database Query Browser](#)
 - [Vision Scripting](#)
- [Designer](#)
 - [Designer Login and Closing](#)
 - [Database Query Browser](#)
- [Alarm Notification](#)
 - [Alarm Notification Attempts](#)
- [Reporting Module](#)
 - [Report Execution](#)

- Saving changes to Gateway Network General Settings .
- Creating, editing, or deleting outgoing connections.
- Approving incoming connections.

Email Settings

- Creating, editing, or deleting an SMTP Profile.

Audit Profile

- Creating, editing, or deleting an Audit Profile.

User Sources

- Creating, editing, or deleting a User Source.
- Creating, editing, or deleting a user.
- Creating, editing, or deleting a role.

The following feature is new in Ignition version **8.1.14**

[Click here](#) to check out the other new features

- User Lockout Events will also be recorded. Note that the audit log will record only the initial lockout event, rather than each failed authentication attempt.

Service Security

- Editing and saving a policy.

Identity Providers

- Creating, editing, or deleting an Identity Provider configuration.
- Making changes to a User Attribute Mapping.
- Creating, editing, or deleting a User Grant.
- Saving changes on a Security Level Rule.

Security Levels

- Creating, editing, and deleting security zones

Security Zones

- Creating, editing, and deleting security zones

Database - Connections

- Creating, editing or deleting a database connection.

Database - Drivers

- Creating, editing or deleting a JDBC driver
- Creating, editing or deleting a Translator

Store and Forward

- Creating, editing or deleting a Store and Forward engine.

Alarming - General

- Saving changes on the Alarming General settings page.

Alarming - Alarm Journal

- Creating, editing or deleting an Alarm Journal Profile

Alarming - Notification

- Creating, editing, or deleting an Alarm Notification Profile.

Schedules

- Creating, editing or deleting a schedule.
- Creating, editing or deleting a holiday.

Tags - Realtime

- Creating, editing or deleting a Realtime Tag Provider

Tags - Historical

- Creating, editing or deleting a Historical Tag Provider.

OPC Client Connections

- Creating, editing or deleting an OPC connection.

OPC UA - Device Connections

- Creating, editing, or deleting a device connection (editing/saving a device connection configuration without making any changes will be recorded as an edit).
- Editing a Modbus address mapping via gateway interface on Modbus device connections.
- Editing DNP3 Aliased Points via gateway interface on DNP3 device connections.
- Editing tags via gateway interface on Omron NJ device connections.
- Adding FINS tags via gateway interface on Omron FIN device connections.

OPC UA - Server Settings

- Editing the OPC UA Settings page.

Enterprise Administration

- Configuring a gateway to be either an Agent or Controller.

Enterprise Administration - Event Thresholds

- Changes made to Event thresholds.

Enterprise Administration - Controller Settings

- Making changes to the Controller Settings page, including uninstalling the controller.

Enterprise Administration - Agent Settings

- Making changes to the Agent Settings page, including uninstalling the agent.

Enterprise Administration - Agent Management

- Creating, editing, deleting an Agent Group.

Enterprise Administration - License Management

- Adding or removing a license from the License Management page.

Enterprise Administration - Agent Tasks

- Creating, editing, or deleting an agent task
- Separate records are taken each time a task executes.

Sequential Function Charts

- Changes made to the SFC Settings page.

Add a Record Manually

- You can also add a record into the audit profile using the function `system.util.audit`.

Remote Gateway Tag Writes

This feature was changed in Ignition version **8.1.34**:

- Actions on tags from remote servers are recorded in the audit log for versions 8.1.16+. The 'System' column shows the originating Gateway name. Note that 'Host', and 'Context' will appear unknown for these records, but auditing events will now include the actor.

Perspective Auditing Actions

Perspective Sessions generate entries in an assigned audit profile. The following actions are recorded in the Audit Profile:

- Tag changes from a component binding.
- Authentication level changes (a user's security level changes).
- Login Request - Indicates a user is requesting to log into an Identity Provider (IdP). The user should have been redirected to the IdP with a login request and Ignition is awaiting the IdP's login response. Note that the user is not logged in until the IdP redirects the user back to Ignition with a login response and Ignition validates the login response.
- Login Response - Records when a login response is received from the IdP. It's possible that a login response will never be received for a login request. For example: if the user bails out of the login flow by closing their web browser before completing the login, Ignition will never receive the login response and will time out the request.
- Logout Request - Indicates a user is requesting to log out of an IdP. The user *may* be redirected to the IdP to log out of their IdP session. Regardless the user will be redirected back to the Perspective Session in a logged out state.
- Logout Response - Records when a logout response is received from the IdP after a user logged out of their IdP session. This event will not occur if the IdP does not support logout or if Ignition is not configured to redirect the user to the IdP for logging out.

The following feature is new in Ignition version **8.1.18**

[Click here](#) to check out the other new features

- Tag changes from a Perspective script. Specifically:
 - Writes, such as, but not limited to, those from `system.tag.writeBlocking`
 - Edits and renaming, such as those caused by `system.tag.configure`
 - Deletions, such as those caused by `system.tag.deleteTags`
 - Moves, such as those caused by `system.tag.move`

Vision Auditing Actions

The Vision project needs an audit profile configured and auditing enabled. Vision Clients will then log records to an assigned audit profile. Here is a list of audit actions that will be tracked in the Ignition auditing system:

Tags

The following Tag related actions generate entries in the audit log. Note that the functions below must originate from the Tag Browser, the Designer's Scripting Console, or Vision component-based scripts.

- Tag Creation - Including tags created with the Tag Editor and the `system.tag.configure` function.
- Tag Deletion - Including those deleted from the Tag Browser's UI and the `system.tag.deleteTags` function.
- Tag Edits - Including edits made to tags from the Tag Editor and the `system.tag.configure` function.
- Moving Tags - Including moves made by drag-and-drop in the Tag Browser or by calling the `system.tag.move` function.
- Tag Renames - Renaming a tag generates an entry.

Vision Tag Writes

Write requests sent from a tag either through a standard Tag Binding, Indirect Tag Binding, or manual entry from the Tag Browser.

Vision Component Database Writes

The system explicitly captures modifications made to database tables through the following methods:

- **SQL Query Bindings** - modifications from the UPDATE Query will be recorded.
- **DB Browse Binding** - modifications made with the Enable Database Writeback area will be recorded.

Vision User Login/Logout

- Logging into a Vision Client will generate an entry in the auditing system, as will logging out of the client.
- Closing the client while logged in is treated as a logout. Note that the entry is only recorded if the client is aware that it is closing, which excludes cases where the client closed unexpectedly.

Database Query Browser

- If the project opened in the Designer has an assigned Audit Profile, then changes made to database tables using the database query browser are automatically recorded to the audit log. "Changes" in this case refer to UPDATE, DELETE, or INSERT statements manually typed and executed from the database query browser.
- Enabling edit mode and applying changes, including typing in new values, adding rows, removing rows, and clearing out fields, are recorded as queries called from the project.

Vision Scripting

The following functions generate entries in the audit log if called from Vision component-based scripts, or from the Designer's Scripting Console.

- [system.db.execSProcCall](#)
- [system.db.runPrepUpdate](#)
- [system.db.runUpdateQuery](#)
- [system.tag.writeBlocking](#)
- [system.tag.writeAsync](#)
- [system.report.executeAndDistribute](#)
- [system.report.executeReport](#)

Designer

Designer Login and Closing

- Opening a project in the Designer that has auditing enabled will also generate a login entry in the auditing system. Note that this occurs when the user opens the project, not when they log in using the Designer's login screen: auditing is project-based, so the user has to select a project that is being edited first.
- Closing the Designer effectively counts as logging off, and will generate a "logout" entry. Similar to vision, should the designer close unexpectedly, then an entry will not be recorded.

Database Query Browser

If the project opened in the [Designer](#) has an assigned Audit Profile, then changes made to [database](#) tables using the [database](#) query browser are automatically recorded to the audit log. "Changes" in this case refer to UPDATE, DELETE, or INSERT statements manually typed and executed from the [database](#) query browser.

Enabling edit mode and applying changes, including typing in new values, adding rows, removing rows, and clearing out fields, are recorded as queries called from the project.

Alarm Notification

Alarm Notification Attempts

Attempts to send out alarm notifications are recorded in the auditing system. Specifically, the Gateway will record when it attempted to send out a notification, as well as if the attempt failed (such as the SMTP server refusing the request). It is important to note that the auditing system can not report failures that occur outside of the Gateway. Thus, if a voice notification fails to send due to some error in the VOIP system, it's possible that the Gateway won't report the VOIP error, but the audit log will have an entry stating that the Gateway attempted to send the notification.

Reporting Module

Report Execution

Reporting Module Reports generate an entry in the auditing system when a report is executed. Thus:

- Reports running on a schedule will generate an entry.
- Report schedules executed on demand will generate an entry.
- Navigating to a Vision window (in either the Designer or a Vision Client) will trigger a report execution, generating an entry in the auditing system.

Related Topics ...

- [Audit Log and Profiles](#)
- [Audit Log Display](#)

Docker Image

Overview

The Ignition Platform offers a container image for use on popular container runtimes such as Docker. You can find the image on Docker Hub at <https://hub.docker.com/r/inductiveautomation/ignition>. Our official image is a Linux container, based on [Ubuntu Linux](#).

The following feature is new in Ignition version **8.1.16**

[Click here](#) to check out the other new features

Ignition's Docker Image build is now multi-architecture and supports linux/amd64, linux/arm64, and linux/arm/v7. Users can `docker pull inductiveautomation/ignition` and get an image that automatically matches their system architecture.

Getting Started

In order to get started with Ignition running as a container, you'll need a container runtime. [Docker Desktop](#), available for Windows and macOS, includes both the [Docker Engine](#) container runtime as well as container tooling and supporting tools such as [Docker Compose](#). If you are running Linux, you'll need to download and install the container runtime and tools such as Docker Compose individually. Guidance for the installation and configuration of Docker (on all platforms) are available here: [Get Docker](#)

Configuration

Various aspects of the Ignition Gateway can be configured through the use of either command-line arguments to the container, or environment variables. This section will detail the available configuration options and how best to use them.

Environment Variables

The available environment variables can either accept explicit values or load the target values from a file. Specifying any of the environment variables below with a `_FILE` suffix will direct the system to take that environment variable value as a path to a file containing the value of interest. This is useful for secrets management systems where you want to keep the values out of the process environment (which might be able to be inspected by individuals with access to `docker inspect <container>`, for example).

Note:

Some of the variables descriptions below make use of a # character. In these cases you can specify the same variables multiple times. Replace the # character with a number.

```
GATEWAY_NETWORK_0_HOST=1.2.3.4
GATEWAY_NETWORK_0_PORT=8088
GATEWAY_NETWORK_0_DESCRIPTION="For tag server"

GATEWAY_NETWORK_1_HOST=5.6.7.8
GATEWAY_NETWORK_1_PORT=8089
GATEWAY_NETWORK_1_DESCRIPTION="For edge server"
```

On this page ...

- [Overview](#)
- [Getting Started](#)
- [Configuration](#)
 - [Environment Variables](#)
 - [Runtime Arguments](#)
 - [Supplemental JVM and Wrapper Arguments](#)
 - [Gateway Arguments](#)
 - [Logging Settings](#)
- [Examples](#)
 - [Deploying an Ephemeral Gateway for Testing](#)
 - [Preserving Gateway State](#)
 - [Automating the Restore of a Gateway Backup](#)
 - [Automate the Commissioning of a Fresh Gateway](#)
 - [Customizing JVM/Wrapper /Gateway Arguments on Container Launch](#)
 - [Using Docker Compose](#)
 - [Integrate Third-Party Modules After an Upgrade](#)
 - [User Migration](#)
- [Licensing](#)

Variable	Requires	Description
TZ	Any	Set to a valid TZ database name, like <code>America/Los_Angeles</code> . See List of tzdata Entries for complete selection.
ACCEPT_INDUCTION_EULA	8.1.7+	Set to <code>Y</code> to accept the Ignition EULA automatically (see Licensing section below).
GATEWAY_RESTORE_DISABLED	8.1.7+	Set to <code>true</code> to have the restored gwbk disabled on startup.
GATEWAY_ADMIN_USERNAME	8.1.8+	Defaults to <code>admin</code> when not specified, used to set value for initial gateway auto-commissioning.

GATEWAY_ADMIN_PASSWORD	8.1.8+	Password value or salted hash to be used for initial gateway auto-commissioning.
GATEWAY_HTTP_PORT	8.1.8+	Defaults to 8088 when not specified, used to override the HTTP port used by the gateway container.
GATEWAY_HTTPS_PORT	8.1.8+	Defaults to 8043 when not specified, used to override the HTTPS port used by the gateway container.
GATEWAY_GAN_PORT	8.1.8+	Defaults to 8060 when not specified, used to override the Gateway Network port used by the gateway container.
IGNITION_EDITION	8.1.8+	Set to standard, edge , or maker , used to set value for initial gateway auto-commissioning.
IGNITION_LICENSE_KEY	8.1.8+	8-character license key (xxxx-xxxx) for leased activation (e.g. Maker Edition), used to set value for initial gateway auto-commissioning <p>The following feature is new in Ignition version 8.1.19 Click here to check out the other new features</p> <p>You can activate multiple 8-character license keys at the same time in a comma-delimited format, such as: ABCD-1234,DCBA-4321,WXYZ-6789,ZYXW-9876.</p> <p>This feature was changed in Ignition version 8.1.20:</p> <p>Changes to this variable will now update the leased activation license configuration after initial commissioning. Previously, they would only be absorbed if an existing leased activation configuration was not present.</p>
IGNITION_ACTIVATION_TOKEN	8.1.8+	Activation token for the license key. <p>The following feature is new in Ignition version 8.1.19 Click here to check out the other new features</p> <p>You can activate multiple 8-character license keys at the same time in a comma-delimited format, such as: abcdef1234...789,bcdef2345...890.</p> <p>This feature was changed in Ignition version 8.1.20:</p> <p>Changes to this variable will now update the leased activation license configuration after initial commissioning. Previously, they would only be absorbed if an existing leased activation configuration was not present.</p>
GATEWAY_NETWORK_#_HOST	8.1.10+	Hostname (or IP) for Outgoing GAN Connection Definition for #.
GATEWAY_NETWORK_#_PORT	8.1.10+	Port Number to use for connection n (defaults to 8060 when not specified)
GATEWAY_NETWORK_#_PINGRATE	8.1.10+	Ping Rate (ms) for connection # (defaults to 1000 when not specified)
GATEWAY_NETWORK_#_PINGMAXMISSED	8.1.10+	Number of missed pings allowed for connection # (defaults to 30 when not specified)
GATEWAY_NETWORK_#_ENABLED	8.1.10+	Set true or false to mark connection # enabled or disabled

GATEWAY_NETWORK #__ENABLE_SSL	8.1.10+	Set true or false to enable SSL on connection # (defaults to true).																
GATEWAY_NETWORK #__WEBSOCKETTIMEOUT	8.1.10+	Timeout (ms) for web socket # (defaults to 10000 when not specified)																
EAM_SETUP_INSTALLEDSELECTICON	8.1.10+	Set Agent or Controller to define the EAM target config (defaults to Agent when not specified)																
EAM_AGENT_CONTROLLER_LLERSERVERNAME	8.1.10+	Gateway Name (not hostname) of EAM Controller to bind to																
EAM_AGENT_SENDSTATSINTERVAL	8.1.10+	Interval (seconds) to send statistics to the EAM Controller (defaults to 5 when not specified)																
EAM_CONTROLLER_ARCHIVEPATH	8.1.10+	Filesystem path to store EAM Archives (gateway backups, etc) (defaults to data/eam_archive when not specified)																
EAM_CONTROLLER_DATABASOURCE	8.1.10+	The database connection name to use for recording agent history.																
EAM_CONTROLLER_ARCHIVELOCATION	8.1.10+	When set to MANUAL, the path set in EAM_CONTROLLER_ARCHIVEPATH will be used. Defaults to AUTOMATIC, if omitted.																
EAM_CONTROLLER_LOWDISKTHRESHOLDMB	8.1.10+	Value is in megabytes. Errors will be reported in this Gateway when the disk used for archiving drops below this value.																
GATEWAY_MODULES_ENABLED	8.1.17+	<p>A comma-delimited set of identifiers that will whitelist the set of built-in modules that will remain installed prior to gateway startup. This feature will help with container re-creation events where you want to ensure only a subset of modules remain enabled.</p> <table border="1"> <thead> <tr> <th>Module Identifier</th><th>Module Filename</th></tr> </thead> <tbody> <tr> <td>alarm-notification</td><td>Alarm Notification-module.modl</td></tr> <tr> <td>allen-bradley-drivers</td><td>Allen-Bradley Drivers-module.modl</td></tr> <tr> <td>bacnet-driver</td><td>BACnet Driver-module.modl</td></tr> <tr> <td>dnp3-driver</td><td>DNP3-Driver.modl</td></tr> <tr> <td>enterprise-administration</td><td>Enterprise Administration-module.modl</td></tr> <tr> <td>iec-61850-driver</td><td> <p>The following feature is new in Ignition version 8.1.25 Click here to check out the other new features</p> <p>IEC 61850 Driver-module.modl</p> </td></tr> <tr> <td>logix-driver</td><td>Logix Driver-module.modl</td></tr> </tbody> </table>	Module Identifier	Module Filename	alarm-notification	Alarm Notification-module.modl	allen-bradley-drivers	Allen-Bradley Drivers-module.modl	bacnet-driver	BACnet Driver-module.modl	dnp3-driver	DNP3-Driver.modl	enterprise-administration	Enterprise Administration-module.modl	iec-61850-driver	<p>The following feature is new in Ignition version 8.1.25 Click here to check out the other new features</p> <p>IEC 61850 Driver-module.modl</p>	logix-driver	Logix Driver-module.modl
Module Identifier	Module Filename																	
alarm-notification	Alarm Notification-module.modl																	
allen-bradley-drivers	Allen-Bradley Drivers-module.modl																	
bacnet-driver	BACnet Driver-module.modl																	
dnp3-driver	DNP3-Driver.modl																	
enterprise-administration	Enterprise Administration-module.modl																	
iec-61850-driver	<p>The following feature is new in Ignition version 8.1.25 Click here to check out the other new features</p> <p>IEC 61850 Driver-module.modl</p>																	
logix-driver	Logix Driver-module.modl																	

		<table border="1"><tr><td>mitsubishi-driver</td><td>Mitsubishi-Driver.modl</td></tr><tr><td>modbus-driver-v2</td><td>Modbus Driver v2-module.modl</td></tr><tr><td>omron-driver</td><td>Omron-Driver.modl</td></tr><tr><td>opc-ua</td><td>OPC-UA-module.modl</td></tr><tr><td>perspective</td><td>Perspective-module.modl</td></tr><tr><td>reporting</td><td>Reporting-module.modl</td></tr><tr><td>serial-support-client</td><td>Serial Support Client-module.modl</td></tr><tr><td>serial-support-gateway</td><td>Serial Support Gateway-module.modl</td></tr><tr><td>sfc</td><td>SFC-module.modl</td></tr><tr><td>siemens-drivers</td><td>Siemens Drivers-module.modl</td></tr><tr><td>sms-notification</td><td>SMS Notification-module.modl</td></tr><tr><td>sql-bridge</td><td>SQL Bridge-module.modl</td></tr><tr><td>symbol-factory</td><td>Symbol Factory-module.modl</td></tr><tr><td>tag-historian</td><td>Tag Historian-module.modl</td></tr><tr><td>udp-tcp-drivers</td><td>UDP and TCP Drivers-module.modl</td></tr><tr><td>vision</td><td>Vision-module.modl</td></tr><tr><td>voice-notification</td><td>Voice Notification-module.modl</td></tr><tr><td>web-browser</td><td>Web Browser Module.modl</td></tr><tr><td>web-developer</td><td>Web Developer Module.modl</td></tr></table>	mitsubishi-driver	Mitsubishi-Driver.modl	modbus-driver-v2	Modbus Driver v2-module.modl	omron-driver	Omron-Driver.modl	opc-ua	OPC-UA-module.modl	perspective	Perspective-module.modl	reporting	Reporting-module.modl	serial-support-client	Serial Support Client-module.modl	serial-support-gateway	Serial Support Gateway-module.modl	sfc	SFC-module.modl	siemens-drivers	Siemens Drivers-module.modl	sms-notification	SMS Notification-module.modl	sql-bridge	SQL Bridge-module.modl	symbol-factory	Symbol Factory-module.modl	tag-historian	Tag Historian-module.modl	udp-tcp-drivers	UDP and TCP Drivers-module.modl	vision	Vision-module.modl	voice-notification	Voice Notification-module.modl	web-browser	Web Browser Module.modl	web-developer	Web Developer Module.modl
mitsubishi-driver	Mitsubishi-Driver.modl																																							
modbus-driver-v2	Modbus Driver v2-module.modl																																							
omron-driver	Omron-Driver.modl																																							
opc-ua	OPC-UA-module.modl																																							
perspective	Perspective-module.modl																																							
reporting	Reporting-module.modl																																							
serial-support-client	Serial Support Client-module.modl																																							
serial-support-gateway	Serial Support Gateway-module.modl																																							
sfc	SFC-module.modl																																							
siemens-drivers	Siemens Drivers-module.modl																																							
sms-notification	SMS Notification-module.modl																																							
sql-bridge	SQL Bridge-module.modl																																							
symbol-factory	Symbol Factory-module.modl																																							
tag-historian	Tag Historian-module.modl																																							
udp-tcp-drivers	UDP and TCP Drivers-module.modl																																							
vision	Vision-module.modl																																							
voice-notification	Voice Notification-module.modl																																							
web-browser	Web Browser Module.modl																																							
web-developer	Web Developer Module.modl																																							
IGNITION_UID	8.1.17+	<p>Numeric user ID for the target user. Passing this variable and <code>IGNITION_GID</code> allows Ignition to run within a container as an <code>ignition</code> user rather than as the <code>root</code> user. When set, the entrypoint will automatically update file ownerships for the Ignition installation on startup to match the target user prior to stepping down from the <code>root</code> user to launch the Gateway.</p> <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"><p><u>This feature was changed in Ignition version 8.1.26.</u></p></div> <p>The default user, as well as default file ownership, has been changed from the <code>root</code> user (UID 0) to a standard <code>ignition</code> user (UID 2003).</p> <p>If upgrading from a version prior to 8.1.26, you may need to take additional steps to migrate your users. Attempting to launch a container against an old data volume without migrating your users can result in file permission errors. You may also continue running your container as a <code>root</code> user by explicitly setting the UID and GID to 0, but this is discouraged due to security concerns. See the User Migration section</p>																																						
IGNITION_GID	8.1.17+	<p>Numeric group ID for the target user. Passing this variable and <code>IGNITION_UID</code> allows Ignition to run within a container as an <code>ignition</code> user rather than as the <code>root</code> user. When set, the entrypoint will automatically update file ownerships for the Ignition installation on startup to match the target user prior to stepping down from the <code>root</code> user to launch the Gateway.</p> <div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"><p><u>This feature was changed in Ignition version 8.1.26.</u></p></div>																																						

		<p>The default user, as well as default file ownership, has been changed from the <code>root</code> user (GID 0) to a standard <code>ignite</code> user (GID 2003).</p> <p>If upgrading from a version prior to 8.1.26, you may need to take additional steps to migrate your users. Attempting to launch a container against an old data volume without migrating your users can result in file permission errors. You may also continue running your container as a <code>root</code> user by explicitly setting the UID and GID to 0, but this is discouraged due to security concerns. See the User Migration section on this page for more details.</p>
DISABLE_QUICKSTART	8.1.23+	Boolean to decide whether or not the "Enable Quick Start" prompt will appear when new Gateway Containers are created. A setting of <code>true</code> will skip the Quick Start prompt.
GATEWAY_NETWORK_#_DESCRIPTION	8.1.26+	<p>The following feature is new in Ignition version 8.1.26 Click here to check out the other new features</p> <p>Documentation about how the connection is used if entered.</p>
GATEWAY_NETWORK_ENABLED	8.1.32+	Set to false to disable using the Gateway Network.
GATEWAY_NETWORK_REQUIRESSL	8.1.32+	If true, only connections that use SSL to encrypt traffic will be allowed. This setting only applies to incoming connections.
GATEWAY_NETWORK_REQUIRETWOFAUTH	8.1.32+	Enforces two-way SSL authentication. If true, you will need to install the remote machine's certificate on this machine, in addition to manual approval of this machine's certificate on the remote machine.
GATEWAY_NETWORK_SENDTHREADS	8.1.32+	The maximum number of threads that will be used to upload messages. Applies to outgoing connections.
GATEWAY_NETWORK_RECEIVETHREADS	8.1.32+	The maximum number of threads that will be used to download messages. Applies to outgoing connections.
GATEWAY_NETWORK_RECEIVEMAX	8.1.32+	Number of received messages that can be held until they are processed by the local system. When this capacity is exceeded, new messages are rejected and errors are reported to the remote Gateway. Applies to incoming connections.
GATEWAY_NETWORK_ALLOWINGCOMING	8.1.32+	If false, only outward connections defined on this Gateway will be allowed.
GATEWAY_NETWORK_SECURITYPOLICY	8.1.32+	Dictates what connections are allowed. If set to 'ApprovedOnly', incoming connections must be approved from the Incoming Connections page. Other options include 'Unrestricted' and 'SpecifiedList'.
GATEWAY_NETWORK_WHITELIST	8.1.32+	If connection policy is set to Specified List, this is the comma-separated list of Gateway names which will be allowed to connect.
GATEWAY_NETWORK_ALLOWEDPROXYHOPS	8.1.32+	The maximum number of proxy hops which could be used to reach the destination Gateway. Any number less than or equal to zero is equivalent to no proxy hops allowed.
GATEWAY_NETWORK_WEBSOCKETSESSIONIDLETIMEOUT	8.1.32+	The maximum number of milliseconds that a websocket is allowed to remain idle before it is closed. This value should always be set higher than outgoing connection ping rates to avoid premature connection termination.

Runtime Arguments

The following are accepted by the container as runtime arguments. *Options* should be followed by a value, *flags* are stand-alone. Note that these arguments are passed after the *image name* in a Docker Run statement.

Options

Option	Requires	Arguments	Description
-n	8.1.0+	<i>String</i>	<i>Gateway name</i> to be applied to the Ignition Gateway on startup
-a	8.1.0+	<i>String</i>	<i>Public web address*</i>
-h	8.1.0+	<i>Integer</i>	<i>Public HTTP port*</i>
-s	8.1.0+	<i>Integer</i>	<i>Public HTTPS port*</i>
-m	8.1.0+	<i>Integer</i>	Max memory for JVM
-r	8.1.7+	<i>String</i>	Path to Gateway Backup for Automated Restore

* - Specify Public address/HTTP/HTTPS ports together, all three must be specified

Flags

Flag	Requires	Description
-d	8.1.0+	<i>Debug mode</i> , applicable to development - attaches port 8000 for remote JVM debugging. Port 8000 will also likely need to be published on the container.

Supplemental JVM and Wrapper Arguments

The following feature is new in Ignition version **8.1.8**

[Click here](#) to check out the other new features

You can also specify JVM and Java Service Wrapper arguments to your Docker containers runtime arguments by adding a double-hyphen to delimit these arguments from the other runtime arguments listed above. Wrapper arguments (starting with `wrapper.*`) will be merged and JVM arguments added to those specified in the `ignition.conf` file.

Gateway Arguments

The following feature is new in Ignition version **8.1.10**

[Click here](#) to check out the other new features

The image also supports the use of gateway arguments, allowing you to modify the `gateway.xml` file when launching a container. Like JVM and wrapper arguments, gateway arguments must be specified after a double-hyphen. Only entries in the file that follow the pattern of "gateway.#" can be modified in this way.

This feature was changed in Ignition version **8.1.16**:

Docker image entrypoint will no longer forcibly recreate the `gateway.xml` file on each launch, allowing for settings adjustments from the gateway web UI to properly persist without static definition in the container configuration.

Logging Settings

Unlike the non-containerized version of Ignition, the Ignition container image is designed to emit logs to `stdout` ("standard out") so they can be leveraged by various logging drivers in the container engine. This is done by directing what would normally be stored in the `logs/wrapper.log` file to `stdout`. As a result, Ignition itself doesn't control rotating the `wrapper.log` based on size, which can result in container logs growing unbounded. You can configure the default logging driver to constrain maximum log sizes. Reference the [Docker logging configuration](#) page to specify logger size limits.

Examples

This section will contain some example run configurations for the Ignition Docker image. Review these examples to learn a bit more about how to run and configure the Ignition Docker image, as well as some typical best practices.

Deploying an Ephemeral Gateway for Testing

The run statement below will launch a container in detached mode -d, publishing port 9088 from the host to port 8088 in the container. The container will be named `ignition-test` and use the `8.1.7` image. Finally, the runtime arguments provided will set the *Gateway name* to `docker-test` with public address of `localhost` on HTTP port 9088 and HTTPS port 9043.

```
docker run -d -p 9088:8088 --name ignition-test inductiveautomation/ignition:8.1.7 \
-n docker-test -a localhost -h 9088 -s 9043
```

Multiline commands

Note that for Linux/macOS/WSL, you can use the backslash for multi-line commands for better readability. You can substitute the backtick ` for Powershell and caret ^ for Windows Command prompt, these characters technically "escape" the newline character, so make sure they are the last character on a given line.

Preserving Gateway State

When using containers, it is common for stateful applications (such as the Ignition Gateway) to leverage volumes to persist that state across image updates, container remove/create cycles, etc. This can be done by applying a named volume to the `/usr/local/bin/ignition/data` path within the container. This is especially important if you need to make a change to your container configuration. Since container configurations are immutable, the only way to change it is to stop/remove the old container and start a new one (with a new configuration).

Similar to the above, but this time with a named volume `gw-data` attached to `/usr/local/bin/ignition/data` within the container. If the volume doesn't already exist, it will be created automatically. Note the additional `--pull` option to make sure that the `latest` tag is always pulled before running--without this, the `latest` tag is only pulled if one isn't already present on your system.

```
docker run -d -p 9088:8088 --name ignition-test \
-v gw-data:/usr/local/bin/ignition/data \
--pull always inductiveautomation/ignition:latest \
-n docker-test -a localhost -h 9088 -s 9043
```

By using named volumes, you're now able to stop and remove the `ignition-test` container and create a new one with a different configuration. As long as you attach your volume, your gateway will start up with all of your tags and projects in their previous state.

Preserving KeyStores

The following feature is new in Ignition version **8.1.12**
[Click here](#) to check out the other new features

Starting in 8.1.12, the following KeyStores are also preserved across image updates.

Gateway Network Certificate

This PKCS #12 KeyStore contains the certificate and private key used for gateway network communications, and is created automatically on Gateway startup if not present. Typically this is maintained independently on a per-installation basis to avoid issues from operations such as gateway backup/ restorations. The situation within a container is somewhat different, and it is usually preferable to track this KeyStore with the rest of the gateway state preserved by a volume. Docker image now redirects, via sym-link, Gateway Network KeyStore creation from `webserver/metro-keystore` to `data/local`. This will allow the underlying image for the container to be upgraded without generating a new Gateway Network certificate (thus breaking existing approved certificates and connections).

SSL Configuration

When SSL is enabled on a Gateway, `ssl.pfx` is created as a PKCS #12 KeyStore with the private key, server certificate, and root/intermediate CA certificates. This file is automatically read by the Gateway on startup to enable SSL. Similar to the [Gateway Network Certificate](#), this resides outside of the `data/` volume since it is intended to persist across a gateway restore operation. Docker image now redirects, via sym-links, SSL KeyStore creation from `webserver/ssl.pfx` and `webserver/csr.pfx` to `data/local`. This allows a Docker container's SSL configuration to persist via the `data` volume without relying on extra bind-mounts. Disabling SSL now recognizes the presence of sym-links when removing the KeyStore and removes the final target of the link, leaving the sym-link in place.

Automating the Restore of a Gateway Backup

The following feature is new in Ignition version **8.1.8**
[Click here](#) to check out the other new features

You can automate the restore of a gateway backup on first-launch of your gateway container. This allows for having a new Ignition Gateway restore to a known initial state automatically, without waiting for the commissioning steps.

To leverage this feature, bind-mount a gateway backup into the container and then use the `-r` runtime argument to specify the location and command the restore. Additionally, supply the `ACCEPT_IGNITION_EULA=Y` environment variable to accept the Ignition EULA (see the [Licensing](#) section below) and bypass that gateway commissioning step.

```
docker run -d -p 9088:8088 --name ignition-test \
-v gw-data:/usr/local/bin/ignition/data \
-v /path/to/gateway.gwbk:/restore.gwbk \
-e ACCEPT_IGNITION_EULA=Y \
inductiveautomation/ignition:8.1.7 \
-n docker-test -a localhost -h 9088 -s 9043 \
-r /restore.gwbk
```

What if I restart the container?

The gateway restore will only apply on a fresh gateway launch. Subsequent restarts of the container will not restore the indicated gateway backup.

Automate the Commissioning of a Fresh Gateway

The following feature is new in Ignition version **8.1.8**

[Click here](#) to check out the other new features

You can automate the commissioning steps that normally require manual user interaction on the very first launch of the Ignition Gateway. By supplying specific [environment variables](#), you can seed the commissioning steps with the required information to start the rest of the Gateway automatically.

```
docker run -d -p 9088:8088 --name ignition-test \
-e ACCEPT_IGNITION_EULA=Y \
-e GATEWAY_ADMIN_PASSWORD=password \
-e IGNITION_EDITION=standard \
inductiveautomation/ignition:8.1.8 \
-n docker-test
```

Reading Environment Variables from a file

You can also specify environment variables in the form of `<env var>_FILE=/path/to/file` in order to have the environment variable value resolved by reading it from a file. This is helpful for secrets management systems within container orchestrators such as Docker Swarm and Kubernetes.

Customizing JVM/Wrapper/Gateway Arguments on Container Launch

This example shows how to leverage the supplemental JVM/wrapper args feature.

```
docker run -d -p 9088:8088 --name args-test \
-e ACCEPT_IGNITION_EULA=Y \
-e GATEWAY_ADMIN_PASSWORD=changeme \
-e IGNITION_EDITION=standard \
inductiveautomation/ignition:8.1.8 \
-n args-test \
-- wrapper.java.initmemory=256 -Dignition.allowunsignedmodules=true \
-XX:MaxGCPauseMillis=200
```

The following feature is new in Ignition version **8.1.10**

[Click here](#) to check out the other new features

The following demonstrates how to enable the gateway's Resolve Host Names and Use Proxy Forwarded Headers settings by modifying the entries in the `gateway.xml`:

```
docker run -d -p 9088:8088 --name ignition-test inductiveautomation/ignition:8.1.10 -n docker-test -a
localhost -h 9088 -s 9043 \
-- gateway.resolveHostNames=true gateway.useProxyForwardedHeader=true
```

Using Docker Compose

A common practice is to leverage [Docker Compose](#) to start your container and bundle the configuration with other services, such as databases and MQTT brokers. Below is an example `docker-compose.yml` file that you can create inside of a new, empty folder. The example below incorporates many of the configurability features mentioned in the various sections above.

Compose Example

```
# Docker Compose Example for inductiveautomation/ignition
# Compose Spec: https://github.com/compose-spec/compose-spec/blob/master/spec.md
---
services:
  # Ignition Gateway
  gateway:
    image: inductiveautomation/ignition:8.1.8
    ports:
      - 9088:8088
      - 9043:8043
    volumes:
      - gw-data:/usr/local/bin/ignition/data
    # env_file: ignition.env # optionally specify variables in a file, or using `environment:` below
    environment:
      - ACCEPT_IGNITION_EULA=Y
      - GATEWAY_ADMIN_USERNAME=admin
      - GATEWAY_ADMIN_PASSWORD_FILE=/run/secrets/gateway-admin-password
      - IGNITION_EDITION=standard
      - TZ=America/Chicago # see https://en.wikipedia.org/wiki/List_of_tz_database_time_zones#List
    secrets:
      - gateway-admin-password
    command: >
      -n docker-test
      -m 1024
      --
      wrapper.java.initmemory=512
      -Dignition.allowunsignedmodules=true

secrets:
  gateway-admin-password:
    file: secrets/GATEWAY_ADMIN_PASSWORD

volumes:
  gw-data:
```

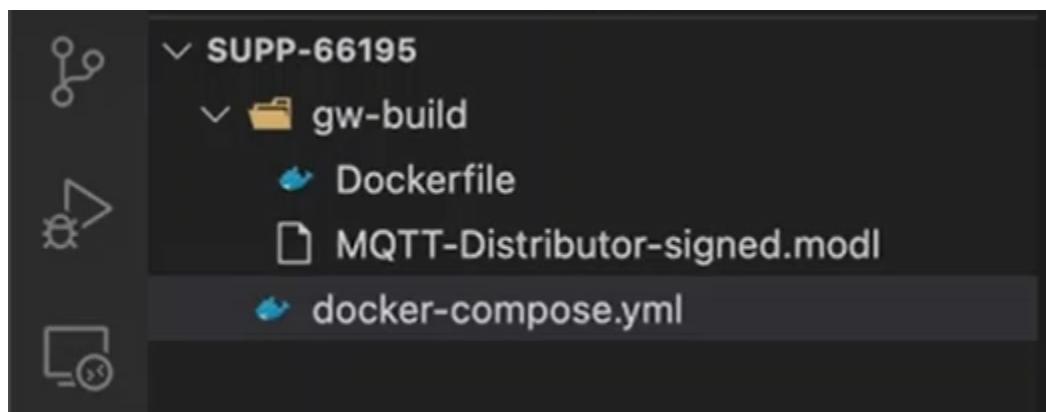
Once defined, you can bring up the solution using ``docker compose`` commands. See the animated example below:



Integrate Third-Party Modules After an Upgrade

When upgrading Ignition in Docker, third-party modules are not included with new images, so they will need to be re-installed. You can do this by creating a Docker derived image and use the derived image to copy these modules to the upgraded container. The example below describes how to conduct a derived image build to integrate third-party modules.

For this example, a module subfolder with the modules we are attempting to integrate will be labeled gw-build. The subfolder will also include a Dockerfile. The docker-compose.yml file needed for this process is located outside of this subfolder. We will be using a docker-compose.yml file similar to the previous example, but instead of using an image tag, we will be using a build tag.



docker-compose.yml

```
services:
  gateway:
    build:
      context: gw-build
      ## Specify the upstream version to derive from for the build argument in the Docker file.
      args:
        IGNITION_VERSION: 8.1.28
    ports:
      - 8088:8088
    volumes:
      - gateway-data:/usr/local/bin/ignition/data
    environment:
```

```

ACCEPT_IGNITION_EULA: "Y"
GATEWAY_ADMIN_PASSWORD: password
IGNITION_EDITION: standard
command: >
  -n Ignition-supp-66195

volumes:
  gateway-data:

```

The Dockerfile within the gw-build folder will define the derived image by containing a build argument for the Ignition version. Notice this does not give a default version. This is because the version we want to build is specified and passed through our docker-compose definition. Next, we will add a command to copy the module files from this folder to the user-lib module folder inside the new Ignition image.

Dockerfile

```

ARG IGNITION_VERSION
FROM inductiveautomation/ignition:${IGNITION_VERSION}

COPY *.modl /usr/local/bin/ignition/user-lib/modules/

```

With these files defined, we can now run `docker compose build` to pull the image, perform all derived image build steps, and create the new image.

Note: Since `docker compose build` needs to be run each time a new module is added, you can alternatively add `pull_policy: build` in the `.yml` file to auto-rebuild each time.

User Migration

User Migration - Docker Compose

The following example uses Docker Compose to demonstrate how to migrate the container user and file ownership from the `root` user to a standard, non-elevated `ignition` user. If you are instead using the command line, see the [next example](#) for migrating users. Your environment may vary, depending on how your container is set up. This container is running Ignition version 8.1.25:

User Migration Example (Docker Compose)

```

# User Migration Example (Docker Compose)
---
services:
  # Ignition Gateway
  gateway:
    image: inductiveautomation/ignition:8.1.25
    volumes:
      - ignition-data:/usr/local/bin/ignition/data
    ports:
      - 8088:8088
    command: >
      -n docker-test

volumes:
  ignition-data:

```

1. Upgrade the container to the latest version by changing the image tag to the version you want to upgrade to:

User Migration Example (Docker Compose)

```

# User Migration Example (Docker Compose)
---
services:
  # Ignition Gateway
  gateway:
    image: inductiveautomation/ignition:8.1.26
    volumes:
      - ignition-data:/usr/local/bin/ignition/data

```

```

ports:
  - 8088:8088
command: >
  -n docker-test

volumes:
  ignition-data:

```

2. Add the `IGNITION_UID=2003` and `IGNITION_GID=2003` environment variables to update the container permissions:

User Migration Example (Docker Compose)

```

# User Migration Example (Docker Compose)
---
services:
  # Ignition Gateway
  gateway:
    image: inductiveautomation/ignition:8.1.26
    volumes:
      - ignition-data:/usr/local/bin/ignition/data
    ports:
      - 8088:8088
    environment:
      - IGNITION_UID=2003
      - IGNITION_GID=2003
    command: >
      -n docker-test

  volumes:
    ignition-data:

```

3. Declare the root user (`user: "0:0"`):

User Migration Example (Docker Compose)

```

# User Migration Example (Docker Compose)
---
services:
  # Ignition Gateway
  gateway:
    image: inductiveautomation/ignition:8.1.26
    volumes:
      - ignition-data:/usr/local/bin/ignition/data
    ports:
      - 8088:8088
    environment:
      - IGNITION_UID=2003
      - IGNITION_GID=2003
    user: "0:0"
    command: >
      -n docker-test

  volumes:
    ignition-data:

```

4. Save the Docker Compose file, then run `docker compose up -d`.

Note: By default, we are using the `-d` option in our Docker command. While you can also run the command without `-d`, the `-d` option runs the container in detached mode, allowing the container to run in the background even after closing the terminal.

5. The container will now use the `root` user when it starts up, but will launch Ignition as a standard `ignition` user. All of the files in the `ignition-data` volume have the proper ownership. We can now remove the `IGNITION_UID` and `IGNITION_GID` environment variables and the `user:"0:0"` override:

User Migration Example (Docker Compose)

```
# User Migration Example (Docker Compose)
---
services:
  # Ignition Gateway
  gateway:
    image: inductiveautomation/ignition:8.1.26
    volumes:
      - ignition-data:/usr/local/bin/ignition/data
    ports:
      - 8088:8088
    command: >
      -n docker-test

volumes:
  ignition-data:
```

6. Finally, we can rerun the `docker compose up -d` command to recreate the container. Since we removed the environment variables and updated container permissions, the container will launch with the default user being a standard `ignition` user.
7. To check if the user successfully migrated over, we see the process listing and who the user is for each process using the command `docker compose exec gateway ps aux` in the terminal. Additionally, you can use the command `docker compose exec gateway whoami` to see who the default user for the container is:

```
$ docker compose exec gateway ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START  TIME COMMAND
ignition     1  0.0  0.0   2500   584 ?        Ss  21:25  0:00 tini -g -- ./ig
ignition    31  0.0  0.0  16348  4084 ?        Sl  21:25  0:00 ./ignition-gate
ignition    49  0.8  1.0 5907440 176192 ?        Sl  21:25  0:07 /usr/local/bin/
ignition  1312  0.0  0.0   8896  3252 pts/0    Rs+ 21:40  0:00 ps aux

$ docker compose exec gateway whoami
ignition
```

User Migration - Command Line

The following example uses the command line to demonstrate how to migrate the container user and file ownership from the `root` user to a standard, non-elevated `ignition` user. If you are instead using Docker Compose, see the [preceding example](#). Your environment may vary, depending on how your container is set up. The following container is called `example-gw` and is running Ignition version 8.1.25:

User Migration Example (Command Line)

```
docker run --name example-gw -v example-gw-data:/usr/local/bin/ignition/data \
-e IGNITION_EDITION=standard -e ACCEPT_IGNITION_EULA=Y -e GATEWAY_ADMIN_PASSWORD=password \
-p 8088:8088 \
inductiveautomation/ignition:8.1.25 \
-n example-gw
```

1. Stop and remove your Docker container using the following commands. Keep in mind that you will need to replace `example-gw` with your container name:

User Migration Example (Command Line)

```
# Stop and remove container
docker stop example-gw
docker rm example-gw
```

2. Modify the container run configuration to declare the root user using `--user 0:0` (The UID and GID will both be 0):

User Migration Example (Command Line)

```
# Declare the root user
docker run --name example-gw --rm -v example-gw-data:/usr/local/bin/ignition/data \
```

```
-p 8088:8088 \
--user 0:0 \
inductiveautomation/ignition:8.1.25 \
-n example-gw
```

3. Update the container permissions by adding the `IGNITION_UID` and `IGNITION_GID` environment variables. You can do this by adding `-e IGNITION_UID=2003 -e IGNITION_GID=2003` after `--user 0:0`:

User Migration Example (Command Line)

```
# Update the container permissions
docker run --name example-gw --rm -v example-gw-data:/usr/local/bin/ignition/data \
-p 8088:8088 \
--user 0:0 -e IGNITION_UID=2003 -e IGNITION_GID=2003 \
inductiveautomation/ignition:8.1.25 \
-n example-gw
```

4. The container's Gateway should now be running under a standard `ignition` user, using a UID of 2003 and a GID of 2003. You can run the following command to count the number of processes running as a standard `ignition` user:

User Migration Example (Command Line)

```
# Count the number of running processes under the ignition user
docker exec example-gw pgrep -u 2003 -c
```

Note:

There should be 3 processes:

- the endpoint
- the Java wrapper
- the Java process

5. File ownership should now also belong to the standard `ignition` user. You can use the following command to verify that there are no files owned by any other user aside from the standard `ignition` user:

User Migration Example (Command Line)

```
# Verify file ownership
docker exec example-gw bash -c "find /usr/local/bin/ignition ! -user 2003 | wc -l"
```

6. Once you have verified file ownership and the user the processes are running under, you can now recreate the container against the new image with the default `ignition` user (UID = 2003) and update your Gateway to the latest version:

User Migration Example (Command Line)

```
# Recreate the container
docker stop example-gw
docker rm example-gw
docker run --name example-gw --rm -v example-gw-data:/usr/local/bin/ignition/data \
-p 8088:8088 \
inductiveautomation/ignition:8.1.<new> \
-n example-gw
```

Licensing

The Inductive Automation software contained in this Docker container image is licensed under the terms of the [Inductive Automation Software License Agreement](#).

By passing the value ACCEPT_IGNITION_EULA=Y, you agree that your use of Inductive Automation software running in this Docker container image is governed by the terms of the [Inductive Automation Software License Agreement](#).

All third party software (whether open or closed) running in this Docker container image is licensed as indicated in this Docker Hub "License" section or in the container itself, usually as a NOTICE or LICENSE file. You agree to comply with any relevant licenses for third party software contained in this Docker container image.